

**BOUNDARY LAWS AND SMOOTH CONTACT RIGIDITY OF PERIODIC  
DISPERSING BILLIARDS  
FULL TECHNICAL MANUSCRIPT**

QIAN QI

**ABSTRACT.** Conditional endpoint laws near a supplied nongrazing periodic billiard polygon determine the complete smooth reflecting contact germs, not merely their Taylor series. The success mass is exponentially small; an exact Jacobi normalization and a two-ended determinant comparison yield a relative physical law on a fixed contact collar. Two positive offsets separate its half-line actions without supplied curvature or local flux amplitudes. A global quadratic inverse, an envelope identity for actual boundary variations and contracted later visits recover the smooth contacts, including flat parts, without analyticity or a closeness assumption. For oblique polygons one law per phase also determines its unreported positive offset. Its four first-moment profiles reconstruct the full limiting density by a uniformly stable rank-two identity, whose nondegeneracy follows from the marked slopes. This reduces the smoothing dimension in finite-preparation reconstruction while retaining the finite-flight error. With independent normalized unit-speed Liouville preparations and exact success tags, the charged record additionally determines the free cell area. We characterize its exact fibers by contact-germ, offset and area equality, and construct a jointly realizable estimator with failures charged and bounded post-acceptance coordinate errors. The area bound retains the flight-length amplification of curvature error. Equal-area remote completions and actual flat perturbations distinguish local rigidity from global completion and from formal-jet determination. Analytic continuation and intrinsic multichannel reconstruction are developed under their respective observation hypotheses.

---

*Date:* September 17, 2026.

*2020 Mathematics Subject Classification.* 37D50, 37C30, 70H09, 53A04, 62G05.

*Key words and phrases.* Dispersing billiards, inverse problems, boundary laws, single-offset density inversion, unknown transverse origins, endpoint association, contact rigidity, analytic rigidity, generic finite-channel rigidity, infinitesimal rigidity, finite observable coordinates, quantized-law stability, lattice holonomy, Poisson boundary experiments, moving support, Le Cam equivalence, observable calibration, global statistical reconstruction.

## CONTENTS

|                                                                        |     |
|------------------------------------------------------------------------|-----|
| 1. Introduction                                                        | 4   |
| <b>Part 1. Periodic boundary laws and smooth contact determination</b> | 12  |
| 2. Relative boundary laws along a periodic itinerary                   | 12  |
| 3. Contact reconstruction along a marked periodic orbit                | 21  |
| 4. Global identification of marked contact germs                       | 29  |
| 5. Smooth contact rigidity from the complete boundary law              | 35  |
| 6. Smooth contact recovery from finite preparations                    | 42  |
| 7. Local determination and exterior scattering data                    | 47  |
| 8. One uncalibrated law at an oblique contact                          | 49  |
| 9. The complete invariant of the charged local experiment              | 54  |
| 10. Reconstruction from first-moment profiles                          | 60  |
| 11. Conditioning of the physical-to-smooth inverse                     | 66  |
| <b>Part 2. Finite-chain foundations and alternating laws</b>           | 68  |
| 12. Geometric setting and the principal relative theorem               | 68  |
| 13. Uniform thresholds and conditional contact metrics                 | 69  |
| 14. Shortest channels in arbitrary periodic configurations             | 70  |
| 15. A parameter-dependent boundary-value and relative-flux lemma       | 72  |
| 16. Full-phase integration and the threshold law                       | 74  |
| 17. Two nonlinear boundary layers                                      | 77  |
| 18. From Jacobi identities to a relative physical law                  | 82  |
| 19. Hyperbolic coordinates and the intrinsic boundary measure          | 85  |
| 20. Boundary energy profiles and nonlinear parity compatibility        | 94  |
| <b>Part 3. Analytic and intrinsic multichannel consequences</b>        | 98  |
| 21. Determination of two independent even contacts                     | 98  |
| 22. Signed endpoint rigidity and the all-order determinant-one inverse | 102 |
| 23. An amplitude-free inverse from one endpoint law                    | 118 |
| 24. Interaction data with unknown transverse origins                   | 131 |
| 25. Intrinsic lifted-channel graphs and analytic table rigidity        | 139 |
| 26. Rank-two holonomy and recovery of the periodic lattice             | 146 |
| 27. Finite analytic signatures and stability of intrinsic gluing       | 148 |
| 28. A nonsymmetric periodic realization with four measured channels    | 152 |
| 29. Finite channel skeletons and generic periodic rigidity             | 157 |
| 30. Finite symmetry and the periodic reconstruction branches           | 164 |
| 31. Infinitesimal rigidity and finite observable coordinates           | 169 |
| <b>Part 4. Boundary information and physical local experiments</b>     | 175 |
| 32. Observation spaces and coarsenings                                 | 175 |
| 33. Small-window information along realizable channels                 | 178 |
| 34. Relative boundary laws for physical experiments                    | 187 |
| 35. Adaptive observation and stopped experiments                       | 189 |
| 36. Boundary information after residual-time integration               | 192 |
| 37. Vector boundary information for a moving-support local family      | 196 |
| 38. Physical realization of anchored finite contact-jet families       | 204 |
| 39. A physical endpoint experiment on fixed laboratory windows         | 206 |

|                |                                                                              |            |
|----------------|------------------------------------------------------------------------------|------------|
| 40.            | Compact local experiments                                                    | 212        |
| 41.            | The complete endpoint–time acquisition experiment at its fastest local scale | 213        |
| 42.            | Two-sided deficiency for the moving-ceiling experiment                       | 216        |
| 43.            | Common domination and the information in exact positions                     | 219        |
| 44.            | Waiting counts and endpoint information after residual-time coarsening       | 222        |
| 45.            | Intrinsic normal–tangent geometry of the count–endpoint limit                | 225        |
| 46.            | The endpoint critical experiment                                             | 227        |
| 47.            | The Abel geometry of the boundary invariant                                  | 230        |
| 48.            | The inverse autoconvolution problem                                          | 233        |
| 49.            | Regularized observation of the complete invariant                            | 233        |
| <b>Part 5.</b> | <b>Finite-resolution inversion and global physical reconstruction</b>        | <b>238</b> |
| 50.            | Quantitative reconstruction from quantized transverse laws                   | 238        |
| 51.            | Physical coordinates and finite-window observation                           | 249        |
| 52.            | Two-flight determination and the finite observation design                   | 252        |
| 53.            | The signed one-flight benchmark and the restriction to long bridges          | 257        |
| 54.            | Direct geometric sampling in the same long-even experiment                   | 258        |
| 55.            | From noisy finite jets to global analytic table loss                         | 260        |
| 56.            | Common physical observations and charged channel calibration                 | 262        |
| 57.            | Single-offset statistical separation and global physical reconstruction      | 265        |
| 58.            | Compatible analytic variations and uniform finite positive designs           | 268        |
| <b>Part 6.</b> | <b>Further formulations and complete comparison proofs</b>                   | <b>271</b> |
| Appendix A.1.  | Further formulations and comparisons                                         | 271        |
| Appendix A.2.  | Collected theorem statements and comparisons                                 | 284        |
| Appendix A.3.  | Information beyond the leading hierarchy                                     | 293        |
| Appendix A.4.  | Statistical determination of an analytic contact                             | 295        |
| Appendix A.5.  | Finite-flight contact recovery and coefficient sensitivity                   | 299        |
| Appendix A.6.  | Reconstruction of boundary profiles from physical counts                     | 301        |
| Appendix A.7.  | Acquisition from integrated physical flux                                    | 305        |
| Appendix A.8.  | Calibration within the smooth physical experiment                            | 308        |
| Appendix A.9.  | Critical boundary experiments and support separation                         | 312        |
| Appendix A.10. | Curvatures invisible to all threshold locations                              | 318        |
| Appendix A.11. | A sharp pairwise inverse through a triple collision                          | 321        |
| Appendix A.12. | Three amplitudes in the constrained physical family                          | 324        |
| Appendix A.13. | Exact inversion from four fixed physical windows                             | 326        |
| Appendix A.14. | Joint curvature and onset accuracy                                           | 329        |
| Appendix A.15. | Reconstruction with unknown smooth remainder functions                       | 331        |
| Appendix A.16. | The sharp cost of unknown smooth remainders                                  | 333        |
| Appendix A.17. | Relation to action, spectral, and statistical inverse problems               | 336        |
| Appendix A.18. | Conditional endpoint metrics and curvature observability                     | 338        |
| Appendix A.19. | Leading-data factorization and quantitative acquisition                      | 342        |
| Appendix A.20. | Local calibration from the onset probabilities                               | 344        |
| Appendix A.21. | A sharp stability scale at the circular configuration                        | 347        |
| Appendix A.22. | Measurable reconstruction on compact parameter sets                          | 349        |
| Appendix A.23. | Locally calibrated reconstruction from unlabelled counts                     | 350        |
| Appendix A.24. | Physical indistinguishability in local count designs                         | 352        |
| Appendix A.25. | Count-only reconstruction from a fixed local bracket                         | 354        |
| Appendix A.26. | Marked onsets and source domains                                             | 356        |

|                |                                                           |     |
|----------------|-----------------------------------------------------------|-----|
| Appendix A.27. | Record topology and differentiated moving tests           | 357 |
| Appendix A.28. | The circular collision hierarchy                          | 359 |
| Appendix A.29. | Stationary flux and the complete marked event             | 361 |
| Appendix A.30. | Localization to actual alternating itineraries            | 362 |
| Appendix A.31. | A boundary-value calculus uniform in the itinerary length | 363 |
| Appendix A.32. | Uniform threshold integration and its response            | 365 |
| Appendix A.33. | The joint impact record and physical sampling cuts        | 367 |
| Appendix A.34. | Mechanical consequences of the threshold coefficients     | 368 |
| Appendix A.35. | Alternating-channel formulation and relative comparison   | 370 |
| Appendix A.36. | Graph and support coordinates                             | 372 |
| References     |                                                           | 372 |

## 1. INTRODUCTION

Does a conditional law supported near one hyperbolic periodic orbit retain the reflecting boundary as a smooth function, or only its Taylor series at the orbit? The distinction is substantial. The mass of a long bridge tends to zero exponentially, while its endpoints can remain in a fixed neighborhood of the contacts. Moreover, a smooth boundary can change on every such neighborhood without changing any contact jet. The main result of this paper identifies the complete smooth contact germs from phase-resolved endpoint laws at two positive excess times. It uses neither analyticity nor an initially chosen local inverse branch. A finite-preparation consequence recovers the complete real profiles on bounded geometric classes, with unsuccessful preparations included in the observation budget.

The proof has two essential interfaces. An exact mixed-twist normalization passes from stationary lengths to physical conditional laws without losing the exponentially small scale. An envelope identity for actual reflecting graphs then passes from the observed half-line actions to smooth boundary functions, including their flat parts. Neither passage follows from inversion of formal action coefficients. We state this mechanism first; analytic continuation and unregistered multichannel reconstruction are subsequent results with different data. The local-completion comparison in Section 7 identifies exactly why these laws determine visited contact germs without supplying a global exterior scattering observation.

**1.1. The marked experiment.** Fix a periodic polygon with  $r$  distinct physical contacts in a planar periodic dispersing billiard. Its lifted vertices  $q_i$ , signed unit tangents  $t_i$ , outward obstacle normals  $n_i$ , phase labels and closing translation are supplied. The polygon satisfies the specular reflection condition, has positive flight lengths and nongrazing incidence, and its open flights are clear of unintended obstacles. These are geometric marks, not measurements of curvature or boundary height. Put

$$L_i = |q_{i+1} - q_i|, \quad e_i = (q_{i+1} - q_i)/L_i, \quad \nu_i = e_i \cdot n_i > 0, \quad p_i = e_i \cdot t_i/\nu_i.$$

The unknown contacts have the graph representation

$$R_i(x) = q_i + \nu_i^{-1} \{t_i x - n_i F_i(x)\}, \quad F_i(0) = F'_i(0) = 0, \quad F''_i(0) > 0. \quad (1.1)$$

The recorded coordinate  $x = \nu_i t_i \cdot (R_i - q_i)$  is a tangent projection; recording it does not reveal  $F_i(x)$ . In particular, neither normal collision coordinates nor unknown arclength conversions are supplied to the inverse.

Throughout, charged preparations are independent draws from normalized unit-speed Liouville measure on the full phase space. The itinerary and acceptance gate, and every success/failure tag, are exact. Only accepted tangent-coordinate readouts may have the stated bounded error.

Not supplying local flux amplitudes therefore does not mean that the reset distribution or detection efficiency is unknown. Conditional-law rigidity does not use the free-area normalization; recovering free area from counts does use this specified preparation measure.

Let  $P$  be a returning oriented period, allowing the physical period to be doubled so that signed frames return, and let  $\mathcal{L}$  be its length. For  $N = nP$  flights starting at phase  $b$ , write  $W_{N,b}(u, v)$  for the stationary bridge length and set

$$E_{N,b} = W_{N,b} - n\mathcal{L} + p_b u - p_b v, \quad D_{N,b} = -W_{N,b,uv}(0, 0) > 0.$$

Two distinct fixed offsets  $0 < d_1 < d_2$  are selected in the small positive time window. The observation at offset  $d$  is the pair of tangent endpoints conditional on the specified itinerary event. It is restricted to a common square on which the physical residual time is positive. The excess  $d$  is fixed as  $N$  increases; it is not sent to zero in the long-flight limit. The exact event and full-phase normalization are specified in Theorem 2.5.

**Theorem 1.1** (Boundary determination by periodic endpoint laws). *For every clear nongrazing marked polygon as above, with positive-curvature  $C^\infty$  reflecting contacts, the following statements hold on a contact collar independent of  $N$ .*

*There are anchored smooth half-line actions  $S_b^-, S_b^+$  and positive amplitudes  $B_b^-, B_b^+$ , with  $B_b^\pm(0) = 1$ , such that, for each separately fixed finite order  $k$ ,*

$$\|E_{N,b} - S_b^- \oplus S_b^+\|_{C^k} + \left\| \frac{-W_{N,b,uv}}{D_{N,b}} - B_b^- \otimes B_b^+ \right\|_{C^k} \leq C_k \tau^N, \quad 0 < \tau < 1. \quad (1.2)$$

*In returning signed Jacobi coordinates the exact reference is*

$$D_{N,b} = \frac{\sinh \chi}{(\mathcal{M}_b)_{12} \sinh(n\chi)}, \quad \text{tr } \mathcal{M}_b = 2 \cosh \chi. \quad (1.3)$$

*The transfer matrix  $\mathcal{M}_b$  is determined by the actual geometric Jacobi operator, not additional observation data.*

*The conditional endpoint densities converge in every fixed  $C^k$  norm on positive interior squares to*

$$f_{b,d}(u, v) = Z_{b,d}^{-1} B_b^-(u) B_b^+(v) \{d - A_b(u) - C_b(v)\}, \quad d = d_1, d_2, \quad (1.4)$$

*where  $A_b(u) = S_b^-(u) - p_b u$ ,  $C_b(v) = S_b^+(v) + p_b v$  and the positive constants  $Z_{b,d}$  are unknown and may differ. Equality of the phase-resolved two-offset laws for two actual contact tuples implies equality of their complete smooth contact germs in (1.1). No curvature input, reflection symmetry, analyticity or closeness between the two candidates is required.*

*Proof.* The relative theorem and the physical measure formula are Theorems 2.3 and 2.5. Lemma 3.1 transfers them to the measured graph coordinates without treating the unknown boundary as a coordinate measurement. Proposition 3.2 extracts the actual actions. The global quadratic comparison in Theorem 4.2 and the signed finite-jet inverse in Theorem 3.4 identify the jets. Theorem 5.3, using the actual envelope once more, identifies the smooth functions. Complete proofs of these statements occupy Sections 2–5. The next subsection explains the two passages at which an absolute asymptotic or a formal inverse would be insufficient.  $\square$

**1.2. From a rare physical law to a smooth boundary.** The mixed twist is the scale of the unconditioned endpoint measure. Indeed, writing  $\beta_{N,b} = -W_{N,b,uv}/D_{N,b}$ , its density on a positive-residual square is exactly

$$\frac{D_{N,b}}{2\pi\mathcal{A}} \beta_{N,b}(u, v) \{d - E_{N,b}(u, v) + p_b u - p_b v\}, \quad (1.5)$$

where  $\mathcal{A}$  is the free cell area. The linear momenta in this formula cannot be discarded by a change of action gauge. Although  $D_{N,b}$  cancels upon conditioning, that cancellation supplies a limit only after the normalized twist  $\beta_{N,b}$  has been controlled. An absolute estimate for  $W_{N,b}$  alone supplies no such control.

The cofactor formula resolves this scale mismatch before a probability quotient is formed. In the returning signed coordinates, if  $H_N$  is the interior Hessian,  $G_N = H_N(0)^{-1}$  and  $\Delta H_N = H_N - H_N(0)$ , then

$$\log \beta_{N,b} = \sum_{i=0}^{N-1} \log \frac{-\ell_{b+i,12}(y_i, y_{i+1})}{L_{b+i}^{-1}} - \log \det(I + G_N \Delta H_N).$$

The geometric perturbation is localized near the two endpoints, so its trace norm is bounded independently of the number of interior sites. Weighted Green estimates separate the two end contributions. The result is a relative limit for the physical measure, not a quadratic conditional model to which a nonlinear correction is later added. The full construction and differentiated determinant comparison are in Section 2.

The density-to-action passage is exact and does not require a known flux amplitude. Suppressing the phase, define

$$\mathfrak{q}(u, v) = \frac{f_{d_1}(u, v) f_{d_2}(0, 0)}{f_{d_2}(u, v) f_{d_1}(0, 0)}, \quad H(u, v) = \frac{d_1 d_2 (1 - \mathfrak{q}(u, v))}{d_2 - d_1 \mathfrak{q}(u, v)}.$$

Cancellation in (1.4) gives  $H(u, v) = A(u) + C(v)$ , whence  $A(u) = H(u, 0)$  and  $C(v) = H(0, v)$ . The denominator is positive on the stipulated square. This operation uses the two exact offsets and the fixed origin, but not acceptance counts, curvature or normal height. It is applied to realizable laws; an arbitrary fitted density need not satisfy this identity.

The action-to-boundary passage is not merely a coefficient calculation. The global Schur-complement comparison determines positive curvatures without choosing a nearby inverse branch. With these quadratic data fixed, the signed response at order  $j \geq 3$  has the form

$$s_j = (I + T_j)(I - T_j)^{-1} q_j + \mathcal{R}_j, \quad (T_j z)_i = \sigma_i^j z_{i+1}, \quad |\sigma_i| < 1,$$

where  $\mathcal{R}_j$  depends only on lower graph jets. The finite stationary envelope, with its terminal contribution estimated before passage to a half-line, proves that these are derivatives of the actual boundary map.

To pass beyond the jets, interpolate the actual graphs from  $F$  to  $G$ . The integrated future-action envelope is

$$\mathcal{S}(G) - \mathcal{S}(F) = \bar{\alpha}(G - F) + K(G - F), \quad \bar{\alpha} \geq \frac{1}{2}. \quad (1.6)$$

Each term in  $K$  evaluates a difference at a later visit  $x_{b,j}(u)$ , with  $|x_{b,j}(u)| \leq a^j |u|$  for a single  $a < 1$  and, crucially, with multiplicative constant one. On functions vanishing through order  $m - 1$ , put

$$\|h\|_{m,R} = \max_b \sup_{0 < |u| \leq R} \frac{|h_b(u)|}{|u|^m}.$$

The actual chord weights give

$$\|\bar{\alpha}^{-1} K\|_{m,R} \leq \frac{6a^m}{1 - a^m} < 1 \quad (1.7)$$

for a sufficiently large finite  $m$ . Equal jets make the norm of  $G - F$  finite; equality of the full actions in (1.6) then forces that difference to vanish. The operator depends on the candidate pair, which is appropriate for a separation proof and is not a directly computable observation-only

inverse. For approximate data, low-degree polynomial alignment is made *before* the weighted estimate is used. This yields

$$\|F - G\|_{C^0([-R, R])} \leq C \max_{b,d} \|f_{b,d}^F - f_{b,d}^G\|_{C^m(Q)} \quad (1.8)$$

on the bounded positive  $C^{m+3}$  classes of Theorem 5.5, for sufficiently small realizable law distance. Thus smooth identity and finite-order stability use the same geometric mechanism, with different uniformity hypotheses.

**1.3. Finite preparations and information beyond Taylor data.** The differentiated norm in (1.8) need not be provided as an observation. Fix a nonempty bounded class  $\mathcal{C}$  of actual marked tables as in Section 6. In addition to the bounded positive  $C^{m+3}$  contact class, impose the finitely many higher graph bounds needed for the forward law through order  $m + s$ , where  $s \geq 1$ . Clearance, incidence, flight-length and residual margins, an upper free-area bound, the period and the two separated offsets are fixed. These priors give

$$\|f_N - f\|_{C^m(Q_+)} \leq C_0 e^{-\omega N}, \quad \|f_N\|_{C^{m+s}(Q_+)} + \|f\|_{C^{m+s}(Q_+)} \leq H, \quad p_N \geq c_0 e^{-\Gamma N}.$$

Here  $Q \Subset Q_+$ , acceptance is gated on  $Q_+$ , and densities normalized on  $Q_+$  are merely restricted, not renormalized, on  $Q$ . Each independent full-phase preparation returns an exact success/failure tag and, on success, an endpoint pair with recording error at most  $\delta$ . Tags, gates and offsets remain exact. There are  $J = 2r$  phase-offset groups.

**Theorem 1.2** (Complete-profile recovery with a charged budget). *In the preceding experiment there is a measurable estimator taking values in actual tables of  $\mathcal{C}$ . For  $B$  preparations in each group, set*

$$\alpha = \frac{s}{2(m+s)+2}, \quad \beta = \frac{\alpha\omega}{\omega + \alpha\Gamma}, \quad \gamma = \frac{s}{m+s+3}, \quad L_B = \log(C_* JB/\zeta).$$

*Choose the returning flight number*

$$N_B = P \left\lfloor \frac{\alpha \log(B/L_B)}{P(\omega + \alpha\Gamma)} \right\rfloor.$$

*For  $0 < \zeta < 1/2$ , under the sufficient-success, bandwidth and small-error conditions of Theorem 6.3,*

$$\sup_{T \in \mathcal{C}} \mathbb{P}_T \left( \|F^{\hat{T}} - F^T\|_{C^0([-R, R])} > C \{(L_B/B)^\beta + \delta^\gamma\} \right) \leq 2\zeta. \quad (1.9)$$

*All  $JB$  preparations, including failures, are charged. At fixed confidence the conditions hold for sufficiently large  $B$  and sufficiently small recording error. The target is the complete real profile on a class-dependent collar, not a finite jet.*

*Proof.* Lemma 6.1 gives an interior kernel estimate in the required differentiated norm. The estimator is first compared with a countable dense subset of the image of actual tables in the product  $C^m$  space. This measurable minimum-distance selection, followed by the triangle inequality, permits (1.8) to be used only between realizable limiting laws. Theorem 6.2 supplies the accepted-sample bound. The binomial success estimate and logarithmic choice of  $N_B$  in Theorem 6.3 give (1.9).  $\square$

For varying confidence, in particular, the requirements  $N_B \geq N_0$ ,  $n_B = \lfloor Bc_0 e^{-\Gamma N_B}/2 \rfloor \geq 2$  and  $Bc_0 e^{-\Gamma N_B} \geq 8 \log(J/\zeta)$  are retained; the upper bandwidth and local comparison restrictions are retained as well. The estimator is an existence construction, not an efficient search algorithm. The rate is a sufficient upper bound, not a minimax threshold. The constants and the order  $m$  are not uniform in a vanishing contraction margin or a varying period. Independence is a

property of the preparations, not a conclusion about a single unreset billiard trajectory. Bounded endpoint errors after acceptance do not cover uncertain gate membership, phase labels or offsets.

The geometric content of the recovered function can be tested within the physical class. Proposition 5.6 constructs actual closed obstacles with fixed free area and marked polygon for which all contact and action Taylor jets agree, but the two-offset laws differ on every sufficiently small contact square. The perturbation at one contact is  $\varepsilon\chi(u)e^{-1/u^2}$ , extended smoothly by zero at  $u = 0$ ; a remote area correction leaves the local actions unchanged. Its real action response is

$$S_b^{-,\varepsilon}(u) - S_b^{-,0}(u) = \varepsilon e^{-1/u^2} \left\{ 1 + O(|u|) + O(e^{-(a^{-2}-1)/u^2}) \right\}.$$

At any fixed nonzero  $u_*$  where  $\chi = 1$ , the profile separation is at least  $\Delta = \varepsilon e^{-1/u_*^2}$ . When the error radius in (1.9) is less than  $\Delta/3$ , nearest-candidate comparison gives a finite charged test with error at either table at most  $2\zeta$ . This separates complete-law information from the formal all-order contact record. It makes no claim about equality of normalized density jets or of full marked length spectra, and it does not give a uniform inexpensive resolution as  $u_* \rightarrow 0$ .

**1.4. One law and an unknown positive offset.** For an oblique marked polygon the calibration can be reduced further. At phase  $b$  let  $d_b > 0$  be fixed during the experiment but unreported; retain just the one conditional law  $f_{b,d_b}$ . The first derivatives of its physical actions are known from the polygon:  $A'_b(0) = -p_b$ ,  $C'_b(0) = p_b$ . They are not unknown curvature data. The anchored quotient of this single law satisfies

$$\mathcal{R}_b(u, v) = \frac{f_{b,d_b}(u, v)f_{b,d_b}(0, 0)}{f_{b,d_b}(u, 0)f_{b,d_b}(0, v)} = 1 - \frac{A_b(u)}{d_b - A_b(u)} \frac{C_b(v)}{d_b - C_b(v)}.$$

In particular  $\partial_{uv}\mathcal{R}_b(0, 0) = p_b^2/d_b^2$ . When  $p_b \neq 0$ , the law determines its own offset; the marked slopes then fix the remaining scalar ambiguity of the two factors.

**Theorem 1.3** (One-law determination without supplied offsets). *In the setting of Theorem 1.1, suppose the marked polygon is oblique at every phase. One limiting endpoint law per phase, at a fixed but unreported positive offset  $d_b$ , determines the offsets and the complete smooth reflecting contact germs. No second offset, acceptance probability, flux amplitude or curvature is supplied. On uniformly oblique bounded positive classes, the joint inverse obeys (8.10) in the same finite  $C^m$  law topology. Independent finite preparations yield (8.11), with  $J = r$  groups, estimating both profiles and offsets with failures charged. Every clear periodic polygon with at least three distinct physical contacts is oblique, so the assertion applies to every such polygon.*

*Proof.* Proposition 8.1 identifies  $d_b = |p_b|/\sqrt{\partial_{uv}\mathcal{R}_b(0, 0)}$ . Theorem 8.2 separates both actual actions, and Theorem 8.3 applies the smooth envelope inverse. Its fixed-slice reconstruction is  $C^m$ -stable without an additional derivative loss. Theorem 8.5 gives the joint finite-preparation statement, including its explicit budget and error conditions. Lemma 8.4 proves the last assertion.  $\square$

Here “unreported” means a fixed nuisance offset in a known positive interval, possibly different at different phases. The acceptance event uses its actual value, and the success tag, itinerary label and endpoint gate are still exact. The theorem does not cover random clock drift or uncertain gating. The marked polygon and frames are still supplied. At normal incidence the displayed mixed derivative vanishes; the general two-offset theorem and the retained known-offset symmetric inverse continue to apply under their respective hypotheses.

This consequence uses the linear physical momenta which an action gauge must not erase. It removes a second observation and numerical offset calibration from the nonsymmetric problem while retaining the same relative physical limit and actual functional inverse. The unknown



normalization and endpoint amplitudes are eliminated algebraically; that elimination alone would not prove either of the analytic interfaces.

**1.5. What is determined when failures are retained.** The same mechanism identifies the fibers of the local experiment, not only an injective contact coordinate map. Let  $\pi_{N,b}$  be the success probability and let  $f_{N,b}$  be the conditional density on its specified gate. The exact phase-volume formula has the central anchor

$$\pi_{N,b} f_{N,b}(0,0) = \frac{D_{N,b} d_b}{2\pi \mathcal{A}},$$

where  $\mathcal{A}$  is the free area of the marked cell. This formula uses independent normalized unit-speed Liouville preparations and exact success tags; it is not an unknown-reset-distribution identity. Both the offset and the Jacobi normalization on the right are determined by the limiting laws. Thus the unconditioned record recovers the one global scalar which conditioning had canceled.

**Theorem 1.4** (The complete charged local invariant). *Under the hypotheses of Theorem 1.3, the germs of the charged reset experiments determine exactly the smooth reflecting contact germs, the fixed offsets and the free cell area. Equality of these three quantities is also sufficient for equality of all sufficiently localized finite-flight records. Conditional limiting laws alone correspond to the first two quantities. One phase's normalized success prefactor, together with these laws, recovers the area.*

*On the bounded classes of Theorem 8.5, a measurable estimator taking values in actual table–offset pairs recovers all three components from  $rB$  charged preparations. With the explicit conditions and returning-length choice in Corollary 9.6, its profile, offset and logarithmic area error is at most*

$$Ct \log(e/t), \quad t = (L_B/B)^\beta + \delta^\gamma,$$

*outside an event of probability at most  $2\zeta$ . Here  $\beta, \gamma$  are the exponents in Theorem 8.5. The count used for area recovery is part of the same record, not an extra preparation group.*

*Proof.* Proposition 9.1 proves the exact anchor. Theorem 9.2 combines it with smooth contact recovery and exact locality to identify both directions of the observation fibers. Lemma 9.3 retains the flight-length amplification of curvature error in the rare normalization. Theorem 9.5 and Corollary 9.6 give the jointly realizable finite construction and its displayed rate.  $\square$

The conclusion separates a local geometric invariant from a global normalizing invariant within one physical experiment. Equal-area remote deformations remain exact fibers of the record. The theorem therefore recovers the full object identifiable by this observation without converting a contact theorem into a whole-table assertion. Its content is the combination of relative physical normalization, actual smooth separation and the unconditioned anchor; the scalar identity by itself is not a substitute for either analytic interface.

**1.6. First-moment profiles and differentiated reconstruction.** The limiting law also permits a reduction in smoothing dimension without a change in the physical record. For a fixed positive gate  $J^2$ , let  $M_0, M_1$  be its marginal density in  $u$  and its first  $v$ -moment profile, and let  $N_0, N_1$  be the analogous functions in  $v$ . The same accepted pairs give the matrix

$$H_f = \mathbb{E}_f \left[ \begin{pmatrix} 1 \\ U \end{pmatrix} (1 \quad V) \right].$$

The opposite marked slopes of the two half-line actions imply  $\det H_f > 0$  on a sufficiently small common gate. The full limiting law is then recovered by

$$f(u, v) = (M_0(u), M_1(u)) H_f^{-1} (N_0(v), N_1(v))^T.$$

This is not reconstruction from ordinary marginals alone. The weighted profiles use paired coordinates, and all moments use the full success gate with its original normalization.

**Theorem 1.5** (Recovery by first-moment profiles). *In the oblique bounded class of Theorem 1.4, the limiting moment profiles and their matrix determine the same contact germs and fixed offsets as the complete limiting law. The charged counts additionally determine the free area. On every fixed class and small common gate as specified in Section 10, the reconstruction of the density from these profiles is Lipschitz in  $C^m$ , without loss of derivatives.*

*Under the same finite forward smoothness assumptions, a measurable estimator with values in actual table–offset pairs satisfies the charged estimates of Theorem 9.5, with its smoothing exponents improved to*

$$\alpha_1 = \frac{s}{2(m+s)+1}, \quad \gamma_1 = \frac{s}{m+s+2}, \quad \beta_1 = \frac{\alpha_1 \omega}{\omega + \alpha_1 \Gamma}.$$

*More precisely, under the count, bandwidth and small-error conditions of Theorem 10.7 and Corollary 10.8,  $rB$  charged resets give local profile and offset error at most  $Ct$  and joint local-and-logarithmic-area error at most  $Ct \log(e/t)$ , where  $t = (L_B/B)^{\beta_1} + \delta^{\gamma_1}$ , with probability at least  $1 - 2\zeta$ . Preparations are normalized unit-speed Liouville resets and success tags are exact.*

*Proof.* Lemma 10.1 proves the moment identity and its  $C^m$  stability. Proposition 10.2 derives its determinant floor from the signed geometric slopes, rather than postulating a nondegenerate statistical model. Corollary 10.4 then applies the actual smooth inverse and the charged area anchor. Proposition 10.5 retains the finite-flight rank defect. Lemma 10.6 proves the weighted one-dimensional kernel bound, and Theorem 10.7 fits the finite image of actual pairs. Corollary 10.8 gives the budget-dependent flight length.  $\square$

The low-rank matrix identity is classical in character; the geometric content here is its forced nondegeneracy, its compatibility with the relative physical limit, and its use before the smooth contact inverse. The finite density is not assumed to have exact rank two. The inherited normal-incidence two-offset theorem and the two-dimensional estimator remain separate valid results, and the improved upper bounds do not assert statistical optimality.

Remark 10.3 gives the explicit affine model  $\det H_f = p^2 R^4 / (9d^2)$  on  $[-R, R]^2$ . It illustrates the dependence on the fixed gate and obliquity; it is not an actual-table realization or a uniform estimate in a shrinking physical collar.

**1.7. Relation to lens and spectral rigidity.** A second comparison is with smooth obstacle scattering. In the corrected planar theorem of Noakes–Stoyanov [53, Theorem 1.1], a finite disjoint union of strictly convex compact domains with  $C^3$  boundaries is determined globally by almost-everywhere travelling times, or by almost-everywhere scattering length spectra. For the former data, one fixes an enclosing circle and records the exit time for each incoming position and direction. For the latter, the data are sets of sojourn times indexed by the incoming and outgoing directions. These are labelled families, not an unlabelled list of lengths. The cited paper supplies a separate two-dimensional proof; its earlier higher-dimensional argument alone does not cover the planar case. Smooth recovery from rich scattering data is therefore not a new category of rigidity supplied by this paper.

The distinctions relevant to Theorem 1.1 are concrete. The exterior geometry there is a finite compact obstacle, whereas here it is a periodic table with a selected recurrent itinerary. An enclosing Euclidean observation frame in the scattering problem does not supply internal reflecting contacts. Our vertices, tangent frames, normals and phase labels are supplied, while the reflecting graphs remain unknown. Our data are two conditional endpoint distributions at each marked phase, not exit times indexed by all exterior incoming rays. Spatially, our smooth conclusion is the collection of contact germs, not the boundaries away from the visited collars.

The smoothness and quantitative priors required for our finite experiment are also stronger than the  $C^3$  uniqueness hypothesis just cited. Thus neither regularity nor spatial extent provides a dominance claim over that theorem.

There is an exact local reason for these differences, not merely a terminological one. Proposition 7.1 proves that all finite-flight laws of the specified experiment are unchanged when the unvisited boundary is altered while the local flight geometry is fixed. With free area fixed, this includes the success probabilities and the failure atom, not only the conditional densities. Conversely the two offset limits determine the complete contact germs by the smooth inverse. Corollary 7.3 realizes both assertions in nonconstant equal-area families of closed obstacles. On a common compact-obstacle realization, their exterior scattering data distinguish the different completions. Equality of our local observation therefore does not supply the global-data equality hypothesis of the planar lens theorem. No assertion in the reverse direction, or of two-way incomparability, is needed.

The contribution is consequently a local sufficiency statement at a singular observation scale. The cofactor identity alone describes a linear reference; the two-ended determinant estimate must show that the *nonlinear physical numerator* remains relatively controlled after its exponentially small reference has been removed. Once the resulting laws yield actual half-line actions, coefficient inversion still leaves every flat perturbation unresolved. The summable stationary envelope and constant-one contraction of actual visits supply the functional step that identifies that remaining information. The finite experiment reaches the finite differentiated topology used by this step while charging the rare successes. This chain, rather than smooth rigidity in isolation, is the mathematical gain. Its ingredients and observation hypotheses are not interchangeable with a supplied global scattering relation. The locality comparison identifies the object determined by the chain; it is not an additional claim of global reconstruction or an independent statistical method.

The comparison with marked length rigidity is a comparison of inverse problems, not an inclusion of data sets. Bálint–De Simoi–Kaloshin–Leguil [1] recover two-periodic contact curvature and marked Lyapunov data from the marked length spectrum of open dispersing billiards. Their Theorem D and Corollary E treat arbitrary periods; Remark 2.3 identifies a limitation of the particular asymptotic relations used there in separating higher-period contact parameters. That remark is not an obstruction theorem for every use of the full spectrum. Here the phase-resolved endpoint experiment supplies different information, and its polygon and frames are fixed marks.

De Simoi–Kaloshin–Leguil [5] determine analytic open billiards from marked lengths under their non-eclipse, symmetry and genericity assumptions. Finamore–Leguil [12] use an enriched marked length spectrum in the finite-horizon Sinai setting. Smooth contact determination from (1.4) removes analyticity in the present observation problem; it does not remove hypotheses from those theorems or assert same-data dominance. The corrected version of Florio–Leguil [11] retains dynamical conjugacy conclusions, not the removed geometric spectral-rigidity assertion. These comparisons do not establish an exhaustive priority claim.

The statistical ingredients in Theorem 1.2 are classical: interior kernel approximation, countable minimum-distance selection and concentration. For uniform kernel estimates see, for example, Giné–Guillou [52]; the exact derivative, gate and bounded-readout estimate needed here is proved in Lemma 6.1. The acquisition theorem changes the admissible input from an exact law to a finite charged record. The geometric identification it uses is the relative physical law together with the actual contracted-visit envelope, rather than a new general method of statistical regularization.

Several further conclusions are developed in the paper without being premises of Theorem 1.1. The analytic Banach inverse controls complete holomorphic germs on a small disc and requires its analytic norm. Whole-obstacle continuation requires connected closed analytic boundaries, fixed placement and visitation of every obstacle. The unregistered alternating-channel problem

instead has weaker local marks, a single-offset extraction and separate finite-channel matching hypotheses; it yields the stated lattice reconstruction and finite ambiguity. The finite-jet stopping estimates retain the propagated curvature error  $C_M(\tau^N + \varepsilon) + L_M E_m$ . None of these statements is combined with smooth contact recovery to assert an assumption-free whole-table inverse.

**1.8. Organization of the proof.** The central route is printed first. Section 2 constructs the periodic Jacobi problem, proves the relative law and derives the exact physical measure. Section 3 identifies measured coordinates and proves two-offset separation, the actual stationary envelope and signed finite-jet inversion; its analytic norm theorem is a further branch of the argument. Section 4 supplies the global positive quadratic comparison and the corrected stopping estimate. Section 5 proves constant-one visit contraction, smooth separation, polynomially aligned stability and the actual flat family. Section 6 then proves the finite experiment. Section 7 proves exact locality of the charged record and gives equal-area completions that separate this observation from exterior scattering data. In particular, the global analytic and lattice results are not hidden inputs to the smooth theorem. Section 8 proves the single-law offset identification, its smooth consequence and the joint charged recovery theorem; it does not replace the two-offset statements at normal incidence. Section 9 identifies the complete invariant of the charged local experiment and reconstructs its area together with its contacts and offsets. The proof keeps the rare-normalization amplification explicit and selects all quantities from one actual table.

Section 10 reconstructs the limiting law from four first-moment profiles, retains the finite-flight rank defect, and proves the charged one-dimensional derivative estimate and its budget balance. Remark 10.3 makes the fixed-gate conditioning restriction explicit.

The subsequent parts contain the finite-chain foundations, the more specialized alternating formulas, and the analytic and multichannel consequences with their complete proofs. The appendices retain the broader formulations, comparison arguments and coordinate conversions. The complete technical manuscript additionally contains the calibration, coarsening and acquisition catalogue. These complementary inverse problems are organized by their observations and targets, rather than counted as separate proofs of the central contact theorem.

## Part 1. Periodic boundary laws and smooth contact determination

### 2. RELATIVE BOUNDARY LAWS ALONG A PERIODIC ITINERARY

The relative normalization is not specific to a normal two-periodic orbit. We prove it for a selected nongrazing periodic orbit of a smooth planar dispersing billiard. The two ends now carry different half-line actions. There is also a linear endpoint contribution to the physical length; it must be restored before a probability law is formed. The signed two-contact inverse in the subsequent sections retains its alternating-channel hypotheses.

**2.1. The geometric Jacobi operator.** Let a periodic orbit have a fixed itinerary, with lifted reflection points  $q_i$ , flight lengths  $L_i > 0$ , curvatures  $\kappa_i > 0$ , and outgoing normal cosines  $\nu_i > 0$ . The indices are periodic modulo an even integer  $P$ : an odd primitive period is doubled so that the signed coordinate frames also return. A lifted period is allowed to end at a lattice translate of its initial point. Put  $L = \sum_{i=0}^{P-1} L_i$ . Every open flight is clear of unintended obstacles. We work with a compact smooth family of such marked orbits, with fixed  $P$  and positive lower bounds for clearance,  $L_i$ ,  $\kappa_i$  and  $\nu_i$ , and finite upper bounds for the flight lengths and each separately used smooth norm. Clearance here refers to portions outside small fixed endpoint neighborhoods; nongrazing incidence controls those neighborhoods. Family derivatives keep the scaled endpoint coordinates and the excess  $d = t - nL$  fixed; the physical observation time moves with the reference period length.

Choose unit-speed signed boundary charts  $r_i(s)$  centered at  $q_i$ . Their tangent orientations alternate relative to the positively oriented obstacle tangents. Write

$$x_i = \nu_i s_i, \quad \ell_i(x, y) = |r_{i+1}(y/\nu_{i+1}) - r_i(x/\nu_i)|, \quad k_i = L_i^{-1}, \quad m_i = 2\kappa_i/\nu_i.$$

The endpoint arguments in  $r_i$  are signed arclengths. Reflection gives numbers  $p_i$  such that

$$\ell_{i,1}(0, 0) = -p_i, \quad \ell_{i-1,2}(0, 0) = p_i.$$

Introduce the normalized edge action

$$h_i(x, y) = \ell_i(x, y) - L_i + p_i x - p_{i+1} y. \quad (2.1)$$

Thus  $h_i(0, 0) = 0$  and  $Dh_i(0, 0) = 0$ . The linear terms telescope; this normalization does not change an interior reflection equation or a mixed endpoint derivative.

**Lemma 2.1** (Coercivity and persistence). *The Hessian of an edge at the reference orbit and the interior Jacobi operator are respectively*

$$D^2 h_i(0, 0) = \begin{pmatrix} k_i + m_i/2 & -k_i \\ -k_i & k_i + m_{i+1}/2 \end{pmatrix}, \quad (H^0 x)_i = (k_{i-1} + k_i + m_i)x_i - k_{i-1}x_{i-1} - k_i x_{i+1}. \quad (2.2)$$

*They are uniformly positive. On a sufficiently small common coordinate box the interior Hessians are symmetric, tridiagonal, and strictly diagonally dominant with a uniform positive margin. The marked orbit and these properties persist under sufficiently small smooth perturbations of the reflecting arcs and of the lattice realization, with the lifted itinerary fixed.*

*Proof.* For a chord with unit direction  $e$  and endpoint tangent  $t$ , the same-end second derivative in signed arclength is  $(1 - (e \cdot t)^2)/L$  plus the curvature contribution. At an outgoing reflection  $e \cdot n = \nu$ , with  $n$  the outward obstacle normal and  $r'' = -\kappa n$ , that contribution is  $\kappa\nu$ ; the incoming end has the same sign. The mixed derivative is  $-[t_i \cdot t_{i+1} - (e \cdot t_i)(e \cdot t_{i+1})]/L$ . The alternating tangent orientations make the two projections onto  $e^\perp$  have the same sign, so it is  $-\nu_i \nu_{i+1}/L_i$ . Dividing the endpoint derivatives by the respective normal cosines proves the edge formula. Summing adjacent edge derivatives proves the second formula.

The edge quadratic form is  $k_i(x - y)^2 + (m_i/2)x^2 + (m_{i+1}/2)y^2$ . The finite and half-line interior forms are therefore bounded below by  $m_* \sum x_i^2$ , where  $m_* := \min_i m_i > 0$ . Continuity, finite periodicity, and compactness preserve a positive row-dominance margin on one smaller box. For persistence, apply the finite-dimensional implicit function theorem to the reflection equations with the prescribed periodic identification of the endpoints. The periodic Hessian has quadratic form  $\sum k_i(x_{i+1} - x_i)^2 + \sum m_i x_i^2$ , hence is invertible. The continued orbit remains in its charts; the clearance and incidence margins keep the chords physical. This argument also covers a repeated period and the prescribed deck translation.  $\square$

## 2.2. The exact reference twist. Set

$$Q_i = \begin{pmatrix} 1 + m_i/k_i & 1/k_i \\ m_i & 1 \end{pmatrix}, \quad \mathcal{M}_b = Q_{b+P-1} \cdots Q_b, \quad \text{tr } \mathcal{M}_b = 2 \cosh \chi. \quad (2.3)$$

Each matrix  $Q_i$  has determinant one and positive entries. The products for different starting phases are conjugate. Since  $(\mathcal{M}_b)_{12}(\mathcal{M}_b)_{21} > 0$  and  $\det \mathcal{M}_b = 1$ , their trace is greater than two. In particular  $\chi > 0$ , and  $\chi$  has a positive lower bound on the compact class.

For  $N = nP$ , let  $W_{N,b}(u, v)$  be the stationary length along  $N$  flights, starting at phase  $b$ , with scaled endpoint coordinates  $u, v$ . Its construction on a fixed box is given below. Define

$$D_{N,b} = -W_{N,b,uv}(0, 0) = \frac{\sinh \chi}{(\mathcal{M}_b)_{12} \sinh(n\chi)}. \quad (2.4)$$

In particular  $D_{N,b} > 0$ , and its decay is comparable to  $e^{-n\chi}$ . This formula is exact, rather than an unspecified exponential equivalent.

**Lemma 2.2** (Transfer and cofactor normalizations). *Formula (2.4) holds. More generally, if  $y$  is the stationary segment and  $H_N(y)$  its interior Hessian, then*

$$-W_{N,b,uv}(u, v) = \frac{\prod_{i=0}^{N-1} [-\ell_{b+i,12}(y_i, y_{i+1})]}{\det H_N(y)}. \quad (2.5)$$

*Proof.* For the reference recurrence write  $z_i = k_{i-1}(x_i - x_{i-1})$ . Then  $z_{i+1} = z_i + m_i x_i$  and  $x_{i+1} = x_i + z_{i+1}/k_i$ , which is precisely the transfer  $Q_i$ . At the initial endpoint  $z_0$  can be defined as  $k_b(x_1 - u) - m_b u$ ; no exterior reflection equation is being imposed. The left boundary derivative of the quadratic stationary action is  $-z_0 - m_b u/2$ . The relation  $v = (\mathcal{M}_b^n)_{11}u + (\mathcal{M}_b^n)_{12}z_0$  therefore gives  $-W_{uv}(0, 0) = 1/(\mathcal{M}_b^n)_{12}$ . Cayley–Hamilton, or diagonalization at the distinct eigenvalues  $e^\chi$  and  $e^{-\chi}$ , gives

$$(\mathcal{M}_b^n)_{12} = (\mathcal{M}_b)_{12} \frac{\sinh(n\chi)}{\sinh \chi}.$$

The first assertion follows. Eliminating the interior variables from the tridiagonal Hessian gives the mixed Schur-complement entry. The corner cofactor is the product of the adjacent mixed entries, with the sign in (2.5). This proves the second assertion, including its normalization at zero.  $\square$

**2.3. A relative theorem on a fixed collar.** Write

$$E_{N,b}(u, v) = W_{N,b}(u, v) - nL + p_b u - p_b v, \quad \beta_{N,b}(u, v) = \frac{-W_{N,b,uv}(u, v)}{D_{N,b}}.$$

The superscripts  $-$  and  $+$  below distinguish the left future tail from the right past tail. They do not indicate signs of the functions.

**Theorem 2.3** (Periodic-itinerary relative boundary law). *For the marked periodic orbits above there is a coordinate interval, independent of  $n$ , on which the finite stationary segments exist uniquely within the prescribed small itinerary collars. There are smooth strictly convex functions  $S_b^-, S_b^+$  and positive smooth functions  $B_b^-, B_b^+$ , with*

$$S_b^\pm(0) = (S_b^\pm)'(0) = 0, \quad B_b^\pm(0) = 1,$$

*such that, for one  $\tau \in (0, 1)$  and each fixed  $k$ ,*

$$\|E_{N,b} - S_b^- \oplus S_b^+\|_{C^k} + \|\beta_{N,b} - B_b^- \otimes B_b^+\|_{C^k} \leq C_k \tau^N, \quad N = nP. \quad (2.6)$$

*The norms include each stipulated fixed number of smooth family derivatives. Constants may depend on the corresponding higher norms; no order-uniform smooth estimate is asserted. The action error has zero value and endpoint differential at zero and is  $O(\tau^N(|u| + |v|)^2)$ . The amplitude error vanishes at zero. For analytic families all four limits are analytic on smaller common complex collars.*

*Proof.* We give the Green estimates, the construction, and the relative normalization explicitly. None of them uses a two-dimensional normalizing chart.

Put  $a_i = k_{i-1} + k_i + m_i$ , and split the reference interior operator as  $H^0 = A(I - T)$ , where  $A = \text{diag}(a_i)$  and the two nonzero entries in row  $i$  of  $T$  are  $k_{i-1}/a_i$  and  $k_i/a_i$ . On a finite interval or a half-line,

$$\|T\|_\infty \leq q < 1, \quad (H^0)^{-1} = \sum_{r \geq 0} T^r A^{-1}, \quad |G_{ij}| \leq Cq^{|i-j|}.$$

The last estimate holds because a nearest-neighbor walk from  $i$  to  $j$  has length at least  $|i - j|$ . Symmetry also gives uniform  $\ell^1$  bounds. The same series is bounded on the exponentially weighted sup norm for any weight  $\rho \in (q, 1)$ , by summing  $q^{|i-j|}\rho^{j-i}$ . For a finite segment one uses the sum of the two endpoint weights  $\rho^i + \rho^{N-i}$ .

These estimates include fixed derivatives in the marked family. A derivative of order  $r$  of a product representing a walk of length  $h$  introduces at most a fixed polynomial in  $h$ , times a bound  $C_r q^{h-r}$ . Alternatively one may differentiate the resolvent identity. A strict enlargement of the exponential base absorbs that polynomial. The finite Green kernel and the left half-line kernel also satisfy

$$|G_{ij}^N - G_{ij}^-| \leq Cq^{2N-i-j} \quad (i, j < N), \quad |G_{ij}^N| \leq Cq^{|i-j|}. \quad (2.7)$$

Indeed a walk contributing to their difference must reach the removed right boundary and return. Its length is at least  $2N - i - j$ . The reflected right version and the fixed differentiated versions follow in the same way, with a slower base when necessary. The diagonal coefficients on retained sites are identical; imposing a Dirichlet boundary deletes walks, not those diagonal coefficients.

Solve the finite reflection equations as their reference linear solution minus  $G^N$  times the nonlinear gradient remainder. That remainder is local, vanishes quadratically at zero, and has derivative  $O(r)$  in a box of radius  $r$ . The weighted Green bounds make this a contraction with constant at most  $1/2$  on a ball of radius  $C(|u| + |v|)$  when the common endpoint interval is sufficiently small. This constructs the physical segment with

$$|y_i| \leq C(|u|\rho^i + |v|\rho^{N-i}).$$

Strict convexity gives uniqueness among all segments in the small box, not just in the contraction ball. The same construction on the future and past half-lines gives  $x_i^-(u)$  and  $x_i^+(v)$ , decaying as  $C|u|\rho^i$  and  $C|v|\rho^i$ . Differentiating the fixed-point equation leaves the same invertible operator on its highest derivative; the other terms are local products of lower derivatives. Induction therefore gives the weighted bounds at every fixed order.

Define the half-line actions by the absolutely convergent sums

$$\begin{aligned} S_b^-(u) &= \sum_{i \geq 0} h_{b+i}(x_i^-(u), x_{i+1}^-(u)), \\ S_b^+(v) &= \sum_{i \geq 0} h_{b-i-1}(x_{i+1}^+(v), x_i^+(v)). \end{aligned} \quad (2.8)$$

Every summand is quadratic in its exponentially decaying arguments. The fixed differentiated sums converge as well. In the first variation of a finite truncation the interior terms cancel. The remaining remote boundary term tends to zero by the weighted estimates, so these sums have the stated endpoint derivatives. The positive edge forms in Lemma 2.1 bound each minimized half-line energy below by a fixed positive multiple of its boundary value squared. Their Hessians at zero are consequently positive. Shrinking the interval proves strict convexity and the stated anchoring.

For the action comparison, glue the two half-lines by  $z_i = x_i^-(u) + x_{N-i}^+(v)$  and overwrite the two endpoint values by  $u, v$ . Locality and  $Dh_i(0, 0) = 0$  imply that the interior residual has  $\ell^1$  norm at most  $C(1 + N)\rho^N$ : every interaction contains one factor from each end. The endpoint overwrites have the same order. The averaged Hessian between  $z$  and the finite stationary segment has uniformly bounded  $\ell^1$  inverse by diagonal dominance and symmetry. Hence  $\|y - z\|_1 \leq C(1 + N)\rho^N$ . Subtracting and differentiating these equations proves the same assertion at each fixed order, with only a polynomial loss in  $N$ . Summing edge interactions and the omitted half-line tails now proves the action part of (2.6).

For the amplitude, let  $G_N = (H_N(0))^{-1}$  and  $\Delta H_N = H_N(y) - H_N(0)$ . The exact cofactor identity gives

$$\log \beta_{N,b} = \sum_{i=0}^{N-1} \log \frac{-\ell_{b+i,12}(y_i, y_{i+1})}{k_{b+i}} - \log \det(I + G_N \Delta H_N). \quad (2.9)$$

In particular the exponentially small  $D_{N,b}$  has already been cancelled exactly. Tridiagonality and endpoint localization give  $\|\Delta H_N\|_1 \leq C(|u| + |v|)$ , where the norm is the trace norm: the sum of absolute matrix entries is an upper bound for that norm. All its fixed derivatives have uniformly bounded trace norms. Take the endpoint interval small enough that  $\|G_N \Delta H_N\| < 1/2$ , uniformly in  $N$ .

On the future half-line define  $\Delta H_b^-(u)$  by the same rule, and put

$$\begin{aligned} \log B_b^-(u) &= \sum_{i \geq 0} \log \frac{-\ell_{b+i,12}(x_i^-(u), x_{i+1}^-(u))}{k_{b+i}} \\ &\quad - \sum_{r \geq 1} \frac{(-1)^{r+1}}{r} \operatorname{tr}[(G_b^- \Delta H_b^-(u))^r]. \end{aligned} \quad (2.10)$$

The analogous reversed formula defines  $B_b^+$ . Each edge sum is absolutely convergent. The perturbation is trace class and the operator norm is less than  $1/2$ , so the trace series and its fixed derivatives converge. Its exponential is positive for real arguments, and equals one at zero. This specifies the determinant normalization completely.

Here is the two-end comparison for (2.9). For  $N \geq 6$  retain the first and last  $r = \lfloor N/3 \rfloor$  interior sites, and delete all other perturbation entries. The trace-norm cost is  $C_k(1+N)^{a_k} \rho^r$  after  $k$  fixed derivatives. On either retained block the gluing estimate replaces the finite segment by its own half-line with exponentially small entrywise sum. The identity  $\det(I + AB) = \det(I + BA)$  reduces a finitely supported perturbation to those two blocks. Equation (2.7) compares the diagonal Green compressions to their half-line compressions; the opposite-end Green blocks are  $O(q^{N-2r})$ , up to a polynomial factor in  $N$ . Multiplication by the bounded trace-norm perturbation preserves these bounds. Extending each retained block to its half-line costs  $O(\rho^r)$  again.

For trace-class  $T, \tilde{T}$  with norms at most  $q_1 < 1$ , the exact telescoping expansion of their powers gives

$$|\operatorname{tr}(T^h - \tilde{T}^h)| \leq h q_1^{h-1} \|T - \tilde{T}\|_1.$$

Summing with coefficient  $1/h$  proves the required continuity of the logarithmic determinant. At each fixed differentiated order the series has an additional fixed polynomial in  $h$ , still summable. Thus the finite logarithm tends relatively to the sum of the two logarithms in (2.10). The edge-product comparison follows from the same gluing and tail bounds. Exponentiation proves the amplitude part of (2.6).

Choose one  $\tau < 1$  slower than the Green, gluing, and block-truncation rates. The strict margins absorb every separately fixed polynomial loss; the constants, not this choice of  $\tau$ , depend on the order. Finitely many shorter segments are absorbed by enlarging those constants. At zero the actions have zero value and differential, and the normalized twists and amplitudes equal one. Taylor's integral remainder gives the additional quadratic vanishing. Finally the contraction, summable perturbations, and logarithmic series extend to uniform smaller complex boxes for analytic families. Their locally uniform convergence proves the analytic assertion.  $\square$



**2.4. Nonlinear transmission and the stable return.** Define the connected endpoint action and the two increasing coordinates by

$$C_{N,b}(u, v) = W_{N,b}(u, v) - W_{N,b}(u, 0) - W_{N,b}(0, v) + W_{N,b}(0, 0),$$

$$\zeta_b^\pm(x) = \int_0^x B_b^\pm(t) dt.$$

**Corollary 2.4** (Relative nonlinear transmission). *On a smaller fixed interval,*

$$-\frac{C_{N,b}(u, v)}{D_{N,b}} = \zeta_b^-(u)\zeta_b^+(v) + R_{N,b}(u, v), \quad \|R_{N,b}\|_{C^k} \leq C_k \tau^N, \quad (2.11)$$

with  $|R_{N,b}(u, v)| \leq C\tau^N|uv|$ . Let  $T_b^-$  be the next-period position on the future stable half-line, and  $T_b^+$  the previous-period position on the past half-line. Then

$$\zeta_b^\pm(T_b^\pm(x)) = e^{-\chi} \zeta_b^\pm(x). \quad (2.12)$$

Each  $\zeta_b^\pm$  is the unique differentiable local linearizing coordinate with value zero and derivative one at zero.

*Proof.* Twice integrating the mixed derivative, with oriented integrals on negative intervals, gives the exact identity

$$-C_{N,b}(u, v)/D_{N,b} = \int_0^u \int_0^v \beta_{N,b}(s, t) dt ds.$$

The relative theorem proves (2.11), including its endpoint vanishing and fixed mixed derivatives. Notice that the unnormalized remainder has the smaller size  $O(D_{N,b}\tau^N|uv|)$ ; an absolute action estimate alone does not give this transmission coefficient.

For the future return concatenate a  $P$ -flight segment with an  $N$ -flight tail. The joining endpoint  $w_N(u, v)$  is uniquely determined by stationarity and tends to  $T_b^-(u)$  with every fixed derivative. The denominator in its implicit equation is the sum of two positive same-end Hessians. Exact stationary differentiation gives

$$-W_{N+P,b,uv}(u, v) = [-W_{N,b,12}(w_N(u, v), v)] \partial_u w_N(u, v).$$

Since  $D_{N+P,b}/D_{N,b} \rightarrow e^{-\chi}$ , relative factorization and cancellation of the positive right factor yield

$$B_b^-(T_b^-(u))(T_b^-)'(u) = e^{-\chi} B_b^-(u).$$

Integration gives the minus identity in (2.12). Reversing the itinerary proves the plus identity. In particular both return maps have derivative  $e^{-\chi}$  at zero and are contractions on a smaller interval. If  $\tilde{\zeta}$  is another normalized differentiable linearizer, then  $h = \tilde{\zeta} \circ (\zeta_b^\pm)^{-1}$  satisfies  $h(\lambda x) = \lambda h(x)$ ,  $h(0) = 0$ ,  $h'(0) = 1$ , where  $\lambda = e^{-\chi}$ . Thus  $h(x) = \lambda^{-n} h(\lambda^n x) \rightarrow x$ , proving uniqueness.  $\square$

**2.5. The physical window law.** The gauge in (2.1) is useful analytically, but it is not a change of the clock. Define the physical boundary actions

$$\Sigma_b^-(u) = S_b^-(u) - p_b u, \quad \Sigma_b^+(v) = S_b^+(v) + p_b v. \quad (2.13)$$

In general they do not have critical points at zero. Their linear parts are the oblique endpoint momenta of the reference orbit.

**Theorem 2.5** (Charged phase-volume law for a repeated itinerary). *Let  $A$  be the free area of the periodic billiard cell and normalize its unit-speed phase volume by  $2\pi A$ . Fix a sufficiently small positive offset interval  $[d_-, d_+]$ . There is a smaller closed endpoint interval  $I$ , independent of  $N = nP$ , such that the following experiment is physical. The first collision is in the phase- $b$  collar; there are the selected  $N$  successive flights and no further collision by time  $t_N = nL + d$ ;*

the two recorded endpoints lie in  $I^2$ . The intervening collisions are required to remain in the prescribed itinerary collars.

Its endpoint measure, before conditioning on success, has density

$$\frac{D_{N,b}}{2\pi A} \beta_{N,b}(u, v) \{d - E_{N,b}(u, v) + p_b u - p_b v\}_+, \quad (u, v) \in I^2. \quad (2.14)$$

On this interval the expression inside braces is bounded below by  $d_-/2$  and above by a constant smaller than either adjacent free-flight length. Write  $Z_{N,b,d}$  for the integral of the numerator after the constant  $D_{N,b}/(2\pi A)$  is removed. The success probability is exactly  $D_{N,b}Z_{N,b,d}/(2\pi A)$  and is bounded above and below by positive constant multiples of  $e^{-n\chi}$ .

The conditional endpoint density converges, in every fixed  $C^k(I^2)$  norm including the stated family derivatives, to

$$f_{b,d}(u, v) = Z_{b,d}^{-1} B_b^-(u) B_b^+(v) \{d - \Sigma_b^-(u) - \Sigma_b^+(v)\}, \quad \|f_{N,b,d} - f_{b,d}\|_{C^k} \leq C_k \tau^N. \quad (2.15)$$

All rejected independent phase preparations remain charged through the success probability. The result is a law for this selected windowed itinerary, not an identification with a maximal-collision event.

*Proof.* Clearance and nongrazing incidence imply that the actual preceding and following free flights of every sufficiently small stationary segment have a common positive lower bound  $l_*$ . Choose  $d_+ < l_*/4$ . The action has a uniform quadratic bound on the fixed box, whereas  $|p_b|$  is uniformly bounded. Shrink  $I$  until  $|E_{N,b} - p_b u + p_b v| \leq d_-/2$ , uniformly in  $N$ . Then the available residual time  $r = t_N - W_{N,b} = d - E_{N,b} + p_b u - p_b v$  lies in  $[d_-/2, 3d_+/2]$ , below  $l_*$ . Any split  $r = s + (r - s)$ ,  $0 \leq s \leq r$ , lies strictly within the adjacent flights (apart from measure-zero collision endpoints). It gives exactly the required first and last collision times. Conversely a successful trajectory has this unique first collision, its unique small stationary bridge, and such a split. Thus no multiplicity is introduced by repeating the word.

If  $x$  is scaled boundary arclength, its canonical tangential momentum is the ordinary arclength momentum divided by  $\nu_b$ . Hence  $dx \wedge d\pi$  is the usual billiard flux form. The endpoint generating action gives  $\pi = -W_{N,b,u}$ , so that the absolute Jacobian to  $(u, v)$  is  $-W_{N,b,uv} > 0$ . The full phase form is this flux times  $ds$ . Integrating over  $0 \leq s \leq r$  proves (2.14), with precisely the normalization  $2\pi A$ . In particular the endpoint gauge terms have been restored.

The residual interval has a positive floor. The finite and limiting amplitudes have positive upper and lower bounds on a common box, by (2.9) and (2.10). Consequently  $Z_{N,b,d}$  and  $Z_{b,d}$  have positive uniform floors and finite uniform upper bounds. The relative theorem compares their unnormalized integrands and fixed derivatives at rate  $C_k \tau^N$ . Integration on the fixed square and division by these floors prove (2.15). Finally (2.4) gives the stated charged probability. At no stage is an absolute action error divided by that probability.  $\square$

## 2.6. A nonnormal orbit and the alternating specialization.

**Proposition 2.6** (A three-obstacle realization). *There is a periodic dispersing table with a nongrazing three-periodic orbit to which Theorems 2.3 and 2.5 apply, and an open family of such realizations with unequal flight lengths and incidence angles.*

*Proof.* In a square cell of side 12 place three unit disks whose centers form an equilateral triangle of side  $D = 5/2$ , centered in the cell. The disks and their translates are disjoint. At each disk choose the point on its radius directed toward the triangle center. These three points form an equilateral triangle of side  $\ell = D - \sqrt{3} > 0$ . Its inscribed three-flight path satisfies reflection: the outward disk normal bisects the two outgoing directions, with normal cosine  $\nu = \sqrt{3}/2$ . The distance from the opposite disk center to a whole side line is  $\sqrt{3}D/2 - 1/2 > 1$ ; hence no third

disk blocks that side. Neighboring cells are much farther away. The flights are therefore clear and the incidence is not normal or grazing.

Take the oriented period  $P = 6$ . Then  $k_i = 1/\ell$ ,  $m_i = 2/\nu$ , and

$$\chi = 6 \operatorname{arcosh}(1 + \ell/\nu).$$

For explicit nonsymmetric realizations, choose radii  $R_0, R_1, R_2$  close to one and prescribe a nearby scalene contact triangle. Place center  $i$  at distance  $R_i$  behind vertex  $i$  along its interior-angle bisector. The triangle then satisfies reflection by construction, has unequal sides and, choosing its three angles distinct, unequal normal cosines. Disjointness and clearance persist by the strict margins just checked. This gives the asserted open family by the same persistence lemma.  $\square$

For a normal alternating channel take  $P = 2$ ,  $L_0 = L_1 = g$  and  $\nu_0 = \nu_1 = 1$ . Then  $p_b = 0$ ,  $\chi = 2\gamma$ , and (2.4) reduces to  $D_{2n,b} = a_b / \sinh(2n\gamma)$ . Reversal identifies the two half-line actions and amplitudes. Passing from arclength to the fixed transverse graph chart changes an amplitude by the corresponding endpoint Jacobian, normalized by its value at zero. Thus the relative law specializes to Theorem 17.2, not to a different normalization of it. The intrinsic scalar coordinate in Corollary 2.4 specializes to Theorem 19.3.

The extension above concerns the geometric forward mechanism and its physical normalization. It neither assumes nor proves that the two boundary actions of a general itinerary can be inverted by the alternating two-contact block formula. The complete signed inverse, analytic-germ results, global lattice reconstruction and finite-preparation bounds that follow keep all their original hypotheses and conclusions.

**2.7. A dimension-independent two-end comparison.** The determinant argument admits a formulation that separates the size of an experiment from the interaction of its two ends. We state it for finite matrices of arbitrary size; padding a matrix by zero does not change any of the quantities below. The trace and Hilbert–Schmidt norms are denoted by  $\|\cdot\|_1$  and  $\|\cdot\|_2$ , respectively. For a parameter domain  $U$ , a  $C^k(U; \mathfrak{S}_p)$  bound means the maximum of these norms over all parameter derivatives of order at most  $k$ . Operator norms will be written without a subscript.

**Proposition 2.7** (Separated perturbation determinants). *Fix  $k \geq 0$  and  $q < 1$ . Relative to two orthogonal end spaces, write*

$$T_N = D_N + E_N, \quad D_N = \operatorname{diag}(T_{N,-}, T_{N,+}), \quad E_N = \begin{pmatrix} 0 & E_{N,-+} \\ E_{N,+-} & 0 \end{pmatrix}.$$

Here  $D_N$  is a block operator, not the scalar reference twist  $D_{N,b}$ . Suppose, uniformly on  $U$  and in  $N$ , that

$$\|D_N + tE_N\| \leq q \quad (0 \leq t \leq 1), \quad \|D_N\|_{C^k(\mathfrak{S}_1)} \leq M_k, \quad \|E_N\|_{C^k(\mathfrak{S}_2)} \leq \varepsilon_N \leq 1.$$

Suppose also that there are functions  $\ell_-, \ell_+$  with bounded  $C^k$  norms such that

$$\|\log \det(I + T_{N,\pm}) - \ell_{\pm}\|_{C^k(U)} \leq \delta_N.$$

All logarithms are defined by the power series at zero. Then

$$\|\log \det(I + T_N) - \ell_- - \ell_+\|_{C^k(U)} \leq C_{k,q,M_k}(\delta_N + \varepsilon_N^2). \quad (2.16)$$

The constant is independent of the dimensions of both end spaces. If  $T_N$  is first replaced by  $\tilde{T}_N$  with  $C^k(\mathfrak{S}_1)$  error  $\eta_N$ , with the same operator margin and bounded  $C^k(\mathfrak{S}_1)$  norms along the intervening segment, an additional  $C\eta_N$  suffices on the right.

*Proof.* For  $\|T\| < 1$  put

$$\mathcal{L}(T) = \sum_{h \geq 1} \frac{(-1)^{h+1}}{h} \operatorname{tr}(T^h).$$

This is the indicated logarithm, whether or not the matrix is symmetric. Along  $T(t) = D + tE$ , let  $R(t) = (I + D + tE)^{-1}$ . Differentiating the convergent series, and then using cyclicity of trace, gives

$$\frac{d}{dt} \mathcal{L}(D + tE) = \operatorname{tr}(R(t)E), \quad \frac{d^2}{dt^2} \mathcal{L}(D + tE) = -\operatorname{tr}(R(t)ER(t)E).$$

The inverse  $R(0)$  is block diagonal and  $E$  is off diagonal. Consequently  $\operatorname{tr}(R(0)E) = 0$ . Taylor's formula is therefore the exact identity

$$\mathcal{L}(D + E) - \mathcal{L}(D) = - \int_0^1 (1-t) \operatorname{tr}(R(t)ER(t)E) dt. \quad (2.17)$$

Since  $\|R(t)\| \leq (1-q)^{-1}$  and  $|\operatorname{tr}(AB)| \leq \|A\|_2 \|B\|_2$ , the undifferentiated absolute value is at most  $\|E\|_2^2 / [2(1-q)^2]$ . For parameter derivatives use  $\partial R = -R\partial(D + tE)R$  repeatedly. Each differentiated integrand in (2.17) still contains the two displayed  $E$  factors, possibly differentiated. Place these in  $\mathfrak{S}_2$  and all other factors in operator norm. The bounds on  $D$ ,  $E$  and  $R$  yield  $C_{k,q,M_k} \varepsilon_N^2$ . The number of product-rule terms depends on  $k$ , not on dimension. The block-diagonal determinant is the product of its two determinants; adding their stipulated errors proves (2.16).

For the last assertion write  $T(t) = T + t(\tilde{T} - T)$  and integrate  $\operatorname{tr}((I + T(t))^{-1}(\tilde{T} - T))$ . In every parameter derivative one factor is a derivative of  $\tilde{T} - T$  in trace norm; all remaining factors have uniformly bounded operator norm. This gives  $C\eta_N$ . These arguments also justify each stated fixed derivative: all the finite-dimensional identities have bounds uniform in dimension.  $\square$

A convenient sufficient condition for the diagonal hypothesis is  $C^k(\mathfrak{S}_1)$  convergence of the padded diagonal blocks to trace-class half-line operators of norm below  $q$ . The same path argument gives convergence of their logarithms. Thus no determinant estimate proportional to the number of interior collisions is needed. The quadratic off-diagonal term expresses the fact that a closed trace must cross from one end to the other and back. A first-order trace-norm estimate remains sufficient when this refinement is not needed.

**Corollary 2.8** (Relative comparison before conditioning). *Let  $Q$  be a fixed compact gate and let the unconditioned endpoint density be exactly*

$$w_N = c d_N \beta_N r_N, \quad c > 0, \quad d_N > 0.$$

*Assume uniform positive lower and upper bounds for  $\beta_N, r_N$  and their limits  $\beta, r$ , uniform  $C^k$  bounds, and*

$$\|\log \beta_N - \log \beta\|_{C^k(Q)} + \|r_N - r\|_{C^k(Q)} \leq e_N.$$

*If  $Z_N = \int_Q \beta_N r_N$  and  $Z = \int_Q \beta r$ , then*

$$\left\| \frac{w_N}{\int_Q w_N} - \frac{\beta r}{Z} \right\|_{C^k(Q)} + \left| \frac{\int_Q w_N}{c d_N} - Z \right| \leq C e_N. \quad (2.18)$$

*The constant is independent of  $d_N$ , which may tend to zero at any rate. In particular, a scalar cofactor identity  $\log \beta_N = b_N - \log \det(I + T_N)$ , with a separated edge sum  $b_N$  and the hypotheses of Proposition 2.7, gives such a comparison with the corresponding end errors.*

*Proof.* Exponentiation is locally Lipschitz on bounded  $C^k$  sets. Products therefore compare  $\beta_N r_N$  with  $\beta r$  with error  $Ce_N$ . Integration compares  $Z_N$  with  $Z$ , and the fixed positive gate gives a positive lower bound for both integrals. Division by these integrals proves the first estimate. The second uses the exact equality  $\int_Q w_N = c d_N Z_N$ . Neither step divides an approximation error by  $d_N$ .  $\square$

For the billiard operator the end spaces are the first and last  $\lfloor N/3 \rfloor$  sites after the middle perturbation has been removed. The Green and gluing bounds in the proof of Theorem 2.3 give the trace-norm deletion error, the two diagonal limits and the Hilbert–Schmidt off-diagonal bound, all exponentially small up to fixed polynomial factors in  $N$ . Shrinking the collar gives the operator margin along the entire segment. Consequently the proposition applies before the cofactor is converted to a probability law. In Corollary 2.8 one takes  $d_N = D_{N,b}$ ,  $c = (2\pi\mathcal{A})^{-1}$ , and  $r_N = d - E_{N,b} + p_b u - p_b v$ . The geometry is responsible for these hypotheses and for the exact physical clock; the operator comparison itself does not use periodicity, a billiard reflection rule, or a lower bound on the total success mass.

### 3. CONTACT RECONSTRUCTION ALONG A MARKED PERIODIC ORBIT

The future and past actions of a periodic orbit need not coincide. Nevertheless their separation in the physical window law permits an inverse which does not use time-reversal symmetry. Two positive offsets separate the actions without a flux model. The derivative of the future action at all collision phases is then a sum of successive boundary visits. Its homogeneous blocks are Cayley transforms of a signed cyclic shift. We prove below that this calculation applies to actual smooth boundaries and gives an inverse of the complete analytic action map.

**3.1. Fixed geometric marks and the observation coordinates.** Fix a clear nongrazing periodic orbit with  $r \geq 2$  collision phases. The collision points are distinct modulo the prescribed lattice, and phase  $i$  has a fixed point  $q_i$ , unit outward obstacle normal  $n_i$  and signed unit tangent  $t_i$ . These data, including the lifted closing translation, will be called its *marked polygon*. Indices are read modulo  $r$ , with that translation understood. Set

$$e_i^0 = \frac{q_{i+1} - q_i}{L_i}, \quad L_i = |q_{i+1} - q_i|, \quad \nu_i = e_i^0 \cdot n_i > 0, \quad p_i = e_i^0 \cdot t_i / \nu_i.$$

The reflection condition is part of the marked polygon. In particular  $e_{i-1}^0 \cdot n_i = -\nu_i$  and  $e_{i-1}^0 \cdot t_i = e_i^0 \cdot t_i$ . No curvature or higher boundary derivative is included in these marks.

Write the unknown arcs as normalized normal graphs

$$R_i(x) = q_i + \nu_i^{-1} \{t_i x - n_i F_i(x)\}, \quad F_i(0) = F'_i(0) = 0, \quad c_i := F''_i(0) > 0. \quad (3.1)$$

The actual curvature is  $\kappa_i = \nu_i c_i$ . The recorded coordinate is the scaled tangent projection  $x = \nu_i t_i \cdot (R_i - q_i)$ , not the unknown normal displacement. Thus recording  $x$  does not supply  $F_i(x)$ . The marks specify the coordinate, and its use in the inverse does not require converting measured arclength by means of an unknown boundary. Put  $k_i = L_i^{-1}$  and

$$\ell_i(x, y) = |R_{i+1}(y) - R_i(x)|, \quad h_i(x, y) = \ell_i(x, y) - L_i + p_i x - p_{i+1} y.$$

These gauges have zero value and differential at the origin. The marked polygon remains a stationary orbit for every anchored graph in (3.1). On a sufficiently small common real collar, small perturbations preserve positive curvature, clearance and nongrazing incidence.

**Lemma 3.1** (Signed Jacobi coordinates). *There are signs  $\epsilon_i \in \{1, -1\}$ , determined by the marked polygon and tangent choices, such that*

$$D^2 h_i(0, 0) = \begin{pmatrix} k_i + c_i & -\epsilon_i k_i \\ -\epsilon_i k_i & k_i + c_{i+1} \end{pmatrix}. \quad (3.2)$$

On the infinite itinerary a diagonal change of signs makes all mixed entries negative. The resulting Jacobi operator has positive masses  $2c_i$ . The future half-line with endpoint  $x_0 = u$  has linearization

$$x'_{b,j}(0) = \sigma_b \sigma_{b+1} \cdots \sigma_{b+j-1}, \quad \sigma_i = \epsilon_i a_i, \quad 0 < a_i < 1. \quad (3.3)$$

The factors  $a_i$  are periodic. They depend smoothly on the positive curvatures, and are uniformly bounded away from one on a compact marked class. In particular  $\Lambda := \prod_{i=0}^{r-1} \sigma_i$  satisfies  $|\Lambda| < 1$ .

The relative and physical laws of Section 2 hold in the graph coordinates, for a returning oriented period  $P$  (which may be taken to be  $2r$ ). In these coordinates write their limits as  $S_b^\pm, B_b^\pm$  and

$$A_b(u) = S_b^-(u) - p_b u, \quad C_b(v) = S_b^+(v) + p_b v. \quad (3.4)$$

Then the physical residual in the limiting law is  $d - A_b(u) - C_b(v)$ . The action and amplitude estimates remain estimates at each separately fixed differentiation order.

*Proof.* At the origin  $R'_i = t_i/\nu_i$  and  $R''_i = -n_i c_i/\nu_i$ . Differentiating a chord twice gives the same-end entry  $k_i + c_i$ . Its mixed entry is

$$-\frac{t_i \cdot t_{i+1} - (e_i^0 \cdot t_i)(e_i^0 \cdot t_{i+1})}{L_i \nu_i \nu_{i+1}}.$$

The two transverse projections have lengths  $\nu_i, \nu_{i+1}$ ; their relative sign defines  $\epsilon_i$ . This proves (3.2). If  $s_0 = 1$  and  $s_{j+1} = \epsilon_{b+j} s_j$ , the substitution  $x_j = s_j y_j$  gives the positive-mass operator

$$(Hy)_j = (k_{b+j-1} + k_{b+j} + 2c_{b+j})y_j - k_{b+j-1}y_{j-1} - k_{b+j}y_{j+1}.$$

Its Green kernel and fixed differentiated bounds follow from the walk expansion in the preceding section; the diagonal change of signs preserves their absolute bounds.

For endpoint  $y_0 = 1$ , finite Dirichlet solutions are strictly positive in their interiors by the discrete maximum principle. Their half-line limit is positive and decays exponentially. Define  $z_{j+1} = k_{b+j}(y_{j+1} - y_j)$ . The recurrence gives  $z_{j+1} - z_j = 2c_{b+j}y_j$ . Since  $z_j \rightarrow 0$  at infinity,  $z_{j+1} = -\sum_{h \geq j+1} 2c_{b+h}y_h < 0$ . Consequently  $0 < y_{j+1} < y_j$ . Uniqueness of a decaying tail shows that its ratio at a site depends only on that phase; call it  $a_i$ . This proves (3.3). Smooth dependence follows by differentiating the uniformly invertible weighted equations. Compactness and finitely many phases give  $\max_i a_i < 1$  uniformly.

The signed charts return after  $2r$  steps, even when they do not return after  $r$ . This does not create  $2r$  independent unknown contact germs. To compare with signed scaled arclength at the initial phase, let  $\phi_b(x)$  be its expression in the graph coordinate, with the same initial orientation. Directly from (3.1),

$$\phi_b(x) = \int_0^x \sqrt{1 + F'_b(s)^2} ds = x + O(x^3).$$

Use the positive branch on the real collar. The endpoint twist is multiplied by  $\phi'_b(u)\phi'_b(v)$ , whose value at zero is one. The exact reference  $D_{N,b}$  is therefore unchanged. The graph future action is the arclength action composed with  $\phi_b$ , plus  $p_b(x - \phi_b(x))$ ; the past action has the opposite gauge correction. The amplitude acquires the ordinary endpoint Jacobian. In particular the graph actions are still anchored and have positive Hessians at zero. Composition on smaller collars preserves the fixed-order relative bounds. The physical actions in (3.4) are simply the physical arclength actions composed with  $\phi_b$ ; no physical time has been changed. This proves the last assertion, without making an unknown coordinate conversion an input to the reconstruction.  $\square$

**3.2. Separation of unequal endpoint actions.** Choose two fixed offsets  $0 < d_1 < d_2$  in the physical positive window, and a common square  $I^2$  on which both residuals have a positive floor. At phase  $b$  the two limiting densities are

$$f_{b,d}(u, v) = Z_{b,d}^{-1} B_b^-(u) B_b^+(v) \{d - A_b(u) - C_b(v)\}, \quad d = d_1, d_2. \quad (3.5)$$

Here  $A_b(0) = C_b(0) = 0$  and  $B_b^\pm(0) = 1$ . The two normalizing constants are unrelated unknown positive numbers.

**Proposition 3.2** (Two-offset action separation). *The two conditional laws in (3.5) determine the complete future and past actions and normalized amplitudes. More precisely, suppressing the phase index, set*

$$Q(u, v) = \frac{f_{d_1}(u, v) f_{d_2}(0, 0)}{f_{d_2}(u, v) f_{d_1}(0, 0)}, \quad H(u, v) = \frac{d_1 d_2 \{1 - Q(u, v)\}}{d_2 - d_1 Q(u, v)}. \quad (3.6)$$

Then

$$A(u) = H(u, 0), \quad C(v) = H(0, v), \quad H(u, v) = A(u) + C(v). \quad (3.7)$$

The gauged actions follow from (3.4), and

$$B^-(u) = \frac{f_d(u, 0)}{f_d(0, 0)} \frac{d}{d - A(u)}, \quad B^+(v) = \frac{f_d(0, v)}{f_d(0, 0)} \frac{d}{d - C(v)}. \quad (3.8)$$

For each fixed  $M \geq 0$  these formulas are locally Lipschitz in  $C^M(I^2)$ , on compact positive-density classes with fixed separated offsets and residual floors. They do not require equality of the two actions, a model for the amplitudes, or an acceptance count.

*Proof.* Put  $h = A(u) + C(v)$ . Cancellation of the amplitudes and the two normalizers gives

$$Q = \frac{1 - h/d_1}{1 - h/d_2}, \quad d_2 - d_1 Q = \frac{d_2(d_2 - d_1)}{d_2 - h} > 0.$$

Solving the first identity for  $h$  proves (3.6)–(3.7). Since  $h$  is bounded on the common square, the second identity gives a positive denominator floor, uniformly on the stated compact class. Taking a single density ratio on each coordinate axis proves (3.8). Continuous densities are fixed everywhere on the positive square by their probability laws, even though finite samples would require estimation.

For stability, slice evaluation is bounded in  $C^M$ . Products and reciprocals are locally Lipschitz in  $C^M$  when the denominators have positive floors. Apply this first to  $Q$ , then to  $H$ , and finally to the two axis formulas. Small perturbations need not themselves satisfy the separation identity. The formulas remain defined on an open neighborhood with the same, possibly halved, denominator floors. The constants are not uniform as  $d_2 - d_1$  or a residual floor tends to zero. No assertion of a minimal number of offsets is made.  $\square$

**3.3. Actual boundary variations and cyclic blocks.** Let  $\mathcal{S}(F) = (S_0^-, \dots, S_{r-1}^-)$  be the future action map in the fixed graph coordinates. A reference half-line starting at phase  $b$  is denoted  $x_{b,j}(u)$ , and its  $j$ th unit chord direction is  $e_{b,j}(u)$ . Define

$$\alpha_b(u) = \frac{e_{b,0}(u) \cdot n_b}{\nu_b}, \quad v_{b,j}(u) = \frac{(e_{b,j}(u) - e_{b,j-1}(u)) \cdot n_{b+j}}{\nu_{b+j}} \quad (j \geq 1). \quad (3.9)$$

At zero these weights are respectively 1 and 2.

**Lemma 3.3** (Stationary boundary visits). *For an anchored smooth variation  $\eta = (\eta_i)$ , on a sufficiently small common collar the derivative of the actual action is*

$$(D\mathcal{S}(F)\eta)_b(u) = \alpha_b(u)\eta_b(u) + \sum_{j \geq 1} v_{b,j}(u)\eta_{b+j}(x_{b,j}(u)). \quad (3.10)$$

The series and its stipulated fixed derivatives converge. The action jet through degree  $M$  depends only on the graph jets through degree  $M$ , for every  $M \geq 2$ . This is a statement about smooth representatives, not only formal Taylor series.

*Proof.* The weighted Green construction of the preceding section applies to the signed graph equations. It gives the actual stationary half-line and bounds  $|x_{b,j}(u)| \leq C|u|\rho^j$ , with the corresponding bounds for each fixed differentiated order. Changing a graph by  $\eta_i$  changes its position by  $-n_i\eta_i/\nu_i$ . Direct differentiation of one chord therefore contributes

$$\frac{e_{b,j} \cdot n_{b+j}}{\nu_{b+j}} \eta_{b+j}(x_{b,j}) - \frac{e_{b,j} \cdot n_{b+j+1}}{\nu_{b+j+1}} \eta_{b+j+1}(x_{b,j+1}).$$

The marked linear gauges have zero variation. Differentiate the sum of the first  $J$  gauged chords along the stationary family. All internal orbit variations cancel. The initial one is zero because its coordinate is fixed. Retain the terminal term

$$h_{b+J-1,2}(x_{b,J-1}, x_{b,J}) D_F x_{b,J}[\eta].$$

The first factor is  $O(|u|\rho^{J-1})$ , since  $Dh(0,0) = 0$ ; the second decays exponentially by the differentiated weighted equations. Thus this term, and each of its separately fixed derivatives, tends to zero. Direct graph variations are summable, because anchored variations are  $O(x^2)$ . Gathering the initial visit once and every internal visit twice gives precisely (3.10).

For completeness, suppose two smooth profiles have the same jets through  $M$ . Their difference obeys  $\eta_i(x) = O(x^{M+1})$  on a smaller collar. Interpolate the profiles there. Their positive quadratic jets agree, and uniform small-collar estimates hold along the interpolation; no global interpolation of obstacles is needed. The envelope formula gives, uniformly along this path,

$$|DS(F)[\eta]_b(u)| \leq C|u|^{M+1} \left( 1 + \sum_{j \geq 1} \rho^{j(M+1)} \right).$$

Integration along the path shows that the two actual actions differ by  $O(|u|^{M+1})$ . Their smoothness, already obtained from the weighted equations, implies equality of the Taylor jets through  $M$ . This jet argument alone does not identify equality of all smooth jets with equality of smooth germs. The latter identification from the full action functions is established by the additional real weighted estimate in Section 5.  $\square$

For a degree  $n \geq 2$  define the signed cyclic shift on  $\mathbb{R}^r$  by

$$(T_n z)_b = \sigma_b^n z_{b+1}, \quad T_n^r = \Lambda^n I, \quad \mathcal{C}_n = (I + T_n)(I - T_n)^{-1}. \quad (3.11)$$

The notation uses the physical  $r$  contact germs, not the doubled list of oriented charts.

**Theorem 3.4** (All-order periodic contact blocks). *Let  $q_n = (F_b^{(n)}(0))_b$  and  $s_n = ((S_b^-)^{(n)}(0))_b$ . The map  $q_2 \mapsto s_2$  is locally invertible throughout the positive-curvature class; its derivative is  $\mathcal{C}_2$ . For every  $n \geq 3$ ,*

$$s_n = \mathcal{C}_n q_n + \mathcal{R}_n(q_2, \dots, q_{n-1}), \quad q_n = (I - T_n)(I + T_n)^{-1}(s_n - \mathcal{R}_n). \quad (3.12)$$

The coefficients also depend on the fixed marked polygon. The formulas hold for actual smooth graphs. They give a local inverse through every finite order, without supplied curvatures.

In the maximum row-sum norm, with  $a_* := \max_b a_b < 1$ ,

$$\max(\|\mathcal{C}_n\|_\infty, \|\mathcal{C}_n^{-1}\|_\infty) \leq \frac{1 + a_*^n}{1 - a_*^n} \quad (n \geq 2), \quad \det \mathcal{C}_n = \frac{1 - (-1)^r \Lambda^n}{1 - \Lambda^n} > 0. \quad (3.13)$$

In particular the determinant is one for an even contact cycle. For an odd cycle it need not be one, but is never zero. These bounds concern the homogeneous blocks, not arbitrary lower-degree smooth couplings.



*Proof.* Vary only the new degree by  $\eta_i(x) = z_i x^n / n!$ , using a smooth cutoff equal to one near zero when necessary. In (3.10) the coefficient of  $u^n / n!$  uses only  $\alpha_b(0) = 1$ ,  $v_{b,j}(0) = 2$ , and (3.3). It is therefore

$$z_b + 2 \sum_{j \geq 1} (\sigma_b \cdots \sigma_{b+j-1})^n z_{b+j} = \left( I + 2 \sum_{j \geq 1} T_n^j \right) z = (I + T_n)(I - T_n)^{-1} z.$$

This proves the derivative formula in degree two. For  $n \geq 3$ , the linearized orbit and the constant visit weights depend only on the quadratic jets, not on the fresh degree. Together with the smooth-jet factorization of Lemma 3.3, integration in that fresh degree gives the affine identity (3.12).

Since  $\|T_n\|_\infty \leq a_*^n < 1$ , the two geometric series for  $(I - T_n)^{-1}$  and  $(I + T_n)^{-1}$  converge. They give the inverse and both estimates in (3.13). For a cyclic shift,  $\det(I - zT_n) = 1 - z^n \Lambda^n$ . Apply this at  $z = 1, -1$  to obtain the displayed determinant;  $|\Lambda| < 1$  makes its numerator and denominator positive. The finite-dimensional inverse function theorem now inverts the degree-two map locally. Successive applications of (3.12) invert all later degrees. At any fixed order this is a smooth inverse on an open finite-jet neighborhood; its constants can be chosen uniformly on compact subsets of that neighborhood. No global uniqueness of the degree-two map is needed or asserted.  $\square$

For a normal two-contact channel, the positive signed factors are  $\sigma_b = \mathbf{r}_b e^{-\gamma}$  and  $\Lambda = e^{-2\gamma}$ . Substitution in (3.11) gives

$$C_n = \begin{pmatrix} \coth(n\gamma) & \mathbf{r}_0^n \operatorname{csch}(n\gamma) \\ \mathbf{r}_1^n \operatorname{csch}(n\gamma) & \coth(n\gamma) \end{pmatrix}.$$

Thus the earlier determinant-one inverse is the two-site specialization. For an odd cycle, replacing the signed  $\sigma_i$  by their absolute values would generally change the odd-degree reconstruction. Doubling the oriented forward period does not justify making that replacement.

**3.4. An inverse on a space of analytic germs.** Fix a real analytic graph tuple  $F^*$ , holomorphic near a closed disc  $D_{R_0}$ , with positive  $c_i^*$ . For  $R < R_0$  put

$$X_{R,q} = \{\eta \in H^\infty(D_R; \mathbb{C}^r) : \eta_i^{(j)}(0) = 0 \ (0 \leq j < q)\}, \quad \|\eta\|_R = \max_i \sup_{D_R} |\eta_i|.$$

The marked polygon is fixed; curvature is allowed to vary in  $X_{R,2}$ . The complex extension of the Euclidean chord uses the bilinear dot product and the square-root branch near the positive reference length.

**Lemma 3.5** (Holomorphic graph half-lines). *For all sufficiently small  $R > 0$  there is an open neighborhood  $\mathcal{U} \subset F^* + X_{R,2}$  on which the graph half-lines and the actual future actions depend holomorphically on  $F$ . Their endpoint maps obey*

$$|x_{b,j}(u)| \leq \rho^j |u| \quad (j \geq 1, u \in D_R), \quad \rho < 1. \quad (3.14)$$

*The action map takes values in  $\mathcal{S}(F^*) + X_{R,2}$ . Its full Fréchet derivative is (3.10). On a smaller neighborhood its coefficients are bounded and  $\inf_{b,u} |\alpha_b(u)| > 0$ .*

*Proof.* We give the function-space argument, since analyticity of finite parameter families would not suffice. Conjugate the reference Jacobi operator by the diagonal signs and use its weighted Green bound. Choose constants

$$\max(a_*, q) < \beta < \rho < \sigma < 1,$$

where  $q < 1$  is a walk-decay base of that operator. For sequences of holomorphic functions vanishing at zero use  $\|y\|_\beta = \sup_{j \geq 1} \beta^{-j} \|y_j\|_R$ . The Green operator is bounded on this space.

For  $\|\eta\|_R < \delta R^2$ ,  $\eta \in X_{R,2}$ , the higher-order Schwarz estimate and protected Cauchy estimates give

$$|\eta_i(z)| \leq \delta |z|^2 \quad (z \in D_R), \quad |\eta_i^{(k)}(z)| \leq C_{\sigma,k} \delta |z|^{2-k} \quad (z \in D_{\sigma R}, \quad 0 \leq k \leq 2). \quad (3.15)$$

For the second estimate write  $\eta_i = z^2 v_i$  and use Cauchy's bounds for  $v_i$  on an intermediate inner disc.

An interior stationarity equation differentiates the two chords with respect to their common interior argument. It uses the initial graph only through  $F_b(u)$ , never through  $F'_b(u)$  at the unprotected boundary argument. Every differentiated graph is evaluated at an interior site in  $D_{\sigma R}$ . Thus neither graph differentiation nor composition at the boundary of  $D_R$  is being treated as a bounded operator on  $H^\infty(D_R)$ .

Subtract the reference linear equation and apply the reference Green operator. On the ball

$$\|x - x^{\text{lin}}\|_\beta \leq C_0(\delta R + R^2), \quad x_{b,j}^{\text{lin}}(u) = (\sigma_b \cdots \sigma_{b+j-1})u,$$

the residual is  $O(\delta R + R^2)$  and its Lipschitz constant in the interior sequence is  $O(\delta + R)$ . Here the  $O(\delta R)$  term is the variation of the quadratic Jacobi coefficients. These estimates follow directly from the chord formulas: the square roots and reciprocal lengths stay in fixed neighborhoods of  $L_i$ ; first and second graph derivatives occur only at protected interior arguments and satisfy (3.15); neighboring weights differ by at most  $\beta^{\pm 1}$ ; products of two coordinates are  $O(R^2)$  in the weighted norm. The fixed analytic reference has bounded higher remainders on those discs. Choose  $C_0$  above twice the residual constant times the Green norm, then  $\delta$  and  $R$  small. The resulting map preserves the ball and has contraction constant below  $1/2$ .

The correction vanishes at  $u = 0$ . Schwarz's estimate gives

$$|x_{b,j}(u)| \leq \{a_*^j + C_0(\delta + R)\beta^j\}|u| \leq \rho^j|u|$$

after a further reduction of  $\delta, R$ . This verifies, rather than assumes, the protected arguments used above. On these protected sets the Cauchy expansions of derivatives and compositions converge uniformly as maps of the Banach variables. All fixed-point iterates are therefore holomorphic, and their locally uniform limit is holomorphic as well. The gauged action summands are bounded by  $CR^2\rho^{2j}$ , so their sum is a holomorphic  $H^\infty(D_R)$ -valued map. Its constant and linear jets are zero.

On a smaller parameter ball, Cauchy's estimate in the Banach parameter gives  $\|D_F x_{b,j}[\eta]\|_R \leq C_R \beta^j \|\eta\|_R$ . The retained terminal term in the proof of Lemma 3.3 consequently vanishes in the same function norm. Direct variations have the summable majorant  $C\rho^{2j}\|\eta\|_R$ , proving the full derivative formula in operator norm. The chord directions are uniformly close to their reference directions on a small disc, so the weights are bounded and  $\alpha_b$  is uniformly close to one. Real members give strictly convex physical arcs on a smaller real collar, by protected Cauchy bounds and the uniqueness of the real weighted construction. No extension of every resulting germ to a closed analytic obstacle is asserted by this local lemma.  $\square$

**Theorem 3.6** (Complete analytic periodic contact inverse). *For a fixed marked polygon and any positive-curvature real analytic contact tuple, the map*

$$\mathcal{S} : F^* + X_{R,2} \longrightarrow \mathcal{S}(F^*) + X_{R,2}$$

*is a biholomorphism between sufficiently small neighborhoods, for all sufficiently small  $R > 0$ . On a smaller convex neighborhood there are constants  $0 < c < C < \infty$  such that*

$$c\|F - \tilde{F}\|_R \leq \|\mathcal{S}(F) - \mathcal{S}(\tilde{F})\|_R \leq C\|F - \tilde{F}\|_R. \quad (3.16)$$

*These bounds include curvature and all higher coefficients, together with their lower-degree couplings. The derivative and its inverse induce mutually inverse operators on  $X_{R,2}/X_{R,N+1}$ ,*

bounded independently of  $N$  in the induced quotient norms. For  $0 < s < R$ ,

$$\sum_{n \geq 2} \frac{s^n}{n!} \max_b |F_b^{(n)}(0) - \tilde{F}_b^{(n)}(0)| \leq \frac{(s/R)^2}{c(1-s/R)} \|\mathcal{S}(F) - \mathcal{S}(\tilde{F})\|_R. \quad (3.17)$$

*Proof.* Fix the base derivative  $L = D\mathcal{S}(F^*)$ . In (3.10) write  $L = A + K$ , where  $A$  is the invertible diagonal multiplier  $\alpha_b$  and  $K$  is the sum of contracted visits. Let  $\|A^{-1}\| \leq a^{-1}$  and  $\sup_{b,j,u} |v_{b,j}(u)| \leq W$ . On  $X_{R,m}$  the Schwarz bound gives

$$\|K|_{X_{R,m}}\| \leq \frac{W\rho^m}{1-\rho^m}.$$

Choose a single finite  $m > 2$  so that  $q_m = a^{-1}W\rho^m/(1-\rho^m) < 1$ . The tail inverse is the norm-convergent operator series

$$(L|_{X_{R,m}})^{-1} = \sum_{j \geq 0} (-A^{-1}K)^j A^{-1}, \quad \|(L|_{X_{R,m}})^{-1}\| \leq \frac{a^{-1}}{1-q_m}.$$

Let  $P_m$  retain the Taylor degrees  $2, \dots, m-1$ . Cauchy's coefficient bounds make it a bounded finite-rank projection. The finite quotient  $P_m L|_{P_m X_{R,2}}$  is block lower triangular, with diagonal blocks  $\mathcal{C}_2, \dots, \mathcal{C}_{m-1}$  by Theorem 3.4. Every block is invertible, including those at odd degree for an odd contact cycle. Hence the quotient has an inverse. Given  $y$ , first solve that finite quotient for  $\eta_{\text{low}}$  and then set

$$\eta_{\text{tail}} = (L|_{X_{R,m}})^{-1}(y - L\eta_{\text{low}}).$$

The residual has zero jets below  $m$ , so this is defined and gives a unique bounded full inverse. This is the contracted-evaluation argument of Proposition 22.13, here with  $q = 2$ ,  $r$  components, and the actual geometric blocks just computed. In particular boundedness of the diagonal blocks alone has not been used to assert boundedness of an infinite triangular inverse.

The action map is holomorphic by Lemma 3.5. Write its increment as  $L\eta + \mathcal{N}(\eta)$ , where  $D\mathcal{N}(0) = 0$ , and shrink the ball until  $\|L^{-1}D\mathcal{N}\| < 1/2$ . The equation  $\eta = L^{-1}y - L^{-1}\mathcal{N}(\eta)$  is then a contraction, with holomorphic locally uniform iterates. It proves the asserted biholomorphism. On a smaller convex ball make  $\|D\mathcal{S} - L\| < (2\|L^{-1}\|)^{-1}$ . Integration along line segments proves (3.16). Conjugation and uniqueness give the real analytic inverse for real data.

Both  $L$  and  $L^{-1}$  preserve every higher vanishing tail: preservation by  $L$  follows from the envelope formula, and preservation by  $L^{-1}$  follows inductively from the invertible homogeneous blocks. Their quotient operator norms are bounded by their full operator norms, independently of  $N$ . Continuity gives the same local assertion at nearby base points. Cauchy's coefficient inequality and (3.16), summed over  $n \geq 2$ , give (3.17). These are analytic norms on a fixed disc, not maximum norms of arbitrarily many unweighted smooth derivatives.  $\square$

### 3.5. Rigidity from physical laws and its observation topology.

**Corollary 3.7** (Periodic contact rigidity and finite-flight inversion). *Fix the marked polygon, the graph-coordinate sensors at every physical phase, and two offsets as above. Their limiting conditional laws locally determine all finite jets of every contact in the smooth category. In the analytic category they determine the complete local contact germs. No curvature, normal graph value, or flux amplitude is supplied.*

*For a fixed  $M \geq 2$ , densities of  $N$ -flight bridges,  $N$  a multiple of a returning oriented period, determine the jets through order  $M$  on the chosen local inverse branch with error at most*

$$C_M\{\tau^N + \epsilon\}, \quad (3.18)$$

when the recorded densities at both offsets and all phases have interior  $C^M$  error at most  $\epsilon$ . The constants are local fixed-order constants. This statement is about error in densities, not acquisition of those densities from a stipulated number of preparations.

If two actual periodic tables have the same fixed lattice and obstacle labels, every obstacle is visited by the marked polygon, and their boundaries are connected closed embedded real analytic curves, equality of these laws implies equality of their full obstacle images, provided the contact germs lie in the stated local analytic inverse neighborhood.

*Proof.* Proposition 3.2 recovers  $A_b, C_b$  and hence  $S_b^\pm$  in the measured graph coordinates. Theorem 3.4 inverts their finite jets locally, beginning with degree two. For analytic germs, equality on a real interval makes the action extensions equal by the identity theorem. Apply Theorem 3.6. In fact the future actions at all phases suffice for the geometric inverse; the recovered past actions provide additional data, not an assumed symmetry.

The graph-coordinate version of Theorem 2.5 gives the density error  $C_M \tau^N$  on the same square at both offsets. The rational extraction is locally Lipschitz there. Before finite-jet inversion subtract the constant and linear Taylor terms of an estimated gauged action; this bounded projection fixes the exact action. Apply the local degree-two inverse and then the finitely many smooth operations in (3.12). They give (3.18).

Finally equality of contact germs gives an open common boundary arc on each visited labelled obstacle. Two connected embedded analytic closed curves which agree on an open arc agree everywhere. Indeed at an endpoint of a continued common arc both curves are locally graphs over their common tangent; the one-variable identity theorem continues equality across that endpoint. Compactness and connectedness then continue around the whole curve. Every obstacle is visited, so all labelled boundary images coincide. The lattice was fixed, and therefore so does the table. This continuation argument gives exact rigidity, not a global quantitative continuation estimate and not reconstruction of an unmarked lattice.  $\square$

**Theorem 3.8** (Conditional stability from real laws). *Fix a positive-curvature analytic base tuple and its marked polygon. There are radii  $0 < s < r_0 < \varrho < R$ , with  $r_0 + 2s < \varrho$ , a bounded neighborhood of that tuple in  $F^* + X_{R,2}$ , and  $C < \infty$ ,  $0 < \theta < 1$ , for which the following holds. Choose the two offsets and the square  $Q = [-s, s]^2$  inside their common positive-density region. For two admissible tuples in this neighborhood set*

$$\epsilon = \max_{b,d \in \{d_1, d_2\}} \|f_{b,d} - \tilde{f}_{b,d}\|_{C^0(Q)}.$$

For  $0 \leq \epsilon \leq \epsilon_0$ ,

$$\|F - \tilde{F}\|_{r_0} \leq C\epsilon^\theta. \quad (3.19)$$

*The analytic prior, local inverse branch, fixed marked coordinates, positive window and radius loss are part of the assertion. For finite-flight densities the right side may instead use  $C(\epsilon + \tau^N)^\theta$  when the corresponding two finite experiments are compared. No total-preparation budget is asserted here.*

*Proof.* Choose small nested radii on which the holomorphic construction holds, and use Theorem 3.6 also on  $D_{r_0}$ . Shrink the neighborhood in the larger-disc norm until all its restrictions belong to that smaller-disc inverse neighborhood. This is possible by bounded restriction; for anchored increments the sharper bound is  $\|\eta\|_{r_0} \leq (r_0/R)^2 \|\eta\|_R$ . The larger-disc action norms have a uniform bound by Lemma 3.5.

The  $C^0$  case of Proposition 3.2 gives  $\max_b \|S_b^- - \tilde{S}_b^-\|_{C^0([-s,s])} \leq C_0\epsilon$ . The momenta cancel because the marked polygon is fixed. Apply Lemma 22.14 to each holomorphic action difference,

using its common larger-disc bound. Explicitly, its exponent can be chosen as

$$A_0 = \frac{2(r_0 + s)}{s}, \quad b_0 = \frac{r_0 + s}{\varrho - s} < 1, \quad \theta = \frac{\log(1/b_0)}{\log(A_0/b_0)}.$$

That lemma yields an  $O(\epsilon^\theta)$  bound for the action differences on  $D_{r_0}$ . The complete inverse bound on that disc proves (3.19). When  $\epsilon = 0$ , use the zero-error case of the same lemma or the identity theorem. The triangle inequality and the uniform finite-flight relative law add  $O(\tau^N)$  before this argument is applied, proving the last assertion. A bounded analytic prior is essential in this argument; unrestricted analytic continuation or same-disc stability from a real norm is not being inferred.  $\square$

### 3.6. A nonnormal analytic realization.

**Proposition 3.9** (A three-contact family with varying curvature). *There are noncircular analytic three-obstacle periodic tables with a marked nongrazing three-contact polygon to which the preceding inverse applies. The marked polygon stays fixed in a family whose three contact curvatures vary independently. With the lattice and obstacle labels fixed, the phase-resolved two-offset laws locally determine every full obstacle image in that family.*

*Proof.* Use the three-disk realization in Proposition 2.6, or a sufficiently small scalene perturbation of its contact triangle. Let  $q_i$  be its vertices and  $n_i$  the inward unit bisectors of the triangle, which are the outward normals of the obstacles there. For unit reference radii the centers are  $q_i - n_i$ . Choose  $t_i$  perpendicular to  $n_i$ . In polar angle  $\vartheta$  about each center, replace the radius by

$$R_i(\vartheta) = 1 + \delta_i(1 - \cos \vartheta)(1 + \eta_i \sin \vartheta),$$

and use radial direction  $n_i \cos \vartheta + t_i \sin \vartheta$ . The value and first derivative of the radius at zero are 1, 0. Thus the contact point and tangent are unchanged. Its second derivative at zero is  $\delta_i$ , and the polar curvature formula gives contact curvature  $1 - \delta_i$ . The three parameters  $\delta_i$  therefore vary the three curvatures independently.

For sufficiently small parameters these are analytic strictly convex closed curves: the polar curvature numerator  $R_i^2 + 2(R_i')^2 - R_i R_i''$  stays positive by continuity from one. Their disjointness, lifted disjointness and selected-flight clearance persist by the strict margins in Proposition 2.6. Reflection on the fixed polygon continues to hold because its normals are fixed. Nonzero generic parameters give noncircular obstacles; unequal choices and a scalene triangle require no axial symmetry. Every obstacle is visited. The local analytic germs belong to the inverse neighborhood after the parameters are made small, so Corollary 3.7 proves the final assertion. The physical period is three; the returning oriented forward period may be six. Only three unknown contact germs are inverted.  $\square$

The new inverse and the alternating inverse have different data requirements. Here the full marked polygon and phase-resolved laws at two offsets are supplied, while curvatures and complete contact germs are recovered locally. The alternating results below need only the gap and a single law of each contact type, and retain their unknown-lattice, finite-ambiguity, calibration and total-preparation conclusions. Neither observation model is substituted for the other. In the general-period experiment the exact success mass remains the one in Theorem 2.5; conditional reconstruction does not make rejected preparations free.

## 4. GLOBAL IDENTIFICATION OF MARKED CONTACT GERMS

The local analytic coordinate theorem of the preceding section controls perturbations of complete germs. Exact identification has a stronger form: for a fixed marked polygon there is no second positive-curvature branch with the same actions. The point is already visible at degree

two. A Schur-complement identity converts the curvature inverse into a strict contraction whose constant is determined by the observed action Hessians. The signed higher-degree recursion then gives exact uniqueness without a closeness assumption. The analytic estimates remain local.

Throughout this section the polygon, graph coordinates and signs are those of Section 3. In particular  $k_i = L_i^{-1} > 0$ ,  $c_i = F_i''(0) > 0$  and  $s_i = (S_i^-)''(0)$ ; indices are modulo the physical contact period  $r$ . We regard  $k$  as fixed and write  $\mathcal{R}_k(c) = s$ . This is a map of quadratic data, independent of the higher contact jets. Its definition uses the actual half-line construction, not an arbitrary solution of a Riccati equation.

#### 4.1. The quadratic inverse.

**Lemma 4.1** (Boundary Schur complement). *For every positive curvature vector  $c$ , the actual half-line action Hessians satisfy  $0 < c_i < s_i$  and*

$$s_i = k_i + c_i - \frac{k_i^2}{k_i + c_{i+1} + s_{i+1}}, \quad a_i = \frac{k_i}{k_i + c_{i+1} + s_{i+1}} \in (0, 1). \quad (4.1)$$

Here  $a_i$  is the unsigned one-step contraction in Lemma 3.1. Conversely, positive vectors  $c, s$  satisfying the first identity have  $s = \mathcal{R}_k(c)$ .

*Proof.* The half-line starting at  $i$  is obtained by adjoining its first edge to the half-line starting at  $i + 1$ . At quadratic order its stationary value is the minimum over  $y$  of

$$\frac{1}{2}(k_i + c_i)u^2 - \epsilon_i k_i u y + \frac{1}{2}(k_i + c_{i+1} + s_{i+1})y^2.$$

The strictly positive half-line Hessian established in the preceding section gives  $s_{i+1} > 0$ . Eliminating  $y$  proves the first identity; the stationary ratio is  $y/u = \epsilon_i a_i$ . The denominator exceeds  $k_i$ , proving both  $a_i < 1$  and  $s_i > c_i$ .

For the converse, define  $a_i$  by the second identity. Substituting the first identity at  $i + 1$  gives

$$\frac{k_i}{a_i} = k_i + k_{i+1} + 2c_{i+1} - k_{i+1}a_{i+1}.$$

Thus the products  $y_j = a_i \cdots a_{i+j-1}u$  solve the sign-conjugated Jacobi equation with boundary value  $u$ . They decay exponentially because there are finitely many phases and every  $a_i < 1$ . Coercivity makes this decaying solution unique: the difference of two such solutions has zero boundary value, and summation of its positive Jacobi quadratic form gives zero. Restoring the signs gives the actual linearized half-line. Its initial action Hessian is  $k_i + c_i - k_i a_i = s_i$ . Consequently these are the actual action Hessians.  $\square$

For  $s \in (0, \infty)^r$  and  $z \in [0, \infty)^r$ , put

$$\begin{aligned} (\Phi_s z)_i &= \left[ s_i - k_i + \frac{k_i^2}{k_i + s_{i+1} + z_{i+1}} \right]_+, \quad [t]_+ = \max(t, 0), \\ q(s) &= \max_i \left( \frac{k_i}{k_i + s_{i+1}} \right)^2 < 1. \end{aligned} \quad (4.2)$$

The positive-part operation specifies an inverse algorithm even for positive data that are not realizable by strictly positive curvatures. It does not add zero-curvature tables to the geometric class.

**Theorem 4.2** (Global curvature coordinates). *The map  $\mathcal{R}_k : (0, \infty)^r \rightarrow (0, \infty)^r$  is one-to-one and a real analytic diffeomorphism onto its open image. That image and the inverse admit the following description.*

For every  $s > 0$ , the map  $\Phi_s$  has a unique fixed point  $z(s)$  in the closed positive orthant. Its iterates from any  $z^{(0)} \geq 0$  converge to  $z(s)$ , with

$$\|z^{(m)} - z(s)\|_\infty \leq \frac{q(s)^m}{1 - q(s)} \|z^{(1)} - z^{(0)}\|_\infty \quad (m \geq 0). \quad (4.3)$$

The first-increment estimate (4.3) is an a priori bound. A residual at the current iterate gives the a posteriori bound

$$\|z^{(m)} - z(s)\|_\infty \leq E_m(s), \quad E_m(s) := \frac{\|\Phi_s z^{(m)} - z^{(m)}\|_\infty}{1 - q(s)}. \quad (4.4)$$

For exact iteration its numerator is  $\|z^{(m+1)} - z^{(m)}\|_\infty$ . Precisely,

$$s \in \mathcal{R}_k((0, \infty)^r) \iff z_i(s) > 0 \text{ for every } i; \quad \mathcal{R}_k^{-1}(s) = z(s). \quad (4.5)$$

For any  $s, \tilde{s} > 0$ , define the observable constant

$$q_* = \max_i \left( \frac{k_i}{k_i + \min(s_{i+1}, \tilde{s}_{i+1})} \right)^2 < 1.$$

Then, including for data outside the geometric image,

$$\|z(s) - z(\tilde{s})\|_\infty \leq \frac{1 + q_*}{1 - q_*} \|s - \tilde{s}\|_\infty. \quad (4.6)$$

*Proof.* For  $z \geq 0$  the expression inside the positive part is strictly smaller than  $s_i$ . Hence  $\Phi_s$  maps the closed orthant into  $\prod_i [0, s_i]$ . Its untruncated expression has derivative in  $z_{i+1}$  equal to  $-k_i^2/(k_i + s_{i+1} + z_{i+1})^2$ , of absolute value at most  $q(s)$ . The positive part is one-Lipschitz. The map is therefore a strict contraction in the maximum norm on the complete closed orthant. The contraction theorem and summation of the successive differences prove existence, uniqueness and (4.3). For any  $v \geq 0$ , the triangle inequality and contraction give

$$\|v - z(s)\|_\infty \leq \|v - \Phi_s v\|_\infty + q(s) \|v - z(s)\|_\infty.$$

This proves (4.4). It also provides a bound for an inexact evaluation: if  $w$  satisfies  $\|w - \Phi_s v\|_\infty \leq \eta$ , replace the residual numerator by  $\|w - v\|_\infty + \eta$ . An error bound for evaluating the map, when used, is part of this certificate and is not assumed to be zero.

If  $s = \mathcal{R}_k(c)$  for  $c > 0$ , Lemma 4.1 says that  $c = \Phi_s c$ . Therefore  $c = z(s)$ , which proves global injectivity. If  $z(s) > 0$ , no positive part is active at the fixed point, so (4.1) holds with  $c = z(s)$ . The converse in that lemma proves the other implication in (4.5). This proves the image characterization without identifying formal positive data with geometric data.

For the two-point estimate, changing  $s_i$  contributes at most  $\|s - \tilde{s}\|_\infty$  to  $\Phi_s z$ ; changing  $s_{i+1}$  in its denominator contributes at most  $q_* \|s - \tilde{s}\|_\infty$ . All intermediate denominators are bounded below by the one used in  $q_*$ . Changing  $z$  contributes at most  $q_* \|z - \tilde{z}\|_\infty$ . Apply these bounds at the two fixed points and move the last term to the left. This proves (4.6).

On the image the fixed point is strictly positive. Locally the positive part can be omitted, and the remaining equation is real analytic in  $(s, c)$ . Its derivative in  $c$  is  $I + T_2$ , where  $(T_2 v)_i = a_i^2 v_{i+1}$ , and has a convergent Neumann inverse because  $\|T_2\|_\infty < 1$ . Thus the inverse is locally real analytic. The forward half-line Hessian is real analytic in  $c$ : for example the positive half-line Jacobi operator with periodic coefficients has a locally uniformly convergent resolvent expansion, and the boundary Schur complement is its analytic matrix element. Differentiating (4.1) also gives

$$(I - T_2) ds = (I + T_2) dc. \quad (4.7)$$

Both factors are invertible. Hence the forward map is a local real analytic diffeomorphism at every  $c > 0$ . Its image is open, and the already proved injectivity makes these local inverses one global inverse on that image.  $\square$

The projection in (4.2) is useful only in the reconstruction algorithm; it plays no role in the analytic inverse near a realized datum. Formula (4.3) also gives a finite sufficient test for membership in the image: if every component of an iterate exceeds either certified error bound, then  $z(s) > 0$ . For the residual bound this test is  $\min_i z_i^{(m)} > E_m(s)$ . It certifies membership in the quadratic image, not extension to a prescribed collection of globally disjoint analytic obstacles. The constants are allowed to deteriorate as an action Hessian tends to zero. No uniform assertion at a grazing or zero-curvature limit is used.

There is a second direct uniqueness check. If two positive curvature vectors have the same  $s$ , subtraction of (4.1) gives

$$c_i - \tilde{c}_i = -a_i \tilde{a}_i (c_{i+1} - \tilde{c}_{i+1}).$$

Iteration around the cycle forces the difference to vanish, since the product of the  $a_i \tilde{a}_i$  is strictly smaller than one. This argument treats odd and even physical periods alike. It uses only quadratic data; the signs in odd higher contact orders still have to be retained in the geometric reconstruction.

For a normal two-contact channel,  $k_0 = k_1 = 1/g$  and  $c_i$  are the physical curvatures. Set  $H = (1 + g^2 s_0 s_1)^{1/2}$ . The recovered curvatures are

$$c_0 = \frac{H\sqrt{s_0/s_1} - 1}{g}, \quad c_1 = \frac{H\sqrt{s_1/s_0} - 1}{g}.$$

These are precisely the alternating curvature formulas of Theorem 23.1. Indeed put  $u_i = g s_i$  and  $C_i = 1 + g c_i$ ; then  $C_0 C_1 = 1 + u_0 u_1$  and  $C_0/C_1 = u_0/u_1$ , which give  $(C_0 - u_0)(C_1 + u_1) = 1$  and its reversed version. These are the two identities in (4.1). Whenever the displayed curvatures are positive, uniqueness identifies them with the fixed point. This is a specialization of the geometric inverse, not a reduction of the two observation models.

Not every positive vector is an action datum for positive curvatures. For example  $r = 2$ ,  $k_0 = k_1 = 1$  and  $s = (1/10, 10)$  give  $z(s) = (0, 109/11)$ . The first component of the untruncated map is negative for every  $z \geq 0$ , and the second fixed-point equation gives the stated value. Thus this  $s$  is excluded by (4.5), even though the stable reconstruction algorithm is defined there.

#### 4.2. Exact rigidity without a local branch choice.

**Theorem 4.3** (Global marked contact identification). *Fix a clear nongrazing marked polygon with distinct physical contact points as in Section 3. For every  $M \geq 2$  the map from positive-curvature smooth contact jets through order  $M$  to the future action jets through that order is one-to-one, with a smooth inverse on its image. No closeness between the candidate jets is required.*

*Consequently, the phase-resolved two-offset limiting laws (3.5), on their common positive coordinate square, identify every finite smooth contact jet. Among positive-curvature analytic contact tuples with this same polygon they identify the complete contact germs, without a local-neighborhood hypothesis. The two offsets share their endpoint amplitude factors as in that formula; the amplitudes, curvatures and graph heights are not supplied.*

*If two actual tables have the same marked polygon, fixed lattice and obstacle labels, their boundaries are connected closed embedded analytic curves, and the polygon visits every labelled obstacle, equality of these laws implies equality of the complete labelled obstacle images. The tables need not be close. The fixed placement and full visitation are part of this conclusion.*



*Proof.* At degree two apply Theorem 4.2. This determines  $c$ , hence the actual  $a_i$  by (4.1), and the signed factors  $\sigma_i = \epsilon_i a_i$ . At degree  $n \geq 3$  use the actual-smooth recursion from Theorem 3.4:

$$q_n = (I - T_n)(I + T_n)^{-1}\{s_n - \mathcal{R}_n(q_2, \dots, q_{n-1})\}.$$

Here  $T_n$  is the signed shift of (3.11), now fully known from the reconstructed quadratic jets and the supplied polygon. Its inverse exists at every degree. Induction determines all higher jets uniquely. The lower-degree remainders are functions of actual smooth jets by Lemma 3.3, so this is not an identification of unrelated formal expansions. The curvature inverse is smooth on its image, and the remaining finite recursion is smooth, proving the finite-order assertion. Arbitrary finite higher jets with  $c > 0$  have local smooth graph representatives; after restricting their collars they satisfy the same local geometric conditions.

Proposition 3.2 extracts the physical actions  $A_i, C_i$  from the two laws. The known momenta give  $S_i^- = A_i + p_i x$ . Thus equality of laws gives equality of every future action jet and hence of every contact jet. In the analytic class the Taylor series of each contact graph converges on some neighborhood of zero. Equality of all its coefficients gives equality of the graph germs. One may restrict two different analytic radii to a smaller common one; no comparison in a prescribed local inverse ball is needed for this exact assertion. In the smooth category this jet argument concludes equality of jets, not equality of germs from flatness alone. Theorem 5.3 supplies the additional functional argument using equality of the complete real actions.

For actual analytic tables this supplies a common open arc on every visited obstacle. The continuation argument in Corollary 3.7 then identifies each entire connected closed boundary image. More explicitly, at a limit point of an interval of common arc both curves have the same tangent. They are analytic graphs over that tangent on a common neighborhood, and their identity continues across the limit point. Compactness and connectedness continue it around the boundary. All obstacles are visited and both lattice and placement were supplied, so the labelled tables coincide.  $\square$

Theorem 4.3 strengthens the exact uniqueness clauses in Corollary 3.7; the latter remains useful for its physical derivation and local estimates. The new exact argument does not replace the complete function-space inverse in Theorem 3.6. That theorem controls all lower-degree couplings in a fixed analytic norm, which coefficient uniqueness alone cannot do. In particular the analytic prior and radius loss in Theorem 3.8 are unchanged.

**Proposition 4.4** (Finite-flight reconstruction on a positive class). *Fix  $M \geq 2$  and a compact class of marked smooth contact data with positive curvature, the fixed-order smooth bounds of the relative law, and the positive-density, residual and separated-offset margins of Proposition 3.2. Use the same fixed polygon. For sufficiently small errors, the two-offset densities of  $N$ -flight bridges, with  $N$  a multiple of a returning oriented period, recover the contact jets through order  $M$  with error at most*

$$C_M(\tau^N + \varepsilon). \quad (4.8)$$

Here  $\varepsilon$  is the maximum interior  $C^M$  error at the two offsets and every phase. The quadratic reconstruction is the fixed point (4.2), not a selected local solution of the curvature equation. All later degrees use the signed recursion.

More precisely, write  $t = (t_2, \dots, t_M)$  for the extracted future action jets and stop the quadratic iteration at  $v = z^{(m)}(t_2)$ . Use the off-fixed-point recursion specified in the proof. For sufficiently small  $\tau^N + \varepsilon + E_m(t_2)$  it is defined and satisfies, in the maximum norm of the derivative coefficients,

$$\|\hat{q}_{\leq M}^{(m)} - q_{\leq M}\|_\infty \leq C_M(\tau^N + \varepsilon) + L_M E_m(t_2) \leq \tilde{C}_M\{\tau^N + \varepsilon + E_m(t_2)\}. \quad (4.9)$$

The constants are uniform on the stipulated compact class, with  $M$  fixed. The first-increment bound in (4.3) may replace  $E_m$ . The factor  $L_M$  includes the sensitivity of all the lower-degree couplings in the higher-jet reconstruction.

*Proof.* The relative physical law first compares the finite-flight densities with their realizable limits, with  $C^M$  error  $O(\tau^N)$  uniformly on the compact class. Stable rational extraction then gives the gauged future action jets with error  $O(\tau^N + \varepsilon)$ . Subtract their constant and linear Taylor terms, as in the preceding section; the true gauged action is fixed by this bounded projection. The true quadratic data have a common positive floor. For sufficiently small errors the estimated  $s$  remain positive, so (4.6) applies with a uniform  $q_* < 1$ . Its output is close to the true  $c$ , whose components also have a positive floor. The reconstructed  $c$  are therefore positive, and the projection is inactive at their fixed point.

At each of the finitely many remaining orders the coefficients, the inverse signed block and the lower-jet remainder are smooth on bounded sets of jets with a positive curvature margin. Inductively their error is  $O(\tau^N + \varepsilon)$ . The finite constants can be bounded on a common compact jet box; the order  $M$  is fixed throughout. This proves (4.8).

We specify the finite-iteration reconstruction, since the error in curvature must pass through the subsequent nonlinear steps. Keep the extracted action vector  $t$  fixed. For  $v > 0$  near  $z(t_2)$  set

$$\begin{aligned}\hat{\sigma}_i(v; t_2) &= \frac{\epsilon_i k_i}{k_i + t_{2,i+1} + v_{i+1}}, \\ (\hat{T}_n(v; t_2)w)_i &= \hat{\sigma}_i(v; t_2)^n w_{i+1}, \\ D_n(v; t_2) &= (I - \hat{T}_n)(I + \hat{T}_n)^{-1}.\end{aligned}\tag{4.10}$$

Here and below the signs are those of the physical cycle. Define  $\mathcal{J}_2(v; t) = v$  and, for  $3 \leq n \leq M$ , define

$$\mathcal{J}_n(v; t) = D_n(v; t_2)\{t_n - \mathcal{R}_n(\mathcal{J}_2(v; t), \dots, \mathcal{J}_{n-1}(v; t))\}.\tag{4.11}$$

The lower-jet remainder is the actual smooth function of Theorem 3.4. When  $v = z(t_2)$ , the Schur identity identifies  $\hat{\sigma}$  with the actual orbit factors, so this is exactly the finite-jet inverse. Away from the fixed point (4.11) is an explicitly chosen algebraic extension of that inverse; no nonstationary tuple is asserted to solve the physical half-line problem. This prescription avoids leaving unspecified how an approximate curvature is used in the later blocks.

Choose a common compact positive curvature box and, inductively, compact boxes for the finitely many lower jets and action coefficients. They can contain all exact reconstructions and their sufficiently small perturbations: smoothness, compactness and the finite number of steps suffice. Enlarge these boxes slightly so the line segments used below remain in them. On these boxes let  $\alpha < 1$  bound  $|\hat{\sigma}_i|$ , and let  $d_*$  bound  $k_i/(k_i + t_{2,i+1} + v_{i+1})^2$ . Then

$$\begin{aligned}\|D_n\|_\infty &\leq b_n := \frac{1 + \alpha^n}{1 - \alpha^n}, \\ \|D_v D_n\| &\leq \beta_n := \frac{2n\alpha^{n-1}d_*}{(1 - \alpha^n)^2}.\end{aligned}\tag{4.12}$$

Indeed  $D_n = 2(I + \hat{T}_n)^{-1} - I$ , and differentiating this identity places  $D_v \hat{T}_n$  between two resolvents. Its norm is at most  $n\alpha^{n-1}d_*$ . Thus no commutation of a matrix variation with the cyclic shift has been assumed.

Let  $U_n$  bound  $\|t_n - \mathcal{R}_n(w_{<n})\|_\infty$  on the chosen boxes and let  $r_n$  bound the derivative norm of  $\mathcal{R}_n$  in the joint maximum norm of its lower jets. Both are finite by the actual smooth-jet factorization. The numbers

$$L_2 = 1, \quad L_n = \max\{L_{n-1}, \beta_n U_n + b_n r_n L_{n-1}\} \quad (3 \leq n \leq M)\tag{4.13}$$

are explicit sufficient propagation constants once these finite smooth bounds are specified. They are not claimed to be known from the noisy law alone. Subtract (4.11) at two curvatures  $v, w$ , keeping  $t$  fixed. The change in its block is at most  $\beta_n U_n \|v - w\|_\infty$ ; the change in the remainder is at most  $b_n r_n$  times the maximum lower-jet difference. Induction therefore proves

$$\|\mathcal{J}_{\leq n}(v; t) - \mathcal{J}_{\leq n}(w; t)\|_\infty \leq L_n \|v - w\|_\infty.$$

Apply this with  $w = z(t_2)$  and use (4.4). The exact-fixed-point reconstruction already differs from the true jets by  $C_M(\tau^N + \varepsilon)$ . Their sum gives (4.9). Taking the combined error small keeps all the iterates used here in the common boxes and in particular keeps  $v$  positive. Alternatively the displayed residual positivity test certifies that  $z(t_2)$  lies in the quadratic image. For any strictly positive exact fixed point the iterates eventually pass that test, because their residuals tend to zero.

Error in an observed density is the input here, not a number of full-phase preparations; the preparation theorems retain their own observation and calibration hypotheses. All constants in this proof are at fixed order. In particular they do not give an order-uniform inverse in the unweighted derivative norm.  $\square$

## 5. SMOOTH CONTACT RIGIDITY FROM THE COMPLETE BOUNDARY LAW

The boundary law determines functions on a real collar, not merely their Taylor coefficients. This distinction removes the analytic hypothesis from local contact identification. After the finite jets have been identified, the actual stationary envelope separates the remaining functions: its initial evaluation is invertible, whereas every later evaluation takes place strictly nearer the contact. A weighted real norm makes the latter part a contraction. The argument uses neither convergence of Taylor series nor a holomorphic extension.

Throughout this section the polygon, the physical phase labels and the graph sensors are those of Section 3. We write  $\mathcal{S}(F) = (S_b^-)_{b=0}^{r-1}$  for the actual future actions, with the same anchored gauges. Contact identification below means equality on a possibly smaller real neighborhood of each marked contact. Continuation to unvisited portions of a smooth obstacle is not part of that conclusion.

### 5.1. A common real collar and successive visits.

**Lemma 5.1** (Uniform real visits). *Let  $F^t$ ,  $0 \leq t \leq 1$ , be a smooth compact family of anchored  $C^\infty$  graph tuples with positive quadratic jets, in the fixed marked coordinates. There are  $R > 0$  and  $a < 1$  such that the actual future half-lines are defined for  $|u| \leq R$  and all  $t$ , and*

$$x_{b,j}^t(u) = X_{b+j-1}^t \circ \cdots \circ X_b^t(u), \quad |x_{b,j}^t(u)| \leq a^j |u|, \quad j \geq 1, \quad (5.1)$$

where  $X_i^t(u) = x_{i,1}^t(u)$ . On this collar the real envelope weights of (3.9) satisfy

$$\frac{1}{2} \leq \alpha_b^t(u) \leq \frac{3}{2}, \quad 1 \leq v_{b,j}^t(u) \leq 3 \quad (j \geq 1). \quad (5.2)$$

The constants can be chosen uniformly on bounded finite-smoothness classes with fixed positive curvature bounds and fixed nongrazing marks. For each stipulated differentiation order only finitely many bounds for the graph derivatives are needed.

*Proof.* The positive Jacobi construction and its weighted implicit equations in Sections 2 and 3 give the unique decaying stationary half-line. Their inverses are uniform on compact positive quadratic classes. Applying the construction with the parameter  $t$  gives a common collar and continuous first derivatives of  $X_i^t$  there. By Lemma 3.1,

$$X_i^t(0) = 0, \quad (X_i^t)'(0) = \sigma_i(t), \quad a_0 := \max_{i,t} |\sigma_i(t)| < 1.$$

Fix  $a_0 < a < 1$  and shrink the collar until  $\sup_{i,t,|u|\leq R} |(X_i^t)'(u)| \leq a$ . A shifted tail solves the same decaying stationary problem at the next physical phase. Uniqueness therefore identifies it with that phase's half-line. Iteration proves the first identity in (5.1); the mean value theorem then proves the second. In particular the multiplicative constant is one. The weaker estimate  $C\rho^j|u|$  alone would not suffice for the large-weight argument below when  $C\rho > 1$ .

Each unit chord is a smooth function of two contact coordinates on a common neighborhood of the clear reference chord. There are only finitely many physical phases. At the origin the two weights are 1 and 2, respectively. Uniform continuity, followed by a further shrink of  $R$ , gives (5.2) for every incident pair of chords, hence for every  $j$ . Signed tangent choices have not been changed; these are normal projection weights of actual chords. For a bounded finite-smoothness class, the same proof uses compactness of its positive quadratic range and uniform bounds in the weighted implicit equations. Bounds through any fixed order follow by successive differentiation with a strict exponential margin. No uniform assertion as curvature or the nongrazing margin vanishes is used.  $\square$

For an integer  $m \geq 2$  define the weighted real norm

$$\|h\|_{m,R} = \max_b \sup_{0 < |u| \leq R} \frac{|h_b(u)|}{|u|^m}. \quad (5.3)$$

It is finite for a  $C^m$  tuple whose derivatives of orders less than  $m$  vanish at zero. It is a norm on such functions, not a norm of their formal series.

**Proposition 5.2** (Separation of the nonlinear remainder). *Suppose that two smooth graph tuples  $F, G$  have the same contact jets through order  $m - 1$ . Apply Lemma 5.1 to their segment  $F^t = F + t(G - F)$  on a common collar. If*

$$\theta_m := \frac{6a^m}{1 - a^m} < 1, \quad (5.4)$$

then

$$\|G - F\|_{m,R} \leq \frac{2}{1 - \theta_m} \|\mathcal{S}(G) - \mathcal{S}(F)\|_{m,R}. \quad (5.5)$$

The same assertion holds with finite  $C^m$  differences whenever the real half-lines and envelope used here have the indicated regularity.

*Proof.* Put  $h = G - F$  and  $g = \mathcal{S}(G) - \mathcal{S}(F)$ . Integrate the actual envelope of Lemma 3.3 along the segment:

$$g_b(u) = \bar{\alpha}_b(u)h_b(u) + (Kh)_b(u), \quad \bar{\alpha}_b(u) = \int_0^1 \alpha_b^t(u) dt, \quad (5.6)$$

where the linear operator  $K$ , with this segment fixed, is

$$(Kh)_b(u) = \int_0^1 \sum_{j \geq 1} v_{b,j}^t(u) h_{b+j}(x_{b,j}^t(u)) dt.$$

The operator is applied to the actual difference, not to a freely chosen formal variation. The initial coefficient is independent of  $h_b$  in this displayed identity once the segment is fixed. For any function in the weighted space,

$$|h_{b+j}(x_{b,j}^t(u))| \leq \|h\|_{m,R} a^{jm} |u|^m.$$

This summable majorant justifies integration and summation and gives

$$\|Kh\|_{m,R} \leq \frac{3a^m}{1 - a^m} \|h\|_{m,R}, \quad \|\bar{\alpha}^{-1}\|_\infty \leq 2.$$

Multiplication by  $\bar{\alpha}^{-1}$  in (5.6) and the triangle inequality prove (5.5). Alternatively, the inverse of  $I + \bar{\alpha}^{-1}K$  is its norm-convergent Neumann series. The coefficients and composition maps of that operator depend on both candidates; this proof does not present that pair-dependent inverse as an algorithm using only the observations.

The finite-jet factorization already proved for actual smooth graphs shows that  $g$  has the same vanishing order as required for this norm. The terminal contribution in the finite envelope has already been sent to zero with uniform parameter bounds before the identity is integrated. It is not dropped by declaring a formal half-line limit.  $\square$

## 5.2. Identification of smooth germs.

**Theorem 5.3** (Smooth marked contact rigidity). *Fix a clear nongrazing marked polygon with distinct physical contacts. Two positive-curvature  $C^\infty$  contact tuples with the same phase-resolved two-offset laws (3.5) have identical contact germs in the measured graph coordinates. No analyticity, curvature input, separately calibrated local flux amplitude or closeness between the candidates is required. In fact equality of their actual future action germs at all physical phases already implies the conclusion.*

*Proof.* The two-offset quotient in Proposition 3.2 recovers the future actions on a common real collar, with the marked physical momenta restored. Thus it suffices to treat equality of  $\mathcal{S}(F)$  and  $\mathcal{S}(G)$  there. Theorem 4.2 first identifies the curvatures without selecting a local inverse branch. The signed recursion of Theorem 3.4 then identifies every finite contact jet. In particular  $h = G - F$  is flat at zero at each physical phase.

The segment between these two actual graph representatives has the same positive quadratic jets. After shrinking the collar it consists of admissible local strictly convex arcs with the fixed stationary polygon. No interpolation of globally closed obstacles is necessary. Lemma 5.1 supplies one  $a < 1$  and one common collar for this segment. Now choose a finite integer  $m \geq 3$  so large that (5.4) holds. Flatness and Taylor's formula make  $\|h\|_{m,R}$  finite. Proposition 5.2 with  $g = 0$  forces this norm to vanish. Hence  $F = G$  on that collar.

Only finitely many jets are used after  $a$  has been chosen. The proof does not infer equality of smooth functions from their Taylor series: it uses equality of their full actions once more, in (5.6). A flat nonzero contact perturbation has a nonzero, possibly flat, action response.  $\square$

For connected closed analytic obstacle boundaries, fixed placement and full visitation still give the whole-obstacle conclusion of Theorem 4.3. For smooth boundaries the new conclusion concerns the visited contact neighborhoods; arbitrary unvisited boundary arcs are not determined by local half-line data. The bounded-holomorphic local inverse of Theorem 3.6 is also retained: equality of real germs and inversion in a specified analytic Banach norm are different assertions.

**5.3. Finite-smoothness control from real densities.** The weighted estimate also gives stability of complete real profiles, rather than only finite jets, under a finite-smoothness prior. We state the derivative loss and the range of uniformity explicitly. Fix a marked polygon, a real interval  $I_0 = [-R_0, R_0]$ , positive numbers  $c_- < c_+$ , and a bound  $K_0$ . For an integer  $m$  let  $\mathcal{K}_m$  consist of anchored graph tuples satisfying

$$c_- \leq F_i''(0) \leq c_+, \quad \max_i \|F_i\|_{C^{m+3}(I_0)} \leq K_0. \quad (5.7)$$

We use a sufficiently small common collar of their fixed clear polygon. Curvature remains positive there by the  $C^3$  bound. Choose  $a < 1$  larger than all the one-step linear contraction factors for the enlarged quadratic box  $[c_-/2, 2c_+]^r$ , and then choose  $m \geq 3$  satisfying (5.4). This order is fixed by the positive marked class; it need not remain bounded in a weak-hyperbolicity limit.

**Lemma 5.4** (Finite-order comparison bounds). *On a sufficiently small common collar  $I = [-R, R]$ , the following bounds hold with constants depending only on the stipulated class. For  $F, G \in \mathcal{K}_m$ , if  $\eta = \|\mathcal{S}(F) - \mathcal{S}(G)\|_{C^m(I)}$ , then*

$$\max_{i, 2 \leq j < m} |F_i^{(j)}(0) - G_i^{(j)}(0)| \leq C\eta. \quad (5.8)$$

*If  $P$  is an anchored polynomial tuple of degree at most  $m - 1$ , and the segment  $G + tP$  stays in an enlarged bounded positive class, then*

$$\|\mathcal{S}(G + P) - \mathcal{S}(G)\|_{C^m(I)} \leq C\|P\|_{C^{m+3}(I_0)}. \quad (5.9)$$

*The enlarged class may be fixed before the small perturbation is chosen.*

*Proof.* For degree two, (4.6) gives a uniform bound because the positive action Hessians have a common lower bound. The signed factors depend smoothly on curvature in the enlarged positive box. At every subsequent degree subtract the two inversions (3.12). Their inverse blocks and derivatives are bounded on that box. The remainders  $\mathcal{R}_j$  are smooth functions of the finitely many lower graph jets on bounded boxes with positive quadratic entries. Their derivatives there are bounded. Induction, using the observed action-jet differences, proves (5.8). This is a two-point comparison; nonsingularity of an everywhere defined Jacobian is not being used as a substitute for global injectivity or a global Lipschitz bound.

For (5.9), use the weighted stationary equations along  $G + tP$ . On a common collar their linearized interior operator is a uniformly small perturbation of the positive Jacobi operator. It has a uniform inverse on an exponentially weighted space. Successively differentiating in the endpoint through order  $m$  puts the same operator on the highest coordinate derivative. The other terms involve previously bounded derivatives and finitely many graph derivatives. A strict margin in the exponential weight absorbs the fixed polynomial factors in the site index. With some  $a < \rho < 1$ , this gives

$$|x_{b,j}^t(u)| \leq C\rho^j|u|, \quad \sup_{|u| \leq R} |\partial_u^k x_{b,j}^t(u)| \leq C_k \rho^j \quad (1 \leq k \leq m),$$

uniformly in  $t$ . Chord differentiation likewise bounds  $\|\alpha_b^t\|_{C^m}$  and  $\|v_{b,j}^t\|_{C^m}$  uniformly in  $b, j, t$ . The stipulated  $C^{m+3}$  graph bound is sufficient for these finitely differentiated equations and weights.

Because  $P(0) = P'(0) = 0$ , the chain rule gives the useful summable bound

$$\|P_{b+j} \circ x_{b,j}^t\|_{C^m([-R, R])} \leq C\rho^{2j}\|P\|_{C^m(I_0)}.$$

For a derivative involving one derivative of  $P$ , its value is  $O(|x_{b,j}^t|)$  and its coordinate-derivative factor decays as  $\rho^j$ . Terms with at least two derivatives of  $P$  have at least two such coordinate factors. The undifferentiated term is quadratic in  $x_{b,j}^t$ . These observations prove the bound for every order at most  $m$ , including the cases where a polynomial derivative is zero. Apply this to the actual envelope with variation  $P$ . The initial term has norm at most  $C\|P\|_{C^m}$  and the later terms have a summable  $C\rho^{2j}\|P\|_{C^m}$  majorant. The terminal term and its fixed derivatives vanish uniformly, as in Lemma 3.3. Thus the parameter derivative of the action has  $C^m$  norm at most  $C\|P\|_{C^{m+3}(I_0)}$ . Integration in  $t$  proves the claim. No optimal derivative count or same-regularity differentiability theorem for composition operators is being used.  $\square$

**Theorem 5.5** (Complete real-profile stability). *For the class (5.7) and  $m$  chosen as above, there are a common  $R > 0$ ,  $\eta_0 > 0$  and  $C < \infty$  such that*

$$\|F - G\|_{C^0([-R, R])} \leq C\|\mathcal{S}(F) - \mathcal{S}(G)\|_{C^m([-R, R])} \quad (5.10)$$

whenever  $F, G \in \mathcal{K}_m$  and the norm on the right is at most  $\eta_0$ . In particular, on compact positive-density classes with fixed separated offsets and residual floors,

$$\|F - G\|_{C^0([-R, R])} \leq C \max_{i, d \in \{d_1, d_2\}} \|f_{i,d}^F - f_{i,d}^G\|_{C^m([-R, R]^2)} \quad (5.11)$$

for sufficiently small right side. These are estimates on realizable pairs of laws. They use a finite real-smoothness prior, not an analytic prior, and do not assert an inverse on arbitrary density tuples.

*Proof.* Let  $\eta$  be the action norm in the statement and set

$$P_i(u) = \sum_{j=2}^{m-1} \frac{F_i^{(j)}(0) - G_i^{(j)}(0)}{j!} u^j, \quad G^\sharp = G + P.$$

Lemma 5.4 gives  $\|P\|_{C^{m+3}(I_0)} \leq C\eta$ . For sufficiently small  $\eta_0$ , both  $G + tP$  and the segment between  $F$  and  $G^\sharp$  remain in a fixed enlarged positive class. Their collars and visit estimates are therefore uniform. The graph tuples  $F, G^\sharp$  have identical jets through order  $m - 1$ . By actual finite-jet factorization so do their actions. Taylor's integral remainder consequently gives

$$\|\mathcal{S}(F) - \mathcal{S}(G^\sharp)\|_{m,R} \leq \frac{1}{m!} \|\mathcal{S}(F) - \mathcal{S}(G^\sharp)\|_{C^m([-R, R])} \leq C\eta.$$

The last inequality uses (5.9) and the definition of  $\eta$ . Applying (5.5) to this pair, and then removing the polynomial correction, proves

$$\|F - G\|_{C^0} \leq R^m \|F - G^\sharp\|_{m,R} + \|P\|_{C^0} \leq C\eta.$$

All collars were fixed from the enlarged class before this comparison. Finally, Proposition 3.2 converts the stated  $C^m$  density error to  $C^m$  action error on the positive square, proving (5.11).  $\square$

The theorem also proves exact contact identification on each of these finite-smoothness classes. Its quantitative conclusion is a  $C^0$  profile estimate from a  $C^m$  law error with a  $C^{m+3}$  graph bound. It is not an analytic-norm estimate, a derivative-free statistical rate, or a theorem uniform in the period. If two realizable candidates fit common finite-flight density records to accuracy  $\varepsilon$  in this same  $C^m$  norm, the relative theorem at order  $m$  gives

$$\|F - G\|_{C^0([-R, R])} \leq C(\tau^N + \varepsilon)$$

whenever the combined error is small and the forward constants are uniform on the class. This uses the triangle inequality through the common records. It does not manufacture density derivative bounds from total variation or assign a preparation budget. The curvature stopping error in Proposition 4.4 remains a separate, explicitly propagated computational error.

#### 5.4. A smooth physical distinction beyond all contact jets.

**Proposition 5.6** (Flat contacts distinguished by boundary laws). *Let a smooth strictly dispersing table have a clear nongrazing marked periodic polygon with distinct contact points. There is a small smooth family of actual tables, with the same polygon and free cell area, whose complete contact Taylor jets at every marked contact agree, but whose two-offset boundary laws do not agree on any sufficiently small common contact square. Every future action also has the same Taylor jets throughout the family. Thus a formal all-order contact record and the complete boundary-law record have different geometric information even in the fixed marked physical class.*

*Proof.* Choose one contact, called  $b$ , and a smooth nonnegative cutoff  $\chi$  supported in its graph chart and equal to one near zero. Replace its graph by

$$F_b^\varepsilon(u) = F_b^0(u) + \varepsilon\psi(u), \quad \psi(u) = \begin{cases} \chi(u)e^{-1/u^2}, & u \neq 0, \\ 0, & u = 0. \end{cases}$$

All other contact graphs are unchanged. The difference is flat but nonzero on every punctured neighborhood. For small  $\varepsilon$ , strict convexity, embeddedness, disjointness and the clearance margins persist. The fixed contact points, tangent frames and curvatures preserve the original periodic polygon exactly.

To keep the free area fixed, choose a smooth normal bump on a boundary arc disjoint from every contact collar, with nonzero first area variation. Such an arc exists after the finitely many collars are shrunk. For a positively oriented obstacle boundary  $R(s)$ , its area is  $\frac{1}{2} \int R \times R' ds$ ; differentiation and integration by parts give first variation  $\int \delta R \times R' ds$ . A nonzero nonnegative outward normal bump therefore has nonzero area derivative. The ordinary implicit function theorem chooses its coefficient smoothly in  $\varepsilon$  so that the obstacle area, hence the free cell area, is unchanged. The correction is outside all the relevant charts and leaves their stationary half-lines unchanged. This constructs actual closed obstacles, rather than only formal local coefficients.

All graph jets remain fixed. The actual finite-jet factorization therefore fixes every action jet as well. The exact linear reference twists, flight lengths, incidence angles and multipliers are also unchanged. Nevertheless the envelope on the perturbed contact detects the difference. On a smaller collar, (5.2) and positivity of  $\psi$  give, for  $\varepsilon > 0$  and  $u \neq 0$ ,

$$S_b^{-,\varepsilon}(u) - S_b^{-,0}(u) \geq \frac{\varepsilon}{2} e^{-1/u^2} > 0.$$

The area-correcting bump contributes nothing to this local action. More precisely, uniformly for a sufficiently small parameter range,

$$S_b^{-,\varepsilon}(u) - S_b^{-,0}(u) = \varepsilon e^{-1/u^2} \left\{ 1 + O(|u|) + O\left(e^{-(a^{-2}-1)/u^2}\right) \right\}. \quad (5.12)$$

Indeed  $\alpha_b^t(u) = 1 + O(|u|)$ , while each later perturbed visit has  $|x_{b,j}^t(u)| \leq a^j |u|$ . Dividing its flat value by  $e^{-1/u^2}$  and using  $a^{-2j} - 1 \geq j(a^{-2} - 1)$  bounds the entire tail by the convergent geometric series with ratio  $e^{-(a^{-2}-1)/u^2}$ . Visits to unperturbed phases contribute zero.

Equality of the two-offset laws would imply equality of these actions by the anchored quotient (3.6), a contradiction. The claim is not equality of the complete fixed-window count or of the normalized density Taylor coefficients: their normalizers can change. It is equality of all formal contact and action data, with the reference geometry and area fixed, in the presence of a nonzero physical-law signal. Nor is it a comparison of full marked length spectra, which are not the observations used here.  $\square$

The flat family explains why the complete-law hypothesis in Theorem 5.3 is substantive. It is not a formal consequence of all-order contact separation. Exact relative normalization first retains the real action functions in the physical law; the contracted visits then identify those functions with actual boundary profiles, including their flat parts. The finite-smoothness estimate records the derivative information needed for stable use of that same mechanism.

**5.5. A phase-dependent inverse for actual functions.** The scalar tail estimate above can be organized by the phases visited by a half-line. This gives an inverse criterion that retains heterogeneous contraction and separates it from the conditioning of finite jets. The criterion applies to actual functions and to an integrated nonlinear envelope; it does not require an analytic power series.

**Proposition 5.7** (Matrix domination of a contracted envelope). *Let  $h_b, g_b$  be continuous functions on  $[-R, R]$ ,  $0 \leq b < r$ , vanishing at zero and with finite weighted order- $m$  norms. Suppose*

$$g_b(u) = \alpha_b(u) h_b(u) + \int_0^1 \sum_{j \geq 1} v_{b,j}^t(u) h_{c(b,j)}(\psi_{b,j}^t(u)) dt, \quad (5.13)$$



where  $\alpha_b \geq \alpha_* > 0$ ,  $|v_{b,j}^t| \leq V_{b,j}$  and  $|\psi_{b,j}^t(u)| \leq \gamma_{b,j}|u|$  with  $0 \leq \gamma_{b,j} \leq 1$ . Define the nonnegative matrix

$$K_m(b, c) = \alpha_*^{-1} \sum_{j: c(b,j)=c} V_{b,j} \gamma_{b,j}^m. \quad (5.14)$$

Assume its entries are finite and its spectral radius is less than one. Writing

$$H_b = \sup_{0 < |u| \leq R} \frac{|h_b(u)|}{|u|^m}, \quad G_b = \sup_{0 < |u| \leq R} \frac{|g_b(u)|}{|u|^m},$$

one has the componentwise estimate

$$H \leq \alpha_*^{-1}(I - K_m)^{-1}G. \quad (5.15)$$

In particular, if  $w_b > 0$  and  $K_m w \leq \theta w$  for some  $\theta < 1$ , then

$$\|h\|_{m,R,w} \leq \frac{\alpha_*^{-1}}{1-\theta} \|g\|_{m,R,w}, \quad \|h\|_{m,R,w} := \max_b \sup_{0 < |u| \leq R} \frac{|h_b(u)|}{w_b |u|^m}. \quad (5.16)$$

Such a weight exists whenever the stated spectral-radius condition holds.

*Proof.* The summability in (5.14) is an absolute majorant for the envelope on the weighted space. Dividing (5.13) by  $\alpha_b(u)|u|^m$  and taking suprema yields  $H \leq \alpha_*^{-1}G + K_m H$ . Iterate this inequality  $n$  times. The finite matrix  $K_m^n$  tends to zero and  $\sum_{j \geq 0} K_m^j = (I - K_m)^{-1}$ , whose entries are nonnegative. Taking the limit proves (5.15). For a weight satisfying  $K_m w \leq \theta w$ , divide the componentwise inequality by  $w_b$  and take a maximum to prove (5.16). Conversely choose any  $\theta \in (\rho(K_m), 1)$  and set  $w = (I - K_m/\theta)^{-1}\mathbf{1}$ . Then  $w \geq \mathbf{1}$  and  $K_m w = \theta(w - \mathbf{1}) \leq \theta w$ . All statements concern a fixed integrated envelope. Its coefficients may depend on both candidates; no observation-only iteration is asserted.  $\square$

**Corollary 5.8** (Phase-weighted smooth separation). *For the graph interpolation in Proposition 5.2, suppose the one-step maps satisfy, on a common collar and for all interpolation parameters,*

$$|(X_i^t)'(u)| \leq \hat{a}_i < 1, \quad 0 < \hat{a}_i, \quad \bar{a} = \left( \prod_{i=0}^{r-1} \hat{a}_i \right)^{1/r}.$$

For any integer  $m \geq 3$  such that  $\bar{a}^m < 1/7$ , two graph tuples with equal jets through degree  $m-1$  satisfy

$$\|G - F\|_{m,R,w} \leq \frac{2}{1 - \vartheta_m} \|\mathcal{S}(G) - \mathcal{S}(F)\|_{m,R,w}, \quad \vartheta_m = \frac{6\bar{a}^m}{1 - \bar{a}^m}, \quad (5.17)$$

where  $w_0 = 1$  and

$$w_b = \left( \prod_{i=0}^{b-1} \frac{\bar{a}}{\hat{a}_i} \right)^m \quad (1 \leq b < r). \quad (5.18)$$

In the unweighted norm the right-side constant is multiplied by  $\kappa(w) := \max_b w_b / \min_b w_b$ .

*Proof.* The mean value theorem and the constant-one composition identity in Lemma 5.1 give

$$|x_{b,j}^t(u)| \leq \left( \prod_{\ell=0}^{j-1} \hat{a}_{b+\ell} \right) |u|.$$

Use  $\alpha_* = 1/2$ ,  $V_{b,j} = 3$  and  $c(b, j) = b + j \pmod{r}$  in Proposition 5.7. If  $(A_m z)_b = \hat{a}_b^m z_{b+1}$ , its domination matrix is

$$K_m = 6 \sum_{j \geq 1} A_m^j = 6A_m(I - A_m)^{-1}.$$

The positive weight (5.18) obeys  $A_m w = \bar{a}^m w$ , including the closing phase, because  $\bar{a}^r = \prod_i \hat{a}_i$ . It follows that  $K_m w = \vartheta_m w$ . In fact conjugation by  $\text{diag}(w)$  changes  $A_m$  into  $\bar{a}^m$  times the cyclic permutation matrix, so its nonnegative geometric sum has spectral radius  $\vartheta_m$ . The hypothesis is exactly  $\vartheta_m < 1$ . The preceding proposition proves (5.17). The norm inequalities

$$\|h\|_{m,R} \leq (\max_b w_b) \|h\|_{m,R,w}, \quad \|g\|_{m,R,w} \leq (\min_b w_b)^{-1} \|g\|_{m,R}$$

prove the unweighted estimate. The actual envelope, its terminal-limit justification and the finite-jet factorization are those already proved; only the estimate of its tail has changed.  $\square$

**Corollary 5.9** (Finite-smoothness inverse with phase bounds). *Fix a bounded positive marked class and an enlarged class for the polynomial alignments of Lemma 5.4. Suppose the latter has common one-step bounds  $\hat{a}_i$  as above. Choose  $m \geq 3$  with  $\bar{a}^m < 1/7$  and retain the  $C^{m+3}$  graph prior and the positive density, gate and offset floors. Then the conclusions (5.10) and (5.11) hold at this order  $m$  on a smaller common collar, for sufficiently small data discrepancy.*

*Proof.* The finite-jet comparison lemma holds at every stipulated finite order: its proof uses the global quadratic inverse, the signed finite blocks and the forward polynomial variation estimate, not the scalar tail threshold. Set  $\eta = \|\mathcal{S}(F) - \mathcal{S}(G)\|_{C^m}$  and form its polynomial correction  $P$  and  $G^\sharp = G + P$  exactly as in Theorem 5.5. Then  $\|P\|_{C^{m+3}} \leq C\eta$ . For small enough  $\eta$ , both required interpolations remain in the fixed enlarged positive class. The actions of  $F$  and  $G^\sharp$  have equal jets through order  $m - 1$ , and Taylor's formula and the forward variation bound give

$$\|\mathcal{S}(F) - \mathcal{S}(G^\sharp)\|_{m,R,w} \leq \frac{C\eta}{\min_b w_b}.$$

Apply (5.17), multiply by  $R^m \max_b w_b$ , and add  $\|P\|_{C^0}$ . This proves the action estimate with a finite class-dependent constant. The two-offset quotient on the same positive gate gives the density estimate. No curvature bound is inferred from the resulting  $C^0$  graph error; curvatures were controlled in the preceding jet alignment.  $\square$

The choice  $\hat{a}_i = a$  recovers the previous sufficient condition and constant. On a heterogeneous class the geometric-mean criterion may allow a smaller order than  $6(\max_i \hat{a}_i)^m / [1 - (\max_i \hat{a}_i)^m] < 1$ . The price is the explicit weight condition number, as well as the unchanged finite-jet constants. Thus this refinement improves a sufficient inverse criterion without asserting that the selected order is optimal or that the unweighted inverse is uniformly better. It also illustrates the reusable content of the smooth argument: invert the finite jets, then dominate the remaining actual-function envelope by a nonnegative visit matrix.

## 6. SMOOTH CONTACT RECOVERY FROM FINITE PREPARATIONS

The preceding stability theorem requires a differentiated density norm. A finite endpoint record does not supply this norm, and a discontinuous histogram cannot be substituted for it. We now construct the missing statistical comparison, for the same marked periodic polygon and the same two controlled offsets. The target is the complete real contact profile, not a truncated jet. The proof regularizes the endpoint data, compares realizable laws, and then applies the geometric inverse. All independent phase preparations, including failures, are counted.

**6.1. A fixed interior observation experiment.** Fix the marked polygon of Section 3, with  $r$  physical phases, a returning oriented period  $P$ , and offsets  $0 < d_1 < d_2$ . Write  $J = 2r$  for the number of phase–offset groups. Choose concentric squares

$$Q = [-R, R]^2 \Subset Q_+ = [-R_+, R_+]^2$$

inside one common positive-residual collar. A success means precisely the selected itinerary event of Theorem 2.5, with its recorded endpoints in  $Q_+$ . This gate is fixed in the calibrated tangent coordinates; it uses no unknown normal graph height. The success/failure tag, including gate membership, is recorded exactly. The bounded errors below affect endpoint readout after acceptance, not selection tags or the two offsets. Conditional densities below are normalized on  $Q_+$ , and compared on  $Q$ . Restricting them to  $Q$  does not renormalize them. The unknown scalar normalizers still cancel in Proposition 3.2.

Let  $\mathcal{C}$  be a nonempty class of actual marked tables whose contact tuples belong to the bounded positive class  $\mathcal{K}_m$  of (5.7). The integer  $m \geq 3$  is fixed so that (5.4) holds on the enlarged quadratic class. Fix also an integer  $s \geq 1$ . Impose the uniform geometric hypotheses of Theorems 2.3 and 2.5 through endpoint order  $m + s$ , including the finitely many higher graph bounds used for that order, and a positive upper bound for the free cell area. Positivity, clearance, incidence, flight-length bounds, and the separated offsets are fixed throughout. These are forward smoothness and geometric priors, not bounds assumed for an observed density error. No analytic prior is imposed. A compact smooth family with the stated positive margins is one admissible choice; the constants use only its finitely stipulated derivative bounds.

For  $T \in \mathcal{C}$ , denote its profiles by  $F^T$ , the limiting density tuple by  $f^T = (f_{b,d}^T)$ , the  $N$ -flight conditional density by  $f_{N,b,d}^T$ , and its success probability by  $p_{N,b,d}^T$ . Here and below  $N$  is a multiple of  $P$ . By shrinking both squares within the stability collar, we obtain constants  $H, C_0, c_0, \omega, \Gamma > 0$  and  $N_0$  such that

$$\sup_{T, b, d, N \geq N_0} (\|f_{N,b,d}^T\|_{C^{m+s}(Q_+)} + \|f_{b,d}^T\|_{C^{m+s}(Q_+)}) \leq H, \quad (6.1)$$

$$\max_{b,d} \|f_{N,b,d}^T - f_{b,d}^T\|_{C^m(Q_+)} \leq C_0 e^{-\omega N}, \quad (6.2)$$

$$p_{N,b,d}^T \geq p_*(N) := c_0 e^{-\Gamma N}. \quad (6.3)$$

We take  $c_0 \leq 1$ . To see the uniformity explicitly, the numerator of the conditional law is the normalized twist times its physical residual. The relative estimates bound its indicated derivatives; a fixed positive square bounds its integral away from zero. Division therefore gives the first two estimates. For the last one the exact reference twist is

$$D_{N,b} = \frac{\sinh \chi}{(\mathcal{M}_b)_{12} \sinh(N\chi/P)}.$$

The positive quadratic box bounds its prefactor below, and bounds  $\chi/P$  above by a finite  $\Gamma$ . The residual integral has a positive lower bound and the area has an upper bound. The exact phase-volume formula gives (6.3). This argument uses relative twist control before conditioning; division of an absolute action error by  $D_{N,b}$  would not give these estimates.

The law comparison in Theorem 5.5 now reads

$$\|F^T - \tilde{F}^T\|_{C^0([-R,R])} \leq C \max_{b,d} \|f_{b,d}^T - \tilde{f}_{b,d}^T\|_{C^m(Q)} \quad (6.4)$$

when the right-hand density distance is at most a fixed  $\eta_0$ . It applies to these restrictions because its proof uses the anchored quotient and positive denominator floors, not a known normalizing constant. A smaller common  $R$  can be chosen once for the class.

**6.2. Recovering the differentiated density norm.** Use the maximum of all derivatives of total order at most  $m$  as the  $C^m(Q)$  norm. Fix a smooth compactly supported kernel  $K$  on  $(-1, 1)^2$  with

$$\int K = 1, \quad \int z^\alpha K(z) dz = 0 \quad (1 \leq |\alpha| < s).$$

Such kernels need not be nonnegative. For completeness, take a positive smooth bump  $w$  on  $(-1, 1)$ , form the positive definite moment matrix  $M_{ij} = \int t^{i+j} w(t) dt$ ,  $0 \leq i, j < s$ , and solve  $Mc = (1, 0, \dots, 0)^t$ . Then  $k(t) = w(t) \sum_{j=0}^{s-1} c_j t^j$  has the required one-dimensional moments, and  $K(z) = k(z_1)k(z_2)$  has the displayed moments. Positive definiteness follows by integrating the square of a nonzero polynomial against  $w$ . We write  $K_h(z) = h^{-2}K(z/h)$ .

**Lemma 6.1** (Interior derivative estimation). *Let each of  $J$  groups contain  $n \geq 2$  independent latent endpoints  $Y_{g,l}$  with density  $q_g$  on  $Q_+$  and  $\|q_g\|_{C^{m+s}(Q_+)} \leq H$ . Observed endpoints may have arbitrary bounded recording errors  $|\tilde{Y}_{g,l} - Y_{g,l}|_\infty \leq \delta$ . Independence of the recording errors is not required. Set*

$$\hat{q}_{g,h}(z) = \frac{1}{n} \sum_{l=1}^n K_h(z - \tilde{Y}_{g,l}).$$

For  $n^{-1} \leq h \leq h_0 < \min(1, R_+ - R)$  and  $0 < \zeta < 1/2$ , put  $L_n = \log(C_1 J n / \zeta)$ , with a sufficiently large fixed  $C_1$ . With probability at least  $1 - \zeta$ ,

$$\max_g \|\hat{q}_{g,h} - q_g\|_{C^m(Q)} \leq C \left\{ h^s + \sqrt{\frac{L_n}{nh^{2m+2}}} + \frac{L_n}{nh^{m+2}} + n^{-2} + \frac{\delta}{h^{m+3}} \right\}. \quad (6.5)$$

The constants depend on the fixed squares, kernel, orders and  $H$ , not on  $n, h, \delta$  or the particular densities.

*Proof.* First set  $\delta = 0$ . On  $Q$  the kernel does not encounter the boundary of  $Q_+$ . Integration by parts gives  $D^\alpha \mathbb{E} \hat{q}_h = K_h * D^\alpha q$  for  $|\alpha| \leq m$ . Taylor's integral remainder through order  $s-1$ , with the moment cancellations, gives a uniform bias at most  $CHh^s$ . Only derivatives through  $m+s$  on  $Q_+$  are used; the zero extension of the density is not asserted to be smooth at the boundary.

For a fixed  $z$  and  $|\alpha| \leq m$ , the summand  $D^\alpha K_h(z - Y)$  is bounded in absolute value by  $Ch^{-m-2}$  and has variance at most  $CHh^{-2m-2}$ . These bounds follow directly by scaling, since

$$\int |D^\alpha K_h(z - y)|^2 q(y) dy \leq Hh^{-2|\alpha|-2} \|D^\alpha K\|_2^2.$$

For independent centered summands bounded by  $U$  with variance at most  $V$ , Bernstein's bound for their average is

$$\mathbb{P} \left( \left| n^{-1} \sum_{l=1}^n Z_l \right| > \sqrt{2Vt/n} + 2Ut/(3n) \right) \leq 2e^{-t}.$$

One may derive it by expanding the exponential and using  $\mathbb{E}|Z_l|^j \leq VU^{j-2}$  for  $j \geq 2$ :  $\mathbb{E}e^{\lambda Z_l} \leq \exp\{\lambda^2 V/[2(1 - |\lambda|U/3)]\}$  for  $|\lambda|U < 3$ . Multiplication of the moment bounds and exponential Markov inequality give the displayed tail estimate, with harmless enlargement of  $U$  for centering.

Cover  $Q$  by a grid of mesh at most  $n^{-2}h^{m+3}$ . Its size is at most  $Cn^4h^{-2m-6}$ . A union bound over this grid, the fixed set of multiindices and the  $J$  groups, with  $t = CL_n$ , gives the square-root and linear terms of (6.5); here  $h \geq n^{-1}$  absorbs the grid logarithm into  $CL_n$ . Both the empirical field and its expectation have Lipschitz constant at most  $Ch^{-m-3}$  after any derivative of order at most  $m$ . The grid error is consequently at most  $Cn^{-2}$ . This proves the asserted uniform stochastic bound. This elementary fixed-bandwidth argument is sufficient here; sharper classical uniform kernel results are discussed in [52].

Finally the mean value theorem, applied to each smooth kernel summand, gives the deterministic comparison

$$\max_g \|\hat{q}_{g,h}^{\text{observed}} - \hat{q}_{g,h}^{\text{latent}}\|_{C^m(Q)} \leq C\delta h^{-m-3}.$$

Adding it proves the result, even for data-dependent recording errors. The estimates need not be densities or nonnegative functions. No nonlinear law identity has yet been applied to them.  $\square$

**6.3. Realizable estimation and complete profile recovery.** Let  $\mathcal{F} = \{f^T|_Q : T \in \mathcal{C}\}$  in the product  $C^m(Q)^J$ , with maximum norm. Choose a countable observation-dense subset  $\{f^{T_j} : j \geq 1\}$  of this image, retaining one actual table for each element. Such a subset exists because  $C^m(Q)^J$  is separable and metrizable; its subspaces are separable. Given an estimated tuple  $\hat{f}$  and a tolerance  $\xi > 0$ , define

$$\hat{T} = T_{j_*}, \quad j_* = \min \left\{ j : \|f^{T_j} - \hat{f}\|_{C^m(Q)} \leq \inf_{l \geq 1} \|f^{T_l} - \hat{f}\|_{C^m(Q)} + \xi \right\}. \quad (6.6)$$

The set is nonempty by the definition of an infimum. This is a measurable estimator: each norm is a supremum of continuous derivative evaluations over a fixed countable dense set, and the remaining operations are countable. This existence construction is not a claim of efficient enumeration of all physical tables. It supplies no unknown curvature or flux amplitude as observation data.

**Theorem 6.2** (Finite-sample recovery of smooth contacts). *Under the stated class assumptions, collect  $n$  successful endpoints at  $N$  flights in each phase–offset group, from independent preparations. Allow recording error at most  $\delta$ , and use (6.6) with the kernel estimates and  $\xi = n^{-1}$ . Let*

$$\alpha = \frac{s}{2(m+s)+2}, \quad \gamma = \frac{s}{m+s+3}, \quad h = \max \left\{ (L_n/n)^{1/[2(m+s)+2]}, \delta^{1/(m+s+3)} \right\}.$$

For  $h \leq h_0$  and sufficiently small combined error, with probability at least  $1 - \zeta$ ,

$$\|F^{\hat{T}} - F^T\|_{C^0([-R,R])} \leq C \left\{ e^{-\omega N} + (L_n/n)^\alpha + \delta^\gamma \right\}. \quad (6.7)$$

No candidate-closeness assumption, analytic prior, supplied density jets, or inversion on arbitrary density tuples is used.

*Proof.* On the event of Lemma 6.1, denote its right-hand bound by  $e$ . The true limiting tuple is within  $e + C_0 e^{-\omega N}$  of the estimates, by (6.2). Density of the countable image makes this also an upper bound for the infimum in (6.6). The triangle inequality gives

$$\|f^{\hat{T}} - f^T\|_{C^m(Q)} \leq 2e + 2C_0 e^{-\omega N} + n^{-1}.$$

Both sides of this comparison are realizable limiting tuples. Apply (6.4) when this bound is small. In particular no factorization of the empirical or finite-flight density is asserted.

To simplify the rate, write  $h_n = (L_n/n)^{1/[2(m+s)+2]}$ . At  $h = h_n$  the square-root term equals  $h_n^s$ ; the linear term is  $h_n^{m+2s} \leq h_n^s$ . Both stochastic terms decrease when  $h$  increases. Also  $h^s \leq h_n^s + \delta^\gamma$  and  $\delta h^{-m-3} \leq \delta^\gamma$ , with the zero-noise term interpreted as zero. For large enough  $n$ ,  $n^{-1} \leq h_n^s$  and  $h \geq n^{-1}$ . These facts prove (6.7).

For repeated independent preparations, each observation is a failure or a successful mark. Its law is the mixture of a failure atom and  $p_{N,b,d}^T f_{N,b,d}^T$ . It can therefore be generated by an independent Bernoulli success sequence and an independent sequence of marks with that conditional density. The first  $n$  successful marks are independent with the required density. No mixing assumption about a single unreset billiard trajectory is being substituted for independent preparations.  $\square$

#### 6.4. A deterministic budget with failures charged.

**Theorem 6.3** (Charged smooth recovery). *Perform exactly  $B$  independent full-phase preparations in each of the  $J$  groups, all at  $N \geq N_0$  flights. Define*

$$n_B = \lfloor Bp_*(N)/2 \rfloor.$$

*Assume  $n_B \geq 2$  and  $Bp_*(N) \geq 8 \log(J/\zeta)$ . Use the first  $n_B$  successes in every group and the preceding estimator; return any fixed member of  $\mathcal{C}$  if a group has fewer successes. Then (6.7) holds with  $n = n_B$  and probability at least  $1 - 2\zeta$ , subject to the same small-error and bandwidth conditions. Every one of the  $JB$  preparations is charged.*

*In particular, put  $L_B = \log(C_2 JB/\zeta)$  and, for large enough  $B$ , choose the returning flight number*

$$N_B = P \left\lfloor \frac{\alpha \log(B/L_B)}{P(\omega + \alpha\Gamma)} \right\rfloor \geq N_0. \quad (6.8)$$

*There is a constant  $C$  such that*

$$\mathbb{P}_T \left( \|F^{\hat{T}} - F^T\|_{C^0([-R, R])} > C\{(L_B/B)^\beta + \delta^\gamma\} \right) \leq 2\zeta, \quad \beta = \frac{\alpha\omega}{\omega + \alpha\Gamma} > 0. \quad (6.9)$$

*The assertion is uniform over the stipulated marked class. Constants, orders and exponents need not be uniform in a varying period or a vanishing hyperbolicity margin.*

*Proof.* The number of successes in one group is binomial with mean  $\mu = Bp_{N,b,d}^T \geq Bp_*(N)$ . The elementary lower-tail bound  $\mathbb{P}(M < \mu/2) \leq e^{-\mu/8}$  implies

$$\mathbb{P}(M < n_B) \leq e^{-Bp_*(N)/8}.$$

For example exponential Markov inequality at parameter  $\log 2$  gives  $\exp\{-\mu(1 - \log 2)/2\}$ , which is no larger than  $e^{-\mu/8}$ . The union bound gives probability at most  $\zeta$  of any insufficient group. By the mixture representation at the end of the preceding proof, the first  $n_B$  marks have the requisite product law independently of the success indicators. Intersecting the sufficient-success event with the density-estimation event loses at most another  $\zeta$ . Arbitrary bounded recording errors are still handled pathwise.

Write  $N_* = \alpha \log(B/L_B)/(\omega + \alpha\Gamma)$ . Then  $N_B \in (N_* - P, N_*]$ , so

$$e^{-\omega N_B} \leq e^{\omega P} (L_B/B)^\beta, \quad n_B \geq cB^{\omega/(\omega + \alpha\Gamma)} L_B^{\alpha\Gamma/(\omega + \alpha\Gamma)}$$

for sufficiently large  $B$ , with the floor absorbed into  $c > 0$ . The logarithm  $L_{n_B}$  is bounded by  $CL_B$  since  $n_B \leq B$ . Substitution gives  $(L_{n_B}/n_B)^\alpha \leq C(L_B/B)^\beta$ , proving (6.9). The sufficient-success condition follows for fixed confidence and large  $B$ ; for varying confidence it remains an explicit requirement. The bandwidth and small-error restrictions also remain explicit, especially when  $\delta$  does not tend to zero.  $\square$

The factor  $\Gamma$  in (6.9) records the cost of rarity. It would be absent if successful bridges were incorrectly treated as free observations. With a bounded period length, the total physical observation time is at most  $CJB(1 + N_B)$ , hence  $O(JB \log B)$  at fixed confidence. The supplied polygon and frames and the two exact offsets are still required. This is a finite-sample result for bounded endpoint recording errors, not arbitrary contamination, uncalibrated clocks, a single correlated trajectory, or a claim of minimax optimality.

The estimate controls complete profiles and hence applies to the actual flat family of Proposition 5.6, when restricted to a small compact parameter interval. To make this consequence precise, fix  $0 < |u_*| < R$  in the region where its cutoff is one. The two tables at parameters 0 and  $\varepsilon > 0$  have profile separation at least  $\Delta = \varepsilon e^{-1/u_*^2}$ . Whenever  $C\{(L_B/B)^\beta + \delta^\gamma\} < \Delta/3$ , (6.9) and comparison of the estimated profile's distance to these two candidates give a test whose

error at either table is at most  $2\zeta$ . The remote area correction does not change this contact value. Thus flat information is accessible to a finite charged experiment at every fixed nonzero scale, although its cost can be very large as  $u_*$  approaches zero. Equal Taylor jets alone still supply no such test. No assertion about equal full marked length spectra, unvisited smooth arcs, or normalized density jets is introduced by this observation.

## 7. LOCAL DETERMINATION AND EXTERIOR SCATTERING DATA

The endpoint law is constructed from actual finite trajectories. We now make precise its dependence on a neighborhood of the marked itinerary. This also fixes the comparison with the global scattering observation in Section 1.7.

**7.1. The local observation record.** Fix the polygon, its marked frames, a sufficiently small common graph interval  $I$ , and the two offsets of Theorem 2.5. Shrink  $I$  once, if necessary, so that the stationary bridges for all returning  $N \geq N_0$  lie in prescribed contact collars with compact closure. There are finitely many such collars modulo the lattice. Their successive chords, together with the short initial and terminal free segments allowed by the time-split construction, lie in a fixed neighborhood  $U$  of the itinerary modulo the lattice. We require the same clearance and free-segment margins as in that theorem. The neighborhood includes the reflecting collars; away from them its relevant trajectory tubes are obstacle-free.

A *local completion* consists of closed obstacles agreeing with the prescribed contact graphs on these collars and satisfying these margins on  $U$ ; the boundary away from  $U$  may vary. The lattice and marks are fixed. This definition concerns actual tables, not independently assigned Taylor coefficients. For a chosen phase  $b$ , returning length  $N$  and offset  $d$ , a reset preparation is drawn from the normalized Liouville measure of that table. Record the endpoint mark  $(u, v)$  on success of the stipulated itinerary event, and a single symbol  $\dagger$  otherwise. Write  $\mu_{N,b,d}$  for the success subprobability measure on the gate  $Q_+$ ,  $p_{N,b,d} = \mu_{N,b,d}(Q_+)$ , and

$$P_{N,b,d} = (1 - p_{N,b,d})\delta_{\dagger} + \mu_{N,b,d}. \quad (7.1)$$

The conditional endpoint density is that of  $\mu_{N,b,d}/p_{N,b,d}$ . The symbol  $\dagger$  contains no trajectory trace or additional data from failed preparations. Both the resetting and this recording convention are essential to the statement below.

**Proposition 7.1** (Exact locality of the charged record). *Let  $\Omega$  and  $\tilde{\Omega}$  be two local completions with the same contact graphs on the prescribed collars, and free areas  $\mathcal{A}$  and  $\tilde{\mathcal{A}}$  in the common cell. Then, for every phase, both selected offsets, and every returning  $N \geq N_0$ ,*

$$\tilde{\mu}_{N,b,d} = \frac{\mathcal{A}}{\tilde{\mathcal{A}}} \mu_{N,b,d}. \quad (7.2)$$

*In particular the conditional finite-flight densities and their limiting laws coincide. If  $\mathcal{A} = \tilde{\mathcal{A}}$ , the complete records (7.1) coincide, and so do all finite products of these independently reset experiments. The assertion is exact, not an asymptotic estimate as  $N \rightarrow \infty$ .*

*Proof.* The single-edge length functions between successive collars are the same in the two tables. The stationary equations and their boundary conditions are consequently identical. The weighted uniqueness in Theorem 2.3, on the chosen common collars, identifies their finite stationary bridges. They remain physical in each table by the common clearance margins. Thus their stationary lengths  $W_{N,b}$ , gauged actions  $E_{N,b}$ , mixed twists and reference normalizations  $D_{N,b}$  agree as functions of the endpoints. Equality holds throughout the endpoint gate, not just to finite order at its origin.

Put  $\beta_{N,b} = (-\partial_{uv}W_{N,b})/D_{N,b}$  and  $r_{N,b,d} = d - E_{N,b} + p_b u - p_b v$ . The exact time-split computation of Theorem 2.5 gives

$$\frac{d\mu_{N,b,d}}{du dv} = \frac{D_{N,b}}{2\pi\mathcal{A}}\beta_{N,b}(u,v)\{r_{N,b,d}(u,v)\}_+. \quad (7.3)$$

Indeed the allowed initial free time ranges from 0 to  $r_{N,b,d}$ ; the free-segment margins ensure that neither an earlier nor an additional reflection truncates this interval. Every successful trajectory is accounted for by this parametrization because the event specifies all intermediate collisions in the collars and the last-collision rule. It does not include other itineraries elsewhere in the table. All factors in (7.3) except the total free area are common to the two completions. This proves (7.2). Integrating gives the same factor for  $p_{N,b,d}$ , so normalization cancels it. The positive success floor ensures that the conditional densities are defined. Their limits agree by the relative law. At equal area the subprobabilities and hence the failure masses agree; independence gives the final product assertion.  $\square$

*Remark 7.2.* If settings are selected from the permitted experiments by the same measurable rule based on earlier recorded outcomes and independent auxiliary randomization, equality also holds for every fixed finite number of preparations. Conditional on the history the next setting has the same distribution and its record kernel is the same; induction proves the assertion. This does not extend it to an unreset trajectory, to recording a rejected ray's path, or to a table-dependent error kernel. A common recording kernel applied to the exact records preserves equality.

**7.2. Complete contact information without remote completion.** Two marked tables will be called *contact equivalent* when their graph functions agree on some common neighborhood of every contact. Use a germ of the local law experiment: endpoints may be restricted to a smaller common positive-residual gate, with its normalization specified anew. This convention is necessary when only contact germs, rather than fixed-radius graphs, are given. In particular, restricting a density without renormalizing its original larger gate is a different operation.

**Corollary 7.3** (Local completion and the determined object). *On the marked smooth class of Theorem 1.1, equality of the two-offset limiting law germs is equivalent to contact equivalence. Here equality of laws is tested on a sufficiently small common gate at every phase. In the equal-area subclass, contact equivalence also gives equality of all charged finite-flight records on sufficiently small common gates.*

*There are nonconstant smooth one-parameter families of closed strictly convex obstacles in a fixed periodic cell, with the same free area and marked polygon, for which these charged records are identical although the obstacle boundaries away from the collars differ. These families can be chosen so that their finite compact obstacle unions in the plane have different almost-everywhere exterior travelling-time data and different almost-everywhere scattering length spectra.*

*Proof.* Equality of the limiting laws gives equality of the anchored action functions by Proposition 3.2. The global quadratic inverse, signed finite-jet inversion, and actual smooth inverse of Theorem 5.3 then identify the graph functions on a common smaller collar. These steps use actual realizable law tuples. Conversely, if the contact germs agree, choose compact collars inside their common domains. The weighted finite and half-line constructions permit a sufficiently small common endpoint gate for which all stationary visits stay inside these collars. Both tables have positive clearance along the fixed polygon, so the short endpoint segments also have common margins. The two offsets are chosen below these common free-flight bounds, as required by the experiment. Proposition 7.1 applies. It proves the limiting assertion, and the finite-record assertion when the areas agree. This is an equivalence of germs of experiments; it asserts no equality on a larger gate outside the common graph domains.



For an actual family take the three-obstacle realization of Proposition 2.6. Choose  $U$  sufficiently small around its contact collars and flight tubes. On an obstacle boundary choose two disjoint open arcs with closures outside  $U$ . Such arcs exist on the side facing away from the interior contact polygon in this explicit configuration. Let  $\psi_1, \psi_2$  be nonzero smooth nonnegative normal bumps supported in these arcs. For small real parameters  $(s, t)$  deform the boundary normally by  $s\psi_1 + t\psi_2$ , leaving all other boundary points unchanged. Disjoint compact embedded boundaries stay disjoint and strictly convex under sufficiently small  $C^2$  perturbations; the lattice translates remain separated as well. The positive distance between the bump supports and the closed trajectory tubes preserves all the required local margins.

Let  $V(s, t)$  be the total obstacle area in the cell. Normal variation of area gives

$$\partial_t V(0, 0) = \int \psi_2 d\sigma > 0.$$

For completeness this derivative follows by using boundary normal coordinates: a strip of signed width  $t\psi_2(\sigma)$  has area  $t \int \psi_2 d\sigma + O(t^2)$ , with a uniform remainder on the compact support. The area is a smooth function of the two parameters. The implicit function theorem gives a smooth  $t = t(s)$  near zero, with  $t(0) = 0$  and  $V(s, t(s)) = V(0, 0)$ . The family is nonconstant, since on the first support its normal displacement is exactly  $s\psi_1$  and the second bump has disjoint support. All contact graphs and marks are unchanged. The free area is constant, so Proposition 7.1 proves equality of every indicated finite record. The observation samples are still drawn from the full table, not from a law postulated only on  $U$ .

Regard the same three obstacles, without their lattice copies, as a compact union  $K_s$  in  $\mathbb{R}^2$ . All  $K_s$  lie in one fixed enclosing circle and satisfy the strictly convex  $C^3$  hypotheses of Noakes–Stoyanov [53, Theorem 1.1]. For  $s \neq s'$  sufficiently small,  $K_s \neq K_{s'}$  by the disjoint-support construction. That theorem implies that their travelling-time families are not equal on a full-measure set of incoming positions and directions. Its separate scattering-length conclusion gives the analogous distinction for incoming/outgoing direction pairs. The periodic records remain identical by the preceding paragraph. This proves the asserted comparison on an explicitly common class of geometric realizations.  $\square$

The two parts of this comparison serve different purposes. In the local experiment, the inverse determines the actual visited contact germs, including flat changes, whereas the construction varies only unvisited boundary pieces and leaves even the finite charged record fixed. The exterior observation distinguishes those pieces by a different data family. Hence equality of the local record cannot simply be substituted for the equality hypothesis in a global lens theorem. There is no claim that local laws distinguish equal full marked length spectra, nor that a known whole obstacle and the supplied marks could not generate its local laws. The comparison sharpens the local determination statement without weakening it and without identifying the two observation models.

## 8. ONE UNCALIBRATED LAW AT AN OBLIQUE CONTACT

The two-offset observation separates unequal actions without using their first derivatives. At an oblique contact those derivatives are already fixed by the geometric marks. They remove both the multiplicative ambiguity of a rank-one factorization and the unknown time scale. This gives a second observation theorem for the same smooth inverse, without changing the two-offset theorem at normal incidence.

**8.1. An anchored factorization and its scale.** Let  $I = [-R, R]$ . Consider a positive  $C^M$  function on a neighborhood of  $I^2$ ,  $M \geq 2$ , of the form

$$f(u, v) = Z^{-1} B_-(u) B_+(v) \{d - A(u) - C(v)\}, \quad (8.1)$$

where  $d > 0$ ,  $Z > 0$ ,  $B_{\pm} > 0$ ,  $B_{\pm}(0) = 1$ , and

$$A(0) = C(0) = 0, \quad A'(0) = -p, \quad C'(0) = p, \quad p \neq 0. \quad (8.2)$$

The signed number  $p$  is known;  $d$ , both amplitudes and  $Z$  are unknown. Assume  $d - A - C > 0$  on  $I^2$ . No equality of  $A$  and  $C$  is imposed. Multiplication of  $f$  by an unknown positive constant is permitted. Define

$$\mathcal{R}_f(u, v) = \frac{f(u, v)f(0, 0)}{f(u, 0)f(0, v)}, \quad \kappa_f = \partial_u \partial_v \mathcal{R}_f(0, 0). \quad (8.3)$$

For a continuous positive density, its law determines its values on the square, and thus these expressions whenever the indicated derivatives exist. Finite samples are treated below, not identified with pointwise density observations.

**Proposition 8.1** (Identification of the unknown offset). *For (8.1)–(8.2),*

$$\mathcal{R}_f = 1 - UV, \quad U = \frac{A}{d - A}, \quad V = \frac{C}{d - C}, \quad \kappa_f = \frac{p^2}{d^2} > 0, \quad d = \frac{|p|}{\sqrt{\kappa_f}}. \quad (8.4)$$

*Consequently the offset is determined by the single conditional law and the marked momentum, without a count or an amplitude model.*

*Proof.* The axis residuals  $d - A(u)$  and  $d - C(v)$  are positive. Cancellation of  $Z$  and the two amplitudes gives

$$\mathcal{R}_f(u, v) = \frac{d\{d - A(u) - C(v)\}}{(d - A(u))(d - C(v))} = 1 - \frac{A(u)C(v)}{(d - A(u))(d - C(v))}.$$

Since  $U(0) = V(0) = 0$ ,  $U'(0) = -p/d$  and  $V'(0) = p/d$ ,  $\kappa_f = -U'(0)V'(0) = p^2/d^2$ . The positive choice of  $d$  fixes the square-root branch. Equivalently  $\kappa_f = \partial_{uv} \log f(0, 0)$ , because  $\mathcal{R}_f$  equals one on both axes. This identity is insensitive to the unknown normalization or separate endpoint factors.  $\square$

*The clock before the relative limit.* For the actual finite-flight density write  $f_N = Z_N^{-1} \beta_N r_N$ , where  $r_N = d - E_N + pu - pv$ . The exact finite-flight counterpart of (8.4) is

$$\partial_{uv} \log f_N(0, 0) = \frac{p^2}{d^2} + \frac{D_N}{d} + \partial_{uv} \log \beta_N(0, 0). \quad (8.5)$$

Indeed  $r_N(0, 0) = d$ ,  $r_{N,u}(0, 0) = p$ ,  $r_{N,v}(0, 0) = -p$  and  $r_{N,uv}(0, 0) = D_N$ . Differentiating  $\log r_N$  proves the formula. The normalization is constant, and the remaining term is the indicated derivative of the amplitude. This identity was pointed out in [54]. The mixed derivative of the anchored finite-density cross ratio has the same value, because that cross ratio equals one on both axes.

The last term in (8.5) has no asserted sign. The relative  $C^2$  theorem and positive floors show that its sum with  $D_N/d$  is  $O(e^{-\omega N})$  on the bounded class; one may decrease  $\omega > 0$  to a common strict exponential margin. In particular, if  $\kappa_* = p_*^2/d_+^2$  and the error is at most  $\kappa_*/2$ , the finite mixed derivative is positive and the direct plug-in clock differs from  $d$  by at most  $Ce^{-\omega N}$ . This follows from the mean value bound for  $x \mapsto |p|x^{-1/2}$  on  $[\kappa_*/2, \infty)$ . It is a biased finite-flight approximation, not the exact identity (8.4). The forward term in (8.12) includes this bias.

The rank-one identity alone would determine  $U, V$  only up to  $(U, V) \mapsto (\lambda U, \lambda^{-1} V)$ . More precisely, if neither factor is identically zero and two pairs give the same product, evaluate at a point where the second factor is nonzero to identify the first factor up to a scalar, and then identify the second. The specified first derivatives in (8.2), after (8.4) has fixed  $d$ , force  $\lambda = 1$ . The following formulas implement that argument without differentiating a whole recovered function one extra time.

Choose one fixed  $a \in I \setminus \{0\}$  sufficiently close to zero that  $A(a)C(a) \neq 0$ . Such an  $a$  exists by (8.2). Using the recovered  $d$ , put

$$v_a(f) = \frac{d}{p} \partial_u \mathcal{R}_f(0, a), \quad u_a(f) = -\frac{d}{p} \partial_v \mathcal{R}_f(a, 0), \quad (8.6)$$

$$U_f(u) = \frac{1 - \mathcal{R}_f(u, a)}{v_a(f)}, \quad V_f(v) = \frac{1 - \mathcal{R}_f(a, v)}{u_a(f)}. \quad (8.7)$$

The quantities  $u_a, v_a$  are signed scalars. They are not positive square roots and must not be replaced by their absolute values.

**Theorem 8.2** (One-law separation with an unreported offset). *The single function  $f$  in (8.1) and the known  $p \neq 0$  determine  $d, A, C, B_-, B_+$  uniquely on  $I$ , with*

$$A_f(u) = \frac{dU_f(u)}{1 + U_f(u)}, \quad C_f(v) = \frac{dV_f(v)}{1 + V_f(v)}, \quad (8.8)$$

$$B_{-,f}(u) = \frac{f(u, 0)}{f(0, 0)}(1 + U_f(u)), \quad B_{+,f}(v) = \frac{f(0, v)}{f(0, 0)}(1 + V_f(v)). \quad (8.9)$$

For  $M \geq 2$  this inverse is locally Lipschitz from  $C^M(I^2)$  to  $\mathbb{R} \times C^M(I)^4$  on classes with uniform bounds for the  $C^M$  norms and positive lower bounds for  $f, \kappa_f, |u_a(f)|, |v_a(f)|, 1 + U_f$ , and  $1 + V_f$ . The formulas extend to an open neighborhood of each such bounded class, even when a perturbed function does not have representation (8.1).

*Proof.* Differentiating  $1 - UV$  at  $(0, a)$  and  $(a, 0)$  gives  $v_a(f) = V(a)$  and  $u_a(f) = U(a)$ . Thus (8.7) recovers the actual factors on all of  $I$ , including their zeros. Solving the definitions of  $U, V$  gives (8.8). Both final denominators are positive, since  $1 + U = d/(d - A)$  and  $1 + V = d/(d - C)$ . The two axis identities in (8.1) give (8.9). This also proves uniqueness: any other representation satisfying the same slopes gives the same offset, factors, actions and normalized amplitudes.

For stability, restriction to a fixed slice is bounded in  $C^M$ . Products and reciprocal functions are locally Lipschitz on bounded  $C^M$  sets with positive denominator floors. Apply these facts first to  $\mathcal{R}_f$ . Evaluation of its mixed derivative at the single point  $(0, 0)$  is a bounded scalar functional when  $M \geq 2$ . The function  $x \mapsto |p|x^{-1/2}$  is Lipschitz on an interval with a positive lower endpoint, so  $d$  is locally Lipschitz. The two first-derivative point evaluations in (8.6) are bounded scalar functionals as well. The stipulated nonzero anchor bounds permit division in (8.7); fixed slices give  $C^M$  functions, with no loss of one derivative. Finally apply the product and reciprocal bounds to (8.8)–(8.9). For small enough perturbations all floors persist, possibly halved. No assertion that the extended inverse produces an actual billiard is made.  $\square$

Uniform anchors follow from geometric priors rather than an assumed conditioning bound on a measured inverse. Suppose  $|p| \geq p_* > 0$ ,  $d_- \leq d \leq d_+$ , and the  $C^2$  norms of  $A, C$  are bounded by  $H_2$ . Choose a fixed  $a > 0$  with  $H_2 a \leq p_*$ , inside the common collar. Taylor's formula gives  $|A(a)|, |C(a)| \geq p_* a/2$ . If the positive axis residuals are bounded above by  $D_*$ , then

$$|u_a|, |v_a| \geq \frac{p_* a}{2D_*}, \quad \kappa_f \geq \frac{p_*^2}{d_+^2}, \quad 1 + U, 1 + V \geq \frac{d_-}{D_*}.$$

The physical relative law gives uniform amplitude and normalized-density bounds on a common positive square. These observations supply all the floors used in Theorem 8.2. Constants need not remain bounded as  $p_*$  or  $a$  tends to zero.

**8.2. Smooth contact determination.** We return to the fixed marked polygon of Section 3. Call it *oblique* if  $p_b \neq 0$  at every physical phase. At each phase select one positive offset  $d_b$ , fixed over the returning lengths and the independent preparations, but not included in the reconstruction data. The offsets may differ between phases. They lie in a stipulated positive physical window; the contact gate and all event tags retain their exact meanings. Thus an unreported offset is a fixed nuisance parameter, not a random offset, an incorrectly recorded success tag, or an uncertain endpoint gate. No new clock-dependent observation is being supplied.

**Theorem 8.3** (Smooth rigidity from one uncalibrated law per phase). *For an oblique clear nongrazing marked polygon with distinct physical contacts and positive-curvature  $C^\infty$  graphs, the collection*

$$(f_{b,d_b} : b = 0, \dots, r-1)$$

*of single-offset limiting endpoint laws determines all the offsets  $d_b$  and the complete smooth reflecting contact germs. Neither the offsets, curvature, flux amplitudes, nor acceptance probabilities are supplied. No symmetry, analyticity or closeness between the candidates is required.*

*On a bounded positive marked class as in Theorem 5.5, choose  $m \geq 3$  so that*

$$\frac{6a^m}{1-a^m} < 1$$

*on the enlarged positive quadratic class. With uniform obliquity, offsets in  $[d_-, d_+]$ , and common positive-residual margins, there are  $R, C, \eta_0 > 0$  such that*

$$\|F - G\|_{C^0([-R, R])} + \|d - \tilde{d}\|_{\ell^\infty} \leq C \max_b \|f_{b,d_b}^F - f_{b,d_b}^G\|_{C^m([-R, R]^2)} \quad (8.10)$$

*when the maximum on the right is at most  $\eta_0$ . These are realizable law pairs; their normalizing constants need not coincide.*

*Proof.* Lemma 3.1 supplies the actual physical law (8.1), with  $A_b = S_b^- - p_b u$  and  $C_b = S_b^+ + p_b v$ . The half-line actions are anchored, so their physical slopes are exactly (8.2). Theorem 8.2 therefore recovers  $d_b$  and both actions at every phase. Add the known  $p_b u$  to  $A_b$  to obtain the future action map  $\mathcal{S}(F)$ . The action version of Theorem 5.3 identifies the actual contact germs. In particular the last step still uses the contracted-visit envelope on flat differences, not just equality of Taylor series.

On the bounded class choose the common collar and fixed anchor using the preceding uniform bounds. The inverse in Theorem 8.2 gives  $C^m$  control of both actions and scalar control of the offsets from the displayed law distance. Apply the action estimate (5.10); its polynomial alignment has already been proved before the weighted norm is used. This gives (8.10). There is no extra derivative loss from clock recovery, since  $m \geq 3$  and the point functional in (8.4) has order only two.  $\square$

**Lemma 8.4** (Normal incidence and distinct physical contacts). *A clear nongrazing periodic billiard polygon with  $r \geq 3$  distinct physical contacts is oblique. If a periodic orbit with distinct physical contacts has a normal reflection, it is the ordinary self-retracing two-contact orbit, with zero closing translation.*

*Proof.* At a normal reflection the outgoing ray is the incoming ray traversed backwards. Its first collision is therefore exactly the preceding lifted contact: the incoming open flight was clear. Consequently the preceding and following physical contacts coincide. For  $r \geq 3$  these are different cyclic positions, contradicting distinctness. The only remaining case is  $r = 2$ . Here the equality of the preceding and following lifted, not merely physical, contact forces zero closing translation. The two flights traverse the same segment in opposite directions; periodicity forces normal reflection at the other endpoint as well. This proves the assertion without a genericity assumption.  $\square$

Thus the one-law theorem applies to every polygon with at least three distinct physical contacts in the main theorem, not just a generic subset of those polygons. At normal incidence  $\kappa_f = 0$ , and the new clock formula is not defined. This is a boundary of this argument, not an impossibility claim about all observations at a normal contact. The general two-calibrated-offset theorem is unchanged. For a self-retracing channel at one *known* offset, the separate symmetric factorization of Theorem 23.1 also remains available. It uses a nonzero action anchor instead of a marked nonzero first derivative and does not identify an unknown offset by (8.4).

The equal-area flat family of Proposition 5.6 can be based on the actual three-contact configuration of Proposition 2.6. Its two distinct tables cannot have equal single-offset limiting laws at every phase, even if the nuisance offsets are allowed to differ: equality would first force equal offsets and then equal smooth germs by Theorem 8.3. Flat contact information therefore survives the removal of the second offset and its calibration. Conversely, identical contact collars with identical offsets still give the exact local-completion invariance of Proposition 7.1. The new theorem changes the local observation, not its target to unvisited boundary arcs.

**8.3. Finite preparations with joint offset recovery.** Let  $\mathcal{U}$  be a nonempty class of pairs  $(T, d)$  consisting of an actual marked table and one fixed offset per phase. Impose the bounded positive contact and forward derivative hypotheses of Section 6, replacing the two known offsets by  $d_b \in [d_-, d_+]$  and adding the obliquity floor  $|p_b| \geq p_*$ . All positive-residual, clearance and free-segment margins hold uniformly over that interval. The finite higher graph bounds are those needed for order  $m + s$ , where  $m$  is the same weighted inverse order and  $s \geq 1$ . The free area has a fixed upper bound.

There are now  $J = r$  groups. A preparation in group  $b$  uses the actual fixed  $d_b$ ; it records the same exact itinerary and  $Q_+$  acceptance tag as before, followed on success by an endpoint pair with bounded post-acceptance error  $\delta$ . Densities normalized on  $Q_+$  are restricted to  $Q \in Q_+$  without renormalization. No additional clock reading, calibration trial or endpoint height is used. Preparations are independent resets, not successive pieces of one orbit.

**Theorem 8.5** (Charged joint recovery of profiles and offsets). *On  $\mathcal{U}$  there is a measurable estimator  $(\hat{T}, \hat{d}) \in \mathcal{U}$  with the following property. Let  $B$  preparations be made in each of the  $J = r$  groups, and set*

$$\alpha = \frac{s}{2(m+s)+2}, \quad \beta = \frac{\alpha\omega}{\omega + \alpha\Gamma}, \quad \gamma = \frac{s}{m+s+3}, \quad L_B = \log(C_*JB/\zeta),$$

where  $\omega, \Gamma, c_0, N_0$  are uniform forward constants. Choose

$$N_B = P \left\lfloor \frac{\alpha \log(B/L_B)}{P(\omega + \alpha\Gamma)} \right\rfloor, \quad n_B = \lfloor Bc_0e^{-\Gamma N_B}/2 \rfloor.$$

For  $0 < \zeta < 1/2$ , assume  $N_B \geq N_0$ ,  $n_B \geq 2$  and  $Bc_0e^{-\Gamma N_B} \geq 8 \log(J/\zeta)$ . Impose also the upper bandwidth and small-error conditions in the proof below. Then

$$\sup_{(T,d) \in \mathcal{U}} \mathbb{P}_{T,d} \left( \|F^{\hat{T}} - F^T\|_{C^0([-R,R])} + \|\hat{d} - d\|_{\ell^\infty} > C\{(L_B/B)^\beta + \delta^\gamma\} \right) \leq 2\zeta. \quad (8.11)$$

All  $rB$  preparations, including failures, are charged. The estimator is an existence construction on the realizable image and the rate is a sufficient upper bound, not an efficiency or minimax assertion.

*Proof.* First check the forward hypotheses with the nuisance parameter included. The offset enters the physical numerator affinely. Its compact range, the positive residual integral and the differentiated relative theorem therefore give, uniformly in  $(T, d)$ ,

$$\|f_{N,b,d_b}\|_{C^{m+s}(Q_+)} + \|f_{b,d_b}\|_{C^{m+s}(Q_+)} \leq H, \quad \|f_{N,b,d_b} - f_{b,d_b}\|_{C^m(Q_+)} \leq C_0e^{-\omega N}.$$

The same exact reference twist and area upper bound as in (6.3) give success probability at least  $c_0 e^{-\Gamma N}$ . Knowledge of the actual value of  $d_b$  is not used by any of these uniform bounds. The limiting densities also have a common positive floor, so Theorem 8.3 applies.

For  $n$  accepted endpoints in each group, use the kernel estimator of Lemma 6.1. With  $L_n = \log(C_1 J n / \zeta)$  and  $n^{-1} \leq h \leq h_0$ , its error relative to the limiting density tuple, in the product  $C^m(Q)$  norm, is at most

$$E_{n,N,h} = C \left\{ e^{-\omega N} + h^s + \sqrt{\frac{L_n}{n h^{2m+2}}} + \frac{L_n}{n h^{m+2}} + n^{-2} + \delta h^{-m-3} \right\} \quad (8.12)$$

with probability at least  $1 - \zeta$ . This lemma allows bounded adversarial readout errors after acceptance; it does not change event tags or replace the fixed offsets by varying offsets.

For clarity, include the nuisance parameter in the realizability step. The subset  $\{(f_{b,d_b}^T|_Q)_b : (T, d) \in \mathcal{U}\}$  of the finite product  $C^m(Q)^J$  is separable. Choose a countable dense subset of that image and one representative  $(T_j, d^j) \in \mathcal{U}$  for each of its points. The infimum of the distances from the fitted tuple to these points is a measurable function. Select the first point whose distance is within  $n^{-2}$  of that infimum. Such a point exists, the selection is measurable, and its distance from the true law tuple is at most  $2E_{n,N,h} + n^{-2}$ . Apply (8.10) only to this selected realizable pair. Require  $2E_{n,N,h} + n^{-2} \leq \eta_0$ ; this is the small-error condition. The output includes the representative's actual offset tuple, not an unconstrained clock estimate attached to an unrelated table.

Take

$$h = \max\{(L_n/n)^{1/(2(m+s)+2)}, \delta^{1/(m+s+3)}, n^{-1}\} \leq h_0.$$

This is the upper bandwidth condition. Substitution into (8.12), with constants enlarged in the small-error regime, gives  $C\{e^{-\omega N} + (L_n/n)^\alpha + \delta^\gamma\}$ . The first  $n_B$  successful marks are independent with their conditional densities: independent reset records admit the Bernoulli/conditional-mark representation used in Section 6. The binomial lower-tail bound and a union bound show that all groups have  $n_B$  successes except on an event of probability at most  $\zeta$ . If a group fails this requirement, choose one fixed representative of  $\mathcal{U}$ ; otherwise apply the preceding selection to the first  $n_B$  successes. This defines an estimator for every record.

Finally the specified returning length satisfies

$$e^{-\omega N_B} \leq C(L_B/B)^\beta, \quad n_B \geq cB^{\omega/(\omega+\alpha\Gamma)} L_B^{\alpha\Gamma/(\omega+\alpha\Gamma)}.$$

The floor changes constants only, under the stipulated lower-count condition. Since  $n_B \leq B$  and  $L_{n_B} \leq CL_B$ , one has  $(L_{n_B}/n_B)^\alpha \leq C(L_B/B)^\beta$ . Adding the two exceptional-event probabilities proves (8.11). At fixed confidence all count conditions hold for large  $B$ ; the bandwidth and small-error conditions still require sufficiently small readout error. They are not suppressed for varying confidence.  $\square$

The two mathematical interfaces of the original theorem remain essential. Without a differentiated *relative* physical limit, (8.4) could not be estimated from rare finite bridges. Without the actual envelope inverse, recovery of the action functions would not identify flat changes of a smooth contact. The new algebra removes a second offset and its supplied numerical calibration; it is not presented as a replacement for those analytic arguments. The marks, exact tags and gates, finite-smoothness priors and independent resetting remain part of the observation problem.

## 9. THE COMPLETE INVARIANT OF THE CHARGED LOCAL EXPERIMENT

The conditional laws determine the reflecting germs and the offsets, but conditioning removes the normalization of phase volume. The charged experiment retains it. We identify the exact

information added by the success probabilities and prove a finite-preparation estimate for the resulting invariant. In particular the reconstructed area will belong to the same actual table as the reconstructed contacts and offsets.

Throughout this section the marked lattice and polygon are fixed, all physical contacts are distinct, and  $p_b \neq 0$ . Offsets  $d_b > 0$  are fixed nuisance parameters as in Section 8. Put  $\vartheta = (T, d)$  and write  $\mathcal{A}(T) > 0$  for the free area of its marked cell. A failure records only  $\dagger$ . Every preparation is an independent draw from normalized Liouville measure. These conventions are unchanged when a success probability is used as inverse data.

**9.1. An exact anchor for phase volume.** For a returning  $N = nP$ , denote the success probability at phase  $b$  by  $\pi_{N,b}^\vartheta$  and its conditional endpoint density by  $f_{N,b}^\vartheta$ . The density is normalized on the actual success gate  $Q_+$ , not on a later interior restriction. The superscript will be suppressed when there is no ambiguity. The reference twist  $D_{N,b}$  is computed in the returning signed coordinates of Lemma 3.1.

**Proposition 9.1** (The phase-volume anchor). *For every permitted positive-residual gate containing the origin,*

$$\pi_{N,b} f_{N,b}(0,0) = \frac{D_{N,b} d_b}{2\pi \mathcal{A}(T)}, \quad \mathcal{A}(T) = \frac{D_{N,b} d_b}{2\pi \pi_{N,b} f_{N,b}(0,0)}. \quad (9.1)$$

*The identities hold at finite  $N$ , without a limiting approximation. Once the limiting conditional laws have determined the contact germs and offsets, the right side contains no supplied curvature or amplitude. It uses the success probability and central conditional density of one finite experiment.*

*Alternatively, at any one fixed phase the additional scalar*

$$\ell_b := \lim_{n \rightarrow \infty} \frac{\pi_{nP,b}}{D_{nP,b}} = \frac{d_b}{2\pi \mathcal{A}(T) f_b(0,0)} \quad (9.2)$$

*exists and determines the area together with the limiting laws. Here  $D_{nP,b}$  is reconstructed from the contact curvatures, not supplied as additional data.*

*Proof.* The success subdensity in (7.3) is

$$\frac{D_{N,b}}{2\pi \mathcal{A}(T)} \beta_{N,b}(u,v) r_{N,b,d_b}(u,v).$$

At the origin  $\beta_{N,b}(0,0) = 1$  exactly, by the definition of the reference twist. The gauged action is zero there, so  $r_{N,b,d_b}(0,0) = d_b$ . The subdensity is also  $\pi_{N,b} f_{N,b}$ , proving (9.1). The relative law gives  $f_{nP,b}(0,0) \rightarrow f_b(0,0) > 0$ ; division in the first identity proves (9.2). Theorem 8.3 determines  $d_b$  and the actual contact germs. Their curvatures and the marked flight lengths determine the Jacobi matrices and hence  $D_{N,b}$  by (2.4). No unknown normalization survives.  $\square$

There is an equivalent raw exponential normalization. With  $\chi$  and  $\mathcal{M}_b$  reconstructed from those same curvatures,

$$\lim_{n \rightarrow \infty} e^{n\chi} \pi_{nP,b} = \frac{d_b \sinh \chi}{\pi \mathcal{A}(T) (\mathcal{M}_b)_{12} f_b(0,0)}. \quad (9.3)$$

This follows because  $e^{n\chi} D_{nP,b} \rightarrow 2 \sinh \chi / (\mathcal{M}_b)_{12}$ . The prefactor, not just the exponential decay exponent, is needed for area recovery. Neither (9.2) nor (9.3) is an observation of an infinite limit from a finite sample. A charged finite-sample construction is given below.

**9.2. The fibers of the observation map.** We use germs of the local experiment in the precise sense of Section 7. Two experiments agree as germs when they agree on all sufficiently small common positive gates and all sufficiently long returning lengths, with the gate normalization specified anew. Equivalently one may restrict to a cofinal family of such gates. This qualification does not assert that equality of contact germs identifies trajectories whose endpoints lie outside their common representatives. Write

$$\mathfrak{J}(\vartheta) = ((F_b^T)_b \text{ as germs}, (d_b)_b, \mathcal{A}(T)).$$

**Theorem 9.2** (Complete local observation invariant). *For actual marked tables in the stated class, equality of the germs of all charged reset records is equivalent to equality of  $\mathfrak{J}(\vartheta)$ . Equality of the single-offset limiting conditional law germs alone is equivalent to equality of the first two components of this invariant.*

*In addition to the limiting law germs, the scalar (9.2) at one phase and one fixed common gate suffices to determine the remaining component. The finite alternative (9.1) also suffices. No remote boundary profile is included in this conclusion.*

*Proof.* Equality of the charged records gives equality of their success masses and, since these are positive, equality of their conditional densities. Taking the relative limit on each common gate and applying Theorem 8.3 identifies the offsets and all contact germs. In particular both candidates have the same quadratic Jacobi operator, and therefore the same  $D_{N,b}$ . Their common success subdensity at the origin in any one finite experiment identifies their areas by (9.1).

Conversely, suppose the three components of  $\mathfrak{J}$  agree. Choose representatives of the common contact germs and shrink the gates so that all stationary bridges and allowed end segments stay in these collars and their clear trajectory tubes. The common margins and weighted uniqueness used in Proposition 7.1 apply. That proposition's proof uses only equality of the offset in each experiment, not that the offsets were reported to the observer. It therefore gives identical success submeasures at every phase and returning length. Equal areas remove its sole multiplicative factor. The failure masses then agree as well. This proves the charged-record assertion.

For limiting conditional laws, the forward implication is exactly Theorem 8.3. The converse follows from the same locality calculation: the area factor cancels upon conditioning, whether or not the areas agree. Finally Proposition 9.1 proves both assertions about one additional scalar record.  $\square$

The area component is independent of the conditional one. A smooth bump on an unvisited boundary arc, disjoint from the fixed collars and trajectory tubes, can have nonzero first area variation. For sufficiently small parameter it preserves strict convexity, clearance and all marks, while changing free area. The conditional experiments remain identical, but their charged success masses differ by the exact area ratio. With the second compensating bump of Corollary 7.3, the area can instead be held exactly fixed. That nonconstant family has the same complete charged-record germ. Thus the theorem describes both the recovered global scalar and the precise surviving freedom in remote completion; it does not infer a whole smooth table from local data.

The equivalence also holds for every fixed finite adaptive composition whose settings are chosen from these local experiments by a common record-based rule and independent randomization. Equality of the one-step record kernels proves this by conditioning on the past. The deterministic one-step experiments are among these compositions, so the converse is unchanged. This statement is about equality of experiment families, not about a sufficient statistic for one fixed finite sample.

**9.3. Sensitivity of the rare normalization.** Area reconstruction requires relative control of the reference twist. An error in curvature cannot be propagated through  $N$  flights with a constant independent of  $N$ . The following estimate retains that cost. Write  $c_i = F_i''(0)$ , fix



$k_i = L_i^{-1}$ , and use the sign-conjugated transfer matrices with masses  $2c_i$ . All derivatives below are with respect to the finite vector  $c$ .

**Lemma 9.3** (Logarithmic twist comparison). *On a fixed compact positive curvature box there is  $C$  such that, for all its  $c, \tilde{c}$ , phases  $b$ , and  $N = nP$ ,*

$$\left| \log \frac{D_{N,b}(\tilde{c})}{D_{N,b}(c)} \right| \leq C(1+N)\|\tilde{c} - c\|_\infty. \quad (9.4)$$

More precisely,

$$\nabla \log D_{N,b} = \{\coth \chi - n \coth(n\chi)\} \nabla \chi - \nabla \log(\mathcal{M}_b)_{12}. \quad (9.5)$$

Along a smooth curvature path, if  $\dot{\chi} \neq 0$ , then  $n^{-1} \partial_t \log D_{nP,b} \rightarrow -\dot{\chi}$ . Thus the flight-length factor cannot in general be deleted from this normalization estimate. This last assertion is not a minimax lower bound.

*Proof.* Every matrix

$$Q_i(c_i) = \begin{pmatrix} 1 + 2c_i/k_i & 1/k_i \\ 2c_i & 1 \end{pmatrix}$$

has determinant one and positive entries. On the compact positive box,  $\chi(c) = \operatorname{arcosh}(\operatorname{tr} \mathcal{M}_b(c)/2)$  is smooth with a positive lower bound, and  $(\mathcal{M}_b)_{12}$  has a positive lower bound. Their first derivatives are bounded. Taking the logarithm of (2.4) and differentiating proves (9.5). Since  $\coth(n\chi) \leq \coth \chi_*$  for  $n \geq 1$ , the gradient is bounded by  $C(1+n)$ . Integrating over the straight segment between the curvature vectors proves (9.4). Dividing the derivative identity by  $n$  proves the final limit. There are admissible directions with  $\dot{\chi} \neq 0$ : increasing one positive curvature makes the trace of the positive monodromy increase strictly. This follows directly by differentiating the displayed positive matrix product. A sufficiently small local curvature bump realizes such a direction while preserving the marked polygon and its open margins.  $\square$

We next turn the identity into an estimate on realizable finite records. Use the bounded class  $\mathcal{U}$  of Section 8, with  $m \geq 3$  chosen so that  $6a^m/(1-a^m) < 1$  on the enlarged positive quadratic class, and the same finite forward derivative bounds. Choose  $Q \subseteq Q_+$  within the common stability collar, with  $0 \in \operatorname{int} Q$ . Densities below are normalized on  $Q_+$  and restricted to  $Q$  without renormalization. Fix one phase  $b_0$ . The forward density floors are uniform for  $N \geq N_0$ , after increasing  $N_0$  if necessary.

**Proposition 9.4** (Stability on the finite-record image). *For  $\vartheta = (T, d)$  and  $\tilde{\vartheta} = (\tilde{T}, \tilde{d})$  in  $\mathcal{U}$ , put*

$$\Delta_f = \max_b \|f_{N,b}^\vartheta - \tilde{f}_{N,b}^{\tilde{\vartheta}}\|_{C^m(Q)},$$

$$\Delta_\pi = \left| \log \frac{\pi_{N,b_0}^{\tilde{\vartheta}}}{\pi_{N,b_0}^\vartheta} \right|, \quad e_N = \Delta_f + e^{-\omega N}.$$

There are class constants  $C, \eta_1 > 0$  such that  $e_N \leq \eta_1$  implies

$$\|F^T - \tilde{F}^{\tilde{T}}\|_{C^0(I)} + \|d - \tilde{d}\|_\infty + \|c - \tilde{c}\|_\infty \leq C e_N, \quad (9.6)$$

$$\left| \log \frac{\mathcal{A}(\tilde{T})}{\mathcal{A}(T)} \right| \leq C(1+N)e_N + \Delta_\pi. \quad (9.7)$$

An upper bound for the areas also gives the corresponding absolute area estimate. No absolute error is divided by an exponentially small twist.

*Proof.* The uniform relative comparison bounds the distance between the two limiting law tuples by  $\Delta_f + 2C_0e^{-\omega N}$ . Make  $\eta_1$  small enough that Theorem 8.3 applies. Its proof first applies the fixed-slice separation to obtain  $C^m$  control of the future actions. Their second derivatives, followed by the global positive quadratic estimate (4.6), control  $c - \tilde{c}$  as well. This proves (9.6).

Apply the finite identity (9.1) at  $b_0$  to both actual candidates. Subtraction of its logarithms gives

$$\begin{aligned} \log \frac{\mathcal{A}(\tilde{T})}{\mathcal{A}(T)} &= \log \frac{D_{N,b_0}(\tilde{c})}{D_{N,b_0}(c)} + \log \frac{\tilde{d}_{b_0}}{d_{b_0}} - \log \frac{\pi_{N,b_0}^{\tilde{\vartheta}}}{\pi_{N,b_0}^{\vartheta}} \\ &\quad - \log \frac{f_{N,b_0}^{\tilde{\vartheta}}(0,0)}{f_{N,b_0}^{\vartheta}(0,0)}. \end{aligned}$$

The positive offset and density floors control the second and fourth terms by  $Ce_N$ . Lemma 9.3 and (9.6) control the first term. This proves (9.7). If both areas are at most  $\mathcal{A}_+$ , the mean value theorem for the exponential gives  $|\mathcal{A}(\tilde{T}) - \mathcal{A}(T)| \leq \mathcal{A}_+ |\log \mathcal{A}(\tilde{T}) - \log \mathcal{A}(T)|$ .  $\square$

**9.4. A jointly realizable estimator.** The count at phase  $b_0$  is already part of the charged record; it does not require another preparation group. To keep the area and contacts compatible, use the finite-record image

$$\mathcal{J}_N(\vartheta) = ((f_{N,b}^{\vartheta}|_Q)_b, \log \pi_{N,b_0}^{\vartheta}) \in C^m(Q)^J \times \mathbb{R}, \quad J = r, \quad (9.8)$$

with the sum of the maximum density norm and the scalar absolute value. For each returning  $N$  choose a countable dense subset of this image and retain one actual representative  $\vartheta_j \in \mathcal{U}$  at each point. Such a subset exists because the ambient metric space is separable. The representatives are not required to exhaust the tables.

For  $B$  resets in each phase, let  $S_b$  be its success count and put  $q_N = c_0e^{-\Gamma N}$  and  $n = \lfloor Bq_N/2 \rfloor$ . On the event  $S_b \geq n$  for every phase, form the interior kernel estimates  $\hat{f}_{N,b}$  from the first  $n$  successful marks and put  $\hat{\pi} = S_{b_0}/B$ . Select the first representative whose distance from  $((\hat{f}_{N,b})_b, \log \hat{\pi})$  is within  $n^{-2}$  of the infimum over the countable image. On the complementary event choose one fixed member of  $\mathcal{U}$ . This defines a measurable estimator  $(\hat{T}, \hat{d})$  for every record, with its area  $\mathcal{A}(\hat{T})$  taken from that same representative. It is not an efficient enumeration algorithm or a projection onto an arbitrary factor-model image.

**Theorem 9.5** (Charged recovery of the complete invariant). *In the stated bounded class, set*

$$\alpha = \frac{s}{2(m+s)+2}, \quad \gamma = \frac{s}{m+s+3}, \quad L_n = \log(C_1 J n / \zeta), \quad 0 < \zeta < \frac{1}{2}.$$

*Suppose  $N \geq N_0$ ,  $n \geq 2$ ,  $Bq_N \geq C_2 \log(4J/\zeta)$  and  $L_n/n \leq 1$ . Use*

$$h = \max\{(L_n/n)^{1/[2(m+s)+2]}, \delta^{1/(m+s+3)}, n^{-1}\} \leq h_0.$$

*For  $e^{-\omega N} + (L_n/n)^\alpha + \delta^\gamma$  sufficiently small, the estimator above satisfies, with probability at least  $1 - 2\zeta$ ,*

$$\begin{aligned} &\|F^{\hat{T}} - F^T\|_{C^0(I)} + \|\hat{d} - d\|_\infty + \left| \log \frac{\mathcal{A}(\hat{T})}{\mathcal{A}(T)} \right| \\ &\leq C(1+N)\{e^{-\omega N} + (L_n/n)^\alpha + \delta^\gamma\}. \end{aligned} \quad (9.9)$$

*All  $JB = rB$  resets, including failures, are charged. Readout errors are bounded by  $\delta$  after acceptance; the tags and gates remain exact. Only the count at  $b_0$  is used as additional geometric information. The estimate is a sufficient upper bound, not an optimality assertion.*

*Proof.* For a binomial count with success probability  $\pi_b \geq q_N$ , Bernstein's inequality gives, for  $t > 0$ , a failure probability at most  $2e^{-t}$  for

$$\left| \frac{S_b/B}{\pi_b} - 1 \right| \leq \sqrt{\frac{2t}{B\pi_b}} + \frac{2t}{3B\pi_b}.$$

Take  $t = \log(4J/\zeta)$ . Increasing the absolute constant  $C_2$  makes the right side at most  $1/2$  simultaneously for all phases, except on an event of probability at most  $\zeta$ . On that event  $S_b \geq B\pi_b/2 \geq n$ , and  $|\log(1+x)| \leq 2|x|$  gives

$$|\log \hat{\pi} - \log \pi_{N,b_0}| \leq C\sqrt{t/(Bq_N)}. \quad (9.10)$$

Independent resetting admits the Bernoulli/conditional-mark representation from Section 6. Thus the first  $n$  successful latent marks have the required conditional law. The density estimate can be justified on their infinite extension even when the finite record has fewer successes; the count event ensures that only actually observed marks are used. No independence between the count error and the density error is required for the union bound.

Lemma 6.1, before adding any physical limit bias, gives a simultaneous finite-density error at most

$$C \left\{ h^s + \sqrt{\frac{L_n}{nh^{2m+2}}} + \frac{L_n}{nh^{m+2}} + n^{-2} + \delta h^{-m-3} \right\}$$

except on an event of probability at most  $\zeta$ . The displayed bandwidth bounds this by  $C\{(L_n/n)^\alpha + \delta^\gamma\}$ . Because  $n$  is comparable to  $Bq_N$  under  $n \geq 2$ , and  $t \leq CL_n$ , (9.10) has this same upper bound:  $\alpha < 1/2$  and  $L_n/n \leq 1$ . The fitted finite image is consequently within  $C\{(L_n/n)^\alpha + \delta^\gamma\}$  of the true image.

Density of the countable image and the near-minimum selection show that the distance between the selected and true images is at most twice this error plus  $n^{-2}$ . Both images correspond to actual pairs in  $\mathcal{U}$ , including their actual areas. Apply Proposition 9.4, adding only now the forward term  $e^{-\omega N}$ . The stipulated small-error condition ensures its applicability. The two exceptional events have total probability at most  $2\zeta$ , proving the theorem. There is no independent numerical area estimate attached to an unrelated completion.  $\square$

**Corollary 9.6** (A budget and readout dependent returning length). *Let*

$$\beta = \frac{\alpha\omega}{\omega + \alpha\Gamma}, \quad L_B = \log(C_*JB/\zeta), \quad t_{B,\delta} = (L_B/B)^\beta + \delta^\gamma < 1,$$

and choose

$$N = P \left\lfloor \frac{\log(1/t_{B,\delta})}{P\omega} \right\rfloor, \quad n = \lfloor Bc_0e^{-\Gamma N}/2 \rfloor. \quad (9.11)$$

Whenever the count, bandwidth, returning-length and small-error conditions of Theorem 9.5 hold, its right side is at most

$$Ct_{B,\delta} \log(e/t_{B,\delta}). \quad (9.12)$$

At zero readout error the returning-length choice is the same as in Theorem 8.5. At positive readout error it stops the length before curvature error is unnecessarily amplified. The original profile-and-offset rate (8.11) is retained separately; the logarithmic factor here pays for the additionally reconstructed area.

*Proof.* Write  $x = (L_B/B)^\beta$  and  $t = t_{B,\delta} \geq x$ . The floor gives  $e^{-\omega N} \leq e^{\omega P}t$  and  $1+N \leq C \log(e/t)$ . Also  $e^{-\Gamma N} \geq t^{\Gamma/\omega}$ , so the lower-count condition gives  $n \geq cBt^{\Gamma/\omega}$ . Since  $n \leq B$ ,  $L_n \leq CL_B$ , and hence

$$(L_n/n)^\alpha \leq C(L_B/B)^\alpha t^{-\alpha\Gamma/\omega} \leq C(L_B/B)^\alpha x^{-\alpha\Gamma/\omega} = Cx \leq Ct.$$

The equality is the defining identity for  $\beta$ . Finally  $\delta^\gamma \leq t$ . Substitution into (9.9) proves (9.12). When  $\delta = 0$ ,  $(\log(1/x))/\omega = \alpha \log(B/L_B)/(\omega + \alpha\Gamma)$ , giving the last comparison.  $\square$

The complete invariant is therefore accessible without numerical clock calibration, supplied curvature or separately calibrated local flux amplitudes. The area assertion uses normalized unit-speed Liouville preparations and exact acceptance tags. Its contact part uses the actual smooth inverse, its volume part uses the unconditioned phase normalization, and their finite reconstruction uses relative rather than absolute rare-event control. The classification and estimate concern this specified marked reset experiment. They neither identify it with an exterior scattering observation nor turn its nontrivial equal-area completion fibers into determined boundary arcs.

## 10. RECONSTRUCTION FROM FIRST-MOMENT PROFILES

The separated limiting law has an additional consequence for the amount of smoothing required by the inverse problem. Four functions of one variable determine the two-variable law. The relevant moment matrix is nonsingular for a geometric reason: the two marked first derivatives have opposite signs. We first prove this assertion on the limiting model, then retain the finite-flight error, and finally construct an estimator on the image of actual tables. All samples in this section are the same paired endpoint records as before.

**10.1. An exact moment identity.** Fix a symmetric interval  $J = [-R_+, R_+]$  and an interior interval  $I = [-R, R] \Subset J$ , both containing zero. For an integrable kernel  $f$  on  $J^2$  define

$$M_j(u) = \int_J v^j f(u, v) dv, \quad N_j(v) = \int_J u^j f(u, v) du, \quad j = 0, 1, \quad (10.1)$$

$$H_f = \int_{J^2} \begin{pmatrix} 1 \\ u \end{pmatrix} (1 \ v) f(u, v) du dv. \quad (10.2)$$

Write  $M(u) = (M_0(u), M_1(u))$  and  $N(v) = (N_0(v), N_1(v))^\top$ . The compression used below is

$$\mathcal{C}_I(f) = (M_0|_I, M_1|_I, N_0|_I, N_1|_I, H_f) \in \mathcal{X}_m := C^m(I)^4 \times \mathbb{R}^{2 \times 2}.$$

Equip this space with the sum of the largest function norm and the matrix operator norm. The integration domain in all five components is the full gate  $J^2$ . In particular,  $H_{f,00} = 1$  for a density normalized on that gate; restriction to  $I$  does not renormalize it. The functions  $M_1, N_1$  retain paired information and are not ordinary marginal densities.

**Lemma 10.1** (Moment reconstruction). *Suppose  $f(u, v) = a(u)^\top C b(v)$ , where  $a, b$  have two components and  $C$  is a  $2 \times 2$  matrix. If  $H_f$  is nonsingular, then*

$$f(u, v) = M(u) H_f^{-1} N(v), \quad (u, v) \in J^2. \quad (10.3)$$

For each integer  $m \geq 0$ , the map

$$\mathcal{R} : (M_0, M_1, N_0, N_1, H) \longmapsto [(u, v) \mapsto M(u) H^{-1} N(v)]$$

is uniformly Lipschitz into  $C^m(I^2)$  on sets with bounded  $\mathcal{X}_m$  norm and  $|\det H| \geq \kappa_* > 0$ . Here  $C^m(I^2)$  uses derivatives of total order at most  $m$ .

*Proof.* Put

$$L = \int_J \begin{pmatrix} 1 \\ u \end{pmatrix} a(u)^\top du, \quad R_0 = \int_J b(v) (1 \ v) dv.$$

Fubini's theorem gives  $H_f = L C R_0$ ,  $M(u) = a(u)^\top C R_0$ , and  $N(v) = L C b(v)$ . Nonsingularity of their square product implies that all three factors are nonsingular. Cancelling them proves (10.3). This argument also proves that the separation rank is exactly two under the hypothesis.

On the stated bounded set, the adjugate formula bounds  $\|H^{-1}\|$ . For a second matrix  $\tilde{H}$  use

$$H^{-1} - \tilde{H}^{-1} = H^{-1}(\tilde{H} - H)\tilde{H}^{-1}.$$

Add and subtract the three products obtained by replacing one factor at a time. For  $i + j \leq m$ , differentiation of a product gives  $M^{(i)}(u)H^{-1}N^{(j)}(v)$ . The same bounds therefore control every required derivative. No derivative of the data above order  $m$  is used.  $\square$

This is a continuous moment version of the classical exact low-rank matrix reconstruction underlying CUR decompositions; see [55]. Its use here requires, in addition, a nondegeneracy bound forced by the billiard law and control of the defect of finite-flight laws from the separated model.

**Proposition 10.2** (Geometric nondegeneracy). *Let the limiting law at one phase be*

$$f(u, v) = Z^{-1}w(u)z(v)\{d - A(u) - C(v)\}, \quad (u, v) \in J^2, \quad (10.4)$$

*with the hypotheses of (8.1)–(8.2). Suppose throughout  $J$  that  $A'$  has sign  $-p$ ,  $C'$  has sign  $p$ , and  $|A'|, |C'| \geq p_*/2 > 0$ . Then  $\det H_f > 0$  and (10.3) holds.*

*On a bounded class with  $|p| \geq p_*$ , fixed  $J$ , uniform positive lower and upper bounds for  $w, z$ , and an upper bound for  $Z$ , the determinant has a positive uniform lower bound. The derivative hypotheses follow by shrinking  $J$  once from the uniform  $C^2$  action bounds and  $A'(0) = -p$ ,  $C'(0) = p$ . They are not additional measured amplitudes or supplied curvature data.*

*Proof.* Use

$$a(u) = \begin{pmatrix} w(u) \\ w(u)A(u) \end{pmatrix}, \quad b(v) = \begin{pmatrix} z(v) \\ z(v)C(v) \end{pmatrix}, \quad C_0 = Z^{-1} \begin{pmatrix} d & -1 \\ -1 & 0 \end{pmatrix}$$

in Lemma 10.1. To compute the determinant without assuming its nonvanishing, set  $W = \int_J w$ ,  $V = \int_J z$ , and let  $\mu_w, \mu_z$  have densities  $w/W, z/V$ . Direct evaluation gives

$$\det H_f = -\frac{W^2 V^2}{Z^2} \text{Cov}_{\mu_w}(u, A(u)) \text{Cov}_{\mu_z}(v, C(v)). \quad (10.5)$$

Indeed the two covariance factors, multiplied by  $W^2, V^2$ , are respectively  $\det L, \det R_0$ , and  $\det C_0 = -Z^{-2}$ . For a probability measure  $\mu$  and independent  $X, Y$  of law  $\mu$ ,

$$\text{Cov}_{\mu}(x, g(x)) = \frac{1}{2} \mathbb{E}[(X - Y)(g(X) - g(Y))].$$

Consequently a derivative of constant sign and magnitude at least  $p_*/2$  gives a covariance of that sign and magnitude at least  $(p_*/2) \text{Var}_{\mu}(x)$ . The two covariances in (10.5) have opposite signs.

For completeness, if  $w_* \leq w \leq w^*$  on  $J$ , then

$$\text{Var}_{\mu_w}(u) \geq \frac{w_*}{2R_+ w^*} \inf_{c \in \mathbb{R}} \int_{-R_+}^{R_+} (u - c)^2 du = \frac{w_*}{w^*} \frac{R_+^2}{3}.$$

The same estimate applies to  $z$ , while  $W, V$  have positive lower bounds. Substitution in (10.5) proves the uniform claim. Finally, if the second action derivatives are bounded by  $K_2$ , choose  $K_2 R_+ \leq p_*/2$ . The mean value theorem gives the asserted sign and magnitude bounds. All constants may deteriorate as the fixed gate shrinks or  $p_*$  tends to zero; no uniformity in those limits is claimed.  $\square$

*Remark 10.3* (Explicit fixed-gate conditioning). On  $J = [-R, R]$ , let  $p \neq 0$  and  $d > 2|p|R$ . The positive normalized rank-two model

$$f(u, v) = \frac{d + pu - pv}{4R^2 d}$$

has moments

$$\mathbb{E}U = \frac{pR^2}{3d}, \quad \mathbb{E}V = -\frac{pR^2}{3d}, \quad \mathbb{E}(UV) = 0.$$

Indeed the odd integrals vanish and  $\int_{-R}^R u^2 du = 2R^3/3$ . Hence

$$H_f = \begin{pmatrix} 1 & -pR^2/(3d) \\ pR^2/(3d) & 0 \end{pmatrix}, \quad \det H_f = \frac{p^2 R^4}{9d^2}. \quad (10.6)$$

Thus even this simple model loses the determinant at fourth order in the gate radius and second order in obliquity. For  $a = pR^2/(3d)$ ,  $H_f^{-1} = \begin{pmatrix} 0 & 1/a \\ -1/a & 1/a^2 \end{pmatrix}$ , so the inverse-matrix norm also diverges. This is an algebraic model, not an assertion that affine half-line actions are realized by a positive-curvature billiard. It illustrates why the fixed-class constants in Proposition 10.2 retain their gate and slope dependence. Rescaling coordinates changes moment units; it does not prove uniform physical stability in a degenerating class. The affine computation was supplied in the author-requested memorandum [56]; it is included here with its derivation.

**Corollary 10.4** (Moment determination of the local invariant). *For the oblique marked class, four limiting moment profiles and their matrix  $H_f$  at each phase determine the offsets and all smooth contact germs. Adding the charged success probabilities as a function of the returning length determines the free area. Conversely, common contact germs, offsets and free area give identical compressed charged local experiments on sufficiently small common gates.*

*Proof.* Choose the common gates as in Proposition 10.2. The moment identity reconstructs each limiting law on  $I^2$ . Theorem 8.3 determines the offsets and smooth contact germs. Their curvatures give  $D_{N,b}$ , and (9.2) then determines the area from one phase. The converse is the locality direction of Theorem 9.2, followed by integration. Applying this argument on a cofinal collection of gates gives the assertion for germs.  $\square$

At finite  $N$  the identity in Lemma 10.1 need not hold for  $f_N$ . Thus the corollary concerns a family of laws and its relative limit, not a sufficient statistic for one finite sample. Nor does it replace paired endpoints by two unpaired marginal samples. For example, at normal incidence the positive model  $Z^{-1}(d - au^2 - cv^2)$  on a symmetric square can have rank two but  $M_1 = N_1 = 0$ , so  $H_f$  is singular. The two-offset results at normal incidence remain available independently of the first-moment inverse.

**10.2. Finite-flight comparison on the physical image.** Fix the bounded class  $\mathcal{U}$  of Proposition 9.4. In particular,  $m \geq 3$  satisfies the explicit later-visit condition on the enlarged positive quadratic class. Choose  $J, I \subseteq J$ , and a contact collar  $I_c \subseteq I$  once so that the preceding proposition and the inherited action-to-contact stability apply uniformly. Densities are normalized on the success gate  $J^2$ ; all integrations below still use that whole gate. The relative physical theorem supplies

$$\max_b \|f_{N,b}^\vartheta - f_b^\vartheta\|_{C^m(J^2)} \leq C_0 e^{-\omega N}, \quad N \geq N_0, \quad (10.7)$$

uniformly over actual pairs  $\vartheta = (T, d) \in \mathcal{U}$ . Shrinking the fixed gates changes class constants, not the charged preparation convention.

**Proposition 10.5** (Stability of compressed finite records). *For two actual pairs in  $\mathcal{U}$  put*

$$\Delta_M = \max_b \|\mathcal{C}_I(f_{N,b}^\vartheta) - \mathcal{C}_I(\tilde{f}_{N,b}^\vartheta)\|_{\mathcal{X}_m}, \quad e_N = \Delta_M + e^{-\omega N}, \quad \Delta_\pi = \left| \log \frac{\pi_{N,b_0}^{\tilde{\vartheta}}}{\pi_{N,b_0}^\vartheta} \right|.$$

There are class constants  $C, \eta_* > 0$  such that  $e_N \leq \eta_*$  implies

$$\|F^T - \tilde{F}^T\|_{C^0(I_c)} + \|d - \tilde{d}\|_\infty + \|c - \tilde{c}\|_\infty \leq Ce_N, \quad (10.8)$$

$$\left| \log \frac{\mathcal{A}(\tilde{T})}{\mathcal{A}(T)} \right| \leq C(1 + N)e_N + \Delta_\pi. \quad (10.9)$$

Here  $c$  is the vector of contact curvatures. The estimate does not assume exact rank two at finite flight length.

*Proof.* Integration in (10.1)–(10.2) is a bounded linear map from  $C^m(J^2)$  to  $\mathcal{X}_m$ . Hence the two limiting compressed tuples are at distance at most  $\Delta_M + Ce^{-\omega N}$ . Proposition 10.2 gives their uniform determinant floor. Lemma 10.1 therefore bounds the distance of the two limiting densities on  $I^2$  by  $Ce_N$ . Equation (10.7) also gives

$$\max_b \|f_{N,b}^\vartheta - \tilde{f}_{N,b}^\vartheta\|_{C^m(I^2)} \leq Ce_N. \quad (10.10)$$

Now apply Proposition 9.4 on these fixed interior and contact collars. Its proof uses the  $C^m$  law norm to recover the actions and their quadratic coefficients before estimating the profiles. In particular it does not infer curvature control from  $C^0$  profile control. The exact same-gate central-density identity and the logarithmic twist comparison then give (10.9).

Equivalently, the preceding argument shows that  $\mathcal{R}(\mathcal{C}_I(f_{N,b}))$  approximates  $f_{N,b}|_{I^2}$  with error  $Ce^{-\omega N}$ , for  $N$  large enough. This is only an approximation: the finite central density in the area identity is the one belonging to the actual candidate, not an exactly identified rank-two surrogate.  $\square$

**10.3. Estimation of the four profiles.** We give the sampling estimate in detail because smoothing a weighted marginal, rather than a two-variable density, changes the exponents. Let  $s \geq 1$  be an integer. Assume the same uniform forward bounds as before, now stated on the full fixed gate:

$$\sup_{N \geq N_0, b, \vartheta \in \mathcal{U}} \|f_{N,b}^\vartheta\|_{C^{m+s}(J^2)} \leq K.$$

All constants below may depend on this class, the fixed intervals,  $m, s$  and the number  $r$  of phases.

Choose  $K_0 \in C_c^{m+1}((-1, 1))$  with integral one and vanishing moments of orders  $1, \dots, s-1$ . Such a kernel can be constructed by multiplying a positive smooth bump by a polynomial of degree at most  $s-1$ : the moment conditions are a nonsingular Gram system for polynomials on the support of that bump. The kernel need not be nonnegative. Let  $h_0 \leq 1$  be smaller than the distance from  $I$  to  $\partial J$ .

For  $n$  accepted pairs at one phase, write their true coordinates as  $(U_i, V_i)$  and their readouts as  $(\tilde{U}_i, \tilde{V}_i)$ , with both coordinate errors at most  $0 \leq \delta \leq 1$ . Define

$$\widehat{M}_j(u) = \frac{1}{nh} \sum_{i=1}^n \tilde{V}_i^j K_0((u - \tilde{U}_i)/h), \quad (10.11)$$

$$\widehat{N}_j(v) = \frac{1}{nh} \sum_{i=1}^n \tilde{U}_i^j K_0((v - \tilde{V}_i)/h), \quad j = 0, 1, \quad (10.12)$$

$$\widehat{H} = \frac{1}{n} \sum_{i=1}^n \begin{pmatrix} 1 \\ \tilde{U}_i \end{pmatrix} (1 \quad \tilde{V}_i). \quad (10.13)$$

Denote the resulting tuple by  $\widehat{\mathcal{C}}_{N,b}$ . There is no inversion of a random nearly singular matrix in this definition.

**Lemma 10.6** (One-dimensional smoothing bound). *For  $n \geq 2$ ,  $n^{-1} \leq h \leq h_0$ , and  $0 < \zeta < 1/2$ , put  $L_n = \log(C_1 r n / \zeta)$ . The first  $n$  accepted latent pairs at each phase may be used in (10.11)–(10.13). With probability at least  $1 - \zeta$ ,*

$$\max_b \|\widehat{\mathcal{C}}_{N,b} - \mathcal{C}_I(f_{N,b})\|_{\mathcal{X}_m} \leq C \left\{ h^s + \sqrt{\frac{L_n}{nh^{2m+1}}} + \frac{L_n}{nh^{m+1}} + n^{-2} + \delta h^{-(m+2)} \right\}. \quad (10.14)$$

*The perturbations may depend arbitrarily on the latent sample. The probability statement is about the independent latent preparations; acceptance and its success/failure tags are exact.*

*Proof.* First set  $\delta = 0$ . For  $0 \leq k \leq m$ , the expectation of  $\widehat{M}_j^{(k)}(u)$  is the convolution of  $M_j^{(k)}$  with the rescaled kernel on the interior interval. Integration by parts is legitimate because the kernel support does not meet  $\partial J$ . The  $C^{m+s}$  bound on  $f_N$  gives the same order bound on  $M_j$ . Taylor's formula and the moment cancellations give a uniform bias  $Ch^s$ . The argument for  $N_j$  is identical.

For a fixed  $u$ , a summand before division by  $n$  has magnitude at most  $Ch^{-k-1}$  and variance at most  $Ch^{-2k-1}$ . Indeed the integral of the square of  $K_0^{(k)}((u - U)/h)$  against the bounded joint density gains one factor  $h$ ;  $|V|^j$  is uniformly bounded on  $J$ . Bernstein's inequality thus gives, with failure probability at most  $2e^{-t}$ , fluctuation at most

$$C \left\{ \sqrt{\frac{t}{nh^{2k+1}}} + \frac{t}{nh^{k+1}} \right\}.$$

Take a grid in  $I$  of mesh at most  $n^{-2}h^{m+2}$ . Both the empirical derivative and its expectation have Lipschitz constants at most  $Ch^{-m-2}$ . Passage from the grid to the whole interval therefore costs at most  $Cn^{-2}$ . Since  $h \geq n^{-1}$ , the grid has at most  $Cn^{m+4}$  points. A union bound over the four profiles,  $r$  phases and  $m+1$  derivative orders, with  $t = C_m L_n$ , proves the stated stochastic bound. The entries of  $\widehat{H}$  are averages of bounded variables and have error at most  $C\{\sqrt{L_n/n} + L_n/n\}$  on the same enlarged event. This is absorbed by the displayed bound when  $h \leq 1$ . Independence of the four profiles is neither assumed nor needed.

Restore the readout errors. The mean value theorem for the kernel and the bounded weights gives, for each summand of a derivative of order  $k$ ,

$$\left| \widetilde{V}_i^j h^{-k-1} K_0^{(k)}((u - \widetilde{U}_i)/h) - V_i^j h^{-k-1} K_0^{(k)}((u - U_i)/h) \right| \leq C\delta\{h^{-k-2} + h^{-k-1}\}.$$

The weights remain bounded by  $C(R_+)$  because  $\delta \leq 1$ . For  $k \leq m$  and  $h \leq 1$  the last bound is  $C\delta h^{-m-2}$ . Averaging does not increase it. The matrix entries have deterministic perturbation at most  $C\delta$ , also absorbed. This comparison holds sample by sample and so requires no distributional assumption on the readout errors.  $\square$

**10.4. Charged recovery with one-dimensional smoothing.** The actual finite image to be fitted is

$$\mathcal{K}_N(\vartheta) = ((\mathcal{C}_I(f_{N,b}^\vartheta))_{b=0}^{r-1}, \log \pi_{N,b_0}^\vartheta) \in \mathcal{X}_m^r \times \mathbb{R}. \quad (10.15)$$

Use the maximum of the tuple norms plus the scalar absolute value. For each returning  $N$ , fix a countable dense subset of this image, and one actual representative  $(T_j, d_j) \in \mathcal{U}$  for each of its points. This is possible because  $C^m(I)$ , and hence the ambient metric space, is separable. Closedness of the image is not required.

Perform  $B$  independent normalized unit-speed Liouville resets at each phase. Put  $q_N = c_0 e^{-\Gamma N}$  and  $n = \lfloor Bq_N/2 \rfloor$ . If every success count  $S_b$  is at least  $n \geq 2$ , use the first  $n$  accepted pairs for the estimated compressed tuples and set  $\widehat{\pi} = S_{b_0}/B$ . Select the first representative whose distance to  $((\widehat{\mathcal{C}}_{N,b})_b, \log \widehat{\pi})$  is within  $n^{-2}$  of the infimum over the countable image. Otherwise select a fixed member of  $\mathcal{U}$ . The infimum of countably many measurable distances is measurable, and a



qualifying index exists by the definition of infimum. Thus this specifies an actual table-offset estimator on every record. Its area is always the area of that same table. An effective enumeration or polynomial-time search is not asserted.

**Theorem 10.7** (Compressed charged reconstruction). *In the stated class set*

$$\alpha_1 = \frac{s}{2(m+s)+1}, \quad \gamma_1 = \frac{s}{m+s+2}. \quad (10.16)$$

Suppose  $N \geq N_0$ ,  $n \geq 2$ ,  $Bq_N \geq C_2 \log(4r/\zeta)$ , and  $L_n/n \leq 1$ . Choose

$$h = \max\{(L_n/n)^{1/[2(m+s)+1]}, \delta^{1/(m+s+2)}, n^{-1}\} \leq h_0.$$

For  $\varepsilon_N = e^{-\omega N} + (L_n/n)^{\alpha_1} + \delta^{\gamma_1}$  sufficiently small, the estimator just defined satisfies, with probability at least  $1 - 2\zeta$ ,

$$\|F^{\hat{T}} - F^T\|_{C^0(I_c)} + \|\hat{d} - d\|_\infty + \|\hat{c} - c\|_\infty \leq C\varepsilon_N, \quad (10.17)$$

$$\left| \log \frac{\mathcal{A}(\hat{T})}{\mathcal{A}(T)} \right| \leq C(1+N)\varepsilon_N. \quad (10.18)$$

All  $rB$  preparations, including failures, are charged. No additional phase group is needed for the area. Readout errors occur after exact acceptance and do not corrupt the success tags.

*Proof.* The count argument in (9.10), with  $t = \log(4r/\zeta)$ , gives simultaneously

$$S_b \geq n \quad \text{for every } b, \quad |\log \hat{\pi} - \log \pi_{N,b_0}| \leq C\sqrt{t/(Bq_N)}$$

except on an event of probability at most  $\zeta$ . The Bernoulli/conditional-mark representation of independent resets gives independent latent marks of density  $f_{N,b}$  at successive successes. They may be extended beyond the finite record when proving a probability bound; on the count event only the first  $n$  actually observed marks are used. This construction does not condition the marks on a possibly mark-dependent noisy acceptance test: acceptance is exact.

Apply Lemma 10.6 to these latent sequences. At the prescribed bandwidth, its right side is at most  $C\{(L_n/n)^{\alpha_1} + \delta^{\gamma_1}\}$ . For the variance term this follows by substituting the first bandwidth lower bound; the envelope exponent becomes  $(m+2s)/[2(m+s)+1] \geq \alpha_1$ . The readout term is bounded using the second lower bound, and the bias is bounded using the maximum formula. The  $n^{-2}$  term is absorbed. Since  $\alpha_1 < 1/2$ ,  $n$  is comparable to  $Bq_N$ , and  $L_n/n \leq 1$ , the count error is absorbed by the same expression.

Density of the countable actual image implies that its infimal distance to the estimated tuple is no larger than the distance of the true image. The selected image is therefore within twice the estimation error plus  $n^{-2}$  of the true image. Both candidates are actual pairs. Proposition 10.5 gives the assertions, with the physical error  $e^{-\omega N}$  added only at this step. The union of the count and smoothing exceptional events has probability at most  $2\zeta$ .  $\square$

**Corollary 10.8** (Budget-dependent flight length). *Let*

$$\beta_1 = \frac{\alpha_1 \omega}{\omega + \alpha_1 \Gamma}, \quad L_B = \log(C_* r B / \zeta), \quad t = (L_B/B)^{\beta_1} + \delta^{\gamma_1} < 1,$$

and choose  $N = P\lceil \log(1/t)/(P\omega) \rceil$ . Whenever the hypotheses of Theorem 10.7 hold, its local error is at most  $Ct$  and its joint local-and-area error is at most

$$Ct \log(e/t). \quad (10.19)$$

*Proof.* The floor gives  $e^{-\omega N} \leq e^{\omega P} t$ ,  $1 + N \leq C \log(e/t)$  and  $n \geq cBt^{\Gamma/\omega}$  under the count conditions. Write  $x = (L_B/B)^{\beta_1} \leq t$ . Since  $L_n \leq CL_B$ ,

$$(L_n/n)^{\alpha_1} \leq C(L_B/B)^{\alpha_1} t^{-\alpha_1 \Gamma/\omega} \leq C(L_B/B)^{\alpha_1} x^{-\alpha_1 \Gamma/\omega} = Cx.$$

The last equality is the definition of  $\beta_1$ . Also  $\delta^{\gamma_1} \leq t$ . Substitution in (10.17)–(10.18) proves the claims.  $\square$

For the same  $m, s, \omega, \Gamma$ , the exponents  $\alpha_1$  and  $\gamma_1$  strictly improve the two-dimensional smoothing exponents  $s/[2(m+s)+2]$  and  $s/(m+s+3)$  in Theorem 9.5. The corresponding budget exponent improves as well, since  $\alpha \mapsto \alpha\omega/(\omega + \alpha\Gamma)$  is strictly increasing. These are comparisons of proved upper bounds on the same oblique bounded class, not minimax assertions or comparisons with lens or marked-length observations. The improvement comes from estimating one-dimensional weighted profiles, not from dividing a small absolute error by a rare success probability. The full smooth inverse and its finite-order later-visit estimate remain essential to the passage from these profiles to the actual reflecting contacts.

## 11. CONDITIONING OF THE PHYSICAL-TO-SMOOTH INVERSE

The inverse has three different quantitative costs: separating the physical actions, inverting their actual smooth response, and acquiring enough successful trajectories. They depend on different geometric margins. We collect them here on the fixed classes used in the preceding results. None of the following sufficient bounds is a lower bound for all possible reconstruction procedures.

**11.1. Contraction, differentiability and phase weights.** For the scalar visit bound  $a < 1$ , the sufficient weighted threshold and its least integer order are exactly

$$\theta_m < 1 \iff a^m < \frac{1}{7}, \quad m_{\min} = \max \left\{ 3, \left\lceil \frac{\log 7}{-\log a} \right\rceil + 1 \right\}. \quad (11.1)$$

At a chosen order the weighted inverse prefactor is  $2/(1 - \theta_m)$ . Minimizing the order need not give a comfortable inverse margin. For any chosen  $0 < \vartheta < 1$ , the choice

$$m \geq \max \left\{ 3, \left\lceil \frac{\log((6 + \vartheta)/\vartheta)}{-\log a} \right\rceil \right\} \quad (11.2)$$

ensures  $\theta_m \leq \vartheta$  and hence a prefactor at most  $2/(1 - \vartheta)$ . These formulas follow by solving  $6a^m/(1 - a^m) \leq \vartheta$ ; the strict inequality in (11.1) explains the floor followed by +1.

For example, the scalar bounds give the following values. The last column uses  $s = 1$  in the one-dimensional smoothing exponent; it is not an asserted optimal statistical rate.

| $a$    | $m_{\min}$ | $\theta_{m_{\min}}$ | $2/(1 - \theta_{m_{\min}})$ | $\alpha_1$ |
|--------|------------|---------------------|-----------------------------|------------|
| 1/2    | 3          | 0.857143            | 14                          | 1/9        |
| 9/10   | 19         | 0.937099            | 31.7961                     | 1/41       |
| 99/100 | 194        | 0.995514            | 445.8342                    | 1/391      |

Taking instead  $\vartheta = 1/2$  gives orders 4, 25, 256, respectively, and weighted prefactor at most 4. It increases the differentiability required of the observed action and the graph prior. The jet-alignment and forward constants also depend on this order; bounding the weighted prefactor alone does not bound the entire inverse.

Corollaries 5.8 and 5.9 replace  $a$  in these sufficient tests by  $\bar{a} = (\prod_i \hat{a}_i)^{1/r}$  when phase-specific bounds are available on the enlarged class. The unweighted prefactor then contains  $\kappa(w)$ . The product of the phasewise upper bounds must not be replaced by the supremum of the products along one orbit without a new uniform estimate. Likewise an exactly computed linear multiplier does not itself provide the nonlinear collar bounds: the latter follow by continuity on the stipulated class. The scalar criterion remains valid and is often simpler to use.

**11.2. Observation margins.** In the two-offset quotient, put  $h = A(u) + C(v)$  and  $\Delta = d_2 - d_1 > 0$ . On the exact image,

$$d_2 - d_1 Q = \frac{d_2 \Delta}{d_2 - h}, \quad \left| \frac{\partial H}{\partial Q} \right| = \frac{d_1 (d_2 - h)^2}{d_2 \Delta}.$$

Thus offset separation and the bounds for the positive residual enter the action inverse explicitly. Positive density floors are also required to form  $Q$ ; its higher derivatives involve the same reciprocal functions. This route does not require a nonzero obliquity  $p$  and therefore retains the normal-incidence case.

For a single law, the inverse additionally uses  $\kappa_f = p^2/d^2$  and the fixed-slice anchors  $u_a, v_a$  of (8.6). The map  $\kappa \mapsto |p|/\sqrt{\kappa}$  has derivative of absolute value  $|p|/(2\kappa^{3/2})$ . The geometric bounds after Theorem 8.2 supply  $\kappa_f \geq p_*^2/d_+^2$  and  $|u_a|, |v_a| \geq p_* a/(2D_*)$ . Their degeneration is independent of the order condition for the smooth envelope.

Moment reconstruction brings its own matrix condition. In the affine model of Remark 10.3, put  $z = pR^2/(3d)$ . Then

$$H = \begin{pmatrix} 1 & -z \\ z & 0 \end{pmatrix}, \quad \det H = z^2, \quad H^{-1} = \begin{pmatrix} 0 & 1/z \\ -1/z & 1/z^2 \end{pmatrix}.$$

The shrinking-gate and small-obliquity losses are visible in the inverse, not just its determinant. This is an algebraic density example, not an asserted positive-curvature billiard realization. Rescaling the moments changes their units and the corresponding error norms; it does not remove the physical conditioning from a comparison in fixed units.

These requirements are separate from the positive Jacobi mass, nongrazing and clearance bounds that construct the stationary trajectories. They are also separate from the uniform  $C^{m+3}$  graph prior used to align jets. On a fixed class with all these margins, their contributions are finite; uniformity is not asserted as any of them degenerates.

**11.3. Preparation budget and the exact finite anchor.** For the moment protocol the conditional-sample and bounded-readout exponents are

$$\alpha_1 = \frac{s}{2(m+s)+1}, \quad \gamma_1 = \frac{s}{m+s+2}.$$

They arise from the complete bound in Lemma 10.6, not just its square-root term. In particular the Bernstein envelope, interpolation-grid and post-acceptance perturbation terms remain part of the estimate. Suppose the uniform finite-flight defect is  $Ce^{-\omega N}$  and the success probability is at least  $ce^{-\Gamma N}$ . For  $B$  independent preparations per phase, the two terms to balance have the form

$$e^{-\omega N}, \quad \left( \frac{L_B e^{\Gamma N}}{B} \right)^{\alpha_1}.$$

Solving their equality gives

$$N \sim \frac{\alpha_1}{\omega + \alpha_1 \Gamma} \log \frac{B}{L_B}, \quad \beta_1 = \frac{\alpha_1 \omega}{\omega + \alpha_1 \Gamma}. \quad (11.3)$$

Rounding to a returning number of flights changes  $N$  by a bounded amount. The count, bandwidth, minimum-flight and small-error conditions in Theorem 10.7 and Corollary 10.8 must still hold. Increasing  $m$  may improve the weighted inverse margin while decreasing  $\alpha_1$ ; rarity decreases it further to  $\beta_1$ . A positive exponent is not a bound for the computational cost of finding the selected actual table.

Counts enter area recovery through the exact finite identity

$$\mathcal{A}(T) = \frac{D_{N,b}(c) d_b}{2\pi \pi_{N,b} f_{N,b}(0,0)}.$$

The count and the fitted actual finite density use the same full gate. Replacing the latter by its rank-two limit would discard a finite-flight term. The logarithmic derivative of  $D_{N,b}$  costs  $1 + N$ , as shown in Section 9; at (11.3) it produces the retained logarithmic area loss. Failures are charged, gate and itinerary tags are exact, and readout perturbations occur after acceptance. The independently reset Liouville experiment, not an arbitrary unknown preparation measure, supplies the normalizer. These are conditions on the stated inverse problem, not additional conclusions of the conditioning calculation.

## Part 2. Finite-chain foundations and alternating laws

### 12. GEOMETRIC SETTING AND THE PRINCIPAL RELATIVE THEOREM

Fix a full-rank lattice  $\Lambda \subset \mathbb{R}^2$  and finitely many bounded  $C^\infty$  strictly convex obstacles  $D_1, \dots, D_s$  with positive curvature. Distinct lattice translates have disjoint closures. Let  $A > 0$  be the free area per lattice cell and prepare the unit-speed billiard with normalized phase volume  $dq d\vartheta / (2\pi A)$  on the torus complement. We observe the collision count and, when specified, the first and last collision positions. Finite horizon, congruence and symmetry are not assumed for the general forward or boundary-law results.

Consider a fixed configuration and a smooth finite-dimensional family  $D_{a,\xi}$  through it, with  $\xi$  in a compact set  $K_0$  contained in a sufficiently small parameter neighborhood. At every fixed derivative order, constants may depend on the corresponding higher norms of this family, on positive separation and curvature bounds, and on the lattice, but not on the collision order.

Let  $\mathcal{E}$  be the finite set, modulo common lattice translation, of oriented closest pairs attaining the minimum gap at the reference configuration. An index  $e = (a, b, \ell)$  represents a segment from  $D_{a,\xi}$  to  $D_{b,\xi} + \ell$ ; reversing it gives  $\bar{e} = (b, a, -\ell)$ . Write

$$\begin{aligned} g_* &= \min_{e \in \mathcal{E}} g_e, & c_{e,b} &= 1 + g_e \kappa_{e,b} \quad (b = 0, 1), \\ c_e &= \sqrt{c_{e,0} c_{e,1}}, & \gamma_e &= \operatorname{arcosh} c_e, & d_{j,e} &= t - j g_e. \end{aligned} \quad (12.1)$$

Here  $\kappa_{e,0}$  and  $\kappa_{e,1}$  are the positive curvatures at the two contact points and need not be equal. A channel event  $E_{j,e}(d)$  consists of the  $j+1$  impacts alternating between the specified facing patches in the time window  $t = j g_e + d$ . These patches are fixed around the reference contacts before observation and define a measurable part of the physical record. The localization theorem below proves that the relevant pairs continue smoothly and that a common onset collar is available.

Parameter derivatives are multi-index derivatives in a fixed Euclidean parameter chart. Unless a statement explicitly says otherwise, every geometric derivative is taken in the *channel-centered* coordinates  $(\xi, d)$ , with

$$t = j g_e(\xi) + d, \quad (12.2)$$

so the offset  $d$  is held fixed while the physical time moves with the channel onset. In particular the mixed geometric derivatives in Theorem 12.1 are not derivatives at fixed physical time. This convention also applies to the parameter derivatives of  $S_b, B_b, \mathcal{F}_p$  and to the uniform constants below. Matrix norms are Euclidean operator norms unless a trace norm is specified.

**12.1. Local multiplier convention.** Suppress the channel index for the local hyperbolic calculations and set

$$\begin{aligned} \varrho &= e^{-\gamma}, & \lambda &= \varrho^2 = e^{-2\gamma}, \\ \mathfrak{r}_b &= \sqrt{c_b / c_{1-b}}, & r_b &= \mathfrak{r}_b \varrho \quad (b = 0, 1). \end{aligned} \quad (12.3)$$

The number  $\mathfrak{r}_b$  is a curvature ratio, whereas  $r_b$  is the derivative at zero of the physical first-hit map  $\varphi_b$ . Thus  $\mathfrak{r}_0 \mathfrak{r}_1 = 1$ ,  $r_0 r_1 = \lambda$ , and  $R_b = \varphi_{1-b} \circ \varphi_b$  has multiplier  $\lambda$ . The integer  $j$  counts

flights;  $n$  iterations of  $R_b$  mean  $j = 2n$ . The symbol  $\rho$  in exponential estimates denotes a freely chosen strict decay margin, not the exact factor  $\rho$ .

**12.2. The nonlinear relative law.** Fix an oriented channel  $e$  and abbreviate  $g = g_e$ ,  $\gamma = \gamma_e$ . Use transverse endpoint coordinates  $u, v$  at its two contacts. Let  $W_j(u, v)$  be the stationary length of its  $j$  alternating flights,  $E_j = W_j - jg$ ,  $d_j^0 = -W_{j,uv}(0, 0)$ , and  $b_j = -W_{j,uv}/d_j^0$ . For  $p = j \bmod 2$  put

$$a_b = \frac{\sqrt{c_0 c_1} \sinh \gamma}{g c_{1-b}}, \quad q_p = \sqrt{a_0 a_p}, \quad d_j^0 = \frac{q_p}{\sinh(j\gamma)}.$$

**Theorem 12.1** (Relative boundary law). *There are positive smooth functions  $B_b$  and strictly convex smooth functions  $S_b$ ,  $b = 0, 1$ , with*

$$S_b(0) = S'_b(0) = 0, \quad S''_b(0) = a_b, \quad B_b(0) = 1,$$

*a common endpoint box, a number  $d_* > 0$ , and  $0 < \tau < 1$  such that, for every fixed derivative order  $k$  and every  $j \geq 1$ ,*

$$\|E_j - S_0 \oplus S_p\|_{C^k} + \|b_j - B_0 \otimes B_p\|_{C^k} \leq C_k \tau^j, \quad p = j \bmod 2. \quad (12.4)$$

*The action difference vanishes to second order at the origin. Define*

$$\mathcal{F}_p(d) = \frac{q_p}{\pi d^2} \int (d - S_0(u) - S_p(v))_+ B_0(u) B_p(v) du dv. \quad (12.5)$$

*Then  $\mathcal{F}_p$  extends smoothly to  $[0, d_*]$ , is positive there, and has value one at zero. The exact selected-channel probabilities obey*

$$\left\| \frac{2A \sinh(j\gamma)}{d^2} \Pr(E_{j,e}(d)) - \mathcal{F}_p(d) \right\|_{C^k([0, d_*])} \leq C_k \tau^j. \quad (12.6)$$

*The norm includes every fixed mixed geometric and offset derivative in the centered convention (12.2). The limiting conditional endpoint-residual law at  $d > 0$  has density proportional to*

$$B_0(u) B_p(v) \mathbf{1}_{\{r > 0, S_0(u) + S_p(v) + r < d\}}.$$

*All assertions hold uniformly on the compact local families specified above. For analytic families the endpoint constructions are analytic.*

*Proof and dependencies.* Lemma 15.2 constructs the finite bridge and its relative determinant. Lemma 17.1 constructs  $S_b, B_b$  with the normalizations above. Theorem 17.2 proves (12.4) by comparing two separated perturbation blocks in trace norm, including mixed derivatives. Theorem 17.3 proves (12.6) using the common Morse domain of Lemma 16.1 and identifies the conditional density. The full proofs occur in Part I.  $\square$

### 13. UNIFORM THRESHOLDS AND CONDITIONAL CONTACT METRICS

**Theorem 13.1** (Uniform thresholds for periodic dispersing configurations). *There is  $\varepsilon_0 > 0$ , independent of every integer  $j \geq 1$ , such that  $\varepsilon_0 < \min_{K_0} g_*/4$  and, for  $0 \leq t - jg_* < \varepsilon_0$ ,*

$$\Pr_{\xi}(N_t \geq j+1) = \Pr_{\xi}(N_t = j+1) = \sum_{e \in \mathcal{E}} \frac{d_{j,e,+}^2}{2A(\xi) \sinh(j\gamma_e(\xi))} [1 + d_{j,e,+} \mathcal{R}_{j,e}(\xi, d_{j,e,+})]. \quad (13.1)$$

*Every fixed mixed derivative of  $\mathcal{R}_{j,e}$  is bounded uniformly in  $j$ ; it has a smooth right extension at offset zero. For  $t \leq jg_*$  the probability is zero. At the selected channel's own window  $t = jg_e + d$ ,  $0 < d < \varepsilon_0$ , the same individual summand is exactly  $\Pr_{\xi}(E_{j,e}(d))$ , whether or not  $e$  minimizes the gap.*

The exponent two is independent of  $j$ , but its coefficient is exponentially small. The essential estimate is relative to that coefficient. The proof uses a weighted Dirichlet construction and a trace-norm estimate for its relative determinant, rather than a long initial-value iteration. The local hypotheses for this construction are verified at every reference configuration above. In the triangular one-obstacle near-circle class,  $\mathcal{E}$  consists of the six unit lattice directions, so (13.1) is its specialization, with all competing onsets.

For the next theorem, lift the first and last collision positions to the fixed contact patches in  $\mathbb{R}^2$ , and denote them by  $Q_0, Q_j$ . Define the two scalar conditional variances

$$V_{j,b}(d) = \frac{1}{d} \left\{ \mathbb{E}(|Q_b|^2 \mid E_{j,e}(d)) - |\mathbb{E}(Q_b \mid E_{j,e}(d))|^2 \right\}, \quad b = 0, j. \quad (13.2)$$

These are translation- and rotation-invariant within the chosen lifts. Their definition requires neither the unknown contact point nor its tangent or curvature. The gap  $g_e$  is identified by its channel onset; quantitative timing calibration is a separate experiment below.

**Theorem 13.2** (Uniform two-variance curvature recovery). *For every odd  $j \geq 1$  the right limits  $v_{j,0}, v_{j,1}$  of  $V_{j,0}(d), V_{j,j}(d)$  exist, and*

$$(v_{j,0}, v_{j,1}) = \frac{g_e \coth(j\gamma_e)}{3 \sinh \gamma_e} \left( \sqrt{\frac{c_{e,1}}{c_{e,0}}}, \sqrt{\frac{c_{e,0}}{c_{e,1}}} \right). \quad (13.3)$$

*Each difference  $V_{j,b}(d) - v_{j,b}$  is  $d$  times a smooth function with uniform fixed-order derivative bounds. At known  $g_e$ , these two limits determine the ordered pair  $(\kappa_{e,0}, \kappa_{e,1})$  without knowing  $A$ . The inverse is given by*

$$\begin{aligned} \frac{\coth(j\gamma)}{\sinh \gamma} &= \frac{3}{g} \sqrt{v_0 v_1}, & r &= \sqrt{v_1/v_0}, \\ \kappa_0 &= \frac{r \cosh \gamma - 1}{g}, & \kappa_1 &= \frac{r^{-1} \cosh \gamma - 1}{g}. \end{aligned} \quad (13.4)$$

*On every compact positive gap–curvature box this observation map and its inverse are Lipschitz, with constants independent of odd  $j$ . They remain uniformly conditioned at  $\kappa_0 = \kappa_1$ . If both measured variances have error at most  $\delta$ , a reconstruction constrained to such a box has curvature error at most  $C(d + \delta)$ .*

Thus the full conditional metric retains information lost by the scalar flux coefficient. Theorem A.18.3 realizes the scalar ambiguity by actual smooth periodic obstacles with the same gap, area and every leading count amplitude, while their endpoint curvatures vary. The two variances distinguish that family. The statement is local contact-curvature recovery, not global shape rigidity. Conditioning on an exponentially rare event is not free: Proposition A.18.2 records the number of independent physical preparations required.

#### 14. SHORTEST CHANNELS IN ARBITRARY PERIODIC CONFIGURATIONS

The relevant localization is a consequence of separated strict convexity, not proximity to a particular model shape.

**Lemma 14.1** (A geometric criterion for the complete event). *At any reference configuration of Section 12, the minimum distance between distinct lifted obstacles is positive and is attained by finitely many oriented pairs modulo common translation. Every such pair has a unique closest segment, positive-definite distance Hessian and positive clearance from the remaining obstacles. Its contact points, length and curvatures continue smoothly under small smooth deformations. There are fixed contact and direction neighborhoods and  $\varepsilon_0 > 0$  such that every physical sequence*

with  $j$  complete flights and total length at most  $jg_* + \varepsilon_0$  alternates in one of these neighborhoods. The constants are independent of  $j$ .

*Proof.* Choose bounded representatives of the finitely many obstacles. The distance between  $D_a$  and  $D_b + \lambda$  tends to infinity with  $|\lambda|$ , uniformly in  $a, b$ . Only finitely many lattice translates can have distance below any fixed bound. Each of those finitely many positive distances is attained; disjoint closures make their minimum positive. Among pairs not attaining this minimum, the next distance is separated from it by a positive amount. Pairs outside the finite list already exceed a larger fixed bound. These statements persist under a small Hausdorff perturbation, with a strict gap to all pairs outside  $\mathcal{E}$ .

The closest displacement vector between two disjoint convex bodies is unique: if two different displacement vectors both minimized its norm, their midpoint would have smaller norm and would still be a displacement. For that vector, each endpoint is a support point in its normal direction. Strict convexity makes those support points unique. Rotate the segment to  $[0, g]$  and use a common transverse coordinate. The facing graphs are  $(-\psi_0(u), u)$  and  $(g + \psi_1(v), v)$ , with  $\psi_b(0) = \psi'_b(0) = 0$  and  $\psi''_b(0) = \kappa_b > 0$ . The distance Hessian at its minimum is

$$\frac{1}{g} \begin{pmatrix} 1 + g\kappa_0 & -1 \\ -1 & 1 + g\kappa_1 \end{pmatrix}. \quad (14.1)$$

Its determinant and its diagonal entries are positive. The implicit function theorem gives smooth contact points and distance. Compactness keeps both eigenvalues uniformly positive and all chosen charts uniform.

A reference minimum segment cannot meet a third obstacle in its open interior: the initial part would give a smaller distance from its first obstacle to that third obstacle. The closed segment also stays a positive distance from every other obstacle, since its endpoints belong only to its own two disjoint bodies and only finitely many other bodies are nearby. Small contact patches and small geometric perturbations preserve this clearance. Strict convexity makes each connecting chord leave the first obstacle and enter the second transversely, with no earlier hit. Thus these local chords are physical first hits.

At a fixed outgoing obstacle, distinct minimizing channels give distinct normal outgoing collision states. If two shared a contact and outgoing normal, their other contacts at the same distance would coincide, contradicting disjointness. The finite collection of states is therefore separated. Choose their neighborhoods so small that normal reflection of a flight arriving along  $e$  can belong to a short outgoing channel only in the neighborhood of  $\bar{e}$ . This property is open and uniform.

Every physical free flight joins distinct lifted obstacles: a straight ray leaving a convex body cannot return to it before another hit. Hence its length is at least  $g_*$ . If a sum of  $j$  such lengths is at most  $jg_* + \varepsilon_0$ , every summand is at most  $g_* + \varepsilon_0$ . The gap to the other pairs restricts that summand to  $\mathcal{E}$ . Away from the unique contact patches the distance to that pair's minimum is bounded below; decrease  $\varepsilon_0$  to exclude this complement. Each short flight is then nearly normal. The separated outgoing states and specular reflection force every following short flight to be the reverse channel. Induction uses the same neighborhoods at each impact and introduces no length-dependent constant. Grazing states and their regular preimages have phase measure zero.  $\square$

The lemma also applies to a specified nonminimal closest chord with positive clearance, for its selected local itinerary. It need not then exhaust the whole count event. All original six near-circle channels satisfy the criterion. For example, the criterion also applies to arbitrarily eccentric ellipses and to noncongruent obstacles on any fixed lattice, provided the indicated disjointness and positive curvature conditions hold; there is no small eccentricity assumption.

For one channel the exact one-flight actions are

$$\ell_b(u, v) = \sqrt{(g + \psi_b(u) + \psi_{1-b}(v))^2 + (v - u)^2}, \quad b = 0, 1. \quad (14.2)$$

Uniformly on a small fixed box,

$$\ell_b(u, v) = g + \frac{c_b u^2 + c_{1-b} v^2 - 2uv}{2g} + O((|u| + |v|)^3), \quad (14.3)$$

$$\ell_b(u, v) - g \geq \psi_b(u) + \psi_{1-b}(v) \geq c_{\text{loc}}(u^2 + v^2). \quad (14.4)$$

The full broken-action Hessian has strictly positive diagonal-dominance margins:  $\kappa_b$  in endpoint rows and  $2\kappa_b$  in interior rows at zero. Every row perturbation involves only neighboring coordinates. The same small box therefore gives strict convexity at every length.

### 15. A PARAMETER-DEPENDENT BOUNDARY-VALUE AND RELATIVE-FLUX LEMMA

Let  $W_{j,e}(\xi, u, v)$  be the stationary value of the sum of (14.2), with  $y_0 = u, y_j = v$ . We construct that value on one common neighborhood. Endpoint Hessians are denoted by  $\mathbf{H}$ ; scalar relative remainders are denoted by  $\mathcal{R}$ .

**Lemma 15.1** (Alternating Jacobi reduction). *Put  $c = \sqrt{c_0 c_1}$ ,  $\gamma = \text{arcosh } c$ ,  $\sigma_i = \sqrt{c_{1-(i \bmod 2)}}$ , and  $D_j = \text{diag}(\sigma_0, \sigma_j)$ . The quadratic endpoint Hessian is*

$$\mathbf{H}_{j,e} = \frac{c \sinh \gamma}{g} D_j^{-1} \begin{pmatrix} \coth(j\gamma) & -\text{csch}(j\gamma) \\ -\text{csch}(j\gamma) & \coth(j\gamma) \end{pmatrix} D_j^{-1}. \quad (15.1)$$

Consequently

$$d_{j,e}^0 := -W_{uv}(\xi, 0, 0) = \frac{c \sinh \gamma}{g \sigma_0 \sigma_j \sinh(j\gamma)}, \quad \sqrt{\det \mathbf{H}_{j,e}} = \frac{c \sinh \gamma}{g \sigma_0 \sigma_j}, \quad \frac{d_{j,e}^0}{\sqrt{\det \mathbf{H}_{j,e}}} = \text{csch}(j\gamma). \quad (15.2)$$

The endpoint eigenvalues have positive uniform upper and lower bounds. The period-two return multiplier is  $e^{2\gamma}$ .

*Proof.* The interior equation is  $y_{i+1} - 2c_{i \bmod 2} y_i + y_{i-1} = 0$ . Set  $y_i = \sigma_i z_i$ . Since  $c_{i \bmod 2} \sigma_i / \sigma_{i+1} = c$ , the recurrence becomes  $z_{i+1} - 2c z_i + z_{i-1} = 0$  and the whole quadratic action becomes

$$\frac{c}{2g} \sum_{i=0}^{j-1} [c(z_i^2 + z_{i+1}^2) - 2z_i z_{i+1}].$$

Its endpoint solution is  $z_i = \sinh((j-i)\gamma) z_0 / \sinh(j\gamma) + \sinh(i\gamma) z_j / \sinh(j\gamma)$ . The first endpoint derivative is  $c(c z_0 - z_1)/g$ . The identity  $c \sinh(j\gamma) - \sinh((j-1)\gamma) = \sinh \gamma \cosh(j\gamma)$ , and its reversed version, give (15.1). Now use  $\coth^2 - \text{csch}^2 = 1$ . The middle matrix has eigenvalues  $\tanh(j\gamma/2)$  and  $\coth(j\gamma/2)$ . Since  $\gamma \geq \gamma_* > 0$  and the  $\sigma_i$  have uniform positive bounds, this proves the eigenvalue bounds. The two-periodic scaling conjugates the two-step recurrence to the square of  $\begin{pmatrix} 2c & -1 \\ 1 & 0 \end{pmatrix}$ . The nonzero one-flight twist  $1/g$  identifies position-pair and collision position-momentum coordinates, so its unstable eigenvalue is the stated physical return multiplier. In particular the formula at  $j = 1$  is exactly (14.1).  $\square$

**Lemma 15.2** (Uniform nonlinear bridge and relative determinant). *Suppose  $g$  and the two graph functions in (14.2) range over a compact smooth parameter family, with  $g, \kappa_0, \kappa_1$  bounded above and bounded away from zero. There is a common endpoint neighborhood on which the*



stationary segment exists uniquely and is smooth. For some fixed  $\rho < 1$ , with  $a = |u| + |v|$  and  $w_i = \rho^i + \rho^{j-i}$ ,

$$y_i^{\text{lin}} = \sigma_i \left[ \frac{\sinh((j-i)\gamma)}{\sinh(j\gamma)} \frac{u}{\sigma_0} + \frac{\sinh(i\gamma)}{\sinh(j\gamma)} \frac{v}{\sigma_j} \right], \quad (15.3)$$

$$|y_i| \leq Caw_i, \quad |y_i - y_i^{\text{lin}}| \leq Ca^2 w_i,$$

$$W_{j,e} - jg = \frac{1}{2}(u, v)H_{j,e}(u, v)^t + O(a^3),$$

$$-W_{uv}(\xi, u, v) = d_{j,e}^0(\xi)b_{j,e}(\xi, u, v), \quad b_{j,e}(\xi, 0, 0) = 1, \quad b_{j,e} = 1 + O(a) > 0. \quad (15.4)$$

Every fixed mixed derivative of  $y_i$  and  $y_i - y_i^{\text{lin}}$  is bounded by  $C_k w_i$  times respectively  $a^{\max(1-|\beta|, 0)}$  and  $a^{\max(2-|\beta|, 0)}$ , where  $\beta$  counts endpoint derivatives. The action remainder and  $b_{j,e}$  have uniform fixed-order bounds. If both graphs are even, the bridge correction is  $O(a^3 w_i)$ , the action remainder is  $O(a^4)$ , and  $b_{j,e} - 1 = O(a^2)$  is even. In an analytic family these bounds hold on a common complex neighborhood and are analytic bounds.

*Proof.* Let  $H_{\text{int}}^0$  be the quadratic interior Hessian. Its inverse, for  $1 \leq i \leq k < j$ , is

$$G_{ik} = \frac{g\sigma_i\sigma_k}{c} \frac{\sinh(i\gamma)\sinh((j-k)\gamma)}{\sinh\gamma\sinh(j\gamma)}, \quad G_{ki} = G_{ik}. \quad (15.5)$$

Substitution verifies the homogeneous recurrence away from  $i = k$ , zero endpoint values and the unit jump at  $i = k$ . In particular  $|G_{ik}| \leq Ce^{-\gamma_*|i-k|}$ . Choose  $e^{-\gamma_*} < \rho_0 < \rho < 1$ . A geometric-series sum yields

$$\sum_k \rho_0^{|i-k|} (\rho^k + \rho^{j-k}) \leq C_\rho w_i. \quad (15.6)$$

For its  $\rho^k$  part, split at  $k = i$  and sum the ratios  $\rho_0/\rho$  and  $\rho_0\rho$ ; the other part follows by reversal. Thus  $G$  is bounded independently of  $j$  on the norm  $\|z\|_w = \max_i |z_i|/w_i$ .

Parameter differentiation of (15.5) has the same property. To see this without a hidden growing dimension, write the hyperbolic quotient as a bounded factor times  $e^{-\gamma(k-i)}$ , using factors  $1 - e^{-2r\gamma}$  for its distances to the endpoints. Derivatives of those factors produce  $r^m e^{-2r\gamma}$ , uniformly bounded for each fixed  $m$ ; derivatives of the first exponential produce  $|i - k|^m e^{-\gamma|i-k|} \leq C_m \rho_0^{|i-k|}$ . All derivatives of the positive  $\sigma_i$  are bounded. The same argument for the endpoint solution gives  $|\partial_\xi^\alpha y_i^{\text{lin}}| \leq C_\alpha a w_i$ . Consequently every fixed parameter derivative of  $G$  is bounded in the weighted norm, with no  $j$  factor.

The interior gradient is  $H_{\text{int}}^0(y - y^{\text{lin}}) + \mathcal{N}_\xi(y)$ , including the fixed endpoint coordinates in the local remainder. Its  $i$ th component depends on only three neighboring coordinates and vanishes quadratically there. Adjacent weights have bounded ratios. On  $\|y\|_w \leq C_0 a$ , its weighted norm is  $O(a^2)$  and its derivative has norm  $O(a)$ . The map  $y \mapsto y^{\text{lin}} - G\mathcal{N}_\xi(y)$  preserves this ball and has Lipschitz constant at most  $1/2$  when  $a$  is uniformly small. This proves existence and the first estimates. The diagonal-dominance bound after (14.4) gives uniqueness throughout the contact box, not just inside the weighted ball.

For clarity, the derivative induction takes place in these same weighted spaces. One derivative of the fixed-point equation gives

$$(I + GD_y \mathcal{N}_\xi) Dy = Dy^{\text{lin}} - (DG)\mathcal{N}_\xi - GD_{\text{explicit}}\mathcal{N}_\xi.$$

Its left inverse has norm at most two. At order  $k$ , the only term containing an order- $k$  derivative of  $y$  is the same left operator; every other term contains derivatives of  $G$ , explicit parameter or endpoint derivatives of the local remainder, and lower derivatives of  $y$ . The local chain rule and bounded neighboring-weight ratios bound their products in  $\|\cdot\|_w$ . Induction proves the estimates at order  $k$ . Parameter derivatives preserve vanishing at zero; Taylor's integral formula in  $(u, v)$

therefore gives precisely the powers of  $a$  stated in the lemma, including the correction of order two.

For every  $p > 0$ ,  $\sum_i w_i^p \leq C_p$ . Sum the cubic action remainders using this inequality. The linear bridge is stationary for the quadratic action with endpoints fixed, so its change on substituting the nonlinear bridge is  $O(a^4)$ . This proves the action estimate. The same summability and the derivative induction prove all its fixed-order bounds.

For the relative estimate put  $b_i = -\ell_{i \bmod 2, uv}(y_i, y_{i+1})$  and  $H_{\text{int}} = H_{\text{int}}^0 + E$ . Schur complementation gives

$$-W_{uv} = \frac{\prod_{i=0}^{j-1} b_i}{\det H_{\text{int}}}. \quad (15.7)$$

Indeed, the endpoint-to-interior entries are  $-b_0, -b_{j-1}$  and the corner entry of the inverse tridiagonal interior matrix is  $b_1 \cdots b_{j-2} / \det H_{\text{int}}$ . At  $j = 1$  the determinant is empty and equals one. The local expansions and endpoint weights give

$$\sum_i |\log(gb_i)| \leq Ca, \quad \|E\|_1 \leq \sum_{i,k} |E_{ik}| \leq Ca, \quad \|(H_{\text{int}}^0)^{-1}\| \leq C.$$

The trace-norm bound uses tridiagonality; each matrix unit has trace norm one. Take the endpoint neighborhood so small that  $\|(H_{\text{int}}^0)^{-1}E\| < 1/2$ . Then

$$\log \frac{\det H_{\text{int}}}{\det H_{\text{int}}^0} = \sum_{m \geq 1} \frac{(-1)^{m+1}}{m} \text{tr}[(H_{\text{int}}^0)^{-1}E]^m = O(a). \quad (15.8)$$

The absolute trace is bounded by one trace norm times  $m - 1$  operator norms. Summing therefore costs  $C\|E\|_1$ , not the matrix dimension. Divide (15.7) by its value at zero and subtract logarithms to obtain (15.4).

Every differentiated  $E$  has uniformly bounded trace norm by the already proved coordinate bounds. The quadratic Hessian and each of its parameter derivatives are uniformly banded with bounded entries; hence derivatives of its inverse have bounded operator norm, by  $\partial H^{-1} = -H^{-1}(\partial H)H^{-1}$  and induction. Termwise differentiation of (15.8) is justified by geometric convergence even after its fixed polynomial factors in  $m$ . Every term has at least one trace-class factor. The differentiated sum of  $\log(gb_i)$  is controlled by the same endpoint weights. This proves all derivative bounds for the logarithmic ratio and its exponential. Positive twists and a positive interior determinant give positivity.

If the graphs are even, the action expansion is even, its nonlinear gradient is cubic, and its Hessian and twist errors are quadratic. The preceding estimates then give the improved orders. Uniqueness implies simultaneous reversal of every coordinate, proving evenness. For a jointly analytic family all local functions extend to a common small complex neighborhood of the fixed compact real parameter set. Shrink it to retain the strict exponential and inverse-norm margins. The same contraction is uniform there. The determinant series remains convergent, and Cauchy estimates on a smaller neighborhood prove the analytic assertion. No assertion about analytic extension is used for a merely smooth family.  $\square$

## 16. FULL-PHASE INTEGRATION AND THE THRESHOLD LAW

Let  $P(\xi)$  be the sum of obstacle perimeters in one cell. The collision section measure is  $ds \, dp / (2P)$ , the mean roof  $\pi A / P$ , and the stationary collision intensity  $P / (\pi A)$ . These identities follow from a flight-tube change of variables: irrational torus lines hit an open obstacle, rational directions form a null angular set, and the positive minimal gap excludes finite-time accumulation.

In particular, in first-impact endpoint coordinates the full phase measure is

$$\frac{-W_{uv}}{2\pi A(\xi)} du dv dr. \quad (16.1)$$

Indeed, first variation gives  $p = -W_u/(ds/du)$ ; the arclength factor cancels in  $|ds dp| = |W_{uv}| du dv$ . The positive relative twist makes the endpoint map injective at fixed  $u$ . The localized stationary path is physical by Lemma 14.1. Thus (16.1) is not a substituted transverse ensemble. The total angle–boundary flux is  $2P$ , and its flight-tube integral is  $2\pi A$ ; this also proves the normalization without an upper roof bound.

**Lemma 16.1** (Uniform smooth radial integration). *There are smooth Morse inverses  $\Phi_{j,e}(\xi, w)$  on a common ball, with all fixed derivatives bounded uniformly in  $j$ , such that*

$$W_{j,e}(\xi, \Phi_{j,e}(\xi, w)) - jg_e = |w|^2/2, \quad D_w \Phi_{j,e}(\xi, 0) = H_{j,e}^{-1/2}.$$

The phase density (16.1) in  $(w, r)$  is

$$\frac{a_{j,e}(\xi, w)}{2\pi A(\xi) \sinh(j\gamma_e)} dw dr, \quad a_{j,e}(\xi, 0) = 1, \quad a_{j,e} = 1 + O(|w|) > 0, \quad (16.2)$$

with the same derivative bounds. If  $a$  is any such uniformly smooth amplitude, then

$$\frac{2}{\pi} \int_{|z|<1} (1 - |z|^2) a(\xi, \sqrt{2d} z) dz = a(\xi, 0) + dH(\xi, d),$$

where  $H$  is smooth for  $d \geq 0$  and has uniform fixed-order bounds. No evenness of  $a$  or of the scatterer is required.

*Proof.* For  $E = W - jg$ , Taylor's integral formula writes  $E(y) = \frac{1}{2} y^t M(y) y$ , where  $M(y) = 2 \int_0^1 (1-s) D^2 E(sy) ds$  is uniformly positive near zero. The map  $y \mapsto M(y)^{1/2} y$  has uniformly invertible differential at zero and small nonlinear differential on a common ball. Its inverse is obtained by contraction, then differentiated to each finite order. This construction uses the uniform endpoint Hessian, not the exponentially small twist as an inverse bound. Inserting the Jacobian and (15.2) into (16.1) gives (16.2).

For the last assertion replace  $a(w)$  by  $(a(w) + a(-w))/2$  in the integral. Every odd Taylor term vanishes. Taylor expansion to any prescribed order, with a sufficiently high differentiated remainder, then yields an expansion in integer powers of  $d$ . One can also set  $r = \sqrt{d}$  first: the integral is a smooth even function of  $r$  with uniform bounds. Successive division of its even Taylor remainder by  $r^2$  shows that it is a smooth function of  $r^2$  on the right. Using two more endpoint derivatives for each division proves the uniform bounds for every fixed derivative of  $H$ .  $\square$

*Proof of Theorem 13.1.* In a channel's window  $jg_e + d$ , each of its  $j$  complete flights is at least  $g_e$ . Hence the allowed first residual time satisfies  $0 < r < d - (W - jg_e) < d < g_*$ . The preceding roof cannot truncate this interval, and the next roof cannot fit after the last hit. The first-impact phase formula (16.1) therefore applies to this selected channel at its own threshold: integrating the residual variable gives exactly that interval length. The coercivity (14.4) implies  $|y_0|^2 + |y_j|^2 + 2 \sum_{0 < i < j} |y_i|^2 \leq Cd$  on its sublevel. It places the entire active set inside the common Morse coordinates, so no chart-boundary term occurs. Its exact probability is

$$\frac{1}{2\pi A \sinh(j\gamma_e)} \int_{|w|^2 < 2d} (d - |w|^2/2) a_{j,e}(\xi, w) dw.$$

Since the quadratic integral is  $\pi d^2$ , Lemma 16.1 proves the channel assertion, including its derivative bounds.

For the complete event let  $t = jg_* + \varepsilon$  with  $0 < \varepsilon < \varepsilon_0$ . No  $j + 2$  hits are possible, since  $t < (j + 1)g_*$ . Lemma 14.1 partitions every contributing  $j$ -flight segment into exactly one of the finitely many oriented channels. Its available excess there is  $d_{j,e} = \varepsilon - j(g_e - g_*) \leq \varepsilon$ . Channels with nonpositive excess contribute zero. Summing the exact channel formulas gives (13.1). The minimum-roof bound proves the negative-side assertion.  $\square$

*Proof of Theorem A.26.1.* Write the mark on a stationary bridge as  $w_{j,e} + V_{j,e}$ , subtracting the normal value separately at each impact. Chord directions and normal reflections are smooth maps of neighboring transverse coordinates, with uniformly positive chord denominators. Lemma 15.2 therefore gives  $|V_{j,e}| \leq C(|u| + |v|)$  and uniform fixed-order bounds for all its derivatives. The normal states now depend on  $\xi$ , but each summand's deviation vanishes at  $(u, v) = 0$  for every  $\xi$ ; parameter differentiation preserves its endpoint localization. The finite chain rule and  $\sum_i (\rho^i + \rho^{j-i}) \leq C$  prove the bounds. Multiplying the Morse amplitude by  $e^{qV_{j,e}}$  and applying Lemma 16.1 proves (A.26.1).

For the convergence assertion, exact parity counting gives

$$w_{j,e} = (j + 1)\bar{f}_e + \mathbf{1}_{\{j \text{ even}\}}(a_{e,0} - a_{e,1})/2. \quad (16.3)$$

For fixed real  $q$  the parity factor is bounded above and below by positive constants, as is the normalized integral of  $a_{j,e}e^{qV}$  on a sufficiently small collar. Also  $\text{csch}(j\gamma_e) \asymp 2e^{-j\gamma_e}$ , uniformly in  $j$ . Thus a channel's normalized term is comparable to  $e^{j(q\bar{f}_e - \gamma_e)}$ . Positivity rules out cancellation: equality in (A.26.2) gives a nondecaying term, and a positive exponent gives a growing term. This proves necessity as well as sufficiency, including the right limit at zero. On a compact subset of the strict domain there is a common margin  $\delta > 0$ . Parameter and source derivatives of the leading factor cost only a polynomial in  $j$ ; offset derivatives act on the uniform smooth remainder. Every fixed derivative is bounded by  $C(1 + j)^m e^{-\delta j}$  and is summable. The vector-source assertion is the same argument with scalar products.  $\square$

**Corollary 16.2** (Moving threshold interfaces). *Within the collar of (13.1), the probability is a  $C^1$  function of the smooth physical variables  $(\xi, t)$ , with smooth pieces separated by  $t = jg_e(\xi)$ . For an orientation  $e$ , along a smooth scalar path  $s \mapsto (\xi(s), t(s))$  entering  $d_{j,e} > 0$  transversely, its contribution has second-derivative jump from the inactive to the active side*

$$\frac{(t'(s) - j \partial_s g_e)^2}{A(\xi) \sinh(j\gamma_e)}. \quad (16.4)$$

*Coincident contributions, including the reversed orientation, add with their respective activation signs. The marked formula has the additional factor  $e^{qw_{j,e}}$ . One-sided and distributional higher derivatives follow by differentiating the positive-part expression, with polynomial-in- $j$  coefficient bounds.*

*Proof.* Each summand is  $d_+^2[C(\xi) + d_+C(\xi)H(\xi, d_+)]$ . It and its first derivatives match at zero. Only  $2C(\partial_s d)^2$  survives in the second-derivative jump. All other terms contain a factor  $d$ . Finite addition proves the claim at intersections. For higher derivatives the ordinary product and chain rules on each side, and the derivatives of  $d_+^2$  in the distributional sense, give the stated formulas; no differentiation of  $g_*$  is used.  $\square$

In the triangular one-obstacle specialization, fix  $R_* \in (0, 1/2)$ . For a concrete change of minimizing mechanism, take  $h_s(\theta) = R_* + s \cos(2\theta)$ . Differentiating the closest distance at the circular normal chord gives  $g'_e(0) = -2 \cos(2 \arg e)$ : the endpoint variation is the normal support variation, since the unperturbed endpoint gradient is zero. The horizontal pair has derivative  $-2$  and the other two pairs derivative 1. Reflection symmetry keeps those two lengths equal. Consequently the horizontal pair minimizes for small positive  $s$ , while the other two minimize

for small negative  $s$ . Formula (13.1) continues through  $s = 0$ . A nonminimal pair contributes exactly while  $j(g_e - g_*) < \varepsilon$ ; the parameter transition scale is therefore  $\varepsilon/j$  at a transverse length crossing. This assertion is about activation of physical channels, not a new family of symbolic itineraries.

For the conditional record, put  $w = \sqrt{2d}z$  and  $r = d\eta$ . The fixed latent domain and its probability measure are

$$\mathcal{U} = \{(z, \eta) : |z| < 1, 0 < \eta < 1 - |z|^2\}, \quad \mathrm{d}\mathbf{m} = \frac{2}{\pi} \mathrm{d}z \mathrm{d}\eta. \quad (16.5)$$

This notation will also be used for the general record response.

## 17. TWO NONLINEAR BOUNDARY LAYERS

Fix one oriented channel and suppress its index and the smooth geometric parameter. All bounds below include each fixed number of parameter derivatives on the compact families specified in Section 12. Shrinking the common contact box is allowed once, independently of the flight number. The functions  $E_j = W_j - jg$  and  $b_j = -W_{j,uv}/d_j^0$  refer to the actual nonlinear stationary action and its normalized twist, not their quadratic models.

**17.1. The half-line construction.** For a starting type  $b \in \{0, 1\}$  put  $\sigma_i^{(b)} = \sqrt{c_{1-((b+i) \bmod 2)}}$ . The quadratic half-line interior Hessian acts on sequences indexed by  $i \geq 1$ , with fixed zero boundary at  $i = 0$ . Its diagonal entries are  $2c_{(b+i) \bmod 2}/g$  and its adjacent entries are  $-1/g$ . Its inverse is

$$G_{ik}^{(b)} = \frac{g\sigma_i^{(b)}\sigma_k^{(b)}}{2c \sinh \gamma} \left( e^{-\gamma|i-k|} - e^{-\gamma(i+k)} \right), \quad i, k \geq 1. \quad (17.1)$$

The formula follows either from the recurrence and its unit jump or from (15.5) by sending the right endpoint to infinity. It is a bounded positive operator on  $\ell^2(\mathbb{N})$ , and is bounded on  $\sup_i \rho^{-i}|x_i|$  whenever  $e^{-\gamma_-} < \rho < 1$ . Every fixed parameter derivative obeys the same assertion after leaving a strict exponential margin.

**Lemma 17.1** (Semi-infinite stationary segments). *For  $|u| < u_*$  there is a unique small stationary half-line segment  $x^{(b)}(u) = (x_i^{(b)}(u))_{i \geq 0}$  with  $x_0^{(b)} = u$  and  $x_i^{(b)} \rightarrow 0$ . It is smooth, and on a common neighborhood*

$$x_i^{(b)}(u) = \frac{\sigma_i^{(b)}}{\sigma_0^{(b)}} e^{-i\gamma} u + O(u^2 \rho^i), \quad (17.2)$$

$$|\partial_\xi^\alpha \partial_u^r x_i^{(b)}(u)| \leq C_{\alpha,r} \rho^i |u|^{\max(1-r,0)}.$$

The boundary action

$$S_b(u) = \sum_{i \geq 0} \left[ \ell_{(b+i) \bmod 2}(x_i^{(b)}(u), x_{i+1}^{(b)}(u)) - g \right] \quad (17.3)$$

is smooth and strictly convex near zero. It satisfies

$$S_b(0) = S'_b(0) = 0, \quad a_b := S''_b(0) = \frac{c \sinh \gamma}{g c_{1-b}} > 0. \quad (17.4)$$

These assertions are analytic on a common smaller neighborhood for jointly analytic graph families.

*Proof.* Use the half-line Green operator in the weighted contraction from Lemma 15.2. The linear solution is the first term of (17.2). The nonlinear gradient is local, vanishes quadratically, and its derivative has norm  $O(|u|)$  on the weighted ball of radius  $C|u|$ . Thus the map has contraction constant at most one half on a common ball. Differentiating its fixed-point equation leaves

the same invertible operator on the highest derivative; every other term is a product of a local derivative and weighted lower ones. The convolution estimate used in (15.6) proves (17.2) by induction.

Uniqueness also holds among all sufficiently small stationary sequences tending to zero, not just among elements of that weighted ball. Their difference solves a homogeneous Jacobi equation whose averaged Hessian is strictly diagonally dominant. Its inverse on bounded sequences is bounded by the reciprocal dominance margin: write the Hessian as its positive diagonal times  $I - T$ , with  $\|T\|_\infty < 1$ . With zero left boundary the difference is zero. This argument applies to every fixed small contact box.

Every summand of (17.3) is  $O(u^2 \rho^{2i})$ . Its differentiated sums converge by the same weighted bounds. Interior stationarity cancels the first-variation terms, leaving only the left endpoint derivative; the boundary term at a truncation tends to zero. At quadratic order this derivative is  $(c_b u - x_1^{(b), \text{lin}})/g$ . Since  $c_b \sigma_0^{(b)}/\sigma_1^{(b)} = c$ , it equals  $a_b u$ . This proves (17.4), hence strict convexity. The uniform complex contraction and Cauchy estimates, as in Lemma 15.2, prove the analytic assertion.  $\square$

Let  $H_b(u)$  be the half-line interior Hessian evaluated on this segment, and write  $\Delta H_b(u) = H_b(u) - H_b(0)$ . It is tridiagonal. Locality and (17.2) imply

$$\|\Delta H_b(u)\|_1 \leq C|u|, \quad \|\partial_\xi^\alpha \partial_u^r \Delta H_b(u)\|_1 \leq C_{\alpha, r}. \quad (17.5)$$

Here and below  $\|\cdot\|_1$  on operators is the trace norm. Indeed, the sum of absolute entries bounds that norm, and the entries are bounded by sums of neighboring endpoint weights. Define

$$U_b(u) = \sum_{i \geq 0} \log \left( g[-\ell_{(b+i) \bmod 2, uv}(x_i^{(b)}, x_{i+1}^{(b)})] \right), \quad (17.6)$$

$$B_b(u) = \exp \left\{ U_b(u) - \sum_{m \geq 1} \frac{(-1)^{m+1}}{m} \text{tr}[(G^{(b)} \Delta H_b(u))^m] \right\}. \quad (17.7)$$

This also defines the Fredholm determinant in the denominator of  $B_b$ : no additional determinant normalization is implicit. The series converges absolutely after requiring  $\|G^{(b)} \Delta H_b\| < 1/2$ . Each trace is bounded by one trace norm times  $m-1$  operator norms. The edge sum converges by endpoint localization. Consequently  $B_b > 0$ ,  $B_b(0) = 1$ , and every fixed mixed derivative is uniformly bounded. Differentiation of the trace series retains a trace-class factor, so its bound does not involve the number of sites of a truncation.

## 17.2. Relative factorization.

**Theorem 17.2** (Nonlinear two-boundary factorization). *There are a common endpoint box and  $\tau \in (0, 1)$  such that for every fixed  $k$ ,*

$$\|E_j - S_0 \oplus S_p\|_{C^k} + \|b_j - B_0 \otimes B_p\|_{C^k} \leq C_k \tau^j, \quad p = j \bmod 2. \quad (17.8)$$

*The norms include mixed geometric-parameter and endpoint derivatives;  $S_0 \oplus S_p$  means  $S_0(u) + S_p(v)$  and  $B_0 \otimes B_p$  means  $B_0(u)B_p(v)$ . The action difference has zero value and first endpoint differential at zero, and therefore is  $O(\tau^j(|u| + |v|)^2)$ . The amplitude difference vanishes at zero. The fixed-order constants, but not  $j$ , may depend on higher boundary norms. Both limits are analytic for analytic families.*

*Proof.* We give the gluing and determinant estimates separately. Reverse a half-line segment starting at type  $p$  at the right endpoint. Its  $i$ th coordinate in a length- $j$  indexing is  $x_{j-i}^{(p)}(v)$ ;

the reversal is consistent because the flight action is symmetric under reversal together with interchange of endpoint types. For  $0 < i < j$  put

$$z_i = x_i^{(0)}(u) + x_{j-i}^{(p)}(v), \quad z_0 = u, \quad z_j = v.$$

The endpoint overwrites cost  $O(\rho^j)$  in every fixed differentiated norm. In each interior stationarity equation the linear terms cancel. The remaining interaction is a sum of products of left and right coordinates or their neighboring values. Its size is bounded by  $C\rho^i\rho^{j-i} = C\rho^j$ . Thus, for the interior residual  $r(z)$ ,

$$\|\partial^\alpha r(z)\|_{\ell^1} \leq C_\alpha(1+j)\rho^j. \quad (17.9)$$

The assertion for derivatives follows from the local chain rule; parameter derivatives of the separated linear exponentials introduce only fixed powers of indices, absorbed by a strict exponential margin.

Let  $y$  be the finite stationary segment already constructed in Lemma 15.2. The identity

$$r(z) - r(y) = \left[ \int_0^1 H_{\text{int}}(y + t(z - y)) \, dt \right] (z - y)$$

has a symmetric strictly diagonally dominant matrix with a uniform positive margin. Its inverse has bounded  $\ell^\infty$  norm by diagonal splitting and a geometric series; symmetry gives the same  $\ell^1$  bound. It follows that  $\|y - z\|_{\ell^1} \leq C(1+j)\rho^j$ . Differentiating this identity leaves the same uniformly invertible operator on the highest difference derivative. All other terms contain lower difference derivatives and uniformly bounded local coefficients. Induction proves the corresponding  $\ell^1$  estimate at every fixed order, with a polynomial in  $j$  multiplying an exponential. Fixing a slower exponential absorbs every such fixed polynomial.

For the action, subtract the two half-line sums. The missing tails are exponentially small. On overlapping edges the difference between the value at the sum and the sum of the two values minus  $g$  is a local interaction bounded by  $C\rho^j$ . The replacement of  $z$  by  $y$  is controlled by its  $\ell^1$  error and bounded local derivatives. The same argument after differentiation proves the first term of (17.8). Both finite and limiting actions have zero value and gradient at zero; Taylor's integral remainder yields the additional quadratic vanishing.

For the normalized twist, the finite cofactor formula gives

$$\log b_j = \sum_{i=0}^{j-1} \log(g[-\ell_{i \bmod 2, uv}(y_i, y_{i+1})]) - \log \det(I + G_j \Delta H_j), \quad (17.10)$$

where  $G_j = (H_{\text{int}}^0)^{-1}$  and  $\Delta H_j = H_{\text{int}} - H_{\text{int}}^0$ . The edge sum converges exponentially to  $U_0(u) + U_p(v)$  by exactly the preceding local interaction and tail estimates. It remains to control the determinant relatively.

Here are the needed finite-dimensional-to-half-line details. For  $j \geq 6$  choose  $r = \lfloor j/3 \rfloor$  and retain the first and last  $r$  interior sites. Set to zero all perturbation entries outside their two diagonal blocks, including crossing entries. The trace-norm error is bounded by

$$C_k(1+j)^{M_k} \rho^r \quad (17.11)$$

after  $k$  fixed derivatives. This uses tridiagonality and the decay of the two endpoint layers, not a dimension times operator-norm estimate. Within each retained block, replacing the finite segment by its own half-line segment has an exponentially small entrywise sum, by (17.9) and the opposite tail.

For a perturbation supported in these blocks, the identity  $\det(I + AB) = \det(I + BA)$  reduces the determinant to that support. The left compression of  $G_j$  differs from the compression of  $G^{(0)}$  by an exponentially small operator norm. The reversed right compression differs in the same way from  $G^{(p)}$ . To check the claim, use (15.5): after expanding hyperbolic sines into exponentials, the

reflection from the remote boundary has a factor  $e^{-\gamma(2j-i-k)}$  for  $i, k \leq r$ . The two off-diagonal block kernels have a factor  $e^{-\gamma(j-2r)}$ . Summing at most  $r^2$  entries and differentiating any fixed number of times only adds a polynomial factor. Consequently the compressed Green matrix differs exponentially from the direct sum of the two compressed half-line operators. The same estimate holds after multiplication by the bounded trace-norm perturbation.

For completeness, if  $T, T'$  are the resulting trace-class operators with norms at most  $q < 1$ , telescoping powers gives

$$|\operatorname{tr}(T^m - (T')^m)| \leq m q^{m-1} \|T - T'\|_1.$$

Summation with coefficient  $1/m$  proves local Lipschitz continuity of the logarithmic determinant in trace norm. Its differentiated series has the same property with a fixed polynomial factor in  $m$ , still summable. This justifies each truncation and comparison above, for all fixed mixed derivatives. Sending the block size to infinity at either end costs  $O(\rho^r)$  again. The direct-sum logarithm is the sum of the half-line logarithms in (17.7).

We have proved exponential convergence of (17.10) to  $\log B_0(u) + \log B_p(v)$ . Exponentiation preserves the bounds, since all logarithms range over one bounded interval. Choose  $\tau < 1$  slower than the finitely many exponential rates in the gluing, truncation, and off-diagonal estimates. Polynomial factors at each fixed derivative order are absorbed in  $C_k$ . The finitely many  $j < 6$  are absorbed by enlarging those constants. At zero the finite and limiting normalized twists are one, which proves the last vanishing assertion. Uniform complex versions of the same estimates prove analyticity. In particular the argument never divides an absolute action error by the exponentially small reference twist.  $\square$

**17.3. Physical limiting laws.** For  $p \in \{0, 1\}$  let  $q_p = \sqrt{a_0 a_p}$  and define, for  $d > 0$ ,

$$\mathcal{F}_p(d) = \frac{q_p}{\pi d^2} \int_{\mathbb{R}^2} (d - S_0(u) - S_p(v))_+ B_0(u) B_p(v) \, du \, dv. \quad (17.12)$$

The integral is restricted to the common contact box; for the collars below its support lies strictly inside that box, so the displayed extension by zero is unambiguous.

**Theorem 17.3** (Fixed-offset long-bridge law). *There is  $d_* > 0$  independent of  $j$  such that  $\mathcal{F}_p$  has a smooth right extension to  $[0, d_*]$ , with  $\mathcal{F}_p(0) = 1$  and  $\mathcal{F}_p > 0$ . On this interval,*

$$\left\| \frac{2A \sinh(j\gamma)}{d^2} \Pr(E_{j,e}(d)) - \mathcal{F}_p(d) \right\|_{C^k} \leq C_k \tau^j, \quad p = j \bmod 2, \quad (17.13)$$

for every fixed mixed parameter and offset derivative order. The conditional endpoint and first-residual-time law converges to

$$\frac{B_0(u) B_p(v)}{I_p(d)} \mathbf{1}_{\{r > 0, S_0(u) + S_p(v) + r < d\}} \, du \, dv \, dr, \quad I_p(d) = \int (d - S_0 - S_p)_+ B_0 B_p \, du \, dv. \quad (17.14)$$

For the scaled variables  $(u/\sqrt{d}, v/\sqrt{d}, r/d)$  the convergence has bounded-Lipschitz error at most  $C\tau^j$ , uniformly for  $0 < d \leq d_*$ . Any fixed number of collisions at the two ends may be included, with their limiting coordinates given by the corresponding half-line segments. The two ends are coupled by the common residual energy constraint; (17.14) is not a product probability.

*Proof.* At parity  $p$ , (15.2) and (17.4) give  $d_j^0 = q_p / \sinh(j\gamma)$ . Full-phase integration therefore gives exactly

$$F_j(d) = \frac{q_p}{\pi d^2} \int (d - E_j)_+ b_j \, du \, dv.$$

To compare this integral down to  $d = 0$ , use the constructive Morse maps of Lemma 16.1 for both  $E_j$  and  $S_0 \oplus S_p$ . Their quadratic matrices are uniformly positive. The matrix  $2 \int_0^1 (1 -$



$s)D^2E(sy)ds$ , its positive square root and the inverse coordinate map depend locally Lipschitz-continuously in each fixed  $C^k$  norm on  $E$ , at the cost of finitely many additional derivatives. This follows by differentiating the defining inverse equation and using the uniform inverse differential; the positive-square-root map has bounded derivatives on the common compact positive cone. Thus their inverse maps and transformed amplitudes differ by  $C_k\tau^j$ . The inverse maps both vanish at zero, so their difference at  $w$  is also  $O(\tau^j|w|)$ .

Put  $w = \sqrt{2d}z$ . The normalized integrals are over the same unit disk with weight  $1 - |z|^2$ . Their odd Taylor terms vanish. Taylor's integral remainder, to two additional endpoint orders for each requested offset derivative, bounds the difference of the smooth functions of  $d$  by  $C_k\tau^j$ . This proves (17.13) including its derivatives at zero; it is not differentiation of a moving sharp boundary without justification. The quadratic integral at zero is  $\pi d^2/q_p$ , proving  $\mathcal{F}_p(0) = 1$ . Positivity and a smaller common collar give a uniform positive lower bound.

For conditional laws retain  $r = d\eta$  and the fixed domain  $\mathcal{U}$  in (16.5). The normalized densities there differ in  $L^1$  by  $C\tau^j$ , since the denominators are uniformly positive. The two endpoint maps, after division by  $\sqrt{d}$ , differ uniformly by  $C\tau^j$ . Couple the common part of the latent densities and use these maps; the remaining mass is  $O(\tau^j)$  on a bounded scaled domain. This proves the bounded-Lipschitz estimate. A fixed number of end coordinates follows from the differentiated gluing estimate and its vanishing at zero. First variations give directions from neighboring coordinates when desired. The explicit density (17.14) follows by returning to  $(u, v, r)$  coordinates.  $\square$

**Corollary 17.4** (Nonlinear weights at a tied onset). *Suppose a fixed finite set of selected ground channels has the same gap. Let  $\gamma_{\min}$  be the least of their exponents. At a fixed  $0 < d \leq d_*$ , along a parity subsequence their conditional weights in the complete maximal-count event converge to*

$$\frac{\mathbf{1}_{\{\gamma_e = \gamma_{\min}\}} \mathcal{F}_{e,p}(d)}{\sum_{e': \gamma_{e'} = \gamma_{\min}} \mathcal{F}_{e',p}(d)}.$$

*The convergence is exponential. At  $d = 0$  the minimizing leading weights are equal, whereas at positive offset the boundary laws supply their corrections.*

*Proof.* Divide the finite sum in (13.1) by  $d^2 e^{-j\gamma_{\min}}/A$ . Theorem 17.3 and  $\operatorname{csch}(j\gamma) = 2e^{-j\gamma}/(1 - e^{-2j\gamma})$  give the limit. All limiting denominators are positive. A finite positive gap between distinct exponents and the factorization rate give an exponential error; if all exponents tie only the latter rate is needed.  $\square$

**Corollary 17.5** (The first poles of the onset series). *For a fixed geometry put  $\mathcal{G}(z, d) = \sum_{j \geq 1} z^j \Pr(E_{j,e}(d))/d^2$ . There is  $\nu > 0$  such that, with  $x = ze^{-\gamma}$ ,*

$$\mathcal{G}(z, d) = \frac{\mathcal{F}_1(d)x + \mathcal{F}_0(d)x^2}{A(1 - x^2)} + \mathcal{H}(z, d), \quad (17.15)$$

*where  $\mathcal{H}$  is holomorphic for  $|z| < e^{\gamma+\nu}$  and smooth in  $d \in [0, d_*]$ . Every fixed geometric and offset derivative has the same assertion on strictly smaller common domains. The positive singularity is a simple pole at  $z = e^\gamma$ , with residue  $-e^\gamma[\mathcal{F}_0(d) + \mathcal{F}_1(d)]/(2A)$ . The possible pole at  $-e^\gamma$  is removable exactly when  $\mathcal{F}_0(d) = \mathcal{F}_1(d)$ .*

*Proof.* The coefficient in question is

$$\frac{e^{-j\gamma}}{A(1 - e^{-2j\gamma})} F_j(d) = \frac{e^{-j\gamma}}{A} \mathcal{F}_{j \bmod 2}(d) + O(e^{-(\gamma+\nu)j}),$$

for any positive  $\nu$  strictly smaller than  $-\log \tau$  and  $2\gamma$ . Fixed parameter derivatives add only fixed powers of  $j$ , absorbed by a smaller strict margin. The error series is normally convergent,

as are its stated derivatives. Summing its even and odd leading terms gives (17.15). Positivity gives the nonzero positive residue. At the negative point the numerator is  $\mathcal{F}_0 - \mathcal{F}_1$ , proving the removability criterion.  $\square$

The series in (17.15) uses a different physical window  $fg + d$  at each coefficient. It is not the generating function of all itineraries in one observation window. Its continuation is a consequence of a nonlinear relative-flux limit for these specified events.

**17.4. Differentiated operator bounds and moving poles.** For use in differentiated forms of the relative factorization, the operator estimates can be recorded in a single common norm convention. Let  $G$  denote either a finite reference Green matrix or a half-line Green operator, and let  $\Delta H$  be the corresponding tridiagonal Hessian perturbation on its stationary segment. For every fixed  $k$ ,

$$\max_{|\alpha| \leq k} \|\partial_\xi^\alpha G\|_{\text{op}} \leq C_k, \quad \max_{|\alpha|+|\beta| \leq k} \|\partial_\xi^\alpha \partial_Y^\beta \Delta H\|_1 \leq C_k, \quad \|G\Delta H\|_{\text{op}} \leq q < 1. \quad (17.16)$$

Here  $Y$  is the finite endpoint vector or the single half-line endpoint. There is no assertion that derivatives of  $G\Delta H$  are themselves small. They are bounded in trace norm, whereas its zeroth-order operator norm is small on the fixed contact box.

For completeness, the first bound follows directly from the Green kernels: their parameter derivatives introduce fixed powers of the indices into exponentially summable kernels, absorbed by the strict exponential margin. The Schur row and column estimates give the operator bound independently of matrix size. For the second, each tridiagonal entry is a local smooth function of neighboring stationary coordinates. The endpoint localization and its differentiated bounds make the sum of its absolute entries bounded. This sum bounds the trace norm. The last inequality follows by choosing the common box once, using  $\Delta H(0) = 0$ . The product rule consequently gives trace-norm bounds for every fixed derivative of  $T = G\Delta H$ .

These estimates also explain the differentiated determinant comparison. Differentiating  $\text{tr}(T^n)$  a fixed  $k$  times produces at most  $C_k n^k$  products with at most  $k$  differentiated factors. One factor is bounded in trace norm and all others in operator norm. The undifferentiated factors contribute the geometric decay  $q^{n-k}$ ; the finitely many  $n < k$  are harmless. In a difference of two such expressions, telescope one factor at a time. The differentiated trace-norm differences from the block comparison in the proof of Theorem 17.2 supply  $C_k \tau^j$ , and summation in  $n$  is uniform. Thus differentiated logarithmic determinants retain a trace-class factor and never acquire a factor equal to the number of interior sites.

In Corollary 17.5, simple poles refer to the undifferentiated series at a fixed geometry. If a parameter moves  $\gamma$ , differentiation of its displayed rational principal part can increase the pole order. After  $k$  parameter derivatives it is at most  $k + 1$ . For example,  $\partial_\xi(1 - ze^{-\gamma})^{-1} = -ze^{-\gamma}(\partial_\xi \gamma)/(1 - ze^{-\gamma})^2$ . The differentiated remainder remains holomorphic on the stated strictly smaller common domains. This separates the regularity of the remainder from the order of the moving principal poles.

## 18. FROM JACOBI IDENTITIES TO A RELATIVE PHYSICAL LAW

The proof of the relative law uses three different types of information. The first is finite-dimensional algebra, the second is localization of nonlinear perturbations, and the third is stability of a physical sublevel integral. We separate these statements before deriving a nonlinear compatibility law from their combination.

**18.1. The classical operator input.** For a periodic discrete Lagrangian orbit, the discrete Hill formula of Bolotin and Treschev [2, Theorem 2.1] relates  $\det(P - I)$  to the periodic action Hessian divided by the product of mixed edge determinants. The Hessian in our endpoint problem instead has Dirichlet boundary conditions. Its relevant scalar identity is the following elementary Schur-complement formula. We include its proof to specify the boundary convention, not as a new determinant identity.

**Proposition 18.1** (Dirichlet cofactor and normalization). *Let  $L(x_0, \dots, x_j) = \sum_{i=0}^{j-1} \ell_i(x_i, x_{i+1})$  have a smooth stationary interior at fixed endpoints  $x_0 = u, x_j = v$ . Suppose its interior Hessian  $H$  is positive definite and tridiagonal, with mixed edge derivatives  $\ell_{i,12} = -J_i < 0$ . For the stationary value  $W(u, v)$ ,*

$$-W_{uv} = \frac{\prod_{i=0}^{j-1} J_i}{\det H}. \quad (18.1)$$

*For a reference interior Hessian  $H^0$  and reference edges  $J_i^0$ , put  $G^0 = (H^0)^{-1}$  and  $\Delta H = H - H^0$ . Then*

$$\frac{-W_{uv}}{(-W_{uv})^0} = \frac{\prod_i (J_i/J_i^0)}{\det(I + G^0 \Delta H)}. \quad (18.2)$$

*If  $H = D(I - T)$ ,  $D$  is diagonal with  $D_{ii} \geq d_0 > 0$ ,  $T$  has only its two adjacent diagonals nonzero, and  $\|T\|_{\ell^\infty \rightarrow \ell^\infty} \leq q < 1$ , then*

$$|(H^{-1})_{ik}| \leq \frac{q^{|i-k|}}{d_0(1-q)}. \quad (18.3)$$

*These estimates have constants independent of the number of sites.*

*Proof.* For  $j \geq 2$ , differentiating the stationary equations eliminates the interior and gives  $W_{uv} = -J_0 J_{j-1} (H^{-1})_{1,j-1}$ . The corner cofactor of a tridiagonal matrix is  $J_1 \cdots J_{j-2}$ , so Cramer's rule proves (18.1). At  $j = 1$  the interior determinant is the empty determinant, equal to one, and the same assertion is immediate. Taking the ratio at the reference point proves (18.2). Finally,  $H^{-1} = \sum_{n \geq 0} T^n D^{-1}$ . The  $(i, k)$  entry of  $T^n$  vanishes for  $n < |i - k|$ , while every entry has absolute value at most  $q^n$ . Summing the remaining series proves (18.3).  $\square$

The separated inverse structure is classical: Meurant's Theorem 2.3 [15] gives inverse entries by two scalar recurrences, and his Theorem 2.4 describes their monotone decrease under strict diagonal dominance. Nabben [16] studies quantitative inverse decay for nonsymmetric tridiagonal and band matrices. Our use of a weighted Green operator belongs to this operator framework. In particular neither the corner cofactor nor exponential inverse decay is claimed here as a new mechanism.

What must additionally be proved for the present probability problem is summability of the *nonlinear perturbation evaluated on the entire stationary bridge*. A bound on each matrix entry, without its location, would not give the required determinant limit. The relevant continuity statement is recorded next.

**Proposition 18.2** (A dimension-independent relative determinant bound). *For trace-class operators  $T, T'$  on a common Hilbert space with  $\|T\|, \|T'\| \leq q < 1$ , define*

$$\mathcal{L}(T) = \sum_{n \geq 1} \frac{(-1)^{n+1}}{n} \operatorname{tr}(T^n).$$

*Then*

$$|\mathcal{L}(T) - \mathcal{L}(T')| \leq (1 - q)^{-1} \|T - T'\|_1. \quad (18.4)$$

Suppose  $T, T'$  depend smoothly on finitely many parameters, their trace-class derivatives through order  $k$  are uniformly bounded, and the corresponding derivatives of  $T - T'$  have trace norm at most  $\epsilon$ . Then

$$\|\mathcal{L}(T) - \mathcal{L}(T')\|_{C^k} \leq C_k \epsilon.$$

The constants depend on  $q$ , the parameter dimension and the stated bounds, not on the dimension of a finite-dimensional truncation.

*Proof.* The algebraic telescoping formula for  $T^n - (T')^n$  and the trace ideal inequality give  $|\operatorname{tr}(T^n - (T')^n)| \leq nq^{n-1}\|T - T'\|_1$ . Division by  $n$  and summation prove the first assertion. After  $k$  parameter derivatives, each product has at most  $k$  differentiated factors. In the difference of the two differentiated products, telescope the factors once more and place the difference factor in trace norm. Every other differentiated factor has a bounded operator norm. All but at most  $k+1$  of the factors still have norm at most  $q$ . There are at most  $C_k n^{k+1}$  terms. The finitely many products with  $n \leq k+1$  are bounded separately; the other terms are bounded by  $C_k \epsilon n^{k+1} q^{n-k-1}$ . Their sum, including  $1/n$ , converges. The same estimates justify differentiating the original series by uniform convergence. No smallness of a differentiated operator has been used.  $\square$

In the billiard application, Lemma 15.2 gives  $|x_i| \leq C(|u|\rho^i + |v|\rho^{j-i})$ , with the corresponding localized derivative bounds. Since a Hessian entry depends on only neighboring coordinates, its perturbation has a uniformly summable entrywise absolute value. Matrix units have trace norm one; hence

$$\|\Delta H_j\|_1 \leq C(|u| + |v|), \quad \|\partial^\alpha \Delta H_j\|_1 \leq C_\alpha.$$

The two-block construction in Theorem 17.2 then compares the normalized finite determinant with the direct sum of the two half-line determinants. Its trace-norm error tends to zero exponentially. Proposition 18.2 converts that specific geometric estimate into convergence of the logarithmic amplitude. Inverse decay supplies access to the blocks; the localized perturbation estimate supplies their summable determinant error.

**18.2. The passage through physical integration.** A second, logically separate requirement is uniformity as the offset vanishes. The following statement isolates this passage. All integrals are over a fixed neighborhood containing the indicated sublevel sets. The parameter set is compact and finite dimensional.

**Proposition 18.3** (Stability of normalized sublevel integration). *Let  $E, \tilde{E}$  be smooth functions of  $y \in \mathbb{R}^2$  and of parameters, with value and  $y$ -gradient zero at  $y = 0$ . Assume their Hessians at zero are uniformly positive and their higher derivatives are uniformly bounded on a common neighborhood. Let  $a, \tilde{a}$  be smooth amplitudes with uniform derivative bounds. There is a common  $d_* > 0$  such that*

$$\mathcal{J}(E, a)(d) = d^{-2} \int (d - E(y))_+ a(y) \, dy$$

*extends smoothly from  $d > 0$  to  $[0, d_*]$ . For each fixed  $k$ , bounds through a sufficiently high finite order  $r(k)$  give*

$$\|\mathcal{J}(E, a) - \mathcal{J}(\tilde{E}, \tilde{a})\|_{C^k} \leq C_k(\|E - \tilde{E}\|_{C^{r(k)}} + \|a - \tilde{a}\|_{C^{r(k)}}). \quad (18.5)$$

*The norms on the left include mixed parameter and offset derivatives.*

*Proof.* Taylor's integral formula gives  $E(y) = y^t M_E(y) y / 2$ , with  $M_E(y) = 2 \int_0^1 (1-s) D^2 E(sy) \, ds$ . After decreasing the common neighborhood, the maps  $y \mapsto M_E(y)^{1/2} y$  and their counterparts for  $\tilde{E}$  have uniformly invertible differentials. Their inverses  $\Phi_E$  are obtained by the contraction used in Lemma 16.1. Subtracting the inverse equations gives their Lipschitz dependence in  $C^0$  on  $E$ ; differentiating and isolating the highest derivative on the same uniformly invertible differential

proves the corresponding finite-order estimates by induction. Smoothness of the positive matrix square root is used only on a compact positive-definite set.

Set  $A_E(w) = a(\Phi_E(w))|\det D\Phi_E(w)|$ . The change of variables  $w = \sqrt{2d}z$  gives the fixed-domain formula

$$\mathcal{J}(E, a)(d) = 2 \int_{|z| < 1} (1 - |z|^2) A_E(\sqrt{2d}z) dz. \quad (18.6)$$

The inverse-map estimates bound  $A_E - A_{\tilde{E}}$  in each fixed  $C^s$  norm by a sufficiently high finite norm of the two input differences. The integral in (18.6), as a function of  $t = \sqrt{2d}$ , is smooth and even. A smooth even function  $h(t)$  has a smooth right representation in  $t^2$  with its first  $k$  derivatives controlled by derivatives of  $h$  through order  $2k$ . For example  $h'(t)/t = \int_0^1 h''(st) ds$ ; iteration, applied also to the even function  $h''$ , proves the assertion and its difference estimate. Parameter derivatives commute with these integrals. Applying it to (18.6) proves the smooth extension and (18.5). All sublevel sets lie inside the common chart by the uniform quadratic lower bound, so there is no additional chart-boundary term.  $\square$

Apply this proposition to  $(E_j, b_j)$  and  $(S_0 \oplus S_p, B_0 \otimes B_p)$ . The exact physical multiplier is  $q_p/(2\pi A \sinh(j\gamma))$ , by (16.1). Thus the exponentially small flux is retained *outside* a uniformly stable normalized integral. In this form the proof yields all fixed-order derivatives on a closed collar and the rare-event-preserving experiment transfer, rather than only an identity at the normal orbit.

The operator identities above, the trace estimate and the smooth integration principle are reusable elementary ingredients. The billiard-specific work is to verify their simultaneous hypotheses for physical stationary records with arbitrary collision order. The resulting law has, in addition to the experiment comparison in Part II, the nonlinear parity consequence established below. This consequence concerns the complete boundary functions at positive offset and does not follow from a finite list of leading amplitudes.

## 19. HYPERBOLIC COORDINATES AND THE INTRINSIC BOUNDARY MEASURE

The relative product law has a classical analytic model. We first calculate that model in physical endpoint coordinates, and then identify the boundary amplitudes intrinsically for the smooth billiards of Theorem 12.1. These are complementary statements: the first assumes normalizing charts; the second constructs a one-dimensional scalar density with the amplitude of the already proved half-line law. The scalar coordinate itself is classical, as recalled in Lemma 19.2. No smooth two-dimensional normal-form theorem is used in this identification.

**19.1. The analytic mixed-boundary calculation.** De Simoi–Kaloshin–Leguil [5], in the normal-form discussion on printed page 10 of arXiv:1905.00890v4, recall the analytic symplectic normal form near a hyperbolic period-two orbit. The local existence statement is distinct from the symmetry assumptions in their geometric inverse theorem; see also their discussion on printed page 13. The following physical endpoint calculation was supplied in the independent memorandum [24]. We include its proof and normalizations. The proposition is conditional on the indicated charts, and is not attributed as a probability theorem to [5].

**Proposition 19.1** (Relative flux in analytic hyperbolic coordinates). *Let*

$$N(s, p) = (\Delta(sp)s, \Delta(sp)^{-1}p), \quad \Delta(0) = \lambda \in (0, 1), \quad (19.1)$$

*where  $\Delta$  is positive and analytic near zero. Let*

$$\mathcal{C}_L(s, p) = (U(s, p), \mathcal{P}(s, p)), \quad \mathcal{C}_R(q, t) = (V(q, t), \mathcal{Q}(q, t))$$

*be analytic canonical charts sending zero to zero, with  $U_s(0, 0)V_t(0, 0) \neq 0$ . For the local branch of  $\mathcal{C}_R \circ N^n \circ \mathcal{C}_L^{-1}$ , prescribe the initial physical position  $u$  and terminal physical position  $v$ . There*

are a fixed box in  $(u, v)$ , an integer  $n_0$ , and  $\sigma \in (0, 1)$  such that this boundary problem has a unique local solution for all  $n \geq n_0$ . Its physical mixed flux  $J_n = |\partial P / \partial v|_u$  satisfies

$$\left\| \frac{J_n(u, v)}{J_n(0, 0)} - \mathcal{B}_L(u) \mathcal{B}_R(v) \right\|_{C^k} \leq C_k \sigma^n \quad (19.2)$$

for every fixed  $k$ , where, writing  $U(\alpha(u), 0) = u$  and  $V(0, \beta(v)) = v$ ,

$$\mathcal{B}_L(u) = \frac{U_s(0, 0)}{U_s(\alpha(u), 0)}, \quad \mathcal{B}_R(v) = \frac{V_t(0, 0)}{V_t(0, \beta(v))}. \quad (19.3)$$

These functions are positive and equal one at zero. If the physical branch has a generating action  $W_n$ , normalized by  $P = -W_{n,u}$  and  $Q = W_{n,v}$ , then there are functions  $\mathcal{S}_L, \mathcal{S}_R$ , vanishing with their first derivatives at zero, for which

$$\|W_n - W_n(0, 0) - \mathcal{S}_L \oplus \mathcal{S}_R\|_{C^k} \leq C_k \sigma^n. \quad (19.4)$$

All assertions hold with mixed parameter derivatives for a supplied jointly smooth family of these functions and charts on common domains, provided their fixed-order norms are uniformly bounded, their projection derivatives are uniformly nonzero, and  $0 < \lambda_- \leq \lambda \leq \lambda_+ < 1$ .

*Proof.* Temporarily prescribe  $s = s_0$  and  $t = p_n$  rather than  $u, v$ . The invariant  $I = sp$  and the orbit endpoints obey

$$I = st\Delta(I)^n, \quad p_0 = t\Delta(I)^n, \quad q = s_n = s\Delta(I)^n. \quad (19.5)$$

For a single map put  $\lambda_+ = \lambda$ . Choose  $\lambda_+ < \mu < \sigma < 1$  and decrease the  $I$  interval so that  $\Delta \leq \mu$ . On a fixed sufficiently small  $(s, t)$  box,  $|st|n\mu^{n-1}\|\Delta'\|_\infty < 1/2$  for every  $n \geq 1$ , since  $\sup_n n\mu^{n-1} < \infty$ . The right side of the first identity in (19.5) maps a common small closed  $I$  interval into itself. It is a contraction there. Its solution satisfies  $|I| \leq C\mu^n$ . The intermediate coordinates are  $(s\Delta(I)^i, t\Delta(I)^{n-i})$ , so the whole orbit remains inside the chosen local chart.

Set

$$a_n = \Delta(I)^n, \quad D_n = 1 - nI \frac{\Delta'(I)}{\Delta(I)}.$$

The contraction estimate gives  $|D_n| \geq 1/2$ . Differentiating (19.5), without dividing by either endpoint, gives

$$T_n := \frac{\partial p_0}{\partial t} \Big|_s = \frac{\partial q}{\partial s} \Big|_t = \frac{\Delta(I)^n}{D_n}. \quad (19.6)$$

For example,  $I_t = sa_n/D_n$  and  $(p_0)_t = a_n + tn\Delta(I)^{n-1}\Delta'(I)I_t = a_n/D_n$ . This also proves the formula on both axes.

Here are the differentiated relative estimates. At every fixed order, derivatives of the function  $\Delta(I)^n$  with  $I$  treated as an independent variable are bounded by  $C_k(1+n)^k\mu^n$ ; a finite number of small  $n$  are absorbed in  $C_k$ . Implicit differentiation of the first equation in (19.5) leaves  $D_n$  on the highest derivative of  $I$ . Its remaining terms contain the preceding derivatives of  $I$ , derivatives of  $st$ , and derivatives of  $\Delta^n$ . Induction therefore gives

$$\|I\|_{C^k} \leq C_k(1+n)^{m_k}\mu^n.$$

These norms can include the supplied parameters. Since  $\log(\Delta(I)/\lambda) = Ih(I)$  with a jointly smooth bounded function  $h$ , the same induction gives

$$\left\| n \log \frac{\Delta(I)}{\lambda} \right\|_{C^k} + \|D_n - 1\|_{C^k} \leq C_k(1+n)^{r_k}\mu^n.$$

The identity at  $I = 0$  holds throughout the parameter family; thus parameter derivatives of  $\lambda$  do not leave an unsuppressed term in this normalized expression. Exponentiating and using  $|D_n| \geq 1/2$  proves

$$\|\lambda^{-n}T_n - 1\|_{C^k} \leq C_k \sigma^n. \quad (19.7)$$

One fixed strict margin  $\mu < \sigma$  absorbs the polynomial at each fixed order; the constants, not the margin, depend on that order.

The physical boundary map in the mixed coordinates is

$$\Phi_n(s, t) = (U(s, ta_n), V(sa_n, t)), \quad \Phi_\infty(s, t) = (U(s, 0), V(0, t)). \quad (19.8)$$

It converges to  $\Phi_\infty$  in every fixed derivative norm at an exponential rate. Choose fixed inner and outer boxes on which the two limiting scalar projections are diffeomorphisms, with the closure of the inner box strictly inside the image of the outer one. Conjugating  $\Phi_n$  by  $\Phi_\infty^{-1}$  gives the identity plus a uniformly small  $C^1$  perturbation. The equation for its inverse is a contraction on a small ball about each point of the inner box, uniformly in that point. This gives a unique inverse on a common nonshrinking physical box for all  $n \geq n_0$ . Subtracting the inverse equations gives its  $C^0$  convergence. Differentiation isolates the same uniformly invertible differential and proves the  $C^k$  convergence by induction. In particular  $(s, t) = \Phi_n^{-1}(u, v)$  tends to  $(\alpha(u), \beta(v))$  at the stated rate.

Put  $\mathcal{D}_n = \det D\Phi_n$ . Canonicity of the initial chart and (19.6) give, on the boundary surface,

$$du \wedge dP = ds \wedge dp_0 = T_n ds \wedge dt, \quad du \wedge dv = \mathcal{D}_n ds \wedge dt.$$

Consequently the exact physical, rather than canonical, flux is

$$J_n(u, v) = \left| \frac{T_n(s, t)}{\mathcal{D}_n(s, t)} \right|, \quad (s, t) = \Phi_n^{-1}(u, v). \quad (19.9)$$

At zero this denominator is exactly

$$\mathcal{D}_n(0, 0) = U_s(0, 0)V_t(0, 0) - \lambda^{2n}U_p(0, 0)V_q(0, 0). \quad (19.10)$$

On the physical box its limit is  $U_s(\alpha(u), 0)V_t(0, \beta(v))$ , uniformly bounded away from zero. Equations (19.7)–(19.10) therefore prove (19.2) and (19.3), including all fixed mixed derivatives. This step is necessary: canonical boundary coordinates are not physical configuration observations.

Finally define

$$\mathcal{S}_L(u) = - \int_0^u \mathcal{P}(\alpha(x), 0) \, dx, \quad \mathcal{S}_R(v) = \int_0^v \mathcal{Q}(0, \beta(y)) \, dy.$$

The physical first variations of  $W_n$  converge to the two displayed integrands, with all fixed derivatives, by (19.8). Integrate the difference of the gradients along the two coordinate segments from  $(0, 0)$  to  $(u, v)$ . The values at the origin fix the additive constants and prove (19.4).  $\square$

For the two-step billiard return the contracting eigenvalue in (19.1) is  $\lambda = e^{-2\gamma}$ , and  $n$  returns mean  $j = 2n$  flights. At contact  $b$  its stable and unstable linear directions in physical position–momentum coordinates are  $(1, -a_b)$  and  $(1, a_b)$ , by (15.1). They are nonvertical, regardless of equality or inequality of the curvatures. Use the same canonical chart at both ends. Write its first derivative as

$$DC(0) = \begin{pmatrix} A & B \\ -a_b A & a_b B \end{pmatrix}.$$

Canonicity gives  $AB = (2a_b)^{-1}$ . Thus (19.10) and (19.9) recover the exact physical normalization

$$J_n(0, 0) = \frac{2a_b \lambda^n}{1 - \lambda^{2n}} = \frac{a_b}{\sinh(2n\gamma)}. \quad (19.11)$$

This is an exact agreement with the Jacobi formula, not merely an agreement of leading exponential orders. If both constructions apply to the same analytic billiard, uniqueness of the normalized limit and evaluation at the other endpoint zero give  $B_b = \mathcal{B}_L = \mathcal{B}_R$ . Their action limits agree as well, with the stated zero normalizations. The normal-form function  $\Delta$  alone does not specify the physical amplitudes: the projection charts enter (19.3) essentially.

**19.2. The classical scalar comparison.** Smooth linearization on a one-dimensional contracting branch is classical. It requires neither a symplectic chart on a neighborhood nor an analytic hypothesis. This is the one-dimensional case of Sternberg's theorem; see [25] and the precise regularity statement in the introduction of Eynard-Bontemps-Navas [26]. We record the elementary product construction in the normalization needed here. The scalar comparison and its application to the width were pointed out in the memorandum [27]. Its role is to separate a classical coordinate from the identification of the physical determinant amplitude.

**Lemma 19.2** (Normalized scalar linearization). *Let  $R_\xi$  be a jointly smooth family on a fixed interval about zero, with parameter  $\xi$  in a compact set, such that*

$$R_\xi(0) = 0, \quad R'_\xi(0) = \lambda_\xi, \quad 0 < \ell \leq R'_\xi(u) \leq q < 1.$$

*The family is understood to be defined on an open neighborhood of the compact parameter set. On a common smaller interval the series*

$$h_\xi(u) = \log \frac{R'_\xi(u)}{\lambda_\xi}, \quad (19.12)$$

$$L_\xi(u) = \sum_{i=0}^{\infty} h_\xi(R_\xi^i(u)), \quad \mathfrak{B}_\xi(u) = \exp L_\xi(u), \quad \mathfrak{z}_\xi(u) = \int_0^u \mathfrak{B}_\xi(x) \, dx$$

*converges with every fixed mixed parameter and endpoint derivative. The coordinate  $\mathfrak{z}_\xi$  is increasing, normalized by  $\mathfrak{z}_\xi(0) = 0$  and  $\mathfrak{z}'_\xi(0) = 1$ , and satisfies*

$$\mathfrak{B}_\xi(R_\xi(u))R'_\xi(u) = \lambda_\xi \mathfrak{B}_\xi(u), \quad \mathfrak{z}_\xi \circ R_\xi = \lambda_\xi \mathfrak{z}_\xi. \quad (19.13)$$

*It is the unique normalized differentiable local linearizer. For every fixed  $k$  and  $\sigma \in (\lambda_+, 1)$ , where  $\lambda_+ = \max_\xi \lambda_\xi$ , one has*

$$\|\lambda_\xi^{-n} R_\xi^n - \mathfrak{z}_\xi\|_{C^k} + \|\lambda_\xi^{-n} (R_\xi^n)' - \mathfrak{B}_\xi\|_{C^k} \leq C_{k,\sigma} \sigma^n. \quad (19.14)$$

*Here and below such norms include any fixed collection of mixed parameter and endpoint derivatives on the common domain.*

*Proof.* Suppress  $\xi$  in compositions, but not in differentiation. Put  $x_n = R_\xi^n(u)$ . Since  $R_\xi(0) = 0$  and  $|R'_\xi| \leq q$ ,  $|x_n| \leq q^n |u|$ . At every fixed positive mixed order  $m$ , differentiation of  $x_{n+1} = R_\xi(x_n)$  leaves  $R'_\xi(x_n)$  multiplying the highest derivative of  $x_n$ . Every other term either contains a lower derivative of  $x_n$ , or is a pure parameter derivative of  $R_\xi$  evaluated at  $x_n$ . The latter vanishes at  $x_n = 0$ , because the fixed point is common to the family. Induction and the mean value theorem therefore give

$$\sup_{|\alpha|+r=m} |\partial_\xi^\alpha \partial_u^r x_n| \leq C_m (1+n)^{M_m} q^n.$$

Indeed, at the inductive step the right-hand forcing in the highest order recurrence is bounded by a polynomial in  $n$  times  $q^n$ ; iteration of that recurrence only increases the polynomial degree. Uniform bounds on the finitely many derivatives of  $R_\xi$  involved follow from compactness. This argument covers endpoint derivatives, parameter derivatives and their mixtures.

The function  $h_\xi$  is smooth,  $h_\xi(0) = 0$ , and all its pure parameter derivatives vanish at zero. The same chain-rule argument bounds each fixed mixed derivative of  $h_\xi(x_n)$  by  $C_m (1+n)^{N_m} q^n$ . Summability proves convergence of (19.12) with all the asserted derivatives. In particular



$L_\xi(0) = 0$  and  $\mathfrak{B}_\xi(0) = 1$ . Shifting the summation index gives  $L_\xi(u) - L_\xi(R_\xi(u)) = h_\xi(u)$ . Exponentiation proves the first identity in (19.13). Integration from zero proves the second, since both sides vanish there.

If  $\tilde{\mathfrak{z}}$  is another normalized differentiable linearizer, set  $H = \tilde{\mathfrak{z}} \circ \mathfrak{z}^{-1}$ . Then  $H(0) = 0$ ,  $H'(0) = 1$ , and  $H(\lambda z) = \lambda H(z)$ . For every  $z$  in a sufficiently small interval,  $H(z) = \lambda^{-n} H(\lambda^n z) \rightarrow z$ . This proves uniqueness.

The exact relation  $R_\xi^n = \mathfrak{z}_\xi^{-1} \circ (\lambda_\xi^n \mathfrak{z}_\xi)$  gives the sharper rate based on  $\lambda_+$  rather than the preliminary contraction bound  $q$ . The inverse function theorem on a common smaller interval makes  $\mathfrak{z}_\xi^{-1}$  jointly smooth. Its value and first derivative at zero are 0 and 1; write it as  $x + x^2 A_\xi(x)$ . Thus

$$\lambda_\xi^{-n} R_\xi^n(u) - \mathfrak{z}_\xi(u) = \lambda_\xi^n \mathfrak{z}_\xi(u)^2 A_\xi(\lambda_\xi^n \mathfrak{z}_\xi(u)).$$

Every fixed mixed derivative is bounded by  $C_k(1+n)^{N_k} \lambda_+^n$ ; the positive lower bound on  $\lambda_\xi$  controls derivatives of its powers. Apply the same argument with one additional endpoint derivative for the density estimate. A fixed strict margin  $\lambda_+ < \sigma$  absorbs the polynomial factors. Finally, the finite product identity

$$\exp \left\{ \sum_{i=0}^{n-1} h_\xi(R_\xi^i(u)) \right\} = \lambda_\xi^{-n} (R_\xi^n)'(u) \quad (19.15)$$

follows directly from the chain rule and agrees with the density limit.  $\square$

The alternating scalar identity is equally classical. Suppose that  $\varphi_b(0) = 0$ ,  $\varphi_b'(0) = r_b > 0$  and  $R_b = \varphi_{1-b} \circ \varphi_b$ , with  $r_0 r_1 = \lambda < 1$ . Let  $\mathfrak{z}_b$  be the normalized scalar linearizer of  $R_b$ . The identity  $\varphi_b \circ R_b = R_{1-b} \circ \varphi_b$  shows that  $r_b^{-1} \mathfrak{z}_{1-b} \circ \varphi_b$  is another normalized linearizer of  $R_b$ . Uniqueness gives

$$\mathfrak{z}_{1-b} \circ \varphi_b = r_b \mathfrak{z}_b, \quad \mathfrak{B}_{1-b}(\varphi_b(u)) \varphi_b'(u) = r_b \mathfrak{B}_b(u).$$

These scalar facts by themselves do not identify the amplitude in a normalized finite physical bridge. That identification is the content of the following theorem and its exact Schur-concatenation proof.

**19.3. A smooth intrinsic characterization.** We now return to the full smooth setting. Write  $W_{j,b}$  for the stationary action starting at type  $b$ , and let

$$\varphi_b(u) = x_1^{(b)}(u), \quad \mathfrak{r}_b = \sqrt{\frac{c_b}{c_{1-b}}}, \quad r_b = \mathfrak{r}_b \varrho, \quad R_b = \varphi_{1-b} \circ \varphi_b, \quad \varrho = e^{-\gamma}, \quad \lambda = \varrho^2.$$

Here  $x^{(b)}$  is the physical half-line in Lemma 17.1. Thus  $\varphi_b$  is the first-hit position along the stable branch, not a freely chosen change of variables. Notice that  $r_0 r_1 = \lambda$  and  $0 < r_b < 1$ . Indeed  $\sqrt{c_b/c_{1-b}} e^{-\gamma} = c_b/(c_0 c_1 + \sqrt{c_0 c_1 (c_0 c_1 - 1)}) < 1$ , since  $c_{1-b} > 1$ .

**Theorem 19.3** (The boundary amplitude as a linearizing density). *On a fixed smaller collar, uniformly for the smooth compact families of Theorem 12.1, the half-line amplitudes satisfy*

$$B_{1-b}(\varphi_b(u)) \varphi_b'(u) = r_b B_b(u), \quad b = 0, 1. \quad (19.16)$$

*The increasing smooth coordinates*

$$\zeta_b(u) = \int_0^u B_b(x) \, dx \quad (19.17)$$

*are normalized by  $\zeta_b(0) = 0$ ,  $\zeta_b'(0) = 1$  and obey*

$$\zeta_{1-b}(\varphi_b(u)) = r_b \zeta_b(u), \quad \zeta_b(R_b(u)) = \lambda \zeta_b(u). \quad (19.18)$$

For each  $b$ ,  $\zeta_b$  is the unique differentiable local coordinate with that normalization which linearizes  $R_b$ . Moreover, for a fixed  $\sigma \in (\lambda_+, 1)$ , where  $\lambda_+$  is the maximum of  $\lambda$  on the compact family,

$$\|\lambda^{-n}R_b^n - \zeta_b\|_{C^k} + \|\lambda^{-n}(R_b^n)' - B_b\|_{C^k} \leq C_k \sigma^n. \quad (19.19)$$

The norms include every fixed mixed parameter and endpoint derivative. No normalizing chart for the two-dimensional return map is a hypothesis.

*Proof.* Put  $o = 1 - b$ . Uniqueness of the half-line tail gives the exact relations

$$S_b(u) = \ell_b(u, \varphi_b(u)) - g + S_o(\varphi_b(u)), \quad \ell_{b,2}(u, \varphi_b(u)) + S'_o(\varphi_b(u)) = 0. \quad (19.20)$$

The second relation also follows from first variation of the middle endpoint. Its  $w$  derivative is  $\ell_{b,22}(u, w) + S''_o(w)$ , uniformly positive on the small box. The implicit function theorem and the already constructed half-line therefore give

$$\varphi'_b(u) = \frac{-\ell_{b,12}(u, \varphi_b(u))}{\ell_{b,22}(u, \varphi_b(u)) + S''_o(\varphi_b(u))} > 0. \quad (19.21)$$

At zero this is  $(c_o + ga_o)^{-1} = \sqrt{c_b/c_o}e^{-\gamma} = r_b$ . After decreasing the common interval, both first-hit maps are contractions into it. All their mixed derivatives are controlled by the same positive denominator. The graph of stable initial physical states is  $P = -S'_b(u)$ : differentiating the first identity in (19.20) gives  $P = -\ell_{b,1}$ , while the terminal momentum is  $\ell_{b,2} = -S'_o$ . The reversed unstable graph is  $P = S'_b(u)$ . This verifies the dynamical interpretation directly from the physical generating action.

We prove the amplitude identity from exact finite concatenation. For a tail of  $j$  flights starting at  $o$ , let  $w_j = w_j(u, v)$  solve

$$\ell_{b,2}(u, w_j) + W_{j,o,1}(w_j, v) = 0.$$

Strict convexity identifies the resulting stationary value with  $W_{j+1,b}(u, v)$ . Set  $H_j = \ell_{b,22}(u, w_j) + W_{j,o,11}(w_j, v)$ . Eliminating the middle endpoint gives the exact Schur identity

$$-W_{j+1,b,12}(u, v) = \frac{-\ell_{b,12}(u, w_j)}{H_j} [-W_{j,o,12}(w_j, v)]. \quad (19.22)$$

The denominator is uniformly positive. Theorem 17.2 and its differentiated action limit imply  $w_j \rightarrow \varphi_b(u)$  and  $H_j \rightarrow \ell_{b,22}(u, \varphi_b(u)) + S''_o(\varphi_b(u))$  on the fixed box. The inverse equations, with the same positive denominator, give this convergence in every fixed mixed norm.

The terminal type of both sides of (19.22) is  $c = (o + j) \bmod 2$ . The exact reference constants are

$$d_{j,o}^0 = \frac{\sqrt{a_o a_c}}{\sinh(j\gamma)}, \quad d_{j+1,b}^0 = \frac{\sqrt{a_b a_c}}{\sinh((j+1)\gamma)}.$$

Their ratio tends to  $\sqrt{a_o/a_b}e^\gamma = r_b^{-1}$ , with all fixed parameter derivatives. Divide (19.22) by  $d_{j+1,b}^0$  and let  $j$  tend to infinity along either parity. Relative factorization and (19.21) give

$$B_b(u)B_c(v) = r_b^{-1}\varphi'_b(u)B_o(\varphi_b(u))B_c(v).$$

Since  $B_c > 0$ , cancellation proves (19.16). Integration from zero proves (19.18). This argument uses the previously proved relative theorem to identify its amplitude; it does not claim that a scalar linearization alone proves the finite-bridge theorem.

For uniqueness, conjugate any other normalized differentiable linearizer by  $\zeta_b^{-1}$ . The resulting function  $h$  has  $h(0) = 0$ ,  $h'(0) = 1$  and  $h(\lambda z) = \lambda h(z)$ . Hence  $h(z) = \lambda^{-n}h(\lambda^n z) \rightarrow z$ , so  $h(z) = z$ . For the quantitative assertion, the exact identity is

$$R_b^n(u) = \zeta_b^{-1}(\lambda^n \zeta_b(u)).$$

Write  $\zeta_b^{-1}(x) = x + x^2 h_b(x)$ , where  $h_b$  is jointly smooth on a common interval. Division by  $\lambda^n$  leaves  $\zeta_b(u) + \lambda^n \zeta_b(u)^2 h_b(\lambda^n \zeta_b(u))$ . At each fixed derivative order this last term, including parameter derivatives of  $\lambda$ , is bounded by  $C_k(1+n)^{m_k} \lambda_+^n$ . Its endpoint derivative proves the second estimate in (19.19); the fixed strict margin gives the displayed rate. This also proves all claimed uniformities without extending the scalar coordinate to a symplectic chart off the invariant branch.  $\square$

**Corollary 19.4** (Physical determinant and scalar product). *For the physical return  $R_b = \varphi_{1-b} \circ \varphi_b$  in Theorem 19.3, put  $h_b(u) = \log(R'_b(u)/\lambda)$ . The physical edge product  $U_b$ , Green operator  $G^{(b)}$  and Hessian perturbation  $\Delta H_b$  of (17.6)–(17.7) satisfy*

$$\frac{\exp U_b(u)}{\det(I + G^{(b)} \Delta H_b(u))} = B_b(u) = \exp \left\{ \sum_{i=0}^{\infty} h_b(R_b^i(u)) \right\}. \quad (19.23)$$

All expressions have the same value 1 at zero, and the series converges in every fixed mixed norm on the common collar. More precisely, the finite scalar product has the exact remainder

$$\log B_b(u) - \sum_{i=0}^{n-1} h_b(R_b^i(u)) = \log B_b(R_b^n(u)). \quad (19.24)$$

For each fixed  $k$  and each  $\sigma \in (\lambda_+, 1)$  the right side has  $C^k$  norm at most  $C_{k,\sigma} \sigma^n$ .

*Proof.* The first equality is the original relative Fredholm determinant construction, with its convergent trace series. Theorem 19.3 identifies its positive density by exact physical concatenation. Composing its two first-flight cocycles gives  $B_b(R_b(u))R'_b(u) = \lambda B_b(u)$ . Let  $\mathfrak{B}_b$  be the scalar density in Lemma 19.2. The positive quotient  $B_b/\mathfrak{B}_b$  is invariant under  $R_b$  and equals one at zero. Iteration to the contracting fixed point makes the quotient identically one. This proves (19.23). Telescoping the logarithmic cocycle proves (19.24) for every finite  $n$ . For the mixed norm bound use  $R_b^n = \zeta_b^{-1}(\lambda^n \zeta_b)$  and write  $\log B_b(\zeta_b^{-1}(x)) = x C_b(x)$  with  $C_b$  jointly smooth. The remainder is then  $\lambda^n \zeta_b(u) C_b(\lambda^n \zeta_b(u))$ . The same differentiated estimate as in (19.14) proves the stated bound.  $\square$

In particular, the infinite scalar product is not introduced as a new linearization theorem. Equation (19.23) is an identification of two normalized constructions for the same physical branch. The finite-bridge trace-norm comparison, the physical reference twist, the first-hit selection and the residual-time integration still have their separate proofs. The scalar product does not replace them.

**Corollary 19.5** (Stable-action width determined by the even laws). *Put  $\widehat{S}_b = S_b \circ \zeta_b^{-1}$ . Then the limiting laws can be written*

$$\mathcal{F}_{bc}(d) = \frac{\sqrt{a_b a_c}}{\pi d^2} \int (d - \widehat{S}_b(x) - \widehat{S}_c(y))_+ \, dx \, dy. \quad (19.25)$$

For small  $E > 0$ , let  $u_{b,-}(E) < 0 < u_{b,+}(E)$  solve  $S_b(u) = E$ . The two even laws determine the normalized widths

$$\begin{aligned} \mathcal{W}_b(E) &:= \sqrt{\frac{a_b}{2}} [\zeta_b(u_{b,+}(E)) - \zeta_b(u_{b,-}(E))] \\ &= \int_0^E x^{-1/2} \mathcal{V}_b(x) \, dx. \end{aligned} \quad (19.26)$$

Thus the smooth invariant is the complete normalized sublevel width of the half-line action in its intrinsic linearizing coordinate.

*Proof.* The change of variables  $x = \zeta_b(u)$  has differential  $dx = B_b(u) du$ . This proves (19.25). The pushforward calculation in (20.1) gives

$$\int_{S_b(u) < E} B_b(u) du = \sqrt{\frac{2}{a_b}} \int_0^E x^{-1/2} \mathcal{V}_b(x) dx.$$

It is also exactly  $\zeta_b(u_{b,+}(E)) - \zeta_b(u_{b,-}(E))$ . Theorem 20.1 determines  $\mathcal{V}_b$  from the even law, proving the assertion. That theorem's proof uses only the half-line formula and Volterra uniqueness, not this corollary.  $\square$

Conversely, on  $E > 0$ ,

$$\mathcal{V}_b(E) = \sqrt{E} \mathcal{W}_b'(E), \quad \mathcal{W}_b(0) = 0, \quad \mathcal{W}_b(E) = 2\sqrt{E} + O(E^{3/2}). \quad (19.27)$$

The first expression extends smoothly to  $E = 0$  with value one. Thus the normalized width and the energy profile are equivalent coordinates on the same invariant, not two independent invariants. The pushforward theorem determines a sublevel measure; the width formula expresses that measure in the classical stable coordinate.

The normalization in (19.26) is useful when the separate curvatures are not supplied. The even laws determine  $\mathcal{W}_b$  without knowing  $a_b$  separately. With  $a_b$  supplied, they also determine the unnormalized coordinate width. When both participating contact graphs are even,  $S_b$  and  $B_b$  are even and  $\zeta_b$  is odd, so this width specifies both inverse branches of  $\widehat{S}_b$ . For a general contact it specifies their difference. The physical contact-graph inverse in Section 21 is a further geometric argument, not an inference that a symmetrized width determines an arbitrary asymmetric boundary.

**19.4. The comparison at the level of the physical theorem.** Proposition 19.1 explains the analytic relative product mechanism, including its physical projection and exact reference normalization. If the limiting endpoint Hessians are positive, Proposition 18.3 also transports its estimates through the normalized residual integral. At an isolated analytic billiard channel the positivity and first-hit property are verified in Sections 14–15. An odd-flight analytic comparison can likewise compose the terminal chart with one physical flight. Thus neither the parity distinction nor the absence of global finite horizon is an obstruction to that local analytic mechanism. They are requirements handled simultaneously in our physical theorem, not separate claims of a new normal-form effect.

Theorem 12.1 does not supply normalizing charts as data or hypotheses. It starts with smooth obstacle graphs and positive geometric margins. Its half-line and finite-bridge contractions, localized trace-ideal estimates, and common sublevel charts prove the mixed parameter and offset bounds on one closed collar, with the exact phase normalization and the physical event selection. A pointwise analytic normal form supplies none of these uniform chart estimates for an arbitrary smooth family. The conditional family clause in Proposition 19.1 states exactly what would be needed to use that route instead. No assertion that such a smooth normal-form route is impossible, or absent from all prior literature, is required here.

Theorem 19.3 makes the geometric meaning of the determinant amplitude explicit in the full smooth class: it is the derivative of the normalized scalar linearizer along the physical stable branch, with an exact first-flight cocycle linking the two contact types. Corollary 19.5 identifies what the complete invariant measures in that coordinate. These conclusions connect the analytic and classical scalar interpretations to the constructed smooth amplitude. Equation (19.23) gives their exact relation; (19.27) records that the width is a coordinate on the previously determined invariant. Neither interpretation is counted as an additional rigidity theorem.

The remaining steps concern information that the forward benchmark does not reconstruct: the Volterra inverse for the complete widths; the explicit independent-contact geometric inverse, its

analytic physical image and its two-flight coordinates; and the regularized binary experiment with its charged smooth-class calibration. The preparation theorem uses certified physical smoothness bounds, not access to  $\Delta$  or to a canonical chart. Its sufficient rate is not claimed to be optimal. The central result is the joint passage from smooth physical records to an intrinsic boundary measure, to its complete observable invariant, and to a finite-preparation inverse. The classical analytic product mechanism is part of that explanation, not the novelty claimed for the package.

**19.5. An exact mixed-parameter benchmark.** The strict exponential margin in Lemma 19.2 is part of its mixed-parameter assertion. The following exact family, recorded in the independent memorandum [29], explains why the multiplier itself cannot in general replace that margin.

**Proposition 19.6** (A strict-margin example). *On a common interval about zero put*

$$R_{\lambda,a}(u) = \frac{\lambda u}{1 + a(1 - \lambda)u}, \quad z_a(u) = \frac{u}{1 + au}, \quad B_a(u) = (1 + au)^{-2}.$$

Let  $0 < \lambda_- \leq \lambda \leq \lambda_+ < 1$  and let  $a$  range over a compact interval containing zero in its interior. For a sufficiently small common endpoint interval these maps are smooth contractions and

$$R_{\lambda,a}^n(u) = \frac{\lambda^n u}{1 + a(1 - \lambda^n)u}, \quad (19.28)$$

$$\lambda^{-n}(R_{\lambda,a}^n)'(u) = [1 + a(1 - \lambda^n)u]^{-2}, \quad (19.29)$$

$$\lambda^{-n}R_{\lambda,a}^n(u) - z_a(u) = \frac{a\lambda^n u^2}{[1 + a(1 - \lambda^n)u](1 + au)}. \quad (19.30)$$

The normalized errors converge with every fixed mixed derivative at rate  $C_{k,\sigma}\sigma^n$  for each  $\lambda_+ < \sigma < 1$ . There is no such bound with  $C\lambda_+^n$  for the mixed derivative  $\partial_\lambda \partial_a$  of (19.30) on a nontrivial endpoint interval.

*Proof.* The denominators are uniformly positive after restricting the endpoint interval. Also  $R_{\lambda,a}'(u) = \lambda/[1 + a(1 - \lambda)u]^2$  is then positive and uniformly less than one. Direct substitution gives  $z_a \circ R_{\lambda,a} = \lambda z_a$  and  $z_a^{-1}(x) = x/(1 - ax)$ . Iteration and one endpoint differentiation prove (19.28)–(19.29). Subtraction proves (19.30). At each fixed order, differentiation gives a polynomial in  $n$  times  $\lambda_+^n$ , multiplied by uniformly bounded rational functions; the lower bound on  $\lambda$  controls the finitely many inverse powers introduced. A strict margin absorbs that polynomial. The density estimate follows either directly or by one additional endpoint derivative.

At  $a = 0$  the mixed derivative of the error is exactly

$$\partial_\lambda \partial_a (\lambda^{-n} R_{\lambda,a}^n(u) - z_a(u)) \Big|_{a=0} = n\lambda^{n-1}u^2.$$

At  $\lambda = \lambda_+$  and any fixed  $u \neq 0$ , its ratio to  $\lambda_+^n$  is  $nu^2/\lambda_+$ , which is unbounded. Finally the exact identity

$$\lambda^{-n}(R_{\lambda,a}^n)'(u)B_a(R_{\lambda,a}^n(u)) = B_a(u)$$

checks the scalar product remainder with no asymptotic expansion.  $\square$

This is a benchmark for mixed estimates, not a billiard realization or an additional linearization theorem. The physical theorem retains its already strict exponential margin and its independent relative finite-bridge proof.

## 20. BOUNDARY ENERGY PROFILES AND NONLINEAR PARITY COMPATIBILITY

Fix one of the general channels of Theorem 12.1 and retain observations from both orientations. For an even number of flights, the first and last impacts have the same contact type; for an odd number they have different types. Write  $\mathcal{F}_{00}, \mathcal{F}_{11}$  for the two even limiting functions and  $\mathcal{F}_{01} = \mathcal{F}_{10}$  for the odd one. Their normalization is that of (12.6), with the same free area  $A$  and multiplier  $\gamma$ . Thus they are scalar functions derived from selected physical event probabilities, not postulated independent boundary densities.

**20.1. An intrinsic positive factorization.** For each contact type put  $a_b = S_b''(0) > 0$  and introduce the signed Morse coordinate and its increasing inverse,

$$t_b(u) = \operatorname{sgn}(u) \sqrt{2S_b(u)}, \quad u = \chi_b(t).$$

The value at  $u = 0$  is defined by smooth extension. Set

$$\begin{aligned} f_b(t) &= B_b(\chi_b(t)) \chi_b'(t), \\ \mathcal{V}_b(x) &= \frac{\sqrt{a_b}}{2} \{f_b(\sqrt{2x}) + f_b(-\sqrt{2x})\}, \quad x \geq 0. \end{aligned} \tag{20.1}$$

We call  $\mathcal{V}_b$  the normalized boundary energy profile. It is a symmetrized weighted energy datum: no assertion that it determines a nonsymmetric contact graph is implicit in this terminology.

**Theorem 20.1** (Nonlinear boundary compatibility). *On a sufficiently small common closed collar the functions  $\mathcal{V}_b$  are smooth, positive, and satisfy  $\mathcal{V}_b(0) = 1$ . For  $b, c \in \{0, 1\}$  the limiting physical laws have the representation*

$$\mathcal{F}_{bc}(d) = \frac{2}{\pi d^2} \int_{\substack{x, y > 0 \\ x+y < d}} (d-x-y) \frac{\mathcal{V}_b(x) \mathcal{V}_c(y)}{\sqrt{xy}} \, dx \, dy. \tag{20.2}$$

Let  $*$  denote Volterra convolution on the positive half-line, used only at arguments in this collar, and define

$$k_b(x) = x^{-1/2} \mathcal{V}_b(x), \quad K_{bc}(d) = \frac{\pi}{2} \frac{d^2}{dd^2} \{d^2 \mathcal{F}_{bc}(d)\}. \tag{20.3}$$

Then  $K_{bc} = k_b * k_c$ ,  $K_{bc}(0) = \pi$ , and

$$K_{01} * K_{01} = K_{00} * K_{11}. \tag{20.4}$$

Within the class of smooth normalized profiles,  $\mathcal{F}_{00}$  and  $\mathcal{F}_{11}$  determine  $\mathcal{V}_0, \mathcal{V}_1$  uniquely and therefore determine the complete odd law  $\mathcal{F}_{01}$  on the same collar. The assertions hold without equal-curvature, even-graph or symmetry assumptions. All the forward identities have the mixed geometric regularity of the relative law.

*Proof.* Strict convexity makes the signed Morse coordinate smooth and increasing near zero, with derivative  $\sqrt{a_b}$  at zero. Consequently  $f_b$  is smooth, positive and has value  $a_b^{-1/2}$ . Its even part is a smooth function of  $t^2$ , by the even-function argument in Proposition 18.3. This proves the asserted properties of  $\mathcal{V}_b$ , including its parameter regularity. Reversing the choice of transverse direction merely interchanges the two terms in (20.1).

Under  $x = S_b(u) = t^2/2$ , the pushforward of  $B_b(u) \, du$  has, for  $x > 0$ , density

$$\frac{f_b(\sqrt{2x}) + f_b(-\sqrt{2x})}{\sqrt{2x}} = \sqrt{\frac{2}{a_b}} x^{-1/2} \mathcal{V}_b(x).$$

Insert this expression twice into (12.5), whose prefactor for types  $b, c$  is  $\sqrt{a_b a_c}/(\pi d^2)$ . The Hessian factors cancel and give (20.2). This calculation also proves equality of the two odd orientations.

It is valid for each even orientation by applying the relative law after interchanging the contact types.

The kernels  $k_b$  are locally integrable. Fubini's theorem rewrites (20.2) as

$$\frac{\pi}{2} d^2 \mathcal{F}_{bc}(d) = \int_0^d (d-s)(k_b * k_c)(s) \, ds.$$

The convolution is continuous down to zero: substitution  $x = ds$  gives

$$(k_b * k_c)(d) = \int_0^1 \frac{\mathcal{V}_b(ds) \mathcal{V}_c(d(1-s))}{\sqrt{s(1-s)}} \, ds,$$

whose value at zero is  $\pi$ . Differentiation proves (20.3). Associativity and commutativity of these locally integrable convolutions, again by Fubini on a bounded simplex, give (20.4). No extension of the physical law beyond the collar is used.

It remains to prove uniqueness, since a formal square root of Taylor coefficients alone would not suffice for smooth functions. Suppose  $k = x^{-1/2}V(x)$  and  $\ell = x^{-1/2}\tilde{V}(x)$ , with smooth  $V(0) = \tilde{V}(0) = 1$ , have  $k * k = \ell * \ell$  on the collar. Put  $f = k - \ell$ ,  $a = k + \ell$  and  $k_*(x) = x^{-1/2}$ . Then  $a * f = 0$  and

$$a(x) = 2x^{-1/2} + x^{1/2}w(x)$$

for a smooth  $w$ . The identity  $k_* * k_* = \pi$  gives

$$(k_* * a)(x) = 2\pi + R(x), \quad R(x) = x \int_0^1 \frac{s^{1/2}}{(1-s)^{1/2}} w(xs) \, ds.$$

Here  $R(0) = 0$  and  $R'$  is bounded. Convolving  $a * f = 0$  with  $k_*$  and differentiating its absolutely continuous integral yields

$$2\pi f + R' * f = 0$$

almost everywhere. Hence  $|f(x)| \leq C \int_0^x |f(s)| \, ds$ . The absolutely continuous function  $\int_0^x |f(s)| \, ds$  has zero initial value, and the elementary Gronwall inequality makes it zero throughout the collar. Thus  $k = \ell$ . Applying this argument to each of  $K_{00}$  and  $K_{11}$  proves uniqueness of both profiles; (20.2) then determines the odd law.  $\square$

**Theorem 20.2** (Stability of the complete boundary energy profile). *Fix  $D > 0$  in the common collar. Let  $V, \tilde{V}$  be two smooth normalized profiles, with  $V(0) = \tilde{V}(0) = 1$  and  $\|V'\|_\infty, \|\tilde{V}'\|_\infty \leq B$  on  $[0, D]$ . Let  $F, \tilde{F}$  be their same-type laws defined by (20.2), and put  $K = (\pi/2)(d^2 F)''$ ,  $\tilde{K} = (\pi/2)(d^2 \tilde{F})''$ . Then*

$$\|V - \tilde{V}\|_{C^0([0, D])} \leq \frac{D}{\pi} e^{BD/2} \|K' - \tilde{K}'\|_{C^0([0, D])} \leq C_{B, D} \|F - \tilde{F}\|_{C^3([0, D])}. \quad (20.5)$$

*Thus the two complete boundary profiles are locally Lipschitz functions of the two same-type laws on their physical image, when the latter are given this specified differentiated norm. The resulting odd law is Lipschitz in  $C^0$  on uniformly bounded profile sets.*

*Proof.* Use  $k = x^{-1/2}V$ ,  $\ell = x^{-1/2}\tilde{V}$ ,  $f = k - \ell$ ,  $a = k + \ell$ ,  $k_* = x^{-1/2}$ , and put  $Q = K - \tilde{K} = a * f$ . The convolution in the preceding proof has the more direct expression

$$(k_* * a)(x) = 2\pi + R(x), \quad R(x) = \int_0^1 \frac{V(xs) + \tilde{V}(xs) - 2}{\sqrt{s(1-s)}} \, ds.$$

It follows that  $R(0) = 0$  and

$$\|R'\|_\infty \leq 2B \int_0^1 \frac{s}{\sqrt{s(1-s)}} \, ds = \pi B.$$

Since  $Q(0) = 0$ , differentiation under a convolution, or integration by parts before differentiation, gives

$$2\pi f + R' * f = (k_* * Q)' = k_* * Q'.$$

The right side is bounded by  $2\sqrt{x}\|Q'\|_\infty$ . The normalized initial values imply that  $f$  extends continuously with  $f(0) = 0$ . Therefore

$$|f(x)| \leq \frac{\sqrt{D}}{\pi} \|Q'\|_\infty + \frac{B}{2} \int_0^x |f(s)| \, ds.$$

Gronwall's inequality gives  $\|f\|_\infty \leq \sqrt{D} e^{BD/2} \|Q'\|_\infty / \pi$ . Multiplication by  $\sqrt{x}$  proves the first estimate. For the second,

$$Q'(d) = \frac{\pi}{2} \{6(F - \tilde{F})'(d) + 6d(F - \tilde{F})''(d) + d^2(F - \tilde{F})'''(d)\},$$

which is bounded by the asserted  $C^3$  norm. Finally the integral (20.2) is a bilinear map of the profiles with a positive kernel of normalized total mass one. Subtracting its two products proves the  $C^0$  Lipschitz assertion for the odd law.  $\square$

This is stability of normalized energy profiles, not of an arbitrary nonsymmetric contact graph. It is an inverse on the image of the forward profile map; no claim that every positive scalar function is a convolution square is needed. The derivative norm in (20.5) is part of the observation statement. The unregularized theorem does not replace it by uniform error of raw probabilities. The finite-preparation theorem in Section A.6 adds regularization and supplied smoothness bounds to address that different data problem. For finite jets, an explicit polynomial inverse is available without solving a function-valued equation.

**20.2. Finite jets and observable identities.** For a smooth right germ  $h$ , write  $[d^n]h = h^{(n)}(0)/n!$ . Fix  $M \geq 0$  and work in the truncated polynomial ring  $\mathbb{R}[z]/(z^{M+1})$ . Introduce

$$\begin{aligned} \mathcal{M}_{bc}^{[M]}(z) &= \sum_{n=0}^M \frac{(n+2)!}{2} [d^n] \mathcal{F}_{bc} z^n, \\ \mathcal{A}_b^{[M]}(z) &= \sum_{n=0}^M \left(\frac{1}{2}\right)_n [x^n] \mathcal{V}_b z^n, \quad \left(\frac{1}{2}\right)_0 = 1. \end{aligned} \tag{20.6}$$

The notation  $(1/2)_n$  denotes the rising factorial. Truncation is important: even for analytic input, no convergence radius of a factorially reweighted infinite series is being asserted.

**Theorem 20.3** (Rank-one jets and stable finite-order inversion). *The transformed jets satisfy*

$$\mathcal{M}_{bc}^{[M]} = \mathcal{A}_b^{[M]} \mathcal{A}_c^{[M]}, \quad (\mathcal{M}_{01}^{[M]})^2 = \mathcal{M}_{00}^{[M]} \mathcal{M}_{11}^{[M]} \pmod{z^{M+1}}. \tag{20.7}$$

*Both profiles' order- $M$  jets are recovered from the two same-type laws by the following triangular recursion. If  $m_{bb,n} = [z^n] \mathcal{M}_{bb}^{[M]}$  and  $\alpha_{b,n} = [z^n] \mathcal{A}_b^{[M]}$ , then*

$$\alpha_{b,0} = 1, \quad \alpha_{b,n} = \frac{1}{2} \left( m_{bb,n} - \sum_{r=1}^{n-1} \alpha_{b,r} \alpha_{b,n-r} \right), \quad [x^n] \mathcal{V}_b = \frac{\alpha_{b,n}}{(1/2)_n}. \tag{20.8}$$

*For fixed  $M$  these reconstruction maps, and the resulting prediction of the odd jet, are Lipschitz on bounded sets of input coefficients. They are smooth with respect to geometric parameters on the compact families of the relative law.*



*Proof.* Write  $v_{b,r} = [x^r]\mathcal{V}_b$ . The beta integral on the simplex gives, for  $r, s \geq 0$ ,

$$\int_{x,y>0, x+y<d} (d-x-y)x^{r-1/2}y^{s-1/2} dx dy = \frac{\pi(1/2)_r(1/2)_s}{(r+s+2)!} d^{r+s+2}.$$

This follows by first setting  $x = dX, y = dY$  and integrating in  $Y$ , then in  $X$ ; both integrals are the ordinary beta integral. Taylor's integral remainders, bounded on the scaled simplex, justify coefficient extraction for smooth rather than only analytic profiles. Thus

$$[d^n]\mathcal{F}_{bc} = \frac{2}{(n+2)!} \sum_{r+s=n} (1/2)_r(1/2)_s v_{b,r} v_{c,s}. \quad (20.9)$$

Multiplying by  $(n+2)!/2$  gives the first identity in (20.7); its determinant identity is an immediate polynomial consequence. All constant coefficients are one. Comparing the degree- $n$  coefficient of  $\mathcal{M}_{bb}^{[M]} = (\mathcal{A}_b^{[M]})^2$  leaves  $2\alpha_{b,n}$  as the only unknown term, proving (20.8). Induction makes every output coefficient a polynomial in finitely many input coefficients, divided only by fixed positive numbers. Its derivative is bounded on each bounded coefficient box, proving the Lipschitz assertion. Finite parameter differentiation of these polynomial identities proves the last assertion without any infinite-order uniformity claim.  $\square$

**Corollary 20.4** (The first two nonlinear compatibility identities). *With primes denoting right offset derivatives at zero,*

$$2\mathcal{F}'_{01} = \mathcal{F}'_{00} + \mathcal{F}'_{11}, \quad (20.10)$$

$$\mathcal{F}''_{01} = \frac{\mathcal{F}''_{00} + \mathcal{F}''_{11}}{2} - \frac{3}{16}(\mathcal{F}'_{00} - \mathcal{F}'_{11})^2. \quad (20.11)$$

*In particular the odd second derivative is at most the mean of the two even second derivatives, with the exact deficit displayed above.*

*Proof.* The first two nonconstant coefficients of  $\mathcal{M}_{bc}$  are  $3\mathcal{F}'_{bc}$  and  $6\mathcal{F}''_{bc}$ . Equate degrees one and two in  $\mathcal{M}_{01}^2 = \mathcal{M}_{00}\mathcal{M}_{11}$  and substitute the degree-one equality in the degree-two equality. This gives the two formulas, including the negative square coefficient.  $\square$

These identities concern nonlinear offset coefficients. The zeroth-order condition merely says that every normalized law has value one. Neither the zeroth-order condition nor the multiplier alone determines the first two coefficients. The geometric variation in Proposition A.3.1 and the realization in Theorem A.3.2 exhibit nonconstant boundary data at fixed leading selected-channel invariants. The compatibility theorem organizes such data without imposing the identical-even hypothesis of the separate graph-jet inverse in the appendix.

**20.3. Finite physical records and a quantitative defect.** The preceding identities are exact statements about the limiting functions. The relative law also controls how they can be read from long but finite records. For  $n \geq 1$  define

$$\begin{aligned} F_{bb,n}(d) &= \frac{2A \sinh(2n\gamma)}{d^2} \Pr(E_{2n,e_b}(d)), \quad e_0 = e, \quad e_1 = \bar{e}, \\ F_{01,n}(d) &= \frac{2A \sinh((2n+1)\gamma)}{d^2} \Pr(E_{2n+1,e}(d)). \end{aligned} \quad (20.12)$$

They are extended to  $d = 0$  by their exact smooth right limits. These definitions use the selected patch labels and the normalization  $A, \gamma$ ; they are not unlabelled binary data with unknown normalizer. Estimation of these auxiliary quantities is a separate observation problem.

**Corollary 20.5** (Finite-record compatibility and profile-jet recovery). *Let  $M$  and the number of geometric derivatives be fixed. Apply the transform (20.6) to the three finite laws in (20.12), obtaining  $\mathcal{M}_{bc,n}^{[M]}$ . With the maximum coefficient norm on truncated polynomials,*

$$\|(\mathcal{M}_{01,n}^{[M]})^2 - \mathcal{M}_{00,n}^{[M]}\mathcal{M}_{11,n}^{[M]}\| \leq C_M \tau^{2n}. \quad (20.13)$$

*The same estimate holds for each fixed mixed geometric derivative. Suppose the measured order- $M$  jets of the two same-type finite laws have coefficient error at most  $\epsilon$ , with their known constant coefficients fixed to one. The recursion (20.8) reconstructs both boundary profile jets, and predicts the odd limiting jet, with error*

$$C_M(\epsilon + \tau^{2n}). \quad (20.14)$$

*For every fixed subcollar  $[0, D] \subset [0, d_*]$ , the finite kernels  $K_{bc,n} = (\pi/2)(d^2 F_{bc,n})''$  also satisfy*

$$\|K_{01,n} * K_{01,n} - K_{00,n} * K_{11,n}\|_{C^0([0,D])} \leq C_D \tau^{2n}.$$

*Proof.* Theorem 12.1, applied to each orientation and parity, bounds  $F_{bc,n} - \mathcal{F}_{bc}$  in every fixed  $C^k$  norm by  $C_k \tau^{2n}$ . Coefficient extraction and the finite transform preserve this order. The determinant in (20.7) is a polynomial in bounded coefficients; subtract its value at the limiting jets to prove (20.13). The product rule proves the parameter derivative assertion. Apply the finite polynomial inverse of Theorem 20.3 to the two measured jets to get (20.14); this is also defined for noisy jets outside the physical image. Finally the relative  $C^2$  estimate bounds  $K_{bc,n} - K_{bc}$  uniformly on  $[0, D]$ . For bounded functions there,  $\|f * g\|_\infty \leq D\|f\|_\infty\|g\|_\infty$ . Telescoping the four products and using (20.4) proves the last claim.  $\square$

The two oriented even laws suffice to reconstruct both normalized boundary energy profiles and to predict the odd law. Odd records are needed for the independent compatibility test, not as a third datum for this reconstruction. Neither conclusion recovers full collision arrays or an arbitrary billiard table. Unlike the fixed-flight parametric inverse below, the compatibility defect and its vanishing rate use the long-bridge relative theorem directly. Section A.6 supplies a regularized finite-preparation version of the full-profile inverse, with the required normalization and smoothness bounds stated explicitly.

### Part 3. Analytic and intrinsic multichannel consequences

#### 21. DETERMINATION OF TWO INDEPENDENT EVEN CONTACTS

The boundary invariant has a geometric inverse on a class substantially larger than the identical-contact class. In this section each facing graph is even in the common transverse coordinate, but the two graphs are otherwise unrelated. Their curvatures may be different or equal. This is an additional hypothesis for geometric determination, not for Theorems 12.1 and 20.1. The gap and the two labelled contact curvatures are fixed. They are not inferred separately from the single multiplier.

Write, for  $b = 0, 1$ ,

$$\psi_b(y) = \frac{\kappa_b}{2} y^2 + \sum_{m \geq 2} \frac{q_{b,2m}}{(2m)!} y^{2m}, \quad c_b = 1 + g\kappa_b, \quad c = \sqrt{c_0 c_1}, \quad \gamma = \operatorname{arcosh} c, \quad \varrho = e^{-\gamma}. \quad (21.1)$$

For smooth graphs the series specifies jets only. Put

$$a_b = \frac{c \sinh \gamma}{g c_{1-b}}, \quad f_{b,k} = [d^k] \mathcal{F}_{bb}(d), \quad q_m = (q_{0,2m}, q_{1,2m})^t, \quad f_{m-1} = (f_{0,m-1}, f_{1,m-1})^t.$$

The two even laws are obtained by observing the two orientations of one selected physical channel. In particular, their indices retain contact labels even when the curvatures coincide.

**Theorem 21.1** (Block-triangular two-contact inverse). *For every  $m \geq 2$ , the vector  $f_{m-1}$  depends only on  $g, \kappa_0, \kappa_1, q_2, \dots, q_m$ , and is affine in  $q_m$ . Its derivative with respect to that last vector is*

$$\frac{\partial f_{m-1}}{\partial q_m} = - \begin{pmatrix} P_m & Q_m \\ Q_m & P_m \end{pmatrix} \text{diag}(K_{m,0}, K_{m,1}), \quad K_{m,b} = \frac{4}{2^m(m+1)(m!)^2 a_b^m}, \quad (21.2)$$

where

$$\begin{aligned} E_m &= \frac{1}{e^{4(m-1)\gamma} - 1} - \frac{1}{e^{4m\gamma} - 1}, \\ O_m &= \frac{1}{2 \sinh(2(m-1)\gamma)} - \frac{1}{2 \sinh(2m\gamma)}, \\ P_m &= \coth(2m\gamma) + 2mE_m, \quad Q_m = \text{csch}(2m\gamma) + 2mO_m. \end{aligned} \quad (21.3)$$

Both eigenvalues of the symmetric factor are positive. More explicitly,

$$P_m + Q_m = \coth(m\gamma) + 2m \left\{ \frac{1}{e^{2(m-1)\gamma} - 1} - \frac{1}{e^{2m\gamma} - 1} \right\}, \quad (21.4)$$

$$P_m - Q_m = m \tanh((m-1)\gamma) - (m-1) \tanh(m\gamma) > 0. \quad (21.5)$$

Consequently, for every finite  $M \geq 2$ , the map

$$(q_2, \dots, q_M) \mapsto (f_1, \dots, f_{M-1}) \quad (21.6)$$

has a block-triangular analytic inverse on its image. The inverse is locally Lipschitz on compact finite-jet sets and positive compact  $g, \kappa_0, \kappa_1$  boxes. These bounds require no separation of  $\kappa_0$  and  $\kappa_1$ , and no separation of the higher jets.

Equality of the two normalized even-law germs, for two pairs of real-analytic even graphs at the same labelled  $g, \kappa_0, \kappa_1$ , implies equality of both graph germs. For smooth compact parameter families, replacing the limiting coefficients by their even finite-flight coefficients recovers the jet vector through order  $2M$  with error at most  $C_M \tau^j$ . No bound uniform in  $M$  is asserted.

### 21.1. The two boundary contributions.

*Proof.* We first establish the last-jet derivative; the resulting block invertibility will then give the inverse. Set  $\mathbf{r}_b = \sqrt{c_b/c_{1-b}}$ . The linear half-line orbit starting at type  $b$  is

$$x_i^{(b)}(u) = t_i^{(b)} u + O(u^3 \rho^i), \quad t_i^{(b)} = \begin{cases} \varrho^i, & i \text{ even}, \\ \mathbf{r}_b \varrho^i, & i \text{ odd}. \end{cases} \quad (21.7)$$

The cubic remainder follows from evenness and the half-line contraction of Lemma 17.1. Notice that  $a_b/a_{1-b} = \mathbf{r}_b^2$ .

For clarity, finite-jet dependence can be proved before making any analyticity assumption on the graphs. Differentiate the stationary half-line equations at zero in their weighted sequence space. At each order, the highest orbit derivative is multiplied by the same positive quadratic Jacobi operator. Its inverse is (17.1); all remaining terms contain lower orbit jets and finite graph jets. This gives an induction in the endpoint degree. The weighted derivative estimates justify the differentiated action sums and the differentiated trace-class determinant series. The coefficient of  $d^{m-1}$  in the residual-time integral uses the action through degree  $2m$  and the amplitude through degree  $2m-2$ . Indeed, scaling the two endpoints by  $\sqrt{d}$  gives degrees  $d^{n/2-1}$  for an action term of endpoint degree  $n$  and  $d^{n/2}$  for an amplitude term. The common Morse coordinates justify this calculation with smooth remainders down to  $d=0$ . Thus only graph jets through order  $2m$  can enter.

Keep all lower jets fixed and vary  $q_{r,2m}$ . The first varying homogeneous term of a flight length is the sum of the corresponding pure endpoint powers divided by  $(2m)!$ . Its mixed derivative is

zero. In the stationary action, a boundary site is counted once and an interior site twice. The envelope identity therefore gives, with  $s_{b,2m} = S_b^{(2m)}(0)$ ,

$$\begin{aligned}\partial_{q_{b,2m}} s_{b,2m} &= 1 + 2 \sum_{k \geq 1} \varrho^{4mk} = \coth(2m\gamma), \\ \partial_{q_{1-b,2m}} s_{b,2m} &= 2\mathfrak{r}_b^{2m} \sum_{k \geq 0} \varrho^{2m(2k+1)} = \mathfrak{r}_b^{2m} \operatorname{csch}(2m\gamma).\end{aligned}\tag{21.8}$$

This can equally be computed on finite stationary segments and then passed to the half-line. Every differentiated sum is absolutely convergent. Orbit corrections of degree at least three cannot change the degree- $2m$  terms displayed here.

Write  $\beta_{b,2m-2} = [u^{2m-2}]B_b(u)$ , using ordinary coefficients. At this degree the edge-product variation vanishes, since the varying length term is pure. The varying interior Hessian is diagonal, with its  $i$ th diagonal entry, at a site of varied type, equal to

$$\frac{2(t_i^{(b)})^{2m-2} u^{2m-2}}{(2m-2)!}.$$

The variation of the stationary orbit starts in degree  $2m-1$ . An unvaried Hessian perturbation starts in degree two. Hence orbit variation does not affect the degree under consideration, and only the linear term of the logarithmic determinant contributes there. The diagonal half-line Green entry is

$$G_{ii}^{(b)} = \frac{1 - \varrho^{2i}}{2a_{(b+i) \bmod 2}}.$$

It follows that

$$\begin{aligned}\partial_{q_{b,2m}} \beta_{b,2m-2} &= -\frac{E_m}{a_b(2m-2)!}, \\ \partial_{q_{1-b,2m}} \beta_{b,2m-2} &= -\frac{\mathfrak{r}_b^{2m} O_m}{a_b(2m-2)!}.\end{aligned}\tag{21.9}$$

For the first formula sum  $(1 - \varrho^{4k})\varrho^{4(m-1)k}$  over  $k \geq 1$ . For the second sum  $(1 - \varrho^{2(2k+1)})\varrho^{(2m-2)(2k+1)}$  over  $k \geq 0$  and use  $a_{1-b}^{-1} = \mathfrak{r}_b^2 a_b^{-1}$ . The sums are respectively  $E_m$  and  $O_m$ . Passing from  $\log B_b$  to  $B_b$  does not change the first varying coefficient: multiplication by a positive-degree unvaried term would raise its degree.

*Residual time and the invertible block.* For the same-type law both endpoint quadratic actions have coefficient  $a_b$ . Under the density proportional to  $(d - a_b(u^2 + v^2)/2)_+$  one has

$$\begin{aligned}\int (d - a_b(u^2 + v^2)/2)_+ du dv &= \frac{\pi d^2}{a_b}, \\ \mathbb{E}(u^{2k}) &= \frac{2 \binom{2k}{k}}{2^k(k+1)(k+2)a_b^k} d^k.\end{aligned}\tag{21.10}$$

These identities follow from the angular average  $\binom{2k}{k}/4^k$  and a one-dimensional radial integration. First variation of the action has no moving-boundary term, because the residual-time weight is zero at the boundary. Its contribution from  $s_{b,2m}$  to  $f_{b,m-1}$  is  $-K_{m,b}s_{b,2m}$ . The amplitude contribution from  $\beta_{b,2m-2}$  has coefficient

$$\frac{4 \binom{2m-2}{m-1}}{2^{m-1}m(m+1)a_b^{m-1}}.$$

Combining these coefficients with (21.8)–(21.9) gives

$$\partial_{q_{b,2m}} f_{b,m-1} = -K_{m,b} P_m, \quad \partial_{q_{1-b,2m}} f_{b,m-1} = -K_{m,b} \mathfrak{r}_b^{2m} Q_m.$$

But  $K_{m,b} \mathfrak{r}_b^{2m} = K_{m,1-b}$ . This proves (21.2), including the unequal-curvature factors. Lower nonlinear action and amplitude terms can multiply a first variation only at higher endpoint degree. Thus the computation remains valid at arbitrary fixed lower even jets. The last-jet derivatives are constant in that jet, proving the asserted affinity.

Both  $E_m$  and  $O_m$  are positive. The identities

$$\coth(2x) + \operatorname{csch}(2x) = \coth x, \quad \coth(2x) - \operatorname{csch}(2x) = \tanh x$$

and the geometric sums give (21.4). For the difference,

$$E_m - O_m = \frac{1}{2} \{ \tanh((m-1)\gamma) - \tanh(m\gamma) \},$$

which proves (21.5). The function  $\tanh x$  is strictly concave for  $x > 0$  and has value zero at zero. Applying strict concavity at  $(m-1)\gamma = ((m-1)/m)m\gamma$  proves the strict inequality in (21.5). Thus

$$\det \frac{\partial f_{m-1}}{\partial q_m} = K_{m,0} K_{m,1} (P_m + Q_m) (P_m - Q_m) > 0.$$

There is no denominator involving  $\kappa_0 - \kappa_1$ .

To write the inverse explicitly, suppose all lower graph jets are already known and let  $R_m$  denote their contribution to  $f_{m-1}$ . For  $z_m = R_m - f_{m-1}$ , solve

$$q_m = \operatorname{diag}(K_{m,0}^{-1}, K_{m,1}^{-1}) \frac{1}{P_m^2 - Q_m^2} \begin{pmatrix} P_m & -Q_m \\ -Q_m & P_m \end{pmatrix} z_m. \quad (21.11)$$

The recursion proves uniqueness at every finite order. Its coefficients are analytic in the positive geometric parameters and in lower finite jets: the operations are local length expansions, inversion of a positive quadratic operator, and convergent geometric sums. On any fixed compact set and finite collection of orders, the positive denominators have positive minima. The recursion and its derivatives therefore give the stated locally Lipschitz inverse.

For analytic graphs, equality of all recovered even jets implies equality of the two convergent graph series; their odd jets vanish by assumption. For smooth families, the differentiated relative theorem gives  $|f_{j,b,k} - f_{b,k}| \leq C_M \tau^j$  through  $k = M-1$ . Apply the finite-order Lipschitz inverse. For large  $j$  its inputs stay in the chosen compact neighborhood. The algebraic recursion is defined for every finite coefficient vector at fixed positive geometric data; the finitely many remaining even  $j$  are covered by increasing the constant. This completes the proof.  $\square$

**Corollary 21.2** (Labelled analytic boundary determination). *In the two-independent-even contact class, the gap, the two labelled contact curvatures, and the two normalized even-law germs determine both real-analytic contact graphs. If their participating obstacles have connected real-analytic closed boundaries, each boundary is determined in its contact frame. The conclusion allows equal contact curvatures and does not require congruent obstacles or identical graphs. It makes no determination of unobserved obstacles or of the lattice.*

*Proof.* The theorem determines each analytic graph germ. To compare two candidate closed boundaries at one contact, parametrize both by oriented arclength from that contact and its fixed tangent. Their curvatures are real-analytic functions on the real line, continued around the closed curves, and agree on an interval. The identity theorem makes them agree everywhere. Uniqueness for the planar Frenet equations with the common initial frame makes the two boundary images identical. Apply the argument at each participating contact. This is exact analytic continuation; no stability of distant boundary points from noisy local data is asserted.  $\square$

At equal curvatures, restricting (21.2) to the diagonal  $q_{0,2m} = q_{1,2m}$  recovers exactly Theorem A.4.1. The new direction is the antisymmetric one, whose nonvanishing coefficient is (21.5). A single identical-contact calculation does not see that coefficient. Without the evenness assumption, odd graph jets also enter the symmetrized energy datum. The theorem does not identify their fibres; the general relative law and full energy-profile inverse remain valid without that assumption.

## 22. SIGNED ENDPOINT RIGIDITY AND THE ALL-ORDER DETERMINANT-ONE INVERSE

The scalar onset laws integrate out the signs of the endpoint coordinates and therefore retain only a symmetrized profile. The physical endpoint record is strictly richer. If the signed endpoint positions are retained, the support recovers the two half-line actions before symmetrization. We prove here that those two actions determine the complete labelled contact jets. The proof is written at the level needed by the inverse theorem: the nonlinear half-line stationarity operator is inverted quantitatively in a weighted Banach space, the infinite action is differentiated through finite truncations, and the first occurrence of each new graph jet is isolated by homogeneous degree.

Fix oriented transverse coordinates at the two labelled contacts. For  $b \in \{0, 1\}$  use an even number of flights and the orientation for which the first and last impacts have type  $b$ . Let  $\Pi_{bb,d}^{\text{end},\infty}$  be the limiting conditional endpoint law after residual time has been integrated out. By Theorem 12.1, on a common endpoint box its interior density is proportional to

$$(d - S_b(u) - S_b(v))_+ B_b(u) B_b(v), \quad (22.1)$$

where  $B_b > 0$ .

**Theorem 22.1** (All-order signed-endpoint rigidity). *Let the two labelled facing contacts be  $C^\infty$  and strictly convex. No reflection symmetry, equality of the two contacts, evenness, or supplied contact curvature is assumed. Suppose the selected channel and its signed endpoint coordinates are retained.*

*There is  $d_* > 0$  such that the two families*

$$\{\Pi_{00,d}^{\text{end},\infty} : 0 < d < d_*\}, \quad \{\Pi_{11,d}^{\text{end},\infty} : 0 < d < d_*\}$$

*determine the half-line actions  $S_0, S_1$  on a fixed contact neighborhood. If physical time is retained, the onset determines the gap  $g$ . Hence the gap, both contact curvatures, and every finite labelled graph jet are determined by the signed endpoint laws.*

*In the fixed oriented contact frames set  $q_{b,n} = \psi_b^{(n)}(0)$ ,  $q_{b,2} = \kappa_b$ , and  $q_n = (q_{0,n}, q_{1,n})^t$ . For each finite  $M \geq 2$  Taylor's formula gives*

$$\psi_b(y) = \sum_{n=2}^M \frac{q_{b,n}}{n!} y^n + R_{b,M}(y), \quad |\partial_y^r R_{b,M}(y)| \leq C_M |y|^{M+1-r}, \quad 0 \leq r \leq M+1. \quad (22.2)$$

*Here the bound uses a  $C^{M+1}$  bound on a common contact interval. For smooth contacts no convergence of an infinite Taylor series is asserted. Equal finite graph jets, including arbitrary smooth remainders, are treated in Lemma 22.4 below. Put*

$$a_b = S_b''(0), \quad s_n = (S_0^{(n)}(0), S_1^{(n)}(0))^t.$$

*Then*

$$\gamma = \text{arsinh}(g\sqrt{a_0 a_1}), \quad c = \cosh \gamma, \quad c_0 = c\sqrt{a_0/a_1}, \quad c_1 = c\sqrt{a_1/a_0}, \quad \kappa_b = \frac{c_b - 1}{g}. \quad (22.3)$$

For every  $n \geq 3$ ,  $s_n$  depends only on the leading geometry and  $q_3, \dots, q_n$ , is affine in  $q_n$ , and

$$D_{q_n} s_n = M_n := \begin{pmatrix} \coth(n\gamma) & \mathfrak{r}_0^n \operatorname{csch}(n\gamma) \\ \mathfrak{r}_1^n \operatorname{csch}(n\gamma) & \coth(n\gamma) \end{pmatrix}, \quad \mathfrak{r}_b = \sqrt{c_b/c_{1-b}}. \quad (22.4)$$

The determinant of  $M_n$  is exactly one. For every fixed  $M$ , the finite signed-jet forward map through order  $M$  is  $C^1$  with an everywhere invertible differential; on compact positive geometric and finite-jet sets its least singular value is bounded away from zero. Its inverse is locally Lipschitz there.

For real-analytic contacts, equality of the two signed endpoint-law germs as  $d \downarrow 0$  and of the physical onset determines the two labelled analytic contact germs. If the participating obstacle boundaries are connected and real analytic, their boundary images are determined in the fixed contact frames. For a smooth compact family and every fixed  $M$ , the corresponding finite even-flight support families recover both graph jets through order  $M$  with error  $O_M(\tau^j)$ .

### 22.1. Support thresholds and the leading geometry.

*Recovery of the actions and the quadratic geometry.* Positivity of  $B_b$  makes the interior support in (22.1) exactly

$$\{(u, v) : S_b(u) + S_b(v) < d\}.$$

Define

$$T_b(u, v) = \inf\{d > 0 : (u, v) \in \operatorname{int} \operatorname{supp} \Pi_{bb,d}^{\operatorname{end}, \infty}\}.$$

Then

$$T_b(u, v) = S_b(u) + S_b(v), \quad S_b(u) = T_b(u, 0), \quad (22.5)$$

because  $S_b(0) = 0$ . Thus the signed support recovers the unsymmetrized action.

For a selected  $j$ -flight itinerary the infimum of physical times at which the near-onset event has positive probability is  $jg$ , hence the time axis gives  $g$ . The half-line Hessian identity

$$a_b = \frac{c \sinh \gamma}{g c_{1-b}}, \quad c = \sqrt{c_0 c_1} = \cosh \gamma$$

gives

$$a_0 a_1 = \frac{\sinh^2 \gamma}{g^2}, \quad \frac{a_0}{a_1} = \frac{c_0}{c_1},$$

which is equivalent to (22.3). No curvature is supplied as an external parameter.  $\square$

**22.2. A quantitative weighted inverse for the nonlinear half-line.** Let  $\ell_b(y, z)$  be the local flight length from a point of type  $b$  with transverse coordinate  $y$  to the facing point of type  $1 - b$  with coordinate  $z$ . For a half-line beginning at type  $b$ , set  $x_0 = u$  and, for  $i \geq 1$ , write

$$\mathcal{F}_{b,i}(x, u; q) := \partial_2 \ell_{b+i-1}(x_{i-1}, x_i) + \partial_1 \ell_{b+i}(x_i, x_{i+1}) = 0, \quad (22.6)$$

where contact types are read modulo two. Fix a compact positive geometric set. Its hyperbolic exponent has a positive minimum  $\gamma_- > 0$ . Choose

$$e^{-\gamma_-} < \rho < 1, \quad (22.7)$$

and define

$$X_\rho = \left\{ x = (x_i)_{i \geq 1} : \|x\|_\rho := \sup_{i \geq 1} \rho^{-i} |x_i| < \infty \right\}.$$

**Lemma 22.2** (Weighted Green bound and nonlinear inverse). *Let  $H_b^0(q) = D_x \mathcal{F}_b(0, 0; q)$  be the quadratic half-line Jacobi operator. Its inverse has kernel*

$$G_{ik}^{(b)} = \frac{g\sigma_i^{(b)}\sigma_k^{(b)}}{2c\sinh\gamma} \left( e^{-\gamma|i-k|} - e^{-\gamma(i+k)} \right), \quad i, k \geq 1. \quad (22.8)$$

*There is  $C_0 < \infty$ , uniform on the compact geometric set, such that*

$$\|G^{(b)}y\|_\rho \leq C_0\|y\|_\rho. \quad (22.9)$$

*After shrinking a common endpoint box, the unique stationary half-line  $x^{(b)}(u; q)$  satisfies*

$$\|x^{(b)}(u; q)\|_\rho \leq C_1|u|, \quad (22.10)$$

*and, with  $H_b(u, q) = D_x \mathcal{F}_b(x^{(b)}(u; q), u; q)$ ,*

$$\|H_b(u, q) - H_b^0(q)\|_{X_\rho \rightarrow X_\rho} \leq C_2|u|. \quad (22.11)$$

*Consequently, after choosing  $u_*$  so that  $C_0C_2u_* \leq 1/2$ ,*

$$\|H_b(u, q)^{-1}\|_{X_\rho \rightarrow X_\rho} \leq 2C_0, \quad |u| \leq u_*. \quad (22.12)$$

*For every fixed finite collection of endpoint and geometric derivatives, the corresponding derivatives of  $x^{(b)}$  are uniformly bounded in  $X_\rho$ .*

*Proof.* The factors  $\sigma_i^{(b)}$  take only two positive values on the compact geometric set, so (22.8) gives  $|G_{ik}^{(b)}| \leq Ce^{-\gamma-|i-k|}$ . If  $|y_k| \leq \|y\|_\rho \rho^k$ , then for  $k \leq i$

$$e^{-\gamma-(i-k)}\rho^k = \rho^i(e^{-\gamma-}/\rho)^{i-k},$$

while for  $k \geq i$

$$e^{-\gamma-(k-i)}\rho^k = \rho^i(\rho e^{-\gamma-})^{k-i}.$$

Both geometric series are uniformly summable because of (22.7); the reflected term in (22.8) is smaller. This proves (22.9).

The nonlinear part of (22.6) is local and vanishes quadratically at the zero orbit. On a weighted ball  $\|x\|_\rho \leq R|u|$ , its derivative differs from  $H_b^0$  by at most  $C(|u| + \|x\|_\rho)$  in  $X_\rho$  operator norm. The Green bound therefore makes the standard fixed-point map a contraction on a common ball and gives (22.10). Substitution of that bound gives (22.11). Factor

$$H_b(u, q) = H_b^0(q)[I + G^{(b)}(H_b(u, q) - H_b^0(q))].$$

The bracket differs from the identity by an operator of norm at most  $1/2$ , so its Neumann series converges and proves (22.12).

Differentiate the fixed-point or stationarity equation. At every order the highest orbit derivative is multiplied by the same operator  $H_b(u, q)$ ; all remaining terms are finite sums of local coefficient derivatives and lower orbit derivatives. The uniform inverse bound, the weighted product estimate, and induction give the last assertion. This is the quantitative version of the half-line construction in Lemma 17.1.  $\square$

**22.3. Derivative bookkeeping, the infinite action, and the envelope limit.** The stationarity equations contain first derivatives of the local flight lengths. Hence differentiating them  $k$  times can expose graph derivatives of order  $k+1$ . We do not infer the jet order of the action from that stationarity count. Instead the action filtration is proved directly from its first variation.

Set

$$\mathcal{A}_b(u, q) = \sum_{i \geq 0} [\ell_{b+i}(x_i^{(b)}(u; q), x_{i+1}^{(b)}(u; q)) - g]. \quad (22.13)$$

The weighted orbit bounds imply absolute and locally uniform convergence of this series and of every fixed finite collection of its mixed derivatives. Its value is  $S_b(u)$ .



**Lemma 22.3** (Finite-truncation envelope identity). *Fix a finite jet coordinate  $q_{r,m}$ , realized by a smooth local graph variation which equals  $y^m/m!$  in a neighborhood containing the statistical contact patch. Then, uniformly for  $|u| \leq u_*$ ,*

$$\partial_{q_{r,m}} S_b(u) = \sum_{i \geq 0} \partial_{q_{r,m}} \ell_{b+i}(x_i^{(b)}(u), x_{i+1}^{(b)}(u)). \quad (22.14)$$

*The series and every fixed finite number of  $u$ -derivatives converge absolutely. In particular no derivative of the stationary orbit appears in the right-hand side.*

*Proof.* Let

$$A_{b,N}(u, q) = \sum_{i=0}^{N-1} [\ell_{b+i}(x_i^{(b)}, x_{i+1}^{(b)}) - g].$$

Differentiate with respect to  $q_{r,m}$ . The coefficients of  $\partial_{q_{r,m}} x_i$  telescope into the stationarity equations for  $1 \leq i \leq N-1$ . They vanish. The only orbit-variation term left is the right boundary term

$$R_N = \partial_2 \ell_{b+N-1}(x_{N-1}^{(b)}, x_N^{(b)}) \partial_{q_{r,m}} x_N^{(b)}.$$

The first factor is  $O(|u|\rho^N)$  by stationarity and the local Taylor expansion at the zero orbit; the second is  $O(\rho^N)$  after any prescribed finite collection of  $u$ -derivatives, by Lemma 22.2. Hence

$$|\partial_u^k R_N| \leq C_k \rho^{2N} \quad (22.15)$$

for every fixed  $k$  (after harmless adjustment at  $u = 0$ ). Thus the boundary term tends to zero uniformly with the derivatives used below.

The direct parameter derivatives of the  $i$ th flight are bounded by a fixed multiple of a summable power of  $\rho^i$ , because both endpoints of that flight are  $O(|u|\rho^i)$  up to a shift of one index. The Weierstrass test therefore permits passage to the limit  $N \rightarrow \infty$  and termwise finite-order differentiation. This proves (22.14) and the asserted tail control.  $\square$

**22.4. Finite smooth remainders and factorization through jets.** The coordinate variation in Lemma 22.3 alone does not justify independence from all smooth remainders. The next lemma treats the remainder as a function, before choosing any finite jet coordinates.

**Lemma 22.4** (Smooth finite-jet factorization). *Fix  $M \geq 2$ . Let  $\psi_0^{[0]}, \psi_1^{[0]}$  and  $\psi_0^{[1]}, \psi_1^{[1]}$  be two smooth strictly convex anchored graph pairs with the same gap and the same graph jets through order  $M$ . On a common sufficiently small interval their half-line actions satisfy*

$$|S_b^{[1]}(u) - S_b^{[0]}(u)| \leq C_M |u|^{M+1} \sum_{i \geq 0} \rho^{(M+1)i} = \frac{C_M}{1 - \rho^{M+1}} |u|^{M+1}, \quad b \in \{0, 1\}. \quad (22.16)$$

*Consequently  $j_0^M S_b^{[1]} = j_0^M S_b^{[0]}$ . The constants and contact interval can be chosen uniformly with positive gap and curvature margins and bounded  $C^{M+2}$  graph norms on a common initial interval. For additional differentiated estimates, the corresponding higher smooth bounds are required. Bounds on a finite list of jet coefficients alone do not supply these functional remainder bounds.*

*Proof.* Put  $\Delta\psi_r = \psi_r^{[1]} - \psi_r^{[0]}$  and  $\psi_{r,t} = \psi_r^{[0]} + t\Delta\psi_r$ ,  $0 \leq t \leq 1$ . Strict convexity and the anchor conditions persist along this interpolation on a common contact interval. Equal quadratic jets and the common gap fix the quadratic half-line operator. The contraction proof of Lemma 22.2, uniformly in  $t$ , gives stationary orbits  $x^{(b,t)}$  with

$$|x_i^{(b,t)}(u)| \leq C|u|\rho^i, \quad i \geq 0, \quad x_0^{(b,t)} = u. \quad (22.17)$$

The same implicit-function argument gives  $C^M$  dependence on  $u$  and smooth dependence on  $t$  at the finite orders used here. In particular  $\|\partial_t x^{(b,t)}\|_\rho$  is uniformly bounded. These assertions

are local statements about graph pairs; no globally realized interpolation of periodic tables is required.

For a flight of starting type  $r$ , write

$$h_{r,t}(y, z) = g + \psi_{r,t}(y) + \psi_{1-r,t}(z), \quad \ell_{r,t}(y, z) = \sqrt{h_{r,t}(y, z)^2 + (z - y)^2}.$$

Taylor's remainder bound for the equal  $M$ -jets gives  $|\Delta\psi_r(y)| \leq C_M|y|^{M+1}$ . At fixed endpoints,

$$\partial_t \ell_{r,t}(y, z) = \frac{h_{r,t}(y, z)}{\ell_{r,t}(y, z)} \{\Delta\psi_r(y) + \Delta\psi_{1-r}(z)\}, \quad |\partial_t \ell_{r,t}(y, z)| \leq C_M(|y|^{M+1} + |z|^{M+1}). \quad (22.18)$$

The ratio is between zero and one since the separation is positive.

Apply finite-truncation differentiation to

$$A_{b,N,t}(u) = \sum_{i=0}^{N-1} [\ell_{b+i,t}(x_i^{(b,t)}, x_{i+1}^{(b,t)}) - g].$$

Interior orbit derivatives cancel by stationarity, and  $\partial_t x_0 = 0$ . Thus its derivative is the sum of the direct variations in (22.18), plus exactly the boundary term

$$B_{N,t} = \partial_2 \ell_{b+N-1,t}(x_{N-1}^{(b,t)}, x_N^{(b,t)}) \partial_t x_N^{(b,t)}.$$

The first factor is  $O(|u|\rho^N)$  and the second is  $O(\rho^N)$ ; hence  $B_{N,t} \rightarrow 0$  uniformly in  $t$  and on the endpoint interval. By (22.17), the direct variations are bounded by  $C_M|u|^{M+1}\rho^{(M+1)i}$ , a summable majorant. Instead of assuming an interchange of infinite differentiation, integrate the exact finite identity in  $t$  and then let  $N \rightarrow \infty$ . Uniform action convergence, the vanishing boundary term and dominated convergence give

$$S_b^{[1]}(u) - S_b^{[0]}(u) = \int_0^1 \sum_{i \geq 0} \partial_t \ell_{b+i,t}(x_i^{(b,t)}, x_{i+1}^{(b,t)}) dt.$$

This proves (22.16). The actions are  $C^M$  by the weighted construction. Taylor's formula then implies equality of their derivatives at zero through order  $M$ : a nonzero Taylor coefficient of degree at most  $M$  would contradict the displayed  $O(|u|^{M+1})$  bound. Every constant used is uniform under the stated functional bounds and positive margins.  $\square$

In particular, the finite map from graph jets to action jets is independent of the chosen smooth representatives. One may compute it on polynomial graph representatives with smooth cutoffs on a common small interval. On a compact positive finite-jet set these representatives and intervals can be chosen with uniform finite smooth bounds. This is the finite-dimensional map used below, not an identification of a smooth function with its formal Taylor series.

A flat variation illustrates the distinction. For sufficiently small  $\epsilon$  and a smooth cutoff equal to one near zero, the two graphs

$$\psi^{[0]}(y) = \frac{\kappa}{2}y^2, \quad \psi^{[1]}(y) = \frac{\kappa}{2}y^2 + \begin{cases} \epsilon e^{-1/y^2}, & y \neq 0, \\ 0, & y = 0 \end{cases}$$

are strictly convex near zero and have identical jets of every finite order, without being equal as functions. Applied for every  $M$ , the lemma shows equality of every finite action jet. It does not assert equality of the action functions or of the smooth boundary images. Analytic boundary determination below uses analyticity only after the finite-jet inverse has been established.

**Lemma 22.5** (Jet filtration and homogeneous isolation). *Let  $m \geq 3$ . Holding the lower graph jets fixed, one has*

$$\partial_{q_r,m} S_b(u) = O(u^m). \quad (22.19)$$

Together with Lemma 22.4, this shows that  $S_b^{(n)}(0)$  depends only on the graph jets through order  $n$ , including independence from smooth remainders and flat variations. For  $m = n$ , the coefficient of  $u^n$  in  $\partial_{q_r,n} S_b(u)$  is obtained by substituting only the linear half-line orbit into the pure endpoint monomials. The error is  $O(u^{n+1})$ , uniformly on compact smooth classes with the required finite derivative bounds. The resulting finite-jet coefficient is uniformly controlled on compact positive finite-jet sets.

*Proof.* For a flight joining transverse coordinates  $y, z$ , let  $h(y, z)$  denote its positive horizontal separation. In the anchored local graph representation its length is

$$\ell(y, z) = \sqrt{h(y, z)^2 + (z - y)^2}.$$

If the varied  $m$ th graph jet occurs at the  $y$  endpoint, then, on the region where the cutoff equals one,

$$\partial_{q_r,m} \ell(y, z) = \frac{h(y, z)}{\ell(y, z)} \frac{y^m}{m!} = \frac{y^m}{m!} + O((|y| + |z|)^{m+2}); \quad (22.20)$$

for a variation at the  $z$  endpoint the same formula holds with  $z^m$ . Indeed  $h/\ell = 1 + O((|y| + |z|)^2)$  at the normal contact. The signs are positive for both facing graphs because increasing either inward graph height increases the horizontal separation.

Lemma 17.1 and Lemma 22.2 give

$$x_i^{(b)}(u) = t_i^{(b)} u + R_i^{(b)}(u), \quad |R_i^{(b)}(u)| \leq C|u|^2 \rho^i, \quad (22.21)$$

with

$$t_i^{(b)} = \begin{cases} e^{-i\gamma}, & i \text{ even}, \\ \mathbf{r}_b e^{-i\gamma}, & i \text{ odd}. \end{cases} \quad (22.22)$$

Hence  $(x_i^{(b)}(u))^m = (t_i^{(b)})^m u^m + O(u^{m+1} \rho^{mi})$ . Insert this and (22.20) into the absolutely convergent envelope sum (22.14). Every summand is  $O(u^m \rho^{mi})$ , proving (22.19). At degree  $m$  only the first term of (22.20) and the linear orbit in (22.22) survive; every nonlinear orbit correction raises the endpoint degree by at least one. Summability is uniform, so the total remainder is  $O(u^{m+1})$ .  $\square$

This lemma is the derivative-order bookkeeping needed by the arbitrary-order inverse. The orbit derivative of order  $k$  may depend on a  $(k+1)$ st graph derivative, but the degree- $n$  action jet cannot: a graph variation of order  $m > n$  starts at endpoint degree  $m$  in the exact envelope formula.

## 22.5. The determinant-one block at every degree.

**Proposition 22.6** (Last-jet block). *For every  $n \geq 3$ , after the lower graph jets are fixed,  $s_n$  is affine in  $q_n$  and its derivative is the matrix  $M_n$  in (22.4). In particular*

$$\det M_n = 1. \quad (22.23)$$

*Proof.* By Lemma 22.5, the degree- $n$  first variation is the pure endpoint-power sum on the linear orbit. A boundary site is counted once and every interior site twice. For a half-line starting at type  $b$ , varying the  $n$ th jet of the starting contact gives

$$\partial_{q_b,n} S_b^{(n)}(0) = 1 + 2 \sum_{k \geq 1} e^{-2n\gamma k} = \frac{1 + e^{-2n\gamma}}{1 - e^{-2n\gamma}} = \coth(n\gamma), \quad (22.24)$$

whereas varying the opposite contact gives

$$\partial_{q_{1-b,n}} S_b^{(n)}(0) = 2\mathbf{r}_b^n \sum_{k \geq 0} e^{-n\gamma(2k+1)} = \mathbf{r}_b^n \operatorname{csch}(n\gamma). \quad (22.25)$$

All multiplicities are explicit and the sums converge absolutely. Interchanging the labels gives the second row of (22.4). Since  $\mathfrak{r}_0 \mathfrak{r}_1 = 1$ ,

$$\det M_n = \coth^2(n\gamma) - (\mathfrak{r}_0 \mathfrak{r}_1)^n \operatorname{csch}^2(n\gamma) = 1.$$

The derivative of  $s_n$  with respect to  $q_n$  is independent of  $q_n$ : the linear orbit depends only on the quadratic geometry and the degree- $n$  homogeneous coefficient above contains no higher occurrence of  $q_n$ . Therefore the dependence is affine.  $\square$

The inverse block is

$$M_n^{-1} = \begin{pmatrix} \coth(n\gamma) & -\mathfrak{r}_0^n \operatorname{csch}(n\gamma) \\ -\mathfrak{r}_1^n \operatorname{csch}(n\gamma) & \coth(n\gamma) \end{pmatrix}. \quad (22.26)$$

If  $R_n$  is the contribution of the already reconstructed lower jets, then

$$q_n = M_n^{-1}(s_n - R_n). \quad (22.27)$$

For fixed maximal order  $M$ , all coefficients and derivatives in this finite recursion are uniformly bounded on compact positive geometric and finite-jet sets.

**Proposition 22.7** (Quantitative tangent isomorphism). *Fix  $g > 0$  and  $M \geq 2$ , and define*

$$\mathcal{A}_M : (\kappa_0, \kappa_1, q_3, \dots, q_M) \mapsto (s_2, s_3, \dots, s_M), \quad s_2 = (a_0, a_1)^t.$$

*Then  $\mathcal{A}_M$  is  $C^1$  and  $D\mathcal{A}_M$  is an isomorphism at every strictly convex labelled contact jet. On every compact positive geometric and finite-jet set there is  $c_M > 0$  such that*

$$\|D\mathcal{A}_M[h]\| \geq c_M \|h\|. \quad (22.28)$$

*Proof.* Lemma 22.4 makes  $\mathcal{A}_M$  well defined independently of smooth representatives. Choose the local polynomial representatives just described. The weighted implicit-function construction and differentiated action bounds give  $C^1$  dependence on their finite jet coordinates; the result is independent of this choice. Formula (22.3) is a smooth inverse of the quadratic map  $(\kappa_0, \kappa_1) \mapsto (a_0, a_1)$  at fixed  $g$ . This gives an invertible leading  $2 \times 2$  diagonal block. For  $n \geq 3$  the new pair  $q_n$  enters  $s_n$  through the diagonal block  $M_n$ , while all dependence on lower jets is strictly below the block diagonal. Hence  $D\mathcal{A}_M$  is block lower triangular with invertible diagonal blocks. The least singular value is a positive continuous function on a compact set, which proves (22.28).  $\square$

*Completion of Theorem 22.1.* The support threshold and quadratic recovery were proved first. Lemma 22.4 proves independence from arbitrary smooth remainders at each finite order. Lemmas 22.2–22.5 and Proposition 22.6 give the all-order recursion (22.27); Proposition 22.7 gives the quantitative tangent inverse.

For analytic contacts, equality of all recovered labelled jets gives equality of the convergent graph germs. Parametrize two candidate connected analytic boundaries by oriented arclength from the recovered contact frame. Their analytic curvature functions agree on an interval and hence everywhere by the identity theorem; uniqueness for the planar Frenet equations gives equal boundary images. This is exact analytic continuation, not a stability claim for distant boundary points.

For the actual finite even-flight law, positivity of the physical flux makes its endpoint support threshold the exact stationary excess action  $E_{j,b}(u, v)$ . The differentiated relative estimate gives

$$\|E_{j,b} - (S_b \oplus S_b)\|_{C^M} \leq C_M \tau^j.$$

Taking the slice  $v = 0$  and applying the finite-order locally Lipschitz inverse proves the  $O_M(\tau^j)$  reconstruction statement.  $\square$

*Remark 22.8* (Oriented intrinsic meaning). Odd graph jets change sign when the chosen transverse orientation is reversed. Equivalently, the recovered data determine the full oriented curvature jet with respect to arclength at each labelled contact. Thus the inverse is intrinsic after the labelled contact point, tangent line, and an orientation of that line have been fixed; no reflection convention is hidden in the theorem.

**22.6. Uniform control of the highest-degree interaction.** The matrices in (22.4) are not arbitrary members of a determinant-one family. Their two shape factors come from positive contact curvatures. This admissibility gives the following order-uniform bounds, observed in [47]. We also record their consequences for actual fixed-lower-jet comparisons and for the block-diagonal operator on Taylor coefficients.

**Proposition 22.9** (Uniform admissible blocks). *Let  $g, \kappa_0, \kappa_1 > 0$ ,  $c_b = 1 + g\kappa_b$ ,  $c = \sqrt{c_0 c_1} = \cosh \gamma$  and  $\mathfrak{r}_b = \sqrt{c_b/c_{1-b}}$ . Set*

$$t = e^{-\gamma}, \quad \vartheta = \frac{1+t^2}{2} < 1, \quad C_\gamma = \frac{3+t^6}{1-t^6}, \quad K_\gamma = \frac{4}{1-t^6}. \quad (22.29)$$

*For the maximum absolute row-sum norm and every integer  $n \geq 3$ ,*

$$\max\{\|M_n\|_\infty, \|M_n^{-1}\|_\infty\} \leq C_\gamma, \quad (22.30)$$

$$\max\{\|M_n - I\|_\infty, \|M_n^{-1} - I\|_\infty\} \leq K_\gamma \vartheta^n. \quad (22.31)$$

*The constants and the exponential rate can be chosen uniformly whenever  $\gamma \geq \gamma_* > 0$ , without a further upper bound on the curvature ratio.*

*If two actual smooth anchored graph pairs have this gap and equal graph jets through order  $n-1$ , then their degree- $n$  graph and action differences satisfy*

$$C_\gamma^{-1} \|\Delta q_n\|_\infty \leq \|\Delta s_n\|_\infty \leq C_\gamma \|\Delta q_n\|_\infty. \quad (22.32)$$

*Here  $C_\gamma$  is independent of  $n$ ; the functional smooth bounds needed to define and differentiate each action are still imposed at the requested finite order.*

*For  $R > 0$ , define the coefficient space*

$$\mathcal{A}_R = \left\{ x = (x_n)_{n \geq 3}, \ x_n \in \mathbb{C}^2 : \|x\|_R = \sum_{n \geq 3} \frac{R^n}{n!} \|x_n\|_\infty < \infty \right\}.$$

*The block-diagonal operator  $(\mathcal{D}x)_n = M_n x_n$  is a bounded isomorphism of  $\mathcal{A}_R$ , with both operator norms at most  $C_\gamma$ . Its corrections to the identity satisfy*

$$\|(\mathcal{D} - I)x\|_{R/\vartheta} + \|(\mathcal{D}^{-1} - I)x\|_{R/\vartheta} \leq 2K_\gamma \|x\|_R. \quad (22.33)$$

*These statements concern the block-diagonal part, not the complete nonlinear graph-to-action map.*

*Proof.* Since  $c_{1-b} > 1$ ,

$$\mathfrak{r}_b = \frac{c}{c_{1-b}} < c, \quad \lambda_b := \mathfrak{r}_b t < ct = \vartheta < 1, \quad \lambda_0 \lambda_1 = t^2. \quad (22.34)$$

Thus the linear visit to the opposite contact has contraction factor  $\lambda_b$ , even when  $\mathfrak{r}_b > 1$ . Summing the endpoint and interior visits as in (22.24)–(22.25) gives, with the sign distinguishing the matrix from its inverse,

$$M_n^{\pm 1} = \frac{1}{1-t^{2n}} \begin{pmatrix} 1+t^{2n} & \pm 2\lambda_0^n \\ \pm 2\lambda_1^n & 1+t^{2n} \end{pmatrix}. \quad (22.35)$$

The inverse formula uses  $\lambda_0\lambda_1 = t^2$  and determinant one. Both matrices have the same absolute row sums. Since  $n \geq 3$ ,

$$\|M_n^{\pm 1}\|_\infty \leq \frac{1 + t^{2n} + 2\vartheta^n}{1 - t^{2n}} \leq \frac{3 + t^6}{1 - t^6}.$$

After subtracting the identity the diagonal numerator is  $2t^{2n}$ . As  $t^2 < \vartheta$ , the row sums are at most  $4\vartheta^n/(1 - t^6)$ . This proves both bounds. If  $\gamma \geq \gamma_*$ , replace  $t$  by  $e^{-\gamma_*}$  in the upper bounds and in  $\vartheta$ ; all resulting constants are uniform.

For actual smooth pairs, Lemma 22.4 first makes each finite action jet a function of the corresponding finite graph jet. Proposition 22.6 then gives  $\Delta s_n = M_n \Delta q_n$ , since the common lower jets fix  $R_n$ . Applying (22.30) in the two directions proves (22.32). No formal-series argument replaces this smooth factorization.

For the coefficient-space assertions, multiply each of the block bounds by  $R^n/n!$  and sum. The inverse is defined coefficient by coefficient and belongs to the same space by the same bound, so it is an actual bounded inverse on  $\mathcal{A}_R$ . Multiplying (22.31) by  $(R/\vartheta)^n/n!$  proves (22.33). In particular, truncating either correction after degree  $N$  makes an operator-norm error on  $\mathcal{A}_R$  of at most  $K_\gamma \vartheta^{N+1}$ . Only absolutely convergent coefficient sums have been used.  $\square$

The location of the remaining conditioning problem is consequently precise. With the leading geometry fixed, the derivative of the full finite triangular map on degrees  $3, \dots, M$ ,  $M \geq 3$ , has the form  $L_M = D_M + N_M$ , where  $D_M = \text{diag}(M_3, \dots, M_M)$  and  $N_M$  is strictly block lower triangular. Hence

$$L_M^{-1} = \left[ \sum_{k=0}^{M-3} (-D_M^{-1} N_M)^k \right] D_M^{-1}. \quad (22.36)$$

The diagonal inverse is bounded independently of  $M$ , but this assertion alone does not bound the displayed sum. It does not control the lower-jet remainders and their derivatives, the finite smooth remainder constants, the density-to-jet operation, or continuation away from the contact. The coefficient-radius gain in (22.33) therefore belongs to the diagonal correction only. The full inverse and finite-flight reconstruction retain their stated fixed-order hypotheses; no unrestricted all-order noisy inverse is being inferred.

**22.7. Analytic inversion of the complete contact action.** The diagonal estimate does not control the powers in (22.36). For analytic contact germs a stronger conclusion follows from a different representation of the *complete* derivative. Stationarity removes the variation of the orbit. The initial contact then contributes an invertible multiplication operator, while every later contact evaluates the boundary variation at a strictly contracted argument. This gives an inverse in an analytic function norm, including all lower-degree couplings. The norm, the local nature of the result, and the restriction on physical collars are specified below; none of the smooth or statistical norms used earlier is silently replaced by it.

Write  $D_R = \{z \in \mathbb{C} : |z| < R\}$  and give  $\mathcal{B}_R = H^\infty(D_R; \mathbb{C}^2)$  the norm  $\|f\|_R = \max_b \sup_{D_R} |f_b|$ . For  $m \geq 3$  let

$$\mathcal{X}_{R,m} = \{f \in \mathcal{B}_R : f_b^{(k)}(0) = 0 \text{ for } 0 \leq k < m, \ b = 0, 1\}, \quad \mathcal{X}_R = \mathcal{X}_{R,3}. \quad (22.37)$$

These are closed complex Banach spaces. Real germs mean the fixed points of  $f(z) \mapsto \overline{f(\bar{z})}$ . Fix  $g > 0$  and a real analytic anchored pair  $\psi^* = (\psi_0^*, \psi_1^*)$ , holomorphic near  $\overline{D_{R_0}}$ , with  $(\psi_b^*)''(0) = \kappa_b > 0$ . The affine space  $\psi^* + \mathcal{X}_R$  fixes the gap and the two quadratic jets. Denote its half-line action map by

$$\mathcal{S}(\psi) = (S_{0,\psi}, S_{1,\psi}).$$

Here the action is the actual stationary sum (17.3), not a formally assigned series.

**Lemma 22.10** (Holomorphic half-lines on a fixed function space). *There are  $0 < R < R_0$ , an open ball  $\mathcal{U}$  about  $\psi^*$  in  $\psi^* + \mathcal{X}_R$ , and  $0 < \rho < 1$  for which both half-line orbits are holomorphic in  $(\psi, u)$ , with*

$$x_{b,0}(u) = u, \quad |x_{b,i}(u)| \leq \rho^i |u| \quad (i \geq 1, u \in D_R, \psi \in \mathcal{U}). \quad (22.38)$$

*The stationary sums define a holomorphic map  $\mathcal{U} \rightarrow \mathcal{S}(\psi^*) + \mathcal{X}_R$ . For real members of  $\mathcal{U}$  these are the physical local actions on a common smaller real collar. Positivity of curvature is required on that smaller collar, not on the whole complex disc or at its boundary.*

*Proof.* We provide the function-space details because pointwise analytic dependence of a finite family would not suffice. Put  $t = e^{-\gamma}$  and  $\lambda_b = \mathbf{r}_b t$ . By (22.34),  $\lambda_* := \max(\lambda_0, \lambda_1) < 1$ . The linear orbit is  $x_{b,i}^{\text{lin}}(u) = \ell_{b,i} u$ , where  $\ell_{b,2k} = t^{2k}$  and  $\ell_{b,2k+1} = \lambda_b t^{2k}$ ; hence  $|\ell_{b,i}| \leq \lambda_*^i$ . Choose

$$\lambda_* < \beta < \rho < \sigma < 1.$$

On sequences of bounded holomorphic functions vanishing at zero use

$$\|y\|_\beta = \sup_{i \geq 1} \beta^{-i} \|y_i\|_{H^\infty(D_R)}.$$

The fixed-reference Green kernel (17.1) is bounded in this norm. Indeed its absolute entries are bounded by a constant times  $t^{|i-k|}$ , and  $\sup_i \sum_{k \geq 1} t^{|i-k|} \beta^{k-i} < \infty$  since  $t < \beta < 1$ .

Consider  $\psi = \psi^* + \eta$  with  $\eta \in \mathcal{X}_R$  and  $\|\eta\|_R < \varepsilon R^3$ . The maximum principle applied to  $\eta_b(z)/z^3$  gives  $|\eta_b(z)| \leq \varepsilon |z|^3$  on  $D_R$ . On  $D_{\sigma R}$ , Cauchy's estimates give uniform third-derivative bounds, and, for  $k \leq 2$ ,

$$|\eta_b^{(k)}(z)| \leq C_{\sigma,k} \varepsilon |z|^{3-k}. \quad (22.39)$$

For example, write  $\eta_b = z^3 v_b$ , use  $\|v_b\|_R \leq \varepsilon$ , and differentiate the product using Cauchy bounds for  $v_b$  on a slightly larger inner disc. The reference pair has the same type of remainder bounds with a fixed constant.

The stationarity equation at an interior site of type  $b_i$  is

$$\frac{h_{i-1} \psi'_{b_i}(x_i) + x_i - x_{i-1}}{\ell_{i-1}} + \frac{h_i \psi'_{b_i}(x_i) + x_i - x_{i+1}}{\ell_i} = 0, \quad (22.40)$$

where  $h_i = g + \psi_{b_i}(x_i) + \psi_{1-b_i}(x_{i+1})$  and  $\ell_i = \sqrt{h_i^2 + (x_{i+1} - x_i)^2}$ , with the square-root branch near  $g$ . Crucially, when  $x_0 = u$  approaches the boundary of  $D_R$ , this equation uses only  $\psi_b(u)$ , never  $\psi'_b(u)$ . Every differentiated graph is evaluated at an interior site, whose argument lies in  $D_{\sigma R}$ . Thus an unbounded differentiation operator on  $H^\infty(D_R)$  is not being used.

Subtract the linear Jacobi equation, whose diagonal is  $2c_{b_i}/g$  and adjacent entries are  $-1/g$ . On the weighted ball  $\|x - x^{\text{lin}}\|_\beta \leq C_0 R^2$ , the nonlinear residual has norm at most  $CR^2$  and Lipschitz constant at most  $CR$  in the interior sequence. These estimates follow by expanding (22.40), using (22.39) at interior arguments and the cubic value remainder at  $u$ . Locality bounds neighboring weights by fixed factors  $\beta^{\pm 1}$ ; products of two such coordinates have norm  $O(R^2)$ . All denominators stay in a fixed neighborhood of  $g$ . Choose  $C_0$  larger than twice the residual constant times the Green norm, then choose  $R$  small. The Green fixed-point map preserves this ball and has contraction constant less than  $1/2$ .

This construction is consistent with the protected inner arguments: for every  $i \geq 1$ , the correction vanishes at zero, so the Schwarz bound gives

$$|x_{b,i}(u)| \leq (\lambda_*^i + C_0 R \beta^i) |u| \leq \rho^i |u|$$

after  $R$  is decreased. The last inequality follows uniformly in  $i$  from  $\lambda_*/\rho + C_0 R \beta/\rho < 1$ . Compositions and graph derivatives at the protected arguments are holomorphic maps of the Banach variables: their Cauchy expansions converge uniformly while these arguments remain

inside  $D_{\sigma R}$ . The fixed-point iterates therefore converge locally uniformly as Banach-valued holomorphic maps. Their limit is holomorphic. Shrink the parameter ball once to keep these bounds and its inverse margin uniform.

Each action summand is bounded by  $CR^2\rho^{2i}$ , so their sum is a locally uniformly convergent holomorphic map into  $\mathcal{B}_R$ . Its constant, linear and quadratic jets are respectively zero, zero and the fixed  $a_b$  of (17.4). Action differences belong to  $\mathcal{X}_R$ . On  $|u| < R/2$ , Cauchy bounds imply that real nearby graphs remain strictly convex and separated. The real stationary sequence just constructed is the unique small half-line of Lemma 17.1, hence gives its actual local action there. No assertion that every perturbed germ extends to an entire periodic obstacle is required for this local statement.  $\square$

**Lemma 22.11** (The complete derivative as weighted boundary visits). *On a smaller ball as in Lemma 22.10, put*

$$w_{b,i}(u) = \frac{g + \psi_{b+i}(x_{b,i}(u)) + \psi_{b+i+1}(x_{b,i+1}(u))}{\ell_{b,i}(u)}, \quad \alpha_b = w_{b,0}, \quad v_{b,i} = w_{b,i-1} + w_{b,i} \quad (i \geq 1),$$

where types are read modulo two. There are constants  $W, a_* > 0$  such that  $|w_{b,i}| \leq W$  and  $|\alpha_b| \geq a_*$  on  $D_R$ . For every  $\eta \in \mathcal{X}_R$ , the full Fréchet derivative is

$$(L_\psi)_b(u) := (D\mathcal{S}(\psi)\eta)_b(u) = \alpha_b(u)\eta_b(u) + \sum_{i \geq 1} v_{b,i}(u)\eta_{b+i}(x_{b,i}(u)). \quad (22.41)$$

Write  $L_\psi = A_\psi + K_\psi$  for the multiplier and the sum. Both preserve every  $\mathcal{X}_{R,m}$ , and

$$\|A_\psi^{-1}\| \leq a_*^{-1}, \quad \|K_\psi|_{\mathcal{X}_{R,m}}\| \leq \frac{2W\rho^m}{1 - \rho^m}. \quad (22.42)$$

The operator  $K_\psi$  is compact on  $\mathcal{X}_R$ ; specifically it has a rank at most  $2(m-3)$  approximation with error at most  $2W\rho^m/((1-\rho)(1-\rho^m))$ .

*Proof.* All length factors are close to  $g$  and  $w_{b,i}$  is close to one on the chosen complex domain, so the stated bounds hold. Differentiate a finite action sum along  $\psi + s\eta$ . The internal orbit derivatives cancel by (22.40); the derivative of  $x_0$  is zero. The only remaining orbit term is

$$\partial_2 \ell_{b+N-1}(x_{b,N-1}, x_{b,N}) D_\psi x_{b,N}[\eta].$$

On a smaller parameter ball, the holomorphic fixed-point construction gives  $\|D_\psi x_{b,N}[\eta]\|_R \leq C_R \beta^N \|\eta\|_R$ . The first factor is at most  $CR\rho^{N-1}$ . The terminal term therefore tends to zero in  $H^\infty(D_R)$ . The direct variation of each flight is exactly  $w_{b,i}\{\eta_{b+i}(x_{b,i}) + \eta_{b+i+1}(x_{b,i+1})\}$ . These variations are norm summable:  $\eta \in \mathcal{X}_R$  and the Schwarz estimate give a majorant  $C\|\eta\|_R \rho^{3i}$ . Passing to the limit in the finite identity proves (22.41). The initial visit occurs once and all internal visits occur twice, as in the finite-jet proof.

For  $\eta \in \mathcal{X}_{R,m}$  the same Schwarz estimate is  $|\eta_b(z)| \leq \|\eta\|_R (|z|/R)^m$ . Summing it at the contracted arguments gives (22.42). Nonvanishing of  $\alpha_b$  makes its reciprocal a bounded holomorphic multiplier. It preserves the order of vanishing.

For compactness, in every interior evaluation replace  $\eta_b$  by its Taylor polynomial through degree  $m-1$ . Each of the  $2(m-3)$  scalar coefficients then multiplies one fixed pair of holomorphic functions, obtained by summing  $v_{b,i}x_{b,i}^n$  over the appropriate parity. Those sums converge in norm. The resulting operator has the stated finite rank. By Cauchy's coefficient bound, the evaluation remainder at  $x_{b,i}$  is at most  $\|\eta\|_R \rho^{im}/(1-\rho^i)$ . Summation proves the asserted error, which tends to zero. Notice that no continuation of the weights or orbits beyond  $D_R$  is used. Compactness here is approximation by finite-rank maps on the same disc; it is not a gain of holomorphic radius for the full derivative.  $\square$



**Theorem 22.12** (Local analytic inversion of the full action map). *For every real analytic positive-curvature anchored pair  $\psi^*$  and gap  $g > 0$  as above, one can choose  $R > 0$  and neighborhoods for which*

$$\mathcal{S} : \psi^* + \mathcal{X}_R \longrightarrow \mathcal{S}(\psi^*) + \mathcal{X}_R$$

*is a biholomorphism between those neighborhoods. In particular, on a sufficiently small convex ball there are constants  $0 < c < C < \infty$  with*

$$c\|\psi - \tilde{\psi}\|_R \leq \|\mathcal{S}(\psi) - \mathcal{S}(\tilde{\psi})\|_R \leq C\|\psi - \tilde{\psi}\|_R. \quad (22.43)$$

*These constants control the complete map, not only pairs with fixed lower jets. For every  $N \geq 3$ , its derivative and inverse induce bounded mutually inverse operators on  $\mathcal{X}_R/\mathcal{X}_{R,N+1}$ , with bounds independent of  $N$  in the quotient norms induced by  $\|\cdot\|_R$ .*

*For  $0 < r < R$ , all contact coefficients of two such pairs satisfy*

$$\sum_{n \geq 3} \frac{r^n}{n!} \max_b |\psi_b^{(n)}(0) - \tilde{\psi}_b^{(n)}(0)| \leq \frac{(r/R)^3}{c(1-r/R)} \|\mathcal{S}(\psi) - \mathcal{S}(\tilde{\psi})\|_R. \quad (22.44)$$

*The same local real analytic inverse allows the gap and curvatures to vary when the gap and both action Hessians are retained as finite coordinates. It concerns local contact germs and a strong norm on their holomorphic extensions, not a weak-density or full-table statistical norm.*

*Proof.* Fix the base point and abbreviate  $L = L_{\psi^*}$ ,  $A = A_{\psi^*}$  and  $K = K_{\psi^*}$ . We construct a bounded inverse for  $L$ , including its lower-order couplings. Choose one finite  $m \geq 4$  so large that

$$q_m := \frac{2W}{a_*} \frac{\rho^m}{1 - \rho^m} < 1. \quad (22.45)$$

On  $\mathcal{X}_{R,m}$ , Lemma 22.11 gives

$$L_{\text{tail}}^{-1} = \sum_{j \geq 0} (-A^{-1}K)^j A^{-1}, \quad \|L_{\text{tail}}^{-1}\| \leq \frac{a_*^{-1}}{1 - q_m}. \quad (22.46)$$

This norm-convergent series acts on actual functions. It is not the finite nilpotent formula with an unbounded number of unspecified factors.

Let  $P_m$  be Taylor projection onto the pairs of polynomials of degrees  $3, \dots, m-1$ . It is bounded on  $\mathcal{X}_R$  by Cauchy's coefficient estimate. The induced finite map

$$L_{\text{low}} = P_m L|_{P_m \mathcal{X}_R}$$

is block lower triangular. Its diagonal blocks are exactly  $M_3, \dots, M_{m-1}$ . One way to see this is to differentiate the actual finite-jet identity proved in Lemma 22.4 and Proposition 22.6. It also follows directly from (22.41):  $w_{b,i}(0) = 1$  and  $x'_{b,i}(0) = \ell_{b,i}$ , so the new degree is the same weighted geometric sum of visits. Analytic continuation of this finite identity, or complex linearity at the real base, gives the complex statement. All these finitely many blocks are invertible. Therefore  $L_{\text{low}}$  has a bounded inverse on its finite-dimensional space, without requiring its strictly lower part to be small.

For an arbitrary  $y \in \mathcal{X}_R$ , define

$$\eta_{\text{low}} = L_{\text{low}}^{-1} P_m y, \quad \eta_{\text{tail}} = L_{\text{tail}}^{-1} (y - L \eta_{\text{low}}). \quad (22.47)$$

The second argument lies in  $\mathcal{X}_{R,m}$  because its projection under  $P_m$  vanishes. These formulas solve  $L\eta = y$  uniquely and boundedly. For example, if  $B = \|L_{\text{low}}^{-1} P_m\|$ , then

$$\|L^{-1}\| \leq B + \frac{a_*^{-1}}{1 - q_m} (1 + \|L\|B). \quad (22.48)$$

Only one fixed finite low-order problem has been inverted. All remaining orders are handled together by (22.46).

For completeness the nonlinear step can be made by contraction rather than a formal inverse expansion. Set  $\mathcal{N}(\eta) = \mathcal{S}(\psi^* + \eta) - \mathcal{S}(\psi^*) - L\eta$ . It is holomorphic and  $D\mathcal{N}(0) = 0$ . On a small ball arrange  $\|L^{-1}D\mathcal{N}\| < 1/2$ . For  $y$  small the equation

$$\eta = L^{-1}y - L^{-1}\mathcal{N}(\eta)$$

preserves this ball and is a contraction. Its iterates converge locally uniformly and holomorphically in  $y$ . This gives the local holomorphic inverse. Shrinking a convex ball also gives  $\|D\mathcal{S} - L\| < (2\|L^{-1}\|)^{-1}$  there. Integration on the segment joining two points proves (22.43), with  $c = (2\|L^{-1}\|)^{-1}$  and a local bound for the upper constant. Real data give real inverses by uniqueness and conjugation symmetry.

Every tail  $\mathcal{X}_{R,N+1}$  is invariant under  $L$ . It is also invariant under  $L^{-1}$ : if the output has zero jets through  $N$ , finite triangular induction with the invertible  $M_n$  makes the input jets through  $N$  zero. Thus  $L$  and  $L^{-1}$  descend to mutually inverse quotient operators. Their norms are bounded by the parent operator norms, uniformly in  $N$ . The same construction, or continuity of the derivative and its inverse on a smaller parameter neighborhood, gives local uniformity in the base pair. This statement is about analytic quotient norms, not maximum norms of unweighted derivatives. Cauchy's estimate at zero and (22.43) then give (22.44) by summing the geometric series.

Finally, let the leading variables vary near their positive base values. The protected-domain construction is holomorphic in  $g, \kappa_0, \kappa_1$  as well. Write  $\psi_b = \kappa_b u^2/2 + v_b$ , with  $v \in \mathcal{X}_R$ , and similarly subtract  $a_b u^2/2$  from each action. The leading map  $(g, \kappa_0, \kappa_1) \mapsto (g, a_0, a_1)$  has the local analytic inverse already given by

$$\gamma = \operatorname{arsinh}(g\sqrt{a_0 a_1}), \quad c_b = \cosh \gamma \sqrt{a_b/a_{1-b}}, \quad \kappa_b = (c_b - 1)/g,$$

with the positive branches near the base. In the finite-variable and  $\mathcal{X}_R$  splitting the derivative is block lower triangular: the finite leading block and the infinite tail block are invertible. The same local contraction proves the stated joint inverse.  $\square$

The theorem supplies a local analytic inverse for the entire action, not just its diagonal response. Its additional input is concrete: the norm controls the holomorphic extensions on one fixed disc. Smallness of a real-interval  $C^k$  error or total variation does not supply that norm. The radius and the inverse constants may depend on the analytic base pair, its gap, curvature margins and extension bounds. The result does not extend a local germ to an entire obstacle, select a discrete global matching branch, or give a uniform acquisition cost. Those remain the separate geometric and statistical statements proved elsewhere. Neither  $K$  nor the full inverse is claimed to improve holomorphic radius; only the finite-rank approximation and the tail bounds stated above are used. In particular the arbitrary triangular-matrix comparison does not refute this result: an actual action derivative has the additional multiplier-and-contracted-visits structure (22.41).

**22.8. Contracted evaluations and conditional stability from real observations.** The complete analytic inverse has two distinct ingredients: a stationary representation of its derivative and a function-space inversion argument. We first isolate the latter. We then combine the inverse with a quantitative restriction estimate to obtain a local conditional bound from real endpoint laws. The conclusion loses holomorphic radius and retains an analytic prior. It therefore does not identify real observation accuracy with the same-disc norm of Theorem 22.12.

**Proposition 22.13** (An inversion criterion for contracted evaluations). *Let  $q \geq 1$ ,  $d \geq 1$ , and let  $X_q$  be the closed subspace of  $H^\infty(D_R; \mathbb{C}^d)$  consisting of functions vanishing to order at least  $q$ ,*

with the maximum component norm. Suppose

$$(Lf)_b(z) = (A(z)f(z))_b + \sum_{i \geq 1} \sum_{c=1}^d w_{bci}(z) f_c(\phi_{bci}(z)), \quad (22.49)$$

where all coefficients and maps are holomorphic on  $D_R$ ,  $\sup_{D_R} \|A^{-1}\| \leq a^{-1}$ ,  $A$  is bounded, and

$$|\phi_{bci}(z)| \leq \rho^i |z|, \quad \sup_{b,z} \sum_c |w_{bci}(z)| \leq W, \quad 0 < \rho < 1.$$

Write  $L = A + K$ . Choose an integer  $m > q$  with  $\nu_m = a^{-1}W\rho^m/(1-\rho^m) < 1$ . Then  $L : X_q \rightarrow X_q$  is boundedly invertible if and only if the finitely many matrices

$$(B_n)_{bc} = A_{bc}(0) + \sum_{i \geq 1} w_{bci}(0) (\phi'_{bci}(0))^n, \quad q \leq n < m, \quad (22.50)$$

are invertible. The operator  $K$  is compact. This criterion applies to the derivative at a point of a holomorphic map between open subsets of  $X_q$  and then yields a local holomorphic inverse of that map.

*Proof.* Higher-order Schwarz estimates give

$$\|K|_{X_n}\| \leq \frac{W\rho^n}{1-\rho^n}, \quad n \geq q.$$

Here  $X_n$  denotes the corresponding order- $n$  vanishing subspace. The series in (22.49) converges in operator norm. Both summands, and multiplication by  $A^{-1}$ , preserve every  $X_n$ . Thus  $L|_{X_m}$  is invertible by the Neumann series, with inverse norm at most  $a^{-1}/(1-\nu_m)$ .

Let  $P_m$  retain Taylor degrees  $q, \dots, m-1$ . Cauchy's coefficient bound makes  $P_m$  a bounded projection on  $X_q$ . Modulo  $X_m$ , the matrix of  $L$  is block lower triangular with diagonal (22.50): only the constant weight and the linear part of an argument contribute to the first nonzero degree of a variation. If those blocks are invertible, solve this finite quotient first and then solve the residual in  $X_m$ , exactly as in (22.47). In particular, if  $C_0 = \|(P_m L|_{P_m X_q})^{-1} P_m\|$ , then

$$\|L^{-1}\| \leq C_0 + \frac{a^{-1}}{1-\nu_m} (1 + \|L\|C_0).$$

Conversely, if the finite quotient were singular, a nonzero polynomial class  $p$  would satisfy  $Lp \in X_m$ . Solving  $Lt = Lp$  inside  $X_m$  would give a nonzero kernel vector  $p - t$ . Consequently invertibility of  $L$  implies the stated finite condition.

To prove compactness, replace each evaluated  $f_c$  by its Taylor polynomial of degrees  $q, \dots, n-1$ . The resulting operator has rank at most  $d(n-q)$ . Cauchy's coefficient bound gives the operator-norm remainder estimate

$$\frac{W\rho^n}{(1-\rho)(1-\rho^n)} \rightarrow 0.$$

The truncation is evaluated only at contracted arguments; convergence of Taylor partial sums on the whole open disc is not asserted. The final assertion follows from the holomorphic inverse function theorem, or from the contraction proof already given above.  $\square$

For the contact action,  $q = 3$ ,  $d = 2$ , the initial multiplier is diagonal, and the only nonzero weight at visit  $i$  is  $v_{b,i}$  in the component of type  $b + i$ . Its bound is  $2W$  in the notation of Lemma 22.11. The matrices  $B_n$  are the actual  $M_n$ . The protected-domain construction and the vanishing terminal term prove that this is indeed the derivative of the physical action; they are not assumptions supplied by functional analysis. The qualitative Fredholm observation suggested in [48] is now also transparent:  $A^{-1}L$  is identity plus compact, and nonresonant jets kill its

analytic kernel. The finite-low-order and tail construction above additionally records an inverse bound. Neither argument makes the inverse improve holomorphic radius.

**Lemma 22.14** (A bounded analytic function from its real restriction). *Fix  $0 < s < r < \varrho < R$  with  $r + 2s < \varrho$ , and put*

$$A_0 = \frac{2(r+s)}{s}, \quad b_0 = \frac{r+s}{\varrho-s} < 1, \quad \vartheta = \frac{\log(1/b_0)}{\log(A_0/b_0)} \in (0, 1). \quad (22.51)$$

*If  $F \in H^\infty(D_R)$ ,  $\|F\|_R \leq M$ , and  $\sup_{[-s,s]} |F| \leq \epsilon \leq M$ , then*

$$\|F\|_r \leq \left(1 + \frac{\varrho}{\varrho-r}\right) M^{1-\vartheta} \epsilon^\vartheta. \quad (22.52)$$

*The assertion includes complex-valued functions and the case  $\epsilon = 0$ .*

*Proof.* For  $n \geq 1$  take the  $n$  roots  $x_k$  of  $T_n(x/s)$ , where  $T_n$  is the Chebyshev polynomial. Their monic nodal polynomial is  $P_n(z) = 2^{1-n} s^n T_n(z/s)$ , and

$$|P'_n(x_k)| \geq n(s/2)^{n-1}.$$

The Lagrange interpolation polynomial  $I_n F$  therefore satisfies  $\|I_n F\|_r \leq \epsilon A_0^{n-1}$ : the numerator of each basis polynomial is a product of  $n-1$  factors of modulus at most  $r+s$ . The residue formula on  $|\zeta| = \varrho$  gives

$$F(z) - I_n F(z) = \frac{1}{2\pi i} \int_{|\zeta|=\varrho} \frac{F(\zeta) P_n(z)}{P_n(\zeta)(\zeta-z)} d\zeta.$$

This holds also at the interpolation nodes by continuity. The bounds  $|P_n(z)| \leq (r+s)^n$  and  $|P_n(\zeta)| \geq (\varrho-s)^n$  imply

$$\|F\|_r \leq \epsilon A_0^{n-1} + \frac{\varrho}{\varrho-r} M b_0^n. \quad (22.53)$$

For  $0 < \epsilon < M$ , choose  $n = \lceil \log(M/\epsilon) / \log(A_0/b_0) \rceil$ . The two terms are bounded by their respective constants times  $M^{1-\vartheta} \epsilon^\vartheta$ . For  $\epsilon = M$ , take  $n = 1$ ; for  $\epsilon = 0$ , let  $n$  tend to infinity. The case  $M = 0$  is immediate. All contours are strictly inside the given holomorphic domain.  $\square$

This is a conditional continuation estimate, not continuity of unrestricted analytic continuation. Its elementary interpolation proof is included to display the constants; the general ill-conditioning of continuation under analytic priors is discussed in [37]. We use only the local estimate just proved, not continuation around an entire obstacle.

**Theorem 22.15** (Conditional contact stability from real endpoint laws). *Fix a positive-curvature real analytic anchored pair and a positive gap. There exist radii  $0 < s < s_+ < r < \varrho < R$ , with  $r + 2s < \varrho$ , and a bounded neighborhood  $\mathcal{U}$  of that pair and gap in the real part of*

$$\mathbb{C} \times \{\psi \in H^\infty(D_R; \mathbb{C}^2) : \psi_b(0) = \psi'_b(0) = 0\},$$

*for which the following assertions hold. The curvatures and gap may vary within  $\mathcal{U}$ . Choose a fixed sufficiently small positive offset  $d$  and the radii so that  $Q_+ = [-s_+, s_+]^2$  is strictly inside every relevant positive-density region. Put  $Q = [-s, s]^2$  and fix  $a \in (-s, s) \setminus \{0\}$ .*

*For each of two admissible members  $(g, \psi)$  and  $(\tilde{g}, \tilde{\psi})$  retain a same-type limiting endpoint density for each type, possibly after separate positive recording factors. On  $Q_+$  these densities have the form*

$$f_b(u, v) = c_b U_b(u) V_b(v) \{d - S_b(u) - S_b(v)\}, \quad c_b, U_b, V_b > 0, \quad (22.54)$$

*and the analogous form for  $\tilde{f}_b$ . Assume common bounds  $m_f \leq f_b, \tilde{f}_b \leq M_f$  and  $q_{f_b}, q_{\tilde{f}_b} \geq q_* > 0$ , where  $q_f$  is the scalar anchor (23.3). The two sets of recording factors need not coincide, and need*

not have holomorphic extensions. The origins, signs, physical units and  $d$  are fixed observation conventions.

Let  $e_g = |g - \tilde{g}|$  and  $e_\infty = \max_b \|f_b - \tilde{f}_b\|_{C^0(Q)}$ . With  $\vartheta$  as in (22.51), there is a finite constant  $C$  such that, for sufficiently small errors,

$$\|\psi - \tilde{\psi}\|_r \leq C(e_g + e_\infty^\vartheta). \quad (22.55)$$

For every  $0 < t < r$  this simultaneously controls all contact coefficients:

$$\sum_{n \geq 2} \frac{t^n}{n!} \max_b |\psi_b^{(n)}(0) - \tilde{\psi}_b^{(n)}(0)| \leq C \frac{(t/r)^2}{1 - t/r} (e_g + e_\infty^\vartheta). \quad (22.56)$$

If, in addition, all four densities have a common Lipschitz bound  $H$  on  $Q_+$ , then, with  $e_1 = \max_b \|f_b - \tilde{f}_b\|_{L^1(Q_+)}$ , one has

$$\|\psi - \tilde{\psi}\|_r \leq C_H(e_g + e_1^{\vartheta/3}). \quad (22.57)$$

In particular  $e_1$  may be replaced by twice the maximum total variation of the recorded laws. For a rectangular grid on  $Q_+$  of side at most  $\delta$ , with the outside as an additional category, it may instead be replaced by

$$\eta = e_{\text{cell}} + 2\sqrt{2}H|Q_+|\delta, \quad e_{\text{cell}} = \max_b \|\mathbf{p}_\delta(P_b) - \mathbf{p}_\delta(\tilde{P}_b)\|_1. \quad (22.58)$$

The constants depend on the fixed local analytic neighborhood, the radii and the stated density/anchor margins, but not on a jet truncation order.

*Proof.* We specify the neighborhood before using small observational error. Lemma 22.10, including its finite leading-variable extension, gives an action map on a disc  $D_R$  near the chosen base. Apply Theorem 22.12 to the same base on a smaller disc to obtain an inverse ball on  $D_r$ , with  $r < R/4$ . Write  $Y_r$  for the space of pairs vanishing to order two. The bounded splitting into the Hessians and the order-three tail turns the last clause of that theorem into a local inverse of

$$(g, \psi) \mapsto (g, \mathcal{S}_g(\psi)) \quad \text{in } \mathbb{C} \times Y_r.$$

Indeed  $|S_b''(0)| \leq 2r^{-2}\|S_b\|_r$ , and subtracting the quadratic polynomial is a bounded operation. Thus no additional measurement of the curvatures or action Hessians is required here. On a smaller convex input ball, its two-point bound has the form

$$\|\psi - \tilde{\psi}\|_r \leq C_r\{e_g + \|\mathcal{S}_g(\psi) - \mathcal{S}_{\tilde{g}}(\tilde{\psi})\|_r\}. \quad (22.59)$$

Intersect the  $D_R$  construction neighborhood with the inverse image of this  $D_r$  ball under restriction. This is an open neighborhood of the base, independent of the observed errors. Shrink it to a bounded ball on which the action norms on  $D_R$  have a common bound  $B$ . Continuity of the holomorphic forward map provides such a  $B$ ; no inverse norm is inferred from the observations. The two constructions of the action commute with restriction: on the smaller real collar both are the unique small physical half-line action, and equality extends holomorphically to  $D_r$ . Choose  $\varrho = R/2$  and then  $s_+ < r/2$  small enough for the common physical collar and fixed offset, followed by  $s < s_+$ . These choices imply  $r + 2s < \varrho$ . Positive contact curvature permits a fixed nonzero anchor with a uniform action-anchor margin. The recorded density bounds remain separate assumptions on the recording experiment.

For (22.54), the four-density identity cancels  $c_b, U_b, V_b$  exactly, just as in (23.6). On the stipulated density range its ratios are Lipschitz in the uniform real norm. Moreover

$$|q_f - q_{\tilde{f}}| = \frac{|R_f(a, a) - R_{\tilde{f}}(a, a)|}{q_f + q_{\tilde{f}}},$$

so the anchor division is uniformly controlled. The final map  $t \mapsto dt/(1+t)$  is Lipschitz for the nonnegative real  $t$  here. Consequently

$$\max_b \sup_{[-s,s]} |S_b - \tilde{S}_b| \leq C_1 e_\infty. \quad (22.60)$$

This argument compares two realizable laws directly; it does not assume that the straight line joining their densities factorizes. Apply Lemma 22.14 componentwise to  $S - \tilde{S}$ , whose  $D_R$  norm is at most  $2B$ . Equations (22.60) and (22.59) give (22.55). Cauchy's coefficient bound on  $D_r$  and summation of a geometric series give (22.56), including odd orders.

For the weaker norm, put  $w = f_b - \tilde{f}_b$  and choose a fixed nonnegative smooth kernel  $k$  supported in the unit ball of  $\mathbb{R}^2$ , with integral one. For  $x \in Q$  and  $0 < h \leq h_0 < \text{dist}(Q, \partial Q_+)$ ,

$$|w(x) - (k_h * w)(x)| \leq CHh, \quad |(k_h * w)(x)| \leq \|k\|_\infty h^{-2} e_1.$$

Take  $h = h_0 e_1^{1/3}$  when  $0 < e_1 \leq 1$ ; for  $e_1 = 0$  let  $h$  decrease to zero. Thus  $e_\infty \leq C_H e_1^{1/3}$ , proving (22.57). Only the stated first-order real bound is used at this step. The restriction of an  $L^1$  difference to  $Q_+$  is at most twice the total variation of the full laws. Finally, the  $L^1$  difference between one density and its cell averages is at most  $\sqrt{2}H|Q_+|\delta$ . The difference of the two cell averages integrates to the sum of the absolute cell-mass differences. Adding the two approximation errors proves  $e_1 \leq \eta$  in (22.58).  $\square$

The neighborhood and outer holomorphic bound in this theorem are a geometric regularity prior, not a consequence of small real data error. The radius loss is essential to the argument. For example,  $F_N(z) = \epsilon(z/R)^N$  has norm  $\epsilon$  on  $D_R$  while all its fixed-order real norms on  $[-s, s]$ ,  $s < R$ , tend to zero; its norm on  $D_r$ ,  $r < R$ , also tends to zero. Thus the estimate does not conflict with the same-disc topology distinction. The exponent is explicit but not asserted optimal. The theorem compares locally admissible candidates and controls their complete germs on an inner disc. It neither constructs a map on arbitrary corrupted densities nor identifies a global matching branch or an unknown contact frame. The histogram bound is deterministic; no uniform preparation cost over arbitrarily small recording efficiencies follows from it. Quantified continuation to entire obstacles and the physical acquisition theorem retain their separate hypotheses and proofs.

### 23. AN AMPLITUDE-FREE INVERSE FROM ONE ENDPOINT LAW

An endpoint position in the plane supplies a point of the unknown boundary. A signed transverse endpoint supplies only one coordinate of that point. The result below concerns the latter datum. It recovers the two half-line actions from one positive-offset endpoint law of each type, without retaining longitudinal positions, residual times, waiting counts, or the unknown flux amplitudes. The gap remains a separate onset datum. This strengthens the support-family inverse of Theorem 22.1: a continuum of offsets is unnecessary when the density, rather than only its support, is retained.

**23.1. A four-density identity.** Fix a known  $d > 0$ . Let  $I = [-r, r]$ , and suppose on a neighborhood of  $I$  that  $S(0) = S'(0) = 0$ ,  $S'' > 0$ , and  $B > 0$ . Shrink  $r$  until  $\max_I S < d/4$ . The endpoint law has on  $I^2$  the positive smooth density

$$f(u, v) = Z^{-1} B(u) B(v) \{d - S(u) - S(v)\}, \quad Z > 0. \quad (23.1)$$

Its continuation outside this interior square may vanish at the moving boundary. A law in this class determines the displayed continuous density uniquely: two continuous versions equal almost everywhere on an open set are equal everywhere there. Thus values on the slices below

are determined by the probability law, although a finite sample does not supply them without estimation.

Choose a fixed  $a \in I \setminus \{0\}$ . For a positive density  $f$  define

$$R_f(u, v) = \frac{f(u, v)f(0, 0)}{f(u, 0)f(0, v)}, \quad (23.2)$$

$$q_f = \sqrt{1 - R_f(a, a)}, \quad T_f(u) = \frac{1 - R_f(u, a)}{q_f}, \quad (23.3)$$

$$\mathcal{S}_d(f)(u) = \frac{dT_f(u)}{1 + T_f(u)}, \quad \mathcal{B}(f)(u) = \frac{f(u, 0)}{f(0, 0)}(1 + T_f(u)). \quad (23.4)$$

The square root in (23.3) is the positive scalar root at the fixed nonzero anchor, not a pointwise square root differentiated at the degenerate minimum  $u = 0$ .

**Theorem 23.1** (Single-offset law inversion). *For (23.1), all quantities in (23.3)–(23.4) are defined and*

$$\mathcal{S}_d(f) = S, \quad \mathcal{B}(f) = B/B(0) \quad \text{on } I. \quad (23.5)$$

*Consequently one same-type signed endpoint law at one known positive offset determines the complete action germ and the normalized amplitude germ. Neither an amplitude model nor its normalizing constant is supplied.*

*Proof.* Set  $t(u) = S(u)/(d - S(u))$ . Direct cancellation in (23.2) gives

$$R_f(u, v) = \frac{d\{d - S(u) - S(v)\}}{(d - S(u))(d - S(v))} = 1 - t(u)t(v). \quad (23.6)$$

Strict convexity and the normalization of  $S$  imply  $S(a) > 0$ , hence  $q_f = t(a) > 0$ . Therefore  $T_f(u) = t(u)$ , and solving this relation for  $S(u)$  gives the first identity in (23.5). Also

$$\frac{f(u, 0)}{f(0, 0)} = \frac{B(u)}{B(0)} \frac{d - S(u)}{d} = \frac{B(u)/B(0)}{1 + t(u)},$$

which proves the second. No symmetry of  $S$  or  $B$  was used. In particular values at negative and positive  $u$  remain separate and all odd derivatives of  $S$  are retained.  $\square$

### 23.2. Fixed-order stability and finite flights.

**Proposition 23.2** (Interior stability of the law inverse). *Fix an integer  $M \geq 0$ , the interval  $I$ , the anchor  $a$ , and  $d > 0$ . On a compact family of densities (23.1) with uniformly positive  $f$ , uniformly bounded  $C^M$  norms, and  $q_f \geq q_* > 0$ , there are  $\varepsilon_*, C_M > 0$  such that*

$$\|\tilde{f} - f\|_{C^M(I^2)} < \varepsilon_* \implies \|\mathcal{S}_d(\tilde{f}) - S\|_{C^M(I)} + \|\mathcal{B}(\tilde{f}) - B/B(0)\|_{C^M(I)} \leq C_M \|\tilde{f} - f\|_{C^M(I^2)}.$$

*The perturbed function need not factorize or have total mass one. It must be defined on a neighborhood of the closed square. The constants are fixed-order constants; they are not uniform as  $M \rightarrow \infty$ ,  $a \rightarrow 0$ , or  $d \downarrow 0$ .*

*Proof.* Evaluation on each fixed slice is a bounded linear map on  $C^M$ . All density denominators in (23.2) have a common positive lower bound. Products and reciprocal functions are locally Lipschitz in  $C^M$  on sets with such bounds, as follows by differentiating  $hh^{-1} = 1$  inductively. Thus the ratios are locally Lipschitz in  $C^M$ . The scalar  $1 - R_f(a, a)$  is at least  $q_*^2$ , so its positive square root has a bounded derivative on a common neighborhood. The denominator  $q_{\tilde{f}}$  is then at least  $q_*/2$ . For the exact density,  $1 + T_f \geq 1$ ; after a further uniform reduction of  $\varepsilon_*$ ,  $1 + T_{\tilde{f}} \geq 1/2$ . Applying the same product and reciprocal bounds to (23.4) proves the claim. This proof avoids differentiation of  $\sqrt{t(u)^2}$  at zero.  $\square$

The hypotheses can be arranged uniformly on any compact smooth contact class with positive curvature and gap bounds. Choose one  $d$  inside its common positive-offset collar, then a common  $I$  on which  $S_b(u) < d/4$ . A fixed nonzero  $a \in I$  has  $S_b(a) \geq ca^2$ , hence a common positive anchor bound. The amplitude and normalization bounds give a positive lower bound for  $f$  on  $I^2$ .

**Corollary 23.3** (Finite-flight law reconstruction). *Fix  $M \geq 2$  on such a compact class. Let  $f_{j,b,d}$  be the successful signed endpoint density for an even  $j$ -flight bridge, in the exact contact transverse chart, at the fixed offset  $d$ . On the selected interior square, for all sufficiently large even  $j$ ,*

$$\|\mathcal{S}_d(f_{j,b,d}) - S_b\|_{C^M(I)} \leq C_M \tau^j. \quad (23.7)$$

*More generally an estimated density  $\hat{f}_{j,b,d}$  with interior  $C^M$  error at most  $\epsilon$ , together with a gap estimate of error at most  $\epsilon_g$ , recovers both labelled graph jets through order  $M$  with error at most*

$$C_M(\tau^j + \epsilon + \epsilon_g), \quad (23.8)$$

*provided the errors are sufficiently small. This is a stability statement in a density norm, not an assertion that total variation alone controls arbitrarily high derivatives.*

*Proof.* In the transverse chart the finite endpoint density is a positive normalizing constant times  $b_j(u, v)(d - E_j(u, v))_+$ . The differentiated relative theorem gives  $b_j = B_b \otimes B_b + O_{C^M}(\tau^j)$  and  $E_j = S_b \oplus S_b + O_{C^M}(\tau^j)$ . On the chosen square the cutoff argument is bounded away from zero. Over the whole common endpoint box the inequality  $|x_+ - y_+| \leq |x - y|$ , the uniform amplitude bounds, and its finite volume give an  $O(\tau^j)$  error for the normalizing integral. That integral has a positive uniform lower bound at fixed  $d$ . Division therefore proves  $\|f_{j,b,d} - f_{b,d}\|_{C^M(I^2)} \leq C_M \tau^j$ . Apply Proposition 23.2.

Before applying the contact-jet recursion, subtract the zeroth and first Taylor terms of the estimated action. This bounded projection fixes the true anchored action, preserves the error bound, and does not change any coefficient of order at least two. Thus a small finite-flight or density error is not interpreted as an unknown change of contact origin or tangent. For completeness, if  $a_b = S_b''(0)$ , the leading inversion is

$$\gamma = \operatorname{arsinh}(g\sqrt{a_0 a_1}), \quad c_b = \cosh \gamma \sqrt{a_b/a_{1-b}}, \quad \kappa_b = (c_b - 1)/g.$$

At each  $3 \leq n \leq M$ , use  $q_n = M_n^{-1}(s_n - R_n)$ , where  $R_n$  is already determined by lower-order graph jets and  $M_n$  is the determinant-one block of (22.4). Each of these finitely many operations is smooth on a neighborhood of the compact positive geometric set. Induction in  $n$  gives (23.8). No infinite-order condition bound is inferred from the determinant identity.  $\square$

### 23.3. The intrinsic data map.

**Theorem 23.4** (Periodic determination from one endpoint law of each type). *For every measured channel  $e$ , retain its gap  $g_e$ , its signed transverse coordinate conventions, and the two conditional endpoint laws  $\Pi_{e,00,d_e}^{\text{end},\infty}$  and  $\Pi_{e,11,d_e}^{\text{end},\infty}$  at one known sufficiently small positive offset  $d_e$ . No longitudinal coordinate, residual time, waiting count, or amplitude is part of these conditional-law data. Then:*

- (1) *For smooth strictly convex contacts the data determine every finite labelled contact jet, without reflection symmetry or supplied curvatures.*
- (2) *For connected real-analytic incident boundaries they determine the complete boundary images in their separate channel frames.*
- (3) *Suppose the data are realized by a labelled periodic dispersing billiard with connected real-analytic strictly convex obstacle boundaries and the marked lifted-channel graph is connected. Assume that this graph contains a rank-two anchoring pair of signature-rigid*



*cycles based at one channel frame and a rooted signature-rigid spanning tree reaching every obstacle orbit, as in Theorem 26.4. Then the data determine the complete marked periodic table and the Euclidean realization of its rank-two lattice, including its Gram form and scale, up to one simultaneous element of  $\text{SE}(2)$ .*

*On a compact class all  $d_e$  can be the same fixed  $d_0 > 0$ .*

*Proof.* Theorem 23.1 recovers each action on a neighborhood of zero. The weighted half-line construction, finite-truncation envelope identity and homogeneous isolation in Theorem 22.1 then recover the labelled graph jets. The density cancellation is not a replacement for that inverse: it supplies its unsymmetrized action input. Analytic continuation gives the complete incident boundary images. The intrinsic gluing theorem matches their framed signatures.

For part (3), fix the common starting channel frame of the anchoring pair, and choose an obstacle orbit  $r$  incident to that frame. A signature-rigid incidence on an anchoring cycle gives a framed point of this obstacle whose complete oriented curvature signature occurs uniquely. Lemma 25.11 therefore permits us to root the supplied spanning tree at  $r$ . This reorients the existing tree; it introduces no channel, incidence, deck displacement, or registration datum. In particular, the analytic image at the new root is already placed in the chosen anchoring frame.

The two signature-rigid cycles determine translation holonomies  $V = (v_1 \ v_2)$  in this same frame. If  $M$  has their two independent marked deck displacements as columns, then Theorem 26.3 gives

$$L = VM^{-1}, \quad L^T L = M^{-T} V^T V M^{-1}.$$

This determines the lattice, including its scale, but does not by itself place obstacle orbits outside the anchoring cycles. Those placements follow by induction along the rerooted tree. At each edge, the already placed parent boundary has a unique framed point with the prescribed oriented signature. Matching that frame fixes the adjacent channel; its recovered analytic image places the child obstacle, with the marked deck displacement interpreted through the same matrix  $L$ . The tree reaches every obstacle orbit, as required by Theorem 26.4.

To verify uniqueness in one gauge, consider two admissible realizations of the data and align their starting channel frames by one element of  $\text{SE}(2)$ . The local inverse and analytic continuation give the same incident boundary images. Unique signature matching along the two cycles gives the same  $V$ , hence the same  $L$ . At each successive tree edge, unique matching at the already common parent frame gives the same child placement in both realizations. Induction identifies every obstacle orbit; the common lattice then identifies all its translates. The realizability hypothesis supplies existence and admissibility of this reconstruction, including compatibility with the remaining measured channels. Changing the starting frame applies one simultaneous Euclidean motion to the entire reconstruction, not independent motions to its obstacles. None of these steps uses ambient position samples or an externally supplied lattice metric. Uniform collar bounds give the last assertion.  $\square$

This theorem concerns an intrinsic, law-valued observation map with specified signed transverse origins and axes, as does the preceding support-family inverse. It does not treat an unknown centering transformation as a common statistical kernel. Section 56 obtains such origins and axes from a charged physical pilot, and Section 57 implements the law inverse by bounded tests after that pilot. Exact longitudinal positions are available in that richer acquisition protocol and permit the independent direct sampling route of Section 54. Neither route is claimed to be statistically equivalent to the other.

**23.4. Finite preparations and complete analytic contact recovery.** We now pass from limiting laws to finite experiments while keeping the complete analytic contact pair as the target. The order of this passage matters. A finite-bridge density does not in general have the separated

form (22.54), and an empirical histogram has still less reason to have that form. We compare realizable limiting laws before applying their inverse. This avoids extending a factorization identity to an arbitrary observed density.

Fix the local analytic neighborhood, radii, squares and anchor of Theorem 22.15. Let  $\mathcal{K}$  be a nonempty class of periodic realizations whose anchored contact pairs and gaps belong to this neighborhood. Require uniform positive curvature, separation, clear-channel and roof margins on the selected itinerary, and  $0 < A_- \leq A(T) \leq A_+ < \infty$ . We use ideal endpoint recording in this subsection. The signed frames and physical units are the observation conventions already specified for the law inverse. The sampling theorem below first treats an exactly calibrated channel; estimated calibration is treated afterwards. The class need not be finite dimensional.

All contact graphs have a common bound on  $D_R$ . Cauchy estimates on a strictly smaller physical collar give the finitely many real derivative bounds needed in the relative construction. By decreasing the analytic neighborhood and the collar, the weighted contractions, diagonal dominance and trace-norm estimates have uniform margins. Thus their endpoint estimates are uniform on  $\mathcal{K}$ ; no differentiation with respect to an infinite-dimensional parameter is being asserted. Fix a compact positive-offset interval  $I_d$  about  $d$ , contained in the common onset collar. Decrease  $Q_+$  once so that it lies strictly inside the positive region for every offset in  $I_d$  and every member of the class.

**Lemma 23.5** (Interior density and acceptance bounds). *There are  $j_0 \geq 2$ ,  $C, H, F, c > 0$ ,  $\omega > 0$ , and  $\Gamma \geq \sup_{T \in \mathcal{K}} \gamma(T) > 0$  such that, for every even  $j \geq j_0$ , both starting types  $b$ , and  $d' \in I_d$ , the actual successful endpoint density  $f_{j,b,d'}^T$  and its limiting density  $f_{b,d'}^T$  satisfy*

$$\|f_{j,b,d'}^T - f_{b,d'}^T\|_{C^0(Q_+)} \leq Ce^{-\omega j}, \quad (23.9)$$

$$\|f_{j,b,d'}^T\|_{\text{Lip}(Q_+)} + \|f_{b,d'}^T\|_{\text{Lip}(Q_+)} \leq H. \quad (23.10)$$

*Their zero extensions are bounded by  $F$  on a common endpoint box, and are zero outside it. Their densities are uniformly Lipschitz in the offset in the uniform norm. The probability  $p_{j,b,d'}^T$  of this selected success in one normalized full-phase preparation obeys*

$$p_{j,b,d'}^T \geq ce^{-\Gamma j}. \quad (23.11)$$

*Proof.* Suppress the indices  $T, b$  and put

$$F_{j,d'} = b_j(d' - E_j)_+, \quad F_{\infty,d'} = B_b(u)B_b(v)(d' - S_b(u) - S_b(v))_+.$$

Integrate over the fixed box containing their supports, and write the normalizers as  $Z_{j,d'}$  and  $Z_{\infty,d'}$ . The relative estimate (12.4) and the Lipschitz property of positive part give  $\|F_{j,d'} - F_{\infty,d'}\|_{\infty} \leq Ce^{-\omega j}$  and the same bound for the normalizer difference. A fixed small square at the origin gives  $Z_{j,d'}, Z_{\infty,d'} \geq z_- > 0$ , after enlarging  $j_0$ . Division proves (23.9), in fact in the uniform norm on the full box. On  $Q_+$  the cutoff is inactive, so the first endpoint-derivative estimate gives (23.10). Bounded numerators and the lower normalizer bound give  $F$ . Also

$$\|F_{j,d'} - F_{j,d''}\|_{\infty} \leq C|d' - d''|, \quad |Z_{j,d'} - Z_{j,d''}| \leq C|d' - d''|.$$

These estimates, and their limiting versions, prove offset continuity without differentiating a moving boundary.

The exact full-phase formula gives

$$p_{j,b,d'}^T = \frac{d_{j,b}^0}{2\pi A(T)} Z_{j,b,d'}, \quad d_{j,b}^0 = \frac{a_b}{\sinh(j\gamma(T))}.$$

The positive lower bound for  $a_b$ ,  $A(T) \leq A_+$ , and  $1/\sinh(j\gamma) \geq 2e^{-j\gamma}$  give (23.11). The itinerary is selected throughout; no nonminimal channel is identified with the complete maximal-count event.  $\square$

**Theorem 23.6** (Complete contact stability from finite bridges). *For two members  $T, \tilde{T} \in \mathcal{K}$ , let  $j, k \geq j_0$  be even, let their physical offsets both equal  $d$ , and put*

$$e_{j,k} = \max_b \|f_{j,b,d}^T - f_{k,b,d}^{\tilde{T}}\|_{C^0(Q)}.$$

*For sufficiently small errors,*

$$\|\psi^T - \psi^{\tilde{T}}\|_r \leq C \left\{ |g_T - g_{\tilde{T}}| + (e_{j,k} + e^{-\omega j} + e^{-\omega k})^\vartheta \right\}. \quad (23.12)$$

*Every  $0 < t < r$  gives the simultaneous coefficient bound*

$$\sum_{n \geq 2} \frac{t^n}{n!} \max_b |(\psi_b^T)^{(n)}(0) - (\psi_b^{\tilde{T}})^{(n)}(0)| \leq \frac{(t/r)^2}{1 - t/r} \|\psi^T - \psi^{\tilde{T}}\|_r. \quad (23.13)$$

*The constants and exponent do not depend on a jet truncation order.*

*Proof.* Apply (23.9) at the two ends of the triangle inequality comparing the limiting densities. Their  $C^0(Q)$  distance is at most  $e_{j,k} + C(e^{-\omega j} + e^{-\omega k})$ . Theorem 22.15 applies to these two limiting laws and gives (23.12). Cauchy's coefficient estimate on  $D_r$  gives (23.13). We have not applied the separated law inverse to either finite density.  $\square$

*A fixed preparation budget.* For the next theorem assume that  $g$  and the two contact frames are exactly calibrated and fixed throughout  $\mathcal{K}$ . All window times and endpoint categories are then programmed without the unknown table. Choose an even  $j$  and perform  $N$  independent normalized full-phase preparations for each starting-type experiment, each of duration  $jg + d$ . A preparation returns failure, or the endpoint category on the prescribed selected itinerary. Partition  $Q_+$  into  $J$  equal squares of side  $h$  and retain its complement as an additional category. In particular, endpoints outside  $Q_+$  are not discarded from the successful-sample denominator. Let  $M_b$  be the number of successes in group  $b$  and  $C_{b,i}$  the number in cell  $i$ . For  $M_b > 0$  set

$$\hat{f}_{b,i} = \frac{C_{b,i}}{M_b h^2}. \quad (23.14)$$

The counters  $M_b$  are used for normalization and may subsequently be forgotten. No ancillary interpretation of the waiting counts is needed.

**Theorem 23.7** (Analytic recovery with every preparation charged). *Under the preceding hypotheses let  $\alpha \in (0, 1/2)$ , set*

$$x = \log \frac{16(J+1)}{\alpha}, \quad m_0 = \frac{c}{2} N e^{-\Gamma j},$$

*and assume  $cN e^{-\Gamma j} \geq 8 \log(8/\alpha)$ . For every  $\xi > 0$  there is a measurable estimator with values in  $\mathcal{K}$  such that, with probability at least  $1 - \alpha$ ,*

$$\|\hat{\psi} - \psi\|_r \leq C \left( e^{-\omega j} + h + \sqrt{\frac{x}{m_0 h^2}} + \frac{x}{m_0 h^2} + \xi \right)^\vartheta, \quad (23.15)$$

*whenever the quantity in parentheses is sufficiently small. The experiment uses exactly  $2N$  preparations and evolution time  $2N(jg + d)$ , including all failures. No finite-dimensional model for the contact functions is imposed.*

*More precisely, for  $L_N = \log(C_* N / \alpha)$  and all sufficiently large  $N / L_N$ , an even choice*

$$j_N = 2 \left\lfloor \frac{\log(N / L_N)}{2(4\omega + \Gamma)} \right\rfloor \geq j_0, \quad h_N \asymp \left( \frac{L_N e^{\Gamma j_N}}{N} \right)^{1/4}, \quad 0 < \xi \leq N^{-1}, \quad (23.16)$$

with  $h_N$  rounded to an equal-square grid, gives

$$\sup_{T \in \mathcal{K}} \Pr_T \left\{ \|\hat{\psi} - \psi^T\|_r > C \left( \frac{L_N}{N} \right)^{\vartheta\omega/(4\omega+\Gamma)} \right\} \leq \alpha. \quad (23.17)$$

The same event controls all coefficients by (23.13). The rate is an upper bound, not an optimality assertion.

*Proof.* The binomial lower-tail estimate and (23.11) give

$$\Pr\{\min_b M_b < m_0\} \leq 2 \exp(-cNe^{-\Gamma^j}/8) \leq \alpha/4. \quad (23.18)$$

For example the usual exponential Markov argument at parameter  $-\log 2$  bounds  $\Pr\{M \leq Np/2\}$  by  $e^{-Np/8}$ ; the displayed bound also holds when the threshold is decreased to  $m_0$ . Conditional on  $M_b = m$ , the successful marks have their actual iid conditional law. This assertion follows directly by factoring the joint density at each fixed set of  $m$  successful indices and summing over those sets. We use a fixed number of attempts, not a stopped experiment conditioned on completing its cap.

For a cell probability  $p_i \leq Fh^2$ , the Bernoulli exponential bound gives, simultaneously in both types and all  $J$  cells, except on an event of probability at most  $4Je^{-x} \leq \alpha/4$ ,

$$\left| \hat{f}_{b,i} - h^{-2} \int_{C_i} f_{j,b,d} \right| \leq \sqrt{\frac{2Fx}{mh^2}} + \frac{2x}{3mh^2}. \quad (23.19)$$

For completeness, for a Bernoulli variable  $X$  of mean  $p$  and  $0 < \lambda < 3$ ,

$$\log \mathbb{E}e^{\lambda(X-p)} \leq p(e^\lambda - 1 - \lambda) \leq \frac{p\lambda^2}{2(1 - \lambda/3)}.$$

The same upper bound with  $|\lambda|$  holds for the lower tail. Applying exponential Markov and optimizing gives  $|\bar{X} - p| \leq \sqrt{2px/m} + 2x/(3m)$  with failure probability at most  $2e^{-x}$ . This proves (23.19). The bound is uniform for every  $m \geq m_0$ ; integrating over these counts and adding (23.18) yields a simultaneous probability at least  $1 - \alpha$ .

Here is an explicit measurable selection construction. The image

$$\{(f_{0,d}^T|_{Q_+}, f_{1,d}^T|_{Q_+}) : T \in \mathcal{K}\} \subset C(Q_+)^2$$

is a separable metric space. Choose once a countable dense subset of this image and a realizing table  $T_l$  for each of its members. Define

$$D_l = \max_{b,i} \left| h^{-2} \int_{C_i} f_{b,d}^{T_l} - \hat{f}_{b,i} \right|.$$

On nonempty success groups take the first  $l$  with  $D_l < \inf_{q \geq 1} D_q + \xi$ , and otherwise take a fixed member of the class. The countable infimum and this first-index rule are measurable. The output contact pair is measurable in the  $D_r$  norm as well, since its range is countable. Existence of this estimator is not an assertion of a finite computational search bound.

Put  $\varepsilon_N = \sqrt{2Fx/(m_0h^2)} + 2x/(3m_0h^2)$ . Density of the selected image, the approximation lemma and (23.19) give  $\inf_l D_l \leq Ce^{-\omega^j} + \varepsilon_N$ . Two triangle inequalities therefore bound the cell-average difference between the fitted and true limiting densities by  $2Ce^{-\omega^j} + 2\varepsilon_N + \xi$ . On each cell, their common Lipschitz bound supplies

$$\|\hat{f}_{b,d}^T - f_{b,d}^T\|_{C^0(Q_+)} \leq 2Ce^{-\omega^j} + 2\varepsilon_N + \xi + 2\sqrt{2}Hh. \quad (23.20)$$

This is a comparison of two realizable limiting laws, not a regularity claim about the empirical histogram. Apply Theorem 22.15 to obtain (23.15).

To prove the rate, put  $a_N = L_N/N$  and  $\zeta = \omega/(4\omega + \Gamma)$ . Rounding  $j_N$  down by less than two gives

$$e^{-\omega j_N} \leq e^{2\omega} a_N^\zeta, \quad \left( \frac{L_N e^{\Gamma j_N}}{N} \right)^{1/4} \leq a_N^\zeta.$$

The rounded grid satisfies  $h_N \asymp (L_N e^{\Gamma j_N}/N)^{1/4}$  and  $J_N \leq C(N/L_N)^{2\zeta}$ . Since  $2\zeta < 1/2$ , enlarging the fixed  $C_*$  ensures  $x \leq CL_N$ . Also  $Ne^{-\Gamma j_N}/L_N \geq (N/L_N)^{4\zeta} \rightarrow \infty$ , so the success-count condition holds. The square-root term in (23.15) is  $O(a_N^\zeta)$  and the linear term is  $O(a_N^{2\zeta})$ . The selection error  $N^{-1}$  is smaller. This proves (23.17) and the asserted charges.  $\square$

The exponent records three different operations: relative finite-flight error, rare-event thinning, and estimation of a two-dimensional density with a first-order real bound. The outer analytic prior and the inner radius determine  $\vartheta$ . The experiment is a sequence of independent full-phase preparations, not a mixing assertion for one continuously observed orbit. The theorem does not assert a rate for global analytic continuation, lattice registration, unknown channel discovery, or arbitrarily inefficient recording. Those operations retain their own hypotheses. In particular the exactly calibrated budget in (23.17) must not be called a calibration-inclusive rate.

*Estimated offset and endpoint coordinates.* The finite-budget estimate also has a calibration-aware form. Consider now variable gaps in  $\mathcal{K}$ . Suppose an independent preliminary experiment has deterministic cap  $B_{\text{pil}}$  and, except with probability  $\beta$ , supplies correct signed channel labels, a gap estimate  $|\hat{g} - g| \leq v_g$ , and observable endpoint projections inducing a Borel map  $\Phi$  with  $\|\Phi - \text{id}\|_\infty \leq r_c$  on the actual endpoint collar. Program the subsequent windows at  $j\hat{g} + d + \epsilon_t$ , with  $|\epsilon_t| \leq v_t$ . Require  $\Delta = jv_g + v_t$  small enough that all actual offsets belong to  $I_d$ , and  $0 \leq r_c \leq h/2$ . The observed variables after this preliminary experiment are fixed maps of the raw record; no unknown true chart is applied by the observer.

**Corollary 23.8** (A charged bound with calibration errors). *Under these hypotheses perform the two groups of  $N$  fresh attempts as above. With  $m_0, x$  and the count condition of Theorem 23.7, there is a measurable estimator satisfying, with probability at least  $1 - \alpha - \beta$ ,*

$$\|\hat{\psi} - \psi\|_r \leq C \left[ v_g + \xi + \left( e^{-\omega j} + \Delta + h + \frac{r_c}{h} + \sqrt{\frac{x}{m_0 h^2}} + \frac{x}{m_0 h^2} + \xi \right)^\vartheta \right]. \quad (23.21)$$

*The complete deterministic preparation charge is  $B_{\text{pil}} + 2N$ .*

*Proof.* Fix a true table and a good preliminary history. For a square cell  $C_i$  of side  $h$ , a change of its category under  $\Phi$  can occur only in the  $r_c$ -neighborhood of its boundary. For  $r_c \leq h/2$  that neighborhood has area at most

$$(h + 2r_c)^2 - (h - 2r_c)^2 = 8hr_c.$$

The global density bound in Lemma 23.5 therefore gives a cell-mass error at most  $8Fhr_c$  and a cell-average error at most  $8Fr_c/h$ . This includes the outer edges of  $Q_+$ ; the outside category is still included in  $M_b$ . No density or Jacobian for  $\Phi_\# P$  is assumed. Its cell masses are at most  $Fh^2 + 8Fhr_c \leq 5Fh^2$ , which suffices for (23.19) with  $F$  replaced by  $5F$ .

Uniform offset continuity adds  $C\Delta$  and the finite-flight lemma adds  $Ce^{-\omega j}$ . The success bound is uniform on  $I_d$ , hence the count and cell concentration estimates hold conditionally on every good preliminary history. Choose countably many representatives dense in the joint image  $(g, f_0|_{Q_+}, f_1|_{Q_+})$ . In the previous first-index selection allow only candidates with  $|g_l - \hat{g}| < v_g + \xi$ . On a good history this set is nonempty and approximates the true density image arbitrarily well: the strict slack  $\xi > 0$  keeps the true image inside the permitted gap strip. The selected gap

differs from the true one by at most  $2v_g + \xi$ . Off that event use any fixed default when necessary. The allowed indices and their residual infimum are measurable.

Repeat (23.20) with the two additional biases  $C\Delta + 8Fr_c/h$ , and apply the two-point analytic inverse with the selected gap error. This proves (23.21) conditionally. Integrate over good histories and add  $\beta$ . The two subsequent groups always use exactly  $2N$  attempts; adding the preliminary cap proves the charge without conditioning on cap completion.  $\square$

For the persistent position-sensor design, the existing calibration theorem 56.2 supplies  $jv_g \leq \varepsilon_{\text{pil}}$  and  $r_c \leq C_\chi \sqrt{\varepsilon_{\text{pil}}}$  under its stated compact marked-class hypotheses, with the explicit cap (56.5). It is a substantive input to this application. With an exact clock, choosing  $\varepsilon_{\text{pil}} \leq c_1 h^4$  makes both  $\Delta$  and  $r_c/h$   $O(h)$ , and  $r_c \leq h/2$  for small  $h$ . The charge is then  $B_{\text{pil}}(j, \varepsilon_{\text{pil}}, \beta) + 2N$ , not  $2N$ . The pilot cost can also be estimated explicitly. This yields a slower rate in the complete preparation budget rather than treating calibration as a free oracle.

**Corollary 23.9** (Complete preparation budget with the position pilot). *Assume in addition the compact marked-class hypotheses of Theorem 56.2, persistent position recording, and an exact clock. The number of marked channels is fixed, and the channel and starting-type labels used by that theorem are retained. There are  $C_0, C_1 > 0$  such that, for  $0 < \alpha < 1/2$  and all sufficiently large  $B/[\log(C_0 B/\alpha)]^2$ , a parameter-independent experiment with at most  $B$  preparations, counting the pilot and every rejected attempt, has a measurable estimator satisfying*

$$\sup_{T \in \mathcal{K}} \Pr_T \left\{ \|\hat{\psi} - \psi^T\|_r > C_1 \left( \frac{[\log(C_0 B/\alpha)]^2}{B} \right)^{\vartheta\omega/(12\omega+\Gamma)} \right\} \leq \alpha. \quad (23.22)$$

Here  $\Gamma$  is enlarged, if necessary, to bound also the hyperbolicity exponent in the pilot mass estimate. The total evolution time is at most  $CB \log(C_0 B/\alpha)$ . All coefficient bounds (23.13) hold on the same event.

*Proof.* Clip the preliminary gap estimate to the known interval  $[g_-, g_+]$  (and use a fixed value there on a failed pilot). This cannot enlarge the good-history error and bounds every programmed time on all histories. Use preliminary failure probability  $\alpha/2$ , and apply Corollary 23.8 with subsequent failure probability  $\alpha/2$ . The pilot has success lower bound  $p_*(j, \varepsilon) = c_* \varepsilon^2 e^{-\Gamma j}$  at a grid point within  $[\varepsilon, 2\varepsilon]$  of onset. Its grid contains at most  $C(j+1)/\varepsilon$  points. Thus the exact integer cap (56.5) satisfies, for  $j \geq j_0$  and  $0 < \varepsilon < \varepsilon_0$ ,

$$B_{\text{pil}}(j, \varepsilon, \alpha/2) \leq C(j+1)e^{\Gamma j} \varepsilon^{-3} \log(C/\alpha). \quad (23.23)$$

The ceiling terms are absorbed since the logarithm is bounded below and  $\varepsilon$  is small.

Choose an equal-square grid with  $h \asymp e^{-\omega j}$ , put  $\varepsilon = c_1 h^4$  with fixed sufficiently small  $c_1 > 0$ , and take

$$N = \left\lceil C_2 e^{\Gamma j} h^{-4} \log \frac{C_2}{\alpha h} \right\rceil, \quad 0 < \xi \leq h^2.$$

For sufficiently large fixed  $C_2$ , the count condition holds and both concentration terms are  $O(h)$  (the linear term is  $O(h^2)$ ). The biases  $e^{-\omega j}$ ,  $\Delta$ , and  $r_c/h$  are  $O(h)$ ; the additive gap error is  $O(h^4/j)$ . Hence (23.21) gives error  $Ch^\vartheta$ . The complete cap satisfies

$$\begin{aligned} B_{\text{pil}} + 2N &\leq C(j+1)e^{\Gamma j} h^{-12} \log \frac{C}{\alpha h} \\ &\leq C(j+1)e^{(12\omega+\Gamma)j} (\log(C/\alpha) + \omega j). \end{aligned}$$

The grid-rounding constants are fixed.

Let  $L = \log(C_0 B / \alpha)$  and choose the even integer

$$j = 2 \left\lfloor \frac{\log\{B/(C_3 L^2)\}}{2(12\omega + \Gamma)} \right\rfloor \geq j_0.$$

Then  $j+1 \leq CL$  and  $\log(C/\alpha) + \omega j \leq CL$ . Taking the fixed  $C_3$  sufficiently large makes the last cap at most  $B$ , including all ceilings. Rounding down by less than two gives  $h \leq C(L^2/B)^{\omega/(12\omega+\Gamma)}$ , which proves (23.22). To bound evolution time on every history, retain the ceiling in the pilot grid. With  $q = j(g_+ - g_-)/\varepsilon$  and  $L_{\text{grid}} = \lceil q \rceil + 2$ , its final time satisfies

$$t_{L_{\text{grid}}} - jg_+ = (\lceil q \rceil - q + 2)\varepsilon \in [2\varepsilon, 3\varepsilon). \quad (23.24)$$

Thus every programmed pilot time is strictly less than  $jg_+ + 3\varepsilon$ . On the successful stopping event used in the pilot proof one has the sharper bound  $T_{e,b} \leq jg_e + 2\varepsilon$ ; that event-specific estimate is not an all-history bound on the prescribed grid. A failed scan may reach  $t_{L_{\text{grid}}}$ . The cap already counts all  $L_{\text{grid}} + 1$  grid points, and is unchanged. Both the pilot times and the programmed later times are bounded by  $C(j+1)$ . Multiplying by the deterministic cap gives the stated evolution-time bound, including failed histories. The labels, grid, tolerance, stopping caps and subsequent window times are either fixed in advance or measurable functions of the preliminary record, as required.  $\square$

The power 12 in this bound comes from the concrete pilot rather than from an asserted intrinsic information barrier: the onset scan costs  $\varepsilon^{-3}$ , while displacement control for histogram cells uses  $\varepsilon \asymp h^4$ . Neither preparation-budget exponent is claimed optimal. This result complements the retained full-table acquisition result 50.10. It recovers complete local analytic germs at a quantitative rate; global continuation, matching and lattice recovery retain their separate hypotheses and stability moduli.

**23.5. Equivariant reconstruction of perturbed densities.** The inverse on the exact factorized family is independent of its anchor. An extension to nearby densities is not. We now specify an extension which commutes with a change of the common transverse orientation, including away from the factorized family. The magnitude of the anchor is fixed once; its sign is part of the coordinate convention, not an additional observation.

Write  $\varsigma(u) = -u$  and  $f^r(u, v) = f(-u, -v)$ . For each nonzero  $\alpha \in I$ , make the anchor in (23.3)–(23.4) explicit:

$$q_\alpha(f) = \sqrt{1 - R_f(\alpha, \alpha)}, \quad T_\alpha(f)(u) = \frac{1 - R_f(u, \alpha)}{q_\alpha(f)}, \quad (23.25)$$

$$S_{d,\alpha}(f) = \frac{dT_\alpha(f)}{1 + T_\alpha(f)}, \quad B_\alpha(f)(u) = \frac{f(u, 0)}{f(0, 0)} \{1 + T_\alpha(f)(u)\}. \quad (23.26)$$

For  $M \geq 0$ , let  $\mathcal{U}_\alpha^M$  be the set of real  $C^M$  functions on a neighborhood of  $I^2$  satisfying

$$\inf_{I^2} f > 0, \quad 1 - R_f(\alpha, \alpha) > 0, \quad \inf_I \{1 + T_\alpha(f)\} > 0. \quad (23.27)$$

All these conditions are open in the restricted  $C^M$  norm. Neither factorization nor a mass normalization is required. Fix  $0 < a < r$ . In orientation sector  $\epsilon \in \{+1, -1\}$  use  $\alpha = \epsilon a$  and set

$$\mathcal{F}_{d,a}(\epsilon, f) = (\epsilon, S_{d,\epsilon a}(f), B_{\epsilon a}(f)).$$

The common reversal acts by  $(\epsilon, f) \mapsto (-\epsilon, f^r)$ , and on the output by changing the tag and precomposing both functions with  $\varsigma$ . For a channel collection the same tag is used for every contact and edge; fixed positive anchor magnitudes may differ between edges.

**Proposition 23.10** (Transported-anchor equivariance and stability). *Reflection is an isometric bijection  $\mathcal{U}_\alpha^M \rightarrow \mathcal{U}_{-\alpha}^M$ , and on this entire domain*

$$\begin{aligned} T_{-\alpha}(f^r) &= T_\alpha(f) \circ \varsigma, \\ S_{d,-\alpha}(f^r) &= S_{d,\alpha}(f) \circ \varsigma, \quad B_{-\alpha}(f^r) = B_\alpha(f) \circ \varsigma. \end{aligned} \quad (23.28)$$

In particular  $\mathcal{F}_{d,a}$  is equivariant, also for nonfactorized perturbations. On sets on which

$$\inf f \geq m > 0, \quad \|f\|_{C^M} \leq K, \quad q_{\epsilon a}(f) \geq q_* > 0, \quad \inf(1 + T_{\epsilon a}(f)) \geq \nu > 0, \quad (23.29)$$

it is Lipschitz in each sector with one constant  $C_M = C(M, I, a, d, m, K, q_*, \nu)$ , the same in the two sectors. The estimate holds for every pair of functions satisfying these bounds, not only for a perturbation compared with an exact model member.

Every compact exact-model family satisfying the hypotheses of Proposition 23.2, together with its reflected copy, has a neighborhood of this form when the anchor is transported. On the exact family the output is  $(\epsilon, S, B/B(0))$ , independently of the admissible anchor magnitude.

*Proof.* Substitution in the four-density ratio gives

$$R_{f^r}(u, v) = R_f(-u, -v), \quad q_{-\alpha}(f^r) = q_\alpha(f).$$

It follows that  $T_{-\alpha}(f^r)(u) = T_\alpha(f)(-u)$ . Substituting in (23.26), including the density ratio on the axis, proves all of (23.28). Each strict inequality in (23.27) is preserved. Reflection is an involution and preserves the supremum of every partial derivative of order at most  $M$ , proving the domain and isometry assertions.

For completeness, all uniform estimates can be proved between the two endpoints, without requiring a segment to remain in the formula domain. Restriction to the axis or to a fixed slice has operator norm at most one for the usual maximum-derivative norm. Leibniz's rule bounds products in  $C^M$ . If  $h \geq m$ , differentiating  $hh^{-1} = 1$  bounds the  $C^M$  norm of  $h^{-1}$  in terms of  $m^{-1}$  and  $\|h\|_{C^M}$ . For two such functions, the identity  $h^{-1} - k^{-1} = (k - h)h^{-1}k^{-1}$  gives the corresponding difference bound. Applying these estimates to  $R_f$  gives

$$\|R_f - R_g\|_{C^M(I^2)} \leq C\|f - g\|_{C^M(I^2)}.$$

At the chosen anchor,

$$|q_\alpha(f) - q_\alpha(g)| = \frac{|R_f(\alpha, \alpha) - R_g(\alpha, \alpha)|}{q_\alpha(f) + q_\alpha(g)} \leq \frac{C}{2q_*}\|f - g\|_{C^M(I^2)}.$$

The lower bound on  $q_\alpha$  now gives a difference bound for  $T_\alpha$ . Its  $C^M$  norm is uniformly bounded by the same estimates. The reciprocal bound using  $\nu$ , followed by the product bound, proves the claims for  $S_{d,\alpha}$  and  $B_\alpha$ . Reflection preserves every bound used, so the constants agree in the two sectors.

For the last assertion, on an exact model  $1 + T_\alpha \geq 1$  and  $q_\alpha = S(\alpha)/(d - S(\alpha)) > 0$ . The compact lower bounds give a common neighborhood with positive density denominators, anchor radical and  $1 + T_\alpha$ . Its reflected copy has exactly the same bounds. Exact recovery and anchor independence follow from (23.6). No differentiated square root at the action minimum is taken.  $\square$

**Corollary 23.11** (Orbit stability and finite-order contact reconstruction). *The transported-anchor action and amplitude reconstruction descends to the common-orientation quotient with the same fixed-order Lipschitz constant. On a compact positive contact class, the finite-jet reconstruction of Corollary 23.3, with this anchor convention, is also equivariant. Its quotient error bound through any fixed order  $M \geq 2$  remains  $C_M(\tau^j + \varepsilon + \varepsilon_g)$ . Here  $\varepsilon$  is the interior  $C^M$  density error and  $\varepsilon_g$  is the gap error, as in that corollary.*



*Proof.* For tagged representatives  $(\epsilon, f)$ ,  $(\delta, g)$ , let  $\eta \in \{0, 1\}$  be the unique value with  $\epsilon = (-1)^\eta \delta$ . Define their orbit distance by

$$\bar{d}_M([\epsilon, f], [\delta, g]) = \|f - g^{r^\eta}\|_{C^M(I^2)}.$$

It is independent of the chosen representatives because reflection is an isometry. Define the output distance by the sum of the two corresponding function norms, using the same tag alignment. Equivariance and the same-sector estimate in Proposition 23.10 prove the quotient estimate with the same constant. For finitely many contact laws use the corresponding sum or maximum product norm throughout.

The anchored projection  $Ph(u) = h(u) - h(0) - uh'(0)$  is bounded on  $C^M(I)$  and satisfies  $P(h \circ \varsigma) = (Ph) \circ \varsigma$ . Reflection fixes the gap and both quadratic action coefficients. It therefore fixes the leading curvatures and  $\gamma$  in (22.3). At degree  $n$ , it multiplies both graph and action coefficients by  $(-1)^n$ . The finite smooth-jet factorization and reflection of the half-line action show that the lower-jet remainder obeys

$$R_n((( -1)^m q_m)_{2 \leq m < n}) = (-1)^n R_n((q_m)_{2 \leq m < n}).$$

One may verify this identity on smooth polynomial representatives; finite smooth-jet factorization makes it independent of their cutoffs and remainders. The block  $M_n$  depends only on the reflection-invariant quadratic geometry. Thus the recursion  $q_n = M_n^{-1}(s_n - R_n)$  commutes with reflection, by induction in  $n$ . Its finitely many operations are smooth on the compact positive set and on a sufficiently small neighborhood of it. Combining their Lipschitz bounds with the density estimate and the finite-flight estimate proves the stated error bound. The alignment argument above then gives the same bound on orientation orbits.  $\square$

The extension just constructed is an action, amplitude and finite-jet reconstruction on its specified neighborhood. An arbitrary perturbed density is not asserted to be the exact law of a periodic billiard, and its output is not asserted to admit analytic continuation or global incidence gluing. The exact global classification still uses the exact law image. All constants remain fixed-order constants.

*Remark 23.12* (Why a fixed signed anchor cannot be used in both sectors). For a fixed signed anchor the identity is instead  $S_{d,a}(f^r) = S_{d,-a}(f) \circ \varsigma$ . This cannot in general be replaced by  $S_{d,a}(f) \circ \varsigma$ . For example, take  $d = 1$  and

$$f_\zeta(u, v) = \frac{2}{\pi}(1 - u^2 - v^2)_+ \{1 + \zeta uv(u + v)\}, \quad 0 < \zeta \leq 1/100.$$

This is a probability density on the unit disk: its perturbation integrates to zero under simultaneous reversal, and its multiplier is positive. On  $[-1/3, 1/3]^2$  it is a positive smooth perturbation of the exact density with  $S(u) = u^2$ . Set  $a = 1/4$ ,  $D = 1/225$ , and  $h = 2\zeta a^3(224/225)$ . Direct substitution yields

$$T_a(f_\zeta^r)(a)^2 = D + h, \quad T_a(f_\zeta)(-a)^2 = \frac{D^2}{D - h}, \quad (D + h) - \frac{D^2}{D - h} = -\frac{h^2}{D - h} \neq 0.$$

Both unsquared quantities are positive. Hence the action outputs differ as well. At  $\zeta = 1/100$  the last difference is  $-49/2092500$ . This is a witness concerning an off-model formula, not physical-table nonidentifiability. Transporting the anchor restores the identity exactly, without removing the perturbation neighborhood.

**23.6. What one density adds to one support.** The family of support thresholds determines the action, but the support at one offset need not do so. The following functional comparison isolates the additional information used by the four-density identity.

**Proposition 23.13** (A support-preserving action change). *On  $[-1, 1]$  let  $d = 1$ ,  $S(u) = u^2$ , and, for  $0 < \epsilon \leq 1/10$ , set*

$$\phi_\epsilon(s) = s + \epsilon s(1-s)(s-1/2), \quad \tilde{S}(u) = \phi_\epsilon(S(u)).$$

*Both actions vanish to second order at zero, are strictly convex, and are different. Nevertheless*

$$S(u) + S(v) < 1 \iff \tilde{S}(u) + \tilde{S}(v) < 1. \quad (23.30)$$

*Their endpoint densities with any positive separable amplitudes cannot have the same four-density invariant on an interior neighborhood of zero.*

*Proof.* For  $0 \leq s \leq 1$ ,

$$\phi'_\epsilon(s) = 1 + \epsilon(-3s^2 + 3s - 1/2) \geq 1 - \epsilon/2 > 0, \quad \phi_\epsilon(1-s) = 1 - \phi_\epsilon(s).$$

Consequently, for  $s, t \in [0, 1]$ ,  $s + t < 1$  is equivalent to  $\phi_\epsilon(t) < \phi_\epsilon(1-s) = 1 - \phi_\epsilon(s)$ . This proves (23.30). Also

$$\tilde{S}''(u) = 2\{\phi'_\epsilon(s) + 2s\phi''_\epsilon(s)\}_{s=u^2}, \quad \phi'_\epsilon(s) + 2s\phi''_\epsilon(s) = 1 + \epsilon(-15s^2 + 9s - 1/2) \geq 1 - 13\epsilon/2 > 0.$$

The zero value and derivative at the origin follow by composition. For  $0 < S(u) < 1/2$  the correction is nonzero, so the actions differ on every neighborhood of zero. On a sufficiently small interior square, (23.6) and the fixed-anchor inverse recover  $S$  from its four-density invariant, irrespective of the positive amplitude. Equality of the invariants would therefore make the two actions equal, a contradiction.  $\square$

This example belongs to the functional action–density class. It is not an assertion that both actions have been realized by physical billiard tables satisfying the global hypotheses. The physical theorem uses the relative-law construction for that realization. The comparison shows precisely why retaining density, rather than just one support, enters Theorem 23.1.

### 23.7. Analytic continuation of a registered boundary germ.

**Lemma 23.14** (A registered analytic germ determines the connected boundary). *Let  $C, C' \subset \mathbb{R}^2$  be connected embedded real-analytic strictly convex closed curves. Suppose there are points  $p \in C$ ,  $p' \in C'$  and an orientation-preserving Euclidean motion  $A$  carrying the oriented point–tangent frame at  $p$  to that at  $p'$ , such that  $A(C)$  and  $C'$  agree on a nonempty boundary arc containing  $p'$ . Then*

$$A(C) = C'.$$

*Equivalently, equality of a nonempty registered analytic boundary germ implies equality of the complete connected boundary image.*

*Proof.* Parametrize  $C$  and  $C'$  by positively oriented arclength from  $p$  and  $p'$ , with lengths  $L$  and  $L'$ . Their positive signed curvature functions  $\kappa$  and  $\kappa'$  are real analytic. Equality of the registered arcs gives equality of the curvature germs at the origins. Analytic continuation on the universal arclength covers therefore gives one real-analytic positive function  $\tilde{\kappa} : \mathbb{R} \rightarrow (0, \infty)$  which is both  $L$ -periodic and  $L'$ -periodic and restricts to the two curvature functions.

Strict convexity gives total turning

$$\int_0^L \tilde{\kappa}(s) ds = 2\pi = \int_0^{L'} \tilde{\kappa}(s) ds.$$

If, say,  $L' < L$ , positivity would give  $\int_{L'}^L \tilde{\kappa} > 0$ , contradicting these two equalities; the opposite inequality is identical. Hence  $L = L'$ .

Finally the planar Frenet system

$$X' = T, \quad T' = \tilde{\kappa}N, \quad N' = -\tilde{\kappa}T$$

has unique solutions for a prescribed oriented point–tangent frame. The motion  $A$  matches those initial data, so the two arclength parametrizations agree on  $[0, L]$ . Their complete boundary images are therefore equal.  $\square$

## 24. INTERACTION DATA WITH UNKNOWN TRANSVERSE ORIGINS

The geometric inverse uses the interaction of the two endpoint coordinates, not the individual flux amplitudes. This distinction also permits the transverse contact origins to be recovered from the law. We prove this for an observation in which the two numerical origins are unknown and successful records are subject to unknown, separate endpoint efficiencies. The signed units and directions of the transverse coordinates, the gap and the excess time remain specified. In particular, an unknown origin is not an unknown nonlinear change of the physical coordinate.

**24.1. The recorded law and its interaction.** Let  $I$  be a bounded interval containing zero. Let  $S$  be smooth on a neighborhood of  $\bar{I}$ , with

$$S(0) = S'(0) = 0, \quad S'' > 0.$$

Fix  $d > 0$  so small that  $\{u \in I : S(u) \leq d\}$  is compactly contained in  $I$ . For unknown  $\xi, \eta \in \mathbb{R}$  and positive smooth functions  $A$  and  $C$  on neighborhoods of  $\xi + \bar{I}$  and  $\eta + \bar{I}$ , respectively, consider the probability density

$$\tilde{f}(x, y) = \frac{A(x)C(y)}{Z} [d - S(x - \xi) - S(y - \eta)]_+, \quad (x, y) \in (\xi + I) \times (\eta + I), \quad (24.1)$$

extended by zero. Here  $Z$  is the positive integral of the numerator. The positive set

$$D = \{(x, y) : S(x - \xi) + S(y - \eta) < d\} \quad (24.2)$$

is understood in the displayed coordinate box. It is bounded, connected, and has interval fibers. Its closure is inside that box. Strict positivity of  $A, C$  is important: no part of the cap is hidden by a zero recording efficiency. The law determines  $D$  and its smooth interior density; a continuous density is determined everywhere in its open positive set by its almost-everywhere class.

For a same-type bridge, the physical limiting density of Theorem 17.3 has this form with  $A(\xi + u) = B(u)$  and  $C(\eta + v) = B(v)$ . Unknown recording probabilities  $e_1(x), e_2(y) > 0$  multiply these factors before normalization. Conversely, since  $B$  is positive on the compact collar, the positive smooth ratios  $A(\xi + u)/B(u)$  and  $C(\eta + v)/B(v)$  can each be rescaled to recording probabilities at most one without changing the normalized law. Thus every pair of effective factors in this class is allowed, and  $A, C$  need not agree. For fixed smooth efficiencies, the corresponding reweighted finite-flight laws have this limit: bounded continuous reweighting preserves weak convergence and the limiting normalizing mass is positive. No uniform bound over efficiencies tending to zero is asserted.

On  $D$  put

$$\mathcal{K}_{\tilde{f}}(x, y) = \partial_x \partial_y \log \tilde{f}(x, y) = \frac{\tilde{f} \tilde{f}_{xy} - \tilde{f}_x \tilde{f}_y}{\tilde{f}^2}. \quad (24.3)$$

This mixed logarithmic derivative is the classical local dependence function of Holland and Wang [40]; the related odds-ratio description of association is developed in [41]. We use those invariants in physical transverse coordinates, not after a transformation to marginal quantiles. The facts needed below are proved directly. No sampling-equivalence or efficiency theorem from association models is imported into the billiard experiment.

**Lemma 24.1** (Recovery of the two transverse origins). *The law (24.1) determines  $\xi$  and  $\eta$ , independently of  $A, C, Z$ . More precisely, with  $H = d - S(x - \xi) - S(y - \eta)$ ,*

$$\mathcal{K}_{\tilde{f}}(x, y) = -\frac{S'(x - \xi)S'(y - \eta)}{H^2}. \quad (24.4)$$

*Among vertical fibers of  $D$ , exactly the fiber at  $x = \xi$  has identically zero interaction. Among horizontal fibers, exactly the one at  $y = \eta$  has that property. If  $(\xi, y_*) \in D$  with  $y_* \neq \eta$ , then*

$$\partial_x \mathcal{K}_{\tilde{f}}(\xi, y_*) = -\frac{S''(0)S'(y_* - \eta)}{[d - S(y_* - \eta)]^2} \neq 0. \quad (24.5)$$

*The analogous horizontal root is simple as well.*

*Proof.* On  $D$ ,  $\log \tilde{f} = \log A + \log C + \log H - \log Z$ . The mixed derivative of the first, second and fourth terms vanishes. Since  $H_{xy} = 0$ , differentiating  $\log H$  gives (24.4). Strict convexity implies that  $S'$  vanishes exactly at zero. Every nonempty vertical fiber is an open interval of positive length, and so contains a point different from  $\eta$ . If  $x \neq \xi$ , the interaction is nonzero at such a point. For  $x = \xi$  it vanishes on the entire fiber. This proves the vertical characterization; interchanging the coordinates proves the horizontal one. Differentiating (24.4) at  $x = \xi$  eliminates the term containing  $S'(0)$  and yields (24.5). All derivatives are taken in the positive interior, not across the boundary of the cap.  $\square$

For the complete positive cap, there is also a support-only centering method: its fiber widths attain their unique maxima at the two origins. Lemma 24.5 gives that comparison. The interaction proof above is useful because its simple roots, and the finite-smooth estimate below, are confined to a positive interior rectangle. Theorem 24.6 and Corollary 24.7 state what survives when only such an interior window is recorded.

**24.2. Recovery of the action, with separate endpoint weights.** After determining the two origins, set  $f^c(u, v) = \tilde{f}(\xi + u, \eta + v)$ . Choose  $r > 0$  with  $[-r, r] \subset I$  and  $\max_{[-r, r]} S < d/4$ , and let  $a \in (-r, r) \setminus \{0\}$ . Define

$$R(u, v) = \frac{f^c(u, v)f^c(0, 0)}{f^c(u, 0)f^c(0, v)}, \quad (24.6)$$

$$q = \sqrt{1 - R(a, a)}, \quad T(u) = \frac{1 - R(u, a)}{q}, \quad \mathcal{S}(\tilde{f})(u) = \frac{dT(u)}{1 + T(u)}. \quad (24.7)$$

**Theorem 24.2** (The action as an invariant of the recorded interaction). *For every law (24.1), the construction above gives  $q > 0$  and  $\mathcal{S}(\tilde{f}) = S$  on  $[-r, r]$ . Hence the law recovers the signed action germ without supplied transverse origins or endpoint weights. Once  $S$  is known, the effective factors have the locally determined ratios*

$$\frac{A(\xi + u)}{A(\xi)} = \frac{f^c(u, 0)}{f^c(0, 0)}(1 + T(u)), \quad \frac{C(\eta + v)}{C(\eta)} = \frac{f^c(0, v)}{f^c(0, 0)}(1 + T(v)). \quad (24.8)$$

*For fixed  $d$ , two centered positive densities in this class have the same action germ if and only if their ratio is a product of a function of  $u$  and a function of  $v$  on some common square about zero. Thus the germ is a complete local invariant for separate endpoint reweighting after the origins are recovered. This does not separate a physical flux factor from an unknown recording factor multiplying it.*

*Proof.* Write  $A_0(u) = A(\xi + u)$  and  $C_0(v) = C(\eta + v)$ . Direct substitution, without assuming  $A_0 = C_0$ , gives

$$R(u, v) = \frac{d[d - S(u) - S(v)]}{[d - S(u)][d - S(v)]}, \quad 1 - R(u, v) = \frac{S(u)}{d - S(u)} \frac{S(v)}{d - S(v)}.$$

Since  $a \neq 0$  and  $S'' > 0$ , the positive scalar root in (24.7) equals  $S(a)/(d - S(a))$ . It follows that  $T(u) = S(u)/(d - S(u))$  for both signs of  $u$ . Solving this identity proves the action formula. In particular, no pointwise square root at the minimum is involved. Substitution in the two axis densities proves (24.8).

Equal actions imply separate factorization of the ratio of the densities by their defining formulas. Conversely, multiplying by a positive function of  $u$  and a positive function of  $v$  leaves  $R$  unchanged. The same nonzero anchor and (24.7) then recover the same action. Equivalently, equality of the two ratios  $R_f, R_h$  on a centered square gives the explicit factorization

$$\frac{f(u, v)}{h(u, v)} = \frac{f(u, 0)}{h(u, 0)} \frac{f(0, v)}{h(0, v)} \frac{h(0, 0)}{f(0, 0)}.$$

This assertion is local in the common positive square. It does not claim that arbitrary prescribed marginal laws are compatible with a given nonrectangular cap, nor that a joint recording factor depending on both coordinates would cancel.  $\square$

**24.3. Finite-order regularity of the extraction.** The identification of a zero fiber is an exact operation on a law. For a local error estimate it is preferable to use simple roots on two fixed slices. This also makes the derivative with respect to the observation explicit and avoids differentiating a set-valued zero-fiber operation.

**Proposition 24.3** (A local differentiable extraction with finite loss). *Fix a law  $\tilde{f}_0$  of the preceding class. For every integer  $m \geq 0$ , there are a closed rectangle  $Q$  compactly contained in its positive set, a number  $r > 0$ , and an open neighborhood  $\mathcal{V}$  of  $\tilde{f}_0|_Q$  in  $C^{m+3}(Q)$  on which maps*

$$f \mapsto (\Xi(f), \Upsilon(f), \mathcal{S}_*(f)) \in \mathbb{R}^2 \times C^m([-r, r])$$

*are  $C^1$ . They equal  $(\xi, \eta, S)$  for laws of the recorded class whose origins and action data are in the indicated local branch. On a sufficiently small convex ball in  $\mathcal{V}$ ,*

$$|\Xi(f) - \Xi(h)| + |\Upsilon(f) - \Upsilon(h)| + \|\mathcal{S}_*(f) - \mathcal{S}_*(h)\|_{C^m([-r, r])} \leq C_m \|f - h\|_{C^{m+3}(Q)}. \quad (24.9)$$

*The constants are local and depend on positive density, simple-root and anchor margins. For a differentiable recorded family, at its base law,*

$$\dot{\xi} = -\frac{\dot{\mathcal{K}}(\xi, y_*)}{\partial_x \mathcal{K}(\xi, y_*)}, \quad \dot{\eta} = -\frac{\dot{\mathcal{K}}(x_*, \eta)}{\partial_y \mathcal{K}(x_*, \eta)}. \quad (24.10)$$

*In particular a stationary interior density has stationary origins and stationary action at every separately fixed smooth order, even when its separate endpoint weights vary.*

*Proof.* Choose noncentral  $x_*, y_*$  close enough to  $\xi_0, \eta_0$  that a closed rectangle  $Q \Subset D_0$  contains  $(\xi_0, \eta_0)$ , both slice anchors, and a small centered square, with positive room to its boundary. Its density has a positive lower bound. The map

$$f \mapsto \mathcal{K}_f = f_{xy}/f - f_x f_y / f^2$$

is  $C^1$  from the open positive part of  $C^{m+3}(Q)$  into  $C^{m+1}(Q)$ . Evaluating at  $(x, y_*)$  or  $(x_*, y)$  gives a  $C^1$  map in  $f$  and the remaining scalar variable. The nonzero derivatives in (24.5) permit the scalar implicit function theorem. It produces unique roots  $\Xi(f)$  near  $\xi_0$  and  $\Upsilon(f)$  near  $\eta_0$ , and both root maps are  $C^1$ . On the model these roots are the actual origins whenever the latter are in those neighborhoods. The exact global identification remains Lemma 24.1; no global root selection for arbitrary nonmodel perturbations is needed.

Shrink the neighborhood and choose  $r$  so that  $(\Xi(f) + [-r, r]) \times (\Upsilon(f) + [-r, r])$  is inside  $Q$  for every  $f$  under consideration. Translation followed by restriction defines

$$f \mapsto f(\Xi(f) + u, \Upsilon(f) + v)$$

as a  $C^1$  map into  $C^{m+1}([-r, r]^2)$ , with the surplus derivatives in the hypothesis. Its differential in a direction  $h$  is the translated  $h$  plus the translated  $f_x D\Xi[h] + f_y D\Upsilon[h]$ . Choose a fixed  $a \neq 0$  in this square. The base value  $1 - R_0(a, a)$  is strictly positive. On a smaller neighborhood every axis denominator, the scalar anchor root, and  $1 + T$  retain positive lower bounds. Products, quotients, slice evaluations and a scalar square root are  $C^1$  in these finite smooth norms. Composing them gives  $\mathcal{S}_*$  by (24.7).

Continuity of the derivative at  $f_0$  bounds its operator norm on a sufficiently small ball. Integration of that derivative on the segment from  $f$  to  $h$  proves (24.9). This uses a finite smooth observation norm; it is not a total-variation estimate. Differentiating the two root equations on their fixed slices proves (24.10). If  $\dot{f} = 0$  in these smooth norms, then  $\dot{\mathcal{K}} = 0$  and both root derivatives vanish. The translated density is then stationary, and differentiation of the scalar anchor and action formulas makes  $\dot{S} = 0$ . Unknown derivatives of  $A, C, Z$  have disappeared through the interaction and ratio identities, not through an assumption that they were known or constant.  $\square$

For a small nonmodel perturbation the extracted action need not have zero first derivative. Subtracting its constant and linear terms fixes that anchoring and changes no coefficient of degree at least two. With a pair of recorded laws and the gap, the finite smooth inverse of Proposition 23.2 therefore gives, for each fixed contact order  $M$ , a local bound for the recovered graph coefficients by the gap error plus the two  $C^{M+3}$  density errors in (24.9). Its constants also depend on the inherited positive geometric bounds. No infinite-order condition bound or uniform estimate near a vanishing recording efficiency follows from this statement.

**24.4. The complete geometric inverse with these nuisances.** For a selected channel  $e$  and each same-type law  $b = 0, 1$ , let  $\tilde{P}_{e,b}$  be its complete translated and separately weighted law as in (24.1). The gaps  $g_e$ , positive offsets  $d_e$ , obstacle and deck marks, and signed transverse length conventions are retained. The two origins and both recording profiles in each law are unknown, and may differ between endpoint types and channels.

**Theorem 24.4** (Geometric determination from uncentered weighted laws). *These recorded laws and the gaps determine every finite labelled contact jet of smooth strictly convex channels. For analytic noncircular tables with a selected skeleton as in Theorem A, the set of compatible tables, after allowing the stated nuisance origins and positive separate recording profiles, is exactly the finite table fiber of the original ideal-law datum. In particular the complete marked lattice is recovered on every branch, and the fiber is a singleton on  $\mathcal{A}_N$ .*

*For a  $C^1$  common-strip analytic-support table family with a persistent selected design and fixed positive offsets, allow also  $C^1$  origins and positive recording profiles with  $C^1$  dependence in each finite smooth norm on the common recorded collars. If the gaps and all recorded interior densities are stationary, the table variation is a common infinitesimal proper Euclidean motion. Conversely, every such table variation admits stationary recorded laws by keeping its intrinsic origins and recording profiles fixed. Thus the projection of the joint data-derivative kernel onto table variations is exactly the common Euclidean tangent space.*

*Proof.* Lemma 24.1 identifies the origins in each record, and Theorem 24.2 recovers both signed action germs. The gap and their quadratic coefficients give

$$\gamma = \operatorname{arsinh}(g\sqrt{a_0 a_1}), \quad c_b = \cosh \gamma \sqrt{a_b/a_{1-b}}, \quad \kappa_b = (c_b - 1)/g, \quad a_b = S_b''(0).$$

The finite-order identities of Lemma 22.4 and Theorem 22.1 then give  $q_n = M_n^{-1}(s_n - R_n)$  successively. The smooth remainder result is still necessary; cancellation of recording factors does not replace it. Every finite signed jet, including its odd terms, is recovered.

For analytic boundaries the resulting contact germs determine the complete ordered channel-frame images. Finite matching and the displacement cochain of Theorem 30.3 enumerate the compatible tables and recover  $L = (v_1 \ v_2)M^{-1}$  on each accepted branch. Every realization of the recorded laws yields these same jets and images, and so belongs to that list. Conversely, each accepted branch places every recovered ordered channel pair by a proper motion. Its local half-line actions and normalized physical amplitudes therefore agree with the original ones; any global phase-volume constant cancels in the conditional law. Reusing the original numerical origins and endpoint recording profiles realizes the complete recorded laws on that branch. This proves both inclusions of the table fibers, not merely finiteness of a list containing them. Asymmetric uniqueness follows from the same matching theorem.

For the differential assertion, work on common positive rectangles about the base contacts. The inherited forward differentiability and the assumed regularity of the profiles give the finite smooth dependence required in Proposition 24.3. The normalizer is differentiable by the positive-part calculation in Lemma 31.1; translations can first be removed in that integral, so the same regular level set applies. Stationarity of the recorded densities gives  $\dot{\xi} = \dot{\eta} = 0$  and  $\dot{S}_b = 0$ . Differentiate the displayed leading recovery and then

$$\dot{q}_n = M_n^{-1}(\dot{s}_n - \dot{R}_n - \dot{M}_n q_n).$$

Finite induction makes every graph-jet variation zero. Analyticity of the support variation on a smaller common strip and the identity theorem give stationary complete channel-frame images, by Lemma 31.3. The actual local congruence branch of Lemma 30.2 makes their relative angular and translation velocities zero. The same finite cochain then makes the lattice derivative zero in one common gauge. Undoing that gauge yields exactly the common infinitesimal Euclidean motion. This argument never infers an injective derivative from exact injectivity.

For the converse choose a common Euclidean motion of the table and keep all nuisance profiles fixed as functions of their intrinsic numerical coordinates. Gaps and the ideal laws are invariant, hence the recorded laws are stationary. This is an existence assertion about a joint kernel vector. Arbitrary changes of the recording profiles need not preserve a law even when the table itself is stationary.  $\square$

**24.5. Relation to the other observations.** The extraction recovers numerical origins, not an unknown clock, channel direction, transverse scale or deck marking. The axes remain the specified oriented transverse axes of each channel; this is not an arbitrary fixed laboratory projection while the channel direction varies. It requires the positive cap not to be cropped before either origin is visible. An inseparable joint recording factor  $w(x, y)$  adds  $\partial_x \partial_y \log w$  to (24.3) and is a different model. Similarly, an arbitrary nonlinear coordinate change alters the physical action; it is not an unobserved translation.

Theorem 24.4 concerns exact laws, with the local finite smooth bound (24.9). It is not a parameter-independent Markov kernel transforming individual uncentered samples into ideal samples, a uniform finite-histogram theorem, or a semiparametric efficiency assertion. The unknown recording factors may drastically reduce the chance of accepting a successful bridge. The preparation budgets and density margins in Part III cannot be transferred uniformly without additional quantitative bounds on those factors. The physical pilot there also determines quantities other than transverse origins, and its proved sensor and clock assumptions remain in force. In this way the same boundary law has distinct roles: its interaction determines geometry, while its amplitude and retained auxiliary records govern acquisition and local statistical information.

**24.6. What the complete support already determines.** There are two different reasons for using an interior interaction. It eliminates separate recording factors, and its local regularity does not require control of a support boundary. The second point should not be confused with

nonidentifiability of the origins from a complete cap. The following observation, communicated in [42], identifies exactly what that support supplies.

**Lemma 24.5** (Centering from complete fiber widths). *In the setting of (24.1), the positive set  $D$  alone determines both origins. Its vertical fiber length has a unique maximum at  $x = \xi$ , and its horizontal fiber length has a unique maximum at  $y = \eta$ . Both maxima are nondegenerate. No symmetry of  $S$  is required.*

*Proof.* For  $0 < q \leq d$ , strict convexity and the compact-sublevel assumption give exactly two inverse branches  $v_-(q) < 0 < v_+(q)$  of  $S(v) = q$ , lying in  $I$ . The inverse function theorem applies to both branches, including in a neighborhood of  $q = d$ , since their derivatives  $S'(v_\pm(q))$  are nonzero and the sublevel has positive room inside  $I$ . Put

$$W(q) = v_+(q) - v_-(q), \quad W'(q) = \frac{1}{S'(v_+(q))} - \frac{1}{S'(v_-(q))} > 0.$$

For every  $x$  with a nonempty vertical fiber, its length is

$$w(x) = W(d - S(x - \xi)). \quad (24.11)$$

This expression has its unique maximum at  $\xi$ , since  $S$  has its unique minimum at zero. Differentiation in a neighborhood of that maximum gives

$$w'(\xi) = 0, \quad w''(\xi) = -W'(d)S''(0) < 0.$$

The horizontal calculation is identical, with  $\eta$  in place of  $\xi$ . Strict positivity of the recording factors makes  $D$  the interior of the support of the recorded law, so these are observable fiber lengths. The width need not be symmetric about its maximum.  $\square$

This comparison concerns centering, not recovery of the entire action from one support: Proposition 23.13 preserves a whole cap while changing its action germ. It is also not a replacement for Proposition 24.3, whose observation norm controls only a fixed rectangle strictly inside the positive set. The support boundary is outside that norm. The four-density cancellation, including unequal separate factors, is already the algebra of Theorem 23.1; it is not an additional independent high-order inverse. The geometric work is done by the relative law and the actual-smooth contact recursion to which the invariant is applied.

**24.7. Recording only a positive interior window.** The extraction can be formulated without an observed cap boundary. Let  $J_x, J_y$  be bounded open intervals with known endpoints and let  $\mathcal{W} = J_x \times J_y$  be a known recording window. The numerical origins  $\xi \in J_x$  and  $\eta \in J_y$  are unknown. On neighborhoods of the translated closed intervals let  $S$  be smooth with  $S(0) = S'(0) = 0$  and  $S'' > 0$ , as above. Assume

$$H(x, y) = d - S(x - \xi) - S(y - \eta) > 0 \quad \text{on } \overline{\mathcal{W}}. \quad (24.12)$$

Only records in  $\mathcal{W}$  are retained. Their law has density

$$p_{\mathcal{W}}(x, y) = \frac{A(x)C(y)}{Z_{\mathcal{W}}} H(x, y) \mathbf{1}_{\mathcal{W}}(x, y), \quad Z_{\mathcal{W}} = \int_{\mathcal{W}} A(x)C(y)H(x, y) \, dx \, dy, \quad (24.13)$$

where  $A, C$  are positive and smooth near the respective closed intervals. They need not agree. Neither the density nor its support outside  $\mathcal{W}$ , nor the discarded mass, is part of this datum. In particular its positive support is just the fixed recording window.

The visibility assumption  $\xi \in J_x, \eta \in J_y$  retains both critical coordinate lines. It specifies a class of observations, not an algorithm for locating a suitable window without prior information. The known signed transverse units, directions and excess time remain those of Section 24. We do not replace them by unknown scales or arbitrary coordinate transformations.



**Theorem 24.6** (Interior-window extraction). *The law (24.13) determines  $\xi$ ,  $\eta$  and the signed germ of  $S$ . A complete cap and an observation of its boundary are not needed. The two effective factor ratios in (24.8) are determined on a small interval about the recovered origins.*

*For each fixed integer  $m \geq 0$ , the extraction has a local  $C^1$  extension from an open neighborhood in  $C^{m+3}(Q)$  to  $\mathbb{R}^2 \times C^m([-r, r])$ , where  $Q \Subset \mathcal{W}$  contains the two slice anchors and the translated square. On a sufficiently small convex neighborhood it satisfies (24.9). The constants depend on the chosen positive density, root, anchor and window margins. They do not require bounds on a density outside  $Q$  or on a discarded cap boundary.*

*Proof.* On the open window the mixed logarithmic derivative is

$$\mathcal{K}_p(x, y) = -\frac{S'(x - \xi)S'(y - \eta)}{H(x, y)^2}.$$

Every vertical fiber is the nonempty interval  $J_y$ . Since  $S'$  has only one zero and  $J_y$  has positive length, exactly the vertical line  $x = \xi$  has identically zero interaction; similarly the horizontal line is  $y = \eta$ . This recovers the origins using no boundary value. The smooth positive representative of a density on the window is uniquely determined by its almost-everywhere class.

Choose a small  $r > 0$  so that  $(\xi + [-r, r]) \times (\eta + [-r, r]) \Subset \mathcal{W}$  and  $\max_{[-r, r]} S < d/4$ . On that square put  $p^c(u, v) = p_{\mathcal{W}}(\xi + u, \eta + v)$  and choose one fixed  $a \in (-r, r) \setminus \{0\}$ . The window normalization is a positive scalar, and hence cancels together with both separate factors:

$$R_p(u, v) = \frac{p^c(u, v)p^c(0, 0)}{p^c(u, 0)p^c(0, v)}, \quad 1 - R_p(u, v) = \frac{S(u)}{d - S(u)} \frac{S(v)}{d - S(v)}.$$

The scalar  $q = \sqrt{1 - R_p(a, a)}$  is positive and equals  $S(a)/(d - S(a))$ . Thus  $T(u) = (1 - R_p(u, a))/q$  and  $S(u) = dT(u)/(1 + T(u))$ , for both signs of  $u$ . Substitution on the axes gives the asserted effective-factor ratios. Nothing in these identities involves  $Z_{\mathcal{W}}$  or the unobserved mass separately.

For the local regularity assertion choose noncentral slices  $x_*, y_*$  inside the window and then a closed  $Q \Subset \mathcal{W}$  containing these slices near their roots and a small centered square. Compactness and (24.12) give a positive density minimum there. The two root slopes are exactly (24.5) and its horizontal counterpart, so they are nonzero. The proof of Proposition 24.3 is local to this  $Q$ : the map  $f \mapsto f_{xy}/f - f_x f_y/f^2$  is  $C^1$  from positive  $C^{m+3}(Q)$  to  $C^{m+1}(Q)$ ; two scalar implicit function theorems give the origin maps; and restriction after their translations, followed by the slice products and scalar anchor inverse, gives a  $C^1$  action map into  $C^m([-r, r])$ . The derivative of the translation contains both  $f_x D\Xi$  and  $f_y D\Upsilon$ , and the scalar anchor is also differentiated. None of these operations evaluates the edge of  $\mathcal{W}$ . A smaller convex ball has bounded derivative, which gives the stated estimate by integration on line segments. This argument applies to arbitrary nearby positive nonmodel inputs as a local extension; exact identification is claimed on the model class.  $\square$

The fixed support of (24.13) cannot perform the centering in Lemma 24.5. For example, begin with an actual channel law and a window strictly inside its cap, with both origins strictly inside the coordinate intervals. Small changes of the numerical origins preserve these strict conditions. The cropped laws then have the same positive support  $\mathcal{W}$  and different origins, while their interactions locate those origins by the theorem. This is a change of the recorded numerical origins of an existing channel, not a claim that an arbitrary functional action is realized by a billiard. The window theorem records the boundary-independent content of the existing interior proof; it does not introduce a new association invariant.

#### 24.8. The geometric fiber after interior cropping.

**Corollary 24.7** (The same table fiber from interior-window laws). *For every selected channel and both same-type laws, retain a known window satisfying (24.12) and containing the respective unknown origins, together with the gaps, marks, signed units, directions and fixed positive offsets. Allow independent positive smooth separate recording factors for every law. Then these window laws determine every finite labelled contact jet of smooth strictly convex channels. On the analytic noncircular class of Theorem A, their compatible table fiber is exactly the ideal-law fiber, including its unknown marked lattices. It is therefore finite, and is a singleton on  $\mathcal{A}_N$ .*

*In the differentiable-family setting of Theorem 24.4, keep the numerical windows fixed and the visibility and positivity conditions locally strict. The projection of the joint window-data derivative kernel onto table variations is exactly the common infinitesimal proper Euclidean tangent space.*

*Proof.* Apply Theorem 24.6 to each record. The gap and the two recovered quadratic action coefficients give the leading geometry by the formulas in Theorem 24.4. The finite-smooth factorization of Lemma 22.4 then makes each higher coefficient an actual graph-jet invariant before the triangular inverse is applied. Thus all finite signed contact jets are fixed. For analytic boundaries these jets give complete ordered channel-frame images by analytic continuation. Every table realizing the window laws therefore gives one of the compatible incidence assignments in Theorem 30.3, with its lattice  $L = (v_1 \ v_2)M^{-1}$ . This is one inclusion in the claimed fiber.

Conversely, each accepted ideal branch places every recovered ordered channel pair by a proper Euclidean motion. Its local physical actions and amplitudes agree with those of the original pair, in the same intrinsic transverse coordinates. Reuse the original numerical origins, windows and recording factors on each such pair. The resulting normalized window density is exactly (24.13). A global phase-volume factor cancels under conditioning. This gives the reverse inclusion; no unobserved portion of a recorded density is being assumed as extra data. The uniqueness conclusion follows on the original properly asymmetric subclass, without altering that class.

For a family, the normalizing integral on the fixed window is particularly simple. In the notation of (24.13), with  $w_t = A_t C_t$ ,

$$\dot{Z}_W = \int_W (\dot{w}H + w\dot{H}), \quad \dot{H} = -\dot{S}(x - \xi) - \dot{S}(y - \eta) + S'(x - \xi)\dot{\xi} + S'(y - \eta)\dot{\eta}.$$

Local smooth bounds on the compact window justify differentiation under the integral and give continuity of the derivative. There is no moving integration boundary. Together with the inherited forward estimates this gives the finite smooth density dependence needed above. Stationary window densities have stationary origins and stationary action jets. Differentiating the leading recovery and then  $q_n = M_n^{-1}(s_n - R_n)$  at each finite order gives zero graph-jet variation, retaining the derivatives of both  $R_n$  and  $M_n$ . The variation of a common-strip analytic-support family is analytic on a smaller strip, so it vanishes on each complete channel-frame image. The actual continuous alignment branches and the displacement cochain then make the table variation zero in one common Euclidean gauge, exactly as in Theorem 24.4. Undoing the gauge gives the asserted tangent space. Conversely, a common proper Euclidean motion with fixed intrinsic numerical origins and recording profiles leaves every window law stationary. This is a statement about the projection of a joint kernel, not identification of all profile variations.  $\square$

The effective factors in this window model are compatible with separate recording probabilities. On the two compact coordinate intervals the physical amplitudes are positive. Extend the logarithms of the smooth positive ratios  $A/B$  and  $C/B$  to the physical collars by smooth cutoffs, and exponentiate. The extensions stay positive and bounded; multiplying each by a sufficiently small positive constant makes it at most one. The two constants cancel in the normalized window law. The additional known crop is itself a product of the two interval indicators.

For fixed recording factors these window laws are also limits of the corresponding recorded finite-flight laws. The ideal endpoint laws converge weakly by Theorem 17.3. Multiplication by the bounded recording weight and  $\mathbf{1}_{\mathcal{W}}$  preserves the limit of integrals against bounded continuous tests: its discontinuities are confined to the four edges, which have zero mass for the absolutely continuous limiting law. The limiting denominator is positive, so normalizing gives (24.13). This is a fixed-observation limit, not a uniform bound over efficiencies or windows with vanishing accepted mass.

In particular, neither interior-window rigidity nor complete-cap rigidity transfers the preparation budgets of Part III to an unconstrained recording class. Multiplying both recording probabilities by  $\varepsilon$  leaves the normalized law unchanged and multiplies the recorded-success probability by  $\varepsilon^2$ . A charged uniform budget therefore still requires a separate positive success bound. The interior smooth estimate, the exact finite-dimensional coordinates, and the finite-histogram theorem retain their different observation norms and hypotheses.

## 25. INTRINSIC LIFTED-CHANNEL GRAPHS AND ANALYTIC TABLE RIGIDITY

The signed one-channel inverse is formulated in a channel-adapted oriented frame. In a periodic billiard a measured channel also carries a deck translation. A global inverse theorem must therefore reconstruct both the relative placement of distinct channel frames and their placement relative to the periodic lattice.

Let  $\Lambda_{\text{abs}} \simeq \mathbb{Z}^2$  be a marked abstract lattice with its fixed Euclidean Gram form and the orientation induced by its marked basis. An oriented Euclidean realization is an orientation-preserving linear isometry

$$\iota : \Lambda_{\text{abs}} \otimes \mathbb{R} \longrightarrow \mathbb{R}^2.$$

The quotient obstacle orbits are labelled  $1, \dots, N$ . An oriented measured channel is indexed exactly as in Section 12:

$$e = (a_e, b_e, \ell_e), \quad \ell_e \in \Lambda_{\text{abs}},$$

and joins a lift of obstacle  $a_e$  to the lift of obstacle  $b_e$  translated by  $\ell_e$ . Its two incidences are denoted  $(e, -)$  and  $(e, +)$ , with

$$r(e, -) = a_e, \quad \nu(e, -) = 0, \quad r(e, +) = b_e, \quad \nu(e, +) = \ell_e. \quad (25.1)$$

Reversing the channel replaces  $(a, b, \ell)$  by  $(b, a, -\ell)$ .

**Definition 25.1** (Intrinsic signed lifted-channel datum). For an oriented channel  $e = (a, b, \ell)$ , its intrinsic datum consists of

- (1) the lifted channel label  $(a, b, \ell)$  in the marked abstract lattice;
- (2) the physical onset  $g_e$ ;
- (3) the two same-type signed endpoint-law germs

$$\{\Pi_{00,d}^{e,\text{end},\infty} : 0 < d < d_e\}, \quad \{\Pi_{11,d}^{e,\text{end},\infty} : 0 < d < d_e\};$$

- (4) the sign convention for the transverse endpoint coordinate.

The law is recorded in the channel-adapted frame in which the first contact is  $(0, 0)$ , the second is  $(g_e, 0)$ , and the oriented transverse axes are those used by the signed endpoint coordinates. No Euclidean placement of this frame relative to the lattice or to any other measured channel is supplied.

Theorem 22.1 and Lemma 23.14 recover two complete oriented analytic boundary images in the edge frame,

$$C_{e,-}, \quad C_{e,+},$$

representing the lifted boundaries at deck offsets 0 and  $\ell_e$ . The remaining inverse problem is a periodic Euclidean gluing problem.

**25.1. Analytic signatures and local matching.** Let  $C$  be a connected oriented real-analytic strictly convex closed curve, parametrized by oriented arclength  $s \in \mathbb{R}/L\mathbb{Z}$ , with signed curvature  $\kappa_C(s) > 0$ . At  $p = C(s_0)$  define

$$\mathbf{j}_C(p) = (\kappa_C(s_0), \kappa'_C(s_0), \kappa''_C(s_0), \dots).$$

Let  $\text{Sym}^+(C) \subset \text{SE}(2)$  be the orientation-preserving Euclidean symmetry group of the boundary image.

**Lemma 25.2** (Coincident analytic signatures are exactly symmetry orbits). *For framed points  $p, q \in C$ ,*

$$\mathbf{j}_C(p) = \mathbf{j}_C(q)$$

*if and only if an element of  $\text{Sym}^+(C)$  carries the oriented framed point  $p$  to the oriented framed point  $q$ .*

*Proof.* The reverse implication is immediate. For the forward implication choose arclength so that  $p = C(0)$  and  $q = C(a)$ . Equality of the complete curvature jets gives equality of the real-analytic germs  $\kappa_C(s)$  and  $\kappa_C(s+a)$ , hence equality for every  $s$  by the identity theorem. Let  $A \in \text{SE}(2)$  send the oriented point-tangent frame at 0 to that at  $a$ . Uniqueness for the planar Frenet system gives  $AC(s) = C(s+a)$  for all  $s$ , so  $A \in \text{Sym}^+(C)$ .  $\square$

Thus a circle has a genuine continuous phase ambiguity and a cyclically symmetric analytic obstacle may have a finite one. These are actual inverse ambiguities, not registration conventions.

**25.2. Admissible periodic gluings.** Let  $\tau_v(x) = x + v$ . An *algebraic periodic gluing* of the intrinsic data is a pair

$$(\iota, (A_e)_{e \in E}), \quad A_e \in \text{SE}(2), \quad (25.2)$$

where  $\iota$  is an oriented Euclidean realization of  $\Lambda_{\text{abs}}$ , such that whenever two incidences  $(e, \sigma), (f, \sigma')$  carry the same quotient obstacle label,

$$\tau_{-\iota(\nu(e, \sigma))} A_e C_{e, \sigma} = \tau_{-\iota(\nu(f, \sigma'))} A_f C_{f, \sigma'} \quad (25.3)$$

as oriented labelled boundary images. Thus every recovered lift is translated back by its deck vector before copies of the same quotient obstacle are identified.

The common values in (25.3) define quotient obstacle images  $C_1, \dots, C_N$ . Extend them by all deck translations  $\iota(\Lambda_{\text{abs}})$ . We call the gluing *admissible* if this periodic family is a dispersing billiard table: distinct translated obstacle closures are disjoint and every measured edge has the prescribed facing contacts, onset and selected closest-pair channel geometry. These are conditions on the reconstructed curves and their candidate placements.

The group  $\text{SE}(2)$  acts by

$$B = (R, t) : \quad \iota \mapsto R\iota, \quad A_e \mapsto B \circ A_e.$$

Both (25.3) and admissibility are invariant. Let  $\mathfrak{G}_{\text{per}}(\mathcal{D})$  be the quotient of the admissible periodic gluings by this action.

**Theorem 25.3** (Intrinsic periodic reconstruction and holonomy classification). *Let two labelled periodic dispersing billiard tables have the same marked abstract Euclidean lattice, the same obstacle-orbit labels, and the same connected measured lifted-channel graph. Assume every obstacle boundary is connected, real analytic and strictly convex, and that the two tables have identical intrinsic signed lifted-channel data for every measured edge. Then:*

- (1) every incident lifted obstacle boundary is recovered as a complete analytic image in its own channel frame;
- (2) the observations determine the admissible periodic gluing space  $\mathfrak{G}_{\text{per}}(\mathcal{D})$  without a common laboratory registration;
- (3) labelled periodic dispersing tables realizing the observations, modulo one simultaneous orientation-preserving Euclidean motion of the lattice and all obstacle images, are in bijection with  $\mathfrak{G}_{\text{per}}(\mathcal{D})$ .

In particular, if  $\mathfrak{G}_{\text{per}}(\mathcal{D})$  is a singleton, the complete labelled periodic table is intrinsically determined up to one global  $\text{SE}(2)$  motion.

For the common-orientation orbit of the entire marked datum, as defined in Definition 25.12, the corresponding classification is modulo  $\text{E}(2)$ , by Theorem 25.14. This assertion concerns one common orientation choice, not pointwise deletion of endpoint signs.

*Proof.* The first assertion follows from Theorem 22.1 and Lemma 23.14. Hence every curve entering (25.3) is determined by the observations.

A table realization supplies its lattice realization  $\iota$  and the actual placement  $A_e$  of every channel frame. The incidence  $(e, \sigma)$  lies on the lift  $D_{r(e, \sigma)} + \iota(\nu(e, \sigma))$ , so translating it back by the deck vector gives the same quotient obstacle image at every incidence with that label. The resulting gluing is admissible, and simultaneous global Euclidean motion changes only its quotient representative.

Conversely an admissible gluing defines

$$C_r = \tau_{-\iota(\nu(e, \sigma))} A_e C_{e, \sigma}$$

for any incidence with  $r(e, \sigma) = r$ . The gluing constraint makes this independent of the choice. The periodic lifts are  $C_r + \iota(\lambda)$ ,  $\lambda \in \Lambda_{\text{abs}}$ . For a measured edge  $e = (a, b, \ell)$ , its placed endpoint curves are exactly  $C_a$  and  $C_b + \iota(\ell)$ , with the observed onset and oriented signed contact convention. Admissibility makes the full periodic family a dispersing billiard table with the declared measured channels. These constructions are inverse modulo the common  $\text{SE}(2)$  action.  $\square$

**25.3. Cycle holonomy and an intrinsic uniqueness criterion.** If two incidences of the same obstacle orbit have contact signatures occurring at unique framed points, Lemma 25.2 gives a unique orientation-preserving Euclidean congruence between their recovered curve copies. We call such an incidence transition *signature-rigid*.

Consider a closed oriented channel cycle

$$c = e_1^{\epsilon_1} \cdots e_m^{\epsilon_m}, \quad \epsilon_i \in \{\pm 1\},$$

for which every required incidence transition is signature-rigid. Its marked deck displacement is

$$\eta_c = \sum_{i=1}^m \epsilon_i \ell_{e_i} \in \Lambda_{\text{abs}}. \quad (25.4)$$

Compose the unique data-determined incidence congruences around the cycle in the starting channel frame and denote the resulting Euclidean motion by  $H_c$ .

**Lemma 25.4** (Periodic cycle-holonomy identity). *For realizable intrinsic data,  $H_c$  has identity rotational part. Hence  $H_c = \tau_{v_c}$  for a data-determined vector  $v_c$ . For every admissible periodic gluing and every placement  $A_{e_0}(x) = R_{e_0}x + t_{e_0}$  of the starting channel frame,*

$$R_{e_0} v_c = \iota(\eta_c). \quad (25.5)$$

Equivalently, after using the global gauge to set the starting frame rotation to the identity,  $\iota(\eta_c) = v_c$ .

*Proof.* For two consecutive incidences of the same quotient obstacle, the gluing equation identifies their placed curve copies after the corresponding deck translations. The unique signature-rigid congruence between the two recovered edge-frame copies is therefore conjugate, by their placements, to translation by the difference of their deck offsets. Compose these identities around the cycle. All intermediate edge placements telescope. The signed deck differences sum to  $\eta_c$ , giving

$$A_{e_0} H_c A_{e_0}^{-1} = \tau_{\iota(\eta_c)}.$$

Conjugation by  $A_{e_0}$  preserves the rotational part, so  $H_c$  is a translation. Comparing translation vectors gives (25.5).  $\square$

**Definition 25.5** (Lattice-anchoring cycle). A signature-rigid cycle is *lattice-anchoring* if  $\eta_c \neq 0$ . Because the marked basis fixes the orientation and Gram form of the abstract lattice, the image of one nonzero vector determines the orientation-preserving realization  $\iota$  relative to the starting channel frame through Lemma 25.4.

**Definition 25.6** (Rooted signature-rigid spanning tree). Fix a quotient obstacle orbit  $r_*$ . A spanning tree rooted at  $r_*$  is *signature-rigid* if its edges can be ordered so that each edge joins an already reached obstacle orbit to one new orbit, and its contact on the already reached orbit has a curvature signature occurring at a unique framed point of that obstacle. For a one-vertex quotient graph the empty tree is allowed.

**Corollary 25.7** (Fully intrinsic periodic table rigidity). *Under the hypotheses of Theorem 25.3, assume the measured network contains a lattice-anchoring cycle and, rooted at an obstacle orbit met by that cycle, a signature-rigid spanning tree. Then  $\mathfrak{G}_{\text{per}}(\mathcal{D})$  is a singleton. Hence the intrinsic signed lifted-channel data determine the complete labelled periodic billiard table up to one simultaneous global element of  $\text{SE}(2)$ . No absolute contact point, absolute tangent frame, inter-channel placement, or lattice orientation relative to a channel frame is supplied in advance.*

*Proof.* Use the global Euclidean gauge to place the starting frame of the anchoring cycle. The cycle places at least one recovered obstacle lift and therefore a root obstacle image. Since  $\eta_c \neq 0$ , Lemma 25.4 fixes the orientation-preserving lattice realization  $\iota$  relative to that frame. Traverse the rooted spanning tree. At each edge, the unique signature on the already placed obstacle fixes the edge-frame transform, and the known deck correction  $\iota(\ell_e)$  fixes the new quotient obstacle image. Induction places all obstacle orbits. Thus there is at most one algebraic gluing modulo the initial global gauge. Since the hypotheses concern data realized by a dispersing table, that unique gluing is admissible. Hence the admissible gluing space is a singleton.  $\square$

**Corollary 25.8** (One-root registration is sufficient once the lattice is known). *Suppose an oriented Euclidean realization of the periodic lattice and the Euclidean placement of one root channel frame relative to that lattice are supplied as ambient calibration. Then a rooted signature-rigid spanning tree determines every obstacle image and every remaining measured channel frame. Thus one root registration suffices; a common laboratory registration of all channels is not required.*

**Remark 25.9** (Symmetry, connectivity, holonomy and admissibility). The sufficient conditions above are not necessary. Finite obstacle symmetries produce finitely many possible incidence matches, and several cycles can eliminate those choices. The exact object controlling uniqueness is the admissible periodic gluing space  $\mathfrak{G}_{\text{per}}(\mathcal{D})$ . It separates four genuine issues: obstacle symmetry, channel-network connectivity, periodic lattice holonomy and global geometric admissibility.

**Corollary 25.10** (Common laboratory registration as a special case). *If the Euclidean placement of the lattice and of every measured channel frame is supplied externally, then the periodic gluing*

is fixed a priori. A family of measured channels meeting every obstacle orbit therefore determines the whole analytic table in that registration, recovering the registered v22 statement as a special case.

**25.4. Transporting the root to an anchoring cycle.** The root of a signature-rigid spanning tree need not be prescribed at the starting frame of an anchoring cycle. The next lemma shows that these two choices are compatible under the existing signature hypotheses. Thus the propagation of obstacle placements requires no additional registration or root-incidence assumption.

**Lemma 25.11** (Rerooting a signature-rigid spanning tree). *Let a rooted signature-rigid spanning tree join a finite collection of connected real-analytic strictly convex obstacle boundaries. Suppose one vertex  $r$  has a framed boundary point at which the complete oriented curvature signature occurs uniquely. Then the same undirected tree, rooted instead at  $r$ , is signature-rigid. In particular, an obstacle orbit met by a signature-rigid anchoring cycle may be chosen as the root.*

*Proof.* First, uniqueness of the complete signature at one framed point of a boundary  $C$  implies

$$\text{Sym}^+(C) = \{\text{id}\}.$$

Indeed, every symmetry carries that framed point to one with the same signature, by Lemma 25.2. Uniqueness makes it fix the point and its oriented unit tangent. An orientation-preserving Euclidean motion fixing such a frame is the identity. The same lemma then shows that the complete signature occurs uniquely at every framed point of  $C$ . Thus uniqueness is a property of this boundary, rather than of a specially chosen incidence on it.

Let  $r_0$  be the original root. Every vertex which has a child in the original rooted tree has a unique-signature contact on an outgoing edge, by Definition 25.6. The preceding argument makes all its framed boundary signatures unique. Write the unique tree path from  $r_0$  to  $r$  as

$$r_0 = v_0, v_1, \dots, v_k = r.$$

Rerooting reverses precisely the edges of this path. On the reversed edge from  $v_i$  to  $v_{i-1}$ , with  $1 \leq i < k$ , the new parent  $v_i$  was the old parent of  $v_{i+1}$  and therefore has unique signatures at every framed point. For  $i = k$ , the same conclusion follows from the hypothesis on  $r$ . Every reversed edge consequently has a unique-signature contact on its new parent. Off the path, the parent-child relations are unchanged and retain their original property. This proves the assertion; when  $k = 0$ , including the one-vertex empty tree, nothing is changed.

Every required incidence transition on a signature-rigid cycle has a unique framed signature by its definition. A vertex met by that cycle therefore satisfies the hypothesis on  $r$ , proving the last assertion.  $\square$

For the rank-two recovery of Theorem 26.4, fix the common starting channel frame of the anchoring pair and choose a root obstacle met there. Lemma 25.11 transports any supplied rooted signature-rigid spanning tree to that obstacle. The two cycle holonomies determine the lattice in this frame, and induction along the rerooted tree then places every obstacle orbit. The same argument applies to the single-offset reconstruction in Theorem 23.4. The tree is only reoriented; no edge, obstacle label, deck displacement, or observation is added or removed.

**25.5. The common-orientation quotient of the marked datum.** There are two different operations on signed observations: changing one orientation convention for the entire experiment, and discarding signs from individual observations. Only the former is involved in passing from oriented to unoriented Euclidean classification. We define it at the level of the complete law-valued datum.

Let  $J(x, y) = (x, -y)$  and  $r(u, v) = (-u, -v)$ . For a signed channel-law collection  $\mathcal{D}$ , write  $\mathcal{R}_0\mathcal{D}$  for the collection obtained by replacing, for every measured edge, both contact types and every retained offset,

$$\Pi_{bb,d}^{e,\text{end},\infty} \quad \text{by} \quad r_*\Pi_{bb,d}^{e,\text{end},\infty}. \quad (25.6)$$

The same involution is used throughout the collection. The obstacle labels, the directed edge labels, the marked deck elements and the onset gaps are unchanged. In particular this operation does not reverse the directed channel  $(a, b, \ell)$ . For a density it reads  $f_{e,b,d}^r(u, v) = f_{e,b,d}(-u, -v)$ . When residual time is present its coordinate is left fixed, and when a full transcript is present its designs, counts, failure symbols and stopping rule are transported together with the same change of endpoint coordinates.

The orientation of a marked lattice requires some care. In the preceding signed formulation  $\iota$  is required to preserve the orientation of the marked basis. A reflection of the physical table does not preserve this sector. To describe the common reflection action, let  $\epsilon \in \{+1, -1\}$  denote the orientation character of  $\iota$ . Write  $\mathcal{D}_\epsilon = (\epsilon, \mathcal{D})$ , and let  $\mathfrak{G}_{\text{per}}^\epsilon(\mathcal{D})$  be the gluing space defined by (25.2)–(25.3), except that the isometry  $\iota$  has orientation character  $\epsilon$ . The placements  $A_e$  are still proper Euclidean motions of the consistently oriented channel frames. The original signed statement is the  $\epsilon = +1$  sector. The proof of its classification applies in either sector: it uses the incidence equations and Euclidean isometry of  $\iota$ , not the sign of its determinant.

The character  $\epsilon$  records the common orientation convention relative to the marked lattice; it is not a supplied lattice angle or a laboratory registration. In an unoriented representative it must be transported with the transverse convention, rather than erased independently.

**Definition 25.12** (Common-orientation datum). On the complete marked signed data define

$$\mathcal{R}(\epsilon, \mathcal{D}) = (-\epsilon, \mathcal{R}_0\mathcal{D}), \quad \mathcal{R}^2 = \text{id}. \quad (25.7)$$

The common-orientation datum is the two-element orbit

$$\overline{\mathcal{D}} = \{(\epsilon, \mathcal{D}), (-\epsilon, \mathcal{R}_0\mathcal{D})\}.$$

Thus it retains all relative sign choices, including their relation to the marked lattice convention, but not a choice of the one common orientation. It is an orbit of the entire collection of probability laws, not a mixture of those laws and not a pushforward under pointwise absolute values.

**Lemma 25.13** (Equivariance of recovery and gluing). *Reflection induces a well-defined involution between the two signed gluing spaces belonging to  $\overline{\mathcal{D}}$ . On representatives it is*

$$(\iota, (A_e), (C_{e,\sigma})) \longmapsto (J\iota, (JA_eJ), (JC_{e,\sigma})). \quad (25.8)$$

*Reflected boundary images are given their canonical positive boundary orientation. The recovered quotient obstacle images are  $JC_r$ , and all labelled deck translations are transported by the same  $J$ .*

*Proof.* A reflected contact graph is  $\psi_b^r(y) = \psi_b(-y)$ . Reflection preserves flight lengths, the reflection law, arclength and the collision flux measure. Sending every endpoint coordinate to its negative therefore gives (25.6); the gap and residual-time coordinate do not change. Consequently the signed inverse and analytic globalization recover exactly the reflected boundary images. Equivalently the half-line actions and their graph jets satisfy

$$S_b^r(u) = S_b(-u), \quad (q_{b,n})^r = (-1)^n q_{b,n}. \quad (25.9)$$

These identities also follow directly by reflecting each finite stationary bridge and then taking the half-line limit. They do not require convergence of an infinite Taylor series for a smooth graph.



Since  $\det J = -1$ , the realization  $J\iota$  lies in the opposite lattice sector while  $JA_e J \in \text{SE}(2)$ . The identity  $J\tau_v = \tau_{Jv}J$  gives, incidence by incidence,

$$\tau_{-J\iota(\nu)}(JA_e J)(JC_{e,\sigma}) = J(\tau_{-\iota(\nu)}A_e C_{e,\sigma}). \quad (25.10)$$

Thus every gluing equation is preserved. Reflection also preserves disjointness, strict convexity, facing contacts, flight lengths and the selected channel geometry; hence it preserves admissibility. Applying the construction twice gives the original gluing.

Finally a different proper representative  $B \circ A_e$ , with lattice realization  $R_B \iota$ , changes the reflected representative by the single proper motion  $JB$ . Formula (25.8) is therefore well defined after the proper Euclidean quotient.  $\square$

**Theorem 25.14** (Unoriented marked periodic classification). *Under the analytic and measured-network hypotheses of Theorem 25.3, let  $\overline{\mathcal{D}}$  be the common-orientation datum of Definition 25.12. Define*

$$\overline{\mathfrak{G}}_{\text{per}}(\overline{\mathcal{D}}) = \left( \mathfrak{G}_{\text{per}}^\epsilon(\mathcal{D}) \sqcup \mathfrak{G}_{\text{per}}^{-\epsilon}(\mathcal{R}_0 \mathcal{D}) \right) / \langle \mathcal{R} \rangle. \quad (25.11)$$

*The labelled periodic tables realizing this datum, modulo one common element of  $E(2)$  acting on the lattice and every obstacle image, are in bijection with  $\overline{\mathfrak{G}}_{\text{per}}(\overline{\mathcal{D}})$ . In particular, a singleton signed gluing space gives a singleton unoriented gluing space and determines the table up to one common Euclidean isometry. The conclusion applies to a single positive offset at each contact type whenever the signed reconstruction is made using the single-offset inverse.*

*Proof.* Choose either representative of  $\overline{\mathcal{D}}$ . The signed classification, in its indicated lattice sector, identifies proper-motion classes of realizing tables with the corresponding signed gluing space. Lemma 25.13 identifies the classifications for the two representatives by the same reflection of the entire table.

More explicitly, if two realizing tables are related by a proper Euclidean motion, the signed classification identifies their signed gluing classes. If they are related by an improper motion, write that motion as  $BJ$  with  $B \in \text{SE}(2)$ . Reflection first interchanges the two signed data representatives by  $\mathcal{R}$ , and  $B$  then identifies their proper classes. Conversely, equality in (25.11) is represented either by equality in one signed gluing space or by equality after its common reflected identification. The inverse signed classification yields respectively a proper or an improper global Euclidean motion of the realizing tables. This proves both directions of the bijection. Reflection maps the two signed gluing spaces bijectively, so the singleton assertion follows. The argument uses only the signed inverse, not the number of offsets used to recover its actions.  $\square$

**Corollary 25.15** (Holonomy and fixed-order stability under the quotient). *The signature-rigid gluing criteria are invariant under the common reflection. If two independent marked cycle columns form  $M$ , their recovered holonomy columns form  $V$ , and  $L = VM^{-1}$ , then*

$$V^r = JV, \quad L^r = JL, \quad (L^r)^\top L^r = L^\top L. \quad (25.12)$$

*Thus the common-orientation quotient retains the recovered Gram form and scale. Any fixed-order stability estimate for the equivariant signed inverse on its exact law image, valid for the representative pairs used in reflection-invariant orbit metrics, descends with the same constant to those metrics. On arbitrary perturbed densities the corresponding action, amplitude and finite-jet assertion uses the transported-anchor extension of Proposition 23.10 and Corollary 23.11. It does not use a fixed signed anchor in both orientation sectors.*

*Proof.* Canonical positive arclength reverses under reflection. Hence the curvature signature transforms as  $(\kappa, \kappa', \kappa'', \dots) \mapsto (\kappa, -\kappa', \kappa'', \dots)$ . Equality and uniqueness of signature matches are preserved. Every incidence congruence is conjugated by  $J$ , as is its product around a cycle. Since  $J\tau_v J = \tau_{Jv}$ , the holonomy vector becomes  $Jv$ ; marked cycle columns are unchanged.

Formula (25.12) follows from  $J^\top J = I$ . A rooted rigid spanning tree is carried to the same rooted rigid combinatorial tree with all its matches reflected.

For the stability assertion on exact laws, let  $F$  be the equivariant signed inverse between the relevant fixed-order data and reconstruction spaces, with reflection-invariant metrics  $d$  and  $d'$ , and suppose  $d'(F(D), F(E)) \leq C_M d(D, E)$  for every representative pair entering the orbit comparison. If the estimate is formulated sectorwise, use only same-sector representative comparisons; their reflected estimates have the same constant. Define the orbit distances by

$$\bar{d}([D], [E]) = \min_{\eta \in \{0,1\}} d(D, \mathcal{R}^\eta E), \quad \bar{d}'([X], [Y]) = \min_{\eta \in \{0,1\}} d'(X, \mathcal{R}^\eta Y).$$

For a sectorwise metric the minima are restricted to the nonempty set of same-sector comparisons. Choose a minimizing  $\eta$  and use equivariance of  $F$ . Then

$$\bar{d}([F(D)], [F(E)]) \leq d'(F(D), F(\mathcal{R}^\eta E)) \leq C_M \bar{d}([D], [E]).$$

For contact jets and densities on symmetric compact boxes, reflection invariance follows from (25.9) and the usual fixed-order supremum norms. This statement retains the order-dependent constants and hypotheses of the signed estimates; it asserts no order-independent stability for analytic continuation. For perturbed densities the same argument applies to the explicitly equivariant extension in (23.28), on its stated domain. The fixed-anchor formula alone has no such off-model equivariance; see Remark 23.12. No arbitrary perturbed density is thereby declared globally realizable.  $\square$

*Remark 25.16* (Law-valued orientation orbits are not folded records). For a continuously distributed endpoint pair, the pushforward under  $(u, v) \mapsto (|u|, |v|)$  has density, on the open positive quadrant,

$$f_{\text{abs}}(a, b) = f(a, b) + f(-a, b) + f(a, -b) + f(-a, -b). \quad (25.13)$$

The coordinate axes have zero mass in this setting. This sum is not the law-valued orbit  $\{f, f \circ r\}$ . Even pushing each observed pair to its orbit under simultaneous reversal is a sample-space coarsening, not an orbit of the complete probability law. Randomly symmetrizing a law is a further operation, and independently reversing different channel collections gives an action of  $\{\pm 1\}^E$ , rather than the single diagonal group used above.

For example, on  $[-1, 1]^2$ , the positive normalized symmetric densities

$$f_1(u, v) = \frac{1}{4}(1 + \alpha(u + v)), \quad f_2(u, v) = \frac{1}{4}(1 + 2\alpha(u + v)), \quad 0 < \alpha < \frac{1}{4},$$

have the same folded density, identically one on  $(0, 1)^2$ , but their law-valued common-reversal orbits are different: the coefficient of  $u + v$  can only change sign under  $r$ , not double. This example distinguishes the observation maps in a class of densities. It is not a claim that these two densities are realized by physical billiard tables, nor a physical nonidentifiability theorem for folded billiard observations. All rigidity statements above use the specified law-valued common-orientation datum.

## 26. RANK-TWO HOLONOMY AND RECOVERY OF THE PERIODIC LATTICE

The preceding intrinsic reconstruction theorem was stated relative to a marked abstract lattice carrying a prescribed Euclidean Gram form. The holonomy identity in fact contains more information. With two independent deck cycles based in one channel frame, the Euclidean realization of the marked lattice, including its metric and absolute scale, is determined by the channel data. We record that stronger form here and use it in the global statistical theorem below.

Let  $\Lambda_{\text{top}} \simeq \mathbb{Z}^2$  be only an oriented marked free abelian group of rank two. No inner product on  $\Lambda_{\text{top}} \otimes \mathbb{R}$  is supplied. A periodic Euclidean realization is an orientation-preserving linear isomorphism

$$L : \Lambda_{\text{top}} \otimes \mathbb{R} \longrightarrow \mathbb{R}^2.$$

It induces the Gram form

$$G_L(\eta, \zeta) = \langle L\eta, L\zeta \rangle_{\mathbb{R}^2}. \quad (26.1)$$

The lifted channel label  $e = (a, b, \ell)$ , with  $\ell \in \Lambda_{\text{top}}$ , remains part of the marked acquisition protocol. Thus the deck class is known combinatorially, but neither its Euclidean length nor its direction is supplied.

An algebraic gluing is now defined exactly as in Section 25, except that the unknown lattice variable is  $L \in \text{GL}^+(2, \mathbb{R})$ , rather than an isometry from a pre-metrized lattice. The incidence relation is

$$\tau_{-L\nu(e, \sigma)} A_e C_{e, \sigma} = \tau_{-L\nu(f, \sigma')} A_f C_{f, \sigma'} \quad (26.2)$$

whenever the two incidences carry the same quotient-obstacle label. Admissibility again means that the periodic translates form a dispersing billiard and that the declared measured edges are the prescribed facing channels. Simultaneous Euclidean motion acts on the whole realization in the usual way.

**Definition 26.1** (Rank-two anchoring pair). Two signature-rigid closed channel cycles  $c_1, c_2$ , based at the same oriented channel frame, form a *rank-two anchoring pair* if their marked deck displacements

$$\eta_i = \sum_{e \in c_i} \epsilon_e \ell_e \in \Lambda_{\text{top}}, \quad i = 1, 2,$$

are linearly independent in  $\Lambda_{\text{top}} \otimes \mathbb{R}$ .

The incidence congruences are determined entirely by the recovered analytic curve copies. Hence, as in Lemma 25.4, their composition around  $c_i$  is a data-determined Euclidean motion  $H_i = \tau_{v_i}$ . The next proposition is the metric-free version of the cycle-holonomy identity.

**Proposition 26.2** (Metric-free cycle holonomy). *For realizable intrinsic data and every signature-rigid cycle  $c$  based at an edge frame whose placement has rotational part  $R_0$ ,*

$$R_0 v_c = L \eta_c. \quad (26.3)$$

*In particular, no Euclidean structure on the abstract deck group is used in deriving the identity.*

*Proof.* The proof of Lemma 25.4 uses only the deck-addition law and Euclidean translations of the realized table. Replacing  $\iota(\nu)$  there by  $L\nu$  leaves the telescoping argument unchanged:

$$A_{e_0} H_c A_{e_0}^{-1} = \tau_{L\eta_c}.$$

The rotational part of  $H_c$  is therefore the identity, and comparison of translation vectors gives (26.3).  $\square$

**Theorem 26.3** (Recovery of lattice shape, scale and Gram form). *Assume the intrinsic signed channel data contain a rank-two anchoring pair  $c_1, c_2$ . Then the data determine the full Euclidean realization of the marked lattice relative to the common starting channel frame, up to the one global rotation already contained in the global  $\text{SE}(2)$  gauge. More precisely, after fixing that gauge so that the starting-frame rotation is the identity, there is a unique linear map*

$$L : \Lambda_{\text{top}} \otimes \mathbb{R} \longrightarrow \mathbb{R}^2, \quad L\eta_i = v_i \quad (i = 1, 2). \quad (26.4)$$

*Consequently the full marked Gram matrix, including absolute scale, is the data-determined matrix*

$$[G_L]_{\mathcal{B}} = M_{\mathcal{B}}^{-T} V^T V M_{\mathcal{B}}^{-1}, \quad (26.5)$$

*where  $M_{\mathcal{B}}$  has as columns the coordinates of  $\eta_1, \eta_2$  in the marked basis  $\mathcal{B}$ , and  $V = (v_1 \ v_2)$ .*

*Proof.* The vectors  $\eta_1, \eta_2$  form a basis of  $\Lambda_{\text{top}} \otimes \mathbb{R}$ . Equation (26.3) therefore prescribes a unique linear map  $L$  once the common starting-frame rotation has been fixed. Realizability implies that  $L$  is invertible and orientation preserving, so  $v_1, v_2$  are linearly independent with the compatible orientation. Formula (26.5) is simply  $G_L = L^T L$  written in the marked basis. A change of the global Euclidean gauge left-multiplies  $L$  by a rotation and hence leaves  $G_L$  unchanged.  $\square$

**Theorem 26.4** (Uncalibrated intrinsic periodic rigidity). *Let two labelled periodic dispersing billiard tables have the same oriented marked deck group  $\Lambda_{\text{top}} \simeq \mathbb{Z}^2$ , the same obstacle-orbit labels and the same connected measured lifted-channel graph. No lattice Gram form, lattice scale, lattice direction, common channel registration or absolute contact frame is supplied. Assume every obstacle boundary is connected, real analytic and strictly convex, the intrinsic signed channel data agree, the measured graph contains a rank-two anchoring pair, and a rooted signature-rigid spanning tree reaches every obstacle orbit. Then the complete marked periodic table and the Euclidean Gram form of its lattice are determined by the observations up to one simultaneous global element of  $\text{SE}(2)$ .*

*Proof.* The signed one-channel inverse and analytic continuation recover every incident obstacle curve in its own channel frame. Fix the common starting frame of the rank-two anchoring pair, and choose an incident obstacle orbit  $r$ . A signature-rigid incidence on an anchoring cycle gives a framed point of this obstacle with a unique complete oriented curvature signature. Lemma 25.11 therefore roots the supplied spanning tree at  $r$ , without changing an edge, a deck label or the observation map. The recovered image at this root is already placed in the chosen anchoring frame.

Theorem 26.3 fixes the complete lattice realization in this same frame. The rerooted signature-rigid spanning tree then propagates the placements of all quotient obstacle images exactly as in Corollary 25.7, with every deck correction interpreted through that one lattice realization. Thus there is at most one algebraic gluing. Existence is not hidden in the word “admissible”: it is provided by the standing hypothesis that the data are generated by an actual periodic dispersing table. Hence that algebraic gluing is admissible and is the unique realization modulo the initial global Euclidean gauge.  $\square$

*Remark 26.5* (What is marked and what is recovered). The deck label  $\ell$  is a combinatorial lift label in the marked covering and is part of the acquisition design; its Euclidean vector is not. Unknown or corrupted deck-class labels constitute a different inverse problem and are not silently inferred here. The classification of all admissible gluings is a bookkeeping theorem; the rigidity content is the rank-two holonomy recovery plus signature-rigid propagation. Theorem 26.4 therefore removes the metric calibration singled out in the preceding formulation while keeping the measurement protocol explicit.

## 27. FINITE ANALYTIC SIGNATURES AND STABILITY OF INTRINSIC GLUING

Exact analytic uniqueness, finite-signature matching, and a compact inverse modulus are three different assertions. We prove them in that order. An outside-arc gap alone does not prove uniqueness inside the correct arc. The local assertion will instead follow from a finite jet embedding, whose order is selected uniformly from the complete analytic signatures.

Fix the marked measured graph, two signature-rigid anchoring cycles with independent marked deck displacements, and a rooted signature-rigid spanning tree. Let  $\mathfrak{R}$  be a compact class of their realizations, modulo one simultaneous global element of  $\text{SE}(2)$ . The designation *signature-rigid* has its meaning in Section 25: each transition used by these cycles and the tree has a unique matching oriented framed point, not merely a unique point after arbitrarily choosing a symmetry

phase. Obstacles not used for such transitions may have exact symmetries; the classification theorem for general gluing spaces remains unchanged.

Assume uniform positive curvature, gap, separation, diameter, free-area and persistent-channel collar bounds. For some  $\rho, B > 0$ , the oriented arclength curvature functions extend holomorphically to strips of width  $\rho$ , with absolute value at most  $B$ . The class is closed in a fixed root-frame gauge and compact in every fixed labelled  $C^q$  topology. A channel margin  $\chi > 0$  excludes switching to a competing closest pair. The data-determined holonomy vectors satisfy

$$\sigma_{\min}(v_1(T) \ v_2(T)) \geq \sigma_{\text{hol}} > 0. \quad (27.1)$$

The compactness assumption includes the lattice and marked placements; it is not inferred from a bound on the curvature of a single obstacle alone. These are the uniform geometric assumptions used below.

A transition obstacle has trivial orientation-preserving symmetry group. Indeed any such nonidentity symmetry carries each boundary point to a different point with the same complete oriented signature, contradicting the designated unique framed match. A nonidentity rigid motion preserving a compact strictly convex body is a rotation about an interior point and has no boundary fixed point. In particular, transition obstacles are not circles. This observation uses the exact signature-rigidity assumption already required for unique global gluing; it is not a new reflection or genericity assumption.

**27.1. Uniform finite embeddings.** Parametrize each transition obstacle by normalized oriented arclength  $x \in \mathbb{R}/\mathbb{Z}$ , with its origin at a designated incidence. The physical arclength derivative is denoted by  $\partial_s$ . Define

$$J_{T,i}^{[m]}(x) = (\kappa_{T,i}(x), \partial_s \kappa_{T,i}(x), \dots, \partial_s^m \kappa_{T,i}(x)).$$

All derivatives with respect to  $x$  below are of this vector-valued map. Uniform upper and lower perimeter bounds make normalized and physical arclength distances equivalent. Fixed positive weights could be placed on the coordinates without changing any assertion.

**Theorem 27.1** (Uniformly selected finite signature embedding). *There are an integer  $m$ , numbers  $v, r, \sigma > 0$ , and  $U, B_2 < \infty$ , uniform on  $\mathfrak{K}$  and its finite list of transition obstacles, such that, writing  $J = J_{T,i}^{[m]}$ ,*

$$|J'(x)| \geq v, \quad |J'(x)| \leq U, \quad |J''(x)| \leq B_2, \quad (27.2)$$

$$|J(x) - J(y)| \geq (v/2) \text{dist}(x, y), \quad \text{dist}(x, y) \leq r, \quad (27.3)$$

$$|J(x) - J(y)| \geq \sigma, \quad \text{dist}(x, y) \geq r. \quad (27.4)$$

*Thus a sufficiently rich finite signature, not necessarily an initially chosen scalar curvature signature, is a uniform embedding of every transition boundary.*

*Proof.* If every derivative of positive order of an analytic curvature function vanished at one point, it would be constant by the identity theorem. Since transition obstacles are not circles, at each  $(T, i, x)$  some component of the derivative of a finite signature is nonzero. The parameter space of these triples is compact: use normalized arclength and the finite obstacle list. Each nonvanishing condition persists on a neighborhood. A finite subcover gives an order  $m_1$  for which  $|(J^{[m_1]})'|$  is bounded below by a positive constant  $v$ . Holomorphic strip bounds give a uniform second derivative bound for this finite vector. Taylor's integral formula then gives

$$|J^{[m_1]}(y) - J^{[m_1]}(x)| \geq v|y - x| - \frac{1}{2} \|(J^{[m_1]})''\|_{\infty} |y - x|^2 \geq (v/2)|y - x|$$

on a uniform short arc. Choose its radius  $r < 1/4$ .

For pairs at normalized arclength distance at least  $r$ , equality of all signature coordinates would, by Lemma 25.2, give a nonidentity orientation-preserving symmetry. No such symmetry exists on a transition obstacle. Each pair is therefore separated by some finite coordinate. The set of  $(T, i, x, y)$  with distance at least  $r$  is compact. A second finite subcover supplies an order  $m \geq m_1$  and a positive minimum separation  $\sigma$ . Adding coordinates preserves the local lower bound already proved for the  $m_1$ -block. The strip bounds supply finite uniform  $U, B_2$  for the enlarged vector. This proves (27.2)–(27.4).  $\square$

The argument also separates a prescribed target arc from its complement: shrink the radius below  $r$ , using (27.3) on the intermediate annulus and (27.4) outside. Therefore an existing outside-arc margin may be retained, but it is not used to infer (27.3). The latter is an independently proved consequence of analytic nonconstancy and compactness.

**Lemma 27.2** (Unique matching for noisy finite signatures). *There are  $\epsilon_0, C > 0$ , uniform on the class, with the following property. Fix a true target  $J(x_0)$ . Suppose a recovered signature curve  $\tilde{J}$ , expressed in the same normalized arclength parameter, and a target vector  $y$  satisfy*

$$\|\tilde{J} - J\|_{C^2} \leq \epsilon, \quad |y - J(x_0)| \leq \epsilon, \quad 0 < \epsilon < \epsilon_0.$$

*Then the minimum of  $|\tilde{J}(x) - y|^2$  on the parameter circle is attained at one point  $\hat{x}$ , and*

$$\text{dist}(\hat{x}, x_0) \leq C\epsilon. \quad (27.5)$$

*Locally this match is Lipschitz in the target vector and in the recovered curve in  $C^1$ . A mere  $C^0$  perturbation of an arbitrary signature curve is not asserted to have a unique nearest point.*

*Proof.* Choose a radius  $a < r/4$  so small that  $2B_2Ua < v^2/16$ , decreasing it if necessary to handle bounded perturbations of the derivatives. Theorem 27.1 gives a positive uniform separation between  $J(x_0)$  and every  $J(x)$  outside this arc, namely at least  $\min(va/2, \sigma)$ . For sufficiently small  $\epsilon$ , no minimizer can lie outside the arc: the value of the residual norm at  $x_0$  is at most  $2\epsilon$ , while outside it is larger.

On the arc put  $f(x) = \frac{1}{2}|\tilde{J}(x) - y|^2$ . Its second derivative is

$$f''(x) = |\tilde{J}'(x)|^2 + (\tilde{J}(x) - y) \cdot \tilde{J}''(x).$$

The derivative lower bound, the choice of  $a$ , and then a sufficiently small  $\epsilon_0$  give  $f'' \geq v^2/8$  throughout the arc. Hence the minimum is unique and interior. At the minimum the residual norm is at most  $2\epsilon$ , so

$$|J(\hat{x}) - J(x_0)| \leq 4\epsilon.$$

The local embedding bound proves (27.5), for example with  $C = 8/v$ . Finally compare the equations  $f'(\hat{x}) = 0$  for two nearby data sets. The derivative of  $f'$  is at least  $v^2/8$ , and its change under a perturbation of  $y$  and  $\tilde{J}$  is bounded by a constant times  $|\Delta y| + \|\Delta \tilde{J}\|_{C^1}$ . The mean value theorem gives the stated local Lipschitz dependence.  $\square$

For example, a scalar curvature with a strict minimum can take each nearby larger value twice inside the correct arc. Adding a finite number of its oriented derivatives as above resolves this failure; declaring the original scalar match unique would not. Exact finite symmetries can instead be handled on a fixed symmetry quotient for a classification problem, but they must not be assigned an arbitrary phase and called a unique Euclidean gluing. The unique-reconstruction theorem uses the signature-rigid transitions specified at the start of the section.

### 27.2. Propagation through the marked graph.

**Lemma 27.3** (Persistence of the prescribed gluing). *On  $\mathfrak{R}$ , the prescribed incidence matches are unique and continuous. More precisely, if their recovered finite signature curves and targets satisfy the hypotheses of Lemma 27.2, and the recovered point–tangent frames are close in  $C^1$ , then the incidence congruences and their finite graph compositions depend locally Lipschitz-continuously on these quantities. The lattice map obeys*

$$L = VM^{-1}, \quad \|\tilde{L} - L\| \leq \|\tilde{V} - V\| \|M^{-1}\|, \quad (27.6)$$

where the fixed invertible matrix  $M$  has the two marked deck vectors as columns and  $V = (v_1 \ v_2)$ . Rank two persists for  $\|\tilde{V} - V\| < \sigma_{\text{hol}}/2$ .

*Proof.* The matching lemma supplies the unique in-arc point; its global separation excludes other arcs. Point and oriented tangent determine the required orientation-preserving rigid motion between the two recovered curve frames. The elementary formula for that motion is smooth in the point and unit tangent data on a fixed orientation chart. Products of a fixed finite number of such motions are locally Lipschitz on bounded point data. Consequently the tree placements and the two cycle holonomies inherit a uniform local bound, with a constant depending on the finite graph. The channel margin prevents a different closest pair from replacing the selected one. Formula (27.6) is direct matrix multiplication by the fixed  $M^{-1}$ , not inversion of a noisy matrix  $V$ . The singular-value perturbation inequality proves rank persistence. Also

$$\|\tilde{L}^t \tilde{L} - L^t L\| \leq 2\|L\| \|\tilde{L} - L\| + \|\tilde{L} - L\|^2,$$

so the Gram form and subsequent deck corrections are stable as well.  $\square$

This lemma concerns already recovered signature curves in the stated norms. It does not turn noisy contact Taylor coefficients into a well-conditioned unrestricted analytic continuation algorithm. The next argument supplies the distinct, non-effective continuity route needed for global statistical consistency.

### 27.3. The compact inverse modulus. Normalize the action coordinates by

$$x_{e,b,m}(T) = \rho^m S_{e,b}^{(m)}(0; T)/m!, \quad m \geq 2,$$

using any fixed positive scale  $\rho$ , and the gaps by  $g_*$ . For a fixed enumeration  $|e|$  of the finite channel set define

$$\begin{aligned} d_{\mathcal{D}}(T, T') &= \sum_{e,b} \sum_{m \geq 2} 2^{-(m+|e|+3)} \min\{1, |x_{e,b,m}(T) - x_{e,b,m}(T')|\} \\ &\quad + \sum_e 2^{-(|e|+2)} \min\{1, |g_e(T) - g_e(T')|/g_*\}. \end{aligned} \quad (27.7)$$

The truncation by one and the summable weights make this metric compatible with coordinatewise convergence of the intrinsic data. This conclusion requires no assertion that the action holomorphic radius equals the curvature holomorphic radius.

**Theorem 27.4** (Uniform inverse modulus). *For every fixed  $q \geq 0$ ,*

$$\omega_q(t) = \sup\{d_q(T, T') : T, T' \in \mathfrak{R}, d_{\mathcal{D}}(T, T') \leq t\} \quad (27.8)$$

*is finite, nondecreasing, and tends to zero as  $t \downarrow 0$ . If  $\mathcal{D}_M$  consists of the gaps and action jets through order  $M$ , then*

$$d_q(T, T') \leq \omega_q(C\|\mathcal{D}_M(T) - \mathcal{D}_M(T')\|_M + C2^{-M}). \quad (27.9)$$

*The finite norm is the maximum norm of the normalized retained coordinates;  $C$  depends only on the finite graph and the chosen weights.*

*Proof.* Equality of the complete action data and gaps gives the full analytic contact germs by Theorem 22.1. Analytic continuation, exact signature-rigid gluing, and Theorem 26.3 determine the marked periodic table in the root-frame gauge. This exact injectivity uses the complete signatures; it does not depend on a possibly noisy finite matching rule. The complete data map is continuous coordinate by coordinate, by the half-line inverse estimates, and hence continuous in (27.7). A continuous injection from compact  $\mathfrak{K}$  into this metric data space has a uniformly continuous inverse on its image. This proves the assertion about  $\omega_q$ . The terms with  $m > M$  in the data metric sum to at most  $C2^{-M}$ ; the retained terms are bounded by the normalized finite maximum norm times  $C$ . Substitution gives (27.9).  $\square$

*Remark 27.5* (Two distinct stability statements). The local finite-signature estimates are proved by immersion, separation and strict convexity of the matching criterion. The global modulus is proved by exact injectivity and compactness. Neither proof substitutes for the other. The tail  $C2^{-M}$  is a product-metric truncation bound, not an exponential analytic-continuation error rate. No effective formula for  $\omega_q$  or a uniform-in-order jet condition number is asserted.

## 28. A NONSYMMETRIC PERIODIC REALIZATION WITH FOUR MEASURED CHANNELS

We give a family in which the contact geometry, the two deck cycles and intrinsic registration can all be evaluated in one construction. The parameters used to construct the table are not supplied to its observer. In particular, the coordinates in the following construction verify realizability; they are not additional observations. The example uses selected channels, not a decomposition of the whole collision-count event.

Put  $n(\theta) = (\cos \theta, \sin \theta)$  and  $t(\theta) = (-\sin \theta, \cos \theta)$ . If  $h + h'' > 0$ , the curve

$$x_h(\theta) = h(\theta)n(\theta) + h'(\theta)t(\theta) \quad (28.1)$$

has derivative  $(h + h'')t$ , outward normal  $n$ , support function  $h$ , and curvature  $(h + h'')^{-1}$ . To see the support assertion directly, for fixed  $\theta_0$  the derivative of  $x_h(\theta) \cdot n(\theta_0)$  is  $(h + h'')(\theta) \sin(\theta_0 - \theta)$ . Its unique maximum occurs at  $\theta = \theta_0$ . Thus (28.1) is the boundary of a strictly convex body, denoted by  $K_h$ .

Consider the two support functions

$$\begin{aligned} h_1(\theta) &= 1 + \epsilon \left( \cos 2\theta + \frac{1}{4} \cos(3\theta + \pi/4) \right), \\ h_2(\theta) &= \frac{3}{4} + \delta \left( \cos 2\theta + \frac{1}{4} \cos(3\theta + \pi/6) \right), \end{aligned} \quad \epsilon, \delta \in [1/200, 1/100]. \quad (28.2)$$

Let  $T$  be a real  $2 \times 2$  matrix with  $\|T - I\| \leq 1/100$ , and let  $|w| \leq 1/100$ . Define

$$L = T \begin{pmatrix} 8 & 0 \\ 0 & 10 \end{pmatrix}, \quad v = T(4, 5)^t + w, \quad D_1 = K_{h_1}, \quad D_2 = v + K_{h_2}. \quad (28.3)$$

The periodic obstacles are  $D_a + Lm$ ,  $a \in \{1, 2\}$ ,  $m \in \mathbb{Z}^2$ . The marked basis of the abstract deck group is denoted by  $\mathbf{e}_1, \mathbf{e}_2$ ; neither column of  $L$  is given as a datum.

**Theorem 28.1** (A uniformly admissible nonsymmetric family). *Every member of (28.2)–(28.3) is a periodic dispersing table with connected analytic boundaries. Every two distinct obstacle closures have distance greater than  $43/10$ . The curvature ranges on the two obstacle types are disjoint:*

$$\frac{20}{21} \leq \kappa_1 \leq \frac{20}{19}, \quad \frac{5}{4} \leq \kappa_2 \leq \frac{10}{7}. \quad (28.4)$$

*Neither obstacle has a nonidentity Euclidean symmetry. Modulo its period translations, the complete table has no Euclidean symmetry either.*



*The four labelled channels*

$$e_{ij} = (1, 2, -ie_1 - je_2), \quad i, j \in \{0, 1\}, \quad (28.5)$$

are unobstructed closest-pair channels for their specified obstacle pairs. Every such closest segment has distance at least  $383/200$  from every third obstacle. Their contacts and gaps depend real analytically on the parameters and have common contact and positive-offset collars. All their incidence transitions are signature-rigid. The cycles

$$c_1 = e_{00}e_{10}^{-1}, \quad c_2 = e_{00}e_{01}^{-1} \quad (28.6)$$

form a rank-two anchoring pair based at  $e_{00}$ ; the edge  $e_{00}$  is a signature-rigid spanning tree rooted at obstacle 1.

*Proof.* For  $h = r + z(\cos 2\theta + \frac{1}{4}\cos(3\theta + \phi))$ ,

$$h + h'' = r - 3z \cos 2\theta - 2z \cos(3\theta + \phi).$$

The curvature radii lie in  $[19/20, 21/20]$  and  $[7/10, 4/5]$ , respectively. This proves strict convexity and (28.4). The support bound gives  $K_{h_a} \subset \overline{B}(0, R)$  with  $R = 81/80$  for both types. At  $T = I, w = 0$ , the set of obstacle centers is

$$\{(8m, 10n) : m, n \in \mathbb{Z}\} \cup \{(4 + 8m, 5 + 10n) : m, n \in \mathbb{Z}\}.$$

Its minimum separation is  $\sqrt{41}$ . Since  $|Tz| \geq (99/100)|z|$ , all distinct perturbed centers have separation at least  $(99/100)\sqrt{41} - 1/100$ . Subtracting  $2R$  gives

$$\text{dist}(D_a + Lm, D_b + Ln) \geq \frac{99}{100}\sqrt{41} - \frac{407}{200} > \frac{43}{10}$$

for distinct obstacles; the last strict inequality follows already from  $\sqrt{41} > 32/5$ . The smallest singular value of  $T$  is positive and  $\det T > 0$ , so the quotient is a genuine flat torus. Its free area is positive. More explicitly, integration of  $\frac{1}{2}h(h + h'')$  gives

$$A = 80 \det T - \frac{25\pi}{16} + \frac{7\pi}{4}(\epsilon^2 + \delta^2) > 0. \quad (28.7)$$

We next exclude symmetry without imposing any contact registration. The first Fourier coefficients of both supports vanish. The vector

$$s(K) = \frac{1}{\pi} \int_0^{2\pi} h_K(\theta) n(\theta) d\theta \quad (28.8)$$

is covariant under translations and orthogonal transformations, so any symmetry of either body fixes the origin. A rotational symmetry through  $\alpha$  would preserve its nonzero modes 2 and 3, and hence satisfy  $e^{2i\alpha} = e^{3i\alpha} = 1$ . Thus it is the identity. Reflection in a line at angle  $\alpha$  would require  $2\alpha \in \pi\mathbb{Z}$  and  $3\alpha + \phi \in \pi\mathbb{Z}$ . Neither  $\phi = \pi/4$  nor  $\phi = \pi/6$  permits this: substituting  $\alpha = k\pi/2$  gives, respectively,  $(6k + 1)/4$  and  $(9k + 1)/6$  as nonintegral multiples of  $\pi$ . The two perimeters,  $2\pi$  and  $3\pi/2$ , differ. A symmetry of the periodic table must therefore preserve obstacle types. Its orthogonal part would be a symmetry of  $K_{h_1}$  and is the identity; its translation belongs to  $L\mathbb{Z}^2$ .

For clearance, first consider the center segment from  $(0, 0)$  to  $(4, 5)$ . Every third reference center has distance at least 4 from the rectangle  $[0, 4] \times [0, 5]$ : this follows directly by separating the integer coordinates of the two displayed center sets. The other three rectangles are reflections of this one. After applying  $T$ , the distance of a third center from the corresponding center segment is at least  $4(99/100)$ . Moving the second-type centers by  $w$  changes a segment by at most  $|w|$  in Hausdorff distance and a third center by at most  $|w|$ . Thus the lower bound becomes  $4(99/100) - 2/100$ . A segment joining points of the two radius- $R$  bodies is contained

in the radius- $R$  neighborhood of the center segment, by convex interpolation of its endpoints. Subtracting a further  $R$  for the third body proves the claimed clearance:

$$4\frac{99}{100} - \frac{2}{100} - 2\frac{81}{80} = \frac{383}{200}. \quad (28.9)$$

This argument covers every lattice translate, not a finite numerical list.

For completeness, the contacts can be found from a scalar equation. Put

$$v_{ij} = T(4 - 8i, 5 - 10j)^t + w, \quad d_{ij} = |v_{ij}|, \quad \vartheta_{ij} = \arg v_{ij},$$

and maximize

$$F_{ij}(\theta) = v_{ij} \cdot n(\theta) - h_1(\theta) - h_2(\theta + \pi). \quad (28.10)$$

The variable  $d_{ij}$  here is a center distance, not the positive time offset  $d$  used elsewhere. The perturbation in  $h_1 + h_2$  from  $7/4$  is at most  $1/40$ . Comparison with  $\theta = \vartheta_{ij}$  excludes  $|\theta - \vartheta_{ij}| \geq \pi/3$  from the set of maxima. On the remaining arc,

$$F''_{ij} \leq -d_{ij}/2 + 1/8 < 0,$$

because  $\|h''_1\| + \|h''_2\| \leq 1/8$ . There is a unique maximum  $\theta_{ij}$  and it depends analytically on the parameters. With

$$p_{ij} = x_{h_1}(\theta_{ij}), \quad q_{ij} = v_{ij} + x_{h_2}(\theta_{ij} + \pi), \quad g_{ij} = F_{ij}(\theta_{ij}), \quad (28.11)$$

the equation  $F'_{ij} = 0$  gives  $q_{ij} - p_{ij} = g_{ij}n(\theta_{ij})$ . The two supporting half-planes show that all cross-body distances are at least  $g_{ij}$ , so these points realize the closest pair. The clearance bound excludes a third impact on this segment. Positive curvature, positive gap and the strict clearance margins give its alternating contact collar. Compactness of the parameter set and the finite number of channels give a common collar and, by the relative theorem, a common positive-offset collar. No finite-horizon assertion is used.

One can also choose gates without using an unknown contact point as an observation. Let  $\vartheta_{ij}^0 = \arg(4 - 8i, 5 - 10j)$ . The center perturbation is less than  $3/40$  and  $d_{ij} > 6.32$ , so  $|\vartheta_{ij} - \vartheta_{ij}^0| < 0.012$ . The equation  $F'_{ij} = 0$  and  $\|h'_1\| + \|h'_2\| \leq 11/200$  give  $|\theta_{ij} - \vartheta_{ij}| < 0.009$ . Hence the first contact is within  $0.062$  of  $n(\vartheta_{ij}^0)$  and the second within  $0.131$  of  $(4 - 8i, 5 - 10j)^t - \frac{3}{4}n(\vartheta_{ij}^0)$ . Fixed planar disks of radius  $1/4$  about these reference points therefore contain all the appropriate contacts with a uniform interior margin. On each obstacle these gates are disjoint. Shrinking the contact collars inside them preserves the strict third-obstacle clearance. These are prescribed selected gates; their common reference coordinates in this proof are not passed to the intrinsic observer as an inter-channel registration.

Finally, Lemma 25.2 and absence of proper obstacle symmetries make every required incidence transition rigid. The two deck sums in (28.6) are  $e_1$  and  $e_2$ , and both cycles start with  $e_{00}$ . The stated tree reaches the only other obstacle orbit. This checks each part of the anchoring and spanning hypotheses in a single realization.  $\square$

**28.1. Recovering the channel frames from the recovered curves.** For a complete convex image  $C$ , let

$$z_k(C) = \frac{1}{2\pi} \int_0^{2\pi} h_C(\theta) e^{-ik\theta} d\theta, \quad k \geq 2.$$

These coefficients do not change under translation. They are derived from the boundary image recovered by analytic continuation, not added to the endpoint-law observations. Identify the oriented plane with  $\mathbb{C}$ , so a unit complex number denotes a rotation.

**Proposition 28.2** (Two-harmonic registration in the realized family). *Let  $C_{e,1}$  and  $C_{e,2}$  be the two complete boundary images recovered in the frame of channel  $e$ , whose origin is its first contact. Set  $e_* = e_{00}$  and  $c_{e,a} = s(C_{e,a})$ . The Euclidean congruence placing the first-type copy of edge  $e$  onto that of  $e_*$  is*

$$U_e = \frac{z_3(C_{e,1})z_2(C_{e_*,1})}{z_3(C_{e_*,1})z_2(C_{e,1})}, \quad A_e(y) = U_e y + c_{e_*,1} - U_e c_{e,1}. \quad (28.12)$$

*In exact data  $|U_e| = 1$ . If  $b_e = A_e(c_{e,2})$ , then the lattice columns and Gram matrix in this one common frame are*

$$V = (b_{00} - b_{10} \quad b_{00} - b_{01}), \quad G = V^T V. \quad (28.13)$$

*The redundant fourth channel satisfies*

$$b_{00} - b_{11} = (b_{00} - b_{10}) + (b_{00} - b_{01}). \quad (28.14)$$

*On this compact family, registration and the displayed lattice recovery are locally Lipschitz in the support functions of the recovered complete curve images, in the uniform norm. This statement does not assert a Lipschitz analytic-continuation map from contact jets or endpoint laws.*

*Proof.* Write the actual edge placement as  $y \mapsto R_e y + p_e$ , with rotation angle  $\theta_e$ . Its first-type image is  $C_{e,1} = R_e^T(K_{h_1} - p_e)$ . Therefore

$$z_k(C_{e,1}) = e^{ik\theta_e} z_k(K_{h_1}), \quad z_2(K_{h_1}) = \epsilon/2, \quad z_3(K_{h_1}) = (\epsilon/8)e^{i\pi/4}.$$

Taking the ratio for mode 3 and dividing by that for mode 2 yields  $e^{i(\theta_e - \theta_{e_*})}$ . This proves the rotation formula without selecting a branch of an argument. Translation covariance of (28.8) proves the formula for  $A_e$ . Indeed  $c_{e,1} = R_e^T(-p_e)$ , and hence

$$A_e(y) = R_{e_*}^T R_e y + R_{e_*}^T(p_e - p_{e_*}), \quad b_{ij} = R_{e_*}^T(v_{ij} - p_{e_*}).$$

Taking differences proves (28.13) and (28.14). These differences are precisely the translation holonomies of the two cycles, since their marked displacement matrix is  $I$ . The copy  $C_{e_*,1}$  and the placed copy  $C_{e_*,2}$ , together with  $V\mathbb{Z}^2$ , recover every obstacle translate.

For stability, the maps  $h \mapsto z_k$  and  $h \mapsto s$  have uniform-norm operator bounds 1 and 2. On the stated family  $|z_2| \geq 1/400$  and  $|z_3| \geq 1/1600$ . The quotients in (28.12) are therefore locally Lipschitz under small uniform perturbations; normalize a perturbed  $U_e$  by its modulus if it is not exactly a rotation. This normalization is smooth near the unit circle. The centers lie in a common bounded set, so the translation formula, the differences in  $V$ , and  $V^T V$  inherit local Lipschitz bounds. This proof concerns the registration stage after the complete curve images have been obtained.  $\square$

## 28.2. The observation-to-table reconstruction in this family.

**Corollary 28.3** (Eight signed laws and four onsets). *There is  $d_0 > 0$ , common to the family, such that the four marked gaps and the two same-type signed limiting endpoint laws at offset  $d_0$  for each of the four channels determine the whole marked table and its Euclidean lattice up to one simultaneous proper Euclidean motion. The parameters  $T, w, \epsilon, \delta$ , the support coefficients, the contact curvatures, and the placements of different channel frames are not supplied. The two actions at a channel are not both even. In particular, replacing signed endpoint laws by their pointwise sign-erased versions is not the observation used in this corollary.*

*Proof.* Choose  $d_0$  inside the uniform positive-offset collar in Theorem 28.1. At each edge and each contact type, form the four-density ratio of Theorem 23.1 on a common positive interior

square. A nonzero anchor recovers  $S_{e,b}$  without knowing  $B_{e,b}$  or the normalization. In particular, if  $a_{e,b} = S''_{e,b}(0)$ , then

$$\gamma_e = \operatorname{arsinh}(g_e \sqrt{a_{e,0} a_{e,1}}), \quad \kappa_{e,b} = \frac{\cosh \gamma_e \sqrt{a_{e,b}/a_{e,1-b}} - 1}{g_e}.$$

For every  $n \geq 3$  the smooth finite-remainder theorem and the homogeneous isolation in Theorem 22.1 give  $q_{e,n} = M_{e,n}^{-1}(s_{e,n} - R_{e,n})$ . The geometric bounds above do not replace that nonlinear inverse; they ensure its hypotheses hold simultaneously at all eight contacts. The recovered analytic germs determine the complete images  $C_{e,a}$ , by Lemma 23.14. Now use Proposition 28.2 to place all edges in the  $e_{00}$  frame and recover  $V$  and both obstacle images. This is also the concrete common-frame implementation of Theorem 23.4.

If both actions in a channel were even, reflecting both transverse graphs would leave the two actions unchanged. The uniqueness of the contact-jet inverse would make both graph jets even. Analyticity would then make each whole obstacle invariant under reflection in its contact normal, contradicting Theorem 28.1. Thus the construction is not reduced to the reflection-symmetric local case.  $\square$

For a fully specified member, take

$$T = \begin{pmatrix} 1 & 1/200 \\ 0 & 1 \end{pmatrix}, \quad w = (1/200, -1/200)^t, \quad \epsilon = 1/125, \quad \delta = 3/400. \quad (28.15)$$

Its four center vectors are  $(4.030, 4.995)$ ,  $(-3.970, 4.995)$ ,  $(3.980, -5.005)$  and  $(-4.020, -5.005)$  in the construction frame. Solving the strictly concave scalar problems (28.10) gives the actual contacts and onsets; these are not assumed to lie on the center directions. The exact recovered Gram matrix, which is independent of the unknown starting-frame rotation, is

$$G = \begin{pmatrix} 64 & 2/5 \\ 2/5 & 40001/400 \end{pmatrix}. \quad (28.16)$$

The finite diagnostic supplied with the source solves all four contact equations and computes (28.12) from the edge-frame support functions. It does not fit the matrix  $G$  directly or give the reconstruction routine the generating matrix  $T$ . The numerical comparison with (28.16) checks this finite geometric computation, not the all-order boundary-law theorem.

### 28.3. Persistence beyond a finite-dimensional model.

**Proposition 28.4** (Analytic neighborhoods of the realization). *The family above has a neighborhood under real-analytic support-function perturbations in which the four selected channels, strict clearance, absence of obstacle symmetry, signature-rigid spanning structure and rank-two anchoring pair persist. One may independently add to  $h_1, h_2$  any real-analytic  $r_1, r_2$  with  $\|r_a\|_{C^2} \leq 10^{-6}$ . Consequently the intrinsic determination applies on an infinite-dimensional analytic neighborhood, not only to tables with the finite-dimensional shape and placement description in (28.2)–(28.3). On each compact analytic subfamily satisfying the regularity conventions of the global physical theorem, its charged long-bridge reconstruction conclusion applies with the same four channel labels.*

*Proof.* Here  $\|r\|_{C^2} = \max_{0 \leq k \leq 2} \|r^{(k)}\|_\infty$ . The perturbation in curvature radius is at most  $2 \cdot 10^{-6}$  and that in the enclosing radius at most  $10^{-6}$ . All separation, curvature and clearance inequalities retain strict margins. The scalar contact maximum remains nondegenerate, so the channels and their gate margins persist. Fourier coefficients change by at most  $10^{-6}$ . In particular the nonzero modes 2 and 3, of minimum moduli  $1/400$  and  $1/1600$ , remain nonzero and exclude every proper rotation.

For reflections, write the arguments of these two coefficients as  $\beta_2, \beta_3$ . A necessary condition is  $2\beta_3 - 3\beta_2 \in \pi\mathbb{Z}$ . In the unperturbed family this quantity is  $\pi/2$  or  $\pi/3$ . Each argument changes by at most  $2 \cdot 10^{-6}/(1/1600) \leq 0.0032$ , so the required congruence remains impossible. The vector (28.8), which may now move, removes the translation part of any candidate symmetry. The two perimeters remain different. Thus signature uniqueness persists, and the same combinatorial cycles and tree prove the intrinsic determination. The two-harmonic formula continues to hold because the same unknown coefficients cancel between copies; it does not require the absence of higher Fourier modes.

For example, fixing any strip width  $\sigma > 0$  permits infinitely many independent sufficiently small perturbations  $a_n \cos n\theta$  and  $b_n \sin n\theta$ , with exponentially summable coefficients. Hence this is genuinely an analytic neighborhood rather than a finite-dimensional reparametrization. On a compact analytic subfamily the positive margins and the hypotheses of the observable calibration and global physical reconstruction theorems give a common collar, the finite separating tests, and the charged preparation caps. Running their budget-indexed stages with the present four labels proves the last assertion. Compactness is used at this last step; the whole bounded analytic neighborhood is not silently asserted to be compact in an unspecified topology.  $\square$

The example thus implements the entire intrinsic route: nonlinear long-bridge laws, amplitude-free action inversion, both contact jets, analytic images, unique registration, and an uncalibrated rank-two lattice. Its explicit registration bounds do not supply an effective rate for the preceding analytic continuation. In the richer planar-position protocol, the direct-position benchmark remains available as well. The two routes continue to use different retained information, even on this same nonsymmetric periodic family.

## 29. FINITE CHANNEL SKELETONS AND GENERIC PERIODIC RIGIDITY

The incidence hypotheses of the intrinsic inverse are not restricted to specially arranged families. We prove that a finite clear channel skeleton exists in every positively separated periodic table, and that its signature-rigidity requirements hold on an open dense set. The resulting inverse uses  $N + 1$  selected channels for  $N$  labelled obstacle orbits. The selection is a geometric existence statement, not an algorithm for discovering unmarked channels from a trajectory.

Write  $K_a = \overline{D}_a$ ,  $1 \leq a \leq N$ , and fix representatives in the plane. Let  $L \in \text{GL}^+(2, \mathbb{R})$  realize the marked deck group  $\mathbb{Z}^2$ . Throughout this section distinct members of

$$\mathcal{B} = \{K_a + Lm : a \in \{1, \dots, N\}, m \in \mathbb{Z}^2\}$$

are disjoint compact strictly convex bodies with smooth boundary and positive curvature. Analyticity will be imposed only when signatures and complete image recovery are used. An edge is counted without its reverse; loops and multiple edges in the quotient graph are allowed.

**29.1. Clear closest segments and descent through an obstruction.** For two distinct bodies  $A, B \in \mathcal{B}$ , let  $d(A, B) > 0$  be their distance and let  $[p_{AB}, q_{AB}]$  be their closest segment. The closest points are unique. Indeed, strict convexity of the squared norm makes the minimizing vector in  $B - A$  unique, and the two supporting faces perpendicular to this vector consist of single points. We call the pair *clear* if its closed closest segment meets no member of  $\mathcal{B} \setminus \{A, B\}$ . Thus a third-body tangency is an obstruction, not a clear pair.

**Lemma 29.1** (Uniform separation and finite distance sublevels). *There is  $\sigma > 0$  such that  $d(A, B) \geq \sigma$  for all distinct bodies in  $\mathcal{B}$ . For every  $D < \infty$ , only finitely many pairs at distance at most  $D$  occur modulo common deck translation. A clear closest segment has positive distance from the union of all third bodies.*

*Proof.* Choose  $R$  with  $K_a \subset \overline{B}(0, R)$  for every  $a$ , and write  $s = s_{\min}(L) > 0$ . Then

$$d(K_a, K_b + L\ell) \geq s|\ell| - 2R. \quad (29.1)$$

For any fixed finite upper distance, the possible  $\ell$  therefore belong to a finite set. Select one pair and use its finite positive distance as an upper bound when taking the minimum over all pairs. The minimum is attained among finitely many disjoint pairs and is positive. This proves the first two assertions. For the last, bodies intersecting the closed unit neighborhood of a fixed compact segment form a finite set, by the same estimate. The distance from the segment to each of these third bodies is positive when the pair is clear; all remaining bodies have distance greater than one. Taking a finite minimum proves the assertion.  $\square$

**Lemma 29.2** (Obstruction descent). *For every pair  $A \neq B$  there is a finite path of clear pairs*

$$A = A_0, A_1, \dots, A_k = B$$

*such that*

$$\sum_{i=1}^k d(A_{i-1}, A_i) \leq d(A, B), \quad k \leq \left\lfloor \frac{d(A, B)}{\sigma} \right\rfloor. \quad (29.2)$$

*In particular, the graph of clear pairs on all lifted obstacles is connected.*

*Proof.* If the pair is not clear, choose a third body  $C$  and  $z \in C \cap [p_{AB}, q_{AB}]$ . Separation implies that  $z$  is in the open segment. With  $d = d(A, B)$ , the two shorter distances obey

$$\begin{aligned} d(A, C) &\leq |p_{AB} - z|, & d(C, B) &\leq |z - q_{AB}|, \\ d(A, C) + d(C, B) &\leq d, & \max\{d(A, C), d(C, B)\} &\leq d - \sigma. \end{aligned} \quad (29.3)$$

The last inequality uses  $|p_{AB} - z|, |z - q_{AB}| \geq \sigma$ . Induct on  $\lfloor d/\sigma \rfloor$ . A blocked pair cannot have  $d < 2\sigma$ ; each pair on the right of (29.3) has a strictly smaller induction index. Concatenate their clear paths. Their total edge distance is at most  $d(A, C) + d(C, B) \leq d$ . Every terminal edge has distance at least  $\sigma$ , proving the bound on  $k$ . Erasing any repeated-vertex subwalk gives a path with the same bounds. This proof also treats tangential obstructions. No length inside an intermediate body is charged: the assertion concerns a graph path, not a physical billiard itinerary or a bound on its travel time.  $\square$

**29.2. A finite skeleton with two independent deck cycles.** An oriented quotient edge is  $e = (a, b, \ell)$ , joining  $K_a$  to  $K_b + L\ell$ . A path lifts uniquely once its starting lift is fixed, and its deck gain is the signed sum of its edge labels. For a closed quotient walk this gain belongs to  $\mathbb{Z}^2$ .

**Theorem 29.3** (Finite clear skeleton). *Every table in the setting above has a set  $E$  of  $N + 1$  clear selected channels whose quotient multigraph is connected and contains two closed walks with linearly independent deck gains. It is the union of a spanning tree and two further edges. All selected contacts have positive third-obstacle clearance and admit fixed gates and alternating contact collars. The same finite labels, tree, gains and gates persist for sufficiently small  $C^2$  perturbations of the supports and small perturbations of  $L$  that preserve the stated setting.*

*More precisely, set*

$$D_0 = \max(\{d(K_1, K_a) : 2 \leq a \leq N\} \cup \{d(K_1, K_1 + Le_i) : i = 1, 2\}). \quad (29.4)$$

*The channels may be selected among the finite set of clear pairs with gap at most  $D_0$ . With  $R, s$  as in Lemma 29.1, their labels satisfy  $|\ell| \leq (D_0 + 2R)/s$ .*

*Proof.* Apply Lemma 29.2 to the  $N - 1$  pairs  $(K_1, K_a)$  and to the two pairs  $(K_1, K_1 + Le_i)$ . Every edge on the resulting finite paths has gap at most  $D_0$ . Project their union to a finite quotient graph  $\Gamma$ . It is connected: the first group of paths reaches every obstacle type. The last two paths project to closed walks based at type 1 whose deck gains are exactly  $e_1$  and  $e_2$ . In particular the

gains of all closed walks in  $\Gamma$  generate the whole deck group, not merely a real one-dimensional subspace. Finiteness and the label bound follow from (29.1).

Choose a spanning tree  $\mathcal{T}$  of  $\Gamma$ , rooted at type 1. For its root-to- $a$  path let  $m_a \in \mathbb{Z}^2$  be the signed deck sum, with  $m_1 = 0$ . For an oriented edge  $e = (a, b, \ell_e)$  define

$$\eta_e = m_a + \ell_e - m_b. \quad (29.5)$$

This is zero for tree edges. The gain of any closed walk is the signed sum of its  $\eta_e$ , since the  $m_a$  terms telescope. Conversely, for a non-tree edge,  $\eta_e$  is the gain of the walk from the root to  $a$  along the tree, then across  $e$ , and back along the tree from  $b$  to the root. The non-tree gains consequently generate the deck-gain group of  $\Gamma$ . Two of them, say  $\eta_{f_1}, \eta_{f_2}$ , are linearly independent over  $\mathbb{R}$ . Take

$$E = \mathcal{T} \cup \{f_1, f_2\}.$$

It has  $N - 1 + 2 = N + 1$  edges, including when  $N = 1$  and the tree is empty. The two gains need not form an integer basis. Their real independence is exactly the hypothesis needed to recover  $L$  from  $L\eta_{f_i}$ ; one must not replace their displacement matrix by the identity.

We verify the geometric assertions rather than identifying a graph edge with an observation without a collar. At a closest pair, use transverse coordinates whose positive directions agree in the plane. The Hessian of the flight length at its normal contact is

$$\begin{pmatrix} \kappa_a + g^{-1} & -g^{-1} \\ -g^{-1} & \kappa_b + g^{-1} \end{pmatrix}, \quad \det = \kappa_a \kappa_b + (\kappa_a + \kappa_b)/g > 0. \quad (29.6)$$

Thus the closest point is nondegenerate. The local implicit construction and strict convexity make the closest contacts continuous under  $C^2$  perturbation and smooth, respectively analytic, along smooth, respectively analytic, finite-dimensional families. Their connecting segments inherit the positive clearance of Lemma 29.1.

Distinct clear incidences on a given lifted obstacle have distinct contact points. Otherwise their closest segments issue from the same point along the same outward normal ray. The nearer endpoint lies on the farther segment and on a third body, contradicting clearness; equality of the two endpoints contradicts disjointness. Since  $E$  is finite, choose disjoint small planar gates around its incident contacts on each representative obstacle, and choose fixed reference gates also at the finitely many lifted endpoints used by  $E$ . Periodic copies at the reference table are assigned their marks before observation. Shrink these gates so that their connecting segments stay in the clear tubes. During a perturbation the finitely many laboratory gates stay fixed; the unknown lattice is not used to relocate them. They are not exact contact observations. After decreasing the perturbation neighborhood they still contain the corresponding contacts with positive interior margins.

There is no unexamined infinity of third obstacles in this persistence argument. On a small parameter neighborhood all bodies have a common radius bound and  $s_{\min}(L)$  has a positive lower bound. Only finitely many translates can approach any of the finitely many selected segments; their positive margins persist. The same argument preserves all-pair separation. Formula (29.6), positive curvature, and the fixed gates give the alternating contact collars. The labels and the two integer gain vectors do not change under this perturbation.  $\square$

**Lemma 29.4** (Locality of a selected clear channel). *The selected-channel conclusions of Theorem 12.1, the signed contact inverse, and the single-offset density inverse apply to every channel of Theorem 29.3. These channels need not attain the minimum gap of the whole table. Each fixed table has a common positive offset collar for this finite channel set. On compact smooth or analytic subfamilies with the uniform higher-norm assumptions of the relative theorem, the collar and every fixed-order bound may be chosen uniformly.*

*Proof.* The event under consideration specifies the same two facing patches and alternation as the selected event of Section 12. Every one of its flights has length at least this pair's gap  $g_e$ . In a window  $jg_e + d$ , the total excess of the  $j$  flights above  $jg_e$ , and both unused end intervals, is at most  $d$ . On fixed patches strict convexity and (29.6) give a positive quadratic lower bound on one-flight excess near the closest points, and a positive lower bound away from that neighborhood. Thus a sufficiently small  $d$  localizes all the selected impacts in the contact collars. Short unused end intervals remain in their clear tubes after the offset collar is shrunk once more.

The weighted Jacobi construction and its determinant comparison use the local positive gap, positive curvatures and twist, not a comparison with another pair's gap. Their resulting bridges remain within the same clear tubes. The half-line construction and the fixed-domain sublevel integration therefore give precisely the selected relative law. Its single-offset density and smooth contact inverse depend on that law and the selected gap, so they apply unchanged. Taking a minimum over finitely many channels proves the fixed-table collar assertion; taking the uniform constants in the stated compact-family hypotheses proves the last one. A  $C^2$  geometric neighborhood alone is not asserted to bound arbitrarily high derivatives or to provide a compact analytic class.  $\square$

**Proposition 29.5** (The channel count for two-cycle reconstruction). *A connected quotient multigraph on  $N$  obstacle types that admits two linearly independent deck-cycle gains has at least  $N + 1$  edges. Consequently Theorem 29.3 attains the smallest channel count for this connected two-cycle reconstruction mechanism.*

*Proof.* If the graph has  $m$  edges, a spanning tree leaves  $m - N + 1$  non-tree edges. As in (29.5), their gains generate every closed walk gain. Real rank two requires  $m - N + 1 \geq 2$ .  $\square$

This is a graph-theoretic count of selected channels, not a lower bound on the number of preparations, the dimension of a probability law, or the information needed by every conceivable inverse observation map.

**29.3. Generic signature rigidity.** Let  $\mathcal{P}_N^\omega$  denote the space of such labelled periodic tables with real-analytic boundaries. A parameter consists of  $L$  and the support functions  $h_a$  of the representatives  $K_a$ ; translations are retained in the first harmonics. Give this space the relative  $C^2$  topology on supports and the ordinary topology on  $L$ . The assertions below hold also with any fixed stronger  $C^r$  support topology. This definition of genericity makes no compactness or common complex-extension assertion.

For a convex body  $K$ , put

$$s(K) = \frac{1}{\pi} \int_0^{2\pi} h_K(\theta) n(\theta) d\theta, \quad z_k(K) = \frac{1}{2\pi} \int_0^{2\pi} h_K(\theta) e^{-ik\theta} d\theta.$$

The vector  $s(K)$  is translation and orthogonal covariant. The coefficients  $z_k$  for  $k \geq 2$  are translation invariant; the centered body  $K - s(K)$  has zero first harmonic.

**Proposition 29.6** (Open dense obstacle asymmetry). *The subset  $\mathcal{A}_N \subset \mathcal{P}_N^\omega$  for which every obstacle has trivial orientation-preserving Euclidean symmetry group is open and dense. It contains the open dense subset  $\mathcal{H}_N$  defined by, for every  $a$ ,*

$$z_2(K_a)z_3(K_a) \neq 0, \quad \text{Im}(z_3(K_a)^2 \overline{z_2(K_a)})^3 \neq 0. \quad (29.7)$$

*Every individual obstacle of a table in  $\mathcal{H}_N$  has no nonidentity Euclidean symmetry, including reflections. No restriction is placed on its other analytic harmonics.*

*Proof.* A symmetry of a bounded body fixes its covariant center. For the centered support, rotation by  $\alpha$  is a symmetry exactly when  $e^{-ik\alpha} z_k = z_k$  for all  $k \geq 2$ . A body has trivial



proper symmetry exactly when the indices of its nonzero coefficients have greatest common divisor one. In that case finitely many such indices already have greatest common divisor one: successive greatest common divisors are positive integers and stabilize. Those finitely many nonzero coefficients remain nonzero after a sufficiently small support perturbation, proving openness of  $\mathcal{A}_N$ .

Nonzero coefficients of orders two and three force a rotational symmetry to be the identity. A reflection across a line at angle  $\alpha$  through the center transforms  $z_k$  to  $e^{-2ik\alpha}\overline{z_k}$ . Invariance of both coefficients would imply

$$2\alpha + \arg z_2 \in \pi\mathbb{Z}, \quad 3\alpha + \arg z_3 \in \pi\mathbb{Z}.$$

Eliminating  $\alpha$  makes  $z_3^2\overline{z_2}^3$  real. Thus (29.7) excludes reflections as well. The conditions are open because Fourier coefficients are continuous in the uniform norm.

To prove density, perturb each support by

$$u_{a,2} \cos 2\theta + v_{a,2} \sin 2\theta + u_{a,3} \cos 3\theta + v_{a,3} \sin 3\theta.$$

These four real parameters vary its second and third complex coefficients independently in a neighborhood. Their zero sets and the zero set of the nonzero real polynomial  $\text{Im}(z_3^2\overline{z_2}^3)$  have empty interior. An arbitrarily small choice therefore satisfies (29.7). Since  $h_a + h_a'' > 0$  and separation are open conditions, the perturbations preserve positive curvature and all-pair disjointness. They are trigonometric polynomials, so they preserve analyticity and can be arbitrarily small in any fixed  $C^r$  norm. The lattice and every other harmonic may be left unchanged. This proves density of  $\mathcal{H}_N$ , hence of  $\mathcal{A}_N$ .  $\square$

**Theorem 29.7** (Generic finite-channel determination). *For every  $T \in \mathcal{A}_N$  there are  $N + 1$  marked selected channels and a positive offset  $d_0$  such that their  $N + 1$  onsets and their  $2N + 2$  same-type signed limiting endpoint laws at  $d_0$  determine the complete labelled periodic table and its Euclidean lattice up to one simultaneous orientation-preserving Euclidean motion. No inter-channel placements, contact curvatures, support coefficients or lattice Gram form are given. The competing realizations may be any tables in  $\mathcal{P}_N^\omega$  realizing these selected data in the contact-law observation class of Theorem 23.4, not only members of  $\mathcal{A}_N$ .*

*The channel labels, gates, spanning tree and two independent cycle gains can be kept fixed in a geometric neighborhood of  $T$ . In particular, this finite-channel conclusion holds on an open dense subset of  $\mathcal{P}_N^\omega$ . For the common-orientation orbit of the entire signed datum the conclusion is modulo one element of  $E(2)$ , not under independent sign erasures in different channels.*

*Proof.* Select  $E$  by Theorem 29.3 and choose a common  $d_0$  by Lemma 29.4. The four-density inverse of Theorem 23.1 recovers both unsymmetrized half-line actions at each channel from its two laws. The smooth finite-remainder inverse of Theorem 22.1 then recovers both complete contact jets. Lemma 23.14 determines the two connected analytic boundary images in that channel's own frame.

Every boundary of  $T$  has trivial proper symmetry. Lemma 25.2 implies that every one of its complete oriented contact signatures occurs at exactly one framed point. Hence all incidence transitions in  $E$  are signature-rigid. Any competing realization has the same complete incident images and inherits this property; asymmetry is not imposed separately on the competitor.

The two fundamental closed walks may be based at the same root obstacle. Choose one incident measured channel as the common starting frame. Unique congruences between root-boundary copies transport each walk's initial and final incidence to that frame. Equivalently, prepend a zero-net-gain backtrack  $ee^{-1}$  when a common initial edge is required. These operations add neither a measured channel nor a deck displacement. They concern compositions of incidence congruences, not realization of a multi-edge closed walk as a specular trajectory. The resulting walks are a

rank-two anchoring pair in the sense of Definition 26.1. The selected tree is signature-rigid and can be rooted at this same obstacle.

Theorem 23.4, or equivalently its common-frame construction, now gives the conclusion. Explicitly, align the starting frames in two realizations. Their unique incidence congruences give the same cycle translations  $v_1, v_2$ . With  $M = (\eta_{f_1} \ \eta_{f_2})$  and  $V = (v_1 \ v_2)$ , they have the same  $L = VM^{-1}$ , including its Gram form and scale. Unique matching along the tree then places every remaining obstacle image in that frame, with deck corrections interpreted through this one  $L$ . Realizability supplies existence and admissibility; the argument proves uniqueness. Persistence follows from the clear-skeleton theorem and openness of  $\mathcal{A}_N$ . The final assertion is the equivariance of Theorem 25.14 applied to the entire marked datum.  $\square$

**29.4. An explicit recovered-image reconstruction on the harmonic locus.** On  $\mathcal{H}_N$  the incidence stage has a finite formula, independent of how many higher harmonics the unknown obstacles possess. This also gives a quantitative statement in a clearly specified topology.

**Proposition 29.8** (Harmonic registration and the displacement cochain). *Suppose  $E = \mathcal{T} \cup \{f_1, f_2\}$  is the skeleton above and every obstacle has nonzero second and third support coefficients. From the complete boundary images already recovered in the separate channel frames, all channel rotations, both lattice columns and all obstacle placements can be computed by the following finite operations.*

*For two copies  $C_s, C_t$  of the same labelled obstacle, the unique proper congruence taking  $C_s$  to  $C_t$  has rotation*

$$U_{s \rightarrow t} = \frac{z_3(C_s)z_2(C_t)}{z_3(C_t)z_2(C_s)}, \quad A_{s \rightarrow t}(y) = U_{s \rightarrow t}y + s(C_t) - U_{s \rightarrow t}s(C_s), \quad (29.8)$$

*where the plane is identified with  $\mathbb{C}$  and  $|U_{s \rightarrow t}| = 1$  in exact data. Propagate these rotations along the tree, starting with one root incidence in a fixed channel frame. Let  $R_e$  be the resulting rotation of edge  $e$  into that common frame, also obtained for the two non-tree edges by matching their initial obstacle copies. Set*

$$d_e = R_e\{s(C_{e,+}) - s(C_{e,-})\}, \quad p_1 = 0, \quad p_a = \sum_{e \in [1,a]_{\mathcal{T}}} \epsilon_e d_e. \quad (29.9)$$

*For  $f_i = (a_i, b_i, \ell_i)$  put*

$$v_i = p_{a_i} + d_{f_i} - p_{b_i}, \quad L = (v_1 \ v_2)M^{-1}, \quad c_a = p_a - Lm_a, \quad M = (\eta_{f_1} \ \eta_{f_2}). \quad (29.10)$$

*Then  $c_a$  are the representative support centers, with  $c_1 = 0$ , and centering and rotating one recovered copy of each type gives its complete boundary at  $c_a$ . Formula (29.10) also recovers  $G = L^t L$ .*

*For fixed labels and tree, this inverse stage is locally Lipschitz in the uniform norms of the support functions of the recovered complete images, uniformly when their radii and centers are bounded, their second and third coefficient moduli have positive lower bounds, and  $\|M^{-1}\|$  is bounded. It is not a Lipschitz assertion from endpoint laws or finite contact jets to entire analytic boundaries.*

*Proof.* For a rotation  $U = e^{i\alpha}$  and translation  $t$ ,

$$z_k(UC + t) = e^{-ik\alpha} z_k(C) \quad (k \geq 2), \quad s(UC + t) = Us(C) + t.$$

If  $C_t = UC_s + t$ , substituting the coefficient relation in (29.8) recovers  $U$ , with no choice of a branch of an argument. The center identity gives the translation. Nonzero modes two and three exclude any competing proper rotation. Starting with the root image, apply this rule to the parent copy in each tree edge; rotating that edge also provides the child's image in the common

orientation. Induction provides all obstacle orientations. The parent type of either non-tree edge is now available, giving its rotation as well. All templates used in this induction are recovered images, not supplied supports of the unknown table.

Let  $c_a$  and  $L$  temporarily denote the true representative centers and lattice in the resulting common orientation, translating so that  $c_1 = 0$ . The two endpoint centers of  $e = (a, b, \ell_e)$  have displacement

$$d_e = c_b + L\ell_e - c_a.$$

Summing on the tree gives  $p_a = c_a + Lm_a$ . Substitution in (29.10) gives  $v_i = L\eta_{f_i}$  and all claimed identities. Thus the formulas use the unknown realization only to prove correctness; their inputs do not contain  $c_a$ ,  $L$ , or an inter-channel registration.

For stability,  $h \mapsto z_k$  has uniform-norm operator bound one and  $h \mapsto s$  has bound two. Quotients in (29.8) are locally Lipschitz away from their stated zero denominators. For perturbed noncongruent images, replace a computed  $U$  by  $U/|U|$ ; this is smooth near the unit circle. Finite tree propagation preserves a local Lipschitz bound. Support functions of bodies in a common radius bound are uniformly Lipschitz in angle, so the small rotation errors have the same bound in the uniform support norm. The remaining operations are finite sums, multiplication by the fixed  $M^{-1}$  and translations by bounded marked deck vectors. They preserve the asserted local bound. Finally

$$\|\tilde{L}^t \tilde{L} - L^t L\| \leq 2\|L\| \|\tilde{L} - L\| + \|\tilde{L} - L\|^2.$$

This controls the recovered Gram form in exactly the stated image topology.  $\square$

Additional acquisition input. The following application uses Theorem 57.3 of the complete technical manuscript, including its uniform-family assumptions and its physical observation space. In particular its calibration stage records planar endpoint positions in a common frame for the two types of each channel, uses the physical clock, and charges all independent preparations, including failures. These are additional acquisition assumptions, not conclusions of the exact law-valued inverse. The present corollary is not used in the proofs of the geometric reconstruction or its differential.

**Corollary 29.9** (Persistent finite-design physical reconstruction). *Let  $\mathfrak{K}$  be a compact analytic subfamily contained in one of the persistent skeleton neighborhoods of Theorem 29.7. Assume the uniform analytic, chart and physical-record hypotheses of Theorem 57.3, with the marked  $N + 1$ -channel design supplied by that neighborhood. Then the acquisition and uniform reconstruction conclusions of that theorem hold for this design, with one common positive offset and all preparation failures charged. In particular its finite capped policies recover the complete labelled table, including the unknown marked lattice, modulo one common proper Euclidean motion. The minimum even flight number may tend to infinity as in its prescribed-budget protocol.*

*Proof.* We apply Theorem 57.3; the task here is to verify its geometric design hypotheses, not to reconstruct its acquisition estimates from exact injectivity. The persistent skeleton fixes connected incidence, a signature-rigid spanning tree, and two independent marked cycle gains. The transition matches are unique on this properly asymmetric neighborhood. The finite list of channels has fixed gates and positive clearance; the stipulated compact-family smooth bounds give common local laws by Lemma 29.4. The integer gain matrix  $M$  is fixed. Since  $L$  is nonsingular throughout the compact class, the smallest singular value of the recovered holonomy matrix  $LM$  has a positive minimum. No Euclidean lattice vector is supplied by this argument.

Theorem 29.7 gives injectivity on this class. The finite analytic signature and compact-inverse results of Section 27 provide the geometric inverse hypotheses used by the cited acquisition theorem. Its remaining uniform analytic, chart and physical-record hypotheses are precisely

those assumed in the statement. Theorem 57.3 therefore supplies the charged same-final-flight pilot, capped fresh-stage estimator, and increasing-order prescribed-budget construction. Applying those conclusions gives the assertion. The design is fixed before acquisition; no new sensor, uncharged selection stage, or common design across unrelated skeleton neighborhoods is introduced.  $\square$

The existence of a skeleton at each table does not select one common design throughout an arbitrary compact set crossing several such neighborhoods. Nor do  $N + 1$  channels mean finitely many scalar observations: an exact endpoint law is infinite-dimensional. The generic theorem removes the need to posit a special incidence architecture within this marked observation model. It does not replace that model by an ordinary marked length spectrum, erase the endpoint signs, or equate it with the richer position experiment. The local nonlinear relative law and the smooth finite-remainder inverse remain the analytic inputs to this broader consequence.

### 30. FINITE SYMMETRY AND THE PERIODIC RECONSTRUCTION BRANCHES

The local signed inverse and analytic image recovery do not require asymmetry. We separate the finite alternatives caused by rotational symmetry from the continuous ambiguity present for a circle. The same distinction will sharpen the differential theorem in the next section. Throughout this section the selected design is fixed and realizable. All competing realizations use the same signed contact-law model on an admissible collar at the prescribed offsets.

**30.1. Congruence choices and a local angular lift.** For a convex body  $K$  let

$$s(K) = \frac{1}{\pi} \int_0^{2\pi} h_K(\theta) n(\theta) d\theta, \quad z_k(K) = \frac{1}{2\pi} \int_0^{2\pi} h_K(\theta) e^{-ik\theta} d\theta.$$

The center  $s$  is proper-Euclidean equivariant, and  $z_k$ ,  $k \geq 2$ , is invariant under translations. Write  $\text{Sym}^+(K)$  for the group of proper Euclidean isometries preserving  $K$ .

**Lemma 30.1** (Finite congruence alternatives). *If  $K$  is noncircular,  $\text{Sym}^+(K)$  is a finite cyclic group of order*

$$m(K) = \gcd\{k \geq 2 : z_k(K) \neq 0\}.$$

*For two properly congruent copies  $C_s, C_t$  of  $K$ , the set of proper congruences between them has exactly  $m(K)$  members. Their translations are determined by their rotations and the two centers.*

*Proof.* Every symmetry fixes  $s(K)$ . After centering it is a rotation. A centered body is a disk precisely when all harmonics of orders  $k \geq 2$  vanish: Fourier uniqueness then makes its support function constant. Otherwise choose one nonzero  $z_k$ . An allowed rotation  $e^{i\alpha}$  must satisfy  $e^{-ik\alpha} = 1$ , so there are finitely many such rotations. A rotation preserves the support function precisely when this equality holds for every nonzero harmonic. The common solutions are the  $m(K)$ th roots of unity. The gcd of the infinite nonempty set is already the gcd of a finite subset, since successive positive integer gcds eventually stabilize. Given one congruence, all others differ by a symmetry of  $K$ . For each selected rotation  $U$ , its translation is  $s(C_t) - Us(C_s)$ .  $\square$

**Lemma 30.2** (Local alignment without a global phase choice). *Let  $C_s(t), C_t(t)$  be  $C^1$  families of properly congruent support images, with noncircular base images. Suppose their actual proper congruences form a continuous family through the base alignment  $U_0 = e^{i\alpha_0}$ . Then one nonzero base harmonic  $z_k$ ,  $k \geq 2$ , determines a  $C^1$  local branch of their rotation and translation through that alignment. If both image derivatives vanish, the derivatives of that rotation and translation vanish. The assertion does not require the symmetry group to be constant along the family.*

*Proof.* Choose  $k$  with  $z_k(C_s(0)) \neq 0$ ; the corresponding coefficients in both copies remain nonzero. Put  $r(t) = z_k(C_t(t))/z_k(C_s(t))$  and  $r_0 = e^{-ik\alpha_0}$ . In a small disk about 1 use the logarithm with  $\text{Log } 1 = 0$ . Along the actual congruent family,

$$U(t) = U_0 \exp\{-k^{-1} \text{Log}(r(t)/r_0)\}, \quad a(t) = s(C_t(t)) - U(t)s(C_s(t)). \quad (30.1)$$

Indeed, choose the continuous actual angular lift through  $\alpha_0$  and restrict it by  $|\alpha(t) - \alpha_0| < \pi/k$ . Then  $\text{Log}(r(t)/r_0) = -ik(\alpha(t) - \alpha_0)$ , proving (30.1). The displayed formula is  $C^1$ , and

$$\frac{\dot{r}}{r} = -ik\dot{\alpha}, \quad \dot{\alpha} = \frac{i}{k} \frac{\dot{r}}{r}. \quad (30.2)$$

The last expression is real along the congruent family. Stationary support images have zero center and harmonic derivatives, so  $\dot{\alpha} = 0$  and  $\dot{a} = 0$ . Other congruence branches can exist, but are separated from the selected angular lift. The branch followed by the actual family is the one required here; a base symmetry which is broken need not extend to nearby pairs. Symmetry breaking does not affect the nonzero harmonic or this argument. No phase is selected globally from unlabelled noisy copies by (30.1).  $\square$

**30.2. A finite reconstruction of the exact fiber.** Let  $E = \mathcal{T} \cup \{f_1, f_2\}$  be a clear skeleton as in Theorem 29.3. Its graph on obstacle labels is connected, and its two independent cycle gains form a fixed invertible integer matrix  $M$ . In addition, form a graph whose vertices are the channels in  $E$ , connecting two channels whenever they contain an occurrence of the same obstacle label. This graph is connected: a path in the obstacle graph gives a chain of incident channels. Choose a rooted spanning tree  $\mathcal{U}$  in this channel graph. For each link  $\nu$  of  $\mathcal{U}$  fix the particular two occurrences of its common obstacle label, denoted  $a(\nu)$ . These are combinatorial choices, including for loops and multiple edges, not supplied poses.

**Theorem 30.3** (Finite fibers for noncircular obstacles). *Suppose every obstacle of a realizable analytic table  $T_0$  is noncircular. Its gaps and same-type signed laws at one known sufficiently small positive offset per channel determine all realizations of the same marked datum, modulo a common proper Euclidean motion, as a finite set. With the tree  $\mathcal{U}$  just chosen its cardinality is at most*

$$\prod_{\nu \in \mathcal{U}} m(K_{a(\nu)}). \quad (30.3)$$

*The bound counts congruence assignments and need not be attained. When every  $m(K_a) = 1$ , the fiber consists of a single table.*

*Proof.* Theorem 23.1, the smooth signed jet inverse, and analytic continuation recover the entire pair of obstacle images  $C_{e,-}, C_{e,+}$  in every channel frame. This step uses no proper asymmetry. It also determines their centers and finite symmetry orders. Every realization of the datum must have these same images.

Fix the rotation of the root channel as the identity. Propagate rotations  $R_e$  along  $\mathcal{U}$ . At a link matching occurrences  $C_s$  and  $C_t$ , Lemma 30.1 supplies exactly  $m(K_{a(\nu)})$  proper congruences  $C_t = UC_s + a$ . For each choice set  $R_t = R_s U^{-1}$ . Thus  $R_t(C_t - s(C_t)) = R_s(C_s - s(C_s))$ . All assignments are finite and their number is bounded by (30.3). Their translations are not propagated as unknown lattice displacements; only the rotations have been fixed so far.

For a complete assignment, require all centered, rotated occurrences of each obstacle to agree, not merely those on  $\mathcal{U}$ . Assignments failing this test are discarded. Denote the resulting centered

representative images by  $\widehat{K}_a$ . Orient the obstacle tree  $\mathcal{T}$  from a root obstacle and define

$$\begin{aligned} D_e &= R_e\{s(C_{e,+}) - s(C_{e,-})\}, & p_1 &= 0, \\ p_a &= \sum_{e \in [1,a]_{\mathcal{T}}} \epsilon_e D_e, & m_a &= \sum_{e \in [1,a]_{\mathcal{T}}} \epsilon_e \ell_e, \\ v_i &= p_{a_i} + D_{f_i} - p_{b_i}, & \eta_i &= m_{a_i} + \ell_{f_i} - m_{b_i} \quad (i = 1, 2), \\ L &= (v_1 \ v_2)M^{-1}, & M &= (\eta_1 \ \eta_2), \\ c_a &= p_a - Lm_a, & K_a &= \widehat{K}_a + c_a. \end{aligned} \tag{30.4}$$

This uses the actual  $M^{-1}$ , without assuming  $\det M = 1$ . The cochain identities give, for every tree edge and both remaining edges,

$$c_b + L\ell_e - c_a = D_e. \tag{30.5}$$

For tree edges this follows directly by telescoping; for  $f_i$  it is  $L\eta_i = v_i$ . Each assignment therefore produces at most one marked candidate in the root-center and root-direction gauge.

Retain a candidate only when  $\det L > 0$ , the lifted obstacles are pairwise disjoint, and the selected pairs have the stipulated clear contacts and collars. These conditions are part of realization, not consequences of the integer gain matrix. They require only finitely many lattice checks once a nonsingular candidate has been produced: if all representative bodies and selected segments lie in a ball of radius  $R$ , a translated object can meet one of them only for marks bounded by a constant times  $R/s_{\min}(L)$ . Exact shape equality is an equality of the already known analytic images. This description is a finite-branch inverse of exact function-valued data, not a numerical decision algorithm for arbitrary finite precision inputs.

Every actual realization gives one of the congruence assignments and satisfies (30.4), hence occurs among the retained candidates. Conversely, for a retained candidate, centered shape agreement and (30.5) place each entire ordered pair exactly as its recovered channel pair, by one proper motion. Its closest segment, oriented transverse convention and local actions and amplitudes therefore agree with those recovered from the datum. The relative law theorem yields exactly the same normalized conditional laws. Global phase-volume normalization cancels under conditioning; it is not extra input here. All selected design conditions have been tested. Thus every retained candidate is a realization. Duplicate candidates can be identified in the common gauge. Finally, realizability makes the set nonempty, and when all symmetry orders are one the upper bound is one.  $\square$

The theorem removes no discrete ambiguity by an arbitrary phase rule. Nor does it assert global uniqueness for all noncircular tables. Its differential counterpart uses the local branch through the actual base alignment and therefore needs only the absence of a continuous rotational stabilizer.

### 30.3. A circular comparison for the same selected design.

**Proposition 30.4** (A circle and a variable marked Gram form). *For one disk orbit of radius  $1/10$ , let*

$$L_t = \begin{pmatrix} 1 & -\sin t \\ 0 & \cos t \end{pmatrix}, \quad |t| < 1/10.$$

*The two selected channels with gains  $e_1, e_2$  are clear on a common neighborhood and have constant gaps and same-type signed limiting conditional laws at a common sufficiently small positive offset. Nevertheless the marked lattice Gram form has a nonzero derivative at zero, and the variation is not a common Euclidean motion.*

*Proof.* Both columns of  $L_t$  have length one. Thus the two disk pairs have gap  $4/5$  and each ordered pair is constant up to its own proper channel motion. The same-type local actions and amplitudes, and hence their conditional laws, are constant. Their common offset and collars can be chosen after clearance is verified.

Put  $\epsilon = \|L_t - I\| = 2|\sin(t/2)| < 1/10$ . For every nonzero integer mark  $m$ ,  $|L_t m| \geq (1 - \epsilon)|m|$ , so distinct disk centers are separated by at least  $9/10$ . Consider either unit center segment from  $0$  to  $L_t e_i$ . If a third center  $L_t m$  had distance at most  $1/2$  from that segment, then  $|L_t m| \leq 3/2$ , whence  $|m| \leq 5/3$ . Apart from the two incident centers,  $m$  is therefore one of the eight nearest nonzero integer neighbors. At  $t = 0$  its distance from  $[0, e_i]$  is at least one. Perturbing that center and segment decreases the distance by at most  $(|m| + 1)\epsilon \leq (\sqrt{2} + 1)/10 < 1/4$ . This contradicts the assumed distance at most  $1/2$ . Consequently every third disk has distance greater than  $2/5$  from the center segment, and also from its shorter boundary-to-boundary closest segment. Uniformly small gates are admissible.

Finally,

$$L_t^T L_t = \begin{pmatrix} 1 & -\sin t \\ -\sin t & 1 \end{pmatrix}, \quad \dot{L}_0 = \begin{pmatrix} 0 & -1 \\ 0 & 0 \end{pmatrix}, \quad \left. \frac{d}{dt} \right|_0 L_t^T L_t = \begin{pmatrix} 0 & -1 \\ -1 & 0 \end{pmatrix}.$$

A common rotation would preserve the Gram form. Thus this is not such a motion. The example concerns the two specified channels; no finite-horizon condition is asserted.  $\square$

This comparison is outside the noncircular class and does not contradict the exact or differential theorem. It exhibits the geometric role of the hypothesis for this design, not a universal failure of reconstruction from other designs or richer observations.

**30.4. A noncircular fiber with exactly two realizations.** Finite ambiguity in Theorem 30.3 occurs inside its noncircular class. The following example was suggested in the referee-style memorandum [39]. We give the all-lattice clearance and completeness arguments, and then show that one additional marked gap distinguishes the two branches. This last observation is specific to the example; it is not a uniform branch-selection assertion for arbitrary noncircular tables.

**Proposition 30.5** (An exact two-table fiber). *Let  $K$  have support function*

$$h(\theta) = \frac{1}{10} + \frac{1}{1000} \cos(4\theta), \quad a = 2^{-1/2}, \quad L_+ = \begin{pmatrix} 1 & a \\ 0 & a \end{pmatrix}, \quad L_- = \begin{pmatrix} 1 & -a \\ 0 & a \end{pmatrix}. \quad (30.6)$$

*The one-obstacle tables  $T_{\pm} = (K, L_{\pm})$  are admissible analytic periodic dispersing tables. For the two channels with fixed deck marks  $e_1, e_2$ , they have the same gaps and the same two same-type signed conditional endpoint laws at every common sufficiently small positive offset. The fiber of this marked datum, modulo one common proper Euclidean motion, consists of exactly these two tables. They have different marked Gram forms, although the intrinsic data derivative at either table is injective modulo common motion on the family class of Theorem 31.5.*

*Proof.* The support is entire and

$$h + h'' = \frac{1}{10} - \frac{15}{1000} \cos(4\theta) \geq \frac{17}{200} > 0. \quad (30.7)$$

It therefore defines a strictly convex analytic body with positive curvature. Its support center is zero, its proper rotation group is exactly  $C_4$ , and  $K \subset \overline{B}(0, R)$  with  $R = 101/1000$ . Both lattice determinants equal  $a > 0$ . Their Gram eigenvalues are  $1 - a, 1 + a$ , so, for any nonzero integer mark  $m$ ,

$$|L_{\pm} m| \geq \sqrt{1 - a} |m| \geq \sqrt{1 - a} > 2R. \quad (30.8)$$

This proves disjointness of every pair of distinct lifted closures, not just of the selected pairs.

The two selected center displacements have length one. In the first channel their direction is 0, while in the second it is  $\pi/4$  for  $T_+$  and  $3\pi/4$  for  $T_-$ . At all these normals  $h' = 0$ . The supporting points therefore lie on the corresponding center segment, and the parallel supporting lines give the normal closest pair. Strict convexity makes this pair unique. Since  $h(\theta + \pi) = h(\theta)$ , the gaps are

$$g_1 = 1 - 2 \left( \frac{1}{10} + \frac{1}{1000} \right) = \frac{399}{500}, \quad g_2 = 1 - 2 \left( \frac{1}{10} - \frac{1}{1000} \right) = \frac{401}{500}. \quad (30.9)$$

We check clearance against the entire lattice. Let  $v_i = L_{\pm} e_i$ . For a lattice center not on the line  $\mathbb{R}v_i$ , its distance from that line is  $|\det(v_i, L_{\pm} m)|/|v_i|$ , a nonzero integer multiple of  $a$ . A lattice center on that line is an integer multiple of  $v_i$ ; unless it is one of the two incident centers, its distance from  $[0, v_i]$  is at least one. Thus every third obstacle has distance at least  $a - R > 3/5$  from the unit center segment. The actual closest segment is a subset of this segment. Uniformly small contact gates and a common positive-offset interval are consequently available in both tables.

The first ordered obstacle pair agrees in the two tables. Rotation by  $\pi/2$  preserves  $K$  and takes  $L_+ e_2 = (a, a)$  to  $L_- e_2 = (-a, a)$ . Hence it identifies their second ordered pairs as well. These are proper motions, preserving the oriented signed contact convention. The local stationary actions and normalized amplitudes, and therefore the two same-type conditional laws, agree in each channel. The overall phase-volume constant cancels under conditioning; equality of the conditional laws does not assert equality of the successful-preparation frequencies.

To prove that the fiber has no additional members, use the signed smooth inverse and analytic continuation in Theorem 30.3. They fix the complete ordered images in both channel frames for every realization of the datum. Fix the first channel direction and the support center. Its center displacement is then  $e_1$  and its centered body is  $K$ . The four possible relative  $C_4$  rotations of the second pair give exactly

$$v_2 \in \{(a, a), (-a, a), (-a, -a), (a, -a)\}. \quad (30.10)$$

Here the marked gain matrix is the identity, so the cochain uniquely gives  $L = (e_1 \ v_2)$ ; no continuous placement variable is left. The determinants of these four candidates are respectively  $a, a, -a, -a$ . Precisely the first two satisfy the positive-orientation condition, and both have already passed separation and clearance. This proves exact two-element completeness of the fiber. The product bound in (30.3) is four in this example and is strictly larger than the number of realizations.

Finally,

$$L_+^t L_+ = \begin{pmatrix} 1 & a \\ a & 1 \end{pmatrix}, \quad L_-^t L_- = \begin{pmatrix} 1 & -a \\ -a & 1 \end{pmatrix}. \quad (30.11)$$

A common Euclidean motion preserves a marked Gram form. Changing integer basis is not permitted when the deck marks are fixed, so the two candidates are distinct in the announced quotient. The body is noncircular and the selected skeleton persists. The differential assertion follows from Theorem 31.5 at each candidate, with its actual local alignment branch. It does not identify the two distinct global branches.  $\square$

### 30.5. A further marked channel separates the two branches.

**Proposition 30.6** (Separation by one additional gap). *For the two tables in Proposition 30.5, the channel with deck mark  $e_1 + e_2$  is clear in both tables. Its closest gaps satisfy*

$$g_{+, e_1 + e_2} > 1 > g_{-, e_1 + e_2}. \quad (30.12)$$



Thus the original two-channel datum together with this one additional marked gap determines the table uniquely within its original exact fiber. Conditional endpoint laws on the additional channel are not needed for this distinction.

*Proof.* Put  $w_{\pm} = L_{\pm}(e_1 + e_2)$  and

$$\ell_+ = |w_+| = \sqrt{2 + \sqrt{2}}, \quad \ell_- = |w_-| = \sqrt{2 - \sqrt{2}}.$$

Since  $K$  and  $K + w_{\pm}$  lie in radius- $R$  disks about their centers, their distance obeys

$$\ell_{\pm} - 2R \leq g_{\pm, e_1 + e_2} \leq \ell_{\pm} + 2R. \quad (30.13)$$

The upper estimate needs no assumption that the closest points lie on the center line. Here  $\ell_+ - 101/500 > 1$  and  $\ell_- + 101/500 < 1$ , which prove (30.12).

It remains to show that these closest segments are genuine clear channels, rather than merely distances between two bodies. For a lattice center  $L_{\pm}(m, n)$  off the line  $\mathbb{R}w_{\pm}$ ,

$$\text{dist}(L_{\pm}(m, n), \mathbb{R}w_{\pm}) = \frac{|\det(w_{\pm}, L_{\pm}(m, n))|}{\ell_{\pm}} = \frac{a|n - m|}{\ell_{\pm}} \geq \frac{a}{\ell_{\pm}}.$$

A center on the line has  $m = n$ . Excluding  $0, w_{\pm}$ , its distance from  $[0, w_{\pm}]$  is at least  $\ell_{\pm}$ . Hence every third center is separated from that center segment by at least

$$c_* := \min\{\ell_-, \ell_+, a/\ell_-, a/\ell_+\} = a/\ell_+ > 2R.$$

For endpoints  $x \in K$  and  $y \in K + w_{\pm}$ , each point of  $[x, y]$  is within  $R$  of  $[0, w_{\pm}]$ : subtract the corresponding convex combination of  $0, w_{\pm}$  and use the radius bound at both ends. This applies in particular to the closest segment, which need not be a subset of the center segment. A third obstacle lies in the radius- $R$  disk about its center. Its distance from the closest segment is therefore at least  $c_* - 2R > 0$ . Positive curvature and strict separation give the unique normal closest pair and allow uniformly small gates. The added channels are clear, completing the proof.  $\square$

The two propositions separate three different statements. The two original function-valued channel observations have a genuinely nontrivial finite exact fiber. At each member the intrinsic derivative has no non-Euclidean kernel. A third, differently marked distance observation resolves this particular discrete ambiguity. Neither the circular comparison nor these propositions assert a failure of the noncircular finite-fiber theorem, a universal minimal channel count, or a uniform statistical selector on the entire noncircular class.

### 31. INFINITESIMAL RIGIDITY AND FINITE OBSERVABLE COORDINATES

The exact single-offset inverse has a differential consequence which does not follow from injectivity alone. We prove it for variations of entire analytic tables, allowing the contacts, obstacle shapes and marked lattice to vary. The analytic hypothesis is imposed on the parameter derivative as well as on each table. At the end of the section, a finite-dimensional model is introduced only for the scalar-coordinate conclusion.

**31.1. Analytic variations and fixed intrinsic coordinates.** For  $\rho > 0$ , let  $\mathcal{H}_{\rho}$  be the real Banach space of bounded  $2\pi$ -periodic holomorphic functions on  $|\text{Im } z| < \rho$  which are real on the real axis, with the supremum norm. A *common-strip analytic-support family* is a map

$$t \longmapsto (h_{1,t}, \dots, h_{N,t}, L_t) \in \mathcal{H}_{\rho}^N \times \text{GL}^+(2, \mathbb{R}) \quad (31.1)$$

which is  $C^1$  on an open parameter interval and represents admissible tables. Families with a parameter in an open subset of  $\mathbb{R}^p$  are defined in the same way. Positive curvature, separation, and the chosen clear-channel design persist near the base parameter. A dot means differentiation

at that parameter. Cauchy's estimates imply  $C^1$  dependence in each real  $C^m$  support norm and holomorphy of  $\dot{h}_i$  on the strip. Changes to moving contact frames preserve these properties on every smaller strip. No common bound for all orders of differentiation is asserted.

Fix a table in  $\mathcal{C}_N$  and a skeleton  $E = \mathcal{T} \cup \{f_1, f_2\}$  from Theorem 29.3. Freeze all obstacle, channel and deck marks, the transverse signs, and one sufficiently small positive offset  $d_e$  for each edge. The same offset may be used on all edges. In the smoothly moving intrinsic frame of each channel, its contacts have coordinates  $(0, 0)$  and  $(g_e, 0)$  and the transverse variable is  $u$ . Fix contact-centered squares  $Q_{e,b} = (-r_{e,b}, r_{e,b})^2$  whose closures lie in the respective positive-density regions, and nonzero scalar anchors  $a_{e,b} \in (-r_{e,b}, r_{e,b})$ . These choices persist on a parameter neighborhood. Write

$$f_{e,b,t} = \frac{dP_{e,b,t}}{du dv}$$

for the limiting conditional density in this frame. The frame is a coordinate convention for the intrinsic experiment, not an observed inter-channel pose.

A common Euclidean gauge can be fixed by putting the support center of a root obstacle at the origin and the oriented axis of a root channel along the first coordinate axis. The support center is the vector  $s(K) = \pi^{-1} \int_0^{2\pi} h_K(\theta) n(\theta) d\theta$ . The root channel has positive length; its direction is smooth. This gives a local  $C^1$  slice for the three-dimensional proper Euclidean action. When needed, restrict the strip after this change of gauge. All conclusions about zero tangent vectors in this slice mean equality of the support-function and lattice derivatives, not a particular parametrization of the boundary curves.

**Lemma 31.1** (Differentiability on an interior square). *For a family (31.1), each gap, intrinsic contact frame and finite contact jet is  $C^1$ . On every fixed positive-density square,  $t \mapsto f_{e,b,t}$  is  $C^1$  in each prescribed finite  $C^m$  norm, after restricting the parameter neighborhood as necessary. Consequently*

$$\left. \frac{d}{dt} \right|_0 \int \phi dP_{e,b,t} = \int_{Q_{e,b}} \phi(u, v) \dot{f}_{e,b}(u, v) du dv, \quad \phi \in C_c^\infty(Q_{e,b}). \quad (31.2)$$

*The derivatives in this statement are at fixed excess time  $d_e$  in the moving intrinsic transverse coordinates.*

*Proof.* For a facing pair, the two endpoint stationarity equations have an invertible Hessian. In contact transverse coordinates the quadratic part of the distance has Hessian

$$\begin{pmatrix} \kappa_0 + g^{-1} & -g^{-1} \\ -g^{-1} & \kappa_1 + g^{-1} \end{pmatrix}, \quad \det = \kappa_0 \kappa_1 + g^{-1}(\kappa_0 + \kappa_1) > 0.$$

The support parametrization and its finite derivatives are  $C^1$  by Cauchy's estimates. The implicit function theorem therefore gives  $C^1$  contact points, gap and frame; it also gives  $C^1$  local graph functions at every finite smooth order. Clearance fixes the selected facing pair and its gates on this neighborhood.

Apply the weighted half-line construction and the differentiated relative law to this family. The contraction constants are uniform on a sufficiently small neighborhood with positive gap and curvature margins. Differentiating the fixed-point equation solves for the parameter derivative with the same bounded inverse; finite endpoint derivatives use the higher smooth bounds already supplied by (31.1). The summable endpoint and terminal estimates of Lemmas 22.2 and 22.3 justify the corresponding action derivatives. For the amplitude, we make the functional-family derivative explicit. Suppress the endpoint type and write

$$H_t^0 = H_t(0), \quad G_t = (H_t^0)^{-1}, \quad \Delta H_t(u) = H_t(u) - H_t^0, \quad \mathcal{T}_t(u) = G_t \Delta H_t(u).$$

The uniformly positive tridiagonal reference operator is  $C^1$  in operator norm. The resolvent identity gives

$$\dot{G} = -G\dot{H}^0G, \quad \Delta\dot{H} = \dot{H}(u) - \dot{H}^0, \quad \dot{\mathcal{T}} = \dot{G}\Delta H + G\Delta\dot{H}. \quad (31.3)$$

In particular the derivative of the moving reference is not omitted. For each fixed endpoint order  $m$ , locality and the weighted half-line estimates imply, on a uniform small neighborhood,

$$\sum_{i,j \geq 1} |\partial_u^r \partial_t^a (\Delta H_t(u))_{ij}| \leq C_m \sum_{i \geq 1} \rho^i < \infty, \quad 0 \leq r \leq m, \quad a = 0, 1.$$

Only neighboring entries occur. The terms without endpoint decay in  $\dot{H}(u)$  cancel those in  $\dot{H}^0$ ; the remaining terms contain an orbit coordinate or its parameter derivative. Fixed index powers are absorbed in the strict exponential margin. Finite entry sums therefore converge with their derivatives uniformly in trace norm. Thus  $\mathcal{T}_t$  is  $C^1$  in trace norm, in each fixed endpoint norm, and  $\|\mathcal{T}_t\| \leq q < 1$  after the collar is shrunk once.

The edge sum  $U_t$  of (17.6) has the same summability:  $g_t[-\ell_{t,uv}(0,0)] = 1$  identically in  $t$ , so its reference derivative cancels before the edges are summed. In the absolutely convergent trace series (17.7), cyclicity and the trace-ideal inequality give

$$\frac{d}{dt} \operatorname{tr}(\mathcal{T}_t^n) = n \operatorname{tr}(\mathcal{T}_t^{n-1} \dot{\mathcal{T}}_t), \quad |\operatorname{tr}(\mathcal{T}_t^{n-1} \dot{\mathcal{T}}_t)| \leq q^{n-1} \|\dot{\mathcal{T}}_t\|_1.$$

Summing this geometric majorant proves

$$\frac{\dot{B}}{B} = \dot{U} - \operatorname{tr}((I + \mathcal{T})^{-1} \dot{\mathcal{T}}). \quad (31.4)$$

Fixed endpoint differentiation introduces only fixed polynomial factors in the series index and bounded resolvent factors, which remain summable. This proves the amplitude's  $C^1$  dependence at each finite smooth order, directly from the same relative construction. The half-line actions have the  $C^1$  dependence already established above.

It remains to check the normalization, since its integral meets a moving boundary. In a fixed box  $U$  containing the selected supports, put  $A_t(u, v) = S_t(u) + S_t(v)$  and  $w_t(u, v) = B_t(u)B_t(v)$ . The normalizer is

$$Z_t = \int_U w_t(d_e - A_t)_+.$$

The level set  $A_t = d_e$  has two-dimensional measure zero. Indeed  $S'_t$  vanishes only at zero, whereas  $d_e > 0$ ; the gradient of  $A_t$  is nonzero at every point of this level set. The positive-part map is Lipschitz. Its difference quotients are bounded by a common bound for the parameter derivative of  $A_t$ . Dominated convergence gives

$$\dot{Z} = \int_U \{\dot{w}(d_e - A)_+ - w\dot{A} \mathbf{1}_{\{A < d_e\}}\}.$$

The same argument as the parameter varies gives continuity of this derivative: the indicators converge almost everywhere off the limiting level set. The positive lower bound for  $Z_t$  persists. On a positive interior square the cutoff is inactive, so the product and quotient rules prove the asserted  $C^1(C^m)$  statement. Integrating against the fixed compactly supported test function proves (31.2). In particular no derivative of a Dirac mass on the cap boundary is introduced into this argument.  $\square$

### 31.2. Linearization of the signed contact inverse.

**Lemma 31.2** (A zero law derivative has zero contact-jet derivative). *In a common-strip analytic-support family of one channel, suppose  $\dot{g} = 0$  and  $\dot{f}_b = 0$  on the chosen interior square for both same-type laws at their fixed known positive offset. Then every labelled anchored graph coefficient has zero derivative. This holds for odd as well as even orders and does not assume a known amplitude.*

*Proof.* We display the differential of the amplitude-free inverse. For one type omit  $b$ , put  $q = \sqrt{1 - R_f(a, a)} > 0$ , and write  $t(u) = (1 - R_f(u, a))/q$ . At a fixed positive density,

$$\dot{R}_f(u, v) = R_f(u, v) \left\{ \frac{\dot{f}(u, v)}{f(u, v)} + \frac{\dot{f}(0, 0)}{f(0, 0)} - \frac{\dot{f}(u, 0)}{f(u, 0)} - \frac{\dot{f}(0, v)}{f(0, v)} \right\}, \quad (31.5)$$

$$\dot{t}(u) = -\frac{\dot{R}_f(u, a)}{q} + \frac{t(u)\dot{R}_f(a, a)}{2q^2}, \quad \dot{S}(u) = \frac{d_e \dot{t}(u)}{(1 + t(u))^2}. \quad (31.6)$$

All these operations are  $C^1$  in any prescribed finite  $C^m$  norm: evaluation on a fixed slice is bounded linear, the denominators are positive, and the sole square root is a scalar evaluated at the nonzero anchor. Lemma 31.1 and Theorem 23.1 allow the ordinary chain rule. Thus  $\dot{f} = 0$  implies  $\dot{S} = 0$ , including every finite derivative at zero. The two signed coordinates have not been identified.

Write  $a_b = S_b''(0)$ . The explicit leading recovery is

$$\gamma = \operatorname{arsinh}(g\sqrt{a_0 a_1}), \quad c_b = \cosh \gamma \sqrt{a_b/a_{1-b}}, \quad \kappa_b = (c_b - 1)/g.$$

It follows that  $\kappa_b = 0$ . At order  $n \geq 3$ , use the actual finite-jet identity

$$s_n = M_n q_n + R_n,$$

where  $R_n$  depends only on lower graph jets and the gap, and  $M_n$  is the invertible block in (22.4). Its validity for smooth representatives, including dependence on finite rather than formal jets, is supplied by Lemma 22.4. Each finite map is smooth, as also follows from the finite Bellman recursion. Differentiating gives

$$\dot{q}_n = M_n^{-1} \{ \dot{s}_n - \dot{R}_n - \dot{M}_n q_n \}. \quad (31.7)$$

Induction, beginning with the leading recovery, makes every term on the right zero. The orders zero and one are fixed by the contact anchors. The induction is carried out separately at each finite  $n$ ; it neither differentiates an infinite inverse series nor uses an order-independent bound on  $M_n^{-1}$ .  $\square$

**Lemma 31.3** (Analytic propagation of a stationary contact variation). *If all anchored contact-graph coefficients of a common-strip analytic-support family have zero derivative and  $\dot{g} = 0$ , the entire incident obstacle images are stationary to first order in that channel frame. Their support functions therefore have zero derivative in every real  $C^m$  norm.*

*Proof.* Express the two bodies in the same moving channel frame. The contact origins are 0 and  $(g, 0)$ , whose derivatives vanish under the hypothesis. The normal directions at the contacts are fixed, and any reversal needed to use the positively oriented tangent is fixed as well. The local graph-to-support conversion is a smooth triangular map at each finite order, with positive curvature as its nondegeneracy condition; see (50.12). Hence every derivative in normal angle of the support variation vanishes at the contact normal.

For clarity, the analyticity needed at this step concerns the variation itself. If  $A_t x = R_{\alpha_t} x + a_t$  is the placement of a channel frame, the support of a body in that frame is obtained from

its global support by an angular translation and a linear first-harmonic term. Its parameter derivative is a sum of the translated  $\dot{h}$ , a multiple of the translated  $h'$ , and first harmonics. Each is holomorphic on any fixed smaller strip. This follows from (31.1), Cauchy's estimates and the  $C^1$  frame dependence, not from pointwise analyticity of  $h_t$  alone. The identity theorem now makes the support variation zero on the connected strip, hence on the entire real normal circle. Its real  $C^m$  norm is zero for each  $m$ . No norm-continuous inverse for analytic continuation is used or concluded.  $\square$

### 31.3. Differentiable incidence matching and the lattice.

**Lemma 31.4** (Finite harmonic registration on the asymmetric locus). *For each obstacle with trivial proper Euclidean symmetry, choose a finite set  $F \subset \{2, 3, \dots\}$  of nonzero support harmonics whose indices have greatest common divisor one. These choices persist locally. Incidence congruences, the rotations of all selected channel frames into one common frame, the lattice  $L$ , and representative obstacle placements are  $C^1$  functions of the complete channel-frame images along common-strip analytic-support families. If all these images are stationary to first order, the reconstructed gauge-fixed table is stationary to first order. For noncircular obstacles, the same conclusions hold on a local alignment branch through the actual base realization, without asserting globally single-valued registration.*

*Proof.* The finite witness exists by Proposition 29.6: otherwise the nonzero centered support harmonics have a common divisor larger than one, or there are none, giving respectively a nontrivial rotation or a circle. The greatest common divisors of successive finite subsets stabilize, so a finite witness suffices when the full divisor is one. Nonzero coefficients remain nonzero on a neighborhood.

Choose integers  $a_k$  with  $\sum_{k \in F} a_k k = 1$ . For two copies  $C_s, C_t$  of this obstacle, with  $C_t = UC_s + a$  and  $U = e^{i\alpha}$ , define

$$r_k = \frac{z_k(C_t)}{z_k(C_s)}, \quad U = \prod_{k \in F} r_k^{-a_k}, \quad a = s(C_t) - Us(C_s), \quad z_k(C) = \frac{1}{2\pi} \int_0^{2\pi} h_C(\theta) e^{-ik\theta} d\theta. \quad (31.8)$$

Indeed  $r_k = e^{-ika}$ , so the product is  $e^{ia}$ . Only integer powers of nonzero scalars occur; no argument or root branch is chosen. Centers and coefficients are bounded linear functionals of the real support function. Formula (31.8) is consequently  $C^1$  on its nonvanishing domain. It recovers the unique proper incidence congruence. For  $F = \{2, 3\}$  it reduces to (29.8). Thus the present argument does not restrict the asymmetric locus to the second-and-third-harmonic subclass.

Propagate these congruences along the finite spanning tree from a fixed root incidence. Match the initial obstacle of either remaining edge to get its rotation as well. This is a finite composition of  $C^1$  operations. In the resulting orientation let  $R_e$  be the rotation of edge  $e$  and set

$$d_e^{\text{cen}} = R_e\{s(C_{e,+}) - s(C_{e,-})\}, \quad p_1 = 0, \quad p_a = \sum_{e \in [1,a]_{\mathcal{T}}} \epsilon_e d_e^{\text{cen}}.$$

Here  $m_a$  denotes the sum of signed deck marks along the same tree path. For  $f_i = (a_i, b_i, \ell_i)$  the marked gain is  $\eta_i = m_{a_i} + \ell_i - m_{b_i}$  and

$$v_i = p_{a_i} + d_{f_i}^{\text{cen}} - p_{b_i}, \quad L = (v_1 \ v_2) M^{-1}, \quad c_a = p_a - L m_a, \quad M = (\eta_1 \ \eta_2). \quad (31.9)$$

These are the displacement-cochain identities of Proposition 29.8, whose telescoping proof depends on the recovered rotations, not on using harmonics two and three in particular. The matrix  $M$  is fixed and invertible; it need not be unimodular. Centering and rotating one recovered copy of each obstacle, then translating it to  $c_a$ , completes the reconstruction in the root-center gauge.

All operations are  $C^1$  along the family. In particular

$$\dot{L} = (\dot{v}_1 \ \dot{v}_2)M^{-1}.$$

If every input support variation is zero, every coefficient and center derivative, every relative-rotation derivative, and then every quantity in (31.9) has zero derivative. The output support variations and lattice variation are zero. This proves the last assertion without assuming a differentiable inverse to analytic continuation.

Local registration at a noncircular table. When the base obstacle has a nontrivial finite symmetry group, the Bezout formula in Lemma 31.4 need not exist. Use instead the local angular lift of Lemma 30.2, separately at each link of a fixed channel-incidence spanning tree. Such branches exist because the actual continuously moving channel frames give congruences between each fixed pair of labelled occurrences. Every noncircular base obstacle has one nonzero harmonic, so these local choices persist along a  $C^1$  family, whether or not its finite symmetry is preserved. Stationary channel-frame support images give stationary relative rotations and translations on these branches. All subsequent center and cochain operations in (31.9) are the same finite  $C^1$  operations. Hence stationary images again give  $\dot{L} = 0$  and stationary representative supports in the common gauge. The actual base placements are used only to specify the local branch for this derivative argument, not supplied as extra observations. The finite alternatives in Theorem 30.3 are not removed globally by this local choice.  $\square$

**Theorem 31.5** (Infinitesimal single-offset periodic rigidity). *Let  $T_t$  be a common-strip analytic-support family through  $T_0 \in \mathcal{C}_N$ , with a persistent selected skeleton and fixed positive offsets as above. The following are equivalent:*

- (1)  $\dot{g}_e = 0$  and  $\dot{f}_{e,b} = 0$  on  $Q_{e,b}$  for every  $e, b$ ;
- (2)  $\dot{g}_e = 0$  and  $\left. \frac{d}{dt} \right|_0 \int \phi dP_{e,b,t} = 0$  for every  $e, b$  and every real  $\phi \in C_c^\infty(Q_{e,b})$ ;
- (3) the table derivative is generated by one common infinitesimal proper Euclidean motion.

*In particular the data derivative is injective in the root-center and root-direction gauge, including variations of the marked lattice.*

*Proof.* The first two assertions are equivalent by Lemma 31.1. For the nontrivial direction, a continuous nonzero  $\dot{f}$  has constant strict sign on some small open ball. A nonnegative smooth function supported in that ball has nonzero integral against  $\dot{f}$ , contradicting the second assertion.

Assume the first assertion. Lemma 31.2 makes all contact-jet variations zero. Lemma 31.3 makes all entire channel-frame support variations zero. Lemma 31.4 then makes the gauge-fixed table variation zero. The  $C^1$  gauge change therefore expresses the original derivative as a common infinitesimal translation and rotation. Explicitly, for constants  $a \in \mathbb{R}^2$  and  $\omega \in \mathbb{R}$ , it has the support and lattice form

$$\dot{h}_i(\theta) = a \cdot n(\theta) - \omega h'_i(\theta), \quad \dot{L} = \omega J L, \quad J = \begin{pmatrix} 0 & -1 \\ 1 & 0 \end{pmatrix}. \quad (31.10)$$

The converse follows from invariance of gaps and signed intrinsic conditional laws under a simultaneous proper Euclidean motion. All parameter derivatives are ordinary derivatives along the stipulated family. Exact injectivity of the undifferentiated map is not used as a substitute for any step of the kernel argument.  $\square$

#### 31.4. Finite scalar coordinates on immersed models.

**Theorem 31.6** (Local coordinates from finitely many gap and law tests). *Let  $\Theta \subset \mathbb{R}^p$  be open,  $p \geq 1$ , and let  $\theta \mapsto T_\theta$  be a  $C^1$  common-strip analytic-support family in the fixed Euclidean gauge. Suppose its table derivative at  $\theta_0$  is injective and  $T_{\theta_0} \in \mathcal{C}_N$ . Use a persistent skeleton and fixed*

positive offsets. Then there are  $p$  scalar observables  $\lambda_1, \dots, \lambda_p$ , each either  $g_e$  or  $\int \phi dP_{e,b}$  for a real  $\phi \in C_c^\infty(Q_{e,b})$ , such that

$$F(\theta) = (\lambda_1(T_\theta), \dots, \lambda_p(T_\theta)) \quad (31.11)$$

is a  $C^1$  diffeomorphism between neighborhoods of  $\theta_0$  and  $F(\theta_0)$ . On a sufficiently small convex parameter ball there are constants  $c, C > 0$  such that

$$c|\theta - \tilde{\theta}| \leq |F(\theta) - F(\tilde{\theta})| \leq C|\theta - \tilde{\theta}|. \quad (31.12)$$

For each fixed  $m$ , after restricting this ball if necessary, the  $C^m$  differences of the gauge-fixed support functions and the matrix difference of  $L$  are bounded by a constant times  $|F(\theta) - F(\tilde{\theta})|$ .

*Proof.* Consider on  $\mathbb{R}^p$  the linear functionals

$$v \mapsto Dg_e(\theta_0)v, \quad v \mapsto \int_{Q_{e,b}} \phi Df_{e,b}(\theta_0)v, \quad \phi \in C_c^\infty(Q_{e,b}).$$

Their common kernel is zero. Indeed a vector in that kernel gives a table variation with vanishing intrinsic derivative by Theorem 31.5; the gauge and the immersion hypothesis then make the vector zero. A set of linear functionals with zero common kernel spans the dual of a finite-dimensional vector space: a proper span would have a nonzero annihilator. Select a basis of  $p$  members from this set. Each selected member is the derivative of an allowed scalar observable, not a derivative evaluation of an unknown density. Lemma 31.1 makes the resulting map  $F$  continuously differentiable. Its derivative  $A = DF(\theta_0)$  is invertible, so the inverse function theorem proves the local coordinate assertion.

There is also a direct bound avoiding any assertion that pointwise nonsingularity alone implies a lower Lipschitz estimate. On a small convex ball, continuity gives  $\|DF(\theta) - A\| \leq (2\|A^{-1}\|)^{-1}$ . For  $z = \theta - \tilde{\theta}$ , integration along the joining segment therefore yields

$$|F(\theta) - F(\tilde{\theta})| \geq |Az| - \frac{|z|}{2\|A^{-1}\|} \geq \frac{|z|}{2\|A^{-1}\|}.$$

The upper bound follows from the local bound for  $DF$ . Finally the  $C^1(\mathcal{H}_\rho)$  family has locally bounded first parameter derivatives in every fixed real  $C^m$  support norm, by Cauchy's estimates on a smaller strip, and bounded lattice derivatives. Integrating these derivatives on the same ball and using the lower bound proves the last assertion.  $\square$

Theorem 31.6 is a model-local coordinate result for exact expectations. It does not claim a fixed finite observation vector for all analytic tables, a dimension bound for an arbitrary compact analytic class, or a statistical efficiency theorem. The full-class finite-histogram reconstruction in the complete technical manuscript uses increasing resolution and its separately stated conditioning hypotheses. The fixed-contact variation bundle and uniform positive designs in Section 58 concern a different question: realizing compatible finite jets and detecting boundary velocities with a finite collection of offsets. Here the gaps, contacts and lattice may move, and a single offset per channel suffices for the differential law inverse. The analytic identity theorem removes the shape kernel; finite harmonic registration removes the relative-frame and lattice kernel; finite-dimensionality is used only in the final selection of scalar tests.

## Part 4. Boundary information and physical local experiments

### 32. OBSERVATION SPACES AND COARSENING

**32.1. From relative laws to physical observation protocols.** The relative comparison retains the probability of obtaining a selected record. In Theorem 34.1, the finite and boundary subprobability laws are compared at each table in common endpoint and residual-time coordinates,

with their missing-event atom included. In unscaled physical positions the same boundary embeddings are used for the two laws at that table. This proves a bound between statistical families; it does not supply a sampling kernel that knows the unknown table's contact center or tangent.

For predetermined physical windows  $t_i$ , the excess  $d_i(T) = t_i - j_i g_e(T)$  is used only in the proof. The acquisition need not observe the gap. For adaptive designs and stopped procedures, Theorem 35.3 compares the complete declared transcript, with a bound of the form

$$\|\mathbf{P}_T^\pi - \mathbf{Q}_T^\pi\|_{\text{TV}} \leq C \mathbb{E}_{\mathbf{P}_T^\pi} \sum_{i=1}^{\mathcal{T}} Y_i \tau^{j(a_i)}.$$

Here  $\mathcal{T}$  is the stopping time,  $a_i$  the chosen design and  $Y_i$  its success indicator. A deterministically capped batch stopping no later than its  $k$ -th success has error at most  $Ck\tau^J$  when all its flight numbers are at least  $J$ . No failed preparation is removed from the cost.

There are two different realizations of endpoint coordinates in this article. The local information theorems in Section 39 concern an anchored registered family: the reported scalar transverse coordinates are literal laboratory coordinates fixed throughout that family, and the programmed windows are  $t_{j,\ell} = jg_0 + d_\ell$ . The local gap parameter is part of the likelihood, not supplied to the acquisition. The following information hierarchy concerns this declared scalar local record and its coarsenings. It does not assert that the local Poisson or Gaussian limit exhausts all information in a two-dimensional position sensor.

The global unknown-contact experiment is specified separately in Section 56. Each channel sensor records both components of its first and last endpoint positions before any transverse projection. A single sensor frame is shared by the two types of that channel, while different channel frames are not registered. The charged pilot of Theorem 56.2 estimates the contact positions, tangent and onset. The subsequent transverse normalization is then a measurable function of those estimates and a fresh physical record. Proposition 56.3 bounds its error for the finite bounded Lipschitz tests used by the global estimator. The parameter-dependent ideal normalization appears only in this error bound, not among the observations.

Thus the same-table physical-to-boundary comparison, the anchored local information experiment, the observable within-channel calibration, and the intrinsic inter-channel gluing have different roles. Their hypotheses are not exchanged. The earlier centered-window and fixed-table coarsening results remain in the auxiliary compendium with their stated preparation costs and calibration assumptions.

**32.2. Observation hierarchy and the information retained by each coarsening.** The same physical preparation generates several statistically distinct experiments. To compare their sigma-fields, fix a parameter-independent policy  $\pi$  with a deterministic preparation cap  $N$ , as in Section 35, and start with its complete stopped physical history. This history retains every chosen design, all collision points, labels and flight times within every preparation, the success/failure sequence, and the stopping time  $T \leq N$ . Independent policy randomness may also be recorded when it belongs to the specified experiment. Full collision histories are encoded in the common physical coordinates as finite marked arrays. Their countable disjoint union, followed by the finite products and disjoint unions imposed by  $T \leq N$ , gives one standard Borel stopped-history space. Geometry changes its law, not its underlying measurable space. No total-variation replacement theorem for this full growing collision array is asserted here.

The per-preparation endpoint-time map is a building block, not a separate stopped experiment. On the prescribed bridge event it retains the first and last signed endpoints and the residual preparation-to-first-impact time; on failure it retains a failure atom. This is the one-record map



for which Theorem 34.1 proves the finite-to-boundary comparison. Applying it separately to the preparations in the stopped history defines the following hierarchy on a single sample space.

- (1) *Full stopped collision history.* The complete history just described retains all within-preparation collision data, designs, failures and the stopping time.
- (2) *Stopped endpoint-time acquisition transcript.* Replace each within-preparation history by its endpoint-time record, retaining every design, every failure, and  $T$ . Recorded policy randomness is retained consistently. Theorem 35.3 compares the laws of this transcript, with every failed preparation charged.
- (3) *Count-endpoint transcript.* Delete residual times from item 2, but retain the complete design list, the success/failure sequence, the successful endpoint pairs, and the stopping time. The failure sequence may equivalently be encoded by waiting counts when the final censored failure run and the stopping information are also retained.
- (4) *Endpoint-output experiment.* Delete waiting counts and residual times, retaining the successful endpoint pairs required by the specified output rule and its cap cemetery symbol. Any output labels are precisely those specified by that rule, rather than an unannounced extra transcript.

Let  $H^\pi$ ,  $T_{\text{ET}}^\pi$ ,  $T_{\text{CE}}^\pi$ , and  $T_{\text{out}}^\pi$  denote these four records. Their deterministic forgetting maps give, after pullback to the full stopped-history space,

$$\sigma(T_{\text{out}}^\pi) \subseteq \sigma(T_{\text{CE}}^\pi) \subseteq \sigma(T_{\text{ET}}^\pi) \subseteq \sigma(H^\pi). \quad (32.1)$$

Common randomized output maps are Markov coarsenings as well; their independent seeds can be adjoined when a sigma-field formulation is needed. The coarsened record is always the output of the same acquisition policy. It is not the history of a different policy obtained by retrospectively withholding observations from the design rule. In particular, retained adaptive designs can encode part of a deleted record. The asymptotic identifications below are those of the protocols and hypotheses in the cited theorems, not a claim that deleting times from every adaptive policy produces the same limiting experiment.

These common Markov coarsenings produce genuinely different asymptotic experiments. At the fastest local geometric scale, the endpoint-time protocol of Theorem 41.3 converges to its Poisson boundary-shift experiment; the positive density at the moving residual-time ceiling yields shape resolution of order  $k^{-1}$ . At that scale the waiting counts are asymptotically ancillary under  $j = o(\sqrt{k})$ , with the other hypotheses of that theorem in force.

After residual time is removed, the ceiling density is integrated into a linearly vanishing endpoint density. The count channel then matters. In the multirate coordinate of Theorem 44.3, the waiting counts yield a Gaussian shift in the hyperbolic-exponent direction at shape scale  $(j\sqrt{k})^{-1}$ , while the iso-hyperbolic endpoint directions are visible at the slower boundary scale  $k\delta^2 \log(1/\delta) \asymp 1$ . Removing the waiting counts in that protocol produces the endpoint-output experiment and its endpoint Gaussian limit.

This hierarchy is consistent with the earlier fixed-table comparison. If  $q_j = e^{-j\gamma}$ , the endpoint-residual record and the endpoint-only record have different critical scales: the former reacts at order  $kq_j$ , whereas the latter reacts at order  $kq_j^2 \log(1/q_j)$ . Thus a weak one-record limit does not determine the information scale, and integrating out residual time is not an innocuous presentation choice.

The distinction between statistical information and physical cost is also maintained. For a selected bridge, success has probability of order  $d^2 e^{-j\gamma}$ , so obtaining  $k$  successful records requires an exponentially larger number of preparations. Failed preparations remain in the stopped transcript whenever that transcript is the experiment under study. Conditioning on success is used only when a theorem explicitly studies the successful mark law.

Finally, the record hierarchy should not be confused with geometric registration. The local statistical realization in Section 39 uses one fixed laboratory chart so that all nearby parameters live on a common observation space. The global rigidity theorem of Section 25, by contrast, does not supply relative Euclidean placements between distinct measured channels; those placements are reconstructed through analytic gluing.

A common choice of orientation is a third, separate issue. The unoriented classification uses the orbit of the entire marked signed law collection under one simultaneous reversal, as specified in Definition 25.12. The reversal transports the lattice orientation convention and all channel signs together; it does not alter deck labels or forget the relative sign choices. This law-valued quotient is not another item in the above hierarchy of Markov coarsenings. In particular, replacing each endpoint pair by its absolute values, averaging a law with its reflection, or independently discarding the orientation of each channel gives different data. The signed inverse is not invoked as an inverse of any of those coarsenings. The precise common-orientation classification and its proof are given in Theorem 25.14.

### 33. SMALL-WINDOW INFORMATION ALONG REALIZABLE CHANNELS

The window inverse and the physical sampling problem use the same boundary measure, but they do not preserve the same information. We make this distinction quantitative within a family of actual periodic tables. The small-window calculation for quadratic functional actions was communicated in [43]. The argument below needs only the leading Taylor coefficient of an actual action. It realizes the effect without requiring an exactly quadratic billiard action, gives a matching two-point testing scale, and transfers that experiment to sufficiently long finite bridges. The recording profiles are fixed as the window shrinks. They may differ between the two alternatives, as permitted by the nuisance model of Theorem 24.6.

**33.1. A normalized window calculation.** Fix  $d > 0$ . Let  $S_i$ ,  $i \in \{0, 1\}$ , be even  $C^4$  functions near zero, with  $S_i(0) = S'_i(0) = 0$ ,  $S''_i > 0$ , and put  $a_i = S''_i(0)$ . For a sufficiently small  $h > 0$  take the window  $(-h, h)^2$  and the separate effective factors  $(d - S_i(u))^{-1}$  and  $(d - S_i(v))^{-1}$ . Positive constants multiplying either factor do not change the conditional law. After the known rescaling  $u = hz$ ,  $v = hw$ , write  $Q_{i,h}$  for that law on  $[-1, 1]^2$ , ignoring edges of measure zero. Define

$$t_i(u) = \frac{S_i(u)}{d - S_i(u)}, \quad I_i(h) = \int_{-1}^1 t_i(hz) \, dz, \quad q_{i,h}(z, w) = \frac{1 - t_i(hz)t_i(hw)}{4 - I_i(h)^2}. \quad (33.1)$$

Indeed the unnormalized window density is  $d^{-1}(1 - t_i(u)t_i(v))$ . Its integral includes the term  $I_i(h)^2$ ; discarding this term changes the leading information coefficient. Throughout this section the squared Hellinger distance is

$$H^2(P, Q) = \int (\sqrt{p} - \sqrt{q})^2, \quad \text{TV}(P, Q) = \frac{1}{2} \int |p - q|.$$

**Lemma 33.1** (Window expansion and quadratic information loss). *Put  $F(z, w) = 1/9 - z^2w^2$ . Uniformly on the rescaled square,*

$$q_{i,h} = \frac{1}{4} + \frac{a_i^2 h^4}{16d^2} F + O(h^6). \quad (33.2)$$

If  $a_0 \neq a_1$ , then

$$\text{TV}(Q_{0,h}, Q_{1,h}) = \Theta(h^4), \quad (33.3)$$

$$H^2(Q_{0,h}, Q_{1,h}) = \frac{7(a_1^2 - a_0^2)^2}{16200d^4} h^8 + O(h^{10}), \quad (33.4)$$

$$\mathbb{E}_{Q_{i,h}} F = \frac{14a_i^2}{2025d^2} h^4 + O(h^6), \quad \text{Var}_{Q_{i,h}} F = \frac{56}{2025} + O(h^4). \quad (33.5)$$

All remainder bounds are uniform on families with a common  $C^4$  bound, positive  $a_i$ , and a common positive window margin.

*Proof.* Evenness and Taylor's theorem give  $S_i(hz) = a_i h^2 z^2 / 2 + O(h^4)$ , uniformly for  $|z| \leq 1$ . Expansion of the reciprocal on a fixed positive neighborhood gives  $t_i(hz) = a_i h^2 z^2 / (2d) + O(h^4)$  and  $I_i(h) = a_i h^2 / (3d) + O(h^4)$ . Expanding the numerator and denominator of (33.1) proves (33.2). The square densities are bounded away from zero for small  $h$ . Consequently

$$\sqrt{q_{1,h}} - \sqrt{q_{0,h}} = \frac{q_{1,h} - q_{0,h}}{\sqrt{q_{1,h}} + \sqrt{q_{0,h}}} = \frac{(a_1^2 - a_0^2)h^4}{16d^2} F + O(h^6).$$

Here the limiting denominator is one. Direct integration gives

$$\int_{[-1,1]^2} F = 0, \quad \int_{[-1,1]^2} F^2 = 4 \left( \frac{1}{25} - \frac{1}{81} \right) = \frac{224}{2025}. \quad (33.6)$$

Squaring and integrating proves (33.4). The nonzero continuous function  $F$  also has a positive integral of its absolute value; this proves (33.3). Integrating  $Fq_{i,h}$  and  $F^2q_{i,h}$  proves (33.5). The stated uniformity follows from the common Taylor and reciprocal bounds. No higher action coefficient is set equal to zero in this proof.  $\square$

**33.2. Actual tables and the attainable discrimination scale.** The model below fixes the lattice and all channel gaps. The only changing geometric quantity is the obstacle shape. Thus a gap observation or a known origin cannot eliminate the information loss by itself.

**Theorem 33.2** (A realizable small-window experiment). *There is a common-strip analytic-support family of noncircular periodic tables with two fixed marked clear channels, the same gaps, and fixed signed contact origins, with the following properties. For each table there are positive smooth separate endpoint recording probabilities, independent of  $h$ , such that all four same-type limiting window laws are  $Q_{s,h}$  of (33.1), with the actual channel action  $S_s$ . Every fixed  $h > 0$  sufficiently small retains the exact geometric fiber of Corollary 24.7.*

More explicitly, let  $R = 1/10$ , let  $s \in [R/128, R/64]$ , put  $L = I_2$ , and let the single obstacle orbit have support function

$$k_s(\theta) = R + s(\cos 4\theta - 1). \quad (33.7)$$

The channels have deck marks  $e_1, e_2$  and gap  $g = 4/5$ . The two types and the two channels have the same even action, with

$$\kappa_s = (R - 16s)^{-1}, \quad a_s = S_s''(0) = \sqrt{\kappa_s^2 + 2\kappa_s/g}. \quad (33.8)$$

For the distinct alternatives  $s_0 = R/128$ ,  $s_1 = R/64$  this gives

$$a_0^2 = 7800/49, \quad a_1^2 = 1900/9, \quad \Delta = a_1^2 - a_0^2 = 22900/441 > 0.$$

Fix one common sufficiently small positive  $d$ . Let  $n = n_h$  independent successful window records be allocated in any fixed way among the four same-type laws. The gaps and channel/type labels may also be retained. Define

$$\mathcal{J} = \frac{7\Delta^2}{4050d^4} > 0. \quad (33.9)$$

For the two specified table/profile alternatives the following hold:

- (1) If  $nh^8 \rightarrow 0$ , their product total-variation distance tends to zero, and the minimum equal-prior error tends to  $1/2$ .
- (2) If  $nh^8 \rightarrow \lambda \in (0, \infty)$ , the minimum equal-prior error converges to  $\Phi(-\sqrt{\lambda\mathcal{J}}/2)$ , where  $\Phi$  is the standard normal distribution function.
- (3) If  $nh^8 \rightarrow \infty$ , a test of the empirical mean of  $F$  has error at most  $\exp(-cnh^8)$  for some  $c > 0$  and all sufficiently small  $h$ . The same mean test attains the limit in the second assertion when its threshold is the midpoint of its two exact means.

Thus  $h^{-8}$  is the attainable successful-record scale for this pair. The upper assertion is a two-point assertion with specified alternatives, not a uniform estimator over unrestricted recording profiles.

*Proof. Geometry and actual actions.* The function in (33.7) is entire and real on the real axis, and

$$k_s + k_s'' = R - s - 15s \cos 4\theta \geq R - 16s \geq 3R/4 > 0.$$

It is the support of a strictly convex noncircular body  $K_s$ , centered at zero and contained in the disk of radius  $R$ , since  $k_s \leq R$ . Distinct  $s$  give noncongruent bodies: their perimeters are  $2\pi(R - s)$ , by integrating the radius of curvature. Nonzero integer center displacements have length at least one, so lifted closures are disjoint. At the horizontal and vertical normals,  $k_s = R$  and  $k_s' = 0$ . The supporting endpoints therefore give the unique closest pairs in  $K_s, K_s + e_i$ , with gap  $1 - 2R = 4/5$ . Every other integer center is at distance at least one from the associated unit axis segment. Every third body is at distance at least  $1 - R$  from that segment, and the closest-contact segment is a subset. Common small collars are clear. No finite-horizon hypothesis is needed for these selected channels.

The quarter-turn symmetry makes the two channels congruent. Reflection in a channel axis sends every transverse coordinate to its negative; the local flight action is invariant under this simultaneous sign change. Uniqueness of the small half-line stationary sequence in Lemma 17.1 therefore gives  $S_s(-u) = S_s(u)$ . The reflection exchanging the two identical obstacles similarly identifies the two types and their amplitudes. These reflection arguments establish symmetry of the functions of a fixed channel; they do not quotient an unknown table by an improper Euclidean motion. The family has a uniform analytic collar and fixed finite smooth bounds by the same half-line lemma. With equal contact curvatures,  $c_0 = c_1 = 1 + g\kappa_s$  and  $\cosh \gamma_s = 1 + g\kappa_s$ ; substituting in (17.4) gives (33.8). This obtains the Taylor coefficient from an actual billiard, without assuming its higher coefficients vanish.

*Recording.* Choose  $d$  in the common positive-offset range and then  $h_0 > 0$  inside the common collar with  $\max_{|u| \leq 2h_0} S_s(u) < d/4$  uniformly in  $s$ . Let  $B_s$  be the positive actual half-line amplitude. On  $[-h_0, h_0]$  choose both recording probabilities to be

$$e_s(u) = \frac{c_*}{B_s(u)(d - S_s(u))}. \quad (33.10)$$

A sufficiently small  $c_* > 0$  works simultaneously over the compact family. For an explicit extension, take a smooth cutoff  $\chi$  equal to one on  $[-h_0, h_0]$  and supported in  $(-2h_0, 2h_0)$ , and set  $e_s = c_* \exp[-\chi \log(B_s(d - S_s))]$  on the gate interval, with the cutoff product extended by zero. The bounded exponent and a smaller  $c_*$  ensure  $0 < e_s \leq 1$  everywhere. These profiles do not depend on  $h$ . Independent recording of the two endpoints cancels  $B_s(u)B_s(v)$  and leaves exactly  $c_*^2 d^{-1}(1 - t_s(u)t_s(v))$  on  $(-h, h)^2$ . Conditioning and rescaling yield (33.1). The fixed origins are zero, the cap is strictly positive on the window, and its critical lines remain inside. The window theorem therefore applies for each  $h > 0$ . The profiles are admissible nuisance alternatives, not profiles estimated or programmed from an unknown table by this theorem.

*Testing.* Symmetry makes every label's successful law the same  $Q_{i,h}$  in these coordinates; fixed label allocations add no distinction between the alternatives. Write  $p_h = q_{0,h}$ ,  $q_h = q_{1,h}$  and  $r_h = (q_h - p_h)/p_h$ . Lemma 33.1 gives

$$\|r_h\|_\infty = O(h^4), \quad \mathbb{E}_{p_h} r_h = 0, \quad \mathbb{E}_{p_h} r_h^2 = \mathcal{J}h^8 + O(h^{10}). \quad (33.11)$$

For  $nh^8 \rightarrow 0$ , affinity tensorization gives  $H^2(p_h^{\otimes n}, q_h^{\otimes n}) \leq nH^2(p_h, q_h)$ ; Cauchy–Schwarz gives  $\text{TV} \leq H$  in our convention. The equal-prior optimal error is  $(1 - \text{TV})/2$ , proving the first assertion.

Suppose  $nh^8 \rightarrow \lambda \in (0, \infty)$  and write the log likelihood ratio as  $\Lambda_h = \sum_{j=1}^n \log(1 + r_h(X_j))$ . Uniform Taylor expansion and (33.11) give

$$\mathbb{E}_{p_h} \log(1 + r_h) = -\frac{1}{2}\mathcal{J}h^8 + O(h^{10}), \quad \text{Var}_{p_h} \log(1 + r_h) = \mathcal{J}h^8 + O(h^{10}).$$

Indeed the cubic error is bounded by  $C\|r_h\|_\infty \mathbb{E}_{p_h} r_h^2 = O(h^{12})$ . The centered summands are uniformly  $O(h^4)$ , so the Lindeberg condition holds and

$$\Lambda_h \implies N(-\lambda\mathcal{J}/2, \lambda\mathcal{J}) \quad \text{under } p_h^{\otimes n}.$$

Under  $q_h$ , integration against  $(1 + r_h)p_h$  reverses the leading mean and leaves the leading variance unchanged; it gives  $N(\lambda\mathcal{J}/2, \lambda\mathcal{J})$ . The likelihood-ratio test at zero is optimal. The limiting distributions are continuous there, so its two error probabilities converge to  $\Phi(-\sqrt{\lambda\mathcal{J}}/2)$ .

Finally set  $\mu_i(h) = \mathbb{E}_{Q_{i,h}} F$ . By (33.5),

$$\mu_1(h) - \mu_0(h) = \frac{14\Delta}{2025d^2}h^4 + O(h^6) > 0.$$

The statistic  $F$  lies in  $[-8/9, 1/9]$ . Testing its mean at  $(\mu_0 + \mu_1)/2$  has each error at most  $\exp[-n(\mu_1 - \mu_0)^2/2]$ , by the bounded-variable exponential inequality. For completeness, that inequality follows by bounding the centered log moment generating function of a variable in an interval of length one by  $t^2/8$ , multiplying the bounds for independent variables, and minimizing in  $t$ . This proves the third assertion. In the critical regime the bounded triangular-array central limit theorem and (33.5) give

$$\frac{\sqrt{n}(\mu_1 - \mu_0)}{2\sqrt{56/2025}} \longrightarrow \frac{\sqrt{\lambda\mathcal{J}}}{2},$$

so this mean test also attains the optimal critical limit. Its threshold uses the exact means of the specified alternatives, not a leading-order threshold claimed uniformly valid for arbitrarily large  $n$ .  $\square$

**33.3. Finite flights and the separately retained mass.** The preceding scale counts successful window records. It neither counts all preparations nor describes a transcript retaining failures. The relative forward estimate permits a precise finite-flight comparison while keeping this distinction intact.

**Corollary 33.3** (Uniform small-window finite-flight transfer). *Choose the gates and recording profiles as in Theorem 33.2. Let  $Q_{s,J,h}$  be the rescaled recorded law for an even  $J$ -flight bridge at time  $Jg + d$ , conditional on recording both endpoints in the window. There are  $J_0$ ,  $C < \infty$  and  $\tau \in (0, 1)$ , independent of  $0 < h \leq h_0$  and  $s$  in the displayed compact interval, such that*

$$\|q_{s,J,h} - q_{s,h}\|_\infty \leq C\tau^J, \quad J \geq J_0. \quad (33.12)$$

After decreasing  $h_0$  and increasing  $J_0$ , both densities are bounded below by a common  $m > 0$  on the rescaled square. In particular,

$$H^2(Q_{s,J,h}, Q_{s,h}) \leq C\tau^{2J}, \quad (33.13)$$

$$\text{TV}\left(\bigotimes_{k=1}^n Q_{s,J,h}^{(k)}, \bigotimes_{k=1}^n Q_{s,h}^{(k)}\right) \leq \min\{1, C\sqrt{n}\tau^J\}. \quad (33.14)$$

Here the superscripts allow any fixed allocation among the four labelled same-type laws. Thus all optimal-error limits in Theorem 33.2 hold for the finite-flight pair if  $n_h\tau^{2J_h} \rightarrow 0$ . The finite-flight mean-test error is at most

$$\exp(-cn_h h^8) + C\sqrt{n_h}\tau^{J_h}. \quad (33.15)$$

For that mean test alone there is also a direct conclusion: whenever  $\tau^{J_h} = o(h^4)$ , its error is at most  $\exp(-c'n_h h^8)$  for all sufficiently small  $h$ , without a product-approximation hypothesis on  $n_h$ .

Let  $\pi_{s,J}$  be the original gated-bridge success probability, before endpoint recording and cropping, and let  $r_{s,J}(h)$  be the conditional probability of the extra recording and crop given that bridge. Then

$$c_- h^2 \leq r_{s,J}(h) \leq c_+ h^2, \quad r_{s,\infty}(h) = \frac{c_*^2 h^2}{dZ_s} \{4 - I_s(h)^2\}, \quad (33.16)$$

where  $Z_s$  is the normalizer of the full limiting same-type endpoint law. For independent repeated preparations at a fixed design, the expected number needed for  $n$  recorded successes is exactly  $n/(\pi_{s,J}r_{s,J}(h))$ . This expectation is not a deterministic preparation-cap guarantee.

*Proof.* Let  $f_{s,J}$  and  $f_s$  be the successful full endpoint densities in the true signed contact chart, before recording. The relative factorization in Theorem 17.2, with its positive fixed-offset normalizer, gives

$$\|f_{s,J} - f_s\|_{C^0([-h_0, h_0]^2)} \leq C\tau^J.$$

This is the density comparison in the proof of Corollary 23.3; the compact analytic family supplies its uniform constants. Shrinking  $h_0$  and increasing  $J_0$  make both densities uniformly positive and bounded on this square. The recording profiles and their reciprocals are bounded there as well. Writing  $w_s(u, v) = e_s(u)e_s(v)$  and  $A_{s,J}(h) = \int_{(-h,h)^2} w_s f_{s,J}$ , we have

$$A_{s,J}(h) \asymp h^2, \quad |A_{s,J}(h) - A_{s,\infty}(h)| \leq Ch^2\tau^J.$$

Subtract the rescaled quotients

$$q_{s,J,h}(z, w) = \frac{h^2 w_s(hz, hw) f_{s,J}(hz, hw)}{A_{s,J}(h)}.$$

The numerator error and normalizer have the same area factor. Their quotient proves (33.12), with no inverse-area loss. The same bounds give the common positive lower density bound.

For any such two densities  $p, q$  on  $[-1, 1]^2$ ,

$$H^2(P, Q) = \int \frac{(p-q)^2}{(\sqrt{p} + \sqrt{q})^2} \leq \frac{1}{4m} \int (p-q)^2 \leq m^{-1} \|p-q\|_\infty^2.$$

This proves (33.13). If  $\mathfrak{a}(P, Q) = \int \sqrt{pq} = 1 - H^2(P, Q)/2$ , independence gives  $\mathfrak{a}(\bigotimes P_k, \bigotimes Q_k) = \prod_k \mathfrak{a}(P_k, Q_k)$ . The inequality  $1 - \prod_k (1 - x_k) \leq \sum_k x_k$ ,  $0 \leq x_k \leq 1$ , therefore gives  $H^2(\bigotimes P_k, \bigotimes Q_k) \leq \sum_k H^2(P_k, Q_k)$ . Together with  $\text{TV} \leq H$ , this proves (33.14), uniformly in the fixed label allocation. Expectations of tests taking values in  $[0, 1]$  differ by at most total variation under each alternative. Taking the infimum over tests proves the assertion about optimal errors; applying the bound to the mean test proves (33.15). The former sufficient

condition  $n_h \tau^{J_h} \rightarrow 0$  and bound  $\exp(-cn_h h^8) + Cn_h \tau^{J_h}$  follow as weaker consequences, since  $n_h \geq 1$  and  $\tau^{J_h} \leq 1$ .

For the direct mean-test bound, put  $D_h = \mu_1(h) - \mu_0(h) = c_0 h^4 + O(h^6)$ , where  $c_0 = 14\Delta/(2025d^2) > 0$  and the means are those of the two ideal laws. The boundedness of  $F$  and (33.12) show that every finite-flight mean differs from its corresponding ideal mean by at most  $C\tau^J$ . If  $\tau^J = o(h^4)$ , the fixed ideal midpoint is at distance at least  $D_h/4$  from each finite-flight mean, on the correct side. The bounded-variable inequality used in the preceding theorem, applied directly to the independent finite-flight variables of range length one, bounds each error by  $\exp(-n_h D_h^2/8)$ . This is  $\exp(-c'n_h h^8)$  for small  $h$ . It compares the performance of this statistic, not the total variation of two growing product experiments.

Finally  $r_{s,J}(h) = A_{s,J}(h)$ . Substitution of the limiting density and (33.10) gives the exact mass formula in (33.16); the positivity bounds give its two-sided estimate. Each independent preparation is accepted with probability  $\pi_{s,J} r_{s,J}(h) > 0$ . The sum of  $n$  independent geometric waiting times has the displayed expectation. No conditioning on completion of a finite preparation cap is used in this calculation.  $\square$

The Hellinger sharpening and the distinction between product transfer and direct testing were communicated in [44]. The choices are joint: for example, taking the smallest admissible even  $J$  with  $\tau^J \leq h/\sqrt{n_h + 1}$  ensures  $n_h \tau^{2J} \rightarrow 0$ . At the critical scale  $n_h h^8 \asymp 1$ , it suffices that  $\tau^{J_h} = o(h^4)$ . For any fixed  $\varepsilon > 0$ , a sufficient explicit choice there is

$$J_h = 2 \left\lceil \frac{(4 + \varepsilon) \log(1/h)}{2|\log \tau|} \right\rceil,$$

with  $J_h$  increased to  $J_0$  if necessary. This improves a sufficient flight-length coefficient; no optimal geometric convergence exponent is asserted. In the supercritical regime a direct bound for the mean test need not approximate its entire growing sample law.

For every fixed window these records still determine their signed action germs and exact table fibers. The attained two-point scale does not supply a uniform noisy inverse over arbitrary profiles, a fixed finite vector recovering all analytic tables, or a calibration rule for unknown windows. It is a quantitative consequence of the same relative boundary law on an explicitly realized geometric family. The association algebra is the inherited one [41]; neither its cancellation nor the normal approximation is offered as a new general association principle.

**33.4. Waiting records, coarsening and a deterministic cap.** We next retain the accepted mass of the same experiment, rather than normalizing it away. Fix one channel and one contact type in Theorem 33.2, and retain its two specified table/profile alternatives  $s_0, s_1$ . Only the acceptance indicator is recorded on a failed preparation; no raw failed-shot position or residual time is available. Independent preparations are made at one even flight number  $J_h \rightarrow \infty$ , fixed before acquisition. Put

$$p_{i,h} = \pi_{s_i, J_h} r_{s_i, J_h}(h), \quad \gamma_i = \operatorname{arcosh}(1 + g\kappa_{s_i}), \quad \rho_h = p_{1,h}/p_{0,h}. \quad (33.17)$$

Let  $W_h$  count preparations up to and including the first accepted window record, and let  $Y_h$  be that rescaled mark. Denote the binary experiments observing  $Y_h$  and  $(W_h, Y_h)$  by  $\mathcal{E}_h$  and  $\mathcal{F}_h$ , respectively. For binary experiments  $\mathcal{U} = (U_0, U_1)$ ,  $\mathcal{V} = (V_0, V_1)$ , use

$$\delta(\mathcal{U}, \mathcal{V}) = \inf_K \max_{i=0,1} \operatorname{TV}(KU_i, V_i),$$

where  $K$  ranges over Markov kernels independent of the unknown alternative. The repeat-to-first-success coarsening calculation below was communicated in [44]. The deterministic-cap conclusion also identifies the testing scale when the experiment can end without an accepted mark.

For a deterministic cap  $b$ , let  $\mathcal{C}_h^{[b]}$  retain the accepted waiting count  $W_h \leq b$  or the cemetery symbol, and let  $\mathcal{B}_h^{[b]}$  retain only  $\mathbf{1}_{\{W_h \leq b\}}$ . Thus the censored count distinguishes acceptance at the last preparation from no acceptance. Write  $\mathcal{F}_h^{[b]}$  for the full capped record,  $a_{i,h}(b) = 1 - (1 - p_{i,h})^b$ , and  $\eta_h = \text{TV}(Q_{s_0, J_h, h}, Q_{s_1, J_h, h})$ . Let  $R_{C,h}^*(b)$  denote the count-only equal-prior risk. We use  $\Delta_{\text{LC}}(\mathcal{U}, \mathcal{V}) = \max\{\delta(\mathcal{U}, \mathcal{V}), \delta(\mathcal{V}, \mathcal{U})\}$ . The count reduction, finite-intensity bit reduction and count-risk formula below were communicated in [45, 46]. They specify which component of this stopped experiment carries its information; they do not supply another geometric inverse.

**Theorem 33.4** (Accepted marks and charged first-acceptance records). *For the specified pair and any even  $J_h \rightarrow \infty$ , as  $h \downarrow 0$ ,*

$$p_{i,h} \asymp h^2 e^{-\gamma_i J_h}, \quad \gamma_1 > \gamma_0 > 0, \quad \rho_h \asymp e^{-(\gamma_1 - \gamma_0) J_h} \rightarrow 0. \quad (33.18)$$

No relation between  $h$  and  $\tau^{J_h}$  is needed for the following conclusions.

- (1) *The optimal equal-prior error for one accepted mark tends to  $1/2$ , whereas that for the record  $(W_h, Y_h)$  tends to zero. More precisely,*

$$\delta(\mathcal{F}_h, \mathcal{E}_h) = 0, \quad \delta(\mathcal{E}_h, \mathcal{F}_h) \rightarrow \frac{1}{2}. \quad (33.19)$$

- (2) *Fix an integer cap  $b_h \geq 1$ . Observe  $(W_h, Y_h)$  if  $W_h \leq b_h$  and a single cemetery outcome  $\partial$  otherwise. Let  $R_h^*(b_h)$  be the optimal equal-prior error in this capped experiment. If  $b_h p_{0,h} \rightarrow \lambda \in [0, \infty)$ , then*

$$R_h^*(b_h) \rightarrow \frac{1}{2} e^{-\lambda}. \quad (33.20)$$

*The test selecting alternative zero exactly when an acceptance occurs before the cap attains this limit. For arbitrary cap sequences,  $R_h^*(b_h) \rightarrow 0$  if and only if  $b_h p_{0,h} \rightarrow \infty$ .*

- (3) *The stopped experiment uses at most  $b_h$  preparations. Its expected charge under alternative  $i$  is exactly*

$$\mathbb{E}_i \min(W_h, b_h) = \frac{1 - (1 - p_{i,h})^{b_h}}{p_{i,h}}, \quad \mathbb{E}_i W_h = p_{i,h}^{-1}. \quad (33.21)$$

*Thus the critical deterministic-cap order for this binary first-acceptance experiment is  $p_{0,h}^{-1} \asymp h^{-2} e^{\gamma_0 J_h}$ ; vanishing error requires a diverging multiple of that scale. This is a scale in preparations, not the  $h^{-8}$  scale in successful marks.*

- (4) *The terminal mark can be removed uniformly over all deterministic caps. More precisely, for every  $b \geq 1$ ,*

$$\delta(\mathcal{F}_h^{[b]}, \mathcal{C}_h^{[b]}) = 0, \quad \delta(\mathcal{C}_h^{[b]}, \mathcal{F}_h^{[b]}) \leq a_{1,h}(b) \eta_h, \quad (33.22)$$

$$0 \leq R_{C,h}^*(b) - R_h^*(b) \leq \frac{1}{2} a_{1,h}(b) \eta_h. \quad (33.23)$$

*In particular the supremum over  $b \geq 1$  of the Le Cam distance tends to zero. The same reduction without a cap has bound  $\eta_h$ . The simulating kernel preserves the censored count and hence the realized preparation charge, not merely its expectation.*

*For sufficiently small  $h$ , suppress  $h$  and put  $p_i = p_{i,h}$ , so that  $0 < p_1 < p_0 < 1$ . Define*

$$K_b = \min \left\{ b, \left\lceil 1 + \frac{\log(p_0/p_1)}{\log((1-p_1)/(1-p_0))} \right\rceil \right\}. \quad (33.24)$$

*Then the exact count-only risk is*

$$R_{C,h}^*(b) = \frac{1}{2} \{ (1-p_0)^{K_b} + 1 - (1-p_1)^{K_b} \}. \quad (33.25)$$



Choosing either hypothesis at a likelihood tie gives the same risk. For comparison, if  $q_i$  are densities of the two terminal mark laws with respect to a common measure  $\mu$ , the exact full-record risk is

$$R_h^*(b) = \frac{1}{2} \left\{ (1 - p_0)^b + \sum_{k=1}^b \int \min\{p_0(1 - p_0)^{k-1}q_0, p_1(1 - p_1)^{k-1}q_1\} d\mu \right\}. \quad (33.26)$$

Thus (33.25) is not asserted to be the exact finite risk of the full record. Bound (33.23) is an absolute error bound, with no relative-error assertion for vanishing risks.

- (5) At finite cap intensity, the acceptance bit already captures the limiting experiment. For every finite cap,

$$\delta(\mathcal{F}_h^{[b]}, \mathcal{B}_h^{[b]}) = 0, \quad \delta(\mathcal{B}_h^{[b]}, \mathcal{F}_h^{[b]}) \leq a_{1,h}(b). \quad (33.27)$$

If  $b_h p_{0,h} \rightarrow \lambda < \infty$ , let  $\mathcal{L}_\lambda$  be the binary experiment on  $\{0, 1\}$  with success probability  $1 - e^{-\lambda}$  under alternative zero and zero under alternative one. Then

$$\begin{aligned} \Delta_{\text{LC}}(\mathcal{F}_h^{[b_h]}, \mathcal{L}_\lambda) &\leq a_{1,h}(b_h) \\ &+ \max\{|a_{0,h}(b_h) - (1 - e^{-\lambda})|, a_{1,h}(b_h)\} \rightarrow 0. \end{aligned} \quad (33.28)$$

This equivalence does not hold uniformly over all caps. For example, if  $b_h p_{1,h} \rightarrow \infty$ , the bit has equal-prior risk tending to  $1/2$ , while the full record has risk tending to zero, and

$$\delta(\mathcal{B}_h^{[b_h]}, \mathcal{F}_h^{[b_h]}) \rightarrow \frac{1}{2}. \quad (33.29)$$

All tests here are for the two specified alternatives; their thresholds may use the corresponding acceptance probabilities. The theorem does not assert a uniform estimator over unknown recording profiles, or an optimal preparation bound for experiments retaining additional failed-shot data.

*Proof.* The contact curvatures of the pair are  $80/7$  and  $40/3$ . Hence  $\gamma_0 = \text{arcosh}(71/7)$  and  $\gamma_1 = \text{arcosh}(35/3) > \gamma_0$ . The positive fixed-offset prefactor in Theorem 17.3 gives  $\pi_{s_i, J} \asymp e^{-\gamma_i J}$  along this even subsequence. The bound (33.16) is uniform in small  $h$  and large  $J$ . Multiplication proves (33.18); in particular both  $p_{i,h}$  tend to zero. These are the original bridge probabilities and the fixed recording profiles of actual tables, not freely assigned geometric-distribution parameters.

Independence of the preparations gives the exact factorization, for  $k \geq 1$  and measurable  $A \subset [-1, 1]^2$ ,

$$\Pr_i(W_h = k, Y_h \in A) = (1 - p_{i,h})^{k-1} p_{i,h} Q_{s_i, J_h, h}(A). \quad (33.30)$$

Thus the waiting count is geometric on  $\{1, 2, \dots\}$  and is independent of its terminal mark under each alternative. Summing (33.30) in  $k$  shows that discarding the count gives exactly the finite-flight conditional mark, without any approximation or selection of a favorable history. By (33.3) and (33.12),

$$\text{TV}(E_{0,h}, E_{1,h}) \leq C(h^4 + \tau^{J_h}) \rightarrow 0. \quad (33.31)$$

Choose  $m_h = \lceil (p_{0,h} p_{1,h})^{-1/2} \rceil$  and select alternative zero when  $W_h \leq m_h$ . The two errors satisfy

$$\Pr_0(W_h > m_h) \leq e^{-\rho_h^{-1/2}}, \quad \Pr_1(W_h \leq m_h) \leq \sqrt{\rho_h} + p_{1,h}. \quad (33.32)$$

The first bound uses  $(1 - p)^m \leq e^{-pm}$ ; the second uses  $1 - (1 - p)^m \leq mp$  and the upper rounding bound for  $m_h$ . Both errors vanish, so  $\text{TV}(F_{0,h}, F_{1,h}) \rightarrow 1$ . The equal-prior error identity proves the first comparison.

Projection discarding  $W_h$  proves  $\delta(\mathcal{F}_h, \mathcal{E}_h) = 0$ . For any kernel  $K$ , contraction and the triangle inequality give

$$\text{TV}(F_{0,h}, F_{1,h}) \leq 2 \max_i \text{TV}(KE_{i,h}, F_{i,h}) + \text{TV}(E_{0,h}, E_{1,h}).$$

Equations (33.31)–(33.32) therefore give the lower limit  $1/2$  for the reverse deficiency. The kernel ignoring its input and drawing from  $(F_{0,h} + F_{1,h})/2$  has maximum error  $\text{TV}(F_{0,h}, F_{1,h})/2 \leq 1/2$  and proves the matching upper bound. Kernels may depend on the specified pair and on  $h$ , but not on which alternative generated the input.

For the deterministic cap, put  $a_{i,h} = 1 - (1 - p_{i,h})^{b_h}$ , the probability of at least one acceptance. Eventually  $p_{1,h} < p_{0,h}$ , so  $a_{1,h} \leq a_{0,h}$ . The two capped laws have respective cemetery masses  $1 - a_{i,h}$ . Testing the event of any acceptance gives a lower bound on their total variation, and their common cemetery mass gives an upper bound:

$$a_{0,h} - a_{1,h} \leq \text{TV}(F_{0,h}^{[b_h]}, F_{1,h}^{[b_h]}) \leq a_{0,h}. \quad (33.33)$$

This bound does not presume equality of the conditional accepted-mark laws. If  $b_h p_{0,h} \rightarrow \lambda < \infty$ , then  $b_h p_{1,h} = \rho_h b_h p_{0,h} \rightarrow 0$  and  $b_h p_{0,h}^2 \rightarrow 0$ . Expanding  $\log(1 - p_{0,h})$  shows  $a_{0,h} \rightarrow 1 - e^{-\lambda}$ , while  $a_{1,h} \rightarrow 0$ . The sandwich proves (33.20). The stated test has exact equal-prior error  $\{(1 - p_{0,h})^{b_h} + 1 - (1 - p_{1,h})^{b_h}\}/2$  and hence attains that limit.

If  $b_h p_{0,h} \rightarrow \infty$ , use the earlier test with the observable threshold  $t_h = \min\{b_h, m_h\}$ . It is determined by the capped record, including its cemetery outcome. Its errors are bounded by  $e^{-p_{0,h} t_h}$  and  $p_{1,h} t_h$ . The first vanishes because  $p_{0,h} t_h = \min\{b_h p_{0,h}, p_{0,h} m_h\} \rightarrow \infty$ ; the second is at most  $\sqrt{\rho_h} + p_{1,h} \rightarrow 0$ . Conversely, if  $b_h p_{0,h}$  does not tend to infinity, a bounded subsequence has a further subsequence converging to some finite  $\lambda$ . On that subsequence (33.20) gives a strictly positive limiting error. This proves the necessity as well as sufficiency of the cap criterion.

Finally the tail-sum identity gives  $\mathbb{E}_i \min(W_h, b_h) = \sum_{k=0}^{b_h-1} (1 - p_{i,h})^k$ , which is (33.21); the infinite geometric sum gives  $\mathbb{E}_i W_h$ . Every failed preparation before stopping is included. The cap scale follows from (33.18) and the proved risk criterion, not from turning an expected waiting time into a deterministic bound.

It remains to identify the two reductions of the full stopped record. For the count reduction, keep  $k \leq b$  and independently append a mark with law  $Q_{s_0, J_h, h}$ ; keep  $\partial$  unchanged. This is one Markov kernel for both alternatives. Formula (33.30) shows it is exact under zero. Under one, its accepted-count components have disjoint supports and weights  $p_{1,h}(1 - p_{1,h})^{k-1}$ . Consequently its error is exactly

$$\sum_{k=1}^b p_{1,h}(1 - p_{1,h})^{k-1} \eta_h = a_{1,h}(b) \eta_h.$$

Projection proves the zero deficiency in the other direction, giving (33.22). For any full-record randomized test, composition with this kernel changes its error under zero by zero and under one by at most  $a_{1,h}(b) \eta_h$ . Averaging the two errors and taking infima gives (33.23); the reverse risk inequality follows by discarding the mark. The bound is uniform in  $b$  since  $a_{1,h}(b) \leq 1$  and (33.31) makes  $\eta_h \rightarrow 0$ . The infinite geometric sum gives the uncapped bound. The kernel retains  $k$  and  $\partial$ , so the charge function, equal to  $k$  on acceptance and  $b$  at  $\partial$ , is preserved pointwise.

For the exact finite risk, the count likelihood ratio at  $k \leq b$  is

$$\frac{p_0}{p_1} \left( \frac{1 - p_0}{1 - p_1} \right)^{k-1}.$$

It strictly decreases in  $k$ , and the likelihood ratio at  $\partial$  is  $((1 - p_0)/(1 - p_1))^b < 1$ . The largest accepted count at which the ratio is at least one is exactly (33.24). The optimal rule therefore selects zero precisely on the accepted counts  $1, \dots, K_b$ . Its two errors are  $(1 - p_0)^{K_b}$  and  $1 - (1 - p_1)^{K_b}$ , proving (33.25). A tie contributes the same amount whichever decision is used. For the full law, use the common dominating mark measure  $\mu = Q_0 + Q_1$ , counting measure on  $\{1, \dots, b\}$  and an additional atom at  $\partial$ . The minimum equal-prior error is half the integral of the pointwise minimum of the two densities. The smaller cemetery mass is  $(1 - p_0)^b$ , and the other

components are precisely those in (33.26). This proves that identity without an assumption that the terminal marks agree or are finitely supported.

For the bit reduction, send zero to  $\partial$ . On input one, draw a pair  $(k, Y)$  from the full law under alternative zero conditional on acceptance by  $b$ . This law is defined because  $a_{0,h}(b) > 0$ ; its weights are  $p_{0,h}(1 - p_{0,h})^{k-1}/a_{0,h}(b)$ , followed by the independent zero-alternative mark. The simulated full law is exact under zero. Under one the cemetery mass is still exact, and the discrepancy on the accepted part is at most its mass  $a_{1,h}(b)$ , proving (33.27). Notice that no small mark distance is needed for this step. Under  $b_h p_{0,h} \rightarrow \lambda < \infty$ , the earlier calculation gives  $a_{0,h} \rightarrow 1 - e^{-\lambda}$  and  $a_{1,h} \rightarrow 0$ . The actual bit laws and the laws of  $\mathcal{L}_\lambda$  share a two-point space, and their maximum total-variation distance is the maximum in (33.28). Composing the stated kernels and using contraction and the triangle inequality proves that bound in both directions. For  $\lambda = 0$  both limiting laws are the point mass at zero; no exceptional case of the argument is needed.

Finally suppose  $b_h p_{1,h} \rightarrow \infty$ . Then both bit acceptance probabilities tend to one, so the bit pairwise distance tends to zero. Also  $b_h p_{0,h} \rightarrow \infty$ ; the proved cap criterion makes the full pairwise distance tend to one. Applying the same contraction and triangle inequality used for (33.19), now to any kernel from the bit experiment to the full experiment, gives the lower limit  $1/2$  in (33.29). Ignoring the bit and drawing the equal mixture of the two full laws gives the matching upper bound. These bit kernels describe comparison of statistical experiments, not equality of realized physical costs: unlike the count kernel, a bit kernel need not retain the observed stopping time. The deterministic budget and charge remain those in (33.21).  $\square$

The two limits describe distinct coarsenings of one physical acquisition record. Conditional normalization suppresses the accepted mass, leaving an  $h^4$  interaction contrast; retaining failures exposes the different hyperbolic exponents. At a fixed preparation cap, the cemetery outcome itself becomes informative. None of these operations changes the exact window invariant or the geometric rigidity theorem. Their statistical conclusions remain specific to the stated records and fixed pair of profiles, which need not be known or shared in the unrestricted nuisance problem.

### 34. RELATIVE BOUNDARY LAWS FOR PHYSICAL EXPERIMENTS

The two-boundary law also approximates finite experiments, rather than only their limiting offset coefficients. It is important to retain the probability of obtaining a selected record in this comparison. We use  $\|P - Q\|_{\text{TV}} = \sup_B |P(B) - Q(B)|$  for probability measures. For finite signed measures,  $|\mu|(\Omega)$  denotes the total variation mass. Constants below are uniform on the compact geometric families of Theorem 13.1. Neither symmetry nor finite horizon is added in this section.

Fix a channel, suppress its index, and put

$$t_j = -W_{j,uv}(0, 0) > 0, \quad p = j \bmod 2.$$

$$E_\infty(u, v) = S_0(u) + S_p(v), \quad b_\infty(u, v) = B_0(u)B_p(v).$$

The physical time from preparation to the first impact is denoted by  $r$ . On a selected event it lies between zero and the unused excess time. At  $d > 0$  use the common coordinates

$$(z_0, z_1, s) = (u/\sqrt{d}, v/\sqrt{d}, r/d), \quad z = (z_0, z_1).$$

The finite and boundary subprobability measures on this space are

$$d\mu_{j,d} = \frac{d^2 t_j}{2\pi A} b_j(\sqrt{d} z) \mathbf{1}_{\{s \geq 0, s+E_j(\sqrt{d} z)/d < 1\}} dz ds, \quad (34.1)$$

$$d\mu_{j,d}^\partial = \frac{d^2 t_j}{2\pi A} b_\infty(\sqrt{d} z) \mathbf{1}_{\{s \geq 0, s+E_\infty(\sqrt{d} z)/d < 1\}} dz ds. \quad (34.2)$$

Their densities are zero outside the common contact box. After decreasing the offset collar, all their nonzero supports lie in a fixed bounded  $z$ -ball and in  $0 \leq s \leq 1$ . The first measure is exactly the full-phase selected-record measure, by the physical flux formula. Write  $p_{j,d} = \mu_{j,d}(\mathbb{R}^3)$  and  $p_{j,d}^\partial = \mu_{j,d}^\partial(\mathbb{R}^3)$ , and normalize them to  $\Pi_{j,d}$  and  $\Pi_{j,d}^\partial$ . Finally set

$$P_{j,d} = (1 - p_{j,d})\delta_\partial + \mu_{j,d}, \quad P_{j,d}^\partial = (1 - p_{j,d}^\partial)\delta_\partial + \mu_{j,d}^\partial.$$

Here  $\partial$  records that the preparation did not produce the selected event. Thus  $P_{j,d}$  is a coarsening of a raw physical record, not an experiment in which unsuccessful preparations have been discarded for free.

**Theorem 34.1** (Uniform transfer to the boundary experiment). *There are  $d_* > 0$ ,  $C < \infty$  and  $0 < \tau < 1$  such that, for all  $j \geq 1$  and  $0 < d \leq d_*$ , the displayed measures are well defined and*

$$p_{j,d} \asymp p_{j,d}^\partial \asymp d^2 e^{-j\gamma}, \quad (34.3)$$

$$|\mu_{j,d} - \mu_{j,d}^\partial|(\mathbb{R}^3) \leq C p_{j,d} \tau^j, \quad \|\Pi_{j,d} - \Pi_{j,d}^\partial\|_{\text{TV}} \leq C \tau^j, \quad (34.4)$$

$$\|P_{j,d} - P_{j,d}^\partial\|_{\text{TV}} \leq C p_{j,d} \tau^j. \quad (34.5)$$

For  $n$  independent preparations, respectively  $k$  independently prepared successful records, these imply

$$\|P_{j,d}^{\otimes n} - (P_{j,d}^\partial)^{\otimes n}\|_{\text{TV}} \leq \min\{1, C n p_{j,d} \tau^j\}, \quad (34.6)$$

$$\|\Pi_{j,d}^{\otimes k} - (\Pi_{j,d}^\partial)^{\otimes k}\|_{\text{TV}} \leq \min\{1, C k \tau^j\}. \quad (34.7)$$

The same bounds hold after applying any common measurable observation map, including retention of the endpoint positions alone. At  $d = 0$ , the conditional statement holds for the corresponding quadratic right limits in the scaled coordinates.

*Proof.* The positive endpoint Hessians and the uniform nonlinear bounds give, on a sufficiently small common box,

$$c|Y|^2 \leq E_j(Y), E_\infty(Y) \leq C|Y|^2, \quad c \leq b_j(Y), b_\infty(Y) \leq C.$$

These inequalities give the asserted common bounded scaled support and upper and lower bounds on the scaled integrals. They also imply that both masses are less than  $1/2$  after decreasing  $d_*$ , since  $t_j$  is uniformly bounded. The Jacobi formula gives  $t_j \asymp e^{-j\gamma}$ , which proves (34.3).

The action difference has zero value and gradient at  $Y = 0$ . Taylor's integral formula and the  $C^2$  factorization estimate imply

$$|E_j(Y) - E_\infty(Y)| \leq C \tau^j |Y|^2, \quad |b_j(Y) - b_\infty(Y)| \leq C \tau^j.$$

Consequently, on the common scaled ball,

$$\left| \frac{E_j(\sqrt{d} z) - E_\infty(\sqrt{d} z)}{d} \right| \leq C \tau^j |z|^2.$$

For each fixed  $z$ , the two allowed  $s$ -fibres have lengths  $(1 - E_j(\sqrt{d} z)/d)_+$  and  $(1 - E_\infty(\sqrt{d} z)/d)_+$ . They start at the same point zero. Since the positive-part map is 1-Lipschitz, the length of their

symmetric difference is at most  $C\tau^j|z|^2$ . Integrating the bounded amplitudes on this difference, and the amplitude difference on the common part, proves

$$|\mu_{j,d} - \mu_{j,d}^\partial|(\mathbb{R}^3) \leq Cd^2t_j\tau^j/A.$$

This proves the first part of (34.4). In particular  $|p_{j,d} - p_{j,d}^\partial|$  obeys the same bound. Dividing by the positive masses and adding the error of their normalizers gives the conditional estimate. Adding the missing-event atom instead gives (34.5). This argument estimates a moving domain by its fibres; it does not differentiate its indicator function.

For probability measures  $P_i, Q_i$ , telescope their product measures, or couple each pair maximally, to obtain

$$\left\| \bigotimes_{i=1}^n P_i - \bigotimes_{i=1}^n Q_i \right\|_{\text{TV}} \leq \sum_{i=1}^n \|P_i - Q_i\|_{\text{TV}}.$$

This proves the two product bounds. A measurable pushforward cannot increase total variation. Lastly  $E_j(\sqrt{d}z)/d$  and its boundary analogue converge to their positive quadratic forms as  $d \downarrow 0$ , and the amplitudes converge to one. The same fibre argument proves  $L^1$  convergence of these bounded densities and of their normalizations. Passing to the limit proves the final assertion.  $\square$

**Corollary 34.2** (Transfer of bounded risks). *For any estimator or test based on the stated coarsened records, and any measurable loss taking values in  $[0, 1]$ , replacing the finite experiment by its boundary model changes its risk by at most the right-hand side of (34.6) or (34.7), respectively. The estimate is uniform in the parameter family. More generally, a fixed design of independent preparations at  $(j_i, d_i)$  has risk error at most  $C \sum_i p_{j_i, d_i} \tau^{j_i}$ .*

*Proof.* Integrate the bounded loss against the difference of the laws, using its layer-cake representation. Apply the product estimate, with possibly different factors, and then take the supremum over parameters.  $\square$

The comparison also applies in unscaled physical endpoint-position and residual-time coordinates, by using the same boundary embeddings for both laws at each parameter. For predetermined physical windows  $t_i$ , set  $d_i(\xi) = t_i - j_i g_e(\xi)$  and restrict the parameter family so that  $0 < d_i(\xi) \leq d_*$ ; the bounds hold uniformly at these offsets. Scaling is then only a device in the proof, not an extra observation of an unknown gap. In a design programmed directly at fixed excess  $d$ , the onset must instead be supplied or calibrated. No onset-calibration cost is removed by the total-variation comparison.

The effective sample size in (34.6) is the expected number  $np_{j,d}$  of selected records. In particular, an experiment with  $np_{j,d} = O(k_j)$  is approximated uniformly whenever  $k_j\tau^j \rightarrow 0$ . This is a simultaneous large-record and large-sample conclusion that separate fixed-flight Taylor expansions do not supply. It concerns the specified endpoints and residual time. The full growing collision array lies on an embedded surface; no total-variation comparison of distinct such surfaces is asserted. Nor does this theorem supply a parameter-free sampling algorithm for an unknown table. It compares the physical statistical families in common observation coordinates and preserves the cost of obtaining a selected event.

### 35. ADAPTIVE OBSERVATION AND STOPPED EXPERIMENTS

The preceding comparison concerns a fixed list of preparations. In an observation procedure the next window may instead depend on the records already obtained, and the procedure may stop after enough successful records. We prove a comparison for the entire stopped transcript. In particular, the failure symbols and the stopping time are not removed from the experiment.

We use the unscaled physical endpoint and residual-time coordinates at the end of Section 34. A design  $a$  specifies a labelled channel, a flight number  $j(a)$  and a physical observation time

$t(a)$ . At geometric parameter  $\xi$  its excess is  $d_\xi(a) = t(a) - j(a)g_\xi$ . Write  $P_{\xi,a}$  for the raw coarsened physical law and  $Q_{\xi,a}$  for the corresponding boundary law. The success probability is  $p_\xi(a) = P_{\xi,a}(\text{success})$ . The chosen designs must lie in the common physical collar:  $0 < d_\xi(a) \leq d_*$ . An unknown excess is used only in proving the comparison; it is not supplied to the observation rule.

The design and record spaces are standard Borel spaces. A policy consists of measurable rules for choosing a design or stopping, using the observed past and auxiliary randomness independent of  $\xi$ . The same rules are used in the two experiments. Conditional on the past and the chosen design, a fresh preparation has law  $P_{\xi,a}$  or  $Q_{\xi,a}$ , respectively. The policy stops after  $T \leq N$  preparations, where  $N$  is a prescribed deterministic cap. The transcript includes every design, every success or failure record, and  $T$ . These conventions include randomized pilot procedures followed by data-dependent designs.

**Lemma 35.1** (A stopped comparison for observation kernels). *Let  $P_{\xi,a}$  and  $Q_{\xi,a}$  be jointly measurable probability kernels on a common record space, and suppose that their total variation distance is at most a nonnegative measurable function  $e_\xi(a)$ . Assume that a measurable maximal coupling is available for each pair of kernels. For any policy as above, let  $\mathbf{P}_\xi^\pi$  and  $\mathbf{Q}_\xi^\pi$  be the laws of its stopped transcript. Then*

$$\|\mathbf{P}_\xi^\pi - \mathbf{Q}_\xi^\pi\|_{\text{TV}} \leq \min \left\{ 1, \mathbb{E}_{\mathbf{P}_\xi^\pi} \sum_{i=1}^T e_\xi(a_i) \right\}. \quad (35.1)$$

*One may interchange  $P$  and  $Q$  in the expectation. The conclusion is unchanged by recording the independent random seed or by applying a common randomized observation map to the transcript.*

*Proof.* Use the same auxiliary randomness for both policies while their pasts agree. They then choose the same design, or both stop. At a common active past, couple the next two records maximally. After their first disagreement continue each marginal according to its own policy and observation kernel. This constructs a coupling of the required transcript laws. Let  $D_i$  be the event that their first disagreement is at preparation  $i$ . Given a common past and its next design  $a_i$ , its conditional probability is at most  $e_\xi(a_i)$ . Therefore

$$\Pr(D_i) \leq \mathbb{E}[\mathbf{1}_{\{\text{pasts agree before } i, T_P \geq i\}} e_\xi(a_i^P)] \leq \mathbb{E}_{\mathbf{P}_\xi^\pi}[\mathbf{1}_{\{T \geq i\}} e_\xi(a_i)].$$

Here  $T_P$  and  $a_i^P$  refer to the physical marginal; no independence of the random designs is used. The events  $D_i$  are disjoint. If none occurs, the two policies also stop together and their complete transcripts agree. For any measurable set  $B$ , a coupling satisfies  $|\Pr(X \in B) - \Pr(Y \in B)| \leq \Pr(X \neq Y)$ . Summing the last display proves (35.1). Interchanging the marginals gives the other expectation. Coupling the randomized output conditionally in the same way proves the last assertion.  $\square$

The coupling inequality is classical; see, for example, [28, Proposition 4.7]. For the physical kernels, a common scalar dominating measure suffices; their planar pushforwards need not have ambient densities.

**Lemma 35.2** (Measurable coupling in physical observation coordinates). *Let  $K$  be a fixed bounded scalar endpoint–time box, and equip  $K_\partial = K \sqcup \{\partial\}$  with the measure  $\nu = \text{Leb}_K + \delta_\partial$ . Suppose  $f_{\xi,a}, g_{\xi,a}$  are jointly measurable probability densities with respect to  $\nu$ . Let  $Y$  be a standard Borel physical record space and let  $\iota_{\xi,a} : K_\partial \rightarrow Y$  be jointly measurable and, for each  $(\xi, a)$ , a Borel isomorphism onto a Borel subset of  $Y$ . Assume that it sends  $\partial$  to the common failure symbol. Then the probability kernels*

$$P_{\xi,a} = (\iota_{\xi,a})_*(f_{\xi,a}\nu), \quad Q_{\xi,a} = (\iota_{\xi,a})_*(g_{\xi,a}\nu)$$

admit a jointly measurable maximal coupling. In particular,

$$\|P_{\xi,a} - Q_{\xi,a}\|_{\text{TV}} = 1 - \int_{K_\partial} \min(f_{\xi,a}, g_{\xi,a}) d\nu. \quad (35.2)$$

No domination of  $P_{\xi,a}$  and  $Q_{\xi,a}$  by ambient Lebesgue measure on  $Y$ , or domination across different geometric parameters, is required.

*Proof.* Suppress  $(\xi, a)$  and put  $h = \min(f, g)$  and  $m = \int h d\nu$ . For a measurable set  $D \subset Y \times Y$  define

$$\begin{aligned} C(D) &= \int_{K_\partial} \mathbf{1}_D(\iota(x), \iota(x)) h(x) d\nu(x) \\ &\quad + \frac{1}{1-m} \int_{K_\partial^2} \mathbf{1}_D(\iota(x), \iota(y)) (f(x) - h(x))(g(y) - h(y)) d\nu(x) d\nu(y), \end{aligned}$$

where the second term is defined to be zero when  $m = 1$ . Both residual measures have mass  $1 - m$ , so  $C$  is a probability measure with marginals  $P$  and  $Q$ . Its dependence on  $(\xi, a)$  is measurable by integration of jointly measurable functions; on  $\{m = 1\}$  the specified zero branch is measurable as well.

The residual measures are supported, up to null sets, on the disjoint sets  $\{f > g\}$  and  $\{g > f\}$ . Injectivity of  $\iota$  makes their images disjoint. Thus the second term puts no mass on the diagonal and the coupling disagrees with probability  $1 - m$ . Conversely, for fixed  $(\xi, a)$  the assumed Borel inverse makes the image of  $\{f > g\}$  measurable, and the difference of its  $P$  and  $Q$  probabilities is

$$\int_{\{f > g\}} (f - g) d\nu = 1 - m.$$

The coupling bound and this measurable set prove (35.2) and maximality.  $\square$

For the billiard kernels,  $K$  is the common unscaled scalar contact box for  $(u, v, r)$ . The finite and boundary densities, including the failure atom, are given by the physical flux formula and its two-boundary limit. The maps  $\iota_{\xi,a}$  insert the two boundary positions in the plane and retain the residual time and the chosen labels. The contact graphs defining the forward model make these maps jointly smooth on each channel component. Their inverses on the physical images are the transverse projections in the corresponding contact frames, together with the unchanged time and labels. Thus they are Borel isomorphisms onto their Borel images. A finite or countable disjoint union handles the flight numbers and channel labels. The preceding lemma therefore supplies the measurable maximal coupling used in Lemma 35.1 for the actual physical records.

When only scalar endpoint coordinates and residual time are observed,  $\iota_{\xi,a}$  is the identity and the same construction applies directly. For planar endpoint positions the successful law is instead supported on an embedded boundary surface; it is not an ambient Lebesgue density. The parameter-dependent embedding is used only to prove a comparison at one fixed table. It is neither an extra observation nor a parameter-free simulator for an unknown table, and does not compare densities on the different boundary surfaces of two distinct tables.

**Theorem 35.3** (Adaptive physical transfer with a success-weighted budget). *For the compact physical families of Theorem 34.1, let  $\pi$  be a policy whose designs are in the common collar on every reachable history. There are constants  $C < \infty$  and  $0 < \tau < 1$ , independent of  $\pi$  and  $N$ , such that*

$$\|P_\xi^\pi - Q_\xi^\pi\|_{\text{TV}} \leq \min\{1, C\mathcal{B}_\xi(\pi)\}, \quad (35.3)$$

$$\mathcal{B}_\xi(\pi) := \mathbb{E}_{P_\xi^\pi} \sum_{i=1}^T p_\xi(a_i) \tau^{j(a_i)} = \mathbb{E}_{P_\xi^\pi} \sum_{i=1}^T Y_i \tau^{j(a_i)}, \quad (35.4)$$

where  $Y_i$  is the indicator of a successful preparation. Hence replacing the physical experiment by the boundary experiment changes any risk for a loss in  $[0, 1]$  by at most  $\min\{1, C\mathcal{B}_\xi(\pi)\}$ . The bounds are uniform in  $\xi$  and survive any common coarsening.

In particular, if  $j(a_i) \geq J$  and the policy stops no later than its  $k$ th success, then

$$\|P_\xi^\pi - Q_\xi^\pi\|_{\text{TV}} \leq \min\{1, Ck\tau^J\}. \quad (35.5)$$

The physical cost remains the number  $T$  of raw preparations, including all failures; it is not replaced by the number of successes.

*Proof.* Theorem 34.1, transported to the common unscaled physical coordinates, supplies  $e_\xi(a) = Cp_\xi(a)\tau^{j(a)}$ . Apply Lemma 35.1. The event  $\{T \geq i\}$  and the design  $a_i$  are known before the  $i$ th fresh observation. Its conditional success probability is  $p_\xi(a_i)$ . Conditional expectation therefore gives

$$\mathbb{E}[\mathbf{1}_{\{T \geq i\}} Y_i \tau^{j(a_i)}] = \mathbb{E}[\mathbf{1}_{\{T \geq i\}} p_\xi(a_i) \tau^{j(a_i)}].$$

Summing the finite list proves (35.4). Under the stopping assumption the realized sum is at most  $k\tau^J$ . For a loss  $0 \leq L \leq 1$ , write  $L = \int_0^1 \mathbf{1}_{\{L > s\}} ds$  and apply the total variation bound. This proves both the risk assertion and the stated uniformities.  $\square$

**Corollary 35.4** (Pilot observations and growing adaptive experiments). *The same bound applies when a charged pilot is kept exact in both experiments and only the subsequent preparations are replaced by their boundary laws. In that case the sum in (35.4) runs over the replaced preparations only, whereas the total cost includes the pilot. If the pilot certifies the common collar on an event of probability at least  $1 - \eta_0$ , the bound becomes*

$$\|P_\xi^\pi - Q_\xi^\pi\|_{\text{TV}} \leq \eta_0 + C\mathbb{E}_{P_\xi^\pi} \left[ \mathbf{1}_{\{\text{good pilot}\}} \sum_{i \in I_{\text{post}}} Y_i \tau^{j(a_i)} \right]. \quad (35.6)$$

Here, on bad pilot histories, the comparison experiment is completed by any specified probability kernels on the same transcript space; no boundary formula outside its collar is assumed. For a sequence of policies, if  $\eta_0 \rightarrow 0$  and the displayed success-weighted expectations tend to zero uniformly on the physical family, every bounded transcript risk has a vanishing comparison error.

*Proof.* Couple the exact pilot identically, including its stopping decision. It has the same law in both experiments. Condition on the entire pilot transcript. On a good transcript apply the preceding proof to the post-pilot policy, and on a bad one bound disagreement by one. Integration over this common pilot law proves (35.6). The conditional-expectation identity for successes holds with the good-pilot indicator, which is measurable before every post-pilot observation. The final assertion follows by taking a uniform bound.  $\square$

For example, any sequence of policies capped after at most  $k_n$  post-pilot successes and using flight numbers at least  $J_n$  has vanishing error if  $k_n \tau^{J_n} \rightarrow 0$  and  $\eta_{0,n} \rightarrow 0$ . This includes adaptive selection of positive physical offsets and early stopping. Neither this criterion nor the exact pilot removes a preparation charge. It is a consequence of the relative physical theorem for complete experiments; it does not replace the separate full-profile regularization argument, or assert an optimal rate for choosing  $J_n$ .

### 36. BOUNDARY INFORMATION AFTER RESIDUAL-TIME INTEGRATION

Residual-time integration preserves the smooth relative law of Theorem 12.1, but it changes the singularity of the finite-versus-boundary experiment. Before integration the tangent success density is positive up to a moving boundary. After integration it vanishes there to first order.



The resulting logarithmic information is a boundary quantity. The theorem below isolates that quantity before the billiard specialization.

The one-dimensional change of rate caused by parameter-dependent support is classical. Akahira [30, 31] and Ibragimov–Has’minskii [32] belong to the general nonregular background, and Smith [33] explicitly treats the linearly vanishing endpoint case. Limits of experiments for models with parameter-dependent support are also developed by Hirano–Porter [34]. We do not use these references as a black-box proof: the form needed here has a moving regular hypersurface, an intrinsic coefficient, and an independent failure atom, and we prove it directly.

**36.1. A linearly vanishing density.** Let  $U \subset \mathbb{R}^m$  be bounded and open. For  $|t| < t_0$  let  $a_t$  and  $w_t$  be real functions which are  $C^3$  jointly on a neighborhood of  $[-t_0, t_0] \times \bar{U}$ . Assume that  $a_t$  is bounded above and bounded away from zero, that  $w_t$  is uniformly negative near  $\partial U$ , and that

$$D = \{x \in U : w_0(x) > 0\} \neq \emptyset, \quad \Sigma = \{x \in U : w_0(x) = 0\}, \quad |\nabla w_0| > 0 \quad \text{on } \Sigma.$$

Thus  $\bar{D} \subset U$  and  $\Sigma = \partial D$  is a compact regular hypersurface. Decreasing  $t_0$  if necessary, suppose

$$f_t(x) = a_t(x)(w_t(x))_+, \quad \int_U f_t(x) dx = 1. \quad (36.1)$$

Normalizing an initially unnormalized density is allowed: a common tubular chart for the moving boundary makes its positive normalizing integral  $C^3$ . Put

$$v(x) = \partial_t w_t(x)|_{t=0}, \quad \mathcal{I}_\Sigma = \int_\Sigma \frac{a_0(x)v(x)^2}{|\nabla w_0(x)|} d\sigma(x). \quad (36.2)$$

We use

$$H^2(P, Q) = \int (\sqrt{dP/d\nu} - \sqrt{dQ/d\nu})^2 d\nu, \quad \|P - Q\|_{\text{TV}} = \sup_B |P(B) - Q(B)|,$$

so  $0 \leq H^2 \leq 2$ , and write  $\Phi$  for the standard normal distribution function.

**Theorem 36.1** (Logarithmic boundary information). *Under (36.1),*

$$H^2(f_t dx, f_0 dx) = \frac{\mathcal{I}_\Sigma}{4} t^2 \log \frac{1}{|t|} + O(t^2), \quad t \rightarrow 0. \quad (36.3)$$

Assume  $\mathcal{I}_\Sigma > 0$ . Let  $t_n \rightarrow 0$ ,  $t_n \neq 0$ , let  $0 < p_n \leq 1$ , and consider  $n$  independent observations from

$$P_{n,0} = (1 - p_n)\delta_{\dagger} + p_n f_0 dx, \quad P_{n,1} = (1 - p_n)\delta_{\dagger} + p_n f_{t_n} dx, \quad (36.4)$$

where the failure symbol  $\dagger$  is retained. If

$$np_n t_n^2 \log(1/|t_n|) \rightarrow b \in [0, \infty], \quad (36.5)$$

then

$$\|P_{n,1}^{\otimes n} - P_{n,0}^{\otimes n}\|_{\text{TV}} \rightarrow 2\Phi(\sqrt{\mathcal{I}_\Sigma b}/2) - 1, \quad (36.6)$$

with the right-hand side interpreted as one when  $b = \infty$ .

If  $0 < b < \infty$ , the log likelihood of the absolutely continuous part of  $P_{n,1}^{\otimes n}$  relative to  $P_{n,0}^{\otimes n}$  converges under the latter law to

$$\mathcal{N}(-\mathcal{I}_\Sigma b/2, \mathcal{I}_\Sigma b), \quad (36.7)$$

and the minimum equal-prior testing error tends to  $\Phi(-\sqrt{\mathcal{I}_\Sigma b}/2)$ .

If  $b < \infty$  (including  $b = 0$ ), the probability that at least one of the  $n$  observations falls in a support-exclusive region tends to zero under either hypothesis. No support-exclusivity assertion is made under the sole assumption  $b = \infty$ ; the conclusion (36.6) in that regime follows from Hellinger affinity and remains valid even when  $np_n t_n^2$  does not tend to zero.

The last paragraph is a scope distinction, not a weakening of the supercritical theorem. At finite information scale,  $np_n t_n^2 = (np_n t_n^2 \log(1/|t_n|))/\log(1/|t_n|) \rightarrow 0$ , so exclusive support is negligible. At supercritical scale that implication is false in general, but it is unnecessary for total-variation separation.

**36.2. Coarea estimates.** For  $x \in D$  define

$$S(x) = \frac{\dot{a}_0(x)}{a_0(x)} + \frac{v(x)}{w_0(x)}.$$

It is integrable under  $f_0 dx$  but, when  $\mathcal{I}_\Sigma > 0$ , has infinite second moment. For  $h > 0$  put  $D_h = \{w_0 \geq h\}$  and  $S_h = S \mathbf{1}_{D_h}$ .

**Lemma 36.2** (Truncated score moments). *As  $h \downarrow 0$ ,*

$$\Pr\{0 < w_0 < h\} = O(h^2), \quad \mathbb{E}_0 S_h = O(h), \quad (36.8)$$

$$\mathbb{E}_0 S_h^2 = \mathcal{I}_\Sigma \log(1/h) + O(1), \quad (36.9)$$

$$\mathbb{E}_0 |S_h|^3 = O(h^{-1}), \quad \mathbb{E}_0 |S_h|^4 = O(h^{-2}). \quad (36.10)$$

For  $|t|$  sufficiently small, the  $f_t$ -mass of its support outside  $D$  and the  $f_0$ -mass of the region where  $f_t = 0$  are both  $O(t^2)$ . The estimates are uniform on compact families with the stated strict positivity and regularity bounds.

*Proof.* Choose  $h_0 > 0$  so that  $|\nabla w_0|$  is bounded away from zero on  $|w_0| \leq 2h_0$ . The flow of  $\nabla w_0/|\nabla w_0|^2$  identifies the level sets  $w_0 = z$  for  $|z| \leq h_0$  with  $\Sigma$ , with uniformly bounded  $C^2$  Jacobians. By coarea,

$$J(z) = \int_{\{w_0=z\}} \frac{a_0 v^2}{|\nabla w_0|} d\sigma = \mathcal{I}_\Sigma + O(|z|).$$

The collar contribution of  $v^2/w_0^2$  to the second moment is therefore  $\int_h^{h_0} J(z) dz/z$ . The remaining terms in  $S^2$  are integrable, which proves (36.9). Since  $f_0 = O(w_0)$  in the collar,

$$\int_{\{0 < w_0 < h\}} f_0 dx = O(h^2), \quad \int_{D_h} |S|^r f_0 dx \leq C \left( 1 + \int_h^{h_0} z^{1-r} dz \right), \quad r = 3, 4.$$

Differentiating  $\int f_t = 1$  at zero gives  $\int_D (\dot{a}_0 w_0 + a_0 v) dx = 0$ , hence  $\mathbb{E}_0 S = 0$ . The omitted collar integral of  $S f_0$  is  $O(h)$ , proving (36.8). Finally  $|w_t - w_0| \leq C|t|$ , so both support-exclusive regions lie in  $|w_0| \leq C|t|$ ; there the relevant density is  $O(|t|)$  and the collar volume is  $O(|t|)$ . Their masses are  $O(t^2)$ .  $\square$

### 36.3. Proof of the information theorem.

*Proof of Theorem 36.1.* Write  $t \neq 0$  and  $\ell = \log(1/|t|)$ . On  $D_h$ , whenever  $h \geq C_0|t|$  for a sufficiently large fixed  $C_0$ , Taylor expansion of  $a_t w_t$  before division by  $a_0 w_0$  gives

$$\frac{f_t}{f_0} = 1 + tS + t^2 R_t, \quad |R_t(x)| \leq C(1 + w_0(x)^{-1}). \quad (36.11)$$

This expansion never differentiates a support indicator.

Take first  $h = C_0|t|$ . The integral of  $f_0 + f_t$  outside the common part  $D_h$  is  $O(t^2)$ . On  $D_h$  set  $z = tS + t^2 R_t$ . With  $C_0$  large and then  $t$  small,  $|z| \leq 1/2$ , and  $(\sqrt{1+z} - 1)^2 = z^2/4 + O(|z|^3)$ . The coarea bounds give

$$\int_{D_h} z^2 f_0 dx = t^2 \mathbb{E}_0 S_h^2 + O(|t|^3 \log(1/|t|)), \quad \int_{D_h} |z|^3 f_0 dx = O(t^2).$$

Together with (36.9) this proves (36.3).

Assume now  $0 < b < \infty$ . In the  $n$ th row write

$$t = t_n, \quad p = p_n, \quad \ell = \log(1/|t|), \quad L = \ell^{1/4}, \quad h = |t|L.$$

Represent a null observation by a Bernoulli success indicator  $Y_i$  of mean  $p$  and, on success, an independent  $X_i$  of density  $f_0$ . With probability tending to one all successful  $X_i$  lie in  $D_h$ , because

$$Cnph^2 = C(npt^2\ell)L^2/\ell = o(1).$$

The same is true under the alternative. In particular the total support-exclusive probability is  $O(npt^2) = o(1)$  in this finite critical regime.

On the high-probability common-support event,

$$\log \frac{dP_{n,1}^{\otimes n}}{dP_{n,0}^{\otimes n}} = t \sum_{i=1}^n Y_i S_h(X_i) - \frac{t^2}{2} \sum_{i=1}^n Y_i S_h(X_i)^2 + o_{P_{n,0}^{\otimes n}}(1). \quad (36.12)$$

Indeed the summed  $t^2 R_t$  contribution is  $O(npt^2) = o(1)$ , the cubic logarithmic remainder has expectation  $O(np|t|^3/h) = O((npt^2\ell)/(L\ell)) = o(1)$ , and replacing the quadratic term by  $t^2 S_h^2$  produces  $o(1)$  by the same coarea estimates.

The mean of the first sum in (36.12) is  $O(np|t|h) = o(1)$  and its variance tends to  $\mathcal{I}_\Sigma b$ . Each centered summand after multiplication by  $t$  is  $O(1/L)$ , so Lindeberg is automatic. The quadratic sum has mean  $\mathcal{I}_\Sigma b + o(1)$  and variance bounded by

$$Cnpt^4 \mathbb{E}_0 |S_h|^4 \leq Cnpt^4 h^{-2} = C(npt^2\ell)/(L^2\ell) = o(1).$$

This proves (36.7).

Let  $Z_n$  be the log likelihood, with  $Z_n = -\infty$  on the null-exclusive set. Even in the presence of an alternative singular part,

$$1 - \|P_{n,1}^{\otimes n} - P_{n,0}^{\otimes n}\|_{\text{TV}} = \mathbb{E}_{P_{n,0}^{\otimes n}} \min\{1, e^{Z_n}\}.$$

The integrand is bounded and continuous on the extended real line. If  $v = \mathcal{I}_\Sigma b$  and  $G \sim \mathcal{N}(0, 1)$ , then

$$\mathbb{E} \min\{1, e^{\sqrt{v}G - v/2}\} = 2\Phi(-\sqrt{v}/2),$$

which proves the finite positive case of (36.6) and the testing-risk assertion.

For  $b = 0$  and  $b = \infty$  use affinity instead. The one-observation affinity of the thinned pair is exactly  $1 - pH^2(f_t, f_0)/2$ . Affinities multiply, and if  $\mathcal{A}(P, Q)$  is the affinity then

$$1 - \mathcal{A}(P, Q) \leq \|P - Q\|_{\text{TV}} \leq \sqrt{1 - \mathcal{A}(P, Q)^2}.$$

Equation (36.3) therefore gives  $\text{TV} \rightarrow 0$  when  $b = 0$  and  $\text{TV} \rightarrow 1$  when  $b = \infty$ , without any assumption on  $npt^2$ . If merely  $b < \infty$ , the support-exclusive probability is  $O(npt^2) = o(1)$  by the same division by  $\ell$  used above. This proves all remaining statements.  $\square$

*Remark 36.3* (Intrinsic coefficient). Replace  $w_t$  by  $\tilde{w}_t = c_t w_t$ ,  $c_t > 0$  smooth, and  $a_t$  by  $a_t/c_t$ . On  $\Sigma$ ,  $\partial_t \tilde{w}_t|_0 = c_0 v$  and  $|\nabla \tilde{w}_0| = c_0 |\nabla w_0|$ , so (36.2) is unchanged. Equivalently, if  $V_\nu = -v/|\nabla w_0|$  is the normal velocity of the support boundary,

$$\mathcal{I}_\Sigma = \int_\Sigma a_0 |\nabla w_0| V_\nu^2 d\sigma.$$

If this integral vanishes, the theorem supplies only  $H^2 = O(t^2)$ ; no nontrivial logarithmic critical scale is asserted. When it is positive, the smooth interior Fisher contribution is lower order.

### 37. VECTOR BOUNDARY INFORMATION FOR A MOVING-SUPPORT LOCAL FAMILY

We formulate the moving-support experiment without assuming absolute continuity with respect to its reference member. The original family is jointly dominated by Lebesgue measure and its failure atom, but a likelihood ratio relative to the reference law need not exist on the support-exclusive part. We separate that mass and apply parameter-independent common-collar censoring. The censored family is dominated by its censored reference member and satisfies a uniform LAN expansion. The original family is uniformly asymptotically equivalent to that representative on every compact local set. We distinguish this equivalence statement from finite-subexperiment convergence to the Gaussian shift and formulate singular information on its identifiable quotient. The further compact-parameter conclusion, proved by a uniform Hellinger modulus in Theorem 40.2, is a separate step. Proposition 43.1 records the common measures explicitly.

Throughout this section we use the Hellinger convention

$$H^2(P, Q) = \int (\sqrt{p} - \sqrt{q})^2 d\mu = 2 - 2 \int \sqrt{pq} d\mu, \quad (37.1)$$

for any common dominating measure  $\mu$ . With this convention  $H^2(P^{\otimes k}, Q^{\otimes k}) \leq kH^2(P, Q)$ .

Let  $\Theta \subset \mathbb{R}^r$  be a neighborhood of the origin. On a fixed bounded open set  $U \subset \mathbb{R}^m$  suppose

$$f_\theta(x) = a_\theta(x)(w_\theta(x))_+, \quad \int_U f_\theta(x) dx = 1. \quad (37.2)$$

Assume that  $a_\theta, w_\theta$  are  $C^3$  jointly in  $(\theta, x)$  on a common compact neighborhood of the relevant supports, uniformly for small  $\theta$ , that  $a_\theta$  is bounded above and bounded away from zero there, and that

$$\Sigma = \{w_0 = 0\} \cap \text{supp } a_0$$

is a compact regular hypersurface with  $|\nabla w_0| \geq c_* > 0$ . After shrinking  $\Theta$ , all moving boundaries lie in one fixed tubular collar. Put

$$V(x) = D_\theta w_\theta(x)|_{\theta=0} \in \mathbb{R}^r, \quad \mathcal{J}_\Sigma = \int_\Sigma \frac{a_0(x)V(x)V(x)^t}{|\nabla w_0(x)|} d\sigma(x). \quad (37.3)$$

Multiplying  $w_\theta$  by a positive smooth factor and dividing  $a_\theta$  by the same factor leaves  $\mathcal{J}_\Sigma$  unchanged.

Let  $\delta_n \downarrow 0$ ,  $0 < p_n \leq 1$ , and assume

$$np_n \delta_n^2 \log(1/\delta_n) \longrightarrow 1. \quad (37.4)$$

For  $h \in \mathbb{R}^r$  define on  $\Omega = U \sqcup \{\dagger\}$

$$P_{n,h} = (1 - p_n)\delta_\dagger + p_n f_{\delta_n h}(x) dx. \quad (37.5)$$

No absolute continuity of  $P_{n,h}$  with respect to  $P_{n,0}$  is assumed.

#### 37.1. Collar estimates and the singular part.

**Lemma 37.1** (Uniform vector collar estimates). *Let*

$$S(x) = D_\theta \log a_\theta(x)|_0 + \frac{V(x)}{w_0(x)}, \quad D = \{w_0 > 0\}.$$

There are  $q_0, C > 0$  such that, for  $0 < q < q_0$  and unit vectors  $u, v \in \mathbb{R}^r$ ,

$$\mathbb{E}_0[(u^t \mathbf{S})(v^t \mathbf{S}) \mathbf{1}_{\{w_0 \geq q\}}] = u^t \mathcal{J}_\Sigma v \log(1/q) + O(1), \quad (37.6)$$

$$\mathbb{E}_0[|u^t \mathbf{S}|^3 \mathbf{1}_{\{w_0 \geq q\}}] \leq Cq^{-1}, \quad (37.7)$$

$$\mathbb{E}_0[|u^t \mathbf{S}|^4 \mathbf{1}_{\{w_0 \geq q\}}] \leq Cq^{-2}, \quad (37.8)$$

$$\|\mathbb{E}_0[\mathbf{S} \mathbf{1}_{\{w_0 \geq q\}}]\| \leq Cq, \quad \mathbb{P}_0(0 < w_0 < q) \leq Cq^2. \quad (37.9)$$

The constants are uniform for the indicated unit directions and on compact families satisfying the displayed regularity bounds.

*Proof.* Use normal coordinates  $(y, s)$  with  $s = w_0$  in a fixed tubular neighborhood of  $\Sigma$ . The Jacobian is  $|\nabla w_0|^{-1} + O(s)$  and the null density is  $a_0 s$ . The only nonintegrable contribution to the second score moment is  $a_0(u^t V)(v^t V)/s$ , whose normal integral gives (37.6). The  $k$ th singular moment has normal integrand  $O(s^{1-k})$ , proving (37.7) and (37.8). The same coordinates give the quadratic collar mass. Differentiating  $\int f_\theta = 1$  at zero gives  $\mathbb{E}_0 \mathbf{S} = 0$  in the improper integrable sense; the omitted collar has first score moment  $O(q)$ .  $\square$

**Lemma 37.2** (Support-exclusive mass). *For every compact  $K \subset \mathbb{R}^r$  there is  $C_K$  such that, uniformly for  $h \in K$  and all large  $n$ ,*

$$P_{n,h}^{\perp 0}(\Omega) \leq C_K p_n \delta_n^2, \quad (37.10)$$

$$P_{n,0}\{f_{\delta_n h} = 0 < f_0\} \leq C_K p_n \delta_n^2. \quad (37.11)$$

Consequently, under either product experiment, the probability of seeing at least one support-exclusive observation is at most  $C_K n p_n \delta_n^2 = o(1)$ .

*Proof.* Joint  $C^1$  regularity gives  $|w_{\delta_n h} - w_0| \leq C_K \delta_n$  on the common collar. A point belonging to exactly one support therefore lies in  $|w_0| \leq C_K \delta_n$ . On that strip the relevant density is  $O(\delta_n)$  and the strip has volume  $O(\delta_n)$ . Multiplication by  $p_n$  gives the one-observation estimates. A union bound and (37.4) give the product estimate.  $\square$

**37.2. A reference-dominated common-collar representative.** Put

$$L_n = (\log(1/\delta_n))^{1/4}, \quad q_n = \delta_n L_n, \quad C_n = \{w_0 \geq q_n\}. \quad (37.12)$$

Introduce a second cemetery symbol  $\partial$  and let  $T_n : \Omega \rightarrow C_n \sqcup \{\dagger, \partial\}$  keep  $\dagger$  and every  $x \in C_n$ , while sending every successful  $x \notin C_n$  to  $\partial$ . Write

$$\bar{P}_{n,h} = T_n P_{n,h} = (1 - p_n) \delta_{\dagger} + p_n f_{\delta_n h} \mathbf{1}_{C_n} dx + p_n r_{n,h} \delta_{\partial}, \quad r_{n,h} = \int_{C_n^c} f_{\delta_n h} dx. \quad (37.13)$$

**Lemma 37.3** (Common-collar asymptotic equivalence). *For every compact  $K \subset \mathbb{R}^r$ ,*

$$\sup_{h \in K} r_{n,h} \leq C_K q_n^2, \quad \sup_{h \in K} n p_n r_{n,h} = o(1). \quad (37.14)$$

For all large  $n$ ,  $\{\bar{P}_{n,h} : h \in K\}$  is dominated by  $\bar{P}_{n,0}$ , and

$$\Delta\left(\{P_{n,h}^{\otimes n} : h \in K\}, \{\bar{P}_{n,h}^{\otimes n} : h \in K\}\right) \leq C_K n p_n q_n^2 = O((\log(1/\delta_n))^{-1/2}) = o(1). \quad (37.15)$$

Both comparison kernels can be chosen independently of  $h$ .

*Proof.* For  $h \in K$ , the moving support lies in  $\{w_0 > -C_K \delta_n\}$ . Integrating  $f_{\delta_n h} = O(w_0 + C_K \delta_n)$  over  $-C_K \delta_n < w_0 < q_n$  gives  $r_{n,h} = O(q_n^2)$ . Equation (37.4) gives the second bound.

On  $C_n$ ,  $w_{\delta_n h} \geq q_n - C_K \delta_n \geq q_n/2$  for all large  $n$ , uniformly on  $K$ , so all successful densities are mutually positive there. Moreover  $r_{n,0} > 0$ : the null collar  $0 < w_0 < q_n$  has positive mass.

Hence the  $\partial$  atom of the null censored law has positive mass, and the whole censored family is dominated by  $\bar{P}_{n,0}$ .

The forward kernel is the parameter-independent censoring map  $T_n$ . For a reverse kernel, keep  $C_n$  and  $\dagger$  fixed and send  $\partial$  to one fixed point of the original successful sample space. This reverse kernel is also parameter independent. Couple it with the original observation before censoring; disagreement is possible only if censoring occurred. A union bound gives probability at most  $np_nr_{n,h}$ , uniformly on  $K$ , proving (37.15).  $\square$

### 37.3. Uniform LAN on the representative.

**Lemma 37.4** (Uniform LAN). *For  $h$  in a fixed compact set define under  $\bar{P}_{n,0}^{\otimes n}$*

$$\Delta_n = \delta_n \sum_{i=1}^n Y_i S(X_i) \mathbf{1}_{C_n}(X_i), \quad (37.16)$$

$$Q_n = \delta_n^2 \sum_{i=1}^n Y_i S(X_i) S(X_i)^t \mathbf{1}_{C_n}(X_i), \quad (37.17)$$

where  $Y_i$  is the success indicator and, conditionally on success under the null,  $X_i$  has density  $f_0$ . Then

$$\log \frac{d\bar{P}_{n,h}^{\otimes n}}{d\bar{P}_{n,0}^{\otimes n}} = h^t \Delta_n - \frac{1}{2} h^t Q_n h + o_{\bar{P}_{n,0}}(1) \quad (37.18)$$

uniformly for  $h$  in compact subsets of  $\mathbb{R}^r$ , and

$$\Delta_n \Rightarrow N_r(0, \mathcal{J}_\Sigma), \quad Q_n \longrightarrow \mathcal{J}_\Sigma \quad \text{in probability.} \quad (37.19)$$

The probability that the atom  $\partial$  occurs in the product experiment is  $o(1)$  uniformly on compact local parameter sets.

*Proof.* The last assertion is Lemma 37.3. On its complement the likelihood ratio uses only the common support  $C_n$  and the unchanged failure atom. Taylor expansion before division by  $w_0$  gives, uniformly for bounded  $h$  and  $x \in C_n$ ,

$$\log \frac{f_{\delta_n h}(x)}{f_0(x)} = \delta_n h^t S(x) - \frac{\delta_n^2}{2} (h^t S(x))^2 + R_{n,h}(x).$$

The regular parameter remainder has expectation  $O(\delta_n^2)$  per successful observation, while the cubic logarithmic remainder is bounded in expectation by  $C_K \delta_n^3 \mathbb{E}_0(1 + \|S\|^3) \mathbf{1}_{C_n}$ . Lemma 37.1 and (37.4) give

$$np_n \delta_n^2 = o(1), \quad np_n \delta_n^3 q_n^{-1} = o(1),$$

which proves (37.18).

The truncated-score mean has norm at most  $Cq_n$ , so  $\|\mathbb{E}\Delta_n\| \leq Cnp_n \delta_n q_n = o(1)$ . Its covariance tends to  $\mathcal{J}_\Sigma$  because  $\log(1/q_n)/\log(1/\delta_n) \rightarrow 1$ . On  $C_n$  every scalar score contribution multiplied by  $\delta_n$  is  $O(L_n^{-1})$ ; Lindeberg therefore holds uniformly on the unit sphere and Cramér–Wold gives the first limit in (37.19). Finally

$$np_n \delta_n^4 \mathbb{E}_0[\|S\|^4 \mathbf{1}_{C_n}] \leq Cnp_n \delta_n^4 q_n^{-2} = o(1),$$

while the mean of  $Q_n$  tends to  $\mathcal{J}_\Sigma$ . Hence  $Q_n \rightarrow \mathcal{J}_\Sigma$  in probability.  $\square$

**Lemma 37.5** (Bounded-sequence contiguity). *If  $h_n$  is any bounded sequence, then  $\bar{P}_{n,h_n}^{\otimes n}$  and  $\bar{P}_{n,0}^{\otimes n}$  are mutually contiguous. The same is true for  $P_{n,h_n}^{\otimes n}$  and  $P_{n,0}^{\otimes n}$ .*

*Proof.* Along every subsequence, take a further subsequence on which  $h_n \rightarrow h$ . Lemma 37.4 gives joint convergence of the log likelihood ratio under the null to  $h^t Z - \frac{1}{2} h^t \mathcal{J}_\Sigma h$  with  $Z \sim N(0, \mathcal{J}_\Sigma)$ . The limiting likelihood is strictly positive and has mean one. The normalized-likelihood argument in Proposition 37.12 therefore applies: truncation gives uniform integrability and forward contiguity, while positivity of the null-law likelihood limit gives reverse contiguity. This argument uses only the null LAN expansion and unit likelihood means; it does not assume a central-sequence limit under the alternative. Since every subsequence admits such a further subsequence, the whole bounded sequence is mutually contiguous. The common-space censoring coupling in that proposition, with the uniform error from Lemma 37.3, transfers the assertion to the original experiments.  $\square$

**37.4. Finite likelihood experiments and normalization.** We record the finite-experiment passage used below, including its normalization hypothesis and both comparison kernels. It is a standard statistical fact; the proof is included to specify exactly what is used. The statement by itself makes no assertion on an infinite parameter set.

**Lemma 37.6** (Finite likelihood experiments). *For each  $n$ , let  $(P_{n,i})_{i=0}^m$  be probability measures on a standard Borel space, with  $P_{n,i} \ll P_{n,0}$ , and put  $L_{n,i} = dP_{n,i}/dP_{n,0}$ . Suppose that, under  $P_{n,0}$ ,*

$$(L_{n,1}, \dots, L_{n,m}) \Rightarrow (L_1, \dots, L_m) \quad \text{under } Q_0, \quad L_i \geq 0, \quad \mathbb{E}_{Q_0} L_i = 1.$$

*Here  $Q_0$  is a probability measure on a standard Borel space, the  $L_i$  are measurable, and  $Q_i = L_i Q_0$ . Then*

$$\Delta(\{P_{n,i} : 0 \leq i \leq m\}, \{Q_i : 0 \leq i \leq m\}) \longrightarrow 0. \quad (37.20)$$

*The kernels in both directions are independent of the unknown index  $i$ . No uniform integrability assumption beyond the stated normalization is needed: it follows from the likelihood convergence and the unit means.*

*Proof.* Write  $\mu_n$  and  $\mu$  for the respective laws of the likelihood vectors on  $[0, \infty)^m$ . Each coordinate has integral one under both measures. For every  $A > 0$ , weak convergence and bounded continuity of  $x \mapsto x_i \wedge A$  give

$$\int (x_i - A)_+ d\mu_n = 1 - \int (x_i \wedge A) d\mu_n \longrightarrow 1 - \int (x_i \wedge A) d\mu = \int (x_i - A)_+ d\mu.$$

Together with  $x_i \mathbf{1}_{\{x_i > 2A\}} \leq 2(x_i - A)_+$ , this proves uniform integrability of every coordinate (the finitely many initial indices are harmless), hence of the  $\ell^1$  norm.

There are couplings  $\pi_n$  of  $\mu_n, \mu$  for which

$$\int \|x - y\|_1 d\pi_n(x, y) \longrightarrow 0. \quad (37.21)$$

For completeness, choose a large cube in  $[0, \infty)^m$  whose outer faces are  $\mu$ -null, and partition it into finitely many rectangles of  $\ell^1$ -diameter at most  $\varepsilon$ , with boundaries that are  $\mu$ -null relative to  $[0, \infty)^m$ . Weak convergence makes all rectangle masses converge. Couple the minimum of the two masses inside each matching rectangle; couple the remaining masses arbitrarily. The matched part costs at most  $\varepsilon$ ; the unmatched part in the fixed cube tends to zero. Contributions involving its complement tend uniformly to zero as the cube grows, by the just-proved uniform integrability. Taking the sample size limit, then enlarging the cube and sending  $\varepsilon$  to zero, and choosing a diagonal sequence, proves (37.21).

Put  $x_0 = y_0 = 1$ . The likelihood-vector statistic identifies the first experiment exactly with  $(x_i \mu_n)_{i=0}^m$ , and the second with  $(y_i \mu)_{i=0}^m$ . Indeed, the forward kernel is the statistic; a reverse kernel is a regular conditional distribution under the reference measure given that statistic.

Multiplication by a likelihood coordinate recovers the corresponding alternative. Standard Borelness ensures the required conditional distributions exist.

Disintegrate  $\pi_n$  in the two directions. For its kernel from  $x$  to  $y$ , the difference between the simulated alternative and  $y_i\mu$  is the  $y$ -marginal of  $(x_i - y_i)\pi_n$ . Since this signed measure has total mass zero,

$$\|K_n(x_i\mu_n) - y_i\mu\|_{\text{TV}} \leq \frac{1}{2} \int |x_i - y_i| d\pi_n.$$

The reversed disintegration gives the identical bound in the other direction. Composing with the preceding statistic and reconstruction kernels proves (37.20). The same coupling and the same disintegrations serve every index, including  $i = 0$ ; no kernel receives the unknown index as input.  $\square$

*Remark 37.7* (Gaussian likelihoods and the scope of the lemma). If  $Z \sim N(0, J)$  with  $J$  positive semidefinite, then, for each fixed  $h$ ,

$$L_h = \exp\{h^t Z - \frac{1}{2} h^t J h\} > 0, \quad \mathbb{E} L_h = 1.$$

Thus joint LAN limits on any finite parameter set satisfy the normalization required in Lemma 37.6, also when  $J$  is singular or zero. One may adjoin the reference index to the finite set and then restrict the resulting kernels. For a moving-support family the lemma is applied to its reference-dominated common-collar representative, not to nonexistent reference likelihoods on support-exclusive observations. Compact-parameter convergence additionally uses the uniform modulus and finite-net argument in Theorem 40.2; it does not follow from the finite lemma alone.

**37.5. The Gaussian limit on the identifiable quotient.** Let

$$H = \text{Ran}(\mathcal{J}_\Sigma), \quad \Pi_H : \mathbb{R}^r \rightarrow H \quad (37.22)$$

be the orthogonal projection. The restriction  $\mathcal{J}_H = \mathcal{J}_\Sigma|_H$  is positive definite. The Gaussian limit experiment is most cleanly written in central-sequence coordinates as

$$\Delta \sim N_H(\mathcal{J}_H \Pi_H h, \mathcal{J}_H). \quad (37.23)$$

Equivalently,  $Z = \Pi_H h + \mathcal{J}_H^{-1/2} G$  with standard Gaussian  $G$  on  $H$ . Directions in  $\ker \mathcal{J}_\Sigma$  are not identifiable at the critical logarithmic scale.

**Theorem 37.8** (Moving-support vector boundary Gaussian limit). *Under (37.2)–(37.4):*

- (1) *On every compact local parameter set, the original experiment is at  $o(1)$  Le Cam distance from the reference-dominated common-collar representative, by (37.15). The representative satisfies the compact- uniform LAN expansion (37.18) and the central limits (37.19).*
- (2) *For every finite collection  $h_1, \dots, h_m$ , the corresponding finite local subexperiment converges in Le Cam distance to the restriction of the Gaussian experiment (37.23) to those parameters. More generally, every bounded local sequence satisfies the contiguity and third-lemma conclusions implied by the LAN expansion. We do not identify this finite-subexperiment statement with a separate assertion of global deficiency convergence for an uncountable parameter set. The stronger conclusion on every fixed compact set is proved in Theorem 40.2.*
- (3) *For every fixed  $h$ ,*

$$\|P_{n,h}^{\otimes n} - P_{n,0}^{\otimes n}\|_{\text{TV}} \longrightarrow 2\Phi\left(\frac{\sqrt{h^t \mathcal{J}_\Sigma h}}{2}\right) - 1. \quad (37.24)$$

*If  $\mathcal{J}_\Sigma$  is positive definite, the reference-model statistic*

$$\hat{h}_n = \mathcal{J}_\Sigma^{-1} \Delta_n \quad (37.25)$$



(with observations outside  $C_n$  contributing zero) satisfies  $\hat{h}_n \Rightarrow N_r(h, \mathcal{J}_\Sigma^{-1})$  uniformly on compact local sets. If  $\mathcal{J}_\Sigma$  is singular, the corresponding estimator of the identifiable coordinate is  $\hat{h}_n^{\text{id}} = \mathcal{J}_\Sigma^+ \Delta_n$ , with limit centered at  $\Pi_{\mathbf{H}} h$  and covariance  $\mathcal{J}_\Sigma^+$ .

- (4) Let  $W$  be positive semidefinite on  $\mathbf{H}$  and  $\ell_M(z) = \min\{z^t W z, M\}$ . The standard local asymptotic minimax theorem applied to the LAN experiment on  $\mathbf{H}$  gives

$$\lim_{R \rightarrow \infty} \liminf_{n \rightarrow \infty} \inf_{T_n} \sup_{\substack{h \in \mathbf{H} \\ \|h\| \leq R}} \mathbb{E}_{n,h} \ell_M(T_n - h) \geq \mathbb{E} \ell_M(Z_0), \quad Z_0 \sim N_{\mathbf{H}}(0, \mathcal{J}_{\mathbf{H}}^{-1}). \quad (37.26)$$

Letting  $M \rightarrow \infty$  gives  $\text{tr}(W \mathcal{J}_{\mathbf{H}}^{-1})$ . When  $\mathcal{J}_\Sigma$  is positive definite, this is the full-parameter bound.

*Proof.* Item 1 is Lemmas 37.3 and 37.4. For any finite parameter set, the joint likelihood-ratio vector is a continuous function of  $\Delta_n, Q_n$  plus a uniform  $o_P(1)$  term. Equations (37.19) therefore give joint weak convergence of the likelihood-ratio vector to the Gaussian likelihood vector in (37.23). Each prelimit likelihood has mean one under the censored reference law; each limiting coordinate has mean one by Remark 37.7. Adjoin the reference parameter to the finite set, apply Lemma 37.6, and then restrict. Lemma 37.3 transfers the resulting two-sided comparison to the original family. Lemma 37.5 gives the bounded-sequence statement.

Le Cam's third lemma translates the limiting central sequence by  $\mathcal{J}_\Sigma h$ . The equal-covariance Gaussian total-variation formula gives (37.24). The estimator statements follow by linear transformation of the central sequence; in the singular case the Moore–Penrose inverse acts only on  $\mathbf{H}$ .

We verify the moment assertion under the original local alternatives. Fix a compact set  $K$  and write  $s = w_0(x)$ . On  $C_n$ , Taylor expansion of  $a_\theta w_\theta$  before division by  $a_0 s$  gives

$$\frac{f_{\delta_n h}(x)}{f_0(x)} = 1 + \delta_n h^t \mathbf{S}(x) + O_K(\delta_n^2(1 + s^{-1})), \quad h \in K. \quad (37.27)$$

The bound is uniform because  $s \geq q_n$  and  $w_{\delta_n h} \geq q_n/2$  there. In particular, the ratio for one successful observation tends uniformly to one on  $C_n$ ; no such assertion is made for the product likelihood. In the normal coordinate  $s$ ,  $\|\mathbf{S}\| = O(1 + s^{-1})$  and  $f_0 dx = O(s) ds d\sigma$ . Thus integration of the remainder in (37.27) against  $\mathbf{S} f_0$  yields

$$\begin{aligned} \mathbb{E}_{n,h} \Delta_n &= np_n \delta_n \mathbb{E}_0[\mathbf{S} \mathbf{1}_{C_n}] + np_n \delta_n^2 \mathbb{E}_0[\mathbf{S} \mathbf{S}^t \mathbf{1}_{C_n}] h \\ &\quad + O_K(np_n \delta_n^3 \log(1/q_n)) = \mathcal{J}_\Sigma h + o_K(1). \end{aligned} \quad (37.28)$$

Here the first term is  $O(np_n \delta_n q_n) = o(1)$  by (37.9); the second converges uniformly on  $K$  by (37.6) and  $\log(1/q_n)/\log(1/\delta_n) \rightarrow 1$ . The displayed remainder also tends to zero. This establishes, in particular, a uniformly bounded mean.

Put  $\xi_{n,i} = \delta_n Y_i \mathbf{S}(X_i) \mathbf{1}_{C_n}(X_i)$ , with value zero for a failure or a successful observation outside  $C_n$ . Under each fixed local alternative these vectors are independent. The uniformly bounded one-observation density ratio and the collar moment estimates give

$$\sum_{i=1}^n \mathbb{E}_{n,h} \|\xi_{n,i}\|^2 \leq C_K np_n \delta_n^2 \log(1/q_n), \quad \sum_{i=1}^n \mathbb{E}_{n,h} \|\xi_{n,i}\|^4 \leq C_K np_n \delta_n^4 q_n^{-2}.$$

Expanding the fourth power of the centered sum coordinate by coordinate therefore gives

$$\mathbb{E}_{n,h} \|\Delta_n - \mathbb{E}_{n,h} \Delta_n\|^4 \leq C_K \left\{ [np_n \delta_n^2 \log(1/q_n)]^2 + np_n \delta_n^4 q_n^{-2} \right\} = O_K(1). \quad (37.29)$$

Indeed, for  $\ell_n = \log(1/\delta_n)$  and  $B_n = np_n \delta_n^2 \ell_n \rightarrow 1$ , the relevant quantities are

$$np_n \delta_n q_n = B_n \ell_n^{-3/4}, \quad np_n \delta_n^2 \log(1/q_n) = B_n \left( 1 - \frac{\log \ell_n}{4\ell_n} \right), \quad np_n \delta_n^4 q_n^{-2} = B_n \ell_n^{-3/2}.$$

Restoring the bounded mean in (37.28) proves  $\sup_{n \geq n_0(K)} \sup_{h \in K} \mathbb{E}_{n,h} \|\Delta_n\|^4 < \infty$ . Each excluded observation contributes zero; the other summands are retained even when censoring occurs elsewhere in the sample. This proves uniform integrability for compact-local quadratic risk, also after applying the fixed linear map  $\mathcal{J}_\Sigma^+$  on  $\mathbf{H}$ . Finally apply the classical finite-dimensional LAN local asymptotic minimax theorem on the identifiable space  $\mathbf{H}$ ; see [32, 34].  $\square$

*Remark 37.9* (Reference-model nature of the central sequence). The statistics (37.16) and (37.25) use the score and information at the reference geometry. They are local-experiment statistics, not globally adaptive estimators for an unknown reference model. The physical theorem below uses them only after a reference channel and a finite local chart have been fixed.

### 37.6. Quantitative stability of the moving boundary.

**Lemma 37.10** (Linearly vanishing support stability). *Let  $f = A(w)_+$  and  $\tilde{f} = \tilde{A}(\tilde{w})_+$  be normalized densities on the same bounded open set. Suppose  $A, \tilde{A}$  are uniformly positive and  $C^1$ , the zero sets of  $w, \tilde{w}$  are uniformly regular in one fixed tubular collar, and*

$$\|w - \tilde{w}\|_{C^1} \leq \varepsilon < 1/2, \quad \|A - \tilde{A}\|_{C^0} \leq \eta.$$

*Then, with the convention (37.1),*

$$H^2(f \, dx, \tilde{f} \, dx) \leq C \left\{ \varepsilon^2 \log \frac{e}{\varepsilon} + \eta^2 \right\}. \quad (37.30)$$

*The estimate is unchanged, up to the uniform constant, if a common positive smooth factor is moved between the amplitude and defining function.*

*Proof.* Use the tubular coordinate  $s = w$ . On  $|s| \leq C\varepsilon$  the total mass of both densities is  $O(\varepsilon^2)$ . On  $s > C\varepsilon$  the supports agree and

$$|\sqrt{Aw} - \sqrt{\tilde{A}\tilde{w}}|^2 \leq C \left( \eta^2 s + \frac{\varepsilon^2}{s} \right).$$

The first term is  $O(\eta^2)$  after integration and the second is  $O(\varepsilon^2 \log(e/\varepsilon))$ . The same estimate holds using  $\tilde{w}$  as the normal coordinate on the other side of the collar; away from the collar all denominators are uniformly bounded. Multiplication of  $w, \tilde{w}$  by the same uniformly positive smooth factor merely changes the tubular Jacobian and amplitude bounds by uniform constants.  $\square$

*Remark 37.11* (Terminology). We reserve “LAN expansion” for the reference-dominated common-collar representative. For the original family we first prove asymptotic equivalence to that representative and the finite-subexperiment Gaussian limit, testing profile, contiguity, and minimax consequences. The uniform moving-boundary modulus then upgrades the Gaussian comparison to two-sided Le Cam convergence on every fixed compact local set in Theorem 40.2; Corollary 40.3 supplies its physical endpoint-output counterpart. The original family already has a common dominating measure. Reference absolute continuity, local asymptotic normality, and convergence of experiments are distinct assertions.

**37.7. Likelihood tilting and compact-local quadratic risk.** We give the likelihood-tilting and moment arguments underlying Lemma 37.5 and the estimator assertions of Theorem 37.8. This also specifies how the unit-mean normalization in Lemma 37.6 is used under local alternatives. The argument concerns the moving-support model (37.5); physical acquisition is handled by the separate stopped-transfer results.

Three distinct estimates enter the conclusion. The density expansion (37.27) gives the alternative mean (37.28); centering under that same alternative then gives (37.29). These are direct moment estimates for the original product law. The normalized likelihoods and the LAN

expansion instead give the weak shifted limit (37.31). Combining that weak limit with the original-law fourth-moment bound yields (37.33). Thus the argument for unbounded quadratic loss does not rely on total-variation comparison alone. The further passage from finite parameter sets to compact experiments uses the uniform modulus in Theorem 40.2.

**Proposition 37.12** (Likelihood tilting and uniform quadratic risk). *Under the hypotheses of Theorem 37.8, let  $h_n \rightarrow h$  be a bounded local sequence. Under the reference-dominated common-collar experiment, set*

$$L_n = \frac{d\bar{P}_{n,h_n}^{\otimes n}}{dP_{n,0}^{\otimes n}}.$$

*The likelihoods  $L_n$  are uniformly integrable under their reference measures. The reference and alternative sequences are mutually contiguous, and, for every bounded continuous  $\varphi$ ,*

$$\mathbb{E}_{\bar{P}_{n,h_n}^{\otimes n}} \varphi(\Delta_n) \longrightarrow \mathbb{E} \left[ e^{h^t Z - h^t \mathcal{J}_\Sigma h/2} \varphi(Z) \right], \quad Z \sim N(0, \mathcal{J}_\Sigma). \quad (37.31)$$

*The limit is the expectation under  $N(\mathcal{J}_\Sigma h, \mathcal{J}_\Sigma)$ , including singular information. The same central-sequence limit holds in the original experiment, with observations outside  $C_n$  contributing zero.*

*For every compact  $K \subset \mathbb{R}^r$ , the original product experiment satisfies*

$$\sup_{n \geq n_0(K)} \sup_{h \in K} \mathbb{E}_{P_{n,h}^{\otimes n}} \|\Delta_n\|^4 < \infty. \quad (37.32)$$

*Write  $J = \mathcal{J}_\Sigma$ ,  $\mathbf{H} = \text{Ran } J$ , and  $\Pi = \Pi_{\mathbf{H}}$ . If  $W$  is positive semidefinite on  $\mathbf{H}$ , then*

$$\sup_{h \in K} \left| \mathbb{E}_{P_{n,h}^{\otimes n}} [(J^+ \Delta_n - \Pi h)^t W (J^+ \Delta_n - \Pi h)] - \text{tr}(W J^+) \right| \longrightarrow 0. \quad (37.33)$$

*Here  $J^+$  acts only on the identifiable space; no risk assertion for an unidentifiable coordinate is made.*

*Proof.* The uniform LAN expansion gives joint convergence under the reference law,

$$(\Delta_n, L_n) \Rightarrow (Z, L), \quad L = e^{h^t Z - h^t \mathcal{J}_\Sigma h/2} > 0.$$

Both  $\mathbb{E} L_n$  and  $\mathbb{E} L$  equal one. For fixed  $A > 0$ , convergence of the bounded continuous truncation gives

$$\mathbb{E}(L_n - A)_+ = 1 - \mathbb{E}(L_n \wedge A) \longrightarrow 1 - \mathbb{E}(L \wedge A) = \mathbb{E}(L - A)_+.$$

Since  $L_n \mathbf{1}_{\{L_n > 2A\}} \leq 2(L_n - A)_+$ , this proves uniform integrability, without any assumption on higher likelihood moments. Apply joint weak convergence to  $(L_n \wedge A) \varphi(\Delta_n)$ , then let  $A \rightarrow \infty$ . The tail bound just proved yields (37.31). Completing the Gaussian exponential, or computing its characteristic function on  $\text{Ran } J$ , identifies the displayed shifted normal law. This calculation does not invert  $J$  on its kernel.

For completeness, let  $P_n$  and  $Q_n = L_n P_n$  denote the censored reference and alternative product measures. For every event  $A_n$ ,

$$Q_n(A_n) \leq A P_n(A_n) + \mathbb{E}_{P_n} [L_n \mathbf{1}_{\{L_n > A\}}].$$

Uniform integrability proves contiguity of  $Q_n$  with respect to  $P_n$ . In the reverse direction, for every  $c > 0$ ,

$$P_n(A_n) \leq P_n\{L_n \leq c\} + c^{-1} Q_n(A_n).$$

First take a continuity value  $c$  for the law of  $L$ , and then send  $c \downarrow 0$ . Positivity of  $L$  proves reverse contiguity. Subsequence extraction proves the assertion for every bounded sequence, not only a convergent one. Embed the original and censored observations in their common space with the extra symbol  $\partial$ . Their laws differ only on an event of probability at most  $C_K n p_n q_n^2 = o(1)$ ,

uniformly on  $K$ . This transfers the contiguity assertions to the original experiments. The central sequence is unchanged by censoring, since every censored observation already contributes zero to its definition.

We next prove the moment estimate directly under the original alternatives. On the common collar  $C_n$ , joint smoothness and positivity give the *one-successful-observation* density bound

$$\left| \frac{f_{\delta_n h}}{f_0} - 1 \right| \leq C_K \delta_n (1 + w_0^{-1}), \quad \sup_{h \in K} \left\| \frac{f_{\delta_n h}}{f_0} - 1 \right\|_{L^\infty(C_n)} = O_K(\delta_n / q_n) = o(1). \quad (37.34)$$

This is not a bound tending to one for the product likelihood. The collar estimates and (37.34) imply

$$\left\| \int_{C_n} S f_{\delta_n h} dx \right\| \leq C q_n + C_K \delta_n \log(1/q_n), \quad (37.35)$$

$$\int_{C_n} \|S\|^2 f_{\delta_n h} dx \leq C_K \log(1/q_n), \quad \int_{C_n} \|S\|^4 f_{\delta_n h} dx \leq C_K q_n^{-2}. \quad (37.36)$$

For the first inequality, the reference truncated mean is  $O(q_n)$ . In normal coordinates  $s = w_0$ , the remaining integrand is bounded by  $C_K \delta_n (s + 2 + s^{-1})$ ; integration from  $q_n$  to a fixed collar width gives the logarithm. The complement of the collar is uniformly regular. The other inequalities follow from the uniformly bounded density ratio and Lemma 37.1.

Put  $\xi_{n,i} = \delta_n Y_i S(X_i) \mathbf{1}_{C_n}(X_i)$ . These vectors are independent under every fixed alternative. Their sum is  $\Delta_n$ , and (37.35) gives

$$\sup_{h \in K} \|\mathbb{E}_{n,h} \Delta_n\| \leq C n p_n \delta_n q_n + C_K n p_n \delta_n^2 \log(1/q_n) = O_K(1).$$

For centered independent vectors in fixed dimension, expansion of the fourth power, coordinate by coordinate, yields

$$\begin{aligned} \mathbb{E}_{n,h} \|\Delta_n - \mathbb{E}_{n,h} \Delta_n\|^4 &\leq C_r \left\{ \left( \sum_{i=1}^n \mathbb{E}_{n,h} \|\xi_{n,i}\|^2 \right)^2 + \sum_{i=1}^n \mathbb{E}_{n,h} \|\xi_{n,i}\|^4 \right\} \\ &\leq C_K \left\{ [n p_n \delta_n^2 \log(1/q_n)]^2 + n p_n \delta_n^4 q_n^{-2} \right\} = O_K(1). \end{aligned}$$

Indeed, with  $q_n = \delta_n [\log(1/\delta_n)]^{1/4}$ , the first mean term  $n p_n \delta_n q_n$  and the final fourth-moment term tend to zero, whereas  $n p_n \delta_n^2 \log(1/q_n) \rightarrow 1$ . This proves (37.32); each excluded observation, rather than the entire product sample on a censoring event, contributes zero.

The tilted central-sequence limit and (37.32) imply convergence of quadratic risks along every convergent local sequence. The Gaussian error  $J^+ Z - \Pi h$ , for  $Z \sim N(Jh, J)$ , has covariance  $J^+$  and mean zero. Its  $W$ -risk is therefore  $\text{tr}(W J^+)$ . If (37.33) were not uniform on  $K$ , a violating sequence would have a convergent subsequence in  $K$ , contradicting the just-proved sequential limit. This proves the compact-uniform assertion.  $\square$

### 38. PHYSICAL REALIZATION OF ANCHORED FINITE CONTACT-JET FAMILIES

The statistical experiment below is formulated in one fixed laboratory chart. We record that the resulting finite-dimensional contact-jet models are realized by actual nearby billiard tables. The contact patch used by the statistical theorem is chosen inside the region where every cutoff is exactly one, so the displayed polynomial jet family is literally the physical graph on the whole observation patch.

**Proposition 38.1** (Anchored realization of finite labelled jets). *Let a smooth periodic dispersing billiard contain a selected labelled closest pair with positive gap  $g_0$ , positive contact curvatures, and a strict clearance margin separating this pair from every competing collision segment outside*

fixed contact neighborhoods. Fix  $M \geq 2$ . Choose one Euclidean laboratory chart in which the reference contacts have common transverse coordinate  $y = 0$  and common tangent direction.

There exist numbers  $0 < u_* < u_{**}$  and a smooth finite-dimensional family of actual periodic billiard tables through the reference table, with coordinates

$$(g, q_{0,2}, \dots, q_{0,M}, q_{1,2}, \dots, q_{1,M}),$$

such that on the whole statistical patch  $|y| < u_*$  the two participating boundaries are exactly

$$(-\psi_0(u), u), \quad (g + \psi_1(v), v), \quad (38.1)$$

with

$$\psi_b(0) = \psi'_b(0) = 0, \quad \psi_b^{(n)}(0) = q_{b,n} \quad (2 \leq n \leq M).$$

The labels, transverse origin, and tangent orientation are the same throughout the family. Strict convexity, embeddedness of every participating obstacle, disjointness of distinct obstacle closures and of their relevant lattice translates, and the selected-channel clearance margin persist uniformly after the parameter neighborhood is made sufficiently small.

The same statement holds if the two contacts lie on different labelled obstacles, on two lattice translates, or in two disjoint contact neighborhoods of one connected periodic obstacle, provided lattice-related perturbations are made equivariantly.

*Proof.* Choose disjoint arclength neighborhoods of the two reference contacts so small that the fixed laboratory graph representation is valid. Select smooth cutoffs  $\chi_b$  supported in  $|y| < u_{**}$  and identically one on a smaller interval  $|y| < 2u_*$ . We choose the statistical endpoint box only afterwards, inside  $|y| < u_*$ . Thus every graph used by the physical statistical theorem lies strictly inside the plateau of the cutoff.

For small parameters set

$$\delta\psi_b(y) = \chi_b(y) \sum_{n=2}^M \frac{\delta q_{b,n}}{n!} y^n. \quad (38.2)$$

On  $|y| < u_*$  this is exactly the displayed polynomial perturbation and hence changes precisely the prescribed finite jet without moving the contact origin or tangent line. A gap variation is realized by a normal displacement of the second patch which is constant on  $|y| < 2u_*$  and is cut off before the boundary leaves the chosen coordinate neighborhood. Its constant value on the statistical patch is absorbed into  $g$  in (38.1).

If the two contacts lie on one connected obstacle, take the two cutoff supports disjoint. The remaining part of the boundary is unchanged. A sufficiently small  $C^1$  perturbation of an embedded compact curve remains embedded; a sufficiently small  $C^2$  perturbation preserves strict positive curvature. The reference table has positive separation of distinct obstacle closures and a strict channel-isolation margin away from the chosen contact patches. On a compact fundamental cell these inequalities persist under small  $C^1$  perturbations. For periodic translates, the unchanged lattice and a uniform obstacle diameter leave only finitely many nearby translates to check. Shrinking the parameter ball therefore preserves global embeddedness, disjointness and channel isolation, including when both patches belong to the same connected obstacle.

Because each cutoff is constant on the statistical patch, the map from bump coefficients to the displayed contact jets has identity Jacobian. The inverse function theorem gives an open finite-dimensional coordinate family.  $\square$

*Remark 38.2* (Analytic versus smooth realization). The proposition serves the finite-dimensional statistical local experiment, which needs only finite smoothness on compact parameter sets. The analytic continuation conclusions of Theorems 22.1 and 25.3 are separate deterministic statements for analytic boundary germs; no compactly supported analytic bump is used there.

### 39. A PHYSICAL ENDPOINT EXPERIMENT ON FIXED LABORATORY WINDOWS

We now realize the boundary-information theorem on a parameter-independent physical observation space. Three points are fixed at the outset. First, all acquisition times are programmed from the reference channel and are unchanged under local alternatives. Second, the statistical theorem is formulated for the anchored registered family of Section 38, in one laboratory chart. Third, we distinguish the sigma-fields used at the successive transfer steps. The finite-to-boundary theorem does not compare the complete growing collision array.

**39.1. The declared physical record hierarchy.** A physical preparation has a full raw collision history, denoted  $\mathcal{R}^{\text{full}}$ , consisting of all collision points, obstacle labels, flight times and collision count observed during that preparation. We use three successively coarser records.

- (1) The *per-preparation endpoint-time record* is the measurable coarsening

$$\mathcal{C}_{j,t}(\mathcal{R}^{\text{full}}) = \begin{cases} (y_{\text{first}}, y_{\text{last}}, r), & \text{if the prescribed bridge succeeds,} \\ \dagger, & \text{otherwise,} \end{cases} \quad (39.1)$$

where  $r$  is the residual preparation-to-first-impact time used in Section 34. This is the record for which the historical finite-to-boundary total-variation theorem is proved.

- (2) A policy produces the *stopped endpoint-time acquisition transcript*: every chosen design, every record (39.1), the success/failure sequence and the stopping time are retained. The adaptive coupling theorem of Section 35 compares the finite and boundary laws of this entire stopped transcript at the declared record level.
- (3) The *endpoint-output experiment* is a further common Markov coarsening which retains the selected endpoint pairs needed for inference and the cap cemetery symbol. The Gaussian information theorem below is a theorem for this endpoint-output sigma-field.

The complete raw collision history  $\mathcal{R}^{\text{full}}$  may contain more information than item 1, and the stopped endpoint-time transcript may contain more information than item 3 through the waiting counts. No theorem below identifies these different sigma-fields.

**39.2. An anchored registered channel in one laboratory chart.** Fix once and for all a Euclidean chart  $(x, y)$ , two obstacle labels, and two facing patches. In this chart the admissible local family has the form

$$\Gamma_{0,\vartheta} = \{(-\psi_{0,\vartheta}(u), u) : |u| < u_*\}, \quad \Gamma_{1,g,\vartheta} = \{(g + \psi_{1,\vartheta}(v), v) : |v| < u_*\}, \quad (39.2)$$

where

$$\psi_{b,\vartheta}(0) = \psi'_{b,\vartheta}(0) = 0, \quad \psi''_{b,\vartheta}(0) > 0. \quad (39.3)$$

The labels, transverse origin and tangent direction are fixed throughout the family. Varying  $g$  translates the second graph in the normal  $x$  direction; it does not re-center the observed transverse coordinate. The local parameter  $\vartheta \in \Theta \subset \mathbb{R}^r$  changes only the anchored graph geometry. For the finite-dimensional theorem below,  $\vartheta$  is a finite labelled contact-jet coordinate realized by Proposition 38.1.

The endpoint observation used after the residual-time coarsening is the same measurable map at every parameter:

$$\mathcal{O}_{bb}(\mathcal{R}^{\text{full}}) = \begin{cases} (y_{\text{first}}, y_{\text{last}}), & \text{if the prescribed } bb \text{ bridge succeeds,} \\ \dagger, & \text{otherwise.} \end{cases} \quad (39.4)$$

Thus endpoint positions are literal laboratory  $y$  coordinates rather than parameter-dependent contact coordinates.

Let  $S_{b,g,\vartheta}$  and  $B_{b,g,\vartheta}$  denote the half-line actions and positive amplitudes constructed in the same fixed transverse coordinate. At  $(g_0, 0)$  write

$$S_b = S_{b,g_0,0}, \quad B_b = B_{b,g_0,0}, \quad \dot{S}_b(u)h = D_{\vartheta}S_{b,g_0,\vartheta}(u)|_0[h].$$

The anchoring gives  $S_{b,g,\vartheta}(0) = S'_{b,g,\vartheta}(0) = 0$  throughout the family.

Fix finitely many same-type designs  $(b_\ell, d_\ell)$  with  $b_\ell \in \{0, 1\}$  and  $0 < d_\ell < d_*$ ,  $1 \leq \ell \leq L$ . A same-type design is always implemented with an *even* number of alternating flights. For even  $j$  its programmed acquisition time is

$$t_{j,\ell} = jg_0 + d_\ell. \quad (39.5)$$

The parity convention is part of the physical design: an odd flight number would terminate at the opposite contact and would instead involve the mixed action  $S_b \oplus S_{1-b}$ .

**39.3. The multirate local parameter and fixed-time expansion.** Let

$$j_n \in 2\mathbb{N}, \quad j_n \longrightarrow \infty, \quad (39.6)$$

and let  $\delta_n \downarrow 0$ . For bounded  $z = (a, h) \in \mathbb{R} \times \mathbb{R}^r$  define

$$g_n(a) = g_0 + \frac{\delta_n}{j_n}a, \quad \vartheta_n(h) = \delta_n h. \quad (39.7)$$

The fixed laboratory time has exact excess

$$t_{j_n,\ell} - j_n g_n(a) = d_\ell - \delta_n a. \quad (39.8)$$

Thus a gap displacement of order  $\delta_n/j_n$  and a shape displacement of order  $\delta_n$  enter the moving endpoint support at the same first order.

For a same-type design  $(b, d)$  define

$$w_{b,d}(u, v) = d - S_b(u) - S_b(v). \quad (39.9)$$

Let  $A_{b,d} > 0$  be the normalized smooth amplitude, so the base conditional endpoint density is  $f_{b,d} = A_{b,d}(w_{b,d})_+$ . Set

$$U_b(u, v)z = a + \dot{S}_b(u)h + \dot{S}_b(v)h \quad (39.10)$$

and

$$\mathcal{K}_{b,d} = \int_{\{S_b(u)+S_b(v)=d\}} \frac{A_{b,d}(u, v) U_b(u, v)^t U_b(u, v)}{|\nabla(S_b \oplus S_b)(u, v)|} d\sigma. \quad (39.11)$$

The derivatives  $\dot{S}_b$  in this display are the historical half-line derivatives in the fixed transverse coordinate at fixed reference gap. The physical acquisition time itself is held fixed. Its gap variation is the explicit term  $-\delta_n a$  in (39.8); the remaining smooth gap dependence of the half-line action is only  $(\delta_n a/j_n)\partial_g S_b$  and belongs to the quantitative remainder below. Thus channel-centered action differentiation is not confused with differentiation of the laboratory time.

**Proposition 39.1** (Quantitative fixed-window reduction). *For every compact local set  $K \subset \mathbb{R}^{1+r}$ , the exact limiting boundary endpoint density at (39.5) can be written*

$$f_{n,z}^{\text{ex}}(u, v) = A_{n,z}^{\text{ex}}(u, v)(w_{n,z}^{\text{ex}}(u, v))_+, \quad z \in K, \quad (39.12)$$

where, uniformly on the common endpoint box,

$$\|w_{n,z}^{\text{ex}} - (w_{b,d} - \delta_n U_b z)\|_{C^2} \leq C_K r_n, \quad r_n := \delta_n^2 + \frac{\delta_n}{j_n}, \quad (39.13)$$

$$\|A_{n,z}^{\text{ex}} - A_{b,d}\|_{C^1} \leq C_K \delta_n. \quad (39.14)$$

Let  $\tilde{f}_{n,z}$  be the normalized ideal density

$$\tilde{f}_{n,z}(u, v) = \frac{A_{b,d}(u, v)(w_{b,d}(u, v) - \delta_n U_b(u, v)z)_+}{\int A_{b,d}(w_{b,d} - \delta_n U_b z)_+ du dv}. \quad (39.15)$$

Then, with the Hellinger convention (37.1),

$$\sup_{z \in K} H^2(f_{n,z}^{\text{ex}}, \tilde{f}_{n,z}) \leq C_K \left\{ r_n^2 \log \frac{e}{r_n} + \delta_n^2 \right\}. \quad (39.16)$$

If

$$k_n \delta_n^2 \log(1/\delta_n) \rightarrow 1, \quad (39.17)$$

then the exact and ideal  $k_n$ -record endpoint experiments have Le Cam distance  $o(1)$ , uniformly on  $K$ .

*Proof.* Equation (39.8) supplies the exact fixed-time contribution  $-\delta_n a$ . Smooth dependence of the half-line action on  $(g, \vartheta)$  gives, with two endpoint derivatives and uniformly on  $K$ ,

$$S_{b,g_n(a),\vartheta_n(h)} = S_b + \delta_n \dot{S}_b h + \frac{\delta_n a}{j_n} \partial_g S_b + O_K(\delta_n^2 + \delta_n^2/j_n^2).$$

Substitution at both endpoints proves (39.13); smooth dependence of the flux amplitude and its normalizing constant gives (39.14).

The ideal normalizing denominator is  $1 + O_K(\delta_n)$ . Lemma 37.10 with support error  $C_K r_n$  and amplitude error  $C_K \delta_n$  proves (39.16). The product error is negligible at (39.17): if  $j_n^{-1} \leq \delta_n$ , then  $r_n \leq 2\delta_n^2$  and  $k_n r_n^2 \log(e/r_n) = O(\delta_n^2)$ ; if  $j_n^{-1} > \delta_n$ , then  $r_n \leq 2\delta_n/j_n$  and the same quantity is  $O(j_n^{-2})$ . The amplitude term is  $k_n \delta_n^2 = O(1/\log(1/\delta_n))$ . Product Hellinger subadditivity and  $\|P - Q\|_{\text{TV}} \leq H(P, Q)$  up to the fixed convention give the conclusion.  $\square$

#### 39.4. Finite positive designs from the tangent inverse.

**Lemma 39.2** (Finite positive-offset design). *Fix a finite labelled contact-jet order  $M$ . Let  $\vartheta$  consist of the two labelled contact jets through order  $M$  at fixed gap in the anchored chart. There exist finitely many same-type designs  $(b_\ell, d_\ell)$  and positive weights  $\alpha_\ell$ ,  $\sum_\ell \alpha_\ell = 1$ , such that*

$$\mathcal{K}_{\text{design}} = \sum_{\ell=1}^L \alpha_\ell \mathcal{K}_{b_\ell, d_\ell} \quad (39.18)$$

*is positive definite on the full local coordinate  $(a, h)$ . The finite design may depend on  $M$  and on the reference geometry; no universal finite design for the infinite jet or for all reference tables is asserted.*

*Proof.* Suppose  $z = (a, h)$  lies in the kernel of  $\mathcal{K}_{b,d}$  for both labels and all sufficiently small  $d > 0$ . Positivity of the weight in (39.11) and continuity imply

$$a + \dot{S}_b(u)h + \dot{S}_b(v)h = 0 \quad \text{whenever} \quad S_b(u) + S_b(v) = d.$$

Take  $v = 0$  and  $d = S_b(u)$ . Anchoring gives  $\dot{S}_b(0)h = 0$ , so  $a + \dot{S}_b(u)h = 0$ . Letting  $u \rightarrow 0$  gives  $a = 0$ ; hence  $\dot{S}_b(u)h = 0$  for both labels near zero. All action-jet variations through order  $M$  vanish, and Proposition 22.7 gives  $h = 0$ . Thus the common kernel is trivial. Every unit vector is detected by some positive design; compactness of the unit sphere selects finitely many. Any positive convex combination of these designs is positive definite.  $\square$



**39.5. A single deterministic cap sequence for all compact local sets.** Let  $p_{n,\ell}^\partial(z)$  be the one-preparation success probability in the boundary model at design  $\ell$ , with the fixed even flight number  $j_n$ . We use one reference-based cap sequence rather than a cap chosen after a compact local set has been specified.

**Lemma 39.3** (Reference-based cap). *Assume*

$$j_n \delta_n \rightarrow 0. \quad (39.19)$$

*Then, for every fixed compact local set  $K$  and each design  $\ell$ ,*

$$\sup_{z \in K} \left| \frac{p_{n,\ell}^\partial(z)}{p_{n,\ell}^\partial(0)} - 1 \right| \rightarrow 0. \quad (39.20)$$

*Choose, before quantifying over  $K$ ,*

$$N_{n,\ell} = \left\lceil \frac{4k_{n,\ell}}{p_{n,\ell}^\partial(0)} \right\rceil. \quad (39.21)$$

*Then for every fixed compact  $K$  there is  $n_K$  such that, for  $n \geq n_K$ ,*

$$\inf_{z \in K} N_{n,\ell} p_{n,\ell}^\partial(z) \geq 2k_{n,\ell}. \quad (39.22)$$

*Consequently*

$$\sup_{z \in K} \Pr^z \{ \text{cap reached before } k_{n,\ell} \text{ successes} \} \leq e^{-ck_{n,\ell}} \quad (39.23)$$

*for all large  $n$ , with the same cap sequence for every compact  $K$ .*

*Proof.* The boundary success law has the form, uniformly on compact local parameter sets,

$$p_{n,\ell}^\partial(z) = C_\ell(z, n) [d_\ell + O_K(\delta_n)]^2 \exp\{-j_n \gamma(z, n)\},$$

where  $C_\ell(z, n)/C_\ell(0, n) = 1 + O_K(\delta_n)$  and  $\gamma(z, n) - \gamma(0, n) = O_K(\delta_n)$ . This is the uniform threshold law and smooth parameter dependence of its prefactor. Therefore

$$\log \frac{p_{n,\ell}^\partial(z)}{p_{n,\ell}^\partial(0)} = O_K(\delta_n) + O_K(j_n \delta_n) = o(1),$$

uniformly on  $K$ , proving (39.20). In particular the ratio is at least  $1/2$  for all  $z \in K$  and all  $n \geq n_K$ . Substituting (39.21) proves (39.22). A standard binomial Chernoff bound then gives (39.23).  $\square$

**39.6. The boundary endpoint-output Gaussian experiment.** Choose a finite design from Lemma 39.2. Let  $k_{n,\ell}$  be target success counts and put  $k_n = \sum_\ell k_{n,\ell}$ . Assume

$$\frac{k_{n,\ell}}{k_n} \rightarrow \alpha_\ell, \quad k_n \delta_n^2 \log(1/\delta_n) \rightarrow 1, \quad j_n \delta_n \rightarrow 0, \quad j_n \in 2\mathbb{N}. \quad (39.24)$$

Use the fixed cap sequence (39.21).

**Theorem 39.4** (Fixed-window multirate endpoint Gaussian experiment). *The capped boundary endpoint-output experiment for  $z = (a, h)$  is, uniformly on every fixed compact local parameter set, at  $o(1)$  Le Cam distance from a dominated uniformly LAN representative with information matrix  $\mathcal{K}_{\text{design}}$ . In particular, its finite local subexperiments have the Gaussian limit with that positive-definite information matrix. Corollary 40.3 proves the stronger two-sided Le Cam convergence on every fixed compact local set, at this same endpoint-output observation level.*

After applying to each design the common-collar censoring of Section 37, the representative has central sequence  $\Delta_n \Rightarrow N_{1+r}(0, \mathcal{K}_{\text{design}})$  and

$$\log \frac{d\bar{\mathbf{Q}}_{n,z}^{\text{end}}}{d\bar{\mathbf{Q}}_{n,0}^{\text{end}}} = z^t \Delta_n - \frac{1}{2} z^t \mathcal{K}_{\text{design}} z + o_P(1) \quad (39.25)$$

uniformly on compact  $z$ -sets. No finite-sample likelihood ratio is asserted for the original endpoint-output family when its support extends beyond the null support.

The gap is visible at rate  $\delta_n/j_n$  and the anchored registered contact geometry at rate  $\delta_n$ . The Gaussian testing profile, reference-model central-sequence estimator and local asymptotic minimax conclusions of Theorem 37.8 hold with  $\mathcal{J}_\Sigma$  replaced by  $\mathcal{K}_{\text{design}}$ .

*Proof.* Before applying the cap, Proposition 39.1 puts each exact boundary endpoint experiment at  $o(1)$  Le Cam distance from the ideal linearly moving-support experiment. Theorem 37.8 applies to the latter, with boundary velocity  $U_b$  and information matrix  $\mathcal{K}_{b,d}$ . Independent design batches add their central sequences and information matrices; the proportions  $k_{n,\ell}/k_n \rightarrow \alpha_\ell$  give  $\mathcal{K}_{\text{design}}$ , which is positive definite by Lemma 39.2. Lemma 39.3 changes the experiment by  $o(1)$  in total variation on every fixed compact set, using the same deterministic cap sequence independently of that set.  $\square$

### 39.7. Transfer to actual long finite bridges at the declared record level.

**Theorem 39.5** (Uniform stopped finite-to-boundary transfer). *Let  $\mathbf{P}_{n,z}^{\text{fin,ET}}$  be the law of the stopped endpoint-time acquisition transcript of item 2 in Section 39.1, under the deterministic caps, fixed physical times and target success counts above. Thus it retains every design, every success or failure record (39.1), and the stopping time. Let  $\mathbf{Q}_{n,z}^{\partial,ET}$  be the boundary transcript generated by the same policy on the same laboratory record space. If*

$$k_n \tau^{j_n} \longrightarrow 0, \quad (39.26)$$

then, for every fixed compact local set  $K$ ,

$$\sup_{z \in K} \|\mathbf{P}_{n,z}^{\text{fin,ET}} - \mathbf{Q}_{n,z}^{\partial,ET}\|_{\text{TV}} \leq C_K k_n \tau^{j_n} = o(1). \quad (39.27)$$

The bound survives every common Markov coarsening. In particular the actual finite-bridge endpoint-output experiment inherits Theorem 39.4.

This theorem makes no total-variation comparison of the full raw collision histories  $\mathcal{R}^{\text{full}}$ . No failed preparation is free. If  $T_{n,\ell}$  is the number of raw preparations in batch  $\ell$ , then

$$\mathbb{E}_z T_{n,\ell} \leq \frac{k_{n,\ell}}{p_{n,\ell}(z)} \leq C_K k_{n,\ell} d_\ell^{-2} e^{j_n \gamma_0} \quad (39.28)$$

uniformly on  $K$  for all large  $n$ ; the deterministic cap gives an absolute upper bound.

*Proof.* For bounded local  $z$ , the excess (39.8) stays in a common compact subcollar and the anchored graphs remain in a compact smooth family. The policy satisfies the deterministic bounded-stopping hypotheses of Theorem 35.3. At each preparation, that theorem uses exactly the per-preparation endpoint-time coarsening of Section 34; it does not compare the full collision array. Its success-weighted coupling gives  $C_K k_n \tau^{j_n}$  for the entire stopped transcript at this declared record level. Total variation contracts under a common Markov kernel, proving the endpoint-output statement. The cost estimate follows from domination by the uncapped negative-binomial stopping time and the uniform success law.  $\square$

*Remark 39.6* (Endpoint information is not complete physical information). The matrix  $\mathcal{K}_{\text{design}}$  characterizes the endpoint-output experiment of item 3. It is not an efficiency bound for every procedure based on the stopped endpoint-time transcript of item 2, and still less for the full raw collision history. The success probability contains the factor  $\sinh(j\gamma)^{-1}$ , so waiting counts can carry a different and possibly faster information channel in directions changing the hyperbolic exponent. Charging preparation cost is not the same as proving those counts ancillary.

### 39.8. A charged pilot for centered windows.

**Corollary 39.7** (Pilot-centered endpoint-output equivalence). *Assume the coarse-bracket hypothesis of Theorem A.20.3. Let  $r_n \downarrow 0$  and let an independent pilot produce  $\hat{g}_n$  with*

$$j_n |\hat{g}_n - g| \leq C r_n^{m+1}$$

*with probability  $1 - \eta_n$ , where  $\eta_n \rightarrow 0$ . Suppose*

$$k_n r_n^{2m+2} \log(1/r_n) \rightarrow 0. \quad (39.29)$$

*After the pilot use the common data-dependent rule  $t_{n,\ell} = j_n \hat{g}_n + d_\ell$  and the same reference-based capped success targets. On the post-pilot endpoint-output sigma-field alone, this experiment is at vanishing Le Cam distance from the oracle centered endpoint experiment using  $t_{n,\ell} = j_n g + d_\ell$ . Including the finite-bridge replacement adds at most*

$$\eta_n + C_K k_n \tau^{j_n} + o(1). \quad (39.30)$$

*The pilot transcript and programmed times are used operationally and their costs are charged, but no equivalence is claimed after those richer random coordinates are added to the statistical output.*

*Proof.* On the good pilot event the post-pilot excess differs from its oracle value by at most  $C r_n^{m+1}$ . Lemma 37.10 gives one-record squared Hellinger distance at most  $C r_n^{2m+2} \log(1/r_n)$ , uniformly over the finite design. Product Hellinger bounds and (39.29) give  $o(1)$  for endpoint output. Add  $\eta_n$  for the bad pilot event and use Theorem 39.5 for the finite-bridge replacement.  $\square$

**39.9. Generic nearby tables.** The anchored model is a physical submodel, not a parameter-dependent coordinate convention. A generic nearby-table family can still be described on one laboratory record space, but translations and rotations of the contact frames become nuisance directions. If  $\chi_{b,\eta}$  maps a fixed laboratory transverse coordinate  $y$  to the contact-centered coordinate for table parameter  $\eta$ , then

$$\mathcal{S}_{b,\eta}(y) = S_{b,\eta}(\chi_{b,\eta}(y))$$

and, at a reference with  $\chi_{b,0}(y) = y$ ,

$$D_\eta \mathcal{S}_b(y)[h] = D_\eta S_b(y)[h] + S'_b(y) D_\eta \chi_b(y)[h]. \quad (39.31)$$

The abstract non-dominated theorem of Section 37 applies once these nuisance coordinates and their full laboratory support velocity are included. We do not identify the anchored information matrix with the generic-table one after silently removing those directions. A charged registration pilot with coordinate error  $\varepsilon_n$  is negligible for endpoint output whenever  $k_n \varepsilon_n^2 \log(1/\varepsilon_n) \rightarrow 0$ , by Lemma 37.10.

## 40. COMPACT LOCAL EXPERIMENTS

The finite-subexperiment conclusions above identify the Gaussian limit. We now establish the separate assertion of convergence in Le Cam distance on an arbitrary fixed compact local parameter set. The additional step is a uniform continuity estimate for the experiments, not a reinterpretation of finite-dimensional likelihood convergence. Throughout, total variation is  $\|P - Q\|_{\text{TV}} = \sup_A |P(A) - Q(A)|$ , and the Hellinger convention is (37.1).

**Lemma 40.1** (Passage from finite sets to compact sets). *Let  $K$  be a compact metric space and let  $\mathbf{E}_n = (P_{n,t} : t \in K)$  and  $\mathbf{E} = (Q_t : t \in K)$  be experiments on standard Borel spaces. Suppose that their restrictions to each finite  $F \subset K$  satisfy  $\Delta(\mathbf{E}_n|_F, \mathbf{E}|_F) \rightarrow 0$ . Put*

$$\omega_n(r) = \sup_{d(s,t) \leq r} \|P_{n,s} - P_{n,t}\|_{\text{TV}}, \quad \omega(r) = \sup_{d(s,t) \leq r} \|Q_s - Q_t\|_{\text{TV}}.$$

If

$$\lim_{r \downarrow 0} \limsup_{n \rightarrow \infty} \omega_n(r) = 0, \quad \lim_{r \downarrow 0} \omega(r) = 0, \quad (40.1)$$

then  $\Delta(\mathbf{E}_n, \mathbf{E}) \rightarrow 0$ . Neither experiment is required to be dominated by one of its own members.

*Proof.* Fix  $r > 0$  and a finite  $r$ -net  $F \subset K$ . For every  $\varepsilon > 0$ , choose a Markov kernel  $R_n$  whose error on  $F$  is at most  $\delta(\mathbf{E}_n|_F, \mathbf{E}|_F) + \varepsilon$ . For  $t \in K$ , choose  $s \in F$  with  $d(s, t) \leq r$ . Contraction of total variation gives

$$\begin{aligned} \|R_n P_{n,t} - Q_t\|_{\text{TV}} &\leq \|P_{n,t} - P_{n,s}\|_{\text{TV}} + \|R_n P_{n,s} - Q_s\|_{\text{TV}} + \|Q_s - Q_t\|_{\text{TV}} \\ &\leq \omega_n(r) + \delta(\mathbf{E}_n|_F, \mathbf{E}|_F) + \varepsilon + \omega(r). \end{aligned}$$

The same argument for a reverse kernel, followed by  $\varepsilon \downarrow 0$ , proves

$$\Delta(\mathbf{E}_n, \mathbf{E}) \leq \Delta(\mathbf{E}_n|_F, \mathbf{E}|_F) + \omega_n(r) + \omega(r). \quad (40.2)$$

First let  $n \rightarrow \infty$  with this finite net fixed, and then let  $r \downarrow 0$ . The net point is used only to estimate the error of a single kernel; it is not an unknown-parameter input to that kernel.  $\square$

**Theorem 40.2** (Compact moving-support Gaussian approximation). *Under the hypotheses of Theorem 37.8, for every compact  $K \subset \mathbb{R}^r$ ,*

$$\Delta\left(\{P_{n,h}^{\otimes n} : h \in K\}, \{N_{\mathbf{H}}(\mathcal{J}_{\mathbf{H}}\Pi_{\mathbf{H}}h, \mathcal{J}_{\mathbf{H}}) : h \in K\}\right) \rightarrow 0. \quad (40.3)$$

Here  $\mathbf{H}$ ,  $\Pi_{\mathbf{H}}$  and  $\mathcal{J}_{\mathbf{H}}$  are as in (37.22)–(37.23). The assertion includes singular information, with the zero-dimensional Gaussian interpreted as the one-point experiment.

*Proof.* Fix  $K$  and a compact neighborhood containing every line segment between its points. Joint smoothness and the common regular collar imply, for  $h, h' \in K$  and  $s = \|h - h'\| \leq 1$ ,

$$\|w_{\delta_n h} - w_{\delta_n h'}\|_{C^1} + \|a_{\delta_n h} - a_{\delta_n h'}\|_{C^0} \leq C_K \delta_n s.$$

Lemma 37.10, with constants uniform in this neighborhood, therefore yields

$$H^2(f_{\delta_n h}, f_{\delta_n h'}) \leq C_K \delta_n^2 s^2 \left(1 + \log \frac{1}{\delta_n} + \log \frac{1}{s}\right), \quad 0 < s \leq 1. \quad (40.4)$$

The failure atoms have identical masses. Hence their one-record squared Hellinger distance is  $p_n$  times the successful-density distance. Product subadditivity and (37.4) give, for all large  $n$ ,

$$H^2(P_{n,h}^{\otimes n}, P_{n,h'}^{\otimes n}) \leq C_K s^2 (1 + \log(1/s)). \quad (40.5)$$

Since  $\|P - Q\|_{\text{TV}} \leq H(P, Q)$  with our convention, this supplies the first modulus in (40.1).

For the displayed Gaussian laws, writing  $v = h - h'$ , direct Gaussian integration gives

$$H^2(Q_h, Q_{h'}) = 2 \left(1 - \exp\{-v^t \mathcal{J}_{\Sigma} v / 8\}\right).$$

This identity is evaluated on  $\mathbf{H}$ ; it does not require an inverse of  $\mathcal{J}_\Sigma$  on its kernel. It supplies the second modulus. The finite-set convergence is item 2 of Theorem 37.8. Apply Lemma 40.1.

In particular, the comparison here is for the original moving-support family. The common-collar representative is needed to identify its finite limits, but no likelihood ratio with respect to the original reference law has been asserted on support-exclusive observations.  $\square$

**Corollary 40.3** (Compact fixed-window endpoint experiments). *With the finite design, rate conditions and reference-based caps of Theorem 39.4, the capped boundary endpoint-output experiment converges in Le Cam distance, on every fixed compact set of  $z = (a, h)$ , to the Gaussian experiment with information  $\mathcal{K}_{\text{design}}$ . The same conclusion holds after restriction to any linear subspace of the local coordinate. Under (39.26), it holds also for the actual finite-bridge endpoint-output experiment, at the same declared observation level.*

*Proof.* Use the normalized ideal densities (39.15) in the independent design batches. Their defining functions and normalized amplitudes satisfy the uniform bounds used in (40.4); the normalization is positive and smooth on a common small parameter neighborhood. Sum the resulting one-record bounds over the  $\sum_\ell k_{n,\ell} = k_n$  records. The critical success scale gives the product modulus (40.5) with  $np_n$  replaced by  $k_n$ . The finite limits have information  $\sum_\ell \alpha_\ell \mathcal{K}_{b_\ell, d_\ell} = \mathcal{K}_{\text{design}}$ , by Theorem 39.4. Lemma 40.1 and Gaussian continuity give compact convergence for the ideal experiment.

Proposition 39.1 transfers it to the exact boundary densities with a uniform  $o(1)$  error. The cap contributes at most  $\sum_\ell e^{-ck_{n,\ell}}$ , by Lemma 39.3. Restriction cannot increase deficiency. Finally, Theorem 39.5 contributes  $C_K k_n \tau^{j_n} = o(1)$  on the common endpoint-output space.  $\square$

*Remark 40.4.* Compact local convergence is not an assertion of uniform equivalence over an unbounded local parameter space or over the complete analytic class of tables. The kernels may depend on the fixed compact set, the reference channel and the declared design, but never on the unknown local parameter.

#### 41. THE COMPLETE ENDPOINT–TIME ACQUISITION EXPERIMENT AT ITS FASTEST LOCAL SCALE

The endpoint-output experiment of Section 39 is a strict coarsening of the record for which the finite-to-boundary theorem is proved. We now restore the residual preparation time and the complete waiting sequence. The resulting local experiment is not Gaussian at its fastest shape scale. Its fastest geometric information is an extreme boundary process. This gives a theorem for the entire stopped endpoint–time acquisition transcript of item 2 in Section 39.1, rather than only for its endpoint-output coarsening.

Fix a finite labelled contact-jet order  $M$  in the anchored realization of Section 38. Choose a finite positive design  $(b_\ell, d_\ell)$ ,  $1 \leq \ell \leq L$ , as in Lemma 39.2. Write

$$w_\ell(u, v) = d_\ell - S_{b_\ell}(u) - S_{b_\ell}(v), \quad D_\ell = \{(u, v) : w_\ell(u, v) > 0\}.$$

The boundary selected-record measure of Section 34, after conditioning on success and returning to unscaled residual time, has density

$$q_{\ell,0}(u, v, r) = \rho_\ell(u, v) \mathbf{1}_{\{0 < r < w_\ell(u, v)\}}, \quad \int_{D_\ell} \rho_\ell w_\ell du dv = 1, \quad (41.1)$$

where  $\rho_\ell$  is positive and smooth on the common endpoint box. Up to the normalizing constant it is the half-line amplitude  $B_{b_\ell}(u)B_{b_\ell}(v)$ . Thus the conditional density has a positive trace at the moving ceiling  $r = w_\ell(u, v)$ ; integrating out  $r$  multiplies the endpoint density by  $w_\ell$  and creates the linear vanishing studied in Section 37.

Let  $k_{n,\ell} \rightarrow \infty$ , put  $k_n = \sum_{\ell} k_{n,\ell}$ , and assume

$$\frac{k_{n,\ell}}{k_n} \rightarrow \alpha_{\ell} > 0, \quad \sum_{\ell=1}^L \alpha_{\ell} = 1. \quad (41.2)$$

Use an even flight number  $j_n \rightarrow \infty$  for every same-type design and the fixed physical windows  $t_{n,\ell} = j_n g_0 + d_{\ell}$ . At the endpoint-time scale the local parameter is

$$g_n(a) = g_0 + \frac{a}{j_n k_n}, \quad \vartheta_n(h) = \frac{h}{k_n}, \quad z = (a, h) \in \mathbb{R}^{1+r_M}. \quad (41.3)$$

For design  $\ell$ , use the support velocity of Section 39,

$$U_{\ell}(u, v)z = a + \dot{S}_{b_{\ell}}(u)h + \dot{S}_{b_{\ell}}(v)h. \quad (41.4)$$

Proposition 39.1, with  $\delta_n = k_n^{-1}$ , gives uniformly on compact  $z$ -sets

$$w_{n,\ell,z}^{\text{ET}}(u, v) = w_{\ell}(u, v) - \frac{1}{k_n} U_{\ell}(u, v)z + o_K(k_n^{-1}) \quad (41.5)$$

in  $C^1$ , while the conditional density trace at the ceiling is  $\rho_{\ell} + O_K(k_n^{-1})$ . The residual gap-dependence of the half-line action is  $O((j_n k_n)^{-1})$  and is absorbed by the remainder.

**41.1. The boundary point-process limit.** For a successful endpoint-time record  $(u, v, r)$ , define its null ceiling slack and scaled slack by

$$s_{\ell} = w_{\ell}(u, v) - r, \quad y = k_n s_{\ell}.$$

Under the reference model  $y \geq 0$ . Under a local alternative the same reference coordinate may be negative when the ceiling expands.

For a compact local set  $K$ , choose

$$R_K > 1 + \max_{1 \leq \ell \leq L} \sup_{(u,v) \in \overline{D}_{\ell}, z \in K} |U_{\ell}(u, v)z|. \quad (41.6)$$

Let  $\mathcal{N}_{n,\ell}$  be the point measure formed from the  $k_{n,\ell}$  successful records by retaining only points with  $|y| \leq R_K$  and recording  $(u, v, y)$ .

**Theorem 41.1** (Endpoint-time Poisson boundary experiment). *Assume (41.2) and let the local parameter be (41.3). Then, uniformly for  $z$  in every fixed compact set  $K$ , the conditional successful endpoint-time experiment is asymptotically equivalent to independent Poisson point processes  $\mathcal{N}_1, \dots, \mathcal{N}_L$  on  $D_{\ell} \times [-R_K, R_K]$  with intensities*

$$d\Lambda_{\ell,z}(u, v, y) = \alpha_{\ell} \rho_{\ell}(u, v) \mathbf{1}_{\{y > U_{\ell}(u, v)z\}} \mathbf{1}_{\{|y| \leq R_K\}} du dv dy. \quad (41.7)$$

*The part of the successful sample outside the displayed boundary layer is asymptotically parameter-free after conditioning on the retained point process. Equivalently, the local experiment is the projective Poisson boundary-shift experiment obtained by increasing  $R_K$ .*

*The finite-strip family has a common dominating Poisson envelope, as shown in Proposition 43.1. It need not be absolutely continuous with respect to the zero-shift member: if  $U_{\ell}z < 0$  on a set of positive  $\rho_{\ell} du dv$ -measure, the alternative assigns positive probability to points in a region where the reference intensity vanishes. The resulting support-exclusive events prevent contiguity in such expanding directions at this scale, and hence prevent a finite-information Gaussian LAN representation there. The obstruction is reference absolute continuity, not common domination.*

*Proof.* Fix one design and suppress  $\ell$ . The level set  $\{w = 0\}$  is smooth for  $d > 0$ . Hence the coarea formula gives, uniformly for fixed  $R$ ,

$$\text{Leb}\{(u, v) : 0 < w(u, v) \leq 2R/k_n\} = O(k_n^{-1}). \quad (41.8)$$

On this endpoint collar the residual interval has length at most  $2R/k_n$ , so its one-record probability is  $O(k_n^{-2})$ . Over  $k_{n,\ell} = O(k_n)$  records the probability of seeing any point from this corner is therefore  $O(k_n^{-1}) = o(1)$ . This removes the codimension-two intersection between the moving ceiling and the fixed face  $r = 0$  from the first-order experiment.

Away from that negligible collar, the condition  $r > 0$  imposes no additional restriction on a fixed  $y$ -strip. By (41.5) and smoothness of the conditional trace, the one-record law in reference coordinates is, uniformly on  $K$ ,

$$\frac{1}{k_n}(\rho(u, v) + o_K(1))\mathbf{1}_{\{y > U(u, v)z + o_K(1)\}} du dv dy \quad (41.9)$$

on  $|y| \leq R_K$ . In particular the probability  $p_{n,R}$  of the whole retained layer is  $O(k_n^{-1})$ . Conditional on entering the layer, the location law converges in total variation to the normalized version of (41.7). The standard binomial–Poisson coupling therefore gives

$$d_{\text{TV}}(\mathcal{L}(\mathcal{N}_{n,\ell}), \text{PPP}(\Lambda_{\ell,z})) \leq Ck_{n,\ell}p_{n,R}^2 + o_K(1) = o_K(1). \quad (41.10)$$

The same estimate for a finite measurable partition of the layer, followed by monotone refinement, gives the point-process statement; independent design batches give the product process.

It remains to verify that no random first-order information is hidden in the bulk. Remove a reference ceiling layer of thickness  $R_K/k_n$ . The choice (41.6) makes the remaining support common to every  $z \in K$ . There the density ratio is a smooth perturbation of the form

$$1 + k_n^{-1}a_z(u, v, r) + o_K(k_n^{-1})$$

with uniformly bounded  $a_z$ . Its centered one-record log-likelihood has variance  $O_K(k_n^{-2})$ , hence the centered sum over  $k_{n,\ell} = O(k_n)$  records has variance  $O_K(k_n^{-1})$  and converges to zero uniformly on  $K$ . The deterministic first-order normalizing term is forced, by total mass one, to equal the negative of the mass gained or lost in the ceiling layer. In the Poisson likelihood this is exactly the exponential compensator  $\exp\{-\Lambda_{\ell,z} + \Lambda_{\ell,0}\}$ . Thus, conditional on the boundary-layer point process, the bulk likelihood ratio tends to one in probability uniformly on  $K$ .

The extraction of  $\mathcal{N}_{n,\ell}$  is a parameter-independent Markov kernel. Conversely, generate the asymptotically parameter-free bulk from the reference conditional law and attach it to the Poissonized boundary layer; (41.10) and the preceding bulk estimate make the reverse deficiency tend to zero. The explicit layer-count, bulk, and corner-safe reverse kernels, with the verified relative-density hypothesis, are given in Theorem 42.2. They prove the asserted two-sided asymptotic equivalence, including exceptional configurations.  $\square$

**Corollary 41.2** (Identification at the  $k^{-1}$  geometric scale). *Choose the positive finite design of Lemma 39.2. Then the map*

$$z \longmapsto (U_\ell(\cdot, \cdot)z)_{\ell=1}^L$$

*is injective. Consequently the Poisson limit experiment of Theorem 41.1 identifies the full finite coordinate  $(a, h)$ : gap is visible at scale  $(j_n k_n)^{-1}$ , while the anchored labelled contact geometry through order  $M$  is visible at scale  $k_n^{-1}$ .*

*Proof.* If  $U_\ell z = 0$  on  $D_\ell$  for every  $\ell$ , it vanishes in particular on the endpoint boundary  $\{S_{b_\ell}(u) + S_{b_\ell}(v) = d_\ell\}$ . Hence  $z^t \mathcal{K}_{\text{design}} z = 0$ . Positive definiteness from Lemma 39.2 gives  $z = 0$ .  $\square$

**41.2. Waiting counts and transfer to finite bridges.** For a fixed design, stop after  $k_{n,\ell}$  successful preparations and let  $T_{n,\ell}$  be the total number of preparations. Before the deterministic cap is applied, the boundary law factorizes exactly into

$$\mathcal{L}_z(T_{n,\ell}) \otimes (\Pi_{n,\ell,z}^{\partial, \text{ET}})^{\otimes k_{n,\ell}}, \quad (41.11)$$

where  $T_{n,\ell}$  is negative binomial with success probability  $p_{n,\ell}^\partial(z)$ . Thus the successful marks are independent of the waiting count.

**Theorem 41.3** (Complete stopped endpoint–time local experiment). *In addition to (41.2), assume*

$$j_n \in 2\mathbb{N}, \quad j_n \rightarrow \infty, \quad j_n = o(\sqrt{k_n}), \quad k_n \tau^{j_n} \rightarrow 0. \quad (41.12)$$

*Use the reference-based deterministic caps of Lemma 39.3, with the target success counts  $k_{n,\ell}$ . Then, uniformly on every fixed compact local parameter set:*

- (1) *the waiting-count component of the boundary stopped transcript is asymptotically ancillary at the scale (41.3);*
- (2) *the entire capped boundary endpoint–time acquisition transcript is at vanishing Le Cam distance from the Poisson boundary experiment of Theorem 41.1, up to a parameter-free factor;*
- (3) *the actual finite-bridge stopped endpoint–time transcript is at vanishing total-variation distance from that boundary transcript, and hence has the same Poisson local experiment.*

*Thus the declared endpoint–time acquisition process has a complete local limit at its fastest geometric scale; this statement retains failures, waiting counts, designs, successful residual times and stopping times. It does not identify the still richer full growing collision array.*

*Proof.* The selected-channel term in Theorem 13.1, (13.1), and its smooth fixed-order parameter dependence give

$$\log \frac{p_{n,\ell}^\partial(z)}{p_{n,\ell}^\partial(0)} = O_K(k_n^{-1}) + O_K(j_n k_n^{-1}) = O_K(j_n/k_n).$$

For a negative-binomial experiment stopped after  $k_{n,\ell}$  successes, the Fisher information for  $\log p$  is  $k_{n,\ell}/(1-p)$ . Hence the squared Hellinger distance between the waiting-count laws is  $O_K(k_n(j_n/k_n)^2) = O_K(j_n^2/k_n) = o(1)$ . This proves asymptotic ancillarity. The same success-ratio estimate is  $1 + o(1)$ , so the single reference-based cap remains adequate and its failure probability is exponentially small.

The exact factorization (41.11), cap negligibility and Theorem 41.1 prove the boundary statement. Finally Theorem 39.5 applies to the complete stopped endpoint–time transcript and gives total-variation error  $O_K(k_n \tau^{j_n}) = o(1)$ . Le Cam distance contracts under composition, which proves the finite-bridge conclusion.  $\square$

*Remark 41.4* (Why the endpoint Gaussian theorem is slower). At the scale  $k_n^{-1}$ , positive ceiling density makes the extreme residual-time layer carry order-one information. Integrating out residual time replaces the positive ceiling density by the linearly vanishing endpoint density  $\rho_\ell w_\ell$ ; the  $k_n^{-1}$  Poisson layer disappears and the next surviving endpoint scale is  $k_n \delta_n^2 \log(1/\delta_n) \asymp 1$ . Thus the Gaussian endpoint experiment of Theorem 39.4 is the slower experiment created by residual-time coarsening.

## 42. TWO-SIDED DEFICIENCY FOR THE MOVING-CEILING EXPERIMENT

We give parameter-independent kernels in both directions for the local endpoint–time experiment. Besides convergence of the ceiling and its trace, the argument uses a uniform first-order relative density bound on the common bulk. We state this bound explicitly: convergence of a trace in  $C^0$  alone would not imply the required product Hellinger estimate.

For one design write  $k = k_n$ ,  $w = w_\ell$ ,  $\rho = \rho_\ell$ , and  $U = U_\ell$ . Let  $K$  be a fixed compact local parameter set containing zero. On a fixed bounded endpoint box  $V$ , the exact successful density has the form

$$q_{n,z}(u, v, r) = \rho_{n,z}(u, v, r) \mathbf{1}_{\{0 < r < w_{n,z}(u,v)\}}.$$



All density factors extend to a common bounded residual-time box, with uniform upper and positive lower bounds. Put  $D = \{w > 0\}$ , whose boundary is regular in the endpoint box, and assume

$$\varepsilon_{n,K} := \sup_{z \in K} \left( k \|w_{n,z} - w + k^{-1} U z\|_{C^1} + \|\rho_{n,z} - \rho\|_{C^0} \right) \longrightarrow 0. \quad (42.1)$$

Here  $\rho = \rho(u, v)$  is independent of  $r$  and  $\int_D \rho w \, du \, dv = 1$ . In addition suppose that on the common bulk specified below,

$$\sup_{z \in K} \left\| \frac{\rho_{n,z}}{\rho_{n,0}} - 1 \right\|_{\infty} \leq C_K / k. \quad (42.2)$$

A bounded expansion  $1 + k^{-1} a_{n,z} + o_K(k^{-1})$  is sufficient, but not necessary, for this bound.

These hypotheses hold for the anchored fixed-window boundary model of Section 41. Its parameter displacement is  $O_K(k^{-1})$ , its normalized smooth density factor is independent of  $r$ , and its reference factor is bounded away from zero. Smooth dependence of that factor and its positive normalizer gives (42.2) uniformly on the box. Equation (41.5) gives the support estimate; the residual gap term is  $O((j_n k_n)^{-1}) = o(k_n^{-1})$ . Thus the additional displayed hypothesis is a verified dependency of the physical application, not a deduction from bare trace convergence.

Choose

$$R > 2 + \sup_{z \in K, (u,v) \in \overline{V}} |U(u, v) z|$$

and define the reference layer

$$L_{n,R} = \{|k(w(u, v) - r)| \leq R\}. \quad (42.3)$$

It is independent of  $z$ . For large  $n$ , the support outside this layer is the common set  $0 < r < w - R/k$ ; the alternative ceilings cannot reach the other side  $r > w + R/k$ . This is the bulk in (42.2).

Let  $p_{n,z} = Q_{n,z}(L_{n,R})$ , let  $A_{n,z}$  be its conditional layer law after  $y = k(w - r)$ , and let  $B_{n,z}$  be the conditional bulk law. The target intensity is the finite measure on  $D \times [-R, R]$  with density

$$d\Lambda_z^R = \rho(u, v) \mathbf{1}_{\{y > U(u, v) z\}} \mathbf{1}_{\{|y| \leq R\}} \, du \, dv \, dy. \quad (42.4)$$

The factor in this formula is the previously defined density  $\rho$ .

**Lemma 42.1** (Uniform layer and bulk bounds). *Uniformly for  $z \in K$ ,*

$$p_{n,z} = O(k^{-1}), \quad (42.5)$$

$$\left| k Q_{n,z} \circ (u, v, r \mapsto u, v, k(w - r))^{-1} \Big|_{L_{n,R}} - \Lambda_z^R \right| (V \times [-R, R]) \leq C_K (\varepsilon_{n,K} + k^{-1}), \quad (42.6)$$

$$H^2(B_{n,z}, B_{n,0}) \leq C_K k^{-2}. \quad (42.7)$$

*The notation in the second line means that the layer submeasure is restricted before pushforward, and that the target intensity is extended by zero from  $D$  to  $V$ . The vertical bars denote total variation mass of a finite signed measure.*

*Proof.* The layer has residual thickness  $2R/k$  over a bounded endpoint set, which gives (42.5). Regularity of  $\{w = 0\}$  implies  $\text{Leb}\{|w| \leq C/k\} = O(k^{-1})$ . On this endpoint strip any physical residual interval involved has length  $O(k^{-1})$ , so its one-record mass is  $O(k^{-2})$ . The part of the target intensity on  $0 < w \leq C/k$  is  $O(k^{-1})$  as well. These estimates handle both the intersection with the fixed face  $r = 0$  and any alternative support created just outside  $D$ .

Off that strip,  $r > 0$  imposes no additional restriction in the reference layer. The change of variable has Jacobian  $k^{-1}$ , and the ceiling condition becomes

$$y > U(u, v) z + e_{n,z}(u, v), \quad \|e_{n,z}\|_{\infty} \leq \varepsilon_{n,K}.$$

Fiber integration bounds the symmetric difference by the endpoint area times  $\varepsilon_{n,K}$ , while the trace-density difference is bounded by the second term of (42.1). Adding the corner error proves (42.6). In particular, the unscaled mass attributable to the support remainder is at most  $C_K \varepsilon_{n,K}/k$ .

On the common bulk write  $\rho_{n,z}/\rho_{n,0} = 1 + e_{n,z}^B$ , with  $\|e_{n,z}^B\|_\infty \leq C_K/k$ . Dividing by the bulk normalizers gives exactly

$$\frac{dB_{n,z}}{dB_{n,0}} = \frac{1 + e_{n,z}^B}{1 + B_{n,0}e_{n,z}^B} = 1 + O_K(k^{-1}).$$

The inequality  $(\sqrt{1+x} - 1)^2 \leq Cx^2$  for small  $x$  proves (42.7). This proof identifies the relative-density hypothesis actually used in the bulk and controls the normalized conditional law, not merely a marginal fluctuation.  $\square$

Let  $\mathbf{E}_{n,K}^R = \{Q_{n,z}^{\otimes k} : z \in K\}$  and  $\mathbf{P}_K^R = \{\text{PPP}(\Lambda_z^R) : z \in K\}$ .

**Theorem 42.2** (Explicit two-sided deficiency). *Under the displayed support, positivity and relative-density hypotheses,*

$$\Delta(\mathbf{E}_{n,K}^R, \mathbf{P}_K^R) \leq C_K(\varepsilon_{n,K} + k^{-1/2}). \quad (42.8)$$

*Both deficiencies admit Markov kernels independent of  $z$ . For finitely many independent batches the bounds add. In particular this proves the two-sided equivalence claimed in Theorem 41.1; it is stronger than convergence of the extracted point process alone.*

*Proof.* The forward kernel retains records in  $L_{n,R}$ , replaces  $r$  by  $y = k(w - r)$ , and discards points outside  $D$ . All operations use reference quantities only. Before the last discard, the retained sample is a binomial point process. Its distance from the Poisson process with the same intensity is at most  $Ckp_{n,z}^2 = O(k^{-1})$ . Poisson coupling and (42.6) bound the further intensity error by  $C_K(\varepsilon_{n,K} + k^{-1})$ . Discarding exterior points has already been accounted for in the corner estimate and cannot increase distance. This proves the forward deficiency bound.

For the reverse kernel, let  $m$  be the number of input Poisson points. If  $m \leq k$ , map each point to  $(u, v, r = w - y/k)$ , draw  $k - m$  independent records from the reference bulk law  $B_{n,0}$ , and interleave the two collections by a uniformly random permutation. If a reconstructed point has  $r \leq 0$ , replace it by one fixed reference record. Such a point has  $0 < w \leq R/k$ , so the probability that this convention is used is  $O(k^{-1})$ . If  $m > k$ , output one fixed  $k$ -sample. The Poisson means are uniformly bounded, hence  $\sup_{z \in K} \Pr_z\{m > k\} = o(k^{-A})$  for every fixed  $A$ . This defines a valid parameter-independent kernel on every configuration, including corner and count exceptions.

For the original iid sample, conditional on its layer count, the layer records and bulk records have respective iid laws  $A_{n,z}$  and  $B_{n,z}$ , are independent of each other, and are uniformly interleaved. The layer count and locations can be coupled with the Poisson input with the forward error just proved. Replacing the conditional bulk by its reference law costs, uniformly in  $m \leq k$ ,

$$H^2(B_{n,z}^{\otimes(k-m)}, B_{n,0}^{\otimes(k-m)}) \leq (k-m)H^2(B_{n,z}, B_{n,0}) \leq C_K/k.$$

Its total variation cost is therefore  $O_K(k^{-1/2})$ . The common permutation does not increase distance. Adding the corner and count tails gives (42.8). Independent-product coupling proves tensorization over finitely many batches. For the proportions in (41.2), scale by total  $k_n$ ; the intensity in a batch is multiplied by  $k_{n,\ell}/k_n$ , converging to  $\alpha_\ell$ . The resulting additional intensity error tends to zero.  $\square$

*Remark 42.3* (The projective Poisson experiment). On  $D \times \mathbb{R}$ , the intensity

$$\rho(u, v) \mathbf{1}_{\{y > U(u, v)z\}} du dv dy$$

is locally finite. Restriction from a larger finite strip to a smaller one is a common Markov kernel. Once  $R$  exceeds every displacement on the compact parameter set, the added positive tail has parameter-independent intensity and the negative tail is empty. It can thus be adjoined as an ancillary factor. This precise projective interpretation does not require an informal limit in the reverse reconstruction kernel.

#### 43. COMMON DOMINATION AND THE INFORMATION IN EXACT POSITIONS

Three properties must be separated: domination of a statistical family by one sigma-finite measure, absolute continuity with respect to a designated reference member, and mutual absolute continuity of the members. Moving support need not prevent the first property. The distinction also separates the scalar-coordinate local experiments from the richer position protocol.

**Proposition 43.1** (Common reference measures). *The family (37.5) is dominated, for every  $n$ , by  $\mu = \text{Leb}|_U + \delta_{\dagger}$ . Its censored family is dominated by  $\text{Leb}|_U + \delta_{\dagger} + \delta_{\partial}$ . The finite products have the corresponding product dominating measures. The original members need not be absolutely continuous with respect to  $P_{n,0}$ ; the common-collar construction establishes this stronger property for the censored members and their censored reference.*

On the fixed finite strip  $B = D \times [-R, R]$ , let

$$d\Lambda_z = \rho(u, v) \mathbf{1}_{\{y > U(u, v)z\}} du dv dy, \quad d\Lambda_* = \rho(u, v) du dv dy.$$

Then the Poisson law  $\mathbf{P}_* = \text{PPP}(\Lambda_*)$  is a common dominating probability measure for the entire finite-strip family, and

$$\frac{d\mathbf{P}_z}{d\mathbf{P}_*}(N) = \exp\{\Lambda_*(B) - \Lambda_z(B)\} \prod_{q \in N} \mathbf{1}_{\{y(q) > U(q)z\}}. \quad (43.1)$$

The empty product is one. Multiplication by a fixed batch proportion gives the same assertion for each batch of (41.7).

*Proof.* The first statements follow by writing the displayed Lebesgue densities and atomic masses. They do not involve division by the reference density. For example, on  $[0, 2]$ ,

$$f_\theta(x) = \frac{2(1 + \theta - x)_+}{(1 + \theta)^2}, \quad 0 \leq \theta < 1,$$

is a family of probability densities, but  $P_\theta((1, 1 + \theta)) = \theta^2/(1 + \theta)^2 > 0$  for  $\theta > 0$ , whereas  $P_0((1, 1 + \theta)) = 0$ .

For the Poisson assertion write  $A_z = \{y > Uz\} \cap B$ . Under  $\mathbf{P}_*$ , restrictions to  $A_z$  and its complement are independent Poisson processes. The probability of no point in the complement is  $\exp\{-\Lambda_*(B \setminus A_z)\}$ . Multiplying its indicator by the reciprocal of that probability has expectation one and leaves the independent process on  $A_z$  unchanged. The resulting law has intensity  $\Lambda_*|_{A_z} = \Lambda_z$ , proving (43.1). The argument includes the zero-intensity and empty-configuration cases.  $\square$

A zero-shift Poisson member need not dominate the others. If  $Uz < 0$  on a set of positive weighted endpoint measure, there is positive intensity in  $\{Uz < y < 0\}$ , where the reference has none. The event of at least one point there has positive alternative probability and zero reference probability. This is a failure of reference absolute continuity, not of common domination. It also excludes a Gaussian shift with mutually absolutely continuous members as the local limit in such a direction. The common-collar Gaussian limit and the moving-ceiling Poisson limit remain different experiments for the reasons supplied by their boundary rates and support events, not by an incorrect classification of their dominating measures.

For a fixed compact parameter set, choose  $R$  beyond all displacements. The positive tail outside the strip is parameter independent and the negative tail is empty. Adjoining the common positive-tail process preserves the common domination just proved, also for the projective experiment of Remark 42.3. This makes no uniform assertion about an unbounded parameter set without a common displacement bound.

**43.1. Matched records and the observed contact type.** We use the convention  $\|P - Q\|_{\text{TV}} = \sup_A |P(A) - Q(A)|$ , so a deficiency belongs to  $[0, 1]$ . For experiments  $\mathbf{E} = (P_\lambda)_{\lambda \in \Theta}$  and  $\mathbf{F} = (Q_\lambda)_{\lambda \in \Theta}$ , write

$$\delta(\mathbf{E}, \mathbf{F}) = \inf_K \sup_{\lambda \in \Theta} \|KP_\lambda - Q_\lambda\|_{\text{TV}}, \quad \Delta(\mathbf{E}, \mathbf{F}) = \max\{\delta(\mathbf{E}, \mathbf{F}), \delta(\mathbf{F}, \mathbf{E})\},$$

where the kernels are independent of the unknown parameter.

Fix an even-flight design starting at contact type  $b$ . Both retained endpoints then have type  $b$ . In this subsection a successful position record is exactly  $(X_{b,\lambda}(u), X_{b,\lambda}(v))$ , and its scalar coarsening is exactly  $(u, v)$ . Residual time is integrated out in both records; intermediate collisions and waiting counts occur in neither. The fixed transverse projection is one parameter-independent kernel. All claims concern these matched records, conditional on success. They do not remove preparation costs from the unconditional acquisition theorems.

**Lemma 43.2** (Dominated input and disjoint output supports). *Suppose  $\Theta$  is uncountable,  $P_\lambda \ll \nu$  for one probability measure  $\nu$ , and there are pairwise disjoint measurable sets  $A_\lambda$  with  $Q_\lambda(A_\lambda) = 1$ . Then  $\delta((P_\lambda), (Q_\lambda)) = 1$ .*

*Proof.* For any kernel  $K$ , put  $Q = K\nu$ . If  $Q(A) = 0$ , nonnegativity gives  $K(x, A) = 0$  for  $\nu$ -almost every  $x$ , hence  $KP_\lambda(A) = 0$  for every  $\lambda$ . Thus  $KP_\lambda \ll Q$ . A probability measure gives positive mass to at most countably many pairwise disjoint measurable sets. Choose  $\lambda$  outside that countable set for  $Q$ . Then  $KP_\lambda(A_\lambda) = 0$ , whereas  $Q_\lambda(A_\lambda) = 1$ . The worst-case error of  $K$  is therefore one. Infimizing over  $K$  proves the assertion. No integration over the uncountable parameter set is used.  $\square$

**Proposition 43.3** (Observed-type dichotomy for anchored homotheties). *Let  $\Theta$  be a nondegenerate compact interval of positive homothety parameters. In a persistent registered analytic facing channel let the obstacle of type  $b_*$  vary as  $K_\lambda = p + \lambda(K - p)$ , with  $p \in \partial K$  its fixed contact point and fixed tangent. Keep the facing obstacle fixed, with positive anchored gap. Choose one admissible even-flight design and common small contact windows. The successful endpoint law has positive smooth flux and a density with respect to the two boundary arclength coordinates. Let  $\mathbf{E}_{\text{sc},b}^{(k)}$  and  $\mathbf{E}_{\text{pos},b}^{(k)}$  denote  $k \geq 1$  independent matched successful records of the preceding definition.*

*If the design starts at  $b = b_*$ , the distinct position laws are mutually singular and admit no common sigma-finite dominating measure. For every finite  $k \geq 1$ ,*

$$\delta(\mathbf{E}_{\text{sc},b_*}^{(k)}, \mathbf{E}_{\text{pos},b_*}^{(k)}) = 1, \quad \delta(\mathbf{E}_{\text{pos},b_*}^{(k)}, \mathbf{E}_{\text{sc},b_*}^{(k)}) = 0. \quad (43.2)$$

*If instead the design starts at the fixed obstacle,  $b = 1 - b_*$ , then*

$$\Delta(\mathbf{E}_{\text{sc},b}^{(k)}, \mathbf{E}_{\text{pos},b}^{(k)}) = 0. \quad (43.3)$$

*In this case the position family has a common dominating measure and a common positive-density region, and distinct members are not mutually singular. In either case the scalar laws are continuous in total variation as  $\lambda$  varies.*

*Proof.* First suppose the observed type is  $b_*$ . Along each ray from  $p$  entering the strictly convex body  $K$ , there is exactly one other boundary point. A homothety multiplies its distance from  $p$  by  $\lambda$ . Thus two distinct homothetic boundaries meet only at  $p$ . Removing the anchor gives pairwise

disjoint Borel curves  $\Gamma_\lambda = \partial K_\lambda \setminus \{p\}$ . The first endpoint has an arclength density and hits  $p$  with probability zero. The event that the first recorded position lies in  $\Gamma_\lambda$  has probability one under that member, also for any finite product of successful records. These events are pairwise disjoint. They prove mutual singularity. A sigma-finite measure has at most countably many disjoint measurable sets of positive measure: intersect with a countable finite-measure cover, and then group by a positive lower bound on the intersection measure. A common dominating measure is therefore impossible.

In the fixed transverse coordinates all scalar laws have densities on one bounded box  $U \subset \mathbb{R}^2$ . They are dominated by normalized Lebesgue measure on  $U$ ; their  $k$ -fold products are dominated by its  $k$ -fold product. Apply Lemma 43.2 to the events just constructed. This gives the first equality in (43.2). The fixed transverse projection gives the reverse deficiency zero. In particular the Le Cam distance is one.

Now suppose  $b = 1 - b_*$ . The observed contact arc is a single fixed graph  $X_b$ , independent of  $\lambda$ . The maps

$$F_b(u, v) = (X_b(u), X_b(v)), \quad T_b(x, x') = (\text{transverse}(x), \text{transverse}(x'))$$

are measurable and mutually inverse on the supported contact boxes. Moreover  $P_{\lambda, b}^{\text{pos}} = (F_b)_* P_{\lambda, b}^{\text{psc}}$ . Extend the maps arbitrarily off those boxes. They are parameter-independent kernels and give (43.3), also on products. The inverse may use the fixed obstacle: it is a known constant of this specified family, not an unknown parameter-dependent boundary.

For either observed type the scalar density is

$$f_{\lambda, b}(u, v) = Z_{\lambda, b}^{-1} A_{\lambda, b}(u, v)(d - E_{J, \lambda, b}(u, v))_+, \quad A_{\lambda, b} > 0. \quad (43.4)$$

The action and flux vary smoothly on the common box, the normalizers are positive, and compactness bounds them away from zero. Uniform continuity and  $|x_+ - y_+| \leq |x - y|$  prove total-variation continuity. Since  $E_{J, \lambda, b}(0, 0) = 0$ , a sufficiently small common square has  $E_{J, \lambda, b} < d/4$  for every  $\lambda$ . On this square all scalar densities are positive. For the fixed observed obstacle,  $(F_b)_* \text{Leb}|_U$  is a common dominating measure, with a common positive-density region. No two position members can be mutually singular. The same continuity argument applies to the limiting factorized endpoint law.  $\square$

*Remark 43.4* (Matched residual-time retention). The same dichotomy holds if residual time is retained in both records. Use  $(u, v, r)$  and  $(X_{b, \lambda}(u), X_{b, \lambda}(v), r)$ , respectively. Scalar densities are dominated on a fixed bounded three-dimensional box. For a fixed observed obstacle use  $F_b \times \text{id}$ ; for a varying observed obstacle use the same first-position disjoint events. Total-variation continuity follows by integrating the difference of ceiling indicators, whose integral is bounded by the ceiling displacement. The common square above, together with  $d/4 < r < d/2$ , is a positive-density core. Comparing a time-retaining record with a time-deleted record is a different question and is not an instance of (43.3).

**43.2. A fixed-obstacle design and nonsymmetric examples.** The dependence on the observed type already occurs for two circular contacts. Set

$$K_0 = \{(x, y) : (x+1)^2 + y^2 \leq 1\}, \quad K_{1, \lambda} = \{(x, y) : (x-1-\lambda)^2 + y^2 \leq \lambda^2\}, \quad \lambda \in [1-\epsilon, 1+\epsilon].$$

Their closest points are  $(0, 0)$  and  $(1, 0)$ , and their gap is one. The second disk is obtained by homothety about  $(1, 0)$ . For small  $\epsilon > 0$ , take  $J = 2$ , physical time  $2 + d$ , and the word  $K_0 \rightarrow K_{1, \lambda} \rightarrow K_0$ . Both observed positions belong to the fixed arc

$$X_0(u) = (-1 + \sqrt{1 - u^2}, u).$$

The finite stationarity equations have a nondegenerate solution at the normal orbit and hence a common smooth solution on a small box. Its excess vanishes at the origin. Formula (43.4)

therefore has the common positive core used in the proof, and  $F_0$  gives exact endpoint-only equivalence. Unobserved intermediate visits to the varying disk do not turn these endpoint-position laws into singular laws. This construction explains why the type qualification in (43.2) cannot be omitted.

The varying-type conclusion does not require circular symmetry. A strictly convex analytic oval with support function

$$h(\theta) = 1 + \frac{1}{50} \cos(2\theta) + \frac{1}{100} \sin(3\theta)$$

has  $h + h'' \geq 43/50 > 0$ . Its nonzero harmonic orders two and three exclude a nonidentity rotational symmetry. Fix a boundary anchor and use nearby homothetic copies, with a facing obstacle kept at positive gap. A sufficiently small parameter interval preserves the local channel. Starting on the oval gives (43.2); starting on the fixed facing obstacle gives (43.3). These are local examples, without an assertion of the global cycle and spanning-tree hypotheses for an unspecified periodic table.

Exact injectivity of the transverse *law* map in Theorem 23.4 is compatible with maximal finite-sample deficiency on the varying-type family. The supremum over an uncountable parameter interval and the finite-sample record are essential. Distribution-level determination, recovery of a particular homothety parameter from an exact position, and reconstruction of an arbitrary analytic obstacle are distinct assertions. The global law inverse remains a statement using both contact types and the separately retained gap.

#### 44. WAITING COUNTS AND ENDPOINT INFORMATION AFTER RESIDUAL-TIME COARSENING

Theorem 41.3 identifies the fastest local experiment when residual time is retained. We now remove residual time but retain the waiting counts. The positive-density ceiling is thereby integrated out. Two information mechanisms remain: the rare-event probability detects the hyperbolic exponent, while the endpoint mark detects motion of its linearly vanishing support. We prove that their product is a Gaussian experiment in *two-sided Le Cam distance on every fixed compact local parameter set*.

Fix one anchored labelled channel and a finite contact-jet model of order  $M \geq 2$ . Let  $\vartheta \in \mathbb{R}^M$  be its shape coordinate at fixed gap, and let  $\gamma(g, \vartheta)$  be the channel hyperbolic exponent. At the reference point  $(g_0, 0)$ , put

$$\lambda = D_{\vartheta} \gamma(g_0, 0), \quad H_{\gamma} = \ker \lambda.$$

The coordinate contains the two contact curvatures, on which  $\gamma$  depends nontrivially, so  $\lambda \neq 0$ . Choose  $v_{\gamma} \in \mathbb{R}^M$  with  $\lambda(v_{\gamma}) = 1$ . This choice gives a fast hyperbolic direction and an iso-hyperbolic tangent space; changing the complement merely changes the local coordinates.

Choose the finite positive design of Lemma 39.2, with target counts  $k_{n,\ell}$ , total  $k_n = \sum_{\ell} k_{n,\ell}$ , and proportions  $k_{n,\ell}/k_n \rightarrow \alpha_{\ell} > 0$ . Let  $\delta_n \downarrow 0$  satisfy

$$k_n \delta_n^2 \log(1/\delta_n) \rightarrow 1. \quad (44.1)$$

For bounded  $(b, a, h) \in \mathbb{R} \times \mathbb{R} \times H_{\gamma}$ , set

$$\eta_n = (j_n \sqrt{k_n})^{-1}, \quad g_n(a) = g_0 + \delta_n a / j_n, \quad \vartheta_n(b, h) = \eta_n b v_{\gamma} + \delta_n h. \quad (44.2)$$

All physical windows remain  $t_{j_n, \ell} = j_n g_0 + d_{\ell}$ , and the caps are the parameter-independent reference caps (39.21). In batch  $\ell$ , retain the waiting count  $T_{n, \ell}$  and the successful laboratory transverse pairs  $(y_{\text{first}}, y_{\text{last}})$  of (39.4), with the cap cemetery convention. This is a common coarsening of the stopped endpoint-time transcript, not of an experiment in which preparation failures are free.

**Lemma 44.1** (Negative-binomial hyperbolic score). *Suppose  $j_n \rightarrow \infty$  and  $j_n \delta_n \rightarrow 0$ . Uniformly on compact  $(b, a, h)$ -sets,*

$$\log \frac{p_{n,\ell}^\partial(b, a, h)}{p_{n,\ell}^\partial(0, 0, 0)} = -b/\sqrt{k_n} + o(k_n^{-1/2}). \quad (44.3)$$

*For the uncapped boundary waiting experiment, define under its reference law*

$$S_{n,\ell} = \frac{k_{n,\ell} - p_{n,\ell}^\partial(0)T_{n,\ell}}{1 - p_{n,\ell}^\partial(0)}, \quad Z_n = -k_n^{-1/2} \sum_{\ell=1}^L S_{n,\ell}.$$

*Then  $Z_n \Rightarrow N(0, 1)$ , and*

$$\log \frac{dW_{n,b,a,h}}{dW_{n,0,0,0}} = bZ_n - \frac{1}{2}b^2 + o_P(1), \quad (44.4)$$

*uniformly on compact sets. The slow coordinates contribute no limiting count score.*

*Proof.* On the fixed positive offset collar the success probability has a smooth positive prefactor times  $\exp\{-j_n \gamma(g, \vartheta)\}$ ; the logarithmic derivatives of the prefactor are uniformly bounded. Taylor expansion of  $\gamma$ , using  $\lambda(h) = 0$ , gives

$$\gamma(g_n(a), \vartheta_n(b, h)) - \gamma(g_0, 0) = \eta_n b + O_K(\delta_n/j_n) + O_K(\delta_n^2 + \delta_n \eta_n + \eta_n^2).$$

Consequently, the remainder in the success log ratio after the term  $-b/\sqrt{k_n}$  is bounded by

$$C_K(\delta_n + \eta_n + j_n \delta_n^2 + j_n \delta_n \eta_n + j_n \eta_n^2).$$

After multiplication by  $\sqrt{k_n}$ , these terms tend to zero:  $\sqrt{k_n} \delta_n \rightarrow 0$ ,  $\sqrt{k_n} \eta_n = 1/j_n \rightarrow 0$ ,  $\sqrt{k_n} j_n \delta_n^2 = (j_n \delta_n)(\sqrt{k_n} \delta_n) \rightarrow 0$ ,  $\sqrt{k_n} j_n \delta_n \eta_n = \delta_n \rightarrow 0$ , and  $\sqrt{k_n} j_n \eta_n^2 = \eta_n \rightarrow 0$ . This proves (44.3) without discarding the mixed fast-slow Taylor terms.

For a negative-binomial waiting time stopped at  $k$  successes, the score for  $\phi = \log p$  is  $(k - pT)/(1 - p)$ , with variance  $k/(1 - p)$ . Represent each waiting time as a sum of independent geometric waiting times. For  $p$  small, the centered one-success log-probability scores have uniformly bounded third absolute moments: the variables  $pT$  have a uniform exponential tail. The Lyapunov bound for the normalized sum is  $O(k_n^{-1/2})$ . Since  $p_{n,\ell}^\partial(0) \rightarrow 0$ , this proves  $Z_n \Rightarrow N(0, 1)$ .

The second and third derivatives of the batch log likelihood are

$$-\frac{(T - k)p}{(1 - p)^2}, \quad -\frac{(T - k)p(1 + p)}{(1 - p)^3},$$

respectively. On the compact local neighborhoods, their orders are  $O_P(k)$ , and the normalized negative second derivative converges to its information. Taylor expansion with (44.3) gives (44.4).  $\square$

**Lemma 44.2** (Uniform comparison of waiting experiments). *For geometric laws  $G_p$  on  $\{1, 2, \dots\}$ , if  $p, q \leq p_* < 1$ ,*

$$H^2(G_p, G_q) \leq \frac{(\log p - \log q)^2}{4(1 - p_*)}. \quad (44.5)$$

*For the sum of  $k$  independent such waiting times, the bound is multiplied by at most  $k$ . Define ideal batch probabilities*

$$p_{n,\ell}^*(b) = p_{n,\ell}^\partial(0) \exp\{-b/\sqrt{k_n}\}.$$

*The uncapped joint waiting experiment is uniformly at  $o(1)$  total variation distance, on compact  $(b, a, h)$ -sets, from the product of the negative-binomial laws with probabilities  $p_{n,\ell}^*(b)$ . This ideal experiment converges in Le Cam distance on every compact  $b$ -set to  $N(b, 1)$ .*

*Proof.* As a curve in  $\ell^2$ , the square-root mass function of  $G_{e^\phi}$  has derivative with squared norm  $1/[4(1 - e^\phi)]$ , one quarter of its log-probability information. Integrating this derivative on the segment between  $\log p$  and  $\log q$  proves (44.5). Product subadditivity and contraction under summation give the waiting-time bound. In particular, this proof uses the exact discrete waiting law, not a Poisson replacement.

By (44.3), the maximum difference between the true and ideal log probabilities is  $o(k_n^{-1/2})$ , uniformly on the compact set. All probabilities are at most  $1/2$  for large  $n$ . Summing (44.5) over all  $k_n$  required successes gives squared Hellinger distance  $o(1)$ , hence the stated total variation comparison.

The count LAN calculation gives joint weak convergence of the ideal likelihood vector on each finite  $b$ -set to  $(\exp\{bZ - b^2/2\})_b$ , with  $Z \sim N(0, 1)$ . The discrete waiting laws have common support and each likelihood has reference mean one. Each Gaussian limiting coordinate also has mean one. Lemma 37.6, adjoining the reference value  $b = 0$  before restriction when necessary, therefore proves the finite-subexperiment limit  $N(b, 1)$ . Its pairwise squared Hellinger distance is bounded by  $C|b - b'|^2$ , uniformly in  $n$ , because the log-probability difference is exactly  $-(b - b')/\sqrt{k_n}$ . The Gaussian family has a continuous Hellinger modulus as well. Lemma 40.1 proves compact Le Cam convergence.  $\square$

Let  $\mathcal{K}_H$  be the restriction of  $\mathcal{K}_{\text{design}}$ , defined in (39.18), to  $\mathbb{R} \times H_\gamma$ . It is positive definite on that space.

**Theorem 44.3** (Two-speed count–endpoint local experiment). *Assume*

$$\begin{aligned} k_n \delta_n^2 \log(1/\delta_n) &\rightarrow 1, & j_n \delta_n &\rightarrow 0, \\ \log(j_n \sqrt{k_n})/j_n^2 &\rightarrow 0, & k_n \tau^{j_n} &\rightarrow 0, & j_n \in 2\mathbb{N}, & j_n \rightarrow \infty. \end{aligned} \quad (44.6)$$

Use (44.2), the fixed laboratory windows and the reference caps. For each fixed compact  $K \subset \mathbb{R} \times \mathbb{R} \times H_\gamma$ , let  $\mathbf{E}_{n,K}^{\partial, CE}$  be the capped boundary experiment retaining waiting counts and the successful laboratory transverse pairs  $(y_{\text{first}}, y_{\text{last}})$  defined by the common map (39.4). Normal endpoint coordinates, residual times and intermediate collisions are not retained in this output. Let  $\mathbf{E}_{n,K}^{\text{fin}, CE}$  be the same coarsening of the actual finite-bridge transcript, with every preparation failure charged. Then

$$\Delta(\mathbf{E}_{n,K}^{\partial, CE}, \mathbf{G}_K) \rightarrow 0, \quad \Delta(\mathbf{E}_{n,K}^{\text{fin}, CE}, \mathbf{G}_K) \rightarrow 0, \quad (44.7)$$

where  $\mathbf{G}_K$  observes, independently,

$$X \sim N(b, 1), \quad Y \sim N(\mathcal{K}_H(a, h), \mathcal{K}_H), \quad (b, a, h) \in K.$$

Thus the limiting information is  $1 \oplus \mathcal{K}_H$ . The hyperbolic shape direction is visible at scale  $(j_n \sqrt{k_n})^{-1}$ , the iso-hyperbolic shape directions at scale  $\delta_n$ , and the gap at scale  $\delta_n/j_n$ . Both comparison kernels are independent of the unknown  $(b, a, h)$ .

*Proof.* First remove the caps. In an iid Bernoulli-mark experiment the successful marks are independent of the waiting times and have their conditional success law. This factorization is exact for the uncapped experiment; the cap error will be added only after the product comparison.

Lemma 44.2 replaces the waiting factor by its ideal  $b$ -only experiment with a uniform  $o(1)$  total variation error. For the endpoint factor, compare the parameter  $(g_n(a), \eta_n b v_\gamma + \delta_n h)$  with  $(g_n(a), \delta_n h)$  at the same physical time. Smooth dependence of the actions and amplitudes changes the defining function in  $C^1$  and the amplitude in  $C^0$  by at most  $C_K \eta_n$ . By Lemma 37.10, the total squared Hellinger error over the successful records is at most

$$C_K k_n \eta_n^2 \log(e/\eta_n) = C_K \frac{1 + \log(j_n \sqrt{k_n})}{j_n^2} = o(1). \quad (44.8)$$



The smooth amplitude term is bounded by  $C_K k_n \eta_n^2$  and is included in this estimate. The remaining endpoint experiment depends only on  $(a, h)$ , and converges on compact sets to its Gaussian shift by Corollary 40.3, restricted to  $\mathbb{R} \times H_\gamma$ .

Take the compact projections of  $K$  onto its count and endpoint coordinates. The product of the two comparison kernels on those projection sets has error at most the sum of the component errors. Restriction from their Cartesian product to  $K$  cannot increase either deficiency. This proves compact two-sided convergence for the uncapped boundary experiment, with independent Gaussian factors; it is stronger than convergence of a fixed finite vector of likelihood ratios.

By (44.3), the true success probabilities are uniformly asymptotic to their reference values on  $K$ . The reference cap therefore has failure probability at most  $\sum_\ell e^{-ck_{n,\ell}} = o(1)$ , as in Lemma 39.3. Coupling to the uncapped acquisition sequence changes the experiment by at most this probability; a reverse kernel may assign any fixed completion on the cap event. This proves the first limit in (44.7).

The common count–endpoint coarsening of the stopped transcript preserves the total variation bound  $C_K k_n \tau^{j_n}$  from Theorem 39.5. Its hypotheses hold here: the local graphs stay in the same compact collar, the fast perturbation tends to zero, and the policy and caps are parameter independent. The same success-weighted adaptive comparison consequently applies, without a free failure convention. The triangle inequality for Le Cam distance proves the second limit.  $\square$

*Remark 44.4* (Likelihoods and the retained record). Count likelihoods are ordinary negative-binomial likelihoods. For endpoint marks, likelihood expansions relative to a reference member are taken only after common-collar censoring. Formula (44.7) instead concerns the original capped experiments and therefore does not assert reference domination where the support moves. Nor does it compare complete raw collision histories.

*Remark 44.5* (Three levels of information). Under the respective rate hypotheses, retaining residual time gives the non-Gaussian Poisson boundary experiment of Theorem 41.3, with shape scale  $k_n^{-1}$  and gap scale  $(j_n k_n)^{-1}$ . Removing residual time but keeping waiting counts gives Theorem 44.3. Removing both gives the endpoint-output Gaussian experiment of Theorem 39.4 and Corollary 40.3. These are different experiments related by common Markov coarsenings, not alternative names for one record.

#### 45. INTRINSIC NORMAL–TANGENT GEOMETRY OF THE COUNT–ENDPOINT LIMIT

The splitting used in Section 44 is a convenient coordinate representative of a canonical quotient construction. We now state the product experiment invariantly and verify directly that the hyperbolic differential is nonzero.

Fix a finite anchored contact-jet model at a reference channel. Let  $V = T_{\vartheta_0} \Theta_M$  be its shape tangent space at fixed gap and let

$$\lambda = d\gamma_{\vartheta_0} \in V^*. \quad (45.1)$$

The two contact curvatures are coordinates of this finite jet model. With the notation of (12.1),

$$c_b = 1 + g\kappa_b, \quad c_0 c_1 = \cosh^2 \gamma, \quad \gamma = \operatorname{arcosh} \sqrt{c_0 c_1}.$$

Consequently, at fixed  $g > 0$ ,

$$\frac{\partial \gamma}{\partial \kappa_0} = \frac{gc_1}{2\sqrt{c_0 c_1} \sinh \gamma} > 0, \quad \frac{\partial \gamma}{\partial \kappa_1} = \frac{gc_0}{2\sqrt{c_0 c_1} \sinh \gamma} > 0. \quad (45.2)$$

In particular  $\lambda \neq 0$  at every dispersing reference channel in the model; no generic non-cancellation hypothesis is needed.

Put

$$H = \ker \lambda, \quad N = V/H. \quad (45.3)$$

The covector  $\lambda$  induces the canonical isomorphism

$$\bar{\lambda} : N \longrightarrow \mathbb{R}, \quad [v] \longmapsto \lambda(v). \quad (45.4)$$

Thus the fast count coordinate is intrinsically an element of the one-dimensional normal quotient  $N$ , normalized by  $\bar{\lambda}$ , not a preferred vector in  $V$ .

Let  $\mathcal{K}$  be the positive endpoint information quadratic form of the fixed positive design on  $\mathbb{R}_{\text{gap}} \oplus V$ . Its slow component is the restricted quadratic form

$$\mathcal{K}_{\text{tan}} = \mathcal{K}|_{\mathbb{R}_{\text{gap}} \oplus H}. \quad (45.5)$$

This restriction is positive definite by the full finite-design theorem.

**Theorem 45.1** (Coordinate-free two-speed information theorem). *Under the rate assumptions of Theorem 44.3, the count-endpoint experiment has the intrinsic product limit*

$$\mathcal{G}_{\bar{b}}^{\text{count}} \otimes \mathcal{G}_{(a,h)}^{\text{end}}, \quad \bar{b} \in N, \quad a \in \mathbb{R}_{\text{gap}}, \quad h \in H, \quad (45.6)$$

where the count factor has quadratic information

$$I_{\text{count}}(\bar{b}) = |\bar{\lambda}(\bar{b})|^2 \quad (45.7)$$

and the endpoint factor has information tensor  $\mathcal{K}_{\text{tan}}$ . The factors are independent.

This formulation is invariant under every  $C^2$  reparameterization of the finite jet model. In particular, choosing any splitting vector  $v_\gamma \in V$  with  $\lambda(v_\gamma) = 1$  identifies  $N$  with the coordinate called  $b$  in Section 44, but the limit experiment does not depend on that choice.

*Proof.* Let  $v_\gamma$  be any representative with  $\lambda(v_\gamma) = 1$ . Theorem 44.3 gives the product limit with count information one in the scalar coefficient of  $v_\gamma$  and endpoint information  $\mathcal{K}_{\text{tan}}$  on the tangent coordinates. If  $v'_\gamma$  is another representative, then  $v'_\gamma - v_\gamma \in H$ . At the fast physical scale  $\eta_n = (j_n \sqrt{k_n})^{-1}$ , replacing one representative by the other adds an  $H$ -displacement of size  $O(\eta_n)$ . The endpoint experiment cannot see this displacement because the proof of Theorem 44.3 gives

$$k_n \eta_n^2 \log(1/\eta_n) \longrightarrow 0.$$

The count likelihood sees only the first differential of  $\gamma$ , hence only  $\lambda(v)$ , and therefore gives the same scalar shift for the two representatives. This proves that the limit descends to the quotient  $N = V/H$  and that its quadratic form is (45.7).

Under a  $C^2$  change of shape coordinates,  $d\gamma$  pulls back as a covector, its kernel maps to the tangent space of the same level set, and  $\mathcal{K}$  transforms as a quadratic tensor. Therefore both (45.7) and (45.5) are coordinate invariant. The exact waiting-count/mark factorization gives independence, as before.  $\square$

**Remark 45.2** (Geometry of the two rates). The faster Gaussian channel is the normal quotient to the level set of the hyperbolic exponent; the slower endpoint channel is tangent to that level set and includes the independently scaled gap direction. A choice of  $v_\gamma$  is analogous to choosing a connection for writing a quotient as a direct sum. It is useful for formulas but is not part of the experiment. Equation (45.2) also closes the previously implicit assertion that the normal quotient is genuinely one dimensional.

**Remark 45.3** (One simultaneous admissible rate regime). The scaling assumptions used by the endpoint, count-endpoint and stopped transfer theorems have a common nonempty regime. For example, let

$$\delta_n = n^{-1}, \quad k_n = \left\lfloor \frac{n^2}{\log n} \right\rfloor, \quad j_n = 2 \lceil C \log n \rceil$$

with  $C > |\log \tau|^{-1}$ . Then

$$k_n \delta_n^2 \log(1/\delta_n) \rightarrow 1, \quad j_n \delta_n \rightarrow 0, \quad k_n \tau^{j_n} \rightarrow 0,$$

and also

$$j_n = o(\sqrt{k_n}), \quad \frac{\log(j_n \sqrt{k_n})}{j_n^2} \rightarrow 0.$$

Thus the stated physical and statistical asymptotics do not rely on an empty intersection of rate conditions.

*Remark 45.4* (Exact source of the cap success-ratio estimate). The success-law input used in Lemma 39.3 is the selected channel summand in Theorem 13.1, equation (13.1), evaluated at the fixed physical window. That formula gives the factor  $\sinh(j\gamma)^{-1}$ , the quadratic offset factor and a smooth remainder with uniform fixed-order parameter derivatives. On a compact local parameter set, Taylor expansion of this exact formula gives the prefactor ratio  $1 + O_K(\delta_n)$  and the exponent variation  $j_n O_K(\delta_n)$  used in (39.20). Hence  $j_n \delta_n \rightarrow 0$  yields the required uniform ratio  $1 + o(1)$  for the single reference-based cap sequence.

#### 46. THE ENDPOINT CRITICAL EXPERIMENT

We apply the preceding theorem to a fixed physical channel. The two hypotheses are its finite-flight law and its boundary model for the *same* table. The observation space, the failure symbol, and the physical preparation cost are those of Section 34. Write  $P_{j,d}^{\text{end}}$  and  $P_{\partial,j,d}^{\text{end}}$  for the pushforwards retaining the two endpoint positions on success and the failure symbol otherwise. Thus residual time is not reported. Their conditional versions are denoted by  $\Pi_{j,d}^{\text{end}}$  and  $\Pi_{\partial,j,d}^{\text{end}}$ . Superscript 0 denotes the corresponding quadratic right-limit laws of Section A.9.

The common whitening in (A.9.1) is used only to prove the comparison for this fixed table. It does not give an observation rule knowledge of an unknown gap or curvature. Total variation and Hellinger distance are unchanged by this common invertible map.

##### 46.1. Exact tangent density and its information. Set

$$q_j = e^{-j\gamma}, \quad \zeta_j = \log \coth(j\gamma/2) = \log \frac{1+q_j}{1-q_j}, \quad r_j = q_j^2 \log(1/q_j). \quad (46.1)$$

After whitening and a common orthogonal rotation, the endpoint density of the finite tangent law is

$$f_{\zeta_j}(x) = \frac{2}{\pi} (1 - e^{-\zeta_j} x_1^2 - e^{\zeta_j} x_2^2)_+, \quad x \in \mathbb{R}^2, \quad (46.2)$$

and the boundary density is  $f_0(x) = 2(1 - |x|^2)_+/\pi$ . The raw pair has the same mass

$$p_{j,d}^0 = \frac{d^2}{2A \sinh(j\gamma)} \quad (46.3)$$

on these successful densities, and the same failure mass  $1 - p_{j,d}^0$ .

**Lemma 46.1** (The endpoint information coefficient). *The densities (46.2) are normalized. In the parameter  $\zeta$  their boundary information coefficient is one, and*

$$H^2(f_{\zeta_j}, f_0) = r_j + O(q_j^2). \quad (46.4)$$

*For every  $j, d$ , the raw tangent pair has identical success-indicator laws, even though its endpoint laws differ. These statements hold for unequal facing curvatures and both parities of  $j$ .*

*Proof.* The eigenvalues of the whitened Hessian are  $\lambda_j = \tanh(j\gamma/2) = e^{-\zeta_j}$  and  $\lambda_j^{-1} = e^{\zeta_j}$ . Integrating residual time in the uniform cap gives (46.2). The change  $(y_1, y_2) = (e^{-\zeta/2}x_1, e^{\zeta/2}x_2)$  has determinant one, and  $\int_{|y|<1} (1 - |y|^2) dy = \pi/2$ ; this proves normalization. The same determinant identity in (A.9.1) gives the common mass (46.3). In particular the two success bits are exactly Bernoulli with the same parameter, not merely equal to leading order in  $q_j$ .

Take  $w_\zeta = 1 - e^{-\zeta}x_1^2 - e^{\zeta}x_2^2$  and  $a_\zeta = 2/\pi$  on a fixed large disk. Then  $v = x_1^2 - x_2^2$  and  $|\nabla w_0| = 2$  on the unit circle. Consequently

$$\mathcal{I}_\Sigma = \frac{1}{\pi} \int_0^{2\pi} \cos^2(2\theta) d\theta = 1.$$

One can also see the logarithm directly. For  $w = 1 - |x|^2$  the score is  $S = (x_1^2 - x_2^2)/w$ . Polar integration gives the exact identities

$$\Pr\{0 < w < h\} = h^2, \quad \mathbb{E}_0[S \mathbf{1}_{\{w \geq h\}}] = 0, \quad \mathbb{E}_0[S^2 \mathbf{1}_{\{w \geq h\}}] = \log(1/h) - \frac{3}{2} + 2h - \frac{h^2}{2}. \quad (46.5)$$

Indeed the last integral is  $\int_h^1 (1-w)^2 dw/w$ ; the angular integral of  $\cos^2(2\theta)$  is  $\pi$ . Theorem 36.1 now gives  $H^2 = \zeta_j^2 \log(1/\zeta_j)/4 + O(\zeta_j^2)$ . Since  $\zeta_j = 2q_j + O(q_j^3)$ , this is (46.4). All dependence on the individual curvatures was removed by a common invertible whitening, not by assuming that the curvatures agree.  $\square$

**Theorem 46.2** (Sharp endpoint testing profile). *Let  $j \rightarrow \infty$ . If  $k_j r_j \rightarrow b \in [0, \infty]$ , then*

$$\|(\Pi_j^{0,\text{end}})^{\otimes k_j} - (\Pi_{\partial,j}^{0,\text{end}})^{\otimes k_j}\|_{\text{TV}} \rightarrow 2\Phi(\sqrt{b}) - 1. \quad (46.6)$$

*If  $n_j p_{j,d_j}^0 r_j \rightarrow b \in [0, \infty]$ , the same limit holds for the raw tangent endpoint laws with  $n_j$  preparations. At  $0 < b < \infty$ , the log likelihood under the boundary hypothesis converges to  $\mathcal{N}(-2b, 4b)$ . The minimum equal-prior testing error tends to  $\Phi(-\sqrt{b})$ . At  $b = \infty$  the distance tends to one and the testing error to zero.*

*In either experiment, vanishing total variation is equivalent to the corresponding effective sample size times  $r_j$  tending to zero. In particular the sharp conditional endpoint scale is  $e^{2j\gamma}/(j\gamma)$ , rather than  $e^{j\gamma}$ .*

*Proof.* Apply Theorem 36.1 with  $t = \zeta_j$ , first with  $p = 1$  and then with  $p = p_{j,d_j}^0$ . The row size in that theorem may of course run through the integers  $k_j$  or  $n_j$ . The information scales satisfy

$$\zeta_j^2 \log(1/\zeta_j) = 4r_j(1 + o(1)).$$

Lemma 46.1 supplies coefficient one. This proves the limits and the likelihood normalization. For necessity in the last assertion, if the effective scale does not tend to zero, pass to a subsequence on which it is bounded below by a fixed positive number. A further subsequence either has a finite positive limit or diverges. The corresponding distance limit is strictly positive. Sufficiency is the case  $b = 0$ .  $\square$

**46.2. Positive-offset physical transfer.** As in Lemma A.9.2, put

$$\omega(d) = \sqrt{d} \quad \text{for general smooth contacts,} \quad \omega(d) = d \quad \text{if both contact graphs are even.}$$

No assumption that the two even contacts coincide is made.

**Theorem 46.3** (The physical endpoint transition). *Suppose  $d_j > 0$  belongs to the common physical collar and*

$$\frac{\omega(d_j)}{r_j} \rightarrow 0. \quad (46.7)$$

If  $k_j r_j \rightarrow b \in [0, \infty]$ , then the actual conditional endpoint laws satisfy

$$\|(\Pi_{j,d_j}^{\text{end}})^{\otimes k_j} - (\Pi_{\partial,j,d_j}^{\text{end}})^{\otimes k_j}\|_{\text{TV}} \rightarrow 2\Phi(\sqrt{b}) - 1. \quad (46.8)$$

If  $n_j p_{j,d_j}^0 r_j \rightarrow b \in [0, \infty]$ , the same limit holds for the actual raw endpoint laws with  $n_j$  independent preparations. The minimum equal-prior errors have the same limits as in Theorem 46.2. Under (46.7), vanishing physical total variation is equivalent to the corresponding effective scale tending to zero. The estimates and conclusions are uniform on the compact geometric families under consideration, including variable geometry and either parity, provided the displayed scaling assumptions hold uniformly.

*Proof.* Projection onto the endpoint observation cannot increase total variation. Lemma A.9.2, telescoping independent products, and the reverse triangle inequality therefore give

$$|\text{TV}_{j,k,d}^{\text{end}} - \text{TV}_{j,k}^{0,\text{end}}| \leq Ck\omega(d), \quad (46.9)$$

$$|\text{TV}_{j,n,d}^{\text{end,raw}} - \text{TV}_{j,n,d}^{0,\text{end,raw}}| \leq Cnp_{j,d}^0\omega(d). \quad (46.10)$$

The symbols in these two inequalities denote the distances of the corresponding simple pairs of product laws. If  $b < \infty$ , their right sides tend to zero by (46.7). Theorem 46.2 proves the assertions in this case.

If  $b = \infty$ , fix  $B > 0$ . The conditional experiment eventually contains its first  $\lfloor B/r_j \rfloor$  observations. Their accumulated tangent error is at most  $CB\omega(d_j)/r_j = o(1)$ , and their limiting distance is  $2\Phi(\sqrt{B}) - 1$ . Projection onto this subsample gives the same lower bound for the lower limit of the full distance. Let  $B \rightarrow \infty$ . For raw preparations use instead  $\lfloor B/(p_{j,d_j}^0 r_j) \rfloor$  observations. Both rounding errors in the effective scales tend to zero. This proves the supercritical assertion without requiring the tangent error accumulated over the entire supercritical sample to vanish.

The same fixed positive subsample proves necessity of a vanishing effective scale, as in Theorem A.9.7. The uniform version uses the uniform tangent estimates and the fact that  $\gamma$  is bounded below by a positive number; the whitened family (46.2) itself is independent of the remaining geometry. Finally use the exact equal-prior testing identity.  $\square$

The positive-offset conditions are nonempty. For any  $\epsilon > 0$ ,  $d_j = r_j^{2+\epsilon}$  in the general smooth case and  $d_j = r_j^{1+\epsilon}$  in the even case satisfy (46.7). The preparation counts are finite for every  $j$ . Since  $p_{j,d}^0 \asymp d^2 e^{-j\gamma}$ , the raw endpoint transition occurs at order

$$d_j^{-2} e^{3j\gamma} / (j\gamma).$$

Failures have not been conditioned away in this statement.

### 46.3. A strict hierarchy of observations.

**Corollary 46.4** (Three distinct comparison experiments). *Under (46.7), suppose  $n_j p_{j,d_j}^0 r_j \rightarrow b \in (0, \infty)$ . For the same two actual physical models and the same  $n_j$  preparations, the respective total variation distances have the limits*

|                                       |                       |         |
|---------------------------------------|-----------------------|---------|
| retained observation on success       | limiting distance     | (46.11) |
| both endpoints and residual time      | 1                     |         |
| both endpoints, without residual time | $2\Phi(\sqrt{b}) - 1$ |         |
| success indicator only                | 0.                    |         |

Every experiment retains the same failed preparations. In the tangent models the last distance is identically zero at every sample size. Conversely, at the full-record critical scale  $n_j p_{j,d_j}^0 q_j \rightarrow c \in (0, \infty)$ , assuming only  $\omega(d_j)/q_j \rightarrow 0$ , the full-record distance tends to  $1 - e^{-2c/\pi}$  while the

endpoint-only distance tends to zero. The analogous assertions hold for a prescribed number of independently prepared successes, with the success-only conditional observation constant.

*Proof.* At the endpoint scale,

$$n_j p_{j,d_j}^0 q_j = \frac{n_j p_{j,d_j}^0 r_j}{q_j \log(1/q_j)} \longrightarrow \infty, \quad \frac{\omega(d_j)}{q_j} = \frac{\omega(d_j)}{r_j} q_j \log(1/q_j) \longrightarrow 0.$$

Theorem A.9.7 gives the first limit in (46.11), and Theorem 46.3 gives the second. For the last, both tangent success-indicator laws are exactly the same. Their actual counterparts differ from these by at most  $C n_j p_{j,d_j}^0 \omega(d_j) = o(1)$  in total variation, by coarsening Lemma A.9.2. This proves the last limit without postulating any rate for a count experiment outside this tangent regime.

At the full-record critical scale the known support calculation gives  $1 - e^{-2c/\pi}$ , whereas  $n_j p_{j,d_j}^0 r_j = (n_j p_{j,d_j}^0 q_j) q_j \log(1/q_j) \rightarrow 0$ . The tangent endpoint distance tends to zero by Theorem 46.2; its physical approximation error is  $O(n_j p_{j,d_j}^0 \omega(d_j)) = o(1)$  under the stated weaker collar condition. The conditional proof replaces  $n_j p_{j,d_j}^0$  by  $k_j$ .  $\square$

The full-record limit is an erasure limit because its tangent laws have equal densities on their common support (Theorem A.9.1). The endpoint limit is instead Gaussian: its supports are still different, but their exclusive mass is only  $O(q_j^2)$ , whereas its Hellinger information is  $q_j^2 \log(1/q_j)(1 + o(1))$ . The success bit loses this tangent shape information entirely through the exact determinant cancellation. Thus neither a one-record weak limit nor a coarsening of the full-record coupling bound identifies the sharp endpoint experiment.

These are intrinsic scales involving the exact hyperbolic multiplier, not a supplied upper convergence certificate  $\tau$ . They establish a sharp observation-dependent validity law in an explicitly nonempty physical tangent regime. They do not identify two unknown tables, assert equivalence of arbitrary noisy smooth-profile experiments, or turn the sufficient full-profile acquisition exponent into a minimax one. The nonlinear relative law on a fixed collar and the complete boundary-profile inverse retain their separate statements and proofs.

#### 47. THE ABEL GEOMETRY OF THE BOUNDARY INVARIANT

The scalar boundary law is obtained by two residual-time integrations of an autoconvolution. Its natural inverse topology is therefore better expressed on the integrated flux than on the law divided by  $d^2$ . For two profiles define

$$H(V, W)(d) = d^2 T(V, W)(d), \quad H_V = H(V, V),$$

where  $T$  is the normalized positive integral (20.2). This notation agrees with (A.6.3) below. For  $H \in C^3([0, D])$  set

$$\begin{aligned} (\mathcal{A}H)(x) &= \frac{\pi}{2} \left\{ H''(0) + \sqrt{x} \int_0^x \frac{H'''(s)}{\sqrt{x-s}} ds \right\}, \quad x > 0, \\ (\mathcal{A}H)(0) &= \frac{\pi}{2} H''(0), \quad |H|_{\mathcal{A}} = \|\mathcal{A}H\|_{C^0([0, D])}. \end{aligned} \tag{47.1}$$

The operator is linear. Equivalently, with  $k_*(x) = x^{-1/2}$  and  $Q = (\pi/2)H''$ , it is  $\sqrt{x}(k_* * Q)'$ . Formula (47.1), not a derivative of a singular expression at zero, defines its endpoint value. The transform is unchanged by adding an affine function to  $H$ .

**Theorem 47.1** (A nonlinear inverse in the Abel flux norm). *Let  $V, W \in C^3([0, D])$ ,  $V(0) = W(0) = 1$ , and suppose  $\|V'\|_\infty, \|W'\|_\infty \leq B$ . Then*

$$\begin{aligned} |H_V - H_W|_{\mathcal{A}} &\leq 2\pi(1 + BD)\|V - W\|_\infty, \\ \|V - W\|_\infty &\leq \frac{1 + BDe^{BD/2}}{2\pi}|H_V - H_W|_{\mathcal{A}}. \end{aligned} \quad (47.2)$$

*In particular, on every bounded normalized positive profile class of Section A.6, the physical same-type integrated law is a bi-Lipschitz coordinate for the full energy profile in this specified data norm. No smallness of  $V - 1$  is assumed.*

*The nullspace of  $|\cdot|_{\mathcal{A}}$  on  $C^3$  consists exactly of affine functions. Consequently it is a norm on  $\{H \in C^3 : H(0) = H'(0) = 0\}$ , which contains all profile-law differences. For arbitrary  $H \in C^3$  one also has*

$$|H|_{\mathcal{A}} \leq \frac{\pi}{2}|H''(0)| + 2\pi\sqrt{D}\sqrt{\|H'' - H''(0)\|_\infty\|H'''\|_\infty}. \quad (47.3)$$

*The same statements apply to  $C^{m-1,1}$  profiles,  $m \geq 4$ .*

*Proof.* Put  $u = V - W$ ,  $f = x^{-1/2}u$  and  $a = x^{-1/2}(V + W)$ . The normalization and the first derivative bounds make  $f$  continuous, with  $f(0) = 0$ . The convolution identities in Section 20 give, for  $Q = (\pi/2)(H_V - H_W)''$ ,

$$Q = a * f, \quad k_* * a = 2\pi + R, \quad R(0) = 0.$$

Indeed,

$$R(x) = \int_0^1 \frac{V(xt) + W(xt) - 2}{\sqrt{t(1-t)}} dt, \quad \|R'\|_\infty \leq \pi B.$$

The constant follows from  $\int_0^1 t\{t(1-t)\}^{-1/2} dt = \pi/2$ . Differentiating  $k_* * Q = (2\pi + R) * f$  gives

$$\mathcal{A}(H_V - H_W)(x) = 2\pi u(x) + \sqrt{x} \int_0^x R'(x-s) \frac{u(s)}{\sqrt{s}} ds. \quad (47.4)$$

All kernels are locally integrable; the differentiation follows first by writing the constant convolution as  $2\pi \int_0^x f$  and then using  $R(0) = 0$ . The displayed identity extends to zero by continuity. Its integral term has absolute value at most  $2\pi Bx\|u\|_\infty$ , proving the first estimate.

Let  $E = |H_V - H_W|_{\mathcal{A}}/(2\pi)$  and  $v(x) = \int_0^x |u(s)|s^{-1/2} ds$ . The same identity implies

$$|u(x)| \leq E + \frac{B}{2}\sqrt{x}v(x), \quad v'(x) \leq Ex^{-1/2} + \frac{B}{2}v(x)$$

almost everywhere. Multiplication by  $e^{-Bx/2}$  and integration, with  $v(0) = 0$ , yield

$$v(x) \leq E \int_0^x e^{B(x-s)/2} s^{-1/2} ds \leq 2E\sqrt{x}e^{Bx/2}.$$

Substitution proves the inverse estimate, including its stated constant. This argument uses the weighted forcing directly; it does not estimate  $f$  by dividing a uniform forcing bound by  $\sqrt{x}$ .

For the nullspace assertion,  $\mathcal{A}H = 0$  first gives  $H''(0) = 0$ . For  $x > 0$  it then gives  $k_* * H''' = 0$ . Convolution once more with  $k_*$ , whose square is the constant  $\pi$ , gives  $\pi(H'' - H''(0)) = 0$ . Thus  $H$  is affine. Conversely every affine function is annihilated by the definition.

To prove (47.3), write  $q = H'' - H''(0)$ ,  $M_0 = \|q\|_\infty$  and  $M_1 = \|q'\|_\infty$ . If either is zero the assertion follows directly. Otherwise  $q(0) = 0$  implies  $\ell = M_0/M_1 \leq D$ . For  $x \leq \ell$ ,  $|(k_* * q')(x)| \leq 2M_1\sqrt{x}$ . For  $x > \ell$ , split the integral at  $x - \ell$ . Integration by parts on  $[0, x - \ell]$ , using  $q(0) = 0$ , bounds that portion by  $2M_0/\sqrt{\ell}$ ; direct integration on the remaining interval bounds it by  $2M_1\sqrt{\ell}$ . Hence

$$\|k_* * q'\|_\infty \leq 4\sqrt{M_0M_1}.$$

Insert this into (47.1). The entire proof uses only the displayed derivatives and locally integrable convolutions, so also applies in the stated  $C^{m-1,1}$  class.  $\square$

The product on the right of (47.3) is an estimate, not the definition of a discrepancy norm. In particular, our regularization will minimize the seminorm of the linear transform  $\mathcal{A}$ , which satisfies the triangle inequality. Treating a geometric mean of two norms as a norm would not give that argument.

**47.1. An exact half-order benchmark.** The following calculation locates the loss between the two neighboring integer derivative norms. It does not identify an optimal statistical experiment or a sharp endpoint Hölder space.

**Proposition 47.2** (Linearization and fixed-class instability). *Let  $F(V) = T(V, V)$ . Its linearization at the constant profile is*

$$DF_1[w](d) = \frac{16}{3\pi d^2} \int_0^d x^{-1/2} (d-x)^{3/2} w(x) \, dx. \quad (47.5)$$

For  $w(x) = x^n$ ,  $n \geq 0$ , this equals  $\lambda_n d^n$ , where

$$\lambda_n = \frac{4(1/2)_n}{(n+2)!} = \frac{4 \binom{2n}{n}}{4^n (n+1)(n+2)} \sim \frac{4}{\sqrt{\pi}} n^{-5/2}. \quad (47.6)$$

At every such monomial the integrated transform is exactly inverted:

$$\mathcal{A}(\lambda_n d^{n+2}) = 2\pi d^n. \quad (47.7)$$

Fix  $D = 1$ ,  $m \geq 4$ ,  $B > 1$  and  $0 < \beta < 1$ . On the single class  $\mathcal{K}_m(B, \beta; 1)$  defined below there is no uniform Lipschitz inverse from the  $C^2$  norm of  $F(V) - F(W)$ , or of  $H_V - H_W$ , to  $\|V - W\|_\infty$ .

*Proof.* Bilinearity gives  $DF_1[w] = 2T(1, w)$ . Integrating out the unperturbed variable in the simplex, with  $\int_0^{d-x} (d-x-y)y^{-1/2} dy = \frac{4}{3}(d-x)^{3/2}$ , gives (47.5). The beta integral, or (20.9), gives the first expression for  $\lambda_n$ . The central binomial identity and Stirling's formula give the other two. For  $n \geq 1$ ,

$$\sqrt{x} (k_* * (x^n)')(x) = n B(n, 1/2) x^n = \frac{n! \sqrt{\pi}}{\Gamma(n+1/2)} x^n.$$

After two derivatives of  $\lambda_n x^{n+2}$ , all factors in (47.1) cancel to  $2\pi$ . For  $n = 0$ , its endpoint term gives the same identity directly.

Choose  $a > 0$  with  $(m+1)a < B-1$  and  $a < 1-\beta$  and, for  $n \geq m$ , put  $V_n^\pm(x) = 1 \pm an^{-m}x^n$ . Every derivative of the perturbation through order  $m$  is bounded by  $an^{r-m}$ . The last such bound controls the Lipschitz constant of the preceding derivative. Thus both signs belong to this same sum-norm class. Bilinearity cancels the quadratic perturbation terms exactly:

$$F(V_n^+) - F(V_n^-) = 2an^{-m}\lambda_n d^n.$$

For the maximum convention on  $C^2$ , the ratio of profile error to law error is  $\{\lambda_n n(n-1)\}^{-1}$ ; for integrated-law error it is  $\{\lambda_n(n+1)(n+2)\}^{-1}$ . Both tend to infinity like a positive multiple of  $\sqrt{n}$ . A sum convention changes these ratios by bounded factors. This proves the assertion on a fixed nontrivial class, rather than on classes with growing smoothness bounds.  $\square$

The monomial and mixed-norm benchmarks were given in the independent memorandum [21]. The two-sided flux norm and its use in the sampling and calibration theorems below make a further observation statement. The alternatives in Proposition 47.2 are abstract profile alternatives; no billiard realization or preparation lower bound is deduced from them. The genuinely physical input remains the relative flux theorem and the geometric meaning of the profiles in (20.1).



## 48. THE INVERSE AUTOCONVOLUTION PROBLEM

After the physical relative law has been proved, the inverse in Section 20 belongs to the theory of nonlinear Volterra equations. The passage from a billiard record to that equation, and the topology of its inversion, must be distinguished from the convolution machinery itself.

Dai and Lamm [18] develop local regularization for inverse autoconvolution, preserving the causal structure and proving convergence for continuous and  $L^2$  data. Their work supplies the relevant regularization context. We do not claim that causal inversion or a regularization principle is new, or apply an unverified theorem from that setting directly to the singular density here.

Hofmann, Werner and Deng [19] consider real  $L^2$  functions supported in a cube, with full or limited convolution data. Their limited-data uniqueness result uses nonnegativity and support reaching the origin; it belongs to a convolution-support framework. Our normalized kernel  $k(x) = x^{-1/2}V(x)$  has support at the origin and is locally integrable, but is not in  $L^2$  there:  $V(0) = 1$  implies  $\int_0^a k(x)^2 dx = \infty$ . Thus neither an  $L^2$  uniqueness theorem nor  $L^2$  ill-posedness is silently substituted for the statements proved here.

There are four specific features of the present inverse. First, the observed scalar law is  $F = T(V, V)$ , with two residual integrations between  $F$  and  $k * k = (\pi/2)(d^2 F)''$ . Second, normalization fixes the leading singularity of  $k$ , so the difference of two candidate kernels is continuous at zero and Abel reduction gives a forced second-kind equation. Third, all convolutions are restricted to a finite collar: reconstruction at energy  $d$  uses no values beyond  $d$ , no full-line Fourier inversion and no analytic continuation. Fourth, the deterministic stability statement uses  $C^3$  law error and uniform profile loss. It is a full-function statement for smooth profiles, rather than a determination of their Taylor series.

These features identify the exact functional setting of Theorem 20.2; they are not a claim of a new general deautoconvolution method. Its geometric content is the physical realization of the singular weighted energy measure and the relative estimate that makes it observable. Its statistical use requires another step. Section A.6 regularizes finite scalar data on a positive-node grid, retaining the preparation cost of the rare selected event. This uses a bounded smoothness class; it does not replace  $C^3$  by  $C^0$  in the unregularized theorem. The exponent obtained there is sufficient and is not asserted to be the optimal derivative loss or minimax exponent. Nor is an arbitrary regularization candidate asserted to be a realizable contact graph.

The Abel norm in Section 47 is a norm of integrated flux differences with zero initial value and slope; on arbitrary flux functions it has an affine nullspace. Its two-sided estimate is on the image of the normalized positive profile map. It does not assert that every scalar law is an autoconvolution square. The refinement in the observation theorem uses this linear discrepancy together with a globally smooth reconstruction and a different variance allocation. The subsequent calibration argument uses the physical identity that an onset error translates the integrated flux. These are specified consequences for this billiard invariant, not claims that Abel inversion, local polynomial regularization or Bernoulli concentration are new.

## 49. REGULARIZED OBSERVATION OF THE COMPLETE INVARIANT

We return to general contacts, without evenness, congruence or equal curvature assumptions. The observations in this section recover the full symmetrized energy profiles. The physical class is not replaced by the finite-dimensional family of the preceding section. The integrated-flux identity supplies two additional derivatives on exactly the profile class used below. This gain and its quantitative consequences were pointed out in the author-requested independent memorandum [22]. We include the proof and the complete modified observation argument; the earlier, weaker constructions remain in Sections A.6–A.8.

Fix  $m \geq 4$ ,  $D > 0$ ,  $B > 0$  and  $\beta > 0$ . Write

$$\|V\|_{r,1} = \sum_{a=0}^r \|V^{(a)}\|_{\infty} + \text{Lip}(V^{(r)}),$$

$$\mathcal{K}_m = \{V \in C^{m-1,1}([0, D]) : V(0) = 1, V \geq \beta, \|V\|_{m-1,1} \leq B\}.$$

The sum convention is part of this definition. The two actual profiles belong to this class. For two such functions define

$$H(V, W)(d) = d^2 T(V, W)(d) = \frac{2}{\pi} \int_{x+y < d} (d-x-y) \frac{V(x)W(y)}{\sqrt{xy}} dx dy, \quad H_V = H(V, V),$$

where the integral is over positive  $x, y$ . The target laws are  $T(\mathcal{V}_b, \mathcal{V}_c)$ . Constants below may depend on  $m, D, B, \beta$  and the specified physical bounds, but not on the accuracy, mesh or flight number.

#### 49.1. The regularity of integrated flux.

**Lemma 49.1** (Two-derivative gain). *For  $V, W \in \mathcal{K}_m$  one has*

$$\|H(V, W)\|_{m+1,1} \leq CB^2, \quad H(V, W)(0) = H(V, W)'(0) = 0. \quad (49.1)$$

*In particular, the candidate fluxes belong to a bounded  $C^{m+1,1}$  class even though the profiles are only bounded in  $C^{m-1,1}$ .*

*Proof.* Fubini's theorem and two integrations of a continuous convolution give

$$Q(d) := H(V, W)''(d) = \frac{2}{\pi} \int_0^1 \frac{V(ds)W(d(1-s))}{\sqrt{s(1-s)}} ds. \quad (49.2)$$

For  $0 \leq r \leq m-1$ , dominated differentiation gives

$$Q^{(r)}(d) = \frac{2}{\pi} \sum_{a=0}^r \binom{r}{a} \int_0^1 \frac{s^a (1-s)^{r-a} V^{(a)}(ds) W^{(r-a)}(d(1-s))}{\sqrt{s(1-s)}} ds.$$

The weight is integrable, including at both endpoints, and every profile derivative displayed is bounded by  $B$ . Thus  $\|Q^{(r)}\|_{\infty} \leq 2^{r+1} B^2$ . Every derivative of either profile through order  $m-1$  is Lipschitz with constant at most  $B$ : for a lower derivative use its next derivative, and for the last use the norm assumption. The difference of each product in the formula at  $d, e$  is consequently bounded by  $CB^2|d-e|$ . It follows that  $Q^{(m-1)}$  is uniformly Lipschitz. Finally

$$H'(d) = \int_0^d Q(s) ds, \quad H(d) = \int_0^d (d-s)Q(s) ds.$$

These identities control the two lower derivatives and prove (49.1). No extension of the profiles across zero, additional derivative observation, or higher profile norm is used.  $\square$

**49.2. Positive-node reconstruction.** Put  $q = m+2$ , let  $L \geq q$ , and write  $h = D/L \leq 1$ . The reconstruction  $J = J_h^{[q]}$  uses the positive nodes  $h, 2h, \dots, Lh$  only. To specify it, set

$$\omega(u) = \begin{cases} e^{-1/(1-u^2)}, & |u| < 1, \\ 0, & |u| \geq 1, \end{cases} \quad \theta_{\ell}(x) = \frac{\omega(x/h - \ell)}{\sum_{r=0}^L \omega(x/h - r)}.$$

For  $0 \leq \ell \leq L$ , interpolate the  $q$  nodal values starting at  $s_{\ell} = \max\{1, \min\{\ell - \lfloor (q-1)/2 \rfloor, L-q+1\}\}$  by a polynomial  $P_{\ell}z$  of degree at most  $q-1$ , and define  $Jz = \sum_{\ell=0}^L \theta_{\ell} P_{\ell}z$ . This is a globally smooth function on the closed interval. Its first stencil is  $h, \dots, qh$ ; the value at zero is extrapolated and is not inserted as an additional observation.

**Lemma 49.2** (Three reconstruction estimates). *Uniformly for the indicated meshes,*

$$|JH_V - H_V|_{\mathcal{A}} \leq Ch^{m-1/2}, \quad V \in \mathcal{K}_m, \quad (49.3)$$

$$|Jf|_{\mathcal{A}} \leq C\|f\|_{C^3}, \quad f \in C^3([0, D]), \quad (49.4)$$

$$|Jz|_{\mathcal{A}} \leq Ch^{-5/2} \max_k |z_k|. \quad (49.5)$$

*The seminorm is the linear Abel seminorm of Theorem 47.1, including its endpoint term.*

*Proof.* At each point at most two weights are positive. A nearest grid index lies within rescaled distance  $1/2$ , so the denominator is bounded below uniformly. Derivatives of the weights of order  $r$  are bounded by  $C_r h^{-r}$ . Every active stencil is within distance  $(q+1)h$ . The rescaled Lagrange polynomials and their fixed-order derivatives are bounded on these compact intervals. Since  $q$  is fixed, the product rule gives  $\|Jz\|_{C^r} \leq Ch^{-r} \max_k |z_k|$  for  $r \leq 3$ . The reconstruction reproduces every polynomial of degree at most  $q-1 = m+1$ . Subtracting the Taylor polynomial of that degree at the evaluation point, using a Lipschitz last derivative, gives

$$\|(JH_V - H_V)''\|_{\infty} \leq Ch^m, \quad \|(JH_V - H_V)'''\|_{\infty} \leq Ch^{m-1}.$$

Lemma 49.1 bounds the remainders uniformly. These estimates are valid at zero and  $D$ , since the rescaled stencils remain bounded there. For  $f \in C^3$ , subtract instead its quadratic Taylor polynomial. Its nodal remainder is  $O(h^3)\|f\|_{C^3}$ , so  $\|Jf\|_{C^3} \leq C\|f\|_{C^3}$ .

Apply the mixed Abel estimate

$$|G|_{\mathcal{A}} \leq \frac{\pi}{2} |G''(0)| + 2\pi\sqrt{D} \sqrt{\|G'' - G''(0)\|_{\infty} \|G'''\|_{\infty}}.$$

The approximation product has square root  $h^{m-1/2}$ , and its endpoint term is  $O(h^m)$ . The nodal-error product has square root  $h^{-5/2}$ , dominating its  $h^{-2}$  endpoint term. The  $C^3$  bound treats the smooth input with no mesh amplification. This proves all three claims. In particular, the gain is used only for limiting candidate fluxes, not to assume two extra derivatives of the physical finite-bridge error.  $\square$

**49.3. Calibrated binary observations.** For the moment suppose  $g, A, \gamma$  are supplied, as well as the selected patch labels, a common collar, and a relative convergence certificate. For even  $j \geq 2$  assume the physical selected means have

$$c_j = 2A \sinh(j\gamma), \quad H_{j,b} = c_j p_{j,b}, \quad 0 \leq H_{j,b}(d) \leq M_0 d^2, \quad \|H_{j,b} - H_{\mathcal{V}_b}\|_{C^3([0,D])} \leq M_1 \tau^j. \quad (49.6)$$

These are consequences of the relative physical law on the supplied compact class. The values  $M_0, M_1$  and  $\tau \in (0, 1)$ , rather than the unknown profiles or finite-flight remainders, are supplied.

**Theorem 49.3** (Complete-profile preparation bound). *For  $L \geq m+2$ ,  $h = D/L \leq 1$ , even  $j \geq 2$  and  $0 < t, \delta \leq 1$ , use*

$$N_{b,k} = \left\lceil 2c_j(M_0 d_k^2 + t/3)t^{-2} \log \frac{4L}{\eta} \right\rceil, \quad d_k = kh, \quad (49.7)$$

*fresh binary preparations in orientation  $b$  at time  $jg + d_k$ . There is a Borel estimator with values in  $\mathcal{K}_m^2$  such that, with probability at least  $1 - \eta$ ,*

$$\max_b \|\widehat{V}_b - \mathcal{V}_b\|_{\infty} \leq C\{h^{m-1/2} + \tau^j + th^{-5/2} + \delta\}. \quad (49.8)$$

*The predicted odd law has at most  $2B$  times this error. The total number of preparations is deterministically bounded by*

$$N_{\text{tot}} \leq 2L + CA \sinh(j\gamma) h^{-1} t^{-2} \log \frac{4L}{\eta}. \quad (49.9)$$

In particular, uniform accuracy  $\varepsilon$  for both profiles and the odd-law prediction, at confidence  $1 - \eta$ , is attained using

$$N_{\text{tot}} \leq C\varepsilon^{-(2+\frac{6}{m-1/2}+\frac{\gamma}{|\log \tau|})} \log \frac{C}{\eta\varepsilon}. \quad (49.10)$$

Every failure is charged. No value at zero or derivative is observed. This is a sufficient preparation bound, not a minimax or computation-time bound.

*Proof.* The class  $\mathcal{K}_m$  is compact in  $C^{m-1}$  by Arzelà–Ascoli and the closed sum-norm bound. For fixed  $h$ , the map  $V \mapsto \mathcal{A}JH_V$  is continuous through its finite vector of nodal values, so its image in  $C^0$  is compact. Choose a finite ordered dictionary  $\mathcal{D}_{h,\delta} \subset \mathcal{K}_m$  whose transformed nodal images form a  $\delta$ -net. For a data vector  $z$ , take the first minimizer over the dictionary of  $|JH_W - Jz|_{\mathcal{A}}$ . Each discrepancy is continuous in the finite data vector. The first-index rule therefore defines a Borel map, including ties. No physical realization of dictionary elements is required for this upper bound.

Let  $Z_{b,k} = c_j \bar{Y}_{b,k}$ . Its mean is  $H_{j,b}(d_k)$  and

$$\text{Var}(Z_{b,k}) = \frac{c_j H_{j,b}(d_k)(1 - p_{j,b}(d_k))}{N_{b,k}} \leq \frac{c_j M_0 d_k^2}{N_{b,k}}.$$

A centered summand is bounded by  $c_j$ . The Bernoulli exponential moment bound consequently gives

$$\Pr\{|Z_{b,k} - H_{j,b}(d_k)| > t\} \leq 2 \exp\left\{-\frac{N_{b,k} t^2}{2c_j(M_0 d_k^2 + t/3)}\right\}.$$

The allocations and a union bound make every nodal error at most  $t$  with probability  $1 - \eta$ . On this event, (49.4) treats the smooth bridge error and (49.5) treats only the centered scalar errors, giving

$$|JZ_b - JH_{\mathcal{V}_b}|_{\mathcal{A}} \leq C(th^{-5/2} + \tau^j).$$

The minimizing property and the dictionary net bound the discrepancy of the reconstructed candidate images by twice this quantity plus  $\delta$ . Apply (49.3) to the true and fitted candidate fluxes. The triangle inequality for the linear seminorm gives their unregularized Abel discrepancy at most  $C(h^{m-1/2} + th^{-5/2} + \tau^j + \delta)$ . The two-sided inverse of Theorem 47.1 proves (49.8). The positive mass-one kernel of  $T$  gives the  $2B$  bound for the odd-law prediction.

Summing (49.7) uses  $\sum d_k^2 \leq D^2 L$  and  $L = D/h$ , and includes each ceiling. This proves (49.9). Put  $\nu = m - 1/2$ . Choose an integer mesh with  $h \asymp \varepsilon^{1/\nu}$ ,  $t = c\varepsilon h^{5/2}$ , and  $\delta = c\varepsilon$ . Take the least even  $j \geq 2$  with  $\tau^j \leq c\varepsilon$ . Then

$$h^{-1}t^{-2} = c^{-2}\varepsilon^{-2}h^{-6}, \quad e^{j\gamma} \leq C\varepsilon^{-\gamma/|\log \tau|}.$$

Integer rounding changes only fixed constants. Substituting in (49.9), and absorbing the smaller ceiling term, proves (49.10).  $\square$

**Corollary 49.4** (A uniform-data conditional modulus). *For  $V, W \in \mathcal{K}_m$ , put  $e = \|H_V - H_W\|_{\infty}$ . Then*

$$\|V - W\|_{\infty} \leq C e^{(m-1/2)/(m+2)}. \quad (49.11)$$

*This is a conditional estimate on a fixed regularity class of forward images. No optimality of its exponent is asserted.*

*Proof.* Use the inverse theorem, insert  $JH_V, JH_W$ , and apply (49.3) and (49.5). This gives  $\|V - W\|_{\infty} \leq C(h^{m-1/2} + eh^{-5/2})$ . For sufficiently small  $e > 0$ , choose an integer mesh with  $h \asymp e^{1/(m+2)}$ . If  $e = 0$ , let  $h \downarrow 0$ ; for larger  $e$ , the bound  $\|V - W\|_{\infty} \leq 2B$  permits an enlargement of  $C$ . These cover all cases.  $\square$

**49.4. Unknown gap, area and multiplier.** We next remove the actual values of  $g, A, \gamma$ , not the supplied labels or class certificates. Assume precisely the pilot model (A.8.5): a supplied positive coarse gap bracket of width at most  $d_c$ , positive compact area and multiplier boxes, and one- and two-flight physical probabilities with exact leading coefficients  $[2A \sinh(r\gamma)]^{-1}$  and unknown  $C^{m-1,1}$  remainders  $F_r$ , normalized at zero and bounded below. Their common norm and positivity constants are supplied. In addition, assume the relative  $C^3$  estimates on  $[0, D_+]$ , a uniform  $C^{3,1}$  bound for  $H_{j,b}$  there, and a fixed  $D < D_+$ . These are the same assumptions as in Theorem A.8.3; the profile class is not replaced by  $\mathcal{K}_{m+2}$ .

**Theorem 49.5** (Self-calibrated complete-profile observation). *Under these supplied smooth physical bounds, for sufficiently small  $\varepsilon > 0$  and  $0 < \eta < 1/4$ , a Borel procedure recovers both complete energy profiles on  $[0, D]$  and their predicted odd law to uniform error  $\varepsilon$ , with probability at least  $1 - \eta$ . It uses a separately charged one-/two-flight pilot and fresh positive-offset bits at one subsequent even flight number. Its deterministic total preparation count is at most*

$$C\varepsilon^{-(2+\frac{6}{m-1/2}+\frac{\gamma_+}{\lceil \log \tau \rceil})} \log \frac{C}{\eta\varepsilon}. \quad (49.12)$$

*The estimate includes the energy origin and does not infer separate contact curvatures or unsymmetrized contact branches.*

*Proof.* We give the plug-in argument because an error in the programmed time must be handled before regularization. Suppose a pilot yields, on an event of probability  $1 - \eta_0$ ,

$$0 \leq \hat{g} - g \leq \rho, \quad |\hat{A}/A - 1| \leq \alpha, \quad |\hat{\gamma} - \gamma| \leq \zeta, \quad j\zeta \leq 1, \quad j\rho \leq D_+ - D.$$

All estimates are projected onto their supplied compact boxes even on bad histories. Program time  $j\hat{g} + d$  and scale a new bit by  $\hat{c}_j = 2\hat{A} \sinh(j\hat{\gamma})$ . Conditionally on the pilot its mean is exactly

$$RH_{j,b}(d + \Delta), \quad \Delta = j(\hat{g} - g), \quad R = \frac{\hat{A} \sinh(j\hat{\gamma})}{A \sinh(j\gamma)}.$$

Since  $\gamma \geq \gamma_- > 0$ , differentiation of  $\log \sinh(j\gamma)$  gives  $|R - 1| \leq C(\alpha + j\zeta)$ . The one-sided gap estimate keeps every translated argument in  $[0, D_+]$ . The  $C^{3,1}$  flux bound gives a  $C^3$  translation error  $O(\Delta)$ , including at zero. Hence

$$\|RH_{j,b}(\cdot + \Delta) - H_{j,b}\|_{C^3([0,D])} \leq C(\tau^j + \alpha + j\zeta + j\rho). \quad (49.13)$$

No factor  $\Delta/d$  occurs, since the bits have not been divided by  $d^2$ .

On the good pilot event the mean is bounded by a supplied constant  $M_H$ . Allocate at each node

$$\left\lceil 2\hat{c}_j(M_H + t/3)t^{-2} \log(4L/\eta_1) \right\rceil$$

fresh preparations. Conditional Bernoulli concentration gives nodal error at most  $t$  with probability  $1 - \eta_1$ . Applying (49.4) to (49.13), and using the same dictionary argument, yields error at most

$$C\{h^{m-1/2} + th^{-5/2} + \delta + \tau^j + \alpha + j\zeta + j\rho\}.$$

The smooth calibration term has no inverse power of  $h$ . The charge on every history is bounded by  $2L + Ce^{j\gamma}h^{-1}t^{-2} \log(4L/\eta_1)$ , since the pilot estimates are always in their compact boxes.

Apply Lemma A.8.2 with  $\xi = c\varepsilon/j$  and  $\eta_0 = \eta/2$ . Its proof, retained in full below, brackets the onset using charged one-flight trials and extrapolates the two unknown smooth remainders from positive nodes; it does not assume an exact-family probability oracle. It supplies  $\alpha + \zeta \leq C\xi$  and  $\rho \leq C\xi^{1+1/m}$  with pilot charge

$$C(j/\varepsilon)^{2+2/m} \log \frac{C \log(Cj/\varepsilon)}{\eta}.$$

Choose the least even  $j$  with  $\tau^j \leq c\varepsilon$ ,  $\eta_1 = \eta/2$ , and the mesh,  $t$  and  $\delta$  from Theorem 49.3. For small enough  $c$  and  $\varepsilon$ , all the displayed pilot constraints hold and the profile error is at most  $\varepsilon$ . Conditioning and the two failure allocations give total failure probability at most  $\eta$ . Finally

$$2 + \frac{6}{m - 1/2} + \frac{\gamma_+}{|\log \tau|} > 2 + \frac{2}{m}.$$

This strict power gap absorbs the pilot cost and its fixed powers of  $\log(1/\varepsilon)$  into (49.12). The allocation is bounded independently of pilot success, proving the unconditional deterministic preparation claim.  $\square$

For  $m = 4$ , the profile-stage accuracy power, excluding the common bridge factor, is  $26/7$  rather than the power 6 of the earlier  $m$ -node Abel construction. Both are sufficient bounds for specified procedures. A larger admissible value of  $\tau$  weakens the displayed bridge factor without changing the billiard. Neither this factor nor the complete exponent is asserted to be an intrinsic information law. The physical finite-dimensional lower bound of Theorem 51.3 and the separate nuisance lower bounds below do not imply a lower bound for this full-profile experiment. The two even laws identify the invariant; separately measured odd records can validate, rather than further identify, its predicted compatibility law.

## Part 5. Finite-resolution inversion and global physical reconstruction

### 50. QUANTITATIVE RECONSTRUCTION FROM QUANTIZED TRANSVERSE LAWS

The exact inverse admits a conditional quantitative version in an observation norm which does not involve derivatives of a measured density. The observation is a finite histogram of signed transverse endpoints, together with an onset estimate for each selected channel. Its resolution is allowed to increase as the required geometric error decreases. In particular, no longitudinal endpoint coordinate is retained. We first give a deterministic estimate between two realizable tables, and then a finite-sample confidence construction. The analytic continuation estimate is deliberately conservative. Its constants record, rather than suppress, the deterioration with the continuation distance and the analytic width; compare [37].

**50.1. Quantified classes and observation distances.** Fix a selected graph  $E = \mathcal{T} \cup \{f_1, f_2\}$  supplied by Theorem 29.3, with  $|E| = N + 1$ . Its obstacle labels, original deck labels, root incidence, tree-path labels  $m_a$ , and the invertible integer gain matrix  $M = (\eta_{f_1} \ \eta_{f_2})$  are fixed. The marked Euclidean realization  $L$  is not supplied. Fix a positive offset  $d_0$  below the common relative-law collar of the selected channels, and nested squares

$$Q_0 = [-r_0, r_0]^2 \Subset Q_1 = [-r_1, r_1]^2, \quad 0 < |a_*| < r_0 < r_1,$$

in the signed contact coordinates of each measured law. The two laws of a channel always refer to the two same-type even bridges. There are  $\mathcal{L} = 2N + 2$  laws in total.

A *quantified analytic class*  $\mathcal{K}$  consists of realizable tables with this persistent design and the following bounds. All gaps, contact curvatures and clearances lie in prescribed positive bounded ranges; the lattice and representative placements are bounded after the fixed root gauge is chosen. Each support function in its contact-centered proper frame extends  $2\pi$ -periodically and holomorphically to

$$|\operatorname{Im} z| < s, \quad |h(z)| \leq B, \quad B \geq 1, \quad (50.1)$$

with prescribed  $s > 0$ . The contact frames use the outward normal and its positively oriented tangent. Thus  $h(0) = h'(0) = 0$  in each individual contact-centered frame; the placement of the second body in a channel also uses its observed gap. For each limiting density  $f$  on  $Q_1$ , assume

$$f \geq m_* > 0, \quad q_f \geq q_* > 0, \quad \|f\|_{C^k(Q_1)} \leq H_k \quad (k \geq 1), \quad (50.2)$$

where  $q_f$  is the scalar anchor of (23.3), with the fixed nonzero coordinate  $a_*$ . Only the finitely many  $H_k$  used in a chosen finite-order estimate are needed. These are a priori forward bounds, not observations of the amplitude function and not bounds on an unknown inverse modulus. They can equally be furnished by quantitative contact bounds and the weighted forward estimates.

For each obstacle label  $a$  fix a finite set  $F_a \subset \{2, 3, \dots\}$  and integers  $b_{a,k}$  such that

$$\sum_{k \in F_a} kb_{a,k} = 1, \quad |z_k(K_a)| \geq \zeta_* > 0 \quad (k \in F_a). \quad (50.3)$$

The Fourier convention is that of (29.8); coefficients of order at least two are translation invariant. The actual coefficients and phases are unknown. No harmonic truncation of an obstacle is imposed.

Every properly asymmetric analytic table has such a finite witness, by the greatest-common-divisor argument in Proposition 29.6. It also has a positive analytic width. After the width is decreased, sufficiently small neighborhoods in the corresponding bounded holomorphic-support norm have the bounds above. This assertion does not give a common analytic width on a  $C^2$  neighborhood. The qualitative open-dense theorem retains its original topology and does not acquire these quantitative hypotheses.

For two labelled realizations put

$$D(T, \tilde{T}) = \inf_{(U, t) \in \text{SE}(2)} \left\{ \|L - U\tilde{L}\| + \max_a d_H(K_a, U\tilde{K}_a + t) \right\}. \quad (50.4)$$

Here representatives and deck marks are fixed; there is no independent change of representatives inside the infimum. On the prescribed bounded class the same estimates control the difference of lattice Gram matrices.

Partition  $Q_1$  into a rectangular grid of cells with side lengths at most  $\delta$  and include its complement as one further category. Half-open cells remove overlap, and cell boundaries have zero probability. Write  $\mathbf{p}_\delta(P)$  for the resulting probability vector. Define

$$\begin{aligned} e_g &= \max_{e \in E} |g_e - \tilde{g}_e|, \\ e_p &= \max_{e, b} \|\mathbf{p}_\delta(P_{e,b}) - \mathbf{p}_\delta(\tilde{P}_{e,b})\|_1, \\ \eta &= e_p + C_{\text{bin}}\delta, \quad C_{\text{bin}} = 4H_1|Q_1|. \end{aligned} \quad (50.5)$$

All laws in this display are limiting laws at  $d_0$ . The inclusion of the complement category keeps the observation a probability vector; the proof uses density regularity only inside  $Q_1$ .

## 50.2. From cells to a finite support jet.

**Lemma 50.1** (Interior interpolation from histogram probabilities). *Let  $f, \tilde{f}$  satisfy (50.2). For every integer  $m \geq 0$  there is a finite constant  $I_m$ , depending only on  $Q_0, Q_1, H_{m+1}$  and a fixed polynomial kernel, such that, when  $\eta \leq 1$ ,*

$$\|f - \tilde{f}\|_{C^m(Q_0)} \leq I_m \eta^{1/(m+3)}. \quad (50.6)$$

*The constant is obtainable by differentiation and integration of finite polynomials and by the displayed forward bounds.*

*Proof.* Let  $\bar{f}_\delta$  be the cell-average function on  $Q_1$ . The mean value inequality gives  $\|f - \bar{f}_\delta\|_{L^1(Q_1)} \leq 2H_1|Q_1|\delta$ . The  $L^1$  difference of two cell-average functions is the sum of the absolute differences of their cell masses. Consequently

$$\|f - \tilde{f}\|_{L^1(Q_1)} \leq \eta. \quad (50.7)$$

Set  $w = f - \tilde{f}$ . Use the product kernel  $K_m(x_1, x_2) = c_m(1 - x_1^2)_+^{m+2}(1 - x_2^2)_+^{m+2}$ , normalized to have integral one. It is sufficiently differentiable for all integrations by parts below; its relevant derivatives and normalization are finite polynomial expressions and beta integrals. For  $h \leq h_0 := \min\{1, (r_1 - r_0)/2\}$  and  $|\alpha| \leq m$ ,

$$\|\partial^\alpha(K_{m,h} * w)\|_{L^\infty(Q_0)} \leq \|\partial^\alpha K_m\|_\infty h^{-2-|\alpha|} \eta.$$

All convolution points are inside  $Q_1$ . Integration by parts and the mean value inequality also give

$$\|\partial^\alpha w - K_{m,h} * \partial^\alpha w\|_{L^\infty(Q_0)} \leq CH_{m+1}h.$$

Choose  $h = h_0\eta^{1/(m+3)}$  and add the two estimates. Since  $h \leq 1$ , the worst power in the first estimate is  $h^{-m-2}$ , proving (50.6). When  $\eta = 0$ , let  $h$  decrease to zero. The a priori derivative bound is essential: the assertion is not an unconditional embedding of total variation into  $C^m$ .  $\square$

The finite-jet inverse in Section 22 has an additional useful description. It makes the constants used here finite computations, without replacing the smooth finite-remainder proof by formal algebra.

**Lemma 50.2** (A finite Bellman recipe for the jet constants). *For each  $m \geq 2$ , the finite action-to-graph inverse and the conversion from a graph jet to its contact-centered support jet are finite smooth formulas with denominators bounded away from zero on the prescribed positive geometric ranges. Their two-point Lipschitz bounds are obtained by a finite arithmetic and differentiation recursion. In particular there is a computable finite  $A_m \geq 1$  such that the support jets recovered from two realizable sets of channel data obey*

$$\sum_{k=0}^m \frac{r^k}{k!} |h^{(k)}(0) - \tilde{h}^{(k)}(0)| \leq A_m \{e_g + \eta^{1/(m+3)}\}, \quad r = \min\{s/8, 1/8\}. \quad (50.8)$$

*Every incidence is covered by the same constant after taking the maximum over the finite design. The constants are not asserted to be bounded uniformly in  $m$ .*

*Proof.* The four-density identities and the scalar positive anchor give a finite  $C^m$  action estimate from (50.6). This can be a two-point estimate on the whole quantified realizable class: compare products, reciprocals and the positive scalar roots directly, using the lower bounds at both endpoints. An interpolating density need not itself factorize. The exact reconstructed actions have zero constant and linear terms.

For completeness, the lower-order remainders in the graph recursion can be computed as follows. Let  $V_b(u)$  be the first interior coordinate of the stationary half-line beginning at  $u$ . Uniqueness of the decaying half-line and the absolutely convergent action yield the local identities

$$\begin{aligned} S_b(u) &= \ell_b(u, V_b(u)) - g + S_{1-b}(V_b(u)), \\ 0 &= \partial_2 \ell_b(u, V_b(u)) + S'_{1-b}(V_b(u)). \end{aligned} \quad (50.9)$$

The derivative of the second equation with respect to  $V_b$  at zero is  $g^{-1} + \kappa_{1-b} + a_{1-b} > 0$ , where  $a_b = S''_b(0)$ . Its linear solution is

$$V'_b(0) = \alpha_b = \mathbf{r}_b e^{-\gamma}, \quad \alpha_0 \alpha_1 = e^{-2\gamma} < 1.$$

Recover the quadratic geometry by (22.3). Suppose all orders below  $n \geq 3$  have been found. In the common transverse channel coordinates the local length used in this computation is explicitly

$$\ell_b(u, v) = \sqrt{(g + \psi_b(u) + \psi_{1-b}(v))^2 + (v - u)^2}.$$

Set the order- $n$  graph and action coefficients provisionally to zero. Solve the second identity in (50.9) through degree  $n - 2$  for the Taylor polynomial of  $V_b$ . This uses only the already determined coefficients, and divides at each step by the positive scalar just displayed. Inserting



that polynomial into the first identity gives a known order- $n$  coefficient  $P_{b,n}$ , expressed in derivative rather than divided-power normalization. The true order- $n$  coefficients satisfy

$$(I - B_n)s_n = (I + B_n)q_n + P_n, \quad B_n = \begin{pmatrix} 0 & \alpha_0^n \\ \alpha_1^n & 0 \end{pmatrix}. \quad (50.10)$$

Indeed the new graph terms in the local flight length are  $(q_{b,n}u^n + q_{1-b,n}v^n)/n!$ , and the new terminal action term is  $s_{1-b,n}v^n/n!$ . They can be evaluated at  $v = \alpha_b u$ . Omitting the unknown degree- $n - 1$  term of  $V_b$  changes the stationary value only by  $O(u^{2n-2})$ , which is of higher order than  $n$ . This last assertion follows by Taylor expansion around the stationary point: the first derivative there vanishes and the displacement is  $O(u^{n-1})$ . The finite-jet factorization and smooth remainders proved in Lemma 22.4 justify using these Taylor identities for the actual functions.

Both  $I - B_n$  and  $I + B_n$  have determinant  $1 - e^{-2n\gamma}$ . Thus

$$q_n = (I + B_n)^{-1}\{(I - B_n)s_n - P_n\}. \quad (50.11)$$

In particular  $(I - B_n)^{-1}(I + B_n)$  is precisely the matrix  $M_n$  of (22.4). Formula (50.11), including the construction of  $P_n$ , involves only finitely many truncated series operations. It is not a numerical truncation of an uncontrolled infinite orbit.

In an individual outward-normal frame the boundary is  $(-\psi(y), y)$ . Its normal angle satisfies  $\psi'(y) = \tan \theta$ , so

$$h(\theta) = -\psi(y(\theta)) \cos \theta + y(\theta) \sin \theta. \quad (50.12)$$

The implicit scalar derivative is  $\psi''(0) = \kappa > 0$ . For order  $n$ , solve for  $y$  through degree  $n - 2$ ; the stationary-value identity in (50.12) shows that the support coefficient through order  $n$  depends on the graph coefficients only through order  $n$ . For  $n = 2$  it is  $h''(0) = 1/\kappa$ ; for  $n \geq 3$  the same quadratic remainder argument as above applies. The channel coordinates use a common transverse axis. On the side whose outward normal is opposite to that of the first side, the positively oriented tangent is also reversed. Its individual graph coefficients are therefore  $(-1)^n q_{b,n}$  when that reversal is required. These known proper changes fix the signs before (50.12) is applied; they introduce no unknown inter-channel pose.

Here is an explicit bounding prescription for  $A_m$ . Bound the input Taylor coefficients by the supplied forward bounds and the contact bounds following from (50.1) and positive curvature. Apply truncated product, reciprocal, square-root and implicit-series recursions to the four-density formula, (22.3), (50.11), and (50.12). At each finite operation use the product difference inequality and, for a reciprocal,  $|x^{-1} - y^{-1}| \leq |x - y|/c^2$  when  $|x|, |y| \geq c$ . The scalar roots use their positive lower bounds. The other denominators are  $g$ , positive leading action curvatures,  $\kappa$ ,  $g^{-1} + \kappa_{1-b} + a_{1-b}$  and  $1 - e^{-2n\gamma}$ ; their lower bounds are fixed, with the last at least  $1 - e^{-6\gamma_-}$ . Elementary hyperbolic functions and their finitely many derivatives are bounded on the prescribed leading-geometry intervals. Summing the resulting finite coefficient bounds with weights  $r^k/k!$  and including  $I_m$  gives  $A_m$ . No supremum over an unknown infinite-dimensional inverse or a compactness separation constant enters this prescription.  $\square$

### 50.3. A quantified continuation of the recovered support jet.

**Lemma 50.3** (Continuation from a finite anchored jet). *Suppose  $h, \tilde{h}$  satisfy (50.1), and put*

$$r = \min\{s/8, 1/8\}, \quad J_s = \lceil \pi/r \rceil + 1, \quad \beta_s = 2^{-J_s}.$$

*If  $\sum_{k=0}^m r^k |h^{(k)}(0) - \tilde{h}^{(k)}(0)|/k! \leq u$ , then*

$$\|h - \tilde{h}\|_{L^\infty(\mathbb{R})} \leq (2B)^{1-\beta_s} \left[ \min \left\{ 2B, u + \frac{8B}{3} 4^{-m-1} \right\} \right]^{\beta_s}. \quad (50.13)$$

For a fixed derivative order  $k$  the right side, multiplied by  $k!(2/r)^k$ , also bounds  $\|(h - \tilde{h})^{(k)}\|_\infty$ .

*Proof.* Write  $F = h - \tilde{h}$  and  $H = 2B$ . The disk of radius  $4r$  about any real point lies strictly inside the given strip. Cauchy's coefficient bound at zero gives on  $|z| \leq r$

$$|F(z)| \leq u + H \sum_{j=m+1}^{\infty} 4^{-j} = u + \frac{8B}{3} 4^{-m-1}.$$

It is also at most  $H$ . If the bound on a disk of radius  $r$  centered on the real line is  $v \leq H$ , the three-circles inequality at radii  $r, 2r, 4r$  bounds the middle disk by  $\sqrt{Hv}$ . One may obtain this inequality directly by comparing the subharmonic function  $\log |F|$  on the annulus with the affine function of  $\log |z|$  having boundary values  $\log v$  and  $\log H$ ; regularize zeros and let the regularization decrease to zero. A disk of radius  $r$  whose center is shifted by at most  $r$  is contained in that middle disk. Iterating along the real line therefore bounds the  $j$ th disk by  $H^{1-2^{-j}} v^{2^{-j}}$ . In each of the two real directions,  $J_s$  steps suffice to cover  $[-\pi, \pi]$ , including a disk of radius  $r/2$  about every point there. Periodicity and Cauchy's derivative bound on the smaller disks prove the assertions. When  $v = 0$ , use the same argument with  $v + \epsilon$  and let  $\epsilon$  decrease to zero.  $\square$

#### 50.4. Registration with arbitrary finite asymmetry witnesses.

**Lemma 50.4** (Bézout registration and cochain stability). *For two congruent properly oriented copies  $C_s, C_t$  of obstacle type  $a$ , set  $p_k(C) = z_k(C)/|z_k(C)|$ . Under (50.3) their unique proper congruence has rotation*

$$U_{s \rightarrow t} = \prod_{k \in F_a} \left( \frac{p_k(C_s)}{p_k(C_t)} \right)^{b_{a,k}} \quad (50.14)$$

and translation  $s(C_t) - U_{s \rightarrow t} s(C_s)$ . Consequently the registration and cochain formulas (29.9)–(29.10) are uniformly Lipschitz in the complete channel-image supports and gaps on the quantified class. Denote one such finite constant by  $C_{\text{reg}}$ . It depends on the design, its integer marks,  $\|M^{-1}\|$ , the prescribed size bounds,  $\zeta_*^{-1}$  and  $\sum_{a,k} |b_{a,k}|$ .

*Proof.* If  $C_t = e^{i\omega} C_s + t$ , then  $p_k(C_t) = e^{-ik\omega} p_k(C_s)$ . Equation (50.3) makes the product in (50.14) equal to  $e^{i\omega}$ , without choosing a root or an argument branch. Centering supplies the displayed translation. The same identity proves uniqueness. For  $|z|, |w| \geq \zeta_*$ ,

$$\left| \frac{z}{|z|} - \frac{w}{|w|} \right| \leq \frac{2|z - w|}{\zeta_*}.$$

Fourier coefficients and the first-harmonic center are bounded linear functionals of a support function. A product of unit complex factors has difference bounded by the sum of the factor differences, with the factor  $|b|$  for a power  $b \in \mathbb{Z}$ . This proves an explicit uniform bound for each rotation and translation. Propagate these bounds through the finitely many tree incidences. The displacement cochain, its two cycle defects and multiplication by the fixed  $M^{-1}$  are finite linear operations with bounded coefficients. The representative correction  $c_a = p_a - Lm_a$  uses the same recovered  $L$  for all labels. The difference of two support functions equals the Hausdorff distance of their convex bodies in the uniform norm: the support inequalities characterize containment in the corresponding parallel body. Rotating an approximate support introduces at most its bounded first derivative times the rotation error, and this derivative is bounded by Cauchy's estimate from (50.1). Adding the gap error for the second contact placement completes the claimed table estimate. Finally  $\|L^t L - \tilde{L}^t \tilde{L}\| \leq (\|L\| + \|\tilde{L}\|)\|L - \tilde{L}\|$  in the common proper frame.  $\square$

The choices  $F_a = \{2, 3\}$  and  $(b_{a,2}, b_{a,3}) = (-1, 1)$  recover the earlier harmonic formula. For example  $F_a = \{4, 9\}$  admits  $(b_{a,4}, b_{a,9}) = (-2, 1)$  and needs neither a second nor a third harmonic. The formula applies to the whole properly asymmetric locus through its local finite witnesses, not only to the explicit two-harmonic sublocus.

**50.5. The deterministic finite-resolution inverse.** Define, for  $e_g \geq 0$  and  $0 \leq \eta \leq 1$ ,

$$\begin{aligned} W_m(e_g, \eta) &= \min \left\{ 2B, A_m(e_g + \eta^{1/(m+3)}) + \frac{8B}{3} 4^{-m-1} \right\}, \\ \mathfrak{W}(e_g, \eta) &= \inf_{m \geq 2} W_m(e_g, \eta), \\ \Omega(e_g, \eta) &= C_{\text{reg}} \left\{ e_g + (2B)^{1-\beta_s} \mathfrak{W}(e_g, \eta)^{\beta_s} \right\}. \end{aligned} \quad (50.15)$$

The constants may be enlarged to cover all channel incidences and the Gram matrix. Each fixed- $m$  expression is already a valid upper bound; computing an infinite infimum is not required.

**Theorem 50.5** (Quantized-law stability of the complete periodic table). *Let  $T, \tilde{T} \in \mathcal{K}$  have the same selected design and satisfy (50.5), with  $\eta \leq 1$ . Then*

$$D(T, \tilde{T}) \leq \Omega(e_g, \eta). \quad (50.16)$$

*The same bound, after the stated enlargement of  $C_{\text{reg}}$ , controls their marked lattice Gram matrices. Moreover  $\Omega(e_g, \eta) \rightarrow 0$  as  $(e_g, \eta) \rightarrow (0, 0)$ . Thus finitely many scalar cell probabilities at every prescribed resolution, and finitely accurate gaps, determine the entire table to a prescribed positive error. This conclusion uses no longitudinal positions, supplied support coefficients, amplitude functions or inter-channel registration. It is not exact determination from a fixed finite list of scalar measurements.*

*Proof.* Lemma 50.1 converts the finite cell discrepancy into a fixed-order interior density error. The finite inverse and support conversion of Lemma 50.2 give the anchored support-jet bound (50.8) for every incidence. Lemma 50.3 yields the complete channel-image bound with  $W_m$  on its right. Lemma 50.4 puts these images in one root frame and recovers the table and lattice with error at most the corresponding fixed- $m$  expression in (50.15). Take the infimum over  $m$ . For convergence, first choose a fixed  $m$  so large that its geometric tail is small, and then decrease  $e_g, \eta$  with this  $m$  held fixed. This order of choices uses no uniform bound on  $A_m$  as  $m$  increases.  $\square$

The estimate gives a finite accuracy prescription, not merely a separation constant obtained by compactness. Given  $\varepsilon > 0$ , set

$$v = \left\{ \frac{\varepsilon}{2C_{\text{reg}}(2B)^{1-\beta_s}} \right\}^{1/\beta_s}.$$

Choose a finite  $m \geq 2$  with  $(8B/3)4^{-m-1} \leq v/2$ , and require

$$e_g \leq \min \left\{ \frac{\varepsilon}{2C_{\text{reg}}}, \frac{v}{4A_m} \right\}, \quad \eta \leq \min \left\{ 1, \left( \frac{v}{4A_m} \right)^{m+3} \right\}. \quad (50.17)$$

Then (50.16) is at most  $\varepsilon$ . Taking  $C_{\text{bin}}\delta \leq \eta/2$  and  $e_p \leq \eta/2$  specifies a finite grid and accuracy for its cell probabilities. The prescription is potentially very ill-conditioned. No sharp Hölder exponent, analytic minimax rate, or uniform Lipschitz law-to-boundary inverse is asserted.

### 50.6. Finite samples, finite flights, and charged acquisition.

**Corollary 50.6** (Histogram confidence reconstruction). *Suppose  $n$  independent successful marks are available for each of the  $\mathcal{L}$  calibrated same-type channels at a fixed positive offset, and retain only their grid categories. Let the grid, including the complement, have  $K_\delta$  categories. If the marks have even-flight laws  $P_{j,e,b}$  satisfying*

$$\|P_{j,e,b} - P_{e,b}\|_{\text{TV}} \leq C_0 \tau^j, \quad 0 < \tau < 1, \quad (50.18)$$

define

$$u_n = \sqrt{K_\delta/n} + \sqrt{2 \log(\mathcal{L}/\alpha)/n}, \quad b_j = 2C_0 \tau^j.$$

*On the event that supplied gap estimates have error at most  $v_g$ , form the set of candidate tables in  $\mathcal{K}$  whose gaps are within  $v_g$  of those estimates and whose limiting histogram vectors are within  $u_n + b_j$  of the empirical vectors in  $\ell^1$ . With conditional-law sampling probability at least  $1 - \alpha$ , the true table belongs to this set, and every candidate  $\tilde{T}$  satisfies*

$$D(T, \tilde{T}) \leq \Omega(2v_g, 2(u_n + b_j) + C_{\text{bin}}\delta) \quad (50.19)$$

*whenever the second argument is at most one. Here  $b_j$  accounts only for finite-flight error in the exact signed contact charts at the true offset  $d_0$ . It does not include timing or chart-calibration error. Programming  $j\hat{g}_e + d_0$  instead gives true offset  $d_0 + j(\hat{g}_e - g_e)$ ; hard cells additionally require a cell-edge crossing bound. These contributions are treated separately in Theorem 50.9. The position pilot of Theorem 56.2 controls  $j|\hat{g}_e - g_e|$  at this same final flight number.*

*If the gap event has failure probability  $\alpha_g$ , the joint failure bound is  $\alpha + \alpha_g$ ; no independence from that event is needed when the sampling bound is uniform. For a compact metrizable  $\mathcal{K}$  on which the finite observation vector is continuous, a measurable approximate minimum-excess-discrepancy selection can be used instead, with its positive tolerance added to each of the two arguments. Neither compactness nor selection is used to define the modulus (50.15).*

*Proof.* For a categorical law with vector  $p$ , the empirical vector satisfies

$$\mathbb{E}\|\hat{p} - p\|_1 \leq \sum_{k=1}^{K_\delta} \sqrt{p_k(1-p_k)/n} \leq \sqrt{K_\delta/n}.$$

Changing one mark changes this norm by at most  $2/n$ . Exposing the marks successively and applying the elementary bounded-interval exponential moment inequality to the resulting martingale differences gives

$$\Pr\{\|\hat{p} - p\|_1 > \sqrt{K_\delta/n} + t\} \leq e^{-nt^2/2}.$$

Here the moment inequality for a centered variable of range length  $c$  is  $\mathbb{E}e^{\lambda X} \leq e^{\lambda^2 c^2/8}$ ; it follows by convexity of the exponential and maximization over the endpoint Bernoulli mean. A union bound proves the assertion simultaneously for all  $\mathcal{L}$  laws. Total variation contracts under the grid map, and the  $\ell^1$  difference of probability vectors is twice their total variation. Thus (50.18) contributes at most  $b_j$ . The true limiting vectors are within  $u_n + b_j$  of their empirical vectors. Every other candidate in the confidence set is therefore within  $2(u_n + b_j)$  of the true limiting vectors, and its gaps are within  $2v_g$  of the true gaps. Apply Theorem 50.5.

For the selection assertion, use a fixed countable dense sequence of the compact metrizable candidate class. Define the nonnegative excess residual

$$R(T) = \max \left\{ 0, \max_e |g_e(T) - \hat{g}_e| - v_g, \max_{e,b} \|\mathbf{p}_\delta(P_{e,b}(T)) - \hat{p}_{e,b}\|_1 - (u_n + b_j) \right\}.$$

On the coverage event  $R(T_{\text{true}}) = 0$ . Continuity makes the infimum of  $R$  over the dense sequence zero as well. The first element whose residual is within a prescribed positive tolerance  $\epsilon_{\text{sel}}$  of the

countable infimum is measurable. Each of its individual gap and histogram constraints is violated by at most  $\epsilon_{\text{sel}}$ , proving the stated addition to each argument. On a compact holomorphic-support subfamily of the persistent design, the uniform limiting density construction gives the required total-variation continuity, including its moving boundary; cell probabilities and the gap coordinates are therefore continuous. This is an existence construction of a selector, not a claim that searching an arbitrary analytic model class is computationally efficient.  $\square$

At fixed  $d_0$ , the uniform relative-law theorem and the normalization argument in Corollary 23.3 give (50.18). Thus the finite-resolution prescription (50.17) permits a finite even  $j$ , a finite grid and a finite  $n$ . Taking  $j$  larger preserves the bias bound and allows the minimum bridge length to diverge along any sequence of target accuracies. Only a quantitative  $C^0$  forward bias bound is needed here; a flight-number estimate uniform over all derivative orders is not used.

There is also a direct charged version for an already calibrated design. Suppose independent preparations have success probability at least  $p_*(j) > 0$  for every selected law. Run at most

$$T_j = \left\lceil \frac{2n + 8 \log(\mathcal{L}/\alpha_c)}{p_*(j)} \right\rceil \quad (50.20)$$

preparations per law, retaining the first  $n$  successful grid categories when possible. The binomial lower-tail bound makes the probability of any cap failure at most  $\alpha_c$ : the expected success count is at least  $2n + 8 \log(\mathcal{L}/\alpha_c)$  and a failure lies below half this mean. The total charge is at most  $\mathcal{L}T_j$ . Apply the preceding concentration argument to the uncapped iid stream of successful categories, then add the cap-failure event. Conditioning on cap completion is neither needed nor asserted to preserve an iid law. On a uniform fixed-offset bridge class one may use its forward lower bound  $p_*(j) = c_* e^{-\Gamma_* j}$ .

The calibrated signed charts, selected labels and onset estimates remain part of the acquisition specification. The argument does not discover an unmarked design or estimate contact frames from these histograms alone. The charged position pilot of Part III is a separate way to acquire such calibration and retains its stated sensor. In contrast, the reconstruction transcript in this section contains only onset estimates and quantized transverse categories. The direct longitudinal graph-interpolation benchmark is not an available estimator on that transcript. This distinction separates a quantitative inverse for the coarsened record from the stronger position experiment, without changing either observation model.

**50.7. Offset and chart errors for hard endpoint categories.** We now allow the contact charts and the programmed onset to be estimated. This is an acquisition result for the same limiting data map, not a change of the inverse in Theorem 50.5. In particular, the exact-offset confidence statement above retains its hypotheses. The distinction between a bounded Lipschitz test and a hard category is relevant: small displacement of an endpoint can change a category, and the probability of that event must be bounded before concentration is applied.

Fix a closed interval  $I_d = [d_0 - a, d_0 + a]$ , where  $0 < a < d_0$ , inside the common positive-offset collar. Let  $U$  be a fixed bounded endpoint box containing the supports of the selected successful laws for offsets in  $I_d$ . The forward construction supplies uniform constants  $W_+, Z_- > 0$  such that, for each law,

$$f_d(x) = Z_d^{-1} w(x) (d - A(x))_+ \mathbf{1}_U(x), \quad 0 \leq w \leq W_+, \quad A \geq 0, \quad Z_d \geq Z_-. \quad (50.21)$$

Here  $x = (u, v)$ ,  $A(x) = S_b(u) + S_b(v)$  and  $w(x) = B_b(u)B_b(v)$ . The bounds are forward bounds on the same quantified class, not measurements of  $A, w$  or  $Z_d$ . For example, on  $Q_1$ ,  $f_{d_0}(0, 0) \leq H_1$  and  $B_b(0) = 1$  give a positive lower bound for  $Z_{d_0}$ ; a smaller positive-offset interval and the uniform forward construction give  $Z_-$ . Alternatively these constants can be read directly from

the positive-offset normalization estimates in Section 56. We use

$$C_{\text{off}} = \frac{W_+|U|}{Z_-}, \quad H_* = \frac{W_+(d_0 + a)}{Z_-}.$$

Both constants can be enlarged uniformly over the finite design. In particular the upper bound  $H_*$  holds for the density extended by zero, including neighborhoods of the outer boundary of  $Q_1$ ; an interior  $C^1$  bound alone is not used to control that exterior neighborhood.

**Lemma 50.7** (Uniform offset continuity). *For  $d, d' \in I_d$ , the limiting endpoint laws satisfy*

$$\|P_d - P_{d'}\|_{\text{TV}} \leq C_{\text{off}}|d - d'|, \quad 0 \leq f_d \leq H_*. \quad (50.22)$$

*The same conclusions hold simultaneously for both types of every selected channel, with the uniform constants above.*

*Proof.* Put  $F_d = w(d - A)_+ \mathbf{1}_U$ . The positive-part function is 1-Lipschitz, so

$$\|F_d - F_{d'}\|_1 \leq W_+|U||d - d'|, \quad |Z_d - Z_{d'}| \leq \|F_d - F_{d'}\|_1.$$

Using  $Z_d, Z_{d'} \geq Z_-$  gives  $\|F_d/Z_d - F_{d'}/Z_{d'}\|_1 \leq 2\|F_d - F_{d'}\|_1/Z_-$ . Total variation is one half of this  $L^1$  norm. Finally  $A \geq 0$  and  $d \leq d_0 + a$  give the density bound. The moving zero set requires no differentiation.  $\square$

Let  $c_\delta : \mathbb{R}^2 \rightarrow \{1, \dots, K_\delta\}$  be the grid map from Section 50, with its complement category. Let  $\Gamma_\delta$  be the union of all cell edges, including the four outer edges of  $Q_1$ . Write

$$V_\delta(r) = |\{x : \text{dist}_\infty(x, \Gamma_\delta) \leq r\}|, \quad r \geq 0. \quad (50.23)$$

This is the area of a finite union of rectangles. It can be computed from the grid, and tends to zero with  $r$  for every fixed grid. For a  $q$ -by- $q$  square grid of side  $\delta = 2r_1/q$ , the union bound gives

$$V_\delta(r) \leq 8r(r_1 + r)(q + 1). \quad (50.24)$$

Indeed there are  $2(q + 1)$  full horizontal or vertical segments of length  $2r_1$ , each with an  $\ell^\infty$  tube of area  $2r(2r_1 + 2r)$ . Their intersections only decrease the union area. This formula includes corners and crossings between  $Q_1$  and its complement.

**Lemma 50.8** (Hard-category perturbations). *Let  $P$  have a density bounded by  $H_*$  on  $\mathbb{R}^2$ , and let  $\Phi$  be a Borel map satisfying  $|\Phi(x) - x|_\infty \leq r$  on its support. Then*

$$\|(c_\delta \circ \Phi)_\# P - (c_\delta)_\# P\|_1 \leq 2 \min\{1, H_* V_\delta(r)\}. \quad (50.25)$$

*Here and below  $\ell^1$  denotes the norm of the finite probability vector. No differentiability or invertibility of  $\Phi$  is required.*

*Proof.* Couple the two categories using the same  $x$ . If they differ, the segment joining  $x$  to  $\Phi(x)$  meets a cell edge; hence  $\text{dist}_\infty(x, \Gamma_\delta) \leq r$ . The mismatch probability is therefore at most  $H_* V_\delta(r)$  and at most one. Under any coupling, the total variation of two finite laws is at most their mismatch probability. Multiplication by two proves the claim. Points exactly on an original edge have zero  $P$ -probability. The image map may put mass on an edge; the half-open convention fixes its category and the same segment argument applies.  $\square$

**50.8. A calibration-aware confidence radius.** Choose the final even flight number  $j$  before taking a pilot. Conditional on its history  $H$ , the post-pilot designs and estimated contact projections are fixed. Suppose an event  $G$ , of probability at least  $1 - \alpha_{\text{cal}}$ , gives correct signed design labels and the bounds

$$\max_e |\hat{g}_e - g_e| \leq v_g, \quad \max_{e,b} \|\hat{x}_{e,b} - x_{e,b}\|_\infty \leq r. \quad (50.26)$$

The second bound is uniform on the common physical endpoint collar. The estimated projections are observable maps of physical positions; the true projections occur only in the analysis. For a fixed true table, a boundary point is determined by its transverse coordinate on this collar, so these maps induce maps  $\Phi_{T,H}$  to which Lemma 50.8 applies. No unknown chart is applied by the statistician.

Program each group at  $t = j\hat{g}_e + d_0 + \omega_{e,b}$ , with  $|\omega_{e,b}| \leq v_t$ . Its true offset is

$$d_{e,b}(H) = d_0 + j(\hat{g}_e - g_e) + \omega_{e,b}, \quad |d_{e,b}(H) - d_0| \leq \Delta := jv_g + v_t. \quad (50.27)$$

Assume  $\Delta \leq a$ . On  $I_d$  use the uniform finite-flight estimate

$$\sup_{e,b,d \in I_d} \|P_{j,e,b,d} - P_{e,b,d}\|_{\text{TV}} \leq C_0 \tau^j. \quad (50.28)$$

This is the positive-offset, full-box normalization consequence of the relative-law estimate, with the same argument as Corollary 23.3. It is uniform in  $d$  on  $I_d$ ; a bound at  $d_0$  alone would not suffice for the random offsets.

**Theorem 50.9** (Histogram reconstruction with estimated calibration). *Under the preceding hypotheses set*

$$b_{j,\delta}^{\text{cal}} = 2C_0 \tau^j + 2C_{\text{off}}(jv_g + v_t) + 2H_* V_\delta(r). \quad (50.29)$$

*For every good pilot history, the probability vector of a fresh successful observed category differs from the true limiting vector at  $d_0$  by at most  $b_{j,\delta}^{\text{cal}}$  in  $\ell^1$ . Retain  $n$  fresh successful categories in each of the  $\mathcal{L}$  groups, and use  $u_n$  from Corollary 50.6. Replace  $b_j$  there by  $b_{j,\delta}^{\text{cal}}$ . With probability at least  $1 - \alpha_{\text{cal}} - \alpha$ , the resulting confidence set contains the true table, and every candidate satisfies*

$$D(T, \tilde{T}) \leq \Omega \left( 2v_g, 2(u_n + b_{j,\delta}^{\text{cal}}) + C_{\text{bin}} \delta \right), \quad (50.30)$$

*provided the second argument is at most one. The same bound controls the marked lattice Gram matrix with the enlargement already specified for  $C_{\text{reg}}$ . On a compact metrizable candidate class with the stated continuity, an approximate measurable selector adds its positive tolerance to each argument.*

*If every good history has preparation success probability at least  $p_*(j) > 0$  in each group, impose the cap (50.20). The failure bound becomes  $\alpha_{\text{cal}} + \alpha + \alpha_c$  and the post-pilot preparation charge is at most  $\mathcal{L}T_j$ . Add the deterministic pilot cap to obtain the total charge. Independence between the pilot events is not needed.*

*Proof.* Fix  $T$  and a good history  $H$ . In the true endpoint coordinates compare successively the finite law at  $d_{e,b}(H)$ , the limiting law at the same offset, and the limiting law at  $d_0$ . Apply the same estimated projection and grid map to each. Data processing, (50.28), and Lemma 50.7 bound the first two vector differences by  $2C_0 \tau^j$  and  $2C_{\text{off}} \Delta$ . Compare the last vector with that formed in the true chart by Lemma 50.8. Its cost is at most  $2H_* V_\delta(r)$ . This comparison is on a fixed true boundary; it does not assert total-variation continuity of ambient position laws on different curves. Summation proves (50.29).

Conditional on the pilot history, the uncapped successful sequences are iid with their actual category laws. The categorical concentration proof of Corollary 50.6 gives a simultaneous event of conditional probability at least  $1 - \alpha$ . On that event the true limiting vectors are within  $u_n + b_{j,\delta}^{\text{cal}}$

of the empirical ones. The gap bound and two triangle inequalities now give the arguments of (50.30). Apply Theorem 50.5. The excess-residual selection argument in the preceding corollary applies without alteration. Integrating the uniform conditional bound over good histories adds at most  $\alpha_{\text{cal}}$ .

For the charged version, couple each capped experiment with its uncapped successful stream. The binomial lower-tail bound used for (50.20) holds conditional on every good history, uniformly with  $p_*(j)$ . Add its union failure probability  $\alpha_c$  before integrating over the pilot. No concentration assertion is made after conditioning on cap completion, and no failure is removed from the preparation count.  $\square$

**50.9. A finite joint choice of grid, pilot, and sample size.** The preceding theorem applies to the existing position pilot without changing its observation space. We make the joint choices explicit to show that refinement of a hard grid is compatible with calibration.

**Corollary 50.10** (Charged acquisition of finite-resolution transverse data). *Let the quantified class be a compact marked analytic class with the persistent position-sensor design of Part III. For every  $\varepsilon > 0$ ,  $\alpha, \alpha_c, \beta \in (0, 1)$  with  $\alpha + \alpha_c + \beta < 1$ , and minimum flight number  $J_0$ , there is a deterministically capped policy whose pilot and subsequent preparations all use one even  $J \geq J_0$ . Its table estimator satisfies*

$$\sup_{T \in \mathcal{K}} \Pr_T \{D(T, \hat{T}) > \varepsilon\} \leq \beta + \alpha + \alpha_c. \quad (50.31)$$

*The reconstruction uses the pilot gap estimates and post-pilot transverse grid categories. Longitudinal positions are used to perform the physical pilot and the observable projection, but are not inputs to the subsequent law-to-table reconstruction. The supplied marks do not contain Euclidean deck vectors or inter-channel poses. The estimate includes the unknown marked Euclidean lattice and its Gram matrix.*

*Proof.* First choose finite tolerances  $e_* > 0$ ,  $0 < \eta_* \leq 1$  by (50.17), so that  $\Omega(e_*, \eta_*) \leq \varepsilon$ . Let

$$q = \max \{1, \lceil 8C_{\text{bin}} r_1 / \eta_* \rceil\}, \quad \delta = 2r_1/q, \quad K_\delta = q^2 + 1. \quad (50.32)$$

Then  $C_{\text{bin}}\delta \leq \eta_*/4$ . Choose a finite even  $J \geq \max\{2, J_0\}$ , above the uniform finite-flight threshold, such that  $2C_0\tau^J \leq \eta_*/32$ . This is the final flight number and is not increased after the pilot.

The calibration theorem gives a constant  $C_\chi \geq 1$  for which a pilot of resolution  $h$  at this same  $J$  has, except with probability  $\beta$ ,

$$J \max_e |\hat{g}_e - g_e| \leq h, \quad \max_{e,b} \|\hat{x}_{e,b} - x_{e,b}\|_\infty \leq C_\chi \sqrt{h}.$$

Choose its allowed upper resolution  $h_0 > 0$  small enough to include all localization and nonzero-denominator requirements of that theorem. Any positive  $h$  at most

$$\min \left\{ \frac{h_0}{2}, \frac{a}{2}, \frac{J e_*}{4}, \frac{\eta_*}{64 C_{\text{off}}}, \left( \frac{r_1}{C_\chi} \right)^2, \left( \frac{\eta_*}{1024 H_* r_1 (q+1) C_\chi} \right)^2 \right\} \quad (50.33)$$

is suitable. Use the exact-clock convention of the physical experiment, so  $v_t = 0$ . Set  $v_g = h/J$  and  $r = C_\chi \sqrt{h}$ . The offset stays in  $I_d$ , and (50.24) with  $r \leq r_1$  gives

$$2v_g \leq e_*/2, \quad 2C_{\text{off}}h \leq \eta_*/32, \quad 2H_* V_\delta(r) \leq 32H_* r r_1 (q+1) \leq \eta_*/32.$$

Thus  $b_{J,\delta}^{\text{cal}} \leq 3\eta_*/32$ . Choose

$$n \geq \left\lceil \frac{64}{\eta_*^2} \left( \sqrt{K_\delta} + \sqrt{2 \log(\mathcal{L}/\alpha)} \right)^2 \right\rceil, \quad 0 < \epsilon_{\text{sel}} \leq \min\{e_*/2, \eta_*/4\}. \quad (50.34)$$



Then  $u_n \leq \eta_*/8$ , and

$$2v_g + \epsilon_{\text{sel}} \leq e_*, \quad 2(u_n + b_{J,\delta}^{\text{cal}}) + C_{\text{bin}}\delta + \epsilon_{\text{sel}} \leq 15\eta_*/16 < \eta_*.$$

The selector in Theorem 50.9 therefore has error at most  $\varepsilon$  on its coverage event. Its monotone modulus uses only nondecreasing finite-order expressions and their infimum.

For the pilot use (56.4)–(56.5) with  $(j, h, \beta) = (J, h, \beta)$ . For the subsequent groups use (50.20); the uniform positive-offset success bound has the form  $p_*(J) = c_*e^{-\Gamma_*J}$ . The total deterministic cap is  $B_{\text{pil}} + \mathcal{L}T_J$ . Each quantity is finite. Every failure, including every failed near-onset pilot preparation, is charged. Apply the preceding theorem with  $\alpha_{\text{cal}} = \beta$ . The specified choices also permit  $J \geq s$  along any sequence of target accuracies indexed by  $s$ , so that the minimum flight number in each newly acquired transcript diverges.  $\square$

For a regular square grid the added chart cost is bounded by a constant times  $r/\delta$  when  $r \leq r_1$ . The pilot gives  $r \leq C_\chi\sqrt{h}$ ; thus its precision is chosen after the grid, rather than assuming that  $h \rightarrow 0$  alone suffices for a simultaneously refined grid. The three distinct errors in (50.29) are the finite-flight error, the amplified onset or clock error, and the cell-edge crossing error. At exact calibration they reduce to the original  $b_j$ . No assertion of a sharp information threshold or an optimal preparation rate is needed for these finite-resolution bounds.

## 51. PHYSICAL COORDINATES AND FINITE-WINDOW OBSERVATION

The block inverse is a statement about actual boundary laws. We now show that its two sets of coordinates vary independently in physical tables while the complete selected leading hierarchy remains fixed. This also gives physical alternatives for a finite-dimensional observation theorem. The dimension is fixed throughout this section; none of its statistical bounds is a lower bound for the full infinite-dimensional profile experiment.

### 51.1. Independent realization at fixed leading data.

**Theorem 51.1** (A physical open image of two-contact jets). *Fix  $M \geq 2$ . For any two curvature radii  $R_0, R_1$  sufficiently close to one, there is a real-analytic family of positively curved periodic billiards on  $3\mathbb{Z} \times 4\mathbb{Z}$ , with  $2(M-1)$  parameters, for which the selected horizontal channel has*

$$g = 1, \quad A = 12 - \pi, \quad \kappa_b = R_b^{-1}.$$

*Both contact graphs are even, but their higher jets are independent. The parameters can be taken to be either*

$$(q_{b,2m})_{b=0,1; 2 \leq m \leq M} \quad \text{or} \quad (f_{b,k})_{b=0,1; 1 \leq k \leq M-1}.$$

*Each of these maps is a local analytic diffeomorphism, with inverse derivative bounded on a sufficiently small compact neighborhood. Every member has the same leading selected count and endpoint-metric coefficients at every collision order, in both orientations. The statement includes  $R_0 = R_1$  and requires no equality of higher jets.*

*Proof.* Use one obstacle per cell and its support function. Put

$$w_0(\theta) = \frac{1 + \cos \theta}{2}, \quad w_1(\theta) = \frac{1 - \cos \theta}{2},$$

and consider the finite-dimensional analytic family

$$\begin{aligned} h(\theta) &= 1 + \alpha \sin^2 \theta + \beta \cos \theta \sin^2 \theta \\ &+ \sum_{m=2}^M \{s_{0m}w_0(\theta) + s_{1m}w_1(\theta)\} \sin^{2m} \theta + z \sin^{2M+2} \theta. \end{aligned} \quad (51.1)$$

At the two horizontal normals the support values equal one, the support derivatives vanish, and the curvature radii are

$$h(0) + h''(0) = 1 + 2\alpha + 2\beta, \quad h(\pi) + h''(\pi) = 1 + 2\alpha - 2\beta.$$

Thus set  $\alpha = (R_0 + R_1 - 2)/4$  and  $\beta = (R_0 - R_1)/4$ . No  $s$  or  $z$  changes these data. The obstacle area is

$$\mathcal{S}(h) = \frac{1}{2} \int_0^{2\pi} (h^2 - (h')^2) d\theta.$$

At the unit disk,  $\partial_z \mathcal{S} = \int_0^{2\pi} \sin^{2M+2} \theta d\theta > 0$ . The analytic implicit-function theorem therefore gives a unique small  $z = z(\alpha, \beta, s)$  for which  $\mathcal{S} = \pi$ . This function is analytic jointly in all the displayed parameters.

For small parameters,  $h + h''$  is positive, so the support function defines a strictly convex analytic closed curve. Its right and left horizontal support points remain  $(1, 0)$  and  $(-1, 0)$ . The selected chord from this obstacle to its translate by  $(3, 0)$  has length one and normal incidence. The unit-disk configuration has positive clearance and a strict gap between its horizontal nearest neighbors and all other translates; these properties persist on a small neighborhood. For the infinitely many distant translates this follows from a uniform bound on the obstacle diameter, leaving only finitely many translates to check. Thus the constructed chords are physical first hits and form the selected ground channel with its reverse. The free cell area is  $12 - \pi$ . The function  $h$  is even in  $\theta$ , so the curve is invariant under reflection across the horizontal axis. Each horizontal contact graph is consequently even in its transverse coordinate. Reflection across the vertical axis is not imposed.

At normal zero the  $s_{0m}$  perturbation starts in order  $2m$ , whereas its  $s_{1m}$  counterpart starts in order  $2m + 2$ . At normal  $\pi$  these roles are interchanged. The area compensator starts in order  $2M + 2$  at both normals. For completeness, the support-envelope calculation in a contact frame gives

$$\frac{\partial q_{b,2m}}{\partial s_{bm}} = -(2m)! \kappa_b^{2m}, \quad \frac{\partial q_{b,2m}}{\partial s_{1-b,m}} = 0. \quad (51.2)$$

Indeed, at the right contact write the boundary as  $x = x(y)$ . Stationarity of  $(h(\theta) - y \sin \theta) / \cos \theta$  at its supporting normal gives, at fixed  $y$ ,  $\partial_s x = \partial_s h(\theta(y)) / \cos \theta(y)$ . Since the radius is  $R_b$ ,  $\theta(y) = \kappa_b y + O(y^3)$ . The contact graph is  $\psi_b = 1 - x$ , giving the negative coefficient in (51.2). At the left contact use the outward horizontal normal and the reflected longitudinal coordinate; the same calculation applies. Even powers make the transverse orientation immaterial.

A perturbation of order  $2m$  cannot change graph jets of smaller order. Derivatives of the compensator also have order  $2M + 2$ . Hence the map from the  $2(M - 1)$  variables  $s$  to the displayed graph jets has a block lower triangular Jacobian, with the invertible blocks (51.2). It is an analytic local diffeomorphism. Compose it with Theorem 21.1 to obtain the asserted statistical coordinates. Compactly restricting the coordinate neighborhood bounds both inverse derivatives. Finally, the quadratic endpoint Hessian and its normalized mixed derivative depend only on  $g, \kappa_0, \kappa_1$ . The full-phase leading count coefficient also uses only  $A$ . These four data are fixed. The endpoint-metric leading coefficients and the selected count leading coefficients are therefore fixed at every order, by (15.1)–(15.2) and Proposition A.19.1. This assertion concerns this channel and its reverse, not the leading data of every other channel.  $\square$

**Corollary 51.2** (Finite-order physical fibres). *For  $N > M \geq 2$ , use the family of Theorem 51.1 with  $N$  in place of  $M$ . In a sufficiently small coordinate neighborhood, each nonempty fibre of the two-law jet map through order  $M - 1$  is an analytic submanifold of dimension  $2(N - M)$ . The full two-law jet map through order  $N - 1$  is locally injective. All these statements hold with the selected leading hierarchy fixed.*

*Proof.* In the statistical coordinates furnished by the theorem the first map is the coordinate projection onto its first  $2(M-1)$  coordinates. Its fibres have the stated dimension; the full map is the coordinate chart itself. These are fibres of specified finite physical families, not a classification of the fibres of arbitrary smooth asymmetric contact geometry.  $\square$

**51.2. A fixed physical experiment.** Fix  $M$  and one compact coordinate neighborhood from the preceding theorem. Write  $n = M - 1$  and  $\vartheta = (f_{b,k})_{b=0,1; 1 \leq k \leq n} \in \mathbb{R}^{2n}$ . The exact family, its labels, and  $g, A, \kappa_0, \kappa_1$  are supplied in the following experiment. A preparation returns the indicator of one selected alternating event, at a specified orientation and time. It does not return an endpoint, a derivative, or a limiting law.

**Theorem 51.3** (Finite-window observation and physical separation). *After shrinking the parameter neighborhood, there are  $h_0 > 0$  and one even integer  $j_0 \geq 2$  such that the  $2n$  physical means*

$$\mathcal{P}(\vartheta) = \left( \Pr_{\vartheta} E_{j_0, b}(\ell h_0) \right)_{b=0,1; 1 \leq \ell \leq n} \quad (51.3)$$

*form a bi-Lipschitz coordinate on a compact parameter ball  $\Theta$  with nonempty interior. Here the physical observation time is  $j_0 g + \ell h_0$ . The constants and the windows depend on the fixed family and  $M$ , but not on the requested accuracy.*

*For  $0 < \varepsilon < \varepsilon_0$  and  $0 < \eta < 1/4$ , a Borel estimator based on independent binary preparations in these windows satisfies*

$$\sup_{\vartheta \in \Theta} \Pr_{\vartheta} \{ \|\hat{\vartheta} - \vartheta\| > \varepsilon \} \leq \eta, \quad N_{\text{tot}} \leq C\varepsilon^{-2} \log(C/\eta).$$

*The count is deterministic and includes all failures. The same order holds for the corresponding finite graph-jet vector.*

*Moreover, for all sufficiently large  $N$ , the minimax squared risk in this fixed finite-window experiment satisfies*

$$\frac{c}{N} \leq \inf_{\hat{\vartheta}, \mathcal{D}_N} \sup_{\vartheta \in \Theta} \mathbb{E}_{\vartheta} \|\hat{\vartheta} - \vartheta\|^2 \leq \frac{C}{N}. \quad (51.4)$$

*The infimum allows adaptive and randomized choices of these windows and at most  $N$  fresh preparations. Its lower bound uses pairs of physical tables in the stated family, with exactly the same selected leading hierarchy.*

*Proof.* Analyticity of the boundary law in the physical parameters and offset, on a common neighborhood, gives

$$\mathcal{F}_{bb}(d; \vartheta) = 1 + \sum_{k=1}^n \vartheta_{b,k} d^k + d^{n+1} R_b(d; \vartheta),$$

where  $R_b$  and its first parameter derivatives are uniformly bounded. At  $d = \ell h$  the Jacobian with respect to  $\vartheta$  is the block diagonal matrix with blocks  $V_h = ((\ell h)^k)_{1 \leq \ell, k \leq n}$ , plus a matrix of norm  $O(h^{n+1})$ . The matrix  $V_h = V_1 \text{diag}(h, \dots, h^n)$  is invertible: its determinant is a nonzero Vandermonde determinant multiplied by  $\prod_{\ell=1}^n \ell$ . Also  $\|V_h^{-1}\| \leq Ch^{-n}$  for  $h \leq 1$ . Choose a fixed  $h = h_0$  so small that  $nh_0$  lies in the common collar and the transformed remainder norm is less than  $1/4$ .

The differentiated relative theorem bounds the difference between  $F_{j,b}$  and  $\mathcal{F}_{bb}$ , including its first parameter derivatives, by  $C\tau^j$ . Choose a fixed even  $j = j_0$  so large that  $\|V_{h_0}^{-1}\|C\tau^{j_0} < 1/4$  (with the fixed block norms included in  $C$ ). The Jacobian of the actual normalized finite law is then invertible. Passing to the physical means multiplies its rows by the positive constants  $(\ell h_0)^2 / (2A \sinh(j_0 \gamma))$ . These constants do not depend on  $\vartheta$ , since all leading geometric data in this family are fixed. Thus the physical Jacobian is invertible as well.

Choose a small convex ball about an interior reference parameter. The Jacobian, after multiplication by its inverse at the center, is uniformly within  $1/2$  of the identity on the ball. Integrating along a line segment shows that the physical mean map is bi-Lipschitz there. Take its compact closure as  $\Theta$ . By further restriction of the offset collar if needed, each of its finitely many probabilities is strictly between zero and one; compactness gives common positive margins. These margins may be very small, but are fixed constants of this experiment.

Take the same number  $L$  of fresh preparations at each of the  $2n$  windows, and let  $Y$  be the vector of sample means. For any  $t > 0$ , Hoeffding's elementary bounded-summand bound and a union bound give

$$\Pr\{\|Y - \mathcal{P}(\vartheta)\|_\infty > t\} \leq 4n \exp(-2Lt^2).$$

Choose a finite ordered net of  $\Theta$  fine enough that its image approximates  $\mathcal{P}(\Theta)$  within  $t$  in Euclidean norm. Minimize the Euclidean discrepancy from  $Y$ , resolving ties by the first index. This is a Borel estimator. On the displayed event its image error is at most  $C_n t$ ; the inverse Lipschitz bound gives parameter error at most  $Ct$ . Set  $t = c\varepsilon$  and  $L = \lceil Ct^{-2} \log(4n/\eta) \rceil$ . This proves the deterministic charge and confidence statement. The graph-jet inverse is Lipschitz on this neighborhood, so the assertion transfers to that vector. For squared risk, use  $L = \lfloor N/(2n) \rfloor$  and a net of mesh  $N^{-1/2}$ . The same deterministic discrepancy inequality and  $\mathbb{E}\|Y - \mathcal{P}(\vartheta)\|^2 \leq C/L$  give the upper bound in (51.4).

For the lower bound choose an interior point  $\vartheta_*$ , a unit vector  $v$ , and the two parameters  $\vartheta_\pm = \vartheta_* \pm aN^{-1/2}v$ . They represent actual analytic obstacles from Theorem 51.1. Every allowed window has a probability in a fixed interval  $[p_*, 1 - p_*]$ . The Bernoulli relative entropy satisfies

$$D(\text{Ber}(p) \parallel \text{Ber}(q)) \leq \frac{(p - q)^2}{q(1 - q)}.$$

The smooth mean map therefore bounds this entropy by  $Ca^2/N$  for each window. In an adaptive experiment, condition on the complete past and any auxiliary random seed. The rule for choosing the next window is the same under the two parameters; the conditional entropy is bounded by the same uniform bound. The chain rule, with padding by parameter-independent observations after an early stop, gives entropy at most  $Ca^2$  for the entire record. Choose  $a > 0$  fixed and small so that the total variation is at most  $1/2$ , using  $\text{TV}^2 \leq D/2$ .

Any estimator induces the test selecting the nearer of the two parameters. On a test error its squared parameter loss is at least  $a^2/N$ . The sum of the two testing error probabilities is at least  $1 - \text{TV} \geq 1/2$ . The maximum risk is therefore at least  $a^2/(4N)$ . This proves the lower bound, including adaptive window choices. All the alternatives keep  $g, A, \kappa_0, \kappa_1$  and the selected leading hierarchy exactly fixed.  $\square$

The fixed-window theorem resolves a physical finite-jet inverse problem, not just a formal coefficient problem: a fixed finite record length and finitely many strictly positive offsets already give local coordinates on independently varying contact jets. The windows and conditioning constants can deteriorate with  $M$ . The theorem neither passes this bound to an increasing-dimensional model nor substitutes its supplied exact family for the unknown smooth remainder experiment of Section 49.

## 52. TWO-FLIGHT DETERMINATION AND THE FINITE OBSERVATION DESIGN

The geometric inverse admits a finite-flight counterpart. This result separates recovery of a fixed contact jet from the relative limit that constructs the general boundary energy invariant. Throughout this section the two contact graphs are individually even, their labels are retained, and  $g, \kappa_0, \kappa_1$  are supplied. Their higher jets need not agree. The finite-flight calculation and

its observation consequence were pointed out in the independent memorandum [23]. We give a direct proof from the stationary action and the physical phase measure.

For  $b \in \{0, 1\}$  define

$$G_b(d) = \frac{2A \sinh(2\gamma)}{d^2} \Pr(E_{2,b}(d)), \quad \xi_{b,k} = [d^k]G_b(d), \quad G_b(0) = 1. \quad (52.1)$$

The input to the exact germ inverse is the normalized germ  $G_b$ . If instead the full unnormalized probability germ  $p_b(d)$  is supplied, its leading coefficient and its normalization are

$$\alpha_2 = \lim_{d \downarrow 0} \frac{p_b(d)}{d^2} = \frac{1}{2A \sinh(2\gamma)}, \quad G_b(d) = \frac{p_b(d)}{\alpha_2 d^2}.$$

This observation concerns exact germs, not a finite-sample estimate of a limit at zero. The finite-family experiment below supplies the fixed free area in its family specification; the general unknown-normalization experiment uses the separately charged smooth-class pilot. No derivative or leading coefficient is supplied as a free statistical observation.

Here two flights means three successive impacts, with itinerary  $b, 1-b, b$ ; the probability also integrates the initial residual time. Put

$$z = c_0 c_1 - 1 = \sinh^2 \gamma > 0, \quad L_b = \frac{g}{2c_b z}, \quad k_m = \frac{4}{2^m(m+1)(m!)^2}, \quad (52.2)$$

$$U_m = (1 + 2z)^m, \quad V_m = 1 + 2mz.$$

The letters  $L_b$  in this section denote inverse-Hessian scales, not flight lengths. As before,  $q_m = (q_{0,2m}, q_{1,2m})^t$ , and write  $\xi_{m-1} = (\xi_{0,m-1}, \xi_{1,m-1})^t$ .

**Theorem 52.1** (Two-flight inverse for independent even contacts). *For every  $m \geq 2$ ,  $\xi_{m-1}$  depends only on the leading geometry and  $q_2, \dots, q_m$ , is affine in  $q_m$ , and has last-jet derivative*

$$D_{q_m} \xi_{m-1} = - \begin{pmatrix} U_m & V_m \\ V_m & U_m \end{pmatrix} \text{diag}(k_m L_0^m, k_m L_1^m). \quad (52.3)$$

The symmetric factor is positive definite. In particular,

$$U_m - V_m = \sum_{r=2}^m \binom{m}{r} (2z)^r > 0. \quad (52.4)$$

For each fixed  $M \geq 2$ , the map

$$(q_2, \dots, q_M) \mapsto (\xi_1, \dots, \xi_{M-1})$$

has a block-triangular analytic inverse on its image. The inverse is locally Lipschitz on compact finite-jet sets and compact positive leading-geometry boxes, including coincident curvatures. Equality of the two oriented normalized two-flight germs for real-analytic even contacts at the same labelled leading geometry determines both graph germs. For connected analytic closed boundaries it determines the two participating boundary images in their contact frames.

*Proof.* Fix  $b$  and put  $o = 1 - b$ . Let  $u, v$  be the endpoints and let  $w$  be the intermediate impact coordinate. The quadratic part of the broken action, after subtracting  $2g$ , is

$$\frac{c_b u^2 + 2c_o w^2 + c_b v^2 - 2w(u + v)}{2g}.$$

Its stationary intermediate point is  $w_0 = (u + v)/(2c_o)$ . The finite implicit-function theorem, or Lemma 15.2, gives a unique smooth actual point  $w = w_0 + O((|u| + |v|)^3)$ . Evenness is used for this cubic order. Eliminating  $w_0$  gives  $E_0(y) = \frac{1}{2} y^t H_b y$ , where  $y = (u, v)^t$  and

$$H_b = \frac{1}{g} \begin{pmatrix} c_b - (2c_o)^{-1} & -(2c_o)^{-1} \\ -(2c_o)^{-1} & c_b - (2c_o)^{-1} \end{pmatrix}, \quad \det H_b = \frac{c_b z}{c_o g^2} = a_b^2. \quad (52.5)$$

For a linear form  $\ell \cdot y$  set  $\nu(\ell) = \ell^t H_b^{-1} \ell$ . Inversion of this two-by-two matrix yields

$$\nu(u) = \nu(v) = L_b(1 + 2z), \quad \nu(w_0) = L_o. \quad (52.6)$$

The notation  $\nu(u)$  denotes the value for the coordinate linear form  $y \mapsto u$ , and similarly for  $v, w_0$ .

Write  $E_2 = W_{2,b} - 2g$  and  $b_2 = -W_{2,b,w_0}/d_2^0$ . The reference mixed derivative is  $d_2^0 = (2gc_o)^{-1} = a_b/\sinh(2\gamma)$ . The exact full-phase formula (16.1), including the residual-time integration, gives

$$G_b(d) = \frac{a_b}{\pi d^2} \int (d - E_2(y))_+ b_2(y) \, dy, \quad I_0(d) := \int (d - E_0(y))_+ \, dy = \frac{\pi d^2}{a_b}. \quad (52.7)$$

For a sufficiently small offset collar the whole active sublevel is inside the physical first-hit charts. Thus this identity has no unobserved endpoint restriction or substituted ensemble.

Keep all lower graph jets fixed. By stationarity, the first variation of the reduced action is obtained by varying the broken action at its stationary intermediate point. The first varying homogeneous length term is a pure endpoint power divided by  $(2m)!$ . The boundary sites are counted once and the interior site twice. Consequently

$$\begin{aligned} \partial_{q_b, 2m} E_2 &= \frac{u^{2m} + v^{2m}}{(2m)!} + O(|y|^{2m+2}), \\ \partial_{q_o, 2m} E_2 &= \frac{2w_0^{2m}}{(2m)!} + O(|y|^{2m+2}). \end{aligned} \quad (52.8)$$

All error estimates here may be differentiated to the required finite orders. Differentiate these identities in both endpoints and divide by  $-d_2^0$ . Since  $d_2^0$  uses only the fixed leading geometry,

$$\begin{aligned} \partial_{q_b, 2m} b_2 &= O(|y|^{2m}), \\ \partial_{q_o, 2m} b_2 &= -\frac{g}{c_o} \frac{w_0^{2m-2}}{(2m-2)!} + O(|y|^{2m}). \end{aligned} \quad (52.9)$$

In particular, omitting the normalized twist variation would omit an essential contribution from the other contact.

We record both moment identities needed for the integration. The change  $y = H_b^{-1/2} x$ , followed by a rotation, sends  $\ell \cdot y$  to  $\sqrt{\nu(\ell)} x_1$ . The angular average of  $\cos^{2k} \theta$  is  $\binom{2k}{k}/4^k$ . Radial integration gives

$$\frac{1}{I_0(d)} \int_{E_0 < d} (\ell \cdot y)^{2m} \, dy = \frac{2(2m)! \nu(\ell)^m}{2^m (m!)^2 (m+1)} d^{m-1}, \quad (52.10)$$

$$\frac{1}{I_0(d)} \int (d - E_0)_+ (\ell \cdot y)^{2k} \, dy = \frac{2 \binom{2k}{k} \nu(\ell)^k}{2^k (k+1)(k+2)} d^k. \quad (52.11)$$

For example, the respective radial integrals before normalization are  $(2d)^{m+1}/(2m+2)$  and  $2^k d^{k+2}/((k+1)(k+2))$ . The residual weight vanishes on the sublevel boundary, so the first action variation contributes the negative unweighted integral, with no additional moving-boundary term.

There is no replacement of a nonlinear domain by a quadratic domain at an uncontrolled order. Indeed, the smooth Morse construction of Lemma 16.1 fixes the integration domain. At scale  $y = \sqrt{d} Y$ , the first action variation has order  $d^{m-1}$  in the normalized law, and the first amplitude variation has the same order. The unvaried nonlinear action and amplitude start two endpoint degrees above their quadratic and constant parts. Every product with either first variation therefore has at least one additional power of  $d$ . Taylor's formula in the common Morse coordinates justifies this degree calculation with smooth remainders at  $d = 0$ .

By (52.6) and (52.10), the own-contact action contribution is  $-k_m L_b^m (1 + 2z)^m$ , and the other-contact action contribution is  $-k_m L_o^m$ . There is no own-contact twist contribution at this order. Substituting  $k = m - 1$  in (52.11), the other-contact twist contribution is

$$-\frac{g}{c_o} \frac{4L_o^{m-1}}{2^m(m+1)m!(m-1)!} = -k_m L_o^m (2mz),$$

since  $g/(c_o L_o) = 2z$ . These are precisely the first row of (52.3). Reversing the labels proves the other.

For completeness, finite-jet dependence does not require analyticity of the original graphs. Differentiate the finite stationary equation at zero. At each degree the highest derivative of  $w$  is multiplied by the same positive scalar  $2c_o/g$ ; the remaining terms involve lower derivatives and only the graph jets through that degree. The normalized law coefficient in question uses the reduced action through degree  $2m$  and its twist through degree  $2m - 2$ . Hence it uses no graph derivative above  $2m$ . The displayed last-jet derivative is independent of that jet and of all lower jets. This proves affinity, also away from the quadratic graph. The finite operations are analytic in their finite-jet arguments and in positive  $g, \kappa_0, \kappa_1$ .

Let  $R_m$  be the value of  $\xi_{m-1}$  with  $q_m = 0$  and the lower jets fixed. The explicit recursive inverse is

$$q_m = \text{diag}((k_m L_0^m)^{-1}, (k_m L_1^m)^{-1}) \frac{1}{U_m^2 - V_m^2} \begin{pmatrix} U_m & -V_m \\ -V_m & U_m \end{pmatrix} (R_m - \xi_{m-1}). \quad (52.12)$$

The binomial identity (52.4) and  $U_m + V_m > 0$  make every denominator positive. A fixed finite collection has a positive minimum on compact positive leading-geometry boxes. The recursion proves the inverse and its local Lipschitz bounds without any factor  $(\kappa_0 - \kappa_1)^{-1}$ . Equality of all recovered jets gives equality of analytic even graph germs. For analytic connected closed boundaries, the arclength-curvature continuation argument in Corollary 21.2 applies to each contact frame and proves the last assertion.  $\square$

At  $g = 1$ ,  $c_0 = c_1 = 2$ ,  $m = 2$ , the block is

$$D_{q_2} \xi_1 = -\frac{1}{1728} \begin{pmatrix} 49 & 13 \\ 13 & 49 \end{pmatrix}.$$

Its antisymmetric eigenvalue has magnitude  $1/48$ . This finite normalization check is consistent with, but does not substitute for, the all-order proof.

**Corollary 52.2** (Finite-order change of invariant coordinates). *At fixed labelled leading geometry and fixed  $M$ , the two-flight coefficient vector  $\Xi = (\xi_1, \dots, \xi_{M-1})$  and the limiting coefficient vector  $f = (f_1, \dots, f_{M-1})$  are related on the even physical jet image by mutually inverse block-triangular analytic maps. On sufficiently small compact finite-jet neighborhoods these changes of coordinates are bi-Lipschitz, also at equal curvatures.*

*Proof.* Let  $\Psi_M : q \mapsto \Xi$  be the map of Theorem 52.1 and let  $\Phi_M : q \mapsto f$  be that of Theorem 21.1. Then  $f = \Phi_M \circ \Psi_M^{-1}(\Xi)$  and  $\Xi = \Psi_M \circ \Phi_M^{-1}(f)$ . Each map is triangular and analytic. On a sufficiently small compact coordinate neighborhood, the maps and their inverses have bounded derivatives. Integration along segments in the relevant coordinate neighborhoods gives the Lipschitz bounds; neither inverse uses a curvature separation.  $\square$

This is an equivalence of fixed finite-jet coordinates, not an equality of  $G_b$  and  $\mathcal{F}_{bb}$  and not a comparison of noisy function-valued experiments. Its constants may depend on  $M$ . For smooth contacts, equality of all jets does not replace the full-profile uniqueness argument of Theorem 20.1.

**Theorem 52.3** (Two-flight physical observation). *Fix  $M \geq 2$  and a physical analytic family of Theorem 51.1. Put  $n = M - 1$ . There are  $h_0 > 0$  and a compact parameter ball  $\Theta$  with nonempty interior such that the  $2n$  actual means*

$$\mathcal{P}_2(\vartheta) = (\Pr_{\vartheta}(E_{2,b}(\ell h_0)))_{b=0,1; 1 \leq \ell \leq n} \quad (52.13)$$

are bi-Lipschitz coordinates. The physical times are  $2g + \ell h_0$ . Thus the finite-window conclusion of Theorem 51.3 holds with  $j_0 = 2$  for every fixed  $M$ .

Using independent binary preparations, a Borel estimator satisfies, for sufficiently small  $\varepsilon > 0$  and  $0 < \eta < 1/4$ ,

$$\sup_{\vartheta \in \Theta} \Pr_{\vartheta} \{ \|\hat{\vartheta} - \vartheta\| > \varepsilon \} \leq \eta, \quad N_{\text{tot}} \leq C\varepsilon^{-2} \log(C/\eta).$$

All failures are charged. For at most  $N$  fresh preparations, with adaptive or randomized choice among these windows, the minimax squared risk lies between  $c/N$  and  $C/N$  for large  $N$ . The same conclusions hold for the corresponding finite graph jets. Constants and offsets may depend on the fixed family and  $M$ ; the flight number does not. The family and its leading data are supplied.

*Proof.* The support realization and Theorem 52.1 give local analytic coordinates  $\Xi$  on the physical family. Analyticity of the finite stationary integral on a common collar yields

$$G_b(d; \Xi) = 1 + \sum_{k=1}^n \Xi_{b,k} d^k + d^{n+1} R_b(d; \Xi),$$

where  $R_b$  and its first parameter derivatives are bounded after compact restriction. The nodal Jacobian at  $d = \ell h$  is block diagonal with blocks  $V_h = ((\ell h)^k)_{\ell,k=1}^n$ , plus  $O(h^{n+1})$ . Since  $V_h = V_1 \text{diag}(h, \dots, h^n)$  and  $V_1$  is invertible,  $\|V_h^{-1}\| \leq Ch^{-n}$ . Choose a fixed  $h = h_0 > 0$  small enough that all offsets lie in the collar and the transformed error has norm less than  $1/2$ . There is no finite-to-infinite bridge error in this argument.

Passing to physical probabilities multiplies rows by the positive constants  $(\ell h_0)^2 / (2A \sinh(2\gamma))$ . They are fixed on the family. Corollary 52.2 transfers invertibility to the original limiting-jet parameter  $\vartheta$ . On a sufficiently small convex parameter ball, the physical Jacobian, multiplied by its inverse at the center, stays within  $1/2$  of the identity. Integrating along line segments gives both Lipschitz bounds. Use the compact closure as  $\Theta$ . Decrease the collar if necessary so that every probability is below one. Positivity and compactness then place the finitely many means in  $[p_*, 1 - p_*]$  for some fixed  $p_* > 0$ .

Here are the statistical details in this direct two-flight design. Take  $L$  independent preparations in each window and let  $Y$  be the vector of sample means. The bounded-summand exponential estimate and a union bound give

$$\Pr \{ \|Y - \mathcal{P}_2(\vartheta)\|_{\infty} > t \} \leq 4ne^{-2Lt^2}.$$

Choose a finite ordered net in  $\Theta$  whose image has mesh at most  $t$ , and minimize its Euclidean discrepancy from  $Y$ , breaking ties by the first index. This is Borel measurable. On the event that all nodal errors are at most  $t$ , its image error is at most  $C_n t$  and its parameter error at most  $Ct$ . Taking  $t = c\varepsilon$  and  $L = \lceil Ct^{-2} \log(4n/\eta) \rceil$  proves the confidence bound and deterministic charge. With  $L = \lfloor N/(2n) \rfloor$  and a net of mesh  $N^{-1/2}$ , the deterministic discrepancy inequality and  $\mathbb{E} \|Y - \mathcal{P}_2(\vartheta)\|^2 \leq C/L$  give risk at most  $C/N$ .

For the lower bound choose an interior  $\vartheta_*$  and a unit vector  $v$ , and use the physical alternatives  $\vartheta_{\pm} = \vartheta_* \pm aN^{-1/2}v$ . They have exactly the same selected leading hierarchy. Smoothness of the



mean map and its probability margins give, uniformly over the allowed windows,

$$D(\text{Ber}(p_+) \| \text{Ber}(p_-)) \leq \frac{(p_+ - p_-)^2}{p_-(1 - p_-)} \leq \frac{Ca^2}{N}.$$

Conditioning on the past and on the procedure's independent random seed, the rule for choosing the next window is the same under both parameters. The entropy chain rule therefore bounds the entropy of the full record by  $Ca^2$ , including adaptive choices. An early stop can be padded by parameter-independent observations. Fix  $a > 0$  small so that total variation is at most  $1/2$ . The test choosing the nearer parameter from any estimator has summed error at least  $1/2$ ; on a test error the squared estimation loss is at least  $a^2/N$ . Hence the maximum risk is at least  $a^2/(4N)$ . Finally the two parameter systems and the graph-jet vector are locally bi-Lipschitz, so the same orders hold in all of them.  $\square$

The preceding theorem supplies the direct finite-flight design. The long-bridge construction in Theorem 51.3 is retained as an alternative construction using limiting coordinates. Neither construction requires  $j_0$  to grow with requested accuracy on its fixed physical family. The present theorem does not assert that two flights are minimal for every restricted family. In particular, the identical-contact one-flight comparison in Section A.17 remains applicable in its stated class. The general, nonsymmetric energy-profile experiment of the next part has neither this supplied finite-dimensional family nor its parametric risk conclusion.

### 53. THE SIGNED ONE-FLIGHT BENCHMARK AND THE RESTRICTION TO LONG BRIDGES

The signed support observation is richer than an integrated probability germ. It therefore has a short-flight inverse different from the two-flight integrated benchmark in the preceding section. Recording this distinction specifies the contribution of the long-bridge results.

**Proposition 53.1** (One-flight signed support inverse). *Suppose mixed-type one-flight endpoint support is observed as a function of physical time in calibrated signed transverse coordinates, with its onset retained. Then the two labelled facing graph germs are determined without symmetry, equal contacts or supplied curvatures.*

*Proof.* On a sufficiently small contact box the one-flight length is

$$\ell(u, v) = \sqrt{(g + \psi_0(u) + \psi_1(v))^2 + (v - u)^2}.$$

The physical flux factor is positive there. Hence the threshold in physical time at which an endpoint pair enters the support is  $\ell(u, v)$ : the entrance residual time can be any sufficiently small positive number, and no record is possible before the flight length. Equivalently, in excess time the support threshold is  $\ell - g$ ; the onset supplies  $g$ . Since  $\psi_0(0) = \psi_1(0) = 0$ , the coordinate slices give

$$\psi_0(u) = \sqrt{\ell(u, 0)^2 - u^2} - g, \quad \psi_1(v) = \sqrt{\ell(0, v)^2 - v^2} - g. \quad (53.1)$$

The positive square-root branch is fixed by  $g + \psi_b > 0$ . Thus the graph functions themselves, and not just finitely many of their jets, are recovered. For analytic boundaries the same continuation and intrinsic gluing questions can subsequently be considered.  $\square$

This proposition does not identify its observations with the same-type limiting laws. It assumes access to a mixed-type one-flight support function  $\ell(u, v)$ . The signed long-bridge theorem instead recovers two half-line actions from the limiting 00 and 11 endpoint laws, and then inverts their interaction with both contact graphs through the order-by-order determinant-one blocks. The two-flight benchmark concerns integrated probability germs and is a third observation family. No Markov ordering between these three families is asserted without a common acquisition model and an actual kernel.

An unrestricted finite-flight consistency theorem would not, by itself, establish a need for long bridges: the identity (53.1) already gives a direct short-flight route when its support data are available. The physical global theorem here has a more specific conclusion. For any prescribed lower bound on the flight number, its pilot and all inference batches use only even same-type records above that bound. Along the consistent reconstruction sequence that lower bound tends to infinity. Its pilot uses shrinking near-onset windows at those same long even flight numbers, and not a hidden one-flight geometric scan. Thus its contact inverse is genuinely implemented in the restricted long-bridge observation family.

The one-flight identity neither yields a relative law uniform on a fixed offset collar as the flight number grows nor describes the information lost by deleting residual time or waiting counts. Those are separate assertions of the relative boundary law, the explicit Poisson equivalence, and the Gaussian coarsening theorems. In particular, the necessity of the long-bridge analysis in this article is tied to these observation and limit restrictions, not to a claim that a billiard contact cannot be recovered from any shorter experiment.

#### 54. DIRECT GEOMETRIC SAMPLING IN THE SAME LONG-EVEN EXPERIMENT

The restriction to long even bridges does not by itself make the half-line inverse necessary when exact planar endpoint positions are observed. We give a comparison within the physical record space (56.1), using its charged pilot and no short or mixed-type flight. The comparison preserves the global reconstruction problem; it identifies the additional content of the single-offset *transverse-law* inverse.

**Proposition 54.1** (Charged direct acquisition of contact jets). *Fix an integer  $M \geq 2$  and a contact class with uniform  $C^{M+1}$  graph bounds and the uniform localization and flux bounds of Lemma 56.1. In particular, the compact analytic class  $\mathfrak{K}$  has these properties. Fix a target error probability  $\beta \in (0, 1)$ , and any prescribed even flight number  $J \geq 2$ . For every sufficiently small  $a > 0$  and pilot-frame tolerance  $\varepsilon > 0$  satisfying  $\varepsilon \leq c_M a$ , there is a finite capped policy, every preparation of which uses  $J$  flights, that produces estimates of the two graph jets with, uniformly on the class,*

$$\max_{e,b} \left| \widehat{\psi}_{e,b}^{(m)}(0) - \psi_{e,b}^{(m)}(0) \right| \leq C_M \{a^{M+1-m} + \varepsilon a^{-m}\}, \quad 0 \leq m \leq M, \quad (54.1)$$

*except on an event of probability at most  $\beta$ . All failed preparations and discarded off-bin successes are charged. No half-line action is estimated or inverted.*

*Proof.* Choose  $M + 1$  distinct fixed numbers  $\xi_i$  in a bounded interval and pairwise disjoint closed bins  $I_i$  about them. The bins have positive length and a positive minimum separation. The observed transverse bins will be  $aI_i$ . Constants below may depend on these fixed choices and on  $M$ .

Run Theorem 56.2 with  $j = J$ , error probability  $\beta/2$ , and a resolution  $h$  such that

$$C\sqrt{h} \leq \varepsilon, \quad h \leq c_M a^2.$$

On its good event the estimated contact frames have error at most  $\varepsilon$ , and  $J|\widehat{g}_e - g_e| \leq h$ . The pilot is charged and its flight number is never increased afterwards.

Choose a fixed sufficiently large  $D_M$  and program  $t_e = J\widehat{g}_e + D_M a^2$ . For small  $a$  these windows lie in the common positive-offset collar. The finite bridge action obeys  $E_J(u, v) \leq C(u^2 + v^2)$  on a common contact box. In each target bin restrict the true first endpoint to a smaller interior sub-bin of length comparable to  $a$ , the second endpoint to an interval of length comparable to  $a$ , and the residual time to an interior interval of length comparable to  $a^2$ . Increasing  $D_M$  first ensures  $r + E_J < t_e - Jg_e$  throughout these sets; taking  $a$  small afterwards keeps them in the

collar. The frame error, after reducing  $c_M$ , ensures that each smaller true-coordinate bin lies in the corresponding observed bin. Positivity of the exact flux and its uniform lower bound give a one-preparation probability at least

$$p_M(J, a) = c_M a^4 e^{-J\gamma_+} > 0 \quad (54.2)$$

for every target bin, conditional on any good pilot history. Alternatively, for a fixed  $J, a, M$ , compactness and strict positivity alone give a positive lower bound; no effective uniform-in- $J$  estimate is needed for the existential comparison.

For each  $(e, b, i)$ , take at most

$$N_M = \left\lceil p_M(J, a)^{-1} \log \frac{4|E|(M+1)}{\beta} \right\rceil$$

fresh preparations at the stated window, stopping when the first endpoint lies in its observed bin. The chance of exhausting any one cap is at most  $e^{-N_M p_M}$ . A union bound over  $2|E|(M+1)$  bins bounds their total failure probability by  $\beta/2$ . The post-pilot preparation cap is  $2|E|(M+1)N_M$ ; success outside the target bin does not evade this count. Every decision is a Borel function of the physical records and the pilot.

For a retained point  $X_i$ , form its measured transverse coordinate  $\hat{y}_i = \hat{t}_e \cdot (X_i - \hat{p}_{e,b})$  and signed normal coordinate  $\hat{z}_i$  in the estimated frame, with the prescribed facing-side sign. The true graph representation and the frame bounds give

$$|\hat{z}_i - \psi_{e,b}(\hat{y}_i)| \leq C\varepsilon.$$

Indeed the transverse-coordinate error is  $O(\varepsilon)$ , the normal error is  $O(\varepsilon)$ , and the graph has a uniform local Lipschitz bound. The normalized observed nodes  $\hat{y}_i/a$  belong to the fixed separated bins  $I_i$ .

Let  $P_M$  interpolate the  $M+1$  measured pairs. The scaled Vandermonde matrix has a uniformly bounded inverse: its entries vary continuously on the compact product of the bins, and its determinant is a product of nonzero node differences with a uniform lower bound. Taylor's formula on the common  $C^{M+1}$  graph patch shows that the discrepancy between the observations and the Taylor polynomial of degree  $M$  is at most  $C_M(a^{M+1} + \varepsilon)$  at every node. Inverting the scaled matrix and converting its degree- $m$  coefficient into a derivative gives

$$|P_M^{(m)}(0) - \psi_{e,b}^{(m)}(0)| \leq C_M a^{-m} (a^{M+1} + \varepsilon).$$

This is (54.1). Default estimates on flagged failures make the rule defined on every transcript. Adding pilot and bin failures completes the proof.  $\square$

Taking  $\varepsilon \leq a^{M+1}$  gives error  $O_M(a)$  simultaneously through order  $M$ . The pilot can achieve this precision at finite cost by reducing  $h$ ; its cost is included rather than supplied externally. Thus, at stage  $s$ , one may first choose any even  $J_s \geq s$ , then choose the geometric resolution and its pilot, all at that same  $J_s$ . The contact graph jets and gaps determine the action jets by the finite smooth forward map. Alternatively their exact infinite collection is injective on the same signature-rigid analytic class. The compact inverse modulus, finite-template matching, and increasing-order diagonal argument therefore give global consistency from these direct graph samples as well. No claim of an effective analytic-continuation rate is needed for this conclusion.

Theorem 23.4 proves a different assertion. Its successful observations retain  $(u, v)$ , not  $(u, \psi_b(u), v, \psi_b(v))$ . The four-density identity reconstructs the unsymmetrized half-line actions from the distribution of those two scalar coordinates, and the determinant-one recursion recovers the contact geometry encoded by repeated alternating returns. Graph interpolation has no graph ordinates to use in that data map. The two results are compared under their actual observations;

a diverging flight number is not offered as an argument excluding direct interpolation in the richer position experiment.

## 55. FROM NOISY FINITE JETS TO GLOBAL ANALYTIC TABLE LOSS

The exact analytic rigidity theorem uses all contact jets, whereas every finite-sample information theorem is first proved in a fixed finite jet model. We now give a precise bridge between those two levels. The result is a compactness/sieve theorem: it does not postulate a uniform-in-order condition number, and it does not perform unstable numerical analytic continuation. Instead it uses injectivity of the intrinsic lifted-channel data map on a compact analytic class to show that sufficiently many accurately recovered finite coordinates resolve the global table uniformly.

Fix the marked abstract Euclidean lattice, obstacle-orbit labels and a connected measured lifted-channel graph. Let  $\mathfrak{K}$  be a compact family of labelled periodic billiard tables, modulo one simultaneous common element of  $\text{SE}(2)$ , with the following properties.

- (1) Every obstacle boundary is connected, real analytic and strictly convex, and the family has uniform positive curvature, separation and channel-collar bounds.
- (2) The chosen lifted channels persist throughout  $\mathfrak{K}$  with the same labels, deck displacements and orientations.
- (3) The admissible periodic intrinsic gluing space of Theorem 25.3 is a singleton throughout  $\mathfrak{K}$ . A sufficient intrinsic model is a fixed signature-rigid lattice-anchoring cycle together with a signature-rigid spanning tree rooted on an obstacle orbit met by that cycle. If an oriented Euclidean realization of the lattice and one root channel frame relative to that lattice are supplied as ambient calibration, the anchoring cycle may instead be omitted.
- (4) The topology of  $\mathfrak{K}$  is strong enough that analytic boundary images and the corresponding half-line actions vary continuously with all fixed finite derivatives. A uniformly bounded holomorphic extension to a fixed complex collar is a sufficient model.

For every integer  $q \geq 0$ , let  $d_q$  be any metric inducing the labelled  $C^q$ -topology on this quotient class.

For a table  $T \in \mathfrak{K}$  and an integer  $M \geq 2$ , define its finite intrinsic action data

$$\mathcal{D}_M(T) = \left( g_e, S_{e,0}^{(m)}(0), S_{e,1}^{(m)}(0) : e \in E, 2 \leq m \leq M \right). \quad (55.1)$$

The fixed lifted-channel labels  $(a_e, b_e, \ell_e)$  are part of the experimental design and are suppressed from this numerical vector. The full sequence  $\mathcal{D}_\infty$ , together with those labels, is determined by the signed endpoint support germs, and Theorem 25.3 makes it injective on  $\mathfrak{K}$ .

**Theorem 55.1** (Uniform finite-coordinate resolution of a compact analytic class). *For every integer  $q \geq 0$  and every  $\varepsilon > 0$ , there exist an integer  $M = M(q, \varepsilon)$  and a number  $\eta = \eta(q, \varepsilon) > 0$  such that for all  $T, T' \in \mathfrak{K}$ ,*

$$\|\mathcal{D}_M(T) - \mathcal{D}_M(T')\| < \eta \implies d_q(T, T') < \varepsilon. \quad (55.2)$$

*Thus global analytic table loss can be controlled by finitely many intrinsic channel action jets, with the required order allowed to depend on the desired global accuracy.*

*Proof.* Suppose the assertion fails for some  $q$  and  $\varepsilon > 0$ . For each integer  $m \geq 2$  there then exist  $T_m, T'_m \in \mathfrak{K}$  such that

$$d_q(T_m, T'_m) \geq \varepsilon, \quad \|\mathcal{D}_m(T_m) - \mathcal{D}_m(T'_m)\| < m^{-1}.$$

By compactness, after passing to a subsequence,

$$T_m \rightarrow T, \quad T'_m \rightarrow T'$$

in  $\mathfrak{K}$ . Fix any finite order  $M$ . For all sufficiently large  $m \geq M$ , continuity of  $\mathcal{D}_M$  and the preceding estimate give

$$\mathcal{D}_M(T) = \mathcal{D}_M(T').$$

Since  $M$  is arbitrary,  $\mathcal{D}_\infty(T) = \mathcal{D}_\infty(T')$ . The all-order inverse and the singleton admissible periodic gluing hypothesis imply that  $T = T'$  in the  $\text{SE}(2)$ -quotient. Continuity of  $d_q$  then gives  $d_q(T_m, T'_m) \rightarrow 0$ , contradicting the fixed lower bound  $\varepsilon$ .  $\square$

This theorem avoids treating analytic continuation of noisy Taylor series as an algorithm. On a fixed compact analytic class there are no globally separated tables whose intrinsic action data remain arbitrarily close through arbitrarily high finite order.

**Theorem 55.2** (Diagonal growing-order reconstruction principle). *Let  $\mathbf{E}_n = \{P_{n,T} : T \in \mathfrak{K}\}$  be any sequence of physical or boundary experiments. Assume that for every fixed  $M$  there exists an estimator  $\widehat{\mathcal{D}}_{n,M}$  of the finite data  $\mathcal{D}_M(T)$  such that, for every  $\eta > 0$ ,*

$$\sup_{T \in \mathfrak{K}} P_{n,T} \left( \|\widehat{\mathcal{D}}_{n,M} - \mathcal{D}_M(T)\| > \eta \right) \longrightarrow 0. \quad (55.3)$$

*Then there exist a single deterministic nondecreasing sequence  $M_n \rightarrow \infty$  and estimators  $\widehat{T}_n \in \mathfrak{K}$  such that for every fixed integer  $q \geq 0$  and every  $\varepsilon > 0$ ,*

$$\sup_{T \in \mathfrak{K}} P_{n,T} \{d_q(\widehat{T}_n, T) > \varepsilon\} \longrightarrow 0. \quad (55.4)$$

*The same sequence  $M_n$  works simultaneously for all fixed  $q$ . It may be chosen to increase arbitrarily slowly. No uniform-in- $M$  conditioning of the determinant-one recursion is required.*

*Proof.* For every integer  $m \geq 1$  and every  $0 \leq q \leq m$ , apply Theorem 55.1 with accuracy  $m^{-1}$  to obtain  $M(m, q)$  and  $\eta(m, q)$ . Put

$$M_m^* = \max_{0 \leq q \leq m} M(m, q), \quad \eta_m^* = \frac{1}{3} \min_{0 \leq q \leq m} \eta(m, q). \quad (55.5)$$

Increase  $M_m^*$  if necessary so that it is strictly increasing. Since closeness of  $\mathcal{D}_{M_m^*}$  implies the same closeness for every lower truncation, an  $3\eta_m^*$ -accurate  $M_m^*$ -vector resolves every metric  $d_q$ ,  $q \leq m$ , to accuracy  $m^{-1}$ .

By (55.3), choose a strictly increasing sequence  $n_m$  so large that for every  $n \geq n_m$ ,

$$\sup_{T \in \mathfrak{K}} P_{n,T} \left( \|\widehat{\mathcal{D}}_{n,M_m^*} - \mathcal{D}_{M_m^*}(T)\| > \eta_m^* \right) < 2^{-m}. \quad (55.6)$$

For  $n_m \leq n < n_{m+1}$ , set  $M_n = M_m^*$ . Choose  $\widehat{T}_n$  to be a measurable  $\eta_m^*$ -approximate minimizer of

$$T' \longmapsto \|\mathcal{D}_{M_m^*}(T') - \widehat{\mathcal{D}}_{n,M_m^*}\|.$$

For example, fix a countable dense subset of the compact metric space  $\mathfrak{K}$  and choose the first element whose criterion is within  $\eta_m^*$  of the infimum; this gives a measurable selection.

On the event in which the finite-data estimator is within  $\eta_m^*$  of the truth, the minimizing property and the triangle inequality give

$$\|\mathcal{D}_{M_m^*}(\widehat{T}_n) - \mathcal{D}_{M_m^*}(T)\| < 3\eta_m^*.$$

Hence, simultaneously for every  $q \leq m$ ,

$$d_q(\widehat{T}_n, T) < m^{-1}.$$

The exceptional probability is at most  $2^{-m}$ . Given a fixed  $q$  and  $\varepsilon > 0$ , take  $m \geq q$  with  $m^{-1} < \varepsilon$  and let  $n \rightarrow \infty$ . This proves (55.4) with one common sequence  $M_n$ .  $\square$

**Corollary 55.3** (How the physical finite-order theorems feed the global bridge). *Whenever a compact analytic class  $\mathfrak{K}$  admits a common finite acquisition atlas for which every fixed finite action-jet vector  $\mathcal{D}_M$  is uniformly consistently estimable, the actual finite-bridge experiments of this paper admit a growing-order global table reconstruction in every fixed  $C^q$  loss by Theorem 55.2.*

*The fixed-order ingredients used to construct such local estimators are the signed-jet tangent isomorphism, finite positive physical designs, the endpoint-time Poisson experiment (or its endpoint Gaussian coarsening), and the stopped finite-to-boundary transfer. The present theorem supplies the missing  $M_n \rightarrow \infty$  and global-loss passage. Sharp rates in  $M_n$  are not claimed.*

*Remark 55.4* (What has and has not been proved). The bridge is stronger than simply placing an exact analytic theorem beside a fixed-dimensional statistical theorem: it gives a genuine growing-order sieve statement and global table loss. It is deliberately weaker than a sharp analytic minimax theorem. Obtaining an explicit optimal relation among analytic radius, jet order, physical preparation cost and global  $C^q$  risk would require uniform quantitative control of the finite-order inverse as  $M \rightarrow \infty$ . No such bound is inserted by assumption here.

## 56. COMMON PHYSICAL OBSERVATIONS AND CHARGED CHANNEL CALIBRATION

A channel coordinate is not an observation until its origin and direction have been supplied or estimated. We construct the latter from the physical records themselves. The construction uses only even, same-type alternating bridges; their flight numbers may be required to exceed any prescribed integer. In particular, no mixed-type one-flight experiment is used.

**56.1. The acquisition space.** Fix the marked channel labels  $e = (a_e, b_e, \ell_e)$  and their transverse sign conventions. A channel detector reports positions in one oriented Euclidean sensor frame for that channel. The same frame is used for both endpoint types of the channel, but the rigid transformations relating *different* channel frames are neither supplied nor observed. A sensor reports both components of a collision position in physical length units. This is the endpoint-position observation, before taking its transverse projection. Unit speed and the physical clock fix the length unit. The marked obstacle and deck labels are acquisition labels, not supplied deck vectors. As in the physical preparation model, each preparation is an independent draw from normalized billiard phase volume. Its distribution may depend on the table; the instructions for taking and recording a preparation do not.

The design and one-preparation record spaces are the standard Borel spaces

$$\mathcal{A} = \{(e, b, j, t) : e \in E, b \in \{0, 1\}, j \geq 2 \text{ even}, t > 0\}, \quad \mathcal{X} = \{\partial\} \sqcup (\mathbb{R}^2 \times \mathbb{R}^2 \times [0, \infty)). \quad (56.1)$$

For a design  $a = (e, b, j, t)$ , monitor the marked collision word up to the programmed time  $t$ . If it is the specified alternating word of  $j+1$  impacts beginning at type  $b$ , record  $(X, Y, r)$ , where  $X, Y$  are the first and last positions in that channel's sensor frame and  $r$  is the time from preparation to the first impact. Otherwise record  $\partial$ . This is a single measurable map of the marked timed trajectory, defined by labels, comparisons with  $t$ , and coordinate readings. It contains no contact point, tangent, gap or unknown boundary graph. All failures, designs and stopping decisions are retained in the transcript. Extra trajectory information used by the apparatus to select the word is then discarded.

The selected labels refer to the facing lifted obstacle pair, not to a parameter-dependent spatial gate. If detector gates are used, they are fixed before observation and contain the persistent contact neighborhoods for the declared class. Near onset, the localization below makes the two implementations agree. Discovery of unlabelled channels is a different observation problem and is not assumed as part of this record map.

Let  $\mathfrak{K}$  be the compact marked analytic class of Section 27. In this section only its uniform positive curvature, gap, area, separation and persistent-channel bounds are needed. Write

$p_{e,0}, p_{e,1}$  for the two contact positions in the channel sensor frame, and put

$$n_e = (p_{e,1} - p_{e,0})/g_e, \quad t_e = R_{\pi/2} n_e.$$

The prescribed sign convention fixes the possible simultaneous sign change of the transverse axes. These quantities are used in the analysis, not in the acquisition instructions. Bounds below are invariant under arbitrary independent rigid motions of the channel sensor frames.

**Lemma 56.1** (Uniform physical localization and success mass). *There are  $h_0, c_*, C_*, \gamma_+ > 0$ , uniform on  $\mathfrak{R}$ , such that a selected  $j$ -flight record is impossible for  $t \leq jg_e$ . If  $0 < t - jg_e \leq 2h \leq 2h_0$ , each recorded endpoint of type  $b$  satisfies*

$$|X - p_{e,b}| + |Y - p_{e,b}| \leq C_* \sqrt{h}. \quad (56.2)$$

For every even  $j \geq 2$  and  $h \leq t - jg_e \leq 2h \leq 2h_0$ , the selected-record probability is at least

$$p_*(j, h) := c_* h^2 e^{-j\gamma_+} > 0. \quad (56.3)$$

The constants can be decreased so that  $p_*(j, h) < 1$ .

*Proof.* Every flight between the specified two lifted obstacles has length at least  $g_e$ . Their unique closest pair is an interior nondegenerate minimum of the two-endpoint distance. In contact arclength coordinates its Hessian is positive definite, uniformly on the compact class. Taylor's formula therefore gives

$$\ell_e(s, s') - g_e \geq c(s^2 + (s')^2)$$

on a common neighborhood. On its complement a positive uniform distance margin follows from uniqueness, channel persistence and compactness. Thus if the sum of the  $j$  flight lengths is less than  $jg_e + 2h$ , for  $h$  sufficiently small every flight is in that neighborhood. Applying the displayed bound to the first and last flights proves (56.2); the constant does not grow with  $j$ . Nonnegative entrance and exit times prove the assertion before onset.

For the lower mass use the exact flux density in (34.1), in unscaled coordinates. Uniformly on the contact box,  $E_j(u, v) \leq C(u^2 + v^2)$ ,  $b_j \geq c > 0$ , and  $-W_{j,uv}(0, 0) \geq ce^{-j\gamma_+}$ . Restrict  $(u, v)$  to a ball of radius  $c\sqrt{h}$  and  $r$  to an interval of length  $ch$ , so that  $r + E_j < h$ . Integrating gives  $ch^2 e^{-j\gamma_+}/A$ . The uniform upper bound on the free area gives (56.3). These bounds follow from the finite bridge and relative-law estimates with strict uniform curvature and hyperbolicity margins; they require no finite-dimensional parametrization of  $\mathfrak{R}$ .  $\square$

**56.2. A finite capped scan.** Choose known bounds  $0 < g_- \leq g_e \leq g_+$  for the class. Fix an even flight number  $j$ , a resolution  $h < h_0$ , and an error probability  $\beta \in (0, 1)$ . All of these are chosen before the pilot. Set

$$t_l = jg_- + lh, \quad 0 \leq l \leq L, \quad L = \lceil j(g_+ - g_-)/h \rceil + 2, \quad N = \left\lceil p_*(j, h)^{-1} \log(2|E|/\beta) \right\rceil. \quad (56.4)$$

For each  $(e, b)$ , scan these times in increasing order, taking at most  $N$  independent preparations at each time, and stop at the first successful record. Denote its time by  $T_{e,b}$  and its first endpoint by  $\hat{p}_{e,b}$ . If a scan exhausts the grid, report pilot failure. The total preparation cap is

$$B_{\text{pil}} = 2|E|(L + 1)N. \quad (56.5)$$

No intermediate failed preparation is omitted from this count.

**Theorem 56.2** (Observable contact-frame and onset calibration). *The preceding policy has, with probability at least  $1 - \beta$  uniformly on  $\mathfrak{R}$ , a successful outcome in every scan and*

$$jg_e < T_{e,b} \leq jg_e + 2h.$$

On this event define

$$\hat{g}_e = \frac{\min_b T_{e,b} - h}{j}, \quad \hat{n}_e = \frac{\hat{p}_{e,1} - \hat{p}_{e,0}}{|\hat{p}_{e,1} - \hat{p}_{e,0}|}, \quad \hat{t}_e = R_{\pi/2} \hat{n}_e. \quad (56.6)$$

For all sufficiently small  $h$ , the denominator is at least  $g_-/2$  and

$$\max_e j|\hat{g}_e - g_e| \leq h, \quad \max_{e,b} |\hat{p}_{e,b} - p_{e,b}| + \max_e |\hat{t}_e - t_e| \leq C\sqrt{h}. \quad (56.7)$$

The policy and estimates are measurable functions of the transcript on (56.1). They require no supplied contact center, tangent, curvature, gap or inter-channel placement.

*Proof.* For each  $(e, b)$  there is a grid point whose true excess is in  $[h, 2h]$ . At that point the probability of  $N$  failures is at most  $(1 - p_*)^N \leq \exp(-Np_*) \leq \beta/(2|E|)$ . If the scan has stopped earlier, its first success cannot precede onset and its time is even smaller. A union bound proves the claimed simultaneous event. In particular, this argument does not require a lower success bound at times far beyond the near-onset collar, nor a decision that a finite run of failures proves zero probability.

The gap bound follows by subtracting the midpoint  $h$  from a number in  $(0, 2h]$ . Lemma 56.1 gives the two position bounds. For vectors  $v, w$  with  $|v| \geq g_-$  and  $|v - w| \leq g_-/2$ ,

$$\left| \frac{v}{|v|} - \frac{w}{|w|} \right| \leq 4|v - w|/g_-.$$

Apply this with the true and estimated contact displacements. Each operation in (56.6) is Borel on the good event. On a failed scan or a small denominator assign a fixed default value and a failure flag. This defines the policy and estimators on every transcript. Rigidly moving a channel sensor moves both estimated contact positions by the same rigid motion and rotates the estimated tangent. Hence the bounds are uniform in these unsupplied sensor placements.  $\square$

The pilot flight number may be any prescribed even  $j$ , however large. Later we set it equal to the already chosen post-pilot flight number  $J$ . The pilot therefore controls  $J|\hat{g}_e - g_e|$ , not a different quantity at an earlier, smaller flight number. Both types are measured in the same channel sensor frame; this does not register different channels.

**56.3. Observable normalization and continuity of the tested laws.** For a fresh successful same-type record set

$$\hat{C}_{e,b}(X, Y, r) = (\hat{t}_e \cdot (X - \hat{p}_{e,b}), \hat{t}_e \cdot (Y - \hat{p}_{e,b}), r). \quad (56.8)$$

Its ideal counterpart  $C_{T,e,b}$  replaces the estimates by the true contact frame. On the pilot event and a common positive-offset collar,

$$|\hat{C}_{e,b}(X, Y, r) - C_{T,e,b}(X, Y, r)| \leq C\sqrt{h}. \quad (56.9)$$

Indeed the projected displacement error is bounded by the position error plus the tangent error times  $|X - p_{e,b}|$ , which is uniformly bounded on the selected collar. No bound on the absolute sensor origin is needed.

Let  $Q_{T,e,b,d}^\infty$  be the centered successful endpoint-time law, with density, extended by zero to a fixed box,

$$q_{T,e,b,d}(u, v, r) = Z_{T,e,b,d}^{-1} B_{e,b}(u) B_{e,b}(v) \mathbf{1}_{\{0 < r < d - S_{e,b}(u) - S_{e,b}(v)\}}. \quad (56.10)$$

For any compact positive-offset interval inside the common collar,  $Z$  is bounded below by a positive constant. If  $S, B, d$  and  $\tilde{S}, \tilde{B}, \tilde{d}$  are two such descriptions, fiber integration and  $|x_+ - y_+| \leq |x - y|$  give

$$\|Q - \tilde{Q}\|_{\text{TV}} \leq C(\|S - \tilde{S}\|_\infty + \|B - \tilde{B}\|_\infty + |d - \tilde{d}|). \quad (56.11)$$



To verify this estimate, integrate the amplitude difference on the common fibers and the uniformly bounded amplitudes on their symmetric difference; the latter length is bounded by the action and offset differences. The same bound applies to the normalizers, after which division by their positive lower bound proves the assertion. The half-line contraction and relative operator estimates imply continuous dependence of  $S, B$  on the table in this fixed collar. Thus the centered laws are continuous in total variation. This is not a claim of total-variation continuity for ambient position laws supported on different boundary curves.

**Proposition 56.3** (Uniform implementation of a finite family of tests). *Let  $(e_r, b_r, d_r, \phi_r)_{r=1}^R$  be a fixed finite family with positive offsets in the interior of the collar and bounded Lipschitz functions  $|\phi_r| \leq 1$ . Program the post-pilot time  $J\hat{g}_{e_r} + d_r$ , where  $J$  is chosen before a pilot with  $j = J$ . Conditional on any good pilot history, a fresh successful record satisfies*

$$\left| \mathbb{E}_T \phi_r(\hat{C}_{e_r, b_r}(X, Y, r)) - Q_{T, e_r, b_r, d_r}^\infty \phi_r \right| \leq C_{\mathcal{F}}(\sqrt{h} + h + \tau^J). \quad (56.12)$$

*The actual statistic and the programmed design are parameter independent functions of the transcript.*

*Proof.* On the good history the true offset differs from  $d_r$  by at most  $h$ . Choose  $h$  below half the smallest offset and below the upper collar margin. The conditional physical-to-boundary comparison (34.4) costs at most  $2C\tau^J$  for a test bounded by one. The same boundary-position embedding is used for both measures at the fixed true table. Applying the *observable* map (56.8) cannot increase that distance. Replacing this map by its ideal counterpart costs at most the Lipschitz constant times  $C\sqrt{h}$ . Finally (56.11) bounds the offset error by  $Ch$ . Summing proves (56.12). The ideal map appears only in this error calculation. It is not a kernel that the experimenter must apply.  $\square$

## 57. SINGLE-OFFSET STATISTICAL SEPARATION AND GLOBAL PHYSICAL RECONSTRUCTION

We construct the estimator on the common observation space of Section 56. The identifiable datum consists of the onsets and the centered conditional laws. The onsets enter the estimation criterion, not merely the programmed times. The single-offset inverse permits every post-pilot test to ignore residual time and both longitudinal coordinates. Only the two endpoint types at one fixed offset per channel are needed. No finite-dimensional model for the entire compact analytic class is imposed.

Fix one  $d_0 > 0$  inside the common collar. For  $M \geq 2$ , write

$$\mathcal{D}_M(T) = (g_e, S_{e,0}^{(m)}(0), S_{e,1}^{(m)}(0) : e \in E, 2 \leq m \leq M).$$

Use the normalized finite-coordinate norm of Theorem 27.4, with a fixed positive gap scale  $g_*$ . The notation  $Q_{T,e,b,d_0}^\infty$  denotes the centered successful endpoint-time law of (56.10); all tests below are independent of its residual-time coordinate.

**Lemma 57.1** (Single-offset mixed onset-law separators). *For every  $M \geq 2$  and  $\eta > 0$ , there are finitely many labelled bounded Lipschitz tests  $(e_r, b_r, \varphi_r)_{r=1}^R$ , with  $|\varphi_r| \leq 1$ , and  $c > 0$  such that*

$$\|\mathcal{D}_M(T) - \mathcal{D}_M(T')\|_M \geq \eta \implies \|F(T) - F(T')\|_\infty \geq 8c, \quad (57.1)$$

*where the continuous moment map is*

$$F(T) = \left( (g_e(T)/g_*)_e, \left( Q_{T,e_r,b_r,d_0}^\infty \varphi_r(u, v) \right)_{r=1}^R \right). \quad (57.2)$$

*All tests use the same fixed offset  $d_0$ . The list and the separation constant depend only on the compact class,  $M$ , and  $\eta$ .*

*Proof.* The set of pairs on the left of (57.1) is compact. For any such pair, either a gap differs, or all gaps agree. In the latter case, equality of both endpoint marginals at  $d_0$  for every channel would, by Theorem 23.1, imply equality of every half-line action germ. It would therefore imply equality of  $\mathcal{D}_M$ , a contradiction. Hence one endpoint law separates the pair. Bounded Lipschitz functions separate finite Borel probability measures on the fixed compact endpoint box: continuous functions do so, and Lipschitz functions are uniformly dense in the continuous functions. Normalize a separating function to absolute value at most one; no uniform bound on its Lipschitz constant is required before the finite list is chosen.

Gap continuity and (56.11), followed by marginalization, make each strict separation persist on a neighborhood of the pair. A finite subcover gives finitely many tests and a positive minimum margin  $8c$ . Include every gap coordinate whether or not it was needed in a particular neighborhood. Add constant tests for any missing  $(e, b)$ , so all  $D = 2|E|$  design groups are represented and  $R \geq D$ . If the compact pair set is empty, any positive  $c$  and these constant tests suffice. This proves the lemma.  $\square$

A conditional law alone is not here asserted to recover the gap. Already at quadratic order, fixing  $a_0 = a_1 = a$  permits

$$\kappa_0(g) = \kappa_1(g) = \frac{\sqrt{1 + g^2 a^2} - 1}{g}.$$

This explains the separate onset coordinate without asserting equality of the full nonlinear laws along this quadratic example.

**57.1. A finite minimum-distance rule.** Choose fixed templates  $T_1, \dots, T_H \in \mathfrak{K}$  whose images form a  $c$ -net of  $F(\mathfrak{K})$ . For each  $(e, b)$ , collect  $k$  fresh successful records after the pilot. Apply the observable transverse projection in (56.8), and evaluate all tests belonging to that group on the same sample. Put

$$\hat{F} = \left( (\hat{g}_e / g_*)_e, \left( k^{-1} \sum_{i=1}^k \varphi_r(\hat{u}_{e_r, b_r, i}, \hat{v}_{e_r, b_r, i}) \right)_{r=1}^R \right).$$

Select

$$\hat{T} = T_{\hat{h}}, \quad \hat{h} = \min_{1 \leq h \leq H} \argmin \|F(T_h) - \hat{F}\|_\infty, \quad \hat{\mathcal{D}}_M = \mathcal{D}_M(\hat{T}). \quad (57.3)$$

On a flagged pilot or batch failure use  $T_1$ . This finite rule is Borel with a fixed tie convention. It does not hide a selection theorem or an unbounded optimization over unknown analytic functions. Compactness gives the library, not an effective rate for generating it from collar constants.

If  $\|\hat{F} - F(T)\|_\infty \leq c$ , a template within  $c$  of  $F(T)$  has criterion at most  $2c$ . Consequently

$$\|F(\hat{T}) - F(T)\|_\infty \leq 3c < 8c, \quad \|\mathcal{D}_M(\hat{T}) - \mathcal{D}_M(T)\|_M < \eta. \quad (57.4)$$

The first norm includes the gap coordinates.

**Theorem 57.2** (Uniform finite-order physical estimation at one offset). *For every  $M \geq 2$ ,  $\eta > 0$ ,  $\delta \in (0, 1)$ , and prescribed minimum flight number  $J_0$ , there is one deterministically capped policy on  $\mathcal{A}, \mathcal{X}$  such that every preparation uses the same chosen even  $J \geq \max(2, J_0)$ , and*

$$\sup_{T \in \mathfrak{K}} \Pr_T \{ \|\hat{\mathcal{D}}_M - \mathcal{D}_M(T)\|_M \geq \eta \} < \delta. \quad (57.5)$$

*No gap, contact frame, inter-channel registration or lattice metric is supplied. After the pilot there are exactly  $2|E|$  design groups, all programmed at  $J\hat{g}_e + d_0$ ; the estimator uses only their transverse endpoint pairs and the pilot summaries. Every preparation is charged, and all design and stopping decisions are measurable.*

*Proof.* Choose the separators,  $c \leq 1$ , and the template library first. Take

$$k \geq 8c^{-2} \log(8R/\delta). \quad (57.6)$$

Choose the final even  $J \geq \max(2, J_0)$  so that the finite-flight bias in Proposition 56.3 is at most  $c/4$  and

$$CDk\tau^J < \delta/8. \quad (57.7)$$

Only then choose  $h > 0$ , below both pilot and offset-collar margins, with

$$C_{\mathcal{F}}(\sqrt{h} + h) \leq c/4, \quad h/(Jg_*) \leq c. \quad (57.8)$$

Run the calibration theorem with  $j = J$ ,  $\beta = \delta/8$ , its finite grid and its deterministic cap  $B_{\text{pil}}$ . Thus the pilot controls  $J|\hat{g} - g|$  for the final flight number, which is not changed later.

On a good pilot history program each group at  $(e, b, J, J\hat{g}_e + d_0)$ , take fresh preparations, and stop at its  $k$ -th success. The true offsets remain in a fixed positive subcollar. Uniform flux gives a success lower bound  $p_0 > 0$ , depending on the fixed design but not on the table or the good history. Impose the per-group cap

$$B = \left\lceil \frac{2k + 8 \log(8D/\delta)}{p_0} \right\rceil. \quad (57.9)$$

Since  $Bp_0 \geq 2k$ , the binomial lower-tail bound gives probability at most  $\exp(-Bp_0/8) \leq \delta/(8D)$  of not obtaining  $k$  successes. The union of batch-cap failures has probability at most  $\delta/8$ .

Conditional on the pilot, successive marks in an uncapped fixed-design success sequence are iid with its successful physical law. The capped and uncapped constructions agree whenever the cap is not hit. Apply concentration to the uncapped sequence and add the cap probability; conditioning on completion of the cap is not used to assert iid sampling. Proposition 56.3, applied to  $\phi_r(u, v, r') = \varphi_r(u, v)$ , gives bias at most  $c/2$ . Hoeffding's inequality on  $[-1, 1]$  and a union bound give

$$\Pr_T \left\{ \max_r |\hat{\mu}_r - \mathbb{E}_T(\hat{\mu}_r \mid \text{pilot})| > c/2 \mid \text{good pilot} \right\} \leq 2R \exp(-kc^2/8) \leq \delta/4.$$

Tests sharing a sample need not be independent for this union bound. On the complementary event the test and gap coordinates are all within  $c$  of  $F(T)$ . Apply (57.4). The sum of pilot, empirical and cap probabilities is at most  $\delta/2$ , hence strictly less than  $\delta$ . The full cap is  $B_{\text{pil}} + DB$ .

The complete-transcript comparison is retained as well. Keep the pilot exact. Conditional on each good history the post-pilot designs are fixed and observable. The stopped comparison of Theorem 35.3 costs at most  $CDk\tau^J$ , since each group stops no later than its  $k$ -th success. Corollary 35.4 adds the pilot exceptional probability when histories are integrated. At each fixed table the finite and boundary position laws use the same physical embedding, and estimated projection is a common transcript map. This proves a distance bound without identifying an unknown exact centering map with an observation. The comparison includes all failures, counts, residual times and stopping decisions, even though the estimator discards some of them. The direct concentration proof already applies to the actual finite physical experiment.  $\square$

## 57.2. Increasing order and a prescribed preparation budget.

**Theorem 57.3** (Uniform global reconstruction from long same-type records). *On the compact marked analytic class of Section 27, there are finite capped policies  $\pi_s$  and table estimators  $\hat{T}_s$  such that every flight number used by  $\pi_s$  is even and at least  $s$ , and, for every fixed  $q \geq 0$  and  $\epsilon > 0$ ,*

$$\sup_{T \in \mathfrak{R}} \Pr_T^{\pi_s} \{d_q(\hat{T}_s, T) > \epsilon\} \longrightarrow 0. \quad (57.10)$$

*The recovered table includes its Euclidean period lattice. Only one simultaneous global element of  $SE(2)$  is quotiented out. After calibration each policy uses the same fixed offset  $d_0$  and only the two endpoint types in each channel.*

*Proof.* Use Theorem 27.4 to choose increasing  $M_s \geq s$  and decreasing  $\eta_s > 0$  with

$$\max_{0 \leq q \leq s} \omega_q(C\eta_s + C2^{-M_s}) < 1/s.$$

Each of the finitely many moduli tends to zero. Apply the preceding theorem with  $(M_s, \eta_s, 2^{-s}, J_0 = s)$  and keep its table-valued template estimator. Except on an event of probability at most  $2^{-s}$ , the finite-data modulus gives  $d_q(\hat{T}_s, T) < 1/s$  for every  $q \leq s$ . This proves the uniform assertion. The exact inverse used by the modulus includes the marked deck holonomies and  $L = VM^{-1}$ , so no lattice Gram form is inserted into the statistical conclusion.

For a prescribed total budget  $N$ , let  $b_s$  be the deterministic preparation cap of  $\pi_s$ , and set

$$C_s = \max\{s, b_1, \dots, b_s\}, \quad s(N) = \max\{s : C_s \leq N\}.$$

For  $N < C_1$  use a fixed default estimate with no preparations. Otherwise run  $\pi_{s(N)}$  afresh and spend no remaining budget. This is a budget-indexed family, not a cumulative execution of all preceding stages. It uses at most  $N$  preparations,  $s(N) \rightarrow \infty$ , and every flight number in its entire acquired transcript is at least  $s(N)$ . This proves the same consistency statement with the claimed flight restriction for the prescribed-budget implementation. A cumulative historical transcript would still contain the earlier flight numbers; no diverging minimum is asserted for that different implementation.  $\square$

The library, separation constants and inverse moduli may be non-effective functions of the compact class. The theorem gives finite measurable acquisition once these constants are fixed, not polynomial-time design generation or a sharp analytic minimax rate. Its caps expose the rare-event cost. The direct-position benchmark also respects a diverging even flight number. The extra intrinsic conclusion is therefore not the exclusion of that benchmark, but the law inverse and its implementation using scalar post-pilot tests at one fixed offset.

## 58. COMPATIBLE ANALYTIC VARIATIONS AND UNIFORM FINITE POSITIVE DESIGNS

The compact analytic class used for global consistency need not be a finite-dimensional manifold. Nor does retaining finitely many target jets make a positive-offset statistical law depend on finitely many parameters. We give a precise finite-rank variation construction for the separate local information assertion. The global estimator of Theorem 57.3 does not depend on this construction.

Fix a compact family  $\mathfrak{K}_0$  with a fixed pattern of coincident and distinct marked contacts on each obstacle. Coincident contacts are identified before specifying variations. Assume the normal angles of distinct contacts on the same obstacle have a positive uniform separation. This is a compact contact stratum; no such stratification is needed in the global separator theorem. Choose an oriented ambient angular coordinate for each obstacle in a fixed root-frame gauge. Its analytic support function is  $h_i(\theta)$ , with radius of curvature  $R_i = h_i + h_i'' > 0$ . At a contact angle  $\theta_a$ , its point is

$$X_i(\theta_a) = h_i(\theta_a)n(\theta_a) + h_i'(\theta_a)t(\theta_a).$$

Consequently a support perturbation with zero value and first derivative at  $\theta_a$  preserves that point and its oriented tangent exactly. It also preserves the gaps of the anchored channel network when this is imposed at every marked incidence. We vary the obstacle shapes in an ambient analytic family of dispersing tables; an arbitrary closed set  $\mathfrak{K}_0$  is not asserted to contain these variations.

**Lemma 58.1** (Uniform analytic Hermite right inverse). *Fix an integer  $M \geq 2$ . For  $l$  distinct contact angles there is a finite-dimensional space  $\mathcal{T}_D$  of real trigonometric polynomials such that the jet evaluation map*

$$E_T : \mathcal{T}_D \longrightarrow \mathbb{R}^{l(M+1)}, \quad E_T f = (f^{(r)}(\theta_a))_{a, 0 \leq r \leq M}$$

*is onto. One may take  $D = (l-1)(M+1) + M$ . In a fixed real Fourier basis,*

$$R_T = E_T^t (E_T E_T^t)^{-1} \quad (58.1)$$

*is a continuous right inverse, uniformly bounded on the compact contact stratum. It depends smoothly on the distinct contact angles.*

*Proof.* For each contact  $a$ , put

$$F_a(\theta) = \prod_{b \neq a} (1 - \cos(\theta - \theta_b))^{M+1}.$$

This polynomial is nonzero at  $\theta_a$  and vanishes to order  $2(M+1)$  at all other contact angles. The  $M+1$  functions  $F_a(\theta) \sin^r(\theta - \theta_a)$ ,  $0 \leq r \leq M$ , have zero jets through order  $M$  at other contacts. At  $\theta_a$  their jet matrix is triangular, with diagonal  $r! F_a(\theta_a) \neq 0$ . Their degrees are at most  $D$ . They therefore interpolate arbitrary jets at this contact, independently of the other contacts. Summing over  $a$  proves surjectivity. The Gram matrix in (58.1) is positive definite and varies smoothly with the contact angles. Its least eigenvalue has a positive minimum on the compact separated-angle set. This proves the remaining assertions.  $\square$

**Proposition 58.2** (Realization of compatible finite contact variations). *Over  $\mathfrak{K}_0$ , assign one labelled oriented graph jet  $(q_{a,2}, \dots, q_{a,M})$  to each distinct physical contact. Repeated incidences use the same jet, transformed to their prescribed oriented frames. These coordinates define a finite-rank vector bundle  $\mathcal{V}_M$  with a continuous norm. Its vectors have analytic realizations through their base tables which preserve every marked contact point, tangent and gap. The realizations vary continuously with the base and smoothly with the fiber variable; on a uniform fiber neighborhood they remain strictly convex separated dispersing tables with the same persistent channels.*

*Proof.* Take the direct sum of the right inverses of Lemma 58.1 over the obstacles. Require the zeroth and first support-jet increments to vanish. The remaining support jets through order  $M$  are local coordinates for graph jets through order  $M$  at a fixed point and tangent. Here is the triangular reason. Since  $dX/d\theta = Rt$ , at the contact the transverse graph coordinate satisfies  $dy/d\theta = R > 0$ . The graph curvature is  $q_2 = R^{-1}$ . Repeated differentiation with  $d/dy = R^{-1}d/d\theta$  at the contact shows that, after lower support jets have been fixed, the coefficient of a variation of  $h^{(m)}$  in  $q_m$  is nonzero (in the convex graph convention it is  $-R^{-m}$ ). Thus the differential from support jets  $h^{(2)}, \dots, h^{(M)}$  to graph jets  $q_2, \dots, q_M$  is triangular and invertible. Changing the oriented transverse sign multiplies the corresponding odd-order coordinates by minus one and does not affect invertibility.

The finite-dimensional inverse function theorem now assigns support-jet increments to each sufficiently small desired graph-jet increment. Extend these increments by  $R_T$ , with the first two jets zero. The result is an actual real-analytic support function, not independent formal contact series. The interpolation imposes all incidences on each obstacle simultaneously, so compatibility is automatic. Uniform curvature and separation margins, compactness and the uniform bounds on  $R_T$  give a common small fiber neighborhood on which the realized tables are admissible. The periodic lattice and obstacle reference placements are kept fixed.

In each choice of labelled frames the fibers are the direct sum of the finite jet-coordinate spaces just described. On overlaps, frame changes and relabellings induce the differentials of the

corresponding jet-coordinate changes. They are continuous invertible matrices and satisfy the cocycle identity because the underlying coordinate changes do. This specifies the bundle and its transition maps. A finite frame atlas and a partition of unity give a continuous fiber norm. Its unit sphere bundle over the compact base is compact. This construction does not give a tangent bundle to an arbitrary closed analytic class; it gives a specified bundle of ambient, compatible analytic variations through it.  $\square$

For  $v \in \mathcal{V}_M(T)$ , let  $\dot{S}_{e,b}[v]$  be the action variation along this realization. Define the positive-offset information form by (39.11) with boundary velocity

$$U_{e,b}(u, w)[v] = \dot{S}_{e,b}(u)[v] + \dot{S}_{e,b}(w)[v].$$

The gap is fixed in this bundle. A gap direction in a separately specified anchored one-channel local family is handled by the additional constant velocity in Lemma 39.2; no independent gap directions are assigned to an arbitrary globally compatible network here.

**Lemma 58.3** (Uniform finite positive design on the specified bundle). *For each fixed  $M$ , there are finitely many labelled same-type positive offsets  $(e_r, b_r, d_r)$ , positive weights  $\alpha_r$  summing to one, and  $c_M > 0$ , independent of the base table in  $\mathfrak{R}_0$ , such that*

$$\sum_r \alpha_r \mathcal{K}_{T, e_r, b_r, d_r}[v, v] \geq c_M \|v\|^2 \quad (v \in \mathcal{V}_M(T)). \quad (58.2)$$

*This statement concerns the explicitly realized finite-rank bundle, not all variations of the infinite-dimensional analytic class.*

*Proof.* If a vector lies in the kernels for both types of every channel and every small positive offset, positivity of the boundary weight implies

$$\dot{S}_{e,b}(u)[v] + \dot{S}_{e,b}(w)[v] = 0 \quad \text{when } S_{e,b}(u) + S_{e,b}(w) = d.$$

Set  $w = 0$  and  $d = S_{e,b}(u) > 0$ . The anchored variation satisfies  $\dot{S}_{e,b}(0)[v] = 0$ , hence every  $\dot{S}_{e,b}$  vanishes near the contact. The fixed-gap finite action-jet inverse in Theorem 22.1 forces the variations of all marked contact jets through order  $M$  to vanish. These jets are the fiber coordinates, so  $v = 0$ .

The forms are continuous in the base, direction and each fixed positive offset. Indeed the realized support families and their finite derivatives are continuous; the half-line inverse transports this continuity to the actions and amplitudes; and each positive-offset level has a regular boundary with a common local coordinate description. The coarea formula then gives continuity of the information integral. Every unit vector in the compact sphere bundle is detected by a design, and this remains true on a neighborhood. A finite subcover supplies finitely many designs. Their positive weighted sum has a positive continuous minimum on that sphere bundle, proving (58.2).  $\square$

The hypotheses about the contact stratum and the compatible analytic realizations belong to this local-design lemma alone. In particular, a finite-rank information argument is not needed to pass from the compact analytic class to the global physical consistency theorem. That passage uses the augmented observable separator lemma and its finite capped policy.

**Acknowledgments.** The nuisance-splicing construction used in Section A.16 was supplied in the author-requested memorandum [17]; the author acknowledges this contribution. The author-requested referee-style memoranda were not commissioned journal reports. Revision records and source provenance are archived separately from the mathematical article. The distinction between finite rotations and infinitesimal symmetry, and the circular-lattice comparison in Section 30, were suggested in the author-requested memorandum [38]. The finite enumeration

of compatible reconstruction branches is developed here from that distinction and the intrinsic cochain inverse. The noncircular two-table construction in Proposition 30.5 was suggested in the subsequent memorandum [39]. The additional-channel separation in Proposition 30.6 is proved here. These memoranda are cited for attribution, not as journal endorsements or proof certificates. The support-width centering comparison in Lemma 24.5 was communicated in the subsequent memorandum [42]. The quadratic functional small-window benchmark was communicated in [43]; Section 33 supplies its actual-table realization, attainable two-point scale and finite-flight comparison. The Hellinger improvement and repeat-to-first-success comparison in Section 33 were communicated in [44]; the deterministic-cap risk is proved here. The stopped-count and finite-intensity bit reductions, and the count-risk formula incorporated in the same theorem, were communicated in the two v54 memoranda [45, 46]; their overlap is acknowledged rather than counted as two independent contributions.

The quadratic Riccati comparison in the author-requested AI-assisted memorandum [50] suggested the global curvature argument developed here. That memorandum is not a commissioned journal report.

The finite-iteration amplification control in the author-requested AI-assisted memorandum [51] led to the explicit residual and higher-jet propagation bounds in Proposition 4.4. This attribution concerns the quantitative correction, not an editorial certification.

## Part 6. Further formulations and complete comparison proofs

### APPENDIX A.1. FURTHER FORMULATIONS AND COMPARISONS

A long billiard bridge is a rare event, but its conditional endpoint law need not be a quadratic record of the orbit it shadows. We prove that on a fixed contact collar it retains nonlinear half-line actions, and that these actions identify the complete smooth reflecting contact germs. The forward statement is relative to an exponentially small endpoint twist. The inverse concerns the actual boundary, including both its curvature and its flat part, not an arbitrary formal generating function. No analytic hypothesis is needed for this local identification. These two statements form the principal mechanism of the paper.

For clarity we first give the conclusion for a supplied periodic polygon. A *marked polygon* consists of the lifted collision points  $q_i$ , their signed tangent and outward-normal frames  $t_i, n_i$ , the phase labels and closing translation, with the specular reflection condition imposed. Write  $L_i = |q_{i+1} - q_i|$ ,  $e_i = (q_{i+1} - q_i)/L_i$  and  $\nu_i = e_i \cdot n_i > 0$ . Unknown contacts are represented by

$$R_i(x) = q_i + \nu_i^{-1} \{t_i x - n_i F_i(x)\}, \quad F_i(0) = F'_i(0) = 0, \quad F''_i(0) > 0.$$

The measured coordinate  $x = \nu_i t_i \cdot (R_i - q_i)$  is a tangent projection. Neither the normal graph height nor the curvature is a mark. Two fixed positive offsets are observed at every physical phase, on a common square with positive residuals. Their limiting densities have the form

$$f_{i,d}(u, v) = Z_{i,d}^{-1} B_i^-(u) B_i^+(v) \{d - A_i(u) - C_i(v)\}, \quad d = d_1, d_2,$$

with anchored actions and amplitudes. The amplitudes are shared across the two offsets, while the normalizers need not agree. The physical momenta are restored in  $A_i = S_i^- - p_i x$  and  $C_i = S_i^+ + p_i x$ , where  $p_i = e_i \cdot t_i / \nu_i$ .

Let  $P$  be a returning oriented period,  $N = nP$ , and  $\mathcal{L}$  the length of that period. Denote by  $W_{N,b}(u, v)$  the stationary length along the selected itinerary and put

$$E_{N,b} = W_{N,b} - n\mathcal{L} + p_b u - p_b v, \quad D_{N,b} = | - W_{N,b,uv}(0, 0) |.$$

For a returning signed frame the mixed derivative has the sign used in the following formula. The exact reference twist is

$$D_{N,b} = \frac{\sinh \chi}{(\mathcal{M}_b)_{12} \sinh(n\chi)},$$

where  $\mathcal{M}_b$  is the positive Jacobi transfer over one returning period in the scaled signed coordinates. It is not additional observation data. The normalized twist, rather than an absolute approximation to  $W_{N,b}$ , is the quantity which survives conditioning.

**Theorem A.1.1** (Identification from periodic boundary laws). *Fix a clear nongrazing marked polygon with finitely many distinct physical contact points in a planar dispersing billiard. On a contact collar independent of  $N$ , for every separately fixed smooth order  $k$ ,*

$$\|E_{N,b} - S_b^- \oplus S_b^+\|_{C^k} + \left\| \frac{-W_{N,b,uv}}{D_{N,b}} - B_b^- \otimes B_b^+ \right\|_{C^k} \leq C_k \tau^N, \quad \tau < 1. \quad (\text{A.1.1})$$

The physical endpoint measure has density

$$\frac{D_{N,b} - W_{N,b,uv}}{2\pi\mathcal{A} D_{N,b}} \{d - E_{N,b} + p_b u - p_b v\}$$

on a common positive-residual square, where  $\mathcal{A}$  is the free cell area. Its conditional densities converge in the same fixed-order topology to  $f_{b,d}$  above. The phase-resolved limiting laws at two separated positive offsets identify the complete positive-curvature  $C^\infty$  contact germs, without supplied curvatures or amplitudes. In particular they identify every finite smooth contact jet. These exact uniqueness conclusions require no closeness between the candidates. For actual connected closed analytic obstacle boundaries, the same laws identify the entire labelled obstacle images when the lattice is fixed and every obstacle is visited by the polygon.

On bounded positive  $C^{m+3}$  contact classes, for a finite order  $m$  depending on the marked contraction bounds, small  $C^m$  real-law error controls the  $C^0$  error of the complete contact profiles on a common real collar. No analytic prior is used in this estimate. At each analytic contact tuple the complete future-action map is also a local biholomorphism in a bounded-holomorphic norm on a sufficiently small complex disc. A bounded local analytic prior gives conditional Hölder stability from real-law error on a smaller disc. At each separately fixed smooth order, finite-flight density error gives the finite-jet error  $C_M(\tau^N + \varepsilon)$  on compact positive classes. If the curvature algorithm is stopped with certified error  $E_m$ , the bound is  $C_M(\tau^N + \varepsilon) + L_M E_m$ , with the subsequent signed recursion evaluated as in Proposition 4.4.

Under the uniform finite-order geometric assumptions of Section 6, the same marked experiment also admits finite-sample recovery. For fixed confidence,  $B$  independent preparations per phase-offset group, and endpoint recording error at most  $\delta$ , a realizable estimator satisfies

$$\|\hat{F} - F\|_{C^0} \leq C\{(\log B/B)^\beta + \delta^\gamma\}$$

with the stated high probability, for positive class-dependent exponents  $\beta, \gamma$  and a logarithmically chosen returning flight number. The exact confidence dependence, priors, estimator, and exponents are given in Theorem 6.3. Failures count toward  $B$ ; no density-derivative observations are assumed.

*Proof.* Theorem 2.3 supplies the relative limit; Theorem 2.5 and Lemma 3.1 put it in the physical measured coordinates. Proposition 3.2 separates the unequal endpoint actions. The global quadratic inverse of Theorem 4.2, followed by the actual signed recursion, proves Theorem 4.3. The nonlinear real remainder estimate in Theorem 5.3 then identifies smooth germs, and Theorem 5.5 gives the finite-smoothness profile estimate. The complete analytic inverse and conditional real-law estimates are Theorems 3.6 and 3.8. Proposition 4.4 proves the fixed-order



finite-flight assertion. The finite-preparation conclusion is Theorem 6.3, using the realizable selection and interior density estimate in the same section. All these proofs are included below.  $\square$

The curvature step explains why exact identification is not confined to one analytic inverse neighborhood. If  $k_i = L_i^{-1}$ ,  $c_i = F_i''(0)$  and  $s_i = (S_i^-)''(0)$ , the boundary Schur complement is

$$s_i = k_i + c_i - \frac{k_i^2}{k_i + c_{i+1} + s_{i+1}}.$$

Rearranged as an equation for  $c$ , it is a strict contraction on the positive orthant after taking positive parts. Its constant is bounded using  $k_i$  and the observed  $s_i$ , not unknown curvature values. The unique positive fixed point determines the signed orbit factors. At degree  $n$  the actual boundary-visit response is then  $(I + T_n)(I - T_n)^{-1}$ , with  $(T_n z)_i = \sigma_i^n z_{i+1}$  and  $|\sigma_i| < 1$ . This identifies each successive contact jet. A finite low-degree inverse together with a norm-convergent high-order tail, rather than coefficient uniqueness alone, gives the analytic norm theorem.

For smooth germs a further use of the complete action is essential. Along the segment between two candidate graphs, the actual envelope has the form

$$\mathcal{S}(G) - \mathcal{S}(F) = \bar{\alpha}(G - F) + K(G - F), \quad \bar{\alpha} \geq \frac{1}{2}.$$

Every term of  $K$  evaluates a graph difference at a later contact with  $|x_{b,j}(u)| \leq a^j |u|$ ,  $a < 1$ . In the real norm  $\|h\|_{m,R} = \max_b \sup_{0 < |u| \leq R} |h_b(u)|/|u|^m$ ,

$$\|\bar{\alpha}^{-1}K\| \leq \frac{6a^m}{1 - a^m} < 1$$

for sufficiently large finite  $m$ . Equal finite jets put the difference in this space, and equality of the actual actions forces it to vanish. Thus the relative law distinguishes even smooth flat contact changes. A finite-jet alignment gives the stated real-profile stability. The weighted comparison is proved along actual graph segments, not assumed for a formal map or presented as a data-only reconstruction algorithm.

**A.1.1. The mechanism and the observation models.** There are three logically different passages. The cofactor identity and two-ended trace-class comparison control the relative physical law before conditioning. The stationary boundary envelope, including its vanishing terminal term, converts its nonlinear actions into contact geometry. Registration and analytic continuation then assemble geometry in the appropriate observation model. The latter passage does not replace the first two. The obstruction in the forward problem is a scale mismatch:  $D_{N,b}$  is exponentially small, so division of an absolute action remainder by  $D_{N,b}$  gives no relative twist estimate. In the proof, the corner cofactor is divided by its reference value *exactly*; the logarithm is then an edge sum minus  $\log \det(I + G_N \Delta H_N)$ . The perturbation has a trace norm bounded by the two endpoint displacements, independently of the number of flights. Reflected Green estimates separate its two ends, and the trace-series telescoping estimate passes to the half-line determinants. This establishes (A.1.1) before probability normalization. It is not an application of weak convergence to a sequence of already conditioned quadratic models.

The inverse has a different obstruction. An abstract generating function need not come from the stated reflecting curves. The stationary envelope is differentiated along actual graph variations: its initial visit has weight one, every interior visit has weight two at the orbit, and its terminal term is estimated before the half-line limit is taken. Interpolation of smooth representatives then shows that the action jet depends on precisely the corresponding finite contact jet. The cyclic blocks are the resulting geometric derivative, not a postulated formal normal form. The full envelope also controls the remainder beyond those jets. Global curvature inversion, the real contracted-visit estimate and, for complete analytic obstacles, continuation complete this relative-law/contact mechanism; none replaces the relative physical estimate.

There is a concrete information comparison within the physical class. The area-preserving analytic family of Theorem A.3.2 has the same gap, free area, contact curvatures and leading endpoint and count data, but different fourth contact jets. The fixed-positive-offset laws distinguish its members by the signed action inverse. Thus the complete conditional law carries strictly more geometric information than that leading record. This example compares those two observations; it makes no assertion about the information of a marked length spectrum.

The smooth class exhibits a sharper distinction. Proposition 5.6 constructs actual area-preserving obstacle perturbations with identical *complete* contact and action Taylor jets, but distinct boundary laws on every sufficiently small positive contact square. The first nonzero real action response is flat, of size  $\varepsilon e^{-1/u^2}(1 + o(1))$ . Consequently no formal all-order contact record can replace the function-valued physical law in the smooth rigidity theorem. This compares specified observations in an actual billiard class, not the full marked-length datum of a different inverse problem.

The amplitudes also retain nonlinear dynamical coordinates. With  $\zeta_b^\pm(x) = \int_0^x B_b^\pm(t) dt$ , the actual stable returns satisfy

$$T_b^\pm = (\zeta_b^\pm)^{-1} \circ (e^{-\chi} \zeta_b^\pm)$$

on a smaller collar, by Corollary 2.4. The two-offset laws recover  $B_b^\pm$ , and the reconstructed quadratic Jacobi data recover  $\chi$ . Hence the same relative record determines these nonlinear returns, not just their linear multiplier. This is a consequence of the proved transmission identity within the marked observation model, not a claim of a reconstruction from lengths alone.

The alternating channel uses less supplied first-order geometry: its gap and signed transverse conventions replace a full marked polygon, and one same-type law at a fixed positive offset suffices at each end. Its normal two-site response is the specialization of the signed cyclic response, but its law-to-action extraction is a different theorem. Its unregistered multichannel reconstruction can recover an unknown lattice. In the marked-polygon conclusion above, in contrast, lattice and placement are supplied, and full visitation turns exact contact identification into whole-boundary identity. No implication that removes these supplied marks is used. The charged finite-preparation results require their calibration and sensor bounds in addition to the geometric inverse; the earlier alternating-channel rates are not being transferred to a different observation model. Section 6 proves a general fixed-period estimate directly. An interior kernel estimate supplies the required  $C^m$  density control from finitely many recorded endpoints. A measurable minimum-distance comparison stays inside the realizable image, and a binomial lower-tail estimate accounts for the exponentially small success probability. This passage is a consequence of the relative law and smooth contact inverse, not an independent claim of a new general density-estimation method. It also gives a finite charged test between the actual equal-jet flat tables at any fixed nonzero profile separation.

The nearby spectral questions concern different observations. Bálint, De Simoi, Kaloshin and Leguil [1], Theorem D and Corollary E, extract general-period Lyapunov information from marked length asymptotics; their Remark 2.3 explains the contact-separation issue for those asymptotics. Here phase-resolved functions and a supplied polygon provide individual action responses. We do not infer them from a marked length spectrum. The enriched spectrum of Finamore–Leguil [12] and the analytic symmetric open-billiard setting of De Simoi–Kaloshin–Leguil [5] likewise have their own data and hypotheses. The exact global uniqueness proved here is for the stated marked boundary-law observation.

The periodic relative and local coordinate statements are recorded next. Their exact local uniqueness clauses are strengthened by Theorem 4.3; their analytic norm estimates retain their stated neighborhoods, priors and radius loss. The subsequent alternating formulations preserve the weaker-data results and their complete proofs.

The smooth contact conclusions in Section 5 use the marked two-offset model just stated. The formulations below retain the separate analytic, unregistered and finite-observation consequences with their own hypotheses and complete proofs.

**A.1.2. Relative transmission along an arbitrary periodic itinerary.** The forward mechanism extends to every selected nongrazing periodic orbit with clear flights and strictly dispersing contacts. Its proof uses the same two-ended normalization, now with periodic Jacobi coefficients and distinct future and past boundary actions. We give the geometric verification rather than assuming a hyperbolic normal form. A three-obstacle orbit and its nonsymmetric perturbations provide realizations outside the normal alternating setting (Proposition 2.6).

Here is the normalization which makes the extension precise. Let  $P$  be an even oriented period,  $L_i$  the flight lengths,  $\kappa_i$  the curvatures, and  $\nu_i$  the outgoing normal cosines. Set  $k_i = L_i^{-1}$ ,  $m_i = 2\kappa_i/\nu_i$  and

$$Q_i = \begin{pmatrix} 1 + m_i/k_i & 1/k_i \\ m_i & 1 \end{pmatrix}, \quad \mathcal{M}_b = Q_{b+P-1} \cdots Q_b, \quad \text{tr } \mathcal{M}_b = 2 \cosh \chi.$$

Positive curvature gives  $\chi > 0$ . For  $N = nP$  flights starting at phase  $b$ , the exact reference twist is

$$D_{N,b} = \frac{\sinh \chi}{(\mathcal{M}_b)_{12} \sinh(n\chi)}.$$

In scaled signed arclength coordinates, let  $p_b$  be the reference endpoint momentum and  $L = \sum L_i$ . Theorem 2.3 proves, on a collar independent of  $n$ ,

$$\begin{aligned} W_{N,b} - nL + p_b u - p_b v &= S_b^-(u) + S_b^+(v) + O_{C^k}(\tau^N), \\ -W_{N,b,uv}/D_{N,b} &= B_b^-(u)B_b^+(v) + O_{C^k}(\tau^N). \end{aligned} \tag{A.1.2}$$

Each differentiation order is fixed separately. The future and past actions need not agree. The physical window law contains  $d - S_b^-(u) - S_b^+(v) + p_b u - p_b v$ ; the endpoint momentum terms cannot be discarded as though the reference collisions were normal. Theorem 2.5 derives this law from normalized phase volume and retains its exact exponentially small acceptance mass.

The same estimate also identifies the nonlinear coefficient of the weak interaction between the ends. If  $C_N = W_N(u, v) - W_N(u, 0) - W_N(0, v) + W_N(0, 0)$  and  $\zeta_b^\pm(x) = \int_0^x B_b^\pm(t) dt$ , then

$$-C_N/D_{N,b} = \zeta_b^-(u)\zeta_b^+(v) + O_{C^k}(\tau^N).$$

The error before normalization is  $O(D_{N,b}\tau^N|uv|)$ . Moreover  $\zeta_b^-$  linearizes the future stable period map, while  $\zeta_b^+$  linearizes the reversed unstable period map, both with multiplier  $e^{-\chi}$  (Corollary 2.4). Scalar linearization and the final double integration are consequences of the relative theorem; they are not substitutes for its geometric determinant estimate.

This extension explains a structural feature of the mechanism: it requires dispersing coercivity, a clear nongrazing itinerary and summable endpoint layers, rather than equal flight lengths or a normal two-orbit. The alternating inversion below exploits the additional geometry of a two-contact pair. All of its complete-jet, analytic-germ, global and finite-observation conclusions are retained. No two-contact inverse is silently applied to the distinct actions of a general periodic itinerary.

**A.1.3. The periodic law as a geometric coordinate.** The two-ended law has a further inverse consequence beyond normal alternating channels. Fix a clear nongrazing periodic polygon with  $r$  distinct marked contacts. Its collision points, tangent lines, phase labels and closing translation are given; its contact curvatures and normal graph germs are unknown. At every physical phase retain the conditional law of the two scaled tangent projections at two fixed,

distinct positive offsets. These are function-valued data on a common interior square, not samples of the unknown normal graph.

**Theorem A.1.2** (Periodic contact reconstruction). *For the fixed marked polygon, the phase-resolved two-offset laws locally determine every finite contact jet in the smooth category, including curvature. At an analytic base tuple they determine the complete contact germs, and the future action map is a local biholomorphism in a bounded-holomorphic norm on a small fixed disc. Under a bounded local analytic prior, real uniform density errors give conditional Hölder bounds on a smaller disc. If a fixed-lattice analytic table has every labelled obstacle visited by the polygon, equality of these laws in the local inverse class determines all of its obstacle images.*

The separation of the two actions does not assume that they agree. Writing the limiting density as

$$f_d(u, v) = Z_d^{-1} B^-(u) B^+(v) \{d - A(u) - C(v)\}, \quad A(0) = C(0) = 0,$$

the quotient  $Q = f_{d_1}(u, v) f_{d_2}(0, 0) / [f_{d_2}(u, v) f_{d_1}(0, 0)]$  gives

$$A(u) + C(v) = \frac{d_1 d_2 (1 - Q)}{d_2 - d_1 Q}.$$

Both flux factors and both normalizers cancel. The oblique linear momenta are then restored when passing from these physical actions to the anchored variational actions.

There is also a geometric reason why all contact orders are invertible. Let  $\sigma_i$  be the signed one-step derivative of the future stable half-line in normalized graph coordinates. Dispersing positivity gives  $|\sigma_i| < 1$ . At contact order  $n$ , the response is

$$\mathcal{C}_n = (I + T_n)(I - T_n)^{-1}, \quad (T_n z)_i = \sigma_i^n z_{i+1}. \quad (\text{A.1.3})$$

This formula is obtained by counting the initial boundary visit once and each internal visit twice in the actual stationary envelope. It is not an identification of a formal series with a smooth action. The block has determinant one for an even cycle; the explicitly computed odd-cycle determinant is different but nonzero, including at odd contact orders. The derivative of the full action is an invertible multiplier plus contracted boundary evaluations. A finite low-degree inverse and one operator-norm tail inverse control all its lower-degree couplings.

*Proof of Theorem A.1.2.* Proposition 3.2 separates the physical actions. Lemma 3.3 proves smooth-jet factorization, and Theorem 3.4 proves (A.1.3) and the local finite-jet inverse, beginning with its invertible curvature block. Lemma 3.5 constructs the actual Banach-valued action map. Its complete inverse, real-observation estimate and exact analytic continuation consequence are respectively Theorem 3.6, Theorem 3.8 and Corollary 3.7. All proofs are given in Section 3. A nonnormal three-contact family with independently varying curvatures is constructed in Proposition 3.9.  $\square$

The relationship with variational and spectral methods is specific. Hill's formula [2] concerns periodic Hessians and monodromy; our forward normalization uses a Dirichlet corner cofactor before conditioning. The general-word length asymptotics and marked Lyapunov recovery of Bálint, De Simoi, Kaloshin and Leguil [1] concern periodic-orbit lengths, rather than endpoint-varying probability laws. Here the fixed marked polygon and the phase-resolved function data make a separate contact inverse possible. We do not identify these data with an ordinary or enriched marked length spectrum. In particular, the enriched-spectrum rigidity theorem of Finamore and Leguil [12] is a different observation model, not a claimed consequence of the present theorem. The matrix identity and the functional-analytic inversion argument are not advanced as separate novel rigidity theories; their geometric realization in the normalized physical law and the actual contact action is the point of the argument.

The following alternating-channel results retain a different advantage: they require only the gap and one law of each contact type, rather than a fixed general polygon and two offsets per phase. They also retain the unknown-lattice reconstruction, finite global ambiguity and charged finite-preparation results under their original hypotheses. The new periodic contact theorem extends the inverse domain without changing those data requirements or replacing their conclusions.

**A.1.4. The alternating channel and its geometric content.** We study the geometry retained by conditional laws of long alternating bridges in a periodic dispersing billiard. The observation records the signed transverse endpoints, but not their normal coordinates. Although the probability of a successful bridge decreases exponentially with its length, the conditional law retains two nonlinear half-line actions. We recover those actions without knowing the flux amplitudes, invert every finite signed contact jet, and reconstruct the entire analytic table and its marked period lattice. Finite rotations of individual obstacles produce only finitely many global matching alternatives. They produce no additional infinitesimal ambiguity when every obstacle is noncircular.

The principal mechanism is a relative nonlinear limit on a collar which does not shrink with the flight number. A quadratic approximation at the periodic orbit would lose the higher contact terms. Moreover, a formal inverse of the action series would not justify the inverse for actual smooth boundaries. The proof therefore combines a weighted half-line construction, trace-norm control of the normalized twist, and a finite Taylor-remainder factorization before any analytic continuation is used. The statistical results concern how this geometric datum can be acquired and how its information changes when further observations are retained.

The analytic germ problem admits a further conclusion for the complete nonlinear map, not just its diagonal blocks. In Theorem 22.12, the half-line actions are local analytic coordinates for anchored analytic contact pairs in a bounded-holomorphic norm on one small fixed disc. The derivative is

$$(D\mathcal{S}(\psi)\eta)_b(u) = \alpha_b(u)\eta_b(u) + \sum_{i \geq 1} v_{b,i}(u)\eta_{b+i}(x_{b,i}(u)), \quad |x_{b,i}(u)| \leq \rho^i|u|, \quad \inf |\alpha_b| > 0.$$

This identity follows from the actual stationary envelope, including its vanishing terminal term. The first term is an invertible multiplier; the remaining evaluations contract high-order vanishing functions. One finite low-jet system and a norm-convergent tail inverse therefore control all lower-degree couplings. A local contraction gives the nonlinear inverse. In the induced analytic quotient norms, its complete linearized finite-jet inverses have bounds independent of the truncation order. The analytic norm is an additional hypothesis: it is not inferred from finitely many density derivatives, total variation, or a finite sample. No continuation to an entire obstacle or uniform acquisition rate is part of this local analytic estimate.

There is also a conditional conclusion in the real observation topology. Theorem 22.15 fixes a local bounded analytic neighborhood and recovers the complete contact germs on a strictly smaller complex disc from uniform real endpoint-density error  $\epsilon$ , with a bound  $C(\epsilon_g + \epsilon^\vartheta)$ ,  $0 < \vartheta < 1$ . Curvatures may vary and are not separately supplied. With a uniform first-order real density bound, total variation or histogram errors give exponent  $\vartheta/3$ . The proof combines amplitude-free extraction on a real square, a displayed interpolation estimate, and the complete analytic inverse on the inner disc. It does not infer same-disc analytic smallness from real data: the outer analytic prior, radius loss, density and anchor margins, and local choice of inverse branch are explicit. No global registration, unknown-origin calibration, or preparation budget is a consequence of this local estimate. Proposition 22.13 isolates the elementary functional-analytic criterion; the protected half-lines and stationary envelope remain its billiard-specific inputs.

**A.1.5. The observation and the main theorem.** A table  $T$  consists of  $N$  labelled strictly convex analytic obstacles  $K_1, \dots, K_N$  and their translates by  $L\mathbb{Z}^2$ , where  $L \in \mathrm{GL}^+(2, \mathbb{R})$  is unknown. Distinct lifted closures are disjoint and all curvatures are positive. The basis of the abstract deck group  $\mathbb{Z}^2$  is marked; its Euclidean realization is not. A selected clear channel  $e = (a, b, \ell)$  joins the closest pair of  $K_a$  and  $K_b + L\ell$ . Each channel has its own oriented contact frame and signed transverse coordinate, with no inter-channel registration. The full geometric conventions are stated in Section 12 and Section 26.

For each channel retain its gap  $g_e$  and, for each starting contact type  $b \in \{0, 1\}$ , the limiting conditional joint law  $P_{e,b}$  of the two signed transverse endpoints of an even bridge at one known positive excess time  $d_e$ . The intrinsic datum is

$$\mathcal{D}_E(T) = (g_e, P_{e,0}, P_{e,1})_{e \in E}. \quad (\text{A.1.4})$$

Obstacle, channel and deck marks and the signed conventions are design information. No longitudinal position, residual time, waiting count, curvature, flux amplitude, Euclidean deck vector or inter-channel pose is supplied. A law is a function-valued datum, not one scalar sample. The gap is retained separately from the conditional laws.

Let  $\mathcal{C}_N$  be the class in which every obstacle is noncircular, and let  $\mathcal{A}_N \subset \mathcal{C}_N$  be the class in which every obstacle has trivial proper Euclidean symmetry. The latter is open and dense in the relative  $C^2$  analytic-support topology (Proposition 29.6). A *common-strip analytic-support family* is  $C^1$  into the Banach space of bounded holomorphic support functions on one fixed complex strip, with  $C^1$  lattice columns; see (31.1). In particular the parameter derivative is itself analytic on a smaller strip. The differential statements below do not merely assume that each parameter slice is analytic.

**Theorem A** (Single-offset periodic rigidity). *For each  $T_0 \in \mathcal{C}_N$  there is a persistent selected design  $E$  of  $N + 1$  channels, consisting of an obstacle spanning tree and two independent marked cycle gains, and a common sufficiently small positive offset, with the following properties.*

- (1) *The datum  $\mathcal{D}_E(T_0)$  determines a finite set of complete labelled periodic tables, including their marked lattices, modulo one common proper Euclidean motion. This set is obtained by finite choices of incidence congruences followed by explicit lattice reconstruction and geometric compatibility tests. If  $T_0 \in \mathcal{A}_N$ , it consists of one table.*
- (2) *For a common-strip analytic-support variation through  $T_0$ , if all gap derivatives and all endpoint-density derivatives on contact-centered positive interior squares vanish, the table variation is one common infinitesimal proper Euclidean motion. Conversely such a motion has zero intrinsic derivative. In a fixed Euclidean gauge the data derivative is injective.*
- (3) *On every immersed  $p$ -dimensional  $C^1$  common-strip analytic-support model in that gauge,  $p \geq 1$ , one can select  $p$  scalar functions, each a gap or an expectation  $\int \phi dP_{e,b}$  with  $\phi \in C_c^\infty$  supported in a positive interior square, which form local  $C^1$  coordinates at  $T_0$ . Their inverse is locally Lipschitz, with local control of every fixed real support norm and of the lattice matrix.*

*Proof.* Theorem 29.3 supplies the persistent design. The signed law inverse gives complete analytic channel-frame images. Theorem 30.3 enumerates the finite matching alternatives and reconstructs their lattices; the asymmetric uniqueness conclusion is also Theorem 29.7. Theorems 31.5 and 31.6 prove the second and third assertions. Their local registration uses Lemma 30.2, not a choice of one globally preferred symmetry branch.  $\square$

The finite bound in the first assertion is a product of orders of obstacle rotation groups along a spanning tree of channel incidences ((30.3)); it is not a sharp degree assertion. The  $N + 1$

count is minimal for the stipulated connected two-cycle architecture, not an information lower bound over arbitrary experiments. The selected design and its marks are not discovered by an unmarked trajectory algorithm. The scalar coordinates in the third assertion may depend on the base table and model; finite dimensionality is not a restriction on the first two assertions.

**A.1.6. Unknown origins and separate recording efficiencies.** The contact-centered formulation isolates the geometric inverse. Its origins can themselves be recovered from the law, even when its two endpoint amplitudes are independently distorted. Suppose the recorded coordinates are  $x = u + \xi$  and  $y = v + \eta$ , with unknown constants  $\xi, \eta$ , and successful records are retained with a positive factor  $e_1(x)e_2(y)$ . The physical signed units and directions are unchanged; neither an unknown clock nor an unknown coordinate scale is included in these constants. Section 24 proves the following extension using the interaction of the two endpoints.

**Theorem B** (Rigidity with unknown transverse origins and weights). *In the geometric setting of Theorem A, retain the same marks, gaps and known positive offsets, but replace every ideal same-type law by its translated law with unknown positive smooth separate endpoint recording factors. Assume that the complete positive cap is recorded and that the numerical coordinates have the specified signed transverse units and directions. Then the laws determine their two transverse origins and the signed action germs. The compatible table fiber is exactly the one in Theorem A(1), including the unknown marked lattices. In a differentiable table-and-recording family of the stated smooth class, no table variation outside the common Euclidean tangent space can be masked by varying those origins or separate recording factors.*

*Proof.* Lemma 24.1 extracts the origins from the zero fibers of  $\partial_x \partial_y \log \tilde{f}$ . Theorem 24.2 then recovers the action by its scalar-anchored four-density ratio. Proposition 24.3 justifies the finite-order differentiable extraction, and Theorem 24.4 proves the exact-fiber and joint derivative kernel assertions through the smooth contact inverse and periodic cochain.  $\square$

The cancellation underlying this extension is a classical association invariant [40, 41]. What the boundary-law mechanism supplies is its specific form

$$\partial_x \partial_y \log \tilde{f} = -\frac{S'(x - \xi)S'(y - \eta)}{[d - S(x - \xi) - S(y - \eta)]^2}.$$

Strict convexity locates the two origins, and the signed contact inverse turns the recovered action into geometry. No equality of the two recording profiles is required. This is an exact-law reconstruction with a local finite smooth stability estimate, not a claim that a fixed finite vector of noisy observations or the old preparation budget suffices for arbitrary unknown recording efficiencies. Theorem A(3) remains a separate statement on its stipulated finite-dimensional ideal-law model.

For the complete-cap observation in this statement, the origins can also be recovered without a density derivative: the vertical and horizontal fiber widths of the positive support have their unique maxima at the two origins (Lemma 24.5). Thus centering a complete cap is not, by itself, a new geometric rigidity mechanism. The benefit of the interaction construction is its dependence on a fixed positive interior rectangle. Theorem 24.6 makes this locality explicit: a known rectangular recording window containing both unknown origins, compactly inside the positive cap, suffices even when the cap boundary and discarded mass are not observed. Its positive support is the same window for every admitted law. The signed action is nevertheless recovered from the interior density, and Corollary 24.7 gives exactly the same analytic table fiber and projected derivative kernel. The visibility of the two origins remains a hypothesis; there is no unmarked search for that window.

The separate-factor cancellation is already the calculation of Theorem 23.1, with no equality of the two factors needed. The interior extension uses this algebra rather than inventing another association invariant. Its regularity is the local  $C^{m+3}$ -to- $C^m$  estimate of Proposition 24.3, now on a window not reaching the cap boundary. A complete support can center the law but need not determine its action; the functional comparison in Proposition 23.13 isolates that distinction. None of these statements provides a total-variation estimate or a uniform preparation bound for an unconstrained recording class.

**A.1.7. Proof of the geometric inverse.** The proof follows one chain. For a bridge with stationary length  $W_j$  put  $E_j = W_j - jg$  and  $b_j = -W_{j,uv}/[-W_{j,uv}(0,0)]$ . On a flight-independent collar, Theorem 17.2 proves

$$(E_j, b_j) \longrightarrow (S_0 \oplus S_p, B_0 \otimes B_p), \quad p = j \bmod 2,$$

in each fixed differentiated norm with an exponential relative error. The twist is normalized before it is compared; no absolute error is divided by an exponentially small quantity. The common-domain sublevel integral then gives the endpoint law of Theorem 17.3.

The scale of the mixed derivative is the obstruction to a direct absolute-error argument. With  $a_b = S_b''(0) > 0$  and  $q_p = \sqrt{a_0 a_p}$ , the quadratic reference is

$$d_j^0 = -W_{j,uv}(0,0) = \frac{q_p}{\sinh(j\gamma)}. \quad (\text{A.1.5})$$

An absolute remainder  $O(\theta^j)$  would give only  $O((\theta e^\gamma)^j)$  after division by this reference; its decay is not automatic. Instead the proof uses the exact cofactor identity

$$\log b_j = \sum_{i=0}^{j-1} \log \{g[-\ell_{i \bmod 2, uv}(y_i, y_{i+1})]\} - \log \det(I + G_j \Delta H_j). \quad (\text{A.1.6})$$

Here  $y$  is the stationary bridge,  $G_j$  is the inverse reference interior Hessian, and  $\Delta H_j$  is the subtracted nonlinear Hessian. The trace-norm estimates sum the localized endpoint layers, whose weights have bounded sum independent of  $j$ , rather than multiplying an unweighted entry bound by the number of collisions. Two separated compressions of the reference Green operator then converge to the half-line operators. The trace-series comparison in (17.11)–(17.12) yields a relative amplitude estimate before the exponentially small scale is removed. The long-flight limit is taken at fixed positive excess  $d$ , not simultaneously with a zero-offset quadratic approximation. This is why it retains nonlinear actions rather than only curvature.

For a same-type law, on a positive interior square,

$$f(u, v) = Z^{-1} B(u) B(v) \{d - S(u) - S(v)\}.$$

The four-density identity of Theorem 23.1 is

$$1 - \frac{f(u, v) f(0, 0)}{f(u, 0) f(0, v)} = t(u) t(v), \quad t(u) = \frac{S(u)}{d - S(u)}.$$

One nonzero scalar anchor fixes the positive normalization of  $t$ ;  $S = dt/(1+t)$  is recovered on the signed interval. The inverse does not take a pointwise square root at the minimum. The gap and the two quadratic action coefficients recover the contact curvatures. Theorem 22.1 then recovers each higher labelled graph coefficient through an invertible two-by-two block after subtracting lower-order terms. Lemma 22.4 proves that this is an identity for actual smooth representatives, including ones differing by flat terms. Odd contact terms remain part of the inverse.

The passage from smooth graphs to finite jets is a separate analytic step in this inverse. For two actual graph pairs with the same gap and the same  $M$ -jet, interpolate the graphs and differentiate a finite stationary action sum. Interior orbit variations cancel; the terminal term



is retained until the weighted half-line estimates make it tend to zero. Integrating this finite identity in the interpolation parameter before taking its limit gives

$$|S_b^{[1]}(u) - S_b^{[0]}(u)| \leq \frac{C_M}{1 - \rho^{M+1}} |u|^{M+1}. \quad (\text{A.1.7})$$

Thus the finite action jet factors through the finite graph jet independently of the smooth representative. At degree  $n \geq 3$ , a boundary visit is counted once and every interior visit twice. The geometric sums give the explicit last-jet block

$$M_n = \begin{pmatrix} \coth(n\gamma) & \mathbf{r}_0^n \operatorname{csch}(n\gamma) \\ \mathbf{r}_1^n \operatorname{csch}(n\gamma) & \coth(n\gamma) \end{pmatrix}, \quad \mathbf{r}_0 \mathbf{r}_1 = 1, \quad \det M_n = 1. \quad (\text{A.1.8})$$

The positive factors  $\mathbf{r}_b$  are the alternating type ratios of the linear half-line orbit in (22.22). Write  $q_n = (\psi_0^{(n)}(0), \psi_1^{(n)}(0))^t$  for the two anchored graph coefficients and  $s_n = (S_0^{(n)}(0), S_1^{(n)}(0))^t$ . After the lower-order contribution  $R_n$  has been computed, the recursion is  $q_n = M_n^{-1}(s_n - R_n)$ . All these operations are performed at a fixed finite order, with the required functional boundary bounds. No infinite-order conditioning estimate is inferred from the determinant. Equality of every smooth contact jet is not identified with equality of arbitrary smooth boundary germs; the subsequent whole-image assertion uses analyticity.

Analytic continuation now gives complete channel-frame obstacle images. For a noncircular obstacle its proper symmetry group is finite. Matching all channel frames therefore gives finitely many candidate orientations, not a continuous family. For each candidate, two marked cycle translations  $v_1, v_2$  determine the unknown lattice by

$$L = (v_1 \ v_2) M^{-1},$$

with the actual fixed integer gain matrix  $M$ ; unimodularity is not assumed. The compatibility tests in Theorem 30.3 retain precisely the realizable candidates. On  $\mathcal{A}_N$  all congruence choices are unique.

For the differential assertion the same chain is differentiated at each fixed finite order. Lemma 31.1 includes the moving-reference amplitude derivative and the moving-cap normalizer. Zero law variation then gives zero signed contact-jet variation. Analyticity of the variation makes each entire channel-frame image stationary. A single nonzero harmonic determines an angular velocity along a local alignment branch, even when the body has finite symmetry. The lattice cochain then eliminates an invisible change of scale or Gram form. Thus the differential assertion uses noncircularity, whereas the global unique matching assertion uses trivial proper symmetry. The circular comparison in Proposition 30.4 exhibits the distinction for this selected-channel design. Proposition 30.5 exhibits a different phenomenon: a noncircular  $C_4$ -symmetric obstacle with exactly two realizations of one two-channel datum and distinct marked Gram forms. This is finite global ambiguity despite an injective local derivative, not a circular exception to the theorem. Proposition 30.6 proves that one further marked clear-channel gap distinguishes those two realizations. Its clearance is proved for the actual closest segment, not assumed from the direction of its center displacement.

**A.1.8. Acquisition, information and related inverse problems.** The deterministic inverse from finite transverse histograms to complete tables is Theorem 50.5. Its quantified analytic, geometric and asymmetry bounds are separate from the hypotheses of Theorem A. Finite contact order is chosen before observational errors are reduced; qualitative analytic uniqueness is not used as a bounded inverse for noisy continuation. Conditional analytic continuation is discussed in [37].

The acquisition theorem uses a richer calibration sensor: noiseless planar endpoint positions, an exact physical clock and preparation outcomes. A same-final-flight pilot is followed by the

coarser transverse-histogram inverse. The proof retains separate finite-flight, flight-amplified timing and internal/outer cell-crossing errors, with all failed preparations charged. It is not calibration from histograms alone. The information results in Part II distinguish endpoint, endpoint-residual-time, count-endpoint and planar-position records; none is a premise of the deterministic periodic inverse.

The common object in these three parts is the unnormalized relative boundary measure. For a same-type law, suppressing the channel index, write

$$\nu_{b,d}(du dv dr) = B_b(u)B_b(v)\mathbf{1}_{\{0 < r < d - S_b(u) - S_b(v)\}} du dv dr. \quad (\text{A.1.9})$$

Its normalization gives the endpoint-residual-time law; integration in  $r$  gives the endpoint law. The preparation probability retains information from the normalization and the exponentially small reference factor that conditional normalization removes. The relative estimates of Theorems 17.2 and 17.3 justify these operations for the actual billiard, not just for an abstract family of cap densities. Part I inverts the geometric interaction in this measure. Part II studies the different likelihoods obtained by retaining or forgetting components of the physical record. Part III controls the passage from finite acquired records to the ideal law and then through the deterministic inverse. The same forward construction is required at all three interfaces; the inverse alone does not supply their statistical hypotheses.

These links specify the scope of the combined treatment. Exact function-valued reconstruction, local interior smooth inversion, model-local finite scalar coordinates and charged finite-sample acquisition are not equivalent experiments. Separate endpoint reweighting preserves the geometric invariant, but generally changes the likelihood and probability of recording a bridge. Cropping to an interior window preserves the contact-germ inverse under its visibility conditions, but does not provide the unrecorded mass. In particular, multiplying both recording probabilities by  $\varepsilon$  keeps the conditional law fixed while multiplying the recorded-success probability by  $\varepsilon^2$ . The success lower bound and density margins used in Part III are therefore additional premises, not consequences of Theorem B or its window extension. Likewise, the finite scalar coordinate theorem concerns a stipulated finite-dimensional ideal-law model, not arbitrary infinite-dimensional efficiency profiles. The structural contribution is the nonlinear relative/smooth inverse; the other parts establish, for their stated records, what is lost by coarsening and what further control permits its acquisition.

A quantitative example ties these distinctions to an actual geometric family. The small-window benchmark in [43] is realized in Theorem 33.2 by noncircular analytic obstacles with fixed lattice and fixed channel gaps. Separate recording profiles, fixed as the half-width  $h$  shrinks, preserve exact identification at each  $h > 0$  while the one-record squared Hellinger distance between two fixed shape alternatives is of order  $h^8$ . For that two-point experiment the successful-record scale  $n \asymp h^{-8}$  is attainable: the critical likelihood limit and a bounded-moment test are both explicit. This is not an optimal uniform estimator over the unrestricted nuisance class. Corollary 33.3 transfers the experiment to actual finite flights with total-variation error  $C\sqrt{n}\tau^J$ , under  $n\tau^{2J} \rightarrow 0$ . The bounded-moment test has a direct exponential error bound whenever  $\tau^J = o(h^4)$ , without requiring product approximation in the supercritical regime. Its normalizing integrals also show an additional recording fraction of order  $h^2$ , separate from the original bridge rarity. For this same pair, Theorem 33.4 shows that one first-acceptance count consistently distinguishes the tables, while its accepted mark alone becomes uninformative: the reverse binary deficiency tends to  $1/2$ . At a deterministic cap  $b_h$ , the optimal error tends to  $e^{-\lambda}/2$  when  $b_h p_{0,h} \rightarrow \lambda < \infty$ , and tends to zero exactly when  $b_h p_{0,h} \rightarrow \infty$ . Here  $p_{0,h} \asymp h^2 e^{-\gamma_0 J_h}$  is the actual acceptance probability. These charged and conditional conclusions concern different records of one specified experiment, not competing universal complexity bounds. Thus the relative forward theorem controls a joint flight/window/sample limit; the exact inverse and the acquisition cost cannot be

interchanged. The table actions in this construction are actual nonlinear actions, not stipulated quadratic functions.

The stopped-record conclusion has a precise reduction within this same pair. Censored waiting counts are asymptotically sufficient uniformly over deterministic caps, with an explicit absolute risk error and an exact count-likelihood threshold. At finite cap intensity the full record converges, in Le Cam distance, to a Bernoulli experiment; at much larger caps the acceptance bit can lose all discrimination while the count remains informative. These reductions are proved in Theorem 33.4. They identify the retained mass, rather than the terminal interaction, as the source of stopped separation. They do not identify the stopped experiment with the conditional-law inverse or supply a sensor-independent preparation bound.

Finamore and Leguil [12] study enriched marked length data for finite-horizon Sinai billiards. De Simoi, Kaloshin and Leguil [5] study marked lengths for analytic open billiards under non-eclipse, symmetry and genericity hypotheses. Neither datum is (A.1.4), and no reduction between these observation maps is used here. The fifth version of Florio and Leguil [11] retains dynamical conjugacy results after removing an earlier geometric spectral-rigidity assertion affected by an error. We use neither that removed assertion nor an identification of conjugacy with the present boundary-law inverse.

Zelditch's inverse spectral theorem provides a useful orbit-local comparison [49]. Theorem 1.1 of that paper determines a simply connected analytic plane domain in the class  $\mathcal{D}_{1,L}$  from its Dirichlet or Neumann spectrum. This class includes an involution reversing a nondegenerate bouncing-ball orbit, multiplicity one for every iterated length, and the additional nonresonance and endpoint conditions specified there. The proof recovers boundary Taylor coefficients from localized wave-trace invariants. The symmetry relates the two boundary graphs. Thus neither localization near a periodic orbit nor all-order boundary recovery, by itself, distinguishes the problem considered here.

The distinction lies in the observation map and in the passage from a long bridge to its nonlinear endpoint action. Our datum is a controlled conditional law obtained from phase-volume preparations, not a Laplace spectrum or a sequence of wave invariants. The selected bridge becomes exponentially rare, but its endpoints remain on a fixed collar at a fixed positive offset. Theorem 17.2 compares the mixed endpoint derivative with its exponentially small reference value in *relative* size. An absolute error bound would not suffice for the physical conditional normalization. The resulting law retains the nonlinear half-line actions; separate-factor cancellation then extracts those actions from the density. It is the relative limit, rather than the cancellation identity alone, that connects an observable law to the geometric inverse.

The signed contact inverse treats both anchored graphs without imposing a symmetry between them. Its finite-order content is a statement about actual smooth graphs: Lemma 22.4 controls functional remainders by finite-truncation stationarity before passage to any jet coordinates. Analyticity enters the subsequent determination of a complete germ and its continuation, not that smooth factorization. The function-space inverse controls the full coupled action map, including its lower-order terms. With an outer analytic prior, its real-observation consequence estimates the complete germ on a smaller disc. The finite preparation theorems further distinguish the exactly calibrated rate from the rate that charges the position pilot. Their minimum-distance construction is an existence result for a measurable estimator; it does not supply a finite running-time bound for computational reconstruction.

There is also an internal comparison for which a geometric separation is proved, rather than inferred from different descriptions of the data. The family in Theorem A.3.2 has the same gap, area, quadratic contact data and leading count amplitudes, while its positive-offset boundary law detects a varying fourth-order contact term. This shows that the present law retains information absent from those specified leading data. It is not a separation theorem between the law and a

full Laplace spectrum, nor between the law and either marked-length datum discussed above. No such comparison is needed for the relative boundary limit or the contact inverse. These distinctions locate the mathematical contribution in the nonlinear relative-action mechanism and its geometric consequences, without treating different observations as interchangeable.

Part I gives the geometric setting and then proves the relative law, the signed contact inverse, analytic image recovery, finite incidence matching and the derivative-kernel theorem in that order. Part II specifies observation spaces and proves the local information results. Part III gives quantitative histogram inversion, calibrated acquisition, and the full-position benchmarks and global estimators. Detailed alternative statements and extended comparisons are collected in Appendix A.2; they are not additional premises in the proof of Theorem A. The remaining appendices and the two-collision companion retain the complete auxiliary proofs.

**A.1.9. Finite experiments and the complete local analytic target.** The relative law and the analytic inverse together give a finite-experiment consequence without truncating the contact geometry. On the local bounded analytic class of Section 23.4, two actual finite-bridge endpoint densities at the fixed offset satisfy

$$\|\psi - \tilde{\psi}\|_r \leq C\{|g - \tilde{g}| + (\|f_j - \tilde{f}_k\|_{C^0(Q)} + e^{-\omega j} + e^{-\omega k})^\vartheta\}.$$

The argument compares their realizable limiting laws; it does not regard a finite density as exactly factorized. In an exactly calibrated selected channel experiment,  $2N$  independent full-phase preparations, including all failures, yield a measurable estimator of the entire contact pair with confidence error

$$C \left( \frac{\log(C_* N/\alpha)}{N} \right)^{\vartheta\omega/(4\omega+\Gamma)}.$$

Here  $e^{-\omega j}$  is a uniform relative error,  $\Gamma$  bounds the acceptance exponent, and the positive offset and analytic prior are fixed. The flight number is even and proportional to  $\log N$ . This is an upper bound for complete germs on a smaller disc, not a claim of statistical optimality or a finite-dimensional parametrization. The proof estimates cell densities, fits a realizable limiting law, and then invokes the complete analytic inverse. It charges the exponential rarity rather than replacing raw preparations by free conditional samples.

Estimated calibration adds the separately visible terms  $jv_g + v_t$  and  $r_c/h$ , where  $h$  is the transverse cell size, and its preliminary preparation cap is added to the total charge (Corollary 23.8). This application substantively uses the existing position-sensor calibration theorem; the exactly calibrated power bound is not asserted for its total preliminary cost. All original smooth, global matching, unknown-origin, differential and acquisition results retain their hypotheses and proofs.

With the persistent position pilot and its compact marked-class hypotheses, Corollary 23.9 further gives error  $C\{[\log(CB/\alpha)]^2/B\}^{\vartheta\omega/(12\omega+\Gamma)}$  with confidence  $1 - \alpha$  and at most  $B$  preparations in total. Its slower power includes the entire onset scan and coordinate calibration. This remains a local analytic-germ estimate, not a quantitative assertion for unmarked discovery or global continuation.

## APPENDIX A.2. COLLECTED THEOREM STATEMENTS AND COMPARISONS

The geometry of a dispersing billiard can be observed through trajectories, lengths, collision positions, or probability laws of selected collision words. These data have different inverse content. An exact planar collision position lies on the unknown boundary and supplies a geometric sample. Its signed transverse projection does not supply the missing normal coordinate. We study the geometry encoded by the distribution of the latter observations for long alternating bridges near a closest pair of strictly convex scatterers.

Our main local result determines both labelled contact jets from the gap and one signed conditional endpoint law of each contact type at a single positive offset. The bridge probability decays exponentially with the number of flights, but its conditional law retains a nonlinear profile formed by repeated returns to both contacts. We construct this profile as a relative limit on a collar independent of the flight number and invert it on the geometric image of the billiard observation map. Neither reflection symmetry, equal contacts, nor supplied curvatures are required. The inverse combines cancellation of the unknown flux amplitude with a finite-remainder analysis of the nonlinear half-line action. This gives a triangular reconstruction of the labelled contact jets at every prescribed finite order. For connected analytic boundaries, intrinsic channel matching then recovers complete obstacles and the Euclidean period lattice under the marked incidence hypotheses stated below.

We also construct globally consistent estimators in an unregistered physical position-sensor experiment, with a charged calibration pilot and only long even bridges. We give a direct graph-sampling proof available in that same richer experiment. Thus the distinction between geometric sampling and the inverse from transverse laws is part of the theorem statements, rather than an inference from the flight restriction. The local information results separately identify the effects of deleting residual times and waiting counts.

**A.2.1. Relative boundary laws.** For a selected closest-pair channel let  $g > 0$  be its gap, let  $\kappa_0, \kappa_1 > 0$  be its labelled contact curvatures, and put

$$c_b = 1 + g\kappa_b, \quad \gamma = \operatorname{arcosh} \sqrt{c_0 c_1}.$$

If  $W_j(u, v)$  is the stationary length of a  $j$ -flight bridge with signed transverse endpoints, define

$$E_j = W_j - jg, \quad b_j = \frac{-W_{j,uv}}{-W_{j,uv}(0, 0)}.$$

There are strictly convex actions  $S_b$  and positive amplitudes  $B_b$ , normalized by  $S_b(0) = S'_b(0) = 0$  and  $B_b(0) = 1$ , such that

$$\|E_j - S_0 \oplus S_p\|_{C^k} + \|b_j - B_0 \otimes B_p\|_{C^k} \leq C_k \tau^j, \quad p = j \bmod 2, \quad 0 < \tau < 1. \quad (\text{A.2.1})$$

The order  $k$  is fixed but arbitrary, including the mixed geometric and offset derivatives stated in Theorem 12.1. The endpoint and offset collars do not shrink with  $j$ . The local forward construction requires neither finite horizon nor global reflection symmetry. Its weighted inverse, differentiated operator comparison, and uniform hyperbolicity margins are proved in Part I.

At excess time  $d > 0$  the successful endpoint-residual-time limit has density proportional to

$$B_0(u)B_p(v)\mathbf{1}_{\{0 < r < d - S_0(u) - S_p(v)\}}.$$

For same-type even bridges, integrating out  $r$  gives a density  $f_{b,d}(u, v)$  proportional to  $(d - S_b(u) - S_b(v))_+ B_b(u)B_b(v)$ . The family of supports as  $d$  varies determines the unsymmetrized action by a threshold slice. The density itself carries more information: only one positive  $d$  is needed. Proposition 23.13 gives a support-preserving action change in the functional density class, showing why one support alone does not supply the same invariant.

**A.2.2. A single-offset inverse and intrinsic periodic determination.** On a small interior square, write

$$R_b(u, v) = \frac{f_{b,d}(u, v)f_{b,d}(0, 0)}{f_{b,d}(u, 0)f_{b,d}(0, v)}.$$

All unknown amplitude and normalization factors cancel, leaving

$$1 - R_b(u, v) = t_b(u)t_b(v), \quad t_b(u) = \frac{S_b(u)}{d - S_b(u)}. \quad (\text{A.2.2})$$

At a fixed nonzero anchor  $a$ , strict convexity gives  $t_b(a) > 0$ . Consequently

$$t_b(u) = \frac{1 - R_b(u, a)}{\sqrt{1 - R_b(a, a)}}, \quad S_b(u) = \frac{dt_b(u)}{1 + t_b(u)}.$$

These are identities of continuous densities determined by the law; point values of a density are not assumed to be individual observations. The fixed-anchor formula has a locally Lipschitz inverse in every fixed interior  $C^M$  norm, uniformly on compact nondegenerate classes. It avoids taking derivatives of a square root at the minimum of  $S_b$ .

The gap is a separate datum, obtained from physical onset  $fg$ . If  $a_b = S_b''(0)$ , then

$$\gamma = \operatorname{arsinh}(g\sqrt{a_0 a_1}), \quad c_b = \cosh \gamma \sqrt{a_b/a_{1-b}}, \quad \kappa_b = (c_b - 1)/g.$$

Write  $q_n = (\psi_0^{(n)}(0), \psi_1^{(n)}(0))^t$  for the next graph jets and  $s_n = (S_0^{(n)}(0), S_1^{(n)}(0))^t$ . For  $n \geq 3$ ,

$$s_n = M_n q_n + R_n, \quad M_n = \begin{pmatrix} \coth(n\gamma) & \mathfrak{r}_0^n \operatorname{csch}(n\gamma) \\ \mathfrak{r}_1^n \operatorname{csch}(n\gamma) & \coth(n\gamma) \end{pmatrix}, \quad \mathfrak{r}_b = \sqrt{c_b/c_{1-b}}, \quad \det M_n = 1, \quad (\text{A.2.3})$$

where  $R_n$  is determined by lower graph jets. Establishing this filtration requires more than the density cancellation. We invert the nonlinear half-line stationarity operator on a weighted sequence space, differentiate finite truncations of its action, control the endpoint tails, and isolate the first homogeneous occurrence of each new graph jet. A finite Taylor-remainder estimate proves that equal graph jets through order  $M$  give equal action jets through order  $M$ , even for smooth graphs differing by flat terms. Smooth boundary images are not identified by their jets; the boundary-image conclusion uses analyticity. The inverse is fixed-order stable; determinant one alone is not a bound uniform in the infinite jet order.

For a periodic table the deck group is a marked abstract  $\mathbb{Z}^2$ , with no supplied Euclidean Gram form. A channel label  $(a, b, \ell)$  specifies its obstacle orbits and abstract deck displacement. Each channel has its own oriented frame; relative placements between channels are not supplied. A signature-rigid transition has a unique framed point with the prescribed complete oriented curvature signature. In the presence of analytic symmetries the general gluing theorem gives a gluing space, not an arbitrarily selected phase.

**Theorem A.2.1** (Determination from single-offset transverse laws). *For a smooth strictly convex channel, its gap and two same-type signed limiting endpoint laws at one known sufficiently small positive offset determine every finite labelled contact jet. No longitudinal position, residual time, waiting count, flux amplitude, or contact curvature is supplied as part of these conditional-law data.*

*For connected real-analytic incident boundaries the data determine their complete boundary images in the separate channel frames. Suppose the marked data are realizable by a periodic table and the measured channel graph contains a rank-two anchoring pair based at one channel frame, in the sense of Definition 26.1, and a rooted signature-rigid spanning tree reaching every obstacle orbit. Then the data determine the entire marked periodic table, including its Euclidean lattice, up to one simultaneous element of  $\operatorname{SE}(2)$ . On a compact class with uniform collars the same positive offset can be used for every channel and both types.*

*Proof.* Theorem 23.1 supplies the action germs. The all-order contact inverse is Theorem 22.1. Analytic continuation and the intrinsic gluing classification are proved in Theorem 25.3; rank-two metric recovery is Theorem 26.3. Propagation along the rooted spanning tree is Theorem 26.4. To use the same frame for these last two steps, apply Lemma 25.11 to root the supplied tree at an obstacle met by the anchoring pair. A signature-rigid incidence on that pair supplies the unique framed signature required by the lemma. Thus the two cycle holonomies and the tree placements

are reconstructed in one common frame, without an additional registration hypothesis. Their composition is Theorem 23.4.  $\square$

The last lattice step is  $L = VM^{-1}$ , where  $V$  consists of recovered translation holonomies and  $M$  of their independent marked deck vectors. The difficult input is the recovery and unique matching of complete framed curve images, not this matrix identity. Likewise, marked deck labels and signed transverse conventions are part of the intrinsic data map, not quantities inferred from an unlabelled trajectory.

**A.2.3. Physical acquisition and the same-experiment comparison.** For global physical acquisition each channel detector reports the first and last planar collision positions in a common oriented sensor frame for that channel's two types, together with residual preparation time or a failure symbol. Frames of distinct channels are not registered. Designs use only the marked channel, type, even flight number, and physical observation time. All preparations and stopping decisions are retained and charged. These common Borel spaces are defined in (56.1). The anchored scalar local experiment is a different, explicitly coarser model; no unknown contact chart is silently treated as observed.

Let  $\mathfrak{K}$  be the compact marked analytic class of Section 27, with uniform positive geometry, holomorphic curvature collars, persistent channels, the preceding uniqueness structure, and nondegenerate rank-two holonomies. Let  $d_q$  be the labelled  $C^q$  distance in a root-frame gauge.

**Theorem A.2.2** (Uniform physical reconstruction at one post-pilot offset). *There are finite deterministically capped policies  $\pi_s$  and measurable table estimators  $\widehat{T}_s$  in the common physical position experiment such that every flight number used by  $\pi_s$  is even and at least  $s$ , and for every fixed  $q \geq 0$  and  $\epsilon > 0$ ,*

$$\sup_{T \in \mathfrak{K}} \Pr_T^{\pi_s} \{d_q(\widehat{T}_s, T) > \epsilon\} \longrightarrow 0.$$

*No gap, contact center, tangent, inter-channel placement, or lattice metric is supplied. After the pilot all channels and both types use one fixed positive offset from the estimated onset, independent of  $s$ . The estimator uses the estimated gaps and bounded tests of the successful transverse endpoint pairs; it does not require their normal coordinates or residual times in its post-pilot criterion.*

*Proof.* Theorem 56.2 gives a charged near-onset pilot at the already selected final even flight number. The single-offset inverse makes gaps and the two endpoint laws at a fixed  $d_0 > 0$  jointly separating. Lemma 57.1 selects finitely many bounded Lipschitz tests, and Theorem 57.2 gives a finite-template estimator with a deterministic preparation cap. The compact inverse modulus and Theorem 57.3 give the conclusion.  $\square$

The theorem does not claim that half-line inversion is necessary in this position experiment. Proposition 54.1 gives a second route using exactly the same long even designs, charged pilot, and position records. Once the frame error is at most  $\epsilon$ , interpolation of  $M + 1$  separated graph samples at scale  $a$  gives

$$|\widehat{\psi}_b^{(m)}(0) - \psi_b^{(m)}(0)| \leq C_M \{a^{M+1-m} + \epsilon a^{-m}\}, \quad m \leq M.$$

Each required bin has positive preparation probability and a finite charged cap. Choosing  $\epsilon \leq a^{M+1}$  gives fixed-order consistency without estimating an action. Arbitrarily large even flight numbers do not exclude this comparison.

The intrinsic theorem has different data: no graph ordinate is available. Its single-offset distributional inverse and its all-order alternating contact recursion give a reconstruction from genuinely coarsened data. Exact law injectivity is not statistical equivalence to noiseless position sampling. Proposition 43.3 identifies the precise dependence on the observed contact type in a

registered homothetic family. For any finite number of successful endpoint pairs, the scalar-to-position deficiency is one when the even design starts on the varying obstacle, while the reverse deficiency is zero. When both observed endpoints instead lie on the fixed facing obstacle, the matched endpoint experiments are equivalent. Residual time is deleted from both records in this comparison; retaining it in both gives the same dichotomy. Intermediate collisions are not silently added to either record. These comparisons do not alter the two-type law inverse above.

The finite separator library and compact inverse moduli may depend non-effectively on  $\mathfrak{R}$ . Neither global estimator asserts an analytic minimax rate or polynomial-time design construction. Separately, Theorem 27.1 supplies a uniform finite signature with local embedding and global separation bounds, and Lemma 27.2 gives unique matching under controlled  $C^2$  perturbations. The compact inverse modulus does not replace that matching proof. Budget-indexed implementations run the selected stage afresh; a cumulative archive of earlier stages is not claimed to have a diverging minimum flight number.

**A.2.4. Local information under coarsening.** In the anchored finite-dimensional family, residual-time retention leaves positive density at a moving ceiling. At shape scale  $k_n^{-1}$ , the successful endpoint-time experiment is equivalent to a Poisson boundary-shift experiment; the corresponding gap scale is  $(j_n k_n)^{-1}$ . Theorem 42.2 gives parameter-independent kernels in both directions, controlling the layer, reconstructed bulk, and the intersection with the fixed residual-time face. Its bulk estimate uses the explicit bounded  $O(k_n^{-1})$  relative-density perturbation, not bare trace convergence.

The finite-strip Poisson family is jointly dominated by a common Poisson envelope. It need not be absolutely continuous with respect to its zero-shift member. The scalar endpoint family is likewise jointly dominated by Lebesgue measure with the required cemetery atoms. Its common-collar censoring produces the stronger property of domination by the censored reference member. Proposition 43.1 gives both reference measures and the Poisson Radon–Nikodym formula explicitly.

Removing residual time produces linear vanishing at the endpoint support boundary and the scale  $k_n \delta_n^2 \log(1/\delta_n) \rightarrow 1$ . The reference-dominated representative has a uniform LAN expansion. Lemma 37.6 identifies its finite Gaussian limits; the uniform moving-boundary modulus and finite-net argument then prove convergence of the original endpoint family in two-sided Le Cam distance on every fixed compact local set, in Theorem 40.2 and Corollary 40.3. Waiting counts identify the normal quotient to a level set of the hyperbolic exponent at scale  $(j_n \sqrt{k_n})^{-1}$ , while endpoint information acts on the tangent directions at their slower scale. Theorem 44.3 gives the compact two-sided Gaussian product approximation for this count–endpoint record. Deleting those counts is a further coarsening. These assertions concern their declared scalar record spaces, not the support-singular planar position experiment or the complete growing collision array.

**A.2.5. Comparison with related inverse problems.** Marked lengths, dynamical conjugacy, and boundary-law determination pose different inverse problems. De Simoi, Kaloshin and Leguil [5] determine analytic open dispersing billiards from marked lengths under their symmetry and genericity hypotheses. Florio and Leguil [11] obtain smooth conjugacies from periodic length data for the hyperbolic systems covered by their hypotheses, including the stated open-billiard application. This is a dynamical conclusion, distinct from identifying obstacle images by a Euclidean isometry.<sup>1</sup> Finamore and Leguil [12] use an enriched marked length spectrum and a CAT(0) construction to recover finite-horizon Sinai billiards up to isometry.

---

<sup>1</sup>We use the revised version v5 of [11]. Its revision notice removes the earlier spectral-rigidity assertion affected by the error in Proposition 3.1; the smooth-conjugacy result is the comparison used here.



The data in Theorem A.2.1 are instead onsets and signed transverse laws of selected channels, with marked deck labels and signature-rigid incidences. The local observation deletes the normal endpoint coordinates while retaining the two labelled contact types. The proof first identifies its nonlinear limiting density from physical bridges, and then recovers the unsymmetrized actions and the actual contact jets. Only the subsequent boundary-image step uses analyticity. This composition does not pass through a supplied normal-form chart or a dynamical conjugacy. Neither the channel laws and incidence data nor the ordinary or enriched marked length spectrum are derived from the other observation map here; the respective hypotheses and conclusions therefore remain distinct.

The cofactor identity and inverse-decay estimates used in the forward argument are classical tools, as detailed in Section 18. Their role here is to control a nonlinear Hessian perturbation in trace norm along the whole physical bridge, before passing through a normalized moving-sublevel integral. The inverse then requires finite-truncation envelope differentiation and a smooth finite-remainder theorem before homogeneous degree isolation yields the contact recursion. These are the analytic and geometric steps connecting the relative law to the determination theorem; the determinant-one identity and the final lattice equation record their explicit inversion, rather than replacing those steps.

Within the present geometry, the mixed-type one-flight support inverse of Proposition 53.1 and the two-flight calculations remain useful benchmarks. The direct-position result above is stronger as a comparison to our physical protocol because it uses the same long even record restriction. The additional inverse content lies in recovering geometry from the single-offset transverse laws, not in prohibiting one short flight while supplying noiseless geometric samples.

Nonregular support-dependent estimation and exceptional normal rates have an established history; see Smith [33]. Meister and Reiss [36] prove equivalence of a nonregular regression experiment to Poisson experiments whose intensity supports encode the regression boundary. A Gaussian/Poisson contrast or Poissonization alone is therefore not asserted as a new principle. The results here identify the billiard-specific nonlinear density, its amplitude-free and all-order contact inverses, and the kernels linking explicitly retained observation levels, with the rare-event preparation cost included.

**A.2.6. Organization of the proof and the collected statements.** The statements and extended comparisons in this appendix retain the alternative local, signature-rigid and observation-specific formulations. The primary proof of Theorem A runs through Part I: relative boundary laws, signed contact inversion, analytic image recovery, finite incidence matching and the derivative kernel. Part II treats the local information experiments. Part III contains finite-resolution inversion, calibrated acquisition and complete physical reconstruction, including the short-flight and direct-position comparisons. The following appendices retain the complete auxiliary derivations.

**A.2.7. A finite-channel theorem for a generic class.** The intrinsic inverse has a consequence for a broad class of periodic tables, without assuming in advance the existence of a suitable measured graph. Let  $\mathcal{P}_N^\omega$  be the space of labelled periodic tables with  $N$  connected real-analytic strictly convex obstacle boundaries of positive curvature and disjoint lifted closures. The marked rank-two deck group is fixed; its Euclidean realization is not observed. Equip support functions with the relative  $C^2$  topology and the lattice matrix with its usual topology.

**Theorem A.2.3** (Finite-channel periodic rigidity). *The tables for which every obstacle has trivial proper Euclidean symmetry form an open dense subset of  $\mathcal{P}_N^\omega$ . For every table in this subset, there is a selection of  $N + 1$  marked channels and a positive offset at which their  $N + 1$*

*onsets and  $2N + 2$  same-type signed limiting endpoint laws determine the complete labelled table and its Euclidean period lattice up to one simultaneous proper Euclidean motion.*

*The selected graph is a spanning tree with two additional edges of independent deck-cycle gain. Its labels, gates and rank persist in a geometric neighborhood. No common laboratory registration, supplied curvatures or lattice metric are used in the inverse.*

The full statement and proof are given in Theorem 29.7. The geometric input is Theorem 29.3, which applies even to symmetric tables. If a closest segment between two lifted obstacles meets a third body, that body splits the pair into two smaller-distance pairs. Uniform positive separation makes this descent finite, and its total terminal edge distance never exceeds the original gap. Applying it to paths reaching the obstacle types and two independent deck translates produces a finite clear graph. A spanning tree and two independent fundamental gains reduce the graph to  $N + 1$  channels. The cycle-space dimension proves that this count is minimal for the connected two-cycle mechanism, not for arbitrary observation maps. The selected channels are genuine local alternating billiard channels; the closed graph walks used for holonomy are not asserted to be specular multi-channel orbits.

The inverse then follows the previously stated analytic chain: relative bridge laws, single-offset amplitude cancellation, smooth finite-remainder contact inversion, and analytic image determination. Proper obstacle asymmetry makes every incidence signature unique. The tree and the two cycles can therefore be registered in one frame. This supplies the architecture required by the more general conditional gluing theorem; it does not weaken that theorem or replace its local proof by graph algebra. On an explicit open dense harmonic sublocus, second and third support coefficients of the *recovered complete images* give all relative rotations. Integrating their center-displacement cochain on the tree and taking its two cycle defects yields both lattice columns and all obstacle placements; see Proposition 29.8.

Genericity refers to the support topology just specified, not to a measure on tables. Fixed gates do not reveal exact contacts, and selecting channels from an unknown unmarked trajectory is a different acquisition problem. Uniform physical reconstruction is obtained on compact analytic subfamilies within a persistent design neighborhood, with the same charged pilot and rare-event caps as before. Image-registration stability alone is not a law-to-entire-boundary rate; Section 50 proves a separate quantitative law-to-table estimate under explicit uniform analytic and forward bounds. The observations remain different from the marked length data of [5, 12]; the theorem is not a removal of hypotheses from those spectral results.

**A.2.8. A quantitative inverse for a coarsened observation.** Exact analytic determination and stable reconstruction from approximate laws are different conclusions. The following theorem passes between them under quantitative analytic bounds. Unlike a stability statement on already recovered curve images, its error is measured directly in finitely many probabilities of quantized transverse endpoint events.

**Theorem A.2.4** (Finite-resolution law-to-table reconstruction). *Fix a persistent  $N + 1$ -channel design and a quantified analytic class as in Section 50. The bounds include a common holomorphic support strip, positive contact and density margins, finite-order forward density bounds, and a finite nonvanishing harmonic witness for the proper asymmetry of each obstacle. There is an explicit conditional modulus  $\Omega(e_g, e_p + C_{\text{bin}}\delta)$  tending to zero with its arguments such that two tables whose onsets differ by at most  $e_g$  and whose signed endpoint histograms of mesh  $\delta$  differ by at most  $e_p$  in  $\ell^1$  satisfy*

$$\inf_{(U,t) \in \text{SE}(2)} \left\{ \|L - U\tilde{L}\| + \max_a d_H(K_a, U\tilde{K}_a + t) \right\} \leq \Omega(e_g, e_p + C_{\text{bin}}\delta).$$

*The same estimate controls the marked lattice Gram matrix. Every positive target accuracy has a finite prescription for the required jet order, grid resolution and probability accuracy. On a calibrated design, finitely many successful even-flight observations give a confidence reconstruction with explicit categorical sampling, finite-flight bias and preparation-cap bounds. No longitudinal endpoint coordinate is used by this reconstruction.*

Theorem 50.5 and Corollary 50.6 give the constants and proofs. The inverse is not assumed quantitatively: a cell-average estimate and an interior interpolation inequality first control a finite density jet. A finite Bellman recursion evaluates the lower-order terms in the determinant-one contact inverse. Contact graphs then give anchored support jets. A chain of three-circles estimates continues these finite jets around the whole boundary, and a Bézout formula for finitely many nonzero harmonics supplies stable registration of all incidences. The resulting modulus is an infimum of fully specified finite-order bounds. To attain a given error one uses one finite order, without estimating the conditioning of an infinite triangular system or appealing to an unspecified compactness separation constant.

The analytic width and its associated bounds are additional quantitative priors; the earlier qualitative theorem retains its full statement. The number of histogram cells grows with accuracy, so the result is not exact recovery from a fixed finite number of scalars. Calibration is specified separately from the reconstruction transcript. The richer position experiment and its direct interpolation benchmark remain distinct. The continuation and empirical-frequency estimates are classical mechanisms, and their severe conditioning is not claimed as a new analytic-continuation principle. The conclusion here is their quantified composition with the nonlinear billiard law-to-contact inverse and the intrinsic periodic gluing construction.

**A.2.9. Calibration and the finite transverse record.** The quantitative inverse also admits a physical acquisition in estimated contact frames. There are three distinct approximation errors. A bridge with  $j$  flights contributes its relative-law error. Programming its time using an estimated gap contributes the amplified error  $j|\hat{g} - g|$ . Replacing an exact transverse coordinate by one with error  $r$  contributes the probability of crossing a grid edge, including an outer edge of the recorded square. These last events are not bounded by applying a Lipschitz test inequality to a cell indicator.

Theorem 50.9 propagates the three errors through the same law-to-table modulus, with explicit confidence and preparation bounds. Corollary 50.10 gives a finite joint prescription for the grid, the final even flight number, the resolution of a pilot using that same flight number, and the number of successful categories. The post-pilot reconstruction retains only the estimated gaps and transverse categories. The richer physical positions used by the pilot and the observable projection remain part of its acquisition specification; they are not silently inserted into the compressed inverse.

The structural input throughout is the nonlinear relative profile and its signed finite-remainder inverse. Offset normalization, categorical coupling, concentration, and analytic continuation propagate this input to different observation levels. Their combination yields the stated geometric and statistical consequences without identifying those observation levels with each other.

**A.2.10. An integrated nonsymmetric realization.** Section 28 constructs a periodic family in which the hypotheses of the intrinsic theorem hold simultaneously. There are two analytic obstacle types with disjoint curvature ranges, four prescribed selected channels, two independent deck cycles based at one channel frame, and a signature-rigid spanning tree. Explicit support functions exclude both proper and reflection symmetries. Uniform separation and third-obstacle clearance estimates hold for every lattice translate, while the lattice shape, scale and the placement of the second obstacle vary.

For this family the registration step has a particularly transparent form. Once the channel laws have yielded complete curve images, their support coefficients of orders two and three determine relative rotations by a single complex quotient. Differences of the registered second-type centers then give the two lattice vectors. These support coefficients are computed from recovered images; they are neither measured ordinates nor additional calibration data. Theorem 28.1, Proposition 28.2, and Corollary 28.3 give the construction, the registration formula, and the complete observation-to-table reconstruction, respectively. The fourth channel supplies a redundant holonomy identity in that same frame. Proposition 28.4 proves persistence under independent analytic boundary perturbations, including higher Fourier modes. Thus the demonstration is not confined to fitting the parameters of two trigonometric supports.

The explicit formulas concern geometric realization and the final matching stage. They do not replace the nonlinear relative limit or the smooth finite-remainder contact inverse, and their registration stability is not an effective analytic-continuation rate. The construction keeps the marked channels, signed conventions and prescribed gates of the original observation map. The richer planar-position benchmark remains distinct.

**A.2.11. Proof architecture and the role of the observation map.** The principal analytic assertion is the relative nonlinear boundary law; the principal inverse assertion is recovery of the two labelled contact jets from signed transverse laws. Their composition precedes the global and statistical applications. We describe the dependencies to distinguish the nonlinear arguments from the algebraic and compactness steps that complete those applications.

*Relative asymptotics.* The finite Jacobi calculation identifies both the endpoint Hessian and the normalization of the mixed derivative. The perturbation of the interior Hessian is summable in trace norm, uniformly in the number of flights. Lemma 15.2 therefore controls a relative determinant, rather than an absolute error divided by an exponentially small reference twist. The half-line construction then separates the two endpoint layers. Gluing the two half-line solutions and estimating the discarded middle block gives the factorization in Theorem 17.2. The uniform sublevel integral and its fixed-domain differentiation yield Theorem 12.1. Every derivative order is finite and fixed; the constants may depend on that order and on the corresponding boundary norms. The offset collar is independent of the flight number. Geometric differentiation is channel-centered, with physical time written as  $t = jg(\xi) + d$ .

*Signed inversion.* The four-density identity in Theorem 23.1 removes the unknown amplitudes at one positive offset. This algebra recovers actions, not graph jets directly. To pass to the contact geometry one must differentiate the half-line action through finite truncations, cancel the interior variations by stationarity, and control the terminal term before taking the limit. The smooth-remainder argument proves finite-jet factorization for actual smooth graphs, including graphs differing by flat terms. Homogeneous degree isolation then gives the triangular recursion of Theorem 22.1, with the explicit two-by-two blocks in (A.2.3). Invertibility is established at each fixed order; no uniform inverse on an unrestricted infinite jet space is inferred from the determinant identity.

*Intrinsic global reconstruction.* Only after the local inverse is proved does analyticity extend the recovered germs to connected boundary images. The gluing problem retains the full oriented signatures and the marked deck labels. A common-frame rank-two anchoring pair recovers the period metric. Lemma 25.11 transports the root of the supplied signature-rigid tree to an obstacle met by that pair. The tree then places every remaining obstacle orbit in the same frame. No new incidence, deck displacement, or observation is needed. Realizability and unique signature matching are hypotheses of this global conclusion, not consequences of the matrix identity for the lattice. The finite-signature embedding and noisy matching theorem supply the local separation

and uniqueness needed before the compact inverse-modulus argument. Neither compactness nor analytic continuation alone gives a numerical reconstruction rate.

*Statistical and physical consequences.* The transfer estimates are used only after the boundary law is available. Finite likelihood-vector convergence and a separate uniform-modulus argument supply the compact-local Gaussian comparisons. The likelihood-tilting and original-alternative moment estimates justify the quadratic-loss statements without assuming that total variation controls unbounded losses. The moving-ceiling comparison uses two actual Markov kernels, together with separate layer, bulk, and corner estimates. Finally, the global physical protocol includes a charged position-based pilot. Its post-pilot transverse statistic is coarser than the total position record, and the direct-position benchmark belongs to that same richer acquisition experiment. The smooth-nuisance envelope in Section A.16 is a distinct comparison model; its lower bound does not assert realizability of every splice by a billiard table. These distinctions are part of the mathematical statements and remain in force throughout the applications.

The complete auxiliary proofs are included in the appendices. No statistical equivalence between different observation maps is used in the deterministic contact or periodic rigidity arguments.

### APPENDIX A.3. INFORMATION BEYOND THE LEADING HIERARCHY

The nonlinear law of the preceding section is not determined by the contact Hessian. We prove this by an analytic deformation with fixed area, gap and contact curvatures. Its higher coefficient can be computed from the half-line action and determinant, without extrapolating an uncontrolled finite-length error.

**A.3.1. The quartic variation of a boundary law.** Assume first that the two facing graphs are identical and even, and write

$$\psi(y) = \frac{\kappa}{2}y^2 + \frac{q}{24}y^4 + O(y^6), \quad c = 1 + g\kappa, \quad \gamma = \operatorname{arcosh} c, \quad a = \frac{\sinh \gamma}{g}, \quad \lambda = e^{-\gamma}. \quad (\text{A.3.1})$$

Here  $q$  is a fourth graph derivative, not a source variable. Fix  $g$  and  $\kappa$  and vary  $q$  in a smooth family. Higher graph jets may vary as well. The two boundary actions and amplitudes now agree; write them as  $S$  and  $B$ . Both limiting parities give one function  $\mathcal{F}$ .

**Proposition A.3.1** (A nonlinear coefficient and its quartic derivative). *Let  $\mathcal{R}_\infty(0) = \partial_d \mathcal{F}(0)$ . At fixed  $g, \kappa$ , its derivative with respect to  $q$  is*

$$\partial_q \mathcal{R}_\infty(0) = -\frac{\cosh(2\gamma) + 2}{12a^2 \sinh(2\gamma)}. \quad (\text{A.3.2})$$

*This derivative does not depend on the graph derivatives of order six or higher. More precisely, if  $S(u) = au^2/2 + tu^4/24 + O(u^6)$  and  $B(u) = 1 + \beta u^2 + O(u^4)$ , then*

$$\partial_q t = \coth(2\gamma), \quad \partial_q \beta = -\frac{1}{4a \sinh(2\gamma)}, \quad \mathcal{R}_\infty(0) = \frac{2\beta}{3a} - \frac{t}{12a^2}. \quad (\text{A.3.3})$$

*Proof.* The linear half-line segment is  $x_i = \lambda^i u$ . Because the quadratic action is stationary with its first endpoint fixed, its change upon substituting the cubic bridge correction begins at degree six. The quartic part of the stationary action is therefore the quartic part of the broken action evaluated on that linear segment. In one flight, the  $q$ -dependent quartic term of the length is  $q(x_i^4 + x_{i+1}^4)/24$ . The endpoint is counted once and every interior site twice. It follows that

$$\partial_q t = 1 + 2 \sum_{i \geq 1} \lambda^{4i} = \frac{1 + \lambda^4}{1 - \lambda^4} = \coth(2\gamma).$$

This use of stationarity may equivalently be obtained by differentiating finite stationary actions and passing to the convergent half-line sum.

To degree two in  $u$ , the  $q$  variation of an edge's mixed derivative vanishes: the quartic term just displayed is a sum of two pure fourth powers. Thus the product term  $U$  in (17.6) has no  $q$ -dependent quadratic coefficient. The interior Hessian variation at that degree is diagonal, with entries  $q\lambda^{2i}u^2$ . Changes in the cubic bridge do not enter degree two, and products of two Hessian perturbations enter only degree four. Consequently the quadratic coefficient of the logarithmic determinant has derivative  $\sum_{i \geq 1} G_{ii}\lambda^{2i}$ . At equal curvatures, (17.1) gives  $G_{ii} = (1 - \lambda^{2i})/(2a)$ . Hence

$$\partial_q \beta = -\frac{1}{2a} \left( \frac{\lambda^2}{1 - \lambda^2} - \frac{\lambda^4}{1 - \lambda^4} \right) = -\frac{1}{4a \sinh(2\gamma)}.$$

The quadratic coefficient of  $B$  equals that of  $\log B$  because both have vanishing linear coefficient by evenness.

Finally use (17.12). The density correction is  $\beta(u^2 + v^2)$ , whose expectation under the quadratic residual-time weight is  $2\beta d/(3a)$ . For the quartic action correction, first variation of the positive-part integral gives the negative integral of  $t(u^4 + v^4)/24$  over  $a(u^2 + v^2)/2 < d$ ; there is no boundary term because the residual weight vanishes there. The two elementary integrals are

$$\int (d - ar^2/2)_+ du dv = \frac{\pi d^2}{a}, \quad \int_{r^2 < 2d/a} (u^4 + v^4) du dv = \frac{2\pi d^3}{a^3}.$$

Taylor's formula in the uniform Morse coordinates justifies the remaining  $O(d^2)$  relative error. This gives the last identity in (A.3.3), and substitution proves (A.3.2). Terms of degree six or higher cannot contribute to this first offset coefficient.  $\square$

For an additional finite-length identity, put  $L_i = (\sinh((j-i)\gamma), \sinh(i\gamma))/\sinh(j\gamma)$ ,  $M_j = H_j^{-1}$  and  $v_i = L_i M_j L_i^t$  in the equal-curvature case. Thus  $L_0 = (1, 0)$  and  $L_j = (0, 1)$ . The same calculation before taking the limit yields

$$\partial_q \mathcal{R}_j(0) = -\frac{1}{3} \sum_{i=1}^{j-1} (G_j)_{ii} v_i - \frac{1}{24} \left( v_0^2 + v_j^2 + 2 \sum_{i=1}^{j-1} v_i^2 \right). \quad (\text{A.3.4})$$

Indeed, the logarithmic amplitude variation is  $-\sum (G_j)_{ii} (L_i Y)^2$ , and its quadratic conditional expectation is  $-d \sum (G_j)_{ii} v_i/3$ . The unweighted fourth moment in the quadratic sublevel, divided by its residual-time integral, is  $dv_i^2$ . These observations prove (A.3.4), including  $j = 1$  with an empty interior sum. Its limit is (A.3.2). This identity provides a finite check of the normalization but is not needed to define the infinite boundary determinant.

### A.3.2. An analytic family invisible to all leading data.

**Theorem A.3.2** (Identical leading metrics and distinct nonlinear laws). *On the lattice  $3\mathbb{Z} \times 4\mathbb{Z}$  there is a real-analytic family of strictly convex scatterers  $D_s$ ,  $|s| < s_*$ , with  $D_0$  the unit disk, for which*

$$g = 1, \quad A = 12 - \pi, \quad \kappa_0 = \kappa_1 = 1$$

*at the two horizontal contacts. The horizontal orientations exhaust the ground-onset event. All leading count amplitudes and all leading endpoint covariance matrices, at every flight number, are independent of  $s$ . Nevertheless their nonlinear boundary coefficient satisfies*

$$\left. \frac{d}{ds} \mathcal{R}_\infty^{(s)}(0) \right|_{s=0} = \frac{\sqrt{3}}{2}. \quad (\text{A.3.5})$$

*It therefore determines  $s$  locally with a bounded inverse derivative. For every sufficiently small fixed positive  $d$ , the limiting normalized probability  $\mathcal{F}^{(s)}(d)$  also varies with  $s$ .*

*Proof.* Consider the support function

$$h_{s,z}(\theta) = 1 + s \sin^4 \theta + z \sin^6 \theta. \quad (\text{A.3.6})$$

Its obstacle area is  $\mathcal{A}(h) = \frac{1}{2} \int_0^{2\pi} (h^2 - (h')^2) d\theta$ . At  $(s, z) = (0, 0)$  its  $z$  derivative is  $\int_0^{2\pi} \sin^6 \theta d\theta = 5\pi/8 > 0$ . The analytic implicit-function theorem gives  $z = z(s)$ ,  $z(0) = 0$ , with  $\mathcal{A}(h_{s,z(s)}) = \pi$  exactly. Since the  $s$  derivative is  $3\pi/4$ , one also has  $z'(0) = -6/5$ .

At  $\theta = 0, \pi$ , both perturbations and their first two derivatives vanish. Hence  $h = 1$ ,  $h' = 0$  and the curvature radius  $h + h'' = 1$ . The support points remain  $(1, 0)$  and  $(-1, 0)$ , their curvatures remain one, and the supporting lines of horizontal neighbors have separation one. These points attain that separation. At the disk the vertical gap is two; all other nonhorizontal gaps are larger. A small analytic parameter interval preserves this strict gap, disjointness, positive curvature, and clearance. Thus only the horizontal pair contributes to the ground onset. Reflection and central symmetry of (A.3.6) make the two facing graphs identical and even. Area and the two contact Hessians are fixed exactly, not merely to first order. Equations (15.2) and (A.18.1) give the assertion about every leading datum.

We compute the first higher graph jet. Let  $r(\theta) = h(\theta) + h''(\theta)$  be the curvature radius. Near the right contact the boundary parametrization has derivatives  $x'(\theta) = -r(\theta) \sin \theta$  and  $y'(\theta) = r(\theta) \cos \theta$ . Since  $r'(0) = 0$ , integrating these expansions and eliminating  $\theta$  gives for the facing graph

$$\psi(y) = \frac{y^2}{2r(0)} + \frac{3r(0) - r''(0)}{24r(0)^4} y^4 + O(y^6).$$

In our family  $r(0) = 1$  and  $r''(0) = h''(0) + h^{(4)}(0) = 24s$ ;  $\sin^6 \theta$  has zero fourth jet and does not affect this identity. Thus  $q_s = 3 - 24s$ . Here  $\gamma = \text{arcosh } 2$  and  $a = \sqrt{3}$ , so  $\cosh(2\gamma) = 7$ ,  $\sinh(2\gamma) = 4\sqrt{3}$ . Multiplying (A.3.2) by  $q'_s = -24$  yields (A.3.5). The inverse-function theorem proves the local scalar reconstruction. Finally the joint analytic expansion  $\mathcal{F}^{(s)}(d) = 1 + d\mathcal{R}_\infty^{(s)}(0) + O(d^2)$  has uniform differentiated remainder. Its  $s$  derivative at zero is  $d\sqrt{3}/2 + O(d^2)$ , nonzero for all sufficiently small positive  $d$ .  $\square$

The family in Theorem A.18.3 varies the two contact curvatures while preserving only their scalar product. The family in Theorem A.3.2 instead keeps both contact curvatures fixed, so even the entire leading conditional metric is unchanged. It is separated by the nonlinear law, not by a relabelling of the leading inverse. Neither construction claims equality of finite-offset probabilities or of marked length spectra. Higher-jet information from other billiard observations has its own inverse theory; the assertion here is its explicit realization in a full-phase statistical threshold law with a collision-order-uniform limiting construction.

#### APPENDIX A.4. STATISTICAL DETERMINATION OF AN ANALYTIC CONTACT

The scalar fourth-jet calculation extends to a triangular inverse at every finite order. In this section the two facing graphs are identical and even. This hypothesis is additional to those of the general threshold theorem; it is satisfied by the analytic obstacles constructed below. Write

$$\psi(y) = \frac{\kappa}{2} y^2 + \sum_{m \geq 2} \frac{q_{2m}}{(2m)!} y^{2m}, \quad a = \frac{\sinh \gamma}{g}, \quad \gamma = \text{arcosh}(1 + g\kappa), \quad \lambda = e^{-\gamma}. \quad (\text{A.4.1})$$

For smooth graphs this notation denotes their Taylor coefficients, not an assertion that the series converges. Both limiting parities coincide. Put

$$f_k = \frac{1}{k!} \partial_d^k \mathcal{F}(0), \quad f_{j,k} = \frac{1}{k!} \partial_d^k F_j(0).$$

The functions  $F_j$  and  $\mathcal{F}$  are the normalized physical probabilities of Theorem 17.3.

**Theorem A.4.1** (Triangular contact-jet recovery). *Fix  $g, \kappa > 0$ . For every  $m \geq 2$ , the coefficient  $f_{m-1}$  depends only on  $g, \kappa, q_4, \dots, q_{2m}$ , is affine in its last argument, and has derivative*

$$\frac{\partial f_{m-1}}{\partial q_{2m}} = -K_m \{\coth(m\gamma) + 2mD_m\} < 0, \quad (\text{A.4.2})$$

where

$$K_m = \frac{4}{2^m(m+1)(m!)^2 a^m}, \quad D_m = \frac{1}{e^{2(m-1)\gamma} - 1} - \frac{1}{e^{2m\gamma} - 1} > 0. \quad (\text{A.4.3})$$

Consequently, for every fixed  $M$ , the map

$$(q_4, \dots, q_{2M}) \mapsto (f_1, \dots, f_{M-1})$$

has a triangular smooth inverse, locally Lipschitz on compact finite-jet sets. It is analytic in finite-jet coordinates. There is no separation assumption on the higher jets. If two real-analytic even facing graphs have the same  $g, \kappa$  and the same germ of  $\mathcal{F}$ , they coincide as graph germs.

For each fixed  $M$ , inserting the finite-bridge coefficients  $(f_{j,1}, \dots, f_{j,M-1})$  in the limiting inverse recovers the actual jet vector with error at most  $C_M \tau^j$ , uniformly on a sufficiently small compact parameter family. The constants may depend on  $M$  and on higher boundary norms; no inverse bound uniform in  $M$  is asserted.

*Proof.* We first justify the finite-jet dependence. Assign degree one to each endpoint coordinate. The coefficient of degree  $n$  of a stationary segment is obtained by differentiating its equation  $n$  times at the zero segment. The highest segment derivative is multiplied by the same quadratic interior Hessian at every step. Its inverse is the Green operator in (17.1); all other terms involve already determined lower derivatives and derivatives of the local length. The weighted bounds in Lemma 17.1 justify this induction also on the half-line. The action and the logarithmic edge and determinant sums can be differentiated term by term, using their summable entrywise bounds. Thus their coefficients through any given degree depend on only finitely many graph derivatives. In particular, the normalized coefficient of  $d^{m-1}$  uses the action through degree  $2m$  and the amplitude through degree  $2m-2$ . This last assertion follows also directly from the smooth Morse construction: after the change  $(u, v) = \sqrt{d}(U, V)$ , an action term of degree  $n$  carries degree  $d^{n/2-1}$  and an amplitude term degree  $d^{n/2}$ . Odd terms vanish under simultaneous sign reversal. Taylor remainders in the common Morse coordinates justify this bookkeeping with differentiated bounds.

Keep the jets below  $2m$  fixed and vary  $q_{2m}$ . In the local length, the first varying homogeneous term is

$$\frac{u^{2m} + v^{2m}}{(2m)!}.$$

It has no mixed derivative. The linear half-line segment is  $x_i = \lambda^i u$ . By the envelope identity for a stationary action, the varying homogeneous term of its stationary value is obtained by substituting this linear segment. This can first be done on finite segments; the weighted derivative estimates pass its absolutely convergent sum to the half-line. If  $s_{2m} = S^{(2m)}(0)$ , then

$$\partial_{q_{2m}} s_{2m} = 1 + 2 \sum_{i \geq 1} \lambda^{2mi} = \coth(m\gamma). \quad (\text{A.4.4})$$

The boundary site is counted once, each interior site twice. Lower nonlinear corrections to  $x_i$  have degree at least three, and cannot alter the displayed degree- $2m$  term.

Write  $\beta_{2m-2} = [u^{2m-2}]B(u)$ , with ordinary rather than factorial coefficients. At degree  $2m-2$  the logarithmic edge product has zero variation. The varying interior Hessian is diagonal at this



degree, with entries

$$\frac{2\lambda^{(2m-2)i}u^{2m-2}}{(2m-2)!}, \quad i \geq 1.$$

The change of the stationary segment itself has degree at least  $2m-1$ . Since the Hessian perturbation for even graphs starts in degree two, that change cannot enter degree  $2m-2$ . For the same reason, only the linear term of the logarithmic determinant contributes at this degree; every product with an unvaried Hessian perturbation has at least two additional degrees. Using  $G_{ii} = (1 - \lambda^{2i})/(2a)$  gives

$$\partial_{q_{2m}}\beta_{2m-2} = -\frac{1}{a(2m-2)!} \sum_{i \geq 1} (1 - \lambda^{2i})\lambda^{(2m-2)i} = -\frac{D_m}{a(2m-2)!}. \quad (\text{A.4.5})$$

Passing from  $\log B$  to  $B$  does not change this first varying coefficient. All products with a lower positive-degree term again lie above the degree under consideration.

It remains to compute how these two variations enter the probability. Set  $k = m-1$ . Under the quadratic residual-time density proportional to  $(d - a(u^2 + v^2)/2)_+$ , polar integration gives

$$\mathbb{E}(u^{2k}) = \frac{2\binom{2k}{k}}{2^k(k+1)(k+2)a^k}d^k, \quad \int (d - a(u^2 + v^2)/2)_+ du dv = \frac{\pi d^2}{a}. \quad (\text{A.4.6})$$

Indeed, the angular average of  $\cos^{2k}\theta$  is  $\binom{2k}{k}/4^k$ , while the ratio of the radial integrals is  $2(2d/a)^k/((k+1)(k+2))$ . First variation with respect to the action gives the negative integral of  $(u^{2m} + v^{2m})/(2m)!$  over the quadratic sublevel. There is no boundary contribution because the residual-time weight vanishes there. Dividing by  $\pi d^2/a$ , its coefficient of  $d^{m-1}$  is  $-K_m$ . The amplitude variation contributes

$$\frac{4\binom{2m-2}{m-1}}{2^{m-1}m(m+1)a^{m-1}}\partial_{q_{2m}}\beta_{2m-2} = -2mK_mD_m.$$

Lower-order terms in the action and amplitude multiply either variation only in higher degrees, so the calculation at the quadratic part is valid for arbitrary fixed lower jets. Equations (A.4.4)–(A.4.5) prove (A.4.2). They also show its independence of all higher jets and of the value of  $q_{2m}$  itself. The finite-jet dependence then proves affinity in that argument.

Inductively write  $f_{m-1} = d_m q_{2m} + P_m(g, \kappa, q_4, \dots, q_{2m-2})$ , with  $d_m$  given by (A.4.2), and solve this identity for  $q_{2m}$ . Each  $d_m$  is nonzero; for a fixed finite set of orders it is bounded away from zero on compact positive gap-curvature boxes. The recursion and its derivatives give the asserted local Lipschitz bounds. The finite-jet operations are analytic: they consist of the local analytic length expansion, inversion of its positive quadratic operator, and absolutely convergent geometric sums. Equality of every recovered jet implies equality of analytic graph germs by convergence of their Taylor series.

Finally Theorem 17.3 gives  $|f_{j,k} - f_k| \leq C_M \tau^j$  for  $k \leq M-1$ . Apply the finite-dimensional local Lipschitz inverse just constructed. For all sufficiently large  $j$  the argument is in its neighborhood. The recursion defines an inverse for arbitrary finite coefficient vectors at fixed  $g, \kappa$ ; the finitely many remaining  $j$  are covered by enlarging the constant on the compact parameter set. This proves the final assertion.  $\square$

**Corollary A.4.2** (Analytic obstacle determination from a selected onset). *Within the identical-even contact class, the gap, the leading selected probabilities, and the germ of their normalized limiting law determine the analytic contact graph. If the participating obstacles have connected real-analytic closed boundaries, this determines each such boundary in its contact frame. It does not determine the lattice or other obstacles which do not participate in this channel.*

*Proof.* Let  $L_j = \lim_{d \downarrow 0} d^{-2} \Pr(E_{j,e}(d))$ . The ratio  $L_1/(2L_2) = \cosh \gamma$  determines  $\kappa = (\cosh \gamma - 1)/g$ . Also  $F_j(d) = \Pr(E_{j,e}(d))/(d^2 L_j)$ , so the normalized datum does not require a supplied area. The theorem determines the contact germ. For the global assertion, parametrize two candidate boundaries by oriented arclength from the common contact. Their signed curvature functions are real analytic on the real line, by continuation around each closed regular curve. They agree on an interval, hence everywhere by the real-analytic identity theorem. Integrating the planar Frenet equations with their common initial frame makes the two boundary images identical. This is exact analytic continuation, not a stability estimate for continuation from noisy local observations.  $\square$

#### A.4.1. Finite-order geometric realization.

**Theorem A.4.3** (Independent statistical jet coordinates at fixed leading data). *For every fixed  $M \geq 2$  there is an  $(M - 1)$ -parameter real-analytic family near the unit disk on  $3\mathbb{Z} \times 4\mathbb{Z}$  which preserves  $g = 1$ ,  $A = 12 - \pi$ , both horizontal contact curvatures, and all collision-order leading count and endpoint-metric coefficients of the selected horizontal ground channel, with its reverse orientation, exactly. Its coordinates can be taken to be  $(q_4, \dots, q_{2M})$ , or equivalently  $(f_1, \dots, f_{M-1})$ . Both coordinate changes have bounded inverse derivatives on a sufficiently small compact parameter neighborhood.*

*Proof.* Use the support functions

$$h_{s,z}(\theta) = 1 + \sum_{m=2}^M s_m \sin^{2m} \theta + z \sin^{2M+2} \theta.$$

At the disk the derivative of the obstacle area with respect to  $z$  is  $\int_0^{2\pi} \sin^{2M+2} \theta d\theta > 0$ . The analytic implicit-function theorem supplies  $z = z(s)$  keeping that area equal to  $\pi$ . All the perturbations vanish to order at least four at  $0, \pi$ . Thus the two horizontal support points and their curvature radii remain fixed, and so do the gap and free area. Positive curvature, clearance, and the strict gap to nonhorizontal translates persist on a small neighborhood, as in Theorem A.3.2.

The compensator vanishes to order  $2M + 2$  and cannot affect graph jets through order  $2M$ . More generally  $q_{2m}$  depends only on  $s_2, \dots, s_m$ , and

$$\partial_{s_m} q_{2m} = -(2m)!. \quad (\text{A.4.7})$$

To see the sign and coefficient without fixing the normal parameter, write the right boundary as  $x = x(y)$ . At its supporting normal  $\theta(y)$ , stationarity of  $(h(\theta) - y \sin \theta)/\cos \theta$  gives, at fixed  $y$ ,  $\partial_{s_m} x = \sin^{2m} \theta(y)/\cos \theta(y)$ . The contact radius is one, so  $\theta(y) = y + O(y^3)$ . Since  $\psi(y) = 1 - x(y)$ , the first varying coefficient is  $-y^{2m}$ , proving (A.4.7). Lower graph coefficients cannot vary because the perturbation starts in that order. The compensator term has still higher order even after its implicit differentiation. The triangular Jacobian is invertible. Combine it with Theorem A.4.1. The constancy of this channel's full leading hierarchy follows from Proposition A.19.1 and the unchanged contact Hessians, including both parities.  $\square$

**A.4.2. The one-flight comparison.** The preceding inverse information is nonlinear, but is not exclusive to long records. This distinction can be made at every order, not just at the fourth jet.

**Proposition A.4.4** (One-flight triangular diagonal). *Let  $\nu_1 = gc/(c^2 - 1)$ , where  $c = 1 + g\kappa$ . For  $m \geq 2$ ,*

$$\partial_{q_{2m}} f_{1,m-1} = -\frac{4\nu_1^m}{2^m(m+1)(m!)^2} < 0. \quad (\text{A.4.8})$$

Thus the normalized one-flight nonlinear germ also determines the analytic even contact graph at known  $g, \kappa$ . For the family of Theorem A.3.2,

$$\partial_s f_{1,1}^{(s)}|_0 = \frac{8}{9}, \quad \partial_s f_1^{(s)}|_0 = \frac{\sqrt{3}}{2}.$$

*Proof.* There are no interior variables or interior determinant at one flight. The inverse quadratic endpoint Hessian has both diagonal entries  $\nu_1$ . The first varying action term is  $(u^{2m} + v^{2m})/(2m)!$  and its mixed derivative is zero. Hence the amplitude does not vary in degree  $2m - 2$ . Transform the quadratic ellipse to a disk. The even moment of either endpoint linear functional is the disk moment with its squared length  $\nu_1$  in place of  $1/a$ . The unweighted sublevel integral used in the proof of (A.4.2) gives exactly (A.4.8). The same finite-jet recursion proves germ determination. For  $m = 2, g = 1, c = 2$ , this derivative is  $-1/27$ ; multiplication by  $q'_s = -24$  gives  $8/9$ . The other identity is (A.3.5).  $\square$

The uniform long-bridge theorem establishes a limiting physical law and relative two-boundary factorization, with uniform convergence of every fixed statistical jet. It is not an assertion that increasing flight number creates contact parameters absent from one-flight nonlinear data. The distinction from the entire leading hierarchy is genuine at arbitrarily high finite jet order, as the realized families show.

#### APPENDIX A.5. FINITE-FLIGHT CONTACT RECOVERY AND COEFFICIENT SENSITIVITY

We continue in the identical-even contact class of Section A.4. The half-line inverse has a finite-flight counterpart with constants uniform in the flight number at each fixed jet order. Write  $Y = (u, v)^\top$  and

$$H_j = a \begin{pmatrix} \coth(j\gamma) & -\operatorname{csch}(j\gamma) \\ -\operatorname{csch}(j\gamma) & \coth(j\gamma) \end{pmatrix}, \quad L_i = \left( \frac{\sinh((j-i)\gamma)}{\sinh(j\gamma)}, \frac{\sinh(i\gamma)}{\sinh(j\gamma)} \right). \quad (\text{A.5.1})$$

For  $0 \leq i \leq j$  put  $\nu_i = L_i H_j^{-1} L_i^\top$ . For interior sites set

$$G_{ii} = \frac{\sinh(i\gamma) \sinh((j-i)\gamma)}{a \sinh(j\gamma)}.$$

Let  $w_0 = w_j = 1$  and  $w_i = 2$  for  $1 \leq i < j$ , and set  $\alpha_m = 2/[2^m(m+1)(m!)^2]$ .

**Theorem A.5.1** (Uniform finite-flight triangular inverse). *At each  $j \geq 1$  and  $m \geq 2$ , the normalized coefficient  $f_{j,m-1}$  depends only on  $g, \kappa, q_4, \dots, q_{2m}$ , is affine in its last argument, and satisfies*

$$-\partial_{q_{2m}} f_{j,m-1} = \alpha_m \left\{ \sum_{i=0}^j w_i \nu_i^m + 4m \sum_{i=1}^{j-1} G_{ii} \nu_i^{m-1} \right\} \geq \frac{2\alpha_m}{a^m} > 0. \quad (\text{A.5.2})$$

Fix  $M$  and a compact finite-jet set with  $g, \kappa$  in a compact positive box. The triangular maps  $J_{j,M} : q \mapsto (f_{j,1}, \dots, f_{j,M-1})$  have locally Lipschitz inverses on neighborhoods of their images whose radii and Lipschitz constants can be chosen independently of  $j$ . In particular, for coefficient errors at most  $\delta$  in such a neighborhood,

$$\|J_{j,M}^{-1}(\hat{f}) - q\| \leq C_M \delta. \quad (\text{A.5.3})$$

Using the limiting inverse instead gives  $C_M(\delta + \tau^j)$ . Every fixed finite-flight nonlinear germ also determines the analytic even contact at known  $g, \kappa$ . No constant uniform in  $M$  is claimed.

*Proof.* The finite-degree argument in the proof of Theorem A.4.1 applies to every finite bridge. To calculate its diagonal, vary only  $q_{2m}$  while fixing lower jets. The first varying term in one flight is  $(u^{2m} + v^{2m})/(2m)!$ . The linear stationary segment is  $x_i = L_i Y$ , so the envelope identity gives

$$\partial_{q_{2m}} W_j^{[2m]}(Y) = \frac{1}{(2m)!} \sum_{i=0}^j w_i (L_i Y)^{2m}.$$

The interior sites occur in two lengths and the endpoints in one. The first varying local length term has zero mixed derivative. Therefore the edge product contributes nothing at degree  $2m - 2$ . The first varying diagonal of the interior Hessian is  $2(L_i Y)^{2m-2}/(2m - 2)!$ . Differentiating its relative determinant gives

$$\partial_{q_{2m}} b_j^{[2m-2]}(Y) = -\frac{2}{(2m - 2)!} \sum_{i=1}^{j-1} G_{ii} (L_i Y)^{2m-2}.$$

The change of the stationary segment starts at degree  $2m - 1$ . The unvaried Hessian perturbation starts at degree two, since the graphs are even. Thus that change, products of Hessian perturbations, and products with lower nonconstant amplitude terms cannot enter the displayed degree. The same observation justifies using the quadratic sublevel to compute the coefficient under consideration.

The quadratic residual integral equals  $\pi d^2/\sqrt{\det H_j}$ . Changing variables from the ellipse to a disk, a linear functional  $L_i Y$  has squared metric length  $\nu_i$ . The unweighted sublevel moment from action variation contributes  $-\alpha_m w_i \nu_i^m$  to the normalized coefficient. The residual-weighted moment from amplitude variation contributes  $-4m\alpha_m G_{ii} \nu_i^{m-1}$ . These constants follow from the angular average  $\binom{2k}{k}/4^k$  and radial integration, exactly as in (A.4.6). The residual weight is zero on the sublevel boundary, so its first variation has no boundary term. Addition proves the identity in (A.5.2).

All its summands are positive. Moreover  $\nu_0 = \nu_j = a^{-1} \coth(j\gamma) \geq a^{-1}$ , proving the lower bound. This also proves that the coefficient is affine in the last jet: finite-degree dependence excludes higher jets, and the displayed derivative is independent of every higher and lower nonquadratic jet. The coefficient maps and their fixed derivatives are uniformly bounded on compact finite-jet sets by the differentiated threshold estimates. One may apply those estimates to the local even polynomial graphs with the specified finite jets on a sufficiently small common contact box; no unprescribed higher jet can enter these coefficients.

Solve successively the equations  $f_{j,m-1} = d_{j,m} q_{2m} + P_{j,m}(q_4, \dots, q_{2m-2})$ . The denominators  $d_{j,m}$  are bounded away from zero uniformly in  $j$  for  $m \leq M$ . The derivatives of the numerators are uniformly bounded on an enlarged compact finite-jet box. Induction gives bounded inverse derivatives, and a common admissible coefficient perturbation radius: at each step choose that radius small enough to keep the recovered lower jets in the enlarged box. There are only  $M - 1$  steps. This proves (A.5.3), without differentiating an infinite vector of jets. The limiting bound follows from Theorem A.4.1 and the triangle inequality. Equality of all recovered Taylor coefficients implies equality of analytic germs. It does not imply equality of merely smooth graphs with identical Taylor jets.  $\square$

For  $j = 1$  the interior sum is empty and  $\nu_0 = \nu_1 = g \cosh \gamma / \sinh^2 \gamma$ , so (A.5.2) reduces to (A.4.8). At each fixed  $m$ , the two separated endpoint layers give

$$\sum_i w_i \nu_i^m \longrightarrow 2a^{-m} \coth(m\gamma), \quad \sum_{i=1}^{j-1} G_{ii} \nu_i^{m-1} \longrightarrow a^{-m} D_m.$$

Indeed the left layer has  $\nu_i \rightarrow a^{-1}e^{-2i\gamma}$  and  $G_{ii} \rightarrow (1 - e^{-2i\gamma})/(2a)$ , and the right layer has its reversed copy. The common geometric majorants supplied by endpoint localization justify both sums. This recovers (A.4.2).

**Proposition A.5.2** (Limiting versus one-flight diagonal sensitivity). *Let  $D_m^\infty = |\partial_{q_{2m}} f_{m-1}|$  and  $D_m^1 = |\partial_{q_{2m}} f_{1,m-1}|$ . Then*

$$\frac{D_m^\infty}{D_m^1} = (\tanh \gamma)^m \{\coth(m\gamma) + 2mD_m\} = (\tanh \gamma)^m \{1 + O_\gamma(me^{-2(m-1)\gamma})\}. \quad (\text{A.5.4})$$

*Proof.* Divide (A.4.2) by (A.4.8) and use  $av_1 = \coth \gamma$ . Expanding  $(e^{2k\gamma} - 1)^{-1} = e^{-2k\gamma}/(1 - e^{-2k\gamma})$  gives the stated remainder at fixed  $\gamma > 0$ .  $\square$

At fixed lower jets and equal absolute error in the corresponding coefficient, the inverse diagonal amplification is the reciprocal of (A.5.4). This comparison is unchanged by using ordinary rather than factorial graph coefficients. It does not compare full triangular condition numbers, the cost of extracting coefficients from noisy probability values, or minimax risks of different experiments. The fixed- $M$  uniformity in Theorem A.5.1 is uniformity in  $j$ , not in the number of recovered jets. Exact analytic continuation in Corollary A.4.2 remains a uniqueness statement, not a noise-stable continuation procedure.

**A.5.1. The channel index in geometric realization.** The unchanged leading data in Theorem A.4.3 are all collision-order leading count and endpoint-metric coefficients of the selected horizontal ground channel, with its reverse orientation. The global free area is also fixed. Other selected channels need not have the same leading data. This distinction can be checked within its  $M = 2$  realization. Write  $h_s = 1 + s \sin^4 \theta + z(s) \sin^6 \theta$ . The area constraint gives

$$z'(0) = -\frac{\int_0^{2\pi} \sin^4 \theta \, d\theta}{\int_0^{2\pi} \sin^6 \theta \, d\theta} = -\frac{6}{5}.$$

For the clear vertical channel on  $3\mathbb{Z} \times 4\mathbb{Z}$ , put  $g_v = 4 - 2h_s(\pi/2)$  and  $r_v = (h_s + h_s'')(\pi/2)$ . Then

$$(g_v(0), g_v'(0)) = (2, 2/5), \quad (r_v(0), r_v'(0)) = (1, 3),$$

$$(c_v(0), c_v'(0)) = (3, -28/5), \quad c_v = 1 + g_v/r_v.$$

These equalities follow by differentiating the displayed trigonometric polynomials. Positive curvature and clearance persist by openness. Thus the oriented one-flight coefficient at its own onset satisfies

$$L_{1,v} = \frac{1}{2(12 - \pi)\sqrt{c_v^2 - 1}}, \quad \left. \frac{d}{ds} \log L_{1,v} \right|_{s=0} = \frac{21}{10}.$$

The horizontal realization is therefore a fibre of the complete leading hierarchy of a specified channel, not a fibre of all channel-resolved observations of the table. The general forward theorem and the full ground-onset count law retain their original scope.

## APPENDIX A.6. RECONSTRUCTION OF BOUNDARY PROFILES FROM PHYSICAL COUNTS

The stability theorem of Section 20 takes a function and three of its derivatives as data. Here the observations are instead finitely many binary reports from independent physical preparations. We recover both full profiles on a closed energy collar, including its origin. No derivatives of empirical proportions are supplied as observations.

Fix  $0 < D \leq d_*$  and an even flight number  $j$ . Write  $e_0 = e$ ,  $e_1 = \bar{e}$  and

$$p_{j,b}(d) = \Pr(E_{j,e_b}(d)), \quad a_j(d) = \frac{2A \sinh(j\gamma)}{d^2}, \quad F_{j,b}(d) = a_j(d)p_{j,b}(d). \quad (\text{A.6.1})$$

The gap  $g$ , free area  $A$ , multiplier parameter  $\gamma$ , and the selected facing patches are supplied. Thus the window  $g + d$  and its normalization are known before sampling. Each preparation reports only whether the selected event occurred; a failure is recorded and charged. This is a labelled binary experiment, not the unlabelled experiment of Section A.13.

**A.6.1. Uniform classes and the statement.** Let  $m \geq 4$ ,  $0 < \beta \leq 1$  and  $B \geq 1$ . Set

$$\|V\|_{m-1,1} = \sum_{r=0}^{m-1} \|V^{(r)}\|_{\infty} + \text{Lip}(V^{(m-1)}),$$

and let  $\mathcal{K} = \mathcal{K}_m(B, \beta; D)$  be the class of real  $C^{m-1,1}([0, D])$  functions with  $V(0) = 1$ ,  $V \geq \beta$  and  $\|V\|_{m-1,1} \leq B$ . The true profiles belong to  $\mathcal{K}$ . Assume supplied uniform constants  $C_0, M$  and  $0 < \tau < 1$  such that, for both orientations and every even  $j \geq 2$ ,

$$\|F_{j,b} - \mathcal{F}_{bb}\|_{C^3([0,D])} \leq C_0 \tau^j, \quad 0 \leq F_{j,b} \leq M. \quad (\text{A.6.2})$$

These are conclusions of Theorem 12.1 on each fixed compact smooth family, with  $B, \beta$  chosen for that family. The theorem below applies to any collection of physical experiments with these common bounds. It does not postulate that every member of  $\mathcal{K}$  is realizable by a billiard. The larger class is used only for regularization. Constants in this section may depend on  $m, D, B, \beta, C_0, M, A, \gamma, \tau$ , not on the accuracy or confidence.

For  $V, W \in \mathcal{K}$  define the explicit bilinear map

$$T(V, W)(d) = \frac{2}{\pi} \int_{s,t>0, s+t<1} \frac{1-s-t}{\sqrt{st}} V(ds) W(dt) \, ds \, dt, \quad T(V, W)(0) = 1. \quad (\text{A.6.3})$$

Thus  $\mathcal{F}_{bc} = T(\mathcal{V}_b, \mathcal{V}_c)$ .

**Theorem A.6.1** (Finite-preparation recovery of full profiles). *Under (A.6.2), choose an integer  $L \geq m$ ,  $h = D/L \leq 1$ , an even  $j \geq 2$ , and  $0 < t, \delta \leq 1$ . At each positive node  $d_k = kh$ ,  $1 \leq k \leq L$ , and for each  $b = 0, 1$ , make*

$$N_{b,k} = \left\lceil 2a_j(d_k)(M + t/3)t^{-2} \log \frac{4L}{\eta} \right\rceil \quad (\text{A.6.4})$$

*independent preparations, where  $0 < \eta < 1/2$ . There is a measurable finite-dictionary estimator  $(\hat{V}_0, \hat{V}_1) \in \mathcal{K}^2$ , defined below, for which with probability at least  $1 - \eta$ ,*

$$\max_{b=0,1} \|\hat{V}_b - \mathcal{V}_b\|_{\infty} \leq C(h^{m-3} + \tau^j + th^{-3} + \delta). \quad (\text{A.6.5})$$

*Its prediction  $T(\hat{V}_0, \hat{V}_1)$  of the odd limiting law has at most  $2B$  times the error on the left. Its deterministic total preparation count satisfies*

$$N_{\text{tot}} \leq 2L + CA \sinh(j\gamma) h^{-2} t^{-2} \log \frac{4L}{\eta}. \quad (\text{A.6.6})$$

*Only positive-offset, even-flight observations are used. Dictionary cardinality and computational running time are not bounded by this preparation estimate.*

**Corollary A.6.2** (An explicit accuracy budget). *For sufficiently small  $\varepsilon > 0$ , the two profiles and the odd limiting law can be recovered in uniform norm to accuracy  $\varepsilon$  with confidence  $1 - \eta$ , using at most*

$$C\varepsilon^{-(2+\frac{8}{m-3}+\frac{\gamma}{|\log \tau|})} \log \frac{C}{\eta\varepsilon} \quad (\text{A.6.7})$$

*independent preparations. The windows remain in the fixed collar, and the even flight number is  $O(\log(1/\varepsilon))$ . This is a sufficient cost in the stated labelled experiment; no optimality over designs or billiard observations is asserted.*

**A.6.2. Positive-node differentiation and regularization.** For each cell  $J_\ell = [\ell h, (\ell + 1)h]$ ,  $0 \leq \ell < L$ , put

$$s_\ell = \max\{1, \min\{\ell, L - m + 1\}\}.$$

Let  $I_h z$  on  $J_\ell$  be the polynomial of degree at most  $m - 1$  interpolating  $z_k$  at the  $m$  nodes  $d_{s_\ell}, \dots, d_{s_\ell + m - 1}$ . We also write  $I_h f$  for interpolation of the sampled values of a function  $f$ . Polynomials on adjacent cells need not have matching derivatives. Their norm is

$$\|P\|_{3,h} = \max_{0 \leq \ell < L} \max_{0 \leq r \leq 3} \|P^{(r)}\|_{C^0(J_\ell)}.$$

The stencil on the first cell uses  $h, \dots, mh$ . There is neither a preparation at zero nor an assumed exact value of  $F_{j,b}(0)$ . The limited extrapolation to zero uses the same fixed stencil bounds.

**Lemma A.6.3** (Three interpolation estimates). *With constants depending only on  $m, D$ , one has*

$$\|I_h f - f\|_{3,h} \leq Ch^{m-3} \|f\|_{m-1,1}, \quad f \in C^{m-1,1}([0, D]), \quad (\text{A.6.8})$$

$$\|I_h f\|_{3,h} \leq C \|f\|_{C^3([0, D])}, \quad f \in C^3([0, D]), \quad (\text{A.6.9})$$

$$\|I_h z\|_{3,h} \leq Ch^{-3} \max_k |z_k|. \quad (\text{A.6.10})$$

*In particular a smooth  $C^3$  model error is not amplified by  $h^{-3}$ , whereas an arbitrary scalar error can be.*

*Proof.* Rescale a cell by  $d = \ell h + hu$ . Its interpolation nodes are distinct integers lying within distance  $m$  of  $[0, 1]$ . There are only finitely many possible relative stencils. Their Lagrange basis polynomials and derivatives through order three are therefore bounded uniformly on  $[0, 1]$ . Scaling back gives  $\|(I_h z)^{(r)}\|_{C^0(J_\ell)} \leq Ch^{-r} \max |z_k|$ ; since  $h \leq 1$ , this proves (A.6.10).

For the first estimate, subtract the degree- $(m - 1)$  Taylor polynomial of  $f$  at  $\ell h$ . The Lipschitz bound on its last derivative gives remainder  $C\|f\|_{m-1,1}h^{m-r}$  after  $r \leq 3$  derivatives on the cell, and remainder  $C\|f\|_{m-1,1}h^m$  at every stencil node. The interpolant reproduces the Taylor polynomial. Applying the scaled basis bound to the nodal remainder and adding the differentiated function remainder proves (A.6.8).

For (A.6.9), instead subtract the quadratic Taylor polynomial. Its interpolant is itself, since  $m \geq 4$ ; its derivatives through order three are bounded by  $C\|f\|_{C^3}$ . Its nodal remainder is bounded by  $Ch^3\|f\|_{C^3}$ , including in the first and last cells. Scaling the remainder gives  $Ch^{3-r}\|f\|_{C^3}$  for  $r \leq 3$ , proving the assertion.  $\square$

**Lemma A.6.4** (Compact regularization class). *The class  $\mathcal{K}$  is compact in  $C^{m-1}([0, D])$ . The functions  $T(V, V)$ ,  $V \in \mathcal{K}$ , have uniformly bounded  $C^{m-1,1}$  norms. For each  $h$  and  $\delta > 0$ , there is a finite ordered set  $\mathcal{D}_{h,\delta} \subset \mathcal{K}$  such that*

$$\inf_{W \in \mathcal{D}_{h,\delta}} \|I_h T(V, V) - I_h T(W, W)\|_{3,h} \leq \delta \quad (V \in \mathcal{K}). \quad (\text{A.6.11})$$

Moreover for  $V, W \in \mathcal{K}$ ,

$$\|V - W\|_\infty \leq C \|T(V, V) - T(W, W)\|_{C^3([0, D])}. \quad (\text{A.6.12})$$

*Proof.* The derivatives through order  $m - 1$  are uniformly bounded and equicontinuous. Successive use of the fundamental theorem of calculus identifies the limits of these derivatives under an Arzelà–Ascoli subsequence. The norm bound, positivity, normalization and Lipschitz bound pass to the limit. This proves compactness.

The weight in (A.6.3) is integrable and has total mass one. Differentiating under the integral through order  $m - 1$  uses bounded products of derivatives of  $V$ , with factors  $s, t \leq 1$ . The same product estimate bounds the Lipschitz seminorm of the last derivative. It gives a uniform  $C^{m-1,1}$

bound. At a fixed grid,  $V \mapsto I_h T(V, V)$  is continuous even for the uniform topology: the forward nodal values are continuous, and interpolation is a linear map between finite-dimensional spaces. Its image is compact, so it admits a finite  $\delta$ -net with representatives in  $\mathcal{K}$ . Fix any order of those representatives in advance.

Finally the proof of Theorem 20.2 uses only normalization, the first-derivative bound and the three derivatives of the forward law. It applies to the present  $C^{m-1,1}$  functions without a smoothness extension. More explicitly, put  $f = x^{-1/2}(V - W)$  and  $a = x^{-1/2}(V + W)$ . Their Abel convolution has the form  $x^{-1/2} * a = 2\pi + R$ , with  $R(0) = 0$  and  $\|R'\|_\infty \leq \pi B$ . If  $Q = (\pi/2)(d^2[T(V, V) - T(W, W)])''$ , then  $2\pi f + R' * f = x^{-1/2} * Q'$ . The normalization makes  $f$  continuous at zero. The Abel bound and Gronwall give  $\|V - W\|_\infty \leq (D/\pi)e^{BD/2}\|Q'\|_\infty$ . The identity

$$Q' = \frac{\pi}{2} (6\Delta F' + 6d\Delta F'' + d^2\Delta F'''), \quad \Delta F = T(V, V) - T(W, W),$$

proves (A.6.12) in the stated regularity class.  $\square$

Given nodal observations  $z = (z_k)_{k=1}^L$ , define

$$\hat{V}(z) = \text{the first minimizer over } W \in \mathcal{D}_{h,\delta} \text{ of } \|I_h T(W, W) - I_h z\|_{3,h}. \quad (\text{A.6.13})$$

Each criterion is continuous in the finite data vector. The prescribed tie rule makes the estimator Borel measurable. Neither the dictionary nor its order depends on the unknown profile or the observed counts. This construction specifies a statistical estimator; it does not claim a tractable enumeration of the dictionary.

### A.6.3. Proof of the acquisition theorem.

*Proof of Theorem A.6.1.* For orientation  $b$  and node  $k$ , let  $\bar{X}_{b,k}$  be the empirical proportion of successes and put  $Z_{b,k} = a_j(d_k)\bar{X}_{b,k}$ . Its expectation is exactly  $F_{j,b}(d_k)$ , and its variance is

$$\text{Var}(Z_{b,k}) = \frac{a_j(d_k)F_{j,b}(d_k)(1 - p_{j,b}(d_k))}{N_{b,k}}. \quad (\text{A.6.14})$$

For a Bernoulli variable  $X$  of mean  $p$ ,  $\log \mathbb{E}e^{s(X-p)} \leq p(e^s - 1 - s)$ . The exponential series gives  $e^s - 1 - s \leq s^2/[2(1 - s/3)]$  for  $0 < s < 3$ ; for the negative exponent use  $e^{-s} - 1 + s \leq s^2/2$ . Independence, exponential Markov bounds, and the parameter  $s = t/(M + t/3)$  for  $X - p$  give

$$\Pr\{|Z_{b,k} - F_{j,b}(d_k)| > t\} \leq 2 \exp \left\{ -\frac{N_{b,k}t^2}{2a_j(d_k)(M + t/3)} \right\}. \quad (\text{A.6.15})$$

The allocation (A.6.4) and a union bound over  $2L$  observations yield an event of probability at least  $1 - \eta$  on which all scalar errors are at most  $t$ .

Work on this event and fix  $b$ . Abbreviate  $\mathcal{F} = \mathcal{F}_{bb}$  and  $\hat{F} = T(\hat{V}_b, \hat{V}_b)$ , where  $\hat{V}_b = \hat{V}((Z_{b,k})_k)$ . The last two bounds of Lemma A.6.3 give

$$\|I_h Z_b - I_h \mathcal{F}\|_{3,h} \leq C(th^{-3} + C_0\tau^j). \quad (\text{A.6.16})$$

In particular the deterministic relative error in (A.6.2) has not been multiplied by  $h^{-3}$ . Choose a dictionary representative for the true profile using (A.6.11). The minimizing property and the triangle inequality imply

$$\|I_h \hat{F} - I_h \mathcal{F}\|_{3,h} \leq \delta + 2\|I_h Z_b - I_h \mathcal{F}\|_{3,h}.$$

The first interpolation bound, applied separately to  $\hat{F}$  and  $\mathcal{F}$ , now gives

$$\|\hat{F} - \mathcal{F}\|_{C^3([0,D])} \leq C(h^{m-3} + th^{-3} + \tau^j + \delta).$$



Both functions are globally  $C^3$ ; the broken norm of their difference therefore bounds each of its derivatives on the whole closed collar. Apply (A.6.12) to prove (A.6.5). The mass-one positive weight in (A.6.3) gives

$$\|T(V_0, V_1) - T(W_0, W_1)\|_\infty \leq B(\|V_0 - W_0\|_\infty + \|V_1 - W_1\|_\infty),$$

which proves the odd-law prediction bound.

To count preparations, the ceilings contribute at most  $2L$ . Since  $\sum_{k=1}^L d_k^{-2} \leq h^{-2} \sum_{k \geq 1} k^{-2} \leq 2h^{-2}$ , summing (A.6.4) gives (A.6.6). All preparations, including failures, have been counted. This completes the proof of Theorem A.6.1.  $\square$

*Proof of Corollary A.6.2.* Choose  $h = D/L$  comparable to a sufficiently small constant times  $\varepsilon^{1/(m-3)}$ , with  $L \geq m$ , and put  $t = c\varepsilon h^3$ ,  $\delta = c\varepsilon$ . Choose the least even  $j \geq 2$  with  $\tau^j \leq c\varepsilon$ . Writing  $\lambda = \lfloor \log \tau \rfloor$ , rounding changes  $j$  by less than two once the lower bound  $j \geq 2$  is inactive. Consequently

$$e^{j\gamma} \leq C\varepsilon^{-\gamma/\lambda}, \quad h^{-2}t^{-2} = c^{-2}\varepsilon^{-2}h^{-8}.$$

A small enough fixed  $c$ , and the corresponding constant in the mesh, make both profile errors and the odd-law prediction error at most  $\varepsilon$ . The logarithm of  $4L/\eta$  is bounded by  $C \log(C/(\eta\varepsilon))$ . Substitute these estimates into (A.6.6); its  $2L$  term is of smaller order. This proves (A.6.7) with a deterministic design.  $\square$

**A.6.4. What the budget measures.** The three accuracy-dependent contributions to the power in (A.6.7) have distinct origins. Scalar accuracy costs  $\varepsilon^{-2}$ . Resolving three derivatives while sampling near a quadratic onset costs  $h^{-8}$ . Suppressing finite-bridge bias costs  $e^{j\gamma}$ , with  $j$  chosen from the relative decay  $\tau^j$ . No equality between  $\tau$  and  $e^{-\gamma}$  is assumed; a slower certified relative rate gives a larger sufficient budget.

The same experiment was implicit in the direct-variance calculation of the referee memorandum [20]. Formula (A.6.14) retains that cost rather than treating the limiting law as a freely observed function. Regularization is needed because the inverse differentiates the data. Its finite smoothness bound is an additional quantitative assumption for acquisition, not a new restriction on the general smooth forward theorem. The recovered objects remain the weighted, symmetrized energy profiles, not arbitrary asymmetric boundary graphs. The two oriented even sequences suffice for their reconstruction and for prediction of the odd law. Separate odd records provide an independent compatibility test.

## APPENDIX A.7. ACQUISITION FROM INTEGRATED PHYSICAL FLUX

The preceding construction controls three integer derivatives after normalizing every observation by  $d^2$ . We now retain the integrated flux and use the linear Abel discrepancy of Section 47. This changes both the approximation loss and the allocation of preparations. The observation is still exactly the selected-event bit; no endpoint coordinate or derivative is added.

Keep  $\mathcal{K}_m(B, \beta; D)$ , the supplied  $g, A, \gamma$ , the labels and (A.6.2). Write

$$c_j = 2A \sinh(j\gamma), \quad H_{j,b}(d) = c_j p_{j,b}(d) = d^2 F_{j,b}(d), \quad H_b = H_{\mathcal{V}_b}. \quad (\text{A.7.1})$$

The relative law gives  $\|H_{j,b} - H_b\|_{C^3} \leq C\tau^j$ . Multiplication by  $d^2$  has not altered that estimate. In contrast, the variance of a scaled empirical integrated flux decreases quadratically at the onset.

**Theorem A.7.1** (Abel-discrepancy acquisition). *Let  $L \geq m$ ,  $h = D/L \leq 1$ ,  $j \geq 2$  be even, and  $0 < t, \delta \leq 1$ . At every positive node  $d_k = kh$  and in each orientation  $b = 0, 1$ , make*

$$N_{b,k} = \left\lceil 2c_j(Md_k^2 + t/3)t^{-2} \log \frac{4L}{\eta} \right\rceil \quad (\text{A.7.2})$$

independent preparations. There is a Borel finite-dictionary estimator  $(\widehat{V}_0, \widehat{V}_1) \in \mathcal{K}^2$  for which, with probability at least  $1 - \eta$ ,

$$\max_b \|\widehat{V}_b - \mathcal{V}_b\|_\infty \leq C\{h^{m-5/2} + \tau^j + th^{-5/2} + \delta\}. \quad (\text{A.7.3})$$

The prediction of the odd limiting law has at most  $2B$  times this error. Every preparation, including every failure, is included in the deterministic bound

$$N_{\text{tot}} \leq 2L + CA \sinh(j\gamma) h^{-1} t^{-2} \log \frac{4L}{\eta}. \quad (\text{A.7.4})$$

No value or derivative at zero is observed. Computational complexity of the finite dictionary is not included in this preparation bound.

**Corollary A.7.2** (A refined preparation bound). *For sufficiently small  $\varepsilon > 0$ , both complete profiles and their predicted odd law are recovered to uniform accuracy  $\varepsilon$ , with confidence  $1 - \eta$ , using at most*

$$C\varepsilon^{-(2+\frac{6}{m-5/2}+\frac{\gamma}{|\log \tau|})} \log \frac{C}{\eta\varepsilon} \quad (\text{A.7.5})$$

preparations. The flight number is even and of order  $\log(1/\varepsilon)$ . This is a sufficient bound for this design, not a minimax assertion. The factor involving  $\tau$  depends on the supplied relative convergence certificate; it is not asserted to be an invariant of the billiard.

**A.7.1. A globally smooth positive-node reconstruction.** The Abel discrepancy is applied to globally smooth functions, not to distributional derivatives across interpolation cells. We give a specific reconstruction with this property. Let

$$\omega(u) = \begin{cases} \exp\{-1/(1-u^2)\}, & |u| < 1, \\ 0, & |u| \geq 1, \end{cases} \quad \theta_\ell(x) = \frac{\omega(x/h - \ell)}{\sum_{r=0}^L \omega(x/h - r)}, \quad 0 \leq \ell \leq L.$$

The denominator is positive on  $[0, D]$ . For each  $\ell$ , let  $P_\ell z$  interpolate the  $m$  positive nodal values starting at

$$s_\ell = \max\{1, \min\{\ell - \lfloor (m-1)/2 \rfloor, L - m + 1\}\}.$$

Thus the nodes are  $s_\ell h, \dots, (s_\ell + m - 1)h$ . Define

$$J_h z = \sum_{\ell=0}^L \theta_\ell P_\ell z. \quad (\text{A.7.6})$$

Every polynomial of degree at most  $m-1$  is reproduced. At a positive grid point, only its central weight is nonzero and its polynomial contains that node, so the nodal values are interpolated. At zero, the first stencil uses  $h, \dots, mh$ ; it extrapolates rather than inserting a known datum. The output is  $C^\infty$  on the closed interval in the usual extension sense.

**Lemma A.7.3** (Smooth reconstruction in the Abel seminorm). *With constants depending only on  $m, D$ , one has*

$$|J_h f - f|_{\mathcal{A}} \leq Ch^{m-5/2} \|f\|_{m-1,1}, \quad f \in C^{m-1,1}, \quad (\text{A.7.7})$$

$$|J_h f|_{\mathcal{A}} \leq C \|f\|_{C^3}, \quad f \in C^3, \quad (\text{A.7.8})$$

$$|J_h z|_{\mathcal{A}} \leq Ch^{-5/2} \max_k |z_k|. \quad (\text{A.7.9})$$

In particular, a smooth bridge error and an arbitrary nodal error have different amplification factors.

*Proof.* At each  $x$  at most two weights are positive. A nearest integer to  $x/h$  is at distance at most  $1/2$ , so the denominator of the weights has a fixed positive lower bound. Its derivatives of order  $r$  are bounded by  $C_r h^{-r}$ , as are those of  $\theta_\ell$ . Each active stencil lies within distance  $(m+1)h$  of  $x$ . After rescaling, there are only finitely many Lagrange basis configurations; their derivatives of every order at most  $m-1$  on the relevant compact intervals are uniformly bounded. The product rule therefore gives

$$\|J_h z\|_{C^r} \leq C h^{-r} \max_k |z_k|, \quad 0 \leq r \leq 3.$$

For  $f \in C^{m-1,1}$ , subtract its degree- $(m-1)$  Taylor polynomial at the point at which a derivative is evaluated. The interpolator reproduces that polynomial; its nodal remainders are bounded by  $C h^m \|f\|_{m-1,1}$ . The differentiated remainder estimate and the product rule give

$$\|J_h f - f\|_{C^r} \leq C h^{m-r} \|f\|_{m-1,1}, \quad r = 2, 3. \quad (\text{A.7.10})$$

This argument is valid at both endpoints because all active stencils remain a bounded rescaled distance away. Subtracting instead the quadratic Taylor polynomial for  $f \in C^3$  gives  $\|J_h f\|_{C^3} \leq C \|f\|_{C^3}$ .

Apply (47.3). In (A.7.10) the product of the second- and third-derivative errors is of order  $h^{2m-5}$ ; its square root is  $h^{m-5/2}$ . The endpoint second-derivative error is  $O(h^{m-2})$ , a smaller bound when  $h \leq 1$ . For nodal noise the two derivative scales are  $h^{-2}$  and  $h^{-3}$ , whose geometric mean is  $h^{-5/2}$ . Finally the  $C^3$  stability just proved bounds the smooth-bias term without any negative power of  $h$ .  $\square$

**A.7.2. The estimator and its full preparation charge.** The maps  $V \mapsto H_V$  take  $\mathcal{K}$  into a bounded  $C^{m-1,1}$  set. This follows from Lemma A.6.4 and multiplication by  $d^2$ . For fixed  $h$ , the map  $V \mapsto \mathcal{A} J_h H_V$  has compact image in  $C^0$ : it is a continuous map of the compact class  $\mathcal{K}$  through a finite-dimensional vector of nodal values. Choose, independently of the unknown profile, a finite ordered set  $\mathcal{D}_{h,\delta}^\mathcal{A} \subset \mathcal{K}$  such that

$$\inf_{W \in \mathcal{D}_{h,\delta}^\mathcal{A}} |J_h H_V - J_h H_W|_\mathcal{A} \leq \delta \quad (V \in \mathcal{K}). \quad (\text{A.7.11})$$

For a data vector  $z$ , let  $\widehat{V}(z)$  be the first minimizer of  $|J_h H_W - J_h z|_\mathcal{A}$  over this ordered set. Each criterion is continuous in  $z$ ; the fixed tie rule makes the estimator Borel. The affine nullspace of  $\mathcal{A}$  causes no problem for this finite minimization or for the norm on actual profile differences.

*Proof of Theorem A.7.1.* Let  $Z_{b,k} = c_j \bar{Y}_{b,k}$ , with  $\bar{Y}$  the empirical success proportion. Its mean is  $H_{j,b}(d_k)$ , and

$$\text{Var}(Z_{b,k}) = \frac{c_j H_{j,b}(d_k)(1 - p_{j,b}(d_k))}{N_{b,k}} \leq \frac{c_j M d_k^2}{N_{b,k}}.$$

The centered summands are bounded by  $c_j$ . The elementary Bernoulli exponential-moment bound used in Section A.6 therefore gives

$$\Pr\{|Z_{b,k} - H_{j,b}(d_k)| > t\} \leq 2 \exp \left\{ -\frac{N_{b,k} t^2}{2c_j(M d_k^2 + t/3)} \right\}.$$

The allocation and a union bound over  $2L$  means give simultaneous nodal error at most  $t$  with probability  $1 - \eta$ .

On this event, Lemma A.7.3 and the relative  $C^3$  bound give

$$|J_h Z_b - J_h H_b|_\mathcal{A} \leq C(th^{-5/2} + \tau^j).$$

The minimizing property and (A.7.11) bound  $|J_h H_{\widehat{V}_b} - J_h H_b|_{\mathcal{A}}$  by twice this quantity plus  $\delta$ . The interpolation error for the two candidate fluxes then gives

$$|H_{\widehat{V}_b} - H_b|_{\mathcal{A}} \leq C(h^{m-5/2} + th^{-5/2} + \tau^j + \delta).$$

Here the triangle inequality is that of the linear Abel seminorm. The inverse estimate (47.2) proves (A.7.3). The positive mass-one bilinear kernel of  $T$  gives the stated odd-law error.

Finally  $\sum_{k=1}^L d_k^2 \leq D^2 L$  and  $L = D/h$ . Summing (A.7.2), including each ceiling, yields (A.7.4). This sums the allocations themselves, not the number of successes or a waiting time conditional on success.  $\square$

*Proof of Corollary A.7.2.* Put  $\nu = m - 5/2 > 0$ . Choose  $L$  so that  $h$  is comparable to a sufficiently small multiple of  $\varepsilon^{1/\nu}$ , take  $t = c\varepsilon h^{5/2}$  and  $\delta = c\varepsilon$ , and choose the least even  $j \geq 2$  with  $\tau^j \leq c\varepsilon$ . All constants can be chosen to control also the odd-law error. Rounding  $j$  changes it by less than two once the lower endpoint is inactive, so  $\sinh(j\gamma) \leq C\varepsilon^{-\gamma/|\log \tau|}$ . Moreover

$$h^{-1}t^{-2} = c^{-2}\varepsilon^{-2}h^{-6}.$$

The logarithmic factor is bounded by  $C \log\{C/(\eta\varepsilon)\}$ , and the ceiling contribution  $2L$  is smaller than the displayed bound. This proves the claim.  $\square$

There are two separate improvements over Corollary A.6.2: the approximation/noise loss is  $5/2$  rather than three, and the sum of integrated-flux variances has order  $h^{-1}$  rather than the  $h^{-2}$  sum produced by dividing first by  $d_k^2$ . Neither improvement follows by merely substituting a mixed-norm inequality in the earlier discrepancy criterion. In particular, when  $m = 4$  the accuracy exponent becomes  $6 + \gamma/|\log \tau|$  rather than  $10 + \gamma/|\log \tau|$ . Both remain sufficient exponents with a supplied convergence certificate.

## APPENDIX A.8. CALIBRATION WITHIN THE SMOOTH PHYSICAL EXPERIMENT

The preceding two acquisition theorems supply the gap, area and multiplier. These quantities cannot be imported from an exact-family inverse when the unknown experiment has arbitrary smooth remainders. We instead construct a calibration procedure inside the same selected physical experiment. Patch labels, a coarse parameter box and common smoothness bounds remain supplied. The target continues to be the symmetrized energy invariant, not an arbitrary asymmetric contact graph.

**A.8.1. Why the integrated flux is stable under calibration.** Let  $D < D_+$  be fixed collars. In this subsection assume the relative  $C^3$  estimates on  $[0, D_+]$  and a supplied uniform  $C^{3,1}$  bound on  $H_{j,b}$  there. These follow on a fixed compact smooth physical family from the relative theorem with four derivatives. Suppose a pilot experiment supplies positive values  $\widehat{A}, \widehat{\gamma}$  and an upper gap estimate  $\widehat{g}$ . For a programmed window  $j\widehat{g} + d$ , put

$$\Delta = j(\widehat{g} - g), \quad R = \frac{\widehat{A} \sinh(j\widehat{\gamma})}{A \sinh(j\gamma)}.$$

The mean of the scaled bit  $2\widehat{A} \sinh(j\widehat{\gamma})Y$  is exactly

$$RH_{j,b}(d + \Delta). \tag{A.8.1}$$

This is a translation of an integrated flux followed by a constant rescaling. It has no factor  $d^{-2}$ . A small calibration error can therefore be treated as a smooth model error in the Abel discrepancy.

**Theorem A.8.1** (Stability of the plug-in experiment). *Assume  $A, \hat{A}$  and  $\gamma, \hat{\gamma}$  lie in fixed positive compact intervals, with the multiplier interval bounded away from zero. On a pilot event suppose*

$$0 \leq \hat{g} - g \leq \rho, \quad |\hat{A}/A - 1| \leq \alpha, \quad |\hat{\gamma} - \gamma| \leq \zeta, \quad j\zeta \leq 1, \quad j\rho \leq D_+ - D.$$

*There is a constant  $M_H$ , determined by the supplied bounds, which bounds the means in (A.8.1) on this event. Conditionally on any pilot outcome in the event, take fresh independent bits at the positive grid windows in both orientations. Set  $\hat{c}_j = 2\hat{A} \sinh(j\hat{\gamma})$  and allocate*

$$N_{b,k} = \left\lceil 2\hat{c}_j(M_H + t/3)t^{-2} \log \frac{4L}{\eta_1} \right\rceil. \quad (\text{A.8.2})$$

*The Abel-dictionary estimator, applied to  $Z_{b,k} = \hat{c}_j \bar{Y}_{b,k}$ , then satisfies, with conditional probability at least  $1 - \eta_1$ ,*

$$\max_b \|\hat{V}_b - \mathcal{V}_b\|_\infty \leq C\{h^{m-5/2} + \tau^j + th^{-5/2} + \delta + \alpha + j\zeta + j\rho\}. \quad (\text{A.8.3})$$

*The allocation has the deterministic bound  $2L + C\hat{c}_j h^{-1} t^{-2} \log(4L/\eta_1)$ . Thus calibration errors in this procedure are not multiplied by an inverse power of the grid spacing. If the pilot event has probability at least  $1 - \eta_0$ , the unconditional success probability is at least  $1 - \eta_0 - \eta_1$ .*

*Proof.* The true offset of the programmed window is  $d + \Delta$ . Applying (A.7.1) at this offset proves (A.8.1). On the fixed multiplier interval,

$$\left| \log \frac{\sinh(j\hat{\gamma})}{\sinh(j\gamma)} \right| \leq j \coth(j\gamma_-) \zeta \leq Cj\zeta.$$

The compact area bounds and  $j\zeta \leq 1$  then give  $|R - 1| \leq C(\alpha + j\zeta)$  and a fixed upper bound for  $R$ . The  $C^{3,1}$  bound and the fundamental theorem of calculus through order two, and the Lipschitz bound at order three, give

$$\|H_{j,b}(\cdot + \Delta) - H_{j,b}\|_{C^3([0,D])} \leq C\Delta.$$

The upper gap estimate ensures that only nonnegative offsets in the supplied collar occur, even when the translated function is evaluated at zero in this analysis. No observation there is made. Together with the relative law this proves

$$\|RH_{j,b}(\cdot + \Delta) - H_b\|_{C^3([0,D])} \leq C(\tau^j + \alpha + j\zeta + j\rho). \quad (\text{A.8.4})$$

It also supplies  $M_H$ . Conditionally on the pilot, the variance of a scaled mean is at most  $\hat{c}_j M_H / N_{b,k}$  and a summand is bounded by  $\hat{c}_j$ . The same Bernoulli bound and union argument as before therefore give nodal error at most  $t$  with conditional probability  $1 - \eta_1$ .

Apply (A.7.8) to (A.8.4), and (A.7.9) only to the centered scalar observations. The dictionary and inverse arguments of Theorem A.7.1 give (A.8.3). Summing the ceilings gives the budget; conditioning and a union bound give the final probability assertion.  $\square$

For comparison, normalizing first by  $d^2$  would give the mean  $R(1 + \Delta/d)^2 F_{j,b}(d + \Delta)$ , the sensitivity formula in [21]. The two calculations concern different processing of the same bits and are consistent. The present proof avoids the singular factor before regularization. In the quadratic benchmark  $H(d) = d^2$ , translation adds only  $2\Delta d + \Delta^2$ , which  $\mathcal{A}$  annihilates exactly. For a general smooth physical flux, (A.8.4) provides the required bound instead of asserting such an exact cancellation.

**A.8.2. A charged calibration procedure with smooth remainders.** We now specify sufficient uniform hypotheses for obtaining the pilot. Let the unknown gap lie in a supplied interval  $[g_-, g_+]$  of width  $w_0 > 0$ , with  $g_- > 0$  and  $w_0 \leq d_c$ . The area and multiplier lie in supplied boxes

$$0 < A_- \leq A \leq A_+ < \infty, \quad 0 < \gamma_- \leq \gamma \leq \gamma_+ < \infty.$$

For one selected orientation let  $P_r(t)$  denote the success probability of the  $r$ -flight event in the physical window  $t$ . For  $r = 1, 2$  assume

$$\begin{aligned} P_r(t) &= 0 \quad (t \leq rg), \\ P_r(rg + d) &= c_r d^2 F_r(d), \quad c_r = \{2A \sinh(r\gamma)\}^{-1}, \quad 0 \leq d \leq d_c, \\ F_r(0) &= 1, \quad F_r \geq \beta_c > 0, \quad \|F_r\|_{m-1,1} \leq B_c. \end{aligned} \quad (\text{A.8.5})$$

All constants in these conditions are supplied, not the functions  $F_r$ . The exact physical leading coefficient and smooth finite-flight law give these hypotheses after choosing a common local collar on a compact smooth family. No equality of the two remainders and no analytic finite-dimensional parameterization are assumed.

**Lemma A.8.2** (Finite-preparation calibration). *Under (A.8.5), for sufficiently small  $\xi > 0$  and  $0 < \eta_0 < 1/2$  there is a measurable procedure using only selected-event bits at flight numbers one and two which outputs  $\hat{g} \in [g_-, g_+]$ ,  $\hat{A} \in [A_-, A_+]$  and  $\hat{\gamma} \in [\gamma_-, \gamma_+]$ . With probability at least  $1 - \eta_0$ ,*

$$0 \leq \hat{g} - g \leq C\xi^{1+1/m}, \quad |\hat{A}/A - 1| + |\hat{\gamma} - \gamma| \leq C\xi. \quad (\text{A.8.6})$$

*Its total number of preparations, including failures, is deterministically at most*

$$C\xi^{-(2+2/m)} \log \frac{C \log(C/\xi)}{\eta_0}. \quad (\text{A.8.7})$$

*The constants are uniform over the specified smooth physical class.*

*Proof.* We first obtain a one-sided gap bracket. The bounds imply a supplied  $c_* > 0$  such that  $P_1(g + d) \geq c_* d^2$  on  $[0, d_c]$ . Fix a target bracket width  $\rho < w_0$ . Starting from  $[l, u]$  of width  $w$ , sample at  $q = l + 3w/4$ . Use

$$N(w) = \left\lceil \frac{16}{c_* w^2} \log \frac{2I}{\eta_0} \right\rceil, \quad I = 1 + \left\lceil \frac{\log(w_0/\rho)}{\log(4/3)} \right\rceil.$$

If any success occurs, replace the interval by  $[l, q]$ ; otherwise replace it by  $[l + w/2, u]$ . A success certifies  $g < q$ , because a success at or below the onset has probability zero. On a correct incoming bracket a false lower update can occur only if  $g \leq l + w/2$ . In that case  $q - g \geq w/4$  and  $q - g \leq w \leq d_c$ , so its probability is at most  $\exp\{-N(w)c_* w^2/16\} \leq \eta_0/(2I)$ . The bracket width decreases by a factor either  $3/4$  or  $1/2$  at each step, irrespective of the observations. Stop at width at most  $\rho$ . There are at most  $I$  steps. A union bound over correct incoming histories gives a correct final bracket with probability  $1 - \eta_0/2$ . Its upper endpoint  $\hat{g}$  then satisfies  $0 \leq \hat{g} - g \leq \rho$ .

The total of  $w^{-2}$  over the steps is at most  $\rho^{-2} \sum_{r \geq 0} (3/4)^{2r}$ : read the widths backward from the last pre-stopping width, which is larger than  $\rho$ . The ceilings add at most  $I$ . Consequently this stage costs at most  $C\rho^{-2} \log\{C \log(C/\rho)/\eta_0\}$  preparations. This bound is deterministic even on the exceptional incorrect-bracket outcomes.

Choose  $s = c_s \xi^{1/m}$  and  $\rho = c_\rho \xi s$  with fixed positive constants small enough that  $ms + 2\rho \leq d_c$  and  $\rho \leq s/4$ . At fresh preparations, for  $r = 1, 2$  and  $1 \leq k \leq m$ , sample the window  $r\hat{g} + ks$ . Let  $\bar{Y}_{r,k}$  be the corresponding proportion and put

$$\tilde{c}_r = \sum_{k=1}^m (-1)^{k-1} \binom{m}{k} \frac{\bar{Y}_{r,k}}{(ks)^2}. \quad (\text{A.8.8})$$

The weights reproduce the value at zero of every polynomial of degree at most  $m - 1$ . Taylor's formula with a Lipschitz highest derivative therefore bounds the extrapolation error for  $c_r F_r$  by  $Cs^m$ . If the gap bracket is correct, writing  $e = \hat{g} - g$  gives the exact mean at the  $k$ th normalized node as

$$c_r \left(1 + \frac{re}{ks}\right)^2 F_r(ks + re).$$

Its difference from  $c_r F_r(ks)$  is at most  $C\rho/s$ ; the finite sum of absolute weights depends only on  $m$ .

For completeness, at these windows  $P_r \leq Cs^2$  uniformly in  $k$ , since  $k \leq m$  and  $0 \leq re \leq 2\rho$ . A summand of  $\bar{Y}_{r,k}/(ks)^2$  is bounded by  $s^{-2}$  and its variance is at most  $Cs^{-2}$ . Thus

$$N_{r,k} = \left\lceil Cs^{-2} \xi^{-2} \log \frac{8m}{\eta_0} \right\rceil$$

with a sufficiently large supplied constant gives simultaneous normalized nodal error at most a fixed small multiple of  $\xi$  with conditional probability  $1 - \eta_0/2$ . This follows from the same Bernoulli exponential-moment inequality; its denominator is bounded by  $Cs^{-2}(1 + \xi)$ . Hence, on the joint good event,

$$|\tilde{c}_r - c_r| \leq C(s^m + \rho/s + \xi) \leq C\xi.$$

Project each  $\tilde{c}_r$  onto the known positive interval of possible  $c_r$  values and call the result  $\hat{c}_r$ . Projection does not increase its distance from the true  $c_r$ . Since  $c_1/(2c_2) = \cosh \gamma$ , project  $\hat{c}_1/(2\hat{c}_2)$  onto  $[\cosh \gamma_-, \cosh \gamma_+]$  and apply arcosh to obtain  $\hat{\gamma}$ . Finally project  $\{2\hat{c}_1 \sinh \hat{\gamma}\}^{-1}$  onto  $[A_-, A_+]$  to obtain  $\hat{A}$ . All these maps are Lipschitz on the specified boxes, since  $\gamma_- > 0$ . They give (A.8.6). The threshold decisions, finite sums, and projections are measurable.

The second stage costs  $Cs^{-2} \xi^{-2} \log(8m/\eta_0)$ . With the selected  $s, \rho$ , both  $s^{-2} \xi^{-2}$  and  $\rho^{-2}$  have order  $\xi^{-(2+2/m)}$ . Adding the two stages proves (A.8.7), including all ceilings and unsuccessful preparations.  $\square$

### A.8.3. The full physical experiment with unknown normalization.

**Theorem A.8.3** (Self-calibrated recovery of boundary profiles). *Consider a class of physical channels satisfying the general relative law, the supplied parameter boxes and calibration conditions (A.8.5), common profile bounds  $\mathcal{K}_m(B, \beta; D_+)$ , and common relative  $C^3$  and finite  $C^{3,1}$  flux bounds on  $[0, D_+]$ . Fix  $0 < D < D_+$ . Only the selected patch labels, these common bounds and the coarse gap bracket are known; the actual  $g, A, \gamma$  are unknown.*

*For sufficiently small  $\varepsilon > 0$  and  $0 < \eta < 1/2$ , a measurable procedure based exclusively on finitely many selected-event bits recovers both full profiles on  $[0, D]$  and predicts the odd law with uniform error at most  $\varepsilon$ , with probability at least  $1 - \eta$ . Its deterministic total preparation count is at most*

$$C\varepsilon^{-\left(2 + \frac{6}{m-5/2} + \frac{\gamma_+}{|\log \tau|}\right)} \log \frac{C}{\eta\varepsilon}. \quad (\text{A.8.9})$$

*The pilot uses flight numbers one and two; the profile stage uses one even flight number of order  $\log(1/\varepsilon)$ . No exact-law oracle or finite-dimensional remainder model is used. The conclusion is a uniform sufficient bound, not a minimax lower bound.*

*Proof.* Choose  $j, h, t, \delta$  as in Corollary A.7.2, with constants small enough to accommodate the plug-in error. In the pilot lemma take  $\xi = c\varepsilon/j$  and  $\eta_0 = \eta/2$ . On its good event,  $\alpha \leq C\xi$ ,  $\zeta \leq C\xi$  and  $\rho \leq C\xi^{1+1/m}$ . Thus

$$\alpha + j\zeta + j\rho \leq Cc\varepsilon$$

for sufficiently small  $\varepsilon$ . The same choice ensures  $j\zeta \leq 1$  and  $j\rho \leq D_+ - D$ . The latter fixed margin is why the theorem is stated on a prescribed subcollar. It never loses the energy origin.

Apply Theorem A.8.1 with fresh observations,  $\eta_1 = \eta/2$ , and a finite dictionary accurate enough also for the  $2B$  odd-law factor. The union bound and (A.8.3) give the error assertion.

The pilot estimates are projected into the supplied boxes even on bad outcomes. Consequently  $\hat{c}_j \leq 2A_+ \sinh(j\gamma_+)$  always, and the profile allocation has the deterministic bound in (A.8.9). Before absorption, the pilot adds at most

$$C(j/\varepsilon)^{2+2/m} \log \frac{C \log(Cj/\varepsilon)}{\eta}.$$

Since  $j \leq C \log(C/\varepsilon)$  and

$$2 + \frac{6}{m-5/2} + \frac{\gamma_+}{|\log \tau|} > 2 + \frac{2}{m},$$

this additional cost is absorbed into (A.8.9), uniformly for small  $\varepsilon$  and  $0 < \eta < 1/2$ . All observations are independent fresh preparations conditionally on the chosen windows; their success or failure is immaterial to the preparation charge.  $\square$

The unknown-normalization theorem is a new observation statement, not an interpretation of the earlier exact-family four-window theorem. Its smooth-remainder calibration is proved above and its pilot cost is included. Relative to a calibrated design controlled by the same supplied parameter box, it does not increase the sufficient accuracy exponent. This does not replace  $\gamma_+$  by an unknown intrinsic complexity, remove the convergence certificate, or remove patch labels. Nor does symmetrized energy recovery become recovery of the two asymmetric branches of a contact graph. Separately observed odd records test the compatibility predicted by the two recovered even laws; they are not a third identifying datum for this reconstruction.

#### APPENDIX A.9. CRITICAL BOUNDARY EXPERIMENTS AND SUPPORT SEPARATION

The relative comparison has a statistical scale which is not captured by weak convergence of a single record. We determine it first for the quadratic experiments, then transfer it to actual positive-offset laws in an explicit joint limit. The calculation allows unequal facing curvatures and either parity of the flight number. Throughout this section the channel and table are fixed; constants are uniform on the compact geometric families specified in Section 12. The two hypotheses compared here are the finite law and its boundary model for that same table, not two unknown billiard tables.

Write  $\iota = j \bmod 2$ ,  $a_0 = S_0''(0)$ ,  $a_\iota = S_\iota''(0)$ , and  $D = \text{diag}(\sqrt{a_0}, \sqrt{a_\iota})$ . Lemma 15.1 gives

$$H_j = DM_j D, \quad H_{\partial, \iota} = D^2, \quad M_j = \begin{pmatrix} \coth(j\gamma) & -\text{csch}(j\gamma) \\ -\text{csch}(j\gamma) & \coth(j\gamma) \end{pmatrix}. \quad (\text{A.9.1})$$

Both determinants equal  $a_0 a_\iota$ . On the common endpoint-residual coordinate space define

$$\mathcal{D}(H) = \{(z, s) \in \mathbb{R}^2 \times \mathbb{R} : s \geq 0, \ s + z^\top H z / 2 < 1\}.$$

Its volume is  $\pi / \sqrt{\det H}$ , as follows by first integrating  $s$  and then making a linear change to a disk. Let  $\Pi_j^0$  and  $\Pi_{\partial, j}^0$  have the common density  $\sqrt{a_0 a_\iota} / \pi$  on  $\mathcal{D}(H_j)$  and  $\mathcal{D}(H_{\partial, \iota})$ , respectively. The second subscript in  $\Pi_{\partial, j}^0$  retains parity. The actual conditional laws  $\Pi_{j, d}$ ,  $\Pi_{\partial, j, d}$  and the raw laws  $P_{j, d}$ ,  $P_{\partial, j, d}$  are those of Theorem 34.1; the boundary raw law retains the finite normalizing twist. Put

$$p_{j, d}^0 = \frac{d^2}{2A \sinh(j\gamma)}, \quad P_{j, d}^0 = (1 - p_{j, d}^0) \delta_{\dagger} + p_{j, d}^0 \Pi_j^0, \quad P_{\partial, j, d}^0 = (1 - p_{j, d}^0) \delta_{\dagger} + p_{j, d}^0 \Pi_{\partial, j}^0. \quad (\text{A.9.2})$$

Shrink the common collar so these are probability measures. We use  $\|P - Q\|_{\text{TV}} = \sup_B |P(B) - Q(B)|$ .



**Theorem A.9.1** (Exact support separation). *For every  $j \geq 1$ , independently of the two facing curvatures except through  $\gamma$ ,*

$$\delta_j := \|\Pi_j^0 - \Pi_{\partial,j}^0\|_{\text{TV}} = \frac{2}{\pi} \arcsin(e^{-j\gamma}). \quad (\text{A.9.3})$$

*For independent observations at any finite list  $(j_i, d_i)_{i=1}^n$ ,*

$$\left\| \bigotimes_i \Pi_{j_i}^0 - \bigotimes_i \Pi_{\partial,j_i}^0 \right\|_{\text{TV}} = 1 - \prod_i (1 - \delta_{j_i}), \quad (\text{A.9.4})$$

$$\left\| \bigotimes_i P_{j_i, d_i}^0 - \bigotimes_i P_{\partial, j_i, d_i}^0 \right\|_{\text{TV}} = 1 - \prod_i (1 - p_{j_i, d_i}^0 \delta_{j_i}). \quad (\text{A.9.5})$$

*As a simple binary experiment, each displayed pair of product laws is exactly equivalent, by measurable randomizations independent of the binary hypothesis, to an erasure experiment: with probability equal to the corresponding product on the right, the observation is an erasure, and otherwise it reveals the hypothesis. Consequently its minimum equal-prior testing error is half that product.*

*Proof.* The common transformation  $x = Dz/\sqrt{2}$  and a subsequent orthogonal rotation reduce every residual section to a disk and an equal-area ellipse. The eigenvalues of  $M_j$  are

$$\lambda = \tanh(j\gamma/2), \quad \lambda^{-1}, \quad 0 < \lambda < 1.$$

For the unit section the ellipse is  $\lambda x^2 + \lambda^{-1} y^2 < 1$ . Its boundary crosses the unit circle in the first quadrant at  $\theta_0 = \arctan \sqrt{\lambda}$ . Polar integration gives the intersection area

$$2 \left\{ \theta_0 + \int_{\theta_0}^{\pi/2} \frac{d\theta}{\lambda \cos^2 \theta + \lambda^{-1} \sin^2 \theta} \right\} = 4\theta_0.$$

For the equality, set  $t = \tan \theta$  in the integral; it becomes  $[\arctan(t/\lambda)]_{\sqrt{\lambda}}^{\infty} = \theta_0$ . Every residual section is obtained by the same dilation for both supports, so integration in  $s$  leaves the overlap fraction  $4\theta_0/\pi$ . Equal densities on the intersection imply  $\delta_j = 1 - 4\theta_0/\pi$ . If  $q = e^{-j\gamma}$ , then  $\lambda = (1 - q)/(1 + q)$  and  $\pi/2 - 2\arctan \sqrt{\lambda} = \arcsin q$ ; this proves (A.9.3). The whitening in (A.9.1) was common to both laws and loses no information. In particular the computation does not require equal curvatures or an even flight number.

For a pair with equal densities on its overlap, the measure with density the minimum of the two densities is exactly their common restriction. In the product the same property holds, and its mass is the product of the individual common masses. This proves (A.9.4). In the raw observation space the common restriction also includes the failure atom; its mass in observation  $i$  is  $1 - p_{j_i, d_i}^0 \delta_{j_i}$ . This proves (A.9.5).

Here is the asserted equivalence, including the reverse randomization. Send the common support to an erasure symbol, the support exclusive to the first law to 0, and the support exclusive to the second to 1. The common restricted measures of the two laws coincide. Given an erasure, draw from their normalized common restriction; given 0 or 1, draw from the corresponding normalized exclusive restriction. These kernels recover the two original laws and do not depend on which of the two hypotheses is true. Zero-mass restrictions may be assigned an arbitrary fixed probability. Under equal prior probabilities there is no error on an exclusive observation, and the posterior probabilities on an erasure remain equal. The minimum error is therefore half the common mass. This construction is for the fixed simple pair specified above; its kernel may depend on the known table and design.  $\square$

The next lemma controls the approximation error in the same topology before allowing the sample size to grow. Define

$$\omega(d) = \sqrt{d}, \quad \text{or } \omega(d) = d \text{ when both contact graphs are even.}$$

Evenness here does not require the two graphs to be identical.

**Lemma A.9.2** (Uniform tangent approximation). *On a common positive collar, uniformly in  $j$  and in the specified compact geometric family,*

$$\|\Pi_{j,d} - \Pi_j^0\|_{\text{TV}} + \|\Pi_{\partial,j,d} - \Pi_{\partial,j}^0\|_{\text{TV}} \leq C\omega(d), \quad (\text{A.9.6})$$

$$\|P_{j,d} - P_{j,d}^0\|_{\text{TV}} + \|P_{\partial,j,d} - P_{\partial,j,d}^0\|_{\text{TV}} \leq Cp_{j,d}^0\omega(d). \quad (\text{A.9.7})$$

*Proof.* Write  $E_j = W_j - jg$  and  $b_j = -W_{j,uv}/[-W_{j,uv}(0)]$ . The uniform Taylor estimates in Lemma 15.2 give, on bounded scaled endpoint sets,

$$\frac{E_j(\sqrt{d}z)}{d} = \frac{1}{2}z^\top H_j z + O(\sqrt{d}|z|^3), \quad b_j(\sqrt{d}z) = 1 + O(\sqrt{d}|z|).$$

All sublevels under discussion lie in a common bounded set by uniform coercivity. At each  $z$ , the actual and quadratic residual intervals start at zero and have lengths  $(1 - E_j(\sqrt{d}z)/d)_+$  and  $(1 - z^\top H_j z/2)_+$ . Their symmetric-difference length is at most  $C\sqrt{d}|z|^3$ . Integrating it and the amplitude error over that bounded set bounds the total variation norm of the unnormalized densities by  $C\sqrt{d}$ . Their normalizing integrals are bounded away from zero; dividing therefore proves the first finite-law estimate.

The same argument applies to  $S_0(u) + S_\ell(v)$  and  $B_0(u)B_\ell(v)$ . Their uniform derivatives follow from Lemma 17.1. If both graphs are even, uniqueness of the stationary segment makes  $E_j$  and  $b_j$  invariant under simultaneous endpoint sign reversal. Their remainders improve to  $O(d|z|^4)$  and  $O(d|z|^2)$ , respectively; the half-line expressions have the same improvement. This proves (A.9.6) with either choice of  $\omega$ .

The successful subprobabilities have the common scalar prefactor  $d^2[-W_{j,uv}(0)]/(2\pi A)$ . The quadratic integral is  $\pi/\sqrt{a_0 a_\ell}$ , and the Hessian identity gives  $[-W_{j,uv}(0)]/\sqrt{a_0 a_\ell} = 1/\sinh(j\gamma)$ . Multiplying the unnormalized estimate by this prefactor bounds the successful-measure difference by  $Cp_{j,d}^0\omega(d)$ . The difference of the failure masses is no larger than that bound. Restoring these atoms proves (A.9.7).  $\square$

**Theorem A.9.3** (Joint physical experiment limit). *For any finite independent schedule in the common collar, the actual positive-offset laws satisfy*

$$\left\| \left\| \bigotimes_i \Pi_{j_i, d_i} - \bigotimes_i \Pi_{\partial, j_i, d_i} \right\|_{\text{TV}} - \left( 1 - \prod_i (1 - \delta_{j_i}) \right) \right\| \leq C \sum_i \omega(d_i), \quad (\text{A.9.8})$$

$$\left\| \left\| \bigotimes_i P_{j_i, d_i} - \bigotimes_i P_{\partial, j_i, d_i} \right\|_{\text{TV}} - \left( 1 - \prod_i (1 - p_{j_i, d_i}^0 \delta_{j_i}) \right) \right\| \leq C \sum_i p_{j_i, d_i}^0 \omega(d_i). \quad (\text{A.9.9})$$

In particular, let  $j \rightarrow \infty$ ,  $d_j > 0$ , and set  $q_j = e^{-j\gamma}$ . For  $k_j$  conditional successes, if

$$k_j q_j \rightarrow b \in [0, \infty], \quad k_j \omega(d_j) \rightarrow 0, \quad (\text{A.9.10})$$

then their total variation distance from the boundary experiment tends to  $1 - e^{-2b/\pi}$ . For  $n_j$  raw preparations the same limit holds if

$$n_j p_{j, d_j}^0 q_j \rightarrow b \in [0, \infty], \quad n_j p_{j, d_j}^0 \omega(d_j) \rightarrow 0. \quad (\text{A.9.11})$$

At  $b = \infty$  the limit is interpreted as one. The minimum equal-prior testing errors tend to  $e^{-2b/\pi}/2$ . These limits hold for unequal facing curvatures and across both parity subsequences.

*Proof.* Telescoping product measures gives  $\|\otimes_i P_i - \otimes_i Q_i\|_{\text{TV}} \leq \sum_i \|P_i - Q_i\|_{\text{TV}}$ . Apply this separately on both sides of each simple binary comparison, use the reverse triangle inequality for their distances, and invoke Lemma A.9.2 and Theorem A.9.1. This proves the finite inequalities.

Now  $\delta_j = (2/\pi)q_j + O(q_j^3)$ . If  $k_j q_j$  has a finite limit,  $k_j \delta_j^2 \rightarrow 0$  and  $k_j \log(1 - \delta_j) \rightarrow -2b/\pi$ ; if it diverges, use  $(1 - \delta_j)^{k_j} \leq e^{-k_j \delta_j}$ . This proves the conditional limit. The raw calculation is identical with  $p_{j,d_j}^0 \delta_j$  in place of  $\delta_j$ ; this quantity tends to zero, and the required first-order sum is  $(2/\pi)n_j p_{j,d_j}^0 q_j (1 + O(q_j^2))$ . The error terms vanish by the second condition in each displayed scaling assumption. For any simple pair  $P, Q$ , integration of the smaller of their two densities gives minimum equal-prior error  $(1 - \|P - Q\|_{\text{TV}})/2$ . This proves the final assertion.  $\square$

The conditions are nonempty. In the general smooth case take  $d_j = q_j^{2+\epsilon}$ , with any fixed  $\epsilon > 0$ , and  $k_j \sim b/q_j$ , or  $n_j \sim b/(p_{j,d_j}^0 q_j)$ , for  $0 < b < \infty$ . In the even case  $d_j = q_j^{1+\epsilon}$  suffices. Thus the critical number of successes is of order  $e^{j\gamma}$ , while the critical number of raw preparations is of order  $d_j^{-2} e^{2j\gamma}$ . These are distinct resources. In the tangent experiment the probability that every observation lies in the common support is the product in Theorem A.9.1; its limiting exponential is the source of the testing profile, rather than a Gaussian fluctuation argument.

More generally, in a triangular schedule, if the largest  $\delta_{j_i}$  tends to zero,  $\sum_i \delta_{j_i} \rightarrow \lambda$ , and  $\sum_i \omega(d_i) \rightarrow 0$ , the conditional distance tends to  $1 - e^{-\lambda}$ . Replace  $\delta_{j_i}$  by  $p_{j_i,d_i}^0 \delta_{j_i}$  and the last sum by  $\sum_i p_{j_i,d_i}^0 \omega(d_i)$  for raw preparations. The logarithm argument in the proof applies without change. The identity maps to the tangent laws and the kernels in Theorem A.9.1 also give two-way randomizations whose errors tend to zero under these same assumptions. This last statement concerns the specified simple binary experiments only.

The uniform relative nonlinear theorem is still needed outside this joint tangent regime: at fixed positive  $d$  it compares the actual nonlinear boundary law, whereas replacing it by its quadratic law would retain an  $O(\omega(d))$  error in each success. The rate  $\tau$  in Theorem 34.1 is not asserted to be optimal. Nor do the preceding identities extend total variation to full embedded growing collision arrays. They show, in an actual physical joint limit, why exponentially small support differences and the number of successful preparations must be controlled together.

**A.9.1. Adaptive critical experiments.** The exact overlap calculation also has a stopped, adaptive form. The relevant overlap mass is generally not a product of unconditional one-step overlaps, because later windows depend on earlier records. We give its exact expression and a physical approximation bound. Throughout this subsection the comparison is the same simple binary pair as in Section A.9: a fixed table's finite law versus its boundary model. The policy is independent of the binary hypothesis, but may use the specified table and the past records.

For a design  $a = (j, d)$  write  $P_a^0 = P_{j,d}^0$  and  $Q_a^0 = P_{\partial,j,d}^0$ . Define

$$h(a) = p_{j,d}^0 \delta_j, \quad \delta_j = \frac{2}{\pi} \arcsin(e^{-j\gamma}), \quad p_{j,d}^0 = \frac{d^2}{2A \sinh(j\gamma)}. \quad (\text{A.9.12})$$

Their densities agree on the common support, including the failure atom. Let  $C_a$  be their common restriction and put  $R_a = C_a/(1 - h(a))$ . The chosen collar ensures  $h(a) < 1$ . For a policy  $\pi$  with cap  $N$ , let  $\mathbf{R}^\pi$  be the law obtained by following the same policy with record kernel  $R_a$ . This auxiliary law is supported on histories that have not revealed the binary hypothesis. Its stopping rule is still the specified rule of  $\pi$ .

**Theorem A.9.4** (Exact stopped overlap and physical transfer). *The tangent stopped transcript laws satisfy*

$$\|\mathbf{P}^{0,\pi} - \mathbf{Q}^{0,\pi}\|_{\text{TV}} = 1 - m_\pi, \quad m_\pi = \mathbb{E}_{\mathbf{R}^\pi} \prod_{i=1}^T (1 - h(a_i)). \quad (\text{A.9.13})$$

As a simple binary experiment, this pair is exactly equivalent by hypothesis-independent randomizations to an erasure experiment with erasure probability  $m_\pi$ . Its minimum equal-prior testing error is  $m_\pi/2$ .

Let  $\mathbf{P}^\pi, \mathbf{Q}^\pi$  instead be the actual positive-offset finite and boundary transcript laws. All reachable designs are assumed to lie in the common collar. Put

$$\varepsilon_\pi = \mathbb{E}_{\mathbf{P}^\pi} \sum_{i=1}^T p_{a_i}^0 \omega(d_i) + \mathbb{E}_{\mathbf{Q}^\pi} \sum_{i=1}^T p_{a_i}^0 \omega(d_i), \quad (\text{A.9.14})$$

where  $\omega(d) = \sqrt{d}$ , or  $\omega(d) = d$  when both contact graphs are even. Uniformly on the compact geometric families under consideration,

$$|\|\mathbf{P}^\pi - \mathbf{Q}^\pi\|_{\text{TV}} - (1 - m_\pi)| \leq C\varepsilon_\pi. \quad (\text{A.9.15})$$

The randomizations in the first assertion concern the specified known simple pair and policy, not an unknown family of billiard tables.

*Proof.* At each design, the measures  $P_a^0, Q_a^0$  have the same restriction  $C_a$  to their overlap; their other restrictions are supported on disjoint sets. A transcript compatible with both hypotheses can therefore contain only common-support records at every step. Conversely, along any such history every observation density agrees under the two hypotheses. The design kernels and stopping probabilities agree as well, since they are functions of the same observed past. Thus the two full transcript densities agree on their overlap. A first exclusive record makes the entire transcript exclusive, regardless of all later decisions.

The common transcript restriction is obtained by using the subprobability kernels  $C_a$  at active steps and the original policy kernels for choosing designs and stopping. Since  $C_a = (1 - h(a))R_a$ , its density relative to  $\mathbf{R}^\pi$  is

$$Z = \prod_{i=1}^T (1 - h(a_i)).$$

This identity follows by induction over the at most  $N$  active steps, keeping the stopping branch at each step; it also covers  $T = 0$  by the empty-product convention. Its total mass is  $m_\pi$ . Equal densities on the overlap give total variation  $1 - m_\pi$ .

Map the overlap to an erasure symbol and each exclusive part to its hypothesis label. For the reverse randomization, given an erasure draw from the normalized common transcript restriction; given a label draw from that hypothesis's normalized exclusive restriction. The common restriction is  $Z\mathbf{R}^\pi$ , not in general  $\mathbf{R}^\pi$  itself. Assign any fixed distribution to a zero-mass branch. These kernels recover both laws and do not depend on the true binary label. With equal priors the posterior remains equal on an erasure and the hypothesis is known otherwise, proving the testing assertion.

By Lemma A.9.2, the actual raw kernel and its own tangent kernel have total variation at most  $Cp_a^0\omega(d)$ , for either hypothesis. Apply Lemma 35.1 separately to these two approximations, under the same policy. The sum of the resulting bounds is  $C\varepsilon_\pi$ . The reverse triangle inequality for total variation then proves (A.9.15).  $\square$

**Corollary A.9.5** (An adaptive critical limit). *Consider a sequence of the preceding experiments and capped policies. Under the corresponding auxiliary laws  $\mathbf{R}^{\pi_n}$ , suppose*

$$\sum_{i=1}^{T_n} h(a_i) \longrightarrow b \in [0, \infty], \quad \max_{i \leq T_n} h(a_i) \longrightarrow 0 \quad \text{in probability}, \quad (\text{A.9.16})$$

where the maximum of an empty list is zero. If  $\varepsilon_{\pi_n} \rightarrow 0$ , the actual stopped transcript distances converge to  $1 - e^{-b}$  and their minimum equal-prior testing errors converge to  $e^{-b}/2$ . For  $b = \infty$ , convergence of the sum to infinity and  $\varepsilon_{\pi_n} \rightarrow 0$  suffice without the maximum condition.

*Proof.* For  $0 \leq x < 1$ , the series of  $-\log(1 - x)$  gives

$$0 \leq -\log(1 - x) - x \leq \frac{x^2}{2(1 - x)}.$$

For finite  $b$ , the sum of squared hazards is at most their maximum times their sum, hence tends to zero in probability. Applying the displayed inequality on the event that the maximum is at most  $1/2$  shows that the product in (A.9.13) converges in probability to  $e^{-b}$ . Products lie in  $[0, 1]$ , so their expectations converge as well: for each  $\epsilon > 0$ , the expectation of their absolute error is at most  $\epsilon$  plus the probability that the error exceeds  $\epsilon$ . If  $b = \infty$ , bound the product by the exponential of minus the hazard sum and use the same boundedness argument. Finally apply (A.9.15) and the binary testing identity.  $\square$

For predetermined windows and a deterministic stopping index, (A.9.13) reduces to (A.9.5). For adaptive windows it retains the observed common-history distribution instead of incorrectly averaging each hazard in isolation. The conditions concern the specified binary comparison; they do not assert optimal design or a minimax lower bound for geometric reconstruction. All transcript laws in this subsection retain the failure atom and the actual number of preparations.  $\square$

**A.9.2. Random accumulated hazard.** The deterministic limit in Corollary A.9.5 has the following distributional form. This clarification is recorded in the independent memorandum [35, Section 4.4].

**Corollary A.9.6** (Random-hazard mixture). *Use the same capped policies, common-history laws and physical approximation budgets as in Theorem A.9.4. Suppose, under  $\mathbf{R}^{\pi_n}$ ,*

$$H_n := \sum_{i \leq T_n} h(a_i) \implies B, \quad M_n := \max_{i \leq T_n} h(a_i) \longrightarrow 0 \quad \text{in probability,}$$

where  $B$  is a finite nonnegative random variable. If  $\varepsilon_{\pi_n} \rightarrow 0$ , then

$$\|\mathbf{P}^{\pi_n} - \mathbf{Q}^{\pi_n}\|_{\text{TV}} \longrightarrow 1 - \mathbb{E}e^{-B}, \quad \inf\{\text{equal-prior testing error}\} \longrightarrow \tfrac{1}{2}\mathbb{E}e^{-B}.$$

*Proof.* The convergence in distribution makes  $H_n$  tight. On  $M_n < 1$ , the logarithmic inequality in the proof of Corollary A.9.5 gives

$$0 \leq -\log \prod_{i \leq T_n} (1 - h(a_i)) - H_n \leq \frac{M_n H_n}{2(1 - M_n)}.$$

The right side tends to zero in probability. Consequently the product converges in distribution to  $e^{-B}$ . It takes values in  $[0, 1]$ , so its expectation tends to  $\mathbb{E}e^{-B}$  by bounded continuous convergence. Apply (A.9.13), (A.9.15) and the exact binary testing identity. No averaging under the physical history law is substituted for averaging under  $\mathbf{R}^{\pi_n}$ .  $\square$

In particular the limit is not determined by the mean of  $B$ . If  $\mathbb{E}B < \infty$ , strict convexity gives  $\mathbb{E}e^{-B} \geq e^{-\mathbb{E}B}$ , with equality only when  $B$  is almost surely constant. This observation does not assert optimality of an adaptive policy or make failed preparations free.

**A.9.3. Completion of the critical transition.** The accumulated remainder conditions in Theorem A.9.3 are useful for heterogeneous finite schedules. For a homogeneous sequence a condition on the error of one observation suffices throughout the entire transition, even when the full accumulated error diverges.

**Theorem A.9.7** (Zero, finite, and infinite effective sample limits). *Let  $j \rightarrow \infty$ ,  $d_j > 0$ ,  $q_j = e^{-j\gamma}$ , and let  $\omega$  have the meaning in Lemma A.9.2. The geometry may vary with  $j$  in a fixed compact family;  $\gamma, A$  and the laws then take their corresponding parameter values. Suppose*

$$\omega(d_j)/q_j \rightarrow 0. \quad (\text{A.9.17})$$

*If  $k_j q_j \rightarrow b \in [0, \infty]$ , then*

$$\|\Pi_{j,d_j}^{\otimes k_j} - \Pi_{\partial,j,d_j}^{\otimes k_j}\|_{\text{TV}} \rightarrow 1 - e^{-2b/\pi}.$$

*If  $n_j p_{j,d_j}^0 q_j \rightarrow b \in [0, \infty]$ , the same limit holds for  $\|P_{j,d_j}^{\otimes n_j} - P_{\partial,j,d_j}^{\otimes n_j}\|_{\text{TV}}$ . At  $b = \infty$  the right side is one. The minimum equal-prior errors converge to  $e^{-2b/\pi}/2$ . Both statements permit unequal facing curvatures and either parity. Under (A.9.17), vanishing total variation is equivalent to  $k_j q_j \rightarrow 0$  in the conditional experiment, and to  $n_j p_{j,d_j}^0 q_j \rightarrow 0$  in the raw experiment.*

*Proof.* If  $b < \infty$ , then  $k_j \omega(d_j) = (k_j q_j) \omega(d_j)/q_j \rightarrow 0$  in the conditional case, and  $n_j p_{j,d_j}^0 \omega(d_j) = (n_j p_{j,d_j}^0 q_j) \omega(d_j)/q_j \rightarrow 0$  in the raw case. Theorem A.9.3 applies, including  $b = 0$ .

Suppose  $b = \infty$  and fix  $B > 0$ . For all sufficiently large  $j$  the conditional sample contains its first  $\lfloor B/q_j \rfloor$  observations. Projection onto them cannot increase total variation. Their effective sample size tends to  $B$  and their accumulated tangent error tends to zero by (A.9.17). Hence the lower limit of the full-sample variation distance is at least  $1 - e^{-2B/\pi}$ . Let  $B \rightarrow \infty$  and use the upper bound one. For raw preparations project instead onto the first  $\lfloor B/(p_{j,d_j}^0 q_j) \rfloor$  observations. Here  $p_{j,d_j}^0 q_j \rightarrow 0$ , so rounding changes their effective sample size by a quantity tending to zero. The identical argument proves the raw assertion. All estimates used above are uniform on the compact geometric family, and its positive lower bound for  $\gamma$  ensures  $q_j \rightarrow 0$  even when the geometry varies. If an effective sample size fails to tend to zero, a subsequence is bounded below by a positive constant. Project on that subsequence onto a fixed positive critical subsample smaller than this constant. Its limiting distance is positive by the finite case, proving necessity in the final equivalences. Sufficiency is the case  $b = 0$ . The testing formula follows from  $\inf(\text{equal-prior error}) = (1 - \|P - Q\|_{\text{TV}})/2$ .  $\square$

For a general smooth contact, (A.9.17) is  $d_j = o(q_j^2)$ ; when both graphs are even it is  $d_j = o(q_j)$ . The finite critical subsamples therefore approach their quadratic experiments in total variation. This is a sharp physical testing profile in an asymptotically tangent regime, not a replacement for the nonlinear fixed-offset theorem or a proof that its rate  $\tau$  is sharp. The supercritical conclusion no longer requires the tangent error of all retained observations to vanish. The subsample argument applies to the known simple pair of laws, and does not supply a simulator for an unknown table.

#### APPENDIX A.10. CURVATURES INVISIBLE TO ALL THRESHOLD LOCATIONS

The following construction separates geometric uncertainty from its statistical measurement. Set  $\theta_r = r\pi/3$ ,  $r = 0, 1, 2$ , and let

$$h_{R,\alpha,\beta,\zeta}(\theta) = R + (1 - \cos 6\theta)b(\theta), \quad b(\theta) = \alpha + \beta \cos 2\theta + \zeta \sin 2\theta. \quad (\text{A.10.1})$$

Here  $R$  ranges over a compact interval in  $(0, 1/2)$ , and the three other parameters are sufficiently small. These are smooth centrally symmetric strictly convex scatterers, generally noncircular.

At each  $\theta_r$  and its opposite,  $h = R$ ,  $h' = 0$ , and

$$\kappa_r^{-1} = h(\theta_r) + h''(\theta_r) = R + 36b(\theta_r). \quad (\text{A.10.2})$$

The map from  $(\alpha, \beta, \zeta)$  to the three  $b(\theta_r)$  is invertible: its rows are  $(1, 1, 0)$ ,  $(1, -1/2, \sqrt{3}/2)$  and  $(1, -1/2, -\sqrt{3}/2)$ .

**Theorem A.10.1** (Equal gaps and identifiable curvature amplitudes). *In (A.10.1) all six shortest gaps are exactly  $g = 1 - 2R$ , independently of  $(\alpha, \beta, \zeta)$ . The entire set of maximal-count threshold locations is therefore the same. Nevertheless, write*

$$C_j = \lim_{d \downarrow 0} d^{-2} \Pr(N_{jg+d} = j + 1), \quad \gamma_r = \text{arcosh}(1 + g\kappa_r).$$

*Together with the observed spacing  $g$ , the unlabelled sequence  $(C_j)_{j \geq 1}$  determines the unordered triple  $(\kappa_0, \kappa_1, \kappa_2)$ , including multiplicities, and the free area  $A$ , without prior knowledge of  $A$ . It obeys*

$$C_j = \frac{1}{A} \sum_{r=0}^2 \text{csch}(j\gamma_r), \quad F_j := \sum_{\substack{k \geq 1 \\ k \text{ odd}}} \mu(k) C_{kj} = \frac{2}{A} \sum_{r=0}^2 e^{-j\gamma_r}, \quad (\text{A.10.3})$$

where  $\mu(k)$  is the arithmetic Möbius function. Both series involved in the inversion are absolutely convergent.

In the subfamily  $\beta = \zeta = 0$ , two coefficients already determine its shape parameter:

$$\alpha = \frac{1}{36} \left\{ \frac{g}{C_1/(2C_2) - 1} - R \right\}. \quad (\text{A.10.4})$$

Threshold locations alone cannot determine this parameter.

*Proof.* For small parameters  $h + h''$  remains positive, which proves smooth strict convexity and (A.10.2). For each nearest lattice vector  $e$  at angle  $\theta_r$ , the two supporting lines have separation  $1 - h(\theta_r) - h(\theta_r + \pi) = g$ . Every displacement from the first body to its translate has projection at least  $g$  in direction  $e$ , and hence length at least  $g$ . Since  $h' = 0$  at the two normals, the supporting points themselves lie on the axis and attain that distance. More distant translates have distance at least  $\sqrt{3} - 2 \sup h > g$  after a uniform decrease of the parameter neighborhood. Thus the global shortest length is exactly  $g$  and all six channels attain it. The three curvatures can vary independently by the displayed invertible matrix.

Central symmetry makes both curvatures of the  $r$ th pair equal. Theorem 13.1 gives the first formula in (A.10.3), with two orientations per pair. Put  $x_r = e^{-\gamma_r} \in (0, 1)$ . Expanding  $\text{csch}(j\gamma_r) = 2 \sum_{m \geq 1, m \text{ odd}} x_r^{mj}$  and interchanging the absolutely convergent sums gives

$$F_j = \frac{2}{A} \sum_r \sum_{\substack{n \geq 1 \\ n \text{ odd}}} x_r^{nj} \sum_{k|n} \mu(k) = \frac{2}{A} \sum_r x_r^j.$$

Absolute convergence follows, for example, by bounding the double sum by  $\sum_{n \geq 1} n x_{\max}^{nj} < \infty$ ; the number of divisors is at most  $n$ . The divisor identity leaves only  $n = 1$ .

After the odd-index transform, recovery is a finite-exponential-sum problem of Prony type [10]. We retain the complete Hankel–Vandermonde argument, rather than replace it by attribution. Let  $x_1, \dots, x_d$  be its distinct nodes,  $d \leq 3$ , with weights  $w_l = 2m_l/A > 0$ , where  $m_l$  are their multiplicities. Then  $F_j = \sum_l w_l x_l^j$ . The Hankel matrices

$$H^{(0)} = (F_{p+q+1})_{0 \leq p, q < d}, \quad H^{(1)} = (F_{p+q+2})_{0 \leq p, q < d}$$

satisfy  $H^{(0)} = V \text{diag}(w_l x_l) V^t$  and  $H^{(1)} = V \text{diag}(w_l x_l^2) V^t$ , where  $V_{pl} = x_l^p$ . The rank of the full Hankel sequence is  $d$ , and  $H^{(0)}$  is positive definite. Thus the generalized eigenvalues of this pair

are exactly  $x_l$ . A Vandermonde system using  $F_1, \dots, F_d$  recovers the  $w_l$ . Since  $\sum_l m_l = 3$ , we recover  $A = 6/(\sum_l w_l)$  and then  $m_l = Aw_l/2$ . Finally  $\kappa_l = (\cosh(-\log x_l) - 1)/g$  proves the curvature claim. No label ordering or global determination of the boundary is asserted.

If  $b \equiv \alpha$ , the three curvatures coincide and  $C_1/C_2 = 2 \cosh \gamma = 2(1 + g/(R + 36\alpha))$ . Solving gives (A.10.4). All threshold locations are independent of  $\alpha$  on a nontrivial open interval, while this ratio varies strictly with it. This proves the claimed nonredundancy.  $\square$

The infinite original sequence enters the Möbius sums; this is an exact identification theorem, not a finite-noisy-data reconstruction guarantee. The special two-amplitude formula needs no such infinite transform. In a circle, by contrast,  $R = (1 - g)/2$  and  $1/C_j$  already satisfies a second-order recurrence. Corollary A.34.2 remains correct; Theorem A.10.1 supplies the additional geometric uncertainty which is absent from that one-parameter model.

**Corollary A.10.2** (Instability selection at equal metric thresholds). *In (A.10.1), conditional on the maximal-count event in a common small collar, the mass of a channel with  $\gamma_r > \min_s \gamma_s$  is at most  $C \exp[-j(\gamma_r - \min_s \gamma_s)]$ , uniformly in the offset. Thus channels with the same exact threshold length can have exponentially different probabilities. The limiting leading-amplitude weights among minimizing channels are equal.*

*Proof.* The relative remainders in (13.1) are bounded above and below by positive constants. Divide a channel's coefficient by one with minimal  $\gamma$  and use the exponential expression for  $\text{csch}$ . At offset zero the equal minimal coefficients give equal leading-amplitude weights. This last assertion concerns amplitudes; it does not identify the fixed-positive-offset limiting corrections.  $\square$

**A.10.1. Area dependence and coalescence of unlabelled amplitudes.** The normalizer recovered in Theorem A.10.1 is not an additional independent shape parameter in its chosen family. The following identity also identifies a loss of sensitivity which is absent from the channel-resolved observation of Theorem 13.2.

**Proposition A.10.3** (The normalizer and the circular differential). *For (A.10.1),*

$$|D| = \pi \left[ R^2 + 2R\alpha - \frac{33}{2}\alpha^2 - \frac{45}{4}(\beta^2 + \zeta^2) \right]. \quad (\text{A.10.5})$$

*The free area  $A = \sqrt{3}/2 - |D|$  is a symmetric function of the gap and the unordered contact-curvature triple. In particular it is not a fourth independent inverse invariant in this three-parameter family.*

*Fix  $R$ , and restrict to  $\alpha = \zeta = 0$ ,  $\beta = s$ . The curvature radii are  $(R + 36s, R - 18s, R - 18s)$  and  $A(s) = A(0) + 45\pi s^2/4$ . For every fixed finite  $J$ ,*

$$\max_{1 \leq j \leq J} |C_j(s) - C_j(0)| = O_J(s^2), \quad \text{dist}_{\text{match}}(\kappa(s), \kappa(0)) \asymp |s|, \quad (\text{A.10.6})$$

*where  $\text{dist}_{\text{match}}$  is the minimum maximum coordinate distance over permutations of the two triples. In contrast, conditional endpoint variances on the three fixed channel patches recover their contact curvatures with a common Lipschitz constant, including at  $s = 0$ .*

*Proof.* For a positively curved support function, the boundary parametrization  $q = hn + h't$  has derivative  $q' = (h + h'')t$ . Green's area identity and integration by parts give

$$|D| = \frac{1}{2} \int_0^{2\pi} (h^2 - (h')^2) d\theta.$$

Write the nonconstant part of  $h$  as

$$\begin{aligned} h &= R + \alpha - \alpha \cos 6\theta + \beta \cos 2\theta - \frac{\beta}{2} \cos 4\theta - \frac{\beta}{2} \cos 8\theta \\ &\quad + \zeta \sin 2\theta + \frac{\zeta}{2} \sin 4\theta - \frac{\zeta}{2} \sin 8\theta. \end{aligned}$$



Orthogonality assigns to a sine or cosine coefficient of order  $k$  its square times  $\pi(1-k^2)/2$ . Thus the  $\alpha$  contribution is  $\pi[(R+\alpha)^2 - 35\alpha^2/2]$ , and the  $\beta, \zeta$  contribution is  $\pi[-3 - 15/4 - 63/4](\beta^2 + \zeta^2)/2$ . This is (A.10.5). Set  $b_r = (\kappa_r^{-1} - R)/36$ . The three values in (A.10.2) give

$$\alpha = \frac{1}{3} \sum_r b_r, \quad \beta^2 + \zeta^2 = \frac{2}{3} \sum_r (b_r - \alpha)^2.$$

Both expressions are symmetric;  $R = (1 - g)/2$  is known. This proves the normalizer assertion without changing its recovery formula.

Let  $\Phi_j(r) = \text{csch}(j \text{arcosh}(1 + g/r))$ . Along the specified one-parameter family,

$$C_j(s) = \frac{\Phi_j(R + 36s) + 2\Phi_j(R - 18s)}{A(s)}.$$

The numerator derivative at zero is  $(36 - 18 - 18)\Phi'_j(R) = 0$ , and  $A'(0) = 0$ . Taylor's theorem on a compact interval proves the first bound in (A.10.6), simultaneously for a fixed finite  $J$ . The largest curvature displacement is comparable to  $|s|$  because  $r \mapsto r^{-1}$  has derivative bounded away from zero there. This proves the matching-distance assertion. The last statement is Theorem 13.2 applied to each fixed channel; opposite curvatures agree by central symmetry, but that equality is not needed by its inverse.  $\square$

Thus exact recovery of a multiset from the infinite amplitude sequence does not imply Lipschitz recovery from any fixed finite amplitude vector at coalescence. The observation underlying the stable theorem is different: it retains physical endpoint positions in a selected channel, rather than only an unlabelled count. Its improvement is therefore not attributed to an algebraic regularization of the same datum. The cost of acquiring the conditional records is included in Proposition A.18.2. The restricted amplitude theorem, including its exact multiplicity recovery and the removal of an initially unknown normalizer, remains a separate conclusion.

#### APPENDIX A.11. A SHARP PAIRWISE INVERSE THROUGH A TRIPLE COLLISION

We now compare two arbitrary nearby curvature triples, rather than a triple with the circular reference alone. In the ambient model with an independent area normalizer, four amplitudes give an inverse without separating repeated nodes. The constrained physical family admits a three-amplitude refinement near  $R = 1/4$ , proved in Section A.12.

Fix  $g, R, A_0 > 0$  and consider the observation model

$$C_j(A, r) = \frac{1}{A} \sum_{i=1}^3 \Phi_j(r_i), \quad \Phi_j(r) = \text{csch}(j \text{arcosh}(1 + g/r)). \quad (\text{A.11.1})$$

Here  $A$  is near  $A_0$  and  $r_i$  are near  $R$ . The equal-gap geometric family of Theorem A.10.1 is a subfamily of this model, with  $r_i = \kappa_i^{-1}$  and its area given by (A.10.5). Allowing  $A$  to vary independently in the observation model proves a stronger nuisance-parameter statement; it does not make area an independent shape coordinate in that geometric family. When  $r_i^{-1}$  is interpreted as an actual contact curvature, each channel is assumed to have equal facing curvatures. Without that equality, the same scalar model describes the effective channel curvatures (A.12.3), not both endpoint curvatures.

**Theorem A.11.1** (Four-amplitude pairwise stability). *There is a neighborhood of  $(A_0, R, R, R)$  on which, for every two members of (A.11.1),*

$$|A - \tilde{A}| \leq C\delta_4, \quad \text{dist}_{\text{match}}(\kappa, \tilde{\kappa}) \leq C\delta_4^{1/3}, \quad \delta_4 = \max_{1 \leq j \leq 4} |C_j - \tilde{C}_j|. \quad (\text{A.11.2})$$

*The inverse determines multiplicities and requires neither channel labels nor a lower bound on the distance between distinct curvatures. In particular the first four amplitudes determine the area*

and unordered triple locally, even at complete coalescence. With deterministic error at most  $\epsilon$  in these four amplitudes, a reconstruction constrained to a sufficiently small compact neighborhood has area error  $C\epsilon$  and matching curvature error  $C\epsilon^{1/3}$ .

*Proof.* Put  $x_i = r_i - R$  and let  $e_1, e_2, e_3$  be their elementary symmetric polynomials. Define

$$p_e(z) = z^3 - e_1 z^2 + e_2 z - e_3.$$

For a small fixed circle  $|z| = b$  around zero,  $\Phi_j(R + z)$  is analytic inside it, with the positive-real branch at  $R$ . If  $e$  is sufficiently small, all roots of  $p_e$  lie strictly inside that circle. The symmetric sum in (A.11.1) has the analytic extension

$$M_j(e) = \frac{1}{2\pi i} \int_{|z|=b} \Phi_j(R + z) \frac{p'_e(z)}{p_e(z)} dz. \quad (\text{A.11.3})$$

This formula counts roots with multiplicity and defines an analytic function on a full coefficient neighborhood, not merely on its real-rooted part. It is real for real  $e$ . Expanding the Newton sums at  $e = 0$  gives

$$M_j(0) = 3\Phi_j(R), \quad \partial_{e_1} M_j = \Phi'_j(R), \quad \partial_{e_2} M_j = -\Phi''_j(R), \quad \partial_{e_3} M_j = \frac{1}{2}\Phi'''_j(R) \quad (e = 0). \quad (\text{A.11.4})$$

For clarity, the first three Newton sums are  $e_1$ ,  $e_1^2 - 2e_2$ , and  $e_1^3 - 3e_1e_2 + 3e_3$ ; all sums of degree at least four have zero first differential at  $e = 0$ . This also follows by differentiating the contour integral.

Consider the map  $\Psi : (A, e_1, e_2, e_3) \mapsto (A^{-1}M_j(e))_{j=1}^4$ . Its Jacobian determinant at  $(A_0, 0, 0, 0)$  is

$$\det D\Psi = \frac{3}{2A_0^5} \det[\Phi_j^{(k)}(R)]_{1 \leq j \leq 4, 0 \leq k \leq 3} > 0. \quad (\text{A.11.5})$$

We verify the strict sign explicitly. Put  $c = 1 + g/R > 1$  and write  $\Phi_j(r) = w(c(r))f_j(c(r))$ , where  $c(r) = 1 + g/r$ ,  $w(c) = (c^2 - 1)^{-1/2}$ , and

$$f_1 = 1, \quad f_2 = \frac{1}{2c}, \quad f_3 = \frac{1}{4c^2 - 1}, \quad f_4 = \frac{1}{4c(2c^2 - 1)}.$$

Column elimination in the product and chain rules gives

$$\det[\Phi_j^{(k)}(R)] = \left(\frac{g}{R^2}\right)^6 (c^2 - 1)^{-2} \det[f_j^{(k)}(c)].$$

The last determinant reduces, using  $f_1 = 1$ , to the three by three matrix of the first, second, and third derivatives of  $f_2, f_3, f_4$ . Differentiating these displayed rational functions and putting the entries over their common denominator yields

$$\det[f_j^{(k)}(c)] = \frac{12(64c^8 + 32c^6 + 116c^4 + 4c^2 + 1)}{c^4(4c^2 - 1)^4(2c^2 - 1)^4}. \quad (\text{A.11.6})$$

Every factor is strictly positive for  $c > 1$ . The column factors in (A.11.4) are  $-3/A_0^2$ ,  $1/A_0$ ,  $-1/A_0$  and  $1/(2A_0)$ , which proves the factor  $3/(2A_0^5)$  in (A.11.5). Thus the analytic inverse-function theorem applies in the full four-dimensional coefficient space. Choose a convex ball in its data image and then a smaller parameter neighborhood whose image lies in that ball. A bounded derivative of the inverse on the ball gives

$$|A - \tilde{A}| + \sum_{k=1}^3 |e_k - \tilde{e}_k| \leq C\delta_4 \quad (\text{A.11.7})$$

for any two members of this smaller neighborhood, including its real-rooted boundary strata. This is the step for which using the symmetric coordinates, rather than differentiating individually labelled roots, is essential.

We give the needed root estimate with multiplicities. For two monic cubics  $p, q$  whose roots lie in a fixed bounded set and whose coefficient distance is  $\epsilon$ , one has  $|p(z) - q(z)| \leq C_0\epsilon$  on a common bounded disk containing the small neighborhoods below. Set  $\rho = (2C_0\epsilon)^{1/3}$ . On the boundary of the union of the open radius- $\rho$  disks about the three roots of  $p$ , every root distance is at least  $\rho$ . Hence  $|p(z)| \geq \rho^3 > |p(z) - q(z)|$  there. Rouché's theorem on each connected component shows that the two polynomials have the same number of roots in that component, counted with multiplicities. If two boundary arcs are tangent, perturb the radius slightly and pass to the limit; the strict inequality persists. Every component consists of at most three intersecting disks and has diameter at most  $6\rho$ . Match the roots inside each component. This proves a matching distance at most  $C\epsilon^{1/3}$ . At  $\epsilon = 0$  the polynomials coincide. The assertion for bounded, not necessarily small,  $\epsilon$  follows by enlarging the constant.

Apply this estimate to (A.11.7). The radii stay in a compact interval bounded away from zero, so  $r \mapsto r^{-1}$  is Lipschitz there. This proves (A.11.2). For noisy data, minimize the maximum discrepancy over a small compact admissible parameter set containing the true member. A minimum exists by continuity. Its exact amplitude vector differs from the true one by at most  $2\epsilon$ , since the true member is itself feasible. Apply the pairwise estimate. Uniqueness of this minimizer is not required for the error bound, nor is an efficient numerical optimization algorithm being assumed in its proof.  $\square$

**A.11.1. Sharpness in the physical geometric family.** For  $a \geq 0$  use the same absolute-amplitude topology as before,  $\|D\|_a = \sup_{j \geq 1} e^{aj}|D_j|$ . Choose a small fixed- $R$  neighborhood in (A.10.1) on which all exponents exceed some  $\gamma_- > a$ .

**Theorem A.11.2** (Optimal pairwise exponent). *On a sufficiently small neighborhood in that physical family,*

$$\text{dist}_{\text{match}}(\kappa, \tilde{\kappa}) \leq C_a \|C - \tilde{C}\|_a^{1/3}.$$

*The exponent 1/3 cannot be increased, even if the entire sequence is observed. On the two-sided path  $\alpha = \zeta = 0$ ,  $\beta = s$ ,*

$$\|C(s) - C(-s)\|_a \asymp |s|^3, \quad \text{dist}_{\text{match}}(\kappa(s), \kappa(-s)) \asymp |s|. \quad (\text{A.11.8})$$

*The circular-reference exponent 1/2 of Theorem A.21.2 remains valid and sharp for its different, reference-point quantifiers.*

*Proof.* The first assertion follows from Theorem A.11.1, since the first four errors are bounded by the sequence norm. For sharpness use exactly the realized path of Proposition A.10.3:

$$r(s) = (R + 36s, R - 18s, R - 18s), \quad A(s) = A_0 + 45\pi s^2/4.$$

Here  $A_0$  is the actual circular free area. Differentiation of  $2e^{-j\gamma}/(1 - e^{-2j\gamma})$ , together with the bounded first three derivatives of  $\gamma(r)$  and  $A(s)^{-1}$ , gives

$$\sup_{|s| \leq s_*} |C_j'''(s)| \leq B(1 + j)^3 e^{-j\gamma_-}.$$

The first derivative at zero vanishes. Subtract the Taylor expansions at  $s$  and  $-s$ : their constant and quadratic terms cancel, as does the zero linear term. The strict margin  $a < \gamma_-$  absorbs the polynomial in  $j$ , proving  $\|C(s) - C(-s)\|_a \leq B_a |s|^3$ .

The first amplitude supplies a nonzero lower bound. Direct differentiation gives

$$\Phi_1(r) = \frac{r}{\sqrt{g(g+2r)}}, \quad \Phi_1'''(r) = \frac{3(3g+r)}{\sqrt{g}(g+2r)^{7/2}} > 0.$$

Since  $36^3 + 2(-18)^3 = 34992$  and  $A(s)$  is even,

$$C_1(s) - C_1(-s) = \frac{11664\Phi_1'''(R)}{A_0}s^3 + O(s^5). \quad (\text{A.11.9})$$

This proves the other sequence bound. In one real dimension sorted matching minimizes the maximum coordinate distance: an inverted pair can be uncrossed without increasing either extreme distance. Sorting the two curvature triples, their middle pair is the largest difference for all sufficiently small  $s > 0$ , and gives

$$\text{dist}_{\text{match}}(\kappa(s), \kappa(-s)) = \frac{36s}{R^2 - 324s^2}.$$

The negative case follows by symmetry. If a pairwise exponent  $\theta > 1/3$  held uniformly, these estimates would imply  $b|s| \leq C|s|^{3\theta}$  for all small nonzero  $s$ , a contradiction. This proves sharpness without replacing the infinite sequence by a finite numerical prefix.  $\square$

The reference and pairwise results describe different singularities of one observation map. Distance from the circular point first detects quadratic dispersion. Comparing two noncircular triples on opposite sides of it can cancel that dispersion as well, leaving a cubic term. The four-amplitude inverse proves that the latter cancellation is the worst local pairwise scale in the stated three-channel model. This positive conclusion does not use relative accuracy at arbitrarily small amplitudes or a separation-dependent Prony bound.

#### APPENDIX A.12. THREE AMPLITUDES IN THE CONSTRAINED PHYSICAL FAMILY

The four-amplitude theorem permits an independent area normalizer. The physical family (A.10.1) has a constraint on that normalizer, and this constraint can reduce the required number of amplitudes. We prove the reduction near  $R = 1/4$ . The four-amplitude result remains available on its full positive parameter range and with an independent area.

Put  $x_i = r_i - R$  for the three contact radii and let  $e_k$  be their elementary symmetric functions. The three values in (A.10.2) give

$$\alpha = \frac{e_1}{108}, \quad \beta^2 + \zeta^2 = \frac{e_1^2}{2916} - \frac{e_2}{972}. \quad (\text{A.12.1})$$

Substituting in (A.10.5) yields the exact identity

$$A(R, e) = A_0(R) - \frac{\pi R}{54}e_1 + \frac{41\pi}{7776}e_1^2 - \frac{5\pi}{432}e_2, \quad A_0(R) = \frac{\sqrt{3}}{2} - \pi R^2. \quad (\text{A.12.2})$$

Indeed, the sum of the three  $x_i$  is  $108\alpha$ , while their squared sum is  $3888\alpha^2 + 1944(\beta^2 + \zeta^2)$ . The identity  $2e_2 = e_1^2 - \sum_i x_i^2$  proves (A.12.1), and collecting its terms in the area formula proves (A.12.2).

The contact radii in this physical family are actual radii at both ends of each channel: central symmetry makes the two facing curvatures equal. This is additional structure not required by the general forward law. Without equality of the facing curvatures, the scalar count coefficient retains only

$$\kappa_e^{\text{eff}} = \frac{\sqrt{(1 + g\kappa_{e,0})(1 + g\kappa_{e,1})} - 1}{g}. \quad (\text{A.12.3})$$

The ambient scalar inverse then recovers these effective parameters, not the two actual endpoint curvatures. The conditional variances in Theorem 13.2 recover the latter without this symmetry.

**Theorem A.12.1** (Physical three-amplitude reconstruction). *There is a neighborhood of  $R = 1/4$ ,  $\alpha = \beta = \zeta = 0$  in (A.10.1) such that the spacing  $g = 1 - 2R$  and the first three unlabelled leading amplitudes determine the unordered contact-curvature triple, including multiplicities, and the area. For two members of a sufficiently small such neighborhood, put*

$$\Delta = |g - \tilde{g}| + \max_{1 \leq j \leq 3} |C_j - \tilde{C}_j|.$$

Then

$$|R - \tilde{R}| + |A - \tilde{A}| + \sum_{k=1}^3 |e_k - \tilde{e}_k| \leq C\Delta, \quad \text{dist}_{\text{match}}(\kappa, \tilde{\kappa}) \leq C\Delta^{1/3}. \quad (\text{A.12.4})$$

The pairwise exponent  $1/3$  is sharp even at fixed  $R = 1/4$  and known  $g$ . Thus no fourth amplitude is required in this constrained family near that reference point. The assertion does not impose a minimum separation between distinct curvatures.

*Proof.* For fixed  $R$ , set  $g = 1 - 2R$ , and use the analytic contour sums  $M_j(R, e)$  of (A.11.3). They depend analytically also on  $R$  in a small real neighborhood of  $1/4$ , with a common complex contour and branches. On a full coefficient neighborhood, define

$$\Psi_3(R, e) = \left( \frac{M_j(R, e)}{A(R, e)} \right)_{j=1}^3.$$

The area stays positive there. At  $e = 0$ , the derivative with respect to the three coefficient coordinates has rows

$$\frac{1}{A_0} \left( \Phi'_j(R) + \frac{\pi R}{18A_0} \Phi_j(R), \quad -\Phi''_j(R) + \frac{5\pi}{144A_0} \Phi_j(R), \quad \frac{1}{2} \Phi'''_j(R) \right). \quad (\text{A.12.5})$$

Derivatives on  $\Phi_j$  here keep  $g$  fixed. The area terms follow from  $M_j(R, 0) = 3\Phi_j(R)$  and  $A_{e_1} = -\pi R/54$ ,  $A_{e_2} = -5\pi/432$ ,  $A_{e_3} = 0$ .

We verify nondegeneracy without assuming it numerically. At  $R = 1/4$  and  $g = 1/2$ , the entries  $\Phi_j^{(k)}(R)/\sqrt{2}$  are

| $j$ | $k = 0$ | $k = 1$  | $k = 2$      | $k = 3$          |
|-----|---------|----------|--------------|------------------|
| 1   | 1/4     | 3/4      | -5/4         | 21/4             |
| 2   | 1/24    | 17/72    | 35/216       | -491/216         |
| 3   | 1/140   | 297/4900 | 36243/171500 | -4458537/6002500 |

These values follow by differentiating  $\Phi_1(r) = r/\sqrt{g(g+2r)}$ ,  $\Phi_2 = \Phi_1/[2(1+g/r)]$  and  $\Phi_3 = \Phi_1/[4(1+g/r)^2 - 1]$ . Substitution into (A.12.5) gives

$$\det D_e \Psi_3(1/4, 0) = -\frac{2\sqrt{2}(15804720A_0 + 64253\pi)}{72930375A_0^4} \neq 0, \quad A_0 = \frac{\sqrt{3}}{2} - \frac{\pi}{16} > 0. \quad (\text{A.12.6})$$

To include changes in the spacing, apply the analytic inverse-function theorem to

$$(R, e) \mapsto (1 - 2R, \Psi_3(R, e)).$$

Its first row is  $(-2, 0, 0, 0)$ , and the determinant in the other three columns is (A.12.6). Hence it is locally invertible in a full four-dimensional coefficient neighborhood, before restriction to real roots. Choose a convex data ball on which the inverse has bounded derivative and then a smaller parameter neighborhood whose image lies in that ball. This gives the coefficient and  $R$  estimates in (A.12.4). Analyticity of (A.12.2) gives the area estimate.

The root-matching proof in Theorem A.11.1 gives  $\text{dist}_{\text{match}}(x, \tilde{x}) \leq C\Delta^{1/3}$ . Adding the change of  $R$  and applying the Lipschitz map  $r \mapsto r^{-1}$  on a compact positive interval proves the curvature estimate, after decreasing the neighborhood so  $\Delta \leq 1$ . This retains all repeated-root strata. On

the physical path  $\alpha = \zeta = 0$ ,  $\beta = s$ , compared with  $-s$ , the gap stays fixed and Theorem A.11.2 gives amplitude discrepancy  $O(|s|^3)$  and matching curvature distance comparable to  $|s|$ . The nonzero first-amplitude cubic coefficient gives the reverse amplitude bound for the first three amplitudes as well. An exponent larger than  $1/3$  would contradict this path.  $\square$

With deterministic error at most  $\delta$  in the observed spacing and these three amplitudes, choose an admissible member minimizing their maximum discrepancy on a sufficiently small compact physical parameter set containing the truth. Its exact data differ from the true data by at most  $2\delta$ . The theorem therefore gives area error  $C\delta$  and matching curvature error  $C\delta^{1/3}$ . This existence argument neither requires a unique minimizer nor asserts a particular efficient optimization algorithm. More importantly, it concerns errors in amplitudes, not yet errors in individual collision counts. We next account for the experiment that estimates those amplitudes.

### APPENDIX A.13. EXACT INVERSION FROM FOUR FIXED PHYSICAL WINDOWS

We now keep the offsets fixed while the requested accuracy improves. This experiment uses the exact analytic support family (A.10.1), not an unspecified smooth remainder in its onset expansion. The true gap is unknown. The reference circle has  $R_c = 1/4$  and  $g_c = 1/2$ . Put  $x_i = r_i - R$ , and let  $e = (e_1, e_2, e_3)$  be the elementary symmetric functions of the three radius increments. Thus  $R = (1 - g)/2$ , the unordered curvatures are  $(R + x_i)^{-1}$ , and  $A = A(R, e)$  is the exact function in (A.12.2). Throughout this section the physical parameter set is a sufficiently small fixed compact neighborhood of the reference member. In particular it contains coalescing and repeated-root configurations. It is not a set of labelled root parameters with a separation condition.

**A.13.1. Analytic descent through coalescence.** For  $j = 1, 2, 3$ , define at positive offset

$$F_j(g, x, d) = d^{-2} \Pr_{g,x} \{N_{jg+d} \geq j + 1\}.$$

These are normalized exact finite-offset probabilities. The analytic extension needed for inversion at a triple root is stronger than continuity in the labelled radii.

**Lemma A.13.1** (Symmetric analytic finite-offset coordinates). *The functions  $F_j$  have jointly analytic extensions to a neighborhood of  $(g_c, 0, 0)$  in the variables  $(g, e, d)$ , before restriction to the real-rooted physical locus. On a smaller such neighborhood,*

$$F_j(g, e, d) = C_j(g, e) + dB_j(g, e, d), \quad (\text{A.13.1})$$

where  $B_j$  is analytic and  $C_j$  is the physical leading-amplitude map in Theorem A.12.1. Every fixed derivative is bounded on a smaller closed neighborhood.

*Proof.* The linear map from  $(\alpha, \beta, \zeta)$  to  $(x_1, x_2, x_3)$  is invertible by (A.10.2). The support family is analytic in these variables and  $g$ . For each of the finitely many channels and flight numbers under consideration, Lemma 15.2 gives the analytic stationary action and its twist on a common complex neighborhood. The positive real endpoint Hessian has an analytic square-root branch on a smaller complex neighborhood. The constructive Morse map in Lemma 16.1 and its inverse are consequently analytic there, by the analytic inverse-function theorem.

After physical residual-time integration and the Morse change of variables, the probability divided by  $d^2$  is an integral over a fixed disk of an analytic amplitude at  $\sqrt{2d}z$ , with weight  $1 - |z|^2$ . Expand the amplitude in its convergent endpoint power series on a strictly smaller polydisk. The series is normally convergent, uniformly in the remaining complex parameters. Integration annihilates all odd total degrees and leaves a normally convergent series in integer powers of  $d$ . This constructs an analytic extension across  $d = 0$ ; it does not assign a nonzero

physical probability below onset. The six channels all have exactly the same gap in the present family, so adding them introduces no moving activation interface.

A rotation of the triangular lattice by  $\pi/3$  cycles the three radius increments. Reflection in the horizontal axis fixes the first and interchanges the other two. These transformations generate their full permutation group. They preserve phase volume and the unlabelled count event, hence  $F_j(g, x, d)$  is symmetric in  $x$  on the real physical open set. The analytic identity principle gives the same symmetry on a permutation-invariant complex polydisk.

We spell out descent at the discriminant. Choose a smaller coefficient polydisk on which all roots of

$$z^3 - e_1 z^2 + e_2 z - e_3$$

stay strictly inside the original root polydisk; this is possible by continuity of the root multiset at  $e = 0$ . Away from the discriminant, local holomorphic root branches define  $F_j$  as a holomorphic function of  $(g, e, d)$ . Symmetry makes this definition independent of the branches and their permutations. It is locally bounded uniformly up to the discriminant, since all roots stay in the smaller root polydisk and the original function is bounded on its closure. The removable-singularity theorem for locally bounded holomorphic functions across an analytic hypersurface therefore extends it jointly across the discriminant. One can apply the theorem in the product domain with  $g, d$  as additional variables, so joint analyticity is retained.

At  $d = 0$ , the physical onset formula identifies the extension with  $C_j(g, e)$  on distinct real-rooted configurations. Such configurations contain real open subsets of the coefficient domain, so uniqueness of analytic extension identifies the two analytic functions everywhere near  $(g_c, 0)$ . Dividing their difference by  $d$ , or using its convergent Taylor series, gives (A.13.1). Cauchy estimates on smaller polydisks give the final boundedness assertion.  $\square$

**A.13.2. The unknown gap as a fourth coordinate.** Fix a positive number  $a$ , to be chosen sufficiently small once and then kept independent of accuracy and confidence. Use the four windows

$$(j, t) = (1, g_c + a), \quad (1, g_c + 2a), \quad (2, 2g_c + a), \quad (3, 3g_c + a). \quad (\text{A.13.2})$$

Their normalizations use the known programmed excesses, not the unknown actual ones. If  $v = g - g_c$ , put

$$Q_a(g, e) = \begin{pmatrix} (a - v)^2 F_1(g, e, a - v)/a^2 \\ (2a - v)^2 F_1(g, e, 2a - v)/(4a^2) \\ (a - 2v)^2 F_2(g, e, a - 2v)/a^2 \\ (a - 3v)^2 F_3(g, e, a - 3v)/a^2 \end{pmatrix}. \quad (\text{A.13.3})$$

On  $|v| < a/6$  these are the normalized physical means of the four binary reports; all actual offsets are positive. Decrease  $a$  and the neighborhood so they lie in the common physical collar.

**Theorem A.13.2** (Four-window physical inverse). *For some fixed  $a > 0$ , the exact map  $Q_a$  is analytically invertible on a full neighborhood of  $(g_c, 0)$  in coefficient coordinates. There is a compact physical neighborhood  $K$  of the circular member such that, for  $\theta, \tilde{\theta} \in K$ ,*

$$|g - \tilde{g}| + \|e - \tilde{e}\|_\infty + |A - \tilde{A}| \leq C \|Q_a(\theta) - Q_a(\tilde{\theta})\|_\infty, \quad (\text{A.13.4})$$

$$\text{dist}_{\text{match}}(\kappa, \tilde{\kappa}) \leq C \|Q_a(\theta) - Q_a(\tilde{\theta})\|_\infty^{1/3}. \quad (\text{A.13.5})$$

*The neighborhood and constants do not depend on a root separation. Only the known family formula, reference window constants, and the fixed neighborhood are supplied. The true gap, area, channel labels, and impact positions are not supplied.*

*Proof.* Let  $D = D_e(C_1, C_2, C_3)(g_c, 0)$ . By (A.12.6),

$$\det D = -\frac{2\sqrt{2}(15804720A_c + 64253\pi)}{72930375A_c^4}, \quad A_c = \frac{\sqrt{3}}{2} - \frac{\pi}{16} > 0. \quad (\text{A.13.6})$$

In particular  $\det D \neq 0$ . The derivation of this exact determinant, including its twelve amplitude derivatives and the physical area constraint, is given in Theorem A.12.1.

For a window of order  $j$  with programmed excess  $la$  at the reference gap, differentiation at  $(g_c, 0)$  gives

$$\begin{aligned} \partial_g \left[ \frac{(la - jv)^2}{(la)^2} F_j(g, e, la - jv) \right] &= -\frac{2jC_j}{la} + O(1), \\ D_e \left[ \frac{(la - jv)^2}{(la)^2} F_j(g, e, la - jv) \right] &= D_e C_j + O(a). \end{aligned}$$

Here Lemma A.13.1 bounds derivatives in the full coefficient neighborhood; estimates only in the labelled roots would not justify the second formula at coalescence.

In  $DQ_a$ , subtract row one from row two, place this difference first, and multiply the gap column by  $a$ . The resulting matrix converges as  $a \downarrow 0$  to

$$\begin{pmatrix} C_1 & 0 & 0 & 0 \\ -2C_1 & D_{11} & D_{12} & D_{13} \\ -4C_2 & D_{21} & D_{22} & D_{23} \\ -6C_3 & D_{31} & D_{32} & D_{33} \end{pmatrix}_{(g_c, 0)}. \quad (\text{A.13.7})$$

Its determinant is  $C_1 \det D \neq 0$ . These row and column operations are invertible for every  $a > 0$ . Fix  $a$  small enough that  $DQ_a$  is nonsingular. The analytic inverse-function theorem now applies to the exact finite-offset map, not to its leading approximation.

Choose a convex data ball with closure inside the image of an inverse chart, and a smaller compact physical neighborhood  $K$  whose image lies in that ball. The inverse has bounded derivative on the ball, giving the gap and coefficient bounds in (A.13.4). The analytic area identity gives its area bound. The root-matching argument of Theorem A.11.1 gives  $\text{dist}_{\text{match}}(x, \tilde{x}) \leq C\|e - \tilde{e}\|_\infty^{1/3}$ . Adding the change in  $R = (1 - g)/2$  and using the bounded derivative of  $r \mapsto r^{-1}$  on the common positive radius interval proves (A.13.5). Both the analytic inverse and the root estimate were established before any exclusion of a repeated root, so no such exclusion is needed.  $\square$

**Corollary A.13.3** (Deterministic-budget fixed-window acquisition). *There are constants  $c, C > 0$  such that for  $0 < \varepsilon, \delta < c$  and  $0 < \eta < 1/4$ , four-window binary observations give a Borel estimator satisfying, uniformly on  $K$ ,*

$$\Pr\{\text{dist}_{\text{match}}(\hat{\kappa}, \kappa) \leq \varepsilon, \quad |\hat{g} - g| \leq \delta\} \geq 1 - \eta, \quad (\text{A.13.8})$$

using a deterministic number of independent preparations bounded by

$$C(\varepsilon^{-6} + \delta^{-2}) \log(C/\eta). \quad (\text{A.13.9})$$

The same estimator has coefficient and area errors at most  $C \min\{\varepsilon^3, \delta\}$  on this event. In particular, curvature-only accuracy  $\varepsilon$  costs at most  $C\varepsilon^{-6} \log(C/\eta)$  and also gives gap and area errors  $O(\varepsilon^3)$ .

*Proof.* Take  $n$  fresh preparations in each window. The centered Bernoulli moment-generating function is at most  $\exp(t^2/8)$ : the second derivative of its log is the variance of a tilted Bernoulli law and is at most  $1/4$ , while its value and first derivative at zero vanish. Exponential Markov



and optimization therefore give  $\Pr\{|\hat{p} - p| > u\} \leq 2e^{-2nu^2}$ . Dividing the four fractions by their fixed design normalizers and taking a union bound gives

$$\Pr\{\|\hat{Q} - Q_a(\theta)\|_\infty > C_a \sqrt{\log(8/\eta)/n}\} \leq \eta.$$

The dependence on  $a$  is a constant because  $a$  is fixed.

Minimize  $\|Q_a(\vartheta) - \hat{Q}\|_\infty$  on  $K$ , using the Borel selection of Lemma A.22.1. This is defined on all outcomes, including data outside the inverse chart. Since the truth lies in  $K$ , its fitted exact data differ from the true data by at most  $2\|\hat{Q} - Q_a(\theta)\|_\infty$ . Apply Theorem A.13.2. Taking

$$n = \left\lceil C \min\{\varepsilon^3, \delta\}^{-2} \log(C/\eta) \right\rceil$$

with a sufficiently large constant proves all stated error bounds. Finally  $\min\{\varepsilon^3, \delta\}^{-2} = \max\{\varepsilon^{-6}, \delta^{-2}\} \leq \varepsilon^{-6} + \delta^{-2}$ , and there are exactly  $4n$  preparations. This specifies a statistical estimator, not an efficient algorithm for exact evaluation of the nonlinear probability map.  $\square$

The fixed-offset analytic descent and the four-window design were suggested in the referee comparison supplied with the two independent reports of September 9, 2026. The preceding proofs incorporate that comparison. The joint lower bound below identifies the additional cost of demanding a finer gap estimate.

#### APPENDIX A.14. JOINT CURVATURE AND ONSET ACCURACY

We specify the class of experiments before giving the lower bound. Fix an integer  $J \geq 3$ , a sufficiently small common collar  $d_* > 0$  for flight numbers at most  $J$ , and the compact physical neighborhood  $K$  of Theorem A.13.2, decreased if necessary. A query chooses a physical time  $t \geq 0$  and  $1 \leq j \leq J$  and returns only  $\mathbf{1}_{\{N_t \geq j+1\}}$  from a fresh independent full-phase preparation. The policy may use previous returned bits and independent randomization, and may stop at a stopping time  $T$ . Its rules are the same for every unknown parameter. On every possible history, every selected query satisfies

$$t - jg(\theta) \leq d_* \quad \text{for every } \theta \in K. \quad (\text{A.14.1})$$

Below-onset queries are allowed and are deterministic failures. We keep  $J, d_*, K$  fixed as the requested errors decrease. The four windows of (A.13.2) belong to this class after the preceding choices are made small enough. A fixed known neighborhood is thus permitted; the true gap is not revealed.

Let  $\mathcal{N}_K(\varepsilon, \delta, \eta)$  be the infimum of  $\sup_{\theta \in K} \mathbb{E}_\theta T$  over these policies and Borel estimators satisfying (A.13.8). All information retained from each trajectory is its one binary report. In particular an unrecorded lower collision count is not an extra observation.

**A.14.1. Testing the unknown onset.** The curvature splitting alternatives of Lemma A.24.1 have the same gap. To measure the price of timing, we need a second pair of physical alternatives.

**Lemma A.14.1** (One-sided entropy for a displacement of the gap). *In the circular subfamily  $e = 0$ , consider gaps  $g_c$  and  $g_c + \Delta$  with  $0 < \Delta$  small. For any common query satisfying (A.14.1), denote its probabilities by  $p_0, p_\Delta$ . Then*

$$D_{\text{KL}}(\text{Ber}(p_\Delta) \parallel \text{Ber}(p_0)) \leq C\Delta^2. \quad (\text{A.14.2})$$

For a stopped policy with  $\mathbb{E}_\Delta T < \infty$ , its complete transcript laws satisfy

$$D_{\text{KL}}(\mathbf{P}_\Delta \parallel \mathbf{P}_0) \leq C\Delta^2 \mathbb{E}_\Delta T. \quad (\text{A.14.3})$$

The constant is independent of how near to onset the queries are.

*Proof.* Put  $d = t - jg_c$  and  $z = j\Delta$ . For the circular subfamily the uniform physical expansion, at these finitely many flight numbers, can be written

$$p_0 = d_+^2 H_j(g_c, d_+), \quad p_\Delta = (d - z)_+^2 H_j(g_c + \Delta, (d - z)_+).$$

Here  $H_j$  is positive, bounded above and below, and has bounded first derivatives on the common closed collar. This follows from Theorem 13.1; the area and all curvatures vary smoothly with  $R = (1 - g)/2$ . Decrease  $d_*$  so all nonzero probabilities are less than  $1/2$ .

If  $d \leq 0$ , both reports are deterministic failures. If  $0 < d \leq z$ , then  $p_\Delta = 0$  and  $p_0 \leq Cd^2 \leq CJd\Delta$ . If  $d > z$ , then

$$|(d - z)^2 - d^2| \leq 2dz \leq 2Jd\Delta,$$

and the mean-value bound for  $H_j$  gives an additional difference at most  $Cd^2(\Delta + z) \leq Cd\Delta$ , since  $d \leq d_*$ . Thus in both positive- $d$  cases

$$|p_\Delta - p_0| \leq Cd\Delta, \quad p_0 \geq cd^2.$$

The Bernoulli entropy inequality used in (A.24.4), including its continuous extension when the first probability is zero, yields

$$D_{\text{KL}}(\text{Ber}(p_\Delta) \| \text{Ber}(p_0)) \leq \frac{(p_\Delta - p_0)^2}{p_0(1 - p_0)} \leq C\Delta^2.$$

The direction of entropy is important: the higher-gap law has no successes where the lower-gap law is identically zero, so no unsupported absolute continuity assertion across the onset is used.

Conditioning on a common finite history and the parameter-independent random seed selects the same query under the two alternatives. The chain rule for the first  $M$  positions of the transcript, padded with stop symbols, gives  $C\Delta^2 \mathbb{E}_\Delta(T \wedge M)$ . Relative entropy on the increasing transcript sigma-fields converges to the entropy of the full transcript, as in the truncation proof of (A.24.4). Monotone convergence proves (A.14.3). A Borel statistic of the transcript adds no relative entropy.  $\square$

**Theorem A.14.2** (Joint local binary minimax order). *For fixed  $J, d_*, K$  as above there are  $c, C, \varepsilon_0, \delta_0 > 0$  such that, for  $0 < \varepsilon < \varepsilon_0$ ,  $0 < \delta < \delta_0$ , and  $0 < \eta < 1/4$ ,*

$$c(\varepsilon^{-6} + \delta^{-2}) \log(1/\eta) \leq \mathcal{N}_K(\varepsilon, \delta, \eta) \leq C(\varepsilon^{-6} + \delta^{-2}) \log(1/\eta). \quad (\text{A.14.4})$$

*The upper bound uses four fixed physical windows and a deterministic preparation count. The lower bound permits adaptive queries and random stopping. It is a statement for this exact physical family and this bounded-flight, local binary experiment.*

*Proof.* The upper bound is Corollary A.13.3; for  $\eta < 1/4$  its logarithm is comparable to  $\log(1/\eta)$ .

We first record the confidence version of the testing step. If a test chooses its first hypothesis on an event  $B$  with  $P(B) \geq 1 - \eta$  and  $Q(B) \leq \eta$ , the log-sum inequality and monotonicity of binary relative entropy on this rectangle give

$$D_{\text{KL}}(P \| Q) \geq (1 - 2\eta) \log \frac{1 - \eta}{\eta} \geq \frac{1}{4} \log(1/\eta), \quad 0 < \eta < 1/4. \quad (\text{A.14.5})$$

For the last inequality use  $1 - 2\eta \geq 1/2$  and  $1 - \eta \geq \sqrt{\eta}$ . The rectangle monotonicity follows by differentiating the binary relative entropy in its two arguments when the first exceeds the second.

For the curvature lower bound, take the physical tables  $R = 1/4$ ,  $\alpha = \zeta = 0$ ,  $\beta = \pm s$  inside  $K$ . Their matching curvature distance is at least  $c_0|s|$ , and Lemma A.24.1 together with the stopped entropy argument gives

$$D_{\text{KL}}(P_s \| P_{-s}) \leq Cs^6 \mathbb{E}_s \sum_{i \leq T} d_{i,+}^2 \leq Cd_*^2 s^6 \mathbb{E}_s T.$$

Common below-onset queries contribute zero. Choose  $s = L\varepsilon$  with  $L$  fixed large enough that the curvature separation exceeds  $2\varepsilon$ . The nearest-alternative test based on the estimated multiset has both errors at most  $\eta$  by (A.13.8) and the triangle inequality for matching distance. Applying (A.14.5) yields

$$\sup_{\theta \in K} \mathbb{E}_\theta T \geq c\varepsilon^{-6} \log(1/\eta).$$

Small enough  $\varepsilon_0$  keeps both alternatives in  $K$ .

For the gap lower bound use the circular tables of Lemma A.14.1 with  $\Delta = 3\delta$ . They remain in  $K$  for small  $\delta_0$ . Choosing the nearer gap to  $\hat{g}$  gives a test with both errors at most  $\eta$ . Apply (A.14.5) in the higher-gap to lower-gap direction and then (A.14.3). It follows that

$$\sup_{\theta \in K} \mathbb{E}_\theta T \geq c\delta^{-2} \log(1/\eta).$$

A policy with infinite worst-case expected cost already satisfies the lower bound; otherwise both stopped entropy arguments apply. Combining the two bounds and using  $\max\{u, v\} \geq (u + v)/2$  proves the lower half of (A.14.4).  $\square$

**Corollary A.14.3** (Curvature loss, finer timing, and shrinking designs). *In the fixed-collar class, curvature-only recovery has optimal order  $\varepsilon^{-6} \log(1/\eta)$ . For each fixed  $m \geq 1$ , the joint loss with curvature tolerance  $\varepsilon$  and gap tolerance  $\delta = \varepsilon^{3+3/m}$  instead has optimal order*

$$\varepsilon^{-(6+6/m)} \log(1/\eta). \quad (\text{A.14.6})$$

*Separately, any all-history design whose positive offsets are at most  $Hh$  and whose flight numbers are at most  $J$ , even with the exact gap supplied, needs worst-case expected cost at least*

$$ch^{-2}\varepsilon^{-6} \log(1/\eta) \quad (\text{A.14.7})$$

*for curvature accuracy  $\varepsilon$  on the fixed-gap splitting subfamily. Consequently the supplied-bracket acquisition theorem has matching upper and lower orders at  $h \asymp \varepsilon^{3/m}$ , including the confidence logarithm.*

*Proof.* The first assertion is the curvature part of the preceding proof and Corollary A.13.3. Substitute the displayed  $\delta$  in (A.14.4) to obtain (A.14.6). For the shrinking class use instead  $\mathbb{E}_s \sum d_{i,+}^2 \leq H^2 h^2 \mathbb{E}_s T$  in the curvature entropy argument; the alternatives have the same gap, so supplying it has no effect. This proves (A.14.7).

In Theorem A.23.2 the supplied bracket is  $|g - g_0| \leq h/4$ . Its pilot offsets are in  $[3h/4, (m + 5/4)h]$ , and all its second-stage offsets are in  $[(L - 3J/4)h, (Lm + 3J/4)h]$  with  $L = J + 1$ , even on an unsuccessful pilot history. They fit such a shrinking collar with a fixed  $H_m$ . Its expected upper bound and (A.14.7) agree after substituting  $h \asymp \varepsilon^{3/m}$ .  $\square$

Thus the fixed-offset comparison improves the curvature-only benchmark, without invalidating the finer joint timing result. Theorem A.14.2 also identifies a timing obstruction in the larger fixed-collar class; its proof does not use the shrinking-design restriction. At the joint cost in (A.14.6), the exact four-window procedure may recover curvature more accurately than requested, because it estimates all symmetric coordinates to the finer gap tolerance. These statements make no assertion about the computational cost of evaluating exact finite-offset probabilities or about a uniform limit  $m \rightarrow \infty$ .

## APPENDIX A.15. RECONSTRUCTION WITH UNKNOWN SMOOTH REMAINDER FUNCTIONS

The exact inverse in Section A.13 uses the known nonlinear family at positive offset. To separate that assumption from the acquisition of leading invariants, we now enlarge the observation model explicitly. The underlying parameter is still  $(g, e) \in K$  and its leading amplitudes are still the

physical  $C_j(g, e)$ . The nuisance functions below need not come from that exact analytic support family.

Fix  $m \geq 1$ , constants  $0 < c_* < C_* < \infty$ , and a sufficiently small  $d_* > 0$ . For each  $j = 1, 2, 3$ , let the unknown function  $H_j$  satisfy

$$\begin{aligned} H_j &\in C^m([0, d_*]), \quad c_* \leq H_j \leq C_*, \quad \|H_j\|_{C^m} \leq C_*, \quad H_j(0) = C_j(g, e), \\ p_j(t) &= (t - jg)_+^2 H_j((t - jg)_+), \quad t - jg \leq d_*. \end{aligned} \quad (\text{A.15.1})$$

Choose  $C_* d_*^2 < 1$ . Each query returns an independent Bernoulli variable with this probability. The fixed neighborhood  $K$ , the common bounds, and a fixed interval  $[L_0, U_0]$  containing the possible gaps are known. No value or derivative of  $H_j$  away from zero is supplied. The class can be chosen to contain the exact physical finite-offset probabilities on a smaller compact neighborhood. It is an observation envelope: we do not assert that every arbitrary nuisance triple in it is realized by a single billiard table. The target area and curvatures are defined from  $(g, e)$  by the same physical leading inverse.

**Theorem A.15.1** (Uniform count recovery over a smooth envelope). *Let  $0 < U_0 - L_0 = W < d_*/12$ . For fixed  $m$  and the preceding bounds, there are  $h_m, C_m > 0$  such that for  $0 < h < h_m$  and  $0 < \eta < 1/4$  a Borel count-only procedure satisfies, uniformly over (A.15.1),*

$$\Pr \left\{ \begin{array}{l} |\hat{g} - g| \leq C_m h^{m+1}, \\ \|\hat{e} - e\|_\infty + |\hat{A} - A| \leq C_m h^m, \\ \text{dist}_{\text{match}}(\hat{\kappa}, \kappa) \leq C_m h^{m/3} \end{array} \right\} \geq 1 - \eta. \quad (\text{A.15.2})$$

*Its number of preparations is bounded on every history by*

$$C_m h^{-(2m+2)} \log(C_m/\eta). \quad (\text{A.15.3})$$

*It uses only flight numbers 1, 2, 3 and programmed times. It supplies its own accuracy-dependent bracket from the fixed interval, and requires no exact finite-offset formula, contact positions, labels, or independently measured area.*

*Proof.* We give a deterministic-budget form of the calibration and extrapolation construction to specify its dependence on the nuisance functions.

First use the bracket search of Lemma A.25.1 with failure allowance  $\eta/3$ . Its proof needs only zero probability below onset and the lower bound  $p_1(g + d) \geq c_* d^2$ , both part of (A.15.1). It returns a midpoint  $g_0 \in [L_0, U_0]$  and, on its good event,  $|g - g_0| \leq h/4$ . Its all-history cost is  $Ch^{-2} \log((2 + \log(W/h))/\eta)$ .

Next take a predetermined

$$n = \left\lceil C_m h^{-(2m+2)} \log(C_m/\eta) \right\rceil$$

fresh preparations at each pilot time  $g_0 + lh$ ,  $1 \leq l \leq m + 1$ . On bracket success all actual offsets lie between  $3h/4$  and  $(m + 5/4)h$ , and the probabilities are between positive multiples of  $h^2$ . The Bernoulli concentration bound proved in Lemma A.23.1, divided by those probabilities, gives relative errors at most  $c_m h^m$  simultaneously, with conditional probability at least  $1 - \eta/3$ , after increasing  $C_m$ .

Interpolate the square roots of the empirical fractions by a polynomial of degree at most  $m$  in the variable  $z = t - g_0$  at nodes  $lh$ . Use its unique-zero-or-zero convention on  $[-h/2, h/2]$  and call the result  $\hat{\tau}$ . This rule is defined and Borel also when a fraction is zero. Put  $\hat{g} = g_0 + \hat{\tau}$ . On every history  $|\hat{g} - g_0| \leq h/2$ . On the good pilot event, write

$$d\sqrt{H_1(d)} = \sqrt{C_1} d\sqrt{H_1(d)/C_1}.$$

The second square-root factor has value one at zero and bounded  $C^m$  norm, uniformly over the envelope. Multiplying its degree- $(m-1)$  Taylor polynomial by  $d$  gives a polynomial of degree  $m$  with a simple zero at onset. The fixed-node Lagrange bounds and the root argument of Lemma A.20.2 therefore give  $|\hat{g} - g| \leq C_m h^{m+1}$ . Only right derivatives through order  $m$  were used; no probability was evaluated below onset for this Taylor argument.

For the last stage put  $L = 4$  and take  $n$  fresh preparations at  $j\hat{g} + Llh$  for  $j = 1, 2, 3$  and  $l = 1, \dots, m$ . Use

$$\omega_l = (-1)^{l-1} \binom{m}{l}, \quad \hat{C}_j = \sum_{l=1}^m \omega_l \frac{\hat{p}_{j,l}}{(Llh)^2}.$$

On bracket success every actual offset of this stage lies in a fixed positive interval of multiples of  $h$ , even if the pilot estimates are inaccurate. On the two good events, its displacement from  $Llh$  is at most  $3C_m h^{m+1}$ . The derivative of  $d^2 H_j(d)$  is  $O_m(h)$  on these intervals, so the resulting normalized timing error is  $O_m(h^m)$ . The identities

$$\sum_l \omega_l = 1, \quad \sum_l \omega_l l^r = 0 \quad (1 \leq r < m)$$

and the bounded  $m$ th derivative of  $H_j$  give extrapolation bias  $O_m(h^m)$ . Conditional Bernoulli concentration on these fresh preparations gives the same stochastic order, with additional failure probability at most  $\eta/3$ . Hence  $\max_j |\hat{C}_j - C_j| \leq C_m h^m$  on their intersection.

Fit  $(\hat{g}, \hat{C}_1, \hat{C}_2, \hat{C}_3)$  on the fixed compact physical leading-data set using Lemma A.22.1. The jointly Lipschitz coefficient inverse of Theorem A.12.1 gives coefficient and area errors  $O_m(h^m)$ , and its cubic root bound gives curvature error  $O_m(h^{m/3})$ . Retain the pilot  $\hat{g}$  as the gap estimate; the fitted physical parameter and the pilot need not have identical gaps. This distinction preserves the finer asserted timing error.

All queries are legitimate even after a failed search. The returned  $g_0$  remains in  $[L_0, U_0]$  and  $|\hat{g} - g_0| \leq h/2$ . Thus pilot positive offsets are at most  $W + (m+1)h$ , and final positive offsets at most  $3(W + h/2) + 4mh$ . The fixed restriction on  $W$  and small enough  $h_m$  place both in  $[0, d_*]$ ; negative offsets are deterministic failures. No waiting time depends on seeing a success. The pilot and final stage have exactly  $(m+1+3m)n$  preparations on every history. As in Theorem A.25.2, their bound absorbs the bracket search cost because  $h^{2m} \log(2 + \log(W/h))$  is bounded. The three conditional failure allowances sum to at most  $\eta$ , proving the result.  $\square$

Taking  $h \asymp \varepsilon^{3/m}$  yields curvature error  $O(\varepsilon)$  and gap error  $O(\varepsilon^{3+3/m})$  with the preparation power  $6 + 6/m$ . This result is uniform over unknown smooth higher-order remainders. The theorem gives an upper bound; the matching envelope lower bound, under the additional strict interior condition, is proved in Theorem A.16.1. In the exact analytic subexperiment, the same joint accuracy has the lower bound of Theorem A.14.2; for curvature loss alone its fixed-window benchmark has power six instead. The two experiments therefore differ in the information supplied to the estimator, even when their stated joint guarantees have the same preparation order.

#### APPENDIX A.16. THE SHARP COST OF UNKNOWN SMOOTH REMAINDERS

We prove a matching lower bound for the smooth-remainder upper bound. The experiment retains three binary reports with the physical leading-amplitude map. Its nuisance functions range over the stated smooth envelope; they are not asserted to be exact finite-offset probabilities of billiard tables. The construction uses a fixed interior norm margin, physical splitting paths for the leading parameters, and nuisance splices that agree outside a shrinking collar. The stopped-transcript argument permits adaptive queries and random stopping throughout.

Use the sum convention  $\|H\|_{C^m} = \sum_{r=0}^m \|H^{(r)}\|_\infty$  in (A.15.1). Fix all its constants, a sufficiently small compact physical neighborhood  $K$  of  $\theta_c = (g_c, e) = (1/2, 0)$ , and the interval  $[L_0, U_0]$  of Theorem A.15.1. The set  $K$  is to contain the two-sided physical splitting paths and the circular gap variations used below. In addition suppose that some  $\rho > 0$  obeys

$$c_* + 4\rho < C_j(\theta_c) < C_* - 4\rho, \quad j = 1, 2, 3. \quad (\text{A.16.1})$$

This condition gives strict positivity and a norm budget around the constant reference nuisances. It can be arranged when the envelope is chosen, but is not a consequence of its non-strict bounds alone.

An admissible policy uses parameter-independent randomization and, after each history, either stops or chooses  $j \in \{1, 2, 3\}$  and a time  $t \geq 0$ . Require  $t - jg(\theta) \leq d_*$  for every  $\theta \in K$  on every history, including histories at which a previous calibration failed. Each fresh preparation returns the bit with probability (A.15.1). The observation is the stopped transcript of these queries and bits, together with the recorded common randomization. For  $0 < \eta < 1/4$ , let  $\mathcal{N}_m^{\text{env}}(\varepsilon, \delta, \eta)$  be the infimum of the worst-case expected number of preparations among such policies and Borel estimators satisfying

$$\Pr_{\theta, H} \{ \text{dist}_{\text{match}}(\hat{\kappa}, \kappa(\theta)) \leq \varepsilon, \quad |\hat{g} - g| \leq \delta \} \geq 1 - \eta \quad (\text{A.16.2})$$

uniformly over the envelope. The two outputs are target estimates; they are not required to be coordinates of one exactly fitted table.

**Theorem A.16.1** (Joint minimax preparation cost over the smooth envelope). *Under these fixed choices and (A.16.1), there are  $c, C, \varepsilon_0, \delta_0 > 0$  such that*

$$\begin{aligned} c(\varepsilon^{-(6+6/m)} + \delta^{-2}) \log(1/\eta) &\leq \mathcal{N}_m^{\text{env}}(\varepsilon, \delta, \eta) \\ &\leq C(\varepsilon^{-(6+6/m)} + \delta^{-2}) \log(1/\eta) \end{aligned} \quad (\text{A.16.3})$$

for  $0 < \varepsilon < \varepsilon_0$ ,  $0 < \delta < \delta_0$  and  $0 < \eta < 1/4$ . The upper bound has a deterministic, all-history preparation budget. The lower bound permits adaptive queries and random stopping. For curvature loss alone the order is  $\varepsilon^{-(6+6/m)} \log(1/\eta)$ , even when the common gap in the lower-bound construction is supplied.

*Proof.* We first construct the curvature alternatives within one fixed envelope, then charge the stopped experiment, and finally treat the timing loss and the upper bound.

*Physical leading parameters.* In the support family of Section A.10, set  $R = 1/4$ ,  $\alpha = \zeta = 0$  and  $\beta = \pm s$ . The two radius multisets are

$$(R + 36s, R - 18s, R - 18s), \quad (R - 36s, R + 18s, R + 18s).$$

Their gap is  $g_c = 1/2$  and their elementary coefficient triples are

$$e^\pm = (0, -972s^2, \pm 11664s^3). \quad (\text{A.16.4})$$

For small  $s > 0$ , sorting the positive reciprocals gives

$$\text{dist}_{\text{match}}(\kappa^+, \kappa^-) = \frac{36s}{R^2 - 324s^2} \asymp s. \quad (\text{A.16.5})$$

For example the middle sorted coordinates differ by the displayed quantity; the two outer differences are smaller when  $s$  is small. The leading amplitudes are analytic in a full coefficient neighborhood by the leading inverse proved above. The two points in (A.16.4) differ by  $O(s^3)$  there, so, with constants independent of  $s$ ,

$$|C_j^+(s) - C_j^-(s)| \leq Cs^3, \quad j = 1, 2, 3. \quad (\text{A.16.6})$$

This uses physical leading parameters, but not a free variation of area or a separated-root inverse.

*Splicing the nuisance functions.* Set  $\psi(u) = (1 - u)_+^{m+1}$  on  $[0, \infty)$ . It is  $C^m$ , equals one at zero, and it and its first  $m$  derivatives vanish at one. Write  $\bar{C}_j = (C_j^+ + C_j^-)/2$  and  $D_j = (C_j^+ - C_j^-)/2$ . Choose a fixed sufficiently large  $L$  and put

$$h = Ls^{3/m}, \quad H_j^\pm(d) = \bar{C}_j \pm D_j\psi(d/h). \quad (\text{A.16.7})$$

These functions have exactly the prescribed leading values at zero. For  $0 \leq r \leq m$ ,

$$\|(H_j^\pm - \bar{C}_j)^{(r)}\|_\infty \leq C_m L^{-r} s^{3(1-r/m)}. \quad (\text{A.16.8})$$

Choose  $L$  first so that the order- $m$  bound is smaller than a specified part of  $\rho$ . Then make  $s$  small so that the sum of all lower-order bounds, and  $|\bar{C}_j - C_j(\theta_c)|$ , use less than the remaining margin. Condition (A.16.1) proves all the positivity and sum-norm bounds in the same fixed envelope. Also  $h < d_*$  for sufficiently small  $s$ . Only the alternatives depend on accuracy; the parameter set and the nuisance norm ball do not.

*Information at every allowed query.* A common query has the same offset  $d = t - jg_c$  under the two alternatives. At  $d \leq 0$  both probabilities are zero, and at  $d \geq h$  they agree exactly. For  $0 < d < h$ ,

$$|p_j^+ - p_j^-| \leq Cd^2 s^3, \quad p_j^- \geq c_* d^2, \quad 1 - p_j^- \geq 1 - C_* d_*^2 > 0.$$

The inequality  $\log x \leq x - 1$  gives the Bernoulli entropy bound  $D_{\text{KL}}(\text{Ber}(p) \parallel \text{Ber}(q)) \leq (p - q)^2 / [q(1 - q)]$  for  $0 < q < 1$ . Consequently every query has

$$D_{\text{KL}}(\text{Ber}(p_j^+) \parallel \text{Ber}(p_j^-)) \leq Cd^2 s^6 \mathbf{1}_{\{0 < d < h\}} \leq Cs^{6+6/m}. \quad (\text{A.16.9})$$

Queries are not restricted to shrinking offsets. It is the nuisance pair itself that makes larger offsets uninformative.

*Stopping and confidence.* Condition on the common randomization and on the same realized finite history. The policy's next query or stop decision is then the same under the two hypotheses. Pad the transcript by a fixed stop symbol after stopping. The entropy chain rule for its first  $M$  positions and (A.16.9) give

$$D_{\text{KL}}(\mathbf{P}_+^{[M]} \parallel \mathbf{P}_-^{[M]}) \leq Cs^{6+6/m} \mathbb{E}_+(T \wedge M).$$

Entropy on the increasing finite-transcript sigma-fields converges to the entropy on the full transcript: one may use its variational formula and conditional approximations to each bounded test function. Monotone convergence on the right therefore yields

$$D_{\text{KL}}(\mathbf{P}_+ \parallel \mathbf{P}_-) \leq Cs^{6+6/m} \mathbb{E}_+ T. \quad (\text{A.16.10})$$

An infinite worst-case expected cost already satisfies the lower bound. No deterministic stopping assumption has been imposed.

Choose  $s$  a fixed sufficiently large multiple of  $\varepsilon$  so that (A.16.5) exceeds  $2\varepsilon$ . Classifying the curvature estimate by its nearer alternative gives a test with both errors at most  $\eta$ . Binary data processing implies

$$D_{\text{KL}}(\mathbf{P}_+ \parallel \mathbf{P}_-) \geq (1 - 2\eta) \log \frac{1 - \eta}{\eta} \geq \frac{1}{4} \log(1/\eta), \quad 0 < \eta < 1/4. \quad (\text{A.16.11})$$

The first inequality follows by minimizing the binary entropy over the two error constraints. For the second,  $1 - 2\eta \geq 1/2$  and  $(1 - \eta)/\eta \geq \eta^{-1/2}$  on this interval. Together with (A.16.10), this proves the curvature lower bound, also for curvature-only loss with supplied gap.

*Timing.* Use the circular leading parameters  $(g_c, 0)$  and  $(g_c + \Delta, 0)$ , and constant nuisance functions  $H_{j,0} = C_j(g_c, 0)$  and  $H_{j,\Delta} = C_j(g_c + \Delta, 0)$ . They are in the same envelope by strict slack and continuity. At a common query, put  $d = t - jg_c$  and  $z = j\Delta$ . Then

$$p_0 = d_+^2 C_j(g_c, 0), \quad p_\Delta = (d - z)_+^2 C_j(g_c + \Delta, 0).$$

Both vanish for  $d \leq 0$ . For  $d > 0$ , splitting the cases  $d \leq z$  and  $d > z$ , expanding the squares and using the bounded derivative of  $C_j(g, 0)$  gives  $|p_\Delta - p_0| \leq Cd\Delta$ . Since  $p_0 \geq cd^2$  and  $1 - p_0$  is bounded away from zero, the entropy in the support-compatible direction satisfies

$$D_{\text{KL}}(\text{Ber}(p_\Delta) \| \text{Ber}(p_0)) \leq C\Delta^2.$$

The argument includes  $p_\Delta = 0$  and does not reverse this entropy direction. Set  $\Delta = 3\delta$  and repeat the stopped chain rule and the classification argument using the gap estimate. It gives the lower cost  $c\delta^{-2} \log(1/\eta)$ . The curvature and timing pairs need not be the same pair. Their two worst-case bounds combine by  $\max(a, b) \geq (a + b)/2$ , proving the lower half of (A.16.3).

*The upper bound.* In Theorem A.15.1, choose

$$h = c_m \min\{\varepsilon^{3/m}, \delta^{1/(m+1)}\}$$

with fixed  $c_m > 0$  sufficiently small. Both target errors then meet (A.16.2). The all-history preparation bound is at most a constant times

$$\max\{\varepsilon^{-(6+6/m)}, \delta^{-2}\} \log(1/\eta),$$

which is bounded by the required sum. The known fixed gap interval and all-history admissibility were part of that theorem. Its pilot gap estimate is retained separately from the fitted coefficients, as required for the finer timing loss. Omitting the timing loss and using  $h = c_m \varepsilon^{3/m}$  proves the curvature-only upper bound and completes the proof.  $\square$

**Corollary A.16.2** (Information in a finite positive-window design). *For any prescribed finite set of query windows whose offsets at  $g_c$  are all strictly positive, there are two different curvature targets in the fixed smooth envelope which give exactly the same law for every stopped experiment using only those windows, including adaptive choices and arbitrary repetitions among them.*

*Proof.* Let  $d_{\min} > 0$  be the smallest such offset. Choose  $s > 0$  in (A.16.7) so small that  $h < d_{\min}$ . The two nuisance functions, and hence their probabilities, agree at every permitted query. Induction on transcript length proves equality on every finite history, and thus on the stopped observation. The curvature targets differ by (A.16.5).  $\square$

This corollary identifies the extra information used in the exact four-window inverse; it does not contradict that inverse. Under the exact analytic physical family, Theorem A.14.2 gives order  $(\varepsilon^{-6} + \delta^{-2}) \log(1/\eta)$ . For unknown  $C^m$  remainders with strict slack, Theorem A.16.1 gives order  $(\varepsilon^{-(6+6/m)} + \delta^{-2}) \log(1/\eta)$ . At the finer timing accuracy  $\delta = \varepsilon^{3+3/m}$  these orders coincide. The envelope alternatives establish an information comparison, not a new geometrically realized lower bound for exact billiard probabilities or a restriction on the general forward law.

#### APPENDIX A.17. RELATION TO ACTION, SPECTRAL, AND STATISTICAL INVERSE PROBLEMS

The analysis starts from a discrete variational problem but ends with a normalized physical probability. Bolotin and Treschev [2] develop Hill-type identities connecting action Hessians and monodromy in discrete and continuous Lagrangian systems. Such determinant identities explain the linear structure used here. The further step in Theorem 17.2 is a relative nonlinear determinant limit with two separated boundary perturbations, followed by full-phase residual-time integration. Propositions 18.1–18.3 separate the exact finite identity, trace-class transport and normalized integration. The additional physical consequence is Theorems 20.1 and 20.2: the two same-type scalar laws determine the complete energy profiles stably in a specified norm, and predict the mixed-type law. The finite-record compatibility defect in Corollary 20.5 is controlled by bridge length. It does not follow from just the value of the linear multiplier.

Zelditch [14] relates successive localized wave-trace invariants to analytic boundary geometry. His construction uses multiple oscillatory boundary integrals and a stationary-phase expansion;



the inverse spectral conclusions concern the specified symmetric classes of simply connected analytic plane domains, including a symmetry reversing a bouncing-ball orbit. This is an important predecessor for the coefficient-to-boundary strategy, not the same observation theorem. Our coefficients arise from a positive full-phase residual-time integral, not a Dirichlet or Neumann wave trace. The highest-jet diagonal in (A.5.2) separates the stationary-action contribution from the relative interior determinant contribution. Exact analytic germ recovery and finite-jet coefficient conditioning are kept distinct from noise-stable global analytic continuation.

Marked-length observations provide a different route to geometry. Bálint, De Simoi, Kaloshin and Leguil [1] recover period-two curvatures and periodic Lyapunov exponents from the marked length spectrum of open dispersing billiards. Osterman [13] studies analytic rigidity by homoclinic asymptotics and normal forms, including recovery with two scatterers supplied. De Simoi, Kaloshin and Leguil [5] prove marked-length determination of analytic chaotic billiards with axial symmetries, under symmetry and genericity hypotheses in a no-eclipse open configuration. Their theorem concerns the geometry of the table, not merely a specified contact jet.

The data in Corollary A.4.2 are instead a selected onset probability and its offset coefficients under normalized phase volume. The identical-even contact hypothesis makes its inverse triangular. Analytic continuation then determines each participating boundary in its contact frame, but supplies neither a lattice nor the other obstacles. There is no implication here between equality of these statistical data and equality of marked length spectra. In particular, the corollary is not a replacement for the global table determination in [5]. Its coefficient formula and its physically realized finite-jet coordinates identify the precise inverse observation being used. Proposition A.4.4 also shows explicitly that this information is not exclusive to long records.

For periodic finite-horizon Sinai billiards, Finamore and Leguil [12] introduce an enriched marked length spectrum and prove isometry from equality of that datum, using geodesic approximations and CAT(0) geometry. Their enriched observation and finite-horizon assumptions should not be identified with the selected count probabilities considered here. Conversely our general boundary-law construction does not require finite horizon.

The exact exponential-sum step in Theorem A.10.1 belongs to the theory of Prony reconstruction. Batenkov and Yomdin [10] analyze local accuracy for confluent Prony systems. The quantitative result here uses a specific feature of the hyperbolic-sine observation: its first four functions have the nonzero Wronskian (A.11.6). They give regular symmetric coefficient coordinates even when all three radii coincide. The classical coefficient-to-root estimate then gives the pairwise exponent  $1/3$ ; the realized physical path proves its sharpness. No new general root perturbation theorem or universally improved Prony algorithm is needed.

The source and onset series use a different physical window at each coefficient. Accordingly their first poles and convergence domains are consequences of the relative boundary law, not assertions about a pressure function for all long itineraries. Return-time and compound Poisson laws such as [6, 9], and the anisotropic statistical frameworks in [3, 4], address other observations and asymptotic regimes. The fixed-window companion [7] retains its separate transversality and moving-level argument. The historical Fourier-inversion analysis [8] remains available with its own explicit hypotheses; those hypotheses are not substituted for the finite-threshold proofs in this article.

Finally, polynomial extrapolation and binomial concentration are classical ingredients of the acquisition statements. Their uniform use here depends on the differentiated physical onset expansion. The self-calibration theorem adds a concrete statistical use of its quadratic vanishing: square-root probabilities have a simple onset zero, independent of the unknown normalizer. Its proof accounts for both the calibration preparations and the subsequent conditional records, and preserves a positive sampling window on every pilot outcome. The sufficient cost is not represented as a minimax rate.

The relative boundary law has an experiment-level consequence beyond coefficient identification. Theorem 34.1 controls total variation for the endpoints and residual time, uniformly down to zero offset, and retains the rare-event mass when unsuccessful preparations are restored. It permits bounded-risk replacement for growing independent samples whenever their expected number of successes times  $\tau^j$  tends to zero. Its proof uses the quadratic vanishing of the action error together with the normalized determinant estimate; an absolute approximation to a vanishing flux would not give this conclusion.

The finite-flight inverse of Theorem A.5.1 and the sensitivity formula (A.5.4) concern fixed jet order and varying collision length. The area constraint (A.12.2) explains why three leading amplitudes suffice in the constrained family while four allow an independent normalizer. Section A.13 uses a different fourth datum: a second first-flight window separates an unknown gap from the three symmetric shape coordinates in the exact physical model.

The minimax statement in Theorem A.14.2 concerns a specified binary experiment, not every observation available from a billiard. Its curvature and timing obstructions are realized by two different physical pairs. The smooth envelope in Theorem A.15.1 removes the supplied exact finite-offset formula while preserving a positive, fully charged reconstruction procedure. The earlier supplied-bracket and fixed-bracket procedures remain valid with their complete proofs. They are neither counterexamples to the fixed-window benchmark nor general lower bounds for all data.

Finally, the support calculation in Theorem A.9.1 is an elementary disk–ellipse comparison after a common whitening. Its role is to identify the constant in a physical testing profile, not to replace the nonlinear relative law. The projection argument in Theorem A.9.7 completes that profile when the full sample’s tangent approximation error no longer tends to zero. At fixed positive offset, the genuinely nonlinear comparison is still supplied by Theorems 12.1 and 34.1.

The smooth-envelope optimum in Theorem A.16.1 uses the additional strict-slack condition and allows free nuisance functions. It is not a lower bound obtained from a pair of exact nonlinear billiard laws. Its attribution to the referee memorandum is separate from the boundary compatibility and profile inverse developed here. Neither the count theorem nor its envelope comparison is used as a substitute for the long-bridge consequence.

The whole-profile count theorem, Theorem A.6.1, occupies a different place from these finite-dimensional benchmarks. It uses the same labelled scalar events as the relative boundary law, not their endpoint positions, and supplies no exact finite-offset probability formula to the estimator. The known normalization and uniform relative estimates allow its finite-bridge bias to be separated from scalar sampling noise. The positive-node regularization makes recovery possible on the full collar without an observed derivative oracle. Its preparation bound records both the cost of differentiating noisy scalar data and the exponentially decreasing event mass. This is the operational interpretation of the boundary energy invariant; it neither subsumes nor borrows the lower bounds of the other experiments.

The self-calibrated profile theorem in Section A.8 is distinct from the calibrated construction just described. It permits unknown  $g, A, \gamma$  because it supplies and charges a new one-/two-flight pilot on a uniform smooth physical class. Patch labels, regularity certificates and coarse boxes remain supplied. Its proof neither applies the exact four-window inverse to a free remainder nor imports the envelope minimax lower bound.

## APPENDIX A.18. CONDITIONAL ENDPOINT METRICS AND CURVATURE OBSERVABILITY

The scalar coefficient in (13.1) contains only the ratio in (15.2). We now retain the metric which appears before that ratio is formed. All expectations in this section are for the actual full-phase preparation, conditioned on  $E_{j,e}(d)$ .

**Proposition A.18.1** (Quadratic endpoint statistics). *In the common transverse endpoint coordinates  $Y = (u, v)$ ,*

$$\mathbb{E}Y = d m_{j,e}(\xi, d), \quad \frac{1}{d} \text{Cov}(Y) = \frac{1}{3} H_{j,e}^{-1} + d B_{j,e}(\xi, d). \quad (\text{A.18.1})$$

*The vector and matrix remainders are smooth for  $d \geq 0$  with uniform fixed-order derivative bounds. For physical positions the scalar variance in (13.2) has the corresponding leading diagonal entry of  $H_{j,e}^{-1}/3$ , with the same  $O(d)$  remainder. In particular the leading metric is determined without a phase-space normalizer or a chosen transverse initial law.*

*Proof.* Integrating the first residual time in (16.2) gives the conditional endpoint density, in  $w = \sqrt{2d}z$  coordinates,

$$\frac{(1 - |z|^2) a_{j,e}(\xi, \sqrt{2d}z)}{\int_{|z| < 1} (1 - |z|^2) a_{j,e}(\xi, \sqrt{2d}z) dz} dz, \quad |z| < 1.$$

Its denominator is uniformly positive. The area and the exponentially small flux coefficient cancel in this conditional expression, but were included in the event probability. At  $d = 0$  the density is  $2(1 - |z|^2)/\pi$ . Direct radial integration gives

$$\mathbb{E}_0 z = 0, \quad \mathbb{E}_0(z z^t) = I_2/6. \quad (\text{A.18.2})$$

For example  $\mathbb{E}_0 |z|^2 = 1/3$  and rotational symmetry divides this between the two coordinates.

The endpoint map is  $Y = \Phi_{j,e}(\xi, w)$ , with linear part  $H_{j,e}^{-1/2} w$ . For any smooth amplitude, integration on this radial domain annihilates its odd Taylor terms. Applied to each component of  $Y a$  and  $Y Y^t a$ , the smooth radial lemma gives

$$\mathbb{E}Y = d m_{j,e}, \quad \mathbb{E}(Y Y^t) = \frac{d}{3} H_{j,e}^{-1} + d^2 \tilde{B}_{j,e}.$$

All statements include smooth right extensions and uniform derivatives: the endpoint maps and amplitudes have the bounds in Lemma 16.1, and each division by  $d$  uses the vanishing constant term of a smooth even integral. Subtracting  $(\mathbb{E}Y)(\mathbb{E}Y)^t$  proves (A.18.1).

The physical endpoint graph has the form  $Q_b = q_b + \mathbf{t} y_b \pm \mathbf{n} \psi_b(y_b)$ , with  $|\mathbf{t}| = |\mathbf{n}| = 1$  and  $\psi_b = O(y_b^2)$ . Its mean is  $q_b + d \tilde{m}_b$ . The leading term in  $|Q_b - q_b|^2$  is  $y_b^2$ , since the two frame vectors are orthogonal. The remaining squared term starts at degree four. Apply the same radial argument and subtract the squared mean displacement. This proves the physical scalar variance formula. In particular no unknown  $q_b$  occurs in the observable itself: it is eliminated by centering the measured positions.  $\square$

*Proof of Theorem 13.2.* The determinant-one middle matrix in (15.1) has inverse obtained by changing the signs of its off-diagonal entries. Thus

$$H_{j,e}^{-1} = \frac{g}{c \sinh \gamma} D_j \begin{pmatrix} \coth(j\gamma) & \text{csch}(j\gamma) \\ \text{csch}(j\gamma) & \coth(j\gamma) \end{pmatrix} D_j. \quad (\text{A.18.3})$$

For odd  $j$ ,  $\sigma_0^2 = c_1$  and  $\sigma_j^2 = c_0$ . Proposition A.18.1 gives (13.3), with its differentiated error estimate. Taking the geometric mean and ratio gives (13.4). The function

$$F_j(\gamma) = \frac{\coth(j\gamma)}{\sinh \gamma}$$

is strictly decreasing from infinity to zero on  $(0, \infty)$ , since

$$-F'_j(\gamma) = \frac{j \text{csch}^2(j\gamma)}{\sinh \gamma} + \frac{\coth(j\gamma) \cosh \gamma}{\sinh^2 \gamma} > 0. \quad (\text{A.18.4})$$

This proves unique recovery of  $\gamma$ , then of the two curvatures. At  $j = 1$ , the inversion is even explicit:  $F_1(\gamma) = c/(c^2 - 1)$ , so if  $P = 3\sqrt{v_0 v_1}/g$ , then  $c = (1 + \sqrt{1 + 4P^2})/(2P)$ .

We verify uniform conditioning, not just nonvanishing of a fixed-length Jacobian. On a compact positive gap-curvature box,  $0 < \gamma_- \leq \gamma \leq \gamma_+ < \infty$ . The second term of (A.18.4) gives a lower bound for  $-F'_j$  independent of  $j$ . Its upper bound is also uniform because  $j \operatorname{csch}^2(j\gamma_-) \leq C$  and  $\coth(j\gamma_-) \leq \coth \gamma_-$ . Both variances are bounded above and bounded away from zero. Taking their ratio and geometric mean is consequently Lipschitz; monotone inversion of  $F_j$ , then the last line of (13.4), is uniformly Lipschitz. Forward differentiation of (13.3) gives the reverse bound with the same type of constants. More explicitly, for data from two points of this box,

$$C^{-1} \|(g, \kappa_0, \kappa_1) - (\tilde{g}, \tilde{\kappa}_0, \tilde{\kappa}_1)\|_\infty \leq \|(g, v_0, v_1) - (\tilde{g}, \tilde{v}_0, \tilde{v}_1)\|_\infty \leq C \|(g, \kappa_0, \kappa_1) - (\tilde{g}, \tilde{\kappa}_0, \tilde{\kappa}_1)\|_\infty. \quad (\text{A.18.5})$$

No factor  $(\kappa_0 - \kappa_1)^{-1}$  appears. Nor do we divide by the exponentially small off-diagonal entry in (A.18.3).

For data outside the exact image, minimize its distance to that image on the compact curvature box at the known  $g$ ; a minimizer exists. The true pair is an admissible comparison, so the selected image point is within twice the data error of the true image. The error is bounded by  $\delta + Cd$  by Proposition A.18.1. Apply the inverse bound in (A.18.5). This proves the last assertion and also specifies a reconstruction for noisy or finite-offset data. A small error in the supplied  $g$  adds its size to the same bound, provided the conditional variances are still those of the stated physical window; we do not replace an onset-location error by a cost-free timing change.  $\square$

The distinction between odd and even lengths is informational. At an even length the first and last impacts are on the same obstacle, and the two diagonal leading variances coincide. At an odd length they sample the two endpoints. The theorem therefore uses the two diagonal entries at an odd length rather than an exponentially small covariance or an assumed equality of endpoint curvatures.

**Proposition A.18.2** (Independent-preparation estimation). *Suppose that  $2K$  independent full-phase preparations are retained upon success of  $E_{j,e}(d)$ , for a fixed odd  $j$ . The two variances have estimators  $\hat{V}_0, \hat{V}_1$ , using only their endpoint positions and the known offset  $d$ , such that*

$$\Pr \left( \max_b |\hat{V}_b - V_{j,b}(d)| > \delta \right) \leq 4 \exp(-cK\delta^2).$$

*A reconstruction as in Theorem 13.2 consequently has error at most  $C[d + \sqrt{\log(4/\eta)/K}]$  with probability at least  $1 - \eta$ . All constants are independent of odd  $j$ . The expected number of independent unconditioned preparations needed for these  $2K$  successes is comparable to*

$$Kd^{-2}e^{j\gamma_e}. \quad (\text{A.18.6})$$

*No independence of successive impacts in one trajectory is assumed.*

*Proof.* Pair the independent successful records and use, at each endpoint,

$$\hat{V}_b = \frac{1}{K} \sum_{k=1}^K \frac{|Q_b^{(2k)} - Q_b^{(2k-1)}|^2}{2d}.$$

Independence gives its expectation as the variance in (13.2). Coercivity confines each endpoint to a set of diameter  $C\sqrt{d}$ , independently of  $j$ . The summands are therefore in one bounded interval  $[0, C_0]$ . For completeness, a random variable in an interval of length  $C_0$  satisfies  $\mathbb{E}e^{t(X - \mathbb{E}X)} \leq e^{t^2 C_0^2/8}$ : the second derivative of the log moment generating function is a variance under a tilted law, bounded by  $C_0^2/4$ , and its value and first derivative vanish at zero. Multiplication over the independent pairs, the exponential Markov inequality, and optimization over  $t$  yield  $2e^{-2K\delta^2/C_0^2}$

at each endpoint. A union bound proves the displayed estimate. The deterministic inverse estimate gives the curvature bound. Rejection from independent preparations gives independent conditional successes, and the expected waiting time is  $2K/\Pr(E_{j,e}(d))$ . Theorem 13.1 and  $\sinh(j\gamma) \asymp e^{j\gamma}$  prove (A.18.6). Uniform conditional conditioning is therefore not a uniform unconditioned sampling cost.  $\square$

**A.18.1. Realized scalar fibers and their separation by records.** At known  $g, A$ , all leading selected-channel count coefficients depend on  $c_0, c_1$  only through  $c_0 c_1$ . Conversely, their first two values recover this product, since their ratio is  $2 \cosh \gamma$ . The following construction shows that this scalar fiber occurs for actual periodic billiards, not just for algebraic local jets.

**Theorem A.18.3** (A geometric fiber of the full leading count hierarchy). *On the rectangular lattice  $3\mathbb{Z} \times 4\mathbb{Z}$  there is a real-analytic family of strictly convex scatterers  $D_s$ ,  $|s| < s_0$ , with  $D_0$  the unit disk, for which the only shortest unoriented channel is horizontal,*

$$g(s) = 1, \quad A(s) = 12 - \pi, \quad c_0(s) = 2e^s, \quad c_1(s) = 2e^{-s}.$$

*All maximal-count threshold locations and every leading coefficient  $\lim_{d \downarrow 0} d^{-2} \Pr_s(N_{j+d} = j+1)$  are independent of  $s$ . The ordered contact curvatures vary. The two endpoint variances at any odd order distinguish  $s$ , with constants uniform in that order.*

*Proof.* Use the support function of the disk,  $h_0 \equiv 1$ . Put

$$\chi_0 = \frac{1}{4}(1 + \cos \theta) \sin^2 \theta, \quad \chi_\pi = \frac{1}{4}(1 - \cos \theta) \sin^2 \theta, \quad \chi_* = \sin^4 \theta.$$

At  $0, \pi$  all three functions and their first derivatives vanish. The second derivatives of  $\chi_0$  at these two points are  $(1, 0)$ , those of  $\chi_\pi$  are  $(0, 1)$ , and those of  $\chi_*$  are  $(0, 0)$ . These explicit trigonometric polynomials permit independent contact jets and a compensating change of area. Set

$$r_0(s) = (2e^s - 1)^{-1}, \quad r_1(s) = (2e^{-s} - 1)^{-1},$$

and

$$h_{s,a} = 1 + (r_0(s) - 1)\chi_0 + (r_1(s) - 1)\chi_\pi + a\chi_*.$$

For small  $s, a$ ,  $h + h'' > 0$ , so this is a smooth strictly convex body. At normals  $0, \pi$ , its support values stay one and its first support derivatives stay zero. Its two curvature radii are exactly  $r_0(s), r_1(s)$ , since the compensator has zero second jet at both normals. The supporting lines between neighboring horizontal translates therefore stay a distance one apart, and their support points attain that distance. The closest horizontal segment is unchanged.

The obstacle area is  $\mathcal{A}(h) = \frac{1}{2} \int_0^{2\pi} [h^2 - (h')^2] d\theta$ . Its derivative in  $a$  at  $(s, a) = (0, 0)$  is  $\int \chi_* d\theta > 0$ . Since  $\int \chi_* d\theta = 3\pi/4$ , the analytic implicit function theorem supplies a real-analytic  $a = a(s)$  with  $a(0) = 0$  and  $\mathcal{A}(h_{s,a(s)}) = \pi$  exactly. This variation has not altered either contact jet. At the disk, the vertical gap is two and every other nonhorizontal pair has still larger gap. The positive gap to the horizontal minimum, disjointness of translates, and clearance persist for small  $s$ . Hence the two horizontal orientations are the entire ground-onset event. Their sum in (13.1) has leading coefficient

$$\frac{1}{(12 - \pi) \sinh(j \operatorname{arccosh} 2)},$$

independent of  $s$ . This assertion concerns leading coefficients, not equality of the finite-offset remainders or all other billiard data. Finally (13.3) gives  $v_{j,1}/v_{j,0} = e^{2s}$  at every odd  $j$ . It separates the family; uniform conditioning follows from Theorem 13.2.  $\square$

The fixed contact patches can be chosen once for this entire family. Even when several minimizing pairs compete, a recorded first-impact patch and its outgoing direction select the oriented channel. On a nonminimal channel one must also use its own window and the itinerary

selection specified in Section 12. Counts in the ground collar alone eventually suppress a fixed gap difference  $\Delta > 0$ , since activation requires  $j\Delta < \varepsilon_0$ . The marked theorem neither changes that condition nor implicitly provides unobserved channel labels.

#### APPENDIX A.19. LEADING-DATA FACTORIZATION AND QUANTITATIVE ACQUISITION

**A.19.1. What increasing collision order does not add at leading order.** For one selected channel write  $A_j = \lim_{d \downarrow 0} d^{-2} \Pr(E_{j,e}(d))$ ; this uppercase subscripted symbol is distinct from the free area  $A$ .

**Proposition A.19.1** (One-flight factorization of the leading data). *At known gap  $g$ , the entire leading selected-channel probability and endpoint-metric hierarchy factors through  $(A_1, v_{1,0}, v_{1,1})$ . More explicitly, the two one-flight variances determine  $\gamma, c_0, c_1$ , and*

$$A_j = A_1 \frac{\sinh \gamma}{\sinh(j\gamma)}, \quad v_{j,b} = v_{1,b} \frac{\coth(j\gamma)}{\coth \gamma} \quad (j \text{ odd}, b = 0, 1). \quad (\text{A.19.1})$$

The full endpoint matrix at either parity is then given by (A.18.3). Also  $A = [2A_1 \sinh \gamma]^{-1}$ , when the normalizer is unknown.

*Proof.* The one-flight case of Theorem 13.2 recovers both curvatures. Divide (13.1) at zero offset by its one-flight version and do the same in (13.3). These are precisely (A.19.1). The displayed Hessian formula covers the even orders as well, including their off-diagonal entries. Solving the first amplitude for  $A$  proves the final assertion.  $\square$

Thus access to long odd records preserves access to the same leading contact parameters; it does not produce new leading inverse coordinates. The collision-order-uniform nonlinear calculus supplies something else: the relative boundary factorization, fixed-offset limiting law and higher-jet separation of Sections 17 and A.3. Its all-order smooth remainder also permits the following controlled bias cancellation. The latter uses classical polynomial extrapolation, rather than a new interpolation algorithm.

**A.19.2. Finite-order bias cancellation.** Fix an integer  $m \geq 1$ . All constants with subscript  $m$  below may depend on this fixed order; no uniform assertion as  $m \rightarrow \infty$  is made. Put

$$\omega_l = (-1)^{l-1} \binom{m}{l}, \quad d_l = lh, \quad 1 \leq l \leq m, \quad (\text{A.19.2})$$

where  $mh < \varepsilon_0$ . Each observation uses the same selected channel at the different physical window  $jg + d_l$ , with  $j$  odd. The true gap and the indicated offsets are assumed known unless otherwise stated. At each window use  $2K$  successful records from independent full-phase preparations, independently across the windows, and form the paired-position estimators  $\hat{V}_{l,b}$  of Proposition A.18.2.

**Theorem A.19.2** (Extrapolated recovery and preparation cost). *Set  $\hat{v}_b^{[m]} = \sum_{l=1}^m \omega_l \hat{V}_{l,b}$ . There is a reconstruction on the compact gap-curvature boxes of Theorem 13.2 such that, with probability at least  $1 - \eta$ ,*

$$\max_b |\hat{\kappa}_b - \kappa_b| \leq C_m \left( h^m + \sqrt{\frac{\log(8m/\eta)}{K}} \right). \quad (\text{A.19.3})$$

The expected total number  $T$  of independent unconditioned preparations is comparable, with constants depending on  $m$ , to  $Kh^{-2}e^{j\gamma}$ . In addition, the reconstruction and the bound

$$T \leq C_m e^{j\gamma} h^{-2} [K + \log(2m/\eta)] \quad (\text{A.19.4})$$

hold together with probability at least  $1 - \eta$ , after adjusting constants in (A.19.3). For each sufficiently small accuracy  $\epsilon > 0$ , this yields the sufficient expected and high-probability acquisition order

$$C_m e^{j\gamma} \epsilon^{-(2+2/m)} \log(C_m/\eta) \quad (\text{A.19.5})$$

for curvature error at most  $\epsilon$ . The experiment counts independent preparations, not independent successive impacts.

*Proof.* The  $m$ th forward difference annihilates every polynomial of degree less than  $m$ . Apply it at the integers  $0, 1, \dots, m$  to obtain

$$\sum_{l=1}^m \omega_l = 1, \quad \sum_{l=1}^m \omega_l l^r = 0 \quad (1 \leq r < m).$$

The smooth right extension of  $V_{j,b}(d)$  in Proposition A.18.1 has uniformly bounded derivatives through order  $m$ . Taylor's integral remainder at zero consequently proves

$$\left| \sum_{l=1}^m \omega_l V_{j,b}(lh) - v_{j,b} \right| \leq C_m h^m. \quad (\text{A.19.6})$$

This step uses the uniform differentiated nonlinear theorem. A pointwise  $O(d)$  statement with a length-dependent constant would not suffice for the stated uniform assertion.

Each paired summand at offset  $lh$  lies in a common bounded interval, by endpoint coercivity and its normalization by  $2lh$ . Proposition A.18.2 and a union bound over the  $m$  windows give simultaneous errors  $C\sqrt{\log(8m/\eta)/K}$  at both endpoints with failure probability at most  $\eta/2$ . Multiplication by  $\sum_l |\omega_l| = 2^m - 1$ , together with (A.19.6), gives the same form for the leading variance estimate. Projection onto the compact image of the leading observation map, followed by its uniform Lipschitz inverse, proves (A.19.3). It also defines the estimator when the extrapolated values do not belong to the exact image.

For the cost, put  $p_l = \Pr(E_{j,e}(lh))$ . The uniform threshold law and a smaller collar give

$$c(lh)^2 e^{-j\gamma} \leq p_l \leq C(lh)^2 e^{-j\gamma}.$$

Waiting for  $2K$  independent successes has mean  $2K/p_l$ . Summing over the fixed number of windows proves the expected-cost assertion. For a tail estimate let  $T_l$  be that waiting time. Then  $\Pr(T_l > N) = \Pr(\text{Bin}(N, p_l) < 2K)$ . If  $Np_l \geq 4K$ , the exponential Markov inequality for the binomial law gives

$$\Pr(\text{Bin}(N, p_l) < Np_l/2) \leq e^{-Np_l/8}.$$

For example this follows by multiplying the Bernoulli moment generating functions and optimizing a negative parameter; the bound  $\Pr(X \leq (1-t)\mathbb{E}X) \leq \exp(-t^2\mathbb{E}X/2)$  for  $0 < t < 1$  follows from  $1+x \leq e^x$  and  $t + (1-t)\log(1-t) \geq t^2/2$ . Choose  $N$  with  $Np_l \geq \max(4K, 8\log(2m/\eta))$ . The preceding lower bound on  $p_l$  and a union bound show that (A.19.4) fails with probability at most  $\eta/2$ . Combining with the estimation event proves the simultaneous assertion; no independence between those two events is needed.

Finally take  $h = c_m \epsilon^{1/m}$ , with a sufficiently small fixed  $c_m$ , and  $K$  at least  $C_m \epsilon^{-2} \log(C_m/\eta)$ . Equation (A.19.3) is then at most  $\epsilon$ , and both cost bounds have order (A.19.5).  $\square$

At  $m = 1$  the sufficient cost is  $e^{j\gamma} \epsilon^{-4} \log(C/\eta)$ , the explicit accuracy–cost interpretation of the original paired estimator. At  $m = 2$  it is  $e^{j\gamma} \epsilon^{-3} \log(C_2/\eta)$ . These are guarantees for the specified procedures, not minimax lower bounds or a statistical advantage of high collision order. The exponential preparation factor persists, and the extrapolation constants and required boundary differentiability grow with  $m$ . There is no independent sensor-noise model hidden in the estimates.

### A.19.3. Gap calibration and the physical window.

**Proposition A.19.3** (Propagation of a timing error). *Suppose the programmed windows are  $j\hat{g} + lh$ , whereas the true gap is  $g = \hat{g} + \Delta g$ . Assume  $|j\Delta g| \leq h/2$  and  $mh + |j\Delta g| < \varepsilon_0$ . Use the programmed offset  $lh$  in each paired-position denominator. The reconstruction based on the supplied gap  $\hat{g}$  then has error, with the same independent-preparation confidence convention,*

$$C_m \left[ h^m + \sqrt{\frac{\log(8m/\eta)}{K}} + \frac{j|\Delta g|}{h} + |\Delta g| \right]. \quad (\text{A.19.7})$$

A sufficient calibration for preserving an  $O(h^m)$  deterministic bias is  $|\Delta g| \leq c_m h^{m+1}/j$ . In particular the original  $m = 1$  procedure requires an  $O(h^2/j)$  gap accuracy by this direct normalization rule.

*Proof.* The actual offset is  $d_l = lh - j\Delta g > 0$ . The expectation of the paired statistic divided by  $2lh$  is

$$\frac{d_l}{lh} V_{j,b}(d_l).$$

Uniform boundedness and a bounded offset derivative imply that its difference from  $V_{j,b}(lh)$  is at most  $C[j|\Delta g|/(lh) + j|\Delta g|]$ . Since  $h$  is bounded above on the chosen collar, the first term absorbs the second after changing  $C_m$ . Sum with the fixed weights in (A.19.2) and use (A.19.6). The random summands remain uniformly bounded because  $d_l/(lh)$  lies between  $1/2$  and  $3/2$ . The concentration proof is unchanged. The joint gap–variance inverse bound (A.18.5) adds  $C|\Delta g|$  for using the supplied gap. This proves (A.19.7) and its stated sufficient calibration. Thus a gap error is not treated as a cost-free shift of the physical observation window.  $\square$

## APPENDIX A.20. LOCAL CALIBRATION FROM THE ONSET PROBABILITIES

The known-gap experiment of Theorem A.19.2 and an experiment which must calibrate that gap are different. We first compute the sensitivity of the former statistic, and then give a count-based calibration procedure whose preparation cost can be included explicitly. Throughout this section  $m$  is fixed and  $j$  is odd. Constants may depend on  $m$  and on the compact geometric family, but not on  $j$ .

**Proposition A.20.1** (The harmonic timing coefficient). *Let  $V(d)$  denote either of the normalized physical endpoint variances at the chosen flight number, and let  $v = V(0)$ . For the programmed-offset normalization in Proposition A.19.3, put  $\tau = j\Delta g$ . The extrapolated population statistic is*

$$W_m(\tau) = \sum_{l=1}^m \omega_l \left( 1 - \frac{\tau}{lh} \right) V(lh - \tau).$$

*Its first timing derivative satisfies*

$$W'_m(0) = -\frac{vH_m}{h} + O_m(1), \quad H_m = \sum_{l=1}^m \frac{1}{l}. \quad (\text{A.20.1})$$

*The remainder is uniform in  $j$ . In particular extrapolation does not cancel the leading normalization error.*

*Proof.* Differentiate the displayed statistic to obtain  $-\sum_l \omega_l [V(lh)/(lh) + V'(lh)]$ . The uniform estimates  $V(d) = v + O(d)$  and  $V'(d) = O(1)$  reduce the leading term to  $-v \sum_l \omega_l / lh$ . The binomial identity

$$\sum_{l=1}^m \frac{(-1)^{l-1} \binom{m}{l}}{l} = \int_0^1 \frac{1 - (1-x)^m}{x} dx = \sum_{l=1}^m \frac{1}{l}$$



proves (A.20.1). The last equality follows on setting  $y = 1 - x$  and summing  $1 + y + \dots + y^{m-1}$ .  $\square$

**A.20.1. A population calibration lemma.** Suppose that a coarse location  $g_0$  is supplied with

$$|j(g - g_0)| \leq h/4. \quad (\text{A.20.2})$$

This is a local bracketing assumption, not a fine calibration at scale  $h^{m+1}/j$ . The cost of locating this coarse bracket is not included in the experiment below. Write  $\tau = j(g - g_0)$  and observe the selected channel at the  $m + 1$  programmed times  $jg_0 + lh$ ,  $1 \leq l \leq m + 1$ . Let  $p_l$  be the respective probabilities. By the uniform threshold law,

$$p_l = K_j(lh - \tau)^2 F_j(lh - \tau), \quad K_j = \frac{1}{2A \sinh(j\gamma)}. \quad (\text{A.20.3})$$

All actual offsets are positive. The common factor  $K_j$  is unknown and will not have to be estimated.

Given positive approximations  $\hat{p}_l$ , interpolate their square roots at  $z = lh$  by a polynomial  $\hat{P}$  of degree at most  $m$ . If it has a unique real zero in  $[-h/2, h/2]$ , define  $\hat{\tau}$  to be that zero. Otherwise set  $\hat{\tau} = 0$ . This convention always keeps  $\hat{\tau}$  in the indicated interval.

**Lemma A.20.2** (Stable extrapolation of the onset zero). *There are  $h_m, \epsilon_m > 0$  such that, if  $0 < h \leq h_m$  and  $\max_l |\hat{p}_l/p_l - 1| \leq \epsilon \leq \epsilon_m$ , then  $\hat{P}$  has exactly one zero in  $[-h/2, h/2]$  and*

$$|\hat{\tau} - \tau| \leq C_m(h^{m+1} + h\epsilon). \quad (\text{A.20.4})$$

*After decreasing  $h_m, \epsilon_m$ , the right side is at most  $h/8$ . The assertion holds uniformly in  $j$  without knowing  $K_j$ .*

*Proof.* Put  $H_j(d) = \sqrt{F_j(d)}$ . Positivity and the uniform differentiated threshold law give  $H_j(0) = 1$  and bounded right derivatives through order  $m$ . At the positive nodal arguments  $d_l = lh - \tau$ ,

$$K_j^{-1/2} \sqrt{p_l} = d_l H_j(d_l).$$

Let  $Q(z)$  be the degree- $m$  Taylor polynomial obtained by multiplying  $z - \tau$  by the degree- $(m - 1)$  right Taylor polynomial of  $H_j$  at zero. Then  $Q(\tau) = 0$ , and  $Q'(z) = 1 + O_m(h)$  for  $|z| \leq h/2$ . Taylor's integral remainder at the positive  $d_l$  shows that the exact nodal values differ from  $Q(lh)$  by at most  $C_m h^{m+1}$ .

For the fixed nodes  $1, \dots, m + 1$ , let  $L_l$  be their Lagrange basis. On  $[-1/2, 1/2]$  both  $\sum |L_l|$  and  $\sum |L'_l|$  are bounded by constants depending only on  $m$ . Therefore interpolation of the nodal Taylor remainders changes the value by at most  $C_m h^{m+1}$  and the  $z$  derivative by at most  $C_m h^m$  throughout  $[-h/2, h/2]$ . This argument uses Taylor coefficients and polynomial evaluation only; it does not evaluate a physical probability at a negative offset.

If the probability errors are at most  $\epsilon$  relatively, the square-root errors are at most  $C\epsilon$  relatively. The normalized nodal square roots have size  $O_m(h)$ . The same Lagrange bounds thus add at most  $C_m h\epsilon$  to the value error and  $C_m \epsilon$  to the derivative error. Hence  $K_j^{-1/2} \hat{P}$  has derivative at least  $1/2$  on the interval for sufficiently small  $h, \epsilon$ , and its value at  $\tau$  is at most  $C_m(h^{m+1} + h\epsilon)$  in absolute value. Since  $|\tau| \leq h/4$ , its values at the two interval endpoints have opposite signs after decreasing these constants' allowed neighborhood. Monotonicity proves existence and uniqueness of the root and the mean-value theorem proves (A.20.4). The common unknown positive factor does not change any zero.  $\square$

### A.20.2. Calibration and conditional recovery in one cost bound.

**Theorem A.20.3** (Locally self-calibrated curvature acquisition). *Assume the coarse bracket (A.20.2), a fixed selected channel, and independent full-phase preparations. For  $0 < h \leq h_m$  and  $0 < \eta < 1/2$ , there is an experiment with the following properties. It estimates the gap by  $\hat{g} = g_0 + \hat{\tau}/j$  and the two contact curvatures by  $\hat{\kappa}_b$ , and, with probability at least  $1 - \eta$ ,*

$$|\hat{g} - g| \leq C_m h^{m+1}/j, \quad \max_b |\hat{\kappa}_b - \kappa_b| \leq C_m h^m. \quad (\text{A.20.5})$$

*Its total number  $T$  of unconditioned preparations, including calibration and conditional-position acquisition, satisfies*

$$\mathbb{E}T \leq C_m e^{j\gamma} h^{-(2m+2)} \log(C_m/\eta), \quad T \leq C_m e^{j\gamma} h^{-(2m+2)} \log(C_m/\eta) \quad (\text{A.20.6})$$

*with the latter bound holding together with (A.20.5) with probability at least  $1 - \eta$ . Thus for curvature accuracy  $\varepsilon$  a sufficient preparation order remains*

$$C_m e^{j\gamma} \varepsilon^{-(2+2/m)} \log(C_m/\eta).$$

*Unlike Theorem A.19.2, this statement includes the fine-gap calibration experiment. It still presupposes the local coarse bracket, fixed channel patches, and noiseless collision recording.*

*Proof.* For each of the  $m + 1$  calibration windows, start independent preparations until  $r$  selected events have occurred. Let  $T_l$  be this stopping time and put  $\hat{p}_l = r/T_l$ . We record the waiting-time estimate used here. For independent Bernoulli trials of probability  $p$ ,

$$\Pr\{|(r/T_r)/p - 1| > \epsilon\} \leq 2e^{-cr\epsilon^2}, \quad \mathbb{E}T_r = r/p, \quad (\text{A.20.7})$$

for  $0 < \epsilon < 1/4$  and  $r \geq C\epsilon^{-2}$ , with absolute constants. Indeed  $T_r > n$  is equivalent to fewer than  $r$  successes in  $n$  trials, and  $T_r \leq n$  is equivalent to at least  $r$  successes. Apply the binomial exponential Markov bound at the two integers nearest  $r/(p(1 - \epsilon))$  and  $r/(p(1 + \epsilon))$ . Rounding errors are absorbed by  $r\epsilon \geq C$ . The relative deviations of the binomial means from  $r$  are comparable to  $\epsilon$ , which gives the exponential bound. The expectation is the sum of  $r$  independent geometric waiting times of mean  $1/p$ .

Choose  $\epsilon = c_m h^m$  with a sufficiently small fixed  $c_m$  and

$$r = \left\lceil C_m h^{-2m} \log(C_m/\eta) \right\rceil.$$

Use (A.20.7) and a union bound over the calibration windows. Except on an event of probability at most  $\eta/4$ , their relative errors are at most  $c_m h^m$ . Lemma A.20.2 then gives  $|\hat{\tau} - \tau| \leq C_m h^{m+1}$  and the first estimate in (A.20.5).

Use fresh independent preparations for the conditional-position stage, at the  $m$  programmed windows  $j\hat{g} + lh$ . At each window collect  $2K$  selected records and form the paired variance estimator of Proposition A.18.2, normalized by the programmed  $lh$ . Take  $K = \lceil C_m h^{-2m} \log(C_m/\eta) \rceil$  and use the extrapolation weights (A.19.2). Conditional on the entire calibration experiment, these new preparations are independent, and all the variance concentration bounds are uniform in the chosen windows. On the calibration success event,

$$\frac{j|g - \hat{g}|}{h} = \frac{|\tau - \hat{\tau}|}{h} \leq C_m h^m.$$

For small  $h$  this is in the timing collar of Proposition A.19.3. That proposition and the conditional concentration bound give the second estimate of (A.20.5), with a further failure probability at most  $\eta/4$ . Use a fixed compact reconstruction box enlarged slightly to include all the supplied nearby gap values.

We check cost also outside the successful calibration event. The root rule always gives  $|\hat{\tau}| \leq h/2$ , whereas  $|\tau| \leq h/4$ . Consequently every actual endpoint-stage offset is

$$lh + \hat{\tau} - \tau \in [h/4, (m + 3/4)h].$$

The calibration offsets are in  $[3h/4, (m + 5/4)h]$ . Decrease  $h_m$  so both ranges lie in the common threshold collar. In both stages every success probability is bounded above and below by fixed multiples of  $h^2 e^{-j\gamma}$ , uniformly over all calibration outcomes. In particular a failed polynomial fit cannot produce a nonpositive offset or an infinite waiting-time expectation. The expectation identity in (A.20.7), conditioning on the pilot in the second stage, gives

$$\mathbb{E}T \leq C_m e^{j\gamma} h^{-2} (r + K).$$

For the high-probability cost, the equivalence with binomial counts and  $\Pr\{\text{Bin}(n, p) < np/2\} \leq e^{-np/8}$  show that waiting for  $r$  or  $2K$  successes costs at most  $Cp^{-1}(r + K + \log(C_m/\eta))$ , simultaneously over the fixed number of windows. These bounds also hold conditionally on the pilot, with the same constants. A union bound, allocating the remaining  $\eta/2$  to the waiting-time tails, proves the simultaneous assertion in (A.20.6). Substituting  $r, K$  proves its order, and  $h = c_m \varepsilon^{1/m}$  gives the final accuracy statement.  $\square$

The calibration gain comes from estimating the zero of the square root of an onset probability, not from cancelling the nonzero harmonic coefficient in (A.20.1). The unknown area and the rare-event prefactor cancel in that zero. The factor  $e^{j\gamma}$  nevertheless persists in the number of preparations needed to locate it. No minimax claim, independence of successive impacts, or uncharged fine timing measurement is used.

#### APPENDIX A.21. A SHARP STABILITY SCALE AT THE CIRCULAR CONFIGURATION

The exact infinite-sequence identification in Theorem A.10.1 does not resolve conditioning at a repeated node. We now give both a sequence-level estimate for the coalescing path and a matching positive estimate relative to the circular configuration. The latter is explicitly a reference-point estimate; it is not asserted for every pair of noncircular triples.

**A.21.1. An observable dispersion statistic.** Use the equal-gap family (A.10.1), with  $R$  fixed and  $g = 1 - 2R$ . Define, from its first three positive amplitudes,

$$m = \frac{2C_2}{C_1}, \quad J = \frac{C_3}{C_1} - f(m), \quad f(x) = \frac{x^2}{4 - x^2}. \quad (\text{A.21.1})$$

For the circular member let  $\kappa_* = 1/R$  and  $x_* = (1 + g\kappa_*)^{-1}$ .

**Theorem A.21.1** (Three-amplitude circular-reference stability). *On a sufficiently small compact positive-curvature neighborhood in the fixed- $R$  family,  $J \geq 0$  and there are constants  $c, C > 0$  such that*

$$c(|m - x_*| + \sqrt{J}) \leq \max_{0 \leq r \leq 2} |\kappa_r - \kappa_*| \leq C(|m - x_*| + \sqrt{J}). \quad (\text{A.21.2})$$

*In particular, if  $C_j^*$  are the circular amplitudes and  $\delta = \max_{1 \leq j \leq 3} |C_j - C_j^*| \leq 1$ , then*

$$\text{dist}_{\text{match}}(\kappa, \kappa_*) \leq C\sqrt{\delta}. \quad (\text{A.21.3})$$

*The statistic  $J$  vanishes exactly when the three contact curvatures coincide. It measures their dispersion without a supplied area normalizer or channel labels.*

*Proof.* Set  $x_r = \text{sech } \gamma_r = (1 + g\kappa_r)^{-1}$  and

$$p_r = \frac{\text{csch } \gamma_r}{\sum_s \text{csch } \gamma_s}.$$

The  $x_r$  lie in a fixed compact interval inside  $(0, 1)$ , while  $p_r \geq p_* > 0$  and  $\sum p_r = 1$ . The identities  $\sinh(2\gamma) = 2 \sinh \gamma \cosh \gamma$  and  $\sinh(3\gamma) = \sinh \gamma (4 \cosh^2 \gamma - 1)$  give

$$m = \sum_r p_r x_r, \quad J = \sum_r p_r f(x_r) - f(m). \quad (\text{A.21.4})$$

On the common interval,

$$f''(x) = \frac{8(4 + 3x^2)}{(4 - x^2)^3}$$

has positive lower and finite upper bounds. Taylor's integral formula at  $m$ , followed by  $\sum p_r (x_r - m) = 0$ , therefore gives

$$\frac{c_f}{2} \sum_r p_r (x_r - m)^2 \leq J \leq \frac{C_f}{2} \sum_r p_r (x_r - m)^2. \quad (\text{A.21.5})$$

This proves nonnegativity and the equality criterion. Since all weights are bounded below, the square root in (A.21.5) is comparable to  $\max_r |x_r - m|$ . Also

$$|m - x_*| \leq \max_r |x_r - x_*|, \quad \max_r |x_r - x_*| \leq |m - x_*| + \max_r |x_r - m|.$$

The reverse comparison follows from  $\max_r |x_r - m| \leq 2 \max_r |x_r - x_*|$ . Finally the scalar change  $\kappa \mapsto (1 + g\kappa)^{-1}$  is bi-Lipschitz on the compact interval. This proves (A.21.2).

All amplitude denominators  $C_1$  are bounded away from zero. At the circle  $m = x_*$  and  $J = 0$ . The two ratios in (A.21.1) are Lipschitz in the first three amplitudes on this set, and  $f$  has bounded derivative. Hence  $|m - x_*| \leq C\delta$  and  $0 \leq J \leq C\delta$ . Substitution gives (A.21.3). Because the reference triple has all entries equal, its matching distance is exactly the maximum coordinate distance, independent of ordering.  $\square$

The theorem quantifies distance to a specified circular table. It does not recover an arbitrary unordered triple from only three numbers, nor does it claim a one-half-Hölder inverse between every two nearby noncircular triples. The exact infinite-sequence theorem and its multiplicity recovery remain separate statements.

**A.21.2. The infinite absolute-amplitude topology.** For a fixed  $a \geq 0$  write

$$\|D\|_a = \sup_{j \geq 1} e^{aj} |D_j|. \quad (\text{A.21.6})$$

Restrict the coalescing path of Proposition A.10.3 to a small closed interval on which all exponents satisfy  $\gamma_r(s) \geq \gamma_- > a$ .

**Theorem A.21.2** (Sharp circular-reference exponent for the full sequence). *Along  $\alpha = \zeta = 0$ ,  $\beta = s$  one has*

$$\|C(s) - C(0)\|_a \leq B_a s^2, \quad \text{dist}_{\text{match}}(\kappa(s), \kappa(0)) \asymp |s|. \quad (\text{A.21.7})$$

*For every table in a sufficiently small fixed- $R$  neighborhood of the circle, the positive estimate*

$$\text{dist}_{\text{match}}(\kappa, \kappa_*) \leq C_a \|C - C^*\|_a^{1/2} \quad (\text{A.21.8})$$

*holds when the right side is small. No exponent larger than one half can replace it throughout that neighborhood, even when the entire sequence is observed. This sharpness concerns the absolute-error norm (A.21.6), with a strict exponential margin.*

*Proof.* On the indicated path the three radii are  $R + 36s$ ,  $R - 18s$ ,  $R - 18s$ , and  $A(s) = A(0) + 45\pi s^2/4$ . With  $\Phi_j(r) = \text{csch}(j \text{arcosh}(1 + g/r))$  the amplitude is

$$C_j(s) = \frac{\Phi_j(R + 36s) + 2\Phi_j(R - 18s)}{A(s)}.$$

Both numerator and denominator have zero first derivative at zero. On the compact interval the derivatives through order two of the radius-to-exponent map and of  $A^{-1}$  are bounded. Differentiate

$$\operatorname{csch}(j\gamma) = \frac{2e^{-j\gamma}}{1 - e^{-2j\gamma}}$$

twice. The denominator is bounded below by  $1 - e^{-2\gamma_-}$ ; each differentiation contributes at most a factor  $C(1 + j)$ . Consequently

$$\sup_{|s| \leq s_*} |C_j''(s)| \leq B(1 + j)^2 e^{-j\gamma_-}. \quad (\text{A.21.9})$$

Taylor's integral remainder and  $C_j'(0) = 0$  imply the first part of (A.21.7), since  $\sup_j (1 + j)^2 e^{-(\gamma_- - a)j} < \infty$ . The curvature displacement is comparable to  $|s|$  because reciprocal radius has a derivative bounded above and away from zero there.

For a general nearby table, the first three absolute amplitude errors are bounded by (A.21.6). Thus Theorem A.21.1 gives (A.21.8); no unstable truncation of a Möbius inversion is used. If an exponent  $\theta > 1/2$  were valid in its place, the specified path would give  $b|s| \leq LB_a^\theta |s|^{2\theta}$  for arbitrarily small nonzero  $s$ , which is impossible. This proves sharpness. In particular, (A.21.8) also supplies a quadratic lower bound for the sequence discrepancy along that path; the order in (A.21.7) is not merely a one-sided diagnosis.  $\square$

Exact identification and stability in an absolute-amplitude norm are different questions. The two theorems give a positive sharp scale for one specified singular reference, while the channel-resolved endpoint observation remains Lipschitz there because it retains different data. None of these assertions treats relative error at arbitrarily small amplitudes as equivalent to (A.21.6).

## APPENDIX A.22. MEASURABLE RECONSTRUCTION ON COMPACT PARAMETER SETS

All compact discrepancy minimizations in the sequel use the same measurable convention. This removes a possible ambiguity at repeated roots or at data having several best fits; uniqueness of a minimizing parameter is unnecessary.

**Lemma A.22.1** (A specified Borel minimizing selection). *Let  $K$  be a nonempty compact metric space, let  $\mathcal{Y}$  be a standard Borel data space, and suppose  $G : \mathcal{Y} \times K \rightarrow \mathbb{R}$  is Borel measurable and continuous in its second argument. There is a Borel map  $\hat{\theta} : \mathcal{Y} \rightarrow K$  such that*

$$G(y, \hat{\theta}(y)) = \min_{\theta \in K} G(y, \theta).$$

*A single choice of finite closed-ball covers of  $K$  specifies such a map for every data value.*

*Proof.* Fix, for each integer  $r \geq 1$ , a finite cover of  $K$  by closed balls  $B_{r,1}, \dots, B_{r,N_r}$  of radius  $2^{-r}$ , with centers in  $K$ , and fix their indicated order. The minimum  $v(y) = \min_K G(y, \cdot)$  is Borel: it equals the infimum of  $G(y, \cdot)$  over any fixed countable dense subset of  $K$ . For any nonempty fixed compact subset  $F \subset K$ , the same argument shows that  $v_F(y) = \min_F G(y, \cdot)$  is Borel. For  $F = \emptyset$ , set  $v_F = +\infty$ .

At the first stage choose the least index  $i_1$  for which  $v_{K \cap B_{1,i_1}}(y) = v(y)$ . At stage  $r$ , after the earlier choices, choose the least  $i_r$  for which

$$v_{K \cap B_{1,i_1} \cap \dots \cap B_{r,i_r}}(y) = v(y).$$

Such a choice always exists because the current nonempty compact set of minimizers is covered by the balls at the next stage. For each fixed finite prefix of indices the compact set in this expression is fixed, so its equality event is Borel. Induction and the finite number of indices at each stage show that every selected index is Borel.

Intersect the minimizer set with the selected balls. These are nested nonempty compact sets with diameter at most  $2^{1-r}$  at stage  $r$ . Their intersection consists of a single point  $\hat{\theta}(y)$ . The centers of the selected balls converge to it, since their distances from this point are at most  $2^{-r}$ . The centers form Borel maps into  $K$ ; their pointwise limit is Borel and is a minimizer by continuity.  $\square$

Apply the lemma with  $G$  equal to the maximum discrepancy between the measured finite vector and the continuous physical observation map. The physical compact set, its covers, and their order are fixed before observations. Every resulting estimator in Theorems A.11.1, A.12.1, and A.23.2 is then a Borel statistic. The proof is an existence and specification result, not an assertion of computational efficiency. It applies equally to compact constrained variance fits.

The polynomial root-or-zero convention in Lemma A.20.2 is Borel as well. One direct verification is to minimize the squared polynomial on a fixed compact subinterval; the minimum is a continuous function of its coefficients. Existence of a zero on any closed rational interval is therefore a closed condition. Existence of two distinct zeros in the prescribed interval is the countable union of conditions asserting zeros in two disjoint closed intervals with rational endpoints, also allowing the two fixed end points. Existence of exactly one zero is consequently Borel. On that set its location is recovered measurably from these interval tests; on its complement the declared value is zero. This includes the zero polynomial, for which uniqueness fails. Thus the pilot and the final compact fit both define statistics on all data outcomes.

#### APPENDIX A.23. LOCALLY CALIBRATED RECONSTRUCTION FROM UNLABELLED COUNTS

This section uses counts only, with neither channel labels nor impact positions. There are three unoriented equal-gap ground channels. For an actual-curvature conclusion we assume equal facing curvatures within each channel. The asymmetry allowed in the general forward theorem is not silently removed there: without the additional equality, the same scalar reconstruction concerns (A.12.3). All preparations are independent and have normalized full phase volume. Collision counts and programmed physical times are recorded without instrument noise. Statistical noise arises from the random preparation.

We consider either a sufficiently small compact subfamily of the physical family in Theorem A.12.1, with  $J = 3$ , or a compact geometric family whose leading data lie in the four-amplitude neighborhood of Theorem A.11.1, with  $J = 4$ . In the latter case the normalizer is treated as an independent nuisance parameter in inversion; the observed probabilities still come from a physical configuration. For  $1 \leq j \leq J$ , the ground-onset law reads

$$p_j(d) := \Pr\{N_{jg+d} \geq j+1\} = d^2\{C_j + dB_j(d)\}, \quad 0 \leq d \leq d_*. \quad (\text{A.23.1})$$

The functions in braces are positive, bounded away from zero, and have bounded derivatives at each fixed order, uniformly on the chosen family. At these offsets the event is equivalently  $N_{jg+d} = j+1$ .

**Lemma A.23.1** (Amplitude error from finitely many Bernoulli experiments). *Suppose  $g$  is known. Fix  $m \geq 1$ , and at each of the  $Jm$  times  $jg + lh$ ,  $1 \leq j \leq J$ ,  $1 \leq l \leq m$ , perform  $n$  independent preparations. Write  $\hat{p}_{j,l}$  for the observed success fraction and put*

$$\omega_l = (-1)^{l-1} \binom{m}{l}, \quad \hat{C}_j = \sum_{l=1}^m \omega_l \frac{\hat{p}_{j,l}}{(lh)^2}.$$

For small  $h$  and  $0 < \eta < 1/2$ , with probability at least  $1 - \eta$ ,

$$\max_{j \leq J} |\hat{C}_j - C_j| \leq C_m \left( h^m + \sqrt{\frac{\log(C_m/\eta)}{nh^2}} + \frac{\log(C_m/\eta)}{nh^2} \right). \quad (\text{A.23.2})$$

*Proof.* The weights satisfy  $\sum_l \omega_l = 1$  and  $\sum_l \omega_l l^k = 0$  for  $1 \leq k < m$ . Taylor's formula for  $p_j(d)/d^2$  in (A.23.1), with its uniformly bounded  $m$ th derivative, gives a population bias at most  $C_m h^m$ . For a Bernoulli variable  $X$  of mean  $p$ , exponential Markov bounds give

$$\Pr\{|\hat{p} - p| > \sqrt{2px/n} + 2x/(3n)\} \leq 2e^{-x}.$$

For completeness, for  $0 < t < 3$  the power series of the exponential and  $\mathbb{E}|X - p|^k \leq \mathbb{E}(X - p)^2 \leq p$  for  $k \geq 2$  give  $\log \mathbb{E}e^{t(X-p)} \leq pt^2/[2(1-t/3)]$ ; the same bound holds for  $-t$ . Optimization for the sum of  $n$  independent variables gives the stated two-sided bound with these harmless enlarged constants. Here  $p_j(lh) \leq C_m h^2$ . Dividing by  $(lh)^2$ , summing the fixed weights, and taking  $x = \log(2Jm/\eta)$  proves the stochastic terms by a union bound. Independence is used between preparations, not between collisions of one trajectory.  $\square$

We now include fine timing instead of assuming it. A coarse bracket  $|g - g_0| \leq h/4$  is supplied. Its acquisition is not charged below. A pilot at order  $j = 1$  suffices for all the finitely many subsequent orders. The safety margin in the second stage will be enlarged to cover every possible pilot outcome.

**Theorem A.23.2** (Count-only acquisition with charged fine calibration). *Fix  $m \geq 1$  and one of the preceding compact geometric families. There are  $h_m > 0$  and  $C_m < \infty$  such that, for  $0 < h \leq h_m$ ,  $0 < \eta < 1/2$  and a supplied bracket  $|g - g_0| \leq h/4$ , an experiment using only unlabelled counts produces estimates  $\hat{g}, \hat{A}, \hat{\kappa}$  with*

$$|\hat{g} - g| \leq C_m h^{m+1}, \quad |\hat{A} - A| \leq C_m h^m, \quad \text{dist}_{\text{match}}(\hat{\kappa}, \kappa) \leq C_m h^{m/3} \quad (\text{A.23.3})$$

*with probability at least  $1 - \eta$ . Its total number  $T$  of raw preparations, including the pilot, satisfies*

$$\mathbb{E}T \leq C_m h^{-(2m+2)} \log(C_m/\eta), \quad T \leq C_m h^{-(2m+2)} \log(C_m/\eta) \quad (\text{A.23.4})$$

*with the second estimate holding jointly with (A.23.3) with probability at least  $1 - \eta$ . Consequently a sufficient raw-preparation order for matching curvature accuracy  $\varepsilon$  is*

$$C_m \varepsilon^{-(6+6/m)} \log(C_m/\eta). \quad (\text{A.23.5})$$

*No channel labels, selected positions, or supplied fine-gap estimate enter this experiment. The result is local and sufficient, not a minimax theorem or a global procedure for locating the initial bracket.*

*Proof.* Put  $\tau = g - g_0$ . At each of the pilot times  $g_0 + lh$ ,  $1 \leq l \leq m+1$ , count successes of  $N_t \geq 2$  until  $r$  successes have occurred, using fresh preparations at every window. Let  $T_l$  be the waiting time and let  $\hat{p}_l = r/T_l$ . Formula (A.23.1) at  $j = 1$  has the form

$$p_1(d) = C_1 d^2 F(d), \quad F(0) = 1,$$

where  $F$  is positive and has the required bounded right derivatives. The unknown  $C_1 > 0$  replaces the unknown factor  $K_j$  in Lemma A.20.2; its proof uses no channel label. Interpolate the square roots of the pilot fractions at  $z = lh$  and use precisely that lemma's root-or-zero rule. Thus  $\hat{g} = g_0 + \hat{\tau}$  with  $|\hat{\tau}| \leq h/2$  on every outcome. Choose

$$r = \left\lceil C_m h^{-2m} \log(C_m/\eta) \right\rceil.$$

The negative-binomial bound (A.20.7) gives relative errors at most  $c_m h^m$ , simultaneously, outside an event of probability  $\eta/4$ . The root estimate then gives  $|\hat{g} - g| \leq C_m h^{m+1}$ . Every pilot offset lies in  $[3h/4, (m+5/4)h]$ , so its success probability is bounded below by  $c_m h^2$ . The expected pilot cost is at most  $C_m h^{-2}(m+1)r$ . The same binomial waiting-time tail used in Theorem A.20.3 bounds the pilot cost by  $C_m h^{-(2m+2)} \log(C_m/\eta)$  except on a further event of probability at most  $\eta/4$ .

For the second stage fix  $L = J + 1$  and perform  $n$  fresh independent preparations at each time  $j\hat{g} + Llh$ , with  $j \leq J$ ,  $l \leq m$ . This stage observes only  $\mathbf{1}_{\{N_l \geq j+1\}}$  and uses the estimator

$$\hat{C}_j = \sum_{l=1}^m \omega_l \frac{\hat{p}_{j,l}}{(Llh)^2}, \quad n = \left\lceil C_m h^{-(2m+2)} \log(C_m/\eta) \right\rceil.$$

The total second-stage cost  $Jmn$  is deterministic. Even if the pilot fails,  $|\hat{g} - g| \leq 3h/4$ . Hence every actual second-stage offset is between  $(L - 3J/4)h$  and  $(Lm + 3J/4)h$ , and in particular exceeds  $h$ . Decreasing  $h_m$  places this entire range in the common collar. Thus unsuccessful pilots cannot produce a nonpositive offset or an undefined normalization.

Condition on the pilot. The fresh fractions have Bernoulli means  $p_j(Llh + j(\hat{g} - g))$ , each bounded by  $C_m h^2$  uniformly over all pilot outcomes. The concentration proof of Lemma A.23.1, with nodes  $Ll$ , therefore still applies. On the successful pilot event, differentiation of (A.23.1) gives  $|p'_j(d)| \leq C_m h$  throughout the relevant offset intervals. The mean-value theorem consequently gives

$$\left| \frac{p_j(Llh + j(\hat{g} - g)) - p_j(Llh)}{(Llh)^2} \right| \leq C_m \frac{|\hat{g} - g|}{h} \leq C_m h^m.$$

This explicitly accounts for the uncanceled denominator/timing error. Taylor extrapolation at the unshifted nodes  $Llh$  has bias  $C_m h^m$ , and the chosen  $n$  makes both concentration terms at most  $C_m h^m$ . A conditional union bound therefore gives  $\max_{j \leq J} |\hat{C}_j - C_j| \leq C_m h^m$ , with an additional failure probability at most  $\eta/2$ .

In the three-amplitude physical case, fit the observations  $(\hat{g}, \hat{C}_1, \hat{C}_2, \hat{C}_3)$  over a sufficiently small compact admissible family containing the true member and apply Theorem A.12.1. Its true spacing discrepancy is  $O_m(h^{m+1})$  and its amplitude discrepancy is  $O_m(h^m)$ , so (A.23.3) follows, retaining the pilot as the stated gap estimate. In the four-amplitude case, augment the coefficient map in Theorem A.11.1 by the coordinate  $g$ . Its derivative has the nonsingular four-by-four block already proved there, and the identity in the  $g$  coordinate. Thus its coefficient inverse is jointly locally Lipschitz in  $g$  and the four amplitudes. The same compact minimum-discrepancy argument and the cubic root bound give the asserted area and curvature estimates. On the failure set arbitrary compact fallback estimates may be used; no success assertion depends on them. Adding the pilot and deterministic second-stage costs proves (A.23.4). Finally choose  $h = c_m \varepsilon^{3/m}$ , with  $c_m$  sufficiently small, to obtain (A.23.5).  $\square$

The exponent in (A.23.5) has two different sources: amplitude acquisition costs  $h^{-(2m+2)}$ , whereas coalescing unlabelled curvatures have matching error  $h^{m/3}$ . The selected-position experiment of Theorem A.20.3 instead has a Lipschitz curvature inverse and sufficient cost  $e^{j\gamma} \varepsilon^{-(2+2/m)}$ . These are different observations and different experiments. The sharp deterministic exponent  $1/3$  does not by itself prove optimality of a statistical sampling rate. All constants here may depend on fixed  $m$ , on  $J \leq 4$ , and on the compact family; the statement makes no uniform extrapolation-order or large- $j$  count-only cost claim.

#### APPENDIX A.24. PHYSICAL INDISTINGUISHABILITY IN LOCAL COUNT DESIGNS

The pairwise cubic inverse has a statistical counterpart in the actual finite-offset count laws. We prove it within the specified Bernoulli query experiment, including adaptive stopping. This result concerns fixed bounded flight numbers; it is logically distinct from the long-bridge experiment of Section A.9.

Fix  $R = 1/4$  in the physical support family (A.10.1), put  $\alpha = 0$ , and let  $\theta = (\beta, \zeta)$  vary near zero. The common gap is  $g = 1/2$ . A query at  $(j, d)$  consists of one fresh full-phase preparation



and the binary report  $\mathbf{1}_{\{N_{jg+d} \geq j+1\}}$ . We do not give the experiment the lower values of  $N_{jg+d}$  or a collision position. Write its probability as  $p_{j,d}(\theta)$ .

**Lemma A.24.1** (Cubic physical indistinguishability). *For each fixed  $J < \infty$  there are  $d_*, s_*, c, C > 0$  such that, for  $1 \leq j \leq J$ ,  $0 < d \leq d_*$ , and  $|s| \leq s_*$ ,*

$$|p_{j,d}(s, 0) - p_{j,d}(-s, 0)| \leq Cd^2|s|^3, \quad cd^2 \leq p_{j,d}(\pm s, 0) \leq Cd^2 < \frac{1}{2}. \quad (\text{A.24.1})$$

*The matching distance between their unordered curvature triples is bounded above and below by positive multiples of  $|s|$  for  $s \neq 0$ .*

*Proof.* A rotation of the triangular lattice and the entire table by  $\pi/3$  preserves the preparation measure and the unlabelled count. On  $(\beta, \zeta)$  it acts as a rotation through  $2\pi/3$ , up to reversal of its sign. This follows directly by replacing  $\vartheta$  by  $\vartheta - \pi/3$  in  $(1 - \cos 6\vartheta)(\beta \cos 2\vartheta + \zeta \sin 2\vartheta)$ . No nonzero linear functional is invariant under that rotation. The gradient of  $p_{j,d}(\theta)$  at zero must therefore vanish.

Theorem 13.1 and the equal-gap property give a smooth extension of  $d^{-2}p_{j,d}(\theta)$  to the closed collar, with bounded third parameter derivatives uniformly for  $j \leq J$ . Taylor expansion along  $(s, 0)$  and  $(-s, 0)$  cancels the constant and quadratic terms; the linear term is zero. This proves the first estimate for the full physical probabilities, not merely their leading amplitudes. Their positive continuous leading coefficients and a smaller common collar prove the second estimates. The two radii multisets are

$$(R + 36s, R - 18s, R - 18s), \quad (R - 36s, R + 18s, R + 18s).$$

For  $s > 0$  sufficiently small, ordered matching of their reciprocals has distance  $36s/(R^2 - 324s^2)$ , as in Theorem A.11.2. Replacing  $s$  by  $|s|$  proves the last assertion.  $\square$

For a stopped adaptive design, a policy chooses the next query and whether to stop using only the previous binary reports and independent randomization. Its rules are the same under the two possible tables. Suppose all its queries have  $j \leq J$  and  $0 \leq d \leq d_*$  on every history. Define its weighted exposure by

$$\mathcal{B} = \max_{\theta \in \{(s,0), (-s,0)\}} \mathbb{E}_\theta \sum_{i=1}^T d_i^2.$$

Queries at  $d = 0$  are deterministic failures and contribute zero. Bounds uniform over a fixed neighborhood mean the same exposure bound holds in particular for each pair used below.

**Theorem A.24.2** (A lower bound for count queries). *There is  $c > 0$  such that any predetermined or stopped adaptive design as above with weighted exposure at most  $B$  uniformly on the local physical family has worst-case expected matching curvature error at least  $c(1 + B)^{-1/6}$ . In particular, if every queried offset satisfies  $d_i \leq Lh$  on every history and  $\sup_\theta \mathbb{E}_\theta T \leq N$ , the lower bound is*

$$c(1 + L^2 h^2 N)^{-1/6}. \quad (\text{A.24.2})$$

*For each fixed confidence level  $1 - \eta_0 > 1/2$ , accuracy  $\varepsilon$  throughout this local family consequently requires*

$$N \geq c_{L,\eta_0} h^{-2} \varepsilon^{-6} \quad (\text{A.24.3})$$

*for sufficiently small  $\varepsilon$ . The conclusion remains valid when the exact common gap is supplied to the procedure.*

*Proof.* For  $0 < p, q < 1$ , the inequality  $\log x \leq x - 1$  gives

$$D_{\text{KL}}(\text{Ber}(p) \parallel \text{Ber}(q)) \leq \frac{(p - q)^2}{q(1 - q)}.$$

Thus Lemma A.24.1 bounds the divergence of a single query under the alternatives  $s$  and  $-s$  by  $Cd^2s^6$ . For a predetermined schedule, the log likelihood is a sum and its expectation is at most  $Cs^6 \sum_i d_i^2$ .

For an adaptive policy, condition on each common observed history and on its random seed. The next query, or stop decision, is then the same in the two experiments. The conditional divergence is again at most  $Cd_i^2s^6$ ; a stop symbol contributes zero. Applying the chain rule to the first  $M$  steps, with the stopped process extended by stop symbols, gives

$$D_{\text{KL}}(\mathbf{P}_s^{[M]} \parallel \mathbf{P}_{-s}^{[M]}) \leq Cs^6 \mathbb{E}_s \sum_{i \leq \min(T, M)} d_i^2.$$

The bound passes to the complete stopped transcript. Indeed relative entropy is the supremum of its finite measurable-partition values; conditioning these partitions on increasing finite histories yields that supremum, while monotone convergence bounds the right side by  $Cs^6B$ . Equivalently, this is the chain rule for a stopped sequence obtained by truncation. A finite expected number of queries ensures that the stop time is finite almost surely; if only weighted exposure is bounded, the same finite-history argument applies to any measurable estimator based on the infinite transcript. Hence

$$D_{\text{KL}}(\mathbf{P}_s \parallel \mathbf{P}_{-s}) \leq Cs^6B. \quad (\text{A.24.4})$$

We recall the testing step to fix constants and topology. Coarsening to any event cannot increase relative entropy, by the log-sum inequality. For Bernoulli laws the second derivative in the first probability of relative entropy is  $1/[x(1-x)] \geq 4$ , so Taylor's formula at equality yields  $D_{\text{KL}}(\text{Ber}(x) \parallel \text{Ber}(y)) \geq 2(x-y)^2$ . Taking the supremum over events therefore gives  $\|P-Q\|_{\text{TV}} \leq \sqrt{D_{\text{KL}}(P \parallel Q)/2}$ . Any test has sum of its two error probabilities at least  $1 - \|P-Q\|_{\text{TV}}$ , directly from the definition of total variation.

The curvature separation  $\Delta_s$  is at least  $c_0|s|$ . Given an estimator, classify it by its nearer alternative, with fixed tie breaking. On a wrong classification its matching error is at least  $\Delta_s/2$ , by the triangle inequality. Its maximum expected error is consequently at least  $\Delta_s(1 - \|\mathbf{P}_s - \mathbf{P}_{-s}\|_{\text{TV}})/4$ . Choose  $s = c_1(1+B)^{-1/6}$ , with  $c_1 > 0$  small enough to stay inside the family and make the variation distance at most  $1/2$  by (A.24.4). This proves the expected-risk bound. The bound  $B \leq L^2h^2N$  proves (A.24.2).

For confidence accuracy choose  $s = c_2\varepsilon$  so that  $\Delta_s > 2\varepsilon$ . Success with probability at least  $1 - \eta_0$  under both alternatives would give a test whose error sum is at most  $2\eta_0$ . Hence its variation distance is at least  $1 - 2\eta_0 > 0$ . The entropy bound forces  $B \geq c_{\eta_0}\varepsilon^{-6}$ , proving (A.24.3). The two alternatives have the same known  $g = 1/2$ , so supplying the gap does not change this test.  $\square$

For the fixed-order design of Theorem A.23.2, all actual offsets are at most  $L_mh$ , including all pilot outcomes, and  $h \asymp \varepsilon^{3/m}$ . Thus the power  $\varepsilon^{-(6+6/m)}$  in its preparation bound has a matching obstruction within this shrinking-offset, bounded-flight-number binary experiment, apart from confidence logarithms. This does not identify a minimax exponent for unrestricted collision observations: nonshrinking offsets, an unbounded choice of flight numbers, lower count values, and position data are different experiments. All constants here are at fixed  $m, J, L_m$ ; sending  $m$  to infinity is not justified.

#### APPENDIX A.25. COUNT-ONLY RECONSTRUCTION FROM A FIXED LOCAL BRACKET

The preceding acquisition theorem presupposes a bracket whose width shrinks with the desired accuracy. We now acquire that bracket using counts, starting from a fixed local interval. A deterministic cap on all later waiting times prevents an unsuccessful search from producing an

infinite preparation cost. The original locally supplied-bracket theorem remains useful, with its sharper description of each stage.

Fix a sufficiently small compact physical neighborhood in Theorem A.12.1. The constants of the forward law on this neighborhood, including a positive lower bound, are regarded as design information. This is local reconstruction in a specified family, not a procedure requiring no prior geometric information. For  $j = 1$  there are known constants  $c_* > 0$ ,  $d_* > 0$  such that

$$\Pr(N_{g+d} \geq 2) \geq c_* d^2 \quad (0 < d \leq d_*), \quad \Pr(N_t \geq 2) = 0 \quad (t \leq g). \quad (\text{A.25.1})$$

These follow from Theorem 13.1 and compactness.

**Lemma A.25.1** (Count acquisition of a coarse bracket). *Suppose  $g \in [L_0, U_0]$ , with fixed width  $0 < W = U_0 - L_0 \leq d_*$ . For  $0 < h < W$  and  $0 < \delta < 1/2$ , a binary count procedure returns  $g_0$  satisfying  $|g - g_0| \leq h/4$  with probability at least  $1 - \delta$ . On every history it uses at most*

$$Ch^{-2} \log \left( \frac{2 + \log(W/h)}{\delta} \right) \quad (\text{A.25.2})$$

*independent preparations. Its queries use only the indicator of two or more collisions at programmed times inside  $[L_0, U_0]$ .*

*Proof.* Maintain an interval  $[L, U]$ , initially  $[L_0, U_0]$ , and stop when its width  $w$  is at most  $h/2$ . At a nonterminal stage query  $t = (L + U)/2$  on

$$n(w) = \left\lceil \frac{16}{c_* w^2} \log \frac{M + 1}{\delta} \right\rceil$$

fresh preparations, where  $M = \lceil \log(2W/h) / \log(4/3) \rceil$ . If at least one query succeeds, replace  $U$  by  $t$ ; otherwise replace  $L$  by  $t - w/4$ . In either case the interval is nested in the previous one and its width is either  $w/2$  or  $3w/4$ . There are at most  $M$  stages, regardless of the observations.

Conditional on an interval containing  $g$ , a success certifies  $g < t$  by (A.25.1). A failure branch can discard  $g$  only if  $g < t - w/4$ . In this case  $t - g > w/4$  and  $t - g \leq W \leq d_*$ , so the probability of seeing no success is at most  $\exp(-n(w)c_* w^2/16) \leq \delta/(M + 1)$ . A union bound over the first possible incorrect discard proves that all retained intervals contain  $g$  with probability at least  $1 - \delta$ . The midpoint of the final interval then has error at most  $h/4$ .

For cost, each nonterminal width exceeds  $h/2$  and contracts by at least a factor  $3/4$ . Summing reciprocal squared widths backwards from the last such stage gives  $\sum w^{-2} \leq Ch^{-2}$ . The ceiling contributions add at most  $M$ , also bounded by  $Ch^{-2}$  on the stated range, with a constant depending on fixed  $W$ . This proves (A.25.2) and its all-history form. There is no waiting for an event of unknown positive probability in this search.  $\square$

**Theorem A.25.2** (Fully charged local count reconstruction). *Fix  $m \geq 1$ , a compact physical neighborhood as above, and a fixed interval  $[L_0, U_0]$  containing its possible gaps and satisfying the hypotheses of Lemma A.25.1. There exist  $h_m > 0$  and  $C_m < \infty$  such that, for  $0 < h \leq h_m$  and  $0 < \eta < 1/2$ , a count-only procedure estimates the gap, area, and unordered curvature triple with*

$$|\hat{g} - g| \leq C_m h^{m+1}, \quad |\hat{A} - A| \leq C_m h^m, \quad \text{dist}_{\text{match}}(\hat{\kappa}, \kappa) \leq C_m h^{m/3} \quad (\text{A.25.3})$$

*with probability at least  $1 - \eta$ . Its total preparation count is bounded on every history by*

$$C_m h^{-(2m+2)} \log(C_m/\eta). \quad (\text{A.25.4})$$

*The procedure uses only binary collision-count reports with flight numbers 1, 2, 3, independent full-phase preparations, and noiseless elapsed-time recording. In particular curvature accuracy  $\varepsilon$  has sufficient total cost  $C_m \varepsilon^{-(6+6/m)} \log(C_m/\eta)$  from a bracket independent of  $\varepsilon$ . The supplied information includes the fixed local family and its uniform bounds, but no accuracy-dependent gap bracket, contact positions, or separately measured area.*

*Proof.* Use Lemma A.25.1 with failure budget  $\eta/4$  to acquire  $g_0$ . Conditional on its success, the hypothesis  $|g - g_0| \leq h/4$  of Theorem A.23.2 holds. All subsequent preparations are fresh. Run that theorem's pilot and amplitude stages with  $J = 3$ , and allocate failure probability  $\eta/2$  to their estimation bounds. Their constants are uniform over all admissible  $g_0$ , so conditioning on the entire first stage is legitimate. On these good events their estimates satisfy (A.25.3).

There is an additional issue on a failed first stage: a programmed pilot window could lie below onset. We therefore modify each pilot waiting time as follows. If its target  $r = \lceil C_m h^{-2m} \log(C_m/\eta) \rceil$  successes have not occurred within

$$B_m = \lceil C_m h^{-(2m+2)} \log(C_m/\eta) \rceil$$

preparations, stop the entire procedure and return a fixed admissible parameter as a fallback. On first-stage success, every pilot probability is at least  $c_m h^2$ , uniformly over its programmed nodes. Choose the constant in  $B_m$  so that  $B_m c_m h^2 \geq 2r$  and  $B_m c_m h^2 \geq 8 \log(4(m+1)/\eta)$ . The exponential Markov bound for a binomial variable gives

$$\Pr\{T_r > B_m\} = \Pr\{\text{Bin}(B_m, p) < r\} \leq e^{-B_m p/8} \leq \eta/[4(m+1)].$$

A union bound over the  $m+1$  pilots charges at most  $\eta/4$  for a cap on a good first stage. Before a cap the procedure agrees exactly with the original one, so this modification does not condition or bias the claimed estimation guarantee. The final amplitude stage has a predetermined number of preparations, whether or not its estimates are good. Its compact minimizer is selected by Lemma A.22.1.

Every history now has a deterministic preparation bound: the first stage has (A.25.2), each of the fixed number of pilots has cap  $B_m$ , and the amplitude stage has the deterministic bound in the proof of Theorem A.23.2. For fixed  $m \geq 1$ , the function  $h^{2m} \log(2 + \log(W/h))$  is bounded for  $0 < h \leq h_m < W$ . Thus the first-stage cost is absorbed into (A.25.4), uniformly in  $\eta < 1/2$ . The three allocated failure probabilities sum to at most  $\eta$ . Finally choose  $h = c_m \varepsilon^{3/m}$ , with fixed  $c_m$  small enough to make the curvature bound at most  $\varepsilon$ .  $\square$

The same initial search and cap construction applies to the independent-normalizer four-amplitude experiment of Theorem A.23.2 on any specified compact physical family satisfying its hypotheses; then  $J = 4$ . The search uses only its positive one-flight onset bound, and no area constraint is needed for that stage.

The first-stage windows are allowed to have offsets larger than  $h$ . Accordingly Theorem A.25.2 is an upper bound for a larger design class than the all-history shrinking-collar class in (A.24.3). We do not infer global minimax optimality for this enlarged class from that lower bound. The weighted-exposure inequality (A.24.4) remains applicable whenever its own query hypotheses hold. The selected-position acquisition theorem has different data and is unchanged by this count-only construction.

## APPENDIX A.26. MARKED ONSETS AND SOURCE DOMAINS

Let  $x_{e,0}, x_{e,1}$  be the normal outgoing states and  $f_\xi$  a smooth collision mark near them. Put  $a_{e,b} = f_\xi(x_{e,b})$ ,  $\bar{f}_e = (a_{e,0} + a_{e,1})/2$ , and  $w_{j,e} = \sum_{i=0}^j a_{e,i \bmod 2}$ .

**Theorem A.26.1** (Marked response and the exact axial source domain). *At the selected channel's own threshold,*

$$Z_{j,e}(\xi, q, d) = \frac{d^2 e^{q w_{j,e}(\xi)}}{2A(\xi) \sinh(j \gamma_e(\xi))} [1 + d \mathcal{R}_{j,e}^f(\xi, q, d)]. \quad (\text{A.26.1})$$

*Every fixed mixed derivative of the remainder is bounded uniformly in  $j$  on compact parameter and source sets. The complete marked event in the ground collar is the sum at its positive channel*

offsets. For fixed  $\xi$ , real  $q$ , and sufficiently small fixed  $d \geq 0$ , the series  $\sum_{j,e} Z_{j,e}(\xi, q, d)/d^2$  converges exactly when

$$\max_{e \in \mathcal{E}}(q\bar{f}_e - \gamma_e) < 0. \quad (\text{A.26.2})$$

Right limits are used at  $d = 0$ . Every fixed mixed derivative is summable on compact subsets of this domain. Finitely many real sources are allowed, with the scalar product in place of  $q\bar{f}_e$ .

This series ranges over different onset windows, not over all itineraries in one unrestricted long window. The source formula, interface jumps and record response are consequences of the same uniform construction, not separate claims of a general limit theory. The exact equal-gap inverse theorem is retained in Section A.10. It uses central symmetry to remove the two-endpoint ambiguity and equal gaps to retain all channels at every collision order. The recovered area is a symmetric function of the three recovered curvatures in that chosen family, not a fourth independent unknown. Section A.10.1 proves this identity and the loss of first-order sensitivity at coalescence.

#### APPENDIX A.27. RECORD TOPOLOGY AND DIFFERENTIATED MOVING TESTS

A record limit is not, by itself, a response theorem. We specify both the metric and the differentiable test class. This also fixes the meaning of the bounded Lipschitz tests at the end of Theorem A.33.1. The unaltered circular estimates in that theorem continue to apply.

In a fixed oriented channel put  $s = \sqrt{2d}$  and use the exact Morse coordinates  $w = sz$ , together with first residual time  $r = d\eta$ . The latent domain is the same  $\mathcal{U}$  as in (16.5). Its unweighted conditional density relative to  $\mathbf{m}$  is

$$f_{j,e,\xi,s}(z, \eta) = \frac{a_{j,e}(\xi, sz)}{\int_{\mathcal{U}} a_{j,e}(\xi, sz) d\mathbf{m}}. \quad (\text{A.27.1})$$

This is a density on common latent coordinates, not a total-variation comparison between physical measures on distinct embedded record surfaces. In general  $f = 1 + O(s)$ ; for the circular unweighted density its evenness gives the stronger  $O(s^2)$  already proved in Theorem A.33.1.

Let  $\mathcal{R}_{j,e}(\xi, z, \eta, s)$  be the array consisting of  $y_i/s$ ,  $(t_i - ig_e)/d$ , and both transverse velocity traces divided by  $s$ , for  $0 \leq i \leq j$ . The transverse frame is that of the closest chord; the axial values have zero transverse velocity. At  $s = 0$  use the continuous limits. Equip each such array with the maximum norm,

$$d_{\infty,j}(\mathcal{R}, \tilde{\mathcal{R}}) = \max_{0 \leq i \leq j} \max\{|Y_i - \tilde{Y}_i|, |T_i - \tilde{T}_i|, |V_i^- - \tilde{V}_i^-|, |V_i^+ - \tilde{V}_i^+|\}. \quad (\text{A.27.2})$$

For tests involving only positions and times omit the other entries. An unnormalized Euclidean norm of the entire growing array is not used. The limiting positions are the linear bridge (15.3) applied to  $H_{j,e}^{-1/2}z$ . Its limiting excess flight times are

$$Q_i = \kappa_{i \bmod 2} Y_i^2 + \kappa_{1-(i \bmod 2)} Y_{i+1}^2 + (Y_{i+1} - Y_i)^2/g_e, \quad \sum_{i < j} Q_i = |z|^2.$$

Thus the limiting time array is  $\eta + \sum_{h < i} Q_h$ .

**Proposition A.27.1** (Differentiated smooth record tests). *Fix any derivative order  $m$ . Let  $\varphi_{j,\xi}$  be smooth tests of these scaled arrays with bounded derivatives through order  $m + 1$ , in operator norm for the maximum-norm array space, and with the corresponding mixed parameter derivatives bounded independently of  $j$ . Then the conditional expectations  $I_j(\xi, s)$  are smooth in  $(\xi, s)$  at  $s = 0$  on the right, and*

$$\sup_{\xi, e, j} |\partial_{\xi}^{\alpha} (I_j(\xi, s) - I_j(\xi, 0))| \leq C_m s, \quad |\alpha| \leq m. \quad (\text{A.27.3})$$

A sufficient test class consists of a fixed finite number of selected record entries and smooth functions with uniformly bounded derivatives. For uniformly bounded Lipschitz tests without differentiability, the same statement without  $\partial_\xi^\alpha$  follows from (A.27.2). The circular unweighted convergence in Theorem A.33.1 retains its stronger  $O(d)$  rate for its positions and times.

*Proof.* The bridge, Morse inverse, chord directions and reflections are smooth in  $\xi, w$ . They vanish to the required order at  $w = 0$ : positions and transverse velocities vanish to first order, and every excess flight length to second order. Taylor's integral formula therefore extends the divisions by  $s$  and  $s^2/2 = d$  smoothly to  $s = 0$ . For times, the estimates are summed along the bridge. The sum of the endpoint-decaying bounds from Lemma 15.2 is independent of  $j$ . Hence  $\mathcal{R}_j$ , all its fixed derivatives, and  $s^{-1}(\mathcal{R}_j(s) - \mathcal{R}_j(0))$  have uniform maximum-norm bounds. The same claims hold for the conditional density (A.27.1); its denominator is bounded below.

Now write the expectation exactly on the fixed domain:

$$I_j(\xi, s) = \int_{\mathcal{U}} f_{j,e,\xi,s}(z, \eta) \varphi_{j,\xi}(\mathcal{R}_{j,e}(\xi, z, \eta, s)) \, d\mathbf{m}.$$

The chain rule in the maximum norm has no dimension factor by the hypothesis on the test derivatives. Dominated differentiation and the fundamental theorem of calculus in  $s$  prove (A.27.3). This differentiates the actual parametrization and density; it does not differentiate a total-variation inequality. For Lipschitz tests, instead bound the difference by the pathwise maximum-norm error plus the density error. In the circular positions–times specialization these errors are  $O(d)$  as already proved. The entire statement also holds under bounded smooth additive sources after replacing the density by its positive normalized factor  $e^{qV_{j,e}}$ , with the same general  $O(s)$  rate.  $\square$

For a physical sampling cut, write the exact excess time of impact  $k$  as  $T_{j,k}(\xi, z, s) = \sum_{i < k} (\ell_i - g_e)/d$ . It has all fixed derivative bounds uniformly in  $j, k$ . A sample at  $t^s = kg_e + d a(\xi)$  precedes that impact precisely when

$$\eta > \psi_{j,k}(\xi, z, s) := a(\xi) - T_{j,k}(\xi, z, s).$$

Here  $a$  is smooth with bounded derivatives; the physical sample moves with the geometry, and is not an extra random coordinate.

**Proposition A.27.2** (A differentiated transverse sampling interface). *Let a smooth test be supported in a fixed compact subset of  $|z| < 1$ , and suppose that on this support, for all parameters under consideration,*

$$\nu \leq \psi_{j,k}(\xi, z, s) \leq 1 - |z|^2 - \nu \tag{A.27.4}$$

for one fixed  $\nu > 0$ . For a smooth test as in Proposition A.27.1, its conditional expectation restricted to the before-impact event is smooth in  $\xi, s$ , with uniform fixed-order bounds. If  $F_j(\xi, z, \eta, s)$  denotes the conditional density times the test, relative to  $dz \, d\eta$ , then its first derivative in any scalar parameter direction is exactly

$$\partial_\xi I_j = \int \left\{ \int_{\psi_{j,k}}^{1-|z|^2} \partial_\xi F_j \, d\eta - F_j(\xi, z, \psi_{j,k}, s) \partial_\xi \psi_{j,k} \right\} dz. \tag{A.27.5}$$

Repeated differentiation gives every fixed higher derivative and the analogue of (A.27.3). The boundary term is part of the formula.

*Proof.* The derivative of the cut function with respect to  $\eta$  is one. Condition (A.27.4) separates it uniformly from the two other boundaries of the latent domain on the support of the test. Its integral is therefore exactly  $\int \int_{\psi_{j,k}}^{1-|z|^2} F_j \, d\eta \, dz$ . Leibniz' rule gives (A.27.5), with no other boundary piece. Smooth bounded density and test derivatives and the uniform bounds for  $T_{j,k}$

justify repeated differentiation under the outer integral. The fundamental theorem in  $s$  gives its differentiated limit estimate. This proves response of the selected moving physical interface, rather than response inferred from its small mismatch strip.  $\square$

The margin is nonempty uniformly even for interior impacts. Take  $a = 1/2$  and support the test on  $|z| < 1/2$ . Every flight excess is nonnegative, and the exact Morse identity gives  $0 \leq T_{j,k} \leq |z|^2$  for every  $k$ . Both margins are therefore at least  $1/4$ . A finite collection of cuts with a fixed strict order and the same margins is handled by the same interval integration. At a fixed physical sample time  $t^s$ , instead,  $a = (t^s - kg_e)/d$  and its derivatives contain factors  $k/d$ . Formula (A.27.5) still holds where the margins hold, but its order- $m$  bounds can grow as  $C_m(1 + k/d)^m$ . No uniform onset response for all fixed schedules is asserted. When margins fail, or for an arbitrary bounded decoder, the convergence theorem alone gives no derivative estimate. The original finite-cut convergence and the smooth additive response remain separate valid statements.

#### APPENDIX A.28. THE CIRCULAR COLLISION HIERARCHY

A time window which admits two collisions need not test dependence between successive complete free flights. Under stationary preparation, its count law can depend only on the marginal law of one flight. The onset of higher collision tails is different: it requires an actual orbit segment, and its probability measures how a single two-dimensional collision state shadows a prescribed sequence of reflections. We determine this hierarchy for a periodic Lorentz gas, including its dependence on the length of the segment.

Let  $\Lambda = \mathbb{Z}(1, 0) + \mathbb{Z}(1/2, \sqrt{3}/2)$  and remove the closed radius- $R$  disk from  $\mathbb{R}^2/\Lambda$ . The particle has unit speed and reflects specularly. Throughout,

$$0 < R_- \leq R \leq R_+ < \frac{1}{2}, \quad K = [R_-, R_+], \quad g_R = 1 - 2R, \quad A_R = \frac{\sqrt{3}}{2} - \pi R^2, \quad \lambda_R = \frac{2R}{A_R}. \quad (\text{A.28.1})$$

Preparation is normalized phase volume  $dq d\vartheta/(2\pi A_R)$ . Write  $N_t$  for the number of collisions in  $(0, t]$ . The outgoing collision section has angular coordinates  $x = (\alpha, \phi)$ , with  $|\phi| < \pi/2$ , and invariant probability

$$d\nu = \frac{\cos \phi}{4\pi} d\alpha d\phi. \quad (\text{A.28.2})$$

Its collision map is  $B_R$ , its physical roof is  $\tau_R$ , and  $S_j \tau_R = \sum_{i=0}^{j-1} \tau_R \circ B_R^i$ . These are the actual mechanical map and roof. In particular, the summands are not independent. Finite horizon is not required: the proof uses the positive lower bound  $\tau_R \geq g_R$ , not an upper bound.

Put

$$c_R = \frac{1-R}{R}, \quad \chi_R = \operatorname{arccosh} c_R, \quad s_R = \sinh \chi_R = \frac{\sqrt{g_R}}{R}, \quad \chi_* = \min_{R \in K} \chi_R > 0. \quad (\text{A.28.3})$$

There are six unit lattice vectors  $e$ . Let  $x_e$  denote the outgoing normal state facing  $e$ ; then  $B_R x_e = x_{-e}$ . In angular coordinates these states do not depend on  $R$ .

##### A.28.1. Unweighted threshold law.

**Theorem A.28.1** (Uniform onset at every collision order). *There are  $\varepsilon_0 > 0$  and  $C < \infty$ , depending on  $K$  but not on  $j$ , such that  $\varepsilon_0 < \min_K g_R/4$  and, for all  $j \geq 1$ ,  $R \in K$  and  $0 < \varepsilon < \varepsilon_0$ ,*

$$\Pr_R\{N_{jg_R+\varepsilon} \geq j+1\} = \Pr_R\{N_{jg_R+\varepsilon} = j+1\} = C_j(R)\varepsilon^2(1+\varepsilon h_j(R, \varepsilon)), \quad C_j(R) = \frac{3\lambda_R}{2R \sinh(j\chi_R)}. \quad (\text{A.28.4})$$

The functions  $h_j$  are smooth up to  $\varepsilon = 0$ . Every fixed mixed derivative of  $h_j$  in  $R$  and  $\varepsilon$  is bounded uniformly in  $j$  and  $R \in K$ , after decreasing  $\varepsilon_0$  if necessary. For  $-\varepsilon_0 < \varepsilon \leq 0$  the probability in (A.28.4) is zero.

The theorem concerns the entire equilibrium event, not a prescribed transverse initial subset. Its threshold collar has width independent of the collision order. The observation times  $jg_R + \varepsilon$  grow with  $j$ ; the events remain the maximal possible collision counts in those windows. This is a uniform extremal regime, not a statement about typical long-time fluctuations.

**A.28.2. Sources and summable response.** Let  $f_R$  be a smooth real-valued function of the outgoing collision state. It is enough that it be smooth on fixed neighborhoods of the six normal states, uniformly on a neighborhood of  $K$ ; choose any bounded extension elsewhere. Let  $F_* \geq 1$  bound  $|f_R|$  on those neighborhoods. On the event in (A.28.4), write  $X_0, \dots, X_j$  for the outgoing states at its successive impacts and set

$$Z_j(R, q, \varepsilon) = \mathbb{E}_R \left[ \mathbf{1}_{\{N_{jg_R+\varepsilon}=j+1\}} \exp \left( q \sum_{i=0}^j f_R(X_i) \right) \right], \quad w_{j,e}(R) = \sum_{i=0}^j f_R(x_{(-1)^i e}).$$

The source is attached to a correlated impact record, not independently to a list of sampled roofs.

**Theorem A.28.2** (Marked threshold expansion). *For every fixed  $q_0 < \infty$ , there are smooth functions  $h_{j,e}(R, q, \varepsilon)$ , uniformly bounded with every fixed mixed derivative for  $R \in K$ ,  $|q| \leq q_0$  and  $0 \leq \varepsilon < \varepsilon_0$ , such that*

$$Z_j(R, q, \varepsilon) = \frac{\lambda_R \varepsilon^2}{4R \sinh(j\chi_R)} \sum_{|e|=1} e^{qw_{j,e}(R)} (1 + \varepsilon h_{j,e}(R, q, \varepsilon)). \quad (\text{A.28.5})$$

All constants are independent of  $j$ . The offset  $\varepsilon$  is held fixed in radius derivatives in this formula. If  $q_0 F_* < \chi_*$ , then for each fixed  $a, b, c \geq 0$ ,

$$\sup_{\substack{R \in K, |q| \leq q_0 \\ 0 \leq \varepsilon < \varepsilon_0}} \left| \partial_R^a \partial_q^b \partial_\varepsilon^c \frac{Z_j(R, q, \varepsilon)}{\varepsilon^2} \right| \leq C_{a,b,c} (1+j)^{a+b} e^{-(\chi_* - q_0 F_*)j}. \quad (\text{A.28.6})$$

The quotient is understood by its smooth right extension at zero. Consequently its series over  $j$  and every fixed mixed derivative converge uniformly on the displayed set. The same conclusions hold for finitely many sources, with  $qf$  replaced by a scalar product and  $q_0 F_*$  by a uniform bound for its absolute value per impact.

Uniformity in  $j$  is the main issue. An initial-value differentiation of  $B_R^j$  would amplify errors exponentially. Instead we parametrize an orbit segment by the transverse positions of its first and last impacts. Its effective action remains uniformly nondegenerate, while the flux density has size  $\sinh(j\chi_R)^{-1}$ . A relative determinant estimate, rather than an absolute error estimate, preserves this scale.

**A.28.3. Consequences and relation to the fixed-window theorem.** Theorem A.28.2 gives the finite-order response at all maximal-count onsets, including the derivative jumps across them. Section A.33 retains collision times, positions, both velocities and sampling-time cuts. Section A.34 compares the hierarchy with independent roofs having exactly the same marginal and recovers the normal periodic-orbit multiplier from the coefficients  $C_j$ .

The accompanying article [7] retains the full proof of the  $T = 11/100$ ,  $R \in [9/20, 47/100]$  three-point count law and its regular-level response. That is a complementary window: its second-collision probability is already positive throughout the radius interval. Here the count



changes support at  $t = jg_R$ , and its second response is singular there. The one-flight onset coefficient suggested by that calculation is the case  $j = 1$  of (A.28.4); the uniform multi-flight determinant and the full marked record follow from the uniform boundary-value argument and the explicit specialization below.

#### APPENDIX A.29. STATIONARY FLUX AND THE COMPLETE MARKED EVENT

For an outgoing boundary state write  $ds = R d\alpha$  and  $p = \sin \phi$ . Its unnormalized flux is  $ds dp$ , of total mass  $4\pi R$ . A flight tube has phase volume  $ds dp dr$ . Except for a null set of directions, a line on the lattice torus meets the open disk in both time directions: irrational lines are dense, and the remaining rational directions are countable. Tangencies and their finite regular preimages have zero flux. The positive gap prevents finite-time accumulation. Thus the flight tubes cover phase volume up to a null set, also when corridors give infinite horizon. It follows that

$$\bar{\tau}_R = \int \tau_R d\nu = \frac{A_R}{2R}, \quad d\mu_R = \lambda_R d\nu(x) ds, \quad 0 \leq s < \tau_R(x). \quad (\text{A.29.1})$$

Reflection preserves flux, and  $B_R$  is invertible and  $\nu$ -preserving on the full-measure regular section.

For completeness, integration of the allowable suspension positions gives

$$\Pr_R(N_t \geq k) = \lambda_R \int [(t - S_{k-1}\tau_R)_+ - (t - S_k\tau_R)_+] d\nu. \quad (\text{A.29.2})$$

Indeed, from  $(x, s)$  the  $k$ th hit occurs at  $S_k\tau_R(x) - s$ . The length of the allowed part of  $[0, \tau_R(x))$  is  $(t - S_{k-1}\tau_R(B_R x))_+ - (t - S_k\tau_R(x))_+$ ; invariance proves (A.29.2). This is a stationary suspension identity, not a renewal theorem. Since  $S_k\tau_R \geq kg_R$ , only finitely many tails are nonzero at each fixed time. Their sum telescopes to  $\mathbb{E}_R N_t = \lambda_R t$ .

A marked version must not shift the first term by invariance while leaving its mark unchanged. We use the following direct form instead.

**Proposition A.29.1** (First-impact coordinates). *Let  $t = jg_R + \varepsilon$ , where  $j \geq 1$  and  $0 < \varepsilon < g_R$ . Let  $y$  be the outgoing state at the first future impact and let  $r$  be its time. For any bounded Borel function  $G$  of the complete impact record,*

$$\mathbb{E}_R[\mathbf{1}_{\{N_t=j+1\}} G] = \lambda_R \int_{S_j\tau_R(y) < t} \int_0^{t-S_j\tau_R(y)} G(\mathcal{P}_{j,R}(y, r; t)) dr d\nu(y). \quad (\text{A.29.3})$$

Here  $\mathcal{P}_{j,R}$  is the actual path record, with impact times  $r + S_i\tau_R(y)$ ,  $0 \leq i \leq j$ , and collision states  $B_R^i y$ . The formula includes the initial and terminal free-flight portions.

*Proof.* Replace the last collision state  $x$  in (A.29.1) by  $y = B_R x$  and the suspension height by the first residual time  $r = \tau_R(x) - s$ . The measure is  $\lambda_R d\nu(y) dr$  with  $0 < r < \tau_R(B_R^{-1}y)$ . The event of at least  $j+1$  hits is  $r + S_j\tau_R(y) \leq t$ . Its upper bound on  $r$  is at most  $\varepsilon < g_R \leq \tau_R(B_R^{-1}y)$ , so no additional truncation occurs. A  $(j+2)$ nd hit would require  $t \geq (j+1)g_R$ , which is impossible. This proves the integration limits and the event identity.

Write  $q_i, v_i^+, v_i^-$  for the impact position and the outgoing and incoming velocities. Starting at  $q_0 - rv_0^-$ , the initial segment lasts  $r$ . Subsequent straight segments and specular reflections are prescribed by  $B_R^i y$ . After the last impact, the path continues for  $t - r - S_j\tau_R(y) \geq 0$ . Both end portions are shorter than  $g_R$  and cannot contain an additional impact. This gives the claimed complete mechanical record, including both velocity traces at each hit.  $\square$

Taking  $G = 1$  recovers

$$P_j(R, \varepsilon) := \Pr_R(N_{jg_R + \varepsilon} = j + 1) = \lambda_R \int (jg_R + \varepsilon - S_j\tau_R)_+ d\nu. \quad (\text{A.29.4})$$

For the source in Theorem A.28.2,  $G$  is independent of  $r$ , so the same formula holds with the correlated weight  $\exp(q \sum_{i=0}^j f_R(B_R^i y))$  inside the integral. For tests involving collision times or physical samples, the inner  $r$  integral in (A.29.3) must be retained.

### APPENDIX A.30. LOCALIZATION TO ACTUAL ALTERNATING ITINERARIES

The smallest distance between distinct disks is  $g_R$ , attained only by the six normal chords between nearest neighbors. The next center distance is  $\sqrt{3}$ . An outgoing ray cannot return to the same convex disk without another reflection.

**Lemma A.30.1** (Uniform localization). *There are a transverse coordinate radius  $a_0 > 0$  and an excess  $\varepsilon_0 > 0$ , independent of  $j \geq 1$  and  $R \in K$ , with the following property. Every regular  $j$ -flight segment satisfying  $S_j \tau_R < j g_R + \varepsilon_0$  alternates between one nearest-neighbor pair. Its impacts lie in the facing arcs  $|y_i| < a_0$  about the normal chord. Every segment in these arcs which satisfies the local specular stationarity equations is an actual, unobstructed billiard segment.*

*Proof.* Every flight has length at least  $g_R$ . If their sum has excess less than  $\varepsilon_0$ , each individual length is at most  $g_R + \varepsilon_0$ . Its target center  $e$  satisfies  $|e| \leq 2R + g_R + \varepsilon_0 = 1 + \varepsilon_0$ ; taking  $\varepsilon_0 < \sqrt{3} - 1$  forces  $|e| = 1$ .

For nearest disks, the minimum of the distance between their boundary points is unique. By compactness in  $R \in K$ , chords whose length is within  $\varepsilon_0$  of  $g_R$  have both endpoints and direction uniformly close to that normal chord. This compactness statement can also be read from the endpoint cosine bound

$$n_0 \cdot v, -n_1 \cdot v \geq \frac{1 - 2R - \tau^2}{2R\tau},$$

whose right side tends to one as  $\tau \downarrow g_R$ , uniformly on  $K$ . Reflection at the target therefore sends the direction close to  $-e$ . Distinct nearest-neighbor directions are separated by  $\pi/3$ . Taking the neighborhoods sufficiently small forces the next short flight to go back to the preceding disk. This argument uses the same neighborhood at every step and hence does not deteriorate with  $j$ .

The closed normal chord between the pair has positive distance from all other disks, uniformly on  $K$ . For the pair  $(0, (1, 0))$ , the other nearest centers have vertical coordinate  $\sqrt{3}/2$ , while  $R \leq R_+ < 1/2$ ; more distant centers are farther away. Small facing arcs preserve this separation. Their connecting chords are outgoing and incoming with normal cosines bounded away from zero. Convexity prevents a second crossing of the originating disk before the target. Thus these are first physical hits, and the stationarity equations, which express the equality of tangential velocity components, are exactly the specular reflection law. There is no grazing cutoff on a preselected preparation: the localization follows from the complete event's excess constraint.  $\square$

Fix the pair  $(0, (1, 0))$  and use a global transverse coordinate  $y$ . The facing boundary points are  $(\sqrt{R^2 - y^2}, y)$  on the first disk and  $(1 - \sqrt{R^2 - y^2}, y)$  on the second. The same symmetric length function applies in both directions:

$$\ell_R(u, v) = \sqrt{(1 - \sqrt{R^2 - u^2} - \sqrt{R^2 - v^2})^2 + (v - u)^2}. \quad (\text{A.30.1})$$

For endpoints  $u = y_0, v = y_j$ , define the local broken-length action

$$\mathcal{L}_{j,R}(y_0, \dots, y_j) = \sum_{i=0}^{j-1} \ell_R(y_i, y_{i+1}). \quad (\text{A.30.2})$$

The interior variables satisfy  $\partial_{y_i} \mathcal{L}_{j,R} = 0$ . We write  $W_{j,R}(u, v)$  for the stationary value, whose existence, uniqueness and uniform estimates are proved next.

The length has the even analytic expansion

$$\ell_R(u, v) = g_R + \frac{c_R(u^2 + v^2) - 2uv}{2g_R} + O((u^2 + v^2)^2). \quad (\text{A.30.3})$$

All derivatives of the remainder have uniform analytic bounds in a fixed small neighborhood for  $R \in K$ . In particular,  $-\partial_{uv}\ell_R(0, 0) = 1/g_R$ . There is also the elementary bound

$$\ell_R(u, v) - g_R \geq R - \sqrt{R^2 - u^2} + R - \sqrt{R^2 - v^2} \geq \frac{u^2 + v^2}{2R}. \quad (\text{A.30.4})$$

Consequently any localized orbit of total excess  $\varepsilon$  satisfies  $y_0^2 + y_j^2 + 2\sum_{i=1}^{j-1} y_i^2 \leq 2R\varepsilon$ . This bounds all interior variables at once, without a factor  $j$ .

#### APPENDIX A.31. A BOUNDARY-VALUE CALCULUS UNIFORM IN THE ITINERARY LENGTH

The circular formulas follow from the general boundary-value lemma. We record their explicit expressions, including the even analytic improvements, to keep this specialization self-contained at the level of its calculations. The uniform contraction and determinant argument are those already proved in Lemma 15.2.

**Lemma A.31.1** (Uniform stationary segment). *For endpoints  $|u|, |v| < a_1$ , with  $a_1 > 0$  independent of  $j$ , the interior equations of (A.30.2) have a unique small solution. The solution and  $W_{j,R}$  are real analytic in  $R, u, v$ . There are  $C < \infty$  and  $0 < \rho < 1$ , independent of  $j$ , such that, writing  $a = |u| + |v|$ ,*

$$y_i = \frac{\sinh((j-i)\chi_R)}{\sinh(j\chi_R)}u + \frac{\sinh(i\chi_R)}{\sinh(j\chi_R)}v + O(a^3(\rho^i + \rho^{j-i})), \quad (\text{A.31.1})$$

$$|y_i| \leq Ca(\rho^i + \rho^{j-i}), \quad 0 \leq i \leq j.$$

*The analytic bounds for the bridge and for  $W_{j,R} - jg_R$ , and every fixed parameter derivative of these functions, are uniform in  $j$ . Moreover*

$$W_{j,R}(u, v) = jg_R + \frac{1}{2}(u, v)H_{j,R}(u, v)^t + O(a^4), \quad (\text{A.31.2})$$

*with uniform analytic remainder bounds.*

*Proof.* Here  $\psi_0 = \psi_1 = R - \sqrt{R^2 - y^2}$  and  $c_0 = c_1 = c_R$ . The endpoint solution of the constant Jacobi equation is the displayed linear term. Formula (15.5) becomes

$$G_{ik} = g_R \frac{\sinh(i\chi_R) \sinh((j-k)\chi_R)}{\sinh \chi_R \sinh(j\chi_R)}, \quad G_{ki} = G_{ik} \quad (1 \leq i \leq k < j). \quad (\text{A.31.3})$$

For any fixed  $e^{-\chi_*} < \rho < 1$  and  $w_i = \rho^i + \rho^{j-i}$ ,

$$\sum_k e^{-\chi_*|i-k|} w_k \leq C_\rho w_i. \quad (\text{A.31.4})$$

Splitting at  $k = i$  proves this by the geometric ratios  $e^{-\chi_*}/\rho$  and  $e^{-\chi_*}\rho$ . The graphs are even and analytic on a common complex neighborhood for  $R \in K$ . The even analytic case of Lemma 15.2 therefore gives the cubic bridge correction, quartic action remainder, uniqueness throughout the fixed small box and all uniform analytic bounds asserted here. Its proof applies with interior row margin  $2/R$  and endpoint margin  $1/R$ , so none of these constants depends on  $j$ . For  $j = 1$  there are no interior variables and (A.30.1) gives the same assertions directly.  $\square$

**Lemma A.31.2** (Effective Hessian and twist). *The endpoint Hessian and mixed derivative at the normal segment are*

$$H_{j,R} = \frac{s_R}{g_R} \begin{pmatrix} \coth(j\chi_R) & -\operatorname{csch}(j\chi_R) \\ -\operatorname{csch}(j\chi_R) & \coth(j\chi_R) \end{pmatrix}, \quad d_{j,R}^0 := -W_{j,R,uv}(0,0) = \frac{s_R}{g_R \sinh(j\chi_R)}. \quad (\text{A.31.5})$$

In particular

$$\sqrt{\det H_{j,R}} = \frac{s_R}{g_R} = \frac{1}{R\sqrt{g_R}}, \quad \frac{d_{j,R}^0}{\sqrt{\det H_{j,R}}} = \frac{1}{\sinh(j\chi_R)}. \quad (\text{A.31.6})$$

Both eigenvalues of  $H_{j,R}$  are bounded above and below by positive constants independent of  $j$  and  $R \in K$ .

*Proof.* For the quadratic action the derivative at its first endpoint is  $(c_R u - y_1^{\text{lin}})/g_R$ . Insert the bridge formula and use  $c_R \sinh(j\chi_R) - \sinh((j-1)\chi_R) = s_R \cosh(j\chi_R)$ . Reversal gives the other endpoint derivative, proving (A.31.5). The identity  $\coth^2 x - \operatorname{csch}^2 x = 1$  proves (A.31.6). The eigenvalues are  $(s_R/g_R) \tanh(j\chi_R/2)$  and  $(s_R/g_R) \coth(j\chi_R/2)$ . Their asserted bounds follow from  $j \geq 1$ ,  $\chi_R \geq \chi_* > 0$  and compactness of  $K$ .  $\square$

An absolute bound  $W_{uv} = -d_{j,R}^0 + O(a^2)$  would be useless for large  $j$ . The next estimate is the relevant replacement.

**Lemma A.31.3** (Relative flux determinant). *On a fixed endpoint neighborhood,*

$$-W_{j,R,uv}(u,v) = d_{j,R}^0 b_{j,R}(u,v), \quad b_{j,R} = 1 + O(u^2 + v^2) > 0. \quad (\text{A.31.7})$$

The functions  $b_{j,R}$  are even under  $(u,v) \mapsto (-u,-v)$  and have uniform analytic bounds, including every fixed radius and endpoint derivative, independently of  $j$ .

*Proof.* The even analytic case of Lemma 15.2 proves the relative estimate and every fixed derivative bound. To identify its exact circular normalization, put  $b_i = -\ell_{R,uv}(y_i, y_{i+1})$  along the stationary segment. Its corner-cofactor identity is

$$-W_{j,R,uv} = \frac{\prod_{i=0}^{j-1} b_i}{\det H_{\text{int}}}, \quad (\text{A.31.8})$$

with empty determinant one when  $j = 1$ . Evenness gives  $\sum_i |\log(g_R b_i)| \leq Ca^2$ . If  $E = H_{\text{int}} - H_{\text{int}}^0$ , the same even improvement gives  $\|E\|_1 \leq Ca^2$  and the exact convergent expansion

$$\log \frac{\det H_{\text{int}}}{\det H_{\text{int}}^0} = \sum_{m \geq 1} \frac{(-1)^{m+1}}{m} \operatorname{tr}[(H_{\text{int}}^0)^{-1} E]^m = O(a^2). \quad (\text{A.31.9})$$

The trace-norm and differentiated estimates for this expansion were proved in the general lemma; no factor equal to the matrix dimension is introduced. Comparing (A.31.8) with its value at zero and using (A.31.5) proves (A.31.7). Uniqueness and simultaneous reversal of the coordinates prove evenness. Positivity follows from the positive twists and interior Hessian.  $\square$

**Proposition A.31.4** (Physical flux in endpoint coordinates). *For each of the six oriented nearest-neighbor segments, the small endpoint coordinates  $(u,v)$  parametrize its outgoing initial collision state one-to-one, and*

$$d\nu = \frac{-W_{j,R,uv}(u,v)}{4\pi R} du dv. \quad (\text{A.31.10})$$

Moreover  $S_{j\tau_R} = W_{j,R}(u,v)$  on this chart.

*Proof.* The first variation of length gives  $\partial_u W = -v_0^+ \cdot q_0'(u)$ . If  $s_0$  is boundary arclength, the outgoing tangential momentum is  $p_0 = -W_u / (ds_0/du)$ . Thus  $|ds_0 dp_0| = |W_{uv}| du dv$ . The mixed derivative is nonzero by Lemma A.31.3, so this is a local coordinate change. At fixed  $u$ ,  $p_0$  is strictly monotone in  $v$  throughout the small box, proving injectivity. The stationary path is physical by Lemma A.30.1; uniqueness identifies its successive states with the iterates of  $B_R$ . Normalization of  $ds_0 dp_0$  by its total mass  $4\pi R$  proves the measure formula. The length identity follows by adding the actual flight lengths.  $\square$

### APPENDIX A.32. UNIFORM THRESHOLD INTEGRATION AND ITS RESPONSE

We now integrate the actual flux on the whole onset event. The dimension of this integration is two, regardless of the number of flights.

**Lemma A.32.1** (Uniform even Morse coordinates). *There is an odd analytic change of variables  $z = \Phi_{j,R}(w)$  on a fixed neighborhood of zero in  $\mathbb{R}^2$ , with uniform inverse and derivative bounds, such that*

$$W_{j,R}(\Phi_{j,R}(w)) - jg_R = \frac{1}{2}|w|^2, \quad D\Phi_{j,R}(0) = H_{j,R}^{-1/2}. \quad (\text{A.32.1})$$

*In these variables the endpoint flux is*

$$d\nu = \frac{1}{4\pi R \sinh(j\chi_R)} a_{j,R}(w) dw, \quad a_{j,R}(w) = 1 + O(|w|^2), \quad (\text{A.32.2})$$

*where  $a_{j,R}$  is even, positive and uniformly analytic, including all fixed parameter derivatives.*

*Proof.* Put  $E(z) = W_{j,R}(z) - jg_R$ . Taylor's integral identity gives

$$E(z) = \frac{1}{2} z^t M(z) z, \quad M(z) = 2 \int_0^1 (1-t) D^2 E(tz) dt.$$

The matrix  $M$  is symmetric, even and analytic, with  $M(0) = H_{j,R}$  and  $M(z) - H_{j,R} = O(|z|^2)$  uniformly. The positive definite square root is analytic on a fixed neighborhood of the uniformly positive matrices  $H_{j,R}$ . Thus  $w = M(z)^{1/2} z$  is an odd analytic map with uniformly invertible differential at zero. Its nonlinear differential is a uniformly small perturbation on a fixed ball. Contraction applied to the equation  $z = H_{j,R}^{-1/2} w + O(|z|^3)$  gives an analytic inverse on a common smaller ball, with uniform derivative estimates. This proves (A.32.1) without a dimension-dependent Morse neighborhood.

Insert (A.31.7) into (A.31.10) and use (A.31.6). The factor  $b_{j,R}(\Phi(w)) \det D\Phi(w) \sqrt{\det H_{j,R}}$  is positive, even, uniformly analytic and equals  $1 + O(|w|^2)$ . It is the required  $a_{j,R}$ . Orientation may be chosen so the local Jacobian is positive.  $\square$

The energy bound (A.30.4) shows that, for a common smaller  $\varepsilon_0$ , the complete sublevel set in each of the six charts lies inside these coordinates. Conversely all of the small ball  $|w|^2 < 2\varepsilon$  gives physical segments. Thus no extraneous cutoff or boundary of a chart enters the following integral.

*Proof of Theorem A.28.1.* By (A.29.4) and (A.32.2), the contribution of one orientation is

$$\frac{\lambda_R}{4\pi R \sinh(j\chi_R)} \int_{|w|^2 < 2\varepsilon} (\varepsilon - |w|^2/2) a_{j,R}(w) dw.$$

The unweighted quadratic integral equals  $\pi\varepsilon^2$ . Indeed, rescale  $w = \sqrt{2\varepsilon} z$  and use  $\int_{|z|<1} (1 - |z|^2) dz = \pi/2$ . All six orientations have the same flux by lattice symmetry. The analytic even amplitude has a convergent expansion in homogeneous even polynomials, with uniform bounds. Integration on the unit disk therefore gives  $1 + \varepsilon h_j(R, \varepsilon)$  after division by the leading term; its

analytic bounds are uniform in  $j$ . This proves (A.28.4) and every stated derivative bound. The roof lower bound proves vanishing on the other side of the onset.  $\square$

**Lemma A.32.2** (Additive marks on the endpoint bridge). *In the  $e$  chart define*

$$V_{j,e}(R, u, v) = \sum_{i=0}^j (f_R(X_i(R, u, v)) - f_R(x_{(-1)^i e})).$$

*Then  $V_{j,e}(R, 0, 0) = 0$ , and  $V_{j,e}$  and every fixed derivative in  $R, u, v$  have bounds independent of  $j$ . In particular  $|V_{j,e}| \leq C(|u| + |v|)$ .*

*Proof.* Each impact position depends on  $y_i$ . The direction of its adjacent flight is the chord divided by (A.30.1); its denominator is at least  $\min_K g_R$ . Interior and final outgoing velocities are obtained by the specular formula  $v^+ = v^- - 2(v^- \cdot n)n$ . These maps are analytic on the fixed arcs. The last velocity uses  $(y_{j-1}, y_j)$ ; no additional endpoint is needed.

Consequently the deviation of  $X_i$  from its normal state is bounded by  $Ca(w_{i-1} + w_i + w_{i+1})$ , omitting unavailable indices. The same weighted bounds hold for every fixed derivative of these deviations: use the complex estimates in Lemma A.31.1, or differentiate the local maps and the bridge. The normal states in angular coordinates are independent of  $R$ . Applying the finite chain rule to  $f_R$  and subtracting its value at the normal state gives these weighted bounds for the summands as well. Since  $\sum_i w_i \leq 2/(1 - \rho)$ , their sum has uniform derivative bounds. The value at zero and the linear bound then follow from the mean value formula.  $\square$

*Proof of Theorem A.28.2.* In one orientation, factor the source into  $e^{qw_{j,e}(R)} e^{qV_{j,e}(R, u, v)}$ . Lemmas A.32.1 and A.32.2 reduce its normalized threshold integral to

$$F_{j,e}(R, q, \varepsilon) = \frac{2}{\pi} \int_{|z| < 1} (1 - |z|^2) a_{j,R}(\sqrt{2\varepsilon} z) \exp(qV_{j,e}(R, \Phi_{j,R}(\sqrt{2\varepsilon} z))) dz. \quad (\text{A.32.3})$$

It equals one at  $\varepsilon = 0$ . Replace the smooth amplitude in this integral by its even part in  $z$ . The integral is unchanged, and every odd Taylor polynomial disappears. Taylor's formula to any prescribed finite order, including its differentiated remainder, shows that  $F_{j,e} = 1 + \varepsilon h_{j,e}$  with all fixed mixed derivative bounds uniform in  $j$ . More explicitly, to control  $c$  derivatives in  $\varepsilon$  use the even Taylor expansion to degree  $2c + 4$  on a slightly larger ball; its integral remainder after the substitution  $w = \sqrt{2\varepsilon}z$  is  $O(\varepsilon^{c+2})$  with the corresponding differentiated bounds. Radius and source derivatives obey the same argument by Lemma A.32.2. This is finite-order smooth calculus and makes no analyticity assumption on  $f_R$ . Multiplying back the constants in (A.32.2) proves (A.28.5).

For every fixed  $a$ , direct differentiation of  $\text{csch}(j\chi_R) = 2e^{-j\chi_R}/(1 - e^{-2j\chi_R})$  gives

$$|\partial_R^a \text{csch}(j\chi_R)| \leq C_a(1 + j)^a e^{-j\chi^*}.$$

Also  $|\partial_R^a w_{j,e}| \leq C_a(j + 1)$  and  $|e^{qw_{j,e}}| \leq e^{q_0 F^*(j+1)}$ . The finite product and chain rules, together with the uniform bounds for  $h_{j,e}$ , prove (A.28.6). Polynomially weighted geometric series converge, proving termwise differentiation and the source-vector version.  $\square$

**A.32.1. The singularity across the onset surface.** Set

$$C_j(R, q) = \frac{\lambda_R}{4R \sinh(j\chi_R)} \sum_{|e|=1} e^{qw_{j,e}(R)}.$$

This is positive for real sources. Formula (A.28.5), with zero on the negative side, is a full expansion of the form  $\varepsilon_+^2 [C_j(R, q) + O(\varepsilon)]$  with smooth right-hand coefficients. Thus it is  $C^1$ ,

and its second  $t$  derivative has jump  $2C_j(R, q)$ . It is not  $C^2$  across any of these onsets. At fixed physical  $t$  and a radius  $R_c$  with  $t = jg_{R_c}$ , the jump of the second radius derivative is

$$8j^2 C_j(R_c, q). \quad (\text{A.32.4})$$

Indeed  $\varepsilon = t - jg_R = 2j(R - R_c)$ ; all other second-derivative terms have a factor  $\varepsilon$  and vanish at the surface. For higher distributional derivatives, differentiation of  $\varepsilon_+^2$  supplies the corresponding delta terms; for example its third  $t$  derivative contributes  $2C_j(R, q)\delta(t - jg_R)$ . At fixed  $t$ , radius differentiation is  $\partial_R + 2j\partial_\varepsilon$  in onset coordinates. Therefore all fixed-order coefficient and one-sided derivative estimates remain polynomially weighted, exponentially summable in  $j$  in the source range of (A.28.6). This statement concerns the onset collars, not singularities of other itinerary families at other times.

### APPENDIX A.33. THE JOINT IMPACT RECORD AND PHYSICAL SAMPLING CUTS

The endpoint parametrization gives more than a count or a marginal roof. It reconstructs all positions, directions and collision times in the window. We describe its onset law explicitly.

Let

$$\mathcal{U} = \{(z, \eta) \in \mathbb{R}^2 \times \mathbb{R} : |z| < 1, 0 < \eta < 1 - |z|^2\}, \quad \text{dm} = \frac{2}{\pi} \mathbf{1}_{\mathcal{U}} \, dz \, d\eta. \quad (\text{A.33.1})$$

This is a probability measure. For  $z \in \mathbb{R}^2$ , set  $(U, V)^t = H_{j,R}^{-1/2} z$  and define

$$\begin{aligned} Y_i(z) &= \frac{\sinh((j-i)\chi_R)}{\sinh(j\chi_R)} U + \frac{\sinh(i\chi_R)}{\sinh(j\chi_R)} V, \quad 0 \leq i \leq j, \\ Q_i(z) &= \frac{Y_i^2 + Y_{i+1}^2}{R} + \frac{(Y_{i+1} - Y_i)^2}{g_R}, \quad 0 \leq i < j. \end{aligned} \quad (\text{A.33.2})$$

These quadratic functions satisfy  $Q_i \geq 0$  and  $\sum_{i=0}^{j-1} Q_i = |z|^2$ .

**Theorem A.33.1** (Marked onset record). *Condition on  $N_{jg_R+\varepsilon} = j+1$ . The orientation  $e$  is exactly uniform on the six nearest lattice vectors. In that orientation, use the exact Morse coordinates  $w$  of the endpoints and the first-impact time  $r$ , and set  $z = w/\sqrt{2\varepsilon}$ ,  $\eta = r/\varepsilon$ . Their law has total variation distance at most  $C\varepsilon$  from  $\mathfrak{m}$ , uniformly in  $j \geq 1$  and  $R \in K$ .*

*On these same coordinates, the transverse impact positions and impact times  $t_i$  satisfy*

$$\frac{y_i}{\sqrt{2\varepsilon}} = Y_i(z) + O(\varepsilon), \quad \frac{t_i - ig_R}{\varepsilon} = \eta + \sum_{h=0}^{i-1} Q_h(z) + O(\varepsilon), \quad 0 \leq i \leq j, \quad (\text{A.33.3})$$

*with errors uniform in the index, itinerary length, radius and  $(z, \eta) \in \mathcal{U}$ . Both velocity traces are obtained from the same bridge by the chord and reflection formulas.*

*Consider any fixed finite number of samples at  $t^s = kg_R + \kappa\varepsilon \in [0, jg_R + \varepsilon]$ , with  $0 \leq k \leq j$  and bounded fixed  $\kappa$ . In the limiting record, the sample precedes impact  $k$  exactly on*

$$\eta + \sum_{h < k} Q_h(z) > \kappa. \quad (\text{A.33.4})$$

*The joint probabilities of the resulting before/after cuts converge with error  $O(\varepsilon)$ , uniformly in  $j$  and  $R$ . Thus the sampling-time interfaces are retained, not suppressed by pointwise differentiation away from event times.*

*Proof.* Keep the  $r$  integral in (A.29.3). The changes  $w = \sqrt{2\varepsilon}z$ ,  $r = \varepsilon\eta$  give exactly the domain  $\mathcal{U}$  and, up to its normalizing constant, density  $a_{j,R}(\sqrt{2\varepsilon}z)$ . By (A.32.2), this density is  $1 + O(\varepsilon)$  uniformly. Since  $\text{vol}\mathcal{U} = \pi/2$ , normalization proves the total variation assertion. Rotation

through multiples of  $\pi/3$  preserves phase volume and the event, proving the exact orientation weights.

The odd Morse inverse has expansion  $\Phi_{j,R}(w) = H_{j,R}^{-1/2}w + O(|w|^3)$ . Insert it into (A.31.1). This proves the first formula in (A.33.3), with the stronger error  $O(\varepsilon(\rho^i + \rho^{j-i}))$ . Expanding each chord with (A.30.3) gives  $\ell_R(y_i, y_{i+1}) = g_R + \varepsilon Q_i(z) + O(\varepsilon^2(\rho^i + \rho^{j-i})^2)$ . The sum of these error bounds is uniform, proving the time formula. The identity for  $\sum Q_i$  is the quadratic endpoint action evaluated at  $H_{j,R}^{-1/2}z$ .

For clarity, in the pair's fixed global transverse direction the leading outgoing transverse velocity at impacts  $i < j$ , divided by  $\sqrt{2\varepsilon}$ , is  $(Y_{i+1} - Y_i)/g_R$ . At the final impact it is  $(Y_j - Y_{j-1})/g_R + 2Y_j/R$ ; the initial incoming value is  $(Y_1 - Y_0)/g_R - 2Y_0/R$ . The longitudinal components and the unscaled velocities are given exactly by the unit chord and specular formulas. Uniform analyticity on the fixed arcs gives their corresponding Taylor bounds. Together with the initial and terminal pieces in Proposition A.29.1, this reconstructs the entire physical path.

Consecutive impacts are separated by at least  $g_R$ , so a sample within  $O(\varepsilon)$  of  $kg_R$  can cross only the  $k$ th impact for small  $\varepsilon$ . Formula (A.33.3) gives (A.33.4). The cut has derivative one in  $\eta$ . For each fixed  $z$ , the set where its left side lies within  $\delta$  of  $\kappa$  has  $\eta$ -length at most  $2\delta$ . Its  $m$ -mass is therefore at most  $4\delta$ . Taking  $\delta = C\varepsilon$  and adding the  $O(\varepsilon)$  density error bounds the mismatch probability. A union bound proves the result for any fixed finite set of cuts. Bounded Lipschitz tests of the scaled positions and times can be added using their uniform pathwise errors. No differentiability of a discontinuous sampling decoder is assumed.  $\square$

For the source-weighted conditional law, the orientation weights tend to  $e^{qw_{j,e}} / \sum_{e'} e^{qw_{j,e'}}$ . Given an orientation, the limiting law is still (A.33.1); its total variation error is  $O(\sqrt{\varepsilon})$  uniformly for bounded sources, by Lemma A.32.2. The integrated source normalization has the stronger  $O(\varepsilon)$  relative error of Theorem A.28.2 because its odd part cancels. The distinction between these two errors is important for a full marked law.

#### APPENDIX A.34. MECHANICAL CONSEQUENCES OF THE THRESHOLD COEFFICIENTS

##### A.34.1. A multi-roof effect not determined by the one-roof law.

**Corollary A.34.1** (The nonrenewal onset exponent). *Fix  $R \in (0, 1/2)$ . Let an independent-roof stationary suspension have the same section roof marginal as the Lorentz gas. Its intensity is therefore the same  $\lambda_R$ . For each fixed  $j \geq 1$ , its maximal count onset satisfies*

$$\lim_{\varepsilon \downarrow 0} \Pr(N_{jg_R + \varepsilon} = j + 1) = \frac{\lambda_R}{(j+1)!} \left( \frac{3}{\sqrt{g_R}} \right)^j \varepsilon^{j+1} (1 + O_j(\varepsilon)). \quad (\text{A.34.1})$$

*For the actual Lorentz gas the exponent is two for every  $j$ , with coefficient (A.28.4). In particular the two laws have different leading powers for every  $j \geq 2$ , despite their identical one-roof marginal and intensity.*

*Proof.* The case  $j = 1$  of the action calculation gives  $\int (g_R + \varepsilon - \tau_R)_+ d\nu = 3\varepsilon^2/(2\sqrt{g_R}) + O(\varepsilon^3)$ . The smooth threshold density proved above permits differentiation. Thus the excess roof  $u = \tau_R - g_R$  has density  $3/\sqrt{g_R} + O(u)$  near  $u = 0$ . For  $j$  independent excess roofs, integrate their product density on  $u_i \geq 0$ ,  $\sum_i u_i < \varepsilon$  with weight  $\varepsilon - \sum_i u_i$ . The leading integral is  $\varepsilon^{j+1}/(j+1)!$ , as follows by integrating successively on this simplex. The error is  $O_j(\varepsilon)$  relative to it. The stationary identity (A.29.2), with  $\varepsilon < g_R$ , supplies the factor  $\lambda_R$  and excludes the next count. This proves (A.34.1). The actual law is Theorem A.28.1.  $\square$



For example the first genuinely multi-roof count has

$$\Pr_R(N_{2g_R+\varepsilon} = 3) = \frac{3\lambda_R R}{4(1-R)\sqrt{g_R}} \varepsilon^2 (1 + O(\varepsilon)). \quad (\text{A.34.2})$$

The independent-roof substitute instead has coefficient  $3\lambda_R/(2g_R)$  multiplying  $\varepsilon^3$ . The stationary identity is the same in both systems; the difference comes from the specular boundary-value geometry of  $S_2\tau_R$ .

#### A.34.2. Recovery of periodic-orbit instability.

**Corollary A.34.2** (Instability in count probabilities). *The unstable multiplier of the normal period-two return map is  $e^{2\chi_R}$ . The measured threshold coefficients satisfy*

$$\lim_{j \rightarrow \infty} \frac{C_{j+1}(R)}{C_j(R)} = e^{-\chi_R}, \quad \frac{C_1(R)}{C_2(R)} = 2 \cosh \chi_R. \quad (\text{A.34.3})$$

*Thus they determine that multiplier. The unweighted coefficient series has the exact representation*

$$\sum_{j \geq 1} C_j(R) z^j = \frac{3\lambda_R}{R} \sum_{m \geq 0} \frac{z e^{-(2m+1)\chi_R}}{1 - z e^{-(2m+1)\chi_R}}, \quad |z| < e^{\chi_R}. \quad (\text{A.34.4})$$

*It extends meromorphically to the complex plane with simple poles at  $z = e^{(2m+1)\chi_R}$ ,  $m \geq 0$ .*

*Proof.* The Jacobi recurrence in the proof of Lemma A.31.1 has transfer matrix  $\begin{pmatrix} 2c_R & -1 \\ 1 & 0 \end{pmatrix}$  and eigenvalues  $e^{\chi_R}, e^{-\chi_R}$ . The local change between successive transverse positions and a collision position–momentum pair is invertible since the one-flight twist is  $1/g_R$ . After two flights the coordinate frame returns to the original orientation, so the physical period-two derivative is conjugate to the square of this transfer matrix. This proves the multiplier claim. The ratios follow from the explicit  $C_j$ .

Expand  $\text{csch}(j\chi) = 2 \sum_{m \geq 0} e^{-(2m+1)j\chi}$  and sum the absolutely convergent geometric series in  $j$ . On every compact set avoiding the displayed poles, the tail in  $m$  converges uniformly, since its terms are  $O(|z|e^{-(2m+1)\chi})$ . This gives the meromorphic continuation. At each indicated point only one summand has a pole and its residue is nonzero, so there is no cancellation.  $\square$

Series (A.34.4) organizes threshold coefficients belonging to different observation windows. It is not the ordinary pressure or the generating function of counts in one fixed window. The recovery of instability here uses equilibrium tail amplitudes; it does not claim that the relation between periodic action Hessians and monodromy is new.

**A.34.3. The retained regular-window response.** The complete companion [7] proves its fixed-window theorem, including the all-order moving-level formula, without using the present threshold asymptotic. We record the terminal-level witness requested for that calculation. At  $T = 11/100$ ,  $R \in [9/20, 47/100]$ , choose  $\alpha = 0$  facing  $(1, 0)$  and

$$\cos \vartheta = \frac{1 - 2R + T^2}{2T(1 - R)}. \quad (\text{A.34.5})$$

The right side lies strictly between zero and one. The point  $(R, 0) + T(\cos \vartheta, \sin \vartheta)$  lies on the target circle, since substitution gives squared distance  $R^2$  from its center. It is an incoming hit because

$$(1 - R) \cos \vartheta - T = \frac{1 - 2R - T^2}{2T} > 0.$$

The outgoing sign and the no-obstruction estimate in the companion make it a genuine first hit. The terminal level is therefore nonempty. Its uniform gradient bound and  $\partial_{RT} < 0$  show directly that the level integral in the companion's second-derivative formula is strictly positive.

This is compatible with the onset singularity: the fixed window remains strictly above  $g_R$  on its specified interval, whereas (A.28.4) crosses the critical level itself.

#### APPENDIX A.35. ALTERNATING-CHANNEL FORMULATION AND RELATIVE COMPARISON

**A.35.1. The local relative and smooth inverse mechanism.** Fix a clear closest channel between two smooth strictly convex reflecting arcs, with gap  $g > 0$ , positive curvatures  $\kappa_0, \kappa_1$ , and a common signed transverse coordinate. In contact frames the arcs are  $(-\psi_0(u), u)$  and  $(g + \psi_1(v), v)$ , with  $\psi_b(0) = \psi'_b(0) = 0$  and  $\psi''_b(0) = \kappa_b$ . The stationary length of  $j$  alternating flights is  $W_j(u, v)$ . Put  $c_b = 1 + g\kappa_b$ ,  $c = \sqrt{c_0 c_1}$ ,  $\gamma = \operatorname{arcosh} c$ , and  $p = j \bmod 2$ . Geometric parameter derivatives, when taken, keep the excess time  $d = t - jg$  fixed in these moving intrinsic coordinates.

**Theorem A.35.1** (Relative boundary law and actual smooth contact inverse). *For a selected clear channel in a smooth periodic dispersing table, there are a flight-independent contact collar, smooth strictly convex actions  $S_b$  and positive smooth amplitudes  $B_b$ , normalized by  $S_b(0) = S'_b(0) = 0$  and  $B_b(0) = 1$ , with the following properties.*

- (1) Write  $a_b = S''_b(0) = c \sinh \gamma / (gc_{1-b})$  and  $d_j^0 = \sqrt{a_0 a_p} / \sinh(j\gamma)$ . At every separately fixed finite differentiation order, on a sufficiently small common collar,

$$\|W_j - jg - (S_0 \oplus S_p)\|_{C^k} + \left\| \frac{-W_{j,uv}}{d_j^0} - B_0 \otimes B_p \right\|_{C^k} \leq C_k \tau^j, \quad 0 < \tau < 1. \quad (\text{A.35.1})$$

The estimates include the stipulated fixed-order family derivatives; the constants may depend on higher smooth norms. No estimate uniform in the differentiation order is asserted.

- (2) At fixed sufficiently small positive excess  $d$ , the limiting same-type conditional endpoint law obtained from the full phase ensemble has positive interior density

$$f_b(u, v) = Z_b^{-1} B_b(u) B_b(v) \{d - S_b(u) - S_b(v)\}. \quad (\text{A.35.2})$$

The gap and one such law of each type determine the two signed action germs and every finite labelled contact jet, without supplied amplitudes, curvatures or normal collision coordinates. The extraction is smooth at each fixed finite order on the positive, nondegenerate domains specified in the density inverse below.

- (3) For  $n \geq 3$ , after lower jets are fixed, write  $s_n = (S_0^{(n)}(0), S_1^{(n)}(0))^t$  and  $q_n = (\psi_0^{(n)}(0), \psi_1^{(n)}(0))^t$ . Then

$$s_n = M_n q_n + R_n, \quad M_n = \begin{pmatrix} \coth(n\gamma) & \mathfrak{r}_0^n \operatorname{csch}(n\gamma) \\ \mathfrak{r}_1^n \operatorname{csch}(n\gamma) & \coth(n\gamma) \end{pmatrix}, \quad \mathfrak{r}_b = \sqrt{c_b / c_{1-b}}, \quad \det M_n = 1. \quad (\text{A.35.3})$$

Here  $R_n$  depends only on the gap and lower graph jets. The triangular inverse holds for actual smooth representatives, including odd contact orders. More precisely, let  $M \geq 2$  and let two anchored graph pairs have the same gap and equal graph jets through order  $M$  at their respective anchored contacts. Under common functional smooth bounds on a sufficiently small collar, their actions satisfy

$$|S_b^{[1]}(u) - S_b^{[0]}(u)| \leq \frac{C_M}{1 - \rho^{M+1}} |u|^{M+1}, \quad 0 < \rho < 1. \quad (\text{A.35.4})$$

*The highest-degree blocks have the additional bounds*

$$\begin{aligned} \max\{\|M_n\|_\infty, \|M_n^{-1}\|_\infty\} &\leq \frac{3+t^6}{1-t^6}, \\ \max\{\|M_n - I\|_\infty, \|M_n^{-1} - I\|_\infty\} &\leq \frac{4\vartheta^n}{1-t^6}, \end{aligned} \quad t = e^{-\gamma}, \quad \vartheta = \frac{1+t^2}{2} < 1. \quad (\text{A.35.5})$$

*Here the norm is the maximum absolute row sum. These bounds are uniform in  $n \geq 3$ , and uniformly over geometries with  $\gamma \geq \gamma_* > 0$ . They do not assert an order-uniform bound for the lower-jet remainders or the complete nonlinear inverse.*

*The local forward and finite-jet conclusions do not require analytic boundaries, equal curvatures, or reflection symmetry.*

*Proof.* Lemma 15.2 constructs the finite stationary bridges. Lemma 17.1 constructs their half-line limits. The two-ended normalized determinant comparison in Theorem 17.2 proves (A.35.1). The full phase formula (16.1), the common-domain integration in Lemma 16.1, and Theorem 17.3 give (A.35.2). Theorem 23.1 and Proposition 23.2 extract the actions and leading curvatures on the positive interior domain. Lemma 22.4 proves (A.35.4) before the last-jet calculation is made; Proposition 22.6 gives (A.35.3) and its inverse. Proposition 22.9 uses the positive-curvature constraint to prove (A.35.5). All of these supporting proofs are printed in the present article. The block bounds are order-uniform; the nonlinear smooth estimates and reconstruction remain at each separately prescribed finite order.  $\square$

**A.35.2. Why the relative estimate is the relevant limit.** The estimates have different roles. To prove the relative estimate one uses, before division by  $d_j^0$ , the exact identity

$$\log b_j = \sum_{i=0}^{j-1} \log\{g[-\ell_{i \bmod 2, uv}(y_i, y_{i+1})]\} - \log \det(I + G_j \Delta H_j), \quad b_j = -W_{j, uv}/d_j^0. \quad (\text{A.35.6})$$

The perturbation  $\Delta H_j$  is localized near two ends; its trace norm is estimated by a summable weight, not by the number of interior sites. The finite Green compressions are then compared to their half-line limits. Fixed powers introduced by differentiation are absorbed by strict exponential margins. For the inverse, one interpolates actual smooth graphs and differentiates finite stationary action sums. Interior variations cancel by stationarity; the terminal term is retained until its decay has been proved. Only then is the finite Taylor recursion used. Thus neither an absolute-error estimate nor a formal power series is substituted for the required nonlinear argument.

The distinction can be stated directly at the level of the observation. On a fixed positive interior square, integrating out the residual time and normalizing gives

$$f_{j,b,d}(u, v) = \frac{b_{j,b}(u, v)\{d - E_{j,b}(u, v)\}}{\int (d - E_{j,b}) + b_{j,b}}.$$

The exponentially small reference twist cancels from this quotient. That cancellation is useful only after  $b_{j,b}$ , not merely  $E_{j,b}$ , has been controlled. The relative estimate supplies the amplitude and normalization bounds needed for the positive-offset limit. In particular, Corollary 23.3 gives, at each fixed order  $M$ , the geometric error  $C_M(\tau^j + \epsilon + \epsilon_g)$  from a finite-flight interior  $C^M$  density error  $\epsilon$  and gap error  $\epsilon_g$ . This connects the relative limit to its inverse without an assumption of order-uniform real differentiation or a bound on the number of preparations required to estimate the density.

## APPENDIX A.36. GRAPH AND SUPPORT COORDINATES

The differential image argument uses a local geometric coordinate change, not the quantitative statistical hypotheses of the full technical manuscript. We record it here so that its proof is part of this technical manuscript.

**Lemma A.36.1** (Finite graph-to-support conversion). *In an outward-normal contact frame let the boundary be  $(-\psi(y), y)$ , where  $\psi(0) = \psi'(0) = 0$  and  $\psi''(0) = \kappa > 0$ . On a small normal-angle interval there is a unique solution  $\psi'(y(\theta)) = \tan \theta$ . Its support function is*

$$h(\theta) = -\psi(y(\theta)) \cos \theta + y(\theta) \sin \theta. \quad (\text{A.36.1})$$

*For each  $m \geq 2$ , the support jet through degree  $m$  is a smooth function of the graph jet through degree  $m$ , on positive curvature domains. The same assertion holds for parameter derivatives in the corresponding finite smooth norms. In particular, zero variations of all anchored graph coefficients imply zero variations of all anchored support coefficients.*

*Proof.* The scalar implicit function theorem applies to  $\psi'(y) = p$  at  $(y, p) = (0, 0)$  since its derivative in  $y$  is  $\kappa$ . Strict convexity identifies this point with the unique local support point for normal  $(\cos \theta, \sin \theta)$ . Put  $H(p) = -\psi(y(p)) + py(p)$ . Stationarity gives  $H'(p) = y(p)$ , so  $h(\theta) = \cos \theta H(\tan \theta)$ . The coefficient of  $p^k$  in  $y(p)$ , for  $k \leq m - 1$ , is obtained by successively solving the scalar equation  $\psi'(y(p)) = p$ ; it depends only on graph coefficients through order  $k + 1$  and divides by the positive number  $\kappa$ . Since  $H(0) = 0$ , integration of  $H' = y$  shows that the jet of  $H$  through order  $m$  needs no higher graph jet. The compositions by  $\tan \theta$  and  $\cos \theta$  preserve this finite-order property. The implicit function theorem and the finite chain rule justify this computation for actual smooth functions, and also justify parameter differentiation. For example  $h(0) = h'(0) = 0$  and  $h''(0) = \kappa^{-1}$ .

When the second contact's outward normal is reversed, its positively oriented tangent must also be reversed. This known proper frame change multiplies its degree- $n$  graph coefficient by  $(-1)^n$  before the same formula is used. It introduces no unknown relative pose. Thus vanishing of all graph-coefficient variations in the common signed channel frame implies vanishing of all support-coefficient variations on both sides.  $\square$

## REFERENCES

- [1] P. Bálint, J. De Simoi, V. Kaloshin and M. Leguil, Marked length spectrum, homoclinic orbits and the geometry of open dispersing billiards, *Comm. Math. Phys.* **374** (2020), no. 3, 1531–1575. doi:10.1007/s00220-019-03448-x.
- [2] S. V. Bolotin and D. V. Treschev, Hill's formula, *Russian Math. Surveys* **65** (2010), no. 2, 191–257. doi:10.1070/RM2010v065n02ABEH004671.
- [3] M. Demers, I. Melbourne and M. Nicol, Martingale approximations and anisotropic Banach spaces with an application to the time-one map of a Lorentz gas, *Nonlinearity* **33** (2020), 4095–4113. doi:10.1088/1361-6544/ab7d22.
- [4] M. F. Demers and H.-K. Zhang, A functional analytic approach to perturbations of the Lorentz gas, *Comm. Math. Phys.* **324** (2013), no. 3, 767–830. doi:10.1007/s00220-013-1820-0.
- [5] J. De Simoi, V. Kaloshin and M. Leguil, Marked length spectral determination of analytic chaotic billiards with axial symmetries, *Invent. Math.* **233** (2023), 829–901. doi:10.1007/s00222-023-01191-8.
- [6] J. Marklof, Entry and return times for semi-flows, *Nonlinearity* **30** (2017), no. 2, 810–824. doi:10.1088/1361-6544/aa518b.
- [7] Q. Qi, *Two-collision counting response in a periodic Lorentz gas*, September 8, 2026, source-pinned companion to this article. The complete original is retained as `two_collision.tex`.
- [8] Q. Qi, *A2: nonempty arithmetic and a checkable Fourier-inversion budget*, Round 33 source reconstruction, September 3, 2026. The original chapter and source identity are retained in the submission history, with the hypotheses printed in that chapter.
- [9] M. Carney, M. Nicol and H.-K. Zhang, Compound Poisson law for hitting times to periodic orbits in two-dimensional hyperbolic systems, *J. Stat. Phys.* **169** (2017), 804–823. doi:10.1007/s10955-017-1893-9.

- [10] D. Batenkov and Y. Yomdin, On the accuracy of solving confluent Prony systems, *SIAM J. Appl. Math.* **73** (2013), no. 1, 134–154. doi:10.1137/110836584.
- [11] A. Florio and M. Leguil, *Smooth conjugacy classes of 3D Axiom A flows*, arXiv:2010.04120v5, June 3, 2021. <https://arxiv.org/abs/2010.04120v5>.
- [12] D. Finamore and M. Leguil, *CAT(0)-approach to the marked length spectral rigidity of Sinai billiards*, arXiv:2510.18983v1, October 21, 2025. <https://arxiv.org/abs/2510.18983v1>.
- [13] O. V. Osterman, On length spectrum rigidity of dispersing billiard systems, *J. Mod. Dyn.* **19** (2023), 847–878. doi:10.3934/jmd.2023025.
- [14] S. Zelditch, Inverse spectral problem for analytic domains, II:  $\mathbb{Z}_2$ -symmetric domains, *Ann. of Math. (2)* **170** (2009), no. 1, 205–269. doi:10.4007/annals.2009.170.205.
- [15] G. Meurant, A review on the inverse of symmetric tridiagonal and block tridiagonal matrices, *SIAM J. Matrix Anal. Appl.* **13** (1992), no. 3, 707–728. doi:10.1137/0613045.
- [16] R. Nabben, Decay rates of the inverse of nonsymmetric tridiagonal and band matrices, *SIAM J. Matrix Anal. Appl.* **20** (1999), no. 3, 820–837. doi:10.1137/S0895479897317259.
- [17] Independent AI referee-style memorandum, *A2 v8: relative boundary laws*, September 9, 2026, and its companion *The smooth-envelope minimax experiment*, repository `TrillionniumFoundation/theta-theory`, review commit `79ff2f96fc2e41899666065355238915f21273b0`.
- [18] Z. Dai and P. K. Lamm, Local regularization for the nonlinear inverse autoconvolution problem, *SIAM J. Numer. Anal.* **46** (2008), no. 2, 832–868. doi:10.1137/070679247.
- [19] B. Hofmann, F. Werner and Y. Deng, *On uniqueness and ill-posedness for the deautoconvolution problem in the multi-dimensional case*, arXiv:2212.06534v1, December 13, 2022. <https://arxiv.org/abs/2212.06534v1>. Related journal publication: doi:10.1088/1361-6420/acd07b.
- [20] Independent AI referee-style memorandum, *Completed A2 v9 nonlinear-boundary revision*, September 10, 2026, repository `TrillionniumFoundation/theta-theory`, review commit `b756f4851669e34c7074ba61b5bcf48689756406`.
- [21] Independent AI referee-style memorandum, *A2 v10: observable boundary profiles*, September 10, 2026, repository `TrillionniumFoundation/theta-theory`, review commit `2b893b931a894dfc9e7730b853f575124ffb5c55`. The report is an author-requested independent review, not a journal editorial decision; its source is retained in the revision history.
- [22] Independent AI referee-style memorandum, *A2 v11: Abel stability and integrated-flux acquisition*, September 10, 2026, repository `TrillionniumFoundation/theta-theory`, review commit `bcde40b514e4dc00a3252f2d0710002a2caec7b`. Section 5 contains the two-derivative gain, the improved sufficient preparation bound and a conditional modulus. This is an author-requested memorandum, not a commissioned journal report or editorial decision.
- [23] Independent AI referee-style memorandum, *A2 v12: independent harsh review and two-flight comparison*, September 10, 2026, repository `TrillionniumFoundation/theta-theory`, review commit `2ae2751f61224b66f314915fd5fc22f6321b606f`. Section 4 gives the all-order two-flight block and the direct fixed-window consequence. An author-requested independent AI memorandum, not a commissioned journal report or an editorial decision.
- [24] Independent AI-assisted referee-style memorandum, *Independent referee report on A2 v13 and analytic normal-form benchmark*, author-requested review, September 10, 2026, `TrillionniumFoundation/theta-theory`, review commit `b3f0ad5843782c221651c9189fc156d20865cab1`.
- [25] S. Sternberg, Local  $C^n$  transformations of the real line, *Duke Math. J.* **24** (1957), 97–102.
- [26] H. Eynard-Bontemps and A. Navas, *On the failure of linearization for germs of  $C^1$  hyperbolic vector fields in dimension one*, arXiv:2212.13646v2, November 5, 2023. <https://arxiv.org/abs/2212.13646v2>. The introduction recalls the classical one-dimensional smooth theorem and distinguishes it from the low-regularity failures studied there.
- [27] Independent AI-assisted referee-style memorandum, *A2 v14: independent review of the intrinsic boundary density*, with the companion *Scalar linearization and width benchmark*, author-requested review, September 11, 2026, `TrillionniumFoundation/theta-theory`, review commit `0c523413ef7ffa2dedd1ab471ce58b3275896a93`. This is an independent AI memorandum, not a journal-commissioned report or an editorial decision.
- [28] D. A. Levin and Y. Peres, with contributions by E. L. Wilmer, *Markov Chains and Mixing Times*, second edition, American Mathematical Society, 2017. <https://pages.uoregon.edu/dlevin/MARKOV/mcmt2e.pdf>.
- [29] Independent AI-assisted referee-style memorandum, *A2 v15: independent harsh review of scalar transport*, author-requested review, September 11, 2026, `TrillionniumFoundation/theta-theory`, review commit `a108adf17c4b9360e340a2d5708b20a375d69283`. The report records the exact strict-margin mixed-parameter benchmark. This is an independent memorandum, not a journal-commissioned report or an editorial decision.

- [30] M. Akahira, Asymptotic theory for estimation of location in non-regular cases, I: Order of convergence of consistent estimators, *Rep. Stat. Appl. Res., JUSE* **22** (1975), 8–26.
- [31] M. Akahira, Asymptotic theory for estimation of location in non-regular cases, II: Bounds of asymptotic distributions of consistent estimators, *Rep. Stat. Appl. Res., JUSE* **22** (1975), 99–115.
- [32] I. A. Ibragimov and R. Z. Has'minskii, *Statistical Estimation: Asymptotic Theory*, Springer, New York, 1981; English reprint 2013. doi:10.1007/978-1-4899-0027-2.
- [33] R. L. Smith, Maximum likelihood estimation in a class of nonregular cases, *Biometrika* **72** (1985), no. 1, 67–90. doi:10.1093/biomet/72.1.67.
- [34] K. Hirano and J. R. Porter, Asymptotic efficiency in parametric structural models with parameter-dependent support, *Econometrica* **71** (2003), no. 5, 1307–1338. doi:10.1111/1468-0262.00451.
- [35] Independent AI-assisted referee-style memorandum, *A2 v16: independent harsh review of adaptive physical experiments*, author-requested review, September 11, 2026, [TrillionniumFoundation/theta-theory](#), review commit `c5fed5340df3147f64cb3479cf1837357f73e7a8`. Section 4.4 records the random-hazard mixture clarification. This is an independent memorandum, not a journal-commissioned report or an editorial decision.
- [36] A. Meister and M. Reiss, *Asymptotic equivalence for nonparametric regression with non-regular errors*, arXiv:1101.5248v1, January 27, 2011. <https://arxiv.org/abs/1101.5248v1>.
- [37] L. N. Trefethen, Quantifying the ill-conditioning of analytic continuation, *BIT Numer. Math.* **60** (2020), 901–915. doi:10.1007/s10543-020-00802-7.
- [38] Independent AI-assisted referee-style memorandum, *A2 revision 48*, September 14, 2026, Section 4, repository [TrillionniumFoundation/theta-theory](#), review commit `a512f70ca77d711bd6c9ee61988f1e8dd4dba790`. An author-requested memorandum, not a commissioned journal report.
- [39] Independent AI-assisted referee-style memorandum, *A2 revision 49*, September 15, 2026, Section 4, repository [TrillionniumFoundation/theta-theory](#), review commit `5051b7789a424656a558179922607d8b55061b28`. An author-requested memorandum, not a commissioned journal report.
- [40] P. W. Holland and Y. J. Wang, *Dependence function for continuous bivariate densities*, *Comm. Statist. Theory Methods* **16** (1987), no. 3, 863–876. doi:10.1080/03610928708829408.
- [41] G. Osius, *Asymptotic inference for semiparametric association models*, *Ann. Statist.* **37** (2009), no. 1, 459–489. doi:10.1214/07-AOS572; arXiv:0903.0702v1.
- [42] *Independent referee-style report on A2, revision 51*, author-requested AI-assisted memorandum, September 15, 2026. Repository [TrillionniumFoundation/theta-theory](#), frozen review memorandum. Cited for the support-width centering observation; not a commissioned journal report or a mathematical certificate.
- [43] *Independent referee-style report on A2, revision 52*, author-requested AI-assisted memorandum, September 15, 2026. Repository [TrillionniumFoundation/theta-theory](#), frozen review memorandum. Cited for the quadratic functional small-window information benchmark; not a commissioned journal report or a mathematical certificate.
- [44] *Independent referee report on A2, revision 53*, author-requested AI-assisted memorandum, September 15, 2026, Sections 4–5. Repository [TrillionniumFoundation/theta-theory](#), frozen review memorandum. Cited for the Hellinger transfer and first-waiting-record comparison; not a commissioned journal evaluation or a mathematical certificate.
- [45] *Independent referee report on A2, revision 54*, author-requested AI-assisted memorandum, September 15, 2026, Section 5. Frozen first v54 memorandum. Cited for the count and bit reductions and the finite count-risk formula; not a commissioned journal evaluation.
- [46] *Second independent referee memorandum on A2, revision 54*, author-requested AI-assisted memorandum, September 15, 2026, Section 4. Frozen second v54 memorandum. Records an overlapping stopped-count comparison and its finite checks; not a commissioned journal evaluation or a separate priority claim.
- [47] *Independent referee-style report on A2, revision 57*, author-requested AI-assisted memorandum, September 15, 2026, Section R57-M3. Repository [TrillionniumFoundation/theta-theory](#), frozen review memorandum. Cited for the admissible highest-degree block bounds; not a commissioned journal report or an editorial decision.
- [48] *Independent referee report on A2, revision 59*, author-requested AI-assisted memorandum, September 16, 2026, Sections R59-M3 and R59-E2. Repository [TrillionniumFoundation/theta-theory](#), frozen memorandum. Cited for the qualitative Fredholm check and the suggestion to isolate the contracted-evaluation structure; not a commissioned journal report.
- [49] S. Zelditch, *Inverse spectral problem for analytic domains, II:  $\mathbb{Z}_2$ -symmetric domains*, *Ann. of Math. (2)* **170** (2009), no. 1, 205–269. doi:10.4007/annals.2009.170.205.

- [50] *Independent referee report on A2, revision 65*, author-requested AI-assisted memorandum, September 16, 2026, Section R65-M5. Repository `TrillionniumFoundation/theta-theory`, frozen memorandum. Cited for the Riccati derivative comparison, not for a global inverse or a commissioned journal decision.
- [51] *Independent referee report on A2, revision 66*, author-requested AI-assisted memorandum, September 16, 2026, Section R66-m1. Repository `TrillionniumFoundation/theta-theory`, frozen memorandum. Cited for the finite-iteration error-propagation control; not a commissioned journal report or a proof certificate.
- [52] E. Giné and A. Guillou, Rates of strong uniform consistency for multivariate kernel density estimators, *Ann. Inst. H. Poincaré Probab. Statist.* **38** (2002), 907–921. doi:10.1016/S0246-0203(02)01128-7.
- [53] L. Noakes and L. Stoyanov, *Lens rigidity in scattering by unions of strictly convex bodies in  $\mathbb{R}^2$* , *SIAM J. Math. Anal.* **52** (2020), 471–480. doi:10.1137/19M1270409. Corrected planar theorem: arXiv:1803.02542v2, Theorem 1.1.
- [54] *Independent referee report on A2, revision 72*, author-requested AI-assisted memorandum, September 17, 2026, Section R72-M5 and Section 7. Repository `TrillionniumFoundation/theta-theory`, frozen memorandum. Cited for the finite-flight clock identity and the request to display the weighted-order condition; not a commissioned journal report or a proof certificate.
- [55] K. Hamm and L. Huang, *CUR Decompositions, Approximations, and Perturbations*, arXiv:1903.09698v2, 2019. arXiv:1903.09698v2.
- [56] *Independent referee report on A2, revision 74*, author-requested AI-assisted memorandum, September 17, 2026, Section R74-M2. Repository `TrillionniumFoundation/theta-theory`, frozen memorandum. Cited for the affine gate-conditioning calculation, not a commissioned journal decision or a mathematical certificate.