

# GENERAL THETA FOUNDATIONS I: COMPATIBLE SIMPLICES AND EXACT STOCHASTIC WIDTH

QIAN QI

**ABSTRACT.** We study finite stochastic realizations of controlled orthogonal experiments. At minimal Hankel rank, equality of row spaces eliminates hidden future directions and identifies every state with a vertex of a compatible simplex. The simplex difference-body identity gives a quantitative occupation bound for minimal-width cuts, even with arbitrarily wider intervening registers. This yields an exact four-state frontier for a class of planar signed-permutation alphabets, including noncommuting ones. Exact minimal-width realization at every horizon is equivalent to a single invariant simplex and hence to a finite permutation action. We further compute a strict separation between word-distortion and enclosure profiles: the former vanish for the antipodal alphabet while minimal-simplex dilation is exactly the ambient dimension. A rational nine-dimensional orthogonal experiment with three command epochs has a cubic two-state optimum strictly above all separate flattening bounds. The full two-state sign-magnitude alternative and the prior arithmetic and finite-bit realization theorems are retained, with explicit certificate-size bounds and a direct comparison to tensor-completion geometry. The resource is nonuniform clocked atomic-row label width.

## 1. INTRODUCTION

A finite-state realization of a numerical experiment must preserve its response to every continuation of the external command word. A low-rank factorization at each cut is necessary, but need not admit transition rows between the cuts. This distinction persists even when the target is orthogonal, the state dimension is fixed, and every one-direction word-distortion profile vanishes. We study the resulting compatibility problem, with the number of available persistent labels as resource.

Fix orthogonal matrices  $U_a$  on  $\mathbb{R}^D$ , a finite seed set  $X$ , a signal  $0 < \rho < 1$ , and a terminal coordinate query  $j$ . A word  $w = a_1 \cdots a_N$  prescribes the binary law

$$(1.1) \quad \mathbb{P}(B = b \mid x, w, j) = \frac{1 + b\rho \langle e_j, U_w x \rangle}{2}, \quad U_w = U_{a_N} \cdots U_{a_1}, \quad b \in \{-1, 1\}.$$

The command word is supplied externally. There is no distinguished probability measure on words. A realization may depend on the horizon and epoch, but all persistent information, including any random selector, must fit in its current label. Its rows may be real and are sampled atomically. This is a nonuniform positive-realization invariant, not ordinary workspace or program length. Definition 2.1 specifies it without an implicit uncharged register.

The principal result identifies a regime in which the observed response space itself is a complete state geometry. Take  $X = \{\pm e_i\}_{i=1}^D$  and allow all coordinate queries. If the identity is a command, every normalized cut Hankel matrix has rank  $D + 1$ . At a cut with exactly that many labels, equality of ranks forces every label's entire future row into the target affine response space. The row normalization then turns a linear-space equality into an affine one. Thus no hidden future coordinate remains, even in a machine with wider cuts before and after it. Exact realization at rank-tight cuts is equivalent to the transport of enclosing simplices by all intervening words

(Theorem 4.1). This conclusion is not asserted for a larger factorization, where hidden future directions can persist.

The simplex description has a quantitative consequence not available from separate cut ranks. Put  $c_D = 2^{-D} \binom{2D}{D}$ . When  $I$  and  $-I$  are commands, any two consecutive rank-tight cuts force the later simplex to contain both signs of the earlier one. The elementary simplex difference-body identity gives a volume factor  $c_D > 1$  for  $D \geq 2$ . Therefore every exact machine has at most

$$(1.2) \quad 1 + \frac{\log(D!\rho^{-D})}{\log c_D}$$

rank-tight cuts, irrespective of the widths between them. This is an obstruction to an entire realization class, rather than a verification of a proposed set of rows. In dimension two it gives an exact frontier: for any signed-permutation alphabet containing  $I, -I$ ,

$$(3/2)^N > 2\rho^{-2} \implies W_{N,0} = 4.$$

The upper bound is a common four-label permutation realization. Such alphabets may be noncommutative. At each fixed qualifying horizon the frontier also persists at sufficiently small positive error. We give no numerical lower bound for that last error interval.

An all-horizon statement strengthens this finite obstruction. Exact rank-tight realization at every horizon, even with complete nonuniform redesign, is equivalent to one invariant full-dimensional simplex containing  $\rho \operatorname{conv} X$  and lying in the coordinate cube. The same machine then works at every horizon, its command rows are permutations, and the generated orthogonal group embeds in  $\mathfrak{S}_{D+1}$  (Theorem 4.5). In the plane and at signal at most  $1/2$ , this characterizes all-horizon three-state feasibility by containment in the symmetry group of an equilateral triangle.

These results also connect the finite theory to the word profiles used in the arithmetic part of the article. For the alphabet  $\{I, -I\}$ , every one-direction orbit has only two points. Hence  $\Gamma_{\mathcal{A}}(k, B) = 0$  and  $\mathfrak{r}_{\mathcal{A}}(k) = 0$  for every  $k \geq 2$ . Nevertheless the common enclosure profile satisfies

$$\lambda_{\mathcal{A}}(D+1, a) = D \quad (0 < a \leq 1/D).$$

Theorem 5.1 proves this equality by barycentric coordinates. Thus the distortion–dilation inequality is genuinely strict, while compatibility has a positive occupation charge. The arithmetic results concern regimes where distortion and enclosure scales can match; they are not consequences of a putative equality of these two invariants in general. This distinction supplies the conceptual link between the finite and asymptotic parts of the paper.

At two states a different geometry is available. For paired seeds and binary responses, the odd part of a legal machine is a bounded rank-one tensor. The complete sign–magnitude alternative retained in Theorems 9.2 and 9.4 treats explicitly listed tables, including zeros, by support reduction and one logarithmic linear system per sign chamber. Integer multiplicative alternatives exclude all machines in that class. Proposition 11.1 prices the chamber count, sparse rays, boundary polynomials and algebraic minimizing rows; finiteness is not confused with polynomial complexity in a succinct horizon.

We give that tensor theory a genuinely multi-epoch orthogonal application. Two rational noncommuting permutation matrices in dimension nine, rational unit seeds, three command epochs and one selected coordinate have exact two-state error

$$\frac{100\rho}{357}t, \quad 500t^3 - 375t^2 + 540t - 124 = 0, \quad 0.26 < t < 0.261.$$

A five-entry Segre circuit proves the lower bound, and legal stochastic rows attain it; these rows are algebraic when  $\rho$  is algebraic. Every separate flattening of the positive three-mode response admits strictly smaller rank-one error. The obstruction therefore requires compatibility across modes, not a single matrix relaxation. This experiment has only its designated query; it is not

an assertion about all nine coordinate queries. The earlier arbitrary-table cubic and commuting orthogonal examples are preserved with their distinct scopes.

Invariant positive cones, tensor exponent relations, monomial linearization, Farkas' lemma and simplex difference-body identities are classical tools. We distinguish their use from the realization results proved here. The comparison in Section 11 includes positive realization [12], matrix infinity-norm approximation [22], rank-one tensor completion [27] and geometric programming [24]. In particular, the cubic circuit itself is not claimed as new. The new conclusions concern rank-tight causal geometry, its occupation and all-horizon consequences, the exact profile separation, and the orthogonal realization of a sharp multi-epoch interval obstruction.

The subsequent sections retain the full two-state, positive-Hankel, Gram, moment, arithmetic and finite-bit theorems and their proofs. The finite-bit implementation has its own resource accounting; it does not make exact algebraic sampling free in a uniform model. The source-bound checks certify specified finite identities and build reproducibility, not the universal proofs or independent priority. The separate analytic dependencies of the wider Theta program are neither assumed nor declared closed by the present finite-realization results.

## 2. THE NUMERICAL EXPERIMENT AND THE REGISTER

Let  $V = \mathbb{R}^D$ , let  $\mathcal{A}$  be a nonempty finite alphabet, and fix orthogonal matrices  $U_a$ ,  $a \in \mathcal{A}$ . Write  $U_w = U_{a_N} \cdots U_{a_1}$  in the order of execution. Let  $X \subset \mathbb{S}^{D-1}$  be a nonempty finite seed set and  $0 < \rho < 1$ . The required law is (1.1). Commands are externally supplied, not chosen by a policy inside the simulator. No distribution on command words is part of the specification. Unless a theorem explicitly restricts the query set, every coordinate query is available. The calibration lemma below uses that full query set; the selected-query example does not invoke it.

**Definition 2.1** (Clocked atomic-row realization). At horizon  $N$ , a machine consists of finite sets  $S_t$  with  $|S_t| = K_t$  for  $0 \leq t \leq N$ , an initialization probability row  $E_x$  on  $S_0$ , stochastic matrices  $T_{t,a} : S_{t-1} \rightarrow S_t$ , and terminal columns  $d_j : S_N \rightarrow [-1, 1]$ . Its numerical mean is

$$E_x T_{1,a_1} \cdots T_{N,a_N} d_j.$$

After the atomic query  $j$ , it draws and holds the binary answer with that mean. The final answer cut has two labels. Its peak is  $\max(2, K_0, \dots, K_N)$ . Available labels count even when unreachable; adding unused normalized rows is permitted.

All persistent randomness, input history and protocol distinctions must be in the current label. Only the external epoch, the prescribed horizon and the read-only row tables are free. Table construction, lookup, arithmetic and exact sampling of known real stochastic rows are uncharged. Matrices, rows and state sets may depend on  $N$ ; this is nonuniform positive-realization width, not uniform computational space. A common-row construction uses the same command matrices at all command epochs but may still have an  $N$ -dependent decoder. It need not answer at earlier stopping lengths.

Let  $W_{N,\varepsilon}$  be the least peak among machines whose every binary row differs from (1.1) by at most  $\varepsilon$  in total variation. In particular  $W_{N,\varepsilon} \geq 2$ . For fixed finite  $K_\bullet$  the optimal actual error  $\mathcal{E}_N(K_\bullet)$  is attained: the row parameters lie in a finite product of compact sets and the error is a finite maximum of continuous functions. For planar results we use  $X = \{\pm e_1, \pm e_2\}$  and include the identity letter. For the retained planar arithmetic results we keep  $\rho \leq 1/10$  and  $\varepsilon < \rho/(2\sqrt{2})$ .

**Lemma 2.2** (Calibration under any external test law). *Fix a seed  $x \in X$  and any law on the complete command word, independent of private initialization. Let  $Y_t = U_{a_t} \cdots U_{a_1} x$ ,  $Z_t = \mathbb{E}[Y_t \mid S_t]$  and  $q_t = \mathbb{E}\|Z_t\|$ . Every machine with row error at most  $\varepsilon$  satisfies*

$$(2.1) \quad q_N \geq \kappa := \rho/\sqrt{D} - 2\varepsilon.$$

*At a boundary preceding an independent complete packet, the next boundary admits the conditional-mean identity used in Lemma 15.2. No single-step independence inside a packet is asserted.*

*Proof.* For a fixed word, every coordinate of  $\mathbb{E}[d(S_N) \mid w] - \rho Y_N$  has modulus at most  $2\varepsilon$ . Since  $\|Y_N\| = 1$ , averaging gives

$$\mathbb{E}\langle Y_N, d(S_N) \rangle \geq \rho - 2\varepsilon\sqrt{D}.$$

On the other hand this pairing equals  $\mathbb{E}\langle Z_N, d(S_N) \rangle$  and is at most  $\sqrt{D}q_N$ . This proves (2.1). Combining terminal columns is a proof device, not a requirement to provide their joint output distribution.  $\square$

### 3. CONTROLLED HANKEL ARRAYS AND COMPATIBLE POSITIVE REALIZATIONS

For a horizon  $N$ , define the probability-valued prefix-suffix matrix

$$(3.1) \quad H_t[(x, u), (v, j, b)] = \frac{1 + b\rho\langle e_j, U_v U_u x \rangle}{2}, \quad |u| = t, \quad |v| = N - t.$$

The two columns indexed by  $b = \pm 1$  sum to one for every fixed  $(v, j)$ . Let  $\mathcal{C}_t$  be the product of these binary probability simplices. Thus a point of  $\mathcal{C}_t$  is a legal collection of responses to all remaining words and queries, not a joint distribution of those answers. The restriction  $L_{t,a} : \mathcal{C}_t \rightarrow \mathcal{C}_{t+1}$  reads the coordinates whose first remaining command is  $a$ .

A normalized positive factorization is  $H_t = E_t B_t$ , where  $E_t$  is row stochastic and every row of  $B_t$  belongs to  $\mathcal{C}_t$ . Its least inner dimension is denoted  $\text{rank}_{\text{st}} H_t$ . The usual nonnegative rank permits arbitrary nonnegative factors and need not impose these simultaneous normalizations. Consequently

$$(3.2) \quad \text{rank } H_t \leq \text{rank}_+ H_t \leq \text{rank}_{\text{st}} H_t \leq K_t$$

for every exact stochastic realization. The last inequality follows by conditioning the remaining experiment on the actual label at cut  $t$ . The factor  $E_t$  contains its state probabilities;  $B_t$  contains its complete future responses, including those of states not reached by a particular prefix.

**Proposition 3.1** (Finite-horizon positive Hankel normal form). *For any finite table of binary responses, a profile  $K_0, \dots, K_N$  is exactly realizable if and only if there are polytopes  $P_t \subset \mathcal{C}_t$ , generated by at most  $K_t$  points, such that*

$$(3.3) \quad H_0[x, \cdot] \in P_0 \quad (x \in X), \quad L_{t,a} P_t \subset P_{t+1} \quad (0 \leq t < N, a \in \mathcal{A}).$$

*The held answer additionally costs two labels. The optimum error for a fixed profile is the least entrywise error from the specified table to one having such a compatible family.*

*Proof.* Given a machine, let  $b_{t,s} \in \mathcal{C}_t$  be the conditional response to every remaining word when its current label is  $s$ . Its stochastic update gives

$$L_{t,a} b_{t,s} = \sum_{s'} T_{t+1,a}(s, s') b_{t+1,s'}.$$

Taking convex hulls and using the initialization proves necessity. Conversely choose generators  $b_{t,s}$ , padding by repetitions when necessary. Express each  $L_{t,a} b_{t,s}$  as a convex combination of the next generators, and use those coefficients as the stochastic row. Express each root row as a convex combination of the initial generators. At time  $N$  the generators themselves give the binary decoder rows. Backward induction on the remaining word proves equality of all responses. The construction is valid on every generator, not only on reachable mixtures. For approximation apply this equivalence to the actual response table of a machine, or to a candidate perturbed table; finite compactness of the stochastic parameter spaces gives attainment.  $\square$

This is the finite, controlled, normalized version of the classical shift-invariant-cone realization principle, not a newly discovered characterization of positive systems. In particular, the cone generated by  $(1, b)$  for  $b \in P_t$  carries the positive restriction maps. Benvenuti–Farina’s Theorems 2–3 give the stationary cone formulation; the multiple cuts here are essential [12]. The same compatibility issue was already explicit in the earlier lift formulation [18]. The normal form has exponentially many suffix coordinates. It is exact, but does not make arbitrary width minimization efficient or eliminate hidden directions by projecting them into the observed  $D$ -space.

**Proposition 3.2** (A finite compatibility certificate). *For proposed generators  $B_t, B_{t+1}$ , put  $c = L_{t,a}b_{t,s}$  and let  $b'_1, \dots, b'_{K_{t+1}}$  be the next generators. The least rowwise uniform defect is*

$$(3.4) \quad \min_{\theta \in \Delta_{K_{t+1}}} \left\| c - \sum_i \theta_i b'_i \right\|_\infty = \max_{\|z\|_1 \leq 1} \left( \langle z, c \rangle - \max_i \langle z, b'_i \rangle \right).$$

*It vanishes exactly when that row has a stochastic extension. If root rows are approximated within  $\delta_0$  and every cut restriction within  $\delta_t$ , the resulting machine has worst-word binary-TV error at most  $\delta_0 + \sum_{t=1}^N \delta_t$ .*

*Proof.* Write the infinity norm as the supremum of its pairing with the unit  $\ell_1$  ball. The simplex and this ball are compact and convex, so finite bilinear minimax interchanges the extrema. Minimizing the pairing over the simplex gives the maximum over its vertices with the displayed sign. For composition, restriction does not increase the infinity norm and stochastic averaging is a contraction. Induction on the suffix length therefore adds the successive errors. Each binary pair is normalized, so its total variation is exactly either coordinate error. Cancellation may reduce the true error; no equality with the local sum is asserted.  $\square$

**Proposition 3.3** (Constant separate ranks for the planar experiment). *For the four planar seeds, an identity command, and  $0 < \rho \leq 1/10$ , at every cut and every horizon,*

$$(3.5) \quad \text{rank } H_t = \text{rank}_+ H_t = \text{rank}_{\text{st}} H_t = 3.$$

*Each factorization embeds in a complete machine when all other cuts are unrestricted. Nevertheless the compatible peak can grow as in Theorems 13.1 and 13.2.*

*Proof.* The entries in (3.1) are affine functions of the two coordinates of  $\rho U_u x$ , so ordinary rank is at most three. Identity prefixes, identity suffixes, both coordinate queries and the four seeds give a rank-three submatrix. For the upper factorization, an equilateral triangle of circumradius  $2\rho$  contains the disk of radius  $\rho$ . Encode  $\rho U_u x$  in its three vertices  $v_i$  by barycentric coordinates. The continuation row at  $v_i$  is  $(1 + b\langle e_j, U_v v_i \rangle)/2$ , which is legal since  $\|v_i\| = 2\rho < 1$ . This is a normalized three-factor decomposition and proves the equalities. To embed it, retain the exact history until this cut, draw its factor label, retain that label and the remaining history afterwards, then use the factor decoder. Both retained histories are finite at a fixed horizon. They are not free registers in a simultaneous three-state realization. The growth assertions follow from the stated all-hidden converses, not from a claimed new separation of signed and positive realization in classical stationary theory.  $\square$

#### 4. RANK-TIGHT COMPATIBILITY AND SIMPLEX GEOMETRY

Throughout this section  $D \geq 2$ ,  $X = \{\pm e_1, \dots, \pm e_D\}$ , all coordinate queries are available, and the orthogonal alphabet contains  $I$ . The signal satisfies  $0 < \rho < 1$ . The conclusions concern exact realization unless a positive tolerance is explicitly specified. Put

$$C_\rho = \rho \text{conv } X, \quad Q_t = \{z \in \mathbb{R}^D : |\langle e_j, U_v z \rangle| \leq 1 \text{ for } |v| = N - t, 1 \leq j \leq D\}.$$

Identity padding gives  $Q_t \subset [-1, 1]^D$ . Define the injective affine map  $J_t : Q_t \rightarrow \mathcal{C}_t$  by

$$J_t(z)_{v,j,b} = (1 + b\langle e_j, U_v z \rangle)/2.$$

The spaces  $\mathcal{C}_t$ , the restrictions  $L_{t,a}$  and the normalized Hankel arrays are those of Section 3. In particular,  $L_{t,a}J_t(z) = J_{t+1}(U_a z)$ .

**Theorem 4.1** (Rank-tight simplex equivalence). *An exact clocked realization with at most  $D + 1$  labels at every command cut exists if and only if there are full-dimensional simplices  $S_t \subset Q_t$ , each with  $D + 1$  vertices, such that*

$$(4.1) \quad C_\rho \subset S_0, \quad U_a S_t \subset S_{t+1} \quad (0 \leq t < N, a \in \mathcal{A}).$$

*More generally, at every cut with  $D + 1$  labels in an arbitrary exact machine, every label's full future-response row is  $J_t(z)$  for a unique  $z \in Q_t$ . If  $s < t$  are two such cuts, their simplices satisfy  $U_w S_s \subset S_t$  for every word of length  $t - s$ , irrespective of all intervening widths.*

*Proof.* Every  $H_t$  has rank  $D + 1$ . Its rows are affine functions of  $\rho U_u x$ , so its rank is at most  $D + 1$ . Identity prefixes and suffixes, the signed coordinate seeds, and the coordinate queries give rank  $D + 1$ . At a cut of this width the machine factorization  $H_t = E_t B_t$  has  $E_t$  of full column rank and  $B_t$  of full row rank. Hence  $\text{rowspan } B_t = \text{rowspan } H_t$ . Write a row of  $B_t$  as a linear combination of rows of  $H_t$ . Normalization of any binary column pair forces the sum of its coefficients to be one. It is consequently an affine combination and equals  $J_t(z)$ . Identity suffixes give uniqueness of  $z$ ; legality gives  $z \in Q_t$ . The  $D + 1$  rows are linearly independent, so their  $z$ 's are affinely independent. The initialization contains  $C_\rho$  in their convex hull. Composing the actual transition rows between two narrow cuts gives

$$L_w B_s = T_w B_t.$$

Injectivity of  $J_t$  gives the asserted containment for every generator, including generators unreachable from a particular seed. No observable projection of a higher-width row has been used.

Conversely express every  $\rho x$  in the vertices of  $S_0$ , and every  $U_a v$  in the vertices of  $S_{t+1}$ , by barycentric coordinates. They are nonnegative and sum to one, hence are legal initialization and transition rows. Decode coordinate  $j$  by the  $j$ th coordinates of the terminal vertices. These lie in  $[-1, 1]$ . Induction gives the conditional mean  $\rho U_w x$  and all required binary laws. The extra held-answer cut has two labels. Rank already excludes fewer than  $D + 1$  labels at any command cut.  $\square$

The rank hypothesis is decisive. For a larger factorization,  $\text{rowspan } B_t$  may strictly contain that of  $H_t$ , and the preceding elimination of hidden future directions is invalid. The underlying cone/factorization principle is classical [12]; the rank-tight conclusion identifies precisely when the observed-space simplex description is necessary as well as sufficient in this nonuniform controlled problem.

**4.1. A finite occupation certificate.** Let

$$c_D = 2^{-D} \binom{2D}{D} > 1 \quad (D \geq 2).$$

We record the elementary simplex case of the difference-body identity; this is classical convex geometry, not a new volume inequality.

**Lemma 4.2** (Difference body of a simplex). *For a full-dimensional simplex  $S$ ,  $\text{vol}(S - S) = \binom{2D}{D} \text{vol}(S)$ . If a convex body  $T$  contains both  $S$  and  $-S$ , then  $\text{vol}(T) \geq c_D \text{vol}(S)$ .*

*Proof.* Affine covariance reduces the identity to  $\Delta = \{x \geq 0 : \sum x_i \leq 1\}$ . Its difference body is given by  $\sum x_i^+ \leq 1$  and  $\sum x_i^- \leq 1$ . An orthant with  $k$  positive coordinates contributes  $1/(k!(D-k)!)$ . Summing its  $\binom{D}{k}$  occurrences and using  $\sum_k \binom{D}{k}^2 = \binom{2D}{D}$  proves the identity, since  $\text{vol}(\Delta) = 1/D!$ . The containment  $(S - S)/2 \subset \text{conv}(S \cup (-S)) \subset T$  proves the inequality.  $\square$

**Theorem 4.3** (Rank-tight occupation budget). *Suppose  $I, -I \in \mathcal{A}$ . In every exact realization, let  $m$  be the number of command cuts whose available width is  $D + 1$ . If  $m \geq 1$ , then*

$$(4.2) \quad c_D^{m-1} \leq D! \rho^{-D}.$$

*All other cuts may have arbitrary finite widths. In particular, an all- $(D + 1)$  profile at horizon  $N$  requires  $c_D^N \leq D! \rho^{-D}$ .*

*Proof.* Order the narrow cuts. Between consecutive ones the identity word and a word with exactly one  $-I$  exist. Theorem 4.1 therefore forces the next simplex to contain both signs of the preceding one. Lemma 4.2 multiplies volume by at least  $c_D$ . At the first narrow cut identity prefixes give  $C_\rho \subset S$ , so its volume is at least  $(2\rho)^D/D!$ . At the last narrow cut identity suffixes give  $S \subset [-1, 1]^D$ , of volume  $2^D$ . Comparing these bounds proves the result. Intermediate registers and their private randomness have disappeared only through their actual composed stochastic rows, not through an uncharged history register.  $\square$

**Corollary 4.4** (An exact four-state frontier). *In dimension two let  $\mathcal{A}$  be any alphabet of signed permutation matrices containing  $I, -I$ . For every  $0 < \rho < 1$  and every horizon with*

$$(4.3) \quad (3/2)^N > 2\rho^{-2},$$

*one has  $W_{N,0} = 4$ . For  $\rho = 1/10$ , all  $N \geq 14$  satisfy this condition. The bound 14 is sufficient; it is not asserted to be the first horizon requiring four labels. For each horizon satisfying (4.3), four labels remain optimal throughout some interval  $0 \leq \varepsilon < \eta_N$ , where  $\eta_N > 0$ .*

*Proof.* Three labels at every cut violate Theorem 4.3; fewer labels are excluded by Hankel rank. Four labels suffice by storing  $x \in \{\pm e_1, \pm e_2\}$  exactly, permuting that label under each command, and decoding by  $\rho x_j$ . This is a common-row construction with no horizon-dependent transition. The integer comparison  $3^{14} > 200 \cdot 2^{14}$  proves the numerical assertion. At fixed  $N$  the space of machines with at most three labels is represented by padding to exactly three and is compact. The continuous worst-row error has a strictly positive minimum  $\eta_N$ , since zero is excluded. This gives the positive-tolerance assertion without assigning an unproved numerical value to  $\eta_N$ .  $\square$

For example, adjoining the swap matrix and  $\text{diag}(1, -1)$  gives genuinely noncommuting commands; their products in opposite orders differ. This is a class-wide causal occupation obstruction beyond width two, followed by an attaining machine, not a test of supplied transition rows. The exact occupation constant need not be sharp.

#### 4.2. From arbitrary horizons to a single permutation action.

**Theorem 4.5** (All-horizon simplex rigidity). *Under the hypotheses at the start of this section, the following are equivalent: every finite horizon admits an exact realization of peak  $D + 1$ ; there is a single full-dimensional simplex  $C_\rho \subset S \subset [-1, 1]^D$  with  $U_a S = S$  for all letters. In the latter case one realization works for all horizons and all command updates are permutations of its  $D + 1$  vertices. The group generated by the commands is finite and embeds in the symmetric group on  $D + 1$  letters.*

*Proof.* Take simplex chains at successively increasing horizons. At every cut identity prefixes and suffixes place their vertices in the compact cube and their convex hulls around  $C_\rho$ . A diagonal subsequence of vertex tuples produces an infinite chain  $S_0, S_1, \dots$  satisfying (4.1). The enclosed crosspolytope prevents degeneracy. Since  $I$  is a letter, this chain is nested. Its Hausdorff limit

$S_\infty = \overline{\bigcup_t S_t}$  is again a simplex with at most  $D+1$  vertices: choose a convergent subsequence of the bounded vertex tuples. It remains full-dimensional. Continuity gives  $U_a S_\infty \subset S_\infty$ . Orthogonality preserves volume, so this containment is equality. Indeed, a proper inclusion of full-dimensional compact convex bodies strictly decreases volume. Each  $U_a$  therefore permutes the vertices. An affine map fixing  $D+1$  independent vertices is the identity, which proves faithfulness and finiteness. Conversely barycentric initialization, the vertex permutations, and coordinate decoders on  $S$  give one exact machine at every horizon. All its suffix responses are legal because every group element preserves  $S \subset [-1, 1]^D$ .  $\square$

In the plane, for  $\rho \leq 1/2$ , this yields a complete criterion for three-state feasibility at all horizons: the generated group must be contained in the symmetry group of some equilateral triangle. Necessity follows by classifying the orthogonal symmetries of a triangle: a nontrivial rotation has order three and forces an equilateral triangle; without such a rotation the group has at most one reflection. For sufficiency a regular triangle of circumradius one and the appropriate orientation contains the disk of radius  $1/2$ . In particular, order-three rotations together with a compatible reflection give a noncommutative three-state class. The theorem does not assert all-width classification.

## 5. AN EXACT SEPARATION OF DISTORTION AND ENCLOSURE

We now link the rank-tight theorem to the inherited word profiles. The obstruction is not a better choice of a packet law: it is information that every one-direction packet law omits.

**Theorem 5.1** (Strict distortion–dilation separation). *For  $\mathcal{A} = \{I, -I\}$  in dimension  $D \geq 2$ ,*

$$\Gamma_{\mathcal{A}}(k, B) = 0 \quad (k \geq 2, B \geq 1), \quad \mathfrak{r}_{\mathcal{A}}(k) = 0 \quad (k \geq 2).$$

*Nevertheless, for every  $0 < a \leq 1/D$  the common enclosure profile is*

$$(5.1) \quad \lambda_{\mathcal{A}}(D+1, a) = D.$$

*Thus the inequality in Theorem 16.1 is strictly unequal:  $0 < \log D$ . At rank-tight cuts the zero packet charges coexist with the positive occupation charge  $\log c_D$  from Theorem 4.3.*

*Proof.* For each unit vector  $u$ , every word sends  $u$  to one of  $u, -u$ . Taking these two vectors as centers makes every scalar product maximum one, for every word law. The definitions of  $D_k$ ,  $\Gamma$  and  $\mathfrak{r}$  give their asserted values. A full-dimensional polytope with at most  $D+1$  vertices is a simplex. Let its vertices be  $v_i$ , and write the barycentric coordinates of zero as  $c_i > 0$ ,  $\sum c_i = 1$ . If  $-S \subset \lambda S$ , uniqueness of barycentric coordinates for  $-v_i/\lambda$  forces

$$(1 + \lambda^{-1})c_i - \lambda^{-1} \geq 0, \quad \text{hence} \quad c_i \geq (1 + \lambda)^{-1}.$$

Summing over  $D+1$  vertices yields  $\lambda \geq D$ . A regular simplex centered at zero with circumradius one has inradius  $1/D$  and obeys  $-S \subset DS$  by the same coordinates  $c_i = 1/(D+1)$ . It is feasible for every stated  $a$ , proving equality. The occupation comparison follows from the previous theorem, not from an invalid reversal of the packet bound.  $\square$

This exact separation preserves rather than replaces the arithmetic results below. They use situations where distortion is positive and can match enclosure scales. Here the two mechanisms are provably inequivalent, and rank-tight Hankel rigidity supplies the missing finite-dimensional geometric obstruction. The value  $D$  is the classical simplex asymmetry calculation; its role here is the exact comparison with the optimized controlled-word profile.

## 6. A THREE-EPOCH ORTHOGONAL CIRCUIT BEYOND SEPARATE FLATTENINGS

The cubic table in Theorem 10.2 can be realized by a noncommuting rational orthogonal alphabet with three actual command epochs. This is a selected-query experiment: only the coordinate specified below is available, not all coordinates of the ambient space. Its dimension is nine. No fixed-dimensional universality claim is made.

Index eight orthonormal vectors by  $z \in \{0, 1\}^3$  and a ninth by  $*$ . For  $a \in \{0, 1\}$  let the permutation matrix  $U_a$  fix  $e_*$  and satisfy

$$(6.1) \quad U_a e_{(z_1, z_2, z_3)} = e_{(z_2, z_3, z_1 \oplus a)}.$$

These are rational orthogonal matrices. They do not commute:  $U_0 U_1 e_{000} = e_{010}$ , whereas  $U_1 U_0 e_{000} = e_{001}$ . Put  $F_z = 1/10, 2/5, 4/5, 2/5$  according as  $|z| = 0, 1, 2, 3$ , and set

$$(6.2) \quad c = \frac{200}{357}, \quad x_z = cF_z, \quad x_* = -\frac{157}{357}.$$

The seed set is  $\{x, -x\}$ , the horizon is three, and the sole available query is  $e_{000}$ . The target mean is  $\rho \langle e_{000}, U_{a_3} U_{a_2} U_{a_1} (\pm x) \rangle$ .

**Theorem 6.1** (An orthogonal cubic optimum). *The vector  $x$  is a rational unit vector. For every  $0 < \rho \leq 1$ , the optimum binary-TV error among machines with two labels at all four command cuts is*

$$(6.3) \quad \frac{100\rho}{357} t, \quad 500t^3 - 375t^2 + 540t - 124 = 0, \quad \frac{13}{50} < t < \frac{261}{1000}.$$

*Legal stochastic rows attain it; they may be chosen algebraic whenever  $\rho$  is algebraic. When  $\rho$  is rational, no machine whose initialization, transitions and decoder are all rational attains it, although rational machines approach it arbitrarily closely. Every separate matricization of its positive-seed three-mode response tensor admits a bounded rank-one approximation with mean error at most  $c\rho/5$ , strictly smaller than the true mean optimum  $c\rho$ . Thus separate flattening bounds cannot certify this optimum.*

*Proof.* One has  $\sum_z F_z^2 = 257/100$ , and

$$\sum_z x_z^2 + x_*^2 = \frac{40000(257/100) + 157^2}{357^2} = 1.$$

Three applications of (6.1) send  $e_z$  to  $e_{z \oplus (a_1, a_2, a_3)}$ . Consequently the positive-seed mean is exactly  $c\rho F_{a_1 a_2 a_3}$ . The scalar odd normal form, proved in Theorem 8.1, makes every legal two-state odd response a bounded decomposable three-mode tensor. The seed and decoder scalars can be absorbed into one mode, so the lower bound applies to every machine, not merely to centered transition rows.

For  $\lambda = c\rho$  the five entries 000, 011, 101, 110, 111 give the identity

$$g_{011} g_{101} g_{110} = g_{000} g_{111}^2.$$

After replacing factors by their absolute values, a mean tolerance  $\delta < 4\lambda/5$  must therefore obey the scaled inequality

$$(4\lambda/5 - \delta)^3 \leq (\lambda/10 + \delta)(2\lambda/5 + \delta)^2.$$

Theorem 10.2 implies  $\delta \geq \lambda t$ ; larger  $\delta$  already satisfy this bound. For attainment take  $u_0 = (1/10 + t)^{1/3}$  and  $u_1 = (2/5 + t)^{1/3}$ , always the positive real cube roots. Initialize with signed coordinate  $\pm\lambda$ , use the slope row  $T(u_a)$  at each of the three epochs, and use terminal decoder  $(1, -1)$ . The resulting response is  $\pm\lambda \prod_{r=1}^3 u_{a_r}$ . Its maximum mean error is  $\lambda t$  by the cubic identity and the weight-one estimate of Theorem 10.2. All factors lie in  $[0, 1]$ .

The integer cubic is primitive and irreducible modulo seven, as checked in the cited theorem. Rational  $\rho > 0$  therefore gives an irrational optimum. Every all-rational finite machine has

rational responses and rational maximum error against this rational target, proving nonattainment. Continuity under rational approximation of finitely many stochastic rows proves approximation from above.

Finally, by symmetry each nontrivial matricization of  $F$  is, up to row/column permutations or transpose,

$$\begin{pmatrix} 1/10 & 2/5 & 2/5 & 4/5 \\ 2/5 & 4/5 & 4/5 & 2/5 \end{pmatrix}.$$

The matrix with both rows  $(1/4, 3/5, 3/5, 3/5)$  has rank one, factors bounded by one, and maximum entry error  $1/5$ . Scaling by  $\lambda$  preserves boundedness. Since  $t > 13/50 > 1/5$ , the claimed strict gap holds simultaneously for every separate flattening. The flattening approximants need not arise from one common tensor; requiring that compatibility is exactly the additional constraint.  $\square$

The displayed monomial relation is a Segre lattice circuit, not a new algebraic identity. The conclusion is stronger than observing that a particular determinant is nonzero: it computes the exact error of a complete, noncommuting, multi-epoch orthogonal experiment while all its individual matrix relaxations have smaller error. The stochastic compiler is essential to its upper bound.

**Proposition 6.2** (Finite word-table embedding). *For any rational table  $F : \{0, 1\}^n \rightarrow \mathbb{Q}$ ,  $n \geq 1$ , there are rational orthogonal commands on  $\mathbb{R}^{2^n+1}$ , rational unit seeds  $\pm x$ , one coordinate query and  $n$  epochs whose positive-seed response is  $\rho c F_w$ , with  $c = 2/(1 + \sum_w F_w^2)$ . The construction is explicit.*

*Proof.* Use the cyclic shift-and-xor permutation on  $n$  bits, fixing an extra coordinate. Put  $x_w = c F_w$  and  $x_* = (1 - \sum_w F_w^2)/(1 + \sum_w F_w^2)$ . Direct expansion gives norm one. After  $n$  commands the permutation is xor by the command word, so the zero coordinate reads  $x_w$ . In particular every target coordinate is bounded by one, without any assumption on the size of  $F_w$ . For  $n = 1$  no noncommutativity is asserted.  $\square$

The dimension is exponential in  $n$  and is explicitly part of the input. This embedding transfers finite-table examples; it is not an efficient construction for a succinct long-horizon process.

## 7. THE TWO-STATE STRUCTURAL ALTERNATIVE

A certificate that a supplied stochastic machine has the correct response is not a certificate that a smaller machine cannot have that response, or that its approximation error is optimal. This distinction is particularly important for a finite experiment with externally chosen commands. Each internal state must support all possible future commands, and a random mixture of whole machines is a larger machine unless its selector is retained in the charged state.

The two-state component develops an exact global error alternative in the smallest nontrivial register class. For antisymmetric binary targets, two labels at every cut reduce without loss to a bounded decomposable tensor. The seed, each command epoch, and the terminal query are its modes. The scalar reduction and its parity obstruction are retained from the preceding revision; their proofs appear in Section 8. The present question is how to determine the distance when all signs are consistent but magnitudes are not.

**Theorem 7.1** (Sign-magnitude certification). *For every finite antisymmetric binary-response table, the optimal two-label uniform error has the following complete description. At each proposed mean tolerance, threshold-active entries determine all factor coordinates which must be nonzero. A finite sign system and one bounded logarithmic linear system per sign chamber decide feasibility. Each rejected chamber has either an impossible positive endpoint or a nonnegative integer multiplicative certificate on at most  $v + 1$  inequalities, where  $v = |V_\delta|$  is the number of essential factor coordinates at the tested threshold  $\delta$ . The sparsity bound is per rejected chamber. For real-algebraic table*

entries, univariate root isolation applied to these certificates determines the optimum and an algebraic stochastic machine attaining it.

Theorem 7.1 is proved in Theorems 9.2 and 9.4. Its quantifiers are global over the prescribed two-state class. It does not require a supplied factorization, a strictly positive target, or a positive optimum. It is not a claim that the union of chambers is convex or that searching them is polynomial in the word horizon. The lower certificates are particularly transparent: weighted multiplication of interval constraints cancels every unknown factor and produces a violated inequality between data.

The exact example in Theorem 10.2 has two different commands, one command epoch, two positive seeds and their negatives, and two queries. Its positive-seed table has rational means  $1/10$ ,  $2/5$ , and  $4/5$ . Its optimal binary-TV error is  $t/2$ , where

$$500t^3 - 375t^2 + 540t - 124 = 0, \quad 13/50 < t < 261/1000.$$

A five-entry certificate proves the lower bound and explicit stochastic rows prove the upper bound. The root has algebraic degree three, so rational tables do not imply a rational optimal machine. This concerns attainment of a positive error at fixed two-state width, not the rationality of an exact nonnegative factorization. The orthogonal example in Proposition 10.3 gives an additional arbitrary-horizon exact error with rational witnesses. The two-bottleneck compatibility frontier from the preceding paper is included separately in Theorem 8.4.

Fixed-sign matrix approximation and monomial logarithms are classical [22, 24]; global Chebyshev approximation also has a direct optimization literature [23]. We use these mechanisms rather than claiming to originate them. Their precise relation to the multi-epoch bounded tensor, the sign boundary, and the stochastic compiler is discussed in Section 12. The claimed finite alternatives are accompanied by proofs, rational certificate verification and worked exact optima, not by a numerical assertion that a local search has found the global minimum.

The article keeps a single operational object throughout: one reset, a finite external command word, and one selected terminal binary query. The complete prior positive-Hankel, all-word Gram, moment, grid, arithmetic-width, and finite-bit results follow the new sections. In particular, the retained entropy and arithmetic growth claims are not reproved by the finite two-state theorem. The new rational-height lemma specifies the bit cost of candidate Gram verification, while the global alternative deliberately states its larger explicit-table input. Free epoch, horizon, transition tables and exact real row sampling are distinguished from charged persistent labels in Definition 2.1.

The orthogonal-word target is equation (1.1), stated in the introduction. The new two-state alternative also allows arbitrary finite odd response tables; such tables are not silently assumed to have an orthogonal realization of a fixed dimension.

## 8. EXACT TWO-LABEL GEOMETRY AND ROBUST COMPATIBILITY

The positive Hankel normal form below retains the whole future array. For two labels, a smaller exact normal form is available. It gives both a class of finite certificates and an explicit error frontier. The statements in this section concern the actual optimized numerical error, not a sum of local defects or an average over a test law.

**8.1. Antisymmetric experiments.** We first allow any finite response table, not necessarily one obtained from orthogonal matrices. Let the seeds be paired symbols  $X = \{+1, -1, \dots, +n, -n\}$ , let  $w \in \mathcal{A}^N$ , and let  $j \in \{1, \dots, d\}$  be the selected query. Suppose the target means satisfy

$$(8.1) \quad f_{-i,w,j} = -f_{i,w,j}, \quad |f_{i,w,j}| \leq 1.$$

The binary law is  $(1 + bf_{i,w,j})/2$ . Write  $\mathcal{E}_2(f)$  for the minimum worst-word binary-TV error of a machine with two available labels at *every* cut  $0, \dots, N$ . The final held answer also has two labels.

Random initialization and arbitrary cut-dependent stochastic matrices are allowed. All extrema below are attained.

**Theorem 8.1** (An exact scalar normal form). *For every table satisfying (8.1),*

$$(8.2) \quad 2\mathcal{E}_2(f) = \min_{|h_i| \leq 1, |\beta_{t,a}| \leq 1, |d_j| \leq 1} \max_{i,w,j} \left| f_{i,w,j} - h_i d_j \prod_{t=1}^N \beta_{t,at} \right|.$$

*Thus the exact numerical optimization is the distance, in the uniform norm, to the set of bounded decomposable tensors with one factor for the positive seed, one for each command epoch, and one for the query. Every feasible tensor on the right has a two-label stochastic realization. For  $N = 0$  the empty product is one.*

*Proof.* Represent a probability row on two labels by its signed coordinate  $r = 2p - 1 \in [-1, 1]$ . Every stochastic update is an affine map  $r \mapsto \alpha + \beta r$  with  $|\alpha| + |\beta| \leq 1$ ; conversely this condition is sufficient. A bounded terminal mean has form  $c_j + d_j r$  with  $|c_j| + |d_j| \leq 1$ . For a word  $w$ , the composite update has the form  $A_w + B_w r$ , where  $B_w = \prod_t \beta_{t,at}$ . If the two paired initialization coordinates are  $r_i, r_{-i}$ , the odd part of the realized mean is therefore

$$\frac{\hat{f}_{i,w,j} - \hat{f}_{-i,w,j}}{2} = h_i B_w d_j, \quad h_i = (r_i - r_{-i})/2.$$

Taking the odd part cannot increase uniform error against the odd target: each resulting residual is half the difference of two original residuals. This operation is realizable without mixing whole machines. Use initialization coordinates  $\pm h_i$ , replace every update by  $r \mapsto \beta_{t,a} r$ , and replace the terminal column by  $d_j r$ . All rows remain stochastic because  $|\beta_{t,a}|, |d_j| \leq 1$ . This proves that the right side is no larger than twice the original optimum, without retaining a symmetrization coin.

Conversely, initialize the positive label with probability  $(1 \pm h_i)/2$ , use the stochastic matrix

$$(8.3) \quad \frac{1}{2} \begin{pmatrix} 1 + \beta_{t,a} & 1 - \beta_{t,a} \\ 1 - \beta_{t,a} & 1 + \beta_{t,a} \end{pmatrix},$$

and terminal means  $(d_j, -d_j)$ . Induction gives the tensor in (8.2). Binary total variation is half the mean error. Compactness of the factor cubes and continuity prove attainment.  $\square$

The scalar description of a two-state stochastic matrix is elementary. The content needed here is that deleting all affine offsets produces an admissible optimizer for an odd target, rather than merely a signed linear approximation. No averaging over an unrecorded global random choice is permitted. Tensor descriptions of stochastic processes are classical; Ohta's reachable/null-space reductions concern a different stationary observation model [16]. Equation (8.2) keeps the prescribed command order, bounded query columns and fixed numerical error objective simultaneously.

**8.2. Balanced sign certificates.** For an entry  $e = (i, a_1, \dots, a_N, j)$ , form its incidence vector  $\iota(e) \in \mathbb{F}_2^M$ , where

$$M = n + N|\mathcal{A}| + d.$$

There is one coordinate for each seed factor, each pair  $(t, a)$ , and each query factor;  $\iota(e)$  marks the factors used by that entry. A set  $C$  of nonzero target entries is *balanced* if  $\sum_{e \in C} \iota(e) = 0$ . Equivalently each factor occurs an even number of times. It is *negative* if  $\prod_{e \in C} f_e < 0$ .

**Theorem 8.2** (A finite sign obstruction and a gap dichotomy). *A negative balanced set  $C$  certifies*

$$(8.4) \quad \mathcal{E}_2(f) \geq \frac{1}{2} \min_{e \in C} |f_e|.$$

Suppose in addition that  $f_e \in \{0, \rho, -\rho\}$  for a fixed  $0 < \rho \leq 1$ . Then exactly one of the following holds:

$$(8.5) \quad \begin{aligned} \mathcal{E}_2(f) &= \rho/2, & \text{the nonzero signs do not factor;} \\ \mathcal{E}_2(f) &\leq \rho/4, & \text{the nonzero signs factor.} \end{aligned}$$

*Factorability, and hence the alternative in (8.5), is decided by a linear system over  $\mathbb{F}_2$ . In the first case there is a negative balanced certificate of at most  $M + 1$  entries. In the second case the proof constructs a legal two-label machine. Complexity is polynomial in the explicit table size; the table can be exponential in the displayed horizon.*

*Proof.* The product of any decomposable real tensor over a balanced set is nonnegative, since each factor has even exponent. If every tensor entry on  $C$  approximated the target within a mean error strictly smaller than  $\min_C |f_e|$ , it would be nonzero with the target sign. Its product would be negative, a contradiction. Apply Theorem 8.1 to obtain (8.4).

Let  $s_e \in \mathbb{F}_2$  encode the nonzero target sign. The signs factor exactly when the system

$$(8.6) \quad \langle \iota(e), z \rangle = s_e \pmod{2} \quad (f_e \neq 0)$$

is consistent. If it is inconsistent, a linear combination of its rows has zero left side and right side one. The corresponding entries are negative and balanced. One may express a contradictory row in a basis of at most  $M$  earlier rows, so at most  $M + 1$  entries suffice. Inequality (8.4) gives  $\mathcal{E}_2(f) \geq \rho/2$ ; the fair-coin decoder attains the reverse inequality.

If (8.6) has a solution, take the seed and epoch factors to be the prescribed signs  $\pm 1$ , and take each query factor to be its prescribed sign times  $\rho/2$ . Every tensor entry then has magnitude  $\rho/2$  and all nonzero targets have the correct sign. Both zero and nonzero entries have mean error at most  $\rho/2$ . The construction (8.3) realizes these factors with deterministic identity/swap command rows and bounded terminal means. Thus  $\mathcal{E}_2(f) \leq \rho/4$ . Ordinary finite-field elimination both decides consistency and records the indicated certificate.  $\square$

The second alternative need not be an equality: a target which is already a bounded decomposable tensor may have zero error. The theorem classifies the strict-improvement threshold, not every smaller optimum. The certificate is an elementary parity identity, not a new nonnegative-rank theorem or a general polynomial-time minimization algorithm for stochastic automata. The difference is useful: an explicit witness certifies the error against *every* two-label machine, whereas an equivalence test checks a supplied machine.

**8.3. Two bottlenecks and an exact finite frontier.** Return to the orthogonal experiment. Put

$$(8.7) \quad X = \{\pm e_1, \pm e_2\}, \quad \mathcal{A} = \{I, R\}, \quad R = \begin{pmatrix} 0 & -1 \\ 1 & 0 \end{pmatrix}, \quad 0 < \rho \leq 1/6.$$

The finite signal range in (8.7) is separate from the smaller sufficient calibration range of the arithmetic theorems. Every target answer probability is in  $[(1 - \rho)/2, (1 + \rho)/2]$ .

**Theorem 8.3** (Two separated bottlenecks). *At any horizon  $N \geq 1$ , suppose a machine for (8.7) has at most two labels at two distinct cuts  $s < t$  in  $\{0, \dots, N\}$ . Its actual worst-word binary-TV error is at least  $\rho/2$ , regardless of the widths at all other cuts. The fair-coin output attains this bound at every such prescribed profile. The same conclusion holds for any positive signal  $\rho \leq 1$ , independently of the upper construction used below.*

*Proof.* Freeze all commands before  $s$  and after  $t$  to  $I$ . Between the two cuts retain the two words  $I^{t-s}$  and  $RI^{t-s-1}$ , with  $R$  in the first executed position. Integrating the intermediate computation gives two stochastic maps between the two endpoint alphabets. Pad a one-label endpoint if necessary. The initial probabilities may depend arbitrarily on the seed; the suffix gives

one bounded column per query. Thus the odd part of this restricted experiment has the one-step scalar form of Theorem 8.1, even when intermediate alphabets are arbitrarily large.

Four entries of its positive-seed target are

$$f_{1,I,1} = \rho, \quad f_{2,I,2} = \rho, \quad f_{1,R,2} = \rho, \quad f_{2,R,1} = -\rho.$$

Each seed, intervening word and query occurs twice. These four entries form a negative balanced set. Taking the odd part of any putative approximation costs no extra error, so (8.4) gives  $\rho/2$ . The zero mean decoder proves sharpness. No reset occurs inside the interval and no command is replaced by a free input symbol: collapsing the interval is a converse argument only.  $\square$

For  $N = 1$  let  $\mathcal{E}(k_0, k_1)$  denote the optimum at the two command-boundary widths. A positive minimum is attained even when the allowed profile is infeasible for exact realization.

**Theorem 8.4** (A complete two-cut error frontier). *In (8.7) at  $N = 1$ , the exact frontier is*

$$(8.8) \quad \begin{array}{c|ccc} \mathcal{E}(k_0, k_1) & k_1 = 1 & k_1 = 2 & k_1 \geq 3 \\ \hline k_0 = 1 & \rho/2 & \rho/2 & \rho/2 \\ k_0 = 2 & \rho/2 & \rho/2 & \rho/4 \\ k_0 \geq 3 & \rho/2 & \rho/4 & 0 \end{array}$$

*Every upper bound has an explicit rational stochastic realization when  $\rho$  is rational. In particular, minimizing either cut separately at tolerance  $\rho/4$  permits two labels, but no single machine can attain both minima simultaneously.*

*Proof.* A one-label bottleneck erases all seed dependence; for the identity word, opposite coordinate means force error at least  $\rho/2$ . Theorem 8.3 treats  $(2, 2)$ . If either cut has two labels, then, for the identity word, the realized four seed means lie on one affine line  $n \cdot y = c$ . If its uniform coordinate error is at most  $\delta$ , then

$$\delta \|n\|_1 \geq \max_{y \in \{\pm \rho e_1, \pm \rho e_2\}} |n \cdot y - c| = \rho \|n\|_\infty + |c| \geq (\rho/2) \|n\|_1.$$

This proves the  $\rho/4$  lower bounds after dividing mean error by two. A degenerate segment lies in such a line as well.

Here are matching machines. For  $a > 0$ , set

$$V_a = \{(a, 0), (0, a), (-a, -a)\}, \quad \beta_a(y) = \frac{1}{3} \left( 1 + \frac{2y_1 - y_2}{a}, 1 + \frac{2y_2 - y_1}{a}, 1 - \frac{y_1 + y_2}{a} \right).$$

Whenever  $\beta_a(y) \geq 0$ , it sums to one and has barycenter  $y$  in  $V_a$ .

For  $(3, 3)$  initialize with  $\beta_{2\rho}(\rho x)$ . For each vertex  $v \in V_{2\rho}$  and command  $U \in \{I, R\}$ , use the row  $\beta_{6\rho}(Uv)$  on  $V_{6\rho}$ . The six rows are nonnegative: substituting the vertices gives  $2(Uv)_1 - (Uv)_2 \geq -6\rho$ , the analogous second inequality, and  $(Uv)_1 + (Uv)_2 \leq 6\rho$ . Use the coordinates of  $V_{6\rho}$  as terminal means. They lie in  $[-1, 1]$  because  $6\rho \leq 1$ . The barycenter identity gives exactly  $\rho Ux$  for each word and seed.

For  $(2, 3)$  let the two initial labels represent  $\pm v$ , where  $v = (\rho/2, \rho/2)$ ; initialize  $e_1, e_2$  at  $+v$  and their negatives at  $-v$ . On command  $U$ , use the row  $\beta_{2\rho}(\pm Uv)$  and decode with  $V_{2\rho}$ . All four rows are nonnegative. The output is  $U[\rho(x_1 + x_2)(1, 1)/2]$ , differing from  $\rho Ux$  in each coordinate by at most  $\rho/2$  for both  $I$  and  $R$ .

For  $(3, 2)$  initialize with  $\beta_{2\rho}(\rho x)$  and decode the two terminal labels as  $(1, 1)$  and  $(-1, -1)$ . For  $v \in V_{2\rho}$  use positive-label probability  $(1 + \ell(Uv))/2$ , where  $\ell(y) = (y_1 + y_2)/2$ . Here  $|\ell(Uv)| \leq 2\rho \leq 1$ . The output mean is  $\rho((Ux)_1 + (Ux)_2)(1, 1)/2$ , again with coordinate error  $\rho/2$ . All remaining entries in the table follow by padding normalized unused rows, or by the zero decoder. Every coefficient just displayed is rational when  $\rho$  is rational.  $\square$

For  $\rho/4 \leq \varepsilon < \rho/2$ , the feasible two-cut profiles are exactly the upward closure of  $(2, 3)$  and  $(3, 2)$ . For  $0 \leq \varepsilon < \rho/4$ , both command cuts need at least three labels; for  $\varepsilon \geq \rho/2$ , one at each command cut suffices. The held answer is always charged separately as two labels. These are finite numerical frontiers, not asymptotic claims about an irrational alphabet. They complement rather than replace the growing-width laws.

**Theorem 8.5** (An arbitrary-horizon profile frontier). *For the alphabet and seeds in (8.7), let  $N \geq 1$  and suppose every command-cut width belongs to  $\{1, 2\} \cup \{4, 5, \dots\}$ . For any  $0 < \rho \leq 1$ , the exact optimized error is*

$$(8.9) \quad \mathcal{E}(K_\bullet) = \begin{cases} \rho/2, & \text{some } K_t = 1, \text{ or at least two cuts have } K_t = 2, \\ \rho/4, & \text{exactly one cut has } K_t = 2 \text{ and none has width one,} \\ 0, & \text{every } K_t \geq 4. \end{cases}$$

*The constructive machines have deterministic initialization and command rows and rational decoder probabilities when  $\rho$  is rational. The excluded width-three profiles are not classified by this statement.*

*Proof.* For a one-label cut freeze every word to the identity; opposite seeds force error  $\rho/2$ . Two two-label cuts are treated by Theorem 8.3. With one such cut, freeze the entire word to the identity. Its four output means lie in a segment, so the line argument in Theorem 8.4 gives error  $\rho/4$ .

For the matching construction without a bottleneck, store which of  $\{\pm e_1, \pm e_2\}$  is the current vector. The command  $R$  permutes these four states, and the decoder is  $\rho$  times the vector. If the unique two-label cut is  $s$ , track this vector until that cut. At cut  $s$  replace the current vector  $y$  by  $Py = (y_1 + y_2)(1, 1)/2$ . There are only two possibilities,  $\pm(1, 1)/2$ , for these four axis vectors. This is a deterministic compression, with no retained old phase. Thereafter four labels track the quarter-rotation orbit of  $Py$ . The first transition out of the bottleneck embeds its two possibilities into these four labels and applies the actual command; subsequent transitions are permutations. At  $s = 0$  perform the compression in initialization; at  $s = N$  decode the two possibilities directly. The decoder multiplies by  $\rho$ . For each axis vector,  $\|y - Py\|_\infty = 1/2$ ; every subsequent quarter rotation preserves this norm. The final mean error is exactly  $\rho/2$ , proving the second line of (8.9). A fair-coin decoder handles the first line. Available excess labels are padded by normalized unused rows. No random branch or clock history is stored for free.  $\square$

**Corollary 8.6** (Stability under arbitrary target perturbations). *If a finite binary response table is perturbed by at most  $\eta$  in uniform row TV, its optimal error at any fixed profile changes by at most  $\eta$ . In particular the two-bottleneck lower bound becomes  $\max(0, \rho/2 - \eta)$  for any such perturbation of (8.7). The construction at  $(2, 3)$  or  $(3, 2)$  has error at most  $\rho/4 + \eta$ . Thus a strict compatibility gap persists whenever  $\eta < \rho/8$ , including full-support perturbations.*

*Proof.* For a fixed machine, the two target errors differ by at most  $\eta$  by the triangle inequality; minimizing in both directions proves the first claim. Apply it to (8.8); the gap between  $\rho/2 - \eta$  and  $\rho/4 + \eta$  is positive exactly when  $\eta < \rho/8$ . Antisymmetry need not survive the perturbation.  $\square$

## 9. COMPLETE SIGN AND MAGNITUDE ALTERNATIVES

Theorem 8.1 reduces the machine problem to a bounded product, but a sign obstruction alone does not determine its distance. We now give an exact alternative at every error threshold, including thresholds at which entries or optimal factors vanish. Throughout this section the complete finite table is part of the input. Put  $s = N + 2$ , write  $\mathcal{I} = \prod_{t=0}^{s-1} [n_t]$  for the positive-seed, epoch and query indices, and let

$$\mathcal{V} = \{(t, i) : 0 \leq t < s, 1 \leq i \leq n_t\}, \quad v = |\mathcal{V}|.$$

To distinguish the exponent incidence data from right-hand-side bases, we use  $\nu_e$  for an entry incidence vector. For  $e \in \mathcal{I}$ , let  $\nu_e \in \{0, 1\}^\mathcal{V}$  mark its  $s$  factors. Products always use one factor in each mode. A mean tolerance is denoted  $\delta$ ; its binary-TV tolerance is  $\delta/2$ . Assume  $0 \leq \delta \leq \max_e |f_e| \leq 1$ .

**9.1. The support is determined by the threshold.** Define the active entries and essential factor coordinates by

$$(9.1) \quad \mathcal{A}_\delta = \{e : |f_e| > \delta\}, \quad V_\delta = \bigcup_{e \in \mathcal{A}_\delta} \text{supp } \nu_e.$$

An entry survives if all its factors belong to  $V_\delta$ ; let  $\mathcal{S}_\delta$  denote these entries. For  $v \in V_\delta$ , put

$$0 < a_v = \max_{e \ni v} (|f_e| - \delta) \leq 1.$$

The maximum may be taken over active entries incident to  $v$ .

**Lemma 9.1** (Threshold support reduction). *There is a bounded product  $g_e = \prod_{v \in e} \theta_v$  with  $\max_e |f_e - g_e| \leq \delta$  if and only if there is one with*

$$\theta_v = 0 \ (v \notin V_\delta), \quad a_v \leq |\theta_v| \leq 1 \ (v \in V_\delta).$$

*If  $\mathcal{A}_\delta$  is empty, the zero product is feasible.*

*Proof.* An active entry has nonzero approximating product of magnitude at least  $|f_e| - \delta$ . Since every factor has magnitude at most one, each factor in that entry is at least this large. This proves the asserted lower bound on essential factors. Set every other factor to zero. Every changed entry was inactive, so its new residual is  $|f_e| \leq \delta$ . All unchanged entries retain their former residual. Conversely the restricted product is an allowed original product. The empty-active case follows by choosing all factors zero. The argument uses the nonstrict tolerance at inactive entries and therefore also covers  $|f_e| = \delta$ .  $\square$

This support reduction is the multilinear counterpart of deleting threshold-inactive rows and columns in rank-one matrix approximation; compare Gillis–Shitov, Lemma 1 [22]. It is not a generic assumption that a minimizing stochastic machine has strictly positive parameters.

**9.2. Sign chambers and logarithmic feasibility.** A sign chamber is a vector  $\sigma \in \{\pm 1\}^{V_\delta}$  satisfying

$$(9.2) \quad \prod_{v \in e} \sigma_v = \text{sign}(f_e) \quad (e \in \mathcal{A}_\delta).$$

These are linear equations over  $\mathbb{F}_2$ . Failure has the balanced certificate of Theorem 8.2. For  $e \in \mathcal{S}_\delta$  write  $\sigma_e = \prod_{v \in e} \sigma_v$  and set

$$(9.3) \quad L_e = \max(0, \sigma_e f_e - \delta), \quad U_e = \sigma_e f_e + \delta.$$

Every surviving entry is retained, including inactive entries all of whose factor coordinates survive. A positive upper endpoint means  $U_e > 0$ . A chamber with some  $U_e \leq 0$  is rejected: essential-factor products are strictly positive. Otherwise form the following finite system, with  $y_v = -\log |\theta_v|$ :

$$(9.4) \quad \begin{aligned} -y_v &\leq 0, & y_v &\leq \log(1/a_v) & (v \in V_\delta), \\ -\nu_e \cdot y &\leq \log U_e & & & (e \in \mathcal{S}_\delta), \\ \nu_e \cdot y &\leq \log(1/L_e) & & & (e \in \mathcal{S}_\delta, L_e > 0). \end{aligned}$$

Here  $\nu_e$  is restricted to  $V_\delta$ . In particular no logarithm of zero is taken. Values  $U_e > 1$  need not be capped; the corresponding constraint is harmless since  $y \geq 0$ .

**Theorem 9.2** (Complete two-state feasibility alternative). *For every antisymmetric finite numerical table and every  $\delta \geq 0$ ,  $\mathcal{E}_2(f) \leq \delta/2$  if and only if either  $\mathcal{A}_\delta = \emptyset$ , or at least one chamber (9.2) has positive  $U_e$  on every surviving entry and a feasible system (9.4). A feasible point gives a legal two-state machine by*

$$\theta_v = \sigma_v e^{-y_v} \quad (v \in V_\delta), \quad \theta_v = 0 \quad (v \notin V_\delta),$$

followed by the stochastic rows in (8.3).

Write (9.4) as  $Cy \leq \log b$ , where  $C$  is an integer matrix and  $b$  is a vector of strictly positive numbers. A chamber is infeasible if and only if there is a vector  $z$  of nonnegative integers satisfying

$$(9.5) \quad z^T C = 0, \quad \prod_i b_i^{z_i} < 1.$$

Such a vector can be chosen with at most  $|V_\delta| + 1$  nonzero coordinates. Thus nonexistence is certified by a sign contradiction, an impossible surviving upper endpoint, or one multiplicative certificate for each remaining chamber. These alternatives concern all two-state machines, not just supplied transition tables.

*Proof.* Use Lemma 9.1. Each nonzero factor has a sign and a positive magnitude. At an active entry its sign must agree with the target. In a fixed chamber, the mean-error constraint is exactly  $L_e \leq \prod_{v \in e} |\theta_v| \leq U_e$ . The positive upper endpoint is necessary, and taking logarithms gives (9.4). The inactive entries not surviving are already satisfied by zero. This proves both directions, including the reconstruction, using Theorem 8.1.

Farkas' alternative applied to  $Cy \leq \log b$  says that infeasibility is equivalent to  $z \geq 0$ ,  $z^T C = 0$ ,  $z \cdot \log b < 0$ . For completeness, normalize a nonzero such vector to have coordinate sum one. Minimize  $z \cdot \log b$  over the compact polytope

$$\{z \geq 0 : C^T z = 0, \mathbf{1}^T z = 1\}.$$

A minimizing extreme point has at most  $\text{rank}(C) + 1 \leq |V_\delta| + 1$  positive coordinates. Its coordinates are rational, since the equality matrix is integral. Clearing denominators and exponentiating proves (9.5). Conversely multiplying the original inequalities with these nonnegative multiplicities would give  $0 \leq \log \prod b_i^{z_i} < 0$ . This proves the alternative without assuming strict feasibility or an interior optimizer.  $\square$

The dualization and the logarithmic substitution are classical [25, §5.8][24, §2 and §8]. The theorem's role is to make their domain exact for the bounded controlled tensor, including zeros, sign changes, and reconstruction of stochastic rows. It is a finite disjunction of convex feasibility problems, not one convex relaxation of the set of machines. Mixing feasible chambers as if the mixing index were a free persistent label would change the resource and is not used.

**Corollary 9.3** (One chamber in two useful regimes). *If  $\delta < \min_e |f_e|$ , first test whether the full sign tensor factors. If it does, a single system (9.4), with  $L_e = |f_e| - \delta$  and  $U_e = |f_e| + \delta$ , decides feasibility. If  $f_e \geq 0$  for every positive-seed entry, one may use the all-positive chamber at every tolerance, including zeros.*

*Proof.* In the first case every entry is active, so all compatible factor-sign choices give the same entry signs and hence the same inequalities. In the second case replace all factors of any approximating product by their absolute values. For  $u \geq 0$  and  $t \in \mathbb{R}$ ,  $|u - |t|| \leq |u - t|$ . Thus this replacement never increases error. Apply the support reduction and Theorem 9.2.  $\square$

**9.3. Algebraic thresholds, not transcendental oracles.** Although (9.4) contains logarithms, it does not require an oracle for arbitrary transcendental comparisons. At an algebraic threshold with algebraic data, every base  $b_i$  is algebraic. In (9.5) integer powers and products are algebraic,

so their order relative to one is decidable by exact real-algebraic arithmetic. Feasible vertices can be obtained just as explicitly.

**Theorem 9.4** (Finite exact optimization by threshold strata). *For a complete table with real-algebraic entries, the optimum  $\mathcal{E}_2(f)$  and a minimizing stochastic machine are real algebraic and can be computed by a finite algorithm using sign elimination, rational polyhedral enumeration, and univariate real-algebraic root isolation. No generic multivariate minimization of the original stochastic rows is required.*

*More specifically, on each open interval between consecutive values in  $\{0\} \cup \{|f_e|\}$ , feasibility is a finite Boolean combination of explicit polynomial inequalities in  $\delta$ . Each nonconstant boundary polynomial is obtained by clearing denominators in a multiplicative certificate (9.5). Endpoints are tested separately by Theorem 9.2.*

*Proof.* The zero product gives the finite upper threshold  $\max |f_e|$ . Between its consecutive absolute-value breakpoints, the active entries, essential factors, surviving entries, and sign equations are constant. In each chamber the sign of every  $U_e = \delta + \sigma_e f_e$  is constant on this interval. The positive lower endpoints are  $|f_e| - \delta$  at active entries, while  $a_v = \max_{e \ni v} |f_e| - \delta$  on essential coordinates. Thus the matrix  $C$  is constant and integral, and each base  $b_i$  is one, a positive affine algebraic function, or its reciprocal.

Enumerate the extreme rays of the rational cone  $\{z \geq 0 : C^T z = 0\}$ , clearing denominators on each ray. There are finitely many. A chamber is feasible exactly when all the resulting products are at least one. Every denominator has fixed positive sign on the interval. Multiplication converts these conditions into polynomial inequalities with algebraic coefficients. Identically zero polynomials impose no additional boundary. Isolate the finitely many roots of the others, test the intervening intervals and the original breakpoints, and choose the smallest feasible threshold. Feasibility is closed globally by compactness of the original factor cube; in particular this smallest value is not merely an unattained infimum. It is algebraic.

At that value apply the exact support/chamber test. Any feasible system (9.4) is a nonempty bounded polytope. It has a vertex determined by a full-rank selection of active linear rows. The inverse of that integer matrix is rational. Consequently every  $y_v$  is a rational linear combination of logarithms of positive algebraic bases. The number  $e^{-y_v}$  is a product of positive rational powers of algebraic numbers, hence algebraic. Candidate vertices and inequalities can be tested after clearing these rational exponents. Select a feasible vertex and apply the stochastic compiler in Theorem 9.2. If the active set is empty, the zero machine is rational and suffices.  $\square$

For effective input, a common primitive element and isolating interval specify the coefficient field, and each entry has explicit coordinates in that field. Positive rational powers always mean the positive real branch. The input comprises defining polynomials, isolating intervals, and field operations for the real-algebraic entries, together with the *explicit* mode table. Rational input is the special case of integer numerator–denominator pairs. Neither chamber enumeration, ray enumeration, the degrees after clearing exponents, nor the table size is bounded here by a polynomial in a succinct horizon. The algorithm is an exact structural alternative, not a complexity claim. For rational thresholds an exact implementation can eliminate the linear variables while retaining each right side as a rational linear combination of logarithms of rational numbers. Its sign is decided by an integer-power rational product, never a floating-point LP tolerance.

## 10. SPARSE EXTREMALITY AND EXACT NUMERICAL FRONTIERS

**10.1. A multiplicative extremality condition.** The alternative in Theorem 9.2 also identifies active constraints at an optimal error. This is not the usual additive alternation condition for a linear approximation space.

**Proposition 10.1** (A sparse tight certificate). *Suppose  $0 < \delta_* := 2\mathcal{E}_2(f) < \min_e |f_e|$  and that the full sign tensor factors. In the resulting single chamber there is a vector  $z \in \mathbb{Z}_+^r \setminus \{0\}$  supported on at most  $v + 1$  inequalities such that*

$$z^T C = 0, \quad \prod_i b_i(\delta_*)^{z_i} = 1.$$

*The product test can be chosen nonconstant in the threshold. Every primal optimizer makes each inequality with  $z_i > 0$  tight. The feasible set in this fixed chamber and stratum is the intersection of all multiplicative-test intervals and the endpoint conditions  $U_e > 0$ . The full optimizer also tests the stratum endpoints separately.*

*Proof.* The rational certificate cone has finitely many extreme rays. All its products must be at least one at  $\delta_*$ . If every nonconstant test were strictly greater than one, continuity would preserve those finitely many inequalities at a slightly smaller threshold in the same stratum; constant tests would be unchanged. This contradicts minimality. Choose a tight nonconstant extreme ray and use the support bound in Theorem 9.2. For a primal feasible point the slacks  $\log b_i - C_i y$  are nonnegative, and their  $z$ -weighted sum is zero. Thus every positively weighted slack vanishes. Each base in this stratum is constant or increasing in  $\delta$ , so each product test defines an interval. This also proves the final assertion.  $\square$

The following example has a zero-crossing threshold larger than one of its target entries. It therefore uses the full support theorem, not an unproved extension of the strictly active special case.

**10.2. An exact cubic optimum with no rational minimizer.** There are two positive seeds, paired with their negatives, two commands in one actual epoch, and two final queries. Index the positive-seed mean tensor by  $(i, a, j) \in \{0, 1\}^3$ . For  $0 < \rho \leq 1$ , prescribe

$$(10.1) \quad f_{iaj} = \rho \begin{cases} 1/10, & i + a + j = 0, \\ 2/5, & i + a + j = 1, \\ 4/5, & i + a + j = 2, \\ 2/5, & i + a + j = 3. \end{cases}$$

All negative-seed means are their negatives. This is a finite binary numerical experiment, with all probabilities in  $[1/10, 9/10]$  when  $\rho = 1$  and in a smaller interval when  $\rho < 1$ . The two commands have different prescribed effects. This particular table is not asserted to be a two-dimensional orthogonal orbit.

**Theorem 10.2** (A cubic error frontier). *Let  $t$  be the unique real root of*

$$(10.2) \quad 500t^3 - 375t^2 + 540t - 124 = 0.$$

*Then  $13/50 < t < 261/1000$  and the actual optimal two-label error of (10.1) is*

$$(10.3) \quad \mathcal{E}_2(f) = \rho t/2.$$

*A five-entry multiplicative certificate proves the lower bound, and explicit algebraic stochastic rows attain it. If  $\rho$  is positive rational, no all-rational stochastic machine with this same profile attains the optimum. Nevertheless, every tolerance strictly above it admits a dyadic machine with the same boundary labels.*

*Proof.* Set  $\rho = 1$  first. Taking absolute values of all factors cannot increase the error against the nonnegative positive-seed table. The five entries used below are 000, 011, 101, 110, 111. Every decomposable tensor satisfies

$$(10.4) \quad g_{011}g_{101}g_{110} = g_{000}g_{111}^2.$$

If its mean error is at most  $\delta < 4/5$ , the three factors on the left are at least  $4/5 - \delta$ , while the two factors on the right are at most  $1/10 + \delta$  and  $2/5 + \delta$ . Necessarily

$$(10.5) \quad (4/5 - \delta)^3 \leq (1/10 + \delta)(2/5 + \delta)^2.$$

The difference of the right side and the left side is the polynomial in (10.2) divided by 250. Its derivative is strictly positive on the real line, since  $1500t^2 - 750t + 540$  has negative discriminant and positive leading coefficient. Direct rational substitution at  $13/50$  and  $261/1000$  gives opposite signs. Consequently (10.5) fails for  $\delta < t$ . This argument does not require  $g_{000}$  to have been active in the threshold sign equations.

For the reverse bound put

$$A = 1/10 + t, \quad B = 2/5 + t, \quad C = 4/5 - t, \quad u_0 = A^{1/3}, \quad u_1 = B^{1/3}.$$

Use the factor  $(u_0, u_1)$  in each of the three modes. The resulting tensor has values  $A, D = A^{2/3}B^{1/3}, C$ , and  $B$  at Hamming weights zero, one, two, and three. Indeed,  $C^3 = AB^2$  is exactly (10.2). The errors at weights zero, two, and three are  $t$ . At weight one,  $D^2 = AC$  and the rational bounds on  $t$  give  $2/5 < D < 1/2$ , so  $|D - 2/5| < 1/10 < t$ . For completeness the weight-one calculation is

$$D^2 = A^{4/3}B^{2/3} = A(AB^2)^{1/3} = AC, \quad \frac{4}{25} < \frac{9}{25} \frac{539}{1000} < AC < \frac{361}{1000} \frac{27}{50} < \frac{1}{4}.$$

All powers use their positive real branches. All factors belong to  $(0, 1)$ . Initialize with signed coordinate  $\pm u_i$ , use the two command rows  $T(u_a)$  of (8.3), and decode by  $(u_j, -u_j)$ . This is an actual two-label machine, not just an unconstrained tensor. Binary TV is half the mean error. For general  $\rho$ , multiply one factor mode by  $\rho$ . The scaled lower constraint is explicitly

$$(4\rho/5 - \delta)^3 \leq (\rho/10 + \delta)(2\rho/5 + \delta)^2;$$

division by  $\rho^3$  proves sharpness.

Modulo seven the polynomial in (10.2) is  $3t^3 + 3t^2 + t + 2$ , with values 2, 2, 5, 1, 1, 2, 1 at the seven residues. It has no root and is irreducible over  $\mathbb{F}_7$ , hence the primitive cubic is irreducible over  $\mathbb{Q}$ . For rational  $\rho > 0$ , the optimal error is therefore irrational. Here all-rational means that every initialization, transition and decoder coefficient is rational. Such a machine at this fixed finite horizon has rational responses and rational maximum error against the rational target. It cannot attain (10.3). The dyadic assertion follows by approximating the finitely many legal rows, or by Theorem 14.4. It asserts existence *below any prescribed larger tolerance*, not equality to every larger real error.  $\square$

This is field sensitivity of a *positive minimum error* in a binary-query problem. It is distinct from the classical dependence of exact nonnegative rank on the factorization field [21]. It also goes beyond the sign criterion: every nonzero positive-seed target in (10.1) has the same sign. For reference, the exact rational interval

$$\frac{260362898747}{10^{12}} < t < \frac{260362898748}{10^{12}}$$

is certified by the signs of the cubic. The source package includes a dyadic upper machine and a strictly violated multiplicative lower certificate at these two endpoints. No numerical local optimizer is used to assert the interval.

For this minimizing machine the all-word residual sums in Theorem 14.1 and (14.6) are, respectively,

$$S_2 = 2\{5t^2 + 3(D - 2/5)^2\}, \quad S_4 = 2\{5t^4 + 3(D - 2/5)^4\}$$

when  $\rho = 1$ . Their positivity detects nonexactness of the supplied machine. It does not prove its optimality. The circuit (10.4) proves the lower bound against every machine; the displayed factors prove the upper bound. These are the distinct roles of candidate verification and global certification.

**10.3. A rational orthogonal example at every horizon.** Let  $A_0 = I$ ,  $A_1 = \begin{pmatrix} 0 & 1 \\ 1 & 0 \end{pmatrix}$ , let the positive seeds be  $(4/5, 3/5)$  and  $(3/5, 4/5)$ , and include their negatives. Request one of the two coordinates after a word of length  $N \geq 1$ , at signal  $0 < \rho \leq 1$ .

**Proposition 10.3** (A magnitude obstruction for distinct orthogonal commands). *The minimum all-two-label error is  $\rho/20$ . More generally, if a profile has at least two labels at every cut and two distinct cuts with exactly two labels, its minimum error is  $\rho/20$ , regardless of the other widths. If a cut has only one label, its minimum error is  $2\rho/5$ .*

*Proof.* Freeze the prefix and suffix outside the two narrow cuts to identity. For cuts  $s < t$ , retain the macro-words  $0^{t-s}$  and  $10^{t-s-1}$  inside them and fix query  $j = 1$ . The endpoint computation has the scalar form of Theorem 8.1. With one fixed query its positive-seed mean matrix is  $\begin{pmatrix} 4\rho/5 & 3\rho/5 \\ 3\rho/5 & 4\rho/5 \end{pmatrix}$ . A decomposable approximant satisfies  $g_{00}g_{11} = g_{01}g_{10}$ . At mean tolerance  $\delta < 3\rho/5$  this forces  $(4\rho/5 - \delta)^2 \leq (3\rho/5 + \delta)^2$ , hence  $\delta \geq \rho/10$ . Larger tolerances already exceed that bound. Store only the sign of the seed, use identity command rows, and give both queries mean  $\pm 7\rho/10$ . Every target mean differs by exactly  $\rho/10$ ; this attains TV error  $\rho/20$  at the whole profile. A one-label cut makes the response independent of the sign of the seed. Identity words and a coordinate of magnitude  $4\rho/5$  force TV error  $2\rho/5$ ; a fair coin attains it.  $\square$

The orthogonal example is a magnitude test beyond sign consistency, not an assertion that its lower bound is invisible to every separate matrix flattening. The genuine tradeoff between separately minimal cuts remains Theorem 8.4. Both examples use actual nonidentical commands and a positive horizon. The identity and swap commute; noncommutativity is supplied instead by Theorem 6.1.

For any of these fixed profiles, perturbing the full target probability table by at most  $\eta$  in row TV changes its optimum by at most  $\eta$ . The triangle inequality proves this in both directions, even when the perturbed target is no longer antisymmetric. Thus their strict error margins are robust statements about all machines, not exact algebraic identities that disappear immediately under perturbation.

## 11. CERTIFICATE SIZE AND COMPARISON WITH TENSOR GEOMETRY

**11.1. Finite output bounds.** Here the input to Theorem 9.2 is explicitly listed. At a threshold  $\delta$ , let  $v = |V_\delta|$  be the number of essential factor coordinates, let  $n$  be the number of table entries, and let  $m \leq 2v + 2n$  be the number of inequalities in one chamber. Sparsity below is per chamber. It is not a bound for the union of chambers. The empty-active case is solved by the zero product and needs no rays.

**Proposition 11.1** (Explicit certificate and algebraic-output bounds). *For  $v \geq 1$ , write  $r_2$  for the rank over  $\mathbb{F}_2$  of the active sign equations. When these equations are consistent, at most  $2^{v-r_2}$  sign assignments need be tested, and induced-pattern identifications can only reduce this number. Each chamber has at most*

$$(11.1) \quad R(m, v) = \sum_{s=1}^{\min(m, v+1)} \binom{m}{s}$$

*extreme rays. Each ray has a primitive integer representative with support at most  $v + 1$ , entries at most  $v!$ , and sum at most  $h_v = (v + 1)!$ . A chamber-complete infeasibility certificate thus uses at most  $2^{v-r_2}$  ray/endpoint records, in addition to its input and sign coverage description. Its ray indices and multiplicities require at most*

$$O(2^{v-r_2}(v + 1)(\log(m + 1) + v \log(v + 1)))$$

*bits. This bound is on the record data, not a polynomial-time search.*

*For rational targets with numerator and denominator bit lengths at most  $\tau$ , every boundary polynomial on a fixed threshold stratum has degree at most  $h_v$  and can be represented by integer coefficients of bit length  $O(h_v(\tau + 1))$ . Including the original breakpoints, there are finitely many such polynomials across at most  $n + 1$  strata, with the chamber and ray counts above. Taking the largest  $v$  over the strata gives a degree bound  $h_v$  for the optimum threshold. A minimizing machine has a shared algebraic-radical description: one isolated real threshold, followed by at most  $v$  products of positive rational powers of its algebraic affine bases. Their common power denominator is at most  $v!$ . If one instead insists on a single number field for all factors, degree at most  $h_v(v!)^v$  suffices for rational input.*

*Proof.* Gaussian elimination gives the sign count. The cone  $\{z \geq 0 : C^T z = 0\}$  is pointed. On an extreme ray with support  $s$ , the restricted columns have rank  $s - 1$ , for otherwise there is a nontrivial perturbation supported there and the ray is not extreme. Thus  $s \leq v + 1$ . A nonzero cofactor vector of an  $(s - 1)$ -row submatrix spans the kernel. Entries of  $C$  are  $0, 1, -1$ , so the Leibniz determinant bound gives every cofactor magnitude at most  $(s - 1)! \leq v!$ . Choose its nonnegative sign and divide by the common integer divisor. Each support supports at most one extreme ray. This proves (11.1), the sum bound, and the record size. If an infeasible chamber exists, Farkas supplies a nonzero dual vector; only then is the normalized dual polytope nonempty. Some extreme ray is strictly negative on the logarithmic right side. Consequently one violating record per rejected chamber is sufficient, although finding all the records may enumerate every ray.

On a fixed stratum every base is one, a positive affine rational function of  $\delta$ , or its reciprocal. Clearing the positive denominators in a ray product gives a difference of two products of affine factors, with total multiplicity at most  $h_v$ . Its degree is at most  $h_v$ . Clear the original rational denominators in the individual factors before multiplying. The coefficient  $\ell_1$  norm of a product is bounded by the product of the factor norms, so its coefficient bit lengths are  $O(h_v(\tau + 1))$ . Zero polynomials add no boundary. Root isolation and separate breakpoint tests give the stated degree bound, using the largest essential-coordinate count.

At a feasible vertex, a nonsingular  $v$ -by- $v$  integer row matrix  $A$  gives  $y = A^{-1} \log b$ . Its determinant has absolute value  $h \leq v!$ , and its adjugate is integral. Hence every positive factor  $e^{-y_i}$  has its  $h$ th power in the field generated by  $\delta$ . Adjoining these  $v$  positive roots enlarges that field by degree at most  $h^v$ , which proves the coarse common-field bound. Row reconstruction is affine in the factors and does not enlarge it.  $\square$

A radical witness can be checked without first constructing that large common field. Check positivity of every base at the isolated threshold, raise each factor inequality to the positive common power  $h$ , and compare the resulting products in the threshold field. The same applies to each surviving entry constraint. Exact field arithmetic, integer powers and real-root isolation are sufficient. Their cost can be exponential in  $v$  and grows with the fully expanded polynomial coefficients; the compressed exponent notation is not a polynomial bit-complexity claim. With algebraic input, first specify a common primitive element, its defining polynomial and isolating interval, and coordinates of each coefficient in that field. The same proof works over this explicitly represented field; norms or resultants reduce the threshold equations to rational polynomials. Field-degree and coefficient-height growth must be charged as well.

For clarity, a feasible set within a fixed sign chamber and threshold stratum is the intersection of all extreme-ray product conditions and the surviving upper-endpoint conditions  $U_e > 0$ . This is the precise interpretation of the final sentence of Proposition 10.1. The full optimizer must also test the stratum endpoints; no strict-interior assumption is being made.

**11.2. The Segre exponent matrix and the added interval problem.** The incidence vectors in Section 9 are columns of the Segre exponent matrix, up to transpose convention. An integer relation between them gives a binomial identity between rank-one tensor entries. Minimal-support

relations are circuits; parity relations describe the real sign constraints. Zero factors delete complete coordinate slices. These structures and the distinction between real and complex completability are treated directly by Kahle–Kubjas–Kummer–Rosen [27, Section 2, especially Proposition 2.2, Remarks 2.3–2.5 and Proposition 2.7]. In particular, cubic, nonsquarefree tensor circuits and sign obstructions are already part of that theory.

The present interval problem does not claim those antecedents as new. Its data are complete bounded response tables with absolute-error intervals whose endpoints move with  $\delta$ . A zero is permitted at an inactive entry, while coordinates forced by active entries must survive even when they also occur in inactive entries. On that support, the upper and lower interval inequalities, factor caps and all real sign chambers are coupled. Farkas’ alternative then certifies the absence of *any* bounded product in that interval box, and the odd-coordinate compiler turns every feasible product into legal stochastic rows. A tensor-completion equality on its own neither provides those moving intervals nor accounts for the charged selector in a mixture of machines.

| Present result                     | Closest antecedent                                                     | Identical mechanism and precise addition                                                                                                                                                 |
|------------------------------------|------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Two-state alternative              | Gillis–Shitov [22]; tensor completion [27]; geometric programming [24] | Support, signs, exponent relations and logarithms are classical. The added result couples the complete error box to an exact stochastic normal form, including zeros and reconstruction. |
| Rank-tight simplex theorem         | Positive realization cones [12]                                        | Cone invariance is classical. Rank equality removes all hidden suffix directions at selected cuts, including separated cuts in variable profiles.                                        |
| Occupation and four-state frontier | Simplex difference-body identity; inherited packet budget              | The volume identity is classical. Its repeated application yields an all-machine narrow-cut budget even when every packet charge is zero.                                                |
| Three-epoch cubic optimum          | Segre circuits [27]; the retained cubic example                        | The relation is classical. Rational orthogonal, noncommuting commands realize it with a sharp stochastic optimum exceeding every separate flattening bound.                              |
| Finite output bounds               | Cofactor bounds, Farkas’ lemma and real-algebraic isolation            | These tools are classical. The counts price this exact chamber-complete algorithm and its minimizing-row representation.                                                                 |

The general algebraic two-state optimizer is proved but not implemented in full. The inherited rational threshold checker and rational bisection are implemented, as is the specialized cubic certificate. The new executable checks verify the rational permutation construction, all its words, the flattening witnesses, the simplex volume identities, and the stated exact numerical frontier bound. They do not implement an arbitrary higher-width optimizer or constitute independent priority certification.

## 12. ENCODING, COMPUTATION, AND THE NEIGHBORING PROBLEMS

The source of an error certificate must be distinguished from the cost of verifying a supplied realization. The new alternatives require the complete numerical tensor. Its size is  $nd \prod_t |\mathcal{A}_t|$ ; it need not be polynomial in a horizon described without its words. The all-word Gram recurrence below, by contrast, accepts an explicitly given orthogonal target and a proposed machine and need not expand the word table. It decides exactness of that candidate, not whether an optimum over candidates has been found.

**Proposition 12.1** (An inherited hardness boundary in a binary interface). *For complete rational antisymmetric tables with one command epoch, one selected query, and a variable seed and command alphabet, deciding  $\mathcal{E}_2(f) \leq \varepsilon$  for rational  $\varepsilon$  is NP-complete. Here  $\varepsilon$  is binary-TV tolerance and  $2\varepsilon$  is mean tolerance. This is a direct transfer of the rank-one matrix infinity-norm decision theorem of Gillis–Shitov [22, Theorem 3], not a new hardness theorem for a fixed-dimensional orthogonal alphabet.*

*Proof.* Given the matrix instance  $(M, k)$  in that theorem, let  $B = \max(1, \|M\|_\infty + k)$  and prescribe  $f_{i,a} = M_{i,a}/B$  and tolerance  $\varepsilon = k/(2B)$ . If  $X = uv^T$  approximates  $M$  within  $k$ , then  $X/B$  has maximum entry magnitude at most one. For a rank-one matrix this maximum is the product of the two factor infinity norms. Rescale one factor to infinity norm one and the other to at most one; the zero matrix is treated separately. The scalar normal form gives the required two-state realization. Conversely its two bounded factors give a rank-one approximation to  $M$ . All input numbers and tables in the reduction have polynomial binary length.

For membership, strip threshold-inactive rows and columns as in Lemma 9.1 and guess their nonzero factor signs. Writing the remaining positive magnitudes as  $h_i\beta_a$  and  $t_i = 1/h_i \geq 1$ , the constraints become

$$(\sigma_{ia}f_{ia} - 2\varepsilon)t_i \leq \beta_a \leq (\sigma_{ia}f_{ia} + 2\varepsilon)t_i, \quad t_i \geq 1, \quad 0 \leq \beta_a \leq 1.$$

This is a rational linear system. On essential columns the active constraints force  $\beta_a > 0$  automatically. A nonempty rational polyhedron of this form has a rational point with polynomial bit length, obtained from an appropriate basic system by Cramer’s rule: after clearing denominators,  $q$  variables and coefficient bit length  $L$  give determinant bit length  $O(q(L + \log(q + 1)))$ . The feasible set is pointed because  $t_i \geq 1$  and  $0 \leq \beta_a \leq 1$ , so a basic feasible point exists. This is the standard rational-polyhedron argument; compare [28, Chapter 10]. Its reciprocals and the guessed signs give a polynomially checkable rational machine. The single query factor is absorbed into the command factors. Thus this restricted decision problem is in NP.  $\square$

This result explains why enumerating all sign chambers must not be called a general polynomial-time algorithm. Even the matrix problem already has a substantial optimization literature. Gillis–Shitov, Lemmas 1–2, remove threshold-inactive coordinates and reduce a fixed-sign matrix instance to linear inequalities using reciprocal variables. Morozov–Smirnov–Zamarashkin analyze alternating minimization and sign-transition graphs for global rank-one Chebyshev approximation [23, Theorem 5.3]. Neither fixed-sign feasibility nor global matrix optimization is introduced here as a discovery. For the multi-epoch tensor, logarithmic magnitudes linearize its higher-order incidence constraints. The construction is the familiar monomial-to-linear conversion of geometric programming; the same reference already describes threshold bisection for relative fitting error [24, §2 and §8]. Our objective remains *absolute* mean error, encoded by its exact asymmetric lower and upper endpoints, not the log-Chebyshev objective obtained by fitting logarithms of the data.

The difference asserted in this revision is the complete finite alternative for the controlled bounded tensor, its zero-boundary handling, its multiplicative infeasibility witnesses against all machines in the two-state class, and the reconstruction of legal stochastic rows. The exact cubic example is an application in which sign consistency alone is insufficient. The existence of such certificates is proved from classical linear alternatives, not from an asserted novel version of Farkas’ lemma. The multiple-cut compatibility tradeoff of Theorem 8.4 and the original all-hidden-state lower bounds are inherited and retained.

| Question                 | Relevant antecedent                                            | Statement used here                                                                   |
|--------------------------|----------------------------------------------------------------|---------------------------------------------------------------------------------------|
| Candidate equality       | Weighted/probabilistic equivalence [20, 14]                    | A source-bound finite-clock Gram test, not an optimum certificate.                    |
| Positive realization     | Invariant cones and controlled suffixes [12, 18]               | Compatible positive Hankel factors retained with all suffixes.                        |
| Fixed-sign approximation | Matrix Chebyshev feasibility [22] and monomial logarithms [24] | Exact multi-mode bounds, zero support, sign chambers and multiplicative alternatives. |
| Global two-state error   | Scalar reduction and sign obstruction in revision 47 [26]      | Complete magnitude certificates and an exact cubic numerical optimum.                 |
| Arithmetic width         | Prior packet and Diophantine proofs [7, 8, 19]                 | Retained, not assigned new exponents by the present certificate theorem.              |

The lower optimization concerns two available labels at each selected command boundary. This two-state optimizer does not classify general three-state or larger hidden realizations. The distinct rank-tight geometric results of Section 4 do apply beyond two states under their stated spanning and orthogonality hypotheses. Common positive Hankel factorization remains the appropriate general model; the one-dimensional odd coordinate is special to two states. Nor does the finite algorithm bound ordinary uniform workspace: it charges neither stored row descriptions nor arithmetic, and it does not turn an algebraic optimum into exact finite-fair-bit sampling. The separate compiler in the retained material states its own assumptions and space accounting.

### 13. RETAINED ARITHMETIC STATEMENTS

The following statements and their proofs later in the article are retained from the preceding development, not consequences of the new two-state alternative.

Our arithmetic conclusions go beyond uniform bad approximation. For  $\alpha \in \mathbb{R}^r$  with  $1, \alpha_1, \dots, \alpha_r$  rationally independent, write

$$(13.1) \quad \omega^*(\alpha) = \limsup_{\|m\|_1 \rightarrow \infty} \frac{-\log \|m \cdot \alpha\|_{\mathbb{R}/\mathbb{Z}}}{\log \|m\|_1}.$$

This is the ordinary dual exponent  $\omega_{1,r}$  of the row vector  $\alpha$  in [17, Definition 1], with an equivalent fixed-dimensional norm; it is not the uniform exponent  $\hat{\omega}_{1,r}$  or the simultaneous exponent  $\omega_{r,1}$ . The alphabet consists of the identity and the rotations through  $2\pi\alpha_i$ . Throughout,  $r$ , the alphabet,  $0 < \rho \leq 1/10$  and  $0 \leq \varepsilon < \rho/(2\sqrt{2})$  are fixed.

**Theorem 13.1** (Minimal Diophantine type). *If  $\omega^*(\alpha) = r$ , then*

$$(13.2) \quad \lim_{N \rightarrow \infty} \frac{\log W_{N,\varepsilon}(\alpha)}{\log N} = \frac{r}{2r+1}.$$

*For Lebesgue-almost every  $\alpha$ , and every fixed  $\delta > 0$ , there are positive constants  $c, C$ , depending on the fixed  $\alpha, \delta, r, \rho, \varepsilon$ , such that, for  $N \geq 2$ ,*

$$(13.3) \quad c \frac{N^{r/(2r+1)}}{(\log(N+2))^{2r(1+\delta)/(2r+1)}} \leq W_{N,\varepsilon}(\alpha) \leq CN^{r/(2r+1)} (\log(N+2))^{2r^2(1+\delta)/(2r+1)}.$$

*The upper inequalities are achieved by exact common-row machines. No uniformity near rational relations or in growing  $r$  is asserted.*

The exponent in (13.2) need not imply a two-sided constant-factor law. The earlier dual badly approximable theorem supplies the stronger  $\Theta$  conclusion on its stated class, and is recovered in Corollary 18.4. Theorem 13.1 instead allows arbitrarily small Diophantine constants along subsequences. The quantitative arithmetic ingredients are classical; the memory conclusions concern their compatible realization and all-hidden-state converse.

There is also a genuine failure of a single exponent, even for one fixed irrational angle.

**Theorem 13.2** (Irrational horizons and Liouville fluctuations). *For every irrational  $\alpha$  there are horizons  $M_n \rightarrow \infty$  such that  $W_{M_n, \varepsilon}(\alpha) = \Theta(M_n^{1/3})$ . If  $\alpha$  is Liouville, then*

$$(13.4) \quad \liminf_{N \rightarrow \infty} \frac{\log W_{N, \varepsilon}(\alpha)}{\log N} = 0, \quad \frac{1}{3} \leq \limsup_{N \rightarrow \infty} \frac{\log W_{N, \varepsilon}(\alpha)}{\log N} \leq \frac{1}{2}.$$

Nevertheless  $W_{N, \varepsilon}(\alpha) \rightarrow \infty$ . All displayed upper bounds have exact realizations.

The upper limit in (13.4) is not determined here. The result establishes oscillation between subpower subsequences and a positive-power subsequence; it does not replace that distinction by a guessed exponent for all irrationals.

#### 14. FINITE CERTIFICATES FOR THE WHOLE NUMERICAL EXPERIMENT

The Hankel normal form uses all suffixes. We now give an equivalent finite certificate whose size is polynomial in a *displayed* horizon and proposed widths. We then treat the actual optimal error, not the sum of local compatibility defects. Throughout this section  $X$  and  $\mathcal{A}$  are nonempty,  $|\mathcal{A}| = m$ ,  $D \geq 1$ , and  $K_0, \dots, K_N \geq 1$ . The held binary answer is a separate two-label cut. A horizon is supplied in unary for statements about input size. Rational numbers have binary encodings; a real algebraic coefficient is supplied by a polynomial and a rational isolating interval. These conventions do not give an oracle for arbitrary real numbers.

In this section  $T_{t,a}$  labels the edge from cut  $t$  to cut  $t+1$ , a reindexing of Definition 2.1. Write  $E_x$  for an initialization row,  $T_{t,a}$  for the  $K_t$ -by- $K_{t+1}$  row-stochastic transition, and  $d_j \in [-1, 1]^{K_N}$  for the terminal mean column. Define the residual mean

$$(14.1) \quad R_{x,w,j} = E_x T_{0,a_1} \cdots T_{N-1,a_N} d_j - \rho \langle e_j, U_{a_N} \cdots U_{a_1} x \rangle.$$

Thus the actual worst-word binary-TV error is  $e(\mathcal{M}) = \frac{1}{2} \max_{x,w,j} |R_{x,w,j}|$ . In particular, centering the numerical response does not turn a bounded stochastic decoder into an unrestricted signed readout.

Set  $n_t = K_t + D$  and form rectangular block matrices and vectors

$$(14.2) \quad A_{t,a} = \text{diag}(T_{t,a}, U_a^\top), \quad u_x = (E_x, -\rho x^\top), \quad v_j = \begin{pmatrix} d_j \\ e_j \end{pmatrix}.$$

Here  $A_{t,a}$  has size  $n_t$  by  $n_{t+1}$ ,  $u_x$  has length  $n_0$ , and  $v_j$  has length  $n_N$ . The order and transpose in this definition are essential for noncommuting commands. Equation (14.1) equals  $u_x A_w v_j$ .

**Theorem 14.1** (A finite all-word certificate). *Define*

$$(14.3) \quad G_0 = \sum_{x \in X} u_x^\top u_x, \quad G_{t+1} = \sum_{a \in \mathcal{A}} A_{t,a}^\top G_t A_{t,a}, \quad S_2 = \sum_{j=1}^D v_j^\top G_N v_j.$$

Then

$$S_2 = \sum_{x \in X} \sum_{w \in \mathcal{A}^N} \sum_{j=1}^D R_{x,w,j}^2.$$

Consequently a legal proposed machine is exact on every word if and only if  $S_2 = 0$ . For rational data this can be checked by exact arithmetic in polynomial time in the displayed data. If  $S_2 > 0$ , a seed, word and query with nonzero residual can also be found in polynomial time, without listing all words.

Exact feasibility of a prescribed profile for an algebraic orthogonal-word target has an existential real formula of polynomial size in  $N, m, D, |X|$  and the explicitly displayed widths. Its equations

may all be taken to have degree at most three after auxiliary arithmetic variables are introduced. Feasibility has a real algebraic witness whenever it has a real witness.

*Proof.* Induction in (14.3) gives  $G_t = \sum_{x, |u|=t} (u_x A_u)^\top (u_x A_u)$ . Pairing with  $v_j$  proves the identity. Every summand is a real square, so zero means equality of all numerical responses, not merely equality on average. No assumption about the ranks of reachable distributions or the observability of individual states is used.

For a witness, compute backwards

$$C_N = \sum_j v_j v_j^\top, \quad C_t = \sum_a A_{t,a} C_{t+1} A_{t,a}^\top.$$

Select  $x$  with  $u_x C_0 u_x^\top > 0$ . Given a current row  $u$  at cut  $t$  with  $u C_t u^\top > 0$ , at least one command satisfies  $u A_{t,a} C_{t+1} A_{t,a}^\top u^\top > 0$ . Append it and replace  $u$  by  $u A_{t,a}$ . At the last cut choose  $j$  with  $u v_j \neq 0$ . The matrices are positive semidefinite by their recursions, so this procedure cannot encounter a dead end. There are  $O(Nm)$  matrix operations at dimensions at most  $\max_t n_t$ . For rational input, Lemma 14.2 below gives the required polynomial intermediate bit length. Exact algebraic sign testing gives the corresponding effective statement for algebraic inputs; no unspecified real-number comparison is invoked.

For existence, treat  $E, T, d, G$  as unknowns. Impose nonnegativity and unit row sums on  $E, T$ , the bounds on  $d$ , the recursions (14.3), and  $S_2 = 0$ . Each recursion entry has degree at most three, and its number of terms is polynomial in the matrix sizes. The Gram variables need no independent positivity constraint: the initial condition and recurrence already force the displayed sum of squares. If target coefficients are algebraic rather than rational, include their defining polynomials and isolating inequalities. Arithmetic circuit variables reduce larger degrees to quadratic equations. Elimination over real closed fields decides existence and permits algebraic sample points. This last conclusion is the usual real-algebraic sampling principle [1], not a polynomial-time optimizer.  $\square$

**Lemma 14.2** (Rational height of the all-word Gram calculation). *Suppose the displayed augmented arrays  $u_x, A_{t,a}, v_j$  contain  $S$  rational entries, each with numerator and positive denominator of binary length at most  $h$ . Let  $d = \max_t (K_t + D)$  and  $m = |\mathcal{A}|$ . All exact Gram and backward-witness calculations have intermediate numerator and denominator bit lengths bounded by*

$$O((N+2)(Sh + h + \log(d+1)) + N \log(m+1) + \log(|X|D+1)).$$

Consequently candidate equality and extraction of one failing word have polynomial bit cost in the displayed rational input and unary horizon. This is a statement about the specified computation, not about optimizing its unknown stochastic rows.

*Proof.* Choose the common denominator  $Q_0$  equal to the product of the input denominators. Then  $\log_2 Q_0 \leq Sh$ , and each input numerator after conversion to this denominator has magnitude at most  $Q_0 2^h$ . An entry of  $G_t$  is a sum of at most  $|X| m^t d^{2t}$  monomials of degree  $2t+2$ . Before reduction its denominator is  $Q_0^{2t+2}$  and the magnitude of its numerator is at most  $|X| m^t d^{2t} (Q_0 2^h)^{2t+2}$ . An entry of the backward matrix has the same estimate with the remaining length, while the final scalar and positivity tests have degree at most  $2N+4$  and at most  $|X| D m^N d^{2N+2}$  summands. Taking logarithms gives the displayed bound. Partial sums and products in the recurrence have no larger order of bit growth. Reduction of a rational fraction cannot enlarge these bounds. There are polynomially many rational operations on matrices of size at most  $d$  and at most  $Nm$  candidate witness edges. Standard integer arithmetic therefore gives the asserted bit bound. For degree  $p$  moments, the same count has degree at most  $p(N+2)$ ; with dimension fixed, the number of monomial coefficients is polynomial in the numerical parameter  $p$  as stated below.  $\square$

The linear-algebraic equivalence mechanism is classical: compare Tzeng's equivalence algorithm [20] and the weighted-automata Hankel and Gram constructions of Balle–Panangaden–Precup [14]. The assertion here makes the finite clock, stochastic constraints, all seed initializations, and bounded query columns simultaneous. It must not be read as a new discovery of polynomial-time weighted-automata equivalence. Also,  $S_2 \approx 0$  in floating point is *not* the exact certificate  $S_2 = 0$ .

**14.1. A hierarchy for the actual optimal error.** Let  $\mathcal{Q} = |X|Dm^N$ . For each even integer  $p \geq 2$ , define

$$(14.4) \quad L_p(\mathcal{M}) = \frac{1}{2} \left( \mathcal{Q}^{-1} \sum_{x,w,j} |R_{x,w,j}|^p \right)^{1/p}, \quad \ell_p(K_\bullet) = \min_{\mathcal{M}} L_p(\mathcal{M}),$$

where every minimization uses the *same* product of stochastic row simplices and bounded decoder cubes for the specified profile. Put  $\mathcal{E}(K_\bullet) = \min_{\mathcal{M}} e(\mathcal{M})$ .

**Theorem 14.3** (Moment certificates and the error frontier). *All the preceding minima are attained. For every even  $p \geq 2$ ,*

$$(14.5) \quad \ell_p(K_\bullet) \leq \mathcal{E}(K_\bullet) \leq \mathcal{Q}^{1/p} \ell_p(K_\bullet).$$

*The numbers  $\ell_p$  increase along the even integers and converge to  $\mathcal{E}$ . The running minima of the right sides of (14.5) decrease to the same value. In particular,  $\ell_2 = 0$  if and only if the profile admits an exact realization. For each fixed  $p$ , its moment sum is computable by a finite homogeneous polynomial recurrence without enumerating words. For fixed  $\max_t(K_t + D)$ , this recurrence has size polynomial in  $p, N, m, D, |X|$ .*

*For algebraic target data the actual optimum  $\mathcal{E}$  and a minimizing machine are algebraic and effectively obtainable. For a finite prescribed width cap, the resulting threshold values give the exact finite-horizon error/width frontier. This is an effective finite optimization statement, not a closed formula or a polynomial-time algorithm for growing width.*

*Proof.* The parameter set is nonempty and compact, and the objectives are continuous. The inequalities for a vector of  $\mathcal{Q}$  numbers are

$$\left( \mathcal{Q}^{-1} \sum |r_i|^p \right)^{1/p} \leq \max_i |r_i| \leq \mathcal{Q}^{1/p} \left( \mathcal{Q}^{-1} \sum |r_i|^p \right)^{1/p}.$$

The first inequality survives minimization. For the second, use a minimizer of  $L_p$  as a candidate for  $\mathcal{E}$ ; no exchange of a maximum over words with a minimum over machines is made. Normalized finite-measure  $L^p$  norms increase with  $p$ , as do their minima. Since  $\mathcal{Q}^{1/p} \rightarrow 1$ , the two sequences converge as asserted, including the case of a zero optimum.

The recurrence is particularly transparent in polynomial form. Let  $z$  be a row variable of the appropriate dimension and set

$$(14.6) \quad F_N(z) = \sum_j (zv_j)^p, \quad F_t(z) = \sum_a F_{t+1}(zA_{t,a}).$$

Then  $\sum_x F_0(u_x)$  is the moment sum. A homogeneous degree- $p$  polynomial in  $n_t$  variables has  $\binom{n_t+p-1}{p}$  coefficients. Repeated polynomial multiplication and linear substitution compute these coefficients in polynomially many arithmetic operations when  $n_t$  is fixed. This is the symmetric-power form of the familiar tensor product of weighted representations. It does not imply a bound polynomial in both  $p$  and unbounded  $n_t$ .

For the last assertion one may instead introduce all finitely many word inequalities  $-2e \leq R_{x,w,j} \leq 2e$  and minimize  $e \in [0, 1]$ . Projection of this compact semialgebraic feasible set gives its semialgebraic attainable-error set. Its left endpoint is algebraic, and real-algebraic sampling gives a machine in the fiber at that endpoint. The horizon and word set are finite; the number of those

inequalities can be exponential. Standard real quantifier elimination proves the claim, with no missing infinite-horizon limit. Apply it successively to the finitely many allowed integer profiles or to padded constant caps. The held two-label output cut is then included in the reported peak.  $\square$

For a *given* machine and tolerance factor  $1 + \delta$ , choosing an even  $p \geq \log Q / \log(1 + \delta)$  gives lower and upper error certificates with ratio at most  $1 + \delta$ . This statement concerns checking the supplied machine, not finding a global moment minimizer. A certificate for a globally optimized bound must certify that optimization as well. A numerical local minimum has no such status. For fixed dimensions, exact rational coefficient growth in (14.6) is bounded by polynomially many input bit lengths, products and multinomial factors; certified root intervals can replace irrational  $p$ th roots in the output. The implementation below uses exact degree-two and degree-four recurrences, not an unreported large-degree optimization.

**14.2. Finite precision without changing boundary widths.** For an integer  $L \geq 1$ , let  $\mathcal{G}_L$  contain all machines whose initial and transition probabilities, and positive-answer probabilities, are integral multiples of  $1/L$ . Let  $e_L = \min_{\mathcal{M} \in \mathcal{G}_L} e(\mathcal{M})$  and set

$$(14.7) \quad C(K_\bullet) = (K_0 - 1) + \sum_{t=1}^N (K_t - 1) + 1.$$

These finite sets are nonempty even when some cuts have only one label.

**Theorem 14.4** (A finite two-sided rational certificate). *For every finite binary target, without any orthogonality assumption,*

$$(14.8) \quad \max\{0, e_L - C(K_\bullet)/L\} \leq \mathcal{E}(K_\bullet) \leq e_L.$$

*Consequently, for every tolerance  $\tau > \mathcal{E}$ , there is a dyadic machine with the same available boundary profile and error strictly less than  $\tau$ . The theorem does not assert a dyadic or rational exact optimizer at the boundary. If two target probability tables differ by at most  $\eta$  uniformly, their optimal errors at the same profile differ by at most  $\eta$ . Thus a strict positive feasibility or infeasibility margin persists under sufficiently small target perturbations.*

*Proof.* Round down the first  $k - 1$  coordinates of any probability row of size  $k$  to multiples of  $1/L$ , and give the remaining mass to the last coordinate. The row stays stochastic and its TV change is at most  $(k - 1)/L$ . Round a terminal positive-answer probability to within  $1/L$ . Couple the two machines at initialization and then at each common command, or telescope the products by TV contraction. For every word the total output change is bounded by (14.7) divided by  $L$ . Applying this construction to a minimizing real machine gives  $e_L \leq \mathcal{E} + C/L$ ; inclusion of the grid in the full parameter set gives  $\mathcal{E} \leq e_L$ . Both use the same labels; no scratch configuration is hidden in that assertion. Choosing  $L = 2^b$  sufficiently large proves the strict-margin assertion. Finally, for any fixed machine its errors against two targets differ by at most  $\eta$  by the triangle inequality. Minimize in both directions.  $\square$

An exhaustive grid search is a terminating certificate algorithm at fixed  $L$ , not an efficient proposal for large instances. If the target is supplied by certified intervals of radius  $\eta$ , the two endpoints in (14.8) can each be enlarged by  $\eta$ . Dyadic row sampling uses  $b$  fresh fair bits at denominator  $2^b$ ; index, table and scratch storage still have to be charged in an implementation. The separate uniform compiler below supplies its own workspace accounting. The rational approximation statement must not be strengthened to exact rational factorization: such field restrictions genuinely matter in nonnegative realization [21].

**Proposition 14.5** (An exact calibrated finite frontier). *Take  $N = 0$ , the four planar seeds  $\{\pm e_1, \pm e_2\}$ , two coordinate queries and  $0 < \rho \leq 1/2$  (this wider signal range is used only for the*

present finite example). If  $k$  is the available initialization alphabet, the optimal binary-TV errors are

$$(14.9) \quad \mathcal{E}(1) = \rho/2, \quad \mathcal{E}(2) = \rho/4, \quad \mathcal{E}(k) = 0 \quad (k \geq 3).$$

The same formula holds for an arbitrary horizon consisting only of identity commands, with  $k = \min_t K_t$ . In all cases the final held answer still costs two labels.

*Proof.* A one-state decoder is one point of the mean square. Approximating  $\pm \rho e_1, \pm \rho e_2$  by a common point has minimum uniform coordinate error  $\rho$ , attained at zero. For two states every seed mean lies on one affine line  $n \cdot y = c$ , with  $n \neq 0$ ; a degenerate segment may be embedded in any such line. If all coordinate errors are at most  $d$ , duality of  $\ell_\infty$  and  $\ell_1$  gives

$$d\|n\|_1 \geq \max_{y \in \{\pm \rho e_1, \pm \rho e_2\}} |n \cdot y - c| = \rho\|n\|_\infty + |c| \geq \frac{\rho}{2}\|n\|_1.$$

The diagonal segment with endpoints  $\pm(\rho/2, \rho/2)$  attains coordinate error  $\rho/2$ : encode  $e_1, e_2$  at its positive endpoint and their negatives at its negative endpoint. Divide mean errors by two to obtain binary TV. Three legal mean vectors  $(2\rho, 0), (-\rho, 2\rho), (-\rho, -2\rho)$  enclose all four targets, so their barycentric initialization gives exact responses. For identity-only words, condition at a narrowest cut to obtain the same one-cut lower bound. Initialize the appropriate construction and carry its label through all cuts, padding unused labels as necessary, to obtain the upper bound.  $\square$

**14.3. A complexity boundary with a different output convention.** The following comparison concerns one *categorical* terminal output, not a selected member of a family of binary queries. Its width is the pre-output memory alphabet; counting the held output would separately add its alphabet size.

**Proposition 14.6** (Normalized factorization and the real field). *Given an explicitly listed rational row-stochastic matrix  $P$ , deciding whether it has a stochastic factorization  $P = ED$  of inner dimension at most  $k$  is  $\exists\mathbb{R}$ -complete. It is already a zero-command numerical realization problem. A rational target can have no rational exact factorization at its minimum real inner dimension.*

*Proof.* Membership follows from the nonnegative variables, row-sum equations and quadratic equations  $P = ED$ . For hardness, remove the zero rows of a nonzero nonnegative rational matrix  $A$  and divide each remaining row by its positive row sum, giving  $P$ . Positive diagonal row scaling and removal of zero rows preserve nonnegative rank. If  $P = WH$  is a nonnegative factorization, put  $h_i = \sum_j H_{ij}$ . For  $h_i > 0$  set  $D_{ij} = H_{ij}/h_i$  and  $E_{xi} = W_{xi}h_i$ . If  $h_i = 0$ , replace that unused decoder row by any probability row and set its  $E$  column to zero. Then  $D\mathbf{1} = \mathbf{1}$  and  $E\mathbf{1} = P\mathbf{1} = \mathbf{1}$ . Thus stochastic inner dimension is exactly nonnegative rank, over either the real or rational field. Shitov's Theorem 2 gives  $\exists\mathbb{R}$ -hardness and his Corollary 3 gives rational matrices with different minimum ranks over the rational and real algebraic fields [21]. Row normalization transfers both statements as just shown. Zero input matrices can be treated separately and do not affect the reduction.  $\square$

This transfers a known complexity theorem; it is not a new hardness proof for nonnegative rank. Nor does it establish the same hardness for the fixed-dimensional orthogonal binary-query subclass. In that subclass Theorem 14.1 supplies the compact exact feasibility formula, while the inherited arithmetic theorems determine specific asymptotic widths. These are different assertions.

## 15. FINITE WORD PROFILES AND COMPATIBLE ENCLOSURES

The results in this section apply to every finite orthogonal alphabet, including noncommuting alphabets. No connectedness, arithmetic condition or spectral gap is assumed. Invariant directions may make a converse profile zero; this is a property of the stated profile, not an omitted hypothesis.

**15.1. An optimized distortion.** For  $B, k \geq 1$  and a probability vector  $p$  on the actual words  $\mathcal{A}^B$ , define

$$(15.1) \quad D_k(p) = 1 - \max_{\substack{u \in \mathbb{S}^{D-1} \\ v_1, \dots, v_k \in \mathbb{S}^{D-1}}} \sum_{w \in \mathcal{A}^B} p_w \max_{1 \leq j \leq k} \langle v_j, U_w u \rangle,$$

$$(15.2) \quad \Gamma_{\mathcal{A}}(k, B) = \max_{p \in \Delta(\mathcal{A}^B)} D_k(p), \quad \chi_{\mathcal{A}}(k, B) = -\log(1 - \Gamma_{\mathcal{A}}(k, B)).$$

We allow  $\chi = +\infty$ . Repeated product matrices can be merged for the optimization: their probability masses add. Their representative words must still be executed one ordinary command at a time.

**Proposition 15.1** (Finite-data properties). *The maxima in (15.1)–(15.2) are attained,  $0 \leq \Gamma_{\mathcal{A}}(k, B) \leq 1$ , and the profile is invariant under simultaneous orthogonal changes of coordinates. It is nonincreasing in  $k$  and unchanged by duplicating a letter. For two indexed orthogonal alphabets  $U, V$ ,*

$$(15.3) \quad |\Gamma_U(k, B) - \Gamma_V(k, B)| \leq B \max_a \|U_a - V_a\|_{\text{op}}.$$

*For a fully listed finite alphabet of real-algebraic matrices and fixed integers  $k, B, D$ ,  $\Gamma_{\mathcal{A}}(k, B)$  and an optimizing word law are real-algebraic and can be determined by real quantifier elimination.*

*Proof.* All parameter spaces are compact and the objective is continuous. Its maximum over directions is convex in  $p$ , so  $D_k$  is continuous and concave. The maximum scalar product is at most one. By choosing a single direction parallel to  $\sum p_w U_w u$ , the inner maximum is nonnegative; hence  $D_k \in [0, 1]$ . Conjugating all directions proves invariance. Additional centers only increase that inner maximum. Product telescoping gives  $\|U_w - V_w\|_{\text{op}} \leq B \max_a \|U_a - V_a\|_{\text{op}}$ , and maxima and minima preserve this uniform bound.

For the last assertion, enumerate the product matrices. Unit-sphere equations, simplex inequalities and finitely many maximum operations define a compact semialgebraic optimization with algebraic coefficients. Maximum operations can be expressed by finite alternatives or auxiliary variables. Quantifier elimination over the real closed field of real-algebraic numbers determines the endpoint of the feasible value interval and a witness [1]. The word table may have  $|\mathcal{A}|^B$  entries; this is not a polynomial-time assertion or a decision procedure for an infinite spectral gap. No logarithm needs to be eliminated: finite budget comparisons can instead use products of  $1 - \Gamma$ .  $\square$

**Lemma 15.2** (Endpoint contraction; inherited mechanism). *Let a complete random word  $W$  with law  $p$  be independent of all information before its start. Internal letters of  $W$  may be correlated. If the endpoint register has at most  $k$  labels, then*

$$(15.4) \quad q' \leq (1 - D_k(p))q.$$

*There is no bound on the intermediate registers.*

*Proof.* Write  $z_s = \mathbb{E}[Y \mid S = s]$  at the start. Composing all internal stochastic rows yields a common endpoint kernel  $T(j \mid s, w)$ . Since the entire word and the machine's fresh randomness are independent of the past,

$$z'_j = \mathbb{E}[U_W z_S \mid S' = j].$$

This is an endpoint identity; it would not generally hold after conditioning on a correlated next letter inside the packet. Choose  $v_j = z'_j / \|z'_j\|$  for nonzero  $z'_j$ , and any unit vector otherwise. Then

$$\begin{aligned} q' &= \sum_s \mathbb{P}(S = s) \sum_w p_w \sum_j T(j \mid s, w) \langle v_j, U_w z_s \rangle \\ &\leq \sum_s \mathbb{P}(S = s) \|z_s\| \max_{\|u\|=1} \sum_w p_w \max_j \langle v_j, U_w u \rangle \leq (1 - D_k(p))q. \end{aligned}$$

Terms with  $z_s = 0$  contribute zero. Replacing an executable assignment by the largest pairing is a relaxation in the correct direction. Arbitrary unobservable hidden coordinates are included in  $s$ ; no statewise observable-vector hypothesis is used. This is the conditional-centroid mechanism of [7, 8], now applied to the optimized finite word datum.  $\square$

For an available profile  $K_\bullet$ , let  $I_0 = 0$  and define

$$(15.5) \quad I_t = \max \left\{ I_{t-1}, \max_{1 \leq B \leq t} [I_{t-B} + \chi_{\mathcal{A}}(K_t, B)] \right\}, \quad \mathcal{I}_N(K_\bullet) = I_N.$$

If  $\Gamma_{\mathcal{A}}(K_t, B) = 1$ , the interval annihilates the conditional-centroid amplitude and is incompatible with any positive terminal calibration. This is the meaning of an infinite charge. The recurrence is the largest sum of charges of disjoint command intervals, with an interval of length  $B$  ending at  $t$  charged by  $\chi_{\mathcal{A}}(K_t, B)$ . The recurrence depends only on the alphabet and the available profile.

**Theorem 15.3** (A nonuniform interval budget). *Every machine with  $\kappa = \rho/\sqrt{D} - 2\varepsilon > 0$  satisfies*

$$(15.6) \quad \mathcal{I}_N(K_\bullet) \leq \log(1/\kappa).$$

*More generally its optimal actual terminal error obeys*

$$(15.7) \quad \mathcal{E}_N(K_\bullet) \geq \frac{1}{2} \left[ \frac{\rho}{\sqrt{D}} - e^{-\mathcal{I}_N(K_\bullet)} \right]_+.$$

*For a uniform peak bound  $K$ , every  $B \leq N$  gives*

$$(15.8) \quad \lfloor N/B \rfloor \chi_{\mathcal{A}}(K, B) \leq \log(1/\kappa).$$

*For a fixed  $B, k$  with  $\Gamma_{\mathcal{A}}(k, B) > 0$ ,*

$$(15.9) \quad \#\{1 \leq t \leq N : K_t \leq k\} \leq B - 1 + \frac{B \log(1/\kappa)}{\chi_{\mathcal{A}}(k, B)}.$$

*If  $\chi = +\infty$ , the last fraction is zero.*

*Proof.* Choose any maximizing disjoint interval family in (15.5). On each selected interval put an independently sampled optimizing complete word from Proposition 15.1. All gaps and the final remainder contain fixed deterministic command words. A deterministic orthogonal command followed by a stochastic update cannot increase the conditional-mean norm, by conditional Jensen. Lemma 15.2 therefore gives

$$q_N \leq \prod_{(t-B, t]} (1 - \Gamma_{\mathcal{A}}(K_t, B)) = e^{-\mathcal{I}_N(K_\bullet)},$$

since  $q_0 = 1$ . The chosen law can depend on the available layer sizes, not on private states. Wordwise correctness requires the same machine to pass it. Lemma 2.2 proves (15.6) and (15.7); taking consecutive intervals proves (15.8).

For (15.9), discard candidate endpoints before  $B$ . Select the earliest remaining endpoint, delete all later candidates less than  $B$  away, and repeat. If  $M$  is the original number and  $J$  the number selected, the explicit count is

$$M - (B - 1) \leq BJ.$$

The associated length- $B$  intervals are disjoint and each contributes at least  $\chi_{\mathcal{A}}(k, B)$ . Applying (15.6) gives the result. No restriction on the other registers was introduced.  $\square$

The cap union estimate below is the geometric certificate already used in the predecessor packet method [7]. The finite incidence formulation specifies how to optimize that certificate on a fully listed word set.

**Proposition 15.4** (Uniform cap certificates). *Suppose a law  $p$  on length- $B$  words satisfies*

$$\sup_{u,v \in \mathbb{S}^{D-1}} p\{w : \|U_w u - v\| < \delta\} \leq b.$$

*Then  $\Gamma_{\mathcal{A}}(k, B) \geq \frac{1}{2}\delta^2(1 - kb)_+$ . For algebraic product matrices and algebraic  $\delta$ , the best such  $b$  is the value of a finite linear program once cap incidence types have been decided by real quantifier elimination.*

*Proof.* The union of  $k$  open caps has probability at most  $kb$ . Outside it,  $1 - \max_j \langle v_j, U_w u \rangle = \frac{1}{2} \min_j \|U_w u - v_j\|^2 \geq \delta^2/2$ . For the algorithmic assertion, a subset of the finite word list either can or cannot be contained in one such cap for some  $u, v$ . This is a finite system of algebraic inequalities, including the sphere equations. For every feasible subset  $A$  impose  $\sum_{w \in A} p_w \leq b$ , along with the simplex constraints, and minimize  $b$ . Every actual incidence set is feasible; conversely a feasible subset is included in an actual incidence set and has no greater mass. Thus these constraints give exactly the desired supremum. The exponentially large enumeration is not concealed.  $\square$

**15.2. A separate exact-synthesis profile.** Fix  $0 < a < 1$ . For  $k$  such that the following family is nonempty, let

$$(15.10) \quad \lambda_{\mathcal{A}}(k, a) = \min_P \inf \{\lambda \geq 1 : U_b P \subseteq \lambda P \text{ for all } b \in \mathcal{A}\},$$

where  $P$  ranges over convex hulls of at most  $k$  points in  $\mathbb{B}^D$  containing  $a\mathbb{B}^D$ . An empty family has value  $+\infty$ .

**Theorem 15.5** (Compatible enclosure synthesis). *The finite value in (15.10) is attained and lies in  $[1, a^{-1}]$ . If  $\rho\lambda_{\mathcal{A}}(k, a)^N \leq a$ , there is an exact machine with  $\max(k, 2)$  labels, common command matrices and a horizon-dependent decoder. Each row has at most  $D + 1$  successors. Algebraic matrices and algebraic  $a$  give an algebraically decidable profile value and algebraic row witnesses. For two indexed alphabets with a common feasible family,*

$$(15.11) \quad |\lambda_U(k, a) - \lambda_V(k, a)| \leq a^{-1} \max_b \|U_b - V_b\|_{\text{op}}.$$

*Proof.* Parameterize  $P$  by  $k$  points in the closed unit ball, allowing repetition. Containment of  $a\mathbb{B}^D$  is closed in these coordinates. Every such  $P$  satisfies  $U_b P \subseteq \mathbb{B}^D \subseteq a^{-1}P$ . The containment conditions are closed on this compact set and  $1 \leq \lambda \leq a^{-1}$ ; hence a minimum exists.

Let its vertices be  $v_s$  and  $\lambda$  its dilation. Initialize with a convex representation of  $ax$ . For each  $b, s$ , represent  $\lambda^{-1}U_b v_s$  by the same vertices and use its convex coefficients as the row  $T_b(s, \cdot)$ . Induction gives

$$\mathbb{E}[v_{S_t} \mid w] = a\lambda^{-t}U_w x.$$

The decoder  $d_j(s) = \rho\lambda^N \langle e_j, v_s \rangle / a$  lies in  $[-1, 1]$  and gives (1.1). These are actual rows throughout the run, not independently chosen checkpoint factorizations. Carathéodory gives  $D + 1$  successors; there are  $|\mathcal{A}|k$  command rows, and their construction and storage remain uncharged.

The defining finite problem is semialgebraic: use convex coefficients for transformed vertices and a universal ball containment condition, then eliminate quantifiers. To prove (15.11), set  $e = \max_b \|U_b - V_b\|_{\text{op}}$ . For every feasible  $P$ ,  $V_b P \subseteq U_b P + e\mathbb{B}^D \subseteq (\lambda + e/a)P$ . Interchanging  $U, V$  completes the proof.  $\square$

The two profiles have different quantifiers. The converse tests every hidden machine; the enclosure theorem builds a particular positive realization. Neither theorem asserts that  $\lambda_{\mathcal{A}}$  is the optimum over hidden machines. This is why finite-scale lower and upper bounds can be used even when they do not match.

**Proposition 15.6** (Bounded word recoding). *Suppose every command of  $\mathcal{A}$  is a product of exactly  $L$  commands of an alphabet  $\mathcal{B}$  on the same space, with the same seeds and queries. Then*

$$\Gamma_{\mathcal{A}}(k, B) \leq \Gamma_{\mathcal{B}}(k, LB), \quad \lambda_{\mathcal{A}}(k, a) \leq \lambda_{\mathcal{B}}(k, a)^L, \quad W_{N, \varepsilon}(\mathcal{A}) \leq W_{LN, \varepsilon}(\mathcal{B}).$$

*An identity letter permits words of length at most  $L$  to be padded. If both alphabets contain an identity and admit such bounded recodings into the other, the existence and value of a logarithmic width exponent agree.*

*Proof.* Push forward a packet law along the word substitution to obtain the first inequality. Iterate the enclosure containments for the second. For the last inequality, compose the  $L$  stochastic rows of a  $\mathcal{B}$  machine into one row for the corresponding  $\mathcal{A}$  command, retaining only the block-end label. All rows may depend on the external epoch. This is an implementation of the atomic  $\mathcal{A}$  interface, not a claim that its physical letters cost  $L$  letters. Applying the two comparisons and monotonicity under final identity padding proves the exponent statement.  $\square$

**Corollary 15.7** (Packing–dilation balance). *Fix a calibrated experiment, an inradius  $a > \rho$ , and positive numbers  $u, v$ . Suppose that for every sufficiently large  $k$  there are a length  $B(k) \leq Ck^u$  and a word law for which*

$$\sup_{x, z \in \mathbb{S}^{D-1}} p\{w : \|U_w x - z\| < ck^{-v}\} \leq (2k)^{-1}.$$

*Suppose also that  $\lambda_{\mathcal{A}}(k, a) \leq \exp(C'k^{-u-2v})$ . Then*

$$W_{N, \varepsilon} = \Theta(N^{1/(u+2v)}).$$

*The same conclusion holds with  $u = 0$  when the lengths are uniformly bounded.*

*Proof.* Proposition 15.4 gives  $\Gamma(k, B(k)) \geq c^2 k^{-2v}/4$ . Equation (15.8) implies  $N \leq C_1 k^{u+2v}$  for a feasible peak  $k$ . Conversely choose  $k$  at a sufficiently large constant times  $N^{1/(u+2v)}$  so that  $C'Nk^{-u-2v} \leq \log(a/\rho)$ , and apply Theorem 15.5. The hypotheses concern finite word orbits and compatible convex containments, and impose no commutativity on the matrices.  $\square$

## 16. A DISTORTION–DILATION COMPARISON

The lower and upper profiles defined in Section 15 do not have identical quantifiers. They nevertheless obey a universal inequality. Define the optimized per-command distortion rate

$$(16.1) \quad \mathfrak{r}_{\mathcal{A}}(k) = \sup_{B \geq 1} \frac{-\log(1 - \Gamma_{\mathcal{A}}(k, B))}{B}.$$

All logarithms in this paper are natural. The value  $+\infty$  has its usual extended meaning.

**Theorem 16.1** (Distortion bounds every common enclosure). *For every finite orthogonal alphabet and every feasible  $k, a$  in (15.10),*

$$(16.2) \quad \mathfrak{r}_{\mathcal{A}}(k) \leq \log \lambda_{\mathcal{A}}(k, a).$$

*In particular, every finite-word certificate gives the lower enclosure bound  $\lambda_{\mathcal{A}}(k, a) \geq (1 - \Gamma_{\mathcal{A}}(k, B))^{-1/B}$ . The assertion does not require commuting commands, a spectral gap, or any Diophantine hypothesis.*

*Proof.* Take any feasible polytope  $P = \text{conv}(v_1, \dots, v_k)$ , contained in the unit ball and containing  $a$  times that ball, with common dilation  $\lambda$ . Use the stochastic rows representing  $\lambda^{-1}U_b v_s$  by the same vertices. Fix a unit seed  $u$  and initialize the mean at  $au$ . This construction makes sense at every length, without imposing a terminal decoder. For each complete word  $w$  of length  $n$  it satisfies

$$\mathbb{E}[v_{S_n} \mid w] = a\lambda^{-n}U_w u.$$

Let  $Y_n = U_w u$  and  $q_n = \mathbb{E} \|\mathbb{E}[Y_n | S_n]\|$ . Pairing with the bounded vector  $v_{S_n}$  gives, under every word law,

$$(16.3) \quad q_n \geq \mathbb{E} \langle Y_n, v_{S_n} \rangle = a\lambda^{-n}.$$

Fix  $B$  and independently repeat an optimizing length- $B$  word law  $m$  times. The packet lemma applies at the  $m$  endpoints, each of which has at most  $k$  labels. Since  $q_0 = 1$ ,

$$a\lambda^{-mB} \leq q_{mB} \leq (1 - \Gamma_{\mathcal{A}}(k, B))^m.$$

Taking  $m$ th roots and letting  $m \rightarrow \infty$  removes the inradius factor. Hence  $\lambda^{-B} \leq 1 - \Gamma_{\mathcal{A}}(k, B)$ . Take the supremum over  $B$  and then the minimum over enclosures. In particular,  $\Gamma = 1$  is incompatible with a finite feasible enclosure at this  $k$ , so no undefined logarithmic operation was used.  $\square$

The removal of the initial inradius loss by repetition is important: a single-horizon application would leave an additive  $\log(1/a)$ . Theorem 16.1 is a comparison, not a duality asserting equality or a characterization of arbitrary hidden width. The Hankel normal form in Proposition 3.1 permits general future-response polytopes, whereas (15.10) uses one common polytope in the observed vector space.

**Corollary 16.2** (Nonuniform enclosure chains). *Suppose  $P_t \subset \mathbb{B}^D$  has at most  $K_t$  generators,  $a\mathbb{B}^D \subset P_0$ , and  $U_b P_{t-1} \subset \lambda_t P_t$  for every command. Then*

$$\mathcal{I}_N(K_\bullet) \leq \log(1/a) + \sum_{t=1}^N \log \lambda_t.$$

*Proof.* Use the cut-dependent convex rows and the initialization with mean  $au$ . The same induction gives mean  $a(\prod_t \lambda_t)^{-1} U_w u$ . Choose the independent packet law attaining the interval recurrence. The pairing argument bounds the final amplitude below by this factor, while the packet theorem bounds it above by  $\exp(-\mathcal{I}_N(K_\bullet))$ . Taking logarithms proves the claim.  $\square$

## 17. EXACT OPTIMIZATION OF PLANAR PACKETS

For a rotation alphabet, let  $z_1, \dots, z_M \in \mathbb{S}^1 \subset \mathbb{C}$  be the distinct product phases of words of length  $B$ , ordered cyclically. A common rotation of all phases can be absorbed into the decoder centers. The initial direction in (15.1) therefore disappears. Any probability distribution on these phases is implemented by a distribution on representative words of the same length.

Let  $\mathfrak{P}_k$  be the finite set of partitions of the cyclically ordered list into at most  $k$  nonempty consecutive arcs. For  $p \in \Delta_M$ , set

$$(17.1) \quad F_k(p) = \max_{\mathcal{P} \in \mathfrak{P}_k} \sum_{A \in \mathcal{P}} \left| \sum_{i \in A} p_i z_i \right|.$$

**Theorem 17.1** (Optimal word laws as a cone program). *For the planar alphabet above,*

$$(17.2) \quad \Gamma_{\mathcal{A}}(k, B) = 1 - \min_{p \in \Delta_M} F_k(p).$$

*Equivalently, minimize  $t$  subject to  $p \in \Delta_M$  and, for every  $\mathcal{P} \in \mathfrak{P}_k$ ,*

$$(17.3) \quad \left\| \sum_{i \in A} p_i (\Re z_i, \Im z_i) \right\|_2 \leq r_{\mathcal{P}, A} \quad (A \in \mathcal{P}), \quad \sum_{A \in \mathcal{P}} r_{\mathcal{P}, A} \leq t.$$

*This is a finite second-order cone program with attained optimum. A fixed law can be evaluated by cyclic segmentation in  $O(kM^3)$  arithmetic operations, including computation of all interval sums. The optimization over laws may have exponentially many partition constraints.*

*Proof.* The decoder centers are on the unit circle; maximizing their inner product is therefore exactly nearest chordal-distance assignment.

For fixed centers  $v_1, \dots, v_k$  assign each phase to a nearest center in chordal distance; equivalently maximize  $\Re(\bar{v}_j z_i)$ . Distinct centers have connected cyclic Voronoi arcs, bounded by the angular midpoints of adjacent centers. Ties can be assigned consecutively; coincident centers can be merged. For the induced partition, the contribution of a block  $A$  is at most  $|\sum_{i \in A} p_i z_i|$ .

Conversely, for any cyclic partition choose the center of each block in the direction of its weighted complex sum, arbitrarily if that sum vanishes. The best assignment to those centers has value at least the sum in (17.1). Thus the maximum over centers equals  $F_k(p)$ . This proves (17.2); introducing the norm epigraphs proves (17.3). Zero weights follow either directly or by continuity from positive weights.

For a fixed cyclic cut, prefix sums give every block vector in constant time. The usual recurrence for the best  $j$  blocks among the first  $m$  points checks at most  $M$  previous split positions. It costs  $O(kM^2)$  for that cut. Repeating for the  $M$  cuts costs  $O(kM^3)$ . This recurrence evaluates a given law; it is not a polynomial-size representation of all constraints in (17.3). Compactness of the simplex and continuity give attainment.  $\square$

Nearest-center partitions and vector quantization are classical [5]. The additional optimization here chooses the distribution of actual command words, and Theorem 15.3 turns its value into a constraint on arbitrary hidden registers. There is no assertion that static quantization alone characterizes the least whole-run width.

**Proposition 17.2** (A finite exact packet certificate). *For the alphabet consisting of identity and a quarter turn, at word length three the four distinct phases are  $1, i, -1, -i$ . Their uniform law is optimal and*

$$\Gamma(1, 3) = 1, \quad \Gamma(2, 3) = 1 - \frac{1}{\sqrt{2}}, \quad \Gamma(3, 3) = \frac{2 - \sqrt{2}}{4}, \quad \Gamma(k, 3) = 0 \quad (k \geq 4).$$

*Proof.* The distortion is concave in the phase law and invariant under its four cyclic rotations. Averaging any law over those rotations proves optimality of the uniform law. For two centers, adjacent pairs give total  $\sqrt{2}/2$ ; a triple and a singleton give only  $1/2$ . For three centers, one pair and two singletons give  $\sqrt{2}/4 + 1/2$ . For one center the sum is zero; four centers fit all phases. Formula (17.2) gives the values. A uniform phase law is realized by the four ordinary words with zero, one, two or three initial quarter turns and then idle commands.  $\square$

For algebraic rotation matrices, both the general quantifier-elimination formulation and the cone program have algebraic coefficients. No finite list of spherical harmonics is used. For merely computable-real entries, finite numerical approximations are diagnostics unless a rigorous error enclosure is supplied; (15.3) then quantifies how an entry enclosure propagates to a packet certificate.

**Corollary 17.3** (Separated phases). *If  $M \geq 2k$  distinct length- $B$  word phases have pairwise circular distance at least  $s \in (0, 1/2]$ , then  $\Gamma_{\mathcal{A}}(k, B) \geq s^2$ .*

*Proof.* Use their uniform law. Each open arc of circular radius  $s/2$  contains at most one phase. At least half the mass is at distance at least  $s/2$  from every one of  $k$  centers. On  $[0, 1/2]$ ,  $1 - \cos(2\pi t) = 2 \sin^2(\pi t) \geq 8t^2$ . The average loss is at least  $s^2$ . This is a lower certificate for (17.2), not necessarily its exact value.  $\square$

## 18. ARITHMETIC PROFILES AT EVERY FINITE SCALE

Now  $\mathcal{A}_\alpha = \{I, R_{2\pi\alpha_1}, \dots, R_{2\pi\alpha_r}\}$ , with  $r$  fixed. Define

$$(18.1) \quad \psi_\alpha(H) = \min_{0 < \|m\|_1 \leq H} \|m \cdot \alpha\|_{\mathbb{R}/\mathbb{Z}}, \quad \eta_\alpha(q) = \max_{1 \leq i \leq r} \|q\alpha_i\|_{\mathbb{R}/\mathbb{Z}}.$$

These are finite functions of the alphabet parameters. The first may vanish in the presence of a rational relation. The upper profile remains usable in that case.

**Theorem 18.1** (Finite arithmetic bounds). *Let  $\kappa = \rho/\sqrt{2} - 2\varepsilon > 0$ . For a machine of peak  $K$ , put  $L = \lceil (2K)^{1/r} \rceil$  and  $B = r(L-1)$ . Then*

$$(18.2) \quad \lfloor N/B \rfloor \psi_\alpha(B)^2 \leq \log(1/\kappa).$$

For every integer  $q \geq 4$ ,

$$(18.3) \quad N\eta_\alpha(q) \leq q^2/100 \implies W_{N,0}(\alpha) \leq q.$$

Consequently every finite planar rotation alphabet satisfies

$$(18.4) \quad W_{N,0}(\alpha) \leq \max(4, \lceil \sqrt{50N} \rceil).$$

All upper bounds use common command rows and hold also after adjoining reflection in the first coordinate axis.

*Proof.* Construct a word of length  $B$  in  $r$  consecutive segments of  $L-1$  positions. In segment  $i$  give  $J_i$  copies of the  $i$ th rotation followed by idle commands, independently and uniformly for  $J_i \in \{0, \dots, L-1\}$ . Differences of count vectors have  $\ell_1$  norm at most  $B$ . If  $\psi_\alpha(B) > 0$ , the  $L^r \geq 2K$  phases are distinct and separated by  $\psi_\alpha(B)$ . Corollary 17.3 and (15.8), with  $-\log(1-x) \geq x$ , give (18.2). If the minimum is zero the displayed inequality is automatic. All letters in a segment are actual commands, not a free count input.

For the upper bound let  $P_q$  be the regular  $q$ -gon inscribed in the unit circle. Choose nearest integers  $p_i$  to  $q\alpha_i$ , and put  $\delta_i = 2\pi(\alpha_i - p_i/q)$ , so  $|\delta_i| \leq \pi/q$ . Elementary chord geometry gives

$$(18.5) \quad R_{2\pi\alpha_i}P_q \subseteq \Lambda_i P_q, \quad \Lambda_i = \frac{\cos(\pi/q - |\delta_i|)}{\cos(\pi/q)}, \quad \log \Lambda_i \leq \frac{4\pi^2\eta_\alpha(q)}{q^2}.$$

Indeed the ray of a rotated vertex meets a neighboring edge at radial coordinate  $\cos(\pi/q)/\cos(\pi/q - |\delta_i|)$ . Integrating the derivative of  $\log \cos(\pi/q - t)$  bounds the logarithm by  $|\delta_i| \tan(\pi/q)$ , and  $\tan(\pi/q) \leq 2\pi/q$  for  $q \geq 4$ .

Take  $\Lambda = \max_i \Lambda_i$ , including one for the identity. Under (18.3),  $\Lambda^N \leq e^{\pi^2/25} < 2$ . Since  $P_q$  contains the half-unit disk, Theorem 15.5 with  $a = 1/2$  applies:  $\rho\Lambda^N \leq 1/5 < 1/2$ . Explicitly, initialize at mean  $x/2$ , represent  $\Lambda^{-1}R_{2\pi\alpha_i}v_s$  by the vertices for each row, and decode by  $2\rho\Lambda^N \langle e_j, v_s \rangle$ . This verifies every numerical word, including any deterministic remainder of a converse packet decomposition.

There are  $(r+1)q$  common command rows and two decoder columns, each row having at most three successors. These counts exclude row-description and arithmetic costs. Since  $\eta_\alpha(q) \leq 1/2$ , (18.4) follows. Reflection preserves  $P_q$ ; its contracted image uses the same vertex set. The reflected lower bound is obtained simply by restricting to rotation-only words.  $\square$

The pair  $(\psi_\alpha, \eta_\alpha)$  is a specialization of the word-distortion and enclosure profiles. It is not assumed to characterize exact finite width. In particular, (18.2) can be weak when long word packets have clusters; the optimized profile of Section 17 is available without that separation relaxation.

**Lemma 18.2** (Denominators under a general power bound). *Suppose for  $\tau \geq r$  and  $b > 0$ ,*

$$(18.6) \quad \|m \cdot \alpha\|_{\mathbb{R}/\mathbb{Z}} \geq b \|m\|_1^{-\tau} \quad (m \neq 0).$$

For every integer  $Q \geq 2$  there is  $q$  with

$$(18.7) \quad cQ^{r/(\tau+1-r)} \leq q \leq Q^r, \quad \eta_\alpha(q) \leq Q^{-1}, \quad c = (b/r^{\tau+1})^{r/(\tau+1-r)}.$$

*Proof.* The  $Q^r + 1$  points  $j\alpha$  in  $Q^r$  equal torus boxes give integers  $p_i, q$  with  $1 \leq q \leq Q^r$  and  $|q\alpha_i - p_i| \leq Q^{-1}$ . For  $M = \lfloor q^{1/r} \rfloor$ , the  $(M+1)^r > q$  integer vectors in  $\{0, \dots, M\}^r$  have a residue collision for  $m \cdot p \pmod{q}$ . Subtracting gives  $m \neq 0$  with  $H = \|m\|_1 \leq rq^{1/r}$  and  $q \mid m \cdot p$ . Thus

$$bH^{-\tau} \leq \|m \cdot \alpha\|_{\mathbb{R}/\mathbb{Z}} \leq H/(qQ), \quad bqQ \leq r^{\tau+1}q^{(\tau+1)/r}.$$

Rearranging yields (18.7). This direct residue argument supplies exactly the quantitative transference needed here, without an unspecified transference constant.  $\square$

**Corollary 18.3** (Power bounds at finite Diophantine type). *Under (18.6), for constants depending on the fixed data,*

$$(18.8) \quad c_1 N^{r/(2\tau+1)} \leq W_{N,\varepsilon}(\alpha) \leq C_1 N^{\min\{1/2, r(\tau+1-r)/(\tau+1+r)\}}.$$

*No matching claim is made when  $\tau > r$ .*

*Proof.* From (18.2),  $N < B + b^{-2}B^{2\tau+1} \log(1/\kappa)$ , while  $B \leq r(2K)^{1/r}$ . This gives the lower power after adjusting finitely many initial horizons. Lemma 18.2 bounds  $N\eta_\alpha(q)/q^2 \leq c^{-2}NQ^{-1-2r/(\tau+1-r)}$ . Choose  $Q$  at constant times  $N^{(\tau+1-r)/(\tau+1+r)}$ , large enough to satisfy (18.3) and  $q \geq 4$ . Then  $q \leq Q^r$  proves the second candidate upper power; (18.4) proves the first.  $\square$

**Corollary 18.4** (The retained badly approximable law). *If (18.6) holds with  $\tau = r$ , then  $W_{N,\varepsilon}(\alpha) = \Theta(N^{r/(2r+1)})$ , also with the reflection extension. This is the previously established matched law [8].*

*Proof.* Substitute  $\tau = r$  in (18.8); the two powers coincide.  $\square$

For  $1, \alpha_1, \dots, \alpha_r$  rationally independent, every optimized finite-block action gap on the circle is zero. Indeed simultaneous approximation supplies unbounded frequencies  $n$  with every  $n\alpha_i$  close to an integer. Each fixed rotation word then acts almost identically on its Fourier mode. For the reflected alphabet use the real cosine mode, which reflection fixes. This is the established almost-invariance argument, not a claim that the finite packet distortion vanishes.

## 19. MINIMAL ORDINARY TYPE AND ALMOST-EVERYWHERE WIDTHS

Uniform bad approximation is substantially stronger than the condition needed to determine the logarithmic width exponent. We first deduce the minimal-type assertion without any measure-theoretic hypothesis.

*Proof of (13.2) in Theorem 13.1.* If  $\omega^*(\alpha) = r$ , then for every  $h > 0$  there is  $b_h > 0$  such that

$$\|m \cdot \alpha\|_{\mathbb{R}/\mathbb{Z}} \geq b_h \|m\|_1^{-r-h} \quad (m \neq 0).$$

Indeed the definition of the limsup gives this inequality for all sufficiently large  $\|m\|_1$ , and rational independence makes the finite remaining minimum positive. Apply Corollary 18.3 with  $\tau = r + h$ . It gives

$$\frac{r}{2r+2h+1} \leq \liminf_N \frac{\log W_{N,\varepsilon}}{\log N} \leq \limsup_N \frac{\log W_{N,\varepsilon}}{\log N} \leq \min \left\{ \frac{1}{2}, \frac{r(1+h)}{2r+1+h} \right\}.$$

For  $0 < h < 1/(2r-1)$  the second branch is smaller. Let  $h \downarrow 0$ . This proves (13.2). No positive uniform constant at exponent  $r$  was used.  $\square$

For completeness we prove a more quantitative metric statement, using only the summable part of metric Diophantine approximation. It does not claim a new metric theorem.

**Lemma 19.1** (A summable logarithmic separation). *For Lebesgue-almost every  $\alpha \in [0, 1]^r$ , and every  $s > 1$ , there is  $b_s > 0$  such that*

$$(19.1) \quad \|m \cdot \alpha\|_{\mathbb{R}/\mathbb{Z}} \geq \frac{b_s}{\|m\|_1^r (\log(2 + \|m\|_1))^s} \quad (m \neq 0).$$

Such a vector has  $\omega^*(\alpha) = r$ .

*Proof.* For a nonzero integer vector  $m$ , the map  $\alpha \mapsto m \cdot \alpha \pmod{1}$  sends Haar measure on the torus to uniform circle measure. There are at most  $C_r H^{r-1}$  integer vectors with  $\|m\|_1 = H$ . Hence the measures of the events

$$\|m \cdot \alpha\|_{\mathbb{R}/\mathbb{Z}} < [\|m\|_1^r (\log(2 + \|m\|_1))^s]^{-1}$$

have finite sum, bounded by a constant times  $\sum_{H \geq 1} [H(\log(2 + H))^s]^{-1}$ . The measure of their tail union tends to zero. Almost every vector violates the unscaled inequality only finitely often. Remove the measure-zero rational hyperplanes and reduce the constant to cover the finite exceptions. Intersect the full-measure sets for rational  $s > 1$ . For a real  $s > 1$ , choose a rational  $1 < s_0 < s$ ; its stronger bound implies the bound with  $s$ , after a harmless change of constant. This gives the assertion on one full-measure set.

This bound implies  $\omega^*(\alpha) \leq r$ . For the reverse inequality, place the  $(M + 1)^r$  values  $m \cdot \alpha \pmod{1}$ ,  $m \in \{0, \dots, M\}^r$ , in fewer than  $(M + 1)^r$  intervals of equal length. A collision gives a nonzero difference of norm at most  $rM$  and distance  $O(M^{-r})$ . As  $M \rightarrow \infty$  the norms of some such differences must tend to infinity, since rational independence excludes an exact relation. Thus  $\omega^*(\alpha) \geq r$ .  $\square$

**Proposition 19.2** (Quantitative bounds under logarithmic separation). *Suppose (19.1) holds for some  $s > 0$  and  $b_s > 0$ . Then there are  $c, C > 0$  such that*

$$(19.2) \quad cN^{r/(2r+1)}(\log(N+2))^{-2rs/(2r+1)} \leq W_{N,\varepsilon}(\alpha) \leq CN^{r/(2r+1)}(\log(N+2))^{2r^2s/(2r+1)}.$$

*Proof.* The packet bound (18.2) gives

$$N < B + b_s^{-2} B^{2r+1} (\log(2 + B))^{2s} \log(1/\kappa).$$

With  $B \leq r(2K)^{1/r}$  this implies  $N \leq C_0 K^{(2r+1)/r} (\log(2 + K))^{2s}$ . If  $K > N$ , the proposed lower bound is already true after reducing its constant. Otherwise  $\log(2 + K) \leq \log(2 + N)$ , and rearrangement proves the first inequality of (19.2).

For the upper inequality repeat the residue argument of Lemma 18.2. It produces  $q \leq Q^r$ ,  $\eta_\alpha(q) \leq Q^{-1}$  and  $H \leq rq^{1/r}$  with

$$\frac{b_s}{H^r (\log(2 + H))^s} \leq \frac{H}{qQ}.$$

Since  $q \leq Q^r$ , it follows for  $Q \geq 2$  that

$$(19.3) \quad q \geq c_0 \frac{Q^r}{(\log(Q+2))^{rs}}.$$

Indeed  $b_s q Q \leq H^{r+1} (\log(2 + H))^s \leq C_r q^{1+1/r} (\log(Q+2))^s$ , and raising the rearranged inequality to the  $r$ th power gives (19.3).

Choose an integer

$$Q = \left\lceil C_2 N^{1/(2r+1)} (\log(N+2))^{2rs/(2r+1)} \right\rceil.$$

For sufficiently large fixed  $C_2$ , (19.3) implies  $N\eta_\alpha(q)/q^2 \leq 1/100$  and  $q \geq 4$ . To see the constant choice explicitly, for large  $N$  one has  $\log(Q+2) \leq C_3 \log(N+2)$  with  $C_3$  fixed; thus the ratio is bounded by a constant times  $C_2^{-(2r+1)}$ . Enlarge the final constant for the finitely many initial horizons. Theorem 18.1 gives an exact machine with  $q \leq Q^r$  labels, proving the upper bound.  $\square$

*Completion of Theorem 13.1.* Apply Lemma 19.1 with  $s = 1 + \delta$  and then Proposition 19.2. This yields (13.3) on one full-measure set for all fixed  $\delta > 0$ . These statements concern a fixed alphabet before the horizon grows. None of their constants is asserted uniform in  $r$  or near a rational relation.  $\square$

**19.1. Rational dependencies without redundant exponents.** The number of listed rotations need not be the relevant arithmetic rank. The following result covers finite rational reparametrizations rather than assigning one exponent to every redundant letter.

**Corollary 19.3** (Finite rational extensions). *Suppose an alphabet contains  $I$  and rotations through angles  $\beta_1, \dots, \beta_d$ , where  $1, \beta_1, \dots, \beta_d$  are rationally independent and  $\omega^*(\beta) = d$ . Suppose every additional angle belongs to  $\mathbb{Q} + \sum_{j=1}^d \mathbb{Q}\beta_j$ . Then*

$$\lim_N \frac{\log W_{N,\varepsilon}}{\log N} = \frac{d}{2d+1}.$$

*The same assertion holds on a full-measure set of  $\beta$  with the logarithmic upper and lower powers of (13.3), replacing  $r$  by  $d$ . Reflection may be added. If all angles are rational, exact width is bounded independently of  $N$ .*

*Proof.* Restriction to the displayed  $d$ -angle subalphabet gives the lower bound. Choose one common denominator  $h$  and integers  $t_i, n_{ij}$  so that every angle is  $t_i/h + \sum_j n_{ij}\gamma_j$  with  $\gamma_j = \beta_j/h$ . Multiplication by  $h$  shows  $\|m \cdot \gamma\|_{\mathbb{R}/\mathbb{Z}} \geq h^{-1} \|m \cdot \beta\|_{\mathbb{R}/\mathbb{Z}}$ . Thus every separating power or logarithmic inequality for  $\beta$  holds for  $\gamma$  with a changed constant.

Take a simultaneous approximant  $p_j/q$  to  $\gamma$  from the proof above. A command's irrational part has residual angle at most  $2\pi(\sum_j |n_{ij}|)/(qQ)$ . The same polygon construction therefore works, after changing only fixed constants. Store in addition the rational phase  $c \in \mathbb{Z}/h\mathbb{Z}$ , updated by  $c \mapsto c + t_i$ . The decoder rotates the stored polygon vector by  $2\pi c/h$ . The peak, including the held answer, is  $\max(2, hq)$ , not an uncharged phase. Reflection sends  $c$  to  $-c$  and conjugates the polygon phase; the regular polygon and the construction are reflection-invariant. The exact expectation identity is unchanged. The finite multiplier  $h$  does not change a logarithmic exponent or the displayed logarithmic powers.

For a purely rational alphabet, choose a common period  $h$  and retain the seed identifier and the accumulated phase. This uses at most  $4h$  command labels for our four planar seeds, plus the held binary answer convention. No infinite-gap hypothesis is involved.  $\square$

## 20. EVERY IRRATIONAL ANGLE HAS CUBIC SCALES

We keep the two-letter alphabet  $\{I, R_{2\pi\alpha}\}$  and the planar resource and error conventions. No Diophantine lower bound is assumed in this section. The needed continued-fraction estimates are classical [3]; we include their short proof to specify the exact finite inequalities being used.

**Lemma 20.1** (Convergent separation). *Let  $p_n/q_n$  be the convergents of an irrational  $\alpha$ , ignoring the finitely many initial indices with repeated denominator. Then*

$$(20.1) \quad \frac{1}{q_{n+1} + q_n} < |q_n\alpha - p_n| < \frac{1}{q_{n+1}}, \quad \min_{1 \leq l < q_n} \|l\alpha\|_{\mathbb{R}/\mathbb{Z}} > \frac{1}{q_n + q_{n-1}} > \frac{1}{2q_n}.$$

*Proof.* The convergent recurrences imply  $p_n q_{n-1} - p_{n-1} q_n = \pm 1$  and alternating signs of  $e_n = q_n \alpha - p_n$ . If  $\xi$  is the next complete quotient, then

$$\alpha = \frac{p_n \xi + p_{n-1}}{q_n \xi + q_{n-1}}, \quad a_{n+1} < \xi < a_{n+1} + 1.$$

It follows that  $|e_n| = 1/(q_n \xi + q_{n-1})$ , proving the first pair.

For  $0 < l < q_n$  and an integer  $p$ , the determinant identity writes  $(p, l) = a(p_n, q_n) + b(p_{n-1}, q_{n-1})$  with integers  $a, b$ . Here  $b \neq 0$ ;  $a, b$  cannot both be positive, and cannot both be negative. If  $a = 0$  the error has modulus at least  $|e_{n-1}|$ . Otherwise they have opposite signs, so  $ae_n$  and  $be_{n-1}$  have the same sign and again  $|l\alpha - p| \geq |e_{n-1}|$ . Minimize over  $p$  and apply the first pair with index  $n - 1$ . This proves the remaining inequalities.  $\square$

**Theorem 20.2** (Universal cubic subsequences). *Set  $\kappa = \rho/\sqrt{2} - 2\varepsilon > 0$ ,  $C_\kappa = 8(1 + \log(1/\kappa))$ , and*

$$M_n = \lceil C_\kappa q_n^3 \rceil, \quad m_\kappa = \lceil 100(C_\kappa + 1) \rceil.$$

*For all sufficiently large  $n$ ,*

$$(20.2) \quad q_n/2 < W_{M_n, \varepsilon}(\alpha) \leq m_\kappa q_n.$$

*The upper bound is exact, uses common command rows and requires no regularity of the partial quotients.*

*Proof.* Use words of length  $B = q_n - 1$  whose number of initial active letters ranges uniformly over  $0, \dots, q_n - 1$ . Lemma 20.1 gives  $q_n$  distinct phases with separation greater than  $1/(2q_n)$ . If the proposed peak  $K$  were at most  $q_n/2$ , Corollary 17.3 and the packet budget would give

$$\left\lfloor \frac{M_n}{q_n - 1} \right\rfloor \frac{1}{4q_n^2} \leq \log(1/\kappa).$$

The left side is at least  $C_\kappa/4 - 1/(4q_n^2)$ , strictly greater than the right side. This proves the lower bound against all hidden-state machines.

For the upper bound use the polygon denominator  $Q_n = m_\kappa q_n$ . Distance to the nearest integer is subadditive, so

$$\|Q_n \alpha\|_{\mathbb{R}/\mathbb{Z}} \leq m_\kappa |q_n \alpha - p_n| < m_\kappa / q_{n+1} \leq m_\kappa / q_n.$$

Consequently  $Q_n^2 / \lceil 100 \|Q_n \alpha\|_{\mathbb{R}/\mathbb{Z}} \rceil > m_\kappa q_n^3 / 100 \geq (C_\kappa + 1) q_n^3 \geq M_n$ . Equation (18.3) gives the exact upper construction. The constants have been chosen for transparency, not optimality.  $\square$

**Proposition 20.3** (Resonant long horizons). *For any irrational  $\alpha$  and integer  $q \geq 4$  with  $N_q = \lfloor q^2 / (100 \|q\alpha\|_{\mathbb{R}/\mathbb{Z}}) \rfloor \geq 1$ , one has  $W_{N_q, 0}(\alpha) \leq q$ . If  $\alpha$  is Liouville, there is a sequence  $N \rightarrow \infty$  on which  $\log W_{N, \varepsilon}(\alpha) / \log N \rightarrow 0$ .*

*Proof.* The first assertion is the finite enclosure implication (18.3). For a Liouville number choose arbitrarily large  $A$  and unbounded denominators  $q$  with  $0 < \|q\alpha\|_{\mathbb{R}/\mathbb{Z}} < q^{-A}$ . The first assertion then supplies horizons  $N_q \geq q^{A+2}/200$  for large  $q$ , with width at most  $q$ . The logarithmic ratio is at most  $1/(A+2) + o(1)$ . A diagonal sequence with  $A \rightarrow \infty$  proves the claim, since widths are at least two.  $\square$

*Proof of Theorem 13.2.* Equation (20.2) and  $M_n = \Theta(q_n^3)$  give the first assertion and a subsequence of logarithmic exponent  $1/3$ . Proposition 20.3 gives the zero lower limit for Liouville angles, and (18.4) gives the upper limit at most  $1/2$ .

It remains to verify unboundedness, which does not follow from a subsequence alone. Fix any proposed width  $K$ . The  $2K$  phases of counts  $0, \dots, 2K - 1$  are distinct for an irrational angle and have some positive minimum separation  $s_K$ . Independent words of length  $2K - 1$  imply  $\lfloor N/(2K - 1) \rfloor s_K^2 \leq \log(1/\kappa)$  for a width- $K$  machine. This fails for all sufficiently large  $N$ . Thus  $W_{N, \varepsilon} \rightarrow \infty$ .  $\square$

**Corollary 20.4** (An explicit fluctuating alphabet). *Let  $\alpha_L = \sum_{j=1}^{\infty} 10^{-j!}$ . For  $q_m = 10^{m!}$  and all large  $m$ , the horizon*

$$N_m = \lfloor q_m^{m+2} / 200 \rfloor$$

*admits an exact machine with at most  $q_m$  labels. The same fixed two-letter alphabet also has the cubic subsequences of Theorem 20.2; hence no single logarithmic width exponent exists. Adding reflection preserves these conclusions.*

*Proof.* The first  $m$  summands have denominator dividing  $q_m$ , and their remaining positive tail is less than  $2 \cdot 10^{-(m+1)!}$ . Thus  $0 < \|q_m \alpha_L\|_{\mathbb{R}/\mathbb{Z}} < 2q_m^{-m}$  for large  $m$ , proving the horizon bound by (18.3). The same tail estimates preclude rationality and show the Liouville property. The general cubic theorem applies. Reflection leaves the upper polygons invariant, while the rotation-only test words give the same lower bounds.  $\square$

These theorems describe different scales of one and the same experiment. A small label count on a long resonant horizon is not a constant-width simulator for all horizons, and it need not persist across the gaps between successive approximants. The exact finite optimum and the Liouville upper logarithmic limit remain separate quantitative questions. None of the proofs substitutes approximate output for the exact upper realizations.

## 21. FINITE RANDOM BITS AND UNIFORM SPACE

Exact real rows and finite-bit algorithms are different resources. We first give a finite perturbation statement and then remove the row table from a concrete uniform construction. Only this section uses a computational model. It does not change the exact atomic-row conclusions proved elsewhere.

**Lemma 21.1** (Dyadic perturbation of a stochastic transducer). *Suppose initialization and every command row have at most  $s$  nonzero entries. They can be replaced by denominator- $2^b$  rows, on the same label sets, with row-TV error at most  $(s-1)2^{-b}$ . Rounding the terminal binary probabilities adds at most  $2^{-b}$ . Consequently every numerical word changes by at most  $((N+1)(s-1)+1)2^{-b}$  in binary total variation. Each rounded row can be sampled with exactly  $b$  independent fair bits.*

*Proof.* Order the  $s$  successors, round the first  $s-1$  probabilities down to multiples of  $2^{-b}$ , and give the remaining mass to the last successor. The total removed mass is at most  $(s-1)2^{-b}$ , which is exactly the row total variation. Couple the initializations and then the command updates, or telescope products using TV contraction, to sum their worst-row errors. A terminal Bernoulli parameter can be rounded down with error at most  $2^{-b}$ . A uniform integer in  $\{0, \dots, 2^b - 1\}$  selects one of the cumulative dyadic intervals. This uses a fixed number of bits, with no retained random tape or rejection-time side channel.  $\square$

For a table implementation the storage is  $O(|\mathcal{A}|Ks(\log K + b))$  bits for common rows of width  $K$ , in addition to initialization and decoder data. Cut-dependent rows can require a factor  $N$  more storage. Lemma 21.1 is the elementary stochastic-rounding principle; it is not claimed as a new finite-bit simulation method. The next construction uses no such table.

**Definition 21.2** (Uniform command-stream model). A fixed finite program receives  $N$  in binary, one of the four planar seeds, exactly  $N$  one-way command symbols, and one final coordinate query. It has a work tape and fresh independent fair bits. The command input cannot be reread. All counters, arithmetic registers, stored parameters, random words and output storage count toward work space. An external source waits while a command is processed; waiting time is not an input or a persistent storage channel. The finite control is fixed independently of  $N$ . The program has no horizon-dependent advice or transition table.

An angle has an admissible digit procedure if a fixed program computes a dyadic approximation within  $2^{-p}$  in polynomial time in  $p$  and  $O(p)$  work bits. The same fixed  $r$  is used throughout. This explicit hypothesis excludes a free oracle for an arbitrary noncomputable real. Algebraic angles, in particular the examples below, meet it by rational bisection with fixed-degree integer arithmetic.

**Theorem 21.3** (A table-free finite-bit compiler). *Fix  $r$  angles with admissible digit procedures, rational  $0 < \rho \leq 1/10$ , and rational  $0 < \varepsilon < 1$ . There is a uniform program in Definition 21.2 whose every binary output row has TV error at most  $\varepsilon$  from the rotation experiment. It uses*

$$(21.1) \quad O_{r,\alpha,\rho}(\log(N+2) + \log(1/\varepsilon)) \text{ work bits}, \quad (N+2)b \text{ fair bits},$$

where  $b = \lceil \log_2(16(N+2)/\varepsilon) \rceil$ . Its total running time is polynomial in  $N + \log(1/\varepsilon)$  for fixed digit procedures. Its clean stochastic register has  $q+1$  labels, including one central label, with  $q = O(\sqrt{N})$ .

If a denominator  $q_0 \geq 4$  satisfies  $N\eta_\alpha(q_0) \leq q_0^2/2000$ , the search may choose  $q \leq q_0$ . Consequently the clean-register upper exponents and logarithmic bounds of Theorem 13.1 hold for this compiler as well, with changed constants. These clean labels do not include its scratch space or deterministic parameter registers; the full work bound is (21.1).

*Proof.* We give the parameter search, real reference rows, dyadic rounding, and finite-space evaluation separately.

*Parameter search.* Put  $p = \lceil \log_2(16000(N+2)/\varepsilon) \rceil$  and obtain dyadic  $\tilde{\alpha}_i$  with error at most  $2^{-p}$ . For  $q = 4, 5, \dots$  compute, by integer arithmetic,

$$e(q) = \max_i \|q\tilde{\alpha}_i\|_{\mathbb{R}/\mathbb{Z}} + q2^{-p}.$$

Stop at the first  $q$  for which  $Ne(q) \leq q^2/1000$ . The test is rational and certifies both  $\eta_\alpha(q)$  and  $\eta_{\tilde{\alpha}}(q)$  by the Lipschitz property of distance to  $\mathbb{Z}$ . It succeeds by  $q = \max(4, \lceil \sqrt{2000(N+2)} \rceil)$ :  $\|q\tilde{\alpha}_i\| \leq 1/2$  and the added term is smaller than one at this bound. If  $q_0$  satisfies the stronger displayed condition, then  $e(q_0) \leq \eta_\alpha(q_0) + 2q_02^{-p}$  and  $2Nq_02^{-p} \leq q_0^2/2000$ . Hence that candidate succeeds. No exact comparison of an irrational angle with a rational sector boundary is performed.

*Reference rows with slack.* Use the regular  $q$ -gon  $v_j = (\cos(2\pi j/q), \sin(2\pi j/q))$  and one additional label  $*$  representing zero. Set  $\Lambda = 1 + 1/(4N)$ . For  $N \geq 1$  the dilation estimate (18.5) and the search condition give

$$\log \Lambda \geq 1/(5N) > 4\pi^2\eta_{\tilde{\alpha}}(q)/q^2.$$

Thus  $\Lambda^{-1}R_{2\pi\tilde{\alpha}_i}v_j$  lies strictly inside the polygon. The case  $N = 0$  is handled by direct rounding of the seed response and a fixed amount of space. Write  $q\tilde{\alpha}_i = n_i + t_i$ , with integer  $n_i$  and  $0 \leq t_i < 1$ , and put  $h = 2\pi/q$ . The command row has successors  $j + n_i$ ,  $j + n_i + 1$  modulo  $q$ , and  $*$ , with respective weights

$$(21.2) \quad A_i = \frac{\sin((1-t_i)h)}{\Lambda \sin h}, \quad B_i = \frac{\sin(t_i h)}{\Lambda \sin h}, \quad 1 - A_i - B_i.$$

They are nonnegative by the dilation bound, and their vector mean is the required contracted rotation. The row at  $*$  stays there. Initialize with vector mean  $x/2$  using the same chord formula with radial factor  $1/2$  and the exact seed angle  $0, 1/4, 1/2, 3/4$ . Decode with mean  $2\rho\Lambda^N \langle e_j, v_s \rangle$ , or zero at  $*$ . Since  $\Lambda^N \leq e^{1/4}$ , all decoder means are strictly inside  $[-1, 1]$ . The usual mean induction realizes every word for  $\tilde{\alpha}$  exactly. These are reference real rows only.

*Certified finite rows.* Compute outward intervals for the first two weights in each row, each of width at most  $2^{-b}$ . Clip their lower endpoints at zero and round down to denominator  $2^b$ ; assign the residue to  $*$ . Relative to the reference row, at most  $4 \cdot 2^{-b}$  mass is moved. Evaluate and round each terminal Bernoulli parameter within  $2 \cdot 2^{-b}$ . TV contraction therefore bounds initialization,  $N$  updates and final decoding by  $(4N+6)2^{-b}$ . Replacing the true angles by their dyadic approximants changes a product rotation by at most  $2\pi N 2^{-p}$  in operator norm, and changes a target binary row by at most  $\pi\rho N 2^{-p}$ . The sum is less than  $\varepsilon$  with the chosen  $p, b$ . Use exactly  $b$  fair bits at each draw, including dummy bits when the central row is deterministic. The random-bit count is  $(N+2)b$ .

*Work and uniform evaluation.* Only  $q$ , the current vertex or central flag, the horizon and command counter, the  $r$  dyadic angles, and  $O(r)$  offsets and dyadic weights are retained. The latter specify circulant rows, not a  $q$ -row table. The denominator search uses integers of  $O(p + \log q)$  bits. Elementary sine and cosine on bounded intervals can be evaluated by fixed-point Taylor recurrences with outward rounding;  $\pi$  can be evaluated by Machin's arctangent identity. There are  $O(b + \log q)$  terms and only a fixed number of live accumulators. Since  $\sin(2\pi/q) \geq 4/q$  for  $q \geq 4$ , division loses only  $O(\log q)$  bits. Repeated squaring evaluates  $\Lambda^N$  with  $O(\log N)$  stages, all intermediate true values between 1 and  $e^{1/4}$ . An additional  $O(\log(N+2) + \log(b+2))$  guard bits bounds accumulated rounding errors. Interval precision can instead be doubled until the stated widths are certified; these error estimates bound the final precision by the same order. No integer of unbounded exact denominator is stored. Every live register has  $O(b + p + \log N)$  bits, and the number of registers is fixed for fixed  $r$ . This proves the work bound and a polynomial time bound, including parameter search. Finally the finite arithmetic upper proofs remain valid with  $1/100$  replaced by  $1/2000$ , since only fixed constants change. Their candidate denominators therefore bound the search.  $\square$

The proof describes a fixed-point implementation, not the numerical behavior of a floating-point library. The accompanying integer-interval reference code certifies each rounded row and uses fresh bounded-length random integers; its Python object overhead is not identified with the Turing work-space constant.

**Corollary 21.4** (Sharp work-space order for explicit alphabets). *Let  $\alpha_i = 2^{i/(r+1)}$ ,  $1 \leq i \leq r$ , and fix rational  $0 < \varepsilon < \rho/(2\sqrt{2})$  and  $0 < \rho \leq 1/10$ . In Definition 21.2, numerical simulation requires and suffices with  $\Theta(\log N)$  work bits. More precisely every fixed program using  $s(N)$  work bits obeys*

$$(21.3) \quad s(N) \geq \frac{r}{2r+1} \log_2 N - O(\log \log(N+2)).$$

*The compiler's clean register has  $O(N^{r/(2r+1)})$  labels. This is not a claim that all its internal fair-bit states retain that label exponent or that the leading coefficient of work space is optimal.*

*Proof.* Put  $\theta = 2^{1/(r+1)}$ . For a nonzero integer vector  $m$  with  $H = \|m\|_1 \geq 1$ , choose the integer  $p$  nearest to  $-\sum_i m_i \theta^i$ . The algebraic integer  $P(\theta) = p + \sum_i m_i \theta^i$  is nonzero: the minimal polynomial  $t^{r+1} - 2$  is Eisenstein and has degree  $r+1$ . Each of its other  $r$  conjugates has modulus at most  $(9/2)H$ , since  $|p| \leq 2H + 1/2$  and  $|\theta^i| < 2$ . Its nonzero integer norm gives

$$\|m \cdot \alpha\|_{\mathbb{R}/\mathbb{Z}} \geq (2/9)^r H^{-r}.$$

This classical algebraic example is credited in [4]; it is not new arithmetic. The inherited badly approximable theorem, Corollary 18.4, now supplies the all-hidden lower bound and the matching candidate denominators for the compiler. Bisection for  $\theta$  and fixed-degree powers meet the digit hypothesis. A fixed program with  $s(N)$  work bits has at most  $C(s(N) + 1)^C 2^{s(N)}$  command-boundary configurations. Its private randomness since the last boundary is integrated into the next stochastic row, so it induces a legal atomic realization of at most that width. Apply the lower bound  $W_{N,\varepsilon} \geq cN^{r/(2r+1)}$  and take logarithms. The polynomial factor accounts for the fixed number of work-tape head positions; it contributes  $O(\log \log N)$  unless the displayed lower bound is already automatic. Theorem 21.3 supplies the upper work bound.  $\square$

## 22. REALIZATION THEORY, ARITHMETIC AND COMPUTATIONAL SEMANTICS

Positive order and cone invariance. In the classical stationary problem one seeks a fixed nonnegative state-space realization of a transfer function. Minimal signed order need not equal minimal positive order. Benvenuti–Farina's Theorems 2–3 formulate positive existence by an invariant polyhedral

cone, either in a minimal linear realization or in a response-sequence space [12]; the 2022 survey concerns positive minimality [13]. Our enclosure construction is such a positive-cone construction after homogenization. The competing machines are different: their rows may vary with the cut and their entire realization may be redesigned for  $N$ . Proposition 3.1 makes the corresponding finite controlled normal form explicit. Classical stationary order cannot be substituted for its least simultaneous profile.

Hankel rank and signed weighted automata. For fixed seed and query the mean series is  $w \mapsto \rho\langle e_j, U_w x \rangle$ , with a signed linear representation of dimension at most  $D$ . Adding the constant probability offset requires at most  $D + 1$  dimensions. Balle–Panangaden–Precup’s Theorem 2 identifies minimal weighted-automaton dimension with infinite Hankel rank [14]. That is the signed, all-length invariant, not stochastic label cardinality. Their Theorem 11 bounds an  $\ell^2$  approximation of absolutely summable rational series. Here the unweighted all-word probability series is not summable, and the objective is the maximum numerical error on a fixed length. A small weighted mean-square error under a chosen word law does not control rare-word conditional error. For an explicitly unweighted finite table, Frobenius error does bound entrywise error, but its optimizer still need not be stochastic or compatible. Equation (3.5) exhibits the precise failure of separate ranks to detect the present growing resource.

Stochastic and hidden Markov realization. Finesso–Grassi–Spreij use pseudo-Hankel nonnegative factorizations and an informational-divergence criterion to approximate a stationary output process at prescribed hidden order [15, Sections III–IV]. Their second realization step is itself evidence that a static factorization is not the complete dynamic specification. Our alphabet is externally controlled rather than generated as an observation process. Seed resets and selected query columns are separately specified; worst-word conditional probabilities replace a divergence between output laws. These are explicit changes of datum and objective, not a claim that Hankel positivity was absent from HMM theory.

Ohta’s reachable and null-space reductions, Theorems 6–9 and Remark 10, identify tensor and generalized-Hankel dimensions in the effective linear space of an HMM [16]. We use the same need to distinguish observable rank from hidden state count, but do not infer a positive stochastic basis from a signed quotient. His Section 4.4 also discusses rotational examples descended from earlier HMM realization work. Rotation-based rank/positive-order separations are therefore not claimed new here. Our quantitative theorems concern all horizon-specific stochastic rows and preserve bounded numerical responses on every controlled word.

Kronecker quasi-uniformity. Dick–Goda–Larcher–Pillichshammer–Suzuki [4], Theorem 3.9, characterize quasi-uniformity of the multidimensional sequence  $n\alpha \pmod{\mathbb{Z}^r}$  by the simultaneous bound  $\eta_\alpha(q) \geq cq^{-1/r}$ . Their Lemma 3.11 bounds ratios of successive best simultaneous denominators; Section 3.3 also contains the vector  $(2^{1/(r+1)}, \dots, 2^{r/(r+1)})$ . These are direct arithmetic antecedents, not new examples of this paper.

Our packet points instead have the form  $m \cdot \alpha \pmod{1}$ . The relevant separation is dual. The residue argument of Lemma 18.2, with an arbitrary nearest-integer numerator and  $\eta_\alpha(q)$  in place of  $Q^{-1}$ , proves directly that a dual bound with exponent  $r$  implies

$$\eta_\alpha(q) \geq br^{-r-1}q^{-1/r}.$$

Thus the cited quasi-uniformity theorem applies to the uniform badly approximable class used in the predecessor. Its denominator control supplies the arithmetic scales for a polygon upper bound. It supplies neither a compatible stochastic transition table nor a lower bound against discarded hidden history. The logarithmic separation of Lemma 19.1 is weaker than a uniform badly approximable inequality; it is proved by a summable tail estimate, not asserted to imply quasi-uniformity.

What is dynamic in the paired profiles. The geometry of nearest-center distortion and convex hull approximation is classical [5, 11]. A static  $k$ -center factorization does not specify how its encoder is calculated without retaining discarded inputs. Conversely, a convex codebook does not force arbitrary hidden states to possess coordinates in that codebook. The two implications in Section 15 avoid both identifications: the converse conditions on the actual state, and the upper theorem provides the same legal rows throughout the run. The packing–dilation balance in Corollary 15.7 is a conditional theorem on independently specified finite profiles, not a claimed equality of static and dynamic optima.

The circle program makes one of these profiles exactly finite, whereas the all-harmonic action-gap optimization of the predecessor remains infinite-dimensional. A finite word length can be exhaustively listed; an unbounded harmonic degree cannot be certified by a finite truncation without a tail theorem. No general action-gap algorithm follows from Proposition 15.1.

Branching programs and probability preservation. At fixed horizon our machine is a probabilistic ordered read-once layered program with arbitrary stochastic rows and bounded real output columns. We distinguish width from total size, uniform description and random-bit complexity. Fourier growth for regular, permutation and width-three branching programs concerns different hypotheses and objectives [9, 10, 6]. Circle-walk harmonic damping and its arithmetic behavior are also established subjects [2]. Our converse can use correlated letters inside independently sampled complete packets; it is not a theorem about average error under fair independent letters.

Preserving a strict cutpoint sign is weaker than preserving a numerical probability. The predecessor’s constant-width sign construction contracts the output amplitude with length and hence does not contradict the present width growth. Here signal and tolerance are fixed, one query is selected after the word, and neither joint answers nor anytime correctness is required. The exact arithmetic upper rows are atomic real kernels. Section 21 gives a separate finite-bit approximation theorem; it does not represent those exact rows by a finite number of fair bits. A row with three successors can still have costly or irrational coefficients.

Conclusions. The finite optimized word profile and the compatible enclosure profile retain quantitative information which qualitative gap data discard. They yield a minimal-type exponent and an almost-everywhere logarithmic estimate, and they exhibit fluctuating numerical width for every Liouville single angle. The reflected versions show that these conclusions survive a particular noncommuting finite extension, without asserting a classification of arbitrary noncommuting nongapped alphabets.

Closed-form finite widths beyond the explicitly solved frontier, the Liouville upper logarithmic limit, optimal enclosure shapes and equality in the distortion–dilation comparison remain distinct quantitative problems. The new profile comparison uses no external spectral gap. For fixed computable planar angles, storing all command counts already gives a classical  $O(\log N)$  finite-accuracy work-space upper bound. We do not claim novelty for that upper order alone. The compiler proved here additionally maintains an arithmetic-scale clean stochastic register, generates circulant dyadic rows without a large advice table, and explicitly prices every arithmetic and random-bit register. The earlier expansion, calibration and representation theorems retain their original hypotheses in the unchanged supporting article.

Certificates and optimization. The Gram identity of Theorem 14.1 is a finite-length version of classical weighted-series equivalence, with stochastic rows and bounded numerical decoders imposed separately. In contrast to all-length  $\ell_2$  truncation, the moment sandwich in Theorem 14.3 targets the maximum error over the finite word set and optimizes over legal stochastic machines. Its global minimum is not supplied by an approximate matrix factorization. The rational grid certificate is complete at every fixed positive margin, but exact rational feasibility is different; the normalized nonnegative-rank example makes that distinction explicit. The new certificate framework does not change any earlier arithmetic exponent or imply a uniform solution of the infinite-gap problem.

## REFERENCES

- [1] S. Basu, R. Pollack and M.-F. Roy, *Algorithms in Real Algebraic Geometry*, 2nd ed., Springer, 2006.
- [2] I. Berkes and B. Borda, Random walks on the circle and Diophantine approximation, *J. London Math. Soc.* **108** (2023), 409–440.
- [3] J. W. S. Cassels, *An Introduction to Diophantine Approximation*, Cambridge Tracts in Mathematics and Mathematical Physics, vol. 45, Cambridge University Press, 1957.
- [4] J. Dick, T. Goda, G. Larcher, F. Pillichshammer and K. Suzuki, On the quasi-uniformity properties of quasi-Monte Carlo point sets and sequences—Part I: Lattices and Kronecker sequences, arXiv:2502.06202v2, version of 13 February 2026, Section 3.3.
- [5] S. Graf and H. Luschgy, *Foundations of Quantization for Probability Distributions*, Lecture Notes in Mathematics, vol. 1730, Springer, 2000.
- [6] C. H. Lee, E. Pyne and S. Vadhan, Fourier growth of regular branching programs, *APPROX/RANDOM 2022*, LIPIcs **245**, Art. 2, 2:1–2:21.
- [7] Q. Qi, *General Theta Foundations I: Sharp Width of Hidden Rotation Experiments*, manuscript, revision 35, 2026; repository publication aabc0731a0c9.
- [8] Q. Qi, *General Theta Foundations I: Alphabet Dynamics and Calibrated Numerical Memory*, manuscript, revision 42, 2026; repository publication b5f9874903a4.
- [9] O. Reingold, T. Steinke and S. Vadhan, Pseudorandomness for regular branching programs via Fourier analysis, *RANDOM 2013*, LNCS **8096**, 655–670.
- [10] T. Steinke, S. Vadhan and A. Wan, Pseudorandomness and Fourier growth bounds for width 3 branching programs, *Theory Comput.* **13** (2017), Art. 12, 1–50.
- [11] M. Vidyasagar, *Hidden Markov Processes: Theory and Applications to Biology*, Princeton University Press, 2014.
- [12] L. Benvenuti and L. Farina, A tutorial on the positive realization problem, *IEEE Trans. Automat. Control* **49** (2004), 651–664, Theorems 2–3 and Example 4.
- [13] L. Benvenuti, Minimal positive realizations: A survey, *Automatica* **143** (2022), Art. 110422.
- [14] B. Balle, P. Panangaden and D. Precup, A canonical form for weighted automata and applications to approximate minimization, arXiv:1501.06841v3, 2015, Theorems 2 and 11.
- [15] L. Finesso, A. Grassi and P. Spreij, Two-step nonnegative matrix factorization algorithm for the approximate realization of hidden Markov models, arXiv:1007.3435v1, 2010, Sections III–IV.
- [16] Y. Ohta, On the realization of hidden Markov models and tensor decomposition, arXiv:2008.11487v1, 2020, Theorems 6–9 and Remark 10.
- [17] Y. Bugeaud and M. Laurent, Exponents of Diophantine approximation, arXiv:math/0611354v1, 2006, Definition 1 and Section 2.
- [18] Q. Qi, *General Theta Foundations I: Compatible Polyhedral Lifts and Quantitative Causal Width*, manuscript, revision 33, 2026; repository publication aab961081253.
- [19] Q. Qi, *General Theta Foundations I: Word Profiles and Arithmetic Fluctuations of Numerical Memory*, manuscript, revision 43, 2026; repository publication 153830f5dc8d.
- [20] W.-G. Tzeng, A polynomial-time algorithm for the equivalence of probabilistic automata, *SIAM J. Comput.* **21** (1992), 216–227, doi:10.1137/0221017.
- [21] Y. Shitov, A universality theorem for nonnegative matrix factorizations, arXiv:1606.09068v2, version of 5 April 2018, Theorem 2 and Corollary 3.
- [22] N. Gillis and Y. Shitov, Low-rank matrix approximation in the infinity norm, *Linear Algebra Appl.* **581** (2019), 367–382; arXiv:1706.00078v1, Lemmas 1–2 and Theorem 3.
- [23] S. Morozov, M. Smirnov and N. Zamarashkin, On the optimal rank-1 approximation of matrices in the Chebyshev norm, *Linear Algebra Appl.* **679** (2023), 4–29; arXiv:2212.01438, Theorem 5.3.
- [24] S. Boyd, S.-J. Kim, L. Vandenberghe and A. Hassibi, A tutorial on geometric programming, *Optim. Eng.* **8** (2007), 67–127, Sections 2 and 8.
- [25] S. Boyd and L. Vandenberghe, *Convex Optimization*, Cambridge University Press, 2004, Section 5.8.
- [26] Q. Qi, *General Theta Foundations I: Robust Compatibility of Numerical Word Realizations*, manuscript, revision 47, 2026; repository publication 98b59f3f45f6.
- [27] T. Kahle, K. Kubjas, M. Kummer and Z. Rosen, The geometry of rank-one tensor completion, *SIAM J. Appl. Algebra Geom.* **1** (2017), 200–221; arXiv:1605.01678v2, Section 2.
- [28] A. Schrijver, *Theory of Linear and Integer Programming*, Wiley, 1986, Chapter 10.